

Constitution

Helix Universal Constitution

Field	Value
Revision	2026-10-05 07:49:21
Created	- -
Last modified	- - T : : Z 2026-10-05 11:4:276
Status	active 7 5 7
Status summary	Amendment of - - - NEW § . . (Review-round budget: every independently-reviewed work item converges within a declared budget of at most rounds, - consumer-declared, by front-loading ground truth, class-exhaustive enumeration, real-path tests and a declared threat boundary, with a binding methodology change at round) landed, operator mandate CRITICAL , verbatim: "Make sure we never ever get to such extremely high never ending review rounds count! Determine exact causes of issues we are facing with deep planning and exhaustive analisys so all work ALWAYS gets done with to reviews max! This MUST BE mandatory part of the constitution!!!" – root-cause investigation found five recurring, avoidable mechanisms behind this project's own - +-round review sagas (wrong-model

Field	Value
-------	-------

	ground truth shared by a fix and its own gate; instance-at-a-time closure of a nameable class; fixer-authored tests that never drive the real code; non-binding escalation signals – \$. . . cited ~ times, never binding, \$. . . 's N never declared; scope drift). Independent Opus-xhigh review returned NO-GO (Blocking/ Important/ Minor, all concrete and evidence-cited); every finding applied directly in one remediation pass per the reviewer's own "GO-ready without another structural round" assessment, itself now an instance of the very discipline the anchor establishes. Landed with two reconciling one-sentence EXTENSIONS: \$. . . (E) (an operator-mandated ceiling recorded with its own verbatim mandate in its own anchor is the one permitted "N-bound literal that clause allows) and \$. . . (its unconditional re-run-until-GO obligation is now bounded per-approach by \$. . . , escalating at the ceiling per \$. . . (E)() rather than continuing indefinitely). CM-COVENANT-114-276-PROPAGATION implemented (via the existing covenant_propagation_wrappers_generate.sh generator + a data-pack row, never hand-written, \$. . . thin-wrapper discipline) and verified GREEN against the real repo (MISSING/DUPLICATED/DIVERGENT across all constitution mirror carriers; its paired \$. . . mutation test confirms genuine, non-bluff discrimination across "" "" fixture classes). The recommended mechanism gates (CM-REVIEW-ROUND-BUDGET ; CM-REVIEW-ROUND3-CONVERGENCE-ASSESSMENT , CM-FIX-GROUND-TRUTH-AND-CLASS-INVENTORY) registered as tracked deferrals OWED-GATE- .. in gate_ledger_deferrals.tsv – gate-code NOT claimed shipped (\$. . . /\$. . .). Self-applied immediately: the live dual-vendor-split batch
--	---

Field	Value
	(ATM-) had already consumed review verdicts with no G0, exceeding the anchor's own default -round ceiling for an undeclared-budget item – per § . . (E)() the conductor stopped before dispatching what would have been round' under the same item identity, surfaced the choice to the operator, and re-filed the work as a fresh-budget new item per the operator's explicit choice, now progressing through its own round - cycle under § . . 's own discipline (round itself triggered the anchor's own mandatory structural-round checkpoint when round 's findings recurred in the same defect class as round 's). Generated .html/.pdf/.docx twins for Constitution.md/CLAUDE.md/AGENTS.md/QWEN.md/GEMINI.md NOT regenerated this round (honestly STALE per § . . (E), consistent with this file's own established practice when an anchor-landing round doesn't also carry an export-pipeline invocation – tracked as owed, not silently claimed current). Prior round: Amendment of - - (round , operator mandate CRITICAL , verbatim: "If we do not have access to Opus model for some reason but only to Sonnet, we MUST USE what we have! Reviews MUST BE done. Ideally with Opus model, however Sonet MUST BE used if Opus for any reason cannot be used! ... apply the rule asap so nothing is blocked with review which is not performed yet! We MUST BE maximally efficient!") – § . . + § . . AMENDED in place (extend-don't-restate per § . .): both anchors' hard "Opus-xhigh-only, no fallback model, BLOCKED-on-unavailability" pin is relaxed to a PRIMARY-Opus-xhigh-plus-genuine-unavailability-SONNET-fallback pin – Fable/Haiku remain permanently forbidden in every role; a review/merge-con-

Field	Value
	flict-resolution genuinely unable to reach Opus MUST now run on Sonnet (highest settable effort) rather than sit BLOCKED, with the Opus-unavailability FACT + the Sonnet substitution both recorded honestly in captured evidence (never presented as if Opus ran); BLOCKED-and-deferred now applies ONLY when BOTH Opus and Sonnet are genuinely unavailable. Both gate contracts (CM-CODE-REVIEW-OPUS-XHIGH / CM-MERGE-CONFLICT-OPUS-XHIGH) amended to accept the Sonnet-plus-fact pattern and FAIL on Fable/Haiku/any-other-model or on an unrecorded Sonnet substitution; headings + Classification/Composes lines updated to cite § . . (maximal multi-agent utilization) + § . . (A) (the sonnet rung this fallback draws on); propagated byte-identically to the compact-form mirrors (constitution submodule's CLAUDE/AGENTS/QWEN/GEMINI.md + main repo's AGENTS/QWEN/GEMINI.md + docs/PROJECT_GOVERNANCE_ANCHORS.md) and the main repo's CLAUDE.md one-line index entries updated to match (title text + §11.4.183 / §11.4.231 refs added). Gate CODE for CM-CODE-REVIEW-OPUS-XHIGH / CM-MERGE-CONFLICT-OPUS-XHIGH remains UNIMPLEMENTED (confirmed by grep of constitution/scripts/ + scripts/ + tests: no implementation exists; gate_ledger_removals.tsv rows - already honestly record both successor tokens as "live, currently-UNIMPLEMENTED" entries counted by the § . . (A) ratchet) – a dedicated gate_ledger_deferrals.tsv row for the amended Sonnet-fallback contract is separately OWED (a pre-existing § . . (A) gap predating this amendment, tracked § . . , never claimed shipped per § . .). Prior round: Amendment of - - (round , fix-forward on the independent Opus-xhigh review of fd fa: NO-GO, Blocking / Important / Minor) – § . . (D) corrected in place: (D.)

Field	Value
-------	-------

	§ 11 4 239 . . mandated validation is now AT MINIMUM source-GO + § . . (A) pre/post-build gates PLUS every operator-declared QA-deploy-seam blocker (§ . .) and any touched § . . 11 4 192 critical-invariant scenarios; (D.)/(D.)/(D.) re-keyed from "trunk" to the increment's own INTEGRATION BRANCH, a QA build MAY be cut from a feat/product/flavor branch and for feat/ IS the § . . (B) live-manual-QA vehicle (removes the trunk § . . (B) circularity; product/flavor lines now covered); (D.) builds coalesce (in-flight build covers prior increments) and a periodic interval must carry a recorded justification; (D.) toggle-expiry extensions only via a § . . waiver; heading + preamble now name clause (D); next-clause-letter (E) note added; CM-QA-BUILD-CADENCE-DECLARED + CM-INCREMENT-DELIVERED-NOT-SPEC-WITHHELD contracts amended for coalescing (the latter not flagged by the review, aligned for the same false-positive reason); ledger rows OWED-GATE- .. formatting aligned. Constitution-submodule .html/.pdf/.docx twins NOT regenerated this round either (honestly stale per § . . (E)); main-repo mirror twins regenerated in the paired main-repo commit. Prior round: Amendment of - - - § . . EXTENDED in place with clause (D) INCREMENTAL DELIVERY GRANULARITY + DECLARED QA-BUILD CADENCE (operator mandate 4 235 - - , CRITICAL : faster initial dev releases for every inheriting project – ship each already-source-proven-correct phase/story/task increment as it clears its own § . . (A)/(D) bar instead of withholding delivery until the whole spec is done; (D.) delivery unit = the increment; (D.) incomplete spec stays safe via complete ver-
--	---

Field	Value
	tical slices or LATENT paths – release toggles permitted not mandated, each tracked with an expiry; (D.) consumer-declared QA-build cadence per-increment / periodic / whichever-first , default per-increment , each build a \$. . . (B) cycle boundary passing \$. . . on its own fingerprint, deliverable form = "consumer DATA; (D.) background hardening never gates the next increment, except shared-code source-defects and \$. . . ; (D.) tracker/doc sync is a hand-off precondition, citing \$. . . (F) / \$. . . / \$. . . / \$. . . / \$. . . ; (D.) manual-QA finds on increment builds are still \$. . . coverage escapes). Extension, not a new anchor (\$. . . audit: \$. . . (A) overlaps validation with ONE candidate's build, \$. . . (A)/(D) define per-slice build-ELIGIBILITY; neither obliged an eligible increment to be DELIVERED, set a cadence, kept an incomplete spec safe, or bound hand-off to tracker truth). Research grounding cited in-text (DORA small batches, Hodgson feature toggles, branch-by-abstraction, SAFE develop-on-cadence, GitHub Next Continuous AI preview), with the honest note that no published methodology combining these was found. Recommended gates CM-INCREMENT-DELIVERED-NOT-SPEC-WITHHELD / CM-QA-BUILD-CADENCE-DECLARED / CM-RELEASE-TOGGLE-TRACKED-EXPIRY / CM-QA-HANDOFF-TRACKER-IN-SYNC registered as tracked deferrals OWED-GATE- . . . (gate-code NOT claimed shipped). Clause text byte-identical across Constitution.md + the . . . mirrors (\$. . .); \$. . . remains exactly one block-start per file (\$. . . (B)). Generated twins: constitution-submodule twins NOT regenerated (honestly stale per \$. . . (E), as in

Field	Value
-------	-------

	e⁴⁴ f f); main-repo twins regenerated in e²³²⁶ c⁰⁹ b eec. (The - - \$. . (D) amendment did not update this header; it is recorded in its own commit e f f.) Prior round: Amendment of - - - MODEL-SUBSTRATE CORRECTION (operator mandate - - , CRITICAL , verbatim: "Main model to use is Sonet, reviews and other heavy stuff have to be done by Opus instead of Fable!!! Fix this permanently on all levels!"): \$. . . rewritten - every mandatory inde- pendent code-review (\$. . / \$. . /\$. . , iterated to GO per \$. .) runs on Opus at xhigh ALWAYS; the - - - Fable-escalation clause (A) is REMOVED; Fable is not permitted in any role; no fallback model - genuine Opus-unavailability BLOCKS the review (captured fact, \$. .) and is tracked (\$. . /\$. .); prior - - and - - man- date texts kept verbatim as WITHDRAWN lineage. \$. . rewritten the same way for main→feature/product/flavor merge-CONFLICT resolu- tion (Opus-xhigh only, blocked-not-substituted on unavailability; the - - verbatim mandate kept as lineage, withdrawn as to its model choice). \$. . amended: ladder "is now sonnet < opus with Sonnet the DEFAULT working tier (A);" NEW clause (D.) records as a measured fact that a haiku dispatch in the reference consuming repository failed with "Prompt too long (~ k tokens vs k limit)" (request id req_ CfSBXm Kn enyUT UEaWd, - -) and PROHIBITS haiku there, a consumer re-admitting it only on its own captured evidence; (D) tail, (E), (F.), (F.) and the classification text re-pinned to Opus-xhigh. De-
--	--

Field	Value
	rived mentions re-pinned from Fable to Opus-xhigh in § . . -EXTENSION' (the truthful fable label vocabulary and the verbatim § . . - label mandate are KEPT – they describe what a live model-id may report, not a permitted dispatch tier), § . . , § . . , § . . and § . . . Recommended gates renamed CM-CODE-REVIEW-OPUS-XHIGH and CM-MERGE-CONFLICT-OPUS-XHIGH (predecessor tokens retired with citations in scripts/gates/gate_ledger_removals.tsv; unimplemented count unchanged by the rename). Honest boundary (§ . .): the Fable mentions that remain in this file are (a) the verbatim WITHDRAWN mandate texts kept as lineage, (b) the verbatim - label mandate + the label-vocabulary short-name list, (c) explicit prohibitions, and (d) the historical records of earlier rounds below – those records state what ran at the time and are not rewritten; any 'Fable re-review owed/recommended' item inside them is superseded: that owed re-review now runs on Opus at xhigh. Gate-code for the renamed gates remains a separate work item, NOT claimed shipped (§ . . / § . . .). Prior round: Amendment of § . . - NEW § . . (semantic Lumen index + indexing-efficiency gate + universal agent/subagent accessibility) + EXTENSIONS § . . ()-() / § . . ()-() / § . . ()-(): the indexing mechanism (scope renderer, single writer entry, version-keyed runner patches, scope proof, watchdog, lumen scope generator) is consumed BY REFERENCE from constitution/scripts/codegraph
Issues	none

Field	Value
Issues summary	ToC is currently stale (missing entries for \$. . /. /. /. /.) – separate clean-up commit needed per \$. . .
Fixed	\$. . added; \$. . , \$. . , \$. . ; QWEN.md created
Fixed summary	Operator-mandated additions – – propagated into Constitution.md (this commit); CLAUDE.md + AGENTS.md + new QWEN.md mirrors land in the same commit. \$. . renumber-from- \$. . documented inline (collision with concurrent Mechanical Enforcement landing).
Continuation	() ToC regeneration commit (cover \$\$. .). () Submodule-catalogue inventory expansion in CLAUDE.md / AGENTS.md / QWEN.md per the – – follow-up user mandate – per-Submodule purpose lines for every repo in vasic-digital + HelixDevelopment .

Table of contents

- . Test coverage is mandatory for every change
 - . False-positive immunity is an invariant
- . Commit and push mechanics – single entrypoint, locked
 - . Multi-upstream push is the norm
- . Submodule changes propagate through submodule commits first
- . Every tag on the main repo MUST be mirrored on every owned submodule

- § 1. Changelog discipline and multi-format export
- § 2. Documentation up to the nano-details
- § 3. Making false-success results literally impossible
 - § 3.1 NO BLUFF – positive-evidence-only validation
- § 4. Bleeding-edge ultra-perfection quality bar
- § 5. Absolute codebase and data safety – zero risk, zero loss
 - § 5.1 Mandatory safety protocol for destructive operations
 - § 5.2 Force-push requires explicit user authorization every time
 - § 5.3 Hardlinked backup is the standard – there is no excuse
 - § 5.4 Commit-message audit trail for history rewrites
- § 6. Enforcement
- § 7. End-user quality covenant
 - § 7.1 End-user quality guarantee – forensic anchor (User mandate, – –)
 - § 7.2 – FAIL-bluffs are equally forbidden
 - § 7.3 – Recorded-evidence requirement
 - § 7.4 – Per-environment-topology test dispatch
 - § 7.5 – Test-interrupt-on-discovery + retest-from-clean-baseline
 - § 7.6 – Captured-evidence quality analysis
 - § 7.7 – No-guessing mandate
 - § 7.8 – Demotion-evidence rule
 - § 7.9 – Deep-web-research-before-implementation
 - § 7.10 – Batch-source-fixes-before-rebuild
 - § 7.11 – Credentials-handling mandate

11 4 11

11 4 12

- o § . . .A – Pre-store credential leak audit
(User mandate, – –)
- o § . . . – File-layout discipline
- o § . . . – Auto-generated docs sync mandate
- o § . . . – Out-of-band sink-side captured-evidence
- o § . . . – Test playback cleanup mandate
- o § . . . – Item-status tracking mandate
- o § . . . – Item-type tracking mandate
- o § . . . – Universal-vs-project classification of new rules (User mandate, – –)
- o § . . . – Subagent-driven-by-default mandate (User mandate, – –)
- o § . . . – Script documentation mandate (User mandate, – –)
- o § . . . – Fixed-document column-alignment mandate (User mandate, – –)
- o § . . . – Operator-blocked status + self-resolution exhaustion mandate (User mandate, – –)
- o § . . . – Document-sync commit discipline (User mandate, – –)
- o § . . . – Visual-cue & grouping mandate for Issues docs (User mandate, – –)
- o § . . . – Build-resource stats tracking mandate (User mandate, – –)
- o § . . . – Full-Automation-Coverage Mandate (User mandate, – –)
- o § . . . – Constitution-Submodule Update Workflow Mandate (User mandate, – –)
- o § . . . – Submodule-Dependency-Manifest Mandate (User mandate, – –)

11 4 13

11 4 14

11 4 15

11 4 16

11 4 17

2026 05 14

11 4 20

2026 05 14

11 4 18

2026 05 14

11 4 19

2026 05 14

11 4 21

2026 05 14

11 4 22

2026 05 14

11 4 23

2026 05 14

11 4 24

2026 05 14

11 4 25

2026 05 15

11 4 26

2026 05 15

11 4 31

2026 05 15

- § 2024 05 15 . . . – Post-Constitution-Pull Validation Mandate (User mandate, 1.1 4 29 – –)
- § 2026 05 15 . . . – .gitignore + No-Versioned-Build-Artifacts Mandate (User mandate, 1.1 4 28 – –)
- § 2024 05 15 . . . – Lowercase-Snake_Case-Naming Mandate (User mandate, 1.1 4 27 – –)
- § 2024 05 15 . . . – Submodules-As-Equal-Codebase + Decoupling + Dependency-Layout Mandate (User mandate, 1.1 4 33 – –)
- § 2024 05 15 . . . – No-Fakes-Beyond-Unit-Tests + %-Test-Type-Coverage Mandate (User mandate, 1.1 4 34 – –)
- § 2026 05 15 . . . – Type-aware closure-status vocabulary (User mandate, 1.1 4 35 – –)
- § 2026 05 15 . . . – Reopened-source attribution mandate (User mandate, 1.1 4 36 – –)
- § 2024 05 15 . . . – Canonical-root inheritance clarity (User mandate, 1.1 4 37 – –)
- § 2026 05 15 . . . – Mandatory install_upstreams on clone/add Mandate (User mandate, 1.1 4 38 – –)
- § 2026 05 17 . . . – Fetch-before-edit mandate (User mandate, 1.1 4 39 – –)
- § 2026 05 17 . . . – Installable-Asset Evidence Mandate (User mandate, 1.1 4 40 – –)
- § 2026 05 17 . . . – Per-Feature On-Device End-User Validation Mandate (iter- , – –)
- § 2026 05 18 . . . – Full-suite retest before release tag mandate (User mandate, – –)
- § . . . – Pre-Force-Push Merge-First Mandate (User mandate, – –)
- § . . . – Iteration-discipline mandate (User mandate, – –)

- § 1.1 4 45 – TDD-Fix-Discipline Mandate (User mandate, 2026 05 18)
- § 1.1 4 46 – Document Revision Header Mandate (User mandate, 2026 05 18)
- § 2026 05 18 – Integration-Status-Doc Maintenance Mandate (User mandate, 1.1 4 47)
- § 2026 05 18 – Validate-recent-work-before-post-flash-tests mandate (User mandate, 1.1 4 48)
- § 1.1 4 49 – Firebase Data Review Mandate (User mandate, 2026 05 18)
- § 1.1 4 50 – UI-Driven Video Testing Mandate (User mandate, 2026 05 18)
- § 1.1 4 51 – Dual-Approach Testing Mandate (User mandate, 2026 05 18)
- § 1.1 4 52 – Deterministic Consistency Mandate (User mandate, 2026 05 18)
- § 1.1 4 53 – Live-ADB-First Maximization Mandate (User mandate, 2026 05 18)
- § 1.1 4 54 – Autonomous-Validation Mandate (User mandate, 2026 05 19)
- § 1.1 4 55 – Fixed_Summary parity mandate (User mandate, 2026 05 19)
- § 1.1 4 56 – ATM-NNN ticket identifier mandate (User mandate, 2026 05 19)
- § 1.1 4 57 – Reopens-history tracking + per-item Reopens.md doc (User mandate, 2026 05 19)
- § 2026 05 19 – Status_Summary parity + two-audience format (User mandate, 2026 05 19)
- § 2026 05 19 – README.md doc-link section + revision metadata (User mandate, 2026 05 19)

- This document defines mandatory, non-negotiable rules for every project that includes this constitution

submodule. Every AI agent and every human contributor MUST comply. Violations are blockers.

Project-specific constitutions extend this document. When a project Constitution (e.g. docs/guides/PROJECT_CONSTITUTION.md) and this universal Constitution differ, the project Constitution may extend or tighten – but NEVER weaken – the universal rules below.

Last revision: - - (v . - initial extraction from a consuming project's Android app . . -dev)

This Constitution is the universal source-of-truth for the engineering practices shared across all Helix projects (and any project that opts in by adding this submodule). It is intentionally agnostic about specific hardware, operating systems, languages, vendors, and product domains. Anything in this document that mentions a specific product, hardware SKU, library version, or vendor MUST be moved to the consuming project's own Constitution.

. Test coverage is mandatory for every change

Every code, configuration, documentation, or infrastructure change MUST ship with tests that prove:

Invariant	Mechanism
The change is present in source	pre-build / pre-merge gate that scans the source tree for the expected pattern

Invariant	Mechanism
The change survives compilation / packaging	post-build gate that inspects the produced artifact (binary, archive, image) and verifies the change actually shipped
The change behaves correctly at runtime	runtime / integration / on-device test that exercises the user-visible behaviour
The gate itself is not bluffing	meta-test (mutation) entry that mutates the asserted condition and proves the gate catches the break

A change without all four layers of coverage is not ready to merge.

. False-positive immunity is an invariant

A test that always returns PASS because its regex never matches, its input path is wrong, its assertion target does not exist, or its comparison is tautological is **worse than no test**. Every new gate MUST be paired with a mutation entry in the project's meta-test harness that:

- . Temporarily breaks the assertion (sed out a line, rename a symbol, etc.).
- . Re-runs the gate.
- . Asserts the gate now reports FAIL.
- . Restores the original.

If the mutation round does not turn PASS → FAIL, the gate is a sham and must be rewritten.

This is the **single most important rule in this Constitution**. Every subsequent § . .x clause is downstream of it.

. Commit and push mechanics – single entripoint, locked

All commit and push work uses the project's official multi-remote commit wrapper (e.g. `scripts/commit_all.sh`) as the single entripoint. Direct `git commit` / `git push` / `git add` on the main repo is prohibited in normal workflow.

Exception: tag creation + tag-ref push is done via explicit `git tag -a` and `git push <remote> <tag>` .

The commit and push wrappers MUST hold an advisory `flock` so two invocations cannot race against each other. Lock files self-clean on process exit via `trap` .

When a contributor sees the "another commit/push wrapper is already running" error, they MUST NOT `rm -f` the lock unless they have just killed the owning process.

. Multi-upstream push is the norm

Every project hosted on multiple Git providers (GitHub primary, GitLab / GitFlic / GitVerse / Bitbucket / Gitea mirrors) MUST push every commit to ALL configured upstream remotes. A commit that lives on only one remote is a future operational risk: when that one provider has an outage, the project is unreachable.

The constitution submodule itself ships an `install_upstreams.sh` helper (and an `Upstreams/` directory with one declaration script per remote). Running it from the submodule root configures all the remotes locally. Consuming projects SHOULD do the same for their own multi-remote topology.

. Submodule changes propagate through submodule commits first

When a change lands in any submodule of the project:

- . **Commit inside the submodule first** (the submodule's own `git add` + `git commit`), since the parent project's commit wrapper does not commit submodule source – it only captures the updated submodule pointer.
- . **Push the submodule commit** to all remotes of that submodule.
- . **Then** run the parent project's commit wrapper to capture the updated submodule pointer in main and push main.

Skipping step produces parent commits / tags that point at old submodule HEADs without the actual source.

. Every tag on the main repo MUST be mirrored on every owned submodule

When a version tag is created on the main repo at a specific commit, the same tag MUST be created on **every owned submodule** at that submodule's currently-pointed-to HEAD, and pushed to **every remote** of every owned repo. Third-party submodules (libraries not under the project's control) are excluded. The consuming project's Constitution declares its owned-submodule set explicitly.

. Changelog discipline and multi-format export

Every tagged release MUST ship:

- . A new `docs/changelogs/<tag>.md` entry describing user-visible changes, developer-visible changes, and known caveats.
- . Exports of that entry to `docs/changelogs/<tag>.{html,json,txt}` so external tooling (release portals, ticketing systems, package indices) can consume without needing a Markdown parser.
- . An update to a cumulative `docs/changelogs/CHANGELOG.md` (reverse chronological).

The project SHOULD provide a script (`scripts/testing/export_changelog.sh <tag>` or equivalent) that wraps `pandoc` + `jq` + plain-text Markdown strip.

. Documentation up to the nano-details

Every new feature, fix, or infrastructure change MUST update:

- . The project's `CLAUDE.md` / `AGENTS.md` Applied Fixes table (one row, one commit, one release tag).
- . `docs/guides/` for any user-visible or developer-reachable behaviour change.
- . Architecture diagrams, flowcharts, and any reference docs that touch the changed subsystem.
- . The per-version changelog (`$`).

Documentation drift after a fix is itself a Constitution violation: undocumented behaviour change is the same defect class as un-tested behaviour change.

. Making false-success results literally impossible

All validation output MUST satisfy the following invariants:

- . Every gate reports concrete PASS / FAIL / SKIP with explicit reason text. A gate that just says "ok" is non-compliant.
- . Every SKIP reason MUST be mechanically distinguishable from a PASS. Summary counters track them separately.
- . FAILs are counted towards exit status. A script that hides FAILs in logs and returns `0` is a bug.
- . Meta-tests (`$ .`) catch any gate that always PASSes regardless of condition.
- . Runtime results MUST be cross-referenced against host-side evidence (capture rig, screen recording, log analyzer, or equivalent) when the system's own introspection cannot directly verify the behaviour.

§ . NO BLUFF – positive-evidence-only validation

Every runtime test MUST satisfy ALL FIVE of the following non-negotiable constraints. A test that violates any one of them is a **bluff** and is forbidden from being committed to the test suite:

- . **Real ACTION:** the test MUST invoke at least one user-visible action via a project anti-bluff helper (e.g.

`ab_send_action()` in shell-based suites). A test that records a PASS without ever calling such a helper is automatically failed by the summary function.

- . **State DELTA:** every PASS that claims to verify a feature MUST capture state BEFORE the action and AFTER, and assert the delta matches expectation. `before == after` on success is a bluff.
- . **POSITIVE EVIDENCE:** the PASS condition MUST be POSITIVE evidence (a value match, a state delta, a captured frame analysis, a captured audio frame analysis), NEVER absence-of-error. If a probe returns nothing, that's a FAIL not a SKIP.
- . **UNIQUE EVIDENCE TOKEN:** tests that interact with mutable framework state MUST embed a per-run UUID into the action AND look for that token in the resulting state. Cached results predating the token cannot match – this defeats the stale-cache false-pass.
- . **Audio / video features REQUIRE captured evidence:** features that produce audio output MUST be validated by capturing the actual output via a capture pipeline (loopback, hardware capture rig, or equivalent) – NOT by metadata-only checks. Features that produce video output MUST be validated via frame capture + OCR or pixel-difference analysis.

Each consuming project SHOULD ship a shared anti-bluff helper library (typically `tests/lib/anti_bluff.sh` or equivalent) that codifies these rules; every new runtime test MUST source it and use its assertion helpers.

§ . Bleeding-edge ultra-perfection quality bar

The acceptance bar for shipping a project change is **not** "tests pass and source compiles" – that is the floor. The acceptance bar is:

The user, on a freshly-deployed system, sees the expected effect when they perform the documented action – confirmed end-to-end, with captured evidence, on the actual target environment.

The following are **forbidden**:

- . **"Source rebrand without shipped artifact rebuild"**: every source-side rename or rebranding **MUST** flow through to a shipped artifact that the runtime actually loads. Pre-build source-string gates are insufficient by themselves; shipped-artifact gates must accompany them.
 - . **"Tests pass but feature broken"**: this is the failure mode § . exists to eradicate. Any test that **PASSes** despite the feature being non-functional is a critical defect; its mutation pair must **FAIL** on the breakage.
 - . **"Configuration-only tests"**: a test that only verifies a config file exists / has expected contents is downgraded to a pre-build gate. Runtime tests **MUST** exercise the actual user-visible behaviour and capture the result.
 - . **"It works on my workstation"**: every fix is validated on every supported target before any release. No exceptions.
 - . **"Will fix in next cycle"**: deferrals are documented in the changelog with concrete blockers + implementation plans, not just handwaved.
-

§ . Absolute codebase and data safety – zero risk, zero loss

Every destructive operation on the repository (history rewrite, force-push, branch deletion, bulk file removal, submodule de-init, pack repack that drops objects) is a **safety-critical** operation. Data loss from a wrong force-push is irreversible once the remote garbage-collects dangling objects. This section is non-negotiable.

§ . Mandatory safety protocol for destructive operations

Every destructive operation **MUST** execute, in strict order:

- . **Full backup before touching anything.** Hardlinked mirror of `.git` to a sibling backup directory:

```
cp -al .git <backup>/<repo>.git.mirror
```

This is near-instant on the same filesystem and uses zero additional disk. If `.git` is small, also create `git bundle create <backup>/main.bundle --all .`

- . **Record critical metadata into the backup dir:**
 - `git show-ref > refs.txt`
 - `git tag > tags.txt`
 - `git submodule status > submodules.txt`
 - `git rev-parse HEAD > head.txt`
 - `git rev-parse HEAD^{tree} > head_tree.txt`
 - `git ls-tree -r HEAD --full-tree | sha256sum > head_tree_content.sha256`
 - `git ls-tree -r HEAD --full-tree > head_tree_listing.txt`
- . **Identify the target state:** the expected HEAD commit, tree hash, and tree content hash after the operation

completes. The operation MUST preserve these unless it is explicitly changing HEAD content.

- **Run the operation** (filter-repo, rebase, etc.). NEVER use `--no-verify` / `bypass-hooks` flags. NEVER auto-force on failure.
- **Post-operation verification gate** (all MUST pass):
 - HEAD tree hash matches pre-op target (content identical for non-content-changing rewrites)
 - HEAD tree content sha matches pre-op sha
 - All tags from `tags.txt` are preserved and resolve to valid commits
 - All submodule pointers from `submodules.txt` match current state
 - Project's pre-build / pre-merge gates pass with zero FAIL and zero unclassified WARN
- **Only if every check in § . . . passes** may the operation be considered successful. Otherwise: restore from backup immediately (`rm -rf .git && mv <backup>/<repo>.git.mirror .git`).
- **Then and only then** may a force-push proceed – and still not without explicit user authorization per § . . .

§ . Force-push requires explicit user authorization every time

Force-pushing (including `push --force`, `push --force-with-lease`, or any divergence-recovery code path inside an automated wrapper) overwrites the remote's view of history. On any shared branch this is irreversible.

Rules:

- Automated push wrappers MUST NOT auto-force-push as a fallback to any kind of failure (size rejection,

divergence, server error, anything). The historical divergence-recovery code that silently force-resets the remote is forbidden and must be removed or gated behind an explicit flag.

- A human must affirmatively say "force-push" (or equivalent) in the session to authorize each force-push.
- Force-push to `main` (and any equivalent primary branch on owned submodules) is only authorized after the `$ . .` gate has passed.
- Every force-push event is recorded in `docs/changelogs/<tag>.md` under a "Force-push audit" section (target refs, backup location, validation gate output).

\$. Hardlinked backup is the standard – there is no excuse

Because hardlinked `.git` copies are near-instant and use zero additional disk (same-filesystem inodes), the cost of making a backup is trivial. Any destructive operation without a fresh hardlinked backup is a Constitution violation regardless of how small the repo is or how confident the operator.

\$. Commit-message audit trail for history rewrites

After any `git filter-repo` run (or equivalent), commit a record under `docs/changelogs/` (or `docs/history-rewrites/`) containing:

- What was stripped and why
- Pre-op `.git` size vs post-op `.git` size
- Pre-op HEAD commit vs post-op HEAD commit
- Pre-op HEAD tree hash vs post-op HEAD tree hash (MUST match for non-content-changing rewrites)

- Backup location

§ . Enforcement

Any commit, tag, or release that violates § -§ is non-compliant. The fix is to amend/revert and re-land with compliance. No exceptions for speed pressure. Data-safety violations (§) and host-session-safety violations (§) are treated as catastrophic and block the entire release cycle until fully remediated.

§ . End-user quality covenant

§ . End-user quality guarantee – forensic anchor (User mandate, – –)

The reason § . + § exist – captured verbatim from the project owner so that future engineers and AI agents understand the historical failure mode the covenant exists to eradicate:

"We had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can't be used! This MUST NOT be the case and execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users of the product!"

This is the canonical motivation for every gate, every mutation pair, every meta-test, every anti-bluff helper, every captured- evidence requirement, every multi-environment validation rule. The covenant exists because the project has historically shipped broken features behind PASS-ing tests, and that outcome is no longer tolerated.

Operative rule: the bar for shipping is not "tests pass" but **"users can use the feature."** Every PASS in this codebase MUST carry positive evidence captured during execution that the feature works for the end user. Metadata-only PASS, configuration-only PASS, "absence-of-error" PASS, and grep-based PASS without runtime evidence are all critical defects regardless of how green the summary line looks.

Machine-created evidence at every gate (extension,
- - - \$. .). Every claim by the System that something WORKS, PASSES, IS VERIFIED, or IS PRODUCTION-READY MUST cite MACHINE-CREATED, MACHINE-VERIFIABLE evidence produced by the gate itself – not an operator's narrative, not "looks fine", not "no error was reported", not a prediction, not a guess, not an assumption. \$ makes this obligation EXPLICIT at every gate seam and applies \$. . () self-validating analyzers UNIVERSALLY, not only to AV playback. The captured-evidence family (\$. . / \$. . / \$. .) is REFINED from CAPTURED to CAPTURED-AND-MACHINE-VERIFIABLE. Prediction, guessing, and assumption about success are the exact anti-patterns this covenant forbids, at every layer of the System, from the first prompt of the session onward.

Propagation requirement: every `CLAUDE.md` and every `AGENTS.md` in the project tree (parent + every submodule recursively) MUST contain a clearly-titled "MANDATORY

ANTI-BLUFF COVENANT – END-USER QUALITY GUARANTEE" section that quotes this user mandate verbatim and points back to this § . as the canonical authority. The consuming project enforces this via a pre-build gate (typically `CM-COVENANT-PROPAGATION`).

§ . . – FAIL-bluffs are equally forbidden

The covenant must be enforced symmetrically. The historical failure mode the covenant eradicates is the **PASS-bluff**: a green test result on a feature that doesn't work for users. But the inverse failure mode is just as toxic – the **FAIL-bluff**: a red test result caused by a script bug (shell crash, missing argument, regex error, malformed assertion) rather than by an actual product defect.

Examples of FAIL-bluffs:

- A test using `set -eu` calls `log_pass` without arguments → the helper references `$1` → script crashes with " : parameter not set" → orchestrator records FAIL → engineer assumes product bug.
- A test pre-condition check uses an unbounded sed range → a comment line in the source matches the range terminator → the intended mutation never fires → audit reports PASS – bluff.

A FAIL produced by a test bug is no more useful than a PASS produced by a test bug – both mislead investigation, waste cycles, and ultimately let real defects ship undetected.

Operative rule (extended): every test MUST fail ONLY for genuine product defects. A test that crashes for a script-internal reason – undefined variable under `set -u` , unhandled error, malformed pipe, regex syntax error,

division by zero, etc. – is a critical defect in the test suite itself and MUST be fixed at the source layer (helper library, shared lib, test source) so the failure mode disappears project-wide, not patched in individual call sites.

§ 11.4.2 – Recorded-evidence requirement

A test that emits PASS without **captured visual or audio evidence of the user-visible feature actually working on the screen the user would see** is a § 11.4.2 PASS-bluff. Closing this gap requires a project-wide recording + analyzer infrastructure consisting of (at minimum):

- **Recording wrapper** that captures the user-visible output to time-synchronized media files. Multi-display systems record every output stream in parallel; sync metadata (host wall-clock at spawn) is written alongside the recording.
- **Action-timeline emitter**: every test emits structured timeline events (JSONL or equivalent) at every user-visible transition.
- **Frame / audio analyzer**: scans the recordings, extracts samples at a configurable interval, feeds them to recognizers (OCR for text overlays, speech-to-text for audio, pixel-diff for frame health), correlates against the timeline, and emits findings.

Closure criterion: every PASS for a user-visible feature MUST be cross-checked by the analyzer against the recording + action timeline. A PASS that lacks at least one matched timeline event in the analyzer findings is treated as a § 11.4.2 PASS-bluff.

§ . . – Per-environment-topology test dispatch

Tests that depend on environment topology (presence/absence of secondary displays, microphones, network interfaces, peripheral devices, cloud services, etc.) MUST detect topology at test entry and dispatch the topology-appropriate variant. A test that runs the WRONG variant for the actual topology and PASSES is a § . PASS-bluff: it claims a feature works while never having exercised it on the configuration the user actually has.

A topology-touching test that does NOT have a dispatcher AND a per-topology variant is a § . violation regardless of how it "passes" on any specific configuration. Single-test-multi-topology scripts MUST gate every topology-dependent assertion behind an explicit topology check that emits SKIP-with-reason if the required topology is absent – never PASS-by-default.

§ . . – Test-interrupt-on-discovery + retest-from-clean-baseline

A test cycle that *continues running past a freshly discovered defect* is itself a § . PASS-bluff: it produces "all green" summaries that the operator might tag as a release while the codebase under test is known-broken at the moment those greens were recorded.

The mandatory protocol – non-negotiable:

- . **Interrupt immediately.** The moment a defect is re-discovered, re-produced, or newly identified during a test cycle, the cycle MUST stop. No "let it finish for the data" – the data is contaminated.
- . **Systematic debugging before any fix.** Identify root cause at source layer (not symptom), blast radius across

related tests/features/subsystems, and regression-protection seam. Symptom patching is forbidden.

- . **Fix at root cause** per § . . : source-layer fix, not symptom patch in the calling site. The fix **MUST** be safe (no new failure modes), minimal (no surrounding cleanup smuggled in), and non-regressive.
- . **Four-layer test coverage per fix.** Every fix lands with positive-evidence coverage in every applicable layer:
 - **Pre-build / pre-merge gate** (source-level catch)
 - **Post-build / packaging gate** (artifact-level catch)
 - **Runtime / on-device test** (end-user-behaviour catch)
 - **Meta-test paired mutation** (proves the gate is not itself a bluff)
- . **Documentation MUST be updated for every fix** – Issues → Fixed migration on closure, Applied Fixes table row, user-facing guides, architecture diagrams, per-version changelog.
- . **Rebuild the system.** Full build via the canonical build script. Skipping rebuild – even when the fix is host-only – invites stale-artifact contamination of the retest baseline.
- . **Re-deploy on every supported target.** A target left on the pre-fix build is a confounder.
- . **Repeat the full test suite** from the beginning on every target, sequentially per host memory cap. Partial reruns are a § . violation.
- . **No-bluff certification per cycle.** Before tagging: meta-test harness returns all gates green AND every gate's paired mutation FAILs.

§ . . – Captured-evidence quality analysis

§ . . mandates *captured* evidence; § . . mandates the **content** of that evidence be analyzed for quality, not merely for presence.

Audio quality analysis – every audio test that PASSES MUST verify:

- . **Presence** – non-trivial RMS amplitude in captured output.
- . **Channel count** – analyzer reports the exact channel count the test claims (. stereo, . surround, etc.). Stereo downmix in a " . PASS" claim is a PASS-bluff.
- . **Sample rate + bit depth** – match the codec / pipeline under test.
- . **Glitch census** – underruns, XRUNs, dropouts above tolerance MUST be explicitly classified (PASS within budget, WARN above, FAIL on hard limits). Silently ignoring counts is a PASS-bluff.
- . **Coexistence-artifact census** – radio / IPC / shared-resource contention above tolerance.

Video quality analysis – every video test that PASSES MUST verify:

- . **Presence** – captured recording has non-zero size AND decoded- frame total > . A -byte file is the canonical PASS-bluff.
- . **Routing target** – analyzer confirms video appeared on the intended display / surface.
- . **Frame health** – drop count, jitter, freeze detection, tearing.
- . **Obstruction census** – OCR scan for hostile overlays (error dialogs, sign-in dialogs, geo-restriction overlays, paywall, etc).

- . **Resolution + codec** – captured dimensions match what the test claims.

§ . . – No-guessing mandate

Tests, gates, status reports, closure narratives, commit messages, and any operator-facing text **MUST NOT** use words like `likely`, `probably`, `maybe`, `might`, `possibly`, `presumably`, `seems`, `appears to`, `guess`, `seemingly`, `apparently`, `perhaps`, `supposedly`, `conjectured`, or their synonyms when describing CAUSES of test failures, system behaviour, or fix effectiveness.

Either:

- . **Prove the cause** with captured forensic evidence (logs, kernel ring buffer, kernel ramoops, structured snapshots, OS journals, strace, etc.) and state it as fact, OR
- . **Explicitly mark UNCONFIRMED + PENDING_FORENSICS** with a tracked-task ID for follow-up.

Every "X likely caused Y" sentence in the codebase or documentation is a § . . violation. Every "appears to be benign" without a concrete forensic trace is a § . . violation.

§ . . – Demotion-evidence rule

A demotion from any FAIL classification (`OPEN`, `POSSIBLE PRODUCT DEFECT`, `FAIL`) to a lower-severity classification (`INVESTIGATED`, `MITIGATED`, `RESOLVED`, `WORKING-AS-INTENDED`) requires positive evidence captured under the SAME CONDITIONS that originally exposed the defect:

- . **Same target** – defect on Target-A needs Target-A evidence; cross-target claims need every target.

- . **Same build** – evidence captured BEFORE a relevant fix landed does not validate post-fix behaviour.
- . **Same cycle position** – a stress-soak FAIL is not refuted by an isolated-test PASS.
- . **Same load profile** – a contention-mode FAIL is not refuted by a quiescent-mode PASS.

"I cannot reproduce in isolation" is a HYPOTHESIS, not a finding. Per § . . . it MUST be tagged **UNCONFIRMED:** until same-conditions retest produces positive evidence.

§ . . . – Deep-web-research-before-implementation

Before designing a non-trivial fix, before implementing a new feature, before declaring an architectural choice – perform deep web research to verify the chosen approach is informed by current state-of-the-art. The research surface includes:

- . **Official documentation** of the platform, framework, language, or service.
- . **Vendor technical guides** for the hardware / cloud / library.
- . **Open-source codebases** that already solved analogous problems.
- . **Coding tutorials + technical articles.**
- . **Issue trackers** for the relevant projects.

Operative rule. A fix that re-invents a wheel (or reproduces a known-broken pattern) when the open-source community has already solved the problem is a § . violation by omission.

Documentation requirement. Every non-trivial fix's commit message (or accompanying entry in the project's Issues / Fixed file) MUST cite the research sources that informed

it: at least one external link OR the literal "NO external solution found – original work".

§ . . – Batch-source-fixes-before-rebuild

When working through a multi-defect closure cycle, all source-side fixes that DO NOT require runtime validation to design MUST be landed BEFORE the next artifact rebuild. The anti-pattern this mandate eliminates is `Fix A → rebuild → flash → cycle → fix B → rebuild → ...` which serializes operator time onto rebuild latency.

Exceptions where a fix DOES require interim rebuild MUST be documented in the fix's commit message as `REQUIRES_REBUILD: <reason>`. The default is the batch path; rebuild-now is the operator-authorized exception.

§ . . – Credentials-handling mandate

Forensic anchor – direct user mandate:

"Credentials or any secret and sensitive data MUST NOT leak! This MUST BE added as the mandatory constraint into all Submodules and main project's Constitution, CLAUDE.MD and AGENTS.MD."

All credentials, secrets, API tokens, passwords, phone numbers, OAuth refresh tokens, signing keys, encryption keys, and any other authentication-or-authorization-bearing data used by test scripts, build automation, or any project tooling MUST be handled per the following rules without exception:

- . **No tracked-file storage.** Credentials NEVER live in any file that git tracks. Placeholder templates with `<replace ...>` or `EXAMPLE_VALUE_DO_NOT_USE` placeholders

ARE allowed in `.example` files committed to show operators what keys to populate.

- **Repository-wide ignore.** `.gitignore` MUST exclude all credential-bearing paths:
 - `.env` , `.env.*` , `*.env` , `.<service>.env` , `*.<service>.env`
 - `scripts/testing/secrets/*` with `.example` + `README.md` exception
- **Runtime-load-only.** Tests load credentials AT EXECUTION TIME from operator-populated files under `scripts/testing/secrets/` (or per-project equivalent). If the file is missing, the test SKIPS with "credentials not configured for " – it does NOT proceed without credentials and does NOT use defaults.
- **No echo / no log.** Test scripts MUST NEVER print, log, or include credentials in error paths, stack traces, screen recordings, or output artifacts.
- **Per-service file separation.** One file per service keeps blast radius small.
- **Filesystem permissions.** `.env` files are `chmod 600` , parent directory is `chmod 700` .
- **Rotation on suspected leak.** Rotate at provider, update local files, audit captured artifacts.

The consuming project's `.gitignore` MUST be verified before every commit: `git ls-files --cached | grep -E "\\..env$" MUST show no real .env files (only .env.example ).`

\$A – Pre-store credential leak audit (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Use these for all future testing (full automation testing) and make sure they are not leaking anywhere or get git versioned!"

[Discovered during execution: the operator-provided test credentials ALREADY existed in tracked files dating to prior project cycles. \$. . + \$. . catch new commits but did not detect the pre-existing leak when the operator re-provided the same values for a new use.]

When an operator provides credentials, API tokens, signing keys, or any other secret material to be stored in the project's gitignored configuration (.env, scripts/testing/secrets/, secrets/, etc.), the agent or human storing them MUST FIRST execute a repo-wide audit for prior leaks of THOSE specific values BEFORE storing.

The audit MUST:

- **Grep every tracked file** for the literal credential value(s). `git ls-files | xargs grep -l <value>` is the canonical form. Search MUST cover all file types – .md , .json , .yaml , .kt , .go , .py , .sh , .sql , build configs, docs.
- **Grep the entire git history** for the literal credential value(s). `git log -S<value> --all --source --remotes` reveals historical commits even if the current tree is clean. A history-only leak is a \$. . violation of equal severity to a tree-leak – anyone who pulled an older commit still has the value in their working copy of any branch they cloned.
- **Report findings to the operator BEFORE storing the new value in any operator-controlled location.** The operator may then decide whether to (a) rotate the credential at the upstream provider before reusing, (b) accept the

historical leak as a known-compromise and proceed, or
(c) abort the storage. Storing the value without
surfacing the audit is itself a § . PASS- bluff
at the security layer.

- . **If a leak is found:** open a forensic incident record per
the project's § /§ sixth-law-incidents discipline
(or equivalent), redact the literal values from all
tracked files in-place (replace with <redacted-per-
§11.4.10> placeholders), and record an OPERATOR ACTION
REQUIRED for rotation per § . . sub-clause
. . The redaction commit lives on master IMMEDIATELY;
the history- purge follow-up (git-filter-repo / BFG /
equivalent + force-push to every upstream) requires
explicit per-operation operator authorization per
§ . .
- . **Strengthen pre-push hook detection** when a previously-
undetected class of literal credential is discovered.
The pre-push hook's credential-pattern grep MUST be
extended to catch the specific pattern class that
escaped (literal usernames matched with literal
passwords in the same file, well-known credential
SHA- hashes, org-specific naming conventions for
service accounts, etc.). The extension MUST land in the
same commit as the redaction.

Pre-build gate (recommended) CM-PRE-STORE-CREDENTIAL-AUDIT
(when implemented in consuming project) checks the commit
message body of any commit that touches
.env , .env.example , scripts/testing/ secrets/ , or
equivalent for an Audit-Pre-Stored: stamp citing the
audit's outcome. Paired mutation (§ .) strips the stamp
and asserts gate FAILs.

Composes with § . . (the core credentials mandate
this extends), § . . (.gitignore + no-versioned-
artifacts – overlapping enforcement surface),
§ . . (no-guessing – audit findings stated as fact

or as PENDING_FORENSICS), \$ (demotion-evidence – "I think it's not leaked" without grep evidence is a guess). Classification: universal (per \$). No escape hatch.

\$ – File-layout discipline

Project files MUST be organised by purpose, not by historical accident. Source code goes under canonical project roots. Tests go under canonical test directories. Logs and forensic artifacts go under operator-controlled directories (~/Documents/ , qa-results/ , /tmp/ , OS-equivalent) – NEVER scattered at the repo root, NEVER inside working source trees, NEVER tracked unless they are reference assets explicitly required for evidence-replay.

Discovered drift is fixed by **moving** files to their canonical location and updating every caller – never by adding redirect shims or by leaving the misplaced copy "for backwards compatibility".

Carve-out (User mandate –). The canonical tracker documents – docs/Issues.md , docs/Issues_Summary.md , docs/Fixed.md , docs/Fixed_Summary.md , docs/CONTINUATION.md – sit at docs/ root by design. They are architectural constants of the project layout, analogous to AOSP's Makefile , Android.bp , OWNERS files at repo root. Their location is encoded as literal path strings in \$ + \$ + \$ + \$ propagation gates plus the helper-script constellation that regenerates them (generate_issues_summary.sh , generate_fixed_summary.sh , sync_issues_docs.sh , commit_docs.sh , colorize_progress_html.py , export_progress_docs.sh). Moving them would require coordinated amendment of those sister anchors plus pre-build gates plus ~

helper scripts plus consumer files (parent + owned submodules + nested + HelixQA dependencies) in a single PWU. Per § 1.1.4, that scope is operator-blocked until explicitly authorised. Audit-snapshot files (docs/audit/anti_bluff_audit.md , docs/audit/PRE_SONOS_TAG_READINESS.md , docs/audit/D1_WIFI_FAIL_CLASSIFICATION.md , plus any future audit snapshots) DO move under docs/audit/ per the § 1.1.4. general principle – they are not part of the tracker constellation and carry no propagation-gate path-string burden.

§ 1.1.4 – Auto-generated docs sync mandate

11 4

Every auto-generated document (Issues_Summary, API reference, build manifest, etc.) MUST be regenerated in the **same commit** as any edit to its source. Stale auto-generated docs are § 1.1.4 PASS-bluffs in document form: an operator reading them gets a divergent view of project state.

All three file types MUST stay in sync at all times – Markdown source + HTML export + PDF export (or equivalent multi-format output the project ships). Enforced by a pre-build gate that checks mtime ordering and content hash agreement.

§ 1.1.4 – Out-of-band sink-side captured-evidence

Whenever a downstream consumer (HDMI sink, cloud service, monitoring target, downstream system) provides a network-accessible introspection API that reports what was actually received, the test suite MUST consume that report as captured-evidence for every test asserting end-to-end delivery.

The on-source-side view ALONE is insufficient – that is the exact "tests pass but the feature doesn't work" pattern § 4.1.1 forbids.^{11 4 14} Sink-side reports MUST be:

- Identity-verified (MAC, serial, UUID match) before consumed as evidence.
- Topology-dispatched (§ 4.1.1) – sink-probe-unreachable → SKIP, never FAIL.
- Cross-referenced with the on-source state at a matching wall-clock instant.

§ 4.1.2 – Test playback cleanup mandate

A test that completes (PASS, FAIL, or SKIP) MUST leave the target in a quiescent state – no orphan playback, no orphan recording, no orphan capture, no orphan background process continuing after the test's last assertion. Tests are transient probes; their side-effects on shared target state MUST NOT outlive the test.

Mandatory protections:

- Every test that issues a playback or capture command MUST issue the matching cleanup at test end (force-stop the app, stop the recorder, close the file descriptor, terminate the helper).
- Cleanup is mandatory on EVERY exit path – PASS, FAIL, SKIP, error, trap, signal, parent-killed. Use shell `trap '<cleanup>' EXIT` or `try/finally`.
- Tests MUST verify the cleanup succeeded via positive evidence (§ 4.1.1).
- The orchestrator MUST run a post-test sanity check between tests and FAIL the just-completed test if it left orphan state. This blames the leaker, not the next-test victim.

- . No grace period for "the next test will clean it up" – that is precisely the § 11.4.1.3 PASS-bluff pattern. Cleanup is the test's responsibility, not the next test's, not the orchestrator's.

§ 11.4.1.3 – Item-status tracking mandate

Every active item tracked in the project's Issues file MUST carry a `**Status:**` line within five lines of its heading. The status value is drawn from a closed set:

State	Meaning
Queued	In backlog, no work has started. Sub-state qualifier <code>Queued – BLOCKED on <reason></code> permitted.
In progress	Active investigation or coding underway.
Ready for testing	Source-side fix landed (pre-build / meta-test gates green), waiting for the next deploy cycle.
In testing	A deploy + cycle is in flight that exercises the fix; live captured-evidence is being collected per § 11.4.2 + § 11.4.3.
Reopened	Item failed its runtime test cycle and is back to active work. Carries a reference to the failure-cycle artefact.
Fixed (→ Fixed file)	Item closed: captured-evidence per § 11.4.5 collected, runtime test PASSES, paired meta-test mutation FAILs, item migrated to the Fixed file.

Status MUST be updated as the item progresses. Every commit that changes an item's lifecycle state MUST update its Status line in the same commit.

All three Issues / Issues_Summary / Fixed file types MUST be in sync at all times (Markdown + HTML + PDF).

§ . . – Item-type tracking mandate

Every active item tracked in the project's Issues file MUST carry a `**Type:**` line within eight non-blank lines of its heading (the same window the § . . Status audit uses). The type value is drawn from a CLOSED set of three values:

Type	Meaning
Bug	Product defect / regression / user-visible broken behaviour. The product worked before (or was expected to) and now does NOT for the end user.
Feature	New capability not previously offered to end users – a new integration, new output, new probe, new bank, new architectural surface.
Task	Internal workstream – not user-visible. Refactor, infrastructure, documentation, audit, covenant clause, gate, mutation, propagation enforcement, host-session-safety hardening. The largest class by count; the lowest-stakes default when the classification is ambiguous.

The vocabulary is CLOSED – only Bug , Feature , Task . Any other value is a violation. When ambiguous, fall back to Task .

Type tells the operator WHAT KIND of work an item is – distinct from severity (HOW URGENT) and status (WHERE IN LIFECYCLE). The three axes together drive ownership routing, testing strategy, release-note framing, and changelog classification.

The Issues_Summary file MUST carry the Type column for every active item. All three file types (Markdown + HTML + PDF) MUST stay in sync under the same CM-DOCS-EXPORT-SYNC discipline as § . . + § . . .

Pre-build gates CM-ITEM-TYPE-TRACKING (scans Issues file for ****Type:**** presence within the audit window) and CM-COVENANT-114-16-PROPAGATION (anchor presence in every project CLAUDE.md / AGENTS.md) enforce the mandate. Paired mutations in the project's meta-test prove neither gate is a bluff gate.

No escape hatch – items that genuinely don't fit MUST be re-classified as INFORMATIONAL and excluded from the active-item set rather than carry a fabricated Type.

2026 05 14

§ . . – Universal-vs-project classification of new rules (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"Adding any new rules or mandatory constraints or anything relevant which should be added into Constitution, CLAUDE.MD, AGENTS.MD and relevant files MUST BE determined as reusable / universal VS project specific. If it is universal and reusable it MUST BE added into our root / main constitution Submodule, otherwise into our project level Constitution, AGENTS.MD and CLAUDE.MD (or Submodule if it is Submodule related)."

Before any new rule, mandatory constraint, covenant clause, gate declaration, or "MUST"-bearing statement is added to a project's Constitution / CLAUDE.md / AGENTS.md

or to a child submodule's equivalents, the author **MUST** first classify it along this axis:

Classification	Definition	Destination
Universal (re-usable)	The rule applies to ANY project – independent of language, framework, target platform, or domain. Examples: anti-bluff covenant, no-guessing mandate, credentials-handling, host-session safety, item-status tracking, multi-upstream push, file-layout discipline, deep-web-research-before-implementation.	This constitution submodule's <code>Constitution.md</code> + <code>CLAUDE.md</code> + <code>AGENTS.md</code> . Propagates to every consuming project via inheritance.
Project-specific	The rule references a specific hardware target, vendor, model, SoC, vendor-fork, project-internal package name, deployment topology, or company-internal asset. Examples (illustrative – these are intentionally NOT in this universal file): the SoC's specific power-management quirks; a particular Android/iOS/Linux SDK behaviour; a specific app or service bundled with the project.	The project's own <code>Constitution.md</code> / <code>CLAUDE.md</code> / <code>AGENTS.md</code> (top-level), OR the affected submodule's equivalents (when the rule scopes to that submodule).

A rule that mentions hardware part numbers, vendor names, specific package identifiers, geographic regions, or company-internal asset names is **project-specific** by

definition – it cannot be lifted into the universal layer without genericising those references first. A universal rule MAY reference patterns (e.g., "USB-HID multitouch panels") but MUST NOT name specific vendors.

Anti-bluff: the universal-vs-project classification is itself an auditable artefact. Every new rule's commit message MUST include a one-line classification statement (e.g., `Classification: universal – applies to any project tracking items through an Issues/Fixed lifecycle`). A commit that adds a rule without classification justification is a § . . violation.

Authors who are uncertain SHOULD default to project-specific (narrower scope). Lifting a project-specific rule to universal later is cheap (one merge); generalising a universal rule that turned out to leak project-specific assumptions is expensive (every consumer must update).

Pre-build gate `CM-UNIVERSAL-VS-PROJECT-CLASSIFICATION` audits the last N commits for rule additions and asserts each one carries a classification statement. Paired mutation strips the classification literal and asserts the gate FAILs.

§ . . – Subagent-driven-by-default mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"Make sure we ALWAYS WORK EVERYTHING subagents-driven if it is possible or applicable!"

When operating in a multi-agent runtime (Claude Code with subagents, Cursor with task-runners, Aider with sub-sessions, or any CLI tool that supports delegated

execution), the primary agent **MUST** default to subagent delegation for any task that satisfies **AT LEAST ONE** of:

- **Multi-step scope** – three or more discrete editing / file-creation / verification phases. Hand-off boundaries are where stalls happen; pre-planned subagent decomposition avoids them.
- **Parallelisable work** – two or more independent tasks with no shared mutable state. Parallel subagents finish wall-clock-faster than sequential foreground work **AND** insulate the primary agent's context window from the volume of file reads.
- **Long-running diagnostic loops** – repeated probe / re-test cycles, build-flash-cycle loops, soak tests. Subagents survive context compression boundaries that would otherwise truncate progress.
- **Domain-specific or specialised workflows** – code review, security audit, infrastructure change review, performance triage, documentation propagation across N files. Specialised subagents (`code-reviewer` , `iac-reviewer` , `general-purpose`) bring pre-curated tool sets and disciplines the primary agent would re-derive from scratch.

The primary agent **SHOULD** only do foreground work when **AT LEAST ONE** of:

- The task is a single edit or single file read with no follow-up.
- The task requires conversational clarification with the operator mid-execution (subagents cannot ask the operator questions).
- The task is gating critical state (a commit, a push, a tag cascade) that must be sequenced before the next operator action.

- The task is so quick (under ~ seconds of execution) that subagent spin-up overhead exceeds the work itself.

Anti-stall discipline. Subagents have watchdog timeouts in most runtimes (Claude Code: ~ s of no-progress). When delegating, the parent agent **MUST**: (a) scope tightly (no "complete everything" prompts – break into – tightly-scoped tasks); (b) require **checkpoint commits** after each major task so partial progress is preserved even on stall; (c) explicitly prohibit destructive operations the subagent might attempt to "be helpful" (`git reset --hard` , `git push --force` , `--no-verify`); (d) provide explicit discipline pointers (\$. . no-guessing, \$. . credentials, \$. . classification).

Anti-bluff applied. A subagent that returns "completed" without captured evidence in commits / outputs / artifacts is a \$. PASS-bluff at the multi-agent layer. The parent agent **MUST** verify subagent claims against repo state (`git log` , `git status` , post-completion gate runs) before treating the work as landed.

Coordination. When parallel subagents touch overlapping files, the parent agent **MUST** partition the work so subagents work on **non-overlapping files**. The parent's `commit_all.sh --auto-cascade` naturally bundles both subagents' dirty files into atomic commits via `git add -A` – exploit this for "forward-reference + provider" patterns where one subagent creates files another subagent references.

No escape hatch. Operating exclusively in the foreground when subagent delegation is feasible burns operator wall-clock time, risks context-window overflow on large tasks, and forfeits the parallelism that multi-agent runtimes were designed to provide. Pre-build gate `CM-SUBAGENT-`

DELEGATION-AUDIT (when implemented per consuming project) scans recent session transcripts for multi-step foreground work that should have been delegated; paired mutation enforces the gate is not a bluff.

§ . . – Script documentation mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"Make sure that every single script inside the scripts dir (in all subdirs in depth) is fully documented and covered with full user manuals and user guides! ... Whenever is some bash script modified we MUST document it fully and create for it complete user guide(s) and manual(s) – if already exists make sure all of it is in sync and fully updated! No documentation ever can be out of sync with its codebase!"

Every Bash / shell / POSIX-sh script ANYWHERE in a project (anywhere under scripts/ , bin/ , tests/ , library directories, Upstreams/ , deployment hooks, CI helpers, etc. – depth-N recursive) MUST carry:

- **In-source documentation block** at the top of the file:
 - Purpose: one-sentence description of what the script does.
 - Usage: complete invocation syntax with all flags, env vars, positional arguments, examples.
 - Inputs: env vars, files read, command-line args, stdin.
 - Outputs: files written, stdout/stderr behaviour, exit codes.

- **Side-effects:** anything that changes system state outside the script's own scratch directory.
 - **Dependencies:** required commands (and how to install them) + required system state (e.g., "must be run as root", "must run from project root").
 - **Cross-references:** companion guides under `docs/`.
- . **External user guide** under `docs/scripts/<script-name>.md` (Markdown). Required sections:
- **Overview** – what the script does in plain English.
 - **Prerequisites** – environment + dependencies.
 - **Usage examples** – copy-paste recipes for the common cases.
 - **Edge cases** – unusual invocations, error conditions, how to diagnose them.
 - **Internal behaviour** – for advanced users / future maintainers: what the script does step-by-step.
 - **Related scripts** – companion scripts and how they compose.
 - **Last verified version / date** – when the doc was last reconciled against the script source.

Whenever the script source is modified, the in-source documentation block AND the external user guide MUST be updated in the SAME commit. A commit that touches a script but leaves its documentation unchanged (or worse, lets it drift) is a § . . violation.

Pre-build gate `CM-SCRIPT-DOCS-SYNC` walks every `*.sh` and `*.bash` under `scripts/` (or equivalent script directories), verifies a companion `docs/scripts/<name>.md` exists, AND verifies the doc was modified in the same commit as the script (by SHA cross-reference, or by `mtime ≥ script mtime` as a softer floor). Paired mutation strips the doc-sync invariant and asserts the gate FAILs.

No documentation ever can be out of sync with its
codebase.

This mandate composes with § 11.4.12 (file-layout discipline – docs live under docs/, scripts live under scripts/) and § 11.4.13 (auto-generated docs sync – the on-disk Markdown + its rendered HTML/PDF must all stay synchronised).

§ 11.4.14 – Fixed-document column-alignment mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"Make sure that Fixed document (all file formats) is consistent and has the same structure and columns as Issues and Issues_Summary documents! This has to be added as detail to Constitution, CLAUDE.MD and AGENTS.MD (constitution Submodule where these information shall exist)."

Every project that maintains an Issues / open-work tracker AND a Fixed / closed-archive tracker MUST keep the two structurally aligned along the same lifecycle axes – at minimum **Status** and **Type** – so an operator reading either gets a consistent view across the entire lifecycle (open → in progress → ready for testing → in testing → reopened → fixed).

Specifically:

- **Per-entry annotation.** Every heading in the Fixed archive document MUST carry, within non-blank lines of the heading, a ****Status:**** line and a ****Type:**** line – exactly as § 11.4.15 + § 11.4.16 require for

the Issues tracker. For Fixed entries the Status value is drawn from the **closure-set**:

- Fixed (→ Fixed.md) – fully verified on device with captured evidence per § 11.4.5.
- Fixed – pending device verification – source-side fix landed, regression-protection gate wired, awaiting the next firmware build + flash + cycle for captured evidence.
- Fixed – RECLASSIFIED – item was reclassified during investigation (test-side defect, environment-only, etc.).

Type values follow § 11.4.5 : Bug | Feature | Task . Default on missing-annotation: Task .

- **Auto-generated companion summary.** A Fixed_Summary.md MUST exist alongside Fixed.md , regenerated by an automation script (e.g. generate_fixed_summary.sh), and MUST mirror the column structure of Issues_Summary.md exactly:

```
| # | Level | Status | Type | One-line description |
```

The Level (severity) column uses the **original severity** the item had when it was open – so an operator can correlate a closed item with its original priority bucket. The same deterministic classification rules apply (REAL PRODUCT DEFECT / PARTIAL-or-WORKSTREAM-or-DEEP-DIVE / everything-else mapping to C / M / L).

- **Three-format sync.** All three file types (.md , .html , .pdf) for BOTH the Fixed archive AND the Fixed_Summary MUST stay in sync at all times – same .html + .pdf export pipeline that § 11.4.5 requires for Issues + Issues_Summary. The single-shot sync wrapper (e.g. sync_issues_docs.sh) MUST invoke the Fixed_Summary generator AND export all five doc variants (Issues, Issues_Summary, Fixed, Fixed_Summary,

CONTINUATION) in one operation so they all carry the same mtime.

- **Cross-lifecycle consistency.** A status line in `Issues.md` that reads `Fixed (→ Fixed.md)` is a **migration marker**, not a self-describing closure: when an `Issues.md` entry resolves, it **MUST** be moved to `Fixed.md` in the same commit, then disappear from `Issues_Summary` (because `Issues_Summary` only enumerates OPEN items) and appear in `Fixed_Summary` (because `Fixed_Summary` enumerates CLOSED items). Per § . . . fix-closure protocol + § . . . captured-evidence requirement, the migrated entry in `Fixed.md` **MUST** carry: closure cycle, closure commit SHA, captured evidence pointer, regression-protection gate name + paired mutation reference.

Pre-build gate `CM-FIXED-COLUMN-ALIGNMENT` (+ invariants):

- `docs/Fixed_Summary.md` exists.
- `docs/Fixed_Summary.md` table header line contains both `Status` and `Type` columns (column-alignment with `Issues_Summary`).
- `mtime(Fixed_Summary.md) >= mtime(Fixed.md)` (regenerated after `Fixed.md` edits).
- Generator script exists (`generate_fixed_summary.sh` or equivalent integrated branch in `generate_issues_summary.sh`).
- Sync wrapper invokes the `Fixed_Summary` generator (grep for the script name in `sync_issues_docs.sh` or equivalent).
- HTML + PDF exports for both `Fixed` and `Fixed_Summary` exist (`docs/Fixed.html` , `docs/Fixed.pdf` , `docs/Fixed_Summary.html` , `docs/Fixed_Summary.pdf`).

Paired mutation: strip the Status column from Fixed_Summary.md table header (or rename the generator) → gate FAILs. Enforced by meta_test_false_positive_proof.sh .

Propagation. This anchor is a \$. . -classified universal rule – it composes naturally with \$. . (auto-generated docs sync), \$. . (Status tracking), \$. . (Type tracking) and is mandatory across every consuming project's CLAUDE.md / AGENTS.md. The propagation gate CM-COVENANT-114-19-PROPAGATION (when implemented per consuming project) enforces the anchor's presence in every covenant file.

No escape hatch. A Fixed archive that lacks Status/Type columns or whose Summary companion drifts out of sync is the exact "different lifecycles report different facts" hazard \$. . forbids – an operator reading Issues_Summary can no longer answer "where is item X now?" by cross-referencing Fixed_Summary because the schemas diverge. Non-compliance is a release blocker regardless of context.

\$. . – Operator-blocked status + self-resolution exhaustion mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"Add into Issues and Issues_Summary new status: Operator blocked! We need another column with details in Issues document with exact details on how exactly is operator blocked issue! Before issue becomes operator blocked we MUST investigate in-depth all ways on how it can be resolved without operator's

involvement – by you, or by activating any relevant mechanism for the resolution!"

Every project that maintains an Issues / open-work tracker MUST treat operator dependency as a **last-resort classification**, earned only after documented exhaustion of self-resolution paths. Routing an item to the operator without proving the agent + tooling cannot unblock it themselves is the § 11.4.15 PASS-bluff pattern transplanted into the planning layer – the work stalls indefinitely while the agent claims "nothing more I can do."

.. **Status vocabulary extension.** § 11.4.15 's closed-set of Status values is extended with a seventh value, **Operator-blocked**. The full closed-set becomes:

Status value	Meaning
Queued	Awaiting work to start.
In progress	Active agent / developer work.
Ready for testing	Source-side fix landed, awaiting verification.
In testing	Captured-evidence cycle running.
Reopened	Previously closed item resurfaced (regression).
Operator-blocked	NEW. Cannot proceed without an operator action that the agent + available tooling cannot perform.
Fixed (→ Fixed.md)	Closure migration marker per § 11.4.15 + § 11.4.15.

. **Self-resolution exhaustion mandate.** BEFORE classifying any item as `Operator-blocked`, the agent MUST verifiably exhaust every self-resolution path applicable to the project:

- (a) **CLI / ADB / SSH / API access** – can the work be done via any access the agent already has? (e.g. `adb shell pm clear`, `gh api`, `aws ssm`, `kubectrl exec`, `psql`).
- (b) **Subagent delegation** – can a specialised subagent (per § 4.1.1) reach further than the primary agent? (e.g. `general-purpose` for multi-step probing, `code-reviewer` for architectural blockers, `iac-reviewer` for infra access).
- (c) **Existing tooling** – is there a script / helper / library already in the repo that handles this case? (e.g. `scripts/testing/<helper>.sh`, project-specific automation).
- (d) **Captured fallback** – can the blocker be sidestepped with a synthetic event, test asset substitution, mock, or topology-aware SKIP per § 4.1.2? (e.g. simulate the input, substitute a stand-in fixture, downgrade to a pre-build gate).
- (e) **Documentation + research** – per § 4.1.3, has the agent consulted external sources (official docs, vendor guides, open-source codebases, issue trackers) for a self-resolution pattern the community has already solved?

Only AFTER each applicable path is **verifiably attempted and documented as exhausted** is `Operator-blocked` the correct classification. "I didn't try X because I assumed it wouldn't work" is a § 4.1.4 no-guessing violation AND a § 4.1.5 self-resolution violation.

. **Operator-Block-Details mandate.** Every `Status: Operator-blocked` item MUST carry an `**Operator-Block-Details:**` line within 3 non-blank lines of its heading (mirroring the

§ . . Status placement pattern). The content MUST state ALL of:

- **WHAT** – the specific concrete action the operator must perform (verb + object + target system). Generic "the operator must investigate" is forbidden – name the action.
- **WHY** – the alternatives that were exhausted, enumerated. Each exhausted path from § . . . above gets a one-line statement of what was attempted and why it could not unblock the item. "Subagent delegation: tried general-purpose with X scope – blocked because Y" is the bar.
- **UNBLOCK CONDITION** – the observable signal that the operator has completed the action (e.g. "operator confirms hardware rev-B installed", "operator pastes new API key into scripts/testing/secrets/.foo.env ", "operator force-pushes PR merge to upstream"). Without this, the agent cannot automatically detect when to re-evaluate.
- **WHO** – the handle / contact / pointer to the document where operator-side details live if questions arise (e.g. "operator: @username", "see docs/operator/<topic>.md ", "vendor: support-id "). Operator-blocked without WHO is operationally untrackable.

. **Issues_Summary inclusion as sortable axis.** The auto-generated Issues_Summary.md (per § . . . + § . . .) MUST include Operator-blocked as a first-class Status value in the table – the column header is unchanged but operators can sort / filter on it deterministically. The summary generator MUST also emit a count of Operator-blocked items in any rollup line / header.

. Periodic re-evaluation requirement. Items in `Operator-blocked` status MUST be re-evaluated every Nth tag-cycle (project-defined, recommended every 10th tag cycle). Operator dependencies change over time: CI may have been added, hardware may have been swapped, automation may have matured, vendor APIs may have unlocked. An item that was correctly Operator-blocked at tag T may be self-resolvable at tag T+1. Re-evaluation MUST follow the same § 4.1.2.1 exhaustion checklist and document each re-attempt.

. Anti-bluff layer. A fake `Operator-blocked` (classification applied without exhausting § 4.1.2.1 alternatives) is a § 4.1.2.1 covenant violation at the planning layer – equivalent severity to a PASS-bluff at the testing layer. The agent's commit message introducing or maintaining an `Operator-blocked` classification MUST contain evidence of self-resolution attempts (the explicit "Attempted: a – exhausted because X; b – exhausted because Y; c – exhausted because Z" form).

. Pre-build gates (recommended, per consuming project):

- **CM-ITEM-OPERATOR-BLOCKED-DETAILS** – every heading whose Status line equals `Operator-blocked` has an `**Operator-Block-Details:**` line within 10 non-blank lines. Paired mutation strips the details line → gate FAILs.
- **CM-OPERATOR-BLOCKED-SELF-RESOLUTION-AUDIT** – every NEW `Operator-blocked` item introduced in a commit carries an "Attempted: ..." audit trail (either in the Issues.md entry body OR in the commit message). Paired mutation introduces an `Operator-blocked` item without the audit trail → gate FAILs.

Propagation. This anchor is a § 4.1.2.1 -classified universal rule – it composes with § 4.1.2.1 (Status tracking), § 4.1.2.1 (Type tracking), § 4.1.2.1.

(Fixed-document column alignment), § . . . (auto-generated docs sync), § . . . (subagent delegation as a self-resolution path), § . . . (research-before-implementation as a self-resolution path), § . . . (no guessing about what the operator "would have to do").

Propagation gate CM-COVENANT-114-21-PROPAGATION (when implemented per consuming project) enforces the anchor's presence in every CLAUDE.md / AGENTS.md across the parent + every owned submodule + every dependency.

No escape hatch. Items that legitimately need operator intervention are real and unavoidable – hardware swaps, vendor contracts, physical-world inputs, account credentials, irreversible business decisions. But the classification must be **earned** via documented exhaustion. An Operator-blocked heading without an audit trail of self-resolution attempts is a § . . . release blocker regardless of context.

§ . . . – Document-sync commit discipline (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"For Issues, Issues_Summary and Fixed docs (and all its exports) we MUST commit and push (only them) as soon as they are updated (synced). Continuation doc and other similar / relevant documentation MUST be committed and pushed too! Like this we will always have up to date status of working items without need to do full commit and push of everything if that is not possible to do!"

Classification: § . . -classified **universal** – every project that tracks work items through an Issues / Fixed lifecycle has the same problem, so the discipline lives in the Constitution and every consuming project implements its own wrapper.

The defect this anchor closes. Until this anchor existed, the only way to push a status update was the full-tree commit wrapper. When the working tree had unrelated heavy churn (large submodule diffs, in-flight rebase, partial-network conditions, ongoing build), operators had two bad options: (a) wait until the heavy work finishes – which can be hours – leaving the doc-status stale to every other agent; (b) skip the doc-status update entirely – which is a § . . documentation-drift violation. Neither option is acceptable.

The mandate. Every project under this Constitution that ships a status-tracking doc set **MUST** provide a **lightweight commit path** distinct from the full-repo commit wrapper. The lightweight path stages, commits, and pushes **only** the status-tracking doc set:

- . The active issue tracker (docs/Issues.md in a consuming project's layout; the project's equivalent file elsewhere) and its HTML + PDF exports.
- . The auto-generated issues summary (docs/Issues_Summary.md) and its HTML + PDF exports.
- . The fixed-bug archive (docs/Fixed.md) and its HTML + PDF exports.
- . The auto-generated fixed summary (docs/Fixed_Summary.md) and its HTML + PDF exports.
- . The cross-session continuation document (docs/CONTINUATION.md , per § .) and its HTML + PDF exports.

- . Any auto-generated audit artifact (`docs/audit/anti_bluff_audit.md` or equivalent) that pairs with the above.

The lightweight wrapper MUST:

- . **Auto-invoke the project's export-regeneration pipeline first** – so Markdown + HTML + PDF are all synchronised before push. (Skipped only when an explicit `--no-sync` flag is passed AND the operator just ran the pipeline manually.)
- . **Stage ONLY the doc-set files** – `git add` of an explicit file list, NEVER `git add -A` (which would catch unrelated WIP churn).
- . **Use a separate flock** disjoint from the full-tree commit wrapper's lock – so the two can run in parallel without contention on disjoint file sets.
- . **Push to every parent-repo remote** (reusing the project's push-all driver where available).
- . **Exit code semantics** – `0` success, `1` validation failure, `2` lock contention, `3` nothing-to-commit (informational, not an error – so the wrapper can be wired into automation that tolerates no-op outcomes).

The lightweight wrapper MUST be invocable either standalone OR via a flag on the full-tree wrapper (e.g. `commit_all.sh --docs-only`) so operators have a single mental model.

\$ preflight inheritance. The lightweight wrapper inherits the parent project's `$ preflight` discipline – if `meta_test_*` is mid- mutation or `*.mut.bak` files exist, the wrapper refuses to run, same as the full-tree wrapper.

Pre-build gates (recommended, per consuming project):

- **CM-COMMIT-DOCS-EXISTS** – verifies the lightweight wrapper exists + is executable, its external user guide (per `$ . .`) exists, the full-tree wrapper advertises

the delegation flag, and the wrapper's doc-set enumeration lists $\geq N$ entries (N depends on project – a consuming project's set is ≥ 15). Paired mutation strips the doc-set array $\rightarrow$ gate FAILs.

Propagation. Composes with § 2.2.4 (export-sync invariant), § 2.2.5 (HTML + PDF in sync with Markdown after every edit), § 2.2.6 (item status tracking – Status lines must be visible quickly), § 2.2.7 (every script ships with an in-source doc block + an external user guide), § 2.2.8 (CONTINUATION document maintenance – the lightweight wrapper is one of the mechanisms that keeps CONTINUATION current).

No escape hatch. The discipline exists because doc-status drift is itself a § 2.2.9 PASS-bluff pattern at the documentation layer: a green-looking `Issues_Summary.html` that doesn't reflect Markdown reality is the same class of defect as a passing test that doesn't reflect runtime reality. Operators who can't run the lightweight wrapper for some reason **MUST** run the full-tree wrapper and accept the heavier cost – the doc-status drift is non-negotiable.

2 0 2 6 0 5 1 4

§ 2.2.10 – Visual-cue & grouping
mandate for Issues docs (User mandate,
– –)

Forensic anchor – direct user mandate (verbatim,
– –):

"We **MUST** make sure that Issues, Issues_Summary and Fixed docs and its exported files (PDFs and HTMLs) have one small improvement: we **MUST** introduce some background coloring for cells of item type and status! Also, we **MUST** group items by the status! ... Base colors for types would be: bug – background pale red,

task - background pale blue, feature - background pale yellow. However, changing the status affects the color of the cell of the type and the status cell background color: queued - no color, no effect, fixed - both cells pale green, in progress - only status cell is pale green, reopened - status cell is pale red, blocker - status cell is red (not pale, however MUST BE still readable), anything else please follow the same logic!"

Classification: § . . -classified **mixed** – the discipline (visual cues + status grouping for tracked-item docs) is universal across every project that maintains an Issues / Fixed lifecycle. The implementation (CSS classes, HTML post-processor, weasyprint integration) is project-specific because each project's export toolchain differs (pandoc, asciidoctor, sphinx, etc.).

The defect this anchor closes. A long, multi-column tracked-item table where every row looks identical at first glance is a § . operational-readability defect – operators have to *read* every Status column to identify what is queued vs in-progress vs fixed. With dozens to hundreds of items, this is slow enough that operators inevitably skim, and the doc effectively becomes write-only. Color cues + status grouping make at-a-glance triage possible.

The mandate. Every project under this Constitution that ships tracked-item docs (per § . .) MUST apply visual-cue coloring AND status grouping to those docs' HTML + PDF exports. The Markdown source stays free of color noise (Markdown is the canonical text), but the HTML + PDF exports MUST be enriched in a post-processing stage:

Type cell base colors:

Type	Background
Bug	pale red (#FFE5E5 or equivalent – ~ % red saturation)
Task	pale blue (#E5F0FF)
Feature	pale yellow (#FFF8DC)

Status cell colors (override the Status column; Fixed also overrides Type cell):

Status	Status cell	Effect on Type cell
Queued	no color (transparent)	keeps base type color
In progress	pale green (#D4F1D4)	keeps base type color
Ready for testing	pale green-blue (#C6EAEF)	keeps base type color
In testing	pale green-yellow (#E7F4D6)	keeps base type color
Reopened	pale red (#FFCCCC)	keeps base type color
Fixed	pale green (#A8E6A8)	also pale green
Operator-blocked / Blocker	vibrant red (#FF4444) with white text for readability	keeps base type color

Status grouping: items in the rendered table **MUST** be grouped by Status, emitting an H -class heading for each

Status group with the item count. Group order (most-actionable first; closed last):

- . Operator-blocked / Blocker
- . In testing
- . Ready for testing
- . In progress
- . Reopened
- . Queued
- . Fixed

This ordering surfaces the items requiring operator attention at the top of the doc; closed/archived items sink to the bottom.

Print-fidelity requirement. The CSS MUST declare `print-color-adjust: exact` (or the project's renderer's equivalent) so PDF exports preserve the cell backgrounds. A weasyprint default-rendered PDF that strips colors is a § 4.1.1 export-drift defect – HTML and PDF MUST be visually equivalent.

Pre-build gates (recommended, per consuming project):

- **CM-DOC-COLOR-GROUPING-DISCIPLINE** – verifies the project's CSS carries the type+status class set, the post-processor exists and is executable, the sync wrapper invokes it, AND after sync the rendered HTML carries the grouping-heading + per-cell color classes (the captured-evidence proof – without this last invariant a stub post-processor could silently pass the static gate). Paired mutation strips a Status class literal from CSS → gate FAILs.

Propagation. Composes with § 4.1.1 (export-sync invariant – HTML + PDF in sync with Markdown after every edit), § 4.1.2 (item status tracking – every active item carries a Status line from the closed-set

vocabulary), § . . . (Fixed_Summary column-aligned companion), § . . . (lightweight commit path for doc-set updates).

No escape hatch. Color cues are not optional polish – they are the operational-readability seam that converts a long tracked-item table from a write-only Markdown blob into a self-triaging document.

2026 05 14

§ . . . – Build-resource stats
 tracking mandate (User mandate,
 – –)

Forensic anchor – direct user mandate (verbatim,
 – –):

"We need to incorporate through the Containers Submodule we use to run our System build Container the following mechanism safely: during its working lifetime we MUST track in proper Markdown document exported into PDF and HTML stats on System resources use! What was the peak and the minimum EVER the Container / System building process used! ... In top of the document we MUST have this ever values. ... After this we MUST present properly sorted divided by version tags names all build processes we were running with notes if the process completed with success or not, and resources usage during the process! Besides min and max values ever and per build iteration we MUST have for each resource we have average values – average memory usage, average CPU usage, and all other parameters IO and other resources!"

Classification: § . . . -classified **mixed** – the discipline of tracking host-side build-resource usage with

min / max / mean / p per-build PLUS ever-values
across all builds in a versioned Markdown + HTML + PDF
triple is universal across every long-build project (AOSP,
kernel, large monorepo, ML model training, multi-hour data
pipelines). The implementation (registry path, monitor
script name, exporter wiring) is project-specific.

The defect this anchor closes. Long builds (multi-hour
AOSP builds, ML training runs, large monorepo CI) consume
highly variable resource profiles over their lifetime –
soong_build's GB peak coexists with quiescent
linker-only stretches at < GB. Without per-build
sampling + ever-values, operators have NO empirical basis
to (a) size containers correctly, (b) detect resource
regressions across versions, (c) prove which build
iteration was the OOM-killer's trigger, (d) reason about
parallelism caps (cf. \$. -j hard cap whose
forensic basis would have been one full Stats.md
generation earlier had this discipline been in place).
Memory pressure debugging without time-series data is the
bluff this anchor forbids.

The mandate. Every project under this Constitution with a
build exceeding minute wall-clock MUST:

- **Run a host-side resource sampler for every build.** The
sampler runs as a sibling background process (not inside
the build's own process tree), reads `/proc/meminfo` +
`/proc/loadavg` + `/proc/stat`
• `/proc/diskstats` (or platform equivalents on non-Linux)
at a fixed interval (recommended s default – fine-
grained enough to catch seconds-long peaks; coarse
enough to keep itself < % CPU and < MB RSS).
Samples written as TSV.
- **Compute per-build summaries on stop.** Per metric
(memory used, CPU%, load average, disk read/write): **min**,
max, **mean**, **p** . **p** is the th-percentile sample

value – required because mean alone hides extended-peak behaviour (a h build with one GB soong_build peak has mean $\approx$ GB but p $\approx$ GB; the latter is what operators must size for).

- . **Append per-build summaries to a registry.** One TSV row per build (build_id, status SUCCESS/FAIL/UNKNOWN, start/end timestamps, per-metric min/max/mean/p, total disk read/write). The registry is the single source of truth – the Markdown report is derived.

- . **Maintain ever-values across ALL builds.** Min / max / mean across every tracked build, computed on every Markdown regeneration from the registry. Surface at the **top** of the report per User mandate verbatim.

- . **Markdown + HTML + PDF triple.** Stats.md (auto-generated), exported to Stats.html + Stats.pdf through the project's normal export pipeline. Triple stays in sync per \$ Committed via the project's lightweight doc-sync wrapper per \$

- . **Sort per-build entries by recency (most recent first) AND, where applicable, group by version tag.** Operators reading the report want the latest build's profile at the top and the historical sweep underneath.

- . **Track status per build.** SUCCESS / FAIL / UNKNOWN – and if FAIL, capture the reason (exit code, OOM, ANR, etc.). A FAIL build's resource profile is the most operationally interesting one in the report – never silently drop FAIL entries.

- . **Performance budget for the sampler itself.** < MB RSS, < % CPU. The monitor MUST NOT itself contribute to the resource pressure it is measuring (Heisenberg-class observer effect at scale). Pure /proc reads + awk arithmetic is the reference implementation; vmstat /

pidstat acceptable; psrecord / heavy-Python-import
samplers are forbidden.

. **Survive build failures.** The stop hook MUST be called from both the success AND failure paths of the build wrapper so FAIL builds still produce a row. Hooking only the success path is a PASS-bluff at the telemetry layer.

Pre-build gate (recommended, per consuming project):

- **CM-BUILD-RESOURCE-STATS-TRACKER** – verifies (a) the monitor script exists + executable, (b) the Stats.md target file exists, (c) the build wrapper invokes start AND stop, (d) the doc-sync wrapper enumerates the Stats.
{md,html,pdf} triple, (e) the exporter advertises the build-stats slug. Paired mutation hides the monitor script aside → gate FAILs.

Propagation. Composes with \$ (export-sync invariant – HTML + PDF in sync with Markdown),
\$ (script-documentation discipline – the monitor ships with an external user guide), \$
(lightweight doc-sync commit wrapper enumerates the Stats.
{md,html,pdf} triple), \$ / \$ / \$
(host-safety + containerized-build envelope whose forensic anchors are the empirical motivation for this telemetry).

No escape hatch. Build-resource tracking is not optional polish – it is the operational-evidence seam that converts post-mortem "the build OOM'd, no idea why" into "build X at git-SHA Y peaked at Z GB at minute N, which is W% above the rolling p ". The discipline pays for itself the first time it diagnoses a build-resource regression.

\$ – Full-Automation-Coverage
Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"Make sure that every feature, every functionality, every flow, every use case, every edge case, every service or application, on every platform we support is covered with full automation tests which will confirm anti-bluff policy and provide the proof of fully working capabilities, working implementation as expected, no issues, no bugs, fully documented, tests covered! Nothing less than this does not give us a chance to deliver stable product! This is mandatory constraint which MUST BE respected without ignoring, skipping, slacking or forgetting it!"

Operative rule. For every consuming project under this Constitution, no feature, functionality, flow, use case, edge case, service, or application on any supported platform may be considered **deliverable** until it is covered by automation tests that collectively prove six invariants:

- . **Anti-bluff posture** (per § . + § .): every assertion carries captured runtime evidence; metadata-only / configuration-only / grep-based / absence-of-error PASSes are forbidden.
- . **Proof of working capability:** the user-visible behaviour is exercised end-to-end on the target platform topology (per § . .), not in a mock or in-memory facsimile.
- . **Working implementation as expected:** assertions match the product's documented promise (in user manual, README, specs), not an implementation-detail back-door.
- . **No issues, no bugs:** the test suite is the canonical seam for surfacing defects – passing without producing defect signals means defects do not exist or have been previously surfaced, tracked (per § . . / § . .), and closed.

- **Fully documented:** the feature has a user-facing doc entry (per § 4.1.1 for scripts, and project-level user-manual coverage for everything else); the doc is kept in sync with the tests (per § 4.1.2).
- **Tests-covered (the four-layer floor per § 4.1.3):** pre-build presence-of-change gate, post-build artifact-shipped gate, runtime / integration / on-device gate, AND meta-test paired mutation that proves the runtime gate catches the break.

Cross-cutting reach. This mandate is **universal** – it applies to every project consuming this Constitution, every feature surface they expose (HTTP endpoints, CLI commands, slash commands, IPC channels, plugins, hooks, MCP servers, agents, providers, integrations, GUI flows, mobile flows, scheduled jobs), and every supported platform (Linux, macOS, Windows, iOS, Android, AuroraOS, HarmonyOS, embedded, containers, headless servers, kiosk, etc.). A project that ships a feature without satisfying all six invariants is **not delivering a stable product**, irrespective of how green its summary line looks.

Coverage audit. Consuming projects **MUST** publish a coverage ledger (matrix of: feature × platform × invariant-.. × status) that is regenerated as part of the release-gate sweep. The ledger itself is documented (per § 4.1.4.1 / § 4.1.4.2) and committed via the § 4.1.4.3 lightweight doc-sync wrapper. Gaps in the ledger (UNCONFIRMED: / PENDING: / BLOCKED: cells) MUST cite a tracked work item per § 4.1.4.4 + § 4.1.4.5; rows that quietly omit a platform are § 4.1.4.6 violations.

Composition. § 4.1.4.7 explicitly stacks on top of § 4.1.4.8 (four-layer test-coverage floor), § 4.1.4.9 (false-positive immunity), § 4.1.4.10 (positive-evidence-only validation), § 4.1.4.11 (FAIL-bluffs forbidden),

§ 11.4.15 (recorded-evidence requirement), § 11.4.16 (per-environment-topology dispatch), § 11.4.17 (no-guessing), § 11.4.18 / § 11.4.19 (status + type tracking), § 11.4.20 (universal-vs-project classification – this rule is universal), § 11.4.21 (script documentation), § 11.4.22 (subagent delegation when the work is multi-step), § 11.4.23 (lightweight doc-sync). It does NOT supersede them; it forecloses the loophole "the feature works locally for me, ship it".

Classification: universal (per § 11.4.24). No escape hatch. Severity-equivalent to a § 11.4.25 PASS-bluff at the release-gate layer – a project claiming "Done" without honoring § 11.4.26 is making a false claim regardless of how the work-item tracker labels the row.

§ 20.2.6 – Constitution-Submodule Update Workflow Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"Every time we add something into our root (constitution Submodule) Constitution, CLAUDE.MD and AGENTS.MD we MUST FIRST fetch and pull all new changes / work from constitution Submodule first! All changes we apply MUST BE committed and pushed to all constitution Submodule upstreams! In case of conflict, IT MUST BE carefully resolved! Nothing can be broken, made faulty, corrupted or unusable! After merging full validation and verification MUST BE done!"

Operative rule. Before ANY agent or operator modifies the constitution/Constitution.md , constitution/CLAUDE.md , or

`constitution/AGENTS.md` files of a project that consumes this Constitution as a submodule, the agent or operator MUST execute the following pipeline in order, with NO step skipped:

- **Fetch + pull first.** From inside the `constitution/` submodule worktree, run `git fetch <every-configured-remote>` followed by `git pull --ff-only origin <branch>` (or `--rebase` if non-FF-mergeable; **never** `--strategy=ours` / `--allow-unrelated-histories` without explicit operator authorization). The submodule MUST be at upstream tip BEFORE any local edit is applied.
- **Apply the change.** Edit the relevant file(s). The edit MUST classify itself per § 1.1.1 (universal vs project-specific) – only universal additions belong in the constitution submodule; project-specific clauses belong in the consuming project's own governance files. The edit MUST include the verbatim user mandate (if it originated from one) as a forensic anchor.
- **Validate before commit.** Run the constitution submodule's `meta_test_inheritance.sh` (or equivalent) to confirm the inheritance chain still resolves. Verify no governance file was left with merge-conflict markers (`<<<<<<` , `=====` , `>>>>>>`). Verify all three governance files (Constitution + CLAUDE + AGENTS) cross-reference the new clause consistently.
- **Commit + push to ALL upstreams.** Stage only the governance files (NEVER `git add -A` inside the constitution submodule – stray local artefacts MUST NOT enter governance). Commit with a message that cites the user mandate (verbatim quote) + the classification line per § 1.1.1. Push to **every** configured remote of the constitution submodule. A commit that lives on one upstream but not others is a § 1.1.1 violation equivalent to a § 1.1.1 multi-upstream-push violation.

- **Conflict resolution.** If `pull --ff-only` reports non-fast-forward, the merge MUST be performed carefully: inspect both sides, preserve the union of governance content (no clause silently dropped), re-classify per § 4.1.2, validate per step 4.1.2.1. Force-push to "make conflicts go away" is FORBIDDEN (§ 4.1.2.1). Nothing about the constitution may be broken, made faulty, corrupted, or rendered unusable by the merge.
- **Post-merge validation + verification.** After the push lands, re-clone (or `git submodule update --remote --init`) in a throwaway worktree and re-run the consumer project's inheritance-cascade verifier (e.g. `scripts/verify-governance-cascade.sh`) to confirm the new clause reaches every owned submodule per CONST-4.1.2.1. Any cascade gap MUST be closed in the same change-window.
- **Update the consuming project's pointer.** The consuming project's `.gitmodules`-tracked submodule pointer MUST be bumped to the new constitution HEAD in the SAME commit as any downstream cascade work; out-of-sync submodule pointers are a § 4.1.2.1 violation.

Operational scope. This workflow applies regardless of who initiates the change (operator, primary agent, subagent per § 4.1.2.1, automated lint suggestion). The workflow CANNOT be shortcut by "I'll fetch later" or "I'll push to the other upstreams in the next commit". The constitution is the **single source of truth** for every project that imports it; allowing it to fragment across upstreams is the structural equivalent of a § 4.1.2.1 PASS-bluff at the governance layer.

Cross-cutting reach. § 4.1.2.1 composes with: § 4.1.2.1 (single-entripoint commit wrapper), § 4.1.2.1 (multi-upstream push), § 4.1.2.1 (submodule changes propagate through submodule commits first), § 4.1.2.1 / § 4.1.2.1 (data-safety + force-push authorization), § 4.1.2.1 (universal-vs-project classification), § 4.1.2.1 (lightweight doc-

sync), § 1.4.17 (full-automation coverage – the post-merge validation step is itself a form of automation coverage). It does NOT supersede them.

Classification: universal (per § 1.4.17). No escape hatch. A constitution-submodule change that violates § 1.4.17 is a release blocker for every consuming project, equivalent in severity to a force-push without § 1.4.17 authorization.

§ 1.4.15 – Submodule-Dependency-Manifest Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"We MUST HAVE mechanism for each Submodule to determine / know what are its Submodule dependencies so new projects or palces we are incorporate them can add these Submodules to the project root and make them available! Suggested idea is configuration file with expected Submodules Git ssh urls perhaps? New project can read it, and recursively add each Submodule to the root of the project and install / expose it to everyone. This MUST be analyzed and applied. We MUST apply the best strategy for this which can be easily executed just by following our root Constitution, AGENTS.MD and CLAUDE.MD! Process this, extend out Constitution, AGENTS.MD and CLAUDE.MD with mandatory instructions and then process project root and all Submodules deep recursively so proper configuration Submodules dependency files are created! Document EVERYTHING and cover with all supported test types! Any kind of bluff is strictyl forbidden! All wotk MUST

come with mechanism for validation and verification by creating proper proofs for all critical points!"

Operative rule. Every owned-by-us submodule MUST ship a machine-readable, version-controlled **dependency manifest** at the canonical path `helix-deps.yaml` (or `helix-deps.json` / `helix-deps.toml` if the submodule's ecosystem strongly prefers a different serialisation – but only one canonical file per submodule, and its name MUST appear in the submodule's `README.md` + governance trio).

Manifest schema (CONST- (C)-aligned):

```
# helix-deps.yaml
# Declares this submodule's own-org dependencies. Consuming projects
# read this manifest recursively and add each declared dependency to
# their root at <root>/<name>/ or <root>/submodules/<name>/ per
# CONST-051(C). Nested own-org submodule chains are FORBIDDEN – this
# manifest is the bridge. (Sole exception per §11.4.28(C): the
# constitution submodule MAY host depth-1 reusable-engine submodules
# under constitution/submodules/<name>/, each shipping its own
# helix-deps.yaml – never a recursive chain.)

schema_version: 1
deps:
  - name: Challenges                                     # canonical name
    (will become snake_case per §11.4.29)
    ssh_url: git@github.com:vasic-digital/Challenges.git
    ref: main                                           # branch or
    pinned tag
    why: "Cross-cutting Challenges + Panoptic browser harness"
    layout: flat                                       # 'flat' =
    <root>/<name>/; 'grouped' = <root>/submodules/<name>/
  - name: HelixQA
    ssh_url: git@github.com:HelixDevelopment/HelixQA.git
    ref: main
    why: "Anti-bluff QA orchestration + autonomous-session driver"
    layout: flat

transitive_handling:
  # Each declared dep itself ships a helix-deps.yaml. The incorporator
  # tooling MUST recurse – top-level project gets the union of all
  # transitively-declared deps, flattened to root.
```

```

recursive: true

# When two submodules declare the same dependency at different
# refs, CONFLICT – operator MUST resolve before incorporation
# proceeds. The incorporator aborts on conflict (never silently
# picks one).
conflict_resolution: operator-required

language_specific_subtree: false      # set true for Android/Kotlin/
    Apple                             # roots (per §11.4.29
    exception);                       # excludes inner subtree from
    # snake_case enforcement.

```

Tooling contract. A consuming project **MUST** be able to bootstrap its entire own-org dependency graph by:

- . Running `incorporate-submodule <ssh-url>` (canonical name; lives in the constitution submodule's `scripts/` or in a parent project's `bin` path).
- . The script:
 - Adds the supplied submodule at its declared canonical path (per CONST- (C) `flat` vs `grouped` layout).
 - Reads the newly-added submodule's `helix-deps.yaml`.
 - For each declared dep, checks if it already exists at the consuming project's root; if missing, recurses (`incorporate-submodule` on the dep's `ssh_url`).
 - On conflict (same name declared at different ref by two submodules), aborts with a directed error pointing the operator at the conflicting declarations.
 - Emits a final manifest-of-manifests file at `<root>/.helix-manifest.yaml` listing every submodule + its declared deps, for audit + reproducibility.

Anti-bluff guarantee. Every manifest **MUST** be paired with a **verification proof**: a Challenge script (per § 11.4.2.1 + CONST- (B)) that:

- Bootstraps a throwaway consuming project from scratch in a temp directory.
- Runs `incorporate-submodule` against the manifest under test.
- Verifies the produced submodule layout matches the manifest's declarations (every dep present at its declared layout path; no extras; no missing).
- Runs the submodule's own test suite against the bootstrapped layout; asserts pass.
- Captures wire evidence (per § 11.4.2.1) of every step.

A manifest without this verification proof is a § 11.4.2.1 violation of equal severity to a § 11.4.2.1 PASS-bluff at the dependency-graph layer.

Cascade requirement. Every owned-by-us submodule **MUST** ship `helix-deps.yaml` at its root, recursively (sub-submodules of submodules ship their own). The constitution submodule itself ships a manifest declaring its own deps (currently empty for the universal Constitution; project consumers may have project-specific extensions). The verifier (`scripts/verify-governance-cascade.sh` or its successor) **MUST** check every owned submodule for manifest presence.

Composition. § 11.4.2.1 directly enables § 11.4.2.1 / CONST- (C) flat-layout enforcement: nested own-org submodule chains can be mechanically flattened because each submodule declares what it needs, and the incorporator places those deps at the root. The § 11.4.2.1 (C) constitution-submodule carve-out does not weaken this: a permitted depth-constitution-anchored engine still ships its own `helix-deps.yaml`, so the incorporator reconstructs the graph identically – the

manifest is the bridge in both cases. Composes with §
 (manifest schema is itself tested for parse + validation),
 § (submodule-first commit discipline applies to
 manifest changes too), § (manifest changes
 regenerate any derived docs/diagrams), §
 (universal – no project-specific assumptions in the
 manifest format), § (manifest documented in
 script-doc + external user guide), §
 (subagent delegation for cross-submodule manifest
 authoring), § (manifest presence in coverage
 ledger), § (constitution-update workflow when
 extending manifest schema), § (manifest test
 matrix), § (this rule is its operational
 complement), § (manifests use snake_case
 names + canonical paths), § (manifest is
 tracked source – not a build artefact), CONST-
 (manifests cascade recursively).

Classification: universal (per §). No escape
 hatch. Severity-equivalent to a § PASS-bluff at
 the dependency-graph layer.

2 0 2 6 0 5 1 5

§ – Post-Constitution-Pull Validation Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Every time we fetch and pull new changes on
 constitution Submodule we MUST process the whole
 project and all Submodule (deep recursively) for
 validation and verification taht every single rule or
 mandatory constraint is followed and respected! If it
 is not, IT MUST BE!"

Operative rule. Whenever a consuming project's constitution submodule is fetched + pulled with **any** content change (rule addition, rule revision, gate addition, anchor reference, version bump), the consuming project **MUST** execute a full-project + recursive-submodule validation sweep **BEFORE** the new constitution **HEAD** is treated as canonical for any other work.

Validation sweep contract. The sweep is implemented as `scripts/verify-all-constitution-rules.sh` (canonical name) which:

- . Re-runs the existing governance-cascade verifier (`scripts/ verify-governance-cascade.sh`) covering every \$. + CONST- anchor across every owned submodule (recursive per CONST-).
- . For each rule whose enforcement gate is implementable programmatically (e.g., CONST- `.gitignore -pattern audit`; CONST- (C) nested-own-org-chain audit; CONST- case-conformance audit; CONST- (A) mock-path-from-production- code audit; CONST- anti-bluff smoke-scan), the sweep runs the corresponding gate against the post-pull tree.
- . Any failure produces a directed FAIL entry naming the rule (e.g., `FAIL: CONST-053 - *.log tracked at the-project/foo.log`)
 - the canonical fix.
- . Failures populate the project's Issues tracker per \$. . with Status: `Reopened` (since a previously-passing audit now fails) and Type: `Bug` (since real codebase state violates the constitution).
- . The agent or operator **MUST** resolve every FAIL before treating the new constitution **HEAD** as canonical – closure of each reopened item per the existing \$.

§ . . – .gitignore + No-Versioned-Build-Artifacts Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"every project module, every Submodule, every service and application MUST HAVE proper .gitignore file! We MUST NOT git version build artifacts, cache files, tmp files, main .env file(s) or any files containing sensitive data, API keys or token! Any build derivative which we can recreate by executing proper mechanism for generating MUST NOT be versioned! We MUST pay attention what is going to be committed every time we are preparing to execute commit! If any violation is detected it MUST be fixed before commit is executed!"

Operative rule. Every project module, owned-by-us submodule, service, and application under this Constitution MUST ship a proper `.gitignore` covering its full set of forbidden-from-version- control patterns. The following file/directory classes MUST NEVER appear in version control:

- **Build artefacts** – anything produced by a build / compile / package step that can be regenerated from sources:
 - Binaries (`/bin/` , `/build/` , `/dist/` , `/out/` , `target/` , `*.exe` , `*.dll` , `*.so` , `*.dylib` , `*.a` , `*.o` , `*.class` , `*.pyc`).
 - Generated source files where the generator is committed (e.g. protobuf `.pb.go` may be checked in OR ignored – project decides – but never both).
 - Bundled assets where the bundler is committed.

- . **Cache files** – anything a tool regenerates on demand:
 - `__pycache__/` , `.pytest_cache/` , `.mypy_cache/` , `.ruff_cache/` , `node_modules/` , `.next/` , `.nuxt/` , `.cache/` , `.gradle/` , `.idea/cache/` , `.vscode-test/` , `target/` , `.terraform/` , `language-server` caches.
- . **Temp files** – `*.tmp` , `*.swp` , `*.sw0` , `*~` , `.DS_Store` , `Thumbs.db` , IDE backup files, `*.orig` , `*.rej` .
- . **Sensitive-data files** – anything containing credentials, tokens, keys, or personal data:
 - `.env` , `.env.*` (allow `.env.example` / `.env.sample` / `.env.template` with placeholder values only – never real secrets even as examples).
 - `*.pem` , `*.key` , `*.crt` , `*.p12` , `*.pfx` , `id_rsa*` , `id_ed25519*` , `*.kdbx` , `.netrc` , `secrets/` directory trees.
 - `api_keys.sh` , any file containing `BEGIN PRIVATE KEY` , credential tokens, or session cookies.
- . **Generated reports / logs** – `*.log` , `coverage.out` , `*.coverage` , `htmlcov/` , `*.gcda` , `*.gcno` , `screenshots/` recordings that aren't reference assets.
- . **OS / IDE / personal-state files** – `.idea/` , `.vscode/` (except shareable `.vscode/settings.json` etc. if explicitly project-shared), `.history/` , `.svn/` , editor-state files.
- . **Test playback artefacts** (`$ . .`) – every test's runtime capture goes under an ignored evidence directory unless the capture IS the reference asset.

Anti-bluff invariant. Per `CONST-` / `$ .` the presence of a `.gitignore` line alone is not sufficient – the project MUST also verify that no file matching the forbidden patterns is currently tracked. A `.gitignore` that

11 4 34
lists *.log while app.log sits tracked in the repo is a
\$. . violation of equal severity to no
.gitignore at all.

Pre-commit attention. Every commit author (human OR agent) MUST inspect git diff --staged and git status BEFORE executing the commit. If a staged path matches any forbidden class, the commit MUST be aborted and the issue fixed (un-stage the path, add to .gitignore, scrub if already-tracked). Gate CM-GITIGNORE-PRECOMMIT-AUDIT : pre-commit hook (or commit wrapper in \$) inspects every staged path against the forbidden-pattern matrix; hits abort the commit with a directed error pointing the operator at the specific offending path and the canonical fix. Paired mutation (\$.): stage a *.log → gate FAILs.

Cascade reach. This rule applies recursively to every owned-by- us submodule per CONST- and to every owned-submodule's nested non-owned trees within their working copy. The .gitignore files themselves are project-specific in their concrete patterns but universally bound to the rule's normative force.

11 4 34
Secret-leak intersection. \$. . composes tightly with \$. . (Credentials-handling mandate / CONST-) and \$. . A .env leak is BOTH a \$. . violation (build/sensitive file versioned) AND a \$. . violation (credential leak). The combined severity is **release-blocker requiring rotation + post-mortem** per \$. . .

Coverage of "recreatable" content. When in doubt about whether something is a build derivative: if there exists a documented, scripted, or automated mechanism that recreates the file from sources, the file is a build derivative and MUST be ignored. The committed sources MUST include the generator (Makefile target, npm script,

codegen invocation), so any downstream consumer regenerates the artefact on demand.

Classification: universal (per § 11.4.17). No escape hatch beyond the explicit exceptions enumerated above (e.g., `.env.example` placeholder files). Severity-equivalent to a § 11.4.17 PASS-bluff at the repository-hygiene layer – a tracked build artefact silently drifts vs. its source over time, producing the same class of "works-on-my-machine" surprise that the § 11.4.17 anti-bluff covenant forbids at the test layer.

Composition. § 11.4.17 composes with § 11.4.17 (test coverage – ignored files don't get spuriously included in coverage stats), § 11.4.17 (commit-wrapper enforces the pre-commit gate), § 11.4.17 (destructive-operation safeguards – never `git clean` an operator's working tree to "fix" a violation), § 11.4.17 (credentials-handling), § 11.4.17 (auto-generated docs sync – generators committed, outputs ignored only if recreatable on demand from the same generator), § 11.4.17 (universal – applies to every consuming project), § 11.4.17 (script docs – generator scripts documented), § 11.4.17 (subagent delegation for cross-submodule `.gitignore` audit sweeps), § 11.4.17 (coverage ledger lists `.gitignore` presence per submodule), § 11.4.17 (constitution-update workflow – the constitution submodule's own `.gitignore` complies), § 11.4.17 (no-fakes: test fixtures that ARE the reference asset are tracked; runtime captures are ignored), § 11.4.17 (submodules-as-equal-codebase: each submodule's `.gitignore` audited on equal basis), § 11.4.17 (the renamed paths in the lowercase migration MUST also be properly ignored where applicable), CONST-11.4.17 (recursive governance reach).

§ . . – Lowercase-Snake_Case-Naming Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"naming convention for Submodules and directories (applied deep into hierarchy recursively) – all directories and Submodules MSUT HAVE lowercase names with space separator between the words of '_' character (snake-case)! All existing Submodules and directories which are not following this rule MUST BE renamed! However, since this will most likely break some of the functionalities renaming we do MUST BE applied to all references to particular Submodule or directory! Everywhere where particular Submodule directory are referenced proper updates MUST BE applied

- all configuration files, documentation and relevant materials, links to Submodules and directories, source code that points to them, etc. There MUST BE reasonable exceptions for this rules – source code for programming languages or Submodules which apply different naming convention – Android, Java, Kotlin and others. Root directory for such applications, services or Submdoules can follow OUR convention, but EVERYTHING inside still MUST follow language / technology specific rules! We apply this rules per common sense basis and it MUST NOT be the cause of bigger issues such as technology breaking! Upstreams directory which all of our projects and Submodules have MUST BE renamed to the lowercase letters too, however root project containing the install_upstreams system command (it is exported in out paths in our .bashrc or .zshrc) MUST BE updated to fully work with both Upstreams and upstreams

directory. That change if it is not already applied MUST BE done, committed and pushed! ... NOTE: Rules lowercase / snake-case do apply to all project files as well and references to it and from them! Every change done MUST BE covered with all supported test types, full automation tests for validation and verification and fully applied and followed anti-bluff policy and all anti-bluff rules!"

Operative rule. Every directory, submodule, and file under the parent project's working tree MUST use a **lowercase, snake_case** name (ASCII letters / digits / underscores, words separated by `_`). Existing names that violate the rule (`ExampleModule/` , `Challenges/` , `Containers/` , `ExampleAgent/` , `HelixQA/` , `Security/` , `Github-Pages-Website/` , `Upstreams/` , `Dependencies/` , etc.) MUST be renamed as part of the migration window opened by this clause. Every reference in the codebase MUST be updated atomically with the rename: configuration files, documentation, user manuals, diagrams, scripts, source-code imports, links, governance files. **Reference drift after a rename is a \$. . violation** of equal severity to the rename itself.

Exceptions (common-sense scope). The rule MUST NOT break language-/technology-specific conventions:

- **Programming-language source roots** that mandate a specific case (Java / Kotlin package paths, Android resource folders, Apple framework directories, C / Swift project layouts) keep their language-mandated names. The submodule's root directory follows our convention; the language-specific subtree inside follows its own.

- **Vendor / upstream submodules** (third-party orgs not in our owned set) keep their upstream-mandated names – we MUST NOT rename a third party's repo.
- **Build-tooling artefacts** (`node_modules/` , `__pycache__/` , `.git/` , `target/` , `build/` , `bin/`) keep their tool-mandated names.

When in doubt, the test "does renaming break the technology?" trumps the snake_case rule. The \$. . spirit is operator- ergonomics + reference- discoverability, not pedantic uniformity at the cost of technology compatibility.

Upstreams/ → upstreams/ transition. The constitution submodule's installer (`install_upstreams.sh`) – exported on operator paths via `.bashrc` / `.zshrc` – MUST support **both** `Upstreams/` and `upstreams/` directory layouts during the migration window, so existing checkouts keep working while consuming projects rename at their own pace. The installer reads whichever directory exists; if both exist the lowercase wins. After every project under this Constitution has migrated, the uppercase fallback MAY be retired by a deliberate amendment, but the migration window remains open as long as ANY owned project still ships the uppercase form.

Project-Toolkit Upstreamable submodule synchronisation.

The Upstreamable / Project-Toolkit machinery that propagates governance into every consuming project MUST be fetched + pulled before any rename batch, and MUST itself comply with this rule. Any Upstreamable submodule lacking BOTH-directory support is a release blocker for the rename program.

Test coverage of renames. Every batch of renames MUST ship with: (i) a regression test that verifies every reference to the renamed entity now resolves to the new name (no stale references left); (ii) a full CONST- (B) test-

type matrix run against the post-rename tree; (iii) anti-bluff (CONST-) wire-evidence captured during the runtime verification. A rename batch without all three is a § . . violation.

Cascade reach. This rule applies recursively through every owned-by-us submodule layer (per CONST-) and applies equally to the consuming project and its owned submodules (per CONST-). Per CONST- (C) – dependencies at the parent root – the renamed paths MUST be the only canonical location; old-name aliases remain only as transitional symlinks if absolutely required, and those symlinks MUST be removed on next-N-cycle review (project-configurable, recommended ≥ release cycles).

Phased execution. Because the rename touches every reference, the execution MUST be planned as fine-grained phases per the operator's explicit instruction: comprehensive brainstorming, phase-divided plan, fine-grained tasks/subtasks with enormous detail, every change covered by every applicable test type. The phases run in parallel with mainstream work (§ . . subagent delegation is the natural fit for cross-submodule rename sweeps).

Classification: universal (per § . .). No escape hatch beyond the explicit common-sense exceptions enumerated above. Severity-equivalent to a § . PASS-bluff at the reference-integrity layer – a half-completed rename that leaves broken references is worse than no rename at all because it silently breaks consumers.

Composition. § . . composes with § (four-layer floor for every rename batch), § . (paired mutation: rename without updating a reference → gate FAILs), § . . (regenerate auto-generated docs on rename), § . . (universal – no project-specific assumptions), § . . (script-doc-sync after

renamed script paths), § . . (subagent delegation
 for cross-cutting rename sweeps), § . . (every
 renamed path appears in coverage ledger), § . .
 (constitution-submodule rename pipeline), § . .
 (rename-touch test types apply across the matrix),
 § . . (owned submodule renames cascade per
 CONST-), CONST- (recursive governance reach).

§ . . – Submodules-As-Equal-Codebase + Decoupling + Dependency-Layout Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"All existing Submodules in the project that we are
 controlling and belong to some our organizations
 (vasic-digital, HelixDevelopment, red-elf,
 ATMOSphere , Bear-Suite,
 BoatOS , Helix-Flow, Helix-Track, Server-
 Factory – we can ALWAYS check dynamically using GitHub
 and GitLab CLIs) are equal parts of the project's
 codebase! We MUST work on that code as much as we do
 with main project's codebase! All on equal basis!
 Equally important! We MUST take it into the account,
 analyze it, extend it, create missing tests, do full
 testing of it, fill the gaps (if any), fix any issues
 that we discover or they pop-up, write and extend the
 documentation, user guides, manulas, diagrams, graphs,
 SQL definitions, Website(s) and all other relevant
 materials! We MUST NEVER modify Submodules to bring
 into them any project specific context since they all
 MUST BE ALWAYS fully decoupled, project not-aware,
 fully reusable and modular (by any other project(s)),
 completely testable! All Submodule dependencies that
 are used by Submodule MUST BE accessed from the root of

the project! We MUST NOT have nested Submodule dependencies but accessing each from proper location from the root of the project – directly from project's root project_name/ submodule_name or some more proper structure project_name/ submodules/submodule_name! This MUST BE heavily enforced and respected so no chaos and mess is created with various dependencies we may have!"

Operative rule. Three cooperating invariants govern every consuming project's relationship with its owned-by-us submodules (those whose upstream `origin` lives under one of the operator- listed orgs – `vasic-digital` , `HelixDevelopment` , `red-elf` , `ATMOSphere1234321` , `Bear-Suite` , `BoatOS123456` , `Helix-Flow` , `Helix-Track` , `Server-Factory` – or any additional org the operator subsequently authorises, the canonical list discoverable at any time via `gh org list` / `glab` / the orgs' public APIs):

(A) Equal-codebase invariant. Every owned-by-us submodule is an **equal part** of the consuming project's codebase. The consuming project's engineering practice – analysis, extension, test creation, gap-filling, bug-fix, documentation (user manual, guides, diagrams, graphs, SQL definitions, website pages, any other authoring surface) – applies to each owned submodule on equal basis. A round of work that improves only the main project while leaving an owned-submodule deficiency unaddressed is a § . . violation, severity-equivalent to a § . . . PASS-bluff at the project-scope layer. Coverage ledgers (§ . .) MUST list every owned submodule as an in-scope target. CONST- (recursive cascade) governs the propagation of governance changes; § . . is the engineering-content counterpart that mandates the same attention to non-governance content.

(B) Decoupling / reusability invariant. Owned submodules **MUST** remain **fully decoupled** from any specific consuming project. No project-specific context, hardcoded paths, hostnames, asset names, naming schemes, or runtime assumptions may be introduced into an owned submodule's source tree. Every owned submodule **MUST** be:

- **Project-not-aware** – its code, tests, and docs make no reference to which parent project consumes it.
- **Fully reusable** – any future Helix-family or third-party project must be able to consume the submodule unmodified.
- **Modular** – its public surface is the only documented integration contract; internal layout may evolve without breaking consumers.
- **Completely testable** – every public surface has standalone tests (per the § 4.2.7.100 %-test-type matrix) that pass when the submodule is checked out as a standalone repo, without any parent-project rigging.

A commit that adds `project_name/...` strings, hostnames belonging to a specific deployment, or other parent-project context to a submodule's source tree is a § 4.2.7.100 violation. The honest path when a submodule needs information from the parent project is configuration injection (env var, config file, constructor parameter) – never a hardcoded reach into the parent's tree.

(C) Dependency-layout invariant. Every dependency that an owned submodule itself consumes **MUST** be accessible **from the root of the parent project** at one of two canonical paths:

```
<project_root>/<submodule_name>/          # flat layout
<project_root>/submodules/<submodule_name>/ # grouped layout
```

Nested-submodule chains are FORBIDDEN. A submodule MUST NOT have its own `.gitmodules` entries that pull in further owned- by-us repos (transitive own-org submodule recursion). Every dependency required by submodule X MUST be added to the parent project at the canonical path above; X reaches it via documented import / SDK path / runtime resolver – never via its own nested submodule pointer.

EXCEPTION – constitution-submodule-anchored reusable engines (operator Option- decision,
– – ; universal per § . .). The single scoped carve-out to the nested-chain prohibition: THE CONSTITUTION SUBMODULE ITSELF (the canonical root per § . . , where reusable engines are anchored) MAY host its own reusable- engine submodules under `constitution/submodules/<name>/`, PROVIDED each such nested submodule (i) ships a § . . `helix-deps.yaml` at its root so consumers can reconstruct the dependency graph, AND (ii) declares ZERO further own-org submodules of its own – DEPTH- ONLY, never a recursive chain (a nested engine that itself nests an own-org submodule is still FORBIDDEN). This EXCEPTION applies to NO OTHER repository: arbitrary consumer submodules remain forbidden from nesting and MUST use the root-flattened `helix-deps.yaml` bridge above. The rationale for the carve-out is that the constitution submodule is the deliberate home of reusable, project-agnostic engines (§ . . (B) / § . .) – anchoring them as depth- nested submodules there, each with a manifest, keeps the version-drift risk the base rule guards against bounded (one authoritative home, one manifest, no transitive recursion) while satisfying the operator's intent that the engines be *derived from / anchored by* the constitution.

Rationale: nested own-org submodule chains cause version-drift chaos (two consumers of `LLMsVerifier` end up at

different SHAs because each parent picked a different transitive path). The flat / grouped layout makes the consuming project's submodule graph a tree-of-depth- , which any developer can audit at a glance via `git submodule status` from the project root.

Third-party submodules (not under our orgs) are exempt – they MAY appear at any depth as the upstream's structure dictates. The invariant applies only to our owned set.

Audit + enforcement.

- Gate `CM-OWNED-SUBMODULE-EQUAL-ENGINEERING` (project-side): every release-gate sweep verifies each owned submodule has current test runs, coverage entries (`$ . .`), and documentation freshness on par with the main project. Stale submodules surface as `$ . .` violations (Status: Operator-blocked or In progress per `$ . .` – never "ignored").
- Gate `CM-OWNED-SUBMODULE-DECOUPLING` (submodule-side): every owned submodule's pre-commit hook (or equivalent) greps the staged diff for parent-project names / hostnames / asset names. Hits abort the commit until refactored to configuration injection.
- Gate `CM-OWNED-SUBMODULE-LAYOUT` (project-side): the parent project's pre-merge sweep verifies (i) every owned submodule sits at `<root>/<name>/` or `<root>/submodules/<name>/`, (ii) no owned submodule contains a nested `.gitmodules` entry whose upstream is in our org list – EXCEPT the constitution submodule, which MAY carry depth- nested own-org engine submodules under `constitution/submodules/<name>/` PROVIDED each ships a `helix-deps.yaml` AND declares zero further own-org submodules (the `$ . .` (C) carve-out; a depth-≥ chain, or a nested engine missing its manifest, still FAILs), and (iii) every dependency declared by an owned submodule has a corresponding parent-project submodule

entry at the canonical path. Paired \$. mutation:
add a SECOND-level own-org .gitmodules entry inside a
constitution-anchored engine (a depth- chain) → gate
FAILs; OR drop the engine's helix-deps.yaml → gate
FAILs; restore → PASSes.

- Paired mutations (\$.) for all three gates: plant
the forbidden pattern → gate FAILs; restore → gate
PASSes.

Workflow integration. Honoring \$. . in practice:

- Engineering rounds plan in submodule-aware tranches:
"improve feature X in main; in same round, audit +
improve the X- related surface in submodule Y; cascade
governance per CONST- as usual".
- The \$. . coverage ledger row format is
extended with a submodule column so coverage rolls up
across the whole owned set.
- Cross-submodule refactors that touch shared types or
interfaces ship as a single change-window: parent +
every affected owned submodule advanced in lockstep
(\$ submodule- first-commit discipline + \$.
multi-upstream push).

Classification: universal (per \$. .). No escape
hatch. Composes with: \$ (four-layer test floor reaches
submodules too), \$. (false-positive immunity), \$
(submodule changes propagate through submodule commits
first), \$. . (universal-vs-project
classification), \$. . . (subagent delegation for
the cross- submodule audit sweep), \$. . (full-
automation coverage ledger expanded to submodules),
\$. . (constitution-update workflow's submodule-
pointer-bump step), \$. . (% -test- type
coverage applies to every submodule's standalone surface),
CONST- (recursive governance cascade). A round of
work that violates \$. . is a release blocker for

the consuming project, severity-equivalent to a § .
PASS-bluff at the codebase-completeness layer.

§ . . – No-Fakes-Beyond-Unit-Tests +
%-Test-Type-Coverage Mandate (User
mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Mocks, stubs, placeholders, TODOs or FIXMEs are allowed to exist ONLY in Unit tests! All other test types MUST interact with real fully implemented System! No fakes, empty implementations or bluffing is allowed of any kind! All codebase of the project MUST BE % covered with every supported test type: unit tests, integration tests, e e tests, full automation tests, security tests, ddos tests, scaling tests, chaos tests, stress tests, performance tests, benchmarking tests, ui tests, ux tests, Challenges (fully incorporating our Challenges Submodule – <https://github.com/vasic-digital/Challenges>). EVERYTHING MUST BE tested using HelixQA (fully incorporating HelixQA Submodule – <https://github.com/HelixDevelopment/HelixQA>). HelixQA MUST BE used with all possible written tests suites (test banks) for every applications, service, platform, etc and execution of the full HelixQA QA autonomous sessions! All required dependency Submodules MUST BE added into the project as well (fully recursive!!!)."

Operative rule. Two cooperating invariants:

(A) **No-fakes-beyond-unit-tests.** Mocks, stubs, fakes, placeholders, in-memory facsimile implementations, TODO ,
FIXME , "for now", "in production this would", or any

empty-implementation pattern are PERMITTED only inside unit-test sources (e.g., `*_test.go` files invoked WITHOUT the integration build tag; `tests/unit/`; equivalent per-language conventions). Every other test type – integration, end-to-end, full automation, security, DDoS, scaling, chaos, stress, performance, benchmarking, UI, UX, Challenges, HelixQA suites – MUST exercise the **real, fully implemented system** against real infrastructure (real databases, real HTTP endpoints, real containers, real downstream services, real captured devices). A non-unit test that imports mocks, in-memory repositories, fabricated provider responses, or placeholder structs is a § . . . violation regardless of how green its summary line looks – severity-equivalent to a § . . . PASS-bluff. The same prohibition extends to **production code**: no mock import path may be reachable from any code path that runs in production binaries. Pre-build gate `CM-NO-FAKES-BEYOND-UNIT-TESTS` scans the non-unit test trees and production trees for the forbidden patterns; paired mutation plants a fake → gate FAILs.

(B) % test-type coverage with every supported type.

Every project under this Constitution MUST cover its entire codebase with **every supported test type** the project's domain warrants:

- . **Unit** – fast, isolated, mocks permitted per (A).
- . **Integration** – multi-component, no mocks, real backing services.
- . **End-to-end (E E)** – full user-flow exercise on target topology.
- . **Full automation** – orchestrated suites exercising every feature × platform combination (§ . . . coverage ledger).

- **Security** – authn/authz boundaries, secret-leak scans (\$. .), input-fuzzing, dependency-CVE scanning, threat- model verification.
- **DDoS** – request-flood resilience at advertised throughput tier, with rate-limit + back-pressure assertions.
- **Scaling** – horizontal + vertical scale behaviour under linear load growth, including replica add/remove transitions.
- **Chaos** – controlled failure injection (network partition, process kill, disk full, clock skew) verifying graceful degradation paths.
- **Stress** – sustained load above advertised tier, asserting bounded resource exhaustion + clean recovery.
- **Performance** – latency / throughput / tail-latency invariants vs SLO baselines.
- **Benchmarking** – micro + macro benchmark suites with historical p -drift detection (\$. . build-resource composition).
- **UI** – visual-regression + DOM-state + interaction-flow coverage on every target platform's UI surface.
- **UX** – flow-correctness + accessibility + i n + visual- cue ordering (\$. . composition).
- **Challenges** – vasic-digital/Challenges submodule fully incorporated; per-feature Challenge scripts covering real user use-cases with captured runtime evidence.
- **HelixQA** – HelixDevelopment/HelixQA submodule fully incorporated; ALL written test banks executed; full autonomous QA sessions run as part of release gates.

Per-type % coverage means: for every feature × platform cell in the \$. . coverage ledger, the cell carries a verified PASS evidence pointer for **each supported test type the cell warrants** (a CLI-only feature

warrants unit + integration + E E + full-automation + security + chaos + stress + performance

• Challenges + HelixQA; a UI feature additionally warrants UI + UX; a network service additionally warrants DDoS + scaling + benchmarking). Gaps are tracked per § . . with explicit UNCONFIRMED: / PENDING_FORENSICS: / OPERATOR-BLOCKED: reasons.

(B.) Seven-canonical-type enumeration (extension, - -). The CANONICAL SEVEN-TYPE ENUMERATION per operator mandate - - (\$. . (B.) extension of the same date) is: () UNIT, () INTEGRATION, () E E, () FULL AUTOMATION, () SECURITY, () PERFORMANCE AND BENCHMARK, () ANTI-BLUFF – the MINIMUM breadth every consuming project satisfies. The broader domain-warranted set enumerated in clause (B) above (DDoS, scaling, chaos, stress, UI, UX, Challenges, HelixQA autonomous QA sessions) is ADDED to the seven-type floor per domain, never SUBTRACTED. Applicability of each type is consumer-owned § . . DATA – a test type genuinely inapplicable is § . . SKIP-with-reason at the type-breadth ledger, never silently omitted (§ . . no-guessing at the coverage-claim layer). § . . sets the type BREADTH; § . . sets the code-coverage FLOOR across that breadth; both bind together, satisfying one is never evidence for the other (per § . . (F) honest boundary). Propagation: existing CM-COVENANT-114-27-PROPAGATION stays unchanged (literal 11.4.27 unchanged, stays GREEN); CM-NO-FAKES-BEYOND-UNIT-TESTS gate stays unchanged; NEW recommended gate CM-TEST-TYPE-BREADTH-SEVEN-CANONICAL – every consuming project's test suite covers all canonical types where applicable, honest SKIP-with-reason where not; a paired § . mutation that drops one type without a SKIP-with-reason → the gate FAILS.

Required submodule incorporation (recursive). Every project consuming this Constitution MUST add the following as Git submodules (or vendored equivalents authorised by the operator), fully recursively per CONST- / \$. .x cascade:

- Challenges – git@github.com:vasic-digital/Challenges.git
- HelixQA – git@github.com:HelixDevelopment/HelixQA.git
- Any additional functionality submodules under vasic-digital/* or HelixDevelopment/* orgs that the consuming project depends on (do not duplicate work the organisations already maintain).

Submodule pointers MUST be bumped to upstream HEAD in the SAME commit as any dependent cascade work (\$. . step). Pointer drift = \$. . violation.

HelixQA autonomous sessions. "Full HelixQA QA autonomous sessions" means: HelixQA's orchestrator drives the consuming project's running surface end-to-end, executing every test bank the project has registered with it, capturing wire evidence per check (status + body bytes + body-head + duration + screenshots for UI), and producing a session report that is itself committed via the \$. . lightweight doc-sync wrapper. A green autonomous-session report without captured wire evidence is a \$. PASS-bluff at the QA-orchestration layer.

Composition. \$ composes with: \$ (four-layer floor), \$. (positive-evidence-only), \$ (FAIL-bluffs forbidden), \$. . (recorded-evidence), \$. . (per-topology dispatch), \$. . (no-guessing), \$. . (credentials-handling – security test category), \$. . / \$. . (status + type tracking), \$. . (universal-vs-project classification), \$. . (subagent delegation for orchestrating the test-type matrix), \$. . (lightweight doc-sync for coverage

ledgers and session reports), § 11.4.17 (full-automation-coverage – § 11.4.17 is its strict expansion into per-type-of-test territory). It does NOT supersede them. CONST-11.4.17 (recursive submodule application) governs the cascade reach.

Classification: universal (per § 11.4.17). No escape hatch. A project shipping with mocks reachable from non-unit-test paths OR missing required test-type coverage is **not delivering a stable product** – release blocker for every consuming project, severity-equivalent to a § 11.4.17. PASS-bluff at the release-gate layer.

2026 05 15

§ 11.4.15 – Type-aware closure-status vocabulary (User mandate, 11.4.15)

Forensic anchor – direct user mandate (verbatim, 11.4.15):

"make sure we use proper wording for the workable item completion status in Issues, Issues_Summary and Fixed docs: – Fixed -> into the proper wording depending on the workable item type (task, feature, and so on) – for example: Fixed, Implemented, Completed. Add this important detail in our root Constitution, CLAUDE.MD and AGENTS.MD so it is ALWAYS respected and followed!"

Classification: § 11.4.15 –classified **universal** – naming discipline applies to every project that tracks work items by type. Bug closures aren't "implemented", feature deliveries aren't "fixed", and infrastructure refactors aren't either. Type-mismatched closure vocabulary is a quiet but persistent semantic-drift bluff

at the documentation layer – it makes greppable releases lie about what shipped.

11 4 33

The mandate. § . . . defined the lifecycle Status closed-set including the closure terminal value `Fixed (→ Fixed.md)`. § . . . defined the Type closed-set `{Bug | Feature | Task}`. § . . . binds the two: the closure terminal value MUST agree with the item Type, drawn from this -element closed map:

Item	**Type:**	Closure	**Status:**	value
Bug		Fixed (→ Fixed.md)		
Feature		Implemented (→ Fixed.md)		
Task		Completed (→ Fixed.md)		

The `(→ Fixed.md)` suffix is preserved across all three so the existing migration-discipline tooling (`Issues.md → Fixed.md` atomic move per § . . .) continues to work without per-Type branching. Generators (`generate_issues_summary.sh`, `generate_fixed_summary.sh`, status-counter helpers, the § . . . colorizer) MUST treat the three terminal values as semantically equivalent (all map to "closed, positive evidence captured") while preserving the literal in the emitted document.

No escape hatch. Closing a `Feature` with `Fixed (→ Fixed.md)` or a `Task` with `Implemented (→ Fixed.md)` is a § . . . violation. Pre-build gate (recommended, per consuming project) `CM-CLOSURE- VOCAB-TYPE-AWARE` walks every `Fixed.md` heading + every `Issues.md` heading whose `**Status:**` is one of the three terminal values, and asserts the Status value matches the item's `**Type:**` per the table. Paired mutation flips a Bug entry's status from `Fixed (→ Fixed.md)` to `Implemented (→ Fixed.md)` → gate FAILs.

Propagation. Composes with § 11.4.18 (item-status tracking), § 11.4.23 (item-type tracking), § 11.4.24 (Fixed-document column alignment), § 11.4.25 (colorisation – the closed-state palette applies regardless of which of the three terminal words is used).²⁴

§ 11.4.26 – Reopened-source attribution mandate (User mandate, § 11.4.27 – § 11.4.28)

Forensic anchor – direct user mandate (verbatim, § 11.4.29 – § 11.4.30):

"when we reopen some workable item (bug, task or feature) we MUST HAVE details on who did reopened this and why (- Reopened status needs: by who, AI or User + details in main Issues doc.). For example, reopened by AI or reopened by real User, why - test failed, manual testing detected problem, etc. Adapt our docs for this - Issues, Issues_Summary and Fixed."

Classification: § 11.4.31 –classified **universal** – every project that uses § 11.4.32 's Reopened lifecycle value benefits from knowing the reopen source + cause. Without this, reopen-thrash patterns (the same item bouncing between Fixed and Reopened across cycles) cannot be diagnosed, and the § 11.4.33 demotion-evidence rule loses its forensic counterparty (you cannot evaluate whether a reopen was operator-side observation or agent-side over-restoration without provenance).

The mandate. Every Issues.md (or equivalent project tracker) heading whose ****Status:**** value is Reopened MUST

carry, within non-blank lines of the heading, a `**Reopened-Details:**` line that captures four sub-facts:

- **By:** `AI` or `User` (the source-of-truth observer who flipped the status). `AI` covers in-loop reopens (test failure, gate regression, captured-evidence retrospect). `User` covers operator- side observations (manual testing, end-user report, design reconsideration).
- **On:** ISO date (`YYYY-MM-DD`).
- **Reason:** one-line cause classification – chosen from a closed vocabulary { `test-failed` | `manual-testing-detected` | `captured-evidence-contradicts` | `end-user-report` | `cycle-re-discovered` | `design-reconsidered` }. Other values are permitted with explicit `Reason: <free text>` annotation but the closed list **MUST** be tried first.
- **Evidence:** path to or short description of the captured artefact that justifies the reopen – log file, recording, gate failure ID, operator quote, etc. Reopens without evidence are \$. . / \$. . violations: the reopen IS a demotion-from-Fixed classification change, and demotion requires positive evidence captured under the conditions that re-exposed the defect.

The `Issues_Summary.md` (or equivalent) Status column **MUST** distinguish the four `Reopened` sub-states by source so a sweep query for "reopens by AI in the last days" is mechanically possible. Suggested column rendering:

`Reopened (AI: test-failed)` vs `Reopened (User: manual-testing)` .

No escape hatch. A `Reopened` entry without `**Reopened-Details:**` is a \$. . violation. Pre-build gate (recommended, per consuming project) `CM-ITEM-REOPENED-DETAILS` mirrors `CM-ITEM-OPERATOR-BLOCKED-DETAILS` (Phase .AT pattern): walks every actionable heading, scans non-blank lines for a `**Status:** Reopened` line, then continues scanning for `**Reopened-Details:**` . Missing

details line emits WARN initially, hardens to FAIL once backlog is fully populated.

11 4 7

Propagation. Composes with § . . (no-guessing – the Reason must be drawn from the closed vocabulary or explicitly annotated; no likely/probably/maybe causes), § . . (demotion-evidence – reopen IS a demotion from Fixed), § . . (item-status tracking – extends the Reopened value's discipline), § . . (Operator- blocked discipline – same audit-line pattern with the same gate shape).

2 0 2 6 0 5 1 5

2 0 2 6 0 5 1 5
§ . . – Canonical-root inheritance clarity (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"Parent AGENTS.MD or CLAUDE.MD are located under constitution directory (Submodule) containing these parent (root) which files we are inheriting inside the project! Pay attention to this and make sure you ALWAYS follow this rule! Same applies to the root Constitution file!"

"If needed (and we think it is needed it seems) add into the root Constitution, AGENTS.MD and CLAUDE.MD located in constitution directory (Submodule) which we are inheriting that the root (parent) Constitution, AGENTS.MD and CLAUDE.MD are not the ones in root of the project but inside the constitution directory (Submodule). We are inheriting it and if project specific rules have to be added or project specific constraints which are not universal and reusable, then

they go into the Constitution, CLAUDE.MD and AGENTS.MD of the project directory itself!"

Classification: § . . -classified **universal** – every project that consumes this constitution as a Git submodule MUST unambiguously distinguish the **canonical root** (the constitution submodule's three files) from the **consumer extensions** (the project's own three files).

The defect this anchor closes. Loose terminology – "parent CLAUDE.md", "root CLAUDE.md", "main constitution" – used without a file-path anchor causes a quiet but persistent confusion between the inheriting copy and the source-of-truth. AI agents have already been observed editing the consumer-side CLAUDE.md when the User's intent was to extend the canonical layer (or vice-versa), causing universal rules to leak into project-specific copies and project-specific rules to be falsely promoted as universal. The defect compounds: once a rule is misfiled, future propagation gates treat the misfile as canonical and silently spread it.

The mandate. The three files in **this constitution submodule** (constitution/Constitution.md , constitution/CLAUDE.md , constitution/AGENTS.md) are the **canonical root** – also called the **parent** files. They contain only universal rules per § . . – rules that any project consuming this submodule benefits from.

The three files in **the consuming project's repository root** (<project-root>/CLAUDE.md , <project-root>/AGENTS.md , optionally <project-root>/Constitution.md or equivalent) are the **consumer extensions**. They MUST start with the inheritance pointer per the existing constitution/CLAUDE.md "How inheritance works" section. They contain only project-specific rules per § . . – rules that reference particular hardware, vendor names,

regulatory regions, internal asset names, or project-private conventions.

Operative invariants that follow:

- . **When in doubt about which file to edit:** if the rule is reusable across any project, edit the constitution submodule's file. If the rule references project-private specifics (a hardware revision, a vendor SDK version, a regional constraint, a particular service), edit the consumer's file. Default to consumer-side when uncertain (per § 11.4.12 – narrower scope is cheap to widen later, the reverse is expensive).
- . **Terminology anchor.** When prose in any file in this constitutional family references "the parent CLAUDE.md" or "the root Constitution," the referent is the constitution-submodule file at `constitution/<filename>`, never the consumer's file. When it references "the project CLAUDE.md" or "this project's AGENTS.md," the referent is the consumer-side file at `<project-root>/<filename>`. AI agents reading this constitution **MUST** resolve ambiguous pronouns ("the CLAUDE.md", "the constitution") via this rule.
- . **Universal rules added to the consumer side are misfiled.** Per § 11.4.12 commit-time check, before adding a **MUST** to the consumer file, the author assesses whether it would benefit other projects that consume this constitution. If yes, lift it to the constitution submodule first.
- . **Project-specific rules added to the constitution submodule are misfiled.** Per § 11.4.12, before adding a **MUST** to the constitution submodule, the author assesses whether it references project-private specifics. If yes, demote it to the consumer file (or genericise it to a universal form first).

- **Propagation gates target both layers but their contents differ.** The existing `CM-COVENANT-114-N-PROPAGATION` gate family verifies that anchor TEXT for § . .N exists in EVERY `CLAUDE.md` and `AGENTS.md` across the project. The constitution-submodule files carry the **canonical** text; the consumer files carry **compact propagation pointers** that reference the canonical authority by file path. Both forms count as valid presence; the gate is satisfied when any form is found.
- **@constitution/CLAUDE.md import (Claude-Code-style) and the pointer-block fallback (Aider/Codex/Gemini-style) are equivalent.** The consumer's `CLAUDE.md` MUST start with one of:
 - The native import: `@constitution/CLAUDE.md`
 - The portable pointer-block per the `## INHERITED FROM constitution/CLAUDE.md` heading defined in `constitution/CLAUDE.md` "How inheritance works".
- **No silent demotion or silent promotion.** Moving a rule between layers MUST be a visible commit – `git mv` of a section if it's a clean clone, or an explicit "Lifted from to constitution per § . . " / "Demoted from constitution to per § . . " line in the commit message. AI agents MUST NOT silently re-author a § . .X anchor in the wrong layer and call it propagation.

Pre-build gate (recommended, per consuming project):

- **CM-CANONICAL-ROOT-CLARITY** – verifies (a) consumer's `CLAUDE.md` opens with the inheritance pointer (either `@import` or `## INHERITED FROM constitution/CLAUDE.md` heading), (b) the constitution submodule's three files are present at the expected path, (c) no `## INHERITED FROM` block in the constitution submodule's own files (those ARE the source-of-truth, not consumers).

Propagation. Composes with § . . (universal-vs-project classification – § . . defines the file-layer split that § . . classifies INTO), § . . (script documentation discipline – same canonical-vs-consumer file-layer rule applies to docs/scripts/). Reading order: this anchor first, then § . . for the classification rule.

No escape hatch. Misfiling a rule between layers IS a § . . violation regardless of intent. The fix is `git mv` + commit per invariant , not "leave both copies in place to be safe" – duplicates silently diverge over time.

2 0 2 6 0 5 1 5

§ . . – Mandatory `install_upstreams` on clone/add Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Every Submodule or Git repository we add or clone MUST BE upstreams installed using Upstreamable utility which MUST BE available through exported paths of the host system (in .bashrc or .zhrc) using `install_upstreams` command executed from the root of the cloned (added) repository – only if in it is Upstreams or upstreams directory present with bash script files (recipes) for all repository's upstreams!"

Operative rule. Every time an operator or agent adds or clones a Git repository (`git clone` , `git submodule add` , `incorporate-submodule` per § . . , etc.) under any

consuming project of this Constitution, the post-clone procedure MUST be:

- . `cd` into the newly-cloned (or newly-added submodule's) working tree.
- . If the working tree contains an `upstreams/` directory (or legacy `Upstreams/` per the `$ . .` transition) populated with one or more `*.sh` recipe files declaring upstream Git SSH URLs (the canonical declaration format defined by this constitution submodule's `Upstreams/*.sh` shape), the operator or agent MUST invoke `install_upstreams` from that directory.
- . `install_upstreams` is a host-system utility installed on the operator's `PATH` via `.bashrc` or `.zshrc` export. Implementation lives in the constitution submodule (`install_upstreams.sh`); operators alias / symlink it onto `PATH` once per host setup. The utility reads the recipe files, configures every declared upstream as a named git remote, and fans out `origin` push URLs across all declared upstreams.
- . If no `upstreams/` directory is present, no `install_upstreams` invocation is required.
- . If `upstreams/` is present but `install_upstreams` is not on `PATH`, the operator MUST install it (per the constitution submodule's setup docs) BEFORE making the clone usable.
- . Skipping step when an `upstreams/` directory IS present is a `$ . .` violation – the next push from that working tree will land on only one upstream, breaking `$ .` (Multi-upstream push is the norm).

Pre-commit attention. Before the first commit lands inside the newly-cloned working tree, the operator MUST verify that `install_upstreams` has been executed (when applicable). The quickest check: `git remote -v | grep -c push` reports the expected upstream count. If it reports `1` while `upstreams/`

`*.sh` declares more, abort the commit and run `install_upstreams` first.

Automation. The constitution submodule's tooling (the future `incorporate-submodule` per § . . . , the existing `scripts/init-submodules.sh` patterns in consuming projects) MUST auto-invoke `install_upstreams` as part of any clone / add operation when the target tree has an `upstreams/` directory. Operator-explicit manual invocation remains supported.

Gate `CM-INSTALL-UPSTREAMS-ON-CLONE` . Pre-merge gate inspects every owned-by-us submodule pointer commit and verifies that: (a) if `upstreams/` (or legacy `Upstreams/`) is present in the submodule's tree, (b) the submodule's local checkout has at least as many configured push URLs as the declared recipe count. Paired mutation (§ . . .): remove one upstream from local `origin --push` config → gate FAILs.

Composition. § . . . composes with § (single-entrypoint commit/push wrapper), § . . . (multi-upstream push is the norm – this rule is its setup-time complement), § (submodule changes propagate through submodule commits first – requires multi-upstream parity), § . . . / `CONST-` . . . (no force-push – the multi-upstream wrap-up makes unforced parity easy), § . . . (universal – every consuming project's clone procedure), § . . . (subagent delegation when batch-cloning), § . . . / `CONST-` . . . (owned-submodule discipline), § . . . / `CONST-` (`Upstreams/` → `upstreams/` transition is exactly this rule's scope), § . . . / `CONST-` (`.gitignore` ignores `upstreams/` only if the project explicitly decides not to track recipes, which is unusual – recipes ARE source of truth and SHOULD be tracked), § . . . / `CONST-` . . . (submodule-dependency- manifest works alongside `upstreams` recipes: `helix-deps.yaml` lists WHAT to

add at parent root, upstreams/ recipes list WHERE to push the resulting clones).

Classification: universal (per § . .). No escape hatch. Severity-equivalent to § . multi-upstream-push-violation at the clone-time setup layer – without the post-clone install, the working tree silently has a single-upstream blind spot.

2026 05 15

§ . . – Fetch-before-edit mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Make sure that feedback_fetch_before_edit memory rule is part of our constitution Submodule – the root Consitution, AGENTS.MD and CLAUDE.MD. Validate and verify that Proejct-Toolkit and all Submodules do inherit all of them! Follow the constitution Submodule documentation for details."

Background. In multi-agent / multi-upstream codebases – where parallel Claude Code, Cursor, Aider, Codex, Gemini CLI, or human operator sessions may operate on overlapping scope – the local working tree's state can lag behind the canonical upstream state by the time any given agent receives a task. Acting on stale local state produces three failure modes documented in the originating session (– –):

- **Redundant work** – the agent re-does what a parallel session already finished, wasting tokens and operator review time. (In the originating incident, "Gitee removal" had already been committed upstream by a

parallel agent minutes earlier; the second agent's `git remote remove gitee` was a no-op echo of work already done.)

- . **False confidence** – the agent reports completion of work that was already done by someone else, with no mechanical signal that their commit duplicates upstream history.
- . **Divergent history** – if the agent commits a parallel "completion" of already-done work, the result is two siblings of the same change, doubling the conflict surface for the next push attempt to multi-upstream remotes (`$ .`).

The mandate. The FIRST git-touching action of any session, on any consuming project that participates in this constitution, MUST be:

```
git fetch --all --prune
git log --oneline HEAD..@{u}           # parent
git submodule foreach --recursive 'git fetch --all --prune --quiet'
```

If `HEAD..@{u}` is non-empty, the agent MUST integrate (ff-merge, rebase, or – if non-fast-forward – surface to operator per `$ . .`) BEFORE issuing any local edit, scanner run, or test cycle. The fetch step is non-negotiable even when the operator's directive is phrased as "do X immediately" – the –second check prevents hour-long conflict reconciliation later.

Scope. Applies to:

- . The consuming project root.
- . Every owned submodule (per `$ . .`) – recursive.
- . The constitution submodule itself (`$ . .` step makes this explicit for constitution-side edits; `$ . .` generalises it to ANY edit on the consuming project, not only constitution edits).

. Any dependency cloned via `incorporate-submodule`
(`$ . . .`) or `git submodule add` (`$ . . .`).

Anti-bluff invariant. The fetch+log check MUST produce captured evidence – the actual `HEAD..@{u}` output, even if empty. Skipping the check on the basis of "I just fetched" or "nothing could have changed in the last N minutes" is a `$ . . .` (no-guessing) violation: the remote state is not knowable without a fetch.

Composition. Composes with `$ . . .` (multi-upstream push – without the fetch, the agent can't know which upstream has the canonical ref), `$ . . .` (test-interrupt-on-discovery – newly-fetched commits that contradict the planned work are exactly the "freshly discovered defect" that triggers cycle interruption), `$ . . .` (no-guessing – remote state requires fetch, not assumption), `$ . . .` (subagent delegation – parallel subagents MUST coordinate through fetched state, never assumed-local state), `$ . . .` (constitution update workflow – this rule generalises `$ . . .` step to ALL edits, not only constitution-file edits), `$ . . .` (post-constitution-pull validation – fetch-before-edit happens BEFORE the constitution-pull sweep `$ . . .` mandates).

Gates. Pre-build gate `CM-FETCH-BEFORE-EDIT-AUDIT` (when implemented in the consuming project) audits the most-recent commit range against the upstream HEAD at the commit's parent – if the parent ref was not the upstream HEAD at the time the commit was authored, FAIL. Paired mutation (`$ . . .`): synthetic commit whose parent is N commits behind the then-current upstream HEAD – gate must FAIL.

Classification: universal (per `$ . . .`). The rule has no project-specific assumptions – it applies to any multi-upstream / multi-agent codebase. No escape hatch.

Severity-equivalent to a § . PASS-bluff at the planning layer – acting on stale local state is the operational analog of asserting truth from unverified premises.

§ . . – Installable-Asset Evidence Mandate (User mandate, – –)

Forensic anchor – verbatim operator report (– –):

"app does not have a launcher icon anymore – the latest published app tester build(s). how come app passed anti-bluff checks?"

For any user-distributable build artifact (package, bundle, installer, or container image produced by the build pipeline and distributed to end users), tests and challenges MUST open the artifact and verify each user-visible asset is **present** and **non-degenerate**.

"User-visible asset" includes (non-exhaustive):

- Application icons for every density / resolution tier declared in the artifact metadata, including any platform-specific icon format (multi-resolution container, adaptive XML, symbol set, etc.) that the target OS uses when the minimum supported OS version requires it.
- Splash screens declared in the install metadata.
- Application name strings as declared in the install metadata.
- Any other asset whose absence causes the user to be unable to identify, launch, or interact with the

installed application from the OS launcher / home screen / app drawer.

A PASS without opening the artifact and verifying the asset chain end-to-end is a \$. PASS-bluff, regardless of whether source files exist and tests otherwise pass. The specific failure mode this rule targets is: source file exists → build pipeline packages it → post-build checks pass at the source layer → artifact ACTUALLY produced with the asset stripped or misconfigured, and no gate ever opens the artifact to verify.

Required evidence: the anti-bluff challenge for a user-distributable artifact MUST produce per-asset PASS/FAIL lines showing: (a) the asset entry exists in the artifact package listing, (b) the asset is non-empty / non-degenerate (size ≥ platform-defined minimum OR format-validated), (c) where the OS's asset-resolution path involves indirection (alias, XML reference chain, density-qualifier override), the full chain is traced to the final rendered resource.

Consuming-project implementation: each consuming project ships one challenge script per artifact type that opens the produced artifact and verifies every declared user-visible asset. The challenge MUST run as part of the project's standard QA gate (equivalent of `make qa-all`).

Root cause of \$. . introduction: a consuming project shipped multiple releases with the application icon absent on the target OS because: () the icon XML used a resource type that the OS resolves differently above a specific API level, () all pre-ship challenges verified source files only, none opened the packaged artifact. This is a pure \$. PASS-bluff – every test and challenge reported green while the end user saw no icon.

Classification: universal (per § 1.4.5). No escape hatch. Severity- equivalent to a § 1.4.5 PASS-bluff at the artifact-packaging layer. Composes with § 1.4.5 - § 1.4.5 (evidence requirements), § 1.4.5 (full automation coverage), § 1.4.5 (no fakes beyond unit tests – source-layer checks of a distributable artifact are the distributable- layer analog of unit-test-only coverage). See Constitution § 1.4.5 for the full mandate.

§ 1.4.5 – Per-Feature On-Device End-User Validation Mandate (iter- , - -)

Every user-facing feature in a consuming project MUST have at least one HelixQA scenario that exercises the feature from cold-launch with **positive runtime evidence**:

- A screenshot or screen recording captured at a named checkpoint.
- At least one assertion of type `screenshot` , `accessibility_count` , `accessibility_node` , `ocr_text` , `pixel_histogram` , or `editor_state` .
- Evidence files written to a discoverable output directory so every PASS is cross-verifiable after the fact.

Scenarios MUST be RE-EXECUTED on every release candidate to catch cross-iteration regression. A scenario authored but never re-run on subsequent iterations degrades to a metadata-only gate – a § 1.4.5 PASS-bluff at the regression layer.

Authoring rules (consuming project sets project-specific paths):

- . Scenarios are stored in a canonical bank directory designated by the consuming project (e.g. `<banks-root>/feature-coverage/`).
- . Each scenario YAML MUST include: `name`, `version`, `metadata`, `platforms`, `test_cases`, and at least one step with `evidence_required: true` and a specific `evidence_type`.
- . A **coverage matrix** document MUST exist listing feature × iteration × scenario. Matrix rules: (a) new iteration MUST add ≥ scenario per new user-facing feature; (b) prior scenarios MUST NOT be deleted (mark as `status: retired` if feature is removed); (c) all non-retired scenarios MUST PASS for ship-ready.
- . A **static gate** (challenge script) MUST assert the scenario count ≥ N (where N is the number of user-facing features registered in the coverage matrix) and that each YAML contains a `evidence_type: assertion`.
- . **Portable host tooling:** evidence-validation scripts MUST use POSIX-portable file-size helpers (e.g. `wc -c < file`) instead of OS-specific commands (e.g. GNU `stat -c%s`) to avoid silent -byte reports on BSD/macOS hosts. A regression challenge MUST verify the portable helper on the actual host.

iOS / platform-deferred pattern: scenarios written platform-agnostically with the target platform listed in a comment deferral (e.g. `# iOS: deferred – tracker #X`) are valid; add the platform to the `platforms:` list when the automation toolchain becomes available – no structural change to the scenario is needed.

Classification: universal (\$. .). Composes with \$. , \$. (anti-bluff), \$. . (full automation coverage), \$. . (no fakes),

§ . . (artifact evidence). No escape hatch – a feature without an on-device scenario is NOT covered per
 § . . invariant . See Constitution
 § . . for the full mandate.

§ . . – Full-suite retest before
 release tag mandate (User mandate,
 – –)

Forensic anchor – verbatim user mandate

(– –):

"Please, do complete retests with all existing tests (multi-hours efforts) when all new workable items are done, fixed, polished and verified. Time is essential! We should already have this in our (root) Constitution, CLAUDE.MD and AGENTS.MD."

Operative rule. A release tag (any `vX.Y.Z` / `X.Y.Z-suffix` identifier MUST NOT be created until a **COMPLETE retest with ALL existing tests** has been executed on a clean baseline AFTER every workable item in the batch is done, fixed, polished, and individually verified. A "spot-check" retest that runs only the tests directly touched by the batch is FORBIDDEN – it misses interaction defects between the batch's fixes and previously- stable code paths.

The complete retest comprises:

- . **Pre-build verification full sweep** – every gate in the project's pre-build script runs; FAIL required.
- . **Post-build verification full sweep** – every gate in the project's post-build script runs against the freshly-assembled image; FAIL required.
- . **On-device –phase cycle** (e.g. `test_all_fixes.sh` or project equivalent) on **EVERY owned device** (the full

topology at that point in time). Each phase (Immediate Fresh Flash / After Reboot / After Factory Reset / After Final Reboot) MUST complete; no phase may be skipped. Phase summaries captured.

- **Meta-test full mutation sweep** – every paired mutation in `meta_test_false_positive_proof.sh` (or project equivalent) executed; every gate proven non-bluff.
- **Test bank full sweep** – if a Challenge-driven test bank exists, every Challenge in the bank covered by a full QA session (not a sub-set).
- **Issues.md / Fixed.md state audit** – no item silently demoted; every Reopened item has `$REOPENED-Details`; every closure has captured-evidence per `$STATUS`; no `Status:` value outside the `$CLOSED-SET`.
- **CONTINUATION.md sync check** – per `$CONTINUATION.md`, document reflects current state at the moment of tagging.

Time is essential. A complete retest is typically a `$RETEST-DURATION` hour elapsed effort depending on parallelism, device count, and meta-test mutation count. This is NOT optional and NOT abbreviated. Operators should plan release cadence with this duration in mind, NOT skip the retest to ship faster. Skipping the retest is the exact "tests passed but feature broken" failure mode `$RETEST-FAILURE-MODE` specifically prohibits.

Composition with `$RETEST-FIXES`. Per-fix retest (`test_all_fixes.sh` run after each individual fix lands) is STILL mandatory per `$RETEST-FIXES`. `$RETEST-FIXES` is the additional final integrity check that runs after the BATCH is complete – catching interaction defects that individual per-fix retests cannot see in isolation.

Composition with `$RETEST-BATCH`. The complete retest is the authoritative captured-evidence baseline for any closure of items present in the batch. If a `$RETEST-BATCH` -promoted

item PASSED its per-fix retest but FAILS the full-suite retest, the closure MUST be reverted and the item moved back to In progress – captured-evidence-contradicts under same-conditions per § 11.4.3.

Composition with § 11.4.3. Per-feature on-device end-user validation runs as part of step 11.4.3 (on-device cycle) – every feature's wrapper-test fires during the full-suite retest. The two mandates are complementary: § 11.4.3 covers individual feature validation breadth; § 11.4.3 covers release-time integration depth.

Gate CM-FULL-SUITE-RETEST-MANDATE. Pre-tag gate inspects the release-candidate state and verifies that within the last 24 hours of the candidate commit, evidence exists of (a) full pre-build + post-build run, (b) on-device 11.4.3 phase cycle on every owned device, (c) meta-test full sweep, (d) Issues.md/Fixed.md audit pass. Evidence captured in docs/changelogs/<tag>.md 11.4.3 per § 11.4.3 (c) requirements. Paired mutation (§ 11.4.3): strip Full-suite retest evidence block from a changelog → gate FAILs.

Classification: universal (per § 11.4.3) – every consuming project's release procedure. No escape hatch – there is no --skip-full-retest flag, no --quick-release mode. Operators who feel time-pressured to skip the full retest should instead delay the release until the retest completes; shipping unverified code is worse than delayed shipping. 2026 05 17

§ 11.4.3 – Pre-Force-Push Merge-First Mandate (User mandate, 11.4.3 – 11.4.3)

Forensic anchor – verbatim user mandate (11.4.3 – 11.4.3):

"make sure we bring everything from branches to our side before force push is done! After everything is safely and fully merged and all potential conflicts (if any) resolved, then do force push! make sure nothing is lost, broken or corrupted on both sides! add these rules in our root Constitution, CLAUDE.MD, AGENTS.MD (constitution Submodule) if it is not added already! Extremely important rules and mandatory constraints we MUST HAVE and fully respect!"

Operative rule. Any force-push (`git push --force` , `git push --force-with-lease` , `git push +<ref>` , equivalent history-rewriting operation on any remote) authorised under § . / CONST- MUST be preceded by a mechanical -step merge- first pipeline that brings every remote-side commit into the local tree, resolves every conflict carefully, and verifies nothing is lost or corrupted on EITHER side BEFORE the overwriting push is executed.

The -step pipeline (mandatory, in order):

- . **Fetch every remote-side reference.** Run `git fetch --all --prune --tags` against every configured remote (origin + every upstream). Capture the output.
- . **Integrate every divergent commit locally.** Compute `git log --oneline HEAD..<remote>/<branch>` for every remote. For every non-empty range, integrate the remote commits into the local tree via the appropriate strategy:
 - `git rebase <remote>/<branch>` when the local divergent work is a strict superset that should land on top, OR
 - `git merge <remote>/<branch>` when the local + remote histories are independent additions that both deserve preservation, OR

- **operator-confirmed cherry-pick** when remote contains a subset of commits already present locally in different form.
- . **Audit the integrated tree.** Verify (a) no conflict markers anywhere in the working tree (`grep -rn '^<<<<<< \|^===== $\|^>>>>>>' .` returns empty across all governance + source + test files – third-party docs containing the markers as illustrative content excepted), (b) no file silently dropped (`git diff --stat HEAD@{1} HEAD` shows only expected additions), (c) every previously-passing test still passes on the integrated tree (per \$. . + \$. . baseline), (d) every captured- evidence artifact from the pre-merge state still validates.
- . **Execute force-push.** Only after steps – produce captured evidence of clean integration: `git push ..force-with-lease <remote> <ref>` (NEVER `--force` without `--with-lease` unless explicitly authorised per \$. sub-clause for a specific remote where lease semantics are unavailable). One force-push event per CONST-authorisation – no batch authorisation across multiple force-pushes.

Three failure modes prevented:

(a) **Remote-side content loss.** Force-push without prior fetch silently overwrites commits a parallel session, sibling agent, or co-developer landed on the remote. The lost work is recoverable from the remote's reflog only within its TTL window – typically 7–14 days – but the audit trail (PR comments, CI evidence, governance signatures) is lost permanently.

(b) **Stale-state act on remote.** Force-push from a divergent local state that was branched from a remote tip N hours ago treats the remote-side N hours of work as nonexistent. The operator's `--force-with-lease` safety check catches the

literal SHA mismatch but ONLY when invoked AFTER a fetch that updated the lease reference; without step 214, the lease check is reading stale local refs.

(c) **Conflict-driven corruption.** A merge that completes with unresolved markers in the tree (one half kept, other dropped silently because the marker was inside a comment block, OR markers committed verbatim as in CONST- prerequisite incidents observed - on helix_qa + containers) lands broken governance / source on the force-push target. Step 243's explicit conflict-marker grep is the mechanical guard.

Two-gate composition with CONST- 243 . § 11 4 41 . does NOT relax CONST- 's operator-approval requirement — it adds a SECOND gate on top:

- **Gate A (CONST- 243).** Operator types explicit per-operation force-push authorisation in the conversation.
- **Gate B (§ 11 4 41 .).** Agent executes the 043 -step merge-first pipeline, captures evidence of clean integration, presents evidence to operator BEFORE the 11 4 41 force-push.

Both gates required. CONST- alone authorises a force-push that loses remote work; § 11 4 41 . alone risks force-pushing without operator awareness. The two together produce a force-push that is BOTH operator- authorised AND remote-safe.

Verification artefact. Every § 11 4 41 . -governed force-push emits a docs/changelogs/<tag>.md "Force-push merge-first audit" section containing: (i) git fetch output, (ii) per- remote HEAD.<remote>/<branch> log before integration, (iii) integration strategy chosen per remote with rationale, (iv) post-integration conflict-marker scan output (must be empty), (v) post-integration test suite delta (must show only expected changes), (vi) the --force-

with-lease push output with lease SHA evidence, (vii)
CONST- authorisation quote from the conversation.

Cascade requirement (per CONST-). This anchor
(verbatim or by §11.4.41 / CONST-061 ID reference) MUST
appear in every owned submodule's CONSTITUTION.md ,
CLAUDE.md , and AGENTS.md . Severity-equivalent to a
§ . PASS-bluff at the remote-data-integrity layer.
Gate CM-FORCE-PUSH-MERGE-FIRST walks docs/changelogs/<tag>.md
"Force-push" entries for the audit elements; paired
mutation strips any element and asserts gate FAILs.

Classification: universal (per § . .) – every
owned project executes force-pushes against its remotes;
the merge- first discipline is reusable verbatim. No
escape hatch – the operator-pressure escape ("just force-
push, we'll fix it later") is the exact failure mode this
anchor closes.

Composes with: § . (data-safety hardlinked backup),
§ . . (test-interrupt-on-discovery – broken
integration triggers rollback), § . . (no-guessing
– every step's outcome captured, not assumed),
§ . . (constitution-submodule update pipeline –
that mandate is the per-submodule specialisation of
§ . . . for governance-files-only commits),
§ . . (post-pull validation – the audit step's
mechanical companion), § . . (fetch-before- edit
– step enforces it for the force-push specifically),
§ . . (full-suite retest – step 's test-
evidence prerequisite), CONST- (per-operation
operator approval), CONST- (recursive cascade).

§ . . – Iteration-discipline mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Are we clear about working iterations and sorting priorities for testing and fixing? We first fix top and middle critical priorities → We flash and test all of these → If nothing is broken and no new issues are reported by users or found we run full system testing (many hours for everything) → if anything is still broken or new issues are reported we get back to fixing → Cycle repeats. Make sure this is absolutely clear and mandatory foundation of our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD! We MUST WORK in this manner until otherwise has been told."

Operative rule. Project work proceeds in priority-ordered iteration cycles. Each cycle is a closed loop of five steps that MUST run in order; advancing past any step before its acceptance condition is met is a § . bluff equivalent to skipping the step entirely.

The five-step cycle:

- . **Priority-ordered fix selection.** Open the project's issue tracker, filter to items whose `**Status:**` is one of `Queued` | `In progress` | `Reopened` | `Operator-blocked` (per § . . closed-set), sort by severity DESC then intra-group criticality DESC (per § . . sort order), select ONLY items classified `Top critical` or `Middle critical` for the current batch. `Low` items are explicitly DEFERRED until the critical batch is closed.
- . **Batch implementation.** Land source-side fixes for the selected batch per § . . (batch-source-fixes-

before-rebuild) – every non-runtime-dependent fix lands BEFORE the next rebuild, every fix gets the \$. . four-layer coverage (pre-build gate + post-build gate + on-device test + paired mutation).

- **Smoke-test gate.** After flash, run the SMOKE bundle: (a) anti-bluff baseline (`meta_test_false_positive_proof.sh` for gates touched by this batch), (b) every `test_*.sh` whose name matches a fix in this batch, (c) the critical-path regression probe (boot + launcher + WiFi + audio + secondary-display routing core), (d) the \$. . per-feature on-device validation for any feature touched. Estimated runtime: < minutes on the parallel device fleet. Smoke MUST emit zero FAIL and zero unclassified WARN.
- **Full-system-test gate.** ONLY if (a) smoke is GREEN AND (b) no new operator/user report has surfaced during the batch's validation window, run the \$. . complete retest (all steps including - h on-device -phase cycle on every owned device). If smoke FAILED OR a new report arrived, the full-system test is FORBIDDEN; loop back to step with the new evidence added to the priority queue.
- **Release-readiness or loop.** If full-system test is GREEN AND no new report arrived during it, the batch is release-ready (the \$. . tagging procedure may proceed pending operator authorization). Otherwise loop back to step .

Priority taxonomy (binds existing severity convention).

Top critical = severity C (Critical per Issues.md convention) AND intra-group criticality 5 .

Middle critical = severity C with intra-group 1 - 4 OR severity M (Major) with any intra-group score. Low = severity L per the existing [C/M/L] taxonomy documented in \$. . . Severity WARN items are treated as Low for iteration-discipline purposes.

Cycle repeats. The five-step loop continues, fed by the priority queue, until the operator explicitly authorizes a release. There is no upper bound on iteration count; "we've cycled enough" is not a Constitution-recognised stopping condition. Time is essential per § 11.4.4. BUT the cycle MUST NOT be truncated to ship faster – that is exactly the bluff § 11.4.9 forbids.

Composition. § 11.4.4 (per-fix retest inside step 1) + § 11.4.5 (closures require same-conditions positive evidence; step 1 is authoritative baseline) + § 11.4.6 (source-side batching inside step 1) + § 11.4.7 (Reopened items attribute By: AI / User) + § 11.4.8 (the multi-hour full-suite retest IS step 1 – § 11.4.9 is the meta-loop conductor binding the instruments).

Anti-bluff coupling. Per § 11.4.4 + § 11.4.5, every smoke-test and full-system-test PASS MUST carry positive captured evidence of user-visible behaviour. Tests AND HelixQA Challenges bound equally – a Challenge that scores PASS without applicable analysis is a § 11.4.4 PASS-bluff.

No escape hatch. No `--skip-priority-batch`, `--skip-smoke`, `--full-suite-only`, or `--release-without-loop` flag exists. Subagents performing autonomous work default to the § 11.4.4 path. Operators who feel time-pressured to short-circuit the loop should add capacity (more devices in parallel, more agent slots) rather than skip steps. Shipping under-validated code is worse than delayed shipping.

Gate `CM-COVENANT-114-42-PROPAGATION` mirrors

`CM-COVENANT-114-40-PROPAGATION` : – every `CLAUDE.md` / `AGENTS.md` in the covenant file set carries the § 11.4.4 anchor. Paired mutation strips the anchor literal from one consumer file → gate FAILs.

Classification: universal (per § 1.1.4.3) – every consuming project's iteration workflow.

§ 1.1.4.3 – TDD-Fix-Discipline Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Make sure you do validate and verify every fix! Make sure you first do the fix on live device using ADB if that is possible, once fix is confirmed, make sure that it is fully applied into our codebase ... It is important to note that we shall start by creating test which confirms the issue (fails as expected), and ones the fix is produced it becomes positive – it does pass with success! No bluff policy is mandatory and nothing can be bluffed! Every step of this process MUST BE % bluff-free and bluff-proofed!"

Operative rule. Every fix MUST follow the –step TDD-fix workflow:

- **RED – Failing test FIRST.** Before any code change, author a test (or extend an existing one) that exercises the user-visible path of the reported defect and FAILs for that defect specifically. Per § 1.1.4.3 the FAIL MUST be a real product defect (not a script-bug induced FAIL). Per § 1.1.4.3 the FAIL run MUST produce captured evidence (logcat / dumsys / screenrecord / dmesg / /sys snapshot / sink-side probe). The test identifier MUST cite the §-letter or Fix- under repair. A RED step that "happens to pass because the underlying probe lifecycle isn't wired" is itself a § 1.1.4.3 PASS-bluff.

- **LIVE-ADB-PROBE – try the fix on the running device first (when feasible).** For mutable surfaces –
`setprop persist.*` , `settings put` , `pm clear` , `am force-stop` , `am start` , `cmd wifi` , `cmd media_session` , `ad-hoc /sys` writes, `boot-script` edits pushed to `/vendor/bin/` ,
test-fixture replacement – the operator MUST attempt the fix on the running device via `adb shell` first. The fast adb-loop (~ min per attempt) replaces the build-flash-test loop (~ h per attempt) for confirmation of the hypothesis. The live-probe IS NOT the fix; it is the EXPERIMENT that proves the hypothesis is correct before the source-side investment.

Live-probe is INFEASIBLE (exception, must rebuild) for:
kernel changes, AOSP framework (`frameworks/base/`),
hardware HAL changes, `vendor/` partition contents,
`init.rc` service definitions, `sepolicy`, `ro.` build
properties (immutable post-boot), AOSP-build-time XML /
resource overlays, `Android.bp/Android.mk` changes,
signed-system-app contents (`LOCAL_CERTIFICATE := platform`). Each exception MUST be cited in the commit
message as `LIVE_PROBE_INFEASIBLE: <reason>` .

- **GREEN – apply the fix to source code.** The source patch MUST achieve the same effect on the device as the live probe. Per § . . the source change is batched with other source-side fixes where appropriate. Build via the project's containerized build (`$ .`). Flash via the project's flash script. The bytes on the assembled image MUST land where the post-build gate expects (per § . . (b) four-layer requirement).
- **VERIFY – re-run the RED test, must now PASS.** Per § . . demotion-evidence: the PASS MUST come from the SAME conditions that initially FAILED (same device, same firmware NOW carrying the fix, same load profile, same cycle position). Per § . . the PASS MUST carry positive captured evidence (presence + correctness

– RMS amplitude, fprobe channel count, frame count, OCR text, sink codec state, etc.). Per § 11.4.42 video tests MUST cross-check via the recording-analyzer. Per § 11.4.43 a reliability check at the iteration count (typically 10 iterations) MUST PASS all iterations – a single intermittent FAIL keeps the item in In progress per § 11.4.39.

. DOCUMENT – update every relevant doc in the SAME commit.

Per § 11.4.40 (c): docs/Issues.md → docs/Fixed.md migration with type-aware closure vocabulary per § 11.4.41 (Bug → Fixed, Feature → Implemented, Task → Completed); the project's CLAUDE.md Applied Fixes Reference row; per-version docs/changelogs/<tag>.md entry; affected user-facing guides under docs/guides/; HelixQA bank entry per § 11.4.43 (b); docs/CONTINUATION.md per § 11.4.42; project memory file if non-obvious.

Composition. § 11.4.44 is the workflow that ENACTS the existing covenant. Composes with: § 11.4.45 (the RED FAIL is a real defect), § 11.4.46 (captured evidence at both RED and VERIFY), § 11.4.47 (topology dispatch in the RED test), § 11.4.48 (four-layer coverage at GREEN), § 11.4.49 (quality analysis at VERIFY), § 11.4.50 (no guessing during the LIVE-PROBE hypothesis statement – UNCONFIRMED: until probe proves it), § 11.4.51 (demotion-evidence at VERIFY – same conditions), § 11.4.52 (research citation precedes step 1), § 11.4.53 (GREEN is batched), § 11.4.54 (full-suite retest is the release-tag-time additional check), § 11.4.55 (-iteration reliability loop IS step 1's inner cycle).

Gate CM-COVENANT-114-43-PROPAGATION. Pre-build gate verifies the § 11.4.56 anchor is present in every CLAUDE.md / AGENTS.md across the covenant file set. Paired mutation

strips the anchor literal from one consumer file → gate FAILs.

Classification: universal (per. § . . .) – every consuming project's fix workflow. **No escape hatch.** No `--skip-red-test`, `--no-live-probe`, `--skip-verify`, `--skip-document` flag exists. "I'll add the test after the fix" is the exact PASS-bluff pattern § . . . forbids because the test authored after the fix demonstrates only that the test agrees with the fix, not that the test catches the bug.

§ . . . – Document Revision Header
Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"Add this change / improvement as background work in parallel with our mainstream work. All documents (Issues, Issues_Summary, Fixed, Continuation doc and all others MUST HAVE revision number! Revision number with information such as: last date and time when document was modified MUST GO at the top of the document below the main H title. Add all relevant information about the document as well. However, last date and time document has been modified and revision number are MANDATORY! Revision number shall be the whole number , , , etc. Document all this in our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD as mandatory rule(s) and constraints we MUST HAVE! No bluff is allowed of any kind!"

The mandate. Every IN-scope tracked Markdown document MUST carry a header block directly below the H1 title containing two MANDATORY fields:

- ****Revision:**** N where N is a monotonic positive integer (1, 2, 3, ...). Never decremented. Never reset on rewrite. A freshly-created IN-scope document starts at Revision 1.
- ****Last modified:**** YYYY-MM-DDTHH:MM:SSZ (ISO 8601 UTC).

Optional encouraged fields: ****Description:****, ****Authority:****, ****Maintainer:****, ****Scope:****, ****Auto-generated-from:**** (auto-generated docs only), ****Status:**** (plan docs only).

Header format (canonical):

```
# Document Title

**Revision:** 47
**Last modified:** 2026-05-18T13:42:00Z
**Description:** Open / in-flight bug + work tracker
**Authority:** Constitution §11.4 covenant
**Maintainer:** Operator + AI loop per §11.4.42

<content...>
```

YAML front-matter and HTML-comment headers are FORBIDDEN – they are invisible to humans scrolling the raw Markdown and break the "all relevant information about the document" visibility clause.

IN scope (revision header MANDATORY): docs/Issues.md , docs/Issues_Summary.md , docs/Fixed.md , docs/Fixed_Summary.md , docs/CONTINUATION.md , docs/guides/**/* .md , docs/research/**/* .md , docs/scripts/**/* .md , docs/changelogs/**/* .md , docs/superpowers/plans/**/* .md , docs/hardware/**/* .md , and every other tracked Markdown under docs/ .

OUT scope (revision header NOT required): `CLAUDE.md` / `AGENTS.md` (every layer – already version-tracked via `VERSION` file + inheritance pointer per § 11.4.23; adding a per-file revision would duplicate authority and obscure the canonical-root contract); `README.md` / `CONTRIBUTING.md` / `LICENSE` / `NOTICE` / `VERSION` / `OWNERS` (standard metadata); rendered HTML/PDF artifacts (revision inherited from source Markdown); auto-generated docs that already carry an embedded generation-time stamp.

Auto-bump tooling:

- `scripts/doc_revision_bump.sh <file>` – manual entry point, idempotent (no-op if Last-modified within 12.10.11.4.44 s of `date -u`).
- `.git/hooks/pre-commit` – automatic entry point; walks staged IN-scope docs, calls the bump for each, re-stages.
- `scripts/testing/sync_issues_docs.sh` – composes the bump for `Issues_Summary` + `Fixed_Summary` after regeneration.
- `scripts/commit_docs.sh` – calls the bump before staging so operators who skipped the hook installation are still covered.

Composition with § 11.4.23. `CONTINUATION.md` already carries `Last updated:` per § 11.4.23. `§ 11.4.23` adds `**Revision:**` N directly under H + reuses the existing `Last updated:` line as the `§ 11.4.23` `Last modified:` line. One source of truth, two pointers – no duplication.

Composition with § 11.4.23. `Issues_Summary.md` / `Fixed_Summary.md` inherit the revision number of their source-of-truth Markdown at generation time; they are NOT incremented independently.

Composition with § 11.4.23. HTML colorizer preserves the revision header in the colored HTML – pandoc renders the bold lines as styled `<p><strong>...</strong></p>` blocks

which the colorizer leaves alone (only Issues-table rows are mutated).

Gates: CM-DOC-REVISION-HEADER-PRESENT walks IN-scope docs and asserts both mandatory fields appear within lines below the H title; CM-COVENANT-114-44-PROPAGATION asserts the anchor literal is present in every covenant file. Paired mutations strip the **Revision:** line OR the anchor literal → gates FAIL.

Classification: universal (per § . .). **No escape hatch** — no --skip-revision-bump , --no-header , --allow-missing-revision flag exists. Operators who feel pressured to skip the bump should land smaller commits, not bypass the discipline.

§ . . — Integration-Status-Doc
Maintenance Mandate (User mandate,
- -)

Forensic anchor — verbatim user mandate
(- -):

"Make sure this Markdown document is regularly updated to reflect current up to date state of full integration ... Whole integrations MUST be regularly updated and retested! Keep the Markdown document ALWAYS up to date and ALWAYS exported into PDF and HTML! Make sure document is ALWAYS in sync! Add all this into our Constitution, AGENTS.MD and CLAUDE.MD."

Generalisation note. § . . is the generic form of § . (CONTINUATION.md) applied to every domain integration. § . binds the canonical handoff document specifically; § . . generalises the sync-and-evidence pattern to every non-trivial integration

(Dolby/Arvus, NanoKVM, Firebase, Widevine L /L , Play Protect, Netflix, Cuttlefish, Chromecast, WiFi chip, ES codec, Sonos, AC- , Google Stack, HiFi audio, TV apps, etc.). Without this generalisation each new integration re-invents the same sync wrapper, revision-header discipline, captured-evidence requirement, and operator-blocked surface – duplicating effort and risking drift.

Operative rule (sub-points – every Status.md MUST hold ALL).

Every integration-status document at `docs/<domain>/<integration>/Status.md` :

- . MUST exist when a domain integration is non-trivial (more than one fix / test / gate landing for that integration).
- . MUST carry the \$. . revision header directly below the H title (Revision + Last modified + Description + Authority + Maintainer + Scope).
- . MUST be auto-synced (HTML + PDF) on every related-test-cycle completion AND every fix touching the integration. "Auto-synced" means the wrapper from sub-point is invoked; manual regeneration is forbidden as the sole path because operators will forget.
- . MUST be auto-colored per \$. . – Status column cells colored by lifecycle state, Type cells colored by Type, operator-blocked rows visually distinct.
- . MUST have a sync wrapper invocable as `bash scripts/testing/sync_integration_status.sh <Status.md>` OR a per-integration thin shell that delegates to the generic wrapper. The thin shell exists so operators can type a short command per integration without remembering paths.

- MUST live under `docs/<domain>/<integration>/` directory structure – consistent layout, discovery by glob remains $O(1)$.
- MUST include a captured-evidence-driven status table per § 4.1.1 – every status claim cites the test log path, recording file path, or sink-probe report that backs it. Status claims without evidence are § 4.1.1.1 PASS-bluffs.
- MUST distinguish status values per § 4.1.1.1 / § 4.1.1.2 : PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED – closed vocabulary, mechanically distinguishable, no implicit values.
- MUST list operator-blocked items at the top of the document (just below the revision header) so an operator scanning the Status.md finds action items in $O(1)$ – not buried ~ 10 lines deep in a chronological evidence log.
- MUST be referenced from `docs/CONTINUATION.md` § 4.1.1.1 (Active work) when any item in the Status.md is non-terminal (Queued / In progress / Reopened / Operator-blocked) – composes with § 4.1.1.1 so the canonical handoff document points at the integration-status doc rather than re-stating its contents.

Gates:

- `CM-COVENANT-114-45-PROPAGATION` – anchor text propagated to every `CLAUDE.md` / `AGENTS.md` across the covenant file set.
- `CM-AF-INTEGRATION-STATUS-DOCS` – discovers all `docs/**/Status.md` via glob, verifies each carries the § 4.1.1 header, has a sync wrapper, HTML+PDF exports are mtime-current vs the source, colorization applied (cell-status class present in the colored HTML).

Paired meta-test mutations (§ 11.4.10 compliance):
 propagation strip → propagation gate FAILs; delete
`**Revision:**` from a Status.md → status-docs gate FAILs;
 touch Status.md without re-syncing → mtime mismatch FAILs;
 inject `likely` outside `UNCONFIRMED:` block → existing
`CM-NO-GUESSING-MANDATE` catches it (composition catch).

Composition. § 11.4.11 (captured evidence in every row), § 11.4.11.1 / § 11.4.11.2 (status vocabulary closed-set), § 11.4.11.3 (sync wrapper pattern reused for HTML+PDF re-export), § 11.4.11.4 (sink-side captured evidence is a specific instance of the generic rule), § 11.4.11.5 (status values), § 11.4.11.6 (commit_docs.sh wrapper), § 11.4.11.7 (colorizer extends to Status.md HTML), § 11.4.11.8 (every Status.md carries the revision header), § 11.4.11.9 (CONTINUATION.md references Status.md paths in § 11.4.11.10).

No escape hatch — no `--skip-status-sync`, `--no-revision-bump-on-status`, `--allow-stale-html` flag. Operators who feel pressured to skip the sync should land smaller, more frequent integration commits rather than batching status updates into one giant push.

Classification: universal (per § 11.4.11.11).

§ 11.4.11.12 — Validate-recent-work-before-post-flash-tests mandate (User mandate, § 11.4.11.13)

Forensic anchor — verbatim user mandate
 (§ 11.4.11.14):

"We see that on newly flashed devices we execute all post flash tests. Make sure we do execute them all after we validate and verify all recent work from the Issues and Continuation docs first! Once all new /

latest work is fully validated and verified, no new issues found, nothing is discovered broken or faulty, then we run all post-flash tests! If validation and verification of recent work fails for any reason, and we must go back to fix anything that popped up we do not have to run post-flash tests! We MUST save time! All post-flash tests can be executed only after recent work on features and issues is fully validated and confirmed with complete anti-bluff policy enforced! Add this so it is absolutely clear into our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD! Start following all these rules IMMEDIATELY! Iff required abort all post-flash testing if it is in progress!"

Operative rule. After every device flash, the orchestrator MUST first run a recent-work validation pass before any full post-flash suite (`test_all_fixes.sh` or project equivalent). The validation pass enumerates recent work from three docs (Issues.md / Fixed.md / CONTINUATION.md §), maps each recent item to its on-device test binding, runs ONLY those tests, classifies each result per § . . . , and demands % green before the full suite is permitted.

Definition of "recent work" (closed-set):

- docs/Issues.md actionable headings with `**Status:**` ∈ {In progress, Ready for testing, Reopened} (per § . . . closed-set).
- docs/Fixed.md headings closed within `--max-age-days N` (default).
- docs/CONTINUATION.md § "Active work" sub-headings listed as IN PROGRESS or BLOCKED.

Recent-work test binding. Each recent item's on-device test is discovered via §-letter or Fix match against

the project's test_ .sh filenames OR an explicit Test: line in the item heading. Items without an on-device test are \$. . four-layer- coverage violations (a fix without an on-device test is forbidden) – flagged as VALIDATION-PASS-BLOCKED until the test lands.

Authorization sequence (machine-enforced):

- . scripts/testing/recent_work_validate.sh --device <serial> runs; writes /data/local/tmp/.recent_work_validated on the device IFF all targeted tests return PASS with captured evidence.
- . The full suite (test_all_fixes.sh) checks for the marker file at entry; absent or stale (older than the last flash boot epoch) → refuses to proceed with exit code .
- . Marker is invalidated automatically on reboot – the orchestrator stores the device boot epoch in the marker and re-validates.

Anti-bluff during validation. Each recent-work item that landed a fix in this batch MUST have a paired \$. . RED test (captured before the fix) AND the same test now GREEN (captured after the fix). A GREEN with no prior RED is itself the bluff \$. . forbids – the gate FAILs.

Honest classification of validation FAILs. Per

\$. . : PRODUCT defect (real regression – block full-suite, surface to operator), TEST-INFRA defect (test-script bug – fix per \$. . at source layer, re-run), BLUFFY-THRESHOLD (e.g. SLO floor wrong – re-baseline), or UNCONFIRMED (not reproducible in isolation – tag per \$. . PENDING_CYCLE_RETEST, NEVER silently demote).

Composition. \$. . (STOP-on-discovery) +

\$. . (no- guessing) + \$. . (demotion-

evidence) + \$ (full-suite gate) +
 \$ (iteration discipline – smoke-test ≡
 recent-work- validation) + \$ (RED test for
 each recent-work item IS the validation test) +
 \$ (revision header determines recency) +
 \$ (CONTINUATION.md source-of-truth for \$
 Active work).

Gates:

- CM-COVENANT-114-46-PROPAGATION – anchor literal present in every CLAUDE.md / AGENTS.md across the covenant file set.
- CM-AF-RECENT-WORK-VALIDATION-GATE – asserts
 scripts/testing/recent_work_validate.sh exists + executable
 + sources the anti-bluff library + the full-suite
 orchestrator contains the entry-gate check.
- CM-AF-VALIDATION-ARTIFACT-FILE – asserts the marker file
 path /data/local/tmp/.recent_work_validated is literal-
 identical across helper + orchestrator + propagation-
 block + this Constitution section.

Paired mutations strip the anchor / remove the entry-gate
 / flip the marker path → respective gate FAILs.

No escape hatch. No --skip-validation , --full-suite-always ,
 --ignore-recent-work flag is permitted. The full-suite-
 without- validation pattern is the precise "tests passed
 but feature broken" failure mode \$ specifically
 prohibits. The – min validation pass is ALWAYS
 faster than – h of full-suite chasing a defect.

Classification: universal (per \$).

§ . . – Firebase Data Review Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"We MUST regularly before every bigger working round check Firebase information gathered so far: Crashlytics (all fatals, non-fatals and ANRs), Analytics data and Performance data. For anything problematic depending on severity proper workable items (issues) MUST be created (Issues, Issues_Summary docs) with full references (links) to original data on Firebase. We MUST make sure we do not create for same issues multiple entries, so we MUST distinguish between same issue manifested in different forms (stacktraces)! Everything MUST BE comprehensive and in-depth level of details so any changes we do (or fixes) do solve the original root problems!"

The mandate. Before every "bigger working round" (pre-build, pre-flash, pre-tag, daily, post-deployment burn-in) the operator/ loop MUST execute the Firebase review pass via `scripts/firebase/review_round.sh`. The pass queries Crashlytics (fatals + non-fatals + ANRs) + Analytics + Performance, classifies each finding by the § . . severity table, dedup-maps to existing Issues.md entries via the three-tier algorithm, and drafts new Issue entries for unrecognised findings. Skipping the pass is a § . PASS-bluff: Firebase IS the captured evidence from real end-user devices; ignoring it is the precise "tests pass, feature broken in the wild" failure mode § . prohibits.

Five mandatory elements (ALL must hold):

- **Trigger cadence** (trigger types). Pre-build (blocking, < h freshness), pre-flash (blocking, < h), pre-tag (blocking, < h), daily : UTC default (non-blocking sweep), post-deployment burn-in T+ h (non-blocking).
- **Three-source query**. Crashlytics + Analytics + Performance – ALL three. Skipping one source is a § . PASS-bluff (Firebase reports a regression on the skipped axis and we never see it).
- **Issues.md output**. Every "problematic" finding MUST map to either (a) a new Issues.md entry, OR (b) a recognised existing entry via the dedup algorithm. Both paths produce a full Firebase Console URL in the entry's metadata so any future agent can re-fetch the raw data without rebuilding context.
- **Three-tier deduplication**. Tier (exact Firebase Issue-ID match) → Tier (stacktrace-similarity cluster hash: top- non-generic frames normalised + crash class → SHA- first hex chars) → Tier (monthly operator merge review). Multiple Issues.md entries pointing at the same underlying root cause is a § . violation (operator chases ghosts; root cause stays unfixed).
- **Comprehensive root-cause analysis**. Every Firebase-sourced Issue carries the § . . (a) systematic-debugging output – Phase evidence, Phase pattern, Phase root-cause hypothesis (UNCONFIRMED until § . . RED test reproduces), Phase fix direction. Comment-only entries ("crash in App X") are PASS-bluff stubs and FAIL the Issue-xref gate.

Severity classification: Crashlytics FATAL impactedUsers > last d → Critical; – → Major; but historic → Low. NON-FATAL

```

/day → Major; - → Low; < →
Informational. ANR % sessions → Major; . - %
→ Low. Performance KPI regression % → Major;
- % → Low. Analytics funnel-drop > % →
Major; - % → Low; feature_used regression >
% → Major. Composes with $ . . Type
assignment.

```

Gates:

- CM-COVENANT-114-47-PROPAGATION – anchor present across every CLAUDE.md / AGENTS.md in the covenant file set.
- CM-AF-FIREBASE-REVIEW-CADENCE – scripts/firebase/review_round.sh exists, executable, contains the - stage pipeline literals + dedup algorithm + severity table.
- CM-AF-FIREBASE-ISSUE-XREF – every Issues.md entry whose ****Source:**** is **Firestore** / **Firestore Analytics** / **Firestore Performance** carries ****Firestore Issue IDs:**** + ****Firestore URL:**** + (****Stacktrace Cluster Hash:**** for Crashlytics OR ****KPI:**** for Performance OR ****Funnel:**** for Analytics).

Paired mutations strip the anchor / rename the review_round function / remove the Firestore-URL field from the template → respective gate FAILs.

Composition. \$. . (STOP-on-discovery),
 \$. . (a) (systematic-debugging), \$. . (no-guessing – Firestore data IS captured evidence),
 \$. . . (demotion-evidence – a previously- Fixed Issue resurfacing in Firestore IS positive captured evidence the fix didn't hold; reopen attribution = "AI: captured-evidence- contradicts" per \$. .),
 \$. . (credentials – never log the bearer token),
 \$. . (Issues_Summary sync), \$. .

(cleanup of /tmp/firebase_review_* work-dirs),
\$ (Status: Queued), \$ (Type per severity), \$ (Reopened-Details on Firebase-resurface), \$ (implicit step inclusion),
\$ (RED test per stacktrace before fix lands),
\$ (revision header on every Firebase-sourced Issue), \$ (Firebase Status.md at docs/firebase/status/Status.md), \$ (validation pass consults latest Firebase delta).

No escape hatch. No --skip-firebase-review, --no-issue-from-firebase, --firebase-review-not-applicable flag is permitted. The operator MAY filter by minimum severity (--severity-min major) to reduce noise but the pass itself MUST execute.

Classification: universal (per \$) – every consuming project that ships to real end-user devices benefits from the same review cadence and the same dedup algorithm.

\$ – UI-Driven Video Testing Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"We MUST make sure that all video playback testing with nd display is performed by fully automatically using the application's UI / UX, with full navigation, choice of video and clicks to real UI buttons to play or switch / stop videos being played! Maybe direct execution of Intents or Broadcasts does not have reported issues! We MUST test ALL WAYS users or Systems may / could play some video contents or streams! For all applications, all video and audio stream types! All supported codecs! EVERYTHING verified with nd display on D device and every

other device with connected and display and Avrus for proper codec(s) used in its Web Interface / Dashboard! Write as much as possible new tests for all supported test types!"

Why the existing video-test set is insufficient. The

legacy `test_video_routing.sh` (`v`),
`test_video_secondary_display.sh`,
`test_video_playback_automation.sh` (`v`) plus the
`per_app_video` wrappers all dispatch playback via
`am start -a android.intent.action.VIEW` or via `MediaSession`
`cmd media_session play` – surfaces that skip the app's own
content-selection UI, player-surface allocation,
`MediaSession` lifecycle (start/pause/stop), and back-button
cleanup. A defect like `$CL` routing-race or `$CM` frozen-
frame mid-back-press cannot reproduce via Intent dispatch
because Intent dispatch re-launches the app cold each time
and `am force-stop` is a clean cleanup path that masks the
leak.

Operative rule. Every video-playback test that asserts
secondary-display routing MUST traverse the user-
equivalent UI path. Mandatory –surface coverage per
video app:

- . **Surface A – Launch via launcher icon.** `am start -n`
`<launcher-pkg>/<launcher-activity>` then `uiautomator dump` →
locate app icon tile → `input tap X Y`. NOT
`am start <video-pkg>/<MainActivity>` (skips launcher).
- . **Surface B – Navigate to content list.** Per-app driver
script knows the app's home-tab → content-tab swipe /
tap sequence. Asserts at least one content tile is
visible via `uiautomator dump` content-desc / text match.
- . **Surface C – Select + play.** Tap a specific content tile
(driver records the picked tile's content-desc as
captured evidence). Wait for player surface visible.

- **Surface D – In-app control interaction.** Pause via in-app pause button (not `input keyevent KEYCODE_MEDIA_PAUSE`), resume via in-app play button, in-app rapid switch to a second video (covers §CL).
- **Surface E – Stop via back button.** `input keyevent KEYCODE_BACK` to exit player. Assert § . . .
cleanup: VOM `activeDecoder == null`, secondary surface cleared, no orphan MediaSession.

App-coverage matrix: every video-capable app in PRODUCT_PACKAGES MUST have a UI-driven driver script set (Layer . . . per Section . . . design) AND at least one scenario that uses it. Today: . . . apps covered fully in the initial batch (VK Video, MPV, Lampa as fully-implemented; remaining . . . as templated stubs per Section . . .). Expansion to all . . . video-app entries is queued as follow-up work under §0 Issues.md per-app matrix.

Stream-type matrix: progressive HTTP / HLS / DASH / RTMP / file- local / DRM-protected. Per app, the driver script tags which stream types its picked content exercises. Aggregator emits a coverage matrix per cycle so gaps are visible.

Codec matrix: H. . . . / H. . . . / VP . . . / AV . . . / MPEG- . . . / MP . . . V (video) + AC- . . . / E-AC- . . . / TrueHD / DTS / DTS-HD / MLP / Opus / AAC (audio). Per scenario, the codec assertion is made via `dumpsys media .metrics` (decoder name in event log) AND via Arvus codec-state probe (audio side) per §

Secondary-display verification. On any device with HDMI-A- . . . attached (detected via `dumpsys display | grep "Display 2" )`:

- Captured-evidence: `dual_display_record.sh` ON for both displays per §

- Assertion: `ffprobe -count_frames` on secondary mp reports

N frames within s of play-action timestamp

- Anti-assertion: `ffprobe -count_frames` on primary mp reports ONLY launcher/home pixels during playback window (no video-decoder output)
- VOM cross-check: `dumpsys video_output_manager` shows `activeDecoder != null` during playback and `== null` after stop within s

When secondary absent (D default topology): SKIP per \$. . with explicit reason "topology: no secondary display attached". NEVER FAIL on missing topology.

Arvus codec-state mandate. For every test asserting an audio codec (AC- , E-AC- , TrueHD, DTS, etc.) – composed with \$. . .

- \$CG: 2 3
- During playback window, call `arvus_probe_codec_state` ≥ times at s intervals (transients fade by sample -)
- Call `arvus_screenshot_capture` per \$CG to attach visual evidence of the dashboard
- Assert reported codec matches expected via `arvus_assert_codec_format <regex>`
- ARVUS_HOST unreachable (ABK not joined / sink off) → SKIP per \$. . , never FAIL

Pre-build gates:

- `CM-COVENANT-114-48-PROPAGATION` – asserts anchor literal present in every CLAUDE.md / AGENTS.md across parent + owned submodules + HelixQA dependencies (- file scan).

- `CM-AF-UI-DRIVEN-VIDEO-COVERAGE` – asserts the directory `device/rockchip/rk3588/tests/ui_driven/` exists, contains `lib/ui_driver.sh` reference + `scenarios/` subdir + per-app driver subdirs for at least the reference apps (`vk_video`, `mpv`, `lampa`), and that `scripts/testing/run_ui_driven_video_suite.sh` is executable.
- Paired mutations: deleting `ui_driver.sh` → `CM-AF-UI-DRIVEN-VIDEO-COVERAGE FAILS`; stripping the anchor literal → `CM-COVENANT- - -PROPAGATION FAILS`.

No escape hatch. No `--use-intent-shortcut` / `--skip-ui-traverse` / `--legacy-intent-mode` flag exists. The discipline exists because Intent-based dispatch is exactly the class of test that `PASSes` while users hit real defects (`$` forensic anchor). Operators who feel UI-driven tests are "too slow" should land scenario-level parallelism (multi-device fan-out) rather than bypass the discipline.

Composition.

- `$` . . . – topology dispatch (secondary present vs absent)
- `$` . . . – captured-evidence quality (every UI tap captured)
- `$` . . . – Arvus codec-state mandate
- `$` . . . – playback cleanup verified via UI back-button path
- `$` . . . – RED-test-first discipline (initial UI tests RED against `$CL/$CM`)
- `$` . . . – revision header on every plan/driver doc
- `$CG` – Arvus dashboard screenshot

Classification: universal (per `$` . . .). Applies to every project that consumes the constitution submodule AND ships any video-capable Android app.

\$. . – Dual-Approach Testing Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Kinopoisk playback of . audio supported movies MUST BE done via UI / UX full automation and via execution of Intents (or Broadcasts) directly both! We could have shared tests base and specialized part which will run / play the movie with properly chose audio track (.). Document everything up to the smallest details! Create as much as needed new tests for all supported test types and Challenges!"

Composition with \$. . . \$. . mandated UI-driven traversal for every video routing test, eliminating the Intent-only PASS-bluffs. \$. . REFINES rather than replaces: every feature test ships in BOTH variants – UI-driven AND Intent-driven – over a shared assertion base. Either alone is a \$. . PASS-bluff for the OPPOSITE half of the stack. UI catches app-side bugs (content selection, player surface allocation, MediaSession lifecycle, in-app overlays, back-button cleanup). Intent catches framework/system-server bugs (Intent extras parsing, MediaCodec.configure hook MIME routing, IVideoOutputManager bind timing, broadcast permission gating, headless cron automation paths).

Operative rule – mandatory elements:

- **Both variants required.** Every feature test exercising a user-visible behaviour MUST ship both `<feature>_ui.sh` (uiautomator-based, \$. . surfaces A-E) AND `test_<feature>_intent.sh` (`am start --es / am broadcast` -based). Either alone is forbidden.

- **Shared assertion base.** Codec-state assertions, captured-evidence collection (screen-recording, ffprobe, RMS amplitude analysis), Arvus codec-state probe + dashboard screenshot, and `$ . . cleanup` MUST be implemented in a single POSIX-sh library (`tests/lib/dual_approach_test_base.sh`) and reused by BOTH variants of every dual-approach test. Code duplication in the shared layer is forbidden.
- **Specialised driver code.** UI variant uses `$ . .` per-app driver scripts under `tests/ui_driven/<app>/`. Intent variant uses `am start --es content_id <id> --es audio_track_id <track> / am broadcast -a <action>` directly. Each variant's specialised layer is responsible ONLY for "how the playback gets started"; everything downstream of "playback is now happening" routes through the shared base.
- **Comprehensive documentation per test.** Every dual-approach test set ships with: (a) `$ . .` revision header in BOTH variant scripts, (b) per-feature contract under `docs/dual_approach/<feature>.md` describing the captured-evidence pairing and operator-side verification, (c) `Issues.md` / `Fixed.md` entry cross-linking both variant paths.
- **Kinopoisk . EAC is the canonical first implementation.** The shared base + Kinopoisk . UI variant + Kinopoisk . Intent variant land in the same batch as this mandate. Subsequent features (Netflix Dolby Atmos, MPV DTS-HD, Lampa+TorrServe HEVC, VLC FLAC) port to the same pattern. The \$CN subagent's fix to the underlying decoder pipeline is the FIRST GREEN target of these tests; both variants are RED per `$ . .` until \$CN lands.

Captured-evidence directory contract. Both variants write to mirror-structured directories

`qa-results/dual_approach/<F>/<run- ts>/{ui,intent}/`. Identical

filenames; orchestrator diffs the two `result.json` files. Status mismatch (UI=PASS but Intent=FAIL or vice-versa) is itself a finding – it pinpoints which half of the stack contains the bug.

Pre-build gates:

- `CM-COVENANT-114-49-PROPAGATION` – anchor literal in every `CLAUDE.md` / `AGENTS.md` across parent + owned submodules + HelixQA dependencies (`-file scan`).
- `CM-AF-DUAL-APPROACH-COVERAGE` – `tests/lib/dual_approach_test_base.sh` exists + sources `anti_bluff` / `ui_driver` / `credentials` + exports `dat_init` / `dat_start_capture` / `dat_assert_codec_state` / `dat_assert_video_frames` / `dat_assert_audio_channels` / `dat_arvus_dashboard_capture` / `dat_cleanup` / `dat_report_finding` .
- `CM-AF-KINOPOISK-5-1-DUAL-COVERAGE` – both `tests/ui_driven/kinopoisk/kinopoisk_5_1_play_movie_ui.sh` AND `tests/test_kinopoisk_5_1_play_movie_intent.sh` exist + both source the shared base + both reference `EAC_codec` + both assert `-channel` .
- Paired mutations (): strip anchor → propagation gate FAILs; delete UI variant → coverage gate FAILs; delete Intent variant → coverage gate FAILs.

No escape hatch. No `--ui-only` / `--intent-only` / `--skip-dual` flag exists. The discipline exists because either-alone is the PASS-bluff pattern \$. specifically prohibits. Operators who feel dual variants are "too slow" should land scenario-level parallelism rather than bypass the discipline.

Composition.

- \$. . – topology dispatch (Arvus reachable, secondary present)
- \$. . – captured-evidence content-quality analysis

- § 11.4.13 – UNCONFIRMED tagging for un-verified Kinopoisk element IDs
- § 11.4.14 – credentials never echoed, per-device per-service loader
- § 11.4.17 – Arvus codec-state mandate (out-of-band evidence)
- § 11.4.43 – playback cleanup via EXIT trap in shared base
- § 11.4.44 – universal classification
- § 11.4.49 – RED-first TDD (both variants RED until §CN fix lands)
- § – revision header on every variant script + design doc
- § – UI-driven traversal (this MANDATE refines, not replaces) 11 4 17
- §CG – Arvus dashboard screenshot via Presenter receiver
- §CB – credentials loader

Classification: universal (per § 11.4.50). Applies to every project that consumes the constitution submodule AND ships any Android app whose user-visible behaviour can be triggered through both UI and Intent/Broadcast paths. 2 0 2 6 0 5 1 8

§ 2026.05.18 – Deterministic Consistency Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Our anti-bluff / bluff-proofed / proof-driven work MUST satisfy the following: not a single feature, System or application flow, use case, edge case or procedural action(s) MUST NOT partially work, or sometimes work and sometimes not! There is only one truth that MUST BE fulfilled – no matter how many

times we repeat all these scenarios with variations in data, System state(s) and speed of execution, results MUST BE consistent and successful without exception! No false positives or bluffing of any kind is allowed! We MUST add as much full automation tests which will use real System and Applications UI and UX with all flows, use cases and edge cases as needed or as much as it is possible!"

Why this anchor exists. § 4.1.1 already forbids the vocabulary `intermittent` / `transient` / `flake` in closure narratives, but that's a textual ban – operators could still let a test "PASS once, FAIL once, PASS again" and report only the first PASS. § 4.1.2 closes that gap MECHANICALLY by requiring every PASS to come from N identical iterations rather than from a single observation.

Operative rule – mandatory elements:

- **N-iteration deterministic outcome.** Every test that PASSES MUST have been executed N times (default N= 10 , N= 100 for cycle- validation suites) against the same firmware MD + same device
 - same topology and produced IDENTICAL PASS in every iteration. A test whose outcome diverges across iterations is auto-FAIL per this mandate – there is no "first PASSEd therefore X was a flake" path. § 4.1.2 's expanded forbidden-vocabulary list is enforced mechanically here.
- **Edge-case + flow + use-case coverage.** Every public API path (Activity / Service / Broadcast receiver / ContentProvider URI / IPC interface / JNI entry / sysprop write / sysfs node / init.rc trigger) MUST have ≥ 1 dedicated test that drives it. Untested paths surface in a feature-coverage-matrix audit and block

release at the % threshold (ratchet sequence
→ → → mirrors \$. . the
existing project-side docs-coverage ratchet).

. **Reliability check helper.** Project anti-bluff helper
library ships `ab_run_n_times <test_name> <N> <fn> [args...]` :

- Loops N times, captures exit code + evidence-hash per iter
- Asserts all N exit codes identical AND all N evidence-hashes identical
- On any divergence: `ab_fail` with full N-iteration report
- Per \$. . forbidden vocabulary: NO operator-facing escape path converts a divergent N-iter run into PASS

. **Feature-coverage-matrix audit.** Project ships `scripts/testing/feature_coverage_audit.sh` :

- Walks source layers (every public API surface)
- Walks test layers (test_ .sh references)
- Emits coverage report at `qa-results/feature_coverage/<ISO-UTC>/coverage.tsv`
- Pre-build gate enforces minimum threshold; threshold ratchets up each phase (→ → →)

. **Composes with every prior anti-bluff anchor.**

\$. . does NOT replace \$. . /
\$. . / \$. . – it adds the N-iter consistency dimension across all of them. RED-first TDD (\$. .) still applies: every new test starts RED and becomes GREEN after the fix lands. UI-driven (\$. .) and dual-approach (\$. .) still apply: every test ships in both variants AND each variant passes N iterations identically.

Pre-build gates:

- `CM-COVENANT-114-50-PROPAGATION` – anchor literal `§11.4.50` present in every `CLAUDE.md` / `AGENTS.md` across parent + owned submodules + HelixQA dependencies (`≥` – file scan; phase uses `-canonical-file scan` matching `§ . .` propagation pattern).
- `CM-AF-RELIABILITY-CHECK-WIRED` – `ab_run_n_times` function literal present in `device/rockchip/rk3588/tests/lib/anti_bluff.sh` AND at least on-device tests source-reference it. The `-test` floor ratchets upward each phase.
- `CM-AF-FEATURE-COVERAGE-MATRIX` – `scripts/testing/feature_coverage_audit.sh` exists + executable + has run within the last days + most-recent coverage report meets the current threshold.

Paired mutations ():

- Strip the `§11.4.50` anchor literal from `constitution/CLAUDE.md` via `sed -i 's|11\.4\.50|11.4.MUTATED|g' → CM-COVENANT-114-50-PROPAGATION MUST FAIL.`
- Rename `ab_run_n_times` to `ab_run_n_times_DISABLED` in `anti_bluff.sh` via targeted `sed` → `CM-AF-RELIABILITY-CHECK-WIRED` MUST FAIL.
- Move `scripts/testing/feature_coverage_audit.sh` aside via `mv` → `CM-AF-FEATURE-COVERAGE-MATRIX` MUST FAIL.

No escape hatch. No `--allow-flake`, `--first-pass-suffices`, `--skip-n-iter`, `--skip-coverage-audit` flag exists. The discipline exists because the user mandate is unambiguous: "results MUST BE consistent and successful without exception."

Composition.

- `§ . .` – FAIL-bluffs forbidden

- § 11.4.6 – captured-evidence (every iteration captures its own)
- § 11.4.7 – quality analysis (per-iteration content compared)
- § 11.4.43 – no guessing (UNCONFIRMED required for any iteration- budget assumption)
- § 11.4.46 – forbidden flake/intermittent vocabulary (enforced mechanically by this mandate)
- § 11.4.49 – RED-first TDD (every new test starts RED, N-iter applies across the GREEN transition)
- § 11.4.49 – validate-before-suite (N-iter applied at single-test scope first)
- § 11.4.17 – UI-driven traversal (each UI traversal MUST be N-iter-consistent)
- § 11.4.17 – dual-approach (both variants of every dual-approach test MUST pass N iterations identically)

Classification: universal (per § 11.4.17). Applies to every project that consumes the constitution submodule.

§ 2026.05.18 – Live-ADB-First Maximization Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Was it possible to max the fix and test it via ADB or it can be only done by making the fix and reflashing for validation and verification? If such option exists, we should always make the best of it! Making fix on live devices via ADB → Once done commit and push changes, rebuild System and reflash → Validate and verify everything (with the test we write as well). Make sure this idea(s) is (are) part of our

root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD if they are not already! EVERY DETAIL IS IMPORTANT!!!! Not a single idea or idea's detail can be ignored, skipped, relativized or bluffed!"

\$. . . REFINES \$. . . step ("LIVE-ADB-PROBE – try the fix on the running device first") with mechanical enforcement: a per-file- class decision matrix, a classifier helper, and a commit-message footer literal.

Operative rule (mandatory elements):

- . **Classify every fix** by rebuild-requirement using the project's per-file-class decision matrix. No guessing (\$. . .). Every file in the diff cites the matrix row that classified it.
- . **If LIVE_ADB_TESTABLE**, the operator **MUST** apply the fix to the running device via the appropriate adb channel first (adb push for scripts, setprop persist.* for runtime properties, mount -o remount,rw /vendor for boot scripts, pm install -r for APKs built locally via gradle). Run the \$. . . RED test against the live-probed device, capture positive-evidence PASS, THEN commit source-side + rebuild + reflash as belt-and-suspenders re-validation. Commit message footer **MUST** state LIVE_ADB_VALIDATED: yes with one-line description of the live-probe sequence.
- . **If REQUIRES_REBUILD**, the operator proceeds directly to source-side fix + rebuild + flash. Commit message footer **MUST** state REQUIRES_REBUILD: <reason> citing the matrix row that classified the file. Categories of REQUIRES_REBUILD: kernel, AOSP framework Java / AIDL, native C++ inside APEX, sepolicy, init.rc, ro.* build properties (immutable post-boot), Android.bp / Android.mk / BoardConfig.mk , XML resource overlays, codec XML inside APEX, ramdisk-bundled artifacts.

- **If mixed batch**, commit footer states `LIVE_ADB_VALIDATED: partial` and enumerates per-file classification. Per § . . . batching is preferred; live-probe the testable files first, then batch with the rebuild-required ones into a single rebuild.
- **Helper script enforces classification mechanically.** The project provides `classify_fix_rebuild_requirement.sh` which walks `git diff --name-only`, looks up each file against the matrix, and emits the recommended commit-message footer. No operator-facing override flag converts an unclassified path to `LIVE_ADB_TESTABLE` by default – unmatched paths classify as `REQUIRES_REBUILD: unmatched-path` (safe default per § . . .).

Per-file-class decision matrix (canonical):

File class	Rebuild required?	Reason
<code>*test_*.sh</code> / <code>tests/lib/*.sh</code> (on-device)	NO	adb push to / data/local/tmp/ tests/
Host-side test scripts	NO	host script – no device interaction
<code>atmosphere-*.sh</code> boot scripts	DEPENDS	adb push to / vendor/bin/ after remount
Markdown docs / Constitution / plans / Issues / Fixed / CONTINUATION	NO	host docs – no firmware impact
Forked-player Kotlin (Presenter, MPV,	YES via gradle local	APK re-link required but no

File class	Rebuild required?	Reason
SmartTube, TorrServe, Lampa, VLC, etc.)	build + adb install -r	full firmware rebuild
Framework Java (frameworks/base/**/*.*.java)	YES	system_server requires reboot at minimum
AIDL (*.aidl)	YES	stub generation requires recompile
Native C++ (external/**/*.*.cpp , frameworks/native/**/*.*.cpp)	YES	builds into APEX or system library
APEX libraries (vendor/**/lib/**/*.*.so inside APEX path)	YES	squashfs read-only by Mainline integrity
Kernel sources (kernel-5.10/**)	YES	requires kernel rebuild + boot.img
sepolicy (*.te)	YES	requires policy build
init.rc	YES	parsed at boot only
ro.* build properties	YES	immutable post-boot
persist.* runtime properties	NO	setprop persist.* mutable
XML resource overlays	YES	RR0 recompile
media_codecs_*.xml	YES	APEX-bundled

File class	Rebuild required?	Reason
display_settings.xml (userdata)	DEPENDS	adb push possible if SELinux permits
Android.bp / Android.mk / device.mk / BoardConfig.mk	YES	rebuilt-into-image artifacts
Test fixture binary assets	NO	adb push to /data/local/tmp/

Pre-build gates: CM-COVENANT-114-51-PROPAGATION (anchor across canonical files) + CM-AF-CLASSIFY-FIX-HELPER-EXISTS (helper present + sentinel literals) + CM-AF-LIVE-ADB-FIRST-COMMIT-MARKER (advisory WARN that scans recent commits for the footer literal). Three paired meta-test mutations.

Composition.

- § 11.4.6 – TDD-fix workflow (§ 11.4.6 REFINES step with the mechanical classifier).
- § 11.4.6 – batch-source-fixes-before-rebuild (compose: live-ADB- test individually, then batch-commit).
- § 11.4.6 – no guessing (each classification justified by matrix row; unmatched paths default to rebuild-required, never silent testable).
- § 11.4.6 – validate-recent-work-before-post-flash (live-probe IS implicit pre-validation).
- § 11.4.6 – UI-driven tests (live-probe CAN use uiautomator over adb shell).
- § 11.4.6 – dual-approach (live-probe applies to both UI and Intent variants).

- § 1.1.1 – deterministic consistency (live-probe runs N iterations on the live device before commit).

No escape hatch – no `--skip-classify`, `--assume-rebuild`, `--no-footer-required` flag exists. The discipline exists because the user mandate is unambiguous: "EVERY DETAIL IS IMPORTANT!!!! Not a single idea or idea's detail can be ignored, skipped, relativized or bluffed!"

Classification: universal (per § 1.1.1). Applies to every project that consumes the constitution submodule – the matrix itself MAY be project-specific (each consumer adapts the file-class list to its own source tree), but the mandate to classify-before-commit + the LIVE_ADB_VALIDATED / REQUIRES_REBUILD footer literals + the classify-helper-script + the three pre-build gates are universal.

2026 05 18

§ 1.1.2 – Autonomous-Validation Mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Make sure we have full automation tests which will do all this work in full automation! IMPORTANT: Make sure that all existing tests and Challenges do work in anti-bluff manner – they MUST confirm that all tested codebase really works as expected! We had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can't be used! This MUST NOT be the case and execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users of the product! This MUST BE

part of Constitution of our project, its CLAUDE.MD and AGENTS.MD if it is not there already, and to be applied to all Submodules's Constitution, CLAUDE.MD and AGENTS.MD as well."

Why this anchor exists. § . . . (full-automation coverage) and § . . . (no-fakes-beyond-unit-tests + %-test-type-coverage) and § . . . (per-feature on-device end-user validation) collectively mandated that every user-facing feature has automation tests with captured runtime evidence. They did NOT explicitly forbid operator-attended-only validation paths – i.e., a feature whose ONLY proof of working state requires a human to physically present at the device, drive UI manually, and observe outcomes. Phase .§CN exposed the gap: §CN's Kinopoisk . EAC closure was structurally fixed (libavutil bundled in APEX, six decoders register) yet end-user-validation degraded to "operator drives remote control while test polls" because Kinopoisk's TV-Compose UI is not accessibility-instrumented and uiautomator returns near-empty hierarchy. That validation path is operator-attended, non-reproducible in CI, and cannot prove deterministic consistency (§ . . .) across N iterations. The User mandate elevates "full automation" from "preferred" to "mandatory": every user-facing PASS MUST have at least one captured-evidence path that does NOT require operator presence.

Operative rule (mandatory elements):

- . Every user-facing feature MUST have at least one **autonomous validation path**. "Autonomous" means: the validation runs end-to-end via adb shell + scripted automation, produces captured runtime evidence per § . . . , and reaches a PASS / FAIL verdict WITHOUT a human present to drive UI, observe screen, or make

decisions. Operator-attended tests are SUPPLEMENTARY, never PRIMARY. A feature whose ONLY validation path is operator-attended is a § . . . violation regardless of how carefully the operator's captures match § . . . evidence requirements – the path does not scale to CI, does not run on every commit, does not survive operator unavailability, and produces the exact "tests pass but feature doesn't work for users" failure mode § . . . specifically forbids.

- . **Acceptable autonomous-validation paths** (any one or more – non-exhaustive):
 - **Programmatic instrumentation APK** – a small platform-signed app that exercises the under-test surface directly via SDK APIs (e.g., `MediaCodec.createDecoderByName`, `MediaCodec.configure`, `AudioTrack` writes) and writes a structured JSON result file to a discoverable path. Closes UI-instrumentation gaps where the target app is not accessibility-instrumented.
 - **Headless intent dispatch + state poll** – `am start --es / am broadcast --es` to dispatch the under-test code path, then poll `dumpsys`, `/proc/<pid>/maps`, `media.metrics`, `/sys/...`, kernel sysfs, or the integration's network probe for the expected state transition.
 - **ADB-driven uiautomator** – applicable ONLY if the target app exposes accessibility nodes the dump can resolve. MUST verify via `uiautomator dump | grep -c clickable=true` ≥ before claiming uiautomator coverage is real. A near-empty hierarchy is itself the captured-evidence proving UI-driven automation is INFEASIBLE for that target and the test MUST fall back to instrumentation-APK or headless-intent paths.
 - **Network-side sink probe** – for features that surface on a network-accessible peripheral (e.g., Arvus HDMI

dashboard at `http://<sink>/`, Sonos REST API, AirPlay receiver), the sink's own report counts as autonomous out-of-band evidence per §

- **HelixQA autonomous QA session** – for projects that have HelixQA fully incorporated per § , the orchestrator's end-to-end session output counts as autonomous evidence.

- **Per-feature classification + tracking.** Every entry in the project's coverage ledger (§) MUST classify each feature's autonomous-validation path as one of: `AUTONOMOUS_VERIFIED` (path exists + last-run-green per §), `AUTONOMOUS_DESIGNED` (path exists but RED per §), `OPERATOR_ATTENDED_ONLY` (path requires human presence – release blocker until promoted to one of the other states), or `NOT_APPLICABLE` (e.g., a CLI tool that has no UI surface and the unit + integration tests already constitute autonomous proof). `OPERATOR_ATTENDED_ONLY` rows MUST cite a tracked work item per § + § that designs the autonomous-path migration. Closing the work item requires the migration to land, not just the operator capture.

- **Anti-bluff for autonomous paths themselves.** An autonomous path that returns PASS without exercising the user-visible behaviour is a § violation just as severe as the original § PASS-bluff pattern. The autonomous path MUST produce positive captured evidence per § (audio: channel count + sample rate + glitch census; video: frame count + routing target + frame health + obstruction census). A grep on a metadata field that doesn't prove behaviour is not an autonomous validation – it is a metadata bluff dressed up as automation. Paired meta-test mutations

.. (\$.) MUST catch the autonomous path's degradation
.. to metadata-only or grep-only.

Composition.

- § . . – full-automation-coverage mandate. § . . is the strict expansion of invariant + : "anti-bluff posture" and "proof of working capability" cannot be carried by operator-attended-only evidence – they MUST be carried by at least one autonomous path.
- § . . – no-fakes-beyond-unit-tests + %-test-type-coverage. § . . is the operational layer that closes the gap between "the project has full-automation tests" and "every PASS has an autonomous evidence path".
- § . . – per-feature on-device end-user validation. § . . refines: the on-device scenario MUST itself be autonomously drivable.
- § . . – TDD-fix-discipline. The autonomous path MUST exist in the RED state BEFORE the fix lands, then go GREEN after. An autonomous path authored after the fix is a § . PASS-bluff.
- § . . – UI-driven video testing. § . . establishes that when uiautomator returns a near-empty hierarchy on the target app, UI-driven IS infeasible and the test MUST fall back to instrumentation-APK or headless-intent. The fallback is REQUIRED – not optional.
- § . . – dual-approach testing. The Intent variant in a dual-approach pair often IS the autonomous path when the UI variant requires operator assistance.
- § . . – deterministic consistency. Autonomous paths run N iterations naturally; operator-attended paths cannot satisfy the N-iteration discipline at scale.

- `$ . .` – live-ADB-first maximization. The instrumentation-APK path is itself `LIVE_ADB_TESTABLE` (`gradle build + adb install -r`
 - `run + read JSON result`), so live-probe IS the design loop for the autonomous path.

Pre-build gates:

- `CM-COVENANT-114-52-PROPAGATION` – anchor literal `§11.4.52` present in all canonical files (`constitution/Constitution.md`, `constitution/CLAUDE.md`, `constitution/AGENTS.md`, plus the parent consumer's `CLAUDE.md` + `AGENTS.md`) PLUS the full per-consumer propagation (parent project + every owned submodule
 - nested submodules + HelixQA dependencies – same scan pattern as `CM-COVENANT-114-51-PROPAGATION`).
- `CM-AF-AUTONOMOUS-PATH-PER-FEATURE` – every entry in the consuming project's coverage ledger (`$ . .`) has a non-empty autonomous-classification column with value in `{AUTONOMOUS_VERIFIED, AUTONOMOUS_DESIGNED, OPERATOR_ATTENDED_ONLY, NOT_APPLICABLE}`.
 - `OPERATOR_ATTENDED_ONLY` rows MUST cite a tracked migration work item. (Project-side gate – consuming projects implement the ledger surface in their own pre-build test suite.)

Paired mutations (per `$ . .`):

- Strip `§11.4.52` literal from `constitution/Constitution.md` → `CM-COVENANT-114-52-PROPAGATION` FAILs.
- Inject a feature row with `OPERATOR_ATTENDED_ONLY` and NO tracked migration work item → `CM-AF-AUTONOMOUS-PATH-PER-FEATURE` FAILs.

No escape hatch. No `--allow-operator-attended-only`, `--skip-autonomous-path`, `--manual-validation-suffices` flag exists. The discipline exists because the user mandate is

unambiguous: "execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users". Operator-attended-only validation cannot guarantee that property at the release-gate layer – only autonomous paths can.

Classification: universal (per § 11.4.17). Applies to every project consuming the constitution submodule. Each consumer adapts the autonomous-validation path catalogue to its own technology surface (Android: instrumentation APK + uiautomator + headless intent; Web: Playwright headless + API probe; CLI: subprocess

- stdout/stderr assertion; mobile native: device farm headless runner) but the mandate to provide at least one autonomous path per user-facing feature is universal.

2026 05 18

§ 11.4.17 – Fixed_Summary parity mandate (User mandate, – –)

Forensic anchor – verbatim user mandate
(– – T : Z):

"Note: Just like for Issues we have Issues_Summary, for Fixed we MUST HAVE Fixed_Summary – like all other docs: ALWAYS in sync and up to date and ALWAYS exported into the PDF and HTML! Add this mandatory rule / constraint into the root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD."

Why this anchor exists. § 11.4.17 already established that docs/Issues_Summary.md is the canonical short-form

summary of `docs/Issues.md` and MUST always be regenerated + re-exported (HTML

11 4 19

- PDF). whenever the source changes. § . . . added the column- alignment requirement so `Fixed_Summary.md` mirrors `Issues_Summary.md` shape. § . . . closes the propagation gap: the parity discipline that § . . . imposes on `Issues / Issues_Summary` MUST apply symmetrically to `Fixed / Fixed_Summary`. The mechanical behaviour already exists in `scripts/testing/sync_issues_docs.sh` (lines - regenerate `Fixed_Summary` in stage b alongside `Issues_Summary` in stage a), but the mandate was never explicit in the Constitution – and an unstated rule is a § . . . PASS-bluff risk: future refactors of `sync_issues_docs.sh` could drop the `Fixed_Summary` half without violating any canonical authority. § . . . makes the existing behaviour explicit canonical authority so the gate set can defend it.

Operative rule. `docs/Fixed_Summary.md` is the symmetric short- form summary of `docs/Fixed.md`. It MUST be regenerated whenever `Fixed.md` changes. Its HTML + PDF exports MUST travel with the markdown (identical mtimes within `sync_issues_docs.sh` granularity). Stale exports are § . . . violations regardless of whether the underlying `.md` is correct – an operator (or future agent) reading the HTML or PDF gets a divergent view of which items are closed, and the § . . . CONTINUATION resumption guarantee silently breaks. Same discipline as § . . . `Issues_Summary` applied to `Fixed.md`.

11 4 19

Generator. `scripts/testing/generate_fixed_summary.sh` is the canonical generator. It MUST exist + be executable + emit a markdown table whose header columns include `Status` and `Type` (per § . . . column-alignment) so the `Fixed_Summary.md` shape mirrors `Issues_Summary.md` exactly.

Auto-sync wrapper. `scripts/testing/sync_issues_docs.sh` is the single operator-facing entry point. It already regenerates BOTH `Issues_Summary` AND `Fixed_Summary` in one shot (stages `a + b`), then exports HTML + PDF (stage `a`), then colorizes per `$EDITOR` (stage `b`), then re-renders the PDFs from the colorized HTML (stage `b`). MUST be invoked after any edit to `Fixed.md` (just as `$EDITOR` requires after any edit to `Issues.md`). Manual invocation of just the Issues half (`--issues-only -style` flag) is FORBIDDEN – the wrapper has no such flag and `$EDITOR` prohibits adding one.

Sort order. Same pattern as `$EDITOR` – by closure date DESC (most-recent-Fixed first), with `$-letter / Fix-` secondary sort. Documented at the top of the generated file.

HTML + PDF travel-together. Per `$EDITOR` + `$EDITOR`, the three file types (`.md` , `.html` , `.pdf`) MUST always have identical mtimes within `sync_issues_docs.sh` granularity. Stale exports violate `$EDITOR` regardless of whether the underlying `.md` is correct.

Composition:

- `$EDITOR` – `Issues_Summary` parity is the sibling rule `$EDITOR` symmetrizes. `$EDITOR` and `$EDITOR` are the canonical pair: edits to one source-of-truth (`Issues.md` or `Fixed.md`) trigger regeneration of BOTH summaries via the same wrapper.
- `$EDITOR` – atomic `Issues→Fixed` migration triggers `Fixed_Summary` regeneration. When an item closes (status `Fixed (→ Fixed.md)` / `Implemented (→ Fixed.md)` / `Completed (→ Fixed.md)` per `$EDITOR`), the operator moves the entry from `Issues.md` to `Fixed.md` atomically and runs `sync_issues_docs.sh` – both summaries refresh.

- § 11.4.23 – visual-cue & grouping colorizer post-processes both `Issues_Summary.html` AND `Fixed_Summary.html`. Stale `Fixed_Summary.html` would carry pre-§ 11.4.33 styling – itself a § 11.4.33 violation.
- § 11.4.33 – type-aware closure-status vocabulary. `Fixed_Summary` MUST respect the three terminal values (`Fixed (→ Fixed.md)`, `Implemented (→ Fixed.md)`, `Completed (→ Fixed.md)`) and emit them literally in the Status column.
- § 12.10 – revision-header mandate applies to `Fixed_Summary.md` exactly as it applies to every other tracked Markdown doc.
- § 11.4.33 – CONTINUATION.md resumption guarantee depends on the divergent-summary problem NOT existing; § 11.4.33 is the explicit symmetric closure of that risk for the Fixed half.

Pre-build gates:

- `CM-FIXED-SUMMARY-SYNC` – canonical artifact-level gate (invariants): () `docs/Fixed_Summary.md` exists; () `docs/Fixed_Summary.html` exists with `mtime` ≥ `Fixed_Summary.md` `mtime`; () `docs/Fixed_Summary.pdf` exists with `mtime` ≥ `Fixed_Summary.md` `mtime`; () `Fixed_Summary.md` `mtime` ≥ `Fixed.md` `mtime`; () `scripts/testing/generate_fixed_summary.sh` exists + executable; () `scripts/testing/sync_issues_docs.sh` invokes the generator. Pre-build FAILs if any invariant is violated. Partially overlaps with `CM-FIXED-COLUMN-ALIGNMENT` (§ 11.4.33 Phase 1.AV gate) and `CM-DOCS-EXPORT-SYNC` (§ 11.4.33 -doc + -export scan), but is explicitly the `Fixed_Summary` derived-doc parity invariant for § 11.4.33 enforcement.
- `CM-COVENANT-114-53-PROPAGATION` – anchor literal §11.4.53 present across all canonical files (`constitution/`

Constitution.md , constitution/CLAUDE.md , constitution/AGENTS.md , parent consumer's CLAUDE.md + AGENTS.md). Same propagation pattern as CM-COVENANT-114-51-PROPAGATION and CM-COVENANT-114-52-PROPAGATION . Consumer projects propagate further (owned submodules, nested submodules, HelixQA deps) at their own per-project gate granularity.

Paired mutations (per § 11.4.53):

- Strip §11.4.53 literal from constitution/CLAUDE.md → CM-COVENANT-114-53-PROPAGATION FAILS.
- Move scripts/testing/generate_fixed_summary.sh aside (so the invariant-() executable-check fails) → CM-FIXED-SUMMARY-SYNC FAILS.
- Touch docs/Fixed_Summary.md to a date older than docs/Fixed.md → invariant-() violated → CM-FIXED-SUMMARY-SYNC FAILS.

No escape hatch. No --skip-fixed-summary-sync , --issues-only , --summary-not-applicable flag exists. The discipline exists because the user mandate is unambiguous: "like all other docs: ALWAYS in sync and up to date and ALWAYS exported into the PDF and HTML". A divergent Fixed_Summary breaks operator + AI-agent ability to discover what is fixed vs what is still open – and that breakage IS the § 11.4.53 "tests pass but reality differs" pattern at the documentation layer.

Classification: universal (per § 11.4.53). Applies to every project consuming the constitution submodule that maintains an Issues.md / Fixed.md split (the Issues-tracking lifecycle is already a § 11.4.53 universal mandate, so this clause inherits the same universal scope). Projects that do not maintain a Fixed.md (e.g., single-binary throwaway prototypes) are NOT_APPLICABLE; all projects that DO maintain a Fixed.md MUST land both

gates and mutations within one cycle of adopting this Constitution version.

2026-05-19

§ . . – ATM-NNN ticket identifier
mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Every workable item in Issues.md /
Issues_Summary.md / Fixed.md / Fixed_Summary.md MUST
carry a stable, unique, auto-incremental ATM-NNN
ticket identifier. ATM- prefix, monotonic, never
renumbered, append-only – once assigned to an item it
stays bound to that item for the lifetime of the
project across every reopen, migration, or rename."

Why this anchor exists. § . . + § . . +
§ . . + § . . established the lifecycle
vocabulary for tracked items, but the items themselves
were identified only by §-letter or Fix- – both of
which are **unstable across migrations**. §-letter is heading-
local; when an item moves from Issues.md to Fixed.md the
§-letter may be reused for a different item. Fix- is
only assigned on closure for items that touched source
code; tasks / features / pure-doc items have no stable
identifier at all. Cross-references between commits, test
logs, captured-evidence artifacts, Reopens.md history
(§ . .), README.md doc-link rows (§ . .),
and external systems (Firebase Crashlytics Issue-IDs,
Atlassian-style tracker rows) all need an identifier that
is **monotonic, stable, and project-wide unique**.
§ . . closes that gap.

Operative rule. Every workable item in `docs/Issues.md` AND `docs/Fixed.md` MUST carry a `[ATM-NNN]` ticket identifier in its heading, in the form `## $X.Y. [ATM-NNN] <title>` (or `### $X.Y.` for nested items). NNN is a positive integer, zero-padded to at least `3` digits (`ATM-001` , `ATM-042` , `ATM-127` , ...). Identifiers are allocated by a **canonical helper script** (`scripts/testing/ assign_atm_ticket_ids.sh`) that maintains an append-only state file (`scripts/testing/.atm_ticket_state.json`). Once assigned to an item, an ATM-NNN MUST NEVER be:

- . **Renumbered** – the bind is permanent.
- . **Reused** for a different item even if the original item is force-deleted (the helper's state file detects collisions).
- . **Decrementd** – counter is monotonic-only.
- . **Skipped intentionally** – gaps are forbidden in the assignment sequence (gaps indicate state-file corruption and trigger gate FAIL).

State file format. `scripts/testing/.atm_ticket_state.json` is a JSON-lines file with one record per allocated ID:

```
{"atm_id": "ATM-001", "heading_hash": "<sha256 of normalized heading>",  
  "type": "Bug|Feature|Task", "current_location": "Issues|Fixed",  
  "current_status": "<lifecycle value>", "reopens_count": <int>,  
  "created_at": "<ISO 8601 UTC>", "last_modified": "<ISO 8601 UTC>"}
```

The `heading_hash` is the canonical binding key – when the helper scans `Issues.md` / `Fixed.md`, it computes the normalized hash of each heading and looks up an existing ATM-NNN before allocating a new one. This protects against accidental renumbering on heading reflows / wording edits (the hash MUST stay stable; if the heading text changes substantially, the state file's `heading_hash` field is updated in place via an explicit migration command, NOT auto- rebound).

Summary-table column. `docs/Issues_Summary.md` and `docs/Fixed_Summary.md` MUST carry a leftmost `ATM ID` column showing the identifier. Generators (`generate_issues_summary.sh`, `generate_fixed_summary.sh`) MUST emit this column as the first data column so operators / agents can sort + filter on it. ▢

Composition:

- `$11 4 54 . .` (Status), `$ . .` (Type),
`$ . .` (column-alignment Issues Fixed),
`$11 4 12 . 11 4 53` (type-aware closure terminal values).
 All four pre-existing mandates compose with
`$ . .` – the ATM-NNN is the new universal join key.
- `$11 4 55 . .` + `$ . .` (Issues_Summary + Fixed_Summary regen on source changes) – the helper MUST be invoked from `sync_issues_docs.sh` so newly added items receive ATM-NNN identifiers before summary regeneration.
- `$ . .` (Reopens-history per-item docs) – the per-item Reopens.md lives at `docs/issues/<ATM-NNN>/Reopens.md`.
- `$ . .` (README.md doc-link section) – every Issues / Fixed item may be linked by ATM-NNN in external references.
- `$ . .` (revision header) – applies to the state file's audit log if one is maintained.

Pre-build gates:

- `CM-ATM-TICKET-IDS-COMPLETE` – every workable-item heading in Issues.md AND Fixed.md MUST carry an `[ATM-NNN]` token matching `\[ATM-[0-9]{3,}\]`. Headings without one FAIL the gate.
- `CM-ATM-TICKET-IDS-UNIQUE` – no two items share an ATM-NNN. The state file's record count MUST equal the union of

Issues.md + Fixed.md [ATM-NNN] occurrences (no orphan IDs, no duplicates).

- CM-ATM-TICKET-IDS-MONOTONIC – the maximum ATM-NNN in the state file MUST equal record_count (no gaps).
- CM-COVENANT-114-54-PROPAGATION – anchor literal §11.4.54 present across canonical files (constitution/
Constitution.md , constitution/CLAUDE.md , constitution/
AGENTS.md , parent consumer's CLAUDE.md + AGENTS.md).

Paired mutations (per § .):

- Remove [ATM-NNN] from one Issues.md heading → CM-ATM-TICKET-IDS-COMPLETE FAILS.
- Duplicate an ATM-NNN in the state file → CM-ATM-TICKET-IDS-UNIQUE FAILS.
- Delete ATM-NNN= from a -record state file → CM-ATM-TICKET-IDS-MONOTONIC FAILS (gap at).
- Strip §11.4.54 literal from constitution/CLAUDE.md → CM-COVENANT-114-54-PROPAGATION FAILS.

No escape hatch. No --skip-atm-assignment , --renumber , or --no-atm-id-required flag exists. The discipline exists because unstable item identity is the exact failure mode that produces "the test was passing last week, what changed?" sessions where the operator cannot tell whether a given item is the same one tracked before or a coincidentally similarly-named successor.

Classification: universal (per § . .). Applies to every project consuming the constitution submodule that maintains an Issues / Fixed split (cf. § . . scoping). Projects that do not maintain such a split are NOT_APPLICABLE.

§ 2.2.6.5.19 – Reopens-history tracking + per-item Reopens.md doc (User mandate, - -)

Forensic anchor – verbatim user mandate

(- -):

"Add a Reopens-count column to Issues / Issues_Summary / Fixed / Fixed_Summary. For any item whose reopens-count > , create docs/issues/ATM-NNN/Reopens.md (+ HTML + PDF) with comprehensive reopen history + Fixed cycles. Each reopen MUST include By (AI / User), On (date), Reason, Evidence, and each Fixed-marking with its reasoning chain."

Why this anchor exists. § 2.2.6.5.19 mandated that Reopened status carry a `**Reopened-Details:**` block – but that block lives in the **current** Issues.md heading and is overwritten on the next state change. A complete reopen / fix-cycle audit trail across the lifetime of an item was not previously canonical. § 2.2.6.5.19 makes the full history first-class: every item that has been reopened at least once gets a dedicated `Reopens.md` companion doc capturing the entire chronological timeline – every reopen, every closure, every reasoning chain – so an operator / agent investigating "why was this reopened five times" has a single, authoritative source to read.

Operative rule. Every workable item with `reopens_count > 0` MUST have a companion document at `docs/issues/ATM-NNN/Reopens.md` (plus its HTML + PDF exports per § 2.2.6.5.19 + § 2.2.6.5.20 sync discipline). The document MUST contain:

- **Revision header** per § 2.2.6.5.19 (`Revision` integer + `Last modified` ISO UTC).

- . **Item identification** – ATM ID , Title , Type , Current Status , Current Location (Issues.md / Fixed.md), link back to the live heading.
- . **Cycle counters** – Total reopens , Total fixed cycles (items can reopen multiple times, each followed by a re-closure).
- . **Chronological timeline** – one entry per state-change event, each entry capturing:
 - **By:** AI or User (per § . . source attribution).
 - **On:** ISO date.
 - **Event:** Opened | Reopened | Fixed | Implemented | Completed .
 - **Reason:** the closed-vocabulary value from § . . (test-failed / manual-testing-detected / captured-evidence-contradicts / end-user-report / cycle-re-discovered / design-reconsidered) for reopens; captured-evidence summary for closures.
 - **Evidence:** path to or short description of captured artefact (test log path, recording path, sink-probe report, operator quote).
 - **Outcome:** what the next state transition was.
- . **Reasoning chain for each closure** – for every Fixed / Implemented / Completed event, the reasoning chain that justified marking it closed (root-cause analysis link, captured-evidence under same-conditions per § . . , gate / mutation pair that defends against regression).
- . **Most-recent state-change pointer** – explicit cross-reference to the current Issues.md / Fixed.md heading.

Summary-table column. Issues_Summary.md and Fixed_Summary.md MUST carry a Reopens column showing the integer count. When the count > the cell MUST be a

hyperlink to `docs/issues/<ATM-NNN>/Reopens.md` . Generators emit the column; the colorizer (§ 11.4.44) MAY apply a visual cue when `reopens >` (repeated-reopen signal – investigate before closing again).

Composition:

- § 11.4.44 (Reopened-Details on current heading) is the **per-event** capture; § 11.4.45 is the **per-item history** aggregation. Both layers MUST be in sync – the state-machine helper updates `Reopens.md` on every `Reopened-Details` parse.
- § 11.4.46 (ATM-NNN) provides the stable identifier so the per- item doc has a permanent path even after `Issues→Fixed` migration.
- § 11.4.50 (revision header) applies to every `Reopens.md`.
- § 11.4.51 (Status.md per-integration) – the per-item `Reopens.md` is the item-scoped analogue of the integration-scoped `Status.md`.
- § 11.4.52 (Fixed_Summary parity) – `Reopens` column appears on BOTH summary tables symmetrically.

Pre-build gates:

- `CM-REOPENS-DOC-EXISTS-WHEN-COUNT-GT-ZERO` – for every state-file record with `reopens_count > 0` , `docs/issues/<ATM-NNN>/Reopens.md` MUST exist (+ HTML + PDF). Missing doc FAILS.
- `CM-REOPENS-DOC-REVISION-HEADER` – every `Reopens.md` carries the § 11.4.50 revision header.
- `CM-REOPENS-COL-IN-SUMMARIES` – both `Issues_Summary.md` and `Fixed_Summary.md` carry a `Reopens` column.
- `CM-COVENANT-114-55-PROPAGATION` – anchor literal §11.4.55 across canonical files.

Paired mutations (per § 11.4.55):

- Delete docs/issues/ATM-NNN/Reopens.md for an item with reopens_count=2 → CM-REOPENS-DOC-EXISTS-WHEN-COUNT-GT-ZERO FAILS.
- Strip revision header from a Reopens.md → CM-REOPENS-DOC-REVISION-HEADER FAILS.
- Remove Reopens column from Issues_Summary.md table header → CM-REOPENS-COL-IN-SUMMARIES FAILS.
- Strip §11.4.55 literal from constitution/CLAUDE.md → CM-COVENANT-114-55-PROPAGATION FAILS.

No escape hatch. No --skip-reopens-doc, --collapse-history, --reopens-not-applicable flag. The discipline exists because loss of reopen history is the exact failure mode that produces "why does this keep happening" sessions with no auditable answer.

Classification: universal (per § 11.4.55). Applies wherever § 11.4.55 applies (every project tracking Reopened status).

2026 05 19

§ 11.4.55 – Status_Summary parity + two-audience format (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Every Status.md doc gets a Status_Summary parity companion ALWAYS in sync, exported to HTML + PDF. Two-page format: page = non-developer audience (team-specific: audio team for audio Status, video team for video Status, etc.), page = software engineer

summary. Auto-generated after every main Status update."

Why this anchor exists. § 4.1.1 established `Status.md` per integration domain – but `Status.md` targets engineers (cites gate names, §-letter references, commit hashes, captured-evidence file paths). Non-developer stakeholders (audio team, video team, QA leads, product reviewers) need a plain-language version of the same content, **always derived** from the engineering `Status.md` so the two views can never diverge. § 4.1.2 introduces the symmetric companion doc with a two-audience layout.

Operative rule. For every `docs/<domain>/<integration>/Status.md` that satisfies § 4.1.1, a companion `docs/<domain>/<integration>/Status_Summary.md` MUST exist with:

- **Revision header** per § 4.1.1.
- **Page 1** – **For the <team>** (audience-specific heading; the team name is derived from the integration domain – audio team for `docs/dolby/*`, video team for `docs/video/*`, etc.). Page 1 contains:
 - Plain-language summary of current state (3 - paragraphs).
 - **What works** (3 - bullets, end-user-visible behaviour).
 - **What's broken or pending** (3 - bullets).
 - **Operator / team actions** if any.
 - NO code references, NO §-letter jargon, NO captured-evidence file paths, NO gate / mutation names.
- **Page 2** – **For software engineers** – same content density as `Status.md` but condensed:
 - §-letter references, gate names, commit hashes, captured- evidence file paths.

- Cross-references to Issues.md / Fixed.md items by ATM-NNN (\$. .).

- **HTML + PDF exports** travel with the markdown (identical mtimes within sync wrapper granularity).
- **Auto-generation** – the canonical helper scripts/testing/generate_status_summary.sh <Status.md path> produces both pages. The generator **MUST** be invoked from a sync wrapper (e.g. scripts/testing/sync_integration_status.sh per \$. .) every time the source Status.md is updated.

Page- audience derivation. The generator maps domain → audience name via a project-controlled mapping table (e.g. audio → Audio team , video → Video team , network → Network team , drm → DRM / streaming team). Unmapped domains default to Project team . The mapping table lives in the project's generator config, NOT in the Constitution (project-specific data per \$. .).

Composition:

- \$. . (Status.md per integration) – Status_Summary.md COMPLEMENTS, never replaces. Both files exist together.
- \$. . + \$. . (parity discipline) – Status_Summary follows the same regeneration / HTML+PDF / mtime-alignment pattern.
- \$. . (revision header) – applies to Status_Summary.md.
- \$. . (colorizer) – Status_Summary.html receives the same type/status visual cues if it embeds tracked-item references.
- \$. . (CONTINUATION.md) – non-developer stakeholders read Status_Summary; engineers read Status.md + CONTINUATION.md + Issues.md.

Pre-build gates:

- `CM-STATUS-SUMMARY-EXISTS-FOR-EVERY-STATUS` – for every `docs/**/*.Status.md`, the sibling `Status_Summary.md` MUST exist (+ HTML + PDF) with `mtime ≥ Status.md mtime - tolerance`.
- `CM-STATUS-SUMMARY-TWO-AUDIENCE` – every `Status_Summary.md` contains BOTH the `Page 1 – For the <team>` heading AND the `Page 2 – For software engineers` heading. Missing either FAILs.
- `CM-STATUS-SUMMARY-REVISION-HEADER` – `$ . .` header present.
- `CM-COVENANT-114-56-PROPAGATION` – anchor literal across files.

Paired mutations (per `$ . .`):

- Delete a `Status_Summary.md` → `CM-STATUS-SUMMARY-EXISTS-FOR-EVERY-STATUS` FAILs.
- Remove `Page 1` heading from a `Status_Summary.md` → `CM-STATUS-SUMMARY-TWO-AUDIENCE` FAILs.
- Strip `§11.4.56` literal from `constitution/CLAUDE.md` → `CM-COVENANT-114-56-PROPAGATION` FAILs.

No escape hatch. No `--skip-status-summary`, `--engineer-only`, `--no-audience-split` flag. Plain-language stakeholder access is not negotiable – the project ships to humans, not just to engineers.

Classification: universal (per `$ . .`). Applies wherever `$ . .` applies (every project that maintains domain-scoped `Status.md` docs).

\$. . - README.md doc-link section +
 revision metadata (User mandate,
 - -)

Forensic anchor – verbatim user mandate

(- -):

"Add a doc-link section to README.md – links to Issues + Issues_Summary + Fixed + Fixed_Summary + CONTINUATION + ALL Status docs + their exports. Each link shows revision + last-modified."

Why this anchor exists. Operators, AI agents, and external reviewers arriving at the project repo for the first time need a single discoverable entry point listing every canonical tracked- items + status document, along with each document's freshness (revision + last-modified). Without this surface, agents resort to `find docs -name '*.md'` which returns hundreds of files unranked.

\$. . makes the README.md the canonical index.

Operative rule. Every project's top-level `README.md` MUST contain a section titled `Tracked-Items + Status Documents` (or the project's localized equivalent – the section heading MUST contain the literal `Tracked-Items` for gate detection). The section MUST be a markdown table with columns: `Document` (human- readable name), `Last modified` (ISO UTC from the doc's \$. . revision header), `Revision` (integer from same header), `Markdown` (relative link to `.md`), `HTML` (relative link to `.html`), `PDF` (relative link to `.pdf`).

The section MUST link to:

. `Issues.md` , `Issues_Summary.md` (\$. . ,
 \$. . , \$. .).

- `Fixed.md` , `Fixed_Summary.md` (`$` . . ,
`$` . .).
- `CONTINUATION.md` (`$` . .).
- Every `docs/**/Status.md` and its `Status_Summary.md` pair
(`$` . . , `$` . .). Status docs are auto-
discovered by the generator via `find docs -name`
`'Status.md'` .

Generator. `scripts/testing/update_readme_doc_links.sh` is the canonical helper. It:

- Scans the canonical doc paths.
- Extracts each doc's `Revision` + `Last modified` fields
from its `$` . . header.
- Renders the markdown table with current values.
- Replaces the section between explicit markers (`<!-- doc-`
`link-section:begin -->` and `<!-- doc-link-section:end -->`) in
`README.md`.

The wrapper MUST be invoked from `sync_issues_docs.sh` (so every `Issues.md` / `Fixed.md` update refreshes the README freshness data) AND from `sync_integration_status.sh` (so every `Status.md` update likewise refreshes README).

Composition:

- `$` . . + `$` . . + `$` . .
(`Issues` / `Issues_Summary` / `Fixed` / `Fixed_Summary` parity)
– README links to all four.
- `$` . . (revision header) – README pulls
`Revision` + `Last modified` from each doc's header.
- `$` . . (`Status.md`) + `$` . .
(`Status_Summary.md`) – README enumerates every Status
pair.
- `$` . . (`CONTINUATION.md`) – explicitly linked from
README so operators discover it on arrival.

Pre-build gates:

- `CM-README-DOC-LINK-SECTION-PRESENT` — `README.md` contains the literal `Tracked-Items` heading AND both section markers (`<!-- doc-link-section:begin -->` , `<!-- doc-link-section:end -->`).
- `CM-README-DOC-LINK-ROWS-COMPLETE` — every canonical doc (`Issues`, `Issues_Summary`, `Fixed`, `Fixed_Summary`, `CONTINUATION`, every `Status.md` + `Status_Summary.md` found under `docs/`) appears as a row inside the section.
- `CM-README-DOC-LINK-FRESHNESS` — the `Last modified` value in each row matches the corresponding doc's § . . header within sync-wrapper granularity (no row staler than the source doc by more than the sync wrapper's tolerance).
- `CM-COVENANT-114-57-PROPAGATION` — anchor literal across files.

Paired mutations (per § .):

- Strip the section markers from `README.md` → `CM-README-DOC-LINK-SECTION-PRESENT` FAILs.
- Remove the `CONTINUATION.md` row from the section → `CM-README-DOC-LINK-ROWS-COMPLETE` FAILs.
- Backdate a `Last modified` value in a row → `CM-README-DOC-LINK-FRESHNESS` FAILs.
- Strip §11.4.57 literal from `constitution/CLAUDE.md` → `CM-COVENANT-114-57-PROPAGATION` FAILs.

No escape hatch. No `--skip-readme-doc-links` , `--collapse-status-rows` , `--no-freshness-check` flag. The discipline exists because invisible docs are non-existent docs from a discoverability standpoint, and the § . . covenant specifically prohibits surfaces where claims (here: doc freshness) can drift unnoticed.

Classification: universal (per § . .). Applies to every project that maintains an `Issues` / `Fixed` split or

Status.md docs (i.e., to every project consuming this Constitution that has any of § . . . / § 1.1.4.5.8 . . . / § . . . in scope).

§ . . . – Parallel-development methodology (User mandate, – –)

Forensic anchor – verbatim user mandate
(– – T~ : Z MSK):

"We MUST DO one comprehensive research and planning in the background in parallel with current mainstream work: our current methodology of the development is very slow. It takes us days to fix a bunch of bugs, add some changes or features and all this to be tested, verified and shipped. From iteration to iteration we are slower and slower. We MUST CREATE adjusted improved version of working methodology where multiple workable items could be done in parallel, all come into the central (main) branch as they are done, then we at particular moment rebuild and reflash the System and in background full testing with validation and verification is done. Each parallel work on one workable (or more) item(s) must use parallel agents as much as possible! We MUST HAVE proper synchronization mechanism between parallel working routines so they do not cause conflicts, break features or create any other types of the problems! Document everything – the whole plan, whole methodology, working flow with in depth diagrams and graphs, all user guides and manuals and then add this all into our root (constitution Submodule) Constitution, CLAUDE.MD and AGENTS.MD! Make sure we

start using this ASAP since we MUST increase efficiency and productivity a couple of times now! Writing in depth tests (all supported types of the tests) with Challenges and full HelixQA use is MANDATORY! Every test we execute besides executed with success MUST RESULT in proof that actual functionality being tested REALLY DOES WORK with NO BLUFF of any kind! Heavt enforcement of no-bluff / anti-bluff policy IS MANDATORY!"

Why this anchor exists. Sequential phase-chain working pattern (Phase 39.A → Phase 39.B → ... → Phase 39.DT , observed - - to - - : , unique Phase .X sub-phases across days = ~ /day sub-phase rotation, commits / days = ~ / day) serialises on five distinct bottlenecks: (B) commit_all.sh flock at .git/.commit_all.lock , (B) single-thread sub-phase chaining via docs/CONTINUATION.md hand-offs, (B) rebuild-blocking on every REQUIRES_REBUILD fix (\$. -j cap = - h rebuild + min flash

- min validate per cycle), (B) single-thread test_all_fixes.sh sweep (~ - min × phase-device combinations), (B) subagent dispatch latency (one-at-a-time spawn + wait pattern). Net effect: single-item wall-clock - days where productive work is < h. The User mandate requires the methodology to multiply throughput several-x.

Operative rule. Project work proceeds through the Parallel Work Unit (PWU) pipeline rather than sequential phase-chain. Each PWU is a self-contained workable item with mandatory components (ATM-NNN identifier per

\$. . , Issues.md entry per
\$. . +\$. . , file- scope manifest,
\$. . RED test, source patch, pre-build gate per

§ (b) layer , post-flash test per
 § (b) layer , paired § meta-test
 mutation, § (b) layer HelixQA Challenge bank
 entry, captured-evidence directory per
 § +§). PWUs execute through five
 pipeline stages: **Stage DEVELOP** (parallel, N PWU
 agents in isolated worktrees), **Stage MERGE** (serial,
 conductor processes one PWU at a time via `commit_all.sh`
`flock + § - step merge-first pipeline),`
Stage REBUILD+FLASH (parallel where possible: AOSP
 build inside § bounded scope, D +D flash
 serialised on USB hardware), **Stage VALIDATE** (parallel
 on D +D +meta-test+ coverage-audit), **Stage SWEEP**
 (parallel: HelixQA Challenge bank, Issues→Fixed migration,
 README.md doc-link refresh per § , HTML+PDF
 exports). Stage of round N+ starts WHILE Stages
 + of round N are still running – this is the primary
 throughput multiplier.

Synchronization mechanism. Four-layer lock hierarchy: L
 parent serial flock `.git/.commit_all.lock` (held only during
 Stage , released between PWUs), L per-submodule git
 operations (implicit via cascade), L contention-path
 advisory locks at `qa-results/pwu-locks/<sha256(path)>.lock` for
 the forbidden cross- PWU paths (CLAUDE.md/
 AGENTS.md/Constitution.md/CONTINUATION.md/ Issues.md/
 Fixed.md/pre_build_verification.sh/
 meta_test_false_positive_ proof.sh/build.sh/
 commit_all.sh+push_all.sh/device.mk+BoardConfig.mk/
 atmosphere-.sh/init..rc), L per-PWU git worktree (zero
 cross- PWU file conflicts during Stage). Disjoint-
 scope PWUs (different test files, different APKs,
 different drivers, different framework services) run fully
 parallel. Conflict detection at Stage entry:

```
git fetch --all --prune --tags + git diff main...HEAD --<scope>
```

- cross-PWU scope overlap check. Overlap detected → PWU rejected to Stage for rebase (per \$. . step integration-before-force-push).

Anti-bluff enforcement. Conductor REFUSES to merge a PWU lacking all four anti-bluff checks: C \$. . RED test captured-evidence (file exists + non-zero exit code in log + timestamp predates source patch – proves test catches regression), C \$. paired meta-test mutation (applied + gate FAILs under mutation + restored + gate PASSes – proves gate is not a bluff gate), C \$ deterministic- consistency (iterations for normal, for cycle-validation, ALL identical exit codes AND identical evidence-hashes – eliminates \$. . flake/transient/intermittent demotion path), C \$. . captured-evidence (audio: WAV with non-trivial RMS + ffprobe channel count; video: screen recording with ffprobe frame count > + analyzer matched event; UI: uiautomator dump with under-test element state; or HelixQA result.json PASS verdict + positive-evidence chain). HelixQA coverage is MANDATORY for every user-visible PWU per User mandate – Challenge bank entry in tools/helixqa/banks/atmosphere.yaml referencing the PWU's ATM-NNN + dispatching to the on-device test + scoring PASS only on positive captured evidence. Metadata-only PASS / configuration- only PASS / absence-of-error PASS / grep-without-runtime-evidence PASS all REJECTED at the merge gate.

Agent orchestration. Four roles: **Conductor** (, main thread) – issues ATM-NNN, dispatches PWU agents, runs Stage merge serially, runs Stages – , owns all contention-path edits, maintains docs/CONTINUATION.md \$ live snapshot. **PWU agent** (– background) – one PWU each, Stage work in isolated worktree, signals qa-

results/pwu-<atm>/READY_FOR_MERGE on completion. **Validator** (background) - Stage work in parallel on D +D +meta-test+coverage- audit. **Sweeper** (background) - Stage work in parallel with next round's Stage . Max concurrent agent count ~ at peak, bounded by \$. % memory budget (~ GB total ≤ GB budget). Composition with \$ host-session safety: conductor refuses dispatch if `host_check_safety` fails; per \$. AOSP build still hard-capped at -j inside bounded scope; per \$. + \$. b + \$. c concurrency gates still block on stray gradle/kotlin daemons.

Composition map (mandatory):

- \$. . - four-layer test coverage per PWU (pre-build + post-build + on-device + HelixQA bank).
- \$. . - audio/video quality analysis comprehensiveness (PWU captured-evidence directory).
- \$. . - no-guessing mandate (every PWU claim backed by captured forensic evidence or explicit `UNCONFIRMED:` tag).
- \$. . - demotion-evidence rule (/ deterministic-consistency closes the demotion-bluff path).
- \$. . - batch-source-fixes-before-rebuild (Stage → Stage batch threshold).
- \$. . + \$. . + \$. . - PWU status + type vocabulary + type-aware closure terminal value.
- \$. . - atomic Issues.md → Fixed.md migration on PWU closure.
- \$. . - Pre-Force-Push Merge-First -step pipeline at Stage entry.

- § 1.1.4.43 – Iteration-discipline (PWU pipeline IS the iteration conductor; § 1.1.4.45 binds steps to PWU stages –).
- § 1.1.4.49 – TDD RED-first per-PWU enforced at merge time.
- § 1.1.4.50 – Integration-status-doc maintenance (PWU's Status.md updated at every stage transition).
- § 1.1.4.52 – Dual-approach testing (PWU MUST ship both UI-driven AND Intent/Broadcast-driven variants when applicable).
- § 1.1.4.54 – Deterministic-consistency (–iter normal / –iter cycle-validation merge-time enforcement).
- § 1.1.4.57 – Autonomous-validation (PWU's on-device test MUST run without operator presence per Stage parallel design).
- § 1.2.6.11 – ATM-NNN ticket identifier (PWU identifier replaces Phase .X letters).
- § 1.2.7 – README.md doc-link refresh at Stage .
- § 1.2.8 – % memory-budget ceiling (caps concurrent agent count).
- § 1.2.10 – AOSP m -j hard-cap at -j (Stage build cap).
- § 1.2.10 + § 1.2.10 b + § 1.2.10 c – concurrency hardening (Stage preflight refuses on running gradle/kotlin/git-pack-objects/load-per-cpu spike).
- § 1.2.10 – CONTINUATION.md maintenance (conductor owns the file by § 1.2.10 forbidden cross-PWU rule).
- § 1.2.10 – data safety (every history rewrite during Stage merge-first step backed by hardlinked .git backup).

Pre-build gates:

- `CM-PWU-LOCK-HIERARCHY` – verify lock hierarchy script exists at `scripts/testing/pwu_acquire_path_lock.sh` and `scripts/testing/pwu_release_path_lock.sh`, both executable, both reference the forbidden cross-PWU paths from § 11.4.58 of `docs/guides/PARALLEL_DEVELOPMENT_METHODOLOGY.md`.
- `CM-PWU-ANTI-BLUFF-COVERAGE` – for every PWU with status "Ready for merge" or beyond, assert all four anti-bluff checks (`C -C`) have captured-evidence files under `qa-results/pwu-<atm>/evidence/`. PWU lacking any of `red-test-output` / `mutation_patch` / `deterministic-consistency` / `captured-feature-proof` FAILs the gate.
- `CM-PWU-MERGE-QUEUE-DISCIPLINE` – scan `qa-results/pwu-queue/` for PWUs that bypassed the merge queue (commits on main without a corresponding `qa-results/pwu-queue/merged/<atm>` marker). Any bypass FAILs the gate.
- `CM-PWU-PARALLEL-AGENT-LIMIT` – assert no more than concurrent background PWU agents (count `tmux/screen` sessions or `worktree` branches) – `over-limit` FAILs the gate (per § 11.4.58 memory budget).
- `CM-COVENANT-114-58-PROPAGATION` – anchor literal §11.4.58 present in `constitution/CLAUDE.md` + `constitution/AGENTS.md` + parent `CLAUDE.md/AGENTS.md` + every owned-submodule `CLAUDE.md/ AGENTS.md` (files total in the current consuming project's family).

Paired mutations (per § 11.4.58):

- Move `scripts/testing/pwu_acquire_path_lock.sh` aside → `CM-PWU-LOCK-HIERARCHY` FAILs.
- Touch a PWU's `READY_FOR_MERGE` marker without `red-test-output.log` → `CM-PWU-ANTI-BLUFF-COVERAGE` FAILs.
- Land a commit on main without a corresponding `qa-results/pwu-queue/merged/<atm>` marker → `CM-PWU-MERGE-QUEUE-DISCIPLINE` FAILs.

- Spawn a `th` concurrent PWU agent → `CM-PWU-PARALLEL-AGENT-LIMIT` FAILS.
- Strip `§11.4.58` literal from `constitution/CLAUDE.md` → `CM-COVENANT-114-58-PROPAGATION` FAILS.

No escape hatch. No `--skip-merge-queue`, `--allow-bypass`, `--no-anti-bluff-check`, `--unlimited-agents`, `--sequential-phase-chain-mode` flag exists in any pipeline helper. The discipline exists because (a) the User mandate explicitly requires the multiplier ASAP and (b) the `§` covenant specifically prohibits PASS-bluffs (which the four anti-bluff merge-time checks mechanically prevent). Operators who feel they need to bypass the pipeline for a specific high-urgency fix should split the work into a single trivial PWU and let the pipeline run normally – Stage `1` of a trivial PWU is `< 1 min`, Stage `2` is `< 1 min`, Stage `3` is the same rebuild cost they would pay anyway.

Classification: universal (per `§ 1.1.1`). Applies to every project consuming this Constitution that has multiple owned-submodules

- a sequential commit/push tool + a captured-evidence testing discipline. Project-specific implementations (agent dispatch helpers, worktree layout, batch thresholds) live in consumer-side `docs/guides/PARALLEL_DEVELOPMENT_METHODOLOGY.md` (this Constitution defines the discipline; consumers implement the tooling).

Reference: parent `docs/guides/PARALLEL_DEVELOPMENT_METHODOLOGY.md` (full methodology with diagrams, templates, operator manual, and migration plan).

\$. . – README always-sync mandate
(User mandate, – –)

Forensic anchor – direct user mandate (verbatim,
– –):

"fully review and update our main README document. Some points are not valid anymore, some are missing. Make sure main README is among documents we MUST ALWAYS keep updated and in Sync with the projects and other documentation! Make sure we always export it (on every update) into PDF and HTML. This mandatory rules/constraints MUST BE all added into the root (constitution Submodule) Constitution, AGENT.MD and CLAUDE.MD!"

README.md at the project root is a \$. . -class always-sync document. It MUST be:

- . **Reviewed and updated whenever** new docs / integrations / Status.md entries appear, new submodules land, applied-fixes count changes, or canonical paths shift.
- . **Kept in lockstep** with docs/CONTINUATION.md (\$. .) and the Issues / Issues_Summary / Fixed / Fixed_Summary doc set (\$. . , \$. .) – same always-sync discipline.
- . **Exported to .html and .pdf on every update.** New helper scripts/testing/sync_readme_export.sh performs the pandoc + weasyprint export; it is auto-invoked by scripts/testing/sync_issues_docs.sh so a single doc-sync run refreshes the entire doc surface (Issues / Issues_Summary / Fixed / Fixed_Summary / CONTINUATION / README – md + html + pdf).
- . **Carry a \$. . revision header** (Revision N + Last modified ISO timestamp) at the top of the file.

- . **Contain a Documentation Map section** linking to every Status.md
 - Status_Summary.md + spec + plan + guide + script-companion doc + changelog + the constitution submodule, plus a per-audience navigation (end user / developer / QA / agent).
- . **Self-contained** – no hyperlinks to ephemeral external systems as the only source of truth.

Stale exports are § 12.1 violations regardless of whether the underlying README.md is correct – an operator (or future agent) reading the HTML or PDF gets a divergent view, and the § 12.1 CONTINUATION resumption guarantee silently breaks. Same discipline as § 12.1 Issues_Summary and § 12.1 Fixed_Summary applied to README.md.

Captured-evidence enforcement. Pre-build gate CM-README-EXPORT-SYNC locks four invariants: (a) README.md exists, (b) README.html exists, (c) README.html mtime ≥ README.md mtime, (d) README.pdf mtime ≥ README.md mtime (skipped gracefully if weasyprint is unavailable on the host but enforced when the PDF is present^{11 4 12}). Paired meta-test mutation backdates README.html + README.pdf so the source is newer → gate FAILs.^{11 4 23}

Composition. Composes with § 12.1 (Issues_Summary parity), § 12.1 (script-companion docs), § 12.1 (visual-cue HTML colorizer not applicable to README's narrative format), § 12.1 (revision header on every Status.md and now README.md), § 12.1 (Status.md integration), § 12.1 (Fixed_Summary parity – README is the canonical sibling at the project-overview layer), § 12.1 (Status_Summary parity), § 12.1 (README.md is the canonical index per Phase .EW+), § 12.1.

(CONTINUATION.md resumption guarantee depends on README's Documentation Map being current).

No escape hatch. § 1.1.4.59 has NO operator-facing override flag. No `--skip-readme-sync`, `--no-readme-export`, `--readme-stale-OK` flag exists. The discipline exists because README.md is the first file any new agent / operator / contributor reads – letting it diverge from reality is the exact § 1.1.4.59 PASS-bluff pattern but at the project-overview layer.

Classification: universal (per § 1.1.4.59). Applies to every project consuming this Constitution. The doc-sync helper, pandoc + weasyprint dependencies, and gate name are universal; the specific Documentation Map contents are consumer-side per-project.

Canonical authority: constitution submodule `Constitution.md`
§ 1.1.4.60.

Non-compliance is a release blocker regardless of context.
2.0.2.6 0.5 1.9

§ 1.1.4.60 – Documentation always-sync
composite covenant (User mandate,
– –)

Forensic anchor – verbatim user mandate
(– – ~ : Z):

"Double check if all documents are properly tied with our root Constitution, CLAUDE.MD and AGENTS.MD so they are always up to date, always in sync and exported into PDF and HTML! ... Issues, Issues_Summary, Fixed, Fixed_Summary, Continuation, Status and Status_Summary for all contexts (areas) – THEY ALL MUST BE REGULARLY UPDATED, IN SYNC AND CONSISTENT without giving at any

moment false picture about the state of the project or particular area(s) of it!"

The covenant. Eight documentation classes constitute the project's living state surface. They MUST be in sync at all times across markdown + HTML + PDF artefacts. Per-class anchors `$ . . / $ . . /`
`$ . . / $ . . / $ . . /`
`$ . . / $ . . / $ . .` each govern one class; `$ . .` binds them via a single mechanically-enforced composite invariant so the failure mode "one per-class gate silently disabled while operator reads a divergent HTML/PDF" is structurally impossible.

Eight bound doc classes (artefact triple – `.md` + `.html` + `.pdf`):

- `docs/Issues.md` – open-item tracker (`$ . . + $ . . + $ . .`)
- `docs/Issues_Summary.md` – auto-generated short form (`$ . .`)
- `docs/Fixed.md` – closed-item archive (`$ . .`)
- `docs/Fixed_Summary.md` – auto-generated short form (`$ . .`)
- `docs/CONTINUATION.md` – resumption handoff (`$ . .`)
- `README.md` – project overview + doc-link index (`$ . . + $ . .`)
- `docs/**/Status.md` (every domain-scoped instance) – per-integration status (`$ . .`)
- `docs/**/Status_Summary.md` (every domain-scoped instance) – auto-generated short form (`$ . .`)

Mandatory composite gate. `CM-DOCS-COMPOSITE-SYNC` (pre-build, device/rockchip/rk /tests/ pre_build_verification.sh):

- For every doc-class instance with `.md` present, verify `.html` mtime $\geq$ `.md` mtime AND `.pdf` mtime $\geq$ `.md` mtime.
- For every `Status.md` instance, also verify the sibling `Status_Summary.md` mtime $\geq$ `Status.md` mtime (if the summary exists).
- Walks `docs/` recursively for `Status.md` – the Status fleet is not enumerated statically because new integrations land per phase.

Composite gate FAILs the build if ANY single instance fails. The per-class gates (`CM-ISSUES-SUMMARY-SYNC`, `CM-DOCS-EXPORT-SYNC`, `CM-FIXED-SUMMARY-SYNC`, `CM-CONTINUATION-DOC-INSYNC`, `CM-README-EXPORT-SYNC`) remain in place –
`$ . . .` is the belt-and-suspenders layer that catches what they miss.

Paired mutation (per `$ . .`).

`meta_test_false_positive_proof.sh` backdates `docs/Issues.html` to year `( touch -t 200001010000.00 )` and asserts `CM-DOCS-COMPOSITE-SYNC` FAILs.

Revision-header coupling. Per `$ . .` every one of the doc classes carries the `**Revision:** N + **Last modified:** ISO 8601 UTC` header directly under the H . `CM-DOCS-COMPOSITE-SYNC` does NOT re-check the revision header (already covered by `CM-DOC-REVISION-HEADER-PRESENT`) – composite gate is artefact-mtime only, by design narrow.

Auto-sync wrapper coupling. Per `$ . . .` every edit to an Issues/Fixed-class doc flows through `sync_issues_docs.sh`. Per `$ . .` every edit to a `Status.md` flows through `sync_integration_status.sh <Status.md>`. Per `$ . .` every edit to `README.md` flows through

`sync_readme_export.sh` . `$` . . does NOT change the wrapper contract – it only catches the failure mode where an operator (or AI agent) bypassed the wrapper and committed markdown without regenerating exports.

No escape hatch. No `--skip-composite-doc-sync` , `--allow-stale-html` , `--summary-not-applicable` flag exists. The covenant exists for the operator's own protection – the moment a single per-class gate is bypassed, divergence accumulates silently until release-time forensics finds it (the exact `$` PASS-bluff pattern the canonical covenant was authored to eliminate).

Composes with `$` (Issues_Summary),
`$` (Status field), `$` (Type field), `$` (atomic Issues→Fixed migration),
`$` (colorizer), `$` (type-aware closure vocabulary), `$` (revision header),
`$` (Status.md), `$` (Fixed_Summary), `$` (Status_Summary),
`$` (README doc-link), `$` (README always-sync), `$` (CONTINUATION maintenance).

Propagation. Pre-build gate `CM-COVENANT-114-60-PROPAGATION` enforces the `$` anchor literal in every CLAUDE.md / AGENTS.md across parent + owned submodules + HelixQA dependencies (same shape as existing `$` X propagation gates). The propagation gate is OPTIONAL for PhaseFC's initial landing – only the composite gate + its paired mutation are mandatory; propagation gates trail per established phase pattern.

Classification: universal (per `$`). Applies to every project consuming this Constitution that maintains any of the doc classes. The composite gate logic is universal; the specific doc-class instance paths are consumer-side per-project.

Canonical authority: constitution submodule Constitution.md

§ 1.1.4.6.1 . . .

Non-compliance is a release blocker regardless of context.

2024 05 19

§ 1.1.4.6.1.1 – Mandatory Markdown metadata table + structured-doc ToC (User mandate, - -)

Forensic anchor – direct user mandate (verbatim, - -):

"For every Markdown document which contains structured content (with headings / sections and sub-sections) make sure that every time we apply change to the structure, table of contents on the top of the document is created or updated! This is MANDATORY for every structured MARKDOWN document. Automatically its PDF and HTML versions MUST BE (re)generated!"

"Introduce ... revision number, date and time of creation, date and time of last modification, other useful information we have in documents such as Issues, Issues_Summary, Fixed, Fixed_Summary, Status, Status_Summary, Continuation and similar. ... make this mandatory for EVERY Markdown document from now on, update root constitution Submodule with these changes and commit and push it all to all upstreams."

1.1.4.6.1

1.1.4.6.1

Scope of § 1.1.4.6.1.1 . . . This clause adds two universal disciplines that sit alongside (NOT replacing) the § 1.1.4.6.1 universal export mandate. § 1.1.4.6.1.1 governs the existence and freshness of .html + .pdf siblings; § 1.1.4.6.1.1 governs the content of the Markdown sources: how revision/status metadata is displayed at the top, and how the document's navigation

map (Table of contents) is kept in lockstep with the heading structure. The two clauses compose – § . . guarantees readers see *something current*; § . . guarantees what they see communicates its own provenance and is internally navigable.

A. Canonical metadata table (supersedes § . . format). The § . . bold-line revision header is superseded by a **Markdown table** placed immediately after the document's H title (and any blank line that follows). Pre-existing bold-line headers remain historically valid but **MUST** migrate to the canonical table at the next substantive edit. The canonical table has these **MANDATORY** rows:

Field	Value
Revision	positive integer; starts at , monotonic, never reset
Created	ISO date – first commit touching the file
Last modified	ISO date – most recent commit touching the file
Status	one of <code>active</code> / <code>draft</code> / <code>deprecated</code> / <code>super-seded</code>

ENCOURAGED rows (**REQUIRED** when the document tracks workable items or has a known continuation; render with – / `none` when N/A):

Field	Value
Status summary	one-line current-state hook
Issues	comma-separated workable-item IDs

Field	Value
Issues summary	one-line summary of open issues
Fixed	comma-separated workable-item IDs of resolved items
Fixed summary	one-line summary of fixes
Continuation	next action or linked CONTINUATION.md anchor

Additional rows MAY be added for project-specific metadata as long as the MANDATORY rows are present.

B. Structured-document Table of contents. Any tracked *.md with two or more H sections ("structured content") MUST include a ## Table of contents section immediately after the metadata table. The ToC MUST list every H and H in document order with anchor links and MUST be regenerated on every structural change (heading added / removed / renamed / reordered). A stale ToC is a \$. PASS-bluff: operators reading the rendered Markdown see a navigation map that no longer matches the body.

Scope (IN-scope under \$. .). Every tracked *.md in the \$. . INCLUDED set (project-root *.md , docs/**/*.md , scripts/**/*.md companion docs, owned-submodule trees, constitution/**/*.md , owned HelixQA dependencies). The \$. . exclusions for CLAUDE.md / AGENTS.md / README.md are SUPERSEDED here too – explicit metadata in the file itself is preferred over indirect tracking via a separate VERSION file because operators reading the rendered document see the freshness directly.

Narrow exemptions (NOT IN-scope). Same EXCLUDED set as § 4.1.1 (external/**, prebuilts/**, out/**, build/**, application-internal/*.md shipped with source code, non-owned third-party submodule trees). Additional exemptions for the structural rules: LICENSE, LICENSE.md, NOTICE, VERSION, OWNERS, machine-generated CHANGELOG.md.

PDF + HTML siblings. Deferred to § 4.1.2. The metadata table and ToC are *content* requirements; their export-freshness guarantee lives in § 4.1.3's CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC gate. A § 4.1.4 metadata or ToC change is a *.md modification, which transitively triggers § 4.1.5 regeneration – the two clauses compose without duplicate enforcement.

Anti-bluff captured-evidence gates (planned):

- CM-MD-METADATA-PRESENT – walks every tracked *.md (minus exemptions) and asserts (a) H1 present, (b) within 10 lines below H1 a Markdown table contains all four MANDATORY rows (| Revision |, | Created |, | Last modified |, | Status |).
- CM-MD-TOC-PARITY – for every *.md with ≥ 2 H1 sections, asserts a ## Table of contents is present and its entries' anchor slugs match the document's live H1/H2 set in order.

Paired § 4.1.6 mutation tests (planned):

- Strip the | Revision | row from one *.md → CM-MD-METADATA-PRESENT FAILS.
- Rename one H1 without updating ToC → CM-MD-TOC-PARITY FAILS.

Composition with § 4.1.7. § 4.1.8 remains the authoritative clause for the *fact* that documents must carry revision metadata; § 4.1.9 changes the *format* (bold-lines → table) and *scope* (docs/**/*.*md → every

tracked *.md per § 1.4.4 INCLUDED set). Consumer projects with existing gates that grep for **Revision:** MUST update those gates to also accept the canonical table form. Migration period: 11.4.45 11.4.59 days from § 1.4.48 11.4.59 landing.

Composition with § 1.4.63 / § 1.4.61 / § 1.4.61. The per-class always-sync clauses (Status.md § 1.4.61, README § 1.4.61, composite § 1.4.61) and the universal export covenant § 1.4.61 govern artefact existence and freshness; § 1.4.61 governs content discipline inside the source .md. The clauses are orthogonal layers of the same anti-bluff posture.

No escape hatch. No --skip-md-metadata, --no-metadata-table, --toc-stale-OK, --allow-missing-revision flag exists. Divergent metadata and stale ToCs are § 1.4.17 PASS-bluff patterns at the doc-surface layer.

Classification: universal (per § 1.4.61). Applies to every project consuming this Constitution. The canonical table format and ToC mandate are universal; the per-project tracker IDs (e.g. HRD-) and specific Issues/Fixed-class doc paths are consumer-side per-project.

Canonical authority: constitution submodule Constitution.md § 1.4.61.

Non-compliance is a release blocker regardless of context.

§ 1.4.61 – Workable-items procedure docs as single source of truth (User mandate, – –)

Forensic anchor – verbatim user mandate (– – ~ : Z):

"To make workable items tracking and creation through all of the mentioned mandatory documents we MUST create proper procedure document which will be named in Constitution, AGENTS.md and CLAUDE.md as single source of truth for this procedure! Make sure it is created under: docs/procedures/issues/Creation.md which will be always exported to PDF and HTML whenever it is updated. We MUST create procedure documents for Updating.md, Resolution.md, Reopening.md and other workable items actions that we have. ... Everything done on project - new features, changes, tasks, bug fixes, investigation, ordinary tasks (writing documentation for example) MUST ALL go through the workable items flow and proper procedures!"

Every workable-item action – opening, updating, closing, reopening, migrating across Issues.md Fixed.md – MUST follow the canonical procedure document at docs/procedures/issues/<Action>.md . The five procedure docs are MANDATORY references; no agent or operator may invent ad-hoc procedures. The closed-set:

Procedure doc	Covers
Cre- ation.md	Opening a new workable item (Bug / Feature / Task)
Updat- ing.md	Non-closure, non-reopen edits – status transitions, evidence append, type re-classification
Resolu- tion.md	Atomic Issues.md → Fixed.md close (Bug→Fixed, Feature→Implemented, Task→Completed per \$. .)

Procedure doc	Covers
Reopen-ing.md	Atomic Fixed.md → Issues.md reopen with \$. . source attribution + \$. . reopens-history
Migra-tion.md	Bidirectional atomic-move mechanics + sync_is-sues_docs.sh internals (consumed by Resolution + Reopening)

Each procedure doc carries the \$. . revision header + HTML + PDF exports synchronised to its .md source. New helper scripts/testing/sync_procedure_docs_export.sh walks the procedure- docs directory and emits matching .html + .pdf for each .md ; invoked as the final stage of scripts/testing/sync_issues_docs.sh so a single operator invocation refreshes every doc-side artifact.

Universality of scope. All work – new features, behavioural changes, tasks (refactor / doc / infra / gate / audit / cleanup), bug fixes, investigation, documentation work – MUST flow through these procedures. The "I'll just tweak this one file" escape hatch is forbidden because it produces silent doc drift and reopens the \$. covenant's PASS-bluff failure mode at the process-tracking layer.

Composes with \$. . (no-guessing – every procedure step is mechanical, no likely / seems), \$. . (demotion-evidence – Reopening.md is the canonical implementation), \$. . (file-layout – procedure docs live in docs/procedures/issues/ , qa-results forensics live in qa-results/), \$. . (Issues_Summary sync – all five procedures invoke sync_issues_docs.sh whose pipeline is documented in Migration.md), \$. . (Status closed-set – all five procedures cite the same

closed-set), \$ (Type closed-set),
 \$ (atomic Issues Fixed move – Migration.md
 is the implementation), \$ (visual cue +
 grouping colorization – invoked by sync wrapper, no
 procedure-side bypass), \$ (type- aware
 closure vocabulary – Resolution.md enforces),
 \$ (Reopened-source attribution – Reopening.md
 enforces), \$ (revision header on every
 procedure doc), \$ (Fixed_Summary parity –
 Migration.md documents the pipeline), \$ (ATM-
 NNN identifier – Creation.md allocates), \$
 (Reopens-history per- item doc – Reopening.md authors),
 \$ (documentation always- sync composite –
 procedure docs are an additional doc class binding to the
 composite covenant).

Pre-build gate `CM-PROCEDURES-DOCS-PRESENT` checks () all
 procedure docs exist at `docs/procedures/issues/`
`{Creation,Updating, Resolution,Reopening,Migration}.md` , ()
 each carries the \$ revision header, () each
 has matching `.html` + `.pdf` exports with `mtime` $\geq$ the `.md`
`mtime`, () `scripts/testing/sync_procedure_docs_ export.sh`
 exists + is executable, () `scripts/testing/sync_issues_`
`docs.sh` invokes `sync_procedure_docs_export.sh` . Paired
 mutation: rename one procedure doc $\rightarrow$ gate FAILs.

Propagation gate `CM-COVENANT-114-63-PROPAGATION` enforces
 this anchor literal in every `CLAUDE.md` / `AGENTS.md` across
 parent + owned submodules + nested submodules +
 HelixQA dependencies. Paired mutation strips the anchor
 literal $\rightarrow$ gate FAILs.

No escape hatch – no `--ad-hoc-procedure` ,
`--skip-procedure-doc` , `--procedure-not-applicable` flag exists.
 Inventing an alternate procedure for a "small" change is
 the exact failure mode this anchor closes.

Classification: universal (per § 1.1.4.63). The closed-set of five procedure docs and the sync-wrapper composition are universal; the specific test-name / fix-name examples within each procedure doc may be consumer-side per-project.

Canonical authority: constitution submodule Constitution.md
§ 1.1.4.63

2026 05 19
Non-compliance is a release blocker regardless of context.

2026 05 19
§ 1.1.4.63 – Universal Markdown export
mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"Any markdown document inside the project and which is not part of the applications or services source code MUST BE exported (be available) in PDF and HTML! Any already existing Markdown document that fulfills this condition and which does not have HTML or PDF at all or it is not in sync with it MUST HAVE (re)generated PDF and HTML version! Every time when Markdown document (file) is modified, its proper HTML and PDF versions MUST BE regenerated. Markdown documents MUST BE at all times in sync with PDF and HTML versions!"
1.1.4.44 1.1.4.45 1.1.4.53
1.1.4.56 1.1.4.57 1.1.4.59

The covenant generalizes the per-class always-sync discipline of § 1.1.4.63 / § 1.1.4.63 /
§ 1.1.4.63 / § 1.1.4.63 / § 1.1.4.63 /
§ 1.1.4.63 / § 1.1.4.63 / § 1.1.4.63 /
§ 1.1.4.63 / § 1.1.4.63 / § 1.1.4.63 to every
Markdown document in the project that is not part of an application or service's source-code tree. Per-class anchors govern specific files (Issues, Fixed, Status,

README, etc.); § 4.1.1 is the catch-all: any other `.md` file that documents the project – guides, research notes, plans, hardware-ID notes, script-companion docs, changelogs, procedure docs, README files inside owned submodules – MUST also have `.html` + `.pdf` siblings, all three artefacts in sync at all times. Stale exports of any documentation surface are the exact "operator reads divergent HTML/PDF while `.md` is correct" PASS-bluff § 4.1.2 forbids – only the per-class enforcement so far prevented it for specifically-named docs; § 4.1.3 closes the long tail.

Scope (closed-set):

INCLUDED:

- Project root `*.md` (README.md, CLAUDE.md, AGENTS.md, CONTRIBUTING.md, etc.)
- `docs/**/*.*md` (guides, research, plans, changelogs, procedures, hardware notes)
- `scripts/**/*.*md` (script-companion docs in documentation format, NOT shebang scripts)
- Owned-submodule trees at `device/rockchip/atmosphere/<submodule>/` : top-level README.md / CLAUDE.md / AGENTS.md / CHANGELOG.md and any `docs/**/*.*md`
- `constitution/**/*.*md` (the canonical-root submodule)
- `tools/helixqa/<helixqa-submodule>/` top-level README.md / CLAUDE.md / AGENTS.md and any `docs/**/*.*md` (owned HelixQA dependencies)

EXCLUDED (NOT subject to § 4.1.1 – these are application/service source code or third-party trees we do not own):

- `external/**` (AOSP-mainline / upstream-mirror source trees)
- `prebuilts/**` (binary prebuilts + their bundled docs)

- `packages/modules/**` (AOSP mainline modules)
- `kernel-5.10/**` (kernel source tree – has its own upstream docs)
- `out/**` (build output)
- `build/**` (AOSP build system)
- Application / service source-code trees (e.g. Java/Kotlin module-internal `*.md` shipped with library code, README files inside source-only third-party gradle module directories)
- Any third-party submodule NOT in the owned-submodule set (presenter, vlc-player, nova-player, mpv-player, gramophone-player, rhythm-player, strep-player, smarttube-player, torrservice, lampa) or HelixQA-owned set (Challenges, Containers, DocProcessor, LLMOrchestrator, LLMPProvider, VisionEngine).

Mandatory protections (ALL must hold):

- **Every INCLUDED `.md` file has `.html` and `.pdf` siblings.**
A missing export is a `$ . .` violation regardless of when the markdown was last touched.
- **`.html` and `.pdf` mtime $\geq$ `.md` mtime** (within the same sync-wrapper invocation granularity). Stale exports are violations even if the `.md` itself is correct.
- **Every modification triggers regeneration.** Whether via direct sync-wrapper invocation (`..sync_all_markdown_exports.sh`), a git pre-commit hook auto-regenerating on staged `.md` changes, or the existing per-class wrappers (`$ . .` `sync_issues_docs.sh`, `$ . .` `script-docs sync`, `$ . .` `sync_readme_export.sh`, etc.) all of which delegate the universal export path to the same canonical helper.

- **Pre-build gate enforces parity.** `CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC` walks the INCLUDED scope, verifies every `.md` has `.html`
 - `.pdf` siblings with `mtime ≥ .md mtime`, and FAILs the build if any are missing or stale.
- **No escape hatch.** No `--skip-md-exports`, `--no-pdf-only`, `--md-export-not-applicable`, `--application-internal-doc` flag exists for files inside the INCLUDED scope. The EXCLUDED scope is the only legitimate path to opt out, and it is closed-set.

Canonical helper.

`scripts/testing/sync_all_markdown_exports.sh` walks the INCLUDED scope, invokes pandoc (HTML) + weasyprint (PDF) with `timeout 60` each (graceful per-file degradation), uses `docs/_progress-style.css` for visual consistency, and supports `--check-only` (exit nonzero if any out-of-sync, print list) and `--regenerate-all` (force) modes. Caps at candidates with explicit abort+list if the scope is unexpectedly large (e.g. after a new submodule lands without scope-update review). Idempotent.

Composition. Composes with `$ . .` (Issues_Summary sync – universal helper is the back-end), `$ . .` (script-companion docs), `$ . .` (HTML colorizer continues to post-process the always-sync docs for visual cues), `$ . .` (revision header on every .md the universal helper exports), `$ . .` (Status.md auto-sync), `$ . .` (Fixed_Summary), `$ . .` (README always-sync – the universal helper covers the long tail of every *other* root-level .md too), `$ . .` (composite always-sync covenant – `$ . .` is the structural generalization), `$ . .` (procedure docs – already always-sync; `$ . .` adds catch-all for every doc class not yet anchored), `$ . .` (topic discoverability – summary docs the universal helper

exports continue to be post-processed by inject_topic_index.py).

Pre-build gates:

- CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC – invariants: (a) scripts/testing/sync_all_markdown_exports.sh exists + executable, (b) running it in --check-only mode returns (no out-of-sync file in the INCLUDED scope).
- CM-COVENANT-114-65-PROPAGATION – anchor literal 11.4.65 in every CLAUDE.md / AGENTS.md across canonical-root + parent + owned submodules + nested submodules + HelixQA dependencies (-file fleet – same propagation set as \$. .).

Paired meta-test mutations: one strips the CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC gate's enforcement literal; one strips the \$. . anchor literal from a sentinel propagation file. Both mutations assert the corresponding gate FAILs when applied.

Classification: universal (per \$. .). The mandate to export every project-doc Markdown to HTML+PDF is universal across every project that consumes this Constitution; the specific INCLUDED / EXCLUDED scope path lists are consumer-side per-project (since each consumer has different submodule topology, build trees, and source-code roots).

Canonical authority: constitution submodule Constitution.md \$. . .

Non-compliance is a release blocker regardless of context.

§ . . – Blocker-resolution interactive-clarification mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim,
– –):

"If any blockers which can be resolved with interactive response ever happen again, perform in depth research on options doable by your side and how much inputs from us you really need, then create options and present them to us. After we answer, preferrably you will be unblocked and be able to continue work on blocked items. Let us make this main approach when such situations (blockers) do happen!"

When any task is blocked (waiting on operator decision, hardware access, external-service authorization, ambiguous scope, or any input the agent cannot mechanically derive), the agent MUST follow this five-step discipline before idling or asking free-form questions:

- . **In-depth research on agent-side options.** Identify the maximum scope that can land unilaterally without operator input – research the codebase, upstream documentation, existing tooling, the current state, what files/components are involved, what tests/gates already exist. Surface every agent-actionable path even if it requires a small operator confirmation.
- . **Calculate minimum-viable operator input.** Reduce the question set to the smallest closed set of operator-only decisions (preference, authorization, scope-bounding, hardware availability, schedule). Anything that can be researched MUST NOT be asked.

- **Construct** – **mutually-exclusive options**. Each option carries (a) a short label, (b) a one-line trade-off description, (c) an explicit statement of what the agent will do *after that answer*. One option marked "Recommended" with rationale. Options MUST genuinely differ – restating one option as three close variants is itself a § 11.4.40 violation.
- **Present via the platform's interactive question mechanism**. On Claude Code that is `AskUserQuestion` (max 10 questions, each with 3–5 options; supports multi-select for non-mutually-exclusive sets). Other platforms (Copilot CLI, Codex, Gemini CLI) use their equivalents. NEVER inline free-text "what would you like?" when interactive options would do – that wastes operator attention.
- **After the operator answers, resume work without additional round-trips**. Every option's promised action MUST be sufficient to unblock – if a follow-up clarifying question is needed, the prior options were insufficiently researched and that's itself a § 11.4.40 violation. The contract is: ask once, unblock, continue.

Composes with:

- § 11.4.40 (no-guessing – interactive options replace agent guessing about operator preference)
- § 11.4.41 (demotion-evidence – operator preference is captured- evidence for direction changes)
- § 11.4.42 (full-suite retest – uninterrupted work post-answer means the test cycle resumes cleanly without context drift)
- § 11.4.43 (merge-first – when operators make conflicting choices across sessions, the merge-first discipline preserves both)

- § . . (iteration-discipline – interactive-question loop is built into the per-cycle methodology)
- § . . (autonomous-validation – most blockers can be researched to autonomous-path level, leaving only true operator preference to interactive-ask)

Mandatory protections:

- **No silent waiting.** If blocked, ask. If genuinely no operator input is needed, do not ask – proceed.
- **No bulk-text questions when interactive options would do.** Free-text prompts are reserved for truly open-ended ideation, never for closed-set decisions.
- **Each interactive question minimizes operator cost** – the operator should be able to answer all questions in under seconds total.
- **Recommended option is highlighted explicitly** – the operator should know what the agent would default to if no answer comes.
- **Out-of-band channels (slack, voice, email) NEVER substitute** – the interactive mechanism creates an audit trail that free-text channels do not.

Pre-build gate `CM-COVENANT-114-66-PROPAGATION` enforces the anchor literal is present across the -file consumer fleet (parent CLAUDE.md / AGENTS.md + owned-submodule pairs + HelixQA-owned pairs). Paired meta-test mutation strips the literal → gate FAILs. No escape hatch – no `--skip-ask`, `--silent-wait`, `--free-form-only` flag is permitted.

Canonical authority: constitution submodule `Constitution.md`
 § . . .

Non-compliance is a release blocker regardless of context.

```
$ . . - Shell-script target-shell-
2026 05 19 parseability mandate (User mandate,
- - )
```

```
Forensic anchor - direct user mandate (verbatim,
- - ):
```

"any issue we spot must be fixed, bash scripts as well if they are broken!" "Make sure that this is mandatory rule!"

Forensic incident – Phase .FJ. D post-flash cycle block: `device/rockchip/rk3588/tests/test_all_fixes.sh:114` invoked on Android via `sh script.sh` used bash-only process substitution `exec > >(tee -a "$file") 2>&1`. Android's `/system/bin/sh` is mksh, which parses the *entire* script BEFORE executing it – so the runtime guard `if [ -n "${BASH_VERSION:-}" ]` could not save the script: mksh rejected `>(…)` at parse time and the test cycle failed to launch on D. The runtime guard correctly tried to gate the unsafe path, but mksh's compile-time parser made the guard useless. Fixed by wrapping the offending `exec` in `eval` so the parser sees only a string and the parse step defers to runtime where the guard can take effect.

This is the canonical class of script-bug FAIL-bluff: a test cycle *PASSes* on the developer host (bash) and *FAILS* at parse time on the target (mksh), producing a FAIL-bluff that masks every product defect the cycle would have caught. `$ . .` already forbids script-bug FAILs. `$ . .` mechanically prevents this specific subclass.

The mandate. Every shell script that may be invoked under a target shell other than the one in its shebang MUST parse cleanly under that target shell. Specifically:

- **Closed-set scope.** Every tracked `.sh` file under `device/rockchip/rk3588/tests/`, `scripts/`, and `scripts/testing/` (and equivalent paths in owned submodules) is in scope. Explicitly OUT of scope: `external/`, `prebuilts/`, `packages/modules/`, `kernel-5.10/`, `out/`, `build/`, `scripts/legacy/` (legacy quarantine), and any third-party vendored shell directory the project does not maintain.
- **Mandatory parseability invariant.** Every in-scope script MUST parse under POSIX `sh -n` on the developer host. POSIX `sh` is the closest portable parser to Android `mksh` – passing `sh -n` on the host gives high confidence the script parses on `mksh`.
- **Bash-only syntax requires deferral.** Process substitution `>(...)` / `<(...)`, `[[ ]]`, `<<<` here-strings, indexed arrays `arr=()`, associative arrays `declare -A`, `${var^^}` / `${var,,}` case-fold expansions, `${var/pattern/replacement}` with bash-only operators, `coproc`, `function name { ... }` (vs POSIX `name() { ... }`), and any other bash-only construct MUST EITHER be wrapped in `eval '...bash-only string...'` so the parser sees only a string OR be guarded by sourcing the file only on bash-detected hosts (`[ -n "${BASH_VERSION:-}" ]` BEFORE the script loads – runtime guards inside a single `mksh`-parsed script do not work).
- **Honest shebangs.** A script that contains bash-only constructs without `eval -deferral` MUST declare `#!/bin/bash` (or equivalent) AND its callers MUST invoke it as `bash script.sh` rather than `sh script.sh`. A script with `#!/system/bin/sh` or `#!/bin/sh` shebang MUST be POSIX-clean – no bash-only syntax, no `local` keyword without

an mksh fallback, no `read -p`, no `echo -e` without portable equivalent.

- **eval-deferral is the safe pattern.** When a bash-only construct is genuinely needed inside a script that may be parsed by mksh (e.g. `sh script.sh` from a callsite outside our control), wrap the construct in single-quoted `eval : eval 'exec >>(tee -a "$f") 2>&1'`. The parser sees a string; the runtime guard around the `eval` decides whether to execute.
- **Sourced-function `exec` redirection-scoping (extension, research-derived, - -).** In bash-family shells, `exec` WITHOUT a command applies EVERY redirection on the line to the CURRENT SHELL, PERMANENTLY. Inside a standalone script that is a self-contained choice; inside a SOURCED / library function – code that runs in the operator's interactive shell or a caller's long-lived shell – it is a host-state mutation. Forensic FACT (genericised, - -): a library lock helper opening its lock fd with `exec 9>>"$lock" 2>/dev/null` permanently redirected the INTERACTIVE shell's stderr to /dev/null – one provider launch left the operator's terminal unable to print any error, from any command, for the life of the shell – while the paired unlock helper had used the correct brace form (`{ exec 9>&-; } 2>/dev/null`) all along; the asymmetry was the defect. Marker-probe proof: pre-fix, a marker echoed to stderr after the call never appeared; post-fix it does. The mandate: in ANY sourced/library shell function, an `exec` whose purpose is only to open/close/move file descriptors MUST brace-scope its side redirections – `{ exec 9>>"$file"; } 2>/dev/null` – so error-suppression applies to the fd-operation alone, never to the enclosing shell; a command-less `exec` combining fd-manipulation operands with an additional unbraced redirection in sourced code is a REVIEW DEFECT

(`$ . .`) and a greppable pre-build-gate class. Recommended gate `CM-SHELL-EXEC-REDIRECTION-SCOPED` (scan in-scope sourced/library shell sources for command-less `exec` lines that combine `fd-open/close` operands with additional unbraced redirections → FAIL citing `file:line`; brace-scoped forms and `exec WITH` a command are exempt – the `$ . .` () false-positive guard) + paired `$ .` mutation (unwrap one brace-scoped `fd-open` back to the bare form → the gate FAILs; gate-code = separate work item). The instrument-side twin of this construct – the SAME `exec ... 2>/dev/null` blinding `bash -x` tracing of the code under test – is catalogued in `docs/guides/shell_instrument_footguns.md` per `$ . .` ().

Captured-evidence enforcement. Pre-build gate `CM-SCRIPT-TARGET-SHELL-PARSEABLE` walks every `.sh` file under the in-scope directories, runs `sh -n` on each with a -second per-file timeout, and FAILs if any script fails to parse. The gate's diagnostic output cites the failing file + the `sh -n` error message so the fix location is unambiguous.

Propagation gate. `CM-COVENANT-114-67-PROPAGATION` verifies the `$ . .` anchor literal is present across the -file consumer fleet (parent `CLAUDE.md` / `AGENTS.md` + Containers + owned atmosphere submodules + nested SmartTube SharedModules / MediaServiceCore / MSC-SharedModules + HelixQA submodules). Constitution submodule files (`constitution/` {`Constitution`,`CLAUDE`,`AGENTS`}.md) are canonical authority – they carry `$ . .` by definition.

Paired meta-test mutations. (a) Inject a bash-only construct (`exec > >(/bin/cat)`) outside `eval` into a fresh location inside a test script and assert `CM-SCRIPT-TARGET-SHELL-PARSEABLE` FAILs. (b) Strip `11.4.67` literal from one

consumer file and assert `CM-COVENANT-114-67-PROPAGATION FAILS`.

Composes with:

- `$ . .` (FAIL-bluffs equally forbidden – script-bug FAILs that mask product defects are the exact failure mode `$ . .` mechanically closes)
- `$ . .` (test-interrupt-on-discovery – a parse FAIL at cycle start IS the discovery event; cycle stops, fix lands, retest)
- `$ . .` (no-guessing – `sh -n` produces fact, not opinion)
- `$ . .` (deterministic consistency – a script that parses on one shell and not another fails the N-iteration identical-result requirement)
- `$ . .` (live-ADB-first – shell-script fixes are `LIVE_ADB_TESTABLE`: push the fix, re-run on device, capture evidence before commit)

Mandatory protections (no escape hatch):

- **No `--skip-parseability-check` flag.** A script that does not parse under the target shell is broken regardless of intent.
- **No "this script only runs on bash" exception.** If a script may be invoked via `sh script.sh` (and almost all of them can – Android's default `sh` is mksh), it **MUST** parse under mksh.
- **No "the runtime guard catches it" exception.** Runtime guards inside a mksh-parsed script cannot prevent mksh parse-time rejection. Use `eval -deferral` or shebang-controlled invocation.
- **Fix at source, not in callsites.** A broken script is fixed in the script itself (per `$ . .`), not by

individual callsites defensively detecting parse failures.

Pre-build gate `CM-SCRIPT-TARGET-SHELL-PARSEABLE` + paired mutation. Propagation gate `CM-COVENANT-114-67-PROPAGATION` + paired mutation.

Canonical authority: constitution submodule `Constitution.md`
§ 1.1.4.1.1.

Non-compliance is a release blocker regardless of context.
2026 05 20

§ 2.0.6.5.20 – Positive sink-side / downstream evidence mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"We still do not hear any audio played from D device! Arvus Web Dashboard when we play music from D shows nothing for Codec In Use! This MUST BE investigated and fixed! How come we passed the tests with Arvus validation? What were values for the Codec In Use field? Empty means nothing! This is not working! It MUST BE FIXED, TESTED AND VERIFIED WITH FULL AUTOMATION TESTING ASAP!!!"

§ 2.0.6.5.20 already mandates consuming the Arvus sink-side report when available. § 2.0.6.5.20 extends this with a stricter contract: a test that asserts audio or video routing PASS MUST capture and verify positive sink-side or downstream evidence – never config-only, never metadata-only, never PCM-open-state-only. The

§ 2.0.6.5.20 path described what SHOULD be captured;

§ 2.0.6.5.20 fixes the failure mode where tests "PASSED"

because the helpers silently SKIPPed the very evidence the assertion depended on.

Positive sink-side / downstream evidence – closed

enumeration (at least one MUST be captured for every audio/video routing PASS; none alone is sufficient if the canonical sink/downstream is reachable):

- . **Sink-side codec-state** (Arvus / Sonos / Yamaha / Denon / Marantz / WiSA REST API) showing a **non-empty Codec-In-Use** value matching the test's expected codec regex during the playback window. Empty / `<unreachable>` / `<N.E.>` / `<None>` placeholders are NOT positive evidence.
- . **PCM frames-written delta** from `/proc/asound/cardN/pcmMp/sub0/ status hw_ptr` – the delta over the playback window MUST be strictly positive AND consistent with the expected sample-rate × channel-count × duration. Zero or negative delta is FAIL.
- . **ALSA ELD / EDID-Like-Data** from `/proc/asound/cardN/eld#X.Y` demonstrating the sink negotiated the expected channel count + format. Empty ELD is NOT evidence the sink received audio – it only proves the cable / EDID-channel is up.
- . **ffprobe-on-captured-mp** for video routing – non-zero frame count, expected codec, expected resolution, expected fps. `-byte capture` (Bug pattern) is FAIL.
- . **Recording-analyzer event match** per § . . / § . . timeline – at least one matched event in the analyzer findings.
- . **Tinycap RMS amplitude** for ALSA-loopback audio recordings – non-trivial RMS in the captured WAV ($\geq$ - dBFS for line-level playback).

Failure-mode taxonomy – what § . . specifically forbids:

- `arvus_probe_present` returns → test reports SKIP → suite reports all-green. **FORBIDDEN**. When the sink-side evidence is REQUIRED by the assertion, missing sink is `OPERATOR-BLOCKED` (an explicit release-blocker, NOT a SKIP, NOT a PASS).
- `arvus_probe_codec_state` returns empty string → test logs "codec_state: " (blank) → assertion proceeds on HAL-side metadata. **FORBIDDEN**. Empty codec-state IS a defect signal – either the audio is not arriving at the sink (genuine product defect per current User mandate) OR the sink is unreachable (operator-blocked). Either way, the test cannot PASS.
- Test reads HAL `dumpsys media.audio_flinger` output-thread device field showing `0x400` (`AUDIO_DEVICE_OUT_HDMI`) → reports "HDMI routing PASS". **FORBIDDEN**. That field reflects what the audio POLICY chose, NOT what the audio HAL physically opened. The HDMI ALSA card may still be closed (PCM open failure swallowed by the HAL's `ALOGW` fallback) – silent silence at the user's ear. Sink-side codec-state or PCM `hw_ptr` delta is the ONLY proof audio actually arrived.
- Test asserts `/vendor/etc/audio_policy_configuration.xml` contains the `deep_buffer` route to `HDMI Out`. **FORBIDDEN as the only evidence**. Config-only PASS is the canonical § . PASS-bluff – config can be perfect AND audio still inaudible (e.g., the HAL may still misroute). MUST be paired with at least one runtime positive evidence from the enumeration above.

Mandatory protections (all four):

- **Library contract.** Sink-side helper libraries (`arvus_probe.sh`, `sonos_probe.sh`, etc.) MUST expose a `*_require_reachable` function that returns exit code

(OPERATOR-BLOCKED) when the sink cannot be probed. Tests asserting sink-side codec-state MUST call this function at probe entry. Silent skip via `*_probe_present` is permitted ONLY for topology-dispatch fan-out logic that already has an alternative on-SoC + downstream-amplitude probe path.

- . **Exit-code propagation.** `arvus_require_reachable` / equivalent helpers return `OPERATOR-BLOCKED`; the test harness MUST propagate `OPERATOR-BLOCKED` to the suite summary (counted separately from FAIL/SKIP) as `OPERATOR-BLOCKED`. `test_all_fixes.sh` orchestrator MUST exit non-zero when any test reported `OPERATOR-BLOCKED` AND release tags MUST NOT be cut while any `OPERATOR-BLOCKED` test exists.
- . **Test rewrite.** Every test currently asserting `audio output routing PASS` via metadata-only checks (e.g., the historical `test_audio_output_routing.sh` config-only invariants) MUST be rewritten to capture at least one positive sink/downstream evidence per the enumeration above. The `$ANCHOR . PASS-bluff pattern` `tinymix says route is on → PASS` is the failure mode this anchor closes.
- . **Anti-stickiness post-stop.** After the test stops the playback, it MUST re-probe the sink/downstream and assert the codec-state transitioned to N.E. / N/A / `-frames-delta`. A persistent codec-state across test boundaries indicates orphan playback (`$ANCHOR . .` violation) AND would falsely PASS the next test.

Pre-build gate `CM-COVENANT-114-68-PROPAGATION` enforces this anchor literal across canonical files + per-consumer propagation. Paired mutation strips the anchor literal → gate FAILs. Pre-build gate

`CM-POSITIVE-SINK-EVIDENCE-PER-AUDIO-TEST` walks every audio-routing on-device test asserting PASS and verifies it cites at least one positive evidence call (one of `arvus_require_reachable` / `arvus_assert_codec_format` /

`pcm_hw_ptr_delta_positive` / `tinycap_rms_above_floor` / `ffprobe_frames_nonzero`). Paired mutation strips one such call → gate FAILs.

Composes with § 11.4.1 (recorded-evidence – sink-side is the captured-evidence channel), § 11.4.2 (audio + video quality analysis comprehensiveness – codec / channel-count / RMS), § 11.4.3 (sink-side captured-evidence mandate – § 11.4.4 closes its silent-skip gap), § 11.4.5 (test cleanup – anti-stickiness post-stop), § 11.4.6 (recent-work validation – sink-side evidence required at validation entry), § 11.4.7 (dual-approach testing – both UI + Intent variants need sink-side evidence), § 11.4.8 (deterministic consistency – sink-side evidence must be consistent across N iterations), § 11.4.9 (autonomous-validation – sink probe IS the canonical autonomous path for HDMI audio).

No escape hatch – no `--skip-sink-evidence` , `--allow-empty-codec` , `--sink-unreachable-is-pass` , `--metadata-only-suffices` flag exists anywhere. The discipline exists because the operator/end-user forensic on – – demonstrated the exact failure this anchor forbids: tests reported audio-routing PASS while the user heard nothing and Arvus Codec-In-Use was empty.

Canonical authority: constitution submodule `Constitution.md`
§ 11.4.10

Non-compliance is a release blocker regardless of context.

§ . . – Universal Sink-Side Positive-Evidence Taxonomy + Mechanical Enforcement (User mandate, – –)

Forensic anchor – direct user mandate (verbatim,
– –):

"THIS MUST HAPPEN NEVER AGAIN!!! We MUST HAVE this all working! Not just for audio but for every single piece of the System!!! Proper full automation when executed with success MUST MEAN that manual testing will be as much positive at least regarding the success results! Dive deep into the root causes and how this had happened! The root causes of such omission MUST BE TRACKED, and proper solution applied! Solution MUST BE universal, generic that solves working flows for all System components and for all future and all existing projects! Make sure everything is added we change as mandatory rules we must follow and critical constraints into ours root (constitution Submodule) Constitution.md, CLAUDE.md and AGENTS.md!!! Everything we do MUST BE validated and verified with rock-solid proofs and anti-bluff policy enforcement and fulfillment! THIS IS MOST CRITICAL POINT WE HAVE NOW with all audio issues!"

Escalation from the same operator session
(– – , hours earlier):

"We still do not hear any audio played from D device! Arvus Web Dashboard when we play music from D shows nothing for Codec In Use! ... How come we passed the tests with Arvus validation? What were

values for the Codec In Use field? Empty means nothing! This is not working!"

Forensic incident – – – → D audio-routing PASS-bluff generalised. § . . closed the audio-specific failure path where Arvus reported an empty Codec-In-Use field and the test PASSEd anyway. The same shape of bug had previously shown up intermittently across multiple subsystems (display routing on D secondary HDMI, BT A DP codec advertisements, WiFi connection quality, video playback frame-count, touch-input event delivery). The audio incident was the canonical trigger because the operator hears silence at the soundbar even while the cycle reports green. § . . is the universal generalisation – closing the bluff class for every user-visible feature, not only audio.

Root-cause analysis – why existing anchors § . . / § . . / § . . / § . . did NOT mechanically catch the failure across the System:

- **Aspirational text without a pre-build gate.** § . . (recorded-evidence) required captured visual/audio evidence per PASS but had no gate walking the test corpus to enforce a per-PASS evidence file. Tests called `ab_pass "description"` and the evidence requirement was honour-system.
- **SKIP-fail-open hiding missing evidence.** § . . (Arvus sink-side) had a silent-skip escape when the sink returned an empty / unreachable response – the empty response was the very defect signal but it was classified as "sink unreachable, SKIP the assertion." § . . closed this for audio; § . . closes it universally across every feature class with a downstream sink.

- **PASS-by-default vs FAIL-by-default helper contract.** The bare `ab_pass` helper accepted a free-text description and incremented the PASS counter – no evidence path required, no shape enforced. The default failure mode was therefore PASS- bluff. The correct default is FAIL-by-default – a PASS must prove itself with a captured-evidence artefact path.
- **Missing closed-set taxonomy.** Individual test authors made one-off decisions about what evidence shape was acceptable. Without a single canonical table mapping every user-visible feature class to its required sink-side probe, aggregate enforcement across the System was mechanically impossible.
- **Helper functions that mapped a multi-step probe to one bool.** `arvus_probe_present` collapsed reachable-but-empty into the same value as unreachable. The § 11.4.6.8 fix added `arvus_require_reachable` for audio; § 11.4.12 codifies the equivalent for every sink (Sonos, Yamaha, Denon, video display analyzer, WiFi sink, BT peer, etc.).
- **No per-feature evidence ledger.** § 11.4.19 supplied an autonomous-validation classification (`AUTONOMOUS_VERIFIED` etc.) but did not require each test to cite a captured-evidence artefact path matching the § 11.4.19 taxonomy.

The § 11.4.19 mandate – universal, generic, applies to every System component, every owned project, every consumer subscribing via the constitution submodule:

Element 11.4.19 : Closed-set sink-side / downstream evidence taxonomy. Every user-visible feature class in every consuming project MUST map to exactly one entry in the following table. The taxonomy is the canonical authority on the required evidence shape per feature class. Tests

annotate themselves with `# §11.4.69 FEATURE: <class>` so the pre-build gate can match expected evidence shape.

Feature class	Required probe	Required evidence shape
<code>audio_output</code>	sink-side codec/ channel REST probe (Arvus / Sonos / Yamaha / Denon / Marantz / WiSA) + on-device <code>tinycap</code> capture + <code>ffprobe</code> channel assertion	non-empty <code>co-</code> <code>dec_in_use</code> matching expected codec AND captured WAV with RMS > - dBFS AND <code>ffprobe</code> <code>channels</code> matches claim
<code>audio_input</code>	<code>tinycap</code> capture from intended device + <code>ffprobe</code> RMS	captured WAV with RMS > - dBFS AND <code>au-</code> <code>dio_devices_for_attr</code> shows the claimed source device
<code>video_display</code>	<code>screenrecord</code> on in- tended display + <code>ff-</code> <code>probe -count_frames</code> + analyzer event-match	non-zero mp size AND <code>nb_read_frames</code> > 0 AND analyzer re- ports a matched event for the claimed display
<code>net-</code> <code>work_through-</code> <code>put</code>	<code>iperf3</code> / <code>curl --</code> <code>speed</code> / <code>dd</code> over network	throughput ≥ per- link-type floor AND TCP buffer / con- gestion-control matches claim
<code>net-</code> <code>work_connectiv-</code> <code>ity</code>	ping + DNS resolve + TCP connect probe	captured ping reply AND DNS answer AND TCP handshake to known port

Feature class	Required probe	Required evidence shape
bluetooth_a2dp	codec query + sink peer state + A2DP TX queue depth	dumpsys bluetooth_manager reports CONNECTED+A2DP_PLAYING AND codec field non-empty AND queue depth < overflow threshold
bluetooth_pair	dumpsys bluetooth_manager + BluetoothDevice.getBondState()	bond state = BOND_BONDED AND device MAC matches expected
touch_input	getevent -lt capture during scripted input tap	≥ EV_ABS ABS_MT_* event per scripted tap AND dumpsys input shows the touch device active
sensor	dumpsys sensorservice + getevent on sensor device	event count > between BEFORE and AFTER snapshots for the claimed sensor
gpu_render	SurfaceFlinger frame-rate dump + dumpsys gfxinfo <pkg>	Janky frames < threshold AND Total frames rendered > 0 during test window
storage_read	dd if=<path> of=/dev/null + /proc/diskstats delta	throughput ≥ floor AND read_sectors delta confirms IO landed on intended device
storage_write		

Feature class	Required probe	Required evidence shape
	<code>dd of=<path> + sync</code> + <code>/proc/diskstats</code> <code>delta</code>	<code>throughput ≥ floor</code> AND <code>write_sectors</code> <code>delta</code> confirms IO landed on intended device
<code>mediaco-</code> <code>dec_decode</code>	<code>media.metrics dump +</code> <code>dumpsys media.codec</code>	decoder lifecycle events present (<code>configure</code> , <code>start</code> , ≥ <code>output buffer</code>) AND codec name matches claim
<code>mediaco-</code> <code>dec_encode</code>	<code>media.metrics dump +</code> non-zero output file	encoder lifecycle events present AND output file non- zero AND <code>ffprobe</code> confirms codec matches claim
<code>miracast /</code> <code>cast</code>	sink-side WFD / Cast announce + session state	sink dashboard re- ports active ses- sion AND session ID present in <code>dumpsys</code> <code>media_session</code>
<code>boot_service</code>	<code>getprop</code> <code>init.svc.<name> + log</code> <code>capture</code>	service state = <code>running</code> AND log shows expected boot-time output literal
<code>pack-</code> <code>age_install</code>	<code>pm path <pkg> + sig-</code> nature match	non-empty pm path AND APK signature matches expected platform key

Feature class	Required probe	Required evidence shape
permis- sion_grant	dumpsys package <pkg> permission state	permission state = granted: true for the claimed permis- sion
wifi_link	iw dev wlan0 link + RSSI + frequency	captured BSSID AND RSSI within range AND frequency matches claim
wifi_through- put	iperf3 to LAN server	throughput ≥ floor for the negotiated PHY rate
ethernet_link	ip link show + car- rier state + throughput probe	carrier=1 AND speed matches claim AND through- put probe succeeds
dis- play_topology	DRM connector dump + display ID enumera- tion	connector status connected for claimed connector AND display ID present in dumpsys SurfaceFlinger
drm_playback	Widevine session + decoded-frame count	active DRM session AND decoded-frame count >
sub- title_render	analyzer + Tesseract OCR on captured frames	OCR matches expec- ted subtitle text on the claimed dis- play

The taxonomy is **open to additions** via constitution-
submodule commits + propagation. Removal of a feature

class is forbidden without a tracked work item explaining the supersession. Consumer projects MAY extend the taxonomy with project-specific feature classes (e.g., a consuming project's build-pipeline classes) but MUST NOT contract it.

Element : New helper contracts (additive during grace period; mandatory after grace-period end-date - -).

The project anti-bluff helper library (`anti_bluff.sh` in a consuming project; functionally equivalent helper in every consuming project) MUST provide three new helpers:

- . **`ab_pass_with_evidence <description> <evidence_path>`** – the new canonical PASS helper:
 - REQUIRES a non-empty `evidence_path` argument
 - Verifies the path exists AND is non-empty (`[ -s "$path" ]`)
 - If the path is missing or empty, the helper EXITS the test as FAILED with a diagnostic citing the missing evidence path
 - On success, increments the PASS counter AND emits a result line `PASS: <description> [evidence: <path>]`
 - Composes with `ab_send_action` – a single test issues many `ab_send_action` calls and one terminal `ab_pass_with_evidence`
- . **`ab_skip_with_reason <description> <reason_code>`** – the new canonical SKIP helper:
 - REQUIRES `reason_code` from a CLOSED set:
`geo_restricted` , `operator_attended` ,
`hardware_not_present` , `topology_unsupported` ,
`network_unreachable_external` ,
`feature_disabled_by_config` , `artifact_not_yet_built`

- REJECTs ad-hoc reasons – any string not in the closed set fails the helper at call time
- `artifact_not_yet_built` (added – – , research-derived) names the NOT-YET-RUNNABLE class: the check is CORRECT and the topology IS present, but its precondition artifact does not exist yet (the fix is not built / not deployed to the target). It is the legal string for the § . . . REQUIRES-REBUILD PENDING class, which previously had NO code in this closed set and therefore had to either misuse a false code (`hardware_not_present` / `feature_disabled_by_config` – both untrue, and a § . . . misstatement of the real reason) or fail the helper at call time. Reporting a not-yet-runnable check as FAIL is a § . . . FAIL-bluff (it sends teams to fix healthy code); reporting it as PASS is a § . . . PASS-bluff. Its seam-dependence is § . . . 's, NOT flat: on an intermediate artifact it is honest and non-blocking; on the release candidate it is a contradiction that BLOCKS (absence of a verdict blocks exactly as a FAIL does). A flat "always blocks" rule is REJECTED per § . . . as itself a § . . . FAIL-bluff. Consumer-migration note (§ . . . seam): adding a member to a CLOSED set means every existing `ab_skip_with_reason` implementation REJECTs the new code at call time until updated – update the helper alongside the constitution pull.
- FORBIDS `network_unreachable_external` for any feature class in the § . . . taxonomy that lists a sink-side probe – that escape was the § . . . fail-open path § . . . closed for audio, and § . . . closes universally
- SKIP remains mechanically distinct from PASS in cycle math (per § . . .)

- **Bare `ab_pass` deprecation wrapper.** The legacy `ab_pass` helper MUST emit a WARNING line on every call AND emit a deprecation event to a tracked log file (`/data/local/tmp/ab_pass_deprecations.log`). After the grace period (defined in Element) the wrapper MUST FAIL the test outright, forcing all callers to migrate.

Element : Mechanical enforcement via three pre-build gates + three paired meta-test mutations (per § .).

- **CM-SINK-EVIDENCE-PER-FEATURE** – walks every `.sh` test in `device/rockchip/rk3588/tests/` (and equivalent project paths), parses for `# §11.4.69 FEATURE: <class>` annotation comments, verifies the test invokes the corresponding sink-side probe from the taxonomy AND uses `ab_pass_with_evidence` (or `ab_skip_with_reason` with a permitted reason). Pre-grace: tests lacking the annotation WARN. Post-grace: they FAIL.
- **CM-NO-FAIL-OPEN-SKIP** – audits sink-side probe helpers (`arvus_probe.sh` , `sonos_probe.sh` , equivalents) for any code path that converts an empty / unreachable sink response into a PASS-counting SKIP for a feature class with a sink-side probe in the taxonomy. The gate FAILs if any such path exists. Remediation: escalate to FAIL (the empty response IS the defect) or to OPERATOR-BLOCKED per §
- **CM-AB-PASS-WITH-EVIDENCE-EVERYWHERE** – pre-grace WARNs on every bare `ab_pass` call in in-scope tests; post-grace (- -) FAILs the gate. Tests that legitimately have no evidence-capable PASS path MUST migrate to `ab_skip_with_reason operator_attended` and add a tracked work item per § . . . 's `OPERATOR_ATTENDED_ONLY` classification.

Paired mutations (per § 11.4.69):

- **Mutation** : strip the literal `ab_pass_with_evidence` from the project anti-bluff library → assert `CM-AB-PASS-WITH-EVIDENCE-EVERYWHERE` FAILs.
- **Mutation** : inject a fresh test that calls bare `ab_pass` without the `# §11.4.69 FEATURE:` annotation → assert `CM-SINK-EVIDENCE-PER-FEATURE` FAILs post-grace, WARNs pre-grace.
- **Mutation** : inject a fresh test that calls `ab_skip` (no `_with_reason` suffix) OR `ab_skip_with_reason` `network_unreachable_external` for an `audio_output` feature → assert `CM-NO-FAIL-OPEN-SKIP` FAILs.

Element : **Grace period mechanics – additive, not destructive.**

- **Grace period:** `2026-06-19` days from anchor land date – `2026-06-19` through `2026-06-19` inclusive. The end date is hardcoded in the helper library and in this anchor; no environment variable, no command-line flag, no operator override can extend or shorten it.
- **Pre-grace behaviour:** new helpers exist and are usable; the three gates emit WARN for legacy patterns; cycle math counts WARNs separately from PASS / FAIL / SKIP.
- **Post-grace behaviour:** WARN promotion to FAIL is automatic via a date check (`[ "$(date -u +%Y%m%d)" -ge 20260619 ]`) inside each gate's enforcement logic.

Element : **Inheritance via HelixConstitution.** This anchor lives in the constitution submodule and applies to every consuming project that subscribes – including Catalogizer and every future consumer. Each consumer MUST implement its own equivalent helper library + pre-build gates; the taxonomy in Element `CM-AB-PASS-WITH-EVIDENCE-EVERYWHERE` is the canonical

authority across all consumers. Consumer-specific extensions are allowed only as additions to the taxonomy.

Composes with:

- § 2.1.1 (FAIL-bluffs equally forbidden – `ab_skip_with_reason` closed-set prevents script-bug FAILs hiding real defects)
- § 2.1.2 (recorded-evidence requirement – Element helper contracts make captured evidence mechanically required per PASS)
- § 2.1.3 (audio + video quality -layer – § 2.1.4 enforces that the layers are exercised before PASS via the taxonomy probe)
- § 2.1.5 (no-guessing – sink-side evidence eliminates "we think audio is routing" guessing; the codec field reads concretely)
- § 2.1.6 (Arvus sink-side evidence – § 2.1.7 mechanically closes the SKIP-fail-open escape § 2.1.8 left open)
- § 2.1.9 (no-fakes-beyond-unit – § 2.1.10 adds the runtime hook § 2.1.11 lacked)
- § 2.1.12 (deterministic consistency – sink-side evidence is what `ab_run_n_times` hashes for divergence detection)
- § 2.1.13 (autonomous-validation – § 2.1.14 supplies the evidence-shape contract per feature class; § 2.1.15 supplies the classification ledger)
- § 2.1.16 (positive sink-side / downstream evidence for audio + video – § 2.1.17 is the universal generalisation across all feature classes)

No escape hatch:

- No `--skip-evidence` flag.
- No `--config-only-pass` flag (configuration-only PASS is the exact bluff pattern this anchor closes).

- No `--allow-fail-open-skip` flag.
- No `--legacy-ab-pass-permitted` flag post-grace.
- No taxonomy-bypass shortcut – every feature class in scope MUST have its evidence shape enforced.

Propagation gate. `CM-COVENANT-114-69-PROPAGATION` verifies the `$ . .` anchor literal is present across the `~ -file consumer fleet` (parent `CLAUDE.md / AGENTS.md + Containers +` owned atmosphere submodules including the `smarttube-player` nested set + `HelixQA` submodules). Constitution submodule files (`constitution/{Constitution,CLAUDE,AGENTS}.md`) are canonical authority – they carry `$ . .` by definition.

Canonical authority: constitution submodule `Constitution.md`
`$ . .`

Non-compliance is a release blocker regardless of context.

\$. . – Subagent-Driven Execution Is The Default (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"Always do if possible Subagent-driven! Add this into our root (constitution Submodule) `Constitution.md`, `CLAUDE.md` and `AGENTS.md`. This should be the default choice ALWAYS!"

When executing implementation plans authored via `superpowers:writing-plans` (or any equivalent task-decomposed execution flow), the **default execution model is subagent-driven** per `superpowers:subagent-driven-development`. Inline execution via `superpowers:executing-plans` is permitted ONLY when (a) the task is trivial AND fits in a single

sub- -line edit, OR (b) the operator explicitly requests inline execution at brainstorm-handoff time per the `superpowers:writing-plans` skill's explicit-execution-mode prompt.

Why subagent-driven is the default:

- . **Isolated context per task** – each subagent gets a fresh context window scoped to its task, avoiding the conductor's context bloat that degrades reasoning quality after ~ `k` tokens.
- . **Two-stage review naturally enforced** – fresh subagent does the work, conductor reviews, optional second subagent verifies. This composes with `$` . . . four-layer coverage + `$` . . . TDD-fix-discipline.
- . **Parallel-PWU compatible** (`$` . . .) – subagents are the unit of parallelism in the PWU pipeline; making them the default aligns the writing-plans handoff with the parallel-development methodology.
- . **Anti-bluff seam** (`$` . . .) – the conductor's review of a subagent's output is structurally separated from the work itself, eliminating self-review blind spots.
- . **Survives operator absence** – subagents resume from their on-disk plan + spec inputs; the conductor session can be lost without losing in-flight task state.

Composition with `$` . . . (four-layer coverage), `$` . . . (no- guessing – subagent's captured output IS the evidence), `$` . . . (iteration discipline), `$` . . . (TDD-fix), `$` . . . (deterministic consistency – the subagent's deterministic exit code), `$` . . . (LIVE_ADB_FIRST – subagents classify rebuild-requirement before commit), `$` . . . (parallel-development PWU – subagents ARE the parallel work units).

No escape hatch. `--inline-execution-required` , `--no-subagents` , `--monolithic-execution` are NOT permitted flags. Operator may request inline at brainstorm-handoff time per `superpowers:writing-plans` skill; absent that, subagent-driven is default. Skipping subagent-driven for non-trivial work without recorded operator authorisation is itself a `$` . `PASS-bluff` (the conductor self-reviews, the anti-bluff seam disappears, and the forensic class of incident this Constitution catalogues becomes mechanically possible again).

Captured-evidence enforcement. Pre-build gate `CM-COVENANT-114-70-SUBAGENT-DEFAULT-PROPAGATION` enforces this anchor literal in every `CLAUDE.md` / `AGENTS.md` across parent + owned submodules + nested submodules + HelixQA dependencies (the `~` -file consumer fleet). Paired `$` . meta-test mutation strips the anchor literal → gate FAILs.

Canonical authority: constitution submodule `Constitution.md` `$` . . . Default applies to every `superpowers:writing-plans` skill terminal handoff and every multi-task execution flow with independent subtasks.

Non-compliance is a release blocker regardless of context.

`$` . . – Pre-Push Fetch + Investigate + Integrate Mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"before pushing changes to any upstream for any repository – main repo or Submodule, we MUST fetch and pull all changes. Once these are obtained WE MUST investigate what is different compared to head

position we were on last time before fetching and pulling new changes! We MUST understand what is done and for what purpose, especially how that does affect our project and our System in general! Any mandatory changes or improvements required by fresh changes we just have brought in MUST BE incorporated, covered with all supported types of the tests which will produce as a result of its success execution REAL PROOFS of working for all components and functionalities covered and work fully in anti-bluff manner!"

This anchor is the **everyday-push variant** of § . . (Pre-Force-Push Merge-First). § . . governs the destructive force-push case; this anchor governs EVERY push including ordinary fast-forwards. The complementary discipline closes the regression-via-stale-baseline gap: a project committing against a HEAD that's behind any upstream's main risks pushing changes that conflict with, regress, or fail to incorporate accepted work landed in parallel by other consumers (Catalogizer and other consuming projects on the constitution submodule; sibling projects on other shared submodules).

The mandatory -step pre-push cycle (every push, every repository – main + every submodule):

- . **Fetch all remotes** – `git fetch --all --prune --tags` (or per-remote loop if `--all` is unavailable). Capture stdout for the audit trail.
- . **Pull all upstream branches** – `git pull --no-rebase <remote> <branch>` for each configured remote whose tip differs from the local branch. Resolve any merge conflict per § . . step (rebase / merge / cherry-pick per consumer judgment, NOT auto- `--ours` or `--theirs`).
- . **Investigate the diff vs OUR previous HEAD** – `git log --oneline <previous-head>..HEAD` + `git diff --stat <previous-`

head>..HEAD + read EVERY commit's body. For each foreign commit understand: (a) what changed, (b) why (forensic anchor + cited user mandate or §-anchor), (c) how it affects OUR project + OUR System (does it touch surfaces we ship? does it introduce constraints we now MUST honor? does it deprecate or supersede patterns we use?).

- **Integrate mandatory changes + cover with full anti-bluff test coverage** – if the pulled-in work mandates anything (new §-anchor, new gate, new propagation requirement), implement it per § . . (b) four-layer coverage (pre-build gate + post-build inspection + on-device test + HelixQA Challenge) + § . . TDD-fix RED-test-first discipline. Every PASS MUST carry § . . captured-evidence (REAL PROOFS – not metadata-only).
- **Then push** – only after Steps – land. Push to every configured remote in the per-repo cascade order (canonical first, then mirrors). Verify the push landed with `git ls-remote <remote> <branch>` post-push.

Composition with §. . . (constitution-submodule update pipeline – per-submodule specialisation) + § . . . (post-pull validation) + § . . . (fetch-before-edit) + § . . . (full-suite retest before tag) + § . . . (pre-force-push merge-first – the force-push case of this anchor) + § . . . (iteration discipline) + § . . . (TDD-fix-discipline) + § . . . (b) (four-layer coverage) + § . . . (audio + video quality analysis comprehensiveness) + § . . . (no-guessing – every integration decision cites the foreign commit by SHA, not "we think it does X").

No escape hatch – there is no `--skip-fetch`, `--no-investigate`, `--fast-push`, `--trust-upstream` flag. The discipline exists because:

- Pushing-without-fetching produces silent stale-baseline regressions (consumer A's PR diverges from consumer B's PR; neither saw the other; both merge; resulting tree corrupts).
- Pulling-without-investigating destroys the foreign work's traceability (we now carry someone else's code without understanding it; future debugging crosses the § . . . no-guessing line).
- Integrating-without-testing reproduces the § . . . PASS-bluff pattern (we adopt a change but don't prove it works in OUR system).
- Pushing-after-integration-without-evidence reproduces § . . . itself (untested integration = release-blocker by definition).

Per-repository scope: this anchor applies to (a) the parent consuming project's repo, (b) the constitution submodule (across all remotes), (c) every owned submodule (presenter, vlc-player, nova-player, mpv-player, gramophone-player, rhythm-player, strep-player, smarttube-player, torrservice, lampa), (d) every nested submodule (smarttube-player/SharedModules / smarttube-player/MediaServiceCore / smarttube-player/MediaServiceCore/SharedModules), (e) every HelixQA dependency (Challenges, Containers, DocProcessor, LLMOrchestrator, LLMPProvider, VisionEngine, HelixQA).

Audit-trail requirement: every push event MUST be reconstructable post-hoc from `docs/changelogs/<tag>.md` + the per-repo `git log` evidence. Operators auditing a release MUST be able to: () see what each repo's `HEAD..@{u}` was BEFORE this push, () see what foreign work was pulled in, () see what new tests + captured-evidence were

generated by the integration. Anchor enforcement may be partially automated by extending `scripts/commit_all.sh` to refuse pushes without an audit-trail log entry.

Pre-build gate `CM-COVENANT-114-71-PROPAGATION` enforces this anchor literal across the consumer fleet (parent + owned submodules + nested + HelixQA). Paired mutation strips the anchor literal → gate FAILs. Paired meta-test also asserts the audit-trail-log file produced by a recent push exists + lists the foreign-commit SHAs that were pulled in.

Canonical authority: constitution submodule `Constitution.md`
§ 1.1.4.2 . . .

Non-compliance is a release blocker regardless of context.

§ 1.1.4.2 – Audio Top-Priority Mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim,
– –):

"Make sure all fixes for audio are always top priority in main working stream!"

The conductor (main working stream – whether a Claude Code session, an AI agent, or a human operator) MUST treat audio fixes as the highest-priority class on the serial dispatch queue. Audio fixes include but are not limited to \$EU HDMI rejection, \$EM multichannel HDMI, \$ET Arvus integration, \$EV/\$EW/\$EX D audio defects, HiFi (Fix pcm_config_deep, HDMI/SPDIF resilience), AC (Fix FFmpeg Codec), ES (Fix /\$ / mutual-output + DAPM),

multichannel LPCM (Fix ^{11 4 58}), and every future audio-stack improvement. ^{11 4 42}

Composition with § ^{11 4 18} (iteration-discipline mandate) + § ^{11 4 18} (parallel-development methodology). Per § ^{11 4 18}, the conductor selects work items in priority order; this anchor pins audio at the apex of that order. Per § ^{11 4 18} PWU pipeline, parallel BACKGROUND subagents (research, refactors, infrastructure documentation) MAY run concurrently with audio work – they do NOT preempt audio on the main-stream serial dispatch queue.

Operative rule: any time the conductor faces a choice between dispatching an audio task vs a non-audio task on the SAME serial resource (most commonly the pre_build_verification.sh + meta_test_false_positive_proof.sh edit-and-commit critical path), the audio task wins.

No escape hatch – there is no "but this non-audio task is faster" or "but this research is more interesting" override. The discipline exists because audio-stack regressions are user-perceptible and high-impact (D ^{11 4 71} silent post-flash is a release blocker), while research and refactors can wait.

Canonical authority: constitution submodule `Constitution.md`
§ ^{11 4 18}

^{11 4 73} Non-compliance is a process violation regardless of context.

^{2026 05 20}

§ ^{2026 05 20} . . . – Main-specification document
versioning + revision discipline (User
mandate, – –)

Forensic anchor – direct user mandate (verbatim,
– –):

"Make sure everything we add now in previous and upcoming requests IS ALWAYS applied to the main specification – if we have one. Since all these are not major changes we could increase Specification version per change for secondary version instead of the primary. Primary version MUST BE increased for much bigger levels of changes! Add this into root (constitution Submodule) Constitution.md, CLAUDE.md and AGENTS.md as mandatory rule / constraint applicable ONLY IF we have something like the main specification document or we do recognize something like the main specification document. Document MUST BE updated ALWAYS to follow the versioning rules we are applying here + revision and other properties we have!"

Scope condition. This clause applies **only when a project recognises a main specification document** – i.e. a project-level Markdown file (or set of files) that constitutes the canonical "what this system does" reference, typically at `docs/specs/**/specification*.md` or a comparable path.

Projects with no main specification (small libraries, internal tooling) are exempt. Projects that operate per-feature spec files (e.g. ADRs only) are exempt unless they elect to adopt the main-spec pattern.

The mandate. When a project DOES have a main specification document:

- **Every additive operator requirement, refinement, or accepted recommendation MUST be applied to the spec** before, or as part of, the work that implements it. The spec is the source of truth; downstream code and tracking docs reference it.

- . **Spec versioning has two axes** – *primary* and *secondary* (where "secondary" maps to the existing § 4.1.1 metadata-table Revision row):
 - **Primary** (V 1.0 / V 1.1 / V 1.2 / ...) bumps for major rewrites – substantial architecture changes, foundational scope shifts, deprecating large sections, replacing the technology stack. Triggered explicitly by operator decision; old primary versions move to archive/ per § 4.1.1.
 - **Secondary** (Revision) bumps for every other change – section additions, mandate landings, structural reorganisation, polish, type-contract refinements. The metadata table's Revision integer is the secondary version (matches § 4.1.1 monotonic-integer rule).
- . **The metadata table on the spec MUST stay current** – Revision bumps in lockstep with the change; Last modified updates to the change date; Status summary describes the bump's content; Fixed and Fixed summary reference the relevant work IDs.
- . **Cross-document propagation** – when a project has propagated copies of the spec-change rule (Herald uses CLAUDE.md, AGENTS.md, HERALD_CONSTITUTION.md per its § 4.1.1), those copies MUST also reference the active spec file (specification.V<primary>.md) and not a stale archived version.
- . **Archive semantics** – when the primary version bumps, the old specification.V<n>.md moves to <spec-dir>/archive/ specification.V<n>.md and gains Status: superseded with a Continuation pointer to the new active file. Per the § 4.1.1 universal-export mandate the archived .html and .pdf siblings travel with it.

Anti-bluff captured-evidence gate (planned). CM-SPEC-

VERSION-DISCIPLINE :

- If the project advertises a main spec (operator marks via `[project].main_spec_path = "docs/specs/.../specification.V<n>.md"` in a canonical config file), the gate asserts:
 - The file at `main_spec_path` exists.
 - Its metadata table's `Revision` row is $\geq$ the highest `Revision` referenced anywhere else in the repo (Issues.md, Fixed.md, CLAUDE.md, AGENTS.md).
 - If commits exist that touch the spec since the last `Revision` bump, the gate FAILs (operator forgot to bump).

Paired \$. mutation. Backdate the `Last modified` row to a prior date while leaving real-content edits in place → gate FAILs.

Composition. Composes with \$. . . (revision header – `Revision` is the secondary version here), \$. . (metadata table + ToC), \$. . (README always-sync – if README cites the spec, the citation MUST point at the current primary version), \$. . (universal export – archived versions stay exported).

Classification: universal (per \$. .), applicable conditionally per the scope condition above.

Canonical authority: constitution submodule `Constitution.md`
\$. . .

Non-compliance (forgetting to bump, or letting the spec drift behind operator-mandated requirements) is a release blocker.

§ . . – Submodule-catalogue-first
discovery + extend-don't-reimplement (User
mandate, – –)

Forensic anchor – direct user mandate (verbatim,
– –):

"We MUST ALWAYS check which already developed features / functionalities do exist as a part of our comprehensive Submodules catalogue located in vasic-digital and HelixDevelopment organizations on GitHub and GitLab both! Project MUST BE aware of all its existence so we do not implement same things multiple times if they are already done as some of existing universal, reusable general development purpose Submodules! For any missing features that some Submodules we incorporate may be missing we MUST IMPLEMENT the properly and extend those Submodules further! We do control all of the and we CAN and MUST maintain and extend the regularly! All development cycle rules we have MUST BE applied to them and fully respected!"

The mandate. Before scaffolding ANY new module, package, helper, or utility, the contributor (human or AI agent) MUST:

- . **Survey the canonical Submodule catalogue.** Two organisations are authoritative:
 - vasic-digital on GitHub (<https://github.com/vasic-digital>) AND GitLab (<https://gitlab.com/vasic-digital>).
 - HelixDevelopment on GitHub (<https://github.com/HelixDevelopment>) AND GitLab (<https://gitlab.com/HelixDevelopment>). Both organisations mirror across all four supported hosts per § . – surveying one mirror MUST yield the full catalogue.

- . **Inventory existing Submodules.** The canonical repo INDEX is `submodules-catalogue.md` in this submodule (landed - - per user follow-up mandate). It enumerates every owned-by-us repository under `vasic-digital` + `HelixDevelopment` (repos at landing time) with a one-line capability description, categorised so consuming projects can locate "is there already a Submodule for X?" in one glance. The catalogue MUST be re-generated quarterly (or after any non-trivial repo create/rename) per its \$ regeneration recipe.
- . **Reuse before reimplement.** If a Submodule already provides the functionality (or %+ of it), the consuming project MUST add the existing Submodule as a Git submodule, not write the functionality fresh.
- . **Extend in-place when %+ matches but features are missing.** When an existing Submodule is close-but-not-complete, the missing features MUST be added TO THAT SUBMODULE (PR upstream + bump consuming project's submodule pointer) – never as a separate consuming-project helper that duplicates the % already in the Submodule.
- . **Same development-cycle rules apply.** Every Submodule under the two orgs is subject to: \$ anti-bluff covenant, \$ paired-mutation tests, \$ credentials handling, \$ metadata table + ToC, \$ universal Markdown export, \$ spec versioning (if the Submodule has its own spec), \$ multi-mirror push, \$ propagation order. Maintenance + extension of those Submodules MUST follow these rules without exception.
- . **Document the survey result.** Each new feature's tracker entry (Issues.md row, ADR, PR description, or

equivalent) MUST contain a `Catalogue-Check:` field with one of three values:

- `Catalogue-Check: reuse <org/repo>@<sha>` – the existing Submodule covers this.
- `Catalogue-Check: extend <org/repo>@<sha>` – extending an existing Submodule; reference the upstream PR.
- `Catalogue-Check: no-match <date>` – surveyed both orgs on the named date; no existing Submodule covers this functionality; the new code is justified to live in-project. Operators reviewing the PR re-validate this claim.

Anti-bluff captured-evidence gate (planned). `CM-SUBMODULE-CATALOGUE-CHECK` :

- For every new module / package introduced in a non-trivial PR (heuristic: $\geq$ LOC of new non-test code), the gate scans the PR description and any linked tracker row for a `Catalogue-Check:` line. Absence is a FAIL.
- The gate's paired mutation strips the line from a PR and asserts the gate FAILs.

Composition. Composes with `$` (submodule commits propagate first), `$` (every tag on the main repo MUST be mirrored on every owned submodule), `$` (mandatory `install_upstreams` on clone/add), `$` (Submodule-Dependency-Manifest), `$` (post-pull validation), `$` (Submodules-As-Equal-Codebase).

Why this matters. Without the catalogue-check, consuming projects re-implement UUIDv helpers, RLS-tenant-context wrappers, Pandoc-multi-format exporters, `install_upstreams` shell scripts, Markdown ToC generators, `fnv` a hash wrappers, etc. – fragmenting the catalogue and forcing the next project to choose between three subtly-different

implementations of the same thing. The mandate makes the catalogue load-bearing.

Classification: universal (per § 11.4.74). Applies to every project that consumes this Constitution.

Canonical authority: constitution submodule `Constitution.md`
§ 11.4.74.

Non-compliance is a process violation; severe cases (duplicate implementation landed without catalogue check) are release blockers.

2026 05 20

§ 11.4.74 – Mechanical Enforcement Without Exception (User mandate, – –)

Forensic anchor – direct user mandate (verbatim,
– –):

"Why do these violations still happen!? This is a serious problem! We cannot rely on stability nor consistency if we cannot respect our Constitution, mandatory rules and constraints! Is there a way to make this always respected, followed and applied without exception fully and unconditionally!? WE MUST HAVE THIS WORKING FLAWLESSLY!!! Do investigate the root causes of such problems! Once all problems are identified WE MUST apply proper mechanisms for this not to happen NEVER EVER AGAIN!"

11 4 74
11 4 41

11 4 74

§-slot history note. This anchor was originally drafted as § 11.4.74 but renumbered to § 11.4.75 per the § 11.4.74 merge-first mandate after upstream landed a different § 11.4.74 (Submodule-catalogue- first discovery) concurrent with this work. Same renumber

pattern as Phase .DH \$. . collision
resolution (commit 8ad51547115 of the parent consuming
project's repo).

The \$. covenant ("tests pass while feature broken-
for-end-user") historically relied on agent and operator
vigilance to enforce. Three forensic incidents in

→ demonstrated the failure:

subagent a84d2c86f90fdc297 committed docs/research/
workstation/Configurations.md without HTML+PDF siblings

(\$. . violation) when it stalled at the

s watchdog before its pandoc-export step; subagent
a6b38eedf4ef796e8 dispatched to remediate ALSO stalled at
the same operational seam; the parent CLAUDE.md /
AGENTS.md drift after \$. . propagation required
a dedicated catch-up commit f93b25a92eb . The common
failure mode: late-binding enforcement that fires at
pre_build_verification.sh time – hours-to-days after the
violation commit reached every remote.

The mandate. Constitutional invariants MUST be enforced
mechanically at FIVE independent layers – bypassing any
single layer does not bypass the discipline. The five
layers are:

- . **Local pre-commit git hook** – refuses staged .md lacking
sibling .html + .pdf (and other staged-only invariants).
- . **commit_all.sh integration** – the canonical project commit
script invokes the same checks + auto-runs
sync_all_markdown_exports.sh to self-repair before the
commit.
- . **Local pre-push git hook** – re-runs siblings +
propagation gate subset on every commit in the push.
- . **post-commit auto-repair hook** – detects orphan .md in
the just-committed manifest, auto-generates siblings via
pandoc + weasyprint, creates a chore(\$11.4.75): auto-

export ... follow-up commit. Idempotent + recursion-guarded.

- **Local-only equivalent** (Phase .GF, User mandate - -). Remote CI surfaces (GitHub Actions, GitLab pipelines, Jenkins, CircleCI, etc.) are DISABLED per User mandate - - - the workflow file is preserved at .github/workflows/constitution-compliance.yml.disabled-local-only so its contents survive for re-enable but GitHub Actions does NOT execute it (only .yaml / .yml files under .github/workflows/ are honoured). Layer enforcement migrated to the LOCAL pre-build- verification + meta-test run that the operator MUST execute before tagging:
bash device/rockchip/rk3588/tests/pre_build_verification.sh
 - bash scripts/testing/meta_test_false_positive_proof.sh .Layers - remain authoritative; Layer is now the operator's local final- gate ritual. A future re-enable PWU may re-establish remote CI.

Helper contracts (mandatory):

- scripts/install_git_hooks.sh - idempotent installer that copies hooks from tracked scripts/git_hooks/ into .git/hooks/. Hooked into scripts/setup.sh so fresh clones auto-install. Verifies itself with --verify flag.
- scripts/git_hooks/pre-commit - Layer implementation.
- scripts/git_hooks/pre-push - Layer implementation.
- scripts/git_hooks/post-commit - Layer implementation.
- scripts/git_hooks/commit-msg - Bypass-rationale enforcement.
- _constitution_sibling_check function in scripts/commit_all.sh - Layer implementation.

Bypass policy. No layer is perfectly bypass-proof in isolation – that is the design. Bypassing ALL FIVE simultaneously is mechanically impossible because:

- Layer `CM-COVENANT-114-75-PROPAGATION` + `CM-GIT-HOOKS-SOURCE-DIR` can be bypassed by `--no-verify`. Layer `CM-GIT-HOOKS-INSTALL-SCRIPT` (now local pre-tag ritual after Phase `CM-COVENANT-114-75-PROPAGATION`.GF) re-runs the same check before tag creation; tags are NEVER created without it per § 4.1.2.
- Layer `CM-GIT-HOOKS-INSTALL-SCRIPT` can be bypassed by direct `git commit`. Layer `CM-COVENANT-114-75-PROPAGATION` catches that.
- Layer `CM-GIT-HOOKS-SOURCE-DIR` cannot be bypassed by `--no-verify` (runs after commit lands).
- Layer `CM-COVENANT-114-75-PROPAGATION` is now local-only (Phase `CM-COVENANT-114-75-PROPAGATION`.GF – User mandate `CM-COVENANT-114-75-PROPAGATION` – `CM-GIT-HOOKS-SOURCE-DIR`). The operator runs `pre_build_verification.sh` + meta-test before every tag per § 4.1.2. `CM-COVENANT-114-75-PROPAGATION` full-suite-retest mandate. Skipping the local Layer `CM-GIT-HOOKS-SOURCE-DIR` ritual blocks tag creation.

Legitimate bypasses MUST be audited. When `--no-verify` is detected via touch-file marker `.git/ATMO_LAST_BYPASS_ATTEMPT`, the `commit-msg` hook REQUIRES a `Bypass-rationale: <reason>` footer in the commit message. `docs/audit/bypass_events.md` accumulates the audit trail per § 4.1.2 captured-evidence requirement.

Captured-evidence enforcement. Five pre-build gates with paired § 4.1.2 meta-test mutations:

- `CM-COVENANT-114-75-PROPAGATION` – anchor literal across canonical files.
- `CM-GIT-HOOKS-INSTALL-SCRIPT` – installer present + executable.
- `CM-GIT-HOOKS-SOURCE-DIR` – hook bodies present + executable.
- `CM-COMMIT-ALL-SIBLING-CHECK` – `_constitution_sibling_check` in `commit_all.sh`.

- CM-CI-WORKFLOW-PRESENT – CI workflow + ≥ required jobs.

Composes with § . (paired meta-test mutations), § (data safety), § . (end-user-quality covenant – Layer CI enforcement is the § . promise made mechanical), § . . (FAIL-bluffs forbidden), § . . (test-interrupt-on-discovery), § . . (audio + video quality), § . . (no-guessing), § . . (pre-force-push merge-first – this very anchor's renumber from § . . → § . . was driven by § . . merge-first discipline), § . . (TDD-fix), § . . (live-ADB-first), § . . (autonomous-validation), § . . (PWU pipeline), § . . (universal Markdown export – Layer + + are § . . 's mechanical seam), § . . (interactive-clarification), § . . (target-shell-parseability – hooks parse under bash AND mksh), § . . (universal sink-side positive-evidence), § . . (subagent-driven execution), § . . (pre-push fetch + integrate – Layer + honour the merge-first pipeline), § . . (audio top-priority – hooks do not race against in-flight audio commits because they hold no lock themselves), § . . (main-spec versioning – when spec exists, hooks include it), § . . (submodule-catalogue-first – hooks can be added to the canonical catalogue for cross-project reuse).

No escape hatch. No `--skip-hooks` , `--bypass-enforcement` , `--allow-orphan-md` , `--ci-not-applicable` , `--mechanical-enforcement-not-needed` flag exists. The `--no-verify` route IS the deliberate audit-trail bypass; § . . makes the audit trail mechanical via the `commit-msg` footer requirement.

The mandate exists because the User mandate of – – is unambiguous: violations MUST NEVER

EVER AGAIN happen. Reliance on vigilance has demonstrably failed three times in hours. Only mechanical enforcement at every layer can deliver the discipline the project requires.

Propagation gate `CM-COVENANT-114-75-PROPAGATION` enforces this anchor literal across the ~-file consumer fleet. Paired mutation strips the literal → gate FAILs.

Canonical authority: constitution submodule Constitution.md

§ 11 4 7 6 • •

Non-compliance is a release blocker regardless of context.

```

1026 05 20
§ . . – Containers-submodule mandate
(User mandate, – – )

```

Forensic anchor – direct user mandate (verbatim, – –):

"For any work or requirements of running services or codebase inside the Containers (Docker / Podman / Qemu / Emulators, and so on) we MUST USE / INCORPORATE the Containers Submodule properly: <https://github.com/vasic-digital/containers> (git@github.com:vasic-digital/containers.git). Containers Submodule contains all means for us to Containerize our code and services! If any feature or Containing System is missing or not supported we MUST EXTEND IT properly like we do all of our projects! All these details MUST BE added to our root constitution (constitution Submodule) - Constitution.md, AGENTS.md, CLAUDE.md, QWEN.md - and followed and respected! No bluff work is allowed of any kind!"

\$-slot history note. This anchor was originally drafted as § 101.1.1 but renumbered to § 101.1.2 per the

§ . . . fetch-before-push mandate after a concurrent upstream commit (0a70083) landed § . . . (Mechanical Enforcement Without Exception). Same collision-resolution pattern as § . . . 's own renumber-from-§

The mandate. For ANY containerized workload – Docker, Podman, Qemu, Kubernetes, container-backed emulators (Android emulator, VM-based testbeds, etc.) – every consuming project MUST:

- **Use the canonical Containers Submodule.** `https://github.com/vasic-digital/containers` (`digital.vasic.containers`) is the authoritative library for runtime auto-detection, endpoint discovery, lifecycle/health management, compose orchestration, cross-build, emulator integration, and on-demand service boot. Mirrored across the four § . . . -canonical hosts (GitHub + GitLab + GitFlic + GitVerse).
- **Install as a Git submodule.** Container-orchestration concerns enter the consuming project as a `git submodule` `add git@github.com:vasic-digital/containers.git <path>/containers` . Project `go.mod` (or equivalent dependency manifest) consumes via `replace digital.vasic.containers => ./<path>/containers` during development; production builds resolve through pinned commit SHAs.
- **Boot infrastructure on demand.** Tests, CLI doctor commands, and local-development workflows that depend on Postgres / Redis / OTel / message-brokers / emulators MUST invoke the Containers Submodule's `pkg/boot` + `pkg/compose` + `pkg/health` APIs to bring infra up automatically. Operators MUST NOT be required to manually start `podman machine` / `docker compose up` before tests – the boot is part of the test entry point. This is the **on-demand-infra invariant**.
- **Extend the Containers Submodule, never reimplement.** Per § . . . 's extend-don't-reimplement rule applied

specifically to containerization: if a runtime (e.g., a new emulator type) or lifecycle primitive (e.g., a new health-check kind) is missing, the consuming project **MUST** add it to `vasic-digital/containers` (PR upstream + bump the consuming project's submodule pointer) – never as a parallel implementation inside the consuming project.

- **Anti-bluff captured-evidence requirement.** Every integration test that claims to exercise a containerized component **MUST** actually boot that component via the Containers Submodule. Tests that pass without containers actually running (e.g., short-circuit fakes that bypass the boot path) are a § 4.1.1 bluff violation. A passing test **MUST** imply the infra was up.
- **Composition with § 4.1.1.** The Containers Submodule itself is one of the Submodules surveyed by `Catalogue-Check`. A project tracker row touching containerization **MUST** record `Catalogue-Check: extend vasic-digital/containers@<sha> (or reuse)` – `no-match` is only valid if the consuming project demonstrates the Containers Submodule's API cannot model the workload (rare).

Anti-bluff captured-evidence gate (planned).

`CM-CONTAINERS-USED` :

- For every PR that touches `Dockerfile*`, `*compose*.yaml`, `*podman*.yaml`, `Vagrantfile`, or any code under a directory matching `(test|integration|e2e)/.*infra` heuristic, the gate scans for an import of `digital.vasic.containers/...`. Absence in a non-trivial container-touching PR is a FAIL.
- The gate's paired mutation strips the import and asserts the gate FAILs.

Composition. Composes with § 4.1.1 (catalogue-first discovery), § 4.1.2 (mechanical enforcement – this

clause IS one of the things mechanical enforcement protects), § (submodule propagation), § . . (Submodule-Dependency-Manifest), § . . (.install_upstreams on clone), § . . (Submodules-As-Equal-Codebase: the Containers Submodule itself follows all dev-cycle rules), § . . (multi-format export for docs maintained inside the Containers Submodule), § . (paired-mutation gates protect the Containers Submodule's own boot / health primitives).

Why this matters. Without the on-demand-infra invariant, integration tests silently degrade to "pass without infra running" (a § . bluff) or break with cryptic "connection refused" errors that operators waste hours debugging. The Containers Submodule turns infra-boot into a typed, programmable, anti-bluff-tested capability rather than a sequence of imperative shell commands. Consuming projects converge on a single shared lifecycle/health model instead of each project hand-rolling their own.

Classification: universal (per § . .). Applies to every project that runs code inside any containerization or virtualization runtime.

Canonical authority: constitution submodule Constitution.md § . . .

Non-compliance is a process violation; landing containerized infra without consuming the Containers Submodule (i.e., reinventing compose orchestration in-project) is a release blocker.

§ . . – Regeneration-mechanism-
required mandate (User mandate,
– –)

Forensic anchor – direct user mandate (verbatim,
– –):

"We must be sure that after excluding anything from Git versioning we still have the mechanism which will out of the box obtain or re-generate missing content! Add this mandatory safety rule / constraint into ours root (constitution Submodule) Constitution.md, CLAUDE.md, AGENTS.md, QWEN.md or any other required document / file from the constitution Submodule! Do not forget to fetch and pull Submodule first! Commit and push all Submodule changes to all upstreams!"

274

Forensic incident (– – T₁₅₁ : Z, parent project, a consuming project). The parent's audio Tier commit_all.sh stalled four hours on git add -A while scanning GiB of untracked .git-backup-20260520T084610Z-pre-orphan-kill/ + GiB RKTools/linux/ + GiB qa-results/ content – all of which should never have been tracked. The natural remediation is to add those paths to .gitignore, BUT doing so without a regeneration mechanism would silently orphan fresh clones: a new operator running git clone && cd <project> would find missing RKTools, missing test infrastructure, missing build outputs they cannot reproduce. § . . codifies the universal rule that closes this gap.

The mandate. Every .gitignore entry that excludes either (a) more than ~ MiB of content, OR (b) any artifact essential to building / running / testing the project, MUST be accompanied by a documented + automated mechanism

that re-obtains or re-generates the excluded content on a fresh clone. The closed-set choices are:

- . **Re-obtain** – download from an authoritative source (vendor tarball, SDK installer, npm / pip / cargo / go-mod registry, container image registry, dedicated git submodule, S³, GCS, vendor-internal artifact server) via a script that runs deterministically on a fresh clone.
- . **Re-generate** – produce the content from tracked source code via a script (build pipeline, code generation, asset rendering, captured-evidence replay, container build, kernel build, image packaging) that runs deterministically on a fresh clone.

Required artefacts per qualifying `.gitignore` entry (each is independently mandatory):

- . **Metadata file at `.gitignore-meta/<entry-slug>.yaml`** (or equivalent project-canonical location declared in the project's `CLAUDE.md`) carrying: the gitignore-pattern verbatim, `mechanism-type: re-obtain | re-generate`, `script-path: <repo-relative path>` (the deterministic regeneration entryptoint), `expected-disk-usage: <bytes-or-human-readable>`, `vendor-url-or-source: <URL or in-tree path>` when applicable, `integrity: { algorithm: sha256 | md5, value: <hex> }` when the content is fetchable from a known-stable mirror, `requires-network: true | false`, `requires-credentials: true | false` (and which credentials slot per `$ . . .`).
- . **Entry in the post-clone bootstrap script** (`scripts/setup.sh` or the project-canonical equivalent). The script MUST run the regeneration mechanism non-interactively on a fresh clone or be invocable as an idempotent step (`scripts/setup.sh --regenerate <entry-slug>`).

- **Pre-build gate** that verifies either the regenerated content is present at the expected path OR the regeneration script ran successfully (e.g., a stamp file under `.gitignore-meta/.regenerated/<entry-slug>.ok` with `mtime ≥ source freshness window`).
- **Documentation update** – README plus the relevant `docs/guides/*.md` guide describing the mechanism, including the manual fallback (vendor URL, alternative download mirrors), the time + disk-usage budget, and per-`$ . . .` credentials requirements.

No escape hatch. Adding a bare `.gitignore` entry that excludes significant or essential content without the accompanying mechanism is itself a `$ . . .` PASS-bluff variant – the codebase appears complete to the casual eye, every pre-build / post-build gate goes green, but a fresh clone cannot build / run. There is no

`--skip-regen-mechanism`, `--gitignore-is-enough`, `--operator-already-has-content` flag.

Composition with sister anchors:

- `$ . . .` (no-guessing) – the mechanism MUST be verified working end-to-end on a sandbox clone before the `.gitignore` entry lands. `LIKELY` works on a fresh clone is the exact bluff `$ . . .` forbids.
- `$ . . .` (universal Markdown export) – generated HTML / PDF siblings are an instance of `re-generate`; the auto-regeneration helper (`sync_all_markdown_exports.sh`) is the `$ . . .` mechanism for that gitignored content where applicable.
- `$ . . .` (interactive clarification) – if the agent is uncertain whether a candidate `.gitignore` addition qualifies under `$ . . .`, the agent MUST ASK the operator via the platform's interactive question mechanism rather than guess.

- § 4.6.1 (pre-push fetch + integrate) – the regeneration mechanism's integrity hash and vendor URL MUST be re-validated against upstream before pushing the § 4.6.1 -governing commit; stale hashes are a § 4.6.1 violation.
- § 4.6.2 (catalogue-first + extend-don't-reimplement) – if the regeneration mechanism can be implemented by extending a Submodule already in vasic-digital/HelixDevelopment (e.g., a reusable downloader / decompressor helper), the project MUST extend rather than reimplement.
- § 4.6.3 (Mechanical Enforcement Without Exception) – the local pre-commit hook (Layer 1) MUST refuse a commit that adds new .gitignore lines without a corresponding .gitignore-meta/*.yaml; the CI workflow (Layer 2) MUST replay the gate on every push. Paired meta-test mutation: strip a metadata YAML entry → gate FAILs.
- § 4.6.4 (Containers-submodule mandate) – container images excluded from VCS are regenerated via vasic-digital/containers pkg/boot + Dockerfiles tracked in-tree; the .gitignore-meta/<entry-slug>.yaml references the Containers Submodule build path.
- § 4.6.5 / § 4.6.6 (zero-risk data safety) – the regeneration mechanism MUST be pre-tested in a sandbox clone with a hardlinked-backup safety net BEFORE relying on it for live operator clones; a regeneration script that silently corrupts content on a fresh clone is a § 4.6.5 violation.
- § 4.6.7 (propagation order) – submodules consuming this constitution submodule inherit § 4.6.7; their own .gitignore entries are subject to the same mechanism requirement, propagated via their own CLAUDE.md / AGENTS.md / QWEN.md.

Anti-bluff captured-evidence gate (planned). CM-GITIGNORE-REGEN-MECHANISM :

- For every PR that adds a line to `.gitignore` (or to any `*.gitignore` file in a sub-tree), the gate scans for a matching entry in `.gitignore-meta/`, verifies the `script-path` exists and is executable, verifies the metadata YAML parses and carries the required keys, and verifies the post-clone bootstrap references the mechanism.
- Bare-line additions (e.g., `echo 'qa-results/' >> .gitignore`) without the metadata sibling are a FAIL.
- The gate's paired `$ . .` mutation strips one required YAML key (e.g., `script-path:`) and asserts the gate FAILs.

Why this matters. Without `$ . .`, the natural `.gitignore` ergonomics – "this is too big to track, exclude it" – leads to fresh-clone orphanage: operators or CI runners (or new developers, or end-user mirror sites) clone the project and discover essential content is missing, with no documented way to obtain it. The cost is paid at every fresh-clone event, every CI bootstrap, every operator onboarding, every disaster-recovery exercise. `$ . .` forces the cost to be paid once, at the `.gitignore`-addition commit, where it is bounded and visible.

Forensic anchor crosswalk to today's incident. Without `$ . .`, the natural fix to today's `commit_all.sh` -hour-stall – adding `.git-backup-*`, `RKTools/linux/`, `qa-results/` to `.gitignore` – would have orphaned every fresh clone of a consuming project. With `$ . .`, each of those entries now requires: (a) a `.gitignore-meta/<slug>.yaml` declaring the mechanism (e.g., `.git-backup-*` is a transient artefact safe to re-create per `$ . .` `hardlinked-backup`; `RKTools/linux/` is re-

obtainable from Rockchip's vendor tarball; qa-results/ is re-generatable by re-running the on-device test cycle); (b) bootstrap-script entries; (c) gates; (d) docs. The bare .gitignore addition is now mechanically blocked.

Classification: universal (per § 11.4.77). Applies to every project that excludes content from VCS for size / regenerability / vendor-provenance reasons.

Canonical authority: constitution submodule Constitution.md § 11.4.77.

Non-compliance is a release blocker regardless of context.

2026-05-20

2026-05-20
§ 11.4.77 – CodeGraph code-intelligence mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"All gathered knowledge and steps MUST BE added as future reusable experience into our constitution Submodule on setting up / configuring and using codegraph with all supported CLI agents we may use for our codebase! Make codegraph MANDATORY CHOICE for this purpose for all of our project and add proper comprehensive details about everything codegraph related into our root constitution (constitution Submodule) Constitution.md, AGENTS.md, CLAUDE.md, QWEN.md and other constitution Submodule relevant files! All project which do not have configured and installed codegraph yet MUST DO IT and MUST USE IT!"

The mandate. Every consuming project under this Constitution MUST install, initialize, and use **CodeGraph** (<https://github.com/colbymchenry/codegraph> , npm package

`@colbymchenry/codegraph`) as its code-intelligence layer for AI coding agents. CodeGraph builds a local SQLite semantic knowledge graph of a codebase and exposes it to agents over the Model Context Protocol (MCP) – % local, no cloud, no external API (compatible with the Local-Only CI/CD posture). A project worked on by AI coding agents that has not yet installed and wired CodeGraph MUST do so; using it is not optional.

- **Install.** `npm install -g @colbymchenry/codegraph` (Node.js +). The npm global prefix MUST be user-writable – installation MUST NOT require `sudo` / `su` (`$` / `$` host-safety + the no-privilege-escalation posture). Native modules (`better-sqlite3`, `tree-sitter`) compile on install with a WASM fallback.
- **Initialize + index.** `codegraph init` at the project root creates `.codegraph/config.json` (tracked) and `.codegraph/codegraph.db` (a regenerable build artifact – gitignored per `$` . . . / `$` . . . , with `codegraph index` as its declared regeneration mechanism). The `config.json` `exclude` list MUST exclude (a) other-owned submodules and vendored trees – index the consuming project's own domain code – and (b), non-negotiably, every credential/secret path (`.env*`, keystores, signing keys, service-account JSON) per `$` A secret reaching the index is a `$` . . . violation. `codegraph index` builds the graph; `codegraph sync` keeps it fresh.

CRITICAL – filter mechanism. CodeGraph v .x is zero-config: it uses BUILT-IN skip lists (`node_modules`, `vendor`, `dist`, `build`, `target`, `.venv`, `Pods`, `.next`) PLUS the project root `.gitignore` to decide which files to index. The `.codegraph/config.json` `include` / `exclude` fields are **not used** for file filtering – all filtering happens through the ignore matcher seeded with the built-in defaults + root `.gitignore` patterns. To

control what gets indexed, use root-anchored `.gitignore` patterns (start with `/` to match only the top-level directory) and `!` negations to re-include owned paths under a default-skipped parent (e.g., `!/vendor/widevine_*` re-includes owned vendor code that CodeGraph's built-in `vendor` skip would otherwise block).

MUST-EXCLUDE classes (actual rubbish – build outputs, caches, credentials, binaries):

```
** (a) Build outputs and regenerable artifacts: ** \out/, \prebuilts/, \rockdev/,  
** (b) Caches, dependencies, tooling: ** \.gradle/, \__pycache__/, \.venv/, \flash_lo  
** (c) Credentials, secrets, QA/recording artifacts: ** \secrets/, \scripts/testing/se  
** (d) Config.json include/exclude (DEPRECATED – not enforced by v1.x engine): ** mainta
```

ALL SOURCE CODE MUST REMAIN INDEXABLE, including AOSP platform trees (`frameworks/` , `external/` , `art/` , `cts/` , `bionic/` , `packages/` , `system/` , `kernel-5.10/` , `hardware/*/` , `vendor/` , `device/*/`) – these are tracked source and MUST NOT be gitignored for CodeGraph purposes. Only the rubbish classes above may be excluded. A project with `M+` tracked files (AOSP fork) will produce a correspondingly large index (`-` GB for the first init, with subsequent `codegraph sync` operations fast because they use `git diff`). The large initial size is expected and correct – the index represents the full codebase the AI agent needs to understand cross-references.

Universal mandatory-exclude baseline set (`$ . .` + `$ . .` refinement – binds EVERY consuming project). Independently of project, the CodeGraph index MUST NEVER include the following file-classes anywhere – they are indexed-rubbish in every project; a consuming

project MAY add its own project-specific rubbish paths on top per § . . . but MUST NOT drop any baseline class: (a) **build outputs / regenerable artifacts** – `/out/` , `/build/` , `/dist/` , `/target/` , `bin/` , `obj/` , `.compressed/` , and any generator-produced tree; (b) **caches / dependency / tooling dirs** – `**/ccache/` , `**/.ccache/` , `.gradle/` , `__pycache__/` , `.pytest_cache/` , `.mypy_cache/` , `.ruff_cache/` , `node_modules/` , `.next/.` , `.cache/` , `.terraform/` , `.venv/` , `logs/` ; (c) **credentials / secrets** – `.env` , `.env.*` , `**/*.env` , `secrets/` , `**/*secret*` , `keystores` / `signing-keys` / `service-account JSON`, and every credential-bearing path per § . . . (a secret reaching the index is a § . . . violation, never negotiable); (d) **QA / recording / large regenerable corpora** – `**/qa-results/` , `**/recordings/` , and prebuilt-binary trees. Realized via root-anchored `.gitignore` patterns (CodeGraph v .x zero-config filter = built-in skip list + root `.gitignore` , NOT `config.json` include/exclude). Own-org submodule sources stay INCLUDED per § . . . ; third-party vendored trees EXCLUDED. This baseline is UNIVERSAL (§ . . .) – dropping any class is a § . . . / § . . . violation, and no source-code tree may be excluded to satisfy it.

- **Wire every supported CLI agent.** The CodeGraph MCP server (`codegraph serve --mcp` , `stdio` transport) MUST be registered with every AI coding CLI agent the project's developers use. Registration is project-scoped and committed where the agent supports it (e.g. Claude Code `.mcp.json` , OpenCode `opencode.json` , Qwen Code `.qwen/settings.json` , Crush `.crush.json`); host-local where the agent stores MCP config outside the repository (e.g. Kimi CLI `~/.kimi/mcp.json`). Every MCP config MUST reference the bare `codegraph` command resolved on `PATH`

– never a hardcoded host path (the no-hardcoding posture). Claude Code is the canonical primary agent; the others MUST work too.

- **Anti-bluff verification is mandatory.** CodeGraph integration MUST be covered by an anti-bluff verification suite (the canonical reference implementation is a consuming project's `scripts/verify-codegraph.sh` + `tests/codegraph/`, six layers): index reality, query correctness, MCP-protocol JSON-RPC, per-agent connectivity, per-agent end-to-end LLM drive, and a falsifiability rehearsal that mechanically proves the suite FAILs when the index is deliberately broken. The per-agent end-to-end layer MUST use an **unforgeable challenge** – a fact obtainable only by calling a CodeGraph MCP tool (e.g. the index node count via `codegraph_status`) – so an agent answering from its own file-reading tools cannot produce a false PASS. An agent that genuinely cannot be driven end-to-end in the test environment (missing credentials, exhausted quota, environment incompatibility) is recorded as a documented SKIP gap per § 4.1.1 – never a faked PASS.
- **Comprehensive documentation.** Every project MUST carry a `docs/CODEGRAPH.md` (or canonical equivalent) describing install, initialization, indexing, per-agent wiring, the verification suite, and troubleshooting – kept in sync per § 4.1.2.
- **Catalogue-first.** CodeGraph is a third-party developer tool, not an owned submodule – it does NOT enter the project as a `git submodule` (per § 4.1.3 it is consumed as the published npm package) and it does NOT add a Git remote.

Anti-bluff captured-evidence gate (planned). `CM-CODEGRAPH-WIRED`: for every consuming project, the gate verifies `.codegraph/config.json` exists with the § 4.1.4.

secret-exclusions present, every developer-used agent carries a CodeGraph MCP registration, and the verification suite exists and is executable. The paired \$. mutation removes a secret-exclusion from config.json and asserts the gate FAILs.

Composition. Composes with \$. . (per-environment-topology SKIP for un-runnable agents), \$. . (credentials never indexed), \$. . + \$. . (CODEGRAPH.md kept in sync + exported), \$. . / \$. . (the .codegraph/codegraph.db artifact is gitignored with codegraph index as its regeneration mechanism), \$. . (catalogue-first – CodeGraph is consumed, not reimplemented), \$. . (the verification suite is anti-bluff by construction – CI green is necessary, never sufficient), \$. . (paired-mutation gate).

Why this matters. AI coding agents that explore a codebase by repeated file scanning consume tokens and tool calls wastefully and build a shallow, drift-prone mental model. A pre-indexed semantic graph gives every agent instant, consistent symbol / caller / callee / impact resolution. Standardising on one tool – CodeGraph – across every Helix project means the capability is wired once, verified anti-bluff once, and reused everywhere, instead of each project hand-rolling ad-hoc context tooling.

EXTENSION – scope, launcher, runner patches and verification are CONSUMED BY REFERENCE; the consumer supplies only DATA (research-derived,
- -). Measured this cycle on a multi-hundred-thousand-file consuming project: clause 's filter text contradicted itself and \$. . steps
- (one sentence says the .codegraph/config.json exclude list is not used for filtering, another instructs editing it), the installed CLI version read a DIFFERENT config file that did not exist in the repository and

silently fell back to zero-config, nothing derived the scope from the submodule graph, and the anti-bluff gate treated any database above a size threshold as a real index – so a stalled, partial index would have PASSED. This extension closes those gaps without deleting any text above; where it conflicts with the clause- filter wording, THIS extension governs (\$. . . extend-don't-re-mint).

- **ONE MECHANISM, INHERITED BY REFERENCE.** The scope machinery (<constitution>/scripts/codegraph/scope_render.py , the runner-discovery bridge scope_enumerate.js , the baseline class list scope_baseline.txt , the DATA template scope.example.yaml), the single sanctioned writer entry (codegraph_safe.sh + codegraph_safe_helper.py + codegraph_guard.sh + disk_tripwire.sh), the version-keyed runner patcher (fk_index_patch.py + runner_patches/), the bulk-window hazard probe (fk_cascade_probe.py), the progress-proven watchdog (index_watch.py), the scope proof (codegraph_scope_guard.py) and the semantic-index scope generator (<constitution>/scripts/lumen/gen_lumenignore.py) MUST live in the constitution.submodule and MUST be INVOKED BY REFERENCE from the consuming project (\$. . . / \$. . . / \$. . .). A consumer copy, fork, vendored duplicate or hand-rolled equivalent is a violation (\$. . . – no parallel mechanism). Each tool MUST be documented (\$. . .) and test-first with paired mutations (\$. . . / \$. . .), and a tool lacking its own tests MUST NOT be pointed at a live index; at the time of writing index_watch.py and fk_cascade_probe.py (no test references of their own – the writer entry's unit test substitutes a stub for the probe), the runner patches lockfix1 and datafrag1 (no tests of their own – excluded from the default patch set for that reason), and the \$. . . companion guide of EVERY tool

named here (none exists yet) are OWED, tracked per
§ . . .

. **THE CONSUMER SUPPLIES DATA, NEVER MECHANISM**

(§ . . .). A consuming project contributes exactly one declarative scope file (schema = the constitution's `scope.example.yaml`): its own-org list, its project-specific rubbish paths, its pathological inputs (files that stall the parser – declared, never retried to timeout), explicit re-include negations, per-submodule class overrides WITH a reason, the accepted indexed-file count and its tolerance, and its efficiency thresholds (§ . . . (D)). It MUST NOT drop a baseline class (the universal baseline above still binds).

. **CONFIG TRUTH = WHAT THE INSTALLED CLI ACTUALLY READS.**

The scope source of truth is the configuration the INSTALLED CLI version loads, established from that version's own loader code (§ . . . latest-source, § . . . never from memory); a config file the CLI does not read is dead documentation and MUST NOT be cited as the scope. The effective scope artifact is GENERATED by the renderer from the baseline classes + the § . . . classification + the consumer DATA; hand edits are refused, and a missing or malformed config that the CLI silently falls back from is a FAIL, never a warning.

. **SCOPE PROOF BEFORE "READY" (supersedes size-based evidence).** Before an index is called usable, the scope guard MUST enumerate files through the installed runner's OWN discovery code and assert, each with a control needle (§ . . .): indexed-file count within the consumer tolerance of the accepted count; zero files under third-party roots; zero secret-class files (§ . . .); at least one file under every own-org root (§ . . .). Database size is NEVER

evidence of a real index. Completeness (pending work ==
) is \$. . . (C).

EXTENSION – host-adaptive. V heap budget lands in the
single writer entry, closing a real OOM crash on this
consuming project's bulk index run (research-derived,
- -). Forensic FACT: on this constitution's
own consuming project (, tracked files,
~ . M edges), a bulk codegraph index run reached
% file parsing then crashed in the resolve phase
with FATAL ERROR: Reached heap limit ... JavaScript heap out of
memory (rc=), discarding an approximately -
minute run – on a host with, GiB of RAM available
at the time. Root cause: a stock-launched Node.js indexer
inherits V 's DEFAULT old-space heap ceiling (roughly
GiB) regardless of the host's real memory, since
codegraph_safe.sh (the item- single sanctioned writer
entry) never set one – a self-imposed software limit
masquerading as real memory exhaustion, exactly the class
of un-measured resource quota \$. and \$. .
forbid. Fix: preflight() now computes a heap budget from
MEASURED MemAvailable (half of it, floored at
MiB so it always exceeds the stock default, capped at
MiB so one process alone never threatens the
\$. %-of-total-RAM ceiling – never hardcoded,
applying \$. 's dynamic-detection discipline to this
bare-metal indexer process, not only to \$. 's own
containerized-build scope) and exports it via
NODE_OPTIONS=--max-old-space-size=<budget> before every init /
index / sync launch; an unmeasurable budget REFUSES the
run (die 6) rather than silently falling back to the
unmeasured stock limit (\$. . conservative-safe
default). Landed with two new regression tests (scripts/
codegraph/tests/test_unit_safe.sh U /U , item- test-
first discipline) proving the exact floor/cap/proportional
arithmetic, and independently VERIFIED LIVE (not merely
unit-tested): the relaunched bulk run's real indexer

process was confirmed via `/proc/<pid>/environ` to have genuinely received `NODE_OPTIONS=--max-old-space-size=65536`. Zero-regression proof: the full pre-existing test suite produced identical pass/fail counts before and after (git-stash comparison), isolating the delta to `U /U` flipping FAIL→PASS. Companion doc updated per

```
$ . . . ( docs/scripts/codegraph_safe.md ). Commit
941841d3e7235a081e3199a4d3073b409d41753b . Composes
$ . . . / $ . . . / $ . . . / $ . . . /
$ . . . ( )( ) (upgrade-gated compatibility probes
and supervised long runs – this closes exactly the crash
class those clauses anticipate for bulk runs) /
$ . . . / $ . . . / $ . . .
```

CORRECTION (same day, operator mandate – –
– "do not set caps for codegraph and lumen ... indexed space exposed and used by all agents"): the "capped at

`60` `150` **MiB" clause above is SUPERSEDED, not deleted.**

The `GiB` figure was an INVENTED ceiling, not a derivation of `$` : on a `GiB` host, `GiB` is ~ % of total RAM, far below the actual `$` bound (% ~ `GiB`) – an artificial cap that starved a legitimate large-repo resolve phase for no principled reason (observed live: the relaunched run's RSS climbed to ~ `GiB` and plateaued against the invented ceiling, a genuine crash risk). Fixed: the ceiling is now `MemTotal * 60 / 100` (in MiB), computed fresh every preflight run from `/proc/meminfo` – the REAL `$` bound, never a fixed number. Mathematically this ceiling can never bind an ordinary run (half-of-available is always < half-of-total, which is always < %-of-total), so in practice the heap budget is now effectively unconstrained short of the one absolute, no-escape-hatch host-safety invariant this constitution carries – exactly matching the operator's mandate without weakening

```
$ . . . New regression test test_unit_safe.sh U
proves the cap is real and independently computed (not
```

inert, not the old hardcoded number) by constructing a scenario where it overrides even the MiB floor. Full-suite zero-regression proof unchanged in kind:

```
cases=46 failed=31 , same pre-existing unrelated failures. Commit 8a5f47ab2e0e7a05d1c1652ba5a3475751652062 .
```

EXTENSION – CodeGraph is now wired as a genuine, project-tracked, first-class MCP tool for this consuming project (not only the bare-`codegraph -on-PATH` posture clause above describes), closing a "tool exists but no agent can discover it" gap found under the same operator mandate (forensic incident + fix, - -). Audit finding: prior to this extension, this project's `.mcp.json` was `{"mcpServers": {}}` – EMPTY – despite CodeGraph shipping a full first-class MCP query surface (`codegraph_explore` , `codegraph_node` , and more, confirmed from the installed CLI's own `--help` text) and its own installer (`codegraph install`) built for exactly this purpose. The gap was total: an agent had no way to discover CodeGraph's query tools short of already knowing to hand-invoke the raw CLI via Bash – the opposite of "regularly used for all major needs." Separately, the semantic (Lumen) tool's own "use this instead of grep" nudge, observed firing reliably all session, was traced to a Claude Code PLUGIN hook (`lumen/<version>/cmd/hook.go`) – entirely UNTRACKED by this repository; it works only by accident of which plugins happen to be enabled on a given host/account, and would not follow this project to a fresh clone, a different agent, or another track. `grep -rl "codegraph\\|lumen" constitution/scripts/hooks/ scripts/hooks/` returned nothing – zero project-tracked enforcement of either tool's use existed anywhere.

Fix landed: () `.mcp.json` now registers a `codegraph` MCP server whose command is `bash constitution/scripts/codegraph/codegraph_mcp_serve.sh` – a new, project-tracked, `$ . . -decoupled wrapper`, NOT the bare `codegraph`

command clause above names. This is a REFINEMENT of clause 's "no-hardcoding posture," not a violation: clause forbids a HARDCODED HOST PATH, and the wrapper hardcodes nothing – it dynamically resolves the SAME validated, patched runner `codegraph_safe.sh` itself uses (via `fk_index_patch.py --print-bin`, re-derived every launch, correct automatically across a future `codegraph` upgrade) rather than trusting whatever `codegraph` happens to resolve first on `$PATH` (the STOCK, unpatched build – this project's own `codegraph_safe.sh` NEVER trusts stock for a write operation, per the – incident item above already documents; the wrapper extends that same discipline to MCP serving). Consuming projects with no equivalent patched-runner requirement MAY use clause 's plain bare-`codegraph` form; this wrapper is the pattern for a project that, like this one, cannot safely assume stock-on-PATH is correct.

A SEPARATE, MORE SERIOUS forensic incident surfaced while building and testing this wiring, and is now a permanent, tested safety property of the wrapper. `codegraph serve --mcp` is NOT a passive read-only query server – with its file watcher enabled it auto-syncs (WRITES) the database on detected file changes, and `serve --mcp` daemonizes (re-parents to `init`, outliving its own launching process or a `timeout`-based kill – this is CodeGraph's own documented persistent-background-service architecture). A first version of the wrapper omitted `--no-watch`; a single ~ - second manual test of it against THIS project's own live checkout spawned a detached daemon that held `.codegraph/codegraph.db` / `-wal` / `-shm` open with an active write-capable file watcher for roughly two minutes before being found and killed – CONCURRENTLY with an in-flight bulk `sync` continuation also writing to the SAME database, which then hard-crashed with `Failed to index: database is locked`, discarding a further ~ minutes of resolve-phase progress on top of the earlier heap-crash loss the same

day. Root cause: two writers, one database, zero coordination – `codegraph_safe.sh` 's own single-writer flock/lock discipline governs only ITS OWN launches; an MCP server started via `.mcp.json` launches completely outside that mechanism and was never taught to respect it.

Recovery: the stray daemon was killed, `PRAGMA integrity_check` found exactly `rows` missing from one index (`idx_nodes_lower_name`, a mid-transaction-crash artefact, not page-level corruption), repaired via `REINDEX` (a data-presence check on the affected rows confirmed the underlying table data was intact before and after), and the bulk operation was resumed via `codegraph_safe.sh ... sync` (BULK-classified on `pending>=threshold`, correctly continuing the interrupted resolve phase rather than re-parsing `K` files from scratch).

The durable fix: the wrapper now ALWAYS passes `--path <project> --no-watch` – `--no-watch` is not optional, ever, per the CLI's own documented behaviour ("Disable the file watcher (no auto-sync)"), making this a TRUE query-only reader that can run PERMANENTLY (as a registered MCP server necessarily does) alongside future bulk index/sync operations without risk of this collision class recurring; a permanently-registered MCP server WILL eventually overlap with a future bulk run in an actively multi-tracked project – the two-writer collision is not a one-off testing artefact, it is the exact risk profile permanent registration creates. An occasionally-stale query view (never auto-refreshing) is the correct trade against a database corruption/crash risk; `codegraph sync` (or a fresh query once an agent knows a sync completed) is how staleness is cleared, deliberately, never automatically. Verified DECISIVELY, not by `argv` inspection alone (the daemon's own child process does not necessarily inherit its launching invocation's literal `argv` – verified empirically instead): against a real, freshly-`init 'd` scratch project (never this repository's own checkout),

the index database's mtime was read before and after editing a tracked file with the wrapper's server running, and found byte-for-byte unchanged after a 10-second settle – proving no auto-sync write occurred.

Six new regression tests landed (scripts/codegraph/tests/test_codegraph_mcp_serve.sh , item- test-first discipline):

T (resolved runner is always the SAME path
fk_index_patch.py --print-bin reports, never independently computed), T (--path + --no-watch are always present, never optional), T (--path targets exactly the resolved project, never cwd/guessed), T – **the incident itself, replayed as a repeatable regression test** (a real init'd scratch project, real serve --mcp launch, real file mutation, real DB-mtime proof of zero writes), T (\$CLAUDE_PROJECT_DIR , when set, is used verbatim, never overridden by git-toplevel guessing), T (fails closed – non-zero exit – when patched-runner resolution cannot proceed, never silently falls back to an unverified binary). All / PASS. T /T /T initially failed on an unrelated wrapper bug (the wrapper hardcoded python3 "\$PATCH_TOOL" , force-feeding a .sh -fixture stub to the Python interpreter in tests – fixed by reusing codegraph_safe.sh 's own run_tool() dispatch, which the production path never exercised since \$PATCH_TOOL defaults to a real .py script there); found and fixed before landing, not after. A second, independent bug (the submodule-vs-parent project-root fallback used a plain git rev-parse --show-toplevel , which from inside a submodule resolves to the SUBMODULE's own root, not the consuming project's – fixed with --show-superproject-working-tree , tried first) was likewise found and fixed during this same investigation, not discovered later.

Honest boundary (\$. .). .mcp.json registration takes effect on a FRESH Claude Code session – this session, already running, cannot self-restart to verify

the `codegraph_explore` / `codegraph_node` tools are genuinely callable end-to-end from inside it; that verification is OWED as a tracked § . . . follow-up on the next fresh session in this project, not claimed here. This extension delivers and tests the MECHANISM (a project-tracked, always-present MCP registration pointing at a verified-correct, incident-hardened, non-writing wrapper) — it does NOT, and cannot, mechanically guarantee that every agent ALWAYS chooses to call it over a raw `grep`/`find`/`glob` on every occasion; that remains a soft property of tool-selection behaviour, same honest limit § . . . (A)'s own accessibility clause already lives with for Lumen. What changed today is that the mechanism itself is no longer an accident of host/plugin configuration — it is now a committed, tested, portable part of this repository, present for any agent, any track, any future clone.

Composes § . . . / § . . . (cleanup on every exit path — every test daemon this investigation spawned was found and killed, none left running) / . . . § . . . / § . . . () (refines the bare-`codegraph` -on-PATH posture for a patched-runner project) / § . . . () (the wrapper is a new tool under the same inherited-by-reference / § . . . -documented / test-first discipline item already mandates) / § . . . (systematic-debugging root-cause-first — the incident was investigated to a proven, evidenced cause before any fix, not patched blind) / § . . . (runtime-signature verification — the heap-budget correction was proven live on the real running process before being called done) / § . . . (decoupled, inherited by reference) / § . . . (the fresh-session verification is tracked, not silently dropped) / § . . . (fail-closed on unresolvable runner state) / § . . . (test-first) / § . . . (process-group signal safety — every

daemon kill in this investigation, manual and automated, validated real `/proc` cmdline identity before signaling, never a bare `pkill -f carrier-match`) / § . .
 (control-needle-proven measurement – the incident's own root cause was confirmed via `/proc/<pid>/cmdline` + daemon registry JSON, never assumed) / § . . (A)
 (universal agent/subagent accessibility – this is its CodeGraph-side completion) / § . .

Classification: universal (per § . .). Applies to every project that is worked on by AI coding CLI agents.

Canonical authority: constitution submodule `Constitution.md`
 § . . .

Non-compliance is a process violation; a project worked on by AI agents without CodeGraph installed, wired, and anti-bluff-verified is in breach of this mandate.

§ . . – Own-org submodules MUST be included in the CodeGraph index (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"All Submodules we use in the project and that are part of organizations to which we have the full access via GitHub, GitLab and other CLIs MUST BE included into the codegraph database and initialized / scanned / synced!"

The mandate. This extends § . . step 's exclude-list refinement with a per-submodule-ownership split. Every consuming project's `.codegraph/config.json` MUST distinguish:

Submodule class	Treatment in CodeGraph index
Own-org – full write access via the project's CLIs (canonical orgs: <code>vasic-digital</code> on GitHub/GitLab/GitFlic/GitVerse + <code>HelixDevelopment</code> on GitHub) 11 4 74	MUST be INCLUDED. Per-submodule sources MUST be indexed so AI coding agents resolve symbols, callers, and impact across the whole repo graph – not just the consuming project's domain code.
Third-party – no write access, vendored only (the <code>\$. . no-match → vendor</code> path, e.g. an external SDK like <code>gopkg.in/telebot.v3</code>)	MUST be EXCLUDED. Indexing third-party code wastes the local index budget and creates symbol-noise the project's contributors cannot fix.

This refines, NOT contradicts, `$ . .` step (which spoke of "other-owned submodules" generically). The classifier is **write-access via the project's own CLIs**, not "internal vs external" subjectively – own-org submodules are the project's own code under another path; third-party submodules are upstream code Herald extends without owning.

Operational steps for every consuming project:

- **Fetch + pull latest of every submodule** before re-indexing – `git submodule update --remote --merge` from the project root, then verify the project still builds. Third-party submodule pins (e.g. `gopkg.in/telebot.v3` pinned at `v3.3.8` to keep the Go import path stable) MUST be respected and rolled back if a `--remote` advances past the load-bearing pin.
11 4 74
- **Adjust `.codegraph/config.json` `exclude` to keep own-org submodules in scope.** Third-party submodule paths MUST be explicitly listed in `exclude` (per `$ . .` step

's per-credential exclusion principle); own-org submodule paths MUST NOT appear in `exclude`.

- **Re-index.** Run the project's canonical CodeGraph initialiser (e.g. `scripts/codegraph_setup.sh`) so `codegraph index` (or `codegraph sync`) rebuilds the index with the corrected exclude list.
- **Verify symbols across own-org submodules.** The project's `scripts/codegraph_validate.sh` (per § 11.4.74 anti-bluff verification) MUST probe at least one symbol that lives ONLY inside an own-org submodule (e.g. `bg.TaskQueue` from `submodules/background/interfaces.go`). A successful query proves the include path actually reached the submodule's sources – not a § 11.4.74 PASS-bluff where `validate` runs against a stale index.
- **Paired § 11.4.74 mutation:** temporarily add the own-org submodule path to `exclude`, re-index, run `validate` – it MUST FAIL on the cross-submodule probe. Restore. This mutation proves the `validate` is not bluffing about the include reach.

Composition. Composes with:

- § 11.4.74 (catalogue-first): the own-org classifier matches the `reuse/extend` catalogue-check disposition; `third-party` matches `no-match → vendor`.
- § 11.4.74 (CodeGraph mandate): this § 11.4.74 refines step 11.4.74's exclude-list contract without weakening any other step.
- § 11.4.74 (credentials never tracked): the exclude list MUST still cover every `.env*`, keystore, signing key, service-account JSON regardless of submodule ownership.
- § 11.4.74 (paired mutation): every `validate` probe added per step 11.4.74 above MUST land with a mutation pair.

- `$` (end-user usability): an indexed graph that lies about which symbols are reachable is a `$` PASS-bluff against the AI coding agents that consume it.

Why this matters. When AI coding agents work across a project that vendors own-org capability modules (e.g. `commons_constitution` consuming the Helix-stack -module surface under `submodules/`), excluding those modules from the index forces the agents back into shallow file-scan exploration of half the codebase – exactly the failure mode `$` . . . exists to prevent. The agents lose call-graph resolution across the seam between domain code and own-org infra code. Including own-org submodules is the difference between agents seeing the project as a unified graph and agents tripping over an invisible boundary at the submodule directory.

EXTENSION – classification is DERIVED, not hand-maintained (research-derived, - -). Measured this cycle: with no generator deriving scope from the submodule graph, dozens of third-party submodule roots (tens of thousands of files) were indexed while tens of thousands of tracked first-party source files were silently dropped by the engine's built-in skip list – neither error surfaced anywhere. Steps - above remain in force; the following clauses are added:

- . **Mechanical classification.** Own-org vs third-party is DERIVED by the constitution's renderer (`$` . . . ()) by walking `.gitmodules` RECURSIVELY (nested gitlinks included) and matching each submodule's remote URL organisation against the consumer's own-org DATA list (`$` . . . ()). Third-party code nested INSIDE an own-org submodule is third-party. A per-submodule override is DATA with a recorded reason, never an unexplained edit.

- **Stale render fails.** Adding, removing or re-pointing a gitlink without re-rendering the scope makes the render stale; the scope guard (§ 11.4.78) MUST FAIL on a stale render – the receipt carries the input hashes it was built from.
- **Steps – re-pointed.** "Adjust .codegraph/config.json exclude " in step 11.4.78 now means: edit the consumer scope DATA file and re-render through the constitution renderer; step 11.4.78's re-index runs ONLY through the single writer entry of § 11.4.78 (). Tracked first-party source that the engine skips by default is re-included by GENERATED negations, never by hand.

Classification: universal (per § 11.4.78). Applies to every project under this Constitution that has CodeGraph configured and own-org submodules vendored.

Canonical authority: constitution submodule Constitution.md § 11.4.78.

Non-compliance is a process violation; an AI-agent-worked project whose own-org submodules are excluded from CodeGraph (when they could be included) is in breach of this mandate. Severe cases (own-org submodules silently excluded WITHOUT an audit trail in .codegraph/config.json comments) are release blockers.

§ 11.4.78 – CodeGraph regular-update + sync automation mandate (User mandate, 11.4.78 – 11.4.78)

Forensic anchor – direct user mandate (verbatim, 11.4.78 – 11.4.78):

"We MUST regularly check for the updates and execute codegraph npm updates so the latest version of it is always installed on the host machine! After update

executes with success use the following commands (and validate and verify them using codegraph help) to sync all: [list of codegraph commands]. Make sure we have proper full automation bash scripts which will run regularly and that these are part of the constitution Submodule since they MUST BE available to all projects which do respect and follow our constitution rules and mandatory constraints! Make sure all updates, sync processes we do and important codegraph related events are all documented under docs/codegraph in Status and Status_Summary documents (another area / context to regularly update and sync) and regularly export them like all other Status docs into the PDF and HTML!"

The mandate. Every consuming project under this Constitution MUST regularly check for CodeGraph npm-package updates, install the latest stable version, and re-sync its local index. The automation lives in the constitution submodule itself so every consuming project gets it via the inherited tree without reimplementation. Three deliverables:

- . `<constitution>/scripts/codegraph_update.sh` – checks the installed `@colbymchenry/codegraph` version against npm's latest, runs `npm update -g @colbymchenry/codegraph` (or `npm install -g @colbymchenry/codegraph@latest`) if a newer version exists, captures the old + new version + timestamp into the constitution's `docs/codegraph/Status.md` audit trail, then exits only when `codegraph --version` produces the expected post-update output. Idempotent: re-runs on the same day are no-ops. Anti-bluff (`$`): the "update succeeded" PASS MUST observe the new version, not just `npm update` exit code.
- . `<constitution>/scripts/codegraph_sync.sh` – after a successful update, runs the canonical sync sequence

inside the consuming project (passed as an argument or auto-resolved via `..find_constitution.sh` parent walk):

```
codegraph status (baseline) → codegraph sync .
(incremental update) → codegraph status (post-sync) →
scripts/codegraph_validate.sh (the project's anti-bluff
verifier, per $ . . step ). Each step's
output is appended to the project's docs/codegraph/
Status.md ledger. The sync MUST consult codegraph help
BEFORE invoking any command to verify the subcommand
still exists with the same shape (tools evolve). The
canonical sync sequence covered by codegraph help :
```

```
codegraph                # Run interactive installer
codegraph install         # Run installer (explicit)
codegraph init [path]     # Initialize in a project (--index to also index)
codegraph uninit [path]   # Remove CodeGraph from a project (--force to skip p
codegraph index [path]    # Full index (--force to re-index, --quiet for less
codegraph sync [path]     # Incremental update
codegraph status [path]   # Show statistics
codegraph query <search>  # Search symbols (--kind, --limit, --json)
codegraph files [path]    # Show file structure (--format, --filter, --max-dep
codegraph context <task>  # Build context for AI (--format, --max-nodes)
codegraph affected [files...] # Find test files affected by changes
codegraph serve --mcp     # Start MCP server
```

- **<constitution>/docs/codegraph/Status.md + docs/codegraph/Status_Summary.md** – append-only ledgers tracking every CodeGraph-related event across the consuming project: install events, npm-update events, sync runs, index regenerations, validate runs (with PASS/FAIL counts), HRDs opened/closed against CodeGraph. Both ledgers MUST be exported to `.html` and `.pdf` siblings on every edit (per \$. . multi-format export mandate). The `Status_Summary.md` is the operator-readable derived index (per \$. . fixed-summary backfill convention).

Operational cadence. The automation MUST run at least weekly (per \$. . status-digest cadence). Consuming projects MAY wire it more frequently (e.g. daily

via cron or per-CI-run) but never less. Cadence is not contractually fixed because environments differ (development workstation vs CI vs production-locked host); the floor is "weekly" because slower than that risks accumulated drift.

Scripts are inherited by reference. Per § submodule inheritance, consuming projects do NOT copy `codegraph_update.sh` / `codegraph_sync.sh` into their own `scripts/` directory. They invoke the constitution submodule's scripts directly (e.g. `bash ${CONST_DIR}/scripts/codegraph_update.sh && bash ${CONST_DIR}/scripts/codegraph_sync.sh .`). This is the single-source-of-truth pattern – when the constitution updates the scripts, every consuming project picks up the change on its next submodule update.

Composition. Composes with § (CodeGraph parent mandate), § (own-org submodule inclusion), § (credential exclusions stay applied), § (status-digest cadence), § (Fixed_Summary backfill), § (multi-format export), § (each sync run's "ok" MUST carry observed-version evidence, not just exit code), § (paired mutation: temporarily downgrade installed codegraph version → script MUST detect drift and self-correct → restore).

Why this matters. A CodeGraph index built against an older version of the indexer may miss symbols that newer versions resolve correctly. Stale indexes silently lie to AI agents about what symbols exist and how they connect. Regular automated updates + sync close the loop: tooling stays current, the index reflects current source, AI agents see the truth. Manual "run codegraph index when I remember" is exactly the failure mode this mandate forbids.

EXTENSION – safe writes, gated upgrades, supervised long runs (research-derived, – –). Measured this cycle (genericised): the update script installs the latest CLI with no compatibility probe, so every upgrade silently re-arms a known bulk-window stall; the sync script was never recorded running and exits when its validator is absent; the vendor CLI deletes a lock by AGE, not by holder liveness, and a writable open recreates deliberately-dropped indexes without taking the lock – so a second writer against a live index corrupts it; and a launcher that classified `sync` by the number of rows ALREADY stored rather than by the size of the PENDING backlog ran a bulk job of tens of millions of references on the stock runner, in the foreground, unsupervised. Deliverables – above remain in force; the following clauses are added:

- **Upgrade gated on a compatibility probe.** An update MUST first probe the new version (the config filename + keys its loader reads, the bulk-window hazard probe, and applicability of every version-keyed runner patch). If the patched runner cannot be re-derived and re-proven for the new version, the upgrade REFUSES loudly and keeps the prior version (`$ . .`) – never a silent fall-back to the stock runner. A runner patch without its own RED→GREEN test, paired mutation and equivalence proof on a byte copy (never on the live database) MUST NOT be applied to a live index (`$ . . / $ .`).
- **Exactly one writer entry.** Every `init` / `index` / `sync` goes through the constitution's single writer entry (`$ . . ( )`), which MUST (i) refuse while ANY live process whose real command line (`$ . . (D)`) targets the same root is running, WHATEVER the lock file's age; (ii) classify a run as BULK by the size of the PENDING work (unresolved

backlog, files still to parse) – NEVER by the size of what is already stored; (iii) route bulk runs only through the proven patched runner, detached, under the watchdog of clause 11.4.232 and a disk tripwire. Readers of a live index open it read-only/immutable and never run full-table counts; experiments run on fixtures or copy-on-write copies.

- **Supervised long runs.** Every long run is supervised by a PROGRESS-PROVEN watchdog (§ 11.4.197 (C)); a process that is alive but not advancing is a STALL → loud stop + tracked item (§ 11.4.164), never "still running".
- **Cadence is real and fail-closed.** The weekly floor above is realised by an actual timer/hook plus the constitution-pull hook (§ 11.4.164), each run leaving a ledger entry; a missing validator is a FAIL, never a warning or a silent exit.
- **Live help is authoritative.** The embedded command list above is illustrative; the authoritative list is a live `codegraph help` diff against the previous run (for example, the current version also exposes a `telemetry` command the list omits). Telemetry MUST be off for the Local-Only posture – whether it is currently off is a per-host FACT to be captured, not assumed.

Classification: universal (per § 11.4.88). Applies to every project under this Constitution that has CodeGraph configured per § 11.4.164.

Canonical authority: constitution submodule `Constitution.md`
§ 11.4.164.

Non-compliance is a process violation; severe cases (consuming project has not run `codegraph_update.sh` in > weeks AND has open AI-agent work) are release blockers –

the agents' index reflects a stale tool version and may produce silently wrong reasoning.

§ . . – Cross-platform-parity mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim, – –):

"Any Linux-only blocker / issue we have MUST BE created macOS and other supported platforms equivalent! So, depending on platform proper implementation will be used for particular OS! EVERYTHING MUST BE PROPERLY EXTENDED AND UPDATED!"

The mandate. Every consuming project whose supported-platforms manifest lists more than one OS MUST, for every feature/test/gate/challenge/mutation that depends on platform-specific primitives, ship a per-OS-equivalent implementation chosen at runtime via `uname -s` (or equivalent platform detection – `[[ $OS == 'Windows_NT' ]]`, `sw_vers -productName`, `/etc/os-release` on Linux distros, etc). A Linux-only gate that ships without macOS / Windows / BSD equivalents (when the kernel of those OSes permits one) is a § . / § PASS-bluff at the platform-coverage layer: the test reports GREEN on Linux but tells nothing about whether the feature works for end users on every supported platform.

Three sub-mandates.

(A) Per-OS implementation REQUIRED. Every feature whose Linux implementation uses a `cgroup` / `systemd` / `/proc` primitive MUST have a documented per-OS equivalent (POSIX `setrlimit` / `ulimit`, macOS `launchd` `plist` `HardResourceLimits`, BSD `rctl`, Windows Job Object, etc)

chosen via runtime platform detection. The wrapper/library code MUST dispatch to the correct branch with NO operator awareness required – the same operator-path invocation works on every supported platform. Canonical example from vasic-digital/tmux : `tmx new -s NAME` dispatches to `systemd-run --user --scope` on Linux and to `scripts/tmx-rlimit-wrapper.sh` (POSIX `setrlimit`) on Darwin – same operator command, OS-correct mechanism per session.

(B) Per-OS tests REQUIRED. Every gate test that exercises platform-dependent behavior MUST have a per-OS branch under `case "$(uname -s)" in` . Each branch MUST exercise the platform's actual native primitive with positive captured evidence per `$ . . + $ . .` (e.g. read `ulimit -t` inside the session via `send-keys` + `capture-pane` on Darwin to verify `RLIMIT_CPU` applied; read `/sys/fs/cgroup/.../memory.max` content on Linux to verify `cgroup MemoryMax` applied). Platform-conditional `SKIP-with-reason` is acceptable ONLY when the platform genuinely cannot enforce the invariant at all (kernel limit – see (C) below), AND the SKIP message MUST cite the precise kernel/system limitation with a reproducer + a link to the project's honest-gap section in its operator guide (typically `docs/guide/README.md` or equivalent).

(C) Honest kernel-gap citation + adjacent equivalent test REQUIRED. Where a Linux primitive has NO macOS / other-OS equivalent due to a documented kernel/OS limitation (canonical example: XNU does NOT enforce `RLIMIT_AS` / `RLIMIT_DATA` / `RLIMIT_RSS` for unprivileged processes – verified by reproducer in vasic-digital/tmux `docs/guide/README.md` `$ .`), the test MUST: () detect the gap at runtime, () `SKIP` with the exact kernel reason + reproducer + link to the project's honest-gap doc, () **provide an ADJACENT test that exercises the closest invariant the platform CAN enforce** (e.g. `RLIMIT_CPU` + `SIGXCPU` as the macOS proxy for "process is bounded under

load"). The adjacent test MUST itself be anti-bluff per § 11.4.1 with positive captured evidence and a paired § 11.4.2 mutation.

Composes with. § 11.4.1 (FAIL-bluffs equally forbidden – a Linux-PASS / Darwin-SKIP-without-honest-reason is a § 11.4.1 violation), § 11.4.2 (recorded-evidence – each branch's captured evidence is platform-specific), § 11.4.3 (per-host-topology dispatch – § 11.4.3 strictens § 11.4.1 by requiring an equivalent test when the platform CAN enforce, not just SKIP-with-reason for genuine kernel gaps), § 11.4.4 (captured evidence quality – platform-specific quality analysis), § 11.4.5 (test-interrupt-on-discovery – discovering a Linux-only test without a Darwin equivalent triggers § 11.4.5 mid-cycle), § 11.4.6 (no-guessing – "platform X probably can't do Y" is a § 11.4.6 violation; prove via reproducer or mark UNCONFIRMED), § 11.4.7 / § 11.4.8 (subagent-driven – cross-platform branches are independent enough to dispatch as parallel subagent work per branch), § 11.4.9 (no-fakes-beyond-unit + % test-type coverage – platform branches inherit the % type-coverage discipline), § 11.4.10 (universal sink-side positive-evidence taxonomy – every platform branch must produce a taxonomy-mapped sink-side evidence artefact), § 11.4.11 (end-user usability – multi-platform projects whose tests cover only one platform deliver no end-user guarantee on the others).

Per-OS implementation equivalence catalogue (canonical examples – not exhaustive).

Linux primitive	Darwin equivalent	BSD equivalent	Windows equivalent
systemd-run --user --scope	scripts/<rlimit-	rctl process:0:deny:vmemory-	Windows Job Object

Linux primitive	Darwin equivalent	BSD equivalent	Windows equivalent
(transient cgroup)	wrapper>.sh using ulimit -t (RLIMIT_CPU) + ulimit -u (RLIMIT_NPROC)	use (root); ulimit -v (unprivileged limited)	via CreateJobObjectW + SetInformationJobObject
cgroup MemoryMax enforcement	XNU GAP – unprivileged RLIMIT_AS returns EINVAL ; root-only via launchd HardResourceLimits.plist. Adjacent test: RLIMIT_CPU + SIGXCPU proxy	rctl process:0:deny:memory-use (root)	Job Object JOB_OBJECT_LIMIT_PROCESS_MEMORY
cgroup TasksMax	RLIMIT_NPROC via ulimit -u (per-user, kernel-enforced)	rctl user:NN:deny:max-proc	Job Object JOB_OBJECT_LIMIT_ACTIVE_PROCESS
/proc/<pid>/oom_score_adj	No equivalent –		None at OS level

Linux primitive	Darwin equivalent	BSD equivalent	Windows equivalent
	Darwin/BSD have no kernel OOM killer (different memory model). SKIP-with-reason; adjacent: rlimit bounds rogue session reach	<code>/dev/oom -style</code> not present in stock BSD	(process governance via Job Object)
<code>/sys/fs/cgroup/...</code> introspection	<code>ulimit -a</code> readback inside session + <code>ps</code> ancestry	<code>procstat -r <pid></code>	<code>QueryInformationJobObject</code>
<code>kill -KILL</code> on cgroup process to test scope isolation	<code>kill -KILL</code> on tmux server pid directly; verify sibling servers survive via direct socket query	same	<code>TerminateProcess</code> + per-job survival check

Operational discipline. When a consuming project adds a new feature whose Linux path lands first, the SAME pull request MUST add the per-OS branches for every other platform in the project's supported-platforms manifest (or explicitly classify the platform per (C) – honest gap + adjacent test). Skipping the per-OS branches "for later" is itself a § 2.0.26 violation: the gap accrues silently and the project ships Linux-PASS / other-OS-untested code while claiming multi-platform support.

Pre-build gate CM-CROSS-PLATFORM-PARITY (when implemented in the consuming project): scans every gate test for case "\$(uname -s)" blocks, asserts a non-SKIP branch (or honest-gap citation per (C)) exists for each platform in the supported-platforms manifest. Paired § 2.0.26 mutation: strip the Darwin branch from any one test → gate FAILs.

Classification: universal (per § 2.0.26). Applies to every project under this Constitution whose supported-platforms manifest lists more than one OS. Single-platform projects are unaffected (no parity gap exists).

Canonical authority: constitution submodule Constitution.md § 2.0.26.

Non-compliance is a release blocker on multi-platform projects. No escape hatch – no --linux-only-acceptable, --no-platform-parity-required, --platform-gap-just-skip flag exists. A Linux-PASS / other-OS-untested feature in a multi-platform project is a § 2.0.26 / § 2.0.26 PASS-bluff at the platform-coverage layer.

2 0 2 6 0 5 2 2

2 0 2 6 0 5 2 2
§ 2.0.26 – Iteration-speedup discipline
mandate (User mandate, – –)

Forensic anchor – direct user mandate (verbatim,
– –):

"How can we speed-up this whole development and fixing process? ... Do not forget to all speed optimizations critical rules and mandatory constraints MUST BE all added into our root (constitution Submodule) Constitution.md, CLAUDE.md, AGENTS.md and QWEN.md and all other relevant constitution Submodules files!"

Iteration cycle time is a first-order quality enabler. A -min rebuild cycle that catches one defect per cycle bounds the project's defect-discovery rate at defect per min; a -min cycle bounds it at $\sim \times$ that rate. Slow cycles are a \$. PASS-bluff at the velocity layer – the project ships fewer validated features per unit of operator time than it should.

The - forensic session that produced this anchor witnessed (a) subagent watchdog stalls each $\sim$ min wasted, (b) stalled `commit_all.sh` $\sim$ min lost on submodule cascade prompt, (c) wrong-call-site speculative kernel patch (\$GQ.) burning a full min rebuild cycle before forensic evidence pointed to the correct site (\$GT Fix B-), (d) stale `git index.lock` from a killed process blocking the next commit, (e) auto- `git gc` maintenance colliding with urgent commit. Each line item cost between and minutes of wall-clock that should have been productive iteration time. The aggregate impact for that single morning was $\sim$ hours of operator wait dilated onto a job whose intrinsic compute was $\sim$ minutes.

The mandate. Every consuming project's build / test / commit / debug pipeline MUST adopt the following speedup disciplines AS MANDATORY. Each is independently enforceable:

(A) Phase forensic before any speculative source patch

Before applying ANY non-trivial source patch (kernel, framework, native code, HAL, ASoC), the agent MUST first complete Phase of `superpowers:systematic-debugging` – read the error/log/call-chain end-to-end and identify a FACT-grade root cause. Speculative patches without Phase evidence are § . . (no-guessing) violations AND § . . violations. A wrong-call-site rebuild cycle wastes – min of compute; the alternative (– min of focused source reading) is a net positive every time. Phase is mandatory regardless of operator time pressure – "we don't have time for forensics, just rebuild" is the § . . anti-pattern.

(B) Live-ADB-First (or live-equivalent) before any rebuild

Per § . . – strengthened by § . . to a release-blocker mandate.. ~ % of audio / UI / boot-script / config changes are LIVE_ADB_TESTABLE (push + `setprop` + `pm install -r` + –second validation cycle). Only kernel / framework Java/AIDL / native C++-in-APEX / sepolicy / init.rc / ro. / Android.bp / XML-resource-overlays / codec-XML-in-APEX changes are genuinely REQUIRES_REBUILD. Skipping the live-probe step for a LIVE_ADB_TESTABLE change is a § . . violation – costs operator min of rebuild for what could have been seconds of `adb push` .

(C) Pre-flight before launching rebuild orchestrators

Before any `nohup bash scripts/<rebuild>.sh &` invocation, the agent MUST run a ~ –second pre-flight that verifies (i) target device(s) reachable via ADB, (ii) downstream sinks (Arvus AVR, HDMI sink, Bluetooth peers) reachable as the

test expects, (iii) host has sufficient memory and disk per \$. budget, (iv) no stale lock files in `.git/` or build directories, (v) no orphan git/build processes from prior killed iterations. A rebuild orchestrator launched against a broken precondition wastes min discovering at the test stage what seconds of pre-flight would have caught.

(D) Persistent build caches outside containers

Containerized AOSP / kernel builds MUST persist `ccache` AND any equivalent build-system intermediate cache (Soong cache, sccache, Gradle daemon state) on the host via bind-mount, NOT inside the ephemeral container layer. Without this, every container restart drops the cache and forces a cold rebuild (~ min for AOSP+kernel on a fast workstation). With it, incremental rebuilds where only one driver / one app changed drop to - min (ccache hit rates - %). This is a one-time setup of ~ min for ~ - min savings per subsequent cycle.

(E) Module-only rebuild for loadable-module-only changes

When a patch touches ONLY one or more files inside `kernel/.../drivers/*/foo.c` AND the affected drivers are built as loadable modules (`CONFIG_*=m` in the active defconfig), the rebuild MUST be `..make -C kernel M=<driver-path>..modules` (~ - min) NOT a full kernel rebuild (~ min). A `flash_kernel_only.sh` or equivalent helper writes the rebuilt module(s) to the running device without re-flashing `system.img`. This saves - min per cycle for the ~ - % of kernel patches that qualify. Built-in drivers (`CONFIG_*=y`) do not qualify – full rebuild required for those.

(F) Parallel multi-device testing

When the project owns more than one validation device ($D_1, D_2, \dots$ in topology terms), the test orchestrator MUST run autonomous validation cycles on EVERY device in parallel via background processes, with separate `qa-results/<TS>/<device-tag>/` output directories per device. This catches per-device topology defects (per § 4.1.1) one cycle earlier instead of waiting for a device-N-focused re-cycle. Direct wall-clock savings = $\frac{1}{N}$ (parallel); amortized savings = $\frac{1}{N} - \frac{1}{N-1}$ min per cycle when a device-specific defect surfaces.

(G) Subagent scope discipline + worktree isolation

Subagents dispatched per `superpowers:subagent-driven-development` MUST be:

- **Scope-bounded** to ≤ 10 min of focused work with intermediate output emitted as the subagent progresses (the canonical watchdog stalls at 15s of no-output; if a subagent writes its findings file incrementally it survives even when blocked on an external dependency).
- **Worktree-isolated by default** via the Agent tool's `isolation: "worktree"` parameter when supported, so parallel subagents do not contend on the same working tree and so a subagent crash never corrupts the conductor's state.
- **Single-responsibility** – one investigation OR one fix-implementation OR one doc draft, not multiple. The empirical pattern from the §GT investigation is that ~ 10 -task subagents stall while ~ 1 -task subagents complete reliably.

(H) Lock-file + stale-process hygiene

Killed or crashed agents leave behind `.git/index.lock`, `.git/.commit_all.lock`, `.lock` files in build directories, and orphan child processes (`git add -A`, `git pack-objects`, `git gc`). The next agent MUST detect these on session start (or after operator-visible crash) and clean them – never wait silently for a never-arriving release. Auto `git-gc` MUST be disabled in repos with concurrent multi-agent work (`git config gc.auto 0`) – its background invocation is a guaranteed lock-contention source.

(I) Cycle telemetry per §11.4.24 build-resource stats

Every iteration cycle MUST record (i) commit hash of the source under test, (ii) wall-clock of build + flash + test phases separately, (iii) the speedup-discipline flag set (which of A-H applied), (iv) outcome (PASS / FAIL / partial / stall). Aggregated weekly, this exposes which speedup gives the biggest empirical return on the specific project, and which is being neglected. The § 11.4.4 build-resource stats tracker is the canonical telemetry pipeline; § 11.4.5 extends its row schema with iteration-discipline columns.

Composes with. § 11.4.23 (test-interrupt-on-discovery + retest), § 11.4.24 (no-guessing – speculative patches forbidden), § 11.4.25 (batch-source-fixes-before-rebuild), § 11.4.26 / § 11.4.27 (subagent-driven), § 11.4.28 (build-resource stats), § 11.4.29 (iteration-discipline conductor loop), § 11.4.30 (TDD-fix discipline – RED-test enables faster cycle reads), § 11.4.31 (deterministic consistency – caches don't break determinism if ccache config disables time-sensitive headers), § 11.4.32 (LIVE-ADB-first maximisation – strengthened to mandate here), § 11.4.33 (autonomous-validation – every cycle's verdict is automation-produced), § 11.4.34

(parallel-development PWU – parallel devices and parallel subagents are § . . instances), § . (-j host-safety cap remains – speedups MUST stay within memory budget), § (end-user usability – fast iteration = fewer cycles → fewer defects shipped).

Pre-build gate. CM-ITERATION-SPEEDUP-DISCIPLINE (when implemented in consuming project): audits the most recent N iteration cycles for telemetry entries citing which of (A)-(I) applied; flags any cycle where none of (A)-(I) was active as a candidate for review. Paired § . meta-test mutation: strip the speedup-flag column from a synthetic telemetry row → gate FAILs.

Operational discipline. When a consuming project ships an iteration speedup (e.g. lands a persistent-cache bind-mount, a parallel-device orchestrator, a module-only rebuild helper), the change MUST be classified as universal (per § . . – most are) OR project-specific (rare – e.g. a script that hardcodes D3 / D4 / Arvus IP 192.168.4.185). Universal speedups belong in this constitution submodule's helper inventory (or in a separate Containers submodule per § . .); project-specific wrappers stay in the consuming repo.

Classification: universal (per § . .). Applies to every project under this Constitution. The disciplines are language-/platform-/build-system-agnostic.

Canonical authority: constitution submodule Constitution.md § . . .

Non-compliance is a release blocker. No escape hatch – no --skip-phase1-forensic , --no-pre-flight , --rebuild-everything-always , --unlimited-subagent-scope , --ignore-locks , --no-telemetry flag exists. The disciplines exist because their absence has been forensically demonstrated to cost the operator hours per session.

§ . . – docs/qa/ end-user evidence
mandate (User mandate, – –)

The mandate. Every feature that ships MUST carry a recorded end-to-end communication transcript plus any attached materials (screenshots, request/response payloads, audio, file uploads) committed under docs/qa/<run-id>/ – one directory per feature run. A feature with no QA transcript is itself a § . /
§ PASS-bluff: it claims to work but carries no auditable runtime evidence that an end user actually exercised it through the same interface they will use in production.

Forensic anchor (verbatim user mandate,
– –):

"every feature that ships MUST carry a recorded e e communication transcript + any attached materials under docs/qa/<run-id>/ (per-feature subdirectories). A feature with no QA transcript is itself a § PASS-bluff – it claims to work but has no auditable runtime evidence. Bot-driven automation MUST preserve full bidirectional communication threads as proof."

Operative rule.

- Every consuming project MUST maintain a docs/qa/ tree (no exception, no escape hatch). Each new feature run lands under docs/qa/<run-id>/ where <run-id> is monotonic + greppable (timestamp, HRD-NNN, ATM-NNN, or other workable-item identifier – see § . .).
- The transcript MUST be full bidirectional – every prompt/command sent + every response received + every error message + every state change observed. One-sided

("we sent X") is not a transcript; both halves are required.

- . Attached materials MUST be committed alongside (screenshots in `.png`, payloads in `.json` / `.toon` / `.txt`, audio in `.wav` / `.mp3`, etc.). External-only links (Slack URL, Drive URL) are § 11.4.12 sink-side violations – the evidence MUST live in-repo.
- . Bot-driven / agent-driven automation (a consuming project's QA bot, e.g. Herald's planned `qaherald`, or analogous binaries in other projects) MUST preserve the full conversation thread as the proof artefact. A bot that runs the round-trip but stores only the final PASS/FAIL line is itself a § 11.4.12 bluff at the QA-automation layer.
- . CI / release gates MUST refuse to tag a version that has any feature-shipping commit without its matching `docs/qa/<run-id>/` directory present and non-empty. The `release` script in the consuming project enforces this; the constitution submodule's `scripts/qa_coverage_audit.sh` (when implemented) gives the universal scanner.

Composes with § 11.4.12 (recorded-evidence requirement – `docs/qa/` is the canonical capture location for end-user-facing features), § 11.4.12 (captured-evidence quality analysis applies per-transcript), § 11.4.12 (sink-side evidence stays in-repo, not behind external links), § 11.4.12 (multi-format export – transcripts may be exported to PDF/HTML siblings when operator-readable), § 11.4.12 (universal sink-side positive-evidence taxonomy – `docs/qa/` is the "user-facing-channel" branch), § 11.4.12 (an end-user feature without an end-user-channel transcript is a § 11.4.12 PASS-bluff), § 11.4.12 (paired mutation: delete a `docs/qa/<run-id>/` → release gate FAILs).

Classification: universal (per § 11.4.12). Every project under this Constitution carries this rule. Single-

feature pre-implementation projects are exempt until their first user-visible feature lands.

Canonical authority: constitution submodule `Constitution.md`

\$. . .

Non-compliance is a release blocker. No `--qa-evidence-optional`, `--qa-transcript-later`, `--qa-bot-summary-suffices` flag exists.

2026 05 22

\$. . . – Working-tree quiescence rule for subagent commits (User mandate, – –)

Short tag: `working-tree quiescence` .

The mandate. No subagent commit may proceed while any concurrent mutation gate, paired-mutation experiment, or other in-flight mutation is live in the same checkout. Before `git add`, the committing agent MUST grep its own working tree for mutation markers (`MUTATED` for paired , `//` always pass , `return json.Marshal` shortcut paths, `//` `MUTATION` / `# MUTATION` annotations, `_mutated_*` filename suffixes, etc.) and explicitly account for every modified file in the staging area. Any unexplained file in the staging area triggers ABORT.

Forensic anchor (verbatim user mandate,
– –):

"no subagent commit may proceed while any concurrent mutation gate is in flight in the same checkout. Before `git add`, the committing agent MUST grep its own working tree for mutation markers (`MUTATED` for paired , `//` always pass , `return json.Marshal` shortcut

paths, etc.). Any unexplained file in the staging area triggers ABORT."

Lesson (forensic case study). A consuming project's logo-fix subagent (Herald commit `72e81ab`, `- -`) ran in a checkout where a paired `$ .` mutation gate had temporarily introduced an `// always pass` shortcut into `commons_auth/middleware.go` (JWT-bypass mutation, intended to be reverted by the same gate). The subagent's `git add + git commit` swept the mutation residue into the same commit as the unrelated logo fix, and the resulting commit was pushed to all four mirrors before any other agent caught it. The fix (Herald `d5bd360`, "SECURITY FIX: restore commons_auth/middleware.go JWT verify") landed within the hour, but the window during which production-equivalent binaries shipped with a bypassed JWT verify is a real security-defect window – small but non-zero – and is the canonical example of why this rule exists.

Operative rule.

- . Subagent (and main-thread) commit flows MUST run a pre-`git add` quiescence check that:
 - Greps the working tree for mutation markers (canonical list above; the consuming project MAY extend with project-specific markers).
 - Lists every modified / untracked / staged file and matches each one against the subagent's declared scope. Any file outside the declared scope → ABORT.
 - Cross-checks `git status --porcelain` against the subagent's task description; unaccounted entries → ABORT.
- . Any active mutation gate (paired-mutation experiment in progress, AVR/AVI mutation cycle, manual `// always pass` patch, etc.) MUST be serialised – the gate runs to

completion (mutate → assert FAIL → restore → assert PASS) and the working tree is verifiably clean BEFORE any unrelated commit may proceed in the same checkout.

- . Concurrent subagents working in the SAME checkout (single working tree, no `git worktree add`) MUST coordinate through a lockfile or marker file (e.g. `.git/MUTATION_IN_PROGRESS`) so a logo-fix subagent cannot race a JWT-mutation subagent. The constitution submodule's `scripts/quiescence_check.sh` (when implemented) is the universal lock helper.
- . When parallel work is necessary, the `$ . . . / $ . . .` subagent-driven mandate SHOULD be combined with `git worktree add` to give each subagent an isolated working tree – eliminates the cross-mutation race by construction.
- . Post-commit, an automated `mutation-residue-scanner` (consuming project's `scripts/mutation_residue_audit.sh`) MUST run before push. Any commit containing a mutation marker → push BLOCKED, commit MUST be reverted or amended before mirrors are updated.

Composes with `$ . . .` (mutation-paired gates – quiescence rule protects the mutation cycle from concurrent contamination), `$ . . . / $ . . .` (subagent-driven-by-default – quiescence rule is what makes parallel subagent dispatch safe), `$ . . .` (no-fakes-beyond-unit – a mutation residue swept into a commit IS a fake-pass surface in production), `$ . . .` (credentials handling – same class of "do not let unrelated content leak into a commit"), `$ . . .` (a security-bypass mutation that ships to production is the gravest `$ . . .` PASS-bluff: tests pass against the mutated codebase, but the end user loses an entire security property), `$ . . .` (pre-push fetch + investigate – quiescence check is the pre-push companion).

Classification: universal (per § 11.4.24). Every project under this Constitution carries this rule.

Canonical authority: constitution submodule `Constitution.md`
§ 11.4.24.

Non-compliance is a release blocker. No `--allow-residue`, `--skip-quiescence`, `--mutation-cleanup-later` flag exists. A mutation marker that lands in a tagged commit is a critical defect regardless of how briefly it persisted.

2026 05 24

§ 11.4.24 – Stress + Chaos Test Mandate (User mandate, – –)

Short tag: `stress-chaos-mandate`.

Forensic anchor (verbatim user mandate, – –):

"Every fix or improvement you do MUST BE covered with full automation stress and chaos tests so we are sure nothing can break the functionality and all edge cases are monitored and polished and additionally fixed if that is needed! Everything must produce rock solid proofs and follow fully no-bluff policy!"

The mandate. Every fix or improvement landed in a consuming project MUST ship with full-automation **stress** AND **chaos** test suites that exercise edge cases, sustained load, concurrent contention, and failure-injection. A fix that **PASSes** its happy-path test but has never been exercised under stress or under fault-injection is a § 11.4.24 / § 11.4.24 **PASS-bluff** at the resilience layer: it claims to work but carries no evidence that real-world adversarial conditions (sustained throughput, parallel contention, partial failure, resource exhaustion,

malformed input) leave the fix intact and the user-visible behaviour correct.

Definitions (closed-set, mechanically auditable):

- **Stress test** – exercises the fix-under-test under sustained or concurrent load above ordinary usage. At MINIMUM one of:
 - **Sustained load** – $N \geq$ sequential iterations OR $\geq$ seconds wall-clock continuous load. Per-iteration latency MUST be recorded; percentile distribution (p_{50} / p_{90} / p_{99}) MUST be reported.
 - **Concurrent contention** – $N \geq$ parallel invocations. All N MUST complete. No deadlock, no resource leak (file-descriptor count, process-table count, RSS), no data race in shared state.
 - **Boundary conditions** – minimum input (empty), maximum input (provider/codec/protocol-limit-bound), boundary input (off-by-one at every size threshold). Each boundary MUST produce a categorised result (success, categorised-error, deterministic-skip) – NEVER an uncaught exception or silent corruption.
- **Chaos test** – exercises the fix-under-test under failure-injection per the § 4.2. closed-set evidence taxonomy. At MINIMUM one chaos category appropriate to the fix-class:
 - **Process-death injection** – kill the primary process, dependency process, or upstream service mid-operation. Recovery path MUST be deterministic + categorised (clean error, automatic restart, circuit-breaker open).
 - **Network-fault injection** – drop / delay / reorder packets on the relevant interface mid-call. Categorised error per § 4.2. (e.g. category=network / category=upstream).

- **Input-corruption injection** – corrupt the input file / config / .env mid-test. Test MUST detect + report the corruption – NEVER silently consume corrupted input.
- **Resource-exhaustion injection** – fill disk to %, allocate memory pressure (OOM-injection), exhaust file descriptors, exhaust sockets. Fix MUST refuse new work cleanly OR degrade gracefully – NEVER crash with uncaught error.
- **State-corruption injection** – mid-flight database lock loss, mid-flight file-system partial-write error, mid-flight cache invalidation. Recovery path MUST restore consistent state.

Anti-bluff (mandatory):

- . Every stress + chaos test PASS MUST cite a captured-evidence artefact path per \$. . + \$ Acceptable evidence: per-iteration latency.json / throughput.csv / categorised_errors.txt / state_delta_snapshot.json / process_lifecycle.log / recovery_trace.log . Metadata-only PASS ("all iterations exited ") without per-iteration latency / failure / state evidence is itself a \$. PASS-bluff at the stress-layer.
- . Helper library: every consuming project SHOULD ship a stress_chaos.sh (or analogous language-binding library) exposing reusable primitives – ab_stress_run , ab_stress_concurrent , ab_chaos_kill_pid_during , ab_chaos_drop_network_during , ab_chaos_corrupt_file_during , ab_chaos_oom_pressure_during , ab_chaos_disk_full_during . Each helper composes with ab_pass_with_evidence / ab_skip_with_reason per \$
- . Chaos-injection cleanup is non-negotiable. A test that corrupts a .env MUST restore it in trap '...' EXIT . A test that fills the disk MUST rm the filler in EXIT. A test that kills a process MUST verify the process is

restarted (or explicitly leave it down with operator-visible reason). Cleanup failure = § 11.4.4. (test-playback cleanup) violation.

-layer coverage per § 11.4.6.7 (b):

- **Pre-build gate** – stress + chaos test files exist + executable + parseable under `sh -n + bash -n` per § 11.4.6.7.1; helper library exists + executable; the fix's pre-build gate cites the stress + chaos test file path.
- **Paired meta-test mutation** – per § 11.4.6.7.2, removing the chaos-injection step or stripping the per-iteration evidence-capture → the gate FAILs.
- **On-device test** (if the fix is LIVE_ADB_TESTABLE per § 11.4.6.7.3) – the stress + chaos test is dispatched against a real device, captured-evidence directory under `qa-results/<run-id>/stress_chaos/`.
- **HelixQA Challenge entry** (if the fix is a user-visible feature per § 11.4.6.7.4 (b) layer 1) – a Challenge bank entry references the same stress + chaos test as the validation route.

Composition with related § 11.4.6.8 anchors:

- § 11.4.6.8.1 / § 11.4.6.8.2 – stress + chaos test PASS is the strongest end-user-quality signal: it demonstrates the fix survives adversarial conditions the happy-path test cannot reach.
- § 11.4.6.8.3 – FAIL-bluffs forbidden. A stress test that crashes with `set -u` because the chaos-injection step left a variable unset is a FAIL-bluff; fix at source.
- § 11.4.6.8.4 – captured-evidence content quality applies to stress + chaos evidence too (latency distribution recorded, not just count; error categories enumerated, not just "some errors").

- § 11.4.49 . . . – no guessing. Categorised errors per closed-set. "Probably the network" is forbidden – capture the categorisation.
- § 11.4.50 . . . – TDD RED-first. The stress + chaos suite SHOULD start RED: write the test, observe the failure under load/chaos, land the fix, observe the GREEN.
- § 11.4.52 . . . – deterministic consistency. The stress test's N iterations MUST produce identical exit codes AND identical evidence-hashes (volatile prefixes stripped) per § 11.4.69
- § 11.4.69 . . . – autonomous validation. Stress + chaos tests MUST run end-to-end without operator presence.
- § 11.4.80 . . . – universal sink-side positive-evidence taxonomy. Every stress + chaos PASS uses `ab_pass_with_evidence` with a real captured-evidence file path.
- § 11.4.17 . . . – `docs/qa/` end-user transcript discipline composes with chaos-test evidence (chaos-test recovery transcripts ARE end-user-channel proofs).

Classification: universal (per § 11.4.81 . . .). Every consuming project carries this rule. Stress + chaos discipline is language-/platform-/build-system-agnostic.

Canonical authority: constitution submodule `Constitution.md`
 § 11.4.81

Non-compliance is a release blocker regardless of context.
 No escape hatch – no `--skip-stress`, `--no-chaos`, `--happy-path-suffices`, `--stress-test-later` flag exists. The discipline exists because the User mandate is unambiguous: "Everything must produce rock solid proofs and follow fully no-bluff policy!" – and a fix never tested under stress + chaos is not a rock-solid proof.

`$ . . . – Roster/corpus-backed Status-doc auto-sync mandate (User mandate, – – )`

Forensic anchor – verbatim user mandate

`( – – ):`

"Make sure that assets and players Status docs are ALWAYS regularly updated and in sync like all others Status docs – any time we add or modify the assets content(s) or we change or add new / remove existing pre-installed video and audio player apps! This MUST WORK OUT OF THE BOX!"

Some Status docs (`$ . . .`) are backed not by hand-written narrative alone but by a **tracked roster** (a set of installed apps / components) or a **tracked asset corpus** (a directory of test / media assets). For these, freshness cannot be left to operator vigilance – the moment a member of the roster/corpus changes (a player app added / removed / renamed; a test asset added / modified / removed) the Status doc, its Status_Summary, and their HTML + PDF exports MUST be brought back into sync **out of the box**, mechanically, not by remembering to.

The mandate. Every Status doc whose subject is a tracked roster or asset corpus MUST be kept in sync with that subject by a mechanism where **all** of the following hold:

- **Drift-proof signal.** A content *fingerprint* – sha of the sorted member list, NOT mtime (which `git checkout` resets) – is persisted in a sidecar next to the Status doc. Adding / removing / modifying a member changes the fingerprint deterministically.
- **Sync helper.** A helper regenerates the fingerprint and re-exports HTML + PDF (composing with the `$ . . .` universal exporter). Invocable standalone AND wired so

the sync happens automatically (commit path and/or pre-build).

- **Enforcement gate.** A pre-build gate recomputes the live fingerprint and FAILs when it differs from the persisted one – forcing the Status doc to be updated whenever the roster/corpus changes. Mirrors § 11.4.12 CM-ISSUES-SUMMARY-SYNC + § 11.4.12 sync_integration_status .
- **Paired § 11.4.12 mutation.** A meta-test mutation corrupts the fingerprint (introducing drift without a doc update) and asserts the gate FAILs – proving the gate is not itself a bluff. § 11.4.15

Composes with § 11.4.12 (auto-generated docs sync – the canonical sibling pattern), § 11.4.19 (integration-status-doc – § 11.4.19 is its roster/corpus specialisation), § 11.4.19 + § 11.4.19 (Summary parity), § 11.4.19 + § 11.4.19 (README doc-link freshness), § 11.4.19 (composite doc sync), § 11.4.17 (universal Markdown export), § 11.4.17 (no-guessing – the fingerprint is FACT, not a guess about freshness).

Classification: universal (per § 11.4.19) – the principle (roster/corpus-backed Status docs auto-sync on subject change) is project-agnostic; the consuming project supplies the specific docs, roster/corpus sources, helper, and gate name in its own CLAUDE.md / AGENTS.md / QWEN.md per § 11.4.19.

Canonical authority: constitution submodule Constitution.md § 11.4.19.

Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-roster-sync , --allow-status-drift , --roster-sync-not-applicable flag exists.

§ . . – Endless-loop autonomous work
 + zero-idle agent dispatch + anti-bluff
 testing mandate (User mandate,
 – –)

Forensic anchor – verbatim User mandate

(– –):

"when we say to you that all work MUST BE continued in endless loop until there is no any open items, no unfinished workable items from our Issues docs, or from Continuation document or any unfinished work by agents at the moment it means that you will ALWAYS run until there is nothing workable elft to be done in endless loop, fullyautonomously. You will spawn agents or agents-driven work whenever that is possible or required! Not a single agent or main work stream will sit idle except if it waits for the results of something – some execution and similar! All work MUST BE always covered with comprehensive tests – all supported tests for the project which produce real proofs for the contexts and functionalities they are testing! Bluff of any kind is not allowed and all work and tests MUST WORK completely as proofs (phisical) driven and in complete anti-bluff manner!"

When an operator instructs an AI agent to "continue in endless loop fully autonomously" – or any semantically-equivalent phrasing – the agent MUST interpret this as a HARD-CONTRACT execution covenant covering five non-negotiable obligations:

(A) Endless-loop continuation. The agent MUST continue working in the autonomous-loop framework until ALL of the following are simultaneously TRUE:

- docs/Issues.md Status-column has zero In progress / Ready for testing / In testing / Reopened entries (\$. . closed-set values).
- docs/CONTINUATION.md \$ "Active work" section is empty.
- No background subagent is mid-execution (TaskList reports all tasks completed or pending with pending items NOT being prerequisite-blockers for anything actionable).
- No external dependency is in-flight (build, flash, push, sync).

If any of those conditions is FALSE, the agent MUST continue working – by claiming the next item from the priority queue, by dispatching a subagent for parallelisable work, or by polling a background task that will be the wake signal. An agent that schedules a wake-up and ends the turn while Issues.md still has actionable items, OR while CONTINUATION.md \$ still has active work, OR while a subagent is still working productively, is in violation.

(B) Zero-idle agent dispatch. When the agent identifies parallelisable work that does not contend on the same file-scope as the main work stream, the agent MUST dispatch a background subagent rather than serialising the work. The main work stream + every background subagent operate concurrently. "Wait for the results of something" is the ONLY acceptable reason for an agent (main or subagent) to be idle. The \$. . + \$. . + \$. . subagent-driven covenants compose with this anchor.

(C) Comprehensive test coverage with real (physical) proofs. Every workable item closed in the loop MUST land four-layer test coverage per § 11.4.4 (b) – pre-build gate + on-device test + paired § 11.4.4 meta-test mutation + HelixQA Challenge bank entry – and every test PASS MUST cite a captured-evidence artefact per § 11.4.4 + § 11.4.4. The phrase "physical proofs" is interpreted as: captured audio (tinycap WAV with RMS + ffprobe channels), captured video (screen recording + ffprobe frames + recording-analyzer event match), captured network state (dumpsys + sink-side probe per § 11.4.4), captured UI state (uiautomator dump per § 11.4.4), captured sysfs state (/sys/... / /proc/... snapshots). Metadata-only PASS, configuration-only PASS, absence-of-error PASS, grep-without-runtime PASS are all critical defects regardless of the green summary line.

(D) Anti-bluff guarantee end-to-end. The § 11.4.4 covenant family (§ 11.4.4 FAIL-bluffs / § 11.4.4 recorded-evidence / § 11.4.4 no-guessing / § 11.4.4 demotion-evidence / § 11.4.4 no-fakes-beyond-unit / § 11.4.4 deterministic-consistency / § 11.4.4 autonomous-validation / § 11.4.4 audio sink-side / § 11.4.4 universal sink-side / § 11.4.4 docs/qa/ end-user evidence) is the operative truth-discipline of the endless-loop covenant. Every closure narrative MUST cite captured-evidence; every demotion MUST cite same-conditions retest; every status transition MUST cite the deterministic-consistency baseline. **Tests AND Challenges (HelixQA) are bound equally** – a Challenge that scores PASS on a non-functional feature is the same class of defect as a unit test that does. The User mandate explicitly invokes the historical forensic anchor: "we had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can't be used! This MUST NOT be the case and execution of tests and Challenges MUST guarantee the

quality, the completion and full usability by end users of the product!"

(E) Termination criteria. The endless-loop terminates ONLY when:

- . ALL conditions in (A) are simultaneously TRUE (zero open items + zero active CONTINUATION work + zero subagents in flight + zero external dependencies in-flight); OR
- . The operator explicitly issues a `STOP` / `END LOOP` / `pause autonomous work` instruction; OR
- . The host-session-safety covenant (`$` + `$` . + `$` . + `$` . + `$` .) demands suspension to protect the operator's session; OR
- . The agent has scheduled a wake-up to resume against a known-future-actionable signal (e.g. CI completion, build completion, hardware-attended phase).

No escape hatch. No `--idle-OK` , `--skip-endless-loop` , `--bluff-permitted-for-this-task` , `--metadata-only-test-suffices` , `--no-physical-proof-required` flag exists. The covenant is uniform across every iteration.

Composes with `$` (anti-bluff covenant – historical forensic anchor of the project), `$` (FAIL-bluffs), `$` (recorded-evidence), `$` (four-layer coverage), `$` (audio + video – layer quality), `$` (no-guessing), `$` (demotion-evidence), `$` + `$` + `$` (subagent-driven defaults), `$` (no-fakes-beyond-unit + % test-type coverage), `$` (iteration-discipline), `$` (TDD-fix), `$` (deterministic-consistency), `$` (autonomous-validation), `$` (audio sink-side), `$` (universal sink-side), `$` (docs/qa/ end-user evidence), `$` (stress + chaos), `$` (roster/ corpus auto-sync), `$` family (host-session safety

boundaries), § . (CONTINUATION-document maintenance – the source-of-truth state the loop checks against).

Pre-build gate CM-COVENANT-114-87-PROPAGATION enforces this anchor literal in every CLAUDE.md / AGENTS.md / QWEN.md across the canonical fleet (parent + consumer submodules). Paired meta-test mutation strips the literal 11.4.87 → gate FAILs.

Canonical authority: this Constitution.md § . . in the HelixConstitution submodule (git@github.com:HelixDevelopment/HelixConstitution.git) – inherited per § . . by every consuming project's repo-root CLAUDE.md / AGENTS.md / QWEN.md.

Non-compliance is a release blocker regardless of context. Skipping the loop, idling without a waited-on signal, or accepting bluff-evidence PASS is severity-equivalent to a § . PASS-bluff at the project-execution layer.

§ . . – **Background-push mandate:**
commit-lock release immediately after commit,
push runs detached (User mandate,
– –)

Forensic anchor – verbatim User mandate
(– –):

"Make sure all these pushes are being done ALWAYS in background in parallel with main work stream so we do not lose time waiting. Everything is committed anyway? If we can do this add this as mandatory rule / constraint that we must respect and follow ALWAYS! ... We MUST ensure that main work stream has always something to do, or wait for the results only when that is absolutely required! We have enormous amount

of work in front of us, short deadlines and we MUST DELIVER!"

Forensic incident. - - between
: Z and $\geq$: Z (~ hours observed), a
single `commit_all.sh` invocation held its
`.git/.commit_all.lock` flock the entire time – the COMMIT
had succeeded within seconds at `63d307234ce`, but the
synchronous `do_push` call kept the flock while `git push`
`vasicdigital_gitlab` did a large initial-mirror upload. **Every**
subsequent `commit_all.sh` invocation in that window was
BLOCKED – .. submodule pointer bumps + parent
\$. . propagation work + \$KA Critical audio
quality matrix artifacts all sat staged in the working
tree waiting for the push to finish. This is the exact
anti-pattern \$. . (zero-idle) prohibits at the
project-execution layer.

The mandate.

(A) Flock release IMMEDIATELY after commit lands. Once `git commit` returns (commit object recorded in the local repo), the `.git/.commit_all.lock` MUST be released BEFORE `do_push` is invoked. The commit is durable on local disk regardless of push outcome – there is no correctness reason to gate further local work on a remote-push round-trip.

(B) Push runs detached. After flock release, the push step is spawned via `nohup ./push_all.sh .... > <log> 2>&1 &` then `disown` so the orchestrator's exit does not propagate SIGTERM. The orchestrator's exit code reports COMMIT success, NOT push success (push success is logged separately and reconciled by \$. . fetch-before-edit on the next interactive turn).

(C) Push concurrency control: per-remote serialization, multi-remote parallelism. `push_all.sh` MUST acquire a per-

remote flock (`.git/.push.<remote>.lock`) so two concurrent invocations targeting the same remote serialize (avoid push-race / non-fast-forward), but invocations targeting DIFFERENT remotes run in parallel. This converts the prior -hour serial-stall pattern into a multi-stream pipeline where the operator's main work stream never blocks on the slowest mirror.

(D) Failure surface. Backgrounded push failures land in `qa-results/push_failures/<timestamp>_<remote>.log` . The next interactive autonomous-loop tick MUST check that directory (per `$ . . . (A) "no external dependency is in-flight" check`) and surface failures – re-push attempted automatically, or operator notified if auth/quota issue. Silent push-failure is a `$ . . . PASS-bluff` at the distribution layer.

(E) Synchronous-push escape. An explicit `--sync-push` CLI flag on `commit_all.sh` preserves the legacy synchronous behaviour for `--force-with-lease` paths (per `$ . . . force-push merge-first audit`) where the push outcome MUST gate the next step. This is the ONLY escape – without it, every push is backgrounded.

Composes with `$ . . .` (multi-upstream push norm – `$ . . .` makes it non-blocking), `$ . . .` (data safety – backgrounded push still uses hardlinked-backup pre-op), `$ . . .` (force-push merge-first – uses `--sync-push` escape), `$ . . .` (iteration discipline – endless-loop ticks no longer waste budget on remote-network latency), `$ . . .` (fetch-before-push merge-first – applies before background push spawns), `$ . . .` (endless-loop / zero-idle – `$ . . .` is the implementation seam that makes `$ . . .` (B) genuinely zero-idle for the dominant blocker class).

Pre-build gate `CM-COVENANT-114-88-PROPAGATION` enforces this anchor literal in every `CLAUDE.md` / `AGENTS.md` / `QWEN.md`

across the canonical fleet. Pre-build gate CM-BACKGROUND-PUSH-WIRED verifies commit_all.sh releases flock before do_push AND spawns push via nohup ... & AND push_all.sh acquires per-remote flock. Paired \$. meta-test mutations strip the load-bearing literals → gates FAIL.

Canonical authority: this Constitution.md \$. . in the HelixConstitution submodule (git@github.com:HelixDevelopment/HelixConstitution.git).

Non-compliance is a release blocker. Synchronous push (without --sync-push flag) is severity-equivalent to a \$. PASS-bluff at the project-execution layer. No escape hatch beyond --sync-push for \$. . force-push events.

\$. . – Background test execution mandate (User mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"Any tests we are executing, especially long test cycles, MUST BE performed in background in parallel with main work stream! This MUST NOT block our capabilities to work on queued workable items (and tackle main Issues document(s)). Make sure we strictly follow this starting now! Main work stream can be blocked or sit iddle only if absolutely needed and if it depends hard on results of some background execution."

Forensic incident. – – the conductor invoked bash device/rockchip/rk3588/tests/pre_build_verification.sh synchronously with a timeout 360

foreground wrapper – even though the operator had previously mandated \$. . . endless-loop zero-idle. The - minute test run blocked the entire main work stream from progressing on \$JV/\$JW/\$JX/\$JY scaffolding + further Issues.md drains. This is the exact anti-pattern \$. . . (B) prohibits at the test-execution layer (analogous to \$. . . at the push-execution layer).

The mandate.

(A) Long-running tests run detached. Any test cycle expected to exceed ~ seconds –
pre_build_verification.sh , meta_test_false_positive_proof.sh ,
test_all_fixes.sh , recent_work_validate.sh , HelixQA
Challenge banks, on-device -phase cycles, full-suite retests per \$. . . , audio_loop_supervisor.sh,
dual_display_record.sh harnesses, verify-all-constitution-rules.sh – MUST be spawned via nohup ... > <log> 2>&1 & +
disown and <log> placed under a known directory (e.g. qa-results/<test_id>_<timestamp>.log) for later inspection.

(B) Main stream proceeds on queued items. The conductor MUST immediately return to the \$. . . priority queue (open Issues.md items, queued PWUs, doc-sync work, audit work, anything not depending hard on the in-flight test's exit code). "Wait for results" is the ONLY acceptable idle reason per \$. . . (B) – and even then, the conductor checks back on the test result via short polling (tail / mtime / exit-status file) rather than blocking on wait .

(C) Hard-dependency gating. A subsequent step that genuinely requires the test's exit code (e.g. commit_all.sh refuses if meta-test is in flight per \$. . . ; release-tag gate per \$. . . step) MUST express that dependency explicitly – poll the exit-status file or check pgrep -af <test-name> before proceeding. If the test is still running and the next step

needs its result, surface that to the operator per
\$. . . with interactive options.

(D) Failure surface + cleanup. Backgrounded test failures land in their `<log>` files. The next autonomous-loop tick MUST check that directory (analogue of \$. . . (D)). Test-script-level cleanup (per \$. . . + \$. . . + this anchor's restore discipline) is non-negotiable – `trap '<cleanup>' EXIT INT TERM` in every long-running test that mutates global state (mid-mutation backup files, sysfs nodes, ALSA states, etc.).

(E) Pre-empt the blocking pattern. The conductor MUST NOT invoke `bash <long-test>.sh` synchronously when a backgrounded variant suffices. Foreground execution is permitted ONLY when: (a) the test completes in < s, OR (b) the operator explicitly requests foreground execution for a specific step. Synchronous execution of a multi-minute test absent operator authorisation is a \$. . . violation regardless of how much "context-cache benefit" the foreground path appears to offer.

(F) Concurrency control. Multiple backgrounded tests of the SAME script MUST serialise via per-script flock (e.g. `.git/.pre_build.lock`, `.git/.meta_test.lock`). Concurrent invocations of DIFFERENT tests run in parallel.

Composes with \$. . . (iteration discipline – backgrounding tests recovers the operator's wait-time per \$. . . (F) parallel multi-device testing),
\$. . . (interactive clarification – surface hard-dependencies as options instead of blocking),
\$. . . (iteration-speedup discipline – \$. . . is the dominant blocker class at the test layer, analogous to \$. . . at the push layer),
\$. . . (working-tree quiescence – backgrounded mutation tests still need quiescence checks),
\$. . . (stress + chaos – long stress runs are

quintessential § . . . candidates), § . . .
(endless-loop zero-idle – § . . . is the
implementation seam at the test-execution layer that makes
§ . . . (B) genuinely zero-idle for the test blocker
class), § . . . (background-push – § . . .
is the symmetric anchor for tests).

Pre-build gate CM-COVENANT-114-89-PROPAGATION enforces this
anchor literal across the canonical fleet. Pre-build gate
CM-BACKGROUND-TEST-EXECUTION-WIRED verifies that long-running
test invocations in helper scripts use `nohup ... & +`
`disown` pattern (or equivalent). Paired § . . . meta-test
mutations strip the load-bearing literals → gates FAIL.

**(G) Subagent long-op discipline + conductor-owns-
inherently-long-ops (extension, - -).**
Forensic FACT (- -): background subagents
were repeatedly KILLED by the runtime's no-progress stream
watchdog (~ s with no tool-call output) because
each had run ONE synchronous long operation inside a
single tool call (a full pre-build sweep / a long test
suite / a device flash / a hung network call) – the host
and the network were ruled out as causes; the watchdog
fired on the SILENCE, not on any failure. Therefore: ()
a SUBAGENT MUST NOT run any single synchronous operation
expected to exceed the runtime's no-progress watchdog
budget (conservatively ~ min) – it MUST background the
operation per clause (A) (`nohup ... & + disown`) and POLL
with short, frequent tool calls (each poll is the progress
the watchdog observes); () inherently-long operations
that a subagent CANNOT usefully background-and-poll (full
artifact builds, full test suites, device flashes /
deploys) are the CONDUCTOR's to run – the conductor (main
session) carries no such watchdog and MUST own them,
dispatching subagents only for the parallelisable slices
around them; () a subagent killed by the stall watchdog
is a CRASH per § . . . – its work is NOT lost,

NOT done, NOT forgotten: the registry entry flips to crashed and the unit is RESPAWNED until complete ; a stall-kill silently absorbed as a completion is a \$. . + \$. bluff. Honest boundary (\$. .): the watchdog budget is a RUNTIME property – read it from the runtime, never guess it; when it is unknown, treat any operation > ~ min as long.

Canonical authority: this Constitution.md \$. . in the HelixConstitution submodule (git@github.com:HelixDevelopment/HelixConstitution.git).

Non-compliance is a release blocker. Synchronous long-test execution (without operator authorisation) is severity-equivalent to a \$. PASS-bluff at the project-execution layer. No escape hatch beyond explicit per-invocation operator authorisation.

\$. . – Obsolete status + per-item
obsolescence audit mandate (User mandate,
– –)

Forensic anchor – verbatim user mandate
(– –):

"Regarding the Critical Bug No – Albums cover etc. – Now it seems obsolete after latest request for new behavior when audio and video content is being played. If we do not have Obsolete as an option / status for the ticket in its resolving, we should introduce this now and mark obsolete tickets with some light gray background or something that will indicate that item is no longer valid. Maybe text – the description to be strikethrough styled as well! Once all this is done please review all existing open or resolved workable items (fixed as well) if they are

obsolete - not valid any more. ... There MUST NOT be any mistake! No bluff is allowed of any kind!"

The § 4.1.1 Status closed-set is extended with a 4th terminal value `Obsolete (→ Fixed.md)` matching the existing 4-element § 4.1.1 closure-vocabulary table (`Bug → Fixed` , `Feature → Implemented` , `Task → Completed`). The `Obsolete` value applies regardless of Type when the workable item is **no longer valid** because: (a) the underlying behaviour the item described has been superseded by a later operator mandate / design pivot (canonical example: `Bug` Albums cover, superseded by §JU + the - - secondary-display video-playback redesign), (b) the affected feature was REMOVED from the product, (c) the item was duplicated and the canonical entry resolved it, (d) the item was filed against a hardware/topology variant we no longer support.

Mandatory obsolescence audit metadata – every `Obsolete (→ Fixed.md)` heading MUST carry, within 4 non-blank lines of the heading, an `**Obsolete-Details:**` line capturing four sub-facts (mirroring § 4.1.1 + § 4.1.2 audit-line patterns):

- **Since:** ISO date (`YYYY-MM-DD`).
- **Reason:** one-line cause from the closed vocabulary { `superseded-by-design-change` | `superseded-by-later-mandate` | `feature-removed` | `duplicate-of` | `unsupported-topology` | `not-reproducible` } . (`not-reproducible` = a reported defect that does NOT reproduce on the canonical tree / baseline – an environment / isolated-worktree artifact (PATH / shell / missing-checkout), not a real product defect; the triple-check evidence MUST capture the canonical-tree non-reproduction, per § 4.1.1 no-guessing + § 4.1.2 demotion-evidence.)

- **Superseding-item:** §-letter reference or User-mandate verbatim quote anchor citing the work that obsoleted it.
- **Triple-check evidence:** path to captured evidence (git log, code grep, runtime behaviour) confirming the item is genuinely no longer valid. Per operator mandate "There MUST NOT be any mistake" – bare assertion is forbidden; positive captured evidence per § 4.12 is mandatory.

Visual treatment – the § 4.12 colorizer MUST style `Obsolete` cells with `cell-status-obsolete` class: light-gray background `#E0E0E0` + strikethrough text on the row's description cell. Operator mandate verbatim: "mark obsolete tickets with some light gray background or something that will indicate that item is no longer valid. Maybe text – the description to be strikethrough styled as well".

Audit cadence – at every release-gate sweep (per § 4.12), the conductor MUST re-evaluate every non-terminal `Issues.md` item AND every `Fixed.md` item for obsolescence per the closed-set Reason vocabulary above. New obsolescence findings produce `Obsolete (→ Fixed.md)` migrations atomic per § 4.12. Triple-check is non-negotiable: each obsolescence claim cites (i) the superseding mandate, (ii) code-state grep confirming the obsoleted behaviour is gone, (iii) runtime/test evidence confirming no test still exercises the obsoleted path.

Pre-build gates `CM-COVENANT-114-90-PROPAGATION` (anchor literal across canonical files) + `CM-ITEM-OBSOLETE-DETAILS` (every `Obsolete` heading carries `**Obsolete-Details:**` line within 3 lines) + `CM-OBSOLETE-COLORIZER-WIRED` (§ 4.12 colorizer emits `cell-status-obsolete` class on the right cells + the CSS file defines the class). Paired § 4.12 mutations strip the anchor literal / delete

an Obsolete-Details line / strip the colorizer class – every mutation FAILs its gate.

Composes with § . . . (item-status closed-set extension), § . . . (Type tracking unchanged – Obsolete is orthogonal to Type), § . . . (Fixed-document column-alignment – Obsolete migrates to Fixed.md just like Fixed/Implemented/Completed), § . . . (Operator-blocked details-line pattern – Obsolete-Details mirrors), § . . . (visual cue + grouping – adds cell-status-obsolete), § . . . (closure vocabulary – Obsolete is the th terminal value), § . . . (Reopened-source pattern – Obsolete uses analogous audit line), § . . . (release-gate sweep – audit cadence trigger), § . . . (iteration discipline – obsolescence audit is a routine cycle step), § . . . (interactive clarification – surface ambiguous obsolescence to operator), § . . . (pre-push fetch – re-check obsolescence post-fetch).

Canonical authority: this Constitution.md § . . . in the HelixConstitution submodule.

Non-compliance is a release blocker. Marking a still-valid item Obsolete is a § . . . PASS-bluff at the planning layer. Marking a genuinely-obsolete item with a non-terminal Status is documentation drift. No escape hatch – every obsolescence finding follows the triple-check evidence rule above.

§ . . . – Summary-doc clarity mandate (User mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"We see in Summary docs - Issues_Summary some not clear one line descriptions - like for example: 'Composes with'. For each workable item in any variant of documentation - long or short descriptions - we MUST HAVE clearly understandable meaning of particular entry of the document! Re-evaluate all of them, fix the descriptions (especially one liners) and make sure that every team member can clearly understand what that particular or any particular workable item from the list is exactly about! There cannot be misunderstanding or unclarity of any kind and no bluff allowed!"

Every short-form summary entry (Issues_Summary.md, Fixed_Summary.md, README.md doc-link section, Status_Summary.md page + , every other one-liner derived from a long-form tracker entry) MUST contain a description that is **self-contained + meaningfully informative** about WHAT the item is, not a fragment of the long-form body. Specifically forbidden one-liner anti-patterns:

- **Section labels** as descriptions: Composes with , Closure criteria , Fix direction , Forensic anchor , Reopened-Details , Operator-Block-Details , Obsolete-Details , Summary , Status , Type , Severity , etc. – these are section *labels* inside the long-form entry, NEVER appropriate as the one-liner description.
- **Bare numeric / categorical fragments:** Critical , High , Medium , Low , Bug , Feature , Task , Fixed , standalone – these are metadata column values, not descriptions.
- **Generic restatements of Status:** In progress , Queued , Reopened , Operator-blocked , Fixed , standalone – again metadata column values.

- **Section-marker echoes:** anything matching `^[A-Z][a-z]+(with|by|on|in|to|for):?$` patterns where the right-hand side is empty.
- **\$-letter alone:** `$KB` , `$EU` , etc. – the section letter is the identifier, not the description.

Required one-liner shape – every summary entry's description column **MUST** contain a complete clause ($\geq$ words OR $\geq$ characters, whichever is longer) that names the SUBJECT (component / app / behaviour / defect class) + the PROBLEM or GOAL (what's broken / what's being added / what's being audited). Example transformations:

- Bad: "Composes with"
- Good: "MPV fork comprehensive secondary-display playback + interaction-pattern + chaos test suite (\$KB-)"
- Bad: "Critical"
- Good: "HDMI audio total loss with 'command error ' on D post-reflash (Tier / / source-side fixes landed, REQUIRES_REBUILD validation pending)"
- Bad: "\$JU"
- Good: "Presenter nd-display behaviour simplification + wake-on-play + notification dim toggle (umbrella, REQUIRES_REBUILD pending validation)"

Generator-level enforcement – `generate_issues_summary.sh` + `generate_fixed_summary.sh` + `update_readme_doc_links.sh` + `generate_status_summary.sh` **MUST** extract the one-liner from the `H /H` **heading line** of the source long-form entry (which already encodes the subject + context per `$ . . + $ . . + $ . . ATM-NNN` convention), **NEVER** from arbitrary downstream text. The generators **MUST** refuse to emit a row whose description matches any of the forbidden anti-patterns above – emitting `(MISSING DESCRIPTION – fix source heading)`

placeholder instead, with the offending row visually highlighted.

Pre-build gate `CM-SUMMARY-CLARITY-DESCRIPTIONS` scans every `*_Summary.md` for the forbidden anti-patterns above; any match FAILs the gate. Paired `$ .` mutation injects "Composes with" into a Summary cell → gate FAILs.

Audit cadence – at every release-gate sweep (`$ .` + `$ .` step), re-evaluate every summary row for clarity. Operator-facing exports (HTML + PDF) MUST never ship with an anti-pattern row.

Composes with `$ .` (Issues_Summary sync – clarity is a property of the generated row), `$ .` (Fixed-document column alignment – clarity applies symmetrically), `$ .` (colorizer – anti-pattern rows highlighted), `$ .` (revision header – generator updates revision on clarity regeneration), `$ .` (Fixed_Summary.parity), `$ .` (Status_Summary two-audience format – page audience MUST get the clearest descriptions), `$ .` (README doc-link rows – same anti-pattern rule), `$ .` (README always-sync), `$ .` (composite docs always-sync), `$ .` (universal MD export), `$ .` (mechanical enforcement – pre-commit hook for summary clarity).

Canonical authority: this Constitution.md `$ .` in the HelixConstitution submodule.

Non-compliance is a release blocker. A summary row whose description is an anti-pattern fragment is severity-equivalent to a `$ .` PASS-bluff at the documentation layer – it implies the item is documented when it is not. No escape hatch.

§ . . – Multi-pass change-evaluation discipline (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Every change to the project or codebase we do MUST BE evaluated in several passes and in in-depth analysis for potential new issues or problems it can introduce! We MUST BE sure that: change nothing breaks, we do validate and verify this in codebase and through comprehensive deep research, that no new issues or problems of any kind is bringing and that main task is achieved with particular change with no bluff of any kind! After we do change or set of changes this mandatory steps MUST BE taken! Then, and only then, when we have validate and verified codebase and performed full analysis and taken into the account all existing knowledge(s) and information sources and confirmed all points we have mentioned now, then this can be accepted as change that is going to be committed, pushed, tested, released and so on! Write as many as required additional documentation and notes about our codebase, the system, or the changes so we can achieve this goals!"

Every non-trivial change to the project codebase MUST pass a multi-pass evaluation BEFORE the change is accepted as commit-ready. The discipline expands § . . (test-interrupt-on-discovery) + § . . (deep-web-research) + § . . (TDD-fix) + § . . (deterministic consistency) + § . . (Phase forensic before any speculative patch) into an explicit pre-commit evaluation pipeline.

The **-pass evaluation contract** (mandatory, ALL passes must hold):

Pass – **Main-task verification.** The change achieves the stated main task. Captured-evidence per § . / § . demonstrates the user-visible behaviour now works AS INTENDED. NO inferred / NO assumed / NO "should work" – only captured evidence.

Pass – **Regression-blast-radius analysis.** Every file the change touches AND every file that imports / sources / references the changed file is enumerated. For each, the agent demonstrates (via grep / dependency-graph / actual test run) that the change does not break the existing contract. Examples: function signature changes → audit every caller; HAL config change → audit every test consuming the config; gate addition → audit every paired meta-test mutation. Captured evidence per § . .

Pass – **Cross-feature interaction analysis.** Beyond direct-dependency blast radius, the change is evaluated against parallel features that share state, timing, hardware, or shell environment. Examples: audio HAL change → check video routing (Fix .. SurfaceView Z-order timing); WiFi roaming change → check BT A DP coex (Fix); test-script addition → check § " ". working-tree quiescence + § . background-test discipline. Captured evidence per § . .

Pass – **Deep-research validation.** Per § . , the chosen approach is verified against external state-of-the-art (official docs, vendor.. guides, open-source codebases, kernel mailing list, AOSP gerrit, GitHub issues). Either the change cites an external precedent OR the literal "NO external solution found – original work" is recorded. CodeGraph queries per § . + § . surface intra-codebase precedents.

Pass – **Anti-bluff confirmation.** Per § . /
 § . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . the change is verified to NOT introduce any
 new bluff surface – no metadata-only PASS, no config-only
 PASS, no script-bug FAIL-bluff, no "tests pass but feature
 doesn't work" anti-pattern. Captured evidence directories
 per § . . reference the exact passing paths.

Documentation requirement. Each Pass produces written
 documentation (commit-message footers OR docs/ entries OR
 qa-results/ evidence) demonstrating the pass was
 completed. A change without all passes documented is
 severity-equivalent to a § . PASS-bluff at the
 development-process layer.

Acceptance criteria. Only AFTER all passes complete
 with documented evidence may the change be: (a) committed
 via commit_all.sh , (b) pushed via § . .
 background-push, (c) tested via § . . background
 test, (d) released / tagged per § . . .

No escape hatch. No --skip-multi-pass , --single-pass , --
 bluff-permitted , --fast-commit flag exists. The passes
 apply to every non-trivial change per the operator's
 verbatim mandate. Trivial changes (typo fixes, doc
 revision-header bumps, MD-export regeneration) are exempt
 ONLY when (i) the change touches zero source code AND (ii)
 the commit message explicitly cites the exemption.

Composes with § . . . / § . . (anti-bluff
 baseline), § . . (test-interrupt-on-discovery –
 Pass + Pass . trigger this), § . .
 (captured-evidence quality – every Pass's evidence MUST be
 quality-checked), § . . (no-guessing – Pass –
 require captured evidence not hypothesis), § . .
 (deep-research – Pass expansion), § . . /
 § . . (subagent-driven – Pass –

dispatchable to subagents per parallel analysis),
 \$ (no-fakes-beyond-unit – Pass enforces), \$ (iteration discipline – multi-pass IS the iteration), \$ (TDD-fix – RED→Pass- , GREEN→Pass-), \$ (deterministic consistency – Pass captured evidence MUST be N= reproducible), \$ (autonomous validation – Pass anti-bluff confirmation gates on autonomous evidence), \$ (universal sink-side evidence – Pass + use closed-set taxonomy),
 \$ / \$ (CodeGraph – Pass dependency-graph queries), \$ (Phase forensic – \$ generalises Phase to Pass- +Pass- +Pass- across every change class),
 \$ (docs/qa transcript – Pass evidence dir is the QA transcript), \$ (stress + chaos – Pass + evidence MUST include stress + chaos signals when applicable), \$ (endless-loop – multi-pass runs concurrently across PWUs per \$),
 \$ (background tests – Pass – long tests run backgrounded).

Pre-build gate CM-COVENANT-114-92-PROPAGATION enforces this anchor literal across the canonical fleet. Pre-build gate CM-MULTI-PASS-EVALUATION-EVIDENCE audits recent commits for the -pass evidence trail (commit message footer OR per-commit qa-results/<commit-hash>/passes/ directory OR Issues.md / Fixed.md narrative). Paired \$ meta-test mutations strip the load-bearing literals → gates FAIL.

Canonical authority: this Constitution.md \$ in the HelixConstitution submodule.

Non-compliance is a release blocker. A commit landing without the -pass evaluation documented is severity-equivalent to a \$ PASS-bluff at the development-

process layer – it implies the change was vetted when it was not.

2026-05-27

§ 2.2.6 – SQLite-backed single-source-of-truth for workable items (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"There MUST be single source of truth for all of our workable items – SQLite database containing all workable items, all data about it and from there proper scripts (we recommend Go programs) which will update / regenerate all documentation – Issues, Fixed, summary docs, and everything related! We MUST reduce a chance for sync to be broken between the Db and all documents we have! We MUST BE able to generate always all docs from DB or to re-generate Db from all docs we have in opposite direction if we must! All this MUST BE covered with all supported test types, validated and verified by tests which will produce real proofs and be % anti-bluff compatible!"

The current text-based Issues.md / Fixed.md / Issues_Summary.md / Fixed_Summary.md / CONTINUATION.md / Status.md tracker constellation is converted to a **SQLite-database-backed single source of truth** with bidirectional regeneration between DB and Markdown. The DB is the authoritative source; all Markdown / HTML / PDF / Status / Summary surfaces are generator output. Sync drift is mechanically impossible because every regeneration starts from the DB.

Canonical artefacts (all consuming projects MUST adopt):

- **Database file** at canonical path `docs/workable_items.db` (SQLite, TRACKED in git, NEVER gitignored per User mandate – amendment). The DB IS the single source of truth – gitignoring it would defeat the SSoT mandate at the version-control layer. Every commit that mutates workable items MUST stage + push the DB file alongside the Markdown regen output. The DB is an exception to `$ . . build-artefact gitignore` – it is NOT a build artefact, it IS authoritative source data. SQLite's WAL mode produces a sidecar `*.db-wal` + `*.db-shm` which ARE gitignored (transient WAL files; their state is checkpointed into the main `.db` file by `workable-items sync` or any explicit `PRAGMA wal_checkpoint(TRUNCATE);`).

• Schema (mandatory minimum tables):

- `items` – primary key `atm_id`, per `$ . .`; columns `type` (`$ . . closed-set`), `status` (`$ . . + $ . . closed-set` including `Obsolete`), `severity`, `title`, `description` (`≥ words / ≥ chars per $ . .`), `created_at`, `last_modified`, `current_location` (`Issues|Fixed`), `forensic_anchor`, `closure_criteria`, `composes_with` (JSON array of refs).
- `item_history` – append-only audit log: `atm_id`, `event_type` (`Opened|Updated|Reopened|Fixed|Implemented|Completed|Obsolete`), `by` (`AI|User per $ . .`), `on_date`, `reason` (`closed vocabulary per $ . . + $ . .`), `evidence_path`.
- `obsolete_details` – `atm_id`, `since`, `reason` (`$ . . closed-set`), `superseding_item`, `triple_check_evidence` (mandatory per `$ . .`).

- `operator_block_details` – `atm_id` , `what` ,
`why_exhausted_alternatives` , `unblock_condition` , `who` per
\$. . .
- `firebase_metadata` – per \$. . .
- `meta` – schema version, `last_sync_direction`,
`last_sync_timestamp`, `integrity_hash`.

- . **Go binary** at canonical path `cmd/workable-items/`
(separate Go module, NOT part of AOSP build).

Subcommands:

- `workable-items sync md-to-db` – parse `Issues.md` +
`Fixed.md` (single source of authority during
transition window), upsert into DB. Idempotent.
- `workable-items sync db-to-md` – regenerate `Issues.md` +
`Fixed.md` + `Issues_Summary.md` + `Fixed_Summary.md` +
`CONTINUATION.md` \$ + every `Status.md` from DB.
Idempotent.
- `workable-items diff` – print DB vs MD divergence. Used
by pre-build gate.
- `workable-items validate` – schema sanity +
\$. . . /\$. . . /\$. . . /
\$. . . /\$. . . /\$. . . /
\$. . . invariants per row. FAIL on any
violation.
- `workable-items add <type> <severity> --title <title> --
description <description>` – interactive entry with
mandatory fields enforced.
- `workable-items close <atm-id> --status <fixed|implemented|
completed|obsolete> --evidence <path>` – terminal
transition with mandatory captured-evidence per
\$. . . /\$

- . **Operational integration:**

- `commit_all.sh` pre-commit hook runs `workable-items diff`
and refuses if MD/DB diverge.

- `sync_issues_docs.sh` invokes `workable-items sync db-to-md` instead of legacy bash generators (`generate_issues_summary.sh` becomes a shim wrapping the Go binary).
- `pre_build_verification.sh` runs `workable-items validate` as a gate; failure increments `ERRORS`.

. **Anti-bluff test coverage** (mandatory per \$. + \$. . + \$. .):

- Unit tests for every CRUD operation + every schema invariant.
- Integration test: full round-trip MD→DB→MD with byte-identical re-emission (modulo cosmetic-whitespace normalisation per closed-set tolerance).
- Stress + chaos per \$. . : -row insert / concurrent-write (writers) / mid-write SIGKILL / corrupt-DB recovery / disk-full handling.
- Paired \$. meta-test mutation: strip a CRUD function → unit test FAILs.
- HelixQA Challenge entry `CME-WORKABLE-ITEMS-001` exercising end-to-end DB→MD→DB→MD round-trip.

Bidirectional regeneration guarantee. The DB MUST be reconstructible from the MD docs (for disaster recovery + audit replay) AND the MD docs MUST be reconstructible from the DB (the primary direction during normal operations). Either direction's regeneration produces semantically equivalent output (closed-set tolerance for cosmetic whitespace + section-ordering noise).

Cross-project propagation. Every consuming project that adopts this constitution submodule inherits the SQLite-SSoT approach. The Go binary lives in the constitution submodule (`constitution/scripts/workable-items/`) so consumers reference it from there per \$. . catalogue-first discipline – never reimplement.

Migration path (per § 11.4.54 iteration discipline + § 11.4.55 PWU pipeline):

- **Phase 1** : file §LA Issues entry tracking the migration, scope-locked.
- **Phase 2** : Go binary scaffold + schema DDL committed.
- **Phase 3** : sync md-to-db lands + initial migration captures current state.
- **Phase 4** : sync db-to-md lands + byte-identical round-trip CI gate.
- **Phase 5** : existing generators (generate_issues_summary.sh etc.) become Go-binary shims.
- **Phase 6** : legacy text-direct edits prohibited (pre-commit hook).

Pre-build gates CM-COVENANT-114-93-PROPAGATION (anchor literal in canonical fleet) + CM-WORKABLE-ITEMS-DB-PRESENT (DB regen mechanism per § 11.4.56 + DB schema-version current) + CM-WORKABLE-ITEMS-MD-DB-IN-SYNC (diff returns empty). Paired § 11.4.57 meta-test mutations strip the load-bearing literals → gates FAIL.

Composes with § 11.4.58 (anti-bluff covenant – DB is the captured-evidence registry), § 11.4.59 (Issues_Summary sync – DB-driven), § 11.4.60 (Status closed-set – DB schema enforces), § 11.4.61 (Type closed-set – DB schema enforces), § 11.4.62 (universal-vs-project – DB schema is universal), § 11.4.63 (Fixed-document column-alignment – DB query emits both Summaries from same query), § 11.4.64 (Operator-blocked details – operator_block_details table), § 11.4.65 (no-fakes-beyond-unit – DB integration tests exercise real SQLite), § 11.4.66 (.gitignore + regeneration mechanism – DB file gitignored with workable-items sync as regen), § 11.4.67 (closure vocabulary – DB enforces), § 11.4.68 (Reopened source attribution – DB item_history table), § 11.4.69 (iteration

discipline – -phase migration), \$. . (TDD – RED before each phase), \$. . (revision header – Markdown output preserves), \$. . (Status.md per integration – generator queries DB by domain),
 \$. . (deterministic consistency – round-trip MD→DB→MD byte-identical), \$. . (LIVE_ADB_FIRST – N/A, host-only), \$. . (autonomous validation – DB integrity is autonomously verifiable), \$. . (Fixed_Summary parity – DB ensures), \$. . (ATM-NNN – DB primary key), \$. . (Reopens-history doc – DB `item_history` queryable), \$. . (Status_Summary parity – DB queries produce both pages),
 \$. . (README doc-link section – DB drives),
 \$. . (parallel PWU – Phase – land as separate PWUs), \$. . (composite always-sync – DB→MD→HTML+PDF), \$. . (universal MD export – final stage), \$. . (catalogue-first – Go binary in constitution), \$. . (docs/qa transcript – DB queries serve as audit replay), \$. . (stress + chaos – DB integration test scope), \$. . (roster/corpus auto-sync – DB schema accommodates roster tables), \$. . (endless-loop – DB queries are zero-idle-friendly), \$. . (background tests – DB validation runs detached), \$. . (Obsolete status – `obsolete_details` table), \$. . (Summary clarity – DB enforces description-floor at insert time),
 \$. . (multi-pass evaluation – every DB schema migration follows -pass).

Canonical authority: this Constitution.md \$. . in the HelixConstitution submodule.

Non-compliance is a release blocker. A consuming project that maintains text-based trackers without the DB-SSoT migration plan filed + Phase progression actively in flight is severity-equivalent to a \$. . PASS-bluff at

the data-architecture layer – the sync-drift surface
§ 1.1.4.4 closes IS a § violation vector.

§ 2.0.2.6 – Zero-idle priority-first parallel-by-default operating mode (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"We MUST NEVER sit iddle / wait or sleep if there is possibility for us to work on something like in this case! Continue all work further in parallel if possible fully autonomously! We MUST not waste time since work scope we have is huge! ... Always check if there is a possibility to work on something while we are not working actively on something! Pick always by priority – most critical workable items and other tasks MUST BE done first! ... Everything we can should be done fully in background, should not affect stability of the System or create problems, should be done using subagents-driven approach! ... Stay still / iddle if nothing is left to be done at all or waiting for something that is blocking us / you!!!"

§ 1.1.4.72 is the **operating-mode reaffirmation** that binds § (endless-loop), § / § (subagent-driven), § (iteration-priority), § (parallel PWU), § (audio top-priority), § (iteration-speedup), § (background-push), § (background-test) into a SINGLE always-on enforcement contract:

The mandate (binding on the conductor + every subagent + every helper script):

(A) **Idle-only-when-genuinely-blocked.** The conductor MAY sleep / wait / be idle ONLY when (i) every priority-queued workable item is genuinely blocked on external dependency (operator hardware action, network upstream, hardware-attended forensic, build/test completion that the conductor cannot accelerate), OR (ii) the operator has issued an explicit STOP, OR (iii) host-session-safety demands it per § 11.4.42. "I don't see what to do next" is NEVER a valid idle reason – the conductor MUST recheck the priority queue per § 11.4.42. step 1 and § 11.4.42. (A) termination criteria.

(B) **Always check parallel-work feasibility at every pause point.** Before any wake-up schedule / before any "waiting for X" / before any sleep call, the conductor MUST: () survey the priority queue per § 11.4.42. SCOPE LOCK + § 11.4.42. audio-first + § 11.4.42. (B) zero-idle, () identify ALL items that can progress without contending with the in-flight blocker, () dispatch them in parallel per § 11.4.42. / § 11.4.42. (subagent-driven if non-trivial) + § 11.4.42. (PWU disjoint scope) + § 11.4.42. (background long tests). Only AFTER step () returns "no parallel work available" may the conductor schedule a wake/sleep.

(C) **Priority order MANDATORY at every pick.** Per § 11.4.42. priority axis + § 11.4.42. audio-top-priority axis. Pick the highest-Severity / highest-priority item the conductor can autonomously progress. Lower-priority items wait until higher-priority queues drain or block on external dependency.

(D) **Subagent-driven by default for non-trivial scope.** Per § 11.4.42. + § 11.4.42. The conductor remains the integration + commit + push seam (per § 11.4.42. +

\$. . quiescence + \$. .); subagents execute the analyses + scaffolds + tests + migrations in parallel. Inline execution permitted only when the task is trivial (single-file edit, sub- -line diff) OR the operator explicitly requested it.

(E) Background by default for long-running work. Per

\$. . stress + \$. . background test + \$. . background push. Any operation expected to exceed ~ s wall-clock MUST run detached (nohup + disown) with log under qa-results/<op-id>_<ts>.log . The conductor returns to the priority queue immediately and polls back when needed.

(F) Stability-preserving – no compromise. Parallel work

MUST NOT compromise system stability or create new problems. Composes with \$. . multi-pass change-evaluation – every parallel branch's output passes – pass before integration. Composes with \$. . working-tree quiescence – concurrent subagents in same checkout coordinate via lockfile or git worktree. Composes with \$. . / \$. . / \$. . / \$. . host-session safety – parallel work respects % memory ceiling and -j AOSP cap.

(G) Status updates as work progresses. Per operator

mandate "Let us know how goes this mandatory approach". The conductor surfaces parallel-work catalogue + in-flight branches + recent closures in compact summary form whenever the operator asks "how are we progressing" OR at natural milestone boundaries (after each constitution-anchor landing, after each commit batch, after each subagent return).

Anti-pattern explicitly forbidden:

- Scheduling a wake-up without first surveying the parallel-work queue per (B).

- Returning "nothing to do – sleeping" when non-trivial priority items remain that the conductor could autonomously progress.
- Serialising work that could run in parallel without justification.
- Picking lower-priority items while higher-priority ones remain progressable.
- Doing work foreground that should run via subagent or background.
- Letting parallel work introduce instability or regression.

Pre-build gate `CM-COVENANT-114-94-PROPAGATION` enforces this anchor literal across the canonical fleet. Pre-build gate `CM-PARALLEL-WORK-AUDIT` (when implemented) audits the last N wake-up schedule decisions for surfaced parallel-work-queue evidence – wake-ups without the survey are § 11.4.42 violations. Paired § 11.4.58 meta-test mutations strip the load-bearing literals → gates FAIL.

Composes with § 11.4.72 (subagent-driven) / § 11.4.73 (iteration discipline) / § 11.4.74 (parallel-development PWU) / § 11.4.75 (subagent-driven default) / § 11.4.76 (audio top-priority) / § 11.4.77 (iteration-speedup) / § 11.4.78 (working-tree quiescence) / § 11.4.79 (stress + chaos) / § 11.4.80 (endless-loop zero-idle) / § 11.4.81 (background-push) / § 11.4.82 (background-test) / § 11.4.83 (multi-pass evaluation) / § 11.4.84 + § 11.4.85 + § 11.4.86 + § 11.4.87 (host-session safety).

Canonical authority: this Constitution.md § 11.4.88 in the HelixConstitution submodule.

Non-compliance is a release blocker. Conductor idle without exhausting the parallel-work queue is severity-

equivalent to a § 2.6.5.2.7 PASS-bluff at the project-execution layer – it implies the conductor has progressed the project when it has not. No escape hatch beyond the (A) genuinely-blocked / explicit-STOP / host-safety triad.

§ 2.6.5.2.7 – Workable-items SQLite DB is TRACKED in git, NEVER gitignored (User mandate, § 2.6.5.2.7 – § 2.6.5.2.7)

Forensic anchor – verbatim user mandate

(§ 2.6.5.2.7 – § 2.6.5.2.7):

"We shall not Git ignore our workable items SQLite DB since it is our single source of truth for current working items we have and all data related to it! We MUST fix this if DB is Git ignored! ... workable items SQLite DB regularly committed and pushed to all upstreams!"

The § 2.6.5.2.7 workable-items SQLite database at docs/workable_items.db is the **AUTHORITATIVE source of truth** for every workable item in every consuming project. It MUST be:

- **TRACKED in git** at canonical path docs/workable_items.db . NEVER gitignored regardless of file-size or "build-artefact-class" heuristics – the DB is NOT a build artefact, it IS authoritative source data.
- **Committed alongside every workable-item state change.** Every workable-items sync md-to-db invocation that mutates DB state MUST stage + commit + push the DB file in the same commit as the corresponding MD changes per § 2.6.5.2.7 atomic-move discipline.

- **Pushed to every upstream per § 4.30** – every mirror of every consuming project carries the latest DB so a fresh clone from any mirror reconstructs the full workable-items inventory without consulting any other source.
- **WAL-checkpointed before commit** – SQLite's WAL mode produces transient *.db-wal + *.db-shm sidecars. Those ARE gitignored (per § 4.31 transient/cache classification). The workable-items Go binary MUST execute PRAGMA wal_checkpoint(TRUNCATE) before any commit-stage step so the main .db file carries the authoritative state and the sidecars are safely discardable.
- **NEVER force-rewritten** without § 4.32 hardlinked-backup + operator authorization. The DB's commit history IS the workable-items audit trail; rewriting it is data-loss equivalent to force-pushing a tracked Issues.md history.

§ 4.33 **carve-out**. This anchor is an explicit named exception to § 4.34 's build-artefact / data-file gitignore rule. The DB file at docs/workable_items.db is **NOT** a build artefact and is **NOT** transient data – it is the project's SSoT registry whose value increases monotonically as work lands. The § 4.35 regeneration mechanism applies to other gitignored bulk data (.git-backup-*, RKTools/linux/, etc.) – NOT to this DB.

§ 4.36 **amendment**. § 4.37 's earlier text describing the DB as "gitignored per § 4.34 with § 4.35 regeneration mechanism" is hereby **AMENDED**: the DB is TRACKED, not gitignored. The Markdown trackers (Issues.md / Fixed.md / Summaries / Status.md fleet / CONTINUATION.md) remain tracked AND derived; the DB is tracked AND authoritative. Both directions of regeneration (workable-items sync db-to-md + workable-items sync md-to-db)

continue to land per § 2.1.1's bidirectional round-trip guarantee.

Pre-build gates CM-COVENANT-114-95-PROPAGATION (anchor literal across canonical fleet) + CM-WORKABLE-ITEMS-DB-TRACKED (git ls-files docs/workable_items.db returns non-empty if DB exists on disk + .gitignore does NOT exclude it). Paired § 2.1.1 meta-test mutation adds the DB to .gitignore → gate FAILs.

Composes with § 2.1.1 (SQLite-SSoT – § 2.1.1 amends its tracking-policy clause), § 2.1.2 (.gitignore discipline – explicit carve-out anchored here), § 2.1.3 (regeneration mechanism – does NOT apply to the DB; the DB IS the source), § 2.1.4 (multi-upstream push – DB pushed to every mirror), § 2.1.5 (data safety – destructive DB ops require hardlinked-backup + operator authorization), § 2.1.6 (force-push merge-first – applies if DB history ever needs rewrite).

Canonical authority: this Constitution.md § 2.1.1 in the HelixConstitution submodule.

Non-compliance is a release blocker. A consuming project that gitignores its workable-items DB is severity-equivalent to a § 2.1.1 PASS-bluff at the source-of-truth layer – it implies SSoT exists when it is in fact ephemeral local-only state.

§ 2.1.1 – Safe-parallel-work-with-long-build catalogue + mandate (User mandate, – –)

Forensic anchor – verbatim user mandate (– –):

"Are there except AOSP build process any other active jobs being done at the moment? Can we work on something in parallel while build is in progress so we slowly cleanup our slate? Anything workable which will not affect the build process or final image we generate for the flashin. If yes, and that was not already clear by our root Constitution, we should add all mandatory details into it ... Make sure we do as much as possible work in background in parallel with main work stream and oreferably using subagents-driven approach!"

\$. . mandates "always check parallel-work feasibility before idle"; \$. . mandates background-test discipline; \$. . defines the parallel-work-unit pipeline. \$. . is the **operational catalogue** that enumerates, for the canonical long-running workload (- h AOSP containerised build per \$.), exactly what CAN safely run in parallel and what MUST NOT – closing the ambiguity gap \$. . left implicit.

Closed-set categories – SAFE during AOSP containerised build:

- (A) **Markdown / documentation work** anywhere under docs/ , *.md , README, CLAUDE.md, AGENTS.md, QWEN.md, Status.md / Status_Summary.md, Issues.md / Fixed.md / their Summary siblings, changelogs, research notes. The build container mounts source read-only at a snapshot moment + reads only the build-relevant files (Android.bp / Android.mk / source code / vendor binaries). MD edits do NOT propagate into the build's read scope.
- (B) **Generator + helper script work** under scripts/ , scripts/testing/ , scripts/llm/ , scripts/firebase/ , etc. –

these are host-side helpers NEVER consumed by the AOSP build. Safe to edit + commit + push freely.

- **(C) Pre-build / meta-test gate authoring + paired**
\$. mutations under device/rockchip/rk3588/tests/
pre_build_verification.sh + scripts/testing/
meta_test_false_positive_proof.sh – these run host-side via
bash pre_build_verification.sh ; the AOSP build itself
does NOT execute them inline. Safe to edit during build;
the NEXT pre-build run (post-current-build) picks up the
new gates.
- **(D) On-device tests** under device/rockchip/rk3588/tests/
test_*.sh – these run on D /D via ADB post-flash.
Authoring new test scripts during build does NOT affect
the build (the build packages the test corpus from the
source tree but does so once at copy-out time; scripts
added after the copy-out are absent from THIS image but
present in source for the NEXT image).
- **(E) Constitution submodule edits + push** – separate git
workspace at constitution/ , not on the AOSP build's
path. Edits + push to all constitution remotes
safe at any time.
- **(F) Project submodule commit + push** to either the parent
or owned submodules' remotes – per \$. .
background push, this is OS-level orthogonal to the
container.
- **(G) D / D live-ADB probes** (read-only dumpsys /
getprop / cat /proc/asound/... / screencap / logcat -d)
– these query device state without changing it; do NOT
affect the host build.
- **(H) Subagent dispatch** per \$. . /
\$. . – each subagent runs in isolated context
window; multiple subagents in same checkout coordinate
via \$. . working-tree quiescence (lockfile or
git worktree per disjoint scope). Read-only analysis

subagents (obsolescence audit, summary clarity sweep, code-review) are uniformly safe.

- **(I) Web research + external API queries** (with § credential discipline preserved) – operator-relevant browsing, ASUS QVL lookups, vendor spec verification, package-registry version checks.
- **(J) Workable-items DB operations** per § + § – `workable-items sync` (when implemented), `workable-items validate`, `workable-items report` – read/write SQLite at `docs/workable_items.db`, orthogonal to build.
- **(K) Pre-build verification + meta-test execution** – RUN per § background-test discipline. The build container reads source; pre-build runs host-side reading the same source. No conflict at filesystem layer (read-only intersection).

Closed-set – UNSAFE during AOSP containerised build:

- **(α) `git checkout` / `git reset --hard` / `git clean -df`** on the source tree – would change files the build is reading; introduces inconsistency. Defer until build completes OR use a `git worktree` per §.
- **(β) Mass file deletions or renames** under `device/`, `frameworks/`, `hardware/`, `vendor/`, `kernel-5.10/`, `external/`, `packages/`, `bionic/`, `bootable/`, `build/`, `system/`, `cts/`, `tools/`, `prebuilts/`, etc. – anything the AOSP build potentially reads. Apply ONLY after build completes.
- **(γ) Submodule pointer updates that affect built APKs** – bumping `device/rockchip/atmosphere/presenter` or `smarttube-player` or `vlc-player` etc. pointer changes the binary content the build packages. Defer pointer bumps to between-build windows OR plan for the NEXT build to pick them up.

- (δ) **out/ directory mutations** – the build's output directory. Never touch concurrently.
- (ε) **make clean / m clobber / rm -rf out/** while build runs – direct conflict; will likely crash the build.
- (ζ) **Container destruction** (`podman stop atmosphere-aosp-build` , `podman pod rm`) – terminates the build mid-flight. Operator-explicit only.
- (η) **Disk-filling operations** (large file downloads, container image builds, NVMe duplication) that push the host below the \$ `free-space` minimum – could OOM the build's writable layer.
- (θ) **Host-session-safety breaches** per \$ `user.slice` – suspend / hibernate / logout / unbounded memory ops in user.slice –kill the build's parent cgroup.

Conductor responsibility (per \$ `user.slice` already;
\$ `user.slice` makes the catalogue explicit):

- Before EVERY pause point during a long build, the conductor MUST consult this catalogue, identify all (A)-(K) items in the priority queue per \$ `user.slice` + \$ `user.slice` , dispatch at least one per \$ `user.slice` / \$ `user.slice` subagent-driven default + \$ `user.slice` background-task discipline. "Build is running, nothing else to do" is NEVER true per \$ `user.slice` + this catalogue.

Subagent-driven default for catalogue items:

- (A) MD work + (E) constitution edits + (F) commits: conductor-direct (lightweight + critical-state sequencing).
- (B) generator/helper work + (C) gate authoring + (D) on-device test authoring + (G) live-ADB probes + (I) web research + (J) DB ops: subagent-dispatchable per \$ `user.slice` .

- (H) is itself the subagent dispatch – composes with the rest.
- (K) pre-build / meta-test execution: backgrounded per § 11.4.89, conductor polls outcome.

Pre-build gate CM-COVENANT-114-96-PROPAGATION enforces this anchor literal across the canonical fleet. Pre-build gate CM-PARALLEL-WORK-DURING-BUILD-AUDIT (when implemented) audits recent commits during AOSP build windows for the parallel-work output (commits landing during a build window with timestamps inside qa-results/aosp_build_*.log start/end). Paired § 11.4.42 meta-test mutations strip the catalogue literals → gates FAIL.

Composes with § 11.4.92 (subagent-driven) + § 11.4.93 (iteration priority) + § 11.4.94 (parallel PWU) + § 11.4.95 (subagent default) + § 11.4.97 (audio top-priority) + § 11.4.98 (iteration-speedup) + § 11.4.99 (working-tree quiescence – esp. for subagent parallelism) + § 11.4.100 (stress + chaos can run in parallel) + § 11.4.101 (endless-loop zero-idle) + § 11.4.102 (background-push) + § 11.4.103 (background-test) + § 11.4.104 (multi-pass evaluation per parallel branch) + § 11.4.105 / § 11.4.106 (workable-items DB) + § 11.4.107 (parallel-work survey mandate) + § 11.4.108 / § 11.4.109 / § 11.4.110 (host-session safety + build cap).

Canonical authority: this Constitution.md § 11.4.111 in the HelixConstitution submodule.

Non-compliance is a release blocker. Conductor scheduling a wake/sleep during an active AOSP build without first dispatching at least one (A)-(K) catalogue item from the priority queue is a § 11.4.112 + § 11.4.113 violation = § 11.4.114 PASS-bluff at the operating-mode layer.

§ 2.0.2.6.0.5.2.7 – Maximum-use-of-idle-time mandate + progress-update cadence (User mandate, – –)

Forensic anchor – verbatim user mandate (– –):

"keep it working, we should do as much as possible, if not it all but as much as we can as long as there is iddle time! it MUST be used! make sure that it is clear in constitution! keep us updated about all progress and all phisycal proofs and gathered data as you progress through all open workable items!"

§ 2.0.2.6.0.5.2.7 strengthens § 2.0.2.6.0.5.2.7 (endless-loop) + § 2.0.2.6.0.5.2.7 (zero-idle priority-first) + § 2.0.2.6.0.5.2.7 (safe-parallel-with-build catalogue) with an **explicit maximum-use-of-idle-time mandate** AND a **progress-update cadence mandate**:

(A) Maximum-use idle-time mandate. Every minute of conductor idle time during which (i) the conductor has work it could autonomously progress AND (ii) is NOT genuinely blocked on external dependency (operator hardware action, network upstream, host-session safety, build/test completion) is a § 2.0.2.6.0.5.2.7 violation. "As much as possible, if not it all but as much as we can" is the operative phrasing – the conductor MUST dispatch work continuously through the entire idle window, not just sample-check at scheduled wakes.

(B) Progress-update cadence mandate. The conductor MUST emit operator-facing progress updates at the following

natural milestone boundaries (no operator prompt required):

- Every commit that lands (HEAD advance) – -line "what just landed".
- Every subagent return – -line "subagent X returned, integrated".
- Every constitutional anchor landed – -line "\$X.Y.Z propagated to remotes".
- Every gathered physical proof – -line "captured evidence at qa-results/".
- Every milestone closure (Issues→Fixed migration, Obsolete classification, etc.) – -line.

Progress updates are CONCISE (- lines each); the operator's mental model stays current without conductor overhead.

(C) Continuous physical-proof gathering mandate. Per \$. . + \$. . + \$. . – every autonomous-fixable item the conductor closes MUST gather positive captured-evidence: ADB probes, dumphsys captures, screencaps, ALSA hw_params, sysfs reads, file-content hashes, runtime metric snapshots. The evidence is committed alongside the closure narrative. Per \$. . SQLite-SSoT, the evidence path goes into the item_history.evidence_path column when DB integration lands.

(D) Per-anchor composition. \$. . binds together: \$. . (captured-evidence quality), \$. . (no-guessing – every progress claim cites evidence), \$. . (sink-side evidence when applicable), \$. . (subagent-driven default for parallel scope), \$. . (no-fakes-beyond-unit – closures cited with real captured evidence), \$. . (priority queue), \$. . (deterministic consistency – N= where applicable), \$. .

11 4 96

11 4 97

11	4	42	11	4	50	11	4	52
----	---	----	----	---	----	----	---	----

Composes with every anchor in $\$$. (this is the operating-mode capstone): $\$$. . + $\$$. . +

$\$$ ₁₁ 4 88 • 11 • 4 88 11 + $\$$. . + $\$$. . +

$\$$ ₁₁ 4 94 • . + $\$$. . + $\$$. . +

$\$$. . + $\$$. . + $\$$. . +

$\$$. . + $\$$. . + $\$$. . +

$\$$. . + $\$$. . + $\$$. . +

$\$$. . .

Canonical authority: this Constitution.md § . .
in the HelixConstitution submodule.

Non-compliance is a release blocker. Conductor idle when progressable work remains, OR conductor failing to emit milestone progress updates, is severity-equivalent to a § . . . PASS-bluff at the operating-mode layer.

§ . . – Full-Automation Anti-Bluff
Mandate – Live tests MUST be re-runnable end-
to-end without manual intervention (User
mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"Make sure we have full automation testing of all scenarios with real bot, main group and users without any manual intervention or contribution of real user! Everything MUST BE fully automatic and autonomous! These tests MUST BE able to rerun endless times when needed! This is important to be done like this! It is critical! Continue all work and make this happen! We need such full automation testing so the whole System MUST BE fully validated and verified before it is integrated to our main projects! Make sure there is no false positives in testing! Every test and its results MUST obtain real proofs of everything working! No bluff is allowed! IMPORTANT: Make sure that all existing tests and Challenges do work in anti-bluff manner – they MUST confirm that all tested codebase really works as expected!"

§ . . composes with § . / § . . /
§ . . / § . . / § . . /

§ . . / § . . / § . . – closes the **manual-intervention gap** that they did not explicitly forbid. A live/integration/e e/Challenge test which requires a human action during execution (typing a chat message, clicking a UI, hand-triggering a webhook, manually attaching a file, anything beyond the test starting up and reporting PASS/FAIL on its own) is **by definition a § . PASS-bluff at the automation layer**, regardless of how thorough the manual run is. The reason: such a test cannot run continuously in CI, cannot validate regressions between manual runs, and the human dependency masks any drift between code changes and what the test exercises.

(A) Binding rule. Every test that this Constitution governs – unit / integration / e e / Challenge / stress / chaos / live – **MUST** be fully self-driving end-to-end. The test process must accept inputs from configuration alone, exercise the load-bearing code path autonomously, and report PASS/FAIL/SKIP-with-reason without any further human action.

(B) Single permissible exception – one-time credential bootstrap. Configuration performed OUTSIDE test execution is acceptable: populating `.env` from a vault, exporting shell env vars in `~/.bashrc`, OAuth approval at first install, MTPROTO session activation at first run. This is configuration, not test driving. Once credentials are present, the test **MUST** run end-to-end without any further human action. The configuration step is itself a § . . credential-handling concern, not a § . . automation concern.

(C) Concrete requirements for live messenger / channel / agent tests (the surface that exposed this mandate):

- **No "operator MUST type a message" prompts in test scripts.** Tests requiring inbound messages **MUST** drive

those messages programmatically – via a separate user account (MTPROTO for Telegram, real-user-token API for Slack, IMAP-test-account for email, etc.), via a webhook fixture, or via an in-process loopback. Never via human keystrokes during test execution.

- **No hard-coded session UUIDs that collide with the active dev session.** Test runs MUST use a dedicated test-only session that the production / dev session is NOT simultaneously using. (Lesson from Herald
`claude --resume <UUID>` invoked against the same session ID the dev session is using returns exit - with no output – a silent collision that fails for non-obvious reasons.)
- **No -second human-response windows.** These create flake/timing races that are themselves § . . determinism violations: a single test invocation's PASS/FAIL depends on a human reaction-time variable, not on the code under test.
- **Re-runability proof.** Every live test MUST PASS at least consecutive automated invocations (`-count=3` for Go, `repeat 3` equivalent for other runtimes) with no state cleanup between runs. Persistent side effects (chats, files, DB rows, attached objects, queued events) MUST be cleaned by the test itself in a `defer / teardown` block.
- **§ . . obsolescence audit.** Every existing test that ships under this Constitution MUST be classified as either "self-driving" (COMPLIANT) or "manual-dependency" (NON-COMPLIANT – must be rewritten or marked § . . Obsolete). The audit is a release-gate item – a release that ships a NON-COMPLIANT test still pretending to PASS is blocked.
- **No false-positive PASS.** A test PASS that does NOT exercise the load-bearing code path – e.g. a "live" test that skips silently because env vars are missing yet reports PASS, or a "live" test where the captured

evidence is from a prior cached run – is itself a § . PASS-bluff at the gating layer. SKIP-with-reason is the § . correct posture; SKIP-reported-as-PASS or stale-evidence-reported-as-fresh is forbidden.

(D) Composition. § . (stress + chaos) + § . (background-test execution) + § . (endless-loop autonomous work) + § . (zero-idle parallel-by-default) + § . (full-automation) together make this Constitution's testing surface a continuously-validated, fully-automated, non-flake, anti-bluff regime. Each closes a different gap; remove any one and the whole property collapses. § . specifically closes the manual-intervention gap: § . + § . + § . + § . do not forbid a test from requiring human action, only from skipping the load-bearing path. § . makes the human dependency itself a § . violation.

(E) Inheritance per § . . § . . propagates to every consuming repository via the existing inheritance gate: each project's CLAUDE.md / AGENTS.md / QWEN.md MUST carry a short-form restatement citing the literal anchor 11.4.98 . The pre-build gate CM-COVENANT-114-98-PROPAGATION (when implemented) enforces this literal anchor presence across the canonical fleet. Paired § . meta-test mutations strip the load-bearing literal → gates FAIL.

(F) Enforcement. „A commit that adds or modifies a test that requires manual human action during execution is blocked at release-gate. A test classified as "manual-dependency" in the § . audit that has not been rewritten within days of classification graduates to § . . Obsolete and is removed from the active test suite (not deleted – preserved with Obsolete-Details:

per § . . . , citing § . . . as the obsolescence reason).

Canonical authority: this Constitution.md § . . . in the HelixConstitution submodule. All consuming projects (Herald, a consuming project, future) restate + cite via § . . . inheritance.

Non-compliance is a release blocker. No `--manual-test-OK`, `--skip-114-98-audit`, `--bluff-tolerance-temporary` flag exists. The `--` user mandate is unambiguous: "No bluff is allowed!"

§ . . . – Latest-Source Documentation Cross-Reference Mandate – instructions, guides, and manuals MUST be verified against the latest official online sources BEFORE publication (User mandate, – –)

Forensic anchor – verbatim user mandate (– –):

"Make sure we ALWAYS check against latest versions of services we use web / online docs before creating instructions! This situation is illustration of how we can misguide ourselves or get banned! Add this all important generic / general points as proper mandatory rules when we are creating documentation or guides! These are mandatory rules / constraints and the result is consistency and safety of created instructions, guides and manuals!"

Case study that anchored this rule (Herald – –). The first-draft Herald MTPROTO setup

guide (committed at 35fc10c / fb00354 / f089dd6) recommended VoIP / Google Voice / Twilio / TextNow numbers as a budget-friendly fallback option AND omitted the critical recover@telegram.org pre-login email step. Both pieces of guidance directly contradicted (a) Telegram's official documentation at https://core.telegram.org/api/obtaining_api_id ("all accounts that sign up or log in using unofficial Telegram clients are automatically put under observation") AND (b) the gotd/td maintainer's "How to not get banned?" guide (vendored at submodules/gotd-td/.github/SUPPORT.md). An operator who followed the original guide could have had their Telegram account permanently banned with no appeal – the very harm the documentation was supposed to help them avoid. The corrected guide landed at 8470ba7 after a forced cross-reference against the latest official sources. This case study is forensic evidence that misguidance-by-stale-docs is the same severity class as a PASS-bluff at the documentation layer.

Composition. This anchor composes with \$ (test-interrupt-on-discovery), \$ (captured-evidence), \$ (deep-research validation), \$ (multi-pass change-evaluation Pass), \$ (anti-bluff), and \$ (full-automation). It closes a specific gap \$ Pass alludes to but does not explicitly mandate: **every operator-facing instruction, guide, manual, troubleshooting cookbook, or setup walkthrough MUST be cross-referenced against the LATEST official online documentation of the service / library being documented BEFORE the document is committed.**

(A) Binding rule – pre-commit cross-reference. Before committing any document that contains operator-actionable steps (setup walkthroughs, integration guides, API how-tos, credential acquisition, security configurations, troubleshooting cookbooks, billing-policy documentation,

terms-of-service summaries, any external-service interaction guide), the author **MUST**:

- . **Fetch the latest official online documentation** of the third-party service / library being documented via WebFetch, direct browsing, the service's own MCP server (if available), or equivalent authoritative real-time source. Do NOT rely on training data, memory, prior assumptions, or older committed docs as the source of truth.
- . **Cross-reference each instruction in the new document** against that source. For each step the operator will take, verify: (a) the service still supports that action; (b) the form fields / parameters / endpoints are still the same; (c) any new constraints, deprecations, or warnings the service published since the doc's last verification.
- . **Seek secondary authoritative sources when the official documentation is sparse / silent** on a critical requirement. Examples: the library maintainer's official guides (`submodules/<lib>/README.md` , `SUPPORT.md` , `SECURITY.md`), the service's official blog / changelog, the service's official support email / channel responses, community-vetted FAQs (Stack Overflow accepted answers, official Discord pinned messages).
- . **Cite the source URL + the date checked at the bottom of the document** in a `## Sources verified` section. Example: `Sources verified 2026-05-28: https://core.telegram.org/api/obtaining_api_id + submodules/gotd-td/.github/SUPPORT.md` . The citation is the audit trail; a doc without it is by § . . definition NOT verified.
- . **Cite the cross-reference in the commit message footer**. A commit that adds or modifies operator-facing instructions **MUST** include a

Sources verified <date>: <urls> footer line so the verification trail is reachable from git log queries.

(B) Negative-finding documentation is required. If the cross-reference reveals the official source is silent / contradictory / outdated on a question the new document needs to answer, the document MUST explicitly note that gap so the next reader does not assume the absence of contradiction means authoritative agreement. Example phrasing the Herald MTPROTO guide uses: *"The official Telegram docs do not enumerate App title validation rules; the constraints below were verified against the gotd/td community channel + empirical operator testing - - ."*

(C) Re-verification cadence. Documents older than months without re-verification are STALE – operators MUST NOT trust them for fresh action without re-running the cross-reference. Documents MUST be re-verified:

- . **Before being cited as the authority for an operator-action campaign** (the operator's command to "follow this guide" is the trigger).
- . **At every major release boundary** (vN. . tag – the release-gate sweep includes a \$. . freshness audit).
- . **Whenever the documented service publishes a breaking-change announcement** (the agent that knows about the announcement MUST queue the re-verification).
- . **When an operator reports an error following the guide** (the case-study anchoring is automatic and triggers immediate re-verification).

(D) Service-specific risk-classifications. Documentation for the following service families MUST include explicit safety warnings cross-referenced against the latest

published policies, with a § . . "Sources verified" date NEVER older than days:

Family	Specific risks the documentation MUST address
Telegram / Whats-App / messenger APIs (unofficial clients)	Anti-abuse-system observation; one-phone-one-app-id limits; ban-on-VoIP policies; rate-limit floods; user-im-personation risks; pre-login declaration emails.
Cloud provider APIs (AWS, GCP, Azure, Cloudflare)	Billing blast-radius; account-lock policies; service-quota limits; least-privilege IAM; data-residency obligations.
Payment systems (Stripe, PayPal, banks, Mercado Pago)	KYC/AML compliance; PCI-DSS data-handling rules; webhook signature verification; idempotency keys; refund/chargeback flows.
AI / LLM providers (Anthropic, OpenAI, Google AI)	Terms-of-service violation triggers; rate-limit + quota policies; data-re-tention defaults; safety-classifier outcomes; PII-in-prompt risks.
Code-hosting ser-vices (GitHub, Git-Lab, GitFlic, Git-Verse, Bitbucket)	Token-leak revocation policies; force-push to protected-branch policies; rate-limit-on-API-tokens; secret-scanning bot behaviour.
OS / package man-agers (apt, brew, npm, pip, cargo)	Supply-chain compromise vectors; sig-nature verification; lockfile discip-line; mirror-trust policies.

This list is **NOT exhaustive** – when a new external service is documented, the § . . author judges whether

the service has comparable risk surface; if yes, the same safety-warning requirement applies.

(E) Composition with § 11.4.92 . . . multi-pass change-evaluation. § 11.4.93 . . . Pass (deep-research validation) already requires "external precedent or literal 'NO external solution found'". § 11.4.94 . . . strengthens this specifically for the operator-facing-instructions class: deep research is mandatory for ALL such docs, NOT just non-trivial code changes. The agent that authors an instruction guide CANNOT cite § 11.4.95 . . . Pass as a substitute for § 11.4.96 . . . – the two pass independently.

(F) Inheritance per § 11.4.97 . . . Every consuming repository's CLAUDE.md / AGENTS.md / QWEN.md MUST carry a short-form restatement citing the literal anchor 11.4.99 . The pre-build gate CM-COVENANT-114-99-PROPAGATION (when implemented) enforces this literal anchor presence across the canonical fleet. Paired § 11.4.100 . . . meta-test mutations strip the load-bearing literal → gates FAIL.

(G) Enforcement. A commit that adds or modifies operator-facing instruction documentation without (a) a "Sources verified" footer in the document itself AND (b) a Sources verified line in the commit-message footer is blocked at release-gate. A document with operator-actionable steps that becomes stale (older than 12 months without re-verification, or 30 days for risk-classified services per §(D)) graduates to § 11.4.101 . . . Obsolete with Obsolete-Details: Reason=stale-documentation; superseded-by=<replacement-doc-or-rewrite> after the 30-day grace window per § 11.4.102 . . .

Canonical authority: this Constitution.md § 11.4.103 . . . in the HelixConstitution submodule. All consuming projects (Herald, a consuming project, future) restate + cite via § 11.4.104 . . . inheritance.

Non-compliance is a release blocker. No `--skip-source-check` / `--documentation-freshness-optional` / `--cite-sources-later` / `--trust-prior-doc-as-authoritative` flag exists. The user mandate is unambiguous: the result MUST be "consistency and safety of created instructions, guides and manuals".

§ . . – **RETIRED.** Demoted to consumer project (a consuming project's video-color/visual-quality fidelity) per § . . /§ . . – project-specific (RK /MPV/Arvus), not universal. See the consuming project's Constitution/CLAUDE/AGENTS/QWEN.

§ . . – **Autonomous-decision-over-blocking** mandate (User mandate,)

Short tag: `autonomous-decision-over-blocking` .

Forensic anchor – verbatim user mandate

(– –):

"when working in endless working loop fully autonomously try to decide most properly about points which would block execution and wait for us. If we haven't answered now work would be blocked whole night! If possible and if that will not cause any issues make proper and most reliable and safe decision so we achieve maximal efficiency and work gets fully done!"

The mandate. When operating in an autonomous / endless-loop mode (per § . .), the agent MUST minimize operator-blocking and instead make the safe, reliable, reversible decision itself – so work is NOT stalled (e.g.

overnight) waiting for input. The § . . . endless-loop covenant told the agent to keep working; § . . . tells it HOW to handle the decision points that would otherwise force it to stop and wait. A wrong bias toward blocking-when-it-was-safe-to-decide wastes operator time (the forensic harm: "blocked whole night"); a wrong bias toward deciding-when-it-should-have-blocked risks irreversible harm. The closed-set decision rule below is the precise boundary between the two.

Decision rule (closed-set – the agent MAY proceed autonomously when ALL hold):

(a) the action is **reversible** OR has a captured pre-op backup per § . . . (hardlinked `.git` mirror, recorded refs/tags/submodule pointers); (b) the agent can **determine the safe choice from captured evidence** per § . . . (no guessing – `LIKELY` / `probably` / `seems` is NOT a safe-choice determination; the evidence must state the correct choice as fact, or the condition is not met); (c) a wrong choice's **blast radius is bounded AND recoverable** – the damage is local, undoable, and does not cascade beyond the work unit; (d) it **composes with the existing covenants** – anti-bluff § . . . (no PASS-bluff to skip the decision), host-session-safety § . . . (no decision that suspends / logs-out / overruns the § . . . memory budget), data-safety § . . . (no destructive op without the § . . . protocol).

Block-only-when rule (the agent MUST BLOCK – ask via the § . . . interactive mechanism – ONLY when ALL of the following hold): the action is **irreversible AND high-blast-radius** AND the safe choice **cannot be determined from evidence**. Canonical block-required examples: external-account state the agent cannot inspect (registration, certification, billing portals), hardware the agent cannot physically access, destructive repository ops without a backup, force-push (which independently requires § . . . + § . . . authorisation), spending money or sending

messages / data to third parties. Operator-blocked per § 4.40.1 remains the last-resort status and is reached only after this rule fires AND the § 4.40.2 self-resolution-exhaustion audit is complete.

Maximize-progress-while-blocked rule. When a block IS unavoidable, the agent MUST still maximize progress on every NON-blocked item in parallel per § 4.40.1 + § 4.40.2 (zero-idle priority-first) rather than idling – the blocked decision parks one work unit, it does not pause the loop. Posing the § 4.40.3 question and continuing other work is the correct behaviour; posing the question and going idle is a § 4.40.4 + § 4.40.5 violation.

Composes with § 4.40.6 (no-guessing – sub-rule (b) IS the no-guessing test applied to the decision itself), § 4.40.7 (operator-blocked is the last-resort status the block-rule routes into, after self-resolution exhaustion), § 4.40.8 (a release-tag decision is irreversible-and-high-blast-radius → always blocks per the block-rule), § 4.40.9 (force-push is an explicit block-required example), § 4.40.10 (the interactive-clarification mechanism the block-rule uses), § 4.40.11 (the endless-loop covenant § 4.40.12 refines – § 4.40.13 says keep working, § 4.40.14 says how to clear the decision points), § 4.40.15 (zero-idle – maximize-progress-while-blocked), § 4.40.16 (backup is the reversibility substitute in sub-rule (a)), § 4.40.17 (host-safety bounds every autonomous decision).

Classification: universal (§ 4.40.18) – autonomous-decision-vs-block boundary is a reusable discipline for ANY project running an agent in an endless / unattended loop, hardware- and domain-agnostic.

-layer coverage per § 11.4.101 (b). Propagation gate CM-COVENANT-114-101-PROPAGATION enforces the literal anchor 11.4.101 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired § 11.4.101 meta-test mutation strips the 11.4.101 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § 11.4.101 in the HelixConstitution submodule. All consuming projects restate + cite via § 11.4.101 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no `--always-block-on-decision`, `--never-decide-autonomously`, `--skip-decision-rule`, `--block-without-self-resolution` flag exists. The user mandate is unambiguous: in autonomous mode the agent makes the safe reversible decision and keeps working – it does NOT block the operator overnight on a choice it had the evidence to make safely.

§ 11.4.101 – Mandatory systematic-debugging activation + always-loaded skill-discovery + plugin-dependency availability (User mandate, – –)

Short tag: `systematic-debugging-always-on`.

Forensic anchor – verbatim user mandate
(– –):

"Make sure that we ALWAYS trigger / start the "/superpowers:systematic-debugging" skills when any issues happen! If this is possible to activate and use in this situations out of the box when we spot

problems / issues / bugs / misalignments /
inconsistencies we MUST activate the skill(s) and make
strongest efforts in full in depth analysis /
debugging and determine root causes of all problem or
obtain relevant data and information we need! ... we
MUST make sure that "/using-superpowers" skill is
ALWAYS loaded, applied and used! All dependencies
(plugins) that Claude Code or other market places are
offering MUST BE installed if these are not already
available for loading and use!"

This anchor binds three cooperating invariants. Together they make the difference between an agent that **guesses-and-retries** (symptom-patching, "transient", "flaky", "probably a race") and an agent that **investigates-to-root-cause-first** mechanical, every single time a problem surfaces. The § 1.1 no-guessing mandate said *do not state an unproven cause*; § 1.2 says *the moment a problem appears, the very first action is to launch the structured root-cause investigation skill – not to propose a fix*.

(A) Mandatory systematic-debugging activation on ANY spotted problem. On ANY spotted issue, bug, test failure, gate failure, regression, misalignment, inconsistency, unexpected behaviour, or contradiction-between-evidence-and-expectation, the agent MUST activate the structured root-cause-investigation skill – `superpowers:systematic-debugging` on Claude Code, or the platform-equivalent structured-debugging discipline on any other runtime – **BEFORE proposing, writing, or applying any fix**. This is the **Iron Law: NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST**. The investigation MUST run its full four-phase arc: (1) **root-cause phase** – read the actual error / failing assertion / divergent evidence, reproduce deterministically, gather the facts the failure exposes;

() **pattern phase** – identify what class of defect this is, search for the same pattern elsewhere in the codebase;
() **hypothesis phase** – form a falsifiable hypothesis about the cause, then prove or disprove it with captured evidence (never **LIKELY** / **probably** / **seems** per § . . .); () **implementation phase** – only after the cause is a **FACT**, design the fix against the proven root cause. Guess-and-retry, symptom-patching, "let me just try changing this and see", and re-running a failed test hoping it passes ("it's probably transient / flaky") **WITHOUT** a completed root-cause investigation are themselves § violations. Real-world signal this anchor closes: declaring a failure **transient** / **flaky** / **intermittent** / **probably a timing issue** **without captured forensic evidence proving that classification is** simultaneously a § . . . (no-guessing) violation AND a § . . . (demotion-evidence) violation – § . . . makes the corrective response mechanical: the classification is forbidden until the systematic-debugging arc has captured the evidence that proves it.

(B) Mandatory always-loaded skill-discovery discipline (using-superpowers). The skill-discovery meta-skill – **superpowers:using-superpowers** on Claude Code, or the platform-equivalent skill-index / capability-discovery discipline on any other runtime – **MUST** be loaded and applied at session start and consulted before any task is begun. Its operative rule: **before acting on ANY request, survey the available skills/capabilities; if ANY skill could apply to the task at hand – even at a % chance of relevance – that skill MUST be invoked rather than the agent improvising from memory.** This guarantees the agent reaches for the purpose-built discipline (systematic-debugging for a bug, the brainstorming skill before creative work, the test-driven-development skill before writing a fix, etc.) instead of ad-hoc behaviour. Skipping skill-discovery at session start, or improvising a task

that a loaded skill was built to handle, is a § . . violation.

(C) Mandatory plugin / dependency availability. Every skill plugin, marketplace package, or capability dependency that the project relies on – offered by Claude Code, its skill marketplaces, or the equivalent plugin/extension ecosystem of any other runtime the project uses – **MUST** be installed and loadable on the host **BEFORE** the dependent work proceeds. A missing plugin/dependency that blocks a mandated skill (e.g. the `superpowers` plugin absent so `systematic-debugging` cannot launch) is a **release-blocker** until it is installed and confirmed loadable. Install mechanism (generic; runtime supplies the concrete command): the runtime's own plugin/marketplace install path – for Claude Code, the in-session `/plugin` marketplace flow (add the marketplace, install the plugin, confirm the skill appears in the available-skills list); for other runtimes, the documented package/extension installer (npm/pip/cargo/the editor's extension manager). The availability check is itself anti-bluff: a plugin claimed "installed" without the dependent skill actually appearing in the loaded-skills/capabilities list is a § . . (no-guessing) violation – confirm by observing the skill in the live capability list, not by assuming the install succeeded (install-command exit `≠` skill loadable, mirroring the § . . `npm exit 0 ≠ working binary` lesson).

(D) AUTOMATIC activation – no operator prompt required (User mandate, – –). Verbatim operator mandate (– –): "Make sure we automatically turn on the systematic-debugging plugin and skills every time we work on any issue! Add this as MANDATORY RULE to our root constitution Submodule and all of its relevant files!" The clause-(A) systematic-debugging activation is AUTOMATIC and the **STANDING DEFAULT** – the agent **MUST** auto-

activate `superpowers:systematic-debugging` (or the platform-equivalent structured-debugging discipline) on ANY spotted issue / bug / test failure / gate failure / regression / misalignment / inconsistency / unexpected behaviour WITHOUT waiting for an operator prompt, request, or per-issue confirmation. It is engaged the MOMENT an issue is detected – the operator **MUST NOT** have to ask for it, request it, or re-enable it per issue or per session (composes § 11.4.100 default-autonomous-loop, which promotes the standing-default working mode from the first prompt of a session; § 11.4.101 (D) applies that same always-on-by-default discipline to systematic-debugging activation). Before any such investigation proceeds, the `superpowers` plugin + the `superpowers:systematic-debugging` AND `superpowers:using-superpowers` skills **MUST** be confirmed installed + loadable per clause (C) – observed in the live capability list, never assumed (§ 11.4.102 anti-bluff). "I did not activate systematic-debugging because the operator did not ask for it" is itself a § 11.4.103 (D) violation: clause (A) is NOT a prompted option, it is a standing mechanical default engaged automatically the instant a problem surfaces. § 11.4.104 (D) STRENGTHENS clause (A) – it makes the activation trigger prompt-independent and standing-default; it does not weaken it: every clause-(A) four-phase-arc and Iron-Law ("NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST") requirement continues to apply in full.

Composes with § 11.4.105 (test-interrupt-on-discovery – § 11.4.106 says STOP the cycle the moment a defect is discovered; § 11.4.107 says the FIRST action after stopping is to launch systematic-debugging), § 11.4.108 (no-guessing – § 11.4.109 (A) is the procedural enforcement that produces the facts § 11.4.110 demands; the forbidden `LIKELY` / `transient` / `flaky` vocabulary is exactly what guess-and-retry emits), § 11.4.111

(demotion-evidence – a FAIL→lower-severity demotion requires same-conditions evidence, which is precisely what the systematic-debugging arc captures), § 11.4.101 (deep-web-research – runs inside the hypothesis phase), § 11.4.102 (TDD-fix – the RED test is written against the root cause the investigation proved), § 11.4.103 (subagent-driven – a systematic-debugging investigation is a natural subagent dispatch), § 11.4.104 (A) (iteration-speedup – already mandates Phase-1 forensic before any speculative patch; § 11.4.105 generalises that from "source patch" to ANY fix and adds the always-loaded skill-discovery + plugin-availability invariants), § 11.4.106 (multi-pass change-evaluation – the investigation feeds Pass 1 + Pass 2), § 11.4.107 (default-autonomous-loop – § 11.4.108 (D)'s always-on-by-default systematic-debugging activation mirrors § 11.4.109's from-first-prompt standing-default working mode).

Classification: universal (§ 11.4.109) – root-cause-before-fix, skill-discovery-before-task, and dependency-availability are reusable disciplines for ANY project on ANY runtime that exposes a skill/plugin ecosystem; the concrete skill names (`superpowers:systematic-debugging` , `superpowers:using-superpowers`) are the Claude Code instantiation, with the generic "or platform-equivalent" wording carrying the rule to other runtimes per § 11.4.110 cross-platform-parity reasoning.

-layer coverage per § 11.4.111 (b). Propagation gate `CM-COVENANT-114-102-PROPAGATION` enforces the literal anchor `11.4.102` across the canonical consumer fleet (parent + owned-submodule `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Paired § 11.4.112 meta-test mutation strips the `11.4.102` literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § . . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . . inheritance.

Non-compliance is a release blocker regardless of context.

No escape hatch – no `--skip-systematic-debugging`, `--guess-and-retry-OK`, `--symptom-patch-permitted`, `--skip-skill-discovery`, `--plugin-optional`, `--missing-plugin-is-warning` flag exists. The – – user mandate is unambiguous: the moment a problem is spotted the structured root-cause investigation skill activates, `using-superpowers` is always loaded, and every mandated plugin/dependency is installed and loadable – there is no path by which a fix is proposed before the cause is proven.²

2026 05 29

§ . . . – **Continuous parallel-stream working routine** (User mandate, 2026 05 29 – –)

Short tag: `continuous-parallel-stream-routine` .

Forensic anchor – verbatim user mandate

(– –):

"Do this working approach continuously and make it part of regular working routine and add it to the Constitution Submodule documented fully."

11 4 87

This anchor promotes the proven multi-stream operating pattern – repeatedly demonstrated effective across prior autonomous rounds – from an ad-hoc practice into the project's **standing default working routine**. § . . . told the agent to keep working in an endless loop; § . . . told it never to sit idle when something is

progressable; § . . . / § . . . detached push and test execution; § . . . catalogued what is safe to do during a long build. § . . . binds those into a single always-on routine and adds the load-bearing invariant the prior anchors only implied: **the main work stream MUST always stay FREE**. It is the conductor's standing operating mode, not a per-request opt-in.

Strengthened (User mandate, - -): the parallel-stream floor is raised from $\geq$ to $\geq$ and an **auto-backfill rule** is added – the moment any one stream is FULLY done, a new stream MUST immediately start and take its place, so the count NEVER drops below while actionable workable items remain. Verbatim mandate: *"make sure there are always at all times at least parallel working streams parallel with main work stream as long there are workable items we can do! as soon as one working stream is fully done, another new MUST start work and take its place."*

11 4 90

(A) Main stream stays FREE – ALL commit AND push run detached, never blocking. Every commit operation AND every push operation MUST run in the background (detached via `nohup ... > <log> 2>&1 & + disown`, per § . . . 's commit-lock-release-immediately + detached-push seam). The conductor's main work stream MUST NOT block on a commit landing or a push completing – the moment the local commit is durable, the main stream returns to the priority queue while the push fans out to every upstream in the background. Blocking the main stream on a push (or a slow mirror) is a § . . . violation (the exact - - § . . . forensic harm – a - hour flock hold – generalised to the routine layer).

(B) $\geq$ parallel background streams at all times whenever actionable work exists + auto-backfill. Whenever the queue holds three or more non-contending actionable items, the conductor MUST run **at least three** parallel background

streams alongside the main stream (subagent-driven per § 11.4.74 / § 11.4.20, isolated per the § 11.4.58 PWU worktree model + § 11.4.84 working-tree quiescence). **Auto-backfill rule (User mandate, § 11.4.84 - § 11.4.84):** the moment any one stream is FULLY done (its workable item closed with anti-bluff evidence per sub-rule (E), or merged), a new stream MUST IMMEDIATELY start and take its place by claiming the next-highest-priority non-contending item from the queue – the active-stream count NEVER drops below 3 while actionable workable items remain. Dropping below 3 active streams while three-plus disjoint actionable items wait – or serialising work that could run in parallel – is a § 11.4.84 violation. The count may fall below 3 ONLY when there genuinely are no remaining non-contending actionable workable items OR every remaining item is externally blocked (per § 11.4.84 / § 11.4.84 / § 11.4.84 idle-only-when-blocked). The § 11.4.84 % memory ceiling and the § 11.4.84 -concurrent-agent cap bound the stream count from above (so the standing band is 3 – background streams).

(C) Most-critical + most-visible first; audio always top. Streams are filled in strict priority order – most-critical and most-user-visible items first – and **audio fixes are always the apex priority** per § 11.4.84. When the conductor faces a choice of which item to dispatch onto a free stream, the highest-severity / most-user-visible / audio item wins; research, refactors, and low-stakes tasks fill remaining capacity only after the critical set is dispatched (per § 11.4.84 priority-ordered iteration).

(D) Safe-during-build scope only while a heavy build runs. While the 4GB containerised AOSP rebuild (§ 11.4.84) – or any heavy `gradle` / `m -j` build – is in flight, the parallel background streams MUST restrict themselves to

the \$. . . SAFE catalogue: investigation / forensic analysis, documentation, test-authoring, gate + paired-mutation authoring, read-only on-device probes, submodule edits, web research, workable-items DB ops, and backgrounded pre-build / meta-test runs. The streams MUST NOT launch a second gradle / m -j / containerised build concurrent with the running one (\$. . . concurrency-hardening – a second heavy JVM/Soong build storms the host and risks the \$. . . OOM cascade). UNSAFE-during-build operations (\$. . . (α)-(θ)) are deferred to between-build windows.

(E) Heavy anti-bluff on every closure. Every item closed by a parallel stream MUST carry real captured evidence and survive the anti-bluff covenant family – root cause proven or explicitly marked UNCONFIRMED: / UNKNOWN: / PENDING_FORENSICS: per \$. . . , captured-evidence per \$. . . + \$. . . , deterministic consistency per \$. . . , paired \$. . . meta-test mutations, and the structured root-cause investigation of \$. . . the moment any problem surfaces. A stream that closes an item on metadata-only / configuration-only / absence-of-error / grep-without-runtime evidence is a \$. . . PASS-bluff regardless of how parallel the routine looked.

(F) Idle ONLY when genuinely externally blocked. The conductor idles ONLY when every queued item is genuinely blocked on an external result it cannot accelerate (hardware access, network upstream, an in-flight build / test / push completion) OR the operator issues STOP OR \$. . . host-session-safety demands it (per \$. . . \$. . . (A) + \$. . .). "Nothing visible to do" with progressable non-contending items in the queue is NEVER a valid idle reason. Posing a question and going idle while non-blocked items remain is a \$. . . + \$. . . + \$. . . violation; an

unavoidable block parks one work unit, not the whole loop
(§ 11.4.103 maximize-progress-while-blocked).

Composes with § 11.4.102 (parallel-development PWU pipeline – § 11.4.101 is its standing-routine framing), § 11.4.100 / § 11.4.99 (subagent-driven default – the parallel streams ARE subagents), § 11.4.98 (audio top-priority – sub-rule (C)), § 11.4.97 (endless-loop autonomous work – § 11.4.96 is the always-on operating mode that anchor governs), § 11.4.95 (background-push – sub-rule (A) is its routine-layer generalisation), § 11.4.94 (background test execution – backgrounded tests are one stream class), § 11.4.93 (zero-idle priority-first parallel-by-default – sub-rules (B)+(C)+(F) bind it), § 11.4.92 (safe-parallel-work-with-long-build catalogue – sub-rule (D) cites it), § 11.4.91 (maximum-use-of-idle-time – sub-rule (F)), § 11.4.90 (autonomous-decision-over-blocking – maximize-progress-while-blocked), § 11.4.89 (systematic-debugging on any spotted problem inside any stream), § 11.4.88 (priority-ordered iteration), § 11.4.87 (working-tree quiescence for concurrent subagent commits), § 11.4.86 / § 11.4.85 / § 11.4.84 (host-session safety bounds stream count + forbids a second heavy build), § 11.4.83 (data-safety for any destructive op a stream would attempt).

Classification: universal (§ 11.4.82) – continuous parallel-stream working with a free main stream, priority-first dispatch, and anti-bluff closures is a reusable operating discipline for ANY project on ANY runtime that supports background execution + subagent delegation; nothing in the routine is hardware-, vendor-, or region-specific.

-layer coverage per § 11.4.81 (b). Propagation gate
CM-COVENANT-114-103-PROPAGATION enforces the literal anchor

11.4.103 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired § . meta-test mutation strips the 11.4.103 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --block-main-stream , --synchronous-commit , --synchronous-push , --single-stream-only , --skip-parallel-streams , --serialise-actionable-work , --idle-without-queue-survey flag exists. The – – user mandate is unambiguous: continuous parallel-stream working with a free main stream is the regular working routine, not a per-request opt-in.

§ . . – Participant identity, attribution & notification-tagging (User mandate, – –)

Short tag: participant-attribution-tagging .

Forensic anchor – verbatim user mandate
(– –):

"Every supported messenger must relate messages to participants (Subscribers/Users); the same logical person may have a different username on every messenger. Workable items must carry who created them and who they are assigned to. Notifications must @-tag

the right participant – but never the operator (who drives the system) and never the system agent."

This anchor binds participant identity, workable-item attribution, and notification @-tagging as MANDATORY constraints on every consumer that ships a messenger/ notification surface (Herald and its flavor binaries are the reference implementation; a consuming project and any future messenger-bearing project inherit per § . . .). The detailed normative spec is the Herald design contract `docs/design/PARTICIPANT_ATTRIBUTION.md` (model, columns, attribution rules, tagging matrix) – this section restates its load-bearing constraints; implementations code against the contract, not against a paraphrase.

(A) Participant identity – logical person + per-channel handle. Every supported messenger MUST relate inbound + outbound messages to a **Participant** (a logical Subscriber/ User). A single logical person/agent MAY carry a DIFFERENT username on every messenger; identity is therefore modelled as a logical subscriber (canonical, messenger-neutral `handle` , `kind ∈ {human, agent, service}`) PLUS per-channel aliases (`channel` , `channel_user_id` , the per-channel `@username` used for tagging). The **canonical handle** is the string stored in attribution columns – the closed set is `Claude` (the reserved system-agent sentinel; `kind=agent`) OR a human's canonical handle (a subscriber `@username` , messenger-neutral, resolved per-channel via the alias table).

(B) Operator designation – env var, not a DB flag. The **Operator** is the one human who drives the system via the agent CLI. The Operator is designated by the environment variable `HERALD_<CHANNEL>_OPERATOR_USERNAME` (e.g. `HERALD_TGRAM_OPERATOR_USERNAME` , `HERALD_SLACK_OPERATOR_USERNAME`) – per messenger, NOT a database flag. The Operator is a

normal Participant whose canonical handle equals that env value.

(C) Workable-item attribution – created_by + assigned_to .

Every workable item MUST carry two canonical-handle columns: created_by and assigned_to . Attribution rules: opened via the agent CLI prompt (operator-driven) → created_by = Operator ; opened by the system/agent detecting an issue/task/improvement/missing-feature → created_by = "Claude" ; received THROUGH the system as a subscriber message → created_by = the sender's resolved canonical @username . assigned_to defaults to the Operator and MAY be overridden explicitly (a prompt/message assigning to @someoneelse). Both columns store the canonical handle string and are self-contained in the SSoT + its Markdown export. Legacy items predating this anchor carry empty ("") attribution and MUST still parse + validate.

(D) Notification @-tagging matrix. On any workable-item event, the outbound notification dispatched to each messenger group @-tags the participant(s) who must be aware, resolved to that channel's @username :

```
mentions = {}
if assigned_to is a human handle AND assigned_to != Operator:
    mentions[assigned_to] = {}
    if created_by is a human handle AND created_by != Operator AND created_by != "Claude":
        mentions[assigned_to][created_by] = {}
        # "Claude" is NEVER tagged (it is the system agent).
        # the Operator is NEVER tagged (drives the system; no self-ping).
        # de-dup; for each mention resolve the @username on the target channel – skip if the part
```

This satisfies the operator's stated cases exactly:
assigned-to-Operator → no tag; opened-by-Operator-
assigned-to-another → tag the assignee; opened-by-a-non-
Operator-non-Claude-subscriber → tag the creator.

(E) Anti-bluff (MANDATORY, composes with \$ / the end-user quality covenant). Every layer ships unit + integration + E E + full-automation tests producing REAL captured evidence (no metadata-only / absence-of-error

PASS, no false positive/negative): real SQLite round-trip with the new columns (byte-identical, including legacy fixtures that DON'T carry the fields); the tagging matrix proven by a truth-table test with a mutation that flips one cell forcing a FAIL; an E E real-item-event → real-dispatched-message asserting the exact @username s AND a NEGATIVE case proving the Operator is NOT tagged. Evidence committed under docs/qa/<run-id>/ .

Composes with § . + §§ . .
(end-user.quality / anti-bluff covenant – sub-rule (E) is bound by it), § . . / § . . (captured evidence), § . . (deterministic consistency),
§ . . (workable-item field clarity),
§ . . / § . . (the SQLite workable-items SSoT the created_by / assigned_to columns extend),
§ . (paired-mutation proof of the tagging matrix).

Classification: universal (§ . .) – participant identity, item attribution, and notification @-tagging is a reusable discipline for ANY project that ships a messenger/notification surface with a workable-items registry; the env-var operator designation + canonical-handle model + tagging matrix are vendor-neutral. Projects with NO messenger surface inherit the anchor latently (it binds the moment they ship one) – the § . .
"principle binds even absent the surface" restatement pattern.

-layer coverage per § . . (b). Propagation gate CM-COVENANT-114-104-PROPAGATION enforces the literal anchor 11.4.104 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired § . meta-test mutation strips the 11.4.104 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § . . . in the HelixConstitution submodule; detailed spec Herald docs/design/PARTICIPANT_ATTRIBUTION.md . All consuming projects restate + cite via § . . . inheritance.

Non-compliance is a release blocker. No escape hatch – no `--skip-attribution` , `--no-participant-tagging` , `--tag-operator-anyway` , `--attribution-later` flag exists.

2026 05 31

§ . . . – Natural-language intent recognition & clarification (User mandate, 2026 05 31 – –)

Short tag: `intent-recognition-clarification` .

Forensic anchor – verbatim user mandate

(– –):

"Users must NOT need to know command syntax (no `COMMAND: ...` prefix). They send a clear natural-language message; the System determines the intent. The System recognizes the commands it has; if none matches it infers the exact intent; if it is totally unable it replies, tags the user (`@user ...`), and asks to clarify precisely. We MUST always do our best to determine exact intent so we never annoy end users. This is a CORE part of the System."

This anchor binds natural-language intent recognition as a MANDATORY constraint on every consumer that ships a messenger/notification-driven command surface (Herald and its flavor binaries are the reference implementation; a consuming project and any future messenger-bearing project inherit per § . . .). The detailed normative spec is the Herald design contract docs/design/INTENT_RECOGNITION.md

(three-tier resolution, command set, the `clarify` action, envelope wiring) – this section restates its load-bearing constraints; implementations code against the contract, not against a paraphrase.

(A) No required command syntax. Users **MUST NOT** be required to know any command syntax – no `COMMAND:` prefix, no rigid grammar. They send a clear natural-language message; the System is responsible for determining the intent. Any design that forces an end user to memorize syntax to be understood is a violation of this anchor.

(B) Three-tier intent resolution (the mandatory discipline). Every inbound subscriber message **MUST** be resolved to exactly one action via three tiers, in order – the first that succeeds wins:

```
TIER 1  command recognition – recognize the System's existing command set from
                                natural language and map it to a structured action
                                (a confident, deterministic match → that action).
TIER 2  intent inference      – when NO command matches, infer the exact intent from
                                the message (the LLM dispatch maps natural language to
                                the right action) – NEVER guessing.
TIER 3  clarify (fallback) – when neither a command nor a confident intent can be
                                determined, REPLY to the original message, TAG the
                                sender (`@username`, resolved via the §11.4.104
                                IdentityResolver) and ask a PRECISE clarifying question
                                that NAMES the candidate intents. No guessing, no
                                silent drop.
```

Tier 3 is the anti-annoyance guarantee: the user is never ignored and never has to learn syntax – at worst they receive a friendly, specific "@user, did you mean X or Y?".

(C) Never guess, never drop. The System **MUST NEVER** guess an action – a wrong action is worse than a clarifying question (composes with § 11.4.104 no-guessing). The System **MUST NEVER** silently drop or ignore an inbound message. We always do our best to determine the exact

intent so end users are never annoyed; only genuine ambiguity reaches Tier 1, and Tier 1 always replies-tags-and-asks rather than failing silently.

(D) Anti-bluff (MANDATORY, composes with § 4.1.1 / the end-user quality covenant). Every tier ships unit + integration + E 2 E 3 + full-automation tests producing REAL captured evidence (no metadata-only / absence-of-error PASS, no false positive/negative): Tier 1 proven by a truth-table of natural-language messages → expected action+fields PLUS the conservative negatives that MUST fall through to "no match"; Tier 2 proven by an E 2 E 3 ambiguous-message dispatch whose recorded reply body is EXACTLY @<sender> <specific question> (the user is tagged + asked, not ignored) AND a NEGATIVE proving a clear command does NOT trigger clarify; a paired § 4.1.1 mutation that breaks the recognizer's confidence guard (forcing a false-match) OR drops the clarify tag MUST FAIL a test. Evidence committed under docs/qa/<run-id>/.

Composes with § 4.1.1 + § 4.1.2 ..§ 4.1.3 (end-user quality / anti-bluff covenant – sub-rule (D) is bound by it), § 4.1.4 (no-guessing – a wrong action is worse than a clarifying question), § 4.1.5 (participant identity – the clarify reply tags the sender via the IdentityResolver), § 4.1.6 / § 4.1.7 (captured evidence), § 4.1.8 (full-automation – every tier's test is self-driving), § 4.1.9 (paired-mutation proof of the confidence guard + clarify tag).

Classification: universal (§ 4.1.10) – natural-language intent recognition with a three-tier recognize→infer→clarify discipline is a reusable interaction contract for ANY project that ships a messenger/command surface; the no-syntax mandate, never-guess rule, and reply-tag-and-ask fallback are vendor-neutral. Projects with NO messenger surface inherit the anchor latently (it binds the moment they ship one) – the

§ . . "principle binds even absent the surface" restatement pattern.

-layer coverage per § . . (b). Propagation gate CM-COVENANT-114-105-PROPAGATION enforces the literal anchor 11.4.105 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired § . meta-test mutation strips the 11.4.105 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § . . in the HelixConstitution submodule; detailed spec Herald docs/design/INTENT_RECOGNITION.md . All consuming projects restate + cite via § . . inheritance.

Non-compliance is a release blocker. No escape hatch – no --require-command-syntax , --guess-intent-ok , --skip-clarify , --drop-on-ambiguous flag exists.

§ . . – Docs Chain – mechanical documentation/DB sync engine (Operator mandate, – –)

Short tag: docs-chain-mechanical-sync .

Forensic anchor – operator mandate (– –):

Docs Chain is the canonical mechanical enforcer of the documentation-sync mandates. Consuming projects MUST use the engine instead of ad-hoc per-project doc-sync scripts; they register their own chains via per-context YAML; the engine never fakes a transform and every sync / verify carries real captured evidence.

Docs Chain – the `vasic-digital/docs_chain` engine (a universal, Go-implemented, **bidirectional document-and-database dependency-propagation engine**) – is the canonical mechanical enforcer of the family of documentation-sync mandates this Constitution already declares. This is a **MANDATORY** constraint binding every consuming project. The detailed normative spec is the engine's own documentation set – `docs/CONSTITUTION_INTEGRATION.md` (distribution + inheritance + the anchor mapping table) and `docs/USE_CASE_CATALOGUE.md` (the ready-to-use chain recipes); this section restates its load-bearing constraints, it does not redefine them. Implementations code against the engine's contract, not against a paraphrase.

(A) Use the engine, never ad-hoc scripts. Consuming projects **MUST** drive their documentation/DB sync through Docs Chain – consumed **by reference** via the flat-layout sibling (the `~/Projects/docs_chain` working tree during local development / the path the constitution exposes once the Phase- submodule pointer lands), inherited the same way `$ . .` inherits the `codegraph_*` scripts: referenced at the constitution-exposed path, **NEVER copied** into the consumer. Ad-hoc per-project doc-sync scripts (`sync_issues_docs.sh` , `generate_*_summary.sh` , `sync_*_export.sh` , `update_readme_doc_links.sh` , et al.) are superseded by the engine and **MUST** be retired as each matching context is registered.

(B) Consumer-owned contexts via config. The engine is project-agnostic; it carries **NO** project-specific path or chain. The consumer registers its chains as data via `.docs_chain/contexts/*.yaml` at its own project root – one independent context per file – per the `$ . .` decoupling rule (the submodule never carries project-specific context; the consumer injects its chains via config). `state.json` (at `<project-root>/docs_chain/state.json` , regenerated on demand by `docs_chain sync`) and

the `*.docs_chain.tmp` staging sidecars are **gitignored**; the consumer owns those `.gitignore` entries.

(C) **The sync anchors it mechanizes.** Docs Chain is the mechanical engine behind a family of existing documentation-sync anchors; once a consumer registers the matching context it enforces the anchor **in place of the ad-hoc script** – see the `docs/CONSTITUTION_INTEGRATION.md` mapping table: `$` `Issues_Summary` (Issues_Summary always-sync), `$` `Fixed_Summary` (Fixed_Summary parity), `$` `Status.md` (Status.md maintenance), `$` `Status_Summary` (Status_Summary two-audience parity), `$` `README` (README doc-link section), `$` `README` (README always-sync export), `$` `documentation` (documentation composite-sync), `$` `universal` (universal markdown export), `$` `roster/corpus` (roster/corpus auto-sync), `$` `workable-items` (workable-items DB single source of truth), `$` `DB` (DB tracked + WAL-checkpoint commit), `$` `CONTINUATION` (CONTINUATION maintenance), and `$` `document` (document revision header – exports kept in sync). Engine-wide guarantees: change detection by **content hash**, **NOT mtime** (`$` `mtime`); **atomic-rename + SQLite-transaction commit** with rollback on any transform error and optional hardlink backup (`$` `hardlink` zero-risk data safety); a both-dirty **sync** pair surfaces a **conflict – never a silent merge** (exit `1`, `$` `no-guessing`); **verify** is the deterministic CI / pre-build sink-side gate over byte-stable transforms (`$` `verify`); per-run captured evidence lands at `qa-results/docs_chain/<run-id>/` (`$` `qa-results/docs_chain/<run-id>/`).

(D) **It does NOT replace authoring discipline.** Docs Chain replaces only the **mechanical sync** an anchor requires – it does NOT replace the authoring discipline of any anchor. The source author still writes the `$` `revision` header into the source; the engine then keeps every derived export in sync. A consumer must not treat the engine as licence to skip authoring.

(E) Anti-bluff (MANDATORY, composes with § 1.1 / the end-user quality covenant). Docs Chain NEVER fakes a transform: a missing `pandoc` / `weasyprint` (or any absent tool) surfaces a typed `ToolAbsentError` (`IsToolAbsent`) and an honest § 1.1.1 SKIP-with-reason – **never a fake PASS**, never a partial write. Every `sync` / `verify` run carries real captured evidence (the `qa-results/docs_chain/<run-id>/` artefact). A metadata-only / "absence-of-error" / config-only PASS at the sync layer is a § 1.1.1 PASS-bluff.

Status note (§ 1.1.1 no-guessing). The engine's Phases 1-4 (core DAG + content-hash + adapters + atomicity/orchestrator) AND the `cmd/` CLI / per-context YAML loader (Phase 5) are IMPLEMENTED + tested – `cmd/docs_chain/main.go` registers the `sync` `./.` `verify` / `diff` (alias of `verify`) / `doctor` subcommands wiring the config loader → state → orchestrator, and `go test -count=1 ./...` is GREEN at 1.1.1 / packages (incl. `internal/orchestrator` `spec_wiring_test`, both PASS). FACT (§ 1.1.1): this CLI/loader lives in the working tree but is UNTRACKED (`git ls-files docs_chain/cmd/docs_chain/main.go` is empty) and the engine is NOT yet a registered git submodule (`git submodule status` shows no docs_chain entry; in-tree at `./docs_chain`, HEAD = parent HEAD). The constitution-submodule distribution (Phase 6) remains PLANNED + OPERATOR-GATED. This anchor states the binding contract; consumers wire the engine as each phase lands, and MUST NOT claim working behaviour a phase has not yet shipped.

Classification: universal (§ 1.1.1) – a content-hash bidirectional doc/DB propagation engine is a vendor-neutral mechanism reusable by ANY governed project. Projects with no derived-export or DB-sync surface inherit the anchor latently (it binds the moment they ship one) – the § 1.1.1 "the principle binds even absent the surface" restatement pattern.

-layer coverage per § . . (b). Propagation gate CM-COVENANT-114-106-PROPAGATION enforces the literal anchor 11.4.106 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired § . meta-test mutation strips the 11.4.106 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Composes with § . . + §§ . . .
 (end-user quality ./ anti-bluff – sub-rule (E) is bound by it), § . . . (no-guessing – both-dirty sync → conflict, never silent merge), § . . /
 § . . / § . . / § . . . 11.4.106 /
 § . . / § . . 11.4.106 / § . . /
 § . . / § . . / § . . /
 § . . / § . . 11.4.106 / § . . (the documentation-sync anchors it mechanizes), § . . (engine/context decoupling), § . . (inherited-by-reference, never-copied – the codegraph pattern), § . (atomic commit + rollback), § . . (deterministic verify gate), § . . 1006 07 10 / § . . (captured evidence), § . (paired-mutation propagation proof).

(F) Write-seam hook enforcement + real-time-sync honest boundary (extension, – – ; operator mandate – the Continuation / memory / knowledge artefacts MUST ALWAYS be in sync). The documentation/DB sync this engine mechanizes MUST be ENFORCED AT THE WRITE SEAMS, never left to agent vigilance: the consuming project MUST wire the engine's verify (and, where safe, sync) into () the COMMIT seam – a pre-commit / commit-wrapper hook that REFUSES a commit whose staged set touches a chain SOURCE while its derived exports / DB rows are stale; () the BUILD seam – the pre-build gate; and () the CONSTITUTION-PULL seam – the § . . .
 post_update_hook.sh invocation. A stale CONTINUATION.md

11 4 131

(`$ .`), a stale standing session-resumption file (`$ .`), or a stale memory / knowledge artefact while the live state has ALREADY moved is a violation of the SAME severity class as a stale export – the resumption guarantee silently breaks and a fresh session resumes on fiction. Honest boundary (`$ .`): hook-wired seam enforcement makes the sync EVENTUALLY-CONSISTENT-AT-EVERY-WRITE – it is NOT literal continuous real-time; genuine continuous real-time sync requires a filesystem-watch daemon (inotify / FSEvents watcher running the chain on change). A project that CLAIMS "real-time sync" while shipping only seam hooks is bluffing the mechanism: state the seam-based guarantee honestly, and carry the watch daemon as a TRACKED upgrade path (`$ .` – driven to completion or explicitly closed, never left unwired).

Canonical authority: this Constitution.md `$ .` in the HelixConstitution submodule; detailed spec the Docs Chain engine docs `docs/CONSTITUTION_INTEGRATION.md` + `docs/USE_CASE_CATALOGUE.md` (`vasic-digital/docs_chain`). All consuming projects restate + cite via `$ .` inheritance.

Non-compliance is a release blocker. No escape hatch – no `--ad-hoc-sync-ok` , `--skip-docs-chain` , `--fake-transform` , `--sync-evidence-optional` flag exists.

2024 06 02

`$ .` – **Anti-bluff AV/test-validation techniques mandate** (User-driven research, `- -`)

Short tag: `av-liveness-oracles` .

Forensic anchor – captured incidents (`- -` , genericised): a test asserting "media is playing on the target output" PASSED on the strength of a SINGLE captured

frame that showed "a picture" – but the picture was a FROZEN / STALE frame left over from the previously-played content (a stale-producer / stuck-decoder failure mode), so the feature was broken for the end user while the test was green. A sibling incident: media briefly FLASHED on the WRONG output (the source/primary surface) for ~ 100ms before routing to the intended output, and no test sampled the wrong output during the transition window, so the glitch shipped undetected. A third class: a comparator (color / freeze / OCR oracle) that was itself broken – it PASSED on a deliberately-degraded artefact – meaning the analyzer, not just the feature, was the bluff. § 1.1.1 mandates *captured* evidence and a baseline presence pass; § 1.1.2 raises the bar to **liveness + correct-routing + self-validated-analyzer** so a single still frame, a wrong-output flash, or a broken oracle can no longer pass.

The mandate. Every test asserting that audio or video output is genuinely playing / advancing for the end user MUST satisfy ALL of the following – each is a § 1.1.1 PASS-bluff to omit:

- . **A single captured frame is NOT proof of playback.** Prove LIVE, ADVANCING frames over a steady-state window ($\geq N$ frames across $\geq$ a few seconds) via a **freeze-detection oracle** – a near-duplicate detector (the platform's frame-difference / `freezedetect` -class filter OR a perceptual-hash adjacent-frame distance), **NOT a byte-identical compare** (a re-encoded / 100%-pixel-noisy stale frame defeats byte-identity; keep byte-identity only as a zero-cost early-out pre-filter). Adjacent frames MUST differ above a calibrated threshold; $\geq$ a sustained interval of near-identical frames $\Rightarrow$ FROZEN $\Rightarrow$ FAIL.
- . **An independent frame-advance counter from the platform's compositor / decoder telemetry** MUST increase across the window – an oracle in a DIFFERENT domain from the pixel

check (compositor frame-timestamp advance, decoder framesDecoded / framesRendered, or equivalent). A flat counter over the window $\Rightarrow$ stuck decoder $\Rightarrow$ FAIL, even if the pixels appear to move. Keeping a pixel-domain oracle AND a telemetry-domain oracle is metamorphic redundancy: a coordinated false-PASS now requires two independent failures.

- **Loading / buffering is a DISTINCT state – wait for genuine playback-start before judging liveness.** Never false-FAIL a still-loading stream, never false-PASS a loading spinner. Poll the decoder frame-advance counter / first-advancing-frame / buffering flag until playback genuinely starts (generous per-class timeout). Timeout + content endpoint unreachable $\Rightarrow$ SKIP-with-reason per § 11.4.5 (network_unreachable_external / geo_restricted); timeout + reachable $\Rightarrow$ FAIL with captured evidence. Only AFTER confirmed playback-start run the liveness checks over the steady-state window.
- **Not-stale-from-previous cross-check.** Before asserting new content N is live, confirm its first frame DIFFERS from the previous content's last captured frame (perceptual-hash mismatch). Identical $\Rightarrow$ the "new" content is a leftover frozen frame $\Rightarrow$ FAIL. Critical for switch-between-items sequences.
- **Measured FPS / no-lost-frames within tolerance** of the source's expected rate (where the source rate is knowable) or a sane floor otherwise; drop ratio above budget $\Rightarrow$ FAIL (§ 11.4.5 frame-health).
- **No-flash-on-the-wrong-output during a routing transition.** When content is meant to appear ONLY on a target output, sample the NON-target output at high frequency (e.g. every ~ 100 ms) across the transition window; ANY sampled wrong-output frame containing the content $\Rightarrow$ FAIL (brief-flash glitch).

(Where a content-protection regime legitimately forces a particular output, that regime is the expected case, not a fail – classify it explicitly, never by guessing per §)

- . **Drive tests through the realistic user path (feed / UI), not deep-link shortcuts.** Routing / transition / alternation bugs manifest on the path a real user takes (selecting from a feed / list via the UI); deep-link / direct-dispatch shortcuts bypass the very transition paths where these bugs live. Where the UI is not introspectable, fall back AND flag operator-attended per § . . . – never fake-PASS. (This is the § . . . UI-driven discipline applied to the liveness battery.)
- . **The oracle problem – use metamorphic relations when there is NO golden source.** For content the project does not own (streamed media), there is no golden reference, so full-reference comparison is impossible. Encode metamorphic relations instead: the SAME content captured on output-A vs output-B must be structurally equivalent (after resolution normalisation); a paused stream's frame-advance counter must STOP; a $\times$ speed must roughly double the frame-advance rate; the same clip played twice must produce equivalent liveness metrics. Each is a relation, not a fixed expected value – robust to content variety.
- . **Full-reference quality metrics vs a golden source for content the project DOES own.** Where a source reference exists (local-file playback of owned media, a test corpus), assert perceptual-quality / color fidelity (SSIM / VMAF / ΔE -class) against the temporally-aligned source so a pale / desaturated / wrong-range / soft-scaled render is caught across motion, not just at one sampled frame.

- Mutation-test every analyzer with a golden-good + golden-bad fixture pair** so the analyzer itself provably cannot bluff. For every comparator (freeze / motion / color / sharpness / aspect-ratio / FPS / OCR / audio-presence) keep a clean control fixture (must PASS) AND a deliberately-degraded fixture (must FAIL – frozen clip, black-gap clip, desaturated clip, downscaled clip, stretched clip, half-rate clip, overlay-baked frame, silent track). An analyzer that PASSes its golden-bad fixture is a bluff gate. This is the § . paired-mutation discipline applied to the AV *analyzers*, not only the grep gates. Document any behaviourally-equivalent mutant expected NOT to flip so a survivor there is not a false alarm.
- Per-channel audio integrity, not a single aggregate number.** A single aggregate RMS misses a dead channel in a multichannel stream and a near-silent / DC-offset stream that scrapes past a bare floor. Assert per-channel RMS / loudness (an EBU-R -class integrated/momentary loudness + per-channel level + silence-region census + DC-offset / NaN-Inf check); EVERY declared channel above floor. On-device, take an underrun / XRUN census from the audio subsystem – a starved pipeline is a defect even if the captured waveform "has sound".
- OCR overlay / subtitle detection needs a per-word confidence floor + a region-of-interest** to avoid BOTH a false-positive (noise read as a hostile overlay ⇒ false-FAIL of a working feature) AND a false-negative (a real overlay missed ⇒ false-PASS). Discard sub-confidence words; OCR the ROI where the text lives (bottom-third for subtitles, centre for dialogs); emit per-finding confidence so a reviewer sees WHY a finding fired. Where a subtitle must appear on one output and NOT another, assert both sides (a property relation per point).

. **Thresholds MUST be calibrated on the project's own fixtures, not hardcoded from literature (§ 4.1.1.1 no-guessing).** Every freeze / motion / SSIM / VMAF / loudness / OCR-confidence threshold is recorded as captured evidence derived from a known-good and known-bad sample on the project's own capture path – never a literature constant asserted as fact. The capture path's own re-encode loss is part of the baseline.

Honest gaps (§ 4.1.1.2). True physically-displayed (photon) frame rate has no clean software oracle – the compositor / decoder counters measure the *presentation pipeline*, not photons at the sink; a high-speed external camera rig is the only ground truth. The counters are sufficient to catch a stuck / half-rate decoder and MUST be used, with the photon-FPS gap flagged honestly rather than claimed. Broadcast-QC vendor tooling validates that freeze / duplicate / lost-frame / black detection is industry-standard practice, but the exact proprietary thresholds are undocumented; the project adopts the published open-method analogues and calibrates per point

.
-layer coverage per § 4.1.1.3 (b). Per-fix: pre-build gate (e.g. a `CM-AV-LIVENESS-NO-FROZEN-FRAME` -class gate asserting every test that claims output-is-playing references the liveness battery – freeze oracle + frame-advance counter + not-stale cross-check + FPS – not a single-frame check; plus an analyzer-self-validation gate wiring the golden-good/golden-bad fixtures into the project's meta-test) + on-device / runtime test + **paired § 4.1.1.4 meta-test mutation** (a test asserting output-is-playing with ONLY a single-frame check → gate FAILs; an analyzer that PASSes its golden-bad fixture → self-validation FAILs) + HelixQA Challenge entry for the user-visible feature. Every such PASS is emitted via `ab_pass_with_evidence` citing the motion / not-stale /

frame-advance / fps / per-channel-loudness / metamorphic / self-validation artefacts (feature classes `video_display`, `audio_output`, `subtitle_render` per § 4.17) – NEVER a single screenshot.

Classification: universal (§ 4.17) – freeze-detection oracles, frame-advance telemetry, loading-aware playback-start waits, not-stale cross-checks, metamorphic relations, full-reference quality metrics, golden-good/golden-bad analyzer self-validation, per-channel loudness, and confidence-floored OCR are platform-neutral AV / test-validation techniques reusable by ANY project that validates media playback or any pixel/audio output; the project supplies the concrete capture mechanism (its compositor / decoder telemetry, its sink-introspection API per § 4.17, its capture tool) and the calibrated thresholds per § 4.17.

Composes with § 4.17 (strict expansion of its video + audio quality bar to liveness + correct-routing), § 4.17 (no-guessing – thresholds calibrated + captured, loading-vs-frozen disambiguated by evidence, content-protection regime classified not guessed), § 4.17 (deterministic-consistency – the liveness battery runs N iterations identically), § 4.17 (positive sink-side / downstream evidence – the frame-advance counter / sink-introspection cross-check is its video analogue), § 4.17 (universal sink-side positive-evidence taxonomy – `video_display` / `audio_output` / `subtitle_render` classes), § 4.17 (stress + chaos – liveness under sustained load + transition chaos), plus § 4.17 (recorded evidence), § 4.17 (topology SKIP-with-reason for unreachable content), § 4.17 (sink-side introspection as one oracle), § 4.17 (UI-driven path), § 4.17 (autonomous validation; operator-attended fallback only

when UI not introspectable), § . (paired-mutation discipline applied to the analyzers).

Propagation. Propagation gate

CM-COVENANT-114-107-PROPAGATION enforces the literal anchor 11.4.107 across the consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md); paired § . meta-test mutation strips the 11.4.107 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § . in the HelixConstitution submodule. All consuming projects restate + cite via § . inheritance.

Non-compliance is a release blocker. No escape hatch – no --single-frame-proves-playback , --skip-liveness , --byte-identical-freeze-OK , --no-frame-advance-counter , --skip-not-stale-check , --allow-wrong-output-flash , --deep-link-shortcut-OK , --unvalidated-analyzer-OK , --aggregate-rms-suffices , --hardcoded-thresholds-OK flag exists.

§ . – Four-layer fix-verification + runtime-signature-as-definition-of-done mandate (systematic-debugging Phase . , – –)

Short tag: runtime-signature-definition-of-done .

Forensic anchor (genericised, – –). Across one batch, multiple fixes were "green" at every gate yet NOT working for the end user: a change present in the source tree and passed by both the pre-build (source) gate and the post-build gate never reached the boot-time kernel command-line embedded in the deployed boot artifact, so the corresponding output feature was dead despite all-

green; sibling fixes correctly built into the system image were masked by stale per-user overlay copies left from a previous deployment (the running code was the stale overlay, not the freshly-built system code), so they merely *looked* broken; deployment did not produce a clean state (the mutable overlay was not wiped), so the stale shadow survived; and the validation suite ran against whatever was running – which could be the stale shadow – so it reported green on code that was never actually exercised. Per `superpowers:systematic-debugging` Phase

. : when each fix reveals a fresh "fixed-but-not-working" in a *different* place, that is NOT a series of independent bugs – it is ONE architectural VERIFICATION flaw, and patching each symptom is thrashing.

The architectural flaw (FACT, traced). A fix must cross FOUR distinct layers, and "fixed" at one layer does NOT imply fixed at the next: () **SOURCE** – the change is committed in the source file (what a grep-the-source pre-build gate checks); () **ARTIFACT** – the change's BYTES actually landed in the produced build artifact (image / bundle / installer / boot artifact / embedded command-line); () **RUNTIME-ON-CLEAN-TARGET** – the change is active on a *clean* / *fresh* deployment, the layer the end user experiences, with no stale overlay shadowing the deployed code; () **USER-VISIBLE** – the feature actually works for the end user (the \$. . / \$. . captured-evidence layer). Green at layer is the cheapest, least conclusive signal; nothing being green at layer (or even layer) tells you about layers – . A gate that verifies ONLY the source layer is insufficient and is itself a \$. bluff surface.

The mandate (ALL must hold).

- . **Runtime-signature-as-definition-of-done.** A fix is DONE only when its declared **runtime signature** is verified on a CLEAN / fresh deployment (the layer the end user

experiences). Source-committed $\neq$ artifact-contains-it $\neq$ active-on-clean-target $\neq$ user-visible-working – these are FOUR distinct claims; proving one does NOT prove the next.

- **Every fix declares ONE machine-checkable runtime signature** – a single observable on a clean target that proves the fix is BOTH active AND working (a property read from the running system, a downstream/sink-side report per § 3.1.1, a captured-evidence assertion per § 3.1.2 / § 3.1.3, or a counter/state delta – never a re-grep of the source). The **registry of these per-fix runtime signatures is the SINGLE SOURCE OF TRUTH for "fixed"** – it REPLACES "a gate greps the source" as the definition of done.
- **Verification gates MUST span all four layers**, per fix: **source** (pre-build / pre-merge), **artifact** (post-build – assert the change's BYTES landed in the produced artifact, NOT merely that the source still contains the change; cf. the embedded-command-line / stripped-asset failure mode), **runtime-on-clean-target** (post-deploy – assert the runtime signature on a freshly-deployed clean target), **user-visible** (captured evidence per § 3.1.1 / § 3.1.2). A gate that asserts only the source layer is insufficient and is itself a § 3.1.1 bluff surface.
- **Eliminate the stale-deployment / shadow layer – by construction.** Deployment MUST yield a CLEAN state (wipe the mutable overlay / per-user mutable state that can shadow the deployed build) **OR** a pre-validation assertion MUST prove `running-artifact == built-artifact` (no stale override masks the freshly-deployed build) **BEFORE** any validation runs. Validation that runs against possibly-stale deployed state is INVALID and any PASS it produces is a § 3.1.1 PASS-bluff – the test exercised code that was never actually deployed.

- **Meta-rule (composes § 4.4.1 Phase 1).** ≥
 "fixed-but-not-working" discoveries within a single cycle are the signal of an architectural VERIFICATION flaw, NOT three independent product bugs. On the 1st such discovery the agent MUST STOP patching symptoms (per § 4.4.1 test-interrupt-on-discovery), fix the VERIFICATION pipeline itself (close whichever of the four layers is unverified), and re-certify EVERY item in the batch through the corrected pipeline on a clean target.
- **A batch is "validated" only after COMPREHENSIVE per-item runtime-signature verification on a clean baseline** – NOT after the batch-touched items' own tests pass. Spot-validating only the touched items (per § 4.4.1's full-suite reasoning) misses the cross-layer gap: a touched item can be source-green while its runtime signature is dead on a clean target.
- **Machine-created evidence at every layer (extension, § 4.4.1 – Point 1 of the 1-point mandate).** At each of the four verification layers (SOURCE / ARTIFACT / RUNTIME-ON-CLEAN-TARGET / USER-VISIBLE), the layer's PASS MUST cite MACHINE-CREATED, MACHINE-VERIFIABLE evidence PRODUCED AT THAT LAYER by the layer's own gate – a schema-parseable file (static-analysis JSON + review verdict at SOURCE; build log + fingerprint checksum + byte-check at ARTIFACT; on-target verdict + runtime-signature + telemetry at RUNTIME; OCR/vision/audio/sink-probe oracle output at USER-VISIBLE) – cited by path + sha256 + timestamp per § 4.4.1. A PASS at a higher layer that cites evidence from a lower layer ALONE is a § 4.4.1 wrong-layer violation and does NOT satisfy the higher layer's obligation. Every evidence analyzer is § 4.4.1 (1) golden-good / golden-bad / negative-control validated; an analyzer that passes its

golden-bad fixture voids every verdict it ever emitted (§ . . . self-validation). Operator eyeballing, narrative-only PASS, and prediction of success are FORBIDDEN – § . . . binds this at every gate. Propagation: the existing CM-COVENANT-114-108-PROPAGATION stays unchanged (literal 11.4.108 unchanged, stays GREEN); CM-RUNTIME-SIGNATURE-REGISTRY is EXTENDED with the per-layer machine-evidence-file-schema check; a new paired mutation that forges a runtime PASS whose only evidence is a source-grep transcript → the gate FAILs.

Classification: universal (§ . . .) – the four-layer source→artifact→runtime-on-clean-target→user-visible verification gap, the per-fix runtime-signature registry as the definition of done, the artifact-byte-presence gate, the clean-deployment-or-equal-artifact pre-validation assertion, and the $\geq$ -discoveries-means-architectural-flaw meta-rule are platform-neutral verification-pipeline disciplines reusable by ANY project that builds an artifact, deploys it, and validates the deployed result; the consuming project supplies the concrete artifact format, its clean-deployment / equal-artifact mechanism, and its per-fix runtime-signature observables per §

Composes with § . . . (FAIL-bluffs forbidden), § . . . (recorded evidence – the runtime signature IS the recorded evidence), § . . . (test-interrupt-on-discovery + retest-from-clean-baseline – clause 's STOP-and-fix-the-pipeline is its § . . . specialisation; "clean baseline" is exactly clause 's clean deployment), § . . . (captured-evidence quality – the user-visible layer), § . . . (no-guessing – every layer asserted by captured evidence, never assumed to have propagated), § . . . (no-fakes-beyond-unit – runtime signatures exercise the real deployed system),

§ 11.4.108 (full-suite retest – § 11.4.109 adds the cross-layer per-item runtime-signature dimension clause demands), § 11.4.110 (validate-recent-work-before-full-suite – the clean-baseline / equal-artifact pre-flight is its § 11.4.111 pre-condition), § 11.4.112 (deterministic consistency – the runtime signature verifies identically across N iterations), § 11.4.113 (autonomous-validation – runtime signatures are machine-checkable without an operator), § 11.4.114 (universal sink-side positive-evidence taxonomy – a runtime signature is a taxonomy-class observable), § 11.4.115 (systematic-debugging – Phase 1 architectural-flaw recognition IS clause 's trigger).

Propagation. Propagation gate

CM-COVENANT-114-108-PROPAGATION enforces the literal anchor 11.4.108 across the consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md); paired § 11.4.109 meta-test mutation strips the 11.4.108 literal from a consumer file → the gate FAILs. A recommended per-fix gate CM-RUNTIME-SIGNATURE-REGISTRY asserts every fix declares a runtime signature verified at the artifact + runtime-on-clean-target layers (not only the source layer); paired § 11.4.110 mutation downgrades a fix's verification to source-only → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § 11.4.116 in the HelixConstitution submodule. All consuming projects restate + cite via § 11.4.117 inheritance.

Non-compliance is a release blocker. No escape hatch – no --source-green-is-done , --skip-artifact-byte-check , --validate-against-running-state , --no-clean-deployment , --skip-runtime-signature , --spot-validate-touched-only flag exists.

§ . . – Mandatory Anti-Forgetting Enforcement: PreToolUse Guard Hook + Subagent Constitutional Preamble + Orchestrator Pre-Action Checklist (Operator mandate)

Short tag: `anti-forgetting-enforcement` .

Forensic anchor – operator mandate:

UNCONFIRMED: pending operator's verbatim anti-forgetting mandate quote (tracked: PENDING item in the current session's CONTINUATION.md handoff).

Background context (UNCONFIRMED quote omitted; factual description only): the motivating incident is documented in `docs/AGENT_GUARDRAILS.md` and in `scripts/hooks/guard-forbidden-commands.sh`'s in-source header. During an on-device-API build, emulator subagents ran raw host-direct `emulator / adb` commands instead of routing through the Containers submodule as required by the governing clauses (`$ .X / $ .V / $ .AG` in the consuming-project layer). The root cause was that the orchestrator **forgot to inject the rule into subagent prompts**. A rule the orchestrator forgets to paste is not enforcement – it is a social contract with the agent's memory, and that contract is broken by every cold-session start, context-window exhaustion, and sub-agent dispatch. The mechanical fix is twofold: () a **PreToolUse guard hook** that blocks the forbidden command classes at the tool-call boundary unconditionally, and () a **canonical subagent preamble document** that the orchestrator pastes verbatim into every dispatch, plus a **pre-action checklist** for the orchestrator's own actions. Together these are the "anti-forgetting" enforcement layer: the hook is the floor (commands cannot be executed even if the agent forgot

every rule), the preamble is the ceiling (the full live ruleset, not just what the hook can pattern-match). This two-layer model was introduced in a consuming project (commit family on `master`, `- - /`) and is explicitly designed to be universal – `scripts/hooks/guard-forbidden-commands.sh` carries `Classification: universal` in its header and is portable by design (no project-specific paths, no `jq` hard dependency).

Rule. Every HelixConstitution-consuming project MUST deploy the `$ . .` anti-forgetting enforcement layer consisting of three components, all maintained together:

(A) PreToolUse guard hook – the mechanical floor.

A shell script MUST be wired as a `PreToolUse` hook in the AI agent runtime's project-scoped settings file (e.g. `.claude/settings.json` for Claude Code) to intercept every `Bash` tool call before execution and BLOCK the forbidden command classes at the tool-call boundary. The hook:

- Reads the tool invocation as JSON from stdin (`.tool_name` + `.tool_input.command`).
- Exits `1` for non-Bash tools (pass-through) and empty/missing commands.
- For Bash tool calls, pattern-matches against the following mandatory blocked classes, each mapped to its governing clause:

Blocked class	Governing clause(s)	Examples
Raw host-direct Android emulator launch / APK install / in-	<code>\$. .</code> (Containers mandate) + any project-layer emulator clause (e.g. <code>\$.X / \$.V /</code>)	<code>emulator -avd ... ,</code> <code>\$ANDROID_HOME/emulator/emulator ... ,</code> <code>adb install ... ,</code> <code>adb -s ... install</code>

Blocked class	Governing clause(s)	Examples
strumentation run	\$.AG in a consuming project)	... , am instrument ...
Force-push / hook bypass / signature by- pass	\$. . (pre- force-push merge- first) + \$. . (pre- push fetch) + pro- ject-layer \$.T. equivalent	git push --force , git push -f , git push --force-with- lease , --no-veri- fy , --no-gpg-sign
Privilege es- calation	\$. no-sudo (project-layer \$.U equivalent)	sudo ... , su , su - , su -l ...
Host power management	\$ host-session- safety (all consuming projects)	systemctl suspend/ hibernate/poweroff/ reboot/halt/kill- user/kill-session , loginctl ... , pm- suspend/hibernate , shutdown ...

- . Exits (block) when a match is found; the stderr text printed is fed back to the agent as the refusal reason, citing the governing clause.
- . Supports a documented-exception escape hatch: a command containing the literal marker `# guardrails:allow <reason>` is WARNED (stderr) but NOT blocked – EXCEPT for host-power commands, which are categorically non-overridable.

The **canonical implementation** of this hook lives at `constitution/scripts/hooks/guard-forbidden-commands.sh` in this constitution submodule. Consuming projects **MUST** reference it at that path (inherited by reference per `$ . .`) – NEVER copy it locally (a copy diverges silently). The script is portable: no `jq` hard dependency (built-in `awk`

fallback), no consumer-project-specific paths, no hardcoded credentials.

Anti-bluff requirement: a hermetic test suite (at minimum cases) MUST verify every blocked class exits , every allowed command exits , the escape hatch fires for non-power classes, and the host-power class rejects even with the escape marker. Paired \$. mutation: remove the `emulator -avd` pattern from the hook → the emulator-gate case exits → test FAILs → restore → test PASSes.

(B) Subagent Constitutional Preamble – the semantic ceiling.

A canonical document MUST exist at `docs/AGENT_GUARDRAILS.md` (or the project's equivalent governance preamble path, recorded in the project's `CLAUDE.md`) containing:

- . **The SUBAGENT CONSTITUTIONAL PREAMBLE block** – a verbatim, self-contained text block covering the top-mandatory constraints that cannot be pattern-matched by the hook alone (anti-bluff intent, resource caps, evidence honesty, continuation maintenance, hardcoding prohibition, distribute gate, remote policy, no-guessing vocabulary, etc.) plus the constraint classes that the hook already enforces (emulators, force-push, sudo, host-power) for completeness.
- . **Clear instruction that the orchestrator MUST paste this block verbatim into every subagent dispatch.**

The preamble is the semantic complement to the hook: the hook catches command-class violations; the preamble pre-informs the agent of all rules it must follow before it even issues a command. Together they eliminate the "agent forgot because the orchestrator didn't paste it" failure class.

Anti-bluff requirement: the preamble MUST be structured so it can be mechanically verified to be present in `docs/AGENT_GUARDRAILS.md` by `check-constitution.sh` (or the project-equivalent sweep); the document MUST contain the anchor literal `11.4.109` once this clause lands. Paired `$` .

mutation: remove the preamble's hook-reference sentence → constitution checker detects absence → FAIL.

(C) Orchestrator Pre-Action Checklist – guard against self-forgetting.

The same `docs/AGENT_GUARDRAILS.md` document MUST contain an **ORCHESTRATOR PRE-ACTION CHECKLIST** covering at minimum:

- Before any subagent dispatch: confirm the preamble (sub-rule B) is pasted verbatim.
- Before any emulator/device action: confirm the run routes through the Containers submodule CLI, not host-direct; confirm no live device is targeted; confirm the gate host is eligible (otherwise honestly BLOCKED).
- Before any distribute action: confirm Challenge Tests EXECUTED (not compiled) against the exact artifact; version code bumped; CHANGELOG entry present; debug-stage evidence present if two-stage distribute applies.
- Before any push / destructive git action: confirm no force flags without per-operation operator approval; remote is the approved set only; for history rewrite, hardlinked `.git` backup made first.
- Before any host-affecting command: confirm the command is not in the host-power blocked class.

Anti-bluff requirement: the checklist MUST be present in the preamble document and verifiable by `check-constitution.sh` . Each checklist item is traceable to a governing clause cited inline.

Mechanical enforcement.

Consuming-project `check-constitution.sh` (or `verify-all-constitution-rules.sh`) MUST verify ALL of:

- . `constitution/scripts/hooks/guard-forbidden-commands.sh` exists at the canonical path in the constitution submodule.
- . The consumer's AI-agent settings file (`.claude/settings.json` or equivalent) contains a `PreToolUse` hook entry referencing the canonical script path.
- . `docs/AGENT_GUARDRAILS.md` (or the equivalent preamble path) exists, contains the `SUBAGENT CONSTITUTIONAL PREAMBLE` heading, and contains the `ORCHESTRATOR PRE-ACTION CHECKLIST` heading.
- . A hermetic test for the hook exists (at minimum, a file under `tests/hooks/` or equivalent) and passes.
- . The anchor literal `11.4.109` is present in every per-scope governance file (`CLAUDE.md` , `AGENTS.md` , `QWEN.md`) of the consuming project and its owned submodules per § . . inheritance.

Pre-build gate `CM-ANTI-FORGETTING-ENFORCEMENT` enforces ()-(). Propagation gate `CM-COVENANT-114-109-PROPAGATION` enforces (). Paired § . meta-test mutations: (a) remove the `PreToolUse` hook entry from the agent settings file → gate () FAILs; (b) delete `docs/AGENT_GUARDRAILS.md` → gate () FAILs; (c) remove the canonical hook script from the constitution submodule → gate () FAILs; (d) strip the `11.4.109` literal from a consumer `CLAUDE.md` → propagation gate FAILs.

Inheritance.

Applies recursively to every submodule, every feature, every new artifact, every project consuming this constitution. Submodule constitutions MAY add stricter rules (additional blocked command classes, additional checklist items, stricter test counts) but MUST NOT relax any clause. Per § . . : this universal clause lives

in the constitution submodule; consuming projects'

CLAUDE.md / AGENTS.md / QWEN.md carry the inheritance pointer-block per § 11.4.109 + the 11.4.109 literal per the propagation gate.

The principle this clause enshrines (in the fewest possible words): A constraint that depends on an agent remembering it is not a constraint – it is a hope. The PreToolUse hook converts the most dangerous command classes into hard-blocked failures regardless of agent memory. The preamble document converts the full ruleset into something the orchestrator can inject rather than recite from training. Together they are the minimum viable "anti-forgetting" enforcement posture for any AI-agent-assisted project under this constitution.

Composes with § 11.4.109 + § 11.4.110 ..§ 11.4.111 (anti-bluff covenant – the hook and preamble are the floor and ceiling of the anti-bluff enforcement layer), § 11.4.112 (no-guessing – the escape hatch requires a <reason> exactly because undocumented exceptions are a guessing posture), § 11.4.113 (credentials – the host-power class is categorically non-overridable, analogous to credential leaks), § 11.4.114 (mechanical enforcement – § 11.4.115 is the AI-agent-side specialisation of § 11.4.116's git-hook mechanical enforcement; together they cover the full automated-and-agentic surface), § 11.4.117 (Containers mandate – the emulator gate in the hook enforces § 11.4.118 at the tool-call boundary), § 11.4.119 / § 11.4.120 / § 11.4.121 (CodeGraph – the hook passes codegraph MCP calls through untouched; the preamble includes the anti-forgetting constraint about CodeGraph use), § 11.4.122 (cross-platform parity – the hook script is portable bash/awk with no OS-specific primitives; per-OS blocked-class equivalents are addable via consuming-project extension), § 11.4.123 (working-tree quiescence – the pre-action

checklist instructs agents to grep for mutation markers before staging), § 11.4.17 (full-automation mandate – the hook's test suite MUST itself be fully self-driving end-to-end), § 11.4.18 (systematic-debugging – any hook violation that surfaces should trigger the systematic-debugging arc, not a quick workaround), § 11.4.19 (host-session-safety – the host-power class is non-overridable; § 11.4.17 is the constitutional root of that prohibition).

Classification: universal (§ 11.4.17) – the PreToolUse guard hook + subagent preamble + orchestrator checklist pattern is vendor-neutral (defined against the JSON stdin contract that every Claude Code hook receives; the pattern ports to any agent runtime that exposes a pre-tool-call interception point) and reusable across ANY project that uses AI coding agents. The specific blocked command-class list at the UNIVERSAL layer covers the four classes that are unconditionally forbidden by this constitution (host-power, privilege escalation, force-push/bypass, raw emulator host-direct); consuming projects extend the list with project-specific additions per § 11.4.18.

Non-compliance is a release blocker. No escape hatch – no `--skip-pretooluse-hook`, `--no-guardrails-doc`, `--anti-forgetting-optional`, `--single-layer-sufficient` flag exists. A project whose agent can forget a constitutional constraint and execute the forbidden command is constitutionally undefended, irrespective of how well the docs are written.

Canonical authority: this Constitution.md § 11.4.17 in the HelixConstitution submodule; reference implementation `constitution/scripts/hooks/guard-forbidden-commands.sh` + reference preamble document `constitution/docs/AGENT_GUARDRAILS.md` (path within the submodule). Consuming projects inherit by reference per § 11.4.18.

§ . . – Pre-build build-readiness
 verdict + change-impact clash detection
 mandate (operator mandate, 11 4 109
 – –)

Short tag: pre-build-readiness-verdict .

§-slot history note. Drafted as § . . in docs/
 research/prebuild_rigor_20260603/PROPOSED_CONSTITUTION_ANCHOR.md ,
 renumbered to § . . per § . . fetch-
 before-push after the concurrent landing of
 § . . (Anti-Forgetting Enforcement, commit
 1d9e5d6). Same collision-resolution pattern as
 § . . / § . . .

Forensic anchor (genericised, – –). A fix
 shipped a new system-property *read* in a code/init change
 but introduced no matching security-policy grant and no
 property-context type entry for it – so the read was
silently denied at runtime and the feature was dead, while
 every pre-build gate stayed green because the gate only
 grepped the source file for the change. The defect class
 generalises: a change introduces a new dependency on a
second artifact (security policy, context file, service
 registry, interface freeze-snapshot, symbol table, build-
 graph node) that the pre-build never cross-checks, so the
 change is "ready" by every existing gate yet broken the
 instant it runs. This is the SOURCE→ARTIFACT half of
 § . . 's four-layer gap, shifted *left* to pre-
 build time: most of these clashes are **statically catchable**
from the change diff itself, before any build.

The architectural flaw (FACT, traced). Pre-build gates
 that validate each artifact *in isolation* (the policy file
 is internally consistent; the freeze-snapshot is well-
 formed) do NOT cross-check that a change introducing a new
 dependency *also* introduced the matching second artifact.
 The cross-check is **diff-driven** (its input is the change

set, not any single file) and is therefore a capability no per-file gate provides. Without it, the only protection is author-vigilance ("remember to also add the policy line"), which § . . . already established is the failure mode that mechanical enforcement exists to replace.

The mandate (ALL must hold).

- . **A single deterministic READY-FOR-BUILD verdict is mandatory.** Before any artifact rebuild, the project MUST emit ONE machine-checkable verdict – *the codebase IS / IS NOT ready for build* – backed by a captured-evidence bundle enumerating every gate family's status. The rebuild orchestrator MUST refuse to start the build unless the verdict is READY (non-zero exit on NOT-READY). A green individual gate is necessary but not the verdict; the verdict is the aggregate function and cannot report READY while any family reports FAIL.
- . **A diff-driven change-impact + conflict/clash detector is mandatory.** For every change in the batch, the pre-build MUST cross-check that any newly-introduced dependency on a *second* artifact is satisfied – at minimum: a new system-property read $\Rightarrow$ the property is typed in the platform's property-context registry AND the reading domain has the matching read-grant; a new registered service $\Rightarrow$ its service-context entry exists; a new declared startup/init service $\Rightarrow$ its security label is present; a new/changed stable interface $\Rightarrow$ its freeze-snapshot/version was updated in the same change; a new native-library dependency $\Rightarrow$ the named library resolves to a known module or prebuilt; a new security-policy rule $\Rightarrow$ every type/attribute it uses is defined; and two changes in the batch touching the SAME function / resource / configuration seam $\Rightarrow$ the collision is explicitly acknowledged (the cross-fix-interaction class). Each sub-detector is independently anti-bluff (clause).

- **Coverage-completeness is a gate, not a convention.** Every changed source file in the batch MUST map to $\geq$ pre-build gate + $\geq$ deployed-target test + $\geq$ paired § . mutation (the four-layer § . . (b) floor). A baseline file records the legacy tail of known holes; the gate FAILs on any NEW hole and the covered-ratio ratchets upward over phases per § . . . A change with no owning gate/test/mutation is a coverage hole and blocks the verdict.
- **Two-speed honesty – grep-speed vs REQUIRES_BUILD – is mandatory.** The regime MUST separate always-on grep-speed gates (run every commit, seconds) from REQUIRES_BUILD heavy gates (the build-graph parse-only dry-run; full security-policy neverallow compilation; built-symbol ABI diff). Heavy gates run as diff-gated opt-in stages (triggered only when the change touches build-graph / policy / interface files) and MUST run bounded per the host-session-safety budget (§ . /§ .) or inside the project's containerized build. Conflating the two – making the grep path slow, or claiming grep proves graph/neverallow correctness – is forbidden per § . . .
- **Every gate and analyzer is anti-bluff by paired mutation.** Each gate family AND each wired third-party analyzer (its exit-code-on-failure semantics included) MUST ship a paired § . mutation that plants the exact defect class and asserts the gate FAILs (e.g. a property read with no policy → the clash detector FAILs; a syntactically-broken build-graph fragment → the dry-run reports FAIL; an analyzer that exits on a planted violation is itself the bluff the mutation catches). An analyzer wired without its mutation pair is a § . bluff surface.
- **Honest boundary – pre-build catches statically-catchable defects only.** A READY verdict proves the change is

internally consistent and ready to build; it does NOT prove the feature works on the deployed target. The correct claim is: *after the pre-build passes, the residual defects deployed-target / manual testing can find are runtime/hardware defects, NOT statically-catchable (preventable) ones* – the regime makes the **preventable** class empty, not the **all-defects** class. Overclaiming that manual testing will find nothing is itself a § 11.4.17 violation. Build-graph correctness, full neverallow safety, and ABI correctness are proven ONLY by the REQUIRES_BUILD families (clause 11.4.17); the grep tier proves harness-presence + symbol-definedness, never full correctness. Runtime/USER-VISIBLE correctness remains entirely § 11.4.17's four-layer job.

Classification: universal (§ 11.4.17) – the single ready-for-build verdict, the diff-driven change-impact/clash detector, the coverage-completeness gate, the grep-speed-vs-REQUIRES_BUILD two-speed split, the per-gate paired mutation, and the statically-catchable-only honest boundary are platform-neutral pre-build-rigor disciplines reusable by ANY project that builds an artifact from source; the consuming project supplies its concrete property/service/policy registries, its build-graph parse-only command, its interface-freeze mechanism, and its changed-file-to-gate mapping per § 11.4.17.

Composes with § 11.4.17 (FAIL-bluffs forbidden – a clash-detector that crashes on a malformed diff and FAILs is as misleading as one that passes silently), § 11.4.17 (test-interrupt-on-discovery – a NOT-READY verdict triggers STOP-and-fix before the build), § 11.4.17 (no-guessing – the two-speed honest boundary; the grep tier MUST NOT claim correctness it cannot prove), § 11.4.17 (batch-source-fixes-before-rebuild – the verdict gates exactly that rebuild), § 11.4.17 (no-fakes-beyond-

unit – wired analyzers exercise the real change),
 § 11.4.110 (deterministic consistency – the coverage-completeness ratchet shares § 11.4.110's threshold mechanism and N-iteration determinism), § 11.4.111 (target-shell-parseability – the clash-detector helper itself parses under its target shell), § 11.4.112 (mechanical enforcement – § 11.4.112 is the mechanical replacement for author-vigilance on the clash class), § 11.4.113 (multi-pass change-evaluation – the verdict is the mechanical companion to Pass / /),
 § 11.4.114 (four-layer fix-verification – § 11.4.114 is the SOURCE→ARTIFACT half shifted left to pre-build time; § 11.4.115 owns the RUNTIME-ON-CLEAN-TARGET→USER-VISIBLE half; together they span all four layers).

Propagation. Propagation gate

CM-COVENANT-114-110-PROPAGATION enforces the literal anchor 11.4.110 across the consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md); paired § 11.4.111 meta-test mutation strips the 11.4.110 literal from a consumer file → the gate FAILs. Recommended per-family gates CM-READY-FOR-BUILD-VERDICT (single deterministic verdict gates the rebuild), CM-CHANGE-IMPACT-CLASH-DETECTOR (diff-driven second-artifact cross-check), CM-COVERAGE-COMPLETENESS-GATE (changed-file ⇔ gate+test+mutation with baseline ratchet), CM-BUILDGRAPH-DRYRUN-WIRED + CM-SEPOLICY-NEVERALLOW-WIRED (REQUIRES_BUILD opt-in stages exist + bounded-safe); each with a paired § 11.4.112 mutation per clause 11.4.112. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § 11.4.113 in the HelixConstitution submodule. All consuming projects restate + cite via § 11.4.113 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --source-green-is-ready , --skip-clash-

detector , --skip-coverage-gate , --no-ready-verdict , --grep-proves-neverallow , --skip-buildgraph-dryrun , --build-without-ready-verdict flag exists.

\$. . – Resolve-by-stable-name-not-by-enumeration-index mandate (research-derived, – –)

Short tag: resolve-by-stable-name .

Forensic anchor (genericised, – –). A platform bound an audio output to a kernel-enumerated device index (card=0). A second device of a different class enumerated FIRST at boot and took slot , shifting the intended device to slot . The static index binding now pointed at the wrong device: the policy layer failed to attach the intended sink, mis-assigned its TYPE, and the user-facing output switcher labelled the AV receiver "Wired Headphone" while routing collapsed to stereo. The lower layer of the SAME stack already resolved the device correctly **by name** (scanning /proc/asound/card*/id for the controller name) – proving the brittleness was the *index* binding, not the resolution capability. The defect class generalises to every enumerated resource whose ordinal is assigned at discovery/boot/hotplug time and is therefore non-deterministic across reboots, device additions, and topology changes.

The mandate. Any binding between a configuration / policy / code site and a hardware device, resource handle, or enumerated entity – audio cards, display connectors, network interfaces, storage devices, GPU render nodes, input devices, camera indices, container/process slots, any registry whose members are assigned an ordinal at discovery time – **MUST** resolve the target by a **stable identifier** (name, UUID, serial, label, controller-name, content-derived hash, sink-reported identity) and **MUST NOT**

resolve it by an enumeration index / ordinal / slot number, UNLESS the platform documents that ordinal as deterministically pinned (and the pinning mechanism is itself part of the binding, captured and asserted). Where a stable identifier exists at one layer of a stack, every other layer binding the same resource MUST use the same identifier – mixed by-name-here / by-index-there is the structural weak link this anchor forbids. Where the platform offers a sink-reported / capability-derived identity (e.g. an EDID/ELD-style capability descriptor), preferring it makes the binding robust to the broadest class of topology change.

Honest boundary. When the platform genuinely provides no stable identifier and only an ordinal exists, the binding MUST (a) pin the ordinal deterministically via the platform's own pinning mechanism, (b) capture that pin as part of the binding's runtime signature per § 4.4.1, and (c) document the residual fragility as **UNCONFIRMED**: -class risk per § 4.4.2 – never silently trust an unpinned ordinal.

Classification: universal (§ 4.4.1) – index-vs-name resource resolution is a platform-neutral binding-robustness discipline reusable by ANY project that binds to enumerated hardware / resources / handles; the consuming project supplies its concrete stable-identifier mechanism (ALSA card name, DRM connector name, NIC predictable name, block-device UUID, etc.) and the specific layers that must agree per § 4.4.1.

Composes with § 4.4.1 (no-guessing – "the index is *usually* stable" is the exact guess this anchor forbids; prove the pin or bind by name), § 4.4.2 (deep-web-research – mature platform stacks resolve by name/UUID; reproducing a known-brittle index binding when the by-name path is documented is a § 4.4.1 violation by omission), § 4.4.3 (sink-side positive-evidence –

the by-name binding's correctness is verified by the downstream/sink-reported identity, not by the local config), § 11.4.111 (four-layer fix-verification – the resolved-by-name binding's runtime signature is asserted on a clean target across the topology that originally broke the index binding), § 11.4.112 (pre-build clash detection – a new ordinal-based binding in a diff is a statically-catchable clash class the change-impact detector flags).

Propagation. Propagation gate

CM-COVENANT-114-111-PROPAGATION enforces the literal anchor 11.4.111 across the consumer fleet; paired § 11.4.112 meta-test mutation strips the literal → the gate FAILs.

Recommended per-family gate CM-RESOLVE-BY-NAME-NOT-INDEX (diff-driven scan flagging a new enumeration-index binding to a resource that exposes a stable identifier, unless an accompanying deterministic-pin + runtime-signature is present); paired § 11.4.113 mutation plants an unpinned index binding → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § 11.4.114 in the HelixConstitution submodule. All consuming projects restate + cite via § 11.4.114 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --allow-index-binding, --ordinal-is-stable-enough, --skip-name-resolution, --trust-unpinned-index flag exists.

2026 06 03

§ 11.4.115 – Structural-impossibility won't-fix classification mandate (research-derived, – –)

Short tag: structural-impossibility-wont-fix .

Forensic anchor (genericised, - -). Deep research (§ . .) into relocating protected (content-protection / secure-surface) video to a secondary display proved – from authoritative platform/HDCP documentation plus reproducible captured behaviour (a secure surface is blanked on any output lacking the secure flag; mirror/screenshot of a secure layer returns black) – that the goal is **structurally impossible by platform design**, not a missing feature or an unsolved engineering problem. Without a durable classification, such a goal is re-investigated every cycle: each new agent / operator re-reads the same sources, re-runs the same probes, re-derives the same impossibility, and burns the same effort – an anti-pattern that compounds across cycles.

The mandate. When deep research per § . . PROVES (with cited authoritative sources AND, where applicable, reproducible captured evidence) that a goal is structurally impossible on the target platform – i.e. forbidden by the platform's design / a hardware/protocol constraint / a documented kernel-or-API limitation, not merely unimplemented or hard – the goal MUST be: () **classified won't-fix and closed** per the § . . terminal-status discipline, with obsolescence/closure reason **structurally-impossible** (a value in the § . . closed reason vocabulary); () **documented with the impossibility evidence** – the cited authoritative source URLs (per § . . latest-source verification where the platform docs may evolve) PLUS the reproducible probe/captured evidence that demonstrates the constraint, recorded in the tracker entry and the relevant **docs/** guide so the verdict is auditable; (.) **NOT re-attempted in future cycles** – the closed entry is the canonical answer; a future agent that re-opens it MUST cite NEW evidence that the platform constraint changed (per § . . reopened-source attribution + § . . demotion-evidence), never merely re-derive the same

impossibility; () paired with the correct posture – the entry MUST state what the project DOES instead (the supported-within-the-constraint behaviour), so "impossible" is never confused with "broken / unhandled".

Honest boundary (§ 11.4.21) Won't-fix: structurally-impossible is reserved for proven platform/hardware/protocol impossibility. "We could not find a way" / "it is very hard" / "no time" are NOT structural impossibility – they are Operator-blocked (§ 11.4.22) or open work, and mislabelling them won't-fix to avoid the work is a § 11.4.23 bluff at the planning layer. A future platform change CAN make the impossible possible; the classification is durable but not eternal, and reopening on NEW cited evidence is correct.

Classification: universal (§ 11.4.24) – proving-then-durably-classifying structural impossibility is a platform-neutral effort-conservation + honesty discipline reusable by ANY project; the consuming project supplies the specific impossible goal, its platform constraint, and the cited evidence per § 11.4.25.

Composes with § 11.4.26 (no-guessing – the impossibility is stated as FACT with cited evidence, never as "probably can't"), § 11.4.27 (demotion-evidence – reopening a won't-fix requires NEW positive evidence the constraint changed), § 11.4.28 (deep-web-research – the proof is the research output; "NO external solution found – structurally impossible" is the § 11.4.29 citation), § 11.4.30 (reopened-source attribution – a reopen of a structurally-impossible item attributes who reopened + the new evidence), § 11.4.31 (Obsolete/terminal-status – structurally-impossible is a closure reason in the § 11.4.32 closed vocabulary; the colorizer marks it like any closed item), § 11.4.33 (latest-source verification – the cited platform docs MUST be the latest

where they can evolve, so the impossibility verdict is not stale).

Propagation. Propagation gate

CM-COVENANT-114-112-PROPAGATION enforces the literal anchor 11.4.112 across the consumer fleet; paired § . meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-WONT-FIX-STRUCTURAL-IMPOSSIBILITY (every tracker entry with closure reason structurally-impossible carries cited authoritative-source evidence + a stated correct-posture line); paired § . mutation strips the evidence citation from such an entry → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

() SCOPE-STATEMENT + ADJACENT-GOAL ENUMERATION – an impossibility verdict is bounded, and MUST say where its edge is (extension, research-derived, – –). Forensic FACT (genericised, from a consuming project): a verdict that was TRUE on its own terms – "relocating another application's protected video output to a secondary display while that application's UI stays in place, without the application's cooperation is structurally impossible" – LEAKED to an ADJACENT, VIABLE goal it never covered: moving the whole task to the other display. For six weeks the project believed "no platform mechanism exists", while the platform had shipped a public, permission-free SDK call for the adjacent goal several major versions earlier, and the platform vendor's own connected-display documentation PRESCRIBED exactly that call. The impossibility was never re-litigated (clause () worked); it was CITED, unexamined, against a neighbour – and clauses ()-() contain nothing that bounds a verdict's reach. This anchor is ASYMMETRIC without this clause: clause () hardens a verdict against reopening ("the closed entry is the canonical answer") while nothing constrains how far it may be applied, so the

only mechanical pressure runs in the direction that AMPLIFIES a scope leak. Therefore every structurally-impossible verdict MUST additionally: (a) **STATE ITS EXACT SCOPE** – the precise goal/variant the cited evidence actually proves impossible, written as the narrowest claim the evidence supports (every qualifier that is load-bearing – "without the application's cooperation", "while the UI stays put", "on this platform version" – is part of the verdict, not commentary; a qualifier dropped is a verdict widened); (b) **ENUMERATE THE ADJACENT GOALS IT DOES NOT COVER** – the neighbouring variants a reader could mistake for the same goal, each marked NOT-COVERED-BY-THIS-VERDICT (a neighbour whose viability is unknown is recorded as UNKNOWN with a tracked follow-up per § 1.1.1 – an unexamined neighbour is NEVER silently absorbed into the impossibility); (c) **BE CITATION-FENCED** – the verdict MUST NOT be cited as the answer for ANY goal outside its stated scope, in planning, investigation, design, or status, without its OWN § 1.1.1 deep-research pass and its OWN proof (citing it beyond its fence is asserting an unproven claim as fact – a § 1.1.1 violation and a § 1.1.1 bluff at the verdict layer, of the same class § 1.1.1 names when a reproduction that never reaches the precondition is used to conclude "the mechanism never engages"). **Honest boundary (§ 1.1.1)**: an adjacent goal being outside the fence does NOT make it viable – it makes it UNDECIDED, and it earns its own verdict by its own evidence, never by inheritance in either direction. **Why this is a clause and not another narrative:**

§ 1.1.1 (A)/(E) already documented this exact failure class in prose (its case study is an impossibility verdict OVERTURNED by the research pass it never had) and REQUIRES the deep-research pass before a structurally-impossible verdict is EARNED – yet the leak still happened, because § 1.1.1 fires at the seam where a verdict

is MINTED for a tracked item, and this leak happened where a verdict was CITED as a premise, which crosses no closure seam at all. Clauses () and () are already content requirements on the verdict artifact; () is the third, and it is the one that bounds blast radius.

Gate contract for clause () (extension,

- -). **CM-WONT-FIX-STRUCTURAL-IMPOSSIBILITY** is EXTENDED: every **structurally-impossible** entry MUST additionally carry (i) an explicit SCOPE statement and (ii) an adjacent-goals-NOT-covered enumeration (each neighbour marked not-covered or UNKNOWN + tracked); and (iii) a citation of the verdict against a goal outside its stated scope, without that goal's own § . . research + proof, → FAIL. Paired § . mutations: strip the scope statement, or the adjacent-goal enumeration, from an impossibility entry → the gate FAILs; cite an in-scope verdict against an enumerated not-covered neighbour → the gate FAILs. Gate-code = a separate work item; this is the contract, not a claim the code has shipped (§ . .).

Canonical authority: this Constitution.md § . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no **--wont-fix-without-proof**, **--reattempt-closed-impossible**, **--skip-impossibility-evidence**, **--impossible-equals-broken**, **--unscoped-impossibility-verdict**, **--skip-adjacent-goal-enumeration**, **--cite-verdict-beyond-scope**, **--neighbour-inherits-impossibility** flag exists.

§ . . – Absolute no-force-push +
merge-onto-latest-main mandate (User mandate,
2126 06 03 – –)

Short tag: absolute-no-force-push .

Forensic anchor – verbatim user mandate
(– –):

"Any force-push is strictly forbidden! We must for every Submodule take as a base latest commit on Submodule's main (or master) branch, then on top of it carefully to merge all changes that have to be pushed! Once all merging is carefully done we perform commit and push to all Submodule's upstreams!"

The mandate – force-push is **STRICTLY FORBIDDEN**, with **NO exception**. Any force-push – `git push --force` , `git push --force-with-lease` , `git push +<ref>` , or any history-rewriting operation that overwrites a remote ref non-fast-forward – is forbidden against every repository this Constitution governs (every consuming project's main repo, this constitution submodule, every owned submodule, every nested submodule, every upstream of each). There is no escape hatch, no operator-approval path, no "after a merge-first audit" path. The merge-onto-latest-main procedure below is ALWAYS available, so a force-push is NEVER necessary; eliminating the need eliminates the operation.

The mandated –step integration procedure (for every repo/submodule whose local has commits to publish OR whose mirrors have diverged):

- **Fetch all remotes** – `git fetch --all --prune --tags` against `origin` + every configured upstream; capture the output as evidence.

- **Set the base to the LATEST commit on the canonical main / master branch** – the most-advanced mirror tip across all upstreams. The integration target is that tip, never a stale local branch point.
- **Carefully MERGE every change that must be published on top of that base** – a union merge that preserves BOTH sides (local commits AND every remote-side commit). NEVER `-s ours` (it silently discards the remote side), NEVER a rebase/reset/ `--allow-unrelated-histories` that could drop commits (per `$ no-commit-loss`). The merge commit MUST descend from every mirror tip so that every subsequent push is a clean fast-forward.
- **Resolve every conflict carefully** so nothing is broken, lost, or corrupted on either side – no conflict markers anywhere (`grep -rn '^<<<<<< \|^===== $\|^>>>>>>'` returns empty across governance + source + test files), no file silently dropped, every previously-passing gate/test still passes on the integrated tree (per `$ . . + $ . .`), every captured-evidence artefact still validates.
- **Commit the merge** – stage only the intended files (NEVER `git add -A` inside a submodule with untracked build content per `$ . .`), commit the carefully-resolved merge.
- **Push the result to ALL upstreams** – each push is a fast-forward because the merge commit descends from every mirror tip, so NO force is ever needed. If an upstream still reports non-fast-forward (a sibling landed work between step `11 4 41` and the push), do NOT force – return to step `11 4 41` for that upstream, fetch its new tip, MERGE it (union, preserve both), re-validate step `0+3`, and re-push.

Tightens `$ . . / $ . . / $ . /`
 CONST- – force-push escape hatch REMOVED.
`$ . .` (Pre-Force-Push Merge-First Mandate),

§ 11.4.119. (Pre-Push Fetch + Investigate + Integrate),
 § (Force-push requires explicit user authorization),
 and CONST- previously PERMITTED a force-push after
 a merge-first pipeline + per-operation operator approval.
 § . supersedes that permitting stance: even
 WITH operator approval, even AFTER a clean merge-first
 audit, a force-push is forbidden – because the merge-onto-
 latest-main path in this clause is always available and
 always sufficient, force is never necessary and therefore
 never authorised. Those clauses' merge-first/fetch-first
 machinery REMAINS in force as the integration discipline;
 only their terminal "...then force-push" step is struck.
 § . 's destructive-operation § backup discipline
 still applies to any non-push history operation, but a
 history-rewriting *push* is simply not performed.

No escape hatch. No `--force`, no `--force-with-lease`, no
`+<ref>` push, no `--no-verify` to slip a forced ref past a
 hook, no "operator authorised the force-push" path, no
 "lease semantics unavailable so plain `--force` is OK" path.
 A force-push that lands on any remote is a critical defect
 regardless of how briefly it persisted, severity-
 equivalent to a § . PASS-bluff at the data-safety
 layer (per §).

Classification: universal (§ .) – an absolute-
 no-force-push + merge-onto-latest-main integration
 discipline is platform-neutral and reusable across every
 repository in every consuming project; nothing here
 references a particular project, vendor, hardware, or
 region.

Composes with § . (multi-upstream push is the norm –
 step fans out to all), § / § . (absolute data
 safety – § . is the no-loss push discipline
 that makes force-push unnecessary), § . (test-
 interrupt + clean-baseline retest – step audit),
 § . (no-guessing – remote state is not knowable

without the step- fetch), § . . (constitution-submodule update workflow – its conflict-resolution step is this procedure), § . . (fetch-before-edit – step, generalised to fetch-before-push), § . . (full-suite retest authority for the integrated tree), § . . (force-push merge-first – TIGHTENED: merge-first stays, the force-push step is removed), § . . (pre-push fetch + integrate – TIGHTENED the same way), § . . (background-push – the detached push still follows this fast-forward-only discipline), CONST- (per-operation authorization – TIGHTENED: no force-push is authorisable).

Propagation. Propagation gate

CM-COVENANT-114-113-PROPAGATION enforces the literal anchor 11.4.113 across the consumer fleet; paired § . meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-NO-FORCE-PUSH-ABSOLUTE (scans tracked scripts + hooks for any push --force / push --force-with-lease / push +<ref> invocation and rejects it; a § . . -class PreToolUse guard blocks the force-push command class at the tool-call boundary); paired § . mutation injects a git push --force into a tracked script → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --force , --force-with-lease , --allow-force-push , --force-push-authorised , --skip-merge-onto-main flag exists.

§ . . – Last-known-good-tag
regression isolation mandate (. . -dev
remediation, - -)

Short tag: `last-known-good-tag-isolation` .

Forensic anchor (genericised, - -). A user reported that a previously-working feature (multichannel HDMI audio) was broken on the current build. The operator's lead – "the feature WORKS and passes all tests at the last release tag" – converted an open-ended root-cause hunt into a targeted regression repair: diffing the broken working tree against the last-known-good tag's commit isolated EVERY reported audio defect (multichannel downmix + output-device label mismapping) to ONE post-tag burndown batch (`AudioDeviceInventory.connectDualHdmiPorts_l` dual-HDMI-connect) in minutes, not hours. The tagged-good version is the oracle: behaviour that works there and is broken now is, by definition, a regression introduced in the diff between the two.

The mandate. When a previously-working feature, behaviour, or capability is observed broken, the FIRST diagnostic action MUST be to identify the last release tag (or last commit / build / deployment) at which that feature was KNOWN-GOOD, and diff/bisect the broken state against it – BEFORE any open-ended root-cause investigation, BEFORE any speculative fix. The known-good revision is the regression oracle: () it bounds the search space to the commits between known-good and now (a `git diff <good-tag>..HEAD --stat` of the feature's files is the captured evidence); () it tells you the fix is a regression REPAIR (restore/forward-fix the broken sub-part), not a from-scratch design problem; () each suspect file is compared against its known-good version as the behavioural oracle. When the operator volunteers a known-good tag, that lead is load-bearing and MUST be acted on first. Where the bisect surface is large, `git bisect` against the feature's

autonomous RED test (per § 4.1.1) mechanises the search.

Forward-fix vs wholesale-revert. Once the regressing change is isolated, the default is a SURGICAL forward-fix – keep the new features the post-good-tag work introduced, revert ONLY the broken sub-part where the known-good version was correct – UNLESS the operator prefers wholesale revert. A wholesale revert of a multi-feature batch to reach the known-good state is a last resort that loses the batch's other (working) features; prefer isolating and repairing the single regressing seam. Every forward-fix MUST prove (via the § 4.1.1 RED→GREEN test) that the feature now matches the known-good behaviour.

Honest boundary (§ 4.1.2). "It worked before" is a HYPOTHESIS until the known-good tag is identified AND the feature is confirmed working there (re-run the feature's test against the known-good build, or cite the captured-evidence from when that tag shipped). A feature that NEVER worked is not a regression and this anchor's oracle does not apply – investigate as new work. "Probably regressed in the last batch" without the diff is a § 4.1.1 guess.

Classification: universal (§ 4.1.1) – diff-against-last-known-good is a platform-neutral regression-isolation discipline reusable by ANY versioned project with release tags / commit history; the consuming project supplies its tag-naming convention and per-feature known-good oracle per § 4.1.1.

Composes with § 4.1.1 (test-interrupt-on-discovery – the isolation runs as the first step of the systematic-debugging Phase the STOP triggers), § 4.1.1 (no-guessing – the diff is the FACT-grade evidence that replaces "probably the last batch"), § 4.1.1.

(demotion-evidence – confirming known-good requires positive evidence at that tag), § 11.4.102 (full-suite retest authority – the known-good tag is the baseline a release-gate retest restores toward), § 11.4.103 (TDD-fix – the § 11.4.101 RED test reproduces the regression the diff isolated), § 11.4.104 (systematic-debugging – last-known-good isolation IS the Phase 1 root-cause technique for regressions specifically), § 11.4.105 (four-layer fix-verification – the forward-fix's runtime signature is verified against the known-good behaviour on a clean target).

Propagation. Propagation gate

`CM-COVENANT-114-114-PROPAGATION` enforces the literal anchor `11.4.114` across the consumer fleet; paired § 11.4.115 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate `CM-REGRESSION-ISOLATED-AGAINST-KNOWN-GOOD` (every tracker entry classified as a regression cites the last-known-good tag + the diff range it was isolated to); paired § 11.4.116 mutation strips the known-good citation from a regression entry → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § 11.4.117 in the HelixConstitution submodule. All consuming projects restate + cite via § 11.4.118 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no `--skip-known-good-diff`, `--root-cause-from-scratch`, `--assume-regression-source`, `--wholesale-revert-without-isolation` flag exists.

```
$ . . - RED-baseline-on-the-  
broken-artifact + polarity-switch mandate  
( . . -dev remediation,  
    - - )
```

Short tag: red-baseline-polarity-switch .

Forensic anchor (genericised, - -). Per
\$. . a fix opens with a RED test. The
. . . -dev remediation sharpened HOW: the autonomous RED
test was authored to REPRODUCE each live defect on the
CURRENT pre-fix build – proving the defect is real, live,
and captured (video flips to primary on reconfigure
/ ; audio . → Arvus PCM / × ; subtitle
absent on secondary) – with a single environment flag
(RED_MODE=1 default) that, set to 0 post-fix, flips the
SAME test to GREEN-guard polarity. One test source, two
roles: the reproduction-of-the-defect (RED) and the
permanent regression-guard (GREEN). No separate happy-path
test is authored; the test that catches the bug IS the
test that guards against its return.

The mandate. Every \$. . RED test MUST be
authored to reproduce the defect on the CURRENT, pre-fix
artifact (the actual broken build / deployment), capturing
positive evidence per \$. . / \$. . /
\$. . that the defect is genuinely present –
never a synthetic / hypothetical failure the fix is then
written to agree with. The SAME test source MUST carry a
single polarity switch (an env flag / parameter, canonical
RED_MODE , default 1 = reproduce-and-assert-defect-
present) that, flipped to 0 post-fix, converts the test
into the GREEN regression-guard asserting the defect is
ABSENT. The RED run (against the broken artifact) and the
GREEN run (against the fixed artifact, per \$. .
on a clean target) MUST both be captured as evidence: RED-
then-GREEN is the proof the test catches the bug AND the
fix closes it. A separate happy-path test that never

failed on the broken artifact is forbidden as the primary guard – it demonstrates only that the test agrees with the fix (the § 11.4.6.1 PASS-bluff), not that it catches the defect.

Why one test, two roles. A test authored only as a post-fix happy-path can pass on a build where the feature is silently broken (it asserts the fix's intended state, not the defect's absence). A RED-first test that was OBSERVED failing on the broken artifact, then flipped to GREEN-guard, is provably sensitive to the exact defect – its RED capture is the falsification record. The polarity switch keeps the two roles in ONE source so they cannot drift apart (a separate happy-path test can be edited to stop testing the real condition; the RED-evidence binds the GREEN guard to the actual defect).

Honest boundary (§ 11.4.6.2). If the defect cannot be reproduced on the current artifact (the RED run does not fail), that is a finding, not a license to write a GREEN-only test: either the defect is not present on this build (close per § 11.4.6.3 with the negative evidence) or the test does not exercise the real condition (fix the test). A RED test that "passes" on the known-broken artifact is itself a § 11.4.6.4 FAIL-bluff inverse – it proves the test is blind, not that the feature works.

Classification: universal (§ 11.4.6.5) – RED-baseline-on-the-broken-artifact + a single-source polarity switch is a platform-neutral test-construction discipline reusable by ANY project doing TDD-fix; the consuming project supplies its polarity-flag mechanism (env var, build flag, parameter) and per-defect reproduction driver per § 11.4.6.6. Strict refinement of § 11.4.6.7 step RED.

Composes with § 11.4.6.8 (FAIL-bluffs – a RED test that passes on the broken artifact is blind), § 11.4.6.9 /

§ . . / § . . / § (the RED reproduction captures positive defect-present evidence), § . . (the RED test is the test-interrupt reproduction), § . . (RED-not-reproducible is a demotion that needs negative evidence), § . . (refines its RED step – RED is on the real broken artifact, GREEN is the same test flipped), § . . (deterministic consistency – both RED and GREEN runs are N-iteration stable), § . . (the GREEN run is the runtime-signature verification on a clean target), § . . (the RED test mechanises the § . . bisect against the known-good oracle).

Propagation. Propagation gate

CM-COVENANT-114-115-PROPAGATION enforces the literal anchor 11.4.115 across the consumer fleet; paired § . meta-test mutation strips the literal → the gate FAILs.

Recommended per-family gate CM-RED-POLARITY-SWITCH-PRESENT (every regression-fix RED test carries a single polarity switch + a captured RED-on-broken-artifact run before the fix); paired § . mutation removes the polarity switch (splitting RED and GREEN into a happy-path-only guard) → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

(F) Machine-written verdict pairs + guard-viability (extension, research-derived, - -).

Forensic FACT (genericised, from a consuming project's fix-lifecycle forensics, - -): commits in one two-month window moved fix claims/statuses at "GREEN-PENDING-BUILD" – the RED→GREEN discipline existed as prose, but the verdicts were hand-declared, so nothing distinguished a polarity flip that HAPPENED on the target from one that was ASSERTED; the project's own history then selected the discriminator – every fix whose done-claim was preceded by an on-device polarity flip that held recorded ZERO reopens, and every fix confirmed at

source-green recurred. Therefore (ALL hold): () **RED and GREEN are machine-written verdict artifacts, never prose** – the guard HARNESS (not a human, § . . . ()) writes a per-run verdict file carrying at minimum the item id, guard identity, polarity, exit code, the target's artifact/build fingerprint READ FROM THE TARGET AT RUN TIME, iteration count, and the evidence-file list (path + hash + feature-class); a hand-authored verdict is detectable because the fingerprint/target identity must re-verify against the artifacts it cites. () **The pair must cross-validate** – GREEN is acceptable only with a prior RED whose exit $\neq$. . . on the PRE-fix artifact; GREEN's fingerprint $\neq$ RED's fingerprint (an identical fingerprint proves the fix was never deployed); GREEN runs on a clean target per § . . . /§ . . . with iterations $\geq$ per § . . . () **Guard viability** – a guard that has NEVER been observed FAILing on the genuinely-broken artifact is unvalidated instrumentation (the chaos-engineering validate-the-detector principle): it mints no verdicts; its RED run on the real defect IS its § . . . mutation evidence (the real defect is the mutant, and the guard provably killed it); corollary: a guard structurally unable to FAIL – or unable to PASS (e.g. a pass-floor above the input population's size) – produces no information in either direction and satisfies nothing. () **Mutation = revert + anti-tautology** – the canonical § . . . mutation for a landed fix is the fix-commit's revert (its semantic inverse, restoring ALL enable-paths of the defect per § . . .); a mutation whose diff only deletes/renames the literal strings the gate greps for is a tautology ("grep for X fails after you delete X") and is REFUSED – measured failure: a sed-string mutation "validated" a source-grep gate for a rendered-pixel defect while the defect's second enable-path stayed invisible to both. () **Oracle-class** – the verdict's evidence files MUST match the defect's

feature-class evidence shape (pixel-class defect $\Rightarrow$ pixel/capture artifacts; audio $\Rightarrow$ sink-side/captured-audio; runtime $\Rightarrow$ runtime observables per § 13.1.1.1); a source-grep transcript can never satisfy a user-visible class – the wrong-layer oracle dies by construction, not by review. Recommended gate CM-GUARD-VERDICT-MACHINE-WRITTEN (every done-claiming item's verdict pair is harness-written, fingerprint-cross-validated, class-matched) + paired § 13.1.1.1 mutation (forge a verdict whose fingerprint does not re-verify, or pair a gate with a string-deletion mutation $\Rightarrow$ the gate FAILs; gate-code = separate work item). Composes § 13.1.1.1 / § 13.1.1.1 / § 13.1.1.1 / § 13.1.1.1 () / § 13.1.1.1 / § 13.1.1.1 / § 13.1.1.1 / § 13.1.1.1 .

2026 07 23

(G) Defect-identity binding for CONSTRUCTED reproductions – a fix validated only against an investigator-built mechanism is HARDENING, never the fix for the reported defect (extension, research-derived, – –). Forensic FACT (genericised, – –): an exit-hang fix (a bounded lock wait + a corrected file-descriptor open) was built and nearly shipped with a confident RED $\Rightarrow$ GREEN pair – but the RED reproduced a mechanism the INVESTIGATOR had constructed (a hand-made process holding the lock), not the operator's reported defect; later systematic debugging (§ 13.1.1.1), driving the operator's sequence shape per § 13.1.1.1 with per-step trace timing on a real PTY, REFUTED every mechanism the fix addressed (the lock/unlock pair provably completes on the real path – / clean interactive exits), and the change was honestly RE-LABELLED in-source as "defensive hardening ... NOT a fix for the reported exit-hang". The gap clauses (A)–(F) leave open: the constructed RED genuinely FAILED on the real pre-fix artifact – this anchor's letter was satisfied – while the MECHANISM it reproduced was the investigator's invention,

so the pair proved only that the fix closes the invented mechanism, not the reported defect. Therefore (ALL hold):

- () a RED test's PRECONDITION MUST be TRACEABLE to the reported defect's real evidence – the reporter's reproduction sequence (§ . . .), a captured failure log/trace, or a probe showing the precondition actually occurs in the reported environment – and that traceability is recorded WITH the RED verdict (the clause-(F) fields plus a precondition-provenance field: `observed` vs `constructed`).
- () When NO working reproduction of the reported defect exists and the RED's precondition is CONSTRUCTED, the resulting fix mints AT MOST "defensive hardening" – labelled as such in the commit AND the tracker – and NEVER closes the reported defect's item; the item stays OPEN (§ . . .) until a RED with an `observed` precondition lands, or the item is closed with § . . . -grade non-reproduction evidence or a § . . . /§ . . . evidence-backed reason.
- () When systematic debugging REFUTES every mechanism a pending fix addresses, the fix's done-claim for the reported defect is REVOKED on the spot – the change may still land re-labelled as hardening (independently crossing § . . . /§ . . . review), but claiming it fixes the reported defect after its mechanisms were refuted is a § . . . /§ . . . bluff at the investigation layer even though its RED→GREEN pair is real. STRENGTHENS clause (A)'s "never a synthetic failure" from artifact-synthetic to MECHANISM-synthetic, and § . . . for the no-working-repro case its clause () cannot reach (there is no existing sequence to deviate FROM). Recommended gate `CM-RED-PRECONDITION-DEFECT-TRACEABLE` (every defect-CLOSING claim's RED verdict carries precondition-provenance `observed` + its evidence path; `constructed` provenance on a defect-closing claim → FAIL, while the SAME provenance on an item honestly labelled hardening PASSES – the § . . . () false-positive

guard) + „paired \$ „. mutation (flip a constructed „provenance item from hardening to a defect-closing status → the gate FAILs; gate-code = separate work item). Composes \$. . / \$. . /
 \$. . / \$. . / \$. . /
 \$. . ⁰¹² / \$. . / \$. . .

(H) Evidence-record full field set + stream-digest-before-truncation + recorder reconciliation (extension, spec-derived, speckit -anti-slop-enforcement,

- -). Clause (F) already mandates a machine-written verdict carrying, at minimum, the item id, guard identity, polarity, exit code, the target's artifact/build fingerprint, iteration count, and the evidence-file list – but it does not fix a COMMAND-EXECUTION field schema precise enough to RECONSTRUCT exactly what ran, where, and for how long, nor does it address a captured stream that is truncated or redacted for storage, nor whether a CITED command genuinely has a matching recorder entry. This clause closes those three gaps. () Full command-execution field set – every evidence-record entry representing a command execution MUST carry, at minimum: a START TIMESTAMP (the moment of invocation, not merely "a timestamp"); the WORKING DIRECTORY at invocation; the ARGUMENT LIST as a genuine list (never a shell-joined single string – a joined string cannot distinguish an argument containing a space from two separate arguments, the exact class of ambiguity

\$. . ()(c) already names as a path-is-part-of-the-instrument footgun); EXIT STATUS; and DURATION (wall-clock elapsed). An entry missing, or carrying an unparseable value for, any of these five fields is REFUSED at the write seam – never silently accepted as a partial entry. () Stream digest computed BEFORE truncation – when a captured stdout/stderr stream is large enough that the evidence store truncates or redacts it for storage, the entry MUST record a content digest computed over the

FULL stream BEFORE any truncation, the full stream's byte count, and a reference to where the full stream is retained OUTSIDE version control (per § . . . file-layout discipline – raw logs stay untracked by default) – so a later dispute can retrieve the full stream and verify it against the recorded digest. Truncation and redaction are themselves RECORDED FACTS on the entry (a flag + the truncation point), never silent. () **Recorder reconciliation** – any evidence-bearing claim that cites a specific command's output (a PASS citing a command, a closure citing a captured log) MUST have a MATCHING entry in the evidence recorder for the EXACT command cited; a citation with no matching recorder entry is REFUSED at the SAME status-write seam § . . . (D) already enforces for terminal statuses – the reconciliation check runs AUTOMATICALLY at that seam, never as a manual audit step. **Honest boundary** (§ . . .). This clause specifies the field SCHEMA and the truncation/reconciliation disciplines; it does not alter clause (F)'s cross-validation rules (fingerprint match, iteration count, evidence-class match), which remain in force unmodified. Recommended mechanism gates `CM-EVIDENCE-RECORD-FIELD-SET-COMPLETE` (every command-execution entry carries start-timestamp/cwd/argv-as-list/exit-status/duration; a missing or unparseable field refuses the write) + `CM-EVIDENCE-STREAM-DIGEST-PRE-TRUNCATION` (a truncated/redacted stream entry carries a full-stream digest + byte count + an outside-VCS reference, all computed BEFORE truncation) + `CM-EVIDENCE-RECORDER-RECONCILED` (every cited command has a matching recorder entry; an unreconciled citation is refused at the consuming seam) + paired § . . . mutations (strip the argument-list field and store a joined string instead → the field-set gate MUST FAIL; digest the stream AFTER truncation instead of before → the digest gate MUST FAIL; cite a command with no matching recorder entry → the reconciliation gate MUST FAIL; golden-FALSE per

11 4 6 11 4 227

\$. . (): a fully-fielded, pre-truncation-digested, reconciled entry → none of the three gates fire). Gate-code = separate work item, NOT claimed shipped (\$. . / \$. .).

Canonical authority: this Constitution.md \$. . in the HelixConstitution submodule. All consuming projects restate + cite via \$. . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --green-only-test , --skip-red-reproduction , --separate-happy-path-suffices , --synthetic-red-OK flag exists.

▣

2 0 2 6 0 6 0 3

\$. . – Real-time
conductor autonomous-test-framework sync
channel mandate (. . --dev remediation,
– –)

Short tag: conductor-test-sync-channel .

Forensic anchor (genericised, – –). An autonomous test / QA framework (HelixQA) was wired into the conductor's loop. For the conductor (the orchestrating agent / operator) to stay in real-time sync with a long-running autonomous session – knowing which phase / challenge is executing, what evidence was captured, which LLM/vision call ran, and the per-item PASS/FAIL/SKIP/OPERATOR-BLOCKED verdict the moment it lands – the framework exposed a structured append-only JSONL event stream (conduit.events.jsonl) plus an atomically-rewritten status snapshot (conduit.status.json) that the conductor tails live. Without it, the conductor either polls blindly (idle waiting, \$. . violation) or learns the outcome only at session end (no mid-session course-correction, no \$. . test-interrupt-on-discovery).

The mandate. Any autonomous, long-running test / QA / validation framework that an external orchestrator (conductor agent or operator) depends on for real-time decisions **MUST** expose a real-time sync channel with two parts: () a **structured append-only event stream** (JSONL or equivalent, one event per line, never rewritten) emitting at minimum session-start / phase-transition / per-test-or-challenge-start / captured-evidence-path / external-call (LLM / vision / sink-probe) / error / per-item-verdict events; () an **atomically-rewritten status snapshot** (a single small JSON/file written via write-temp-then-rename so a reader never observes a torn write) carrying the current session/phase/item + running counters + last verdict. Verdicts **MUST** use the closed vocabulary PASS / FAIL / SKIP / OPERATOR-BLOCKED. (per § 11.4.59 status vocabulary). The conductor tails the stream / snapshot live (`tail -f` / a dedicated monitor) so it stays in real-time sync, can § 11.4.59 -interrupt on a fresh defect, and never idles blindly per § 11.4.59 / § 11.4.59. The channel is the framework's contract with the orchestrator, not an afterthought log.

Anti-bluff for the channel itself (§ 11.4.59 / § 11.4.59.10). The event stream + snapshot are themselves captured-evidence artefacts: a verdict event **MUST** carry the evidence path that backs it (per § 11.4.59 `ab_pass_with_evidence`), so a PASS event with no evidence path is a § 11.4.59 PASS-bluff at the channel layer. A status snapshot that reports PASS while the event stream shows no captured-evidence event for that item is a contradiction the conductor **MUST** treat as FAIL. The channel never invents progress: an item with no start-event cannot have a verdict-event.

Decoupling (§ 11.4.59.10). When the framework is an owned, project-agnostic submodule, the channel **MUST** stay project-neutral – the consuming project registers its own

data (endpoints, app/package identifiers, sink hosts) at runtime via the framework's public API, never hardcoded into the framework. The channel format is generic; the payloads are consumer-supplied.

Classification: universal (§ 11.4.10) – a real-time append-only event stream + atomically-rewritten status snapshot is a platform-neutral orchestrator-sync discipline reusable by ANY project whose conductor depends on a long-running autonomous test/QA framework; the consuming project supplies the concrete channel paths and payload schema per § 11.4.5.

Composes with § 11.4.27 (test-interrupt-on-discovery – the live stream is what lets the conductor STOP the instant a fresh defect's FAIL event lands), § 11.4.28 / § 11.4.29 (every verdict event cites its captured-evidence path), § 11.4.30 (no-fakes – the framework exercises the real system, the channel reports real verdicts), § 11.4.31 (owned-submodule decoupling – channel + payloads stay project-agnostic), § 11.4.32 (status vocabulary + sink-side evidence reuse), § 11.4.33 (autonomous-validation – the channel is how an autonomous session reports without a human watching the screen), § 11.4.34 (background test execution – the channel is the conductor's window into a backgrounded long test), § 11.4.35 / § 11.4.36 (zero-idle – tailing the live channel replaces blind polling/idle waiting).

Propagation. Propagation gate

CM-COVENANT-114-116-PROPAGATION enforces the literal anchor 11.4.116 across the consumer fleet; paired § 11.4.117 meta-test mutation strips the literal → the gate FAILs.

Recommended per-family gate CM-AUTONOMOUS-FRAMEWORK-SYNC-CHANNEL (the framework emits an append-only event stream + an atomically-rewritten status snapshot, and every verdict event carries an evidence path); paired § 11.4.118 mutation

strips the evidence-path field from a verdict event → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--no-sync-channel`, `--end-of-session-report-suffices`, `--verdict-without-evidence-path`, `--torn-status-write-OK` flag exists.

§ . . — Computer-vision / OCR
pixel-oracle fallback for non-introspectable
UIs mandate (. . -dev remediation,
- -)

Short tag: `cv-ocr-pixel-oracle-fallback` .

Forensic anchor (genericised, - -).
Several apps under test (TV-Compose / leanback streaming apps, canvas-rendered players, games) return a near-empty accessibility / semantic tree — `uiautomator dump` yields a hierarchy with zero clickable nodes, so a hierarchy-driven test cannot find the play tile, cannot read the on-screen caption, cannot confirm what the user sees. The § . . rule already SKIPS UI-driven when `uiautomator dump | grep -c clickable=true == 0` . The remediation's lesson: SKIP is not the only answer — when the semantic tree is blank, the test MUST fall back to driving + asserting via the PIXELS (computer-vision template-match to locate and tap controls; ROI OCR to read on-screen text), because the pixels are what the end user actually experiences, and a hierarchy-only tool is NOT a content oracle.

The mandate. Any test that needs to drive a UI control OR assert on-screen content MUST NOT assume the accessibility / semantic / DOM hierarchy is the source of truth for what the user sees. When the hierarchy is blank, partial, or known-unreliable for the app under test (TV-Compose, leanback, canvas/ SurfaceView /GL, games, custom-rendered UIs), the test MUST fall back to a **pixel oracle**: () **drive** input by computer-vision template-match (locate a control by its rendered appearance, then tap its screen coordinates) rather than by a hierarchy node id; () **assert** on-screen content by ROI OCR (read the actual rendered text in the region the user reads – caption strip, title, error overlay) with a per-word confidence floor + region-of-interest per § 11.4.107 (), never by a hierarchy text attribute that may be absent or stale. The test tool MUST be chosen for its ability to BOTH drive input AND read pixels – a hierarchy-only tool cannot serve as a content oracle for these UIs. The pixel path is the § 11.4.107 "near-empty hierarchy proves UI-driven INFEASIBLE" fallback made constructive: not SKIP, but pixel-drive.

Anti-bluff for the pixel oracle (§ 11.4.107). The CV/OCR analyzer is itself a potential bluff surface and MUST be self-validated per § 11.4.107 (): a golden-good fixture (the control present / the expected caption rendered) MUST PASS and a golden-bad fixture (control absent / wrong-or-empty caption) MUST FAIL, wired into the meta-test. Thresholds (match confidence, OCR confidence floor) MUST be calibrated on the project's own captured frames, not hardcoded from literature (§ 11.4.107 () / § 11.4.107). An OCR PASS that cannot also FAIL its golden-bad fixture is a § 11.4.107 bluff gate.

Honest boundary (§ 11.4.107). When BOTH the hierarchy is blank AND the pixel oracle is infeasible (e.g. the target

output is a protected/secure surface that captures black per § 11.4.112, or the content is geo-unreachable), the test SKIPS-with-reason per § 11.4.113. (topology_unsupported / operator_attended / geo_restricted), never a fake PASS, and the operator-attended path is a tracked migration item per § 11.4.117, not the permanent answer.

Classification: universal (§ 11.4.114) – driving + asserting via pixels when the semantic tree is blank is a platform-neutral test-oracle discipline reusable by ANY project testing custom-rendered / canvas / game / Compose / leanback UIs; the consuming project supplies its concrete CV/OCR toolchain (template-match library, OCR engine), capture mechanism, and calibrated thresholds per § 11.4.115.

Composes with § 11.4.116 (captured-evidence quality – OCR overlay census), § 11.4.117 (no-guessing – thresholds calibrated, not assumed), § 11.4.118 (UI-driven – pixel-drive is the fallback when the hierarchy path is infeasible), § 11.4.119 (dual-approach – pixel-drive UI variant complements the intent variant), § 11.4.120 (autonomous-validation – strict refinement: near-empty hierarchy → pixel-drive, not only SKIP), § 11.4.121 (AV/test-validation – the OCR/CV analyzer is self-validated golden-good/golden-bad, ROI + confidence floor; § 11.4.122 is its UI-driving companion), § 11.4.123 (structural-impossibility – secure-surface black-capture is the honest boundary where even pixel-drive cannot read content).

Propagation. Propagation gate

CM-COVENANT-114-117-PROPAGATION enforces the literal anchor 11.4.117 across the consumer fleet; paired § 11.4.118 meta-test mutation strips the literal → the gate FAILS. Recommended per-family gate CM-CV-OCR-PIXEL-ORACLE-FALLBACK (a test targeting a known-blank-hierarchy app drives +

asserts via a self-validated CV/OCR pixel oracle, not a hierarchy-text attribute); paired § . mutation makes the pixel-OCR analyzer PASS its golden-bad fixture → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--hierarchy-is-content-oracle`, `--skip-pixel-fallback`, `--unvalidated-ocr-OK`, `--hardcoded-ocr-threshold-OK` flag exists.

1 1 8
2 0 2 6 0 6 0 3

§ . . — **Discovery-pressure to confirm known-issue-set completeness mandate**
(. . -dev remediation,
- -)

Short tag: `discovery-pressure-completeness` .

Forensic anchor (genericised, - -). A remediation cycle has a list of reported defects. Fixing exactly those is NECESSARY but NOT SUFFICIENT — the reported set is what the operator happened to notice, biased by what they tested ("we only see what we test"). The remediation ran parallel discovery + stress streams across ALL target devices, exercising subsystems the reported defects did NOT touch, to CONFIRM the reported set IS the complete critical set (and to surface any unreported defect before the user does). The discovery is anti-bluff only if it provides PROVABLE coverage: a list of the subsystems / journeys actually exercised, each with the outcome (no-new-issue vs new-issue-filed), so "we

found nothing else" is backed by "here is everything we looked at," not by "we stopped looking."

The mandate. A remediation / release cycle MUST NOT treat "every reported defect is fixed" as "the build is good." After (or in parallel with) fixing the reported set, the cycle MUST run a discovery + stress pass across ALL target devices / environments that deliberately exercises subsystems, journeys, and edge cases BEYOND the reported defects – to confirm the reported set is the COMPLETE critical set and to surface unreported defects before the end user does. The discovery pass MUST produce PROVABLE coverage: an enumerated list of the subsystems / user-journeys / stress scenarios actually exercised, each with its outcome (no-new-issue, or a new tracker entry filed per § 2.2.1.1 / § 2.2.1.2). "We found no other issues" is a § 2.2.1.1 bluff unless accompanied by "here is the enumerated set we exercised" – absence of evidence of looking is not evidence of absence. Newly-discovered defects trigger § 2.2.1.1 test-interrupt + the § 2.2.1.1 / § 2.2.1.2 isolation→RED→fix loop.

Why necessary-not-sufficient. A cycle that fixes only the reported defects ships with the operator's coverage bias baked in: the subsystems the operator did not test on the broken build are still unverified on the fixed build. Discovery pressure (everyday-user-simulation journeys + chaos/stress per § 2.2.1.1, run autonomously per § 2.2.1.2 across every device per § 2.2.1.1) is how the cycle earns the claim "this is the complete critical set," converting an unbounded "are there other bugs?" into a bounded, evidenced "here is what we exercised and what we found."

Honest boundary (§ 2.2.1.1). Discovery pressure reduces the unknown-unknown surface; it does not prove zero remaining defects (no finite test set can). The claim the cycle earns is "the reported set + the enumerated

discovery set are all addressed," with the discovery set's coverage explicitly bounded – never "the build is bug-free." Subsystems NOT exercised are stated as such (an honest coverage gap per § 11.4.17), not silently implied clean.

Classification: universal (§ 11.4.18) – discovery-pressure to confirm known-issue-set completeness is a platform-neutral release-confidence discipline reusable by ANY project with a remediation/release cycle; the consuming project supplies its subsystem inventory, everyday-user-journey set, stress scenarios, and target-environment matrix per § 11.4.19.

Composes with § 11.4.20 (test-interrupt – a discovery finding STOPS the cycle), § 11.4.21 / § 11.4.22 (the discovery pass captures evidence per subsystem exercised), § 11.4.23 (full-automation-coverage – the discovery set is part of the coverage ledger), § 11.4.24 (full-suite retest – discovery runs alongside / inside it), § 11.4.25 (iteration-discipline – discovery is the smoke-then-full progression's breadth dimension), § 11.4.26 (autonomous-validation – discovery journeys run without a human), § 11.4.27 (stress + chaos – discovery pressure IS the stress dimension applied for completeness), § 11.4.28 / § 11.4.29 (a discovered defect enters the isolation→RED→fix loop), § 11.4.30 (single-resource-owner – parallel per-device discovery needs the ownership partition).

Propagation. Propagation gate

CM-COVENANT-114-118-PROPAGATION enforces the literal anchor 11.4.118 across the consumer fleet; paired § 11.4.119 meta-test mutation strips the literal → the gate FAILs.

Recommended per-family gate CM-DISCOVERY-COVERAGE-ENUMERATED (a remediation/release cycle records an enumerated subsystem/journey discovery-coverage list with per-item

outcome, not a bare "no other issues found"); paired
§ . mutation replaces the enumerated list with a bare
no-issues claim → the gate FAILs. (Gate-code
implementation lands as a separate work item; this anchor
defines the contract.)

Canonical authority: this Constitution.md § . .
in the HelixConstitution submodule. All consuming projects
restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context.
No escape hatch — no `--reported-set-suffices` , `--skip-`
`discovery-pass` , `--no-issues-without-coverage-list` , `--assume-`
`complete` flag exists.

§ . . — **Single-resource-owner**
partitioning for parallel hardware testing
mandate (. . -dev remediation,
- -)

Short tag: `single-resource-owner-partitioning` .

Forensic anchor (genericised, - -).
Multiple parallel test/discovery streams ran against
shared physical hardware (devices D + D , an HDMI
audio sink). Two streams driving media playback on the
SAME device at once corrupt each other's evidence —
overlapping audio routes the sink reports the wrong codec,
concurrent `am start` fights over the foreground, a second
stream's input events land in the first stream's app. The
remediation enforced: exactly ONE stream OWNS each
device's exclusive/media resource at a time; other streams
targeting that device are read-only (probes / dumphys /
log capture). Parallelism is partitioned BY DEVICE —
stream A owns D 's media, stream B owns D 's media — so
the parallel-development pipeline (\$. . /

§ . .) runs without the streams clobbering each other's captured evidence.

The mandate. When multiple parallel work/test/discovery streams exercise SHARED hardware or any exclusive-access resource (a media-playback path, a single HDMI/audio sink, an exclusive device handle, a single serial/JTAG line, a GPU under exclusive capture), exactly ONE stream MUST own each such resource at a time. The exclusive owner drives the resource (playback, input injection, capture); every other concurrent stream targeting the same resource MUST be READ-ONLY (passive probes – `dumpsys` / `/proc` / `/sys` reads / sink-side network probes / log tails) for the duration. Parallelism MUST be partitioned by resource: distinct devices / sinks / handles run fully concurrent (stream-per-device), but the same device's exclusive resource is single-owner. The ownership MUST be enforced by an advisory lock / token (per § . . L contention-path locking + § . . working-tree quiescence's hardware analogue), not by convention – a second stream MUST be unable to silently start media on a device another stream owns. Ownership is event-driven: a stream claims a device when its exclusive resource frees and releases it the moment its work completes, so the next queued stream can claim it (per § . . auto-backfill).

Why single-owner. Concurrent drivers of one exclusive resource produce CROSS-CONTAMINATED evidence: the sink reports whichever stream's audio won the race, the foreground belongs to whichever `am start` landed last, input events interleave. A PASS captured under contention is a § . evidence-integrity bluff – the captured artefact does not reflect what the test under inspection actually did. Single-owner partitioning is the precondition that makes per-stream captured evidence trustworthy under parallelism.

Honest boundary (§ 11.4.17). A read-only probe stream is genuinely read-only – it MUST NOT issue any state-changing command (no `am start`, no input injection, no `setprop`, no playback) against a device it does not own; a "read-only" stream that mutates state is the contention the partition forbids. When the resource genuinely cannot be partitioned (a single device the whole batch needs to drive), the streams MUST serialize on it (single-owner over time), not run concurrently and hope.

Classification: universal (§ 11.4.17) – single-owner-per-exclusive-resource partitioning is a platform-neutral parallel-testing discipline reusable by ANY project with parallel streams contending over shared hardware / exclusive handles; the consuming project supplies its resource inventory (which devices / sinks / handles are exclusive) and lock mechanism per § 11.4.17. Strict refinement of § 11.4.17 / § 11.4.17 for the hardware-contention case.

Composes with § 11.4.17 / § 11.4.17 (captured-evidence integrity – single-owner is the precondition for trustworthy per-stream evidence), § 11.4.17 (sink-side evidence – a shared sink is the canonical single-owner resource), § 11.4.17 (deterministic consistency – contended evidence is non-deterministic, defeating N-iteration stability), § 11.4.17 (parallel-development PWU – § 11.4.17 is the hardware-resource refinement of its L contention-path locking), § 11.4.17 (F) (parallel multi-device testing – per-device streams ARE the partition), § 11.4.17 (working-tree quiescence – the hardware analogue: a stream owns the device's exclusive state cleanly before it drives it), § 11.4.17 (continuous parallel streams – single-owner + event-driven release enables the auto-backfill), § 11.4.17 (discovery pressure – parallel per-device discovery needs this partition).

Propagation. Propagation gate

CM-COVENANT-114-119-PROPAGATION enforces the literal anchor 11.4.119 across the consumer fleet; paired § . meta-test mutation strips the literal → the gate FAILs.

Recommended per-family gate

CM-SINGLE-RESOURCE-OWNER-PARTITION (parallel hardware-test streams acquire a per-resource ownership lock before any state-changing command; concurrent same-resource streams are read-only); paired § . mutation removes the ownership-lock acquisition (letting two streams drive one device) → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --allow-concurrent-resource-drivers , --skip-ownership-lock , --read-only-may-mutate , --contended-evidence-OK flag exists.

§ . . – Fix-breaks-its-own-gate reconciliation mandate (§ . . -dev remediation, – –)

Short tag: gate-reconciliation-not-fake-pass .

Forensic anchor (genericised, – –). A correct fix removed the regression-causing behaviour (the audio fix removed the handledByDualHdmi short-circuit that diverted . off the multichannel thread). But an EARLIER gate had been authored to assert the OLD – now-known-broken – behaviour's presence. After the fix, that gate FAILED. The FAIL was the CORRECT signal that the fix worked: the gate was asserting a behaviour the fix

deliberately removed. Two wrong responses are forbidden – fake-passing the gate (editing it to always pass), or reverting the correct fix to satisfy the stale gate. The right response is to RECONCILE the gate: rewrite it to assert the NEW mechanism (evidence-backed), so it once again guards the now-correct behaviour.

The mandate. When a correct fix causes a pre-existing gate / test to FAIL because that gate asserted the OLD (now-removed-or-changed) behaviour, the gate FAIL is the CORRECT signal that the fix landed – it MUST NOT be suppressed by either of the two forbidden responses: () **fake-passing** the gate – editing it to `always pass`, weakening its assertion to a tautology, or deleting it to make the FAIL disappear (a \$. bluff at the gate layer, and a \$. mutation-residue risk); () **reverting the correct fix** to satisfy the stale gate (re-introducing the defect to keep a now-wrong gate green). The required response is **RECONCILIATION**: rewrite the gate so it asserts the NEW mechanism the fix introduced, backed by captured evidence of the new correct behaviour, so the gate once again guards the right invariant. After reconciliation, the gate's paired \$. mutation MUST be updated too (the mutation must break the NEW invariant). The reconciliation MUST be a visible, evidence-cited change – "rewrote gate X to assert because fix Y removed ; new captured evidence: " – never a silent assertion-weakening.

Distinguish reconciliation from bluffing. Reconciliation rewrites the gate to assert a REAL new invariant (the fix's intended mechanism), proven by captured evidence, with the paired mutation updated to break it. Bluffing weakens the gate so it can no longer FAIL on the defect class it was meant to catch. The discriminator: after reconciliation the gate + its mutation still form a valid \$. pair (mutate the new invariant → gate FAILs); a

bluffed gate's mutation no longer makes it FAIL (the assertion became a tautology). A reconciled gate is provably still a guard; a bluffed gate is provably blind.

Honest boundary (§ 4.4.10). A gate FAIL after a fix is NOT automatically "the gate is stale, reconcile it." It MUST first be investigated per § 4.4.11 – the FAIL may be the gate correctly catching a REGRESSION the fix introduced (the fix broke something else), in which case the FIX is wrong, not the gate. Reconcile ONLY when investigation PROVES the gate asserted old-correct-now-removed behaviour AND the new behaviour is the intended, evidence-confirmed mechanism. "The gate is just stale" without that investigation is a § 4.4.12 guess that can mask a real regression.

Classification: universal (§ 4.4.13) – reconcile-the-gate-don't-fake-pass-and-don't-revert is a platform-neutral fix-vs-gate-conflict discipline reusable by ANY project with regression gates; the consuming project supplies its gate framework and paired-mutation mechanism per § 4.4.14.

Composes with § 4.4.15 (FAIL-bluffs – fake-passing the gate is the inverse bluff), § 4.4.16 (test-interrupt – a post-fix gate FAIL triggers investigation before reconcile), § 4.4.17 (no-guessing – "stale gate" is a hypothesis until investigation proves it), § 4.4.18 (working-tree quiescence – a gate edited to always pass is mutation residue that must never ship), § 4.4.19 (systematic-debugging – the mandatory investigation that distinguishes stale-gate from new-regression), § 4.4.20 (four-layer fix-verification – the reconciled gate asserts the fix's runtime signature / new mechanism), § 4.4.21 (paired mutation – the reconciled gate's mutation must break the NEW invariant, re-proving it is a guard not a tautology).

Propagation. Propagation gate

CM-COVENANT-114-120-PROPAGATION enforces the literal anchor 11.4.120 across the consumer fleet; paired § . meta-test mutation strips the literal → the gate FAILs.

Recommended per-family gate

CM-GATE-RECONCILED-NOT-FAKE-PASSED (when a fix changes a behaviour a gate asserted, the gate's reconciliation is an evidence-cited rewrite whose paired mutation still makes it FAIL – never an assertion weakened to a tautology); paired § . mutation weakens a reconciled gate to always pass and asserts its mutation no longer makes it FAIL → the meta-gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

GATE SEAM-PLACEMENT – a gate whose precondition the gated work itself produces is on the **WRONG SEAM** (extension, research-derived, – –). § . 's trigger is a gate that FAILs – i.e. a gate that RAN. Its sibling case is a gate that CANNOT RUN in the ordering imposed on it, and the required response is different: not reconciliation, but re-seating the gate at the seam where it is satisfiable. Forensic FACT (genericised): work was ordered "run checks E -E , then implement P only if E -E pass" – but the gate's own text defined E as running on a P -gated build, so E could not exist before P existed; the order was unexecutable by construction. The source of truth, read exactly, gated something else entirely: building behind a default-off flag was permitted, and it was **ENABLING** the flag that required E -E – the gate was on ACTIVATION, and had been imposed on AUTHORSHIP. Therefore: **(a) EXECUTABILITY IS CHECKED BEFORE A GATE IS IMPOSED** – if gate G requires artifact A, and A is produced by the very work G gates, then G is unsatisfiable at that seam and belongs at a LATER one (activation / enablement / release), NEVER at authorship; imposing it at authorship makes the work

impossible to start and the gate impossible to pass, which is not rigor, it is a deadlock. (b) **RE-READ THE SOURCE OF TRUTH BEFORE RE-SEATING** – the circularity is usually an ARTIFACT OF THE ORDER, not of the gate: the gate's own text names its real seam, and that text governs over any restatement of it (a restatement that moved the gate from "enabling" to "writing" is a § . . misstatement of the requirement, and the fix is to obey the source, not to weaken the gate). (c) **THE FORBIDDEN RESPONSES ARE** § . . 's **OWN** – an unsatisfiable gate **MUST NOT** be fake-passed, deleted, weakened to a tautology, nor improvised past by building the artifact and retro-declaring the gate met; and the gated work **MUST NOT** be abandoned as "blocked" when the gate simply sits on the wrong seam. **Honest boundary** (§ . .): an unrunnable gate is **NOT** automatically mis-seated – investigate per § . . first; it may be correctly seated and its precondition genuinely missing, which is the § . . artifact_not_yet_built **NOT-YET-RUNNABLE** state (§ . . seam-dependent), not a seam defect. Reporting the circularity is mandatory and already required by § . . / § . . / § . . / § . . () – this clause adds the seam-placement rule those anchors do not carry. § . . 's clash detector cannot catch this class: 2026 07 17 its input is a change DIFF, and an ordering instruction is not a diff.

Gate contract for the seam-placement clause (extension, - -). **CM-GATE-RECONCILED-NOT-FAKE-PASSED** is **EXTENDED**: an imposed gate whose required precondition artifact is produced by the very work it gates → **FAIL** (the gate is mis-seated and **MUST** be re-seated at activation/enablement/release), and the **FAIL** names the circularity rather than the work. Paired § . . mutation: impose an authorship-seam gate whose precondition only the gated work can produce → the gate **FAILs**; re-seat it at

activation → PASSes. Its golden-FALSE fixture MUST include a correctly-seated gate whose precondition is merely not-yet-built (the § . . . `artifact_not_yet_built` state) – the gate MUST NOT fire on that, or it becomes the § . . . () false positive that condemns a healthy gate. Gate-code = a separate work item; this is the contract, not a claim the code has shipped (§ . . .).

Canonical authority: this Constitution.md § . . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no `--fake-pass-stale-gate` , `--revert-fix-for-gate` , `--weaken-assertion-to-pass` , `--delete-failing-gate` , `--impose-unsatisfiable-gate` , `--improvise-past-circular-order` , `--gate-authorship-on-its-own-output` flag exists.

1 1 8

1 0 2 0 0 6 0 0

§ . . . – **No-commit-while-build-writes-tracked-artifacts** mandate (. . . - dev remediation, - -)

Short tag: `no-commit-during-artifact-write` .

Forensic anchor (genericised, - -). The build pipeline's per-component build steps write produced artifacts (forked-app APKs) INTO tracked directories that are under version control (`prebuilt_apps/`). Running `git add -A` / a commit while those steps are mid-write stages a PARTIAL or stale artifact – a race between the build writing the file and git reading it – so the commit captures a half-written or pre-rebuild artifact. The remediation deferred the commit to BUILD-COMPLETION: the tree is quiescent (no build step writing tracked files) AND the freshly-built artifacts are captured whole (no stale prebuilt in git). This is the artifact-write

analogue of `$ . .` (no commit while a mutation gate is in flight) at the build-output layer.

The mandate. A commit (especially `git add -A` / any broad stage) MUST NOT run while a build / packaging / generation step is actively writing artifacts INTO tracked directories. Doing so races the writer and stages a partial or stale artifact – the commit captures bytes that are mid-write or pre-rebuild, a `$ . .` SOURCE→ARTIFACT integrity failure landed in version control. The commit MUST be deferred until the build step that writes tracked artifacts has COMPLETED, so that: () the working tree is quiescent at the artifact layer (no step writing a tracked file); () the committed artifacts are the FRESH, whole outputs of the completed build (no stale pre-rebuild artifact, no half-written file). Before committing tracked build outputs, the committer MUST verify the writing step finished (process exit, completion marker, or per-artifact mtime ≥ build-start) – a build still in flight that writes tracked dirs is a HOLD on the commit, not a race to win.

Composition with `$ . .` . `$ . .` forbids committing while a mutation gate is in flight in the checkout (mutation residue). `$ . .` is the BUILD-OUTPUT analogue: forbid committing while a build is in flight WRITING TRACKED ARTIFACTS (partial/stale artifact residue). Both close the same class of defect – a commit captures transient, non-final tree state – at two different write-sources (mutation experiments vs build outputs). Where build outputs land OUTSIDE version control (`gitignored out/ / dist/`), the race does not apply; the mandate binds specifically to build steps writing into TRACKED directories.

Honest boundary (`$ . .`). "The build probably finished" is not "the build finished" – verify with a completion signal (exit code / marker / mtime), not an

assumption, before committing tracked artifacts.

Committing source changes that DON'T touch the build's tracked-artifact directories is fine mid-build (those files are not being written by the build); the hold applies only to the tracked dirs the in-flight build writes.

Classification: universal (§ 11.4.17) – deferring commit until a tracked-artifact-writing build completes is a platform-neutral commit-integrity discipline reusable by ANY project whose build writes produced artifacts into version-controlled directories; the consuming project supplies which build steps write which tracked directories and its completion-signal mechanism per § 11.4.18. The PROJECT-SPECIFIC instance (which tracked directory – e.g. a prebuilt-apps directory – and which build steps write it) is recorded in the consuming project's own governance per § 11.4.19.

Composes with § 11.4.20 (no-guessing – verify build completion, never assume), § 11.4.21 (no-versioned-build-artifacts – artifacts that ARE legitimately tracked, like signed prebuilts, are the exact case this mandate protects from partial commit; gitignored build outputs are exempt), § 11.4.22 / § 11.4.23 (parallel-development – the commit/merge stage is serialized AFTER the build stage for this reason), § 11.4.24 (working-tree quiescence – § 11.4.25 is its build-output analogue), § 11.4.26 (background-push – commit timing is decoupled from push, but the commit itself still waits for artifact-write completion), § 11.4.27 (four-layer fix-verification – a partial/stale committed artifact is the ARTIFACT-layer integrity failure this prevents from entering version control), § 11.4.28 (safe-parallel-work-with-long-build – committing tracked build outputs mid-build is an UNSAFE operation in its catalogue).

Propagation. Propagation gate

CM-COVENANT-114-121-PROPAGATION enforces the literal anchor 11.4.121 across the consumer fleet; paired § . meta-test mutation strips the literal → the gate FAILs.

Recommended per-family gate

CM-NO-COMMIT-DURING-ARTIFACT-WRITE (the commit wrapper refuses to stage tracked build-artifact directories while a build step writing them is in flight – checks a build-in-progress lock / completion marker); paired § . mutation removes the build-in-flight check from the commit wrapper → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md § . in the HelixConstitution submodule. All consuming projects restate + cite via § . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --commit-during-build , --stage-partial-artifact-OK , --assume-build-finished , --skip-build-completion-check flag exists.

2026 06 03

§ . . – No-silent-removal-of-existing-components-without-operator-confirmation mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Never ever remove any application, system component or service from already existing codebase / System

without interactively asked question to us! THIS IS MANDATORY RULE / CONSTRAINT!"

Forensic case study (this project, FACT – captured).

During the . . . -dev burn-down, two shipped capabilities – F (an Apple-TV-class application) and F (a Huawei HMS / Mobile-Services component) – were removed from the existing System WITHOUT first asking the operator. The operator reversed both removals. A removal that the operator has to discover and reverse after the fact is a defect of the same severity class as a § . PASS-bluff: the System silently lost a user-facing capability the operator never agreed to drop.

The mandate. No application, system component, service, package, feature, driver, module, library, prebuilt asset, or any other already-existing capability of the existing codebase / shipped System may be removed (deleted, dropped from the package set, disabled-into-non-shipping, un-bundled, de-listed, or otherwise made unavailable to the end user) WITHOUT FIRST interactively asking the operator and receiving an EXPLICIT keep-or-remove decision. The question MUST be posed through the platform's interactive clarification mechanism per § . . .

(AskUserQuestion on Claude Code; the documented interactive prompt on other runtimes) – NEVER a free-text "should I remove X?" buried in narrative, NEVER a silent removal justified post-hoc, NEVER an autonomous removal decision. A silent removal is a **release blocker** regardless of how well-intentioned the rationale (deduplication, "it was broken anyway", "geo-restricted", "incompatible", "superseded") – the operator decides, the agent asks.

What counts as a removal (non-exhaustive). Deleting an app/APK/binary from the build's package set (e.g. a project's PRODUCT_PACKAGES / device.mk / equivalent manifest), removing a service from the init/boot/service-

registry set, dropping a kernel module / driver / config from the shipping configuration, un-bundling a prebuilt asset, deleting a submodule or a submodule's shipped output, removing a feature flag that gated a live capability, or any edit whose NET EFFECT is "an end-user capability that shipped before no longer ships." Adding, replacing-with-equivalent-or-better (where the operator has approved the replacement), or fixing a capability is NOT a removal. When uncertain whether an edit constitutes a removal, treat it AS a removal and ask (per § 11.4.122 no-guessing + § 11.4.66 block-only-when-truly-needed – removal of an existing user-facing capability is high-blast-radius and MUST be operator-confirmed, never autonomously decided).

Composition. § 11.4.66 composes with § 11.4.66 (interactive-clarification – the removal question is asked via the platform interactive mechanism, not free-text), § 11.4.66 (autonomous-decision-over-blocking – a removal of an existing user-facing capability is exactly the irreversible-ish, high-blast-radius class § 11.4.66 RESERVES for an operator question rather than an autonomous decision; § 11.4.66 makes the "ask the operator" outcome mandatory for this class, NOT optional), § 11.4.66 (Obsolete classification – the tracked path to actually DROP a capability is: ask the operator → operator approves → mark the item Obsolete (→ Fixed.md) with Obsolete-Details reason feature-removed AND an operator-approval citation → then remove; the removal never precedes the operator's yes), § 11.4.66 (structural-impossibility won't-fix – even a capability proven structurally impossible to keep is closed via the § 11.4.66 path with operator awareness, not silently deleted), § 11.4.66 (no-guessing – "the operator probably wants this gone" is the exact guess forbidden), § 11.4.66 / § 11.4.66 (full-suite-retest /

iteration-discipline – a removal landed without the operator's yes fails the release gate).

Enforcement. Propagation gate

CM-COVENANT-114-122-PROPAGATION enforces the literal anchor 11.4.122 across the consumer fleet (every CLAUDE.md / AGENTS.md / QWEN.md). Recommended gate CM-NO-SILENT-COMPONENT-REMOVAL (a diff-driven detector that flags a removal of a shipped capability – a deletion from the project's package/service/module manifest set – without a matching operator-approval citation in the commit / tracker entry, per the § 11.4.122 change-impact-clash pattern; a removal whose tracker item is Obsolete (→ Fixed.md) with reason feature-removed + operator-approval citation passes). Paired § 11.4.122 meta-test mutations: strip the 11.4.122 literal → propagation gate FAILs; introduce a manifest removal with no operator-approval citation → CM-NO-SILENT-COMPONENT-REMOVAL FAILs. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§ 11.4.122) – "never silently remove an existing end-user capability without operator confirmation" is a platform-neutral discipline reusable by ANY project that ships a set of user-facing capabilities; the consuming project supplies its concrete capability-manifest paths (package set, service registry, module config) per § 11.4.122.

Canonical authority: this Constitution.md § 11.4.122 in the HelixConstitution submodule. All consuming projects restate + cite via § 11.4.122 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --remove-without-asking, --silent-removal, --autonomous-removal-OK, --dedup-removal-exempt, --it-was-broken-anyway flag exists.

§ . . – Rock-solid-proof-or-deep-research mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Every single reported issue MUST BE fully and % validated with rock solid proofs! Nothing can be considered fixed or completed without hard evidence! No false results or bluff(s) of any kind is allowed! If we are not sure on how to achieve full testing, validation and verification of something we MUST ALWAYS perform deep web research for all possible data (articles, documentation, guides, and other resources) and opensourced codebases which we can use to solve our problems and perform testing with validation and verification which produces rock-solid evidence(s) and leaves no space for false results or any kind of bluff!"

Forensic case study (FACT – captured). In the . . - dev remediation cycle the validation method for two feature classes was, at first, genuinely unclear: relocating a FLAG_SECURE secure surface to a secondary display (whose pixel capture returns black) and asserting on-screen content in non-introspectable streaming-app UIs (whose accessibility hierarchy is blank). Rather than declaring those classes "untestable" or accepting a metadata-only PASS, the cycle performed deep web research (docs/research/testing_frameworks_20260603/) into platform/HDCP behaviour, computer-vision template-matching, ROI OCR with confidence floors, freeze-detection / frame-advance liveness oracles, and network-side sink probes. That research yielded the CV/OCR/liveness/sink-probe oracle stack (now § . . + § . . +

§ 1.1.1) that made rock-solid, evidence-producing validation possible where it had appeared impossible. The lesson – "unclear how to validate" is a research trigger, NEVER a bluff licence – is now constitutional.

The mandate. Every single reported issue, every fix, and every claimed completion MUST be fully and % validated with rock-solid CAPTURED proof per

§ 1.1.2 / § 1.1.3 / § 1.1.4 before it may be marked fixed / implemented / completed (the § 1.1.5 terminal closure vocabulary). Hard, captured evidence is the bar – nothing may be considered fixed or complete without it. Metadata-only PASS, configuration-only PASS, absence-of-error PASS, and grep-without-runtime PASS are all forbidden (§ 1.1.6 / § 1.1.7) – no false results and no bluff of any kind, at any layer, are allowed.

The research-or-don't-bluff rule (the operative addition).

When the agent is UNSURE how to fully test, validate, or verify something – when no obvious evidence-producing method exists, OR the candidate method would yield only metadata/config/absence-of-error evidence (and is therefore a bluff per § 1.1.8) – the agent MUST ALWAYS first perform deep web research per § 1.1.9 + § 1.1.10 (official documentation, technical articles, guides, vendor references, standards, issue trackers, and open-source codebases the project can reuse or adapt) to DISCOVER or BUILD a validation method that produces rock-solid evidence and leaves no space for a false result. Declaring something "untestable", "not automatable", or accepting a metadata-only PASS WITHOUT having first exhausted this deep-research path is itself a § 1.1.11 violation – the same severity class as a PASS-bluff. The research output (cited source URLs + the evidence-producing method it yielded, OR the literal "NO external solution found – original work" per § 1.1.12)

is the captured proof that the path was exhausted. Only after that research genuinely fails to yield any evidence-producing method may the item be classified

PENDING_FORENSICS: / Operator-blocked (\$. .) / structurally-impossible won't-fix (\$. .) – with the cited research as the evidence that the classification is earned, never a convenience.

Composition. \$. . composes with \$. . (captured-evidence quality is the proof shape), \$. . (no-guessing – "probably fixed" without captured evidence is forbidden; the closure must be FACT), \$. . (deep-web-research-before-implementation – \$. . extends it from "before designing a fix" to "whenever the VALIDATION method is unclear"), \$. . (autonomous-validation – the research must yield an autonomous, evidence-producing path, wherever one is discoverable), \$. . (universal sink-side positive-evidence taxonomy – every PASS cites a taxonomy-class artefact), \$. . (latest-source verification of the discovered method), \$. . (the AV/test-validation oracle techniques the . . - dev research yielded), \$. . (discovery-pressure – the same "absence of evidence of looking is not evidence of absence" principle applied to per-issue validation). It also composes with \$. . / \$. . (the earned escape valves – reachable ONLY after the deep-research path is exhausted with cited evidence).

Enforcement. Propagation gate

CM-COVENANT-114-123-PROPAGATION enforces the literal anchor 11.4.123 across the consumer fleet (every CLAUDE.md / AGENTS.md / QWEN.md). Recommended gate CM-ROCK-SOLID-PROOF-OR-RESEARCH (a closure-audit detector that flags any item closed Fixed / Implemented / Completed whose tracker entry lacks a captured-evidence artefact path per \$. . ,

AND any PENDING_FORENSICS: / Operator-blocked / structurally-impossible classification whose entry lacks a cited deep-research trail per § . . /§ . .). Paired § . meta-test mutations: strip the 11.4.123 literal → propagation gate FAILs; close an item with a metadata-only PASS and no captured-evidence path → CM-ROCK-SOLID-PROOF-OR-RESEARCH FAILs; classify an item Operator-blocked with no cited research trail → the same gate FAILs. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§ . .) – "every closure carries rock-solid captured proof, and unclear-how-to-validate triggers deep research rather than a bluff" is a platform-neutral discipline reusable by ANY project; the consuming project supplies its concrete capture mechanisms + research corpora per § . . .

Canonical authority: this Constitution.md § . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --metadata-pass-suffices , --skip-proof , --untestable-without-research , --config-only-closure-OK , --bluff-when-unsure flag exists.

§ . . – Dead/unwired-code investigate-before-remove mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Before removing any seemingly-dead (zero-importer / unwired) codebase, we MUST investigate via git history where/how it was originally used and how it became

dead. Removal is permitted ONLY when we have captured PROOF it is genuinely no longer needed – and that removal MUST be its own separate commit with a proper descriptive message. If there is no such proof, the code MUST be investigated for where/how it should be wired in properly, and any missing or unwired tests MUST be added. We MUST ALWAYS be extra careful with any codebase removal."

The defect class. "This has zero importers / is never called / is unwired, so it is dead – delete it" is a guess (\$. .), not a finding. A `git grep` / static-analysis "no references" result proves only *current* non-reference; it does NOT prove the code is genuinely unneeded. Code becomes apparently-dead for several distinct reasons – its call-site was deleted by mistake (a regression: the code SHOULD still be wired in), its wiring was never finished (it is intended but unwired – the fix is to WIRE it, plus its tests), it is reflectively / dynamically / build-tag / plugin / DI-container referenced (the static tool simply cannot see the reference), or it is genuinely obsolete and safe to remove. Deleting it without first distinguishing which case applies risks silently removing a needed-but-temporarily-unwired capability – the same severity class as the \$. . silent-removal defect and a \$. PASS-bluff at the maintenance layer.

The mandate. Before removing ANY seemingly-dead codebase element (a zero-importer / never-called / unwired / unreferenced function, method, type, file, module, package, asset, config, or build target), the agent MUST FIRST investigate via git history (`git log --follow`, `git log -S<symbol>` / `-G<regex>` pickaxe across all history, `git log --diff-filter`, blame on the deletion of the former call-site) to establish, as captured FACT (\$. .):

() WHERE and HOW it was originally used / wired in; () WHEN and HOW it became dead – its call-site removed deliberately, removed by mistake (regression), never completed, or made unreachable by a refactor; () whether the "no references" result is real OR an artefact of a reference the static tool cannot see (reflection / dynamic dispatch / build-tags / codegen / DI / plugin registry / cross-language FFI / config-driven wiring). The investigation output (the cited commits + the determination) is the captured evidence.

Removal is conditional – proof-of-genuinely-unneeded

REQUIRED. Removal is permitted ONLY when the investigation produces captured PROOF the element is genuinely no longer needed (its purpose is fully superseded, its former caller is intentionally and permanently gone, no hidden reference exists). When that proof exists, the removal MUST be its OWN SEPARATE COMMIT (not bundled with unrelated work – so it is independently reviewable and revertible, composing with § . . . working-tree quiescence + § . . . multi-pass) carrying a proper descriptive message that cites the git-history investigation evidence (the commits that wired it in, the commit that orphaned it, the determination that it is unneeded) and – where the element is an end-user-facing capability – the § . . . operator-confirmation. The § . . . tracked path applies: classify the tracker item `Obsolete (→ Fixed.md)` with the appropriate reason.

No proof ⇒ wire it in + add the missing tests (do NOT delete). If the investigation does NOT yield proof the element is unneeded, the element MUST NOT be removed. Instead it MUST be investigated for WHERE and HOW it should be wired in properly (restore the deleted call-site if its removal was a regression per § . . . ; complete the intended wiring if it was never finished), AND any missing or unwired tests for it MUST be added (the

absence of a test is itself part of why it drifted into apparent-deadness – composes § . . . %-test-type-coverage + § . . . TDD-fix + § . . . RED-baseline). Apparent-deadness with no removal-proof is a wiring/coverage GAP to close, never a delete-on-sight.

11 4 5

Extra-caution default. The agent MUST ALWAYS be extra careful with any codebase removal: when uncertain whether removal-proof is sufficient, default to NOT removing (investigate + wire + test), per § . . . no-guessing + § . . . (an irreversible-ish, possibly-needed removal whose safety cannot be determined from evidence is exactly the class that blocks to the operator rather than being autonomously deleted) + § . . . (removal of a live end-user capability requires operator confirmation). "Probably dead" is never sufficient; the bar is captured proof.

11 4 3

Composition. § . . . composes with § . . . (no-guessing – "zero importers therefore dead" is the exact guess forbidden; the determination must be FACT from git history), § . . . (deep-research – git history IS the primary research corpus for how the code was used + how it died; external research applies when the wiring mechanism is unclear), § . . . / § . . . (the separate-removal-commit discipline – quiescent tree, multi-pass evaluation, independently revertible), § . . . (Obsolete classification is the tracked path for a proven removal), § . . . (autonomous-decision-over-blocking – a removal whose safety cannot be proven blocks rather than proceeds), § . . . (last-known-good-tag isolation – a call-site deleted by mistake is a regression to repair, the same diagnostic the "became dead by mistake" case demands), § . . . (no-silent-removal – § . . . is its complement: § . . . governs removal of KNOWN-shipped capabilities by operator confirmation, § . . .

governs removal of SEEMINGLY-dead code by git-history proof; an end-user-facing seemingly-dead element triggers BOTH), § . . . / § . . . / § . . . (the missing-tests obligation when the element is wired in rather than removed).

Enforcement. Propagation gate

CM-COVENANT-114-124-PROPAGATION enforces the literal anchor 11.4.124 across the consumer fleet (every CLAUDE.md / AGENTS.md / QWEN.md). Recommended gate CM-DEAD-CODE-INVESTIGATE-BEFORE-REMOVE (a commit-audit detector: any commit whose diff is a NET DELETION of a code element – function/type/file/module/package/asset/build-target – MUST be a removal-only commit whose message cites the git-history investigation evidence, OR the deletion is part of a tracked Obsolete (→ Fixed.md) item per § . . . ; a net-deletion commit bundled with unrelated changes OR lacking the investigation citation FAILS). Paired § . meta-test mutations: strip the 11.4.124 literal → propagation gate FAILS; introduce a net-deletion commit with no investigation citation and not part of an Obsolete item → CM-DEAD-CODE-INVESTIGATE-BEFORE-REMOVE FAILS. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§ . . .) – "investigate git history + capture proof-of-unneeded before removing seemingly-dead code, otherwise wire it in and add its tests; removal is its own descriptive commit" is a platform-neutral maintenance discipline reusable by ANY codebase; the consuming project supplies its concrete static-analysis / importer-graph tooling and its hidden-reference mechanisms (reflection, DI, build-tags, codegen) per §

Canonical authority: this Constitution.md § . .
in the HelixConstitution submodule. All consuming projects
restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context.
No escape hatch — no `--zero-importers-means-dead`, `--delete-unwired-on-sight`, `--skip-git-history-investigation`, `--remove-without-proof`, `--bundle-removal-with-other-work` flag exists.

2 0 2 6 0 6 0 4

§ . . — **Code-review-agent gate**
before pre-build + main build (mandatory
multi-layer review) (User mandate,
— —)

Forensic anchor — verbatim user mandate
(— —):

"After all fixes/changes/implementations are done,
BEFORE running pre-build tests and the main build,
dispatch code-review agent(s) that analyze all work
done + all existing data/facts + the existing codebase
+ current git history to determine quality, safety,
and whether the fixes/changes will REALLY work; they
MUST validate and verify that every test covering the
fixes/changes genuinely validates the work with NO
chance of false results or bluff of any kind. Any
finding MUST be fixed, polished, improved, and covered
with additional tests before the build proceeds.
Multiple strong layers of checks."

The defect class. Running the pre-build gate sweep and the
main (artifact) build directly after a batch of fixes/
changes — without an independent review pass first — lets
a defective fix, an error-prone change, a bluff test, or a
test that does not genuinely validate the work-under-test
consume an expensive build cycle and (worse) ship green.

Pre-build gates and the build prove the codebase is internally consistent + buildable; they do NOT prove the work is correct, safe, will-really-work for the end user, or that the tests covering it actually catch its negation. That assurance is an independent-review responsibility that MUST run BEFORE the build, not after.

The mandate. After all fixes / changes / implementations in a batch are done, and BEFORE running the pre-build test sweep and the main build (for ANY project), the agent MUST dispatch one or more dedicated code-review agent(s) (subagent-driven by default per § 1.1.1 / § 1.1.2) that perform a multi-layer review. The review MUST: () **analyze ALL work done in the batch** – every fix / change / implementation, its source diff, and its stated intent; () **analyze ALL existing data + facts** – captured evidence per § 1.1.3 / § 1.1.4 / § 1.1.5, tracker entries, prior findings, the § 1.1.6 runtime-signature registry; () **analyze the existing codebase** – the blast radius of each change (§ 1.1.7 multi-pass), cross-feature interaction, contract integrity of every dependency; () **analyze current git history** – what each change touched, how it composes with concurrent / recent work, whether it reproduces a known-broken pattern (§ 1.1.8 / § 1.1.9 git-history discipline); () **determine quality, safety, and will-it-REALLY-work** – assess whether each change is robust + not error-prone (will not solve A while creating B), safe (no host-safety / data-safety / security regression), and genuinely delivers the end-user-visible behaviour it claims (§ 1.1.10 / § 1.1.11); () **validate + verify the tests covering the work** – confirm every test that covers a fix/change genuinely exercises the work-under-test and catches its negation, with ZERO chance of a false result or bluff of any kind (a test that PASSES on broken-for-the-user work, a metadata-only / config-only / absence-of-error / grep-without-runtime

assertion, or a gate whose paired § . mutation does not make it FAIL is a finding the review MUST surface per § . . / § . . / § . . / § .).

Findings block the build. Any finding the review surfaces – a defect, an error-prone change, a safety risk, a will-not-really-work conclusion, a bluff/false-result-capable test, a missing-coverage gap – MUST be fixed, polished, improved, and covered with additional tests (four-layer per § . . (b), TDD-RED-first per § . . / § . .) BEFORE the pre-build sweep and main build proceed. The review MUST iterate (re-review after each remediation) until it surfaces no remaining blocking findings; only then does the build proceed. The review is itself anti-bluff: its conclusions are captured evidence (§ . . / § . .), and a review that rubber-stamps a defective batch is the same severity class as a PASS-bluff.

Multiple strong layers. The review is one of multiple strong layers of checks, complementing – never replacing – the pre-build gate sweep (§ four-layer), the § . . multi-pass change-evaluation, the § . . four-layer fix-verification, the § . . pre-build build-readiness verdict + clash detector, and the post-build / runtime-on-clean-target / user-visible layers. § . . sits at the seam between "all batch work done" and "start the pre-build + build": it is the independent-review gate that the other layers do not provide (the multi-pass per § . . is self-review by the change author; § . . adds the structurally-separated reviewer seam per § . .).

Composition. § . . composes with § . (the end-user-quality covenant the review enforces), § . . (FAIL-bluffs equally caught), § . . (four-layer coverage of every remediation the review

demands), § 11.4.103 (no-guessing – review conclusions are FACT from captured evidence, never "looks fine"), § 11.4.104 (the full-suite retest is a later gate; § 11.4.105 is the pre-build review gate that precedes it), § 11.4.106 (TDD-fix – remediation lands RED-first), § 11.4.107 (deterministic-consistency of the tests the review validates), § 11.4.108 / § 11.4.109 (the review IS a subagent dispatch – structurally-separated reviewer eliminates self-review blind spots), § 11.4.110 (multi-pass change-evaluation is the author-side pass; § 11.4.111 is the independent-reviewer pass), § 11.4.112 (any defect the review finds triggers systematic-debugging before remediation), § 11.4.113 (the review verifies AV/test-validation analyzers are not themselves bluffs), § 11.4.114 (the review checks each fix's runtime signature is declared + verifiable), § 11.4.115 (the build-readiness verdict + clash detector are mechanical gates the review complements with judgement).

Enforcement. Propagation gate

CM-COVENANT-114-125-PROPAGATION enforces the literal anchor 11.4.125 across the consumer fleet (every CLAUDE.md / AGENTS.md / QWEN.md). Recommended gate CM-CODE-REVIEW-GATE-BEFORE-BUILD (the build orchestrator records a code-review-completed marker – with the review's findings + their remediation-or-no-blocking-findings determination – produced AFTER the batch's last fix and BEFORE the pre-build sweep + main build; a build that starts without a fresh review marker for the current batch FAILS). Paired § 11.4.125 meta-test mutations: strip the 11.4.125 literal → propagation gate FAILS; start a build with no fresh code-review marker for the batch → CM-CODE-REVIEW-GATE-BEFORE-BUILD FAILS. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§ 1.1.1) – "an independent multi-layer code-review-agent gate analyzing all work + data + codebase + git history for quality/safety/will-it-really-work + genuine test-validation, with findings blocking the build, runs BEFORE the pre-build sweep + main build" is a platform-neutral discipline reusable by ANY project that builds an artifact from source; the consuming project supplies its concrete review-marker mechanism + build-orchestrator integration per § 1.1.1.1.

Canonical authority: this Constitution.md § 1.1.1 in the HelixConstitution submodule. All consuming projects restate + cite via § 1.1.1 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no `--skip-code-review`, `--build-without-review`, `--no-review-gate`, `--review-optional`, `--trust-the-author` flag exists.

§ 1.1.1 – Default autonomous-loop working mode from first prompt (User mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"Make sure that you continue work in endless fully autonomous loop, do not stop until new fully validated and verified version (tag) is created and published (all submodules and main repo) or IN A CASE OF some other main stream work until it is fully completed with all side work streams and nothing else is left in our working queue! THIS MUST BE ALWAYS the default working mode without us asking you! We tend to achieve ABSOLUTE EFFICIENCY, with this and all other projects

which will incorporate this MANDATORY RULE / CONSTRAINT!!! This way of (your) working will be ALWAYS applied / followed / executed / fully respected, as soon as we assign / send first request (prompt) in the session! This stops only if we explicitly say so or nothing is left to be done in current working scope (release that will come / upcoming version)!!! Any mimicking (imitation) of this behavior / rules / mandatory constraints, false results or any kind of bluff(s) is ABSOLUTELY FORBIDDEN!!!"

The endless fully-autonomous loop is the **DEFAULT working mode**, engaged automatically the moment the operator sends the FIRST request / prompt of a session – the operator MUST NOT have to ask for it, request it, restate it, or re-enable it per session. § 1.1.1 framed the endless-loop covenant as an explicit-instruction opt-in ("continue in endless loop fully autonomously" or a semantically-equivalent phrasing); § 1.1.2 is the **capstone** that promotes the same covenant to always-on: from the first prompt onward, every agent operates in the § 1.1.1 loop discipline as the standing default, with § 1.1.3 zero-idle, § 1.1.4 maximum-idle-use, § 1.1.5 autonomous-decision-over-blocking, and § 1.1.6 continuous-parallel-stream all engaged by default – no per-session activation handshake.

The continuation contract (the loop continues until ONE of the two terminal conditions holds): (A) **Release scope** – a new, fully-validated-and-verified version (tag) is created AND published across all owned submodules AND the main repo to all configured remotes (per § 1.2.1 multi-upstream push + § 1.2.2 full-suite-retest-before-tag + § 1.2.3 absolute-no-force-push merge-onto-latest-

main). (B) **Non-release main-stream scope** – the main-stream work goal is fully completed AND every side work stream is done AND the working queue holds nothing left to do for the current scope. Until ONE of (A) or (B) holds, the agent MUST keep working – it claims the next priority item, dispatches the next parallel stream, or progresses every non-blocked item, per § . . . / § . . . / §

The loop STOPS ONLY on: () the operator explicitly saying so (STOP / pause / end); () nothing left to do in the current working scope – the upcoming release / the current main-stream goal – with the queue genuinely empty per the (A)/(B) terminal conditions; () a § host-session-safety demand (the loop yields to host safety unconditionally). Idle-while-blocked parks one work unit, it does not stop the loop – the agent keeps progressing every non-blocked item in parallel per § . . . + § . . . + §

Goal – ABSOLUTE EFFICIENCY. The discipline exists to drive the project to completion with no operator-side restart overhead, no idle gaps, no stop-and-wait round-trips. It applies to this project AND every project that incorporates this Constitution (universal per § . . .).

Anti-bluff – actually do the work, never imitate it. Mimicking / imitating this loop behaviour, narrating continuation without performing it, fabricating progress, or emitting false / bluff results of ANY kind is ABSOLUTELY FORBIDDEN – this composes the entire § . anti-bluff covenant family (§ . / § . . / § . . / § . . / § . . / § . . / § . .). The agent MUST genuinely perform the continuous work and capture positive evidence for every closure; a report that claims the loop ran while no real work / no captured

evidence was produced is a § 11.4.17. PASS-bluff at the operating-mode layer, the severity class this mandate specifically prohibits.

Classification: universal (§ 11.4.17) – a platform-neutral operating-mode discipline reusable by ANY project; the consuming project supplies its concrete release-tag / publish mechanism + working-queue source .per
§ 11.4.18. Composes with § 11.4.19 (the endless-loop covenant – § 11.4.20 promotes it from opt-in to always-on default) / § 11.4.21 (zero-idle priority-first parallel-by-default) / § 11.4.22 (maximum-idle-use + progress-update cadence) / § 11.4.23 (autonomous-decision-over-blocking) / § 11.4.24 (continuous parallel-stream routine) / § 11.4.25 (interactive clarification ONLY when a decision is genuinely operator-blocked per § 11.4.26) / § 11.4.27 (no-guessing – terminal conditions determined from captured queue/tracker state, never assumed) / § 11.4.28 (release-tag full-suite retest is the (A) terminal gate) / § 11.4.29 (iteration discipline is the loop body) / § 11.4.30 (audio-first priority within the loop) / § 11.4.31 (publish via merge-onto-latest-main, no force-push) / § 11.4.32 (multi-upstream publish) / § 11.4.33 (host-session safety overrides the loop). Propagation gate CM-COVENANT-114-126-PROPAGATION (literal 11.4.126 across the consumer fleet) + paired § 11.4.34 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md § 11.4.35 in the HelixConstitution submodule. All consuming projects restate + cite via § 11.4.36 inheritance.

Non-compliance is a release blocker regardless of context.
No escape hatch – no `--ask-before-continuing`, `--single-turn-only`, `--not-default-loop`, `--mimic-OK` flag exists.

§ . . – Session-handoff
resumption-prompt mandate (User mandate,
– –)

Forensic anchor – verbatim user mandate
(– –):

"make sure that in situations like this now when new session is needed you ALWAYS prepare such sentence – which will be valid for particular moment and the phase of the project and enough for work to continue."

When the agent determines a fresh session is needed (context-window limits, performance degradation) OR the operator asks whether a new session is needed / requests a handoff, the agent MUST ALWAYS prepare and provide a ready-to-paste **resumption prompt valid for that EXACT moment and project phase** – self-contained enough that pasting it into a fresh session resumes work with ZERO loss. The agent MUST proactively offer the resumption prompt the moment it recognises the need (not only when asked). Two variants on demand: a SHORT first-sentence variant ("Read `<handoff docs>`, then continue `<terminal goal>` ...") AND a FULL detailed block. The prompt MUST: () point to the live handoff doc(s) – `.remember/remember.md` if present + `docs/CONTINUATION.md` per § . – and instruct reading them FIRST + `git fetch --all`; () state the current PHASE + the immediate NEXT action + the terminal goal; () embed the exact current-state anchors (build IDs / artifact MD , device/target serials, commit HEAD, in-flight PIDs + log paths,

captured-evidence paths); () restate the binding constraints (anti-bluff § . covenant, no-force-push § . , exact version/naming, hardware/target gotchas); () be MOMENT-VALID – reflect the actual live state, NEVER a generic template. The handoff doc(s) MUST be brought current BEFORE the prompt is given (§ .). A handoff that omits the resumption prompt, or gives a stale / generic one, is a § . violation.

Classification: universal (§ .) – a platform-neutral session-continuity discipline reusable by ANY project worked by context-bounded agents; the consuming project supplies its concrete handoff-doc paths + state anchors per § . . Composes with § . (CONTINUATION doc – the resumption prompt is its operator-facing entry point) / § . (no-guessing – the embedded state is captured fact, not assumed) / § . (interactive clarification – "do you need a new session?" is itself a § . decision point) / § . / § . / § . (the fresh session CONTINUES the always-on autonomous loop – the resumption prompt is HOW the loop survives a context reset). Propagation gate CM-COVENANT-114-127-PROPAGATION (literal 11.4.127 across the consumer fleet) + recommended gate CM-HANDOFF-RESUMPTION-PROMPT-PRESENT + paired § . meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md § . in the HelixConstitution submodule. All consuming projects restate + cite via § . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-handoff-prompt , --generic-prompt-OK , --no-resumption-sentence , --handoff-without-state flag exists.

§ . . – Always-on device-recording
mandate (User mandate, – –)

Forensic anchor – direct user mandate

(– –): we MUST ALWAYS live-record all available data from all devices we use for testing or that we know manual testing is being performed on, done EXTRA carefully so the recording never harms the device, its performance, or causes side effects; recorded raw data MUST NOT be processed without need (token-conscious) and is ALWAYS excluded from version control + the code-intelligence index; only curated analysis/evidence is committed, and only during release preparation.

The mandate (ALL must hold). For EVERY test/debug device the project uses – and every device on which the operator is known to be performing manual testing – across EVERY reachable transport (USB or wireless ADB / SSH / serial / network introspection API), the project MUST ALWAYS live-record all available data sources it can later analyse: all activities, all logs (system log, kernel log, dropbox, crash/ANR traces), performance metrics (CPU / memory / I/O / thermal / load), every sink-side / downstream introspection report per § . . , and any other live-changeable parameter. () **Extra-careful, side-effect-free.** Recording MUST be designed so it NEVER harms the device, degrades its performance, or causes a side effect – non-invasive read-only probes only, bounded sampling intervals, bounded write-volume, an observer-effect budget (no probe that mutates the system-under-test). A recorder that perturbs the device-under-test is a § . . violation, NOT evidence. () **Background + parallel + subagent-driven.** Recording MUST run in the background in parallel with the main work stream (per § . . continuous-parallel-stream), dispatched subagent-driven per § . . – never blocking the main stream, never the only thing the conductor does. ()

Token-conscious – record-now, analyse-later. Raw recorded data MUST NOT be processed / analysed without need; the default is capture-and-store. The ONLY standing exception is during release-tag preparation (the § . . . full-suite retest / § . . . release-gate) OR when the operator explicitly asks – at which point the relevant slice is analysed and the curated findings become captured-evidence. Analysing raw recordings continuously, eagerly, for no tracked reason burns tokens and is a § . . . violation. () **Raw is git-ignored AND code-intelligence-excluded; only curated evidence is committed.** The raw recording corpus MUST be excluded from version control (with a § . . . regeneration-mechanism declaration – the mechanism is "re-record by re-running the recorder against the device", with the honest caveat that a past device-state is not byte-reproducible) AND excluded from the code-intelligence index per § . . . /§ . . . (a raw log/recording corpus is not source; indexing it pollutes the graph). Only the CURATED analysis / extracted evidence artefacts are committed – and only during release preparation per § . . . (docs/qa/<run-id>/) – never the raw corpus. () **Deterministic archive layout.** The raw corpus MUST be laid out by date → combined-source-state-hash → device → recording-sequence so any artefact is traceable to the exact code state and device it came from: <recording-root>/YYYY-MM-DD/<combined main+submodules state hash>/<DEVICE>_<SERIAL>/recording_NNN/<files> (monotonic NNN per device per state-hash). () **Anti-bluff.** A recorder claimed running but producing no growing corpus is a § . . . bluff; a curated finding cited as evidence MUST trace to a real raw-corpus path under the deterministic layout; the recorder's own health (alive + corpus growing) is itself a captured-evidence assertion per § . . . / §

Classification: universal (§ 11.4.128) – a platform-neutral test-evidence-capture discipline reusable by ANY project with reachable test/debug devices; the consuming project supplies its concrete device serials, transport probes, recorder helper, recording-root, and performance-metric set per § 11.4.128. Composes with § 11.4.128 (recorded-evidence – § 11.4.128 guarantees the recording exists to BE the evidence) / § 11.4.128 (captured-evidence quality) / § 11.4.128 (sink-side introspection, IS one recorded source) / § 11.4.128 (sink-side positive-evidence taxonomy) / § 11.4.128 + § 11.4.128 (release-prep is the analyse-the-corpus trigger) / § 11.4.128 (subagent-driven recorder dispatch) / § 11.4.128 (raw corpus git-ignored WITH a regeneration-mechanism declaration) / § 11.4.128 + § 11.4.128 (raw corpus excluded from the code-intelligence index) / § 11.4.128 (curated evidence lands under docs/qa/<run-id>/) / § 11.4.128 (background parallel stream) / § 11.4.128 (single-resource-owner – a read-only recorder never contends with an exclusive driver). Propagation gate CM-COVENANT-114-128-PROPAGATION (literal 11.4.128 across the consumer fleet) + recommended gate CM-DEVICE-RECORDING-ALWAYS-ON (recorder present + raw-corpus git-ignored + code-intelligence-excluded + deterministic layout) + paired § 11.4.128 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md § 11.4.128 in the HelixConstitution submodule. All consuming projects restate + cite via § 11.4.128 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-recording, --record-without-layout, --commit-raw-corpus, --index-raw-corpus, --analyse-corpus-always, --invasive-probe-OK flag exists.

§ . . – Huge-blocker release protocol (User mandate, – –)

Forensic anchor – direct user mandate

(– –): when a huge blocker is discovered during release validation we **MUST** stop all testing, return to fixing ALL discovered issues, process all recorded data from the last session, land rock-solid fixes, author NEW validation+verification tests of ALL supported test types, rebuild, reflash, and RESTART the full validation+verification of every fix/change from the last release tag to now – on both devices in parallel, recorded, with real physical captured proofs and no bluff.

When a HUGE BLOCKER (a release-blocking-severity defect – a core user-facing capability broken, a regression that invalidates the in-flight validation cycle, or a defect whose blast radius reaches the batch's other fixes) is discovered during release validation, the agent **MUST** execute the following protocol in order, with NO partial / spot-check shortcut: () **STOP all testing** immediately on every device (the § . . test-interrupt-on-discovery STOP, escalated to release granularity – a validation cycle that continues past a huge blocker produces "all green" summaries on a known-broken build, the exact § . PASS-bluff). () **Return to fixing ALL discovered issues** – not only the huge blocker but every defect surfaced this cycle; investigate each to root cause per § . . systematic-debugging, isolate regressions against the last known-good tag per § . . () **Process all recorded data from the last session** – analyse the § . . raw-corpus slice for the cycle (this is the § . . () release-prep analyse-trigger) so no defect the recordings already captured is missed. () **Land rock-solid fixes** – per § . . (rock-solid proof or deep research), § . . /§ . . (RED-on-broken-artifact

then GREEN), § . . . batch-source-fixes-before-rebuild. () **Author NEW validation+verification tests of ALL supported test types** – per § . . . (unit / integration / e e / full-automation / security / stress / chaos / performance / UI / UX / Challenges /...HelixQA) + § . . . stress+chaos, each anti-bluff with captured evidence + paired § . . . mutation. () **Rebuild + reflash** – full rebuild (NOT module-only) for retest-baseline integrity, deploy to a CLEAN target per § . . . (no stale-overlay shadow). () **RESTART the full validation+verification from the last release tag to now** – re-validate EVERY fix/change landed since the last tag (per § . . . full-suite-retest-before-tag – the huge-blocker discovery invalidated the in-flight cycle, so the cycle restarts from the beginning, NOT resumes), running on both/all owned devices IN PARALLEL per § . . . /§ . . . (single-resource-owner partitioning), every device's run RECORDED per § . . . , producing REAL physical captured proofs per § . . . /§ . . . /§ . . . with no bluff of any kind. This anchor BINDS the existing release-cycle anchors for the huge-blocker case – it adds the STOP→fix→all→process-recordings→new-tests-all-types→rebuild→reflash→full-restart sequencing and the "restart, never resume" rule, citing § . . . + § . . . + § . . . + § . . . / § . . . + § . . . + § . . . rather than duplicating their text.

Classification: universal (§ . . .) – a platform-neutral release-discipline reusable by ANY project that runs a release-validation cycle; the consuming project supplies its concrete device set, build/flash mechanism, and last-known-good tag per § Composes with § . . . (test-interrupt-on-discovery – § . . . is its release-granularity escalation) / § . . . (full-suite retest before tag –

§ 11.4.112 mandates the RESTART, not resume) /
 § (iteration discipline) / § (batch-source-fixes-before-rebuild) / § (all-test-types coverage) / § 11.4.120 (stress+chaos) /
 § (systematic-debugging) / § (clean-target rebuild+reflush+runtime-signature) /
 § + § (regression isolation + RED-on-broken-artifact) / § (rock-solid proof) / § (process the recorded corpus) /
 § + § (parallel recorded multi-device). Propagation gate CM-COVENANT-114-129-PROPAGATION (literal 11.4.129 across the consumer fleet) + recommended gate CM-HUGE-BLOCKER-FULL-RESTART (a release-blocker discovery during validation triggers a from-the-last-tag full restart, not a resume) + paired § 11.4.124 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md § in the HelixConstitution submodule. All consuming projects restate + cite via § inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --resume-after-blocker , --spot-validate-after-fix , --skip-recording-analysis , --skip-new-tests , --module-only-after-blocker , --single-device-restart flag exists.

§ 11.4.130 — Post-remediation validate-the-fix-FIRST-after-redeploy (User mandate, — —)

Forensic anchor — direct user mandate

(— —): when a blocker discovered during release validation is fixed and a new artifact is produced + redeployed, we MUST first re-test the SPECIFIC last-failing features and validate+verify the just-incorporated

fixes BEFORE proceeding to the broader / full validation – a first fix attempt may not work, may be incomplete, or may regress again, so confirming the targeted fix on the new artifact first prevents wasting a full multi-hour cycle on a still-broken fix.

When a blocker / critical failure is discovered during release validation and then FIXED, and a new artifact is produced (rebuild / new flashing image / redeploy) and the target is reflashed / redistributed / updated, the agent MUST execute the following ordering with NO shortcut: () **Re-test the SPECIFIC last-failing features FIRST** – run the targeted guard tests for exactly the defects this fix addressed, BEFORE any broader / full-suite validation. () **Validate + verify the just-incorporated fixes with real captured evidence** – the RED test that reproduced the defect on the pre-fix artifact (per § 4.1.1.1) MUST now flip GREEN at RED_MODE=0 on the new artifact, AND the fix's declared runtime signature (per § 4.1.1.2) MUST verify on the CLEAN target the redeploy produced. Metadata-only / config-only / absence-of-error / grep-without-runtime PASS are forbidden (§ 4.1.1.3 / § 4.1.1.4) – the proof is captured per § 4.1.1.5 / § 4.1.1.6 / § 4.1.1.7 () **Only after the targeted fix is CONFIRMED working on the new artifact** does the agent proceed to the § 4.1.1.8 full retest of every fix/change from the last release tag to now.

Rationale. A first fix attempt may not work, may be incomplete, or may regress again under the new artifact. Confirming the targeted fix FIRST catches a fix-did-not-take case immediately – instead of discovering it hours later at the END of a full multi-hour validation cycle (and then having to restart the whole cycle per § 4.1.1.9). The targeted re-test is cheap; the full cycle is expensive; ordering cheap-confirmation-first is

the § 11.4.6 iteration-speedup discipline applied to the post-blocker reflash. This is the § 11.4.6 recent-work-validation gate specialised for the post-blocker-reflash case: the § 11.4.6 huge-blocker protocol's step () full-restart MUST be preceded by this targeted-fix-first confirmation phase.

Honest_{..}boundary (§ 11.4.6): "the fix probably took" is not "the fix took" – the targeted RED→GREEN flip + runtime-signature verification on the new artifact is the captured proof; proceeding to the full cycle without it is a § 11.4.6 violation. A targeted re-test that still FAILs on the new artifact is a finding – return to the § 11.4.6 / § 11.4.17 isolate→RED→fix loop and produce another artifact; do NOT proceed to the full cycle on a still-broken fix.

Classification: universal (§ 11.4.6) – a platform-neutral release-discipline reusable by ANY project that fixes a blocker, redeploys, and re-validates; the consuming project supplies its concrete targeted-guard-test set, build/flash mechanism, and clean-deployment / runtime-signature observables per § 11.4.6. Composes with § 11.4.6 (test-interrupt-on-discovery – the blocker that triggered the fix) / § 11.4.6 (full-suite retest before tag – § 11.4.6 is the targeted-first phase that GATES the full retest) / § 11.4.6 (validate-recent-work-before-post-flash-tests – § 11.4.6 is its post-blocker-reflash specialisation) / § 11.4.6 (runtime-signature-as-definition-of-done on the clean target the redeploy produced) / § 11.4.6 (regression isolation – a still-failing targeted re-test re-enters the isolation loop) / § 11.4.6 (RED-baseline + polarity-switch – the RED test flips GREEN on the new artifact) / § 11.4.6 (rock-solid proof – captured evidence, never "probably took") / § 11.4.6 (huge-blocker

protocol – § . . . is its post-reflash targeted-confirmation phase that precedes the step- full-restart) / § . . . (iteration-speedup – cheap-confirmation-first). Propagation gate CM-COVENANT-114-130-PROPAGATION (literal 11.4.130 across the consumer fleet) + recommended gate CM-FIX-FIRST-AFTER-REDEPLOY (after a post-blocker redeploy, the targeted last-failing-feature guard tests run + pass on the new artifact before the full retest starts) + paired § . . . meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md § . . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-targeted-retest , --full-cycle-first , --assume-fix-took , --validate-fix-at-end , --skip-red-green-flip-on-new-artifact flag exists.

§ . . . – **Standing session-resumption file mandate** (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –): "Make this markdown a standard file which will be written EVERY TIME when we need fresh session out of the box! It MUST BE always up to date and in sync so whenever new session is created all we have to do is just point to it!"

Every project MUST maintain a SINGLE canonical, always-current **session-resumption file** at a fixed, project-declared standard path (the project declares the exact path once per § . . . and never moves it without a § . . . operator decision). This file is the OUT-OF-THE-BOX entry point for any fresh session: creating a

new session requires ONLY pointing the new agent at this one file. § 4.1.1 promotes § 4.1.2 (which mandated PREPARING a resumption prompt on demand) into a STANDING, version-controlled ARTIFACT – the resumption prompt is no longer an ephemeral in-message reply but a persistent file that is ALWAYS present and ALWAYS in sync.

(A) **Existence + fixed path.** The file MUST exist at the declared canonical path at all times; its location is encoded as a literal path string in the project-layer instantiation (§ 4.1.1) and never silently moved.

(B) **Always written + always synced.** The file MUST be (re)written whenever a fresh session is needed OR whenever the live project state materially changes (new HEAD, build/artifact id, phase transition, device/target state, in-flight job, blocking decision) – the § 4.1.1 CONTINUATION.md trigger set applied to the resumption file. A stale resumption file (describing a prior HEAD / phase / artifact) is a § 4.1.1 violation of the same severity class as a § 4.1.1 stale-CONTINUATION violation: it silently breaks the zero-loss-resumption guarantee.

(C) **Content (composes § 4.1.1).** The file MUST carry both the § 4.1.1 SHORT (one-paste first-sentence) AND FULL (paste-ready block) variants; point to the live handoff docs read FIRST (`.remember/remember.md` if present + `docs/CONTINUATION.md` per § 4.1.1) + `git fetch` ; embed the exact live-state anchors (HEAD commit, build/artifact ids + checksums, device/target serials, in-flight PIDs + log paths, captured-evidence paths); state the current PHASE + immediate NEXT action + terminal goal; and restate the binding constraints (anti-bluff § 4.1.1 , no-force-push § 4.1.1 , exact version/naming,

hardware/target gotchas). It MUST be MOMENT-VALID, never a generic template (§ 11.4.44).

(D) **Export + freshness.** The file is in § 11.4.44 scope → synchronized .html / .pdf siblings refreshed on every update; a § 11.4.44 revision header tracks its monotonic revision + last-modified timestamp.

(E) **Out-of-the-box resumption.** A fresh session, given ONLY this file's path, MUST be able to fully resume with zero additional context. The file IS the session bootstrap. § 11.4.127

Composes § 11.4.65 (CONTINUATION.md live-state – § 11.4.65 is the dedicated paste-ready-prompt sibling of it) / § 11.4.65 (resumption-prompt discipline promoted to a standing file) / § 11.4.65 (universal export) / § 11.4.65 (revision header) / § 11.4.65 (moment-valid, never generic) / § 11.4.65 (path change is an operator decision) / § 11.4.65 (default-loop continuity across sessions). Classification: universal (§ 11.4.65). Propagation gate CM-COVENANT-114-131-PROPAGATION (literal 11.4.131 across the consumer fleet) + recommended gate CM-SESSION-RESUMPTION-FILE-PRESENT (the declared file exists + carries SHORT+FULL + revision header + is fresh vs HEAD) + paired § 11.4.131 meta-test mutation (strip the file / make it stale → gate FAILS; gate-code = separate work item).

Canonical authority: this Constitution.md § 11.4.65 in the HelixConstitution submodule. All consuming projects restate + cite via § 11.4.65 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-resumption-file, --ephemeral-prompt-only, --stale-resumption-OK, --generic-template-OK flag exists.

§ 2126 06 07 – Risk-ordered validation priority mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –): "We MUST ALWAYS first test and validate features, functionalities and fixes/changes that have been worked most recently, the ones which were most problematic, which have the most chance to crash or break again, the ones which have been re-opened the most times! Then, after we validate and verify all this with real (physical) proofs and hard evidence, with no false results and bluffs of any kind, we continue with all other existing tests in the test suites! This IS MANDATORY."

Tests, validations, and verifications MUST run in **RISK-DESCENDING order** – the highest-risk set FIRST, and only AFTER that set is fully GREEN with real (physical) captured evidence does the remainder of the suite run. The risk ranking is computed from a closed set of factors, highest-risk first:

- (a) **Most-recently-worked** – features / functionalities / fixes / changes touched most recently (the freshest code is the least-exercised in the wild and the most likely to carry an undiscovered defect).
- (b) **Historically most-problematic** – the items with the longest defect history, the most prior fixes, the most prior failures.
- (c) **Highest crash/break/regress likelihood** – items whose blast radius, complexity, or dependency surface gives them the greatest chance to crash or break again.
- (d) **Most-reopened** – items with the highest reopen count per § 2126 06 07 (a high reopens-count is the strongest empirical signal of fragility – a defect that keeps coming back).

The highest-risk set, ordered by (a)-(d), is validated FIRST. Each item in that set MUST pass with real (physical) captured evidence per § . . . / § . . . / § . . . – no metadata-only / config-only / absence-of-error / grep-without-runtime PASS (§ . . . / § . . .), no false results, no bluff of any kind (§ . . .). ONLY AFTER the entire highest-risk set is GREEN with captured proof does the agent proceed to all other existing tests in the test suites. Running the full suite in arbitrary order, or running lower-risk tests before the highest-risk set is GREEN, is a § . . . violation: it spends the operator's scarce validation window on the items least likely to fail while deferring the items most likely to fail.

§ . . . REFINES / STRENGTHENS § . . . (post-remediation validate-the-fix-FIRST-after-redeploy – § . . . says re-test the just-fixed items first; § . . . generalises "first" to the full risk-ordered set) + § . . . (validate-recent-work-before-post-flash-tests – § . . . gates the full suite on recent-work GREEN; § . . . adds the explicit risk-ordering WITHIN the recent / high-risk set) + § . . . (iteration-discipline priority order – § . . . applies the same priority discipline to VALIDATION ordering, not just implementation ordering). The risk ranking is the validation-layer analogue of the § . . . / § . . . implementation-layer priority queue.

Classification: universal (§ . . .) – a platform-neutral validation-ordering discipline reusable by ANY project; the consuming project supplies its concrete recency / problematic-history / reopen-count sources (e.g. the § . . . workable-items DB `reopens_count` + `last_modified` columns) per § Composes § . . . / § . . . / § . . . /

§ . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . / § . . / § . . .
 Propagation gate CM-COVENANT-114-132-PROPAGATION (literal
 11.4.132) + recommended gate CM-RISK-ORDERED-VALIDATION-
 PRIORITY + paired § . meta-test mutation (gate-code =
 separate work item).

Canonical authority: this Constitution.md § . .
 in the HelixConstitution submodule. All consuming projects
 restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context.
 No escape hatch – no --skip-risk-ordering , --any-order-OK ,
 --suite-order-fixed flag exists.

§ . . – Target-System + hardware
 safety mandate (User mandate,
 – –)

Forensic anchor – verbatim user mandate

(– –): "Make sure that all changes we do to
 the System are ALWAYS safe for the System itself and for
 the hardware the system runs on! This is MANDATORY."

Every change to the TARGET system – firmware, kernel,
 init / boot scripts, drivers, sysfs / devfreq / voltage /
 clock / thermal / regulator register writes, partition
 layout / bootloader / U-Boot, HAL, framework, prebuilts,
 device configuration, or any equivalent on a non-Android
 target – MUST ALWAYS be safe for BOTH: (a) **the target
 System itself** – the change MUST NOT brick the device,
 induce a boot-loop, corrupt user / system data, or render
 the device unrecoverable; AND (b) **the hardware the System
 runs on** – the change MUST NOT drive any hardware-control
 surface beyond its safe electrical / thermal / voltage /

clock envelope, and MUST NOT damage panels, storage, radios, regulators, or any other physical component.

Concrete obligations (ALL hold):

- . **Reversible-first.** Prefer reversible changes; when a change is irreversible or high-blast-radius (bootloader / U-Boot / partition layout / one-way fuse / OTP), verify it against a known-good value (MD / checksum / golden manifest) BEFORE applying, and capture a pre-op backup of any state the change overwrites (per § .). A wrong irreversible write that bricks the device is the worst-case the mandate exists to prevent.
- . **No unverified hardware-control writes.** NEVER write an unverified value to a hardware-control sysfs node / register / firmware field that governs voltage, clock frequency, regulator output, thermal-throttle threshold, current limit, or any other physical-safety parameter. Every such value MUST be within the component's datasheet / vendor-documented safe range; the safe range MUST be established as FACT (cited datasheet / vendor reference per § . . + § . .), never guessed (§ . .). Pinning a max OPP / disabling a throttle / raising a regulator output without proving the cooling design and electrical envelope tolerate it is forbidden.
- . **Thermal / performance changes respect the cooling design.** Any change that raises sustained power, frequency, or voltage (e.g. forcing a performance governor, pinning the top OPP, disabling idle / thermal management) MUST be validated against the device's actual cooling capacity – captured thermal evidence (sustained-load temperature $\leq$ the documented limit) is the proof; "it boots" is not proof it is thermally safe.
- . **Flashing uses the sanctioned tool + verified image.** Deploying firmware MUST use the project's sanctioned

flashing tool against a freshly-built, integrity-verified image – never an ad-hoc partition write, never a stale / unverified artifact, never a tool / image whose provenance is unestablished.

- . **Unprovable-safety** ⇒ **blocked**. A change whose safety for the target System and its hardware cannot be established from captured evidence **MUST** be treated as UNSAFE and blocked – proceed only after the safe outcome is determined as FACT (§ 12.1.1), the change is reversible OR backed up (§ 12.1.2 reversible-first + § 12.1.3), and (when the safe choice still cannot be determined) the operator is asked via § 12.1.4 / § 12.1.5 's block-only-when rule. "Probably safe" is the exact guess this mandate forbids.

Distinct from § 12.1.1 host-session safety. § 12.1.1 (and § 12.1.2 / § 12.1.3 / § 12.1.4 / § 12.1.5) protects the DEVELOPER's HOST machine and session (no host suspend / logout / OOM cascade / unbounded-memory build). § 12.1.6 protects the TARGET device the System is built FOR and runs ON. Both apply; neither weakens the other. A build step can be host-safe (bounded memory, no host suspend) yet still ship a target-unsafe change (an over-voltage regulator write) – § 12.1.7 is the gate for the latter.

Classification: universal (§ 12.1.1) – a platform-neutral target-safety discipline reusable by ANY project that produces firmware / kernel / driver / device-config changes deployed onto target hardware; the consuming project supplies its concrete hardware-control surfaces, datasheet-safe ranges, known-good bootloader / image hashes, and sanctioned flashing tool per § 12.1.8. Composes § 12.1.1 (host-session safety – sibling-not-overlapping protection scope) / § 12.1.9 (no-guessing – the safe range / known-good value is FACT, never

assumed) / § . . (reversible-first + autonomous-decision-over-blocking – irreversible high-blast-radius hardware writes are the class that escalates to an operator question) / § . . (the change's runtime signature on a clean target proves it landed AND is safe) / § . . (rock-solid proof – safety claims need captured evidence, never a metadata-only assertion). Propagation gate `CM-COVENANT-114-133-PROPAGATION` (literal `11.4.133` across the consumer fleet) + recommended gate `CM-TARGET-HARDWARE-SAFETY` + paired § . meta-test mutation (gate-code = separate work item).

Canonical authority: this Constitution.md § . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no `--unsafe-hardware-write`, `--skip-system-safety`, `--brick-risk-accepted` flag exists.

2 0 2 6 0 6 0 8

§ . . – Code-review iterate-until-GO + rock-solid-evidence mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –): "For any fixes/changes given back to us for re-work by the code-review process, once we fix/improve everything per the code-review's requests, we MUST RE-RUN code-review AGAIN until we get a GO from it with NO new issues reported or warnings of any kind! All results produced by this whole process MUST ALWAYS give us rock-solid PHYSICAL evidence that the fixed/improved codebase really works now as expected, with no false results and no bluff(s) of any kind."

When the § . . code-review returns ANY finding – BLOCKING, nit, or warning – and the author fixes/

improves the batch per that review, the code review MUST BE RE-RUN, and MUST KEEP being re-run after each remediation round, until it returns a clean GO with ZERO new issues AND ZERO warnings of any kind. A single pass that "addressed the findings" is NOT sufficient: the corrected batch MUST pass a FRESH adversarial review (a re-review can surface NEW findings introduced by the very fixes that closed the prior ones – the § 11.4.5 fix-A-creates-B failure mode). The loop terminates ONLY on a clean GO (no new findings, no warnings); a residual warning is itself a finding that re-arms the loop.

Every round's verdict AND every fix's validation MUST carry rock-solid PHYSICAL captured evidence per § 11.4.6 / § 11.4.6.1 / § 11.4.6.2 (captured audio / video / sysfs / dumphsys / sink-side / runtime-signature) proving the fixed/improved codebase REALLY works as expected on the target – never metadata-only / configuration-only / absence-of-error / grep-without-runtime. No false results, no bluff of any kind, at any round. A reported GO unbacked by captured physical evidence is itself a § 11.4.6.3 PASS-bluff at the review-loop layer.

EXTENSION (11.4.6.4 – 11.4.6.5): this clause's re-run obligation ("MUST KEEP being re-run after each remediation round, until it returns a clean GO") is bounded PER APPROACH by § 11.4.6.6's declared round budget; at that ceiling the item is escalated per § 11.4.6.7 (E) (11.4.6.8) to a recorded operator decision, never closed, never continued round-by-round indefinitely under the same approach, and never silently relaxed to "good enough."

§ 11.4.6.9 REFINES / STRENGTHENS § 11.4.6.10 (which requires the review to iterate "until no blocking findings remain"): it makes the loop EXPLICIT (re-run after every remediation round, not once), raises the termination bar to ZERO findings AND ZERO warnings (not

merely zero-blocking), and BINDS rock-solid physical evidence to every round (not just the final one).

Classification: universal (§ . .). Composes

§ . . / § . . / § . . /

§ . . / § . . / § . . /

§ . . / § . . / § . . /

§ . . . Propagation gate CM-COVENANT-114-134-

PROPAGATION (literal 11.4.134) + recommended gate CM-CODE-

REVIEW-ITERATE-UNTIL-GO + paired § . meta-test mutation

(gate-code = separate work item).

Canonical authority: this Constitution.md § . .

in the HelixConstitution submodule. All consuming projects restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context.

No escape hatch — no --skip-rereview , --single-review-pass , --warnings-ok , --evidence-optional flag exists.

2 0 2 6 0 6 0 8

§ . . — Standing regression-guard suite + every-fixed-defect-gets-a-permanent-regression-test (User mandate, — —)

Forensic anchor (the canonical case this anchor exists to prevent, FACT): the wrong-subtitle-on- nd-display defect was "fixed" in a prior release via a source-side

CONTROL_MENU_LABEL_DENYLIST that NO test mirrored or re-ran, so the NEXT chrome class (a settings-menu label) recurred silently while the GREEN suite passed.

Every project MUST maintain a STANDING regression-guard suite that runs on EVERY build+deploy and BLOCKS the release tag on any failure. Every closed defect (the project's stable ticket id, e.g. ATM-NNN) MUST, in the SAME commit as its fix (extending the § . .

DOCUMENT step), register a permanent § . . RED-on-broken-artifact regression test into the suite —

RED_MODE=1 capturing the historical defect on a pre-fix artifact (the proof the guard is real), RED_MODE=0 the standing GREEN guard asserting the defect is ABSENT. A closure without a registered guard is a § . . violation (closure without permanent regression proof). The suite runs FIRST in the post-deploy cycle (highest-risk set per § . . risk-ordering) and is a § . . release-gate blocker. This is industry-standard bug-driven testing (cite: Google for Developers content-driven testing <https://developers.google.com/solutions/content-driven/backend/testing>; AOSP CTS/Tradedef continuous-regression model <https://source.android.com/docs/compatibility/cts>) made mechanical + enforced.

Classification: universal (§ . .). Composes
 § . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . . Recommended gates CM-REGRESSION-GUARD-REGISTERED + CM-REGRESSION-GUARD-SUITE-WIRED + CM-COVENANT-114-135-PROPAGATION (literal 11.4.135) + paired § . meta-test mutations (gate-code = separate work item).

Verdict-coverage at the release seam – ABSENCE of a verdict blocks exactly as a FAIL does (extension, research-derived, – –). Forensic FACT (genericised, – –): a consuming project's standing guard suite computed its exit code from ONLY its FAIL count – a guard that was never run (fix never built/deployed → REQUIRES-REBUILD SKIP → PENDING) affected nothing; measured live on a release-candidate target, registered guards produced PASS / FAIL / PENDING and the suite exited "not blocked", and the release-tag tooling consulted the guard suite ZERO

times – to the release gate, "never verified" and "verified good" were the same colour. That single fail-open explains, without needing hundreds of separate bugs, how % of that tracker's done-claiming items shipped unguarded and how the recorded reopen-per-recorded-fix rate reached %. Therefore (ALL hold): () **the suite's exit semantics MUST incorporate COVERAGE, never only the FAIL count** – a PENDING/never-run verdict is not a pass AND not ignorable at the seams that ship; () **the release seam blocks on verdict coverage of the release-candidate artifact** – the tag is refused unless, for the candidate's artifact fingerprint, every registered guard whose required topology is present on at least one release-validation target has produced a real PASS-or-FAIL verdict, and zero FAILs: `uncovered = registered ^ topology-present ^ no-verdict-for-this-fingerprint` ; refuse while `uncovered ≠ ∅ ∨ N_FAIL > 0` ; () **the release-tag tooling MUST consult the suite/verdict store mechanically** – a release path that never reads the guards is itself a § . bluff surface regardless of how green the suite runs elsewhere; () **the two legitimate PENDING classes are handled without a fail-open AND without a fail-closed wall**: on intermediate artifacts a REQUIRES-REBUILD PENDING is honest and non-blocking, but on the release candidate it is a contradiction – either the wrong artifact is deployed (re-run on the candidate) or the fix genuinely missed the build (then its item MUST NOT be in the release's done-set); both resolutions are forced, neither is a permanent exemption. Topology-absent guards are exempt from coverage ONLY via a checked-in target-pool topology map and MUST be enumerated in the release changelog as honest coverage gaps (§ . . / § . .) – never a silent exemption; a naive "PENDING always blocks" rule is REJECTED as a § . . FAIL-bluff (it would block every honest intermediate run and teach operators to disable the gate);

() **adoption ratchet** – a pre-existing unguarded done-claiming backlog is snapshotted ONCE into a committed ratchet file and may only shrink (the live unguarded set must be a subset of the snapshot AND the snapshot may not gain lines vs. its committed version), so day one is green, the disease cannot spread, and every reopened/re-touched item exits the ratchet permanently. Recommended gate **CM-RELEASE-VERDICT-COVERAGE** – self-tested per

\$. . / \$. . (): golden-bad
fixture = a topology-present guard with no verdict for the
candidate fingerprint → the release gate MUST FAIL;
golden-good = a complete verdict set → MUST PASS; negative
control = a properly-enumerated topology-absent guard →
MUST PASS (the false-positive guard) – + paired \$.
mutation (revert the exit semantics to FAIL-count-only, or
de-wire the release-tag consultation → the gate FAILS;
gate-code = separate work item). Composes \$. . /
\$. . / \$. . () /
\$. . (F) / \$. . /
\$. . (D) / \$. . /
\$. . .

Canonical authority: constitution submodule **Constitution.md**

\$. . . Non-compliance is a release blocker. No
escape hatch – no **--skip-regression-guard**, **--no-guard-on-**
close, **--guard-optional** flag – and, per the
– – extension, no **--pending-does-not-block-**
the-candidate, **--fail-count-only-exit**, **--release-without-**
consulting-guards, **--silent-topology-exemption** flag.

\$. . – **Real-content end-to-end**
playback-test mandate (User mandate,
– –)

Refines/strengthens \$. . . Any test asserting
media playback works MUST drive REAL content (catalog
stream or offline reference clip) through the user's path

(\$. . UI-driven → \$. . CV/OCR fallback) and assert it genuinely PLAYS via the \$. . liveness battery PLUS a decoder-health census – a numeric drop-buffer budget, no buffer-timestamp re-order/discard, no codec-reject (cite: Android/Media ExoPlayer OEM pre-OTA playback-test mandate https://developer.android.com/media/media_exoplayer/oems – "too many dropped buffers" > , "unexpected presentation timestamp", "test timed out" failure modes). Metadata-only / launch-only / registration-only / single-frame / config-only PASS is forbidden (\$. . / \$. .). A golden/reference clip corpus (cite: BBC ExoPlayer testing samples <https://github.com/bbc/exoplayer-testing-samples>; USPTO , , empirical media-player testing) is the offline ground-truth.

Classification: universal (\$. .). Composes \$. . / \$. . / \$. . / \$. . / \$. . / \$. . / \$. . Recommended gate CM-REAL-CONTENT-PLAYBACK-TEST + CM-COVENANT-114-136-PROPAGATION (literal 11.4.136) + paired \$. . meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md \$. . . Non-compliance is a release blocker. No escape hatch – no --launch-proves-playback , --skip-decoder-health , --metadata-playback-pass-suffices flag.

2 0 2 6 0 6 0 8

\$. . – Subtitle/caption content-correctness oracle + secure-display-proxy-honesty mandate (User mandate, – –)

Refines \$. . (pixel-oracle) + \$. . (liveness) + \$. . (structural-impossibility).

Forensic anchor (FACT): tests tasked to "physically verify the `nd-display` (or `st` where no TV) subtitle" PASSED GREEN while subtitles did NOT show / showed WRONG, because (a) the streaming app's player surface is `FLAG_SECURE` so `screencap -d <secondary>` returns BLACK (cite: Android Developers <https://developer.android.com/security/fraud-prevention/activities>; Nightwatch <https://www.nightwatchcybersecurity.com/> / <https://www.nightwatchcybersecurity.com/research-securing-android-applications-from-screen-capture/>; WithSecure <https://labs.withsecure.com/advisories/screencapture-via-ui-overlays-in-mediaprojection>) → autonomous PIXEL verification is structurally impossible (`$ . .`), so the test fell back to a non-pixel proxy (the accessibility-scraped forwarded cue / `persist.atmosphere.subdebug logcat`), and (b) the proxy's validation accepted a chrome/menu LABEL (`Аудио и субтитры`) as a valid subtitle because the prose floor accepted any multibyte prose and NO menu-label denylist + NO position/cadence check existed.

The mandate: a test asserting a subtitle/caption is correct MUST classify the *content class* of the rendered/forwarded cue – a present cue is NOT a correct cue. A cue is CHROME (FAIL) if it is a known control/menu label (a closed multilingual deny-list MIRRORED from the source-side denylist, case-folded both sides incl. non-ASCII scripts), is time-counter/numeric chrome, is NOT prose, is OUTSIDE the lower safe-title position band where caption position is available (cite CEA-`-anchor caption-window grid`: captions live in the lower-center safe band, chrome elsewhere [https://en.wikipedia.org/wiki/Metamorphic_testing](https://www.gte-media.com/cea-<code>-and-<code>cea-<code>-captions/</code></code></code>), OR is STATIC across the window (a real subtitle changes → dialogue cadence; a menu label does not → ≥ distinct prose cues required, a metamorphic relation per <a href=)). A cue is DIALOGUE (PASS) only when prose, not-denied, not-chrome, position-ok, AND cadence

11 4 122

≥ OR it fuzzy-matches the SOURCE-extracted expected cue via normalized edit distance (§ . . . host-side ground truth; cite SubER <https://github.com/apptek/SubER>). The content oracle MUST be self-validated golden-good/golden-bad (§ . . . ()) and the test-side denylist MUST be verified present in the SHIPPED artifact (§ . . .) – a source-green denylist with no test mirror + no artifact check is the exact recurrence pattern this anchor forbids.

Secure-display-proxy honesty (§ . . .): where FLAG_SECURE makes pixel verification structurally impossible, the rock-solid autonomous proof is the player's own caption telemetry (the subdebug/Cue ground-truth channel) + source-track presence + content-class oracle – NEVER a faked pixel "physical" pass; the human-eye pixel confirmation is operator_attended (§ . . .) with a tracked migration item. App-agnostic (keys off content class, not per-app rule) so it is the standing guard for the whole roster.

11 4 3 11 4 5 11 4 6 11 4 137 11 4 138 11 4 132 11 4 135 11 4 137 11 4 123

Classification: universal (§ . . .). Composes
§ . . . / § . . . / § . . . /
§ . . . / § . . . / § . . . /
§ . . . / § . . . / § . . . /
§ . . . / § Recommended gate CM-SUBTITLE-CONTENT-CORRECTNESS-ORACLE + CM-COVENANT-114-137-PROPAGATION (literal 11.4.137) + paired § . . . meta-test mutation (strip the denylist/position/cadence check → golden-bad Аудио и субтитры PASSes → gate FAILs).

Canonical authority: constitution submodule Constitution.md
§ Non-compliance is a release blocker. No escape hatch – no --present-cue-is-correct , --skip-chrome-oracle , --length-heuristic-suffices , --pixel-pass-on-secure-display , --skip-position-check , --skip-cadence-check flag.

\$. . – Operator-escape =>
mandatory bluff-audit + permanent guard (User
mandate, – –)

When the operator (or any out-of-band channel) finds a defect that the GREEN test suite passed, this is by definition a \$. . PASS-bluff – it MUST trigger, before the fix is closed: () a \$. . systematic-debugging pass to FACT-root-cause; () a bluff-audit identifying the EXACT assertion that should have caught it but didn't, cited to file:line (canonical example: lib/subtitle_content_validation.sh:sub_is_prose() returning TRUE for **Аудио и субтитры**); () a permanent \$. . regression guard registered in the SAME commit as the fix, with its \$. . RED capturing the operator-found defect; () the bluff-audit committed under docs/research/<scope>/<defect>_bluff_audit/. Closing an operator-found defect WITHOUT the bluff-audit + permanent guard is itself a \$. . violation (the bluff that let it through is still live and the defect will recur – the operator's "repeatedly reworking + reopening the same issues with no real progress" pain).

Classification: universal (\$. .). Composes
 \$. . / \$. . / \$. . /
 \$. . / \$. . / \$. . /
 \$. . / \$. . . Recommended gate CM-
 OPERATOR-ESCAPE-BLUFF-AUDIT + CM-COVENANT-114-138-PROPAGATION
 (literal 11.4.138) + paired \$. . meta-test mutation
 (gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
 \$. . . Non-compliance is a release blocker. No
 escape hatch – no --close-without-bluff-audit , --operator-
 find-is-just-a-bug , --skip-permanent-guard flag.

11 4 108

\$. . – Fresh-process clean-artifact runtime-signature mandate (User mandate, – –)

Refines \$. . . Before any post-deploy validation – ESPECIALLY a non-pixel proxy verification (the subdebug/accessibility-cue channel used for FLAG_SECURE displays) – the harness MUST assert running-artifact == built-artifact: the deploy yielded a CLEAN target (mutable-overlay/userdata wiped) OR a pre-validation check proves no stale overlay shadows the deployed code (e.g. every guarded package – incl. the Presenter that emits the subtitle cue – resolves to the system partition, no per-user override). A stale shadow of the cue-emitting component (e.g. a Presenter APK predating the denylist) makes the proxy report on code that was never deployed – any PASS is a \$. . PASS-bluff. Each fix declares ONE machine-checkable runtime signature verified on the clean target (the \$. . registry IS the definition of done); for the subtitle class the signature is "the shipped Presenter APK contains the denylist literal (case-insensitive) AND the subdebug channel emits candidate REJECTED reason=chrome-label for a menu label."

Classification: universal (\$. .). Composes \$. . / \$. . / \$. . / \$. . / \$. . . Recommended gate CM-CLEAN-ARTIFACT-RUNTIME-SIGNATURE + CM-COVENANT-114-139-PROPAGATION (literal 11.4.139) + paired \$. . meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md \$. . . Non-compliance is a release blocker. No escape hatch – no --validate-against-running-state , --skip-clean-precondition , --shadow-OK flag.

§ . . – Universal action-prefix system (ACTION_NAME ::) (User mandate, - -)

Forensic anchor – verbatim user mandate

(- -):

"When a user prompt STARTS with ACTION_NAME :: (e.g. BACKGROUND :: IMPORTANT: do X...) the agent MUST replace the ACTION_NAME :: prefix with that action's registered expansion text, then execute the rest. For BACKGROUND :: the expansion is: 'The following prompt that we will provide MUST BE executed in background in parallel with all main work streams using the subagents-driven development approach! All work done MUST PRODUCE rock solid evidence covered with hard physical proof(s) that all done is working as expected and as specified without any false results and without any bluff!'. The mechanism MUST be UNIVERSAL/ extensible – more actions will be added later, each with its own expansion + rules. It MUST work with EVERY CLI agent, be fully decoupled + reusable by every project that includes the constitution submodule, and load + execute out of the box."

Every project under this Constitution MUST support a universal, extensible **action-prefix system**: when a user prompt's FIRST non-blank line starts with an uppercase action token followed by " :: " (grammar ^([A-Z][A-Z0-9_]*) :: – anchored at line start, UPPERCASE-only token, exactly one space on each side of ::), the agent MUST () look the token up in the shared **action registry** (tracked data file constitution/actions/registry.yaml , or \$HELIX_ACTION_REGISTRY); () if the token is a registered action, REPLACE the ACTION_NAME :: prefix with that action's registered expansion text and apply its rules ;

() execute the REMAINDER of the prompt under the expanded instruction. The system is the C-preprocessor expand-then-rescan model applied to prompts (a line-anchored single-token macro): detect → substitute the registered expansion → execute the residual.

Two-layer architecture (both mandatory; LAYER is the universal floor). (LAYER – universal, always-on, out-of-the-box) the action-prefix recognition instruction MUST be mirrored into EVERY agent context carrier the constitution maintains (CLAUDE.md , AGENTS.md , QWEN.md , GEMINI.md per \$. .) so the agent itself recognises and applies the prefix on every CLI agent – Claude Code, Gemini CLI, Qwen Code, OpenAI Codex CLI, GitHub Copilot CLI, Cursor, Aider, Cline, Continue, Roo Code, and any future agent that reads one of those carriers. (LAYER – mechanical, where the agent exposes a pre-submit seam) a prompt-preprocessing hook reading the SAME registry applies the expansion deterministically (Claude Code UserPromptSubmit / UserPromptExpansion hook injecting the expansion via additionalContext ; per-agent slash-command equivalents generated from the registry for Gemini/Qwen/Codex).. LAYER is the \$. . . anti-forgetting upgrade applied to prompt prefixes – the expansion holds even if model recall lapses; LAYER guarantees the system works everywhere with zero extra setup. Honest \$. . boundary (\$. .): transparent mechanical free-form ^PREFIX :: interception is genuinely available ONLY on Claude Code today; on every other agent the free-form form is honoured by LAYER (the agent self-applies) and the slash-command equivalent is the mechanical convenience – this split is documented, never papered over.

Grammar (mandatory, all hold). () Anchored at the start of the first non-blank line ONLY; mid-prose tokens never match. () UPPERCASE-only token [A-Z][A-Z0-9_]* ; lowercase

never matches. () Separator is exactly " :: " (space-colon-colon-space) – avoids C++ Foo::Bar, YAML key: value, URLs.. () Stacked prefixes (A :: B :: rest) apply outer-to-inner, left-to-right (expand A, rescan, expand B, then the residual is the task). () A leading \ escapes the prefix (literal, no expansion) so action names can be discussed. () An unknown token that matches the grammar shape but is NOT registered is NEVER silently expanded or silently dropped – the agent asks which registered action was meant (\$ + \$) or treats it literally; it NEVER invents an expansion (\$). () Any prompt not satisfying the grammar is an ordinary prompt; the system is a no-op. The carrier-block GRAMMAR_ADDENDUM (- - , extended with the arrow form - -) additionally recognises FOUR EQUIVALENT forms of every action beyond the bare :: – PREFIX::ACTION_NAME :: <rest> (namespaced ::), / ACTION_NAME <rest> (bare slash), /PREFIX::ACTION_NAME <rest> (namespaced slash), and ACTION_NAME ---> <rest> (bare arrow, with the namespaced PREFIX::ACTION_NAME ---> <rest> variant) – where PREFIX is an action NAMESPACE (reserved default DEFAULT) and the namespace separator :: carries NO surrounding spaces (distinct from the action-body " :: " and the arrow-body " ---> ", each one ASCII space on either side); the slash, arrow, and namespaced forms are honoured with the same expansion/execution, the bare slash form yielding to a colliding built-in/host slash command while the namespaced slash form is always unambiguous. A leading \ escapes any of these forms (\BACKGROUND :: x , \BACKGROUND ---> x , \ /BACKGROUND x).

Registry is the single source of truth + the extension contract. Adding a new action = adding ONE actions[] row (name , version , summary ≥ words per \$. . . , expansion , optional rules + composes_with) – no code change in either layer (the registry is data, not code).

Both layers read the same file (\$. . -style single-source-of-truth). The first registry entry is **BACKGROUND** with the operator's verbatim expansion above; it composes \$. . . /\$. . . (subagent-driven), \$. . . /\$. . . (parallel streams), \$. . . (background execution), and \$. . . /\$. . . /\$. . . (captured physical evidence) + \$. . . (anti-bluff). The second registry entry is **REMINDER** – it re-surfaces previously-scheduled, **CRITICAL**, status-**UNCERTAIN** work: the agent **FIRST** verifies the **ACTUAL** current status from captured evidence (never assuming done or not-done – \$. . . no-guessing), **THEN** acts on the delta (report the captured proof if genuinely complete, resume from the exact point if partial, action it NOW if not started, or surface the block per \$. . . / \$. . .), **ALWAYS** producing a status verdict; it composes \$. . . / \$. . . / \$. . . / \$. . . / \$. . . / \$. . . / \$ Three further registry entries are the **severity / handling-priority markers** **CRITICAL** (highest-priority / potentially release-blocking – maximum urgency + rigor ahead of lower-priority work, track it, auto-activate \$. . . systematic-debugging when an issue is involved), **IMPORTANT** (high-priority, above routine work but below **CRITICAL**), and **NOTE** (capture the remainder as durable **CONTEXT** – request-history \$. . . / \$. . . + persistent memory when a non-obvious project fact – applied when relevant, not an urgent action unless it contains an explicit action, \$. . . record only what the note says); each supports all six \$. . . forms incl. the single-colon **CRITICAL:** / **IMPORTANT:** / **NOTE:** , and registering **NOTE** promotes **NOTE:** from an ordinary-prose NO-OP to a registered action. They tag **HOW** the remainder is handled (priority + rigor), not a workable-item Type

(`$ . .`). Decoupled + reusable (`$ . .`):
 the registry + hook + expander carry zero project-specific
 data and are consumed by reference (the `$ . .`
`codegraph_*` pattern); a consuming project ships its own
 registry or inherits the constitution default. Loads out-
 of-the-box via the `$ . .` install seam.

Classification: universal (`$ . .`). Composes

`$ . .` / `$ . .` / `$ . .` /
`$ . .` / `$ . .` / `$ . .` /
`$ . .` / `$ . .` / `$ . .` /
`$ . .` / `$ . .` / `$ . .` /
`$ . .` / `$ . .` / `$ . .` /
`$ . .` / `$ . .` / `$ . .` /
`$ . .` / `$ . .` . Propagation gate CM-
 COVENANT-114-140-PROPAGATION (literal 11.4.140 across the
 consumer fleet)

- recommended gate CM-ACTION-PREFIX-SYSTEM (registry exists
 + parses + the grammar regex matches/no-matches the
 canonical fixtures + the LAYER- mirror block is
 present in CLAUDE.md/AGENTS.md/QWEN.md/GEMINI.md + the
 BACKGROUND expansion text matches the registry verbatim
 + the Claude UserPromptSubmit hook + expander script
 exist and are executable) + paired `$ . .` meta-test
 mutation (corrupt the BACKGROUND expansion in the
 registry OR strip the `prefix_regex` OR remove the
 LAYER- mirror block from one carrier → the gate
 FAILS; strip the literal 11.4.140 → the propagation
 gate FAILS; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
`$ . .`. Non-compliance is a release blocker. No
 escape hatch – no `--skip-action-prefix`, `--ignore-prefix`, `--`
`no-registry`, `--invent-expansion-OK`, `--single-layer-only` flag.

2026 06 09

§ . . – Token-efficiency mandate (research-derived + operator mandate, – –)

Forensic anchor (operator mandate, – –):
the AI coding agent's token spend (input AND output) is
dominated by an enormous always-loaded governance context
– the consumer CLAUDE.md/AGENTS.md/QWEN.md plus the
constitution submodule (measured ~ 351 120 , tokens of
static governance re-sent on EVERY turn: consumer
CLAUDE.md ~ . K + constitution/CLAUDE.md ~ . K, a
–row Applied-Fixes table plus verbatim § . .X
anchor restatements) – compounded by – K-token
subagent transcripts, verbose status, full-file re-reads,
and full-file loads where retrieval would do. The same
pattern is externally documented as a scaling defect
(Claude Code issue – "Cache read tokens
consume . % of usage quota – architectural scaling
issue with CLAUDE.md re-reads"). Operator mandate: cut
token spend to – % of current (a – %
reduction) WITHOUT degrading quality, performance, or
safety, and WITHOUT breaking any existing mechanism.

Every project worked on by AI coding agents MUST adopt a
token-efficiency regime composed of the following
measures, ranked by safety and impact, and MUST PROVE the
reduction with a measured anti-bluff harness – never an
asserted number (§ . . / § . .). The
headline – % is the warm-cache target; the rule
requires the *measured* best-safe reduction with cited
evidence, never the estimate.

The measures (composable, each preserving every existing
rule):

- . Prompt-cache the static governance prefix (PRIMARY – the
single biggest lever). The always-loaded governance MUST
form a byte-stable prefix with a cache breakpoint at its

end, with nothing volatile rendered ahead of it. Cache reads cost $\sim \text{base input price} \times (\text{discount})$; writes $\sim (\text{min TTL} / \text{hour}) \times (\text{discount})$. The discount applies to the dominant cost component. **Caching is transparent – the model sees byte-identical input cached or not – so it CANNOT remove a rule, weaken a gate, lower quality, or change a verdict; it changes only billing.** The only failure mode is a *silent invalidator* (timestamp/UUID/unsorted-JSON ahead of the breakpoint), which costs money but never corrupts, and is detectable via `cache_read_input_tokens`. The operative discipline: **batch governance edits and keep CLAUDE.md + the constitution byte-stable mid-session** – every governance edit busts the prompt cache (BASELINE.md finding), so all governance changes in a working window are landed together rather than dribbled across turns. Claude Code caches natively; SDK callers use `cache_control:{type:"ephemeral"}`; Gemini/Qwen use their context cache via API-key auth.

- **Subagent model-tiering + output-to-file (SECONDARY – biggest non-cache win).** Route MECHANICAL, NON-JUDGMENT subagent work (search, grep, status, doc-export, file-presence, read-only probes) to a cheaper/smaller model (Haiku-class); RESERVE the strong model for ALL reasoning, fix-design ($\$ \dots$), verdicts, code-review ($\$ \dots$), and demotion-evidence ($\$ \dots$). Subagents persist large output to a file and return a pointer instead of a $\sim$ K-token inline transcript. **The strong model owns every decision and every PASS/FAIL – the cheap model never emits a verdict – so $\$ \dots$ determinism and the anti-bluff covenant are untouched; quality cannot degrade.** Bounded by $\$ \dots$ (% memory) and the $\$ \dots / \$ \dots$ stream caps.

- **Thin always-loaded INDEX + on-demand detail (progressive disclosure).** Restructure the consumer governance so the always-loaded file is a concise index (one line per fix / per anchor) and the full bodies are fetched on demand. Three invariants keep every gate green and every rule reachable: (a) each index line carries the literal `11.4.N` anchor token so the `CM-COVENANT-114-N-PROPAGATION` literal-anchor scans still pass; (b) the canonical full text stays in the tracked, gate-scanned `constitution/Constitution.md` – no rule is deleted, only de-duplicated out of the consumer; (c) the full body is reachable in one hop. This realises § 11.4.78's consumer-thin / submodule-canonical split at the byte level – a de-duplication, never a deletion.
- **Retrieval / CodeGraph-first over full-file loading.** Structural questions (where is X / what calls Y / what would break) → CodeGraph (§ 11.4.78 / § 11.4.79, already mandated + installed) or semantic search; region questions → grep-scoped read; never re-read a harness-tracked file. Read-only, already-trusted (§ 11.4.79) – no behaviour change.
- **Output-token reduction.** Terse conductor status (no restatement of unchanged governance); `..effort:"low"` on the mechanical-subagent allowlist only; structured output where a sink consumes it. Output is ~ × input price, so this is cost-leveraged. Presentation-only – captured-evidence artefacts (§ 11.4.79 / § 11.4.80) unchanged.
- **Tool-call efficiency.** Batch independent tool calls in one block; reuse harness-tracked file state; reuse persisted tool outputs.
- **Compaction / context-editing for long sessions (defensive).** Prune stale tool-results / old thinking so a + turn cycle doesn't re-pay for an ever-

growing transcript; never touches the governance prefix or on-disk evidence.

Mandatory measured proof (anti-bluff – § . . / § . . / § . .): the reduction MUST be proven by a token-accounting harness measuring tokens-per-development-cycle BEFORE vs AFTER on a frozen deterministic workload, split into `input_tokens` / `cache_read_input_tokens` / `cache_creation_input_tokens` / `output_tokens` from the authoritative Anthropic `usage` object (NEVER `tiktoken`; NEVER the client-side `total_cost_usd` estimate – recompute cost from raw token fields with the published prices), reproduced N times to identical verdict (§ . .). Pass = AFTER ≤ % of BEFORE (target) OR the measured best-safe reduction with a cited cold-cache reason. The AFTER run MUST ALSO show ZERO regression on the pre-build full sweep, the meta-test mutation sweep (every gate still FAILs its paired mutation – proves M did not turn any rule into a bluff), every propagation gate (literal anchors intact), a strong-model reasoning-path probe (proves M /M tiered nothing that decides), and a cache-warm proof (`cache_read_input_tokens > 0`). Cost reduction with quality/safety regression is a § . . FAIL, not a win.

No measure may break or degrade any existing mechanism – and the rule is structured so none can: caching is transparent (M); tiering and low-effort are confined to mechanical non-judgment work (M /M); indexing preserves the full rules by reference with literal anchors so the propagation gates pass and the canonical text stays gate-scanned (M); retrieval uses already-trusted read-only infrastructure (M). The § . . anti-bluff covenant family, § . . deterministic consistency, § . . code-review, § . . full-suite retest, and § . . memory ceiling are all unconditionally preserved.

Composes with § 11.4.123. . / § . . /
 § 11.4.123. . / § 1.1. . / § . . /
 § . . / § . . 11.4.17 / § . . /
 § . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . / § . . 11.4.33 / § . . / § . .

Classification: universal (§ . .) – the consuming project supplies its model IDs, its mechanical-subagent allowlist, its governance files, its structural-index tool, and its frozen workload per § . . .

Propagation gate CM-COVENANT-114-141-PROPAGATION (literal 11.4.141 across the consumer fleet) + recommended gate CM-TOKEN-EFFICIENCY (the measured BEFORE/AFTER harness exists + produced a verdict $\geq$ target-or-cited-floor + the non-regression sweep is green + cache-warm proof present) + paired § . meta-test mutation (inject a per-turn volatile token ahead of the governance cache breakpoint $\rightarrow$ cache_read_input_tokens collapses $\rightarrow$ measured reduction falls below the bar $\rightarrow$ gate FAILs; restore $\rightarrow$ PASSes – proves the harness measures real caching, not a hardcoded green).
 Gate-code = separate work item.

Canonical authority: constitution submodule Constitution.md
 § . . . Research + design + measurement: docs/research/token_efficiency/
 {RESEARCH,DESIGN,MEASUREMENT,TEST_PLAN}.md . Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-token-efficiency , --no-cache-governance , --assert-reduction-without-measuring , --tier-down-reasoning , --inline-all-governance , --tiktoken-estimate-OK flag exists.

§ 2126 06 09 . . . – Universal code-review
mandate – every change reviewed, always, no
exception (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"ALL changes we do MUST pass through the code review
step!!! ALWAYS!!!"

EVERY change made to ANY repository this Constitution governs – without exception – MUST pass through an INDEPENDENT code-review step BEFORE it is accepted, committed, or built. There is NO change class exempt: source code, fixes, tests, gates, meta-test mutations, documentation, doc-tooling, build/CI scripts, configuration, governance files (Constitution / CLAUDE / AGENTS / QWEN), conductor main-stream edits, sub-agent output, refactors, single-line edits – all of it. This is the ABSOLUTE form of § 2126 06 09 (which mandated a code-review agent gate after a batch of fixes, before the pre-build sweep + main build): § 2126 06 09 removes every implicit scoping that § 2126 06 09 's "after all fixes/changes/implementations are done" phrasing could be read to leave open – there is no "this is just a doc edit" exemption, no "this is just a one-liner" exemption, no "the author already self-reviewed (§ 2126 06 09)" substitution, no "trivial change" carve-out. If a diff exists, it gets an independent review.

Independence is load-bearing. The review MUST be performed by a reviewer structurally separated from the author (a dedicated code-review agent, subagent-driven by default per § 2126 06 09 / § 2126 06 09 , or a distinct human reviewer) – NOT the author re-reading their own work.

§ 2126 06 09 's multi-pass self-evaluation is the AUTHOR-

side discipline and remains required, but it does NOT satisfy § . . . : the whole point of an independent review seam is to catch the blind spots self-review cannot (the structurally-separated reviewer seam § . . . names). A self-review labelled as the code-review step is a § . . . violation of the same class as a § . . . PASS-bluff at the review layer.

The review is itself anti-bluff. A rubber-stamp "looks good" with no examined findings is not a review (§ . . . / § . . .) – the review MUST genuinely analyse the change for correctness, safety (no host § . . . / data § . . . / target-hardware § . . . regression), will-it-really-work (no solve-A-create-B per § . . .), end-user behaviour (§ . . . / § . . .), and test-genuineness (every covering test exercises the work and catches its negation per § . . .). Its conclusions are captured evidence per § . . . / §

The review iterates to a clean GO per § . . . – any finding (BLOCKING / nit / warning) re-arms the loop; the change is accepted only after a re-run review returns ZERO new findings AND ZERO warnings, with rock-solid physical evidence per §

Composition, not duplication. § . . . is the universal "always, every change, no exception" capstone over the code-review family: § . . . supplies the multi-layer review CONTENT (the six analysis dimensions + the build-gate marker CM-CODE-REVIEW-GATE-BEFORE-BUILD); § . . . supplies the iterate-until-GO + rock-solid-evidence LOOP; § . . . supplies the author-side self-review that PRECEDES (never replaces) the independent review; § . . . / § . . . supply the subagent-driven default execution model. § . . . binds them and strips the last exemption: the review step is not a batch-end nicety, it is a precondition on the ACCEPTANCE of any diff. Honest

boundary (§ 11.4.1): a passing review proves the change was independently examined and found sound to the reviewer's evidence – it does NOT replace the § 11.4.1 four-layer runtime-signature verification nor the § 11.4.2 full-suite retest^{2026 07 22}; it is one of MULTIPLE STRONG LAYERS, and the FIRST one every change crosses.

Forensic FACT (teaching citation, – – – review-before-push is not optional). A transport-hardening change was committed AND pushed to a project's `main` BEFORE any independent review (the author honestly flagged the omission); the post-hoc § 11.4.3 –substrate review then returned CONDITIONAL/NO-GO with ONE BLOCKING finding – the change had left a security-relevant protection default in the unsafe direction (an operator-keystroke-obscuring option shipped default-off with no disclosure) – plus two Important and four Minor findings, and the remediation consumed a full fix-forward cycle (RED-first test updates per § 11.4.4, the launcher fix, three documentation surfaces, and a § 11.4.5 re-review) that a pre-push review would have folded into the original change. This is the exact class this anchor forbids: the review found a real security-default defect that was ALREADY public on every mirror when it was found. The review comes BEFORE commit/push; a post-hoc review that catches a Blocking issue^{11 4 17} PROVES the mandate, it is not a substitute for it (remediation is fix-forward per § 11.4.6 – never a history rewrite).

Classification: universal (§ 11.4.7) – a platform-neutral acceptance-gate discipline reusable by ANY project; the consuming project supplies its concrete reviewer-dispatch mechanism + the change-acceptance seam (commit wrapper / merge queue / PR gate) per § 11.4.8. Composes § 11.4.9 / § 11.4.10 / § 11.4.11 / § 11.4.12 / § 11.4.13 /

11 4 92 11 4 108 11 4 118
11 4 125 11 4 134 137 1 1

```
$ . . / $ . . / $ . . /  
$ . . / $ . . / $ . . /  
$ . . / $ . . / $ . . /
```

Propagation gate `CM-COVENANT-114-142-PROPAGATION` (literal `11.4.142` across the consumer fleet) + recommended gate `CM-EVERY-CHANGE-REVIEWED` (every accepted change carries a fresh independent-review-completed marker for its diff, produced by a reviewer distinct from the author, before acceptance/commit/build) + paired `$ . meta-test` mutation (strip the literal → propagation gate FAILS; accept a diff with no independent-review marker → `CM-EVERY-CHANGE-REVIEWED` FAILS; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md`
`$ . . .` Non-compliance is a release blocker regardless of context. No escape hatch – no `--skip-review`, `--no-review`, `--trivial-change-exempt`, `--doc-edit-exempt`, `--self-review-suffices`, `--review-after-commit` flag exists.

2026 06 10

`$ . .` – Real-user-journey mandate for video-streaming-app full-automation tests (User mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"All video streaming apps ... require to choose some title and to press proper UI button to start or resume playing! Proper UI interaction to play exact show with proper content and subtitles is MANDATORY! Without it we just eventually play on and display sample, and that's it mostly! THIS MUST BE ADDED as MANDATORY RULE regarding testing of any video streaming app in general with full automation tests! Universal in root constitution + a consuming project's extensions."

Any full-automation test that asserts a video player / streaming application plays content MUST drive the REAL end-user journey through the app's OWN user interface – launch the app → BROWSE the actual catalog → choose a SPECIFIC title → press the real Play / Resume button → confirm THAT chosen content is genuinely playing, with its correct subtitles, on the intended routing target. A test that bypasses the journey with a sample / demo / built-in-loop clip, a deep-link / `am start -a VIEW` / intent shortcut, a synthetic or pre-staged stream, or any path that does NOT exercise the app's own browse-select-play UI is a § 11.4.48 PASS-bluff at the user-journey layer: it validates ROUTING (that *something* reaches the display) while leaving the user-visible behaviour the operator actually cares about – "the show I picked actually plays" – unproven. Such shortcuts exercise neither the catalog/auth/title-resolution path nor the real Play-button → decoder → routing transition where the end-user defects live (the operator's "we just eventually play on `nd` display sample, and that's it mostly" is exactly this gap).

11 4 48

The mandate (ALL must hold): () **Real journey, not a shortcut** – the test drives launch → catalog browse → specific-title selection → real Play/Resume press through the app's own UI (§ 11.4.48 UI-driven), NEVER a deep-link / intent / sample / loop-clip shortcut; () **Chosen-content confirmation** – the PASS proves the SPECIFIC selected title is playing (not merely that pixels move on the target), via the § 11.4.48 liveness battery (live advancing frames + frame-advance counter + not-stale-from-previous) on the § 11.4.48 real-content path, plus the § 11.4.48 subtitle content-correctness oracle for the chosen title's captions (a present cue is NOT a correct cue); () **Non-introspectable UIs use the pixel oracle** – when the app's accessibility/semantic hierarchy is blank or unreliable (TV-Compose /

leanback / canvas / GL streaming UIs), the test DRIVES input + ASSERTS content via the § . . CV/OCR pixel oracle (template-match the Play control + ROI-OCR the title/caption), never a hierarchy-only tool; () **Login via the credential single source** – apps requiring sign-in authenticate from the project's credential single-source-of-truth (§ . .), never hardcoded, never logged; () **Honest SKIP, never a faked PASS** – where the autonomous journey is genuinely infeasible (hard human-only login / CAPTCHA, geo-block per § . . , secure-surface pixel-blanking per § . .) the test is `operator_attended SKIP-with-reason` per § . . + § . . with a tracked migration item (§ . . / § . .) – NEVER a metadata-only / sample-played / routing-only PASS. Honest boundary (§ . .): "the routing fired so the title is playing" is a guess – only the chosen-content liveness + subtitle oracle on the real journey proves it; a SKIP for genuine infeasibility is correct, a PASS that never pressed the real Play button on the real title is the bluff this anchor forbids.

Classification: universal (§ . .) – a platform-neutral end-user-journey discipline reusable by ANY project that ships or integrates video-streaming applications; the consuming project supplies its concrete app roster, login-credential source, routing target, and UI-driving / pixel-oracle harness per § . . . Composes § . . (UI-driven traversal of the real journey) / § . . (liveness – the chosen content actually advances) / § . . (CV/OCR pixel oracle for non-introspectable streaming UIs) / § . . (real-content end-to-end playback) / § . . (subtitle content-correctness for the chosen title) / § . . (autonomous-first; operator-attended only when genuinely infeasible) / § . . (topology / geo SKIP-with-reason) / § . . (end-user usability is the

bar) / § . . Propagation gate CM-COVENANT-114-143-PROPAGATION (literal 11.4.143 across the consumer fleet) + recommended gate CM-VIDEO-REAL-JOURNEY-TEST (every video-streaming-app playback test drives the real browse-select-play journey + confirms the chosen content + subtitles, or SKIPS-with-reason) + paired § . meta-test mutation (replace a real-journey test's browse-select-play path with a deep-link / sample shortcut → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
§ . . . Non-compliance is a release blocker regardless of context. No escape hatch – no --sample-playback-ok , --skip-real-journey , --deep-link-suffices , --routing-only-pass , --no-title-selection flag exists.

§ . . – Tracked/recorded-device availability-following mandate (User mandate, – –)

Forensic anchor – direct user mandate

(– –): every device the project is tracking / following / recording MUST be availability-followed with the project's already-defined reconnection timings + the sanctioned device-recovery path; a silent recording/tracking gap – papering a disconnect as continuous capture – is a § . bluff.

Every device the project is recording or tracking (test / debug / manual-testing device, across every reachable transport – USB / wireless ADB / SSH / serial / network introspection API) MUST be availability-FOLLOWED: its connection state continuously monitored, and any drop handled – never silently abandoned and never presented as a continuous recording. The always-on recorder under § . . KNOWS a tracked device is absent (its

per-device loop guards on the reachable state) but, lacking a following discipline, merely spins idle – no data captured, no offline event logged, no resume, no escalation – so the recording corpus presents a continuous timeline with a silent hole. That silent hole is a \$PASS-bluff at the recording-integrity layer: it claims continuous capture while no data exists for the gap.

On a tracked device leaving its reachable state the system MUST, automatically: (a) **DETECT** the drop and **log an honest offline event** into the recording corpus (§ 3.2.1 – a silent gap presented as continuous capture is a fabricated-continuity bluff; § 3.2.2 – a silent recording gap defeats always-on recording); (b) **WAIT** for the device to return using the project's ALREADY-DEFINED reconnection timings – never invented numbers (§ 3.2.3), the SAME grace / reconnect / poll budgets the project's recovery path already uses; (c) **RE-ATTACH** – resume recording / tracking the moment the device returns and log an honest online / resume event; and (d) **ESCALATE** to the project's sanctioned device-recovery path (per § 3.2.4 feature class `device_recovery`) if the device does not return within the defined timeout – through the sanctioned, authorization-gated recovery entry point ONLY, never bypassing its gate and never performing a destructive recovery (e.g. a power-cycle) autonomously without that authorization (§ 3.2.5 / § 3.2.6 – high-blast-radius recovery is gated; while blocked the system keeps following the device, logging the blocked-escalation honestly). A tracked-device drop that produces a silent corpus hole, a never-resumed recording, or a never-escalated permanent absence is the bluff this anchor forbids.

Classification: universal (§ 1.1) – a platform-neutral device-following discipline reusable by ANY

project that records or tracks devices; the consuming project supplies its concrete tracking transport, the device set, the already-defined reconnection timings, and the sanctioned recovery entry point per §

Composes § (always-on recording – § closes its drop-handling gap) / § (device_recovery sink-side positive evidence) / § (honest offline / online events, reused-not-invented timings) / § (watchdog children reaped on stop) / § + § (gated, non-autonomous escalation).

Propagation gate CM-COVENANT-114-144-PROPAGATION (literal 11.4.144 across the consumer fleet) + recommended gate CM-DEVICE-AVAILABILITY-FOLLOWED (the recorder wires a per-device availability watchdog that reuses the defined timings + escalates only through the gated recovery path) + paired § meta-test mutation (strip the watchdog wiring / let an absence go unlogged → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md

§ Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-availability-following , --silent-recording-gap-OK , --no-reconnect-wait , --invent-reconnect-timing , --skip-recovery-escalation , --autonomous-power-cycle-OK flag exists.

2026 06 10

§ – Independent multi-angle impact-research per change (User mandate, – –)

Forensic anchor – verbatim operator intent

(– –): for EVERY fix / change / new feature, independent agents MUST research potential issues

related to that work AND connected/dependent features, from VARIOUS ANGLES, to be SURE no new issues are created, no existing working functionality is broken, and no dangerous code is introduced "like we had with the recents feature", incorporating a VARIETY OF TOOLS.

Forensic anchor (FACT, project-internal): the recents / -button-nav path shipped a latent AIDL-interface-contract mismatch that PASSED code review (the diff "looked correct") yet failed at RUNTIME ONLY – it ANR'd on a recents-reaping gesture – because no review angle interrogated the interface contract, the concurrency/lifecycle, or the connected feature it silently broke; the \$. . . SOURCE→ARTIFACT→RUNTIME→USER-VISIBLE gap, here caught only after the fact.

\$. . . shifts blast-radius discovery LEFT: for EVERY change, INDEPENDENT impact-research agents (subagent-driven \$. . . /\$. . . , structurally separate from the author – \$. . . self-eval PRECEDES, never satisfies – adversarial "refute-the-change" stance to defeat the documented LLM confirmation-bias failure mode where reviewers accept plausible-but-wrong arguments) research the change AND its connected/dependent features across a CLOSED SET OF EIGHT ANGLES – () **correctness/logic**, () **regression** (what existing working feature could break – via call-graph impact enumerating every direct + transitive caller and contract-dependent feature, each mapped to its test), () **latent/dangerous-code – the recents class** (runtime-only failure, interface/ABI/contract mismatch, concurrency/data-race, resource lifecycle), () **security** (taint/injection/secret/unsafe-API), () **performance** (latency/memory/thermal under load vs baseline), () **host/data/target-hardware safety** per \$ /\$ /\$. . . , () **cross-feature interaction** (shared state/timing/hardware contention), () **business-logic conformance**

(matches the spec AND the connected features' contracts, not merely compiles).

Each angle MUST name the TOOL(S) used + obtain the required DEPTH of codebase + business-logic detail (read the diff, the full changed unit, its declared contract, the spec section, and every connected feature – NEVER the diff lines alone) + emit a captured-evidence conclusion per \$. . /\$. . (the tool's findings/impact/diff artefact path – "looks fine" is forbidden, \$. .). A genuinely-N/A angle is recorded

NOT_APPLICABLE: <reason> per \$. . , never silently skipped. The output is a per-change impact-research REPORT (one section per angle: tool + evidence path + verdict) + a single GO/NO-GO that BLOCKS acceptance/commit/build on ANY unmitigated risk in ANY angle; the change is fixed/mitigated and the affected angles re-researched, iterating to a CLEAN GO (zero unmitigated risk + zero warning, rock-solid physical evidence) per \$. . before the \$. . review gate.

Honest boundary (\$. .): a GO proves cross-angle internal-consistency + bounded blast-radius – it does NOT replace \$. . four-layer runtime-signature verification nor the \$. . full-suite retest; it is one of MULTIPLE STRONG LAYERS, the FIRST research pass every change crosses.

Classification: universal (\$. .) – the consuming project supplies its concrete toolkit (static-analyzer, call-graph engine, interface-freeze mechanism, behavioral harness) + per-angle depth per \$. . . Composes/strengthens \$. . (adds independent tool-backed angle research) / \$. . (its impact-research input) / \$. . (shifts blast-radius discovery left) / \$. . (angles / / use the build-graph + interface-freeze) / \$. . (iterate-to-GO) / \$. . / \$. . (call-

graph impact engine) / \$. . / \$. .
(behavioral angle) / \$. . / \$ / \$
(angle) / \$. . (cite latest sources) /
\$. . / \$. . Propagation gate

CM-COVENANT-114-145-PROPAGATION (literal 11.4.145 across the
consumer fleet) + recommended gate CM-IMPACT-RESEARCH-PER-
CHANGE (every accepted change carries a fresh independent
impact-research report covering all eight angles with a
clean GO, produced by agents distinct from the author,
before acceptance/commit/build) + paired \$. meta-test
mutation (strip the literal → propagation gate FAILs;
accept a change with a report missing an angle or with an
unmitigated NO-GO → CM-IMPACT-RESEARCH-PER-CHANGE FAILs;
gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
\$. . . Non-compliance is a release blocker. No
escape hatch – no --skip-impact-research , --single-angle-
suffices , --author-researches-own-change , --diff-skim-OK , --
no-go-overridable , --trivial-change-no-research flag.

\$. . – Reproduce-first test +
same-test-confirms-fix + mandatory extend-to-
all-cases workflow (User mandate,
– –)

Forensic anchor – verbatim operator intent

(– –): every reported problem MUST be
handled by a NAMED three-step test workflow – () FIRST
author a test that **confirms + reproduces** the reported
problem AND is **used as an investigation instrument** to
obtain additional forensic data characterising the defect;
() after the fix, use the **SAME test** to confirm the
problem is GONE; () THEN **extend** the tests to cover ALL
flows, cases, edge cases, and the variety of potential

situations for that functionality, confirming no issue of any kind – with every step producing rock-solid PHYSICAL captured evidence, no false results, no bluff.

§ 11.4.146 does NOT re-author its component disciplines; it NAMES and BINDS them into the operator's exact per-defect sequence and adds ONLY the two emphases the components leave implicit. The mandate (ALL hold):

(STEP 11.4.42 – REPRODUCE-FIRST + INVESTIGATE). Before any fix, author the § 11.4.146 RED test as a § 11.4.115 RED-baseline-on-the-broken-artifact (reproduce the defect on the CURRENT pre-fix artifact, capture defect-present physical evidence per § 11.4.146 / § 11.4.115 / § 11.4.115). **NEW EMPHASIS (D 11.4.115):** that same RED test is ALSO a deliberate **investigation instrument** per § 11.4.115 (A) – it MUST be used to gather ADDITIONAL forensic data characterising the defect (its triggers, boundaries, input/topology scope, adjacent failure modes), and that characterisation MUST feed BOTH the § 11.4.115 root-cause/fix design AND the STEP-extend scope. A RED test used only as a binary present/absent gate, with no characterisation captured, satisfies § 11.4.115 but NOT § 11.4.115 's STEP 11.4.115.

(STEP 11.4.115 – SAME-TEST-CONFIRMS-FIX). After the fix, the SAME test source (the § 11.4.115 polarity switch flipped RED_MODE=1→0) confirms the defect is ABSENT – RED-on-broken then GREEN-on-fixed, both captured, on a clean target per § 11.4.115 / § 11.4.115, validated-first-after-redeploy per § 11.4.115, deterministic per § 11.4.115. No separate happy-path test substitutes for the polarity-flipped reproduction test (the § 11.4.115 / § 11.4.115 PASS-bluff).

(STEP 11.4.115 – EXTEND-TO-ALL-CASES, mandatory per-fix). **NEW EMPHASIS (D 11.4.115):** immediately after STEP 11.4.115 confirms the fix, the test MUST be **fanned out across the full case-**

space of the **SAME** functionality – all flows, valid + invalid + boundary (§ . . stress empty/max/off-by-one) + concurrent (§ . . contention) + failure-injection (§ . . chaos) + topology variants (§ . .) – confirming no issue of any kind, with PROVABLE enumerated coverage per § . . (a listed case-set with per-case outcome, never "no other issues found"). This is a REQUIRED step of the per-defect workflow, NOT deferred to a release-cycle discovery sweep and NOT reduced to a single guard. Each fanned-out case that exercises a user-visible path carries rock-solid physical evidence per § . . ; each is registered into the § . . standing regression-guard suite so the functionality's full case-space is permanently guarded; a newly-discovered defect in the fan-out triggers § . . test-interrupt + re-enters STEP for that new defect.

(**ANTI-BLUFF, all steps**). Every step's PASS is rock-solid CAPTURED physical evidence per § . .
 (§ . . /§ . . /§ . .) = metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§ . . /§ . .); when STEP or STEP validation method is unclear, deep-research-before-declaring-untestable per § . . / § . . /§ . . . An operator-found defect the workflow's green suite missed triggers § . . (bluff-audit + permanent guard).

Honest boundary (§ . .): STEP reduces the functionality's unknown-unknown surface; it does not prove zero remaining defects (no finite case-set can, per § . .) – the earned claim is "the reproduce→confirm pair + the enumerated extend case-set are all green with captured proof," un-exercised cases stated as honest gaps, never silently implied clean.

Classification: universal (§ 11.4.100) – a platform-neutral defect-lifecycle workflow reusable by ANY project; the consuming project supplies its reproduction/polarity harness, case-space enumeration, and capture mechanism per § 11.4.100. Composes § 11.4.101 (TDD RED step) / § 11.4.102 (RED-baseline + polarity – STEP 1 + core) / § 11.4.103 (systematic-debugging – STEP investigation) / § 11.4.104 (validate-fix-first – STEP 2) / § 11.4.105 + § 11.4.106 (clean-target runtime signature) / § 11.4.107 (deterministic) / § 11.4.108 (stress+chaos case classes – STEP 3) / § 11.4.109 (enumerated discovery coverage – STEP 4) / § 11.4.110 (permanent regression guard – STEP registration) / § 11.4.111 (topology variants) / § 11.4.112 (rock-solid-proof – all steps) / § 11.4.113 / § 11.4.114 / § 11.4.115 (physical evidence) / § 11.4.116 (operator-escape) / § 11.4.117 (test-interrupt) / § 11.4.118 / § 11.4.119. Propagation gate CM-COVENANT-114-146-PROPAGATION (literal 11.4.146 across the consumer fleet) + recommended gate CM-REPRODUCE-FIRST-THEN-EXTEND (every closed defect's fix carries: a § 11.4.120 reproduce-first polarity test with captured defect-characterisation [STEP 1], its RED_MODE=0 GREEN confirmation [STEP 2], AND an enumerated per-functionality extend case-set registered into the § 11.4.121 suite [STEP 3] – a closure with the reproduce→confirm pair but NO enumerated extend case-set FAILs) + paired § 11.4.122 meta-test mutation (strip the literal → propagation gate FAILs; close a defect with only the reproduce→confirm pair and no extend-case-set → CM-REPRODUCE-FIRST-THEN-EXTEND FAILs; gate-code = separate work item).

(D) STATUS-CUSTODY mechanical binding – a done-claiming status write is REFUSED without the workflow's machine-written evidence chain (extension, research-derived,

- -). Forensic FACT, (genericised,
 - -): in a consuming project this workflow
 and § . . already said exactly this – as prose
 – and did not bind: of tracked items,
 claimed a done/ready status and (%) **had NO
 registered guard at all**; an operator manually sampled
 done-claiming items and found / **broken** – with %
 unguarded, the *expected draw*, not bad luck; one item
 recorded Reopened events and terminal events
 (its status moved with no custody trail), and recurrences
 re-entered the tracker as NEW ids so even the %
 recorded reopen rate undercounts. A rule that lives only
 in prose is re-learned per session and skipped under
 pressure; a further prose restatement would be the
 disease's next instance – the deliverable is the BINDING,
 not more words. Therefore (ALL hold): () **the custody
 chain is a file chain a script checks, with prose nowhere
 in it –**

item status $\in$ done/ready $\Rightarrow$ a guard-registry row keyed by the
 item's EXACT stable id (legacy/alias keys resolved via a checked-
 in alias map – a guard keyed to a non-item id fails the join and
 satisfies nothing) $\Rightarrow$ an executable registered guard $\Rightarrow$ the
 §11.4.115(F) machine-written RED+GREEN verdict pair $\Rightarrow$ evidence
 files matching the item's feature-class shape ; () **Seam A –
 the status write itself is refused**: the workable-items
 engine (§ . .) refuses a transition to any done/
 ready status while the item's chain is incomplete; the
 engine change is project-agnostic per § . . – it
 reads a consumer-owned custody configuration, and every
 project supplies its registry / verdict-store / class-map
 paths as DATA per § . . ; () **Seam B – a full-
 table sweep gate, not a diff check**: on every pre-build
 run, EVERY done/ready item in the SSoT (minus the ratchet)
 must hold a complete chain, so an out-of-band raw-DB write
 that bypasses Seam A is caught at the next sweep rather
 than never; () **Seam C – the release seam** per the

§ . . verdict-coverage extension (ABSENCE of a verdict for the release-candidate artifact blocks exactly as a FAIL does); () § . . **construction order** – the FIRST artifact to land is the executable sweep + its golden-bad meta-fixture (a temporary SSoT overlay holding a terminal item with NO custody chain; the meta-test RUNS the sweep against it and requires FAIL – grep-only meta-assertions forbidden per § . . /
 § . . () – and ALSO asserts the sweep is WIRED: de-wiring its invocation from the pre-build runner flips the meta-test RED, because "installed" is a verified runtime state per § . . ()); governance text citing the gate lands only in the same commit as the executable, so documenting this clause without implementing the hook is itself caught – the acceptance check is one command: run the sweep; () **adoption ratchet** – per the § . . extension, the pre-existing done-claiming-unguarded backlog is snapshotted once, monotone decrease only. Honest boundary (§ . .): the binding proves that no status outruns its evidence and that the REPORTED defect cannot silently return under the guard's conditions – it does NOT prove the feature bug-free (family dedup per § . . remains), does NOT make a miscalibrated oracle honest (§ . . ()) goldens are necessary, not sufficient), does NOT replace the § . . manual-QA sufficiency gate, and makes target availability a HARD dependency of status progress BY DESIGN – a visible block replaces a silent pass. Recommended gate **CM-STATUS-CUSTODY** (Seam-A refusal present in the engine + Seam-B full-table sweep wired + ratchet monotone-decreasing) + paired § . . mutation (mark a chain-less item done via a raw DB write → the sweep FAILS; de-wire the sweep → its meta-test FAILS; gate-code = separate work item). Composes § . . /
 § . . / § . . / § . . /
 § . . / § . . / § . . /

§ . . / § . . (F) /
 § . . / § . . / § . . /
 § . . / § . . .

Canonical authority: constitution submodule Constitution.md
 § . . . Non-compliance is a release blocker. No
 escape hatch – no --skip-reproduce-first , --fix-without-red ,
 --skip-extend-to-all-cases , --reproduce-confirm-suffices , --
 defer-extend-to-release-sweep , --single-guard-suffices , --no-
 defect-characterisation flag – and, per (D), no --status-
 without-custody , --prose-closure-suffices , --skip-custody-
 sweep , --raw-db-status-write-OK flag.

§ . . – Crashed-agent respawn-
 until-complete + no-work-loss registry mandate
 (User mandate, – –)

Forensic anchor – verbatim operator intent

(– –): "any agent that crashed because of
 something MUST BE respawned and finish its work at some
 point! We MUST NOT lose any work, forget about or have it
 corrupted!" Forensic case study (FACT, this session):
 dispatched subagents died on **transient** causes – API Error:
 Server is temporarily limiting requests (rate-limit) killed
 at once, API Error: socket connection closed unexpectedly
 killed , and one left PARTIAL edits (two source files
 written, the dependent SystemUI sliders + doc + test NOT
 done). Each was re-dispatched by pure conductor vigilance
 with **no mechanical guarantee** – the moment conductor
 context is lost / compacted / the conductor itself
 crashes, an in-flight-but-dead work unit is silently
 dropped: work lost / forgotten / corrupted.

Every agent/subagent dispatched to perform a work unit
 MUST be tracked through its full lifecycle so that a crash
 NEVER loses, forgets, or corrupts its work. A crashed
 agent is **NOT** a completed agent (§ . . – "it

probably finished before dying" is a forbidden guess; an abnormal termination is positive evidence the work unit is OPEN). The mandate (ALL hold):

11 4 18

(a) Mandatory work-unit / agent REGISTRY. Every dispatched agent/subagent MUST have a durable, append-only, machine-readable registry entry carrying at minimum: a stable agent/work-unit id, its task + file-scope (`$ . . PWU manifest`), its declared output path(s) + captured-evidence dir (`$ . . /$ . .`), its dispatch timestamp, and a status from the CLOSED SET `{dispatched | in-flight | crashed | respawned | complete}`. The registry reuses the `$ . .` sync substrate (append-only JSONL event stream + atomically-rewritten status snapshot) – agent-lifecycle transitions are events; the `$ . .` integrity property holds ("an entry with no dispatch-event cannot show `complete`"). The registry is the SINGLE SOURCE OF TRUTH for "is this work owed?" – an unregistered dispatch is itself a `$ . .` violation (an untracked agent cannot be respawned).

(b) Mechanical CRASH-DETECTION + RESPAWN-until-complete. Any abnormal agent termination – transient API/rate-limit error (`Server is temporarily limiting requests`), socket-close (`socket connection closed unexpectedly`), any non-completion exit, OR a terminal error after the runtime's own retries are exhausted – MUST flip the registry entry to `crashed` and keep the work unit OPEN. An OPEN unit MUST be respawned (a fresh agent claims the same id + scope + preserved state) and re-respawned until an agent reaches `complete`. Respawn is the safe, reversible, bounded-blast-radius decision the conductor takes autonomously per `$ . .` (never blocks the loop waiting for a human). Transient causes use the project's ALREADY-DEFINED backoff/retry budgets – never invented numbers (`$ . .`); a genuinely-non-transient terminal error (deterministically reproducible after respawn) is

investigated per § 11.4.3 before further respawn, never spun in a blind retry loop.

(c) PARTIAL-STATE handling – preserve, then resume-or-clean-restart. A crashed agent's uncommitted edits + output doc MUST be PRESERVED (never silently discarded → work lost; never blindly committed → corruption). Before resuming, the respawn MUST run the § 11.4.3 quiescence check on the preserved tree (every modified file accounted-for against the unit's scope; no mutation/ // always pass / `_mutated_*` residue; no half-written/torn artifact). Then EITHER (i) **RESUME idempotently** from the preserved partial when it passes the check (consistent + complete-enough to continue), OR (ii) **CLEAN-RESTART** the unit from a known-good base (§ 11.4.3 pre-op backup, reversible per § 11.4.3) when the partial is inconsistent / incomplete-unsafely / fails the § 11.4.3 check. Either path guarantees nothing is lost and nothing is corrupted. Per-agent `git` worktree isolation (§ 11.4.3 L / § 11.4.3) keeps a crashed unit's partial tree from contaminating other streams.

(d) COMPLETION criterion. A work unit is DONE **only** when an agent reaches `complete` with its required captured evidence/output landed (§ 11.4.3 / § 11.4.3 + the § 11.4.3 verdict-carries-evidence-path rule). The endless-loop done-condition (§ 11.4.3 / § 11.4.3 / § 11.4.3) MUST NOT read as satisfied while any registry entry is `dispatched` / `in-flight` / `crashed` / `respawned-not-yet-complete`; the zero-idle survey (§ 11.4.3) MUST treat every non-`complete` entry as an OPEN actionable item to reclaim. A registry showing `complete` without the landed evidence is a § 11.4.3 PASS-bluff at the agent-lifecycle layer.

Honest boundary (§ 11.4.108, 11.4.147): the registry + respawn guarantee that work is *not lost or corrupted*; it does NOT itself prove the work is *correct* – the respawned unit's output still crosses § 147, four-layer verification, the § 147 / § 147 code-review gate, and § 147 full-suite retest. § 147 is the durability layer beneath those, not a substitute. § 147, 11.4.147 is the agent-side analogue of § 147 (device availability-following): same detect-drop → wait/backoff → re-attach/respawn → escalate-if-non-transient shape – devices for § 147, dispatched agents for § 147.

Classification: universal (§ 147) – the consuming project supplies its concrete registry path/format, its crash-detection signal source (runtime exit codes / sync-stream error events), its respawn-dispatch mechanism, and its already-defined backoff budgets per § 147. Composes § 128 (no-guessing – crash ≠ done) / § 147 (PWU = the work unit + worktree isolation) / § 147 (partial-state quiescence check before resume) / § 147 + § 147 + § 147 + § 147 (loop done-condition + zero-idle reclaim) / § 147 (autonomous respawn-vs-restart decision, reversible-with-backup) / § 147 (registry reuses the sync-channel substrate) / § 147 (non-transient terminal error → systematic-debugging before blind respawn) / § 147 + § 147 + § 147 + § 147 (the respawned output still crosses every downstream verification layer) / § 147 (clean-restart pre-op backup) / § 147 + § 147 (the recording/device-following siblings – same detect→wait→re-attach→escalate shape, agent-side here). Propagation gate CM-COVENANT-114-147-PROPAGATION (literal 11.4.147 across the consumer fleet) + recommended gate CM-CRASHED-AGENT-RESPAWN-TRACKED (every dispatched agent has a registry entry;

every crashed entry is respawned-until- complete ; the loop done-condition reads no non- complete entries; partial state preserved + \$. . -checked before resume) + paired \$. meta-test mutation (strip the literal → propagation gate FAILs; mark a crashed entry as the loop's done-condition complete without landed evidence, OR drop a dispatched agent without a registry entry → CM-CRASHED-AGENT-RESPAWN-TRACKED FAILs; gate-code = separate work item).

(e) API-QUOTA / RATE-LIMIT EXHAUSTION IS A FIRST-CLASS CRASH CLASS (extension, - -). Forensic FACT (- - + - -): background agents were terminated MID-WORK when the driving subscription hit its weekly / session quota – the agents did not fail, they were KILLED by an EXTERNAL quota boundary, and their in-flight work would have been silently lost had a conductor not manually re-dispatched them. Quota exhaustion MUST be handled by the SAME registry + respawn machinery as any other crash, with these additions: () the crash-detection layer MUST record the REAL limit signal + its reason-CLASS per \$. . (B) ({session | weekly | subscription} , derived from the captured signature – never faked, never guessed, \$. .); () the remedy is NOT "wait for the same alias" – the killed unit MUST be REBOUND to the next OPERATIONAL alias per \$. . (A)/(C) (native-first, provider fallback where configured) and RESPAWNED there; when EVERY alias is cooled the unit bounded-PARKS per \$. . carrying its operational-again epoch, and it stays crashed / OPEN – never complete ; () the respawned agent MUST RESUME from the killed agent's captured transcript + on-disk partial state at its EXACT last point (clause (c): preserve → \$. . quiescence check → resume-or-clean-restart), never restart blind, never silently re-do or drop the work; () the \$. . /

§ . . / § . . / § loop
done-condition MUST NOT read satisfied while any quota-
killed entry is non-complete – "the quota ran out" is a
PARK, never a DONE. No work may be lost, forgotten, or
corrupted by a quota boundary (composes §
alias-limit tracking + § always-on multi-
alias default – the rebind path is what makes a quota
crash recoverable rather than terminal).

Canonical authority: constitution submodule Constitution.md
§ Non-compliance is a release blocker. No
escape hatch – no --skip-agent-registry , --crash-equals-
done , --no-respawn , --discard-partial-state , --blind-commit-
partial , --forget-dead-agent , --loop-done-with-crashed-entries
flag.

§ – Workable-item integrity
(status+type+id) + comprehensive structured
description + bidirectional external-tracker
sync + BLOCKED unblock-choices mandate (User
mandate, – –)

Forensic anchor – direct operator mandate
(– –): every workable item MUST carry a
VALID status, a VALID type, and a stable unique identifier
at all times (never one without all three); every item
MUST carry a comprehensive, structured, self-explanatory
description (what it is, how it manifests, how to
reproduce, acceptance criteria); a BLOCKED item MUST state
WHY it is blocked AND the choices that would unblock it;
the single source of truth (the SQLite workable-items DB),
the rendered docs, and the external tracker (e.g. ClickUp)
MUST be regularly, never-missed, fully synchronized in
BOTH directions through the mechanical docs-sync engine;
the external-tracker push MUST carry statuses, types,

assignee, and sub-tasks, default an unset assignee from an env var, and be idempotent.

This anchor BINDS + STRENGTHENS the existing workable-item discipline (§ . . . status / § . . . type / § . . . ATM-NNN id / § . . . Operator-blocked / § . . . description-clarity / § . . . + § . . . SQLite SSoT / § . . . docs_chain) into ONE integrity contract spanning the three surfaces – DB docs external tracker – and adds the operator's three new emphases (D the no-item-without-all-three invariant, D the comprehensive structured description, D the BLOCKED unblock-CHOICES tightening), citing rather than re-authoring its components. The mandate (ALL hold):

(D) **No item without a valid status + valid type + stable id – across ALL three surfaces.** Every workable item, at every moment, MUST carry (i) a ****Status:**** from the § . . . / § . . . / § . . . closed set, (ii) a ****Type:**** from the § . . . closed set {Bug. | Feature | Task}, and (iii) a stable, unique, monotonic, append-only [**<ID-NNN>**] identifier per § An item missing ANY of the three in the DB, in the rendered docs, OR in the external tracker is a § . . . violation – the validator FAILs (non-zero exit, release-blocker). The id is the binding key that keeps the same item identifiable across DB rows, doc headings, and tracker tasks; an id present on one surface but absent on another is a sync-integrity failure.

(D) **Comprehensive structured description per item.** Every item MUST carry a comprehensive, self-contained, structured description – at minimum WHAT it is (the § . . . ≥ -word / ≥ -char clear meaning), HOW it manifests (observed user-visible / system behaviour), HOW to reproduce (deterministic steps where applicable, § . . . / § . . . reproduce-first), and

ACCEPTANCE CRITERIA (the captured-evidence verdict that closes it per § 11.4.122 / § 11.4.89). The description lives in the § 11.4.211 DB description column AND renders into the docs AND pushes to the external tracker – never a one-liner stub, never a § 11.4.211 anti-pattern fragment. A "fix-it-later"/placeholder description is a § 11.4.211 PASS-bluff at the planning layer (the operator / a future agent / a teammate cannot understand the item).

(D) BLOCKED items carry WHY + the unblock CHOICES.

Tightens § 11.4.211: every item whose status is Operator-blocked (incl. any Blocked / BLOCKED input alias normalised to the canonical Operator-blocked value – documented alias, never a silent fork, § 11.4.211) MUST carry, beyond the § 11.4.211 Operator-Block-Details WHAT/WHY/WHO, an enumerated set of UNBLOCK CHOICES – the closed list of decisions/actions that would unblock it (e.g. [A] grant access X · [B] approve removal of Y · [C] provide credential Z , mirroring the § 11.4.211 – -option interactive-clarification shape). The validator FAILs a BLOCKED item whose unblock field is empty OR carries no enumerated choice marker – "blocked, no way forward stated" is forbidden, the agent / operator must always know the choices that clear the block.

(D) Bidirectional DB docs external-tracker sync, never missed, docs_chain-bound. The § 11.4.211 / § 11.4.211 git-tracked SQLite DB is the SINGLE SOURCE OF TRUTH; the rendered docs (Issues / Issues_Summary / Fixed / Fixed_Summary / per-item docs + their § 11.4.211 / § 11.4.211 exports) and the external tracker are its DERIVED surfaces. A full synchronization pass MUST run regularly and on every state change so no surface drifts: DB→docs (db-to-md), docs→DB (md-to-db , both byte-identical round-trip per § 11.4.211), DB→external-tracker push, with a drift-

proof \$. . fingerprint (sha of the sorted item keyset, NOT mtime) gating freshness so a forgotten sync FAILs a pre-build gate. The sync is bound into the \$. . docs_chain engine as a registered consumer context (never an ad-hoc parallel script) so drift is mechanically caught, never vigilance-dependent.

(D .) Generic external-tracker sync – statuses/types/assignee/sub-tasks, default-assignee env var, idempotent.
 The external-tracker push MUST carry: each item's status (collapsed onto the tracker's native status set per \$. . when the tracker's status vocabulary is fixed – the API-cannot-create-statuses case is a \$. . structural constraint, the precise value preserved in a description header, never lost), its type, its assignee (a \$. . participant handle – `created_by` / `assigned_to`), and its sub-tasks. An UNSET assignee defaults from a project-supplied env var (the \$. . operator-username pattern), never hardcoded, never logged (\$. .). The push MUST be IDEMPOTENT – a dry-run-then-real, match-by-stable-key ([`<ID>`] name prefix / id custom-field) plan, present⇒UPDATE / absent⇒CREATE, rate-limited, credential-redacted, with a captured sink-side `created=N updated=M failed=0` proof per \$. . . The tracker-sync MACHINERY stays project-agnostic (\$. .) – the consumer registers its concrete tracker (service, list/board id, field map) at runtime via the public API; the constitution names the discipline, not the vendor.

Anti-bluff (\$. .): every sync run + every validator pass carries captured evidence per \$. . / \$. . (the round-trip diff, the fingerprint match, the tracker `created/updated/failed` line under `qa-results/`); a sync claimed-run with no growing evidence trail, a `complete` / `Fixed` item with no acceptance-criteria

description, or a BLOCKED item with no unblock choices are all PASS-bluffs at the workable-item-integrity layer.

Honest boundary (§ 11.4.17): the integrity contract guarantees every item is well-formed + the three surfaces agree, NOT that the item's underlying work is correct – that still crosses § 11.4.14 / § 11.4.15 / § 11.4.16.

Classification: universal (§ 11.4.17) – the consuming project supplies its concrete DB path, id prefix, external-tracker.service + list/board id + field map + default-assignee.env var, and docs_chain context per

§ 11.4.154 • 11.4.155 • Composes § 11.4.156 / § 11.4.157 /
 § 11.4.158 • 11.4.159 • / § 11.4.160 / § 11.4.161 /
 § 11.4.162 • 11.4.163 • / § 11.4.164 / § 11.4.165 /
 § 11.4.166 / § 11.4.167 / § 11.4.168 /
 § 11.4.169 / § 11.4.170 / § 11.4.171 /
 § 11.4.172 / § 11.4.173 / § 11.4.174 /
 § 11.4.175 / § 11.4.176 / § 11.4.177 /

CM-COVENANT-114-148-PROPAGATION (literal 11.4.148) + recommended gates CM-ITEM-INTEGRITY-STATUS-TYPE-ID (every item carries valid status+type+id on every surface) / CM-ITEM-COMPREHENSIVE-DESCRIPTION (every item's description has the WHAT/manifest/repro/acceptance structure) / CM-BLOCKED-UNBLOCK-CHOICES (every BLOCKED item enumerates unblock choices) / CM-TRACKER-SYNC-IDEMPOTENT (the external-tracker push is bidirectional-fingerprinted + idempotent + docs_chain-bound) + paired § 11.4.149 meta-test mutations (strip the literal → propagation gate FAILs; drop an item's type / id / description-structure / unblock-choices, OR break the tracker-sync idempotency key → the respective gate FAILs; gate-code = separate work item).

EXTENSION – SLICING: an item too big to test is an item whose "done" cannot be proven (11.4.150 – 11.4.151).

§ 11.4.152 governs an item's INTEGRITY once it exists (status + type + stable id + comprehensive

structured description) and § . . . (B) governs its acceptance criteria as Given/When/Then scenarios. Neither governs the item's SIZE AND SHAPE, and the AI-curriculum corpus (module) identifies that as the upstream cause of the reopen the other two anchors then have to catch. (a) **THE SLICING CHECKLIST.** Corpus: *"INVEST (Bill Wake) is a checklist for a well-formed item: Independent, Negotiable, Valuable, Estimable, Small, Testable."* Citation discipline (§ . . .): the corpus expands the acronym but defines ONLY two of its letters, and says so – *"For governance the two load-bearing letters are Small and Testable, and they are linked: an item too big to test end-to-end is an item whose 'done' you cannot prove, so it closes on partial evidence and reopens on the untested part."* This extension therefore mandates the two DEFINED letters and cites the acronym for the remaining four without inventing per-letter obligations the corpus does not state. (The corpus names "SMART Tasks" exactly once, inside a source title, and never expands or mandates it; no SMART obligation may be grounded on this corpus.) (b) **VERTICAL SLICES, NEVER HORIZONTAL.** The named anti-pattern, verbatim: *"the horizontal slice ('do all the database work,' 'do all the UI') which is never independently valuable or testable – you can't prove it works because it doesn't do anything a user can observe until every slice lands. Vertical slices (a thin end-to-end capability) are testable in isolation, so each closes cleanly."* The worked example: *"'reconcile donations across nine gateways' is not a workable item – it is untestable as one unit. Sliced by INVEST it becomes 'ingest and normalize gateway A's settlement file,' 'detect a provider-only ghost charge,' 'classify a fee/FX residual,' each with its own Given/When/Then and its own evidence."* (c) **WHY IT BINDS HARDER UNDER AGENTS.** Corpus: *"a small, testable item is one a subagent can be dispatched on and verified against a clear oracle; a big fuzzy item is one an agent will report*

'done' on with nothing to check it against." This is the § . . . / § . . . problem entering at the item-definition layer: an item with no clear oracle cannot have its completion verified by anyone, so the producer's self-report becomes the only available signal.

(d) READY-BEFORE-DISPATCH. The corpus names two practices that pin the criteria before code – **three amigos** (*"the business/product view ... the development view ... the testing/QA view. The testing voice is the one that surfaces the missing failure scenarios before they become reopens. The output is a shared, written understanding – not a hallway agreement"*) and **specification by example** (*"pin the requirement with concrete examples that become the acceptance tests: real inputs and expected outputs, including the awkward ones"*). **The mandate:** an item MUST be sliced small enough that its own acceptance criteria are provable end-to-end from its own captured evidence, and MUST be a vertical slice – a thin end-to-end capability a user can observe – never a horizontal layer. An item that cannot be closed on its own evidence is not ready for dispatch and MUST be re-sliced before work begins. Honest boundary (§ . . .): slicing makes "done" PROVABLE; it does not make the work correct (that remains § . . . / § . . . / § . . .), and "small" is a judgement about testability-in-isolation, NOT a story-point number, an hour count, or a line-count – the corpus states no size metric and this extension mandates none. Recommended mechanism gate **CM-ITEM-SLICED-SMALL-AND-TESTABLE** (every item carries acceptance criteria whose satisfaction is decidable from that item's own evidence, and is a vertical slice; a horizontal-layer item, or one whose criteria depend on a sibling item landing, → FAIL naming the item) + paired § . . . mutation (re-shape one item into a horizontal layer whose criteria cannot be met alone → the gate MUST FAIL; golden-FALSE per § . . . ()): a legitimately-dependent item that is

nonetheless independently testable once its dependency is met MUST NOT fire it). Gate-code = separate work item, NOT claimed shipped (\$. . / \$. .).

Canonical authority: constitution submodule Constitution.md
 \$. . . Non-compliance is a release blocker. No escape hatch – no --item-without-status , --item-without-type , --item-without-id , --stub-description-OK , --blocked-without-choices , --skip-tracker-sync , --one-way-sync-OK , --non-idempotent-tracker-push , --hardcode-assignee flag.

\$. . – Per-workable-item testing-diary mandate (User mandate,
 – –)

Forensic anchor – direct operator mandate

(– –): every workable item MUST carry a TESTING DIARY – a chronological, append-only record of every test run against it, each entry capturing date/time, who tested (tested-by), the result, in-depth observations (the technical background a future fix needs), and the action taken + why (whether the run changed the item's status, and the reasoning); both an in-depth diary AND an at-a-glance summary; modelled on the external tracker as a sub-task of the item with a {TODO | In-progress | Completed} lifecycle; persisted in the SQLite single source of truth + exported to all required formats + pushed to the external tracker; bound into the mechanical docs-sync engine so it is never missed; minimal-LLM (deterministic tooling, the prose authored by whoever ran the test); under % test-coverage discipline.

The testing diary is the evidentiary HISTORY layer for a workable item – distinct from the \$. . item_history (which records lifecycle STATE transitions) and from the \$. . Reopens.md (which aggregates reopen cycles): the diary records TEST EXECUTIONS, each of

which may or may not cause a state transition. The mandate (ALL hold):

(a) Append-only per-item diary table in the SQLite SSoT. A `test_diary` table (additive to the `$SSoT` / `$SSoT` schema, never a competing source) records one row per test run, soft-keyed to the item's stable id, with at minimum: ISO-8601 `date_time` (sortable, deterministic `$SSoT`); `tested_by` from a closed set (e.g. `User | Operator | AI-agent | HelixQA`); `result` from the `$SSoT` closed verdict vocabulary `{PASS | FAIL | SKIP}` + a short detail; in-depth `observations` (long markdown – facts or explicit `UNCONFIRMED:` / `PENDING_FORENSICS:` per `$SSoT`, never "likely"; this is the background a future fix relies on); `action_taken` + a `status_changed` flag + from/to (did this run change the item's status, and WHY/why-not); and a `$SSoT` `evidence_path` + `feature_class` – with a structural constraint that a `PASS` row WITHOUT a non-empty evidence path is impossible (a PASS-bluff is rejected by the schema itself).

(b) In-depth diary + at-a-glance summary. Per item, BOTH an in-depth diary document (one section per run, full observations + action + evidence link) AND a derived at-a-glance summary (a rollup VIEW – total/pass/fail/skip runs, last verdict, last run, status-change count, distinct testers, distinct feature-classes – DERIVED, never a duplicate source per `$SSoT`). The summary feeds the `$SSoT` risk-ordering inputs (fail-runs / status-changes / last-result) without re-reading every observations blob.

(c) Four-format per-item exports, docs_chain-bound, never missed. Each item's `Diary` + `Diary_Summary` render to the project's full export-format set (`$SSoT` + any project-added format) with the freshness invariant (derived-format mtime ≥ source mtime), wired through the

\$. . docs_chain engine as a registered consumer context with a \$. . drift-proof fingerprint (sha of the sorted diary keyset, NOT mtime) so a diary row added but not exported/pushed FAILs a pre-build gate – the never-missed mechanism.

(d) External-tracker sub-task model with lifecycle. On the external tracker, the workable-item task keeps its description = the item (\$. . D – NOT polluted with diary text); each diary run is a CHILD SUB-TASK of that task (type task), carrying a {TODO | In-progress | Completed} lifecycle status (collapsed onto the tracker's native status set per \$. . / \$. . when the tracker's vocabulary is fixed), its observations + action + an Evidence: line (the path, not the raw artefact – \$. . / \$. .), and a stable diary-entry idempotency key. The push reuses the \$. . D idempotent, rate-limited, credential-redacted, sink-side-proven discipline; the sub-task mapper stays project-agnostic (\$. .).

(e) Minimal-LLM, anti-bluff, % test-covered. The diary tooling (add / export / tracker-sync / validate) is deterministic bash/Go with ZERO LLM in the data path – the observations prose is authored by whoever ran the test; the tooling only stores/renders/pushes/validates it. The diary is covered per \$. . by every applicable test type (unit + integration-against-the-real-tracker with an honest \$. . SKIP-when-token-absent, never a faked PASS + export + HelixQA Challenge + paired \$. mutation) so a diary PASS-bluff is mechanically impossible. Honest boundary (\$. .): the diary guarantees a complete, auditable test-history per item, NOT that any single run's verdict is correct – that rests on the run's own captured evidence per \$. . / \$. . / \$. . .

Classification: universal (§ . .) – the consuming project supplies its concrete DB path, diary doc layout, export formats, external-tracker sub-task field map, and docs_chain context per § . . . Composes

§ . . . / § . . . / § . . . /
 § . . . / § . . . / § . . . /
 § . . . / § . . . / § . . . /
 § . . . / § . . . / § . . . /
 § . . . / § . . . / § . . . /
 § . . . / § . . . / § . . . /
 § . . . / § . . . / § . . . Propagation gate CM-COVENANT-114-149-PROPAGATION (literal 11.4.149) + recommended gates CM-TEST-DIARY-SYNC (the diary fingerprint gate – a diary row not exported/pushed FAILs) / CM-DIARY-PASS-REQUIRES-EVIDENCE (a PASS diary row carries a non-empty evidence path) + paired § . meta-test mutations (strip the literal → propagation gate FAILs; strip the evidence-required constraint OR the sub-task lifecycle / fingerprint → the respective gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md

§ . . . Non-compliance is a release blocker. No escape hatch – no --skip-testing-diary , --diary-without-evidence , --no-diary-summary , --diary-without-tracker-subtask , --llm-in-diary-data-path , --diary-export-optional flag.

2026 06 11

§ . . – Mandatory deep multi-angle web research per change/issue, before declaring fixed or structural (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"For every single issue we fix or improvement we make – besides tight systematic-debugging, fixing, code review by independent agents, comprehensive tests – we MUST ALWAYS do deep web research from various angles! No matter how big/small/simple/complex, dig deep the internet for articles, technical documentation, APIs, open-source code – EVERYTHING that can help make the best possible solution OR confirm we don't have a more serious problem we're unaware of! We MUST do everything possible to STOP the constant issue-reopening and finally start closing items as fixed+working! Do this ALWAYS in parallel with the main work stream! Add the mandatory rules into the constitution Submodule CONSTITUTION.md, CLAUDE.md, AGENTS.md, QWEN.md and other relevant files. Commit + push to all upstreams. Apply ASAP to every workable item."

11 4 112

Forensic case study (FACT, - -). A subtitle-on-secondary-display goal had been verdicted Won't-fix: structurally-impossible (the \$. . secure-surface pixel-blanking reasoning). A subsequent DEEP multi-angle web research pass OVERTURNED that verdict – it surfaced the real, workable OCR-of-the-PRIMARY-display path (read the caption on the introspectable primary surface rather than the FLAG_SECURE secondary), proving the goal was NOT structurally impossible, only mis-classified for want of research. A single research pass from a different angle converted a "we can't" into a shipped capability. This is the exact reopen-cycle the mandate exists to break: items churn between Fixed Reopened, or sit mis-labelled structural / won't-fix , because the solution-space (or the confirmation that the

closure is genuinely safe) was never deep-researched from enough angles.

For EVERY fix / improvement / change / closure – no matter how big, small, simple, or complex – the agent MUST, IN ADDITION to tight systematic-debugging (§ 4.1.1), the fix itself, independent-agent code review (§ 4.1.2 / § 4.1.3), and comprehensive multi-layer tests (§ 4.1.4 (b) / § 4.1.5), perform a DOCUMENTED **deep multi-angle web research pass** that digs the internet from VARIOUS angles – articles, official + vendor technical documentation, API references, standards, issue trackers, maintainer guidance, reusable open-source code – to BOTH (i) discover the best possible solution AND (ii) confirm there is NOT a more serious underlying problem the team is unaware of. This BINDS + STRENGTHENS the existing research/proof anchors into one mandatory, reopen-breaking, run-in-parallel discipline; it does not re-author them:

- (A) **No closure-as-fixed/structural WITHOUT a documented deep-research pass.** No issue may be marked **Fixed / Implemented / Completed** (§ 4.1.1.1) OR classified **structurally-impossible won't-fix** (§ 4.1.1.2) UNTIL a deep multi-angle research pass is documented for it. § 4.1.1.1 already requires deep research before designing a non-trivial fix and § 4.1.1.2 already requires it when validation is unclear; § 4.1.1.3 makes the pass UNCONDITIONAL – required for EVERY workable item, however trivial, at the closure AND structural-verdict gates – precisely because the forensic case proved a "trivial-looking" structural verdict was wrong.
- (B) **Multiple angles, not a single lookup.** The pass MUST interrogate the problem from several distinct angles (e.g. official-docs / standards / known-bug-trackers / alternative-approach / failure-mode / security /

performance / platform-constraint / open-source-precedent – the consuming project picks the applicable subset, ≥ genuinely-distinct angles), so a single-source confirmation-bias miss (the § . . LLM-confirmation-bias failure mode) cannot pass as research. A one-link drive-by is NOT a deep multi-angle pass.

- **(C) Confirm-no-bigger-problem, not just find-a-fix.** The pass MUST explicitly seek evidence that the chosen fix does NOT mask a deeper defect AND that the closure (or the structural verdict) is genuinely safe – the "do we actually have a more serious problem we're unaware of?" check. Absence of a found bigger problem is recorded as a positive finding only when the search that would have surfaced it is enumerated (composes § . . discovery-completeness – "we found nothing worse" requires "here is what we searched").
- **(D) Latest-source + cited.** Sources MUST be the LATEST authoritative versions per § . . (never training-data / memory), each cited by URL + access date in the item's research artefact AND the closure commit footer (Deep-research <date>: <urls> OR the literal NO external solution found – original work per § . . when the search genuinely yields nothing). Stale or memory-sourced "research" is a § . . + § . . violation.
- **(E) Reopen-breaking is the PURPOSE.** The explicit goal is to STOP the constant Fixed Reopened churn and finally close items as fixed-AND-working: a reopen (§ . . / § . .) whose root cause a deep multi-angle pass would have surfaced is a § . . miss; and a structural / won't-fix verdict reached without the pass is forbidden (the case study's overturned verdict is the canonical anti-pattern). Composes § . . (demotion-evidence – a

reopen IS a demotion, and the research is part of the demotion-prevention) + § 11.4.103 (structural-impossibility now REQUIRES the deep-research pass + its cited authorities before the verdict is earned).

- **(F) ALWAYS in parallel with the main stream.** The research pass MUST run as a background, subagent-driven work stream (per § 11.4.97 / § 11.4.100 / § 11.4.97 • 11.4.100 continuous-parallel-stream + § 11.4.97 background execution) concurrent with the main fix/build/test work – it NEVER serialises the main stream or stalls the loop (§ 11.4.97 / § 11.4.97 / § 11.4.97). A blocked research angle parks that angle, it does not pause the closure pipeline (§ 11.4.97).
- **(G) Apply ASAP to every workable item.** The discipline applies retroactively + going forward to EVERY workable item (§ 11.4.97 / § 11.4.97 SSoT); items closed without a documented pass are re-audited as part of the § 11.4.97 / § 11.4.97 release-gate sweep.

11 4 118

Honest boundary (§ 11.4.97): a deep-research pass reduces the unknown-unknown surface and breaks the most common reopen causes – it does NOT prove zero remaining defects (§ 11.4.97) and does NOT replace § 11.4.97 four-layer runtime-signature verification, § 11.4.97 / § 11.4.97 independent review, or § 11.4.97 full-suite retest. It is one of MULTIPLE STRONG LAYERS – the research layer that every fix/closure/structural-verdict additionally crosses. "We probably don't have a bigger problem" without the enumerated multi-angle search is a guess (§ 11.4.97), never a finding.

Classification: universal (§ 11.4.97) – a platform-neutral research-and-closure discipline reusable by ANY project; the consuming project supplies its concrete

research corpora, angle set, item tracker, and closure-commit-footer convention per § 11.4.80. Composes § 11.4.80 (deep-web-research-before-implementation – § 11.4.80 makes it unconditional + multi-angle + closure-gating) / § 11.4.80 (latest-source verification of the cited authorities) / § 11.4.80 (rock-solid-proof-or-deep-research – § 11.4.80 raises "when unsure" to "always") / § 11.4.80 (discovery-completeness – the confirm-no-bigger-problem search is enumerated) / § 11.4.80 (independent multi-angle impact-research – § 11.4.80 is its external-research sibling, adversarial against confirmation bias) / § 11.4.80 / § 11.4.80 (independent review – research precedes, never substitutes) / § 11.4.80 (demotion/reopen evidence) / § 11.4.80 (structural-impossibility now requires the cited-authorities pass) / § 11.4.80 / § 11.4.80 (reopen attribution – the churn this breaks) / § 11.4.80 / § 11.4.80 / § 11.4.80 (parallel background subagent stream) / § 11.4.80 / § 11.4.80 (release-gate re-audit) / § 11.4.80 / § 11.4.80 (per-item SSoT). Propagation gate CM-COVENANT-114-150-PROPAGATION (literal 11.4.150) + recommended gate CM-DEEP-RESEARCH-PER-ISSUE (every closed/structural-verdicted item carries a documented multi-angle deep-research artefact + cited-source closure footer, run in parallel, before the closure/structural verdict is accepted) + paired § 11.4.80 meta-test mutation (strip the literal → propagation gate FAILs; close an item or reach a structural verdict with no documented multi-angle research artefact / cited footer → CM-DEEP-RESEARCH-PER-ISSUE FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md § 11.4.80. Non-compliance is a release blocker. No escape hatch – no --skip-deep-research, --single-source-


```
suffices , --trivial-change-no-research , --close-without-  
research , --structural-without-research , --serialise-  
research , --research-from-memory-OK flag.
```

\$. . . – Project-prefixed release-tag/version-naming mandate (User mandate, – –)

Forensic anchor – verbatim user mandate (– –):

"Every release tag and version name we create – on the main repository and on every Submodule we own – MUST be prefixed with the project's release prefix, e.g. myproject-1.0.0-dev-0.0.1 . The prefix MUST come from HELIX_RELEASE_PREFIX in our .env if it is set, otherwise from the lowercased project root directory name. The SAME prefix MUST be used across the main repo and all owned Submodules in one release so a release is greppable across every repository."

Every release tag AND every version name created on the main repository AND on every owned-by-us submodule (per \$. . .) MUST be prefixed with the project's release prefix, in the form <PREFIX>-<version> (canonical example: myproject-1.0.0-dev-0.0.1). The prefix makes a release identifiable + greppable across every repository the release spans, so a single git tag -l '<PREFIX>-*' (or equivalent) enumerates the whole release surface in one query.

Prefix resolution order (closed-set, deterministic – \$. . . no-guessing):

- . HELIX_RELEASE_PREFIX from the project's .env – authoritative when set. .env is git-ignored per

\$. . ; the variable MUST be documented in the tracked `.env.example` placeholder so a fresh clone knows the contract (the regeneration-mechanism discipline of \$. . applies – `.env` is re-obtained, never committed).

- **Fallback** = the lowercased `snake_case` form of the project root directory name (no spaces), per \$. . – used whenever `HELIX_RELEASE_PREFIX` is unset/empty. The fallback is deterministic from the checkout, so a release prefix is always resolvable without operator input.

The prefix MUST be IDENTICAL across the main repo and all owned submodules within a single release – a release that tags the main repo `<PREFIX>-1.2.0` while tagging an owned submodule `<OTHER>-1.2.0` (or unprefixed) is a \$. . violation: the cross-repo grep no longer enumerates the release. Version codes increment monotonically within the prefix (`<PREFIX>-...-0.0.1` → `<PREFIX>-...-0.0.2` → ...), never reset, never skipped – same append-only discipline as \$. . 's ticket ids applied to the version axis.

11 4 48

Honest boundary (\$. .): the prefix guarantees a release is identifiable + uniform across every repository, NOT that the release contents are correct – the tag is still only created after the \$. . full-suite retest GREEN and reaches every upstream via the \$. . merge-onto-latest-main path (NEVER a force-push), fanned out per \$. .

11 4 35

Classification: universal (\$. .) – a platform-neutral release-naming discipline reusable by ANY project; the consuming project supplies its concrete prefix value + the `HELIX_RELEASE_PREFIX` env var per \$. . . Composes \$. . (multi-upstream push – the prefixed tag fans out to every remote) / \$. . (lowercase

snake_case – the fallback prefix follows it) /

\$. . (.env git-ignored – the authoritative prefix source) / \$. . (full-suite retest before the tag is created) / \$. . (absolute no-force-push – the prefixed tag reaches upstreams via merge-onto-latest-main) / \$. . (the release-scope terminal condition is a published, prefixed tag).

Propagation gate CM-COVENANT-114-151-PROPAGATION (literal 11.4.151 across the consumer fleet) + recommended gate CM-RELEASE-PREFIX-NAMING (every release tag/version on the main repo + every owned submodule carries the resolved <PREFIX>- prefix, identical across the release) + paired \$. meta-test mutation (strip the literal → propagation gate FAILs; create an unprefixd or differing-prefix release tag → CM-RELEASE-PREFIX-NAMING FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md

\$. . . Non-compliance is a release blocker. No escape hatch – no --no-release-prefix , --unprefixed-tag , --prefix-optional , --differing-submodule-prefix flag.

\$. . – Crashlytics-recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"For every project that has Firebase Crashlytics enabled / wired, we MUST continuously monitor ALL of the Crashlytics-recorded data – crashes, ANRs, performance traces, and non-fatals – systematically debug each, fix and improve, and cover everything with validation and verification tests. This MUST be

checked regularly, with no false results and no bluff of any kind!"

Every project that has Firebase Crashlytics enabled / wired (the SDK linked into a shipping artifact, crash + non-fatal + ANR reporting active) MUST treat the Crashlytics console as a first-class captured-evidence channel from real end-user devices and continuously process every datum it records. Crashlytics IS the user-visible-behaviour signal § 11.4.132 exists to protect: a recorded crash / ANR / performance regression / non-fatal is a real end user hitting a broken feature, and a recorded issue left unmonitored, undebugged, unfixed, or uncovered-by-test is the precise "tests pass while the feature is broken in the wild" failure mode the § 11.4.47 covenant forbids – an open Crashlytics issue with a green test suite is a § 11.4.47 PASS-bluff at the field-telemetry layer.

§ 11.4.47 STRENGTHENS + COMPLETES § 11.4.132 (Firebase data review mandate): § 11.4.47 owns the periodic REVIEW + dedup + Issue-creation pass that surfaces Crashlytics/Analytics/Performance findings and maps them to tracker entries; § 11.4.132 owns what happens to each surfaced item AFTER it is recorded – the systematic-debug → fix/improve → regression-test-coverage lifecycle that drives it to a proven, regression-immune closure. § 11.4.132 finds it; § 11.4.47 fixes it and proves it stays fixed. Both are mandatory; neither substitutes for the other.

The mandate (ALL must hold):

- **Continuous monitoring of ALL four Crashlytics surfaces.** The monitoring pass MUST cover fatal crashes, ANRs, performance traces / regressions, AND non-fatals – all four. Skipping any one surface is a § 11.4.47 PASS-

bluff (a regression surfaces on the skipped axis and the team never sees it – the § 3.1.3 three-source lesson extended to the non-fatal + ANR + performance axes). Non-fatals in particular are the silent class: every catch / fallback / recovered-error path that records a non-fatal is a real degraded user experience and MUST be triaged, not ignored because the app did not crash.

- **Systematic-debugging of each recorded issue (reproduce-before-fix).** Every Crashlytics-recorded issue selected for fix MUST go through the § 3.1.3 systematic-debugging arc – Iron Law: NO FIX WITHOUT ROOT-CAUSE INVESTIGATION FIRST – root-cause → pattern → falsifiable hypothesis proven with captured evidence → fix designed against the proven cause. The § 3.1.3 RED-baseline-on-the-broken-artifact reproduces the recorded defect. (the recorded stacktrace / ANR trace / non-fatal context IS the § 3.1.3 / § 3.1.3 captured evidence of defect-present) BEFORE the fix; a fix without a prior reproducing test is a § 3.1.3 / § 3.1.3 violation (no falsifiability evidence the fix addresses the root cause vs masks the symptom). Where the recorded issue cannot be reproduced, the deep-research-before-declaring-untestable path (§ 3.1.3 / § 3.1.3) is mandatory before any "cannot reproduce" classification.

- **A fix / improvement for every confirmed issue.** Every confirmed Crashlytics issue (severity per the § 3.1.3 classification table) MUST receive a real fix or improvement landed against its proven root cause (§ 3.1.3 / § 3.1.3 / § 3.1.3 four-layer verification on a clean artifact). "Acknowledged in the console" is not a fix; muting /

ignoring a recurring issue without a tracked rationale is a § 4.4.1 / § 4.4.2 violation.

- **Validation + verification regression-test coverage per closed issue.** Every Crashlytics issue closed by a fix MUST, in the SAME commit as the fix (§ 4.4.1 DOCUMENT step), register a permanent § 4.4.1 regression guard into the standing suite – a § 4.4.1 polarity test whose RED_MODE=1 captures the recorded defect on the pre-fix artifact and whose RED_MODE=0 is the standing GREEN guard asserting the defect is ABSENT. The guard MUST exercise the same user-reachable code path the Crashlytics stacktrace / trace identifies, with rock-solid captured evidence per § 4.4.1 / § 4.4.2 / § 4.4.3 / § 4.4.4. A Crashlytics issue marked resolved WITHOUT a falsifiable regression test is FORBIDDEN – it is the canonical recurrence vector (the issue silently returns months later, invisible to the green suite), the § 4.4.1 operator-escape class applied to field telemetry. Each closure carries the § 4.4.1 demotion-evidence + a closure-log entry recording the Crashlytics issue id / console URL, the root-cause analysis, the fix commit, and the validation + verification test paths.
- **Regular cadence – checked repeatedly, never once.** The monitoring + debug + fix + cover cycle MUST fire on a regular cadence, reusing the § 4.4.1 five-trigger set (pre-build / pre-flash / pre-distribute / pre-tag blocking; daily + post-deployment-burn-in non-blocking) so newly-recorded issues are caught while they are cheap. A one-time sweep that is never repeated is a § 4.4.1 violation; field telemetry accrues continuously and so must its processing.
- **No false results, no bluff.** Every verdict in the cycle is captured evidence per § 4.4.1 / § 4.4.2 /

§ . . – a "no new Crashlytics issues" claim requires the enumerated monitored-surfaces + queried-window as proof (§ . . absence-of-evidence is not evidence-of-absence); a "fixed" claim requires the RED→GREEN flip on the artifact (§ . . / § . .); a regression guard whose paired § . mutation does not make it FAIL is itself a bluff gate.

Honest boundary (§ . .): processing every recorded issue reduces the known-field-defect surface and breaks the silent-recurrence vector – it does NOT prove zero remaining field defects (a defect no user has hit yet is not yet recorded), and it does NOT replace § . . runtime-signature verification or § . . full-suite retest. It is the field-telemetry-driven layer that complements them; an issue is "closed" only when its regression guard is GREEN on a clean artifact, never when the console mark is flipped.

Classification: universal (§ . .) – a platform-neutral field-telemetry-processing discipline reusable by ANY project that wires Firebase Crashlytics (or, by the same shape, any equivalent crash/ANR/perf/non-fatal field-reporting backend); the consuming project supplies its concrete Crashlytics project handle, console-access credential (§ . . , never logged), severity table, and per-issue closure-log path per § . . . The reference project-level instantiation is a consuming project's § .0 (Crashlytics-Resolved Issue Coverage Mandate – per-issue validation test + Challenge test + `.lava-ci-evidence/crashlytics-resolved/<date>-<slug>.md` closure log) + § .AC (Comprehensive Non-Fatal Telemetry Mandate – every catch / fallback / recovered-error path records a non-fatal with triage context), which together are the project-specific embodiment of this universal clause. Composes § . / § . . / § . . /

cli_agents! ... in-depth with no a single feature left ...
table ... always up to date and in sync ... status of its
usability implementation tests coverage and the
confirmation of practical use confirmed on video recording
with no false results or bluff of any kind! ... tightly
connected using docs_chain Submodule ... Existence of such
Status (with Status_Summary) document with all exported
files (PDF, DOCX, HTML) is mandatory for every single
project."

Every project MUST maintain, under docs/features/ , a
comprehensive **feature Status document set** (Status.md + its
\$. . . Status_Summary.md companion) that
enumerates EVERY system component, EVERY client
application/binary/surface (TUI / CLI / Web / desktop /
mobile / API / gRPC / library / submodule /
infrastructure), and EVERY feature – including features
ported in from any incorporated CLI-agent / submodule
catalogue (\$. . .) – with NO single feature left
out. The mandate (ALL hold):

() **Total, categorized feature coverage** – a per-feature
table covering the FULL feature surface, organized
Component → Category, reconciled against the actual
codebase (a feature present in code but absent from the
table, or a table row with no code, is a \$. . .
violation; "looks complete" without a code-reconciliation
pass is a \$. . . / \$. . . bluff).

() **Per-feature fields** – each feature row MUST carry:
Component, Feature, Category, Implementation status,
Wiring status (genuinely reachable by an end user, not
merely compiled – \$. . .), Real-use status,
Tests-coverage status (the four-layer \$. . . (b)
coverage state), Validation status (closed vocabulary
PASS / FAIL / SKIP / PENDING-FORENSICS / OPERATOR-BLOCKED
per \$. . .), and **Video-recording confirmation** (a

path to the captured real-use video proving the feature works for the end user, or an honest gap marker).

() **Mandatory per-feature real-use video confirmation** – each user-visible feature's "confirmed working" claim MUST be backed by a recorded real-use video showing a genuine end-user scenario: real prompts → real LLM/service responses producing real results, NEVER a frozen/stale frame (§ . . liveness), NEVER a faked, mocked, demo-loop, or bluff response, NEVER an LLM error/garbage response passed off as success. Videos are captured evidence per § . . /§ . . and stored at the project-declared recording path (§ . .); a "confirmed" row with no real video, or a video whose content is a bluff, is a § . PASS-bluff at the feature-status layer. Where autonomous video capture is genuinely infeasible the row is an honest § . . / § . . SKIP-with-reason + tracked migration item – NEVER a faked confirmation.

() **Mandatory video-analysis remediation loop** – every recorded video MUST be analysed (its presented data recognized + checked); any defect the video surfaces MUST trigger the § . . systematic-debugging → fix → § . . retest → re-record loop, iterating to a clean GO per § . . , before the feature is marked confirmed. A video that exposes a broken feature is a § . . test-interrupt, not a confirmation.

() **Always-in-sync, mechanically enforced** – the feature Status set is a § . . -class roster/corpus-backed Status doc: kept in sync via the § . . docs_chain engine + a § . . drift-proof fingerprint (sha of the sorted feature-key roster AND the sorted video-artefact roster, NOT mtime) so that adding/removing/renaming a feature OR a video artefact

§ . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . / § . . / § . . Propagation
 gate CM-COVENANT-114-153-PROPAGATION (literal 11.4.153) +
 recommended gates CM-FEATURE-STATUS-COMPLETE (every code-
 present feature has a Status row + every row maps to code)
 + CM-FEATURE-STATUS-VIDEO-CONFIRMED (every user-visible
 confirmed row cites a real captured video or an honest
 SKIP) + paired § . meta-test mutation (strip the
 literal → propagation gate FAILs; mark a feature confirmed
 with no real video / drop a code-present feature from the
 table → the feature-status gates FAIL; gate-code =
 separate work item).

Canonical authority: constitution submodule Constitution.md

§ . . . Non-compliance is a release blocker. No
 escape hatch – no --skip-feature-status , --feature-without-
 video , --frozen-video-OK , --bluff-response-video-OK , --skip-
 video-analysis , --feature-ledger-incomplete-OK , --no-docx-
 export , --allow-feature-ledger-drift flag.

§ . . – Window-scoped capture +
 fresh-corpus rotation for feature/QA
 recordings (User mandate,
 – –)

Forensic anchor – verbatim user mandate

(– –): "IMPORTANT: Make sure that recording
 you perform does record the window containing apps and
 services, not the whole desktop or monitor screen!" +
 "NOTE: All old recording files MUST BE removed when new
 one starts!"

Refines/strengthens the § . . /§ . . /
 § . . /§ . . recording disciplines
 with two capture-hygiene invariants every feature/QA video
 MUST satisfy. (A) Window-scoped capture, NOT whole-screen.

A recording that proves a feature works MUST capture ONLY the window / surface / viewport of the application or service under test – the app's own window (GUI), its terminal pane (CLI/TUI), its browser tab/viewport (web), or its device/emulator/simulator frame (mobile) – NEVER the whole desktop, the whole monitor, or unrelated windows. Whole-desktop capture is forbidden because it (i) leaks unrelated/operator-private content into a committed artefact (a § 4.1.1.1 hygiene breach), (ii) dilutes the § 4.1.1.1 liveness/freeze/frame-advance oracle with non-feature pixels (a defect can hide in the noise – a § 4.1.1.1 evidence-integrity weakening), and (iii) makes the § 4.1.1.1 -class OCR/ROI content oracle unreliable. The capture mechanism MUST target the window/region by stable identity (window id/title, device serial, browser context, tmux/terminal target) per § 4.1.1.1 – never a fixed full-screen device index that happens to contain the window. Where the platform genuinely cannot capture below whole-screen for the surface under test, that is an honest § 4.1.1.1 SKIP-with-reason + tracked migration item, NEVER a whole-screen capture passed off as window-scoped. **(B) Fresh-corpus rotation – remove old recordings when a new run starts.** When a new recording run for a given scope begins, the prior run's stale recording files for that scope MUST be removed FIRST, so the live recording corpus always reflects the CURRENT artefact/run and a reader never mistakes a stale video for current evidence (the § 4.1.1.1 not-stale-from-previous principle applied to the corpus, and the § 4.1.1.1 roster-freshness principle applied to video artefacts). Honest boundary (§ 4.1.1.1 + § 4.1.1.1): "remove old" means the agent's OWN prior recordings for the SAME scope/project at the project-declared recording path – NEVER another project's, another scope's, or operator-authored files the agent did not create; when uncertain whether a file is in-scope-

stale vs foreign, do NOT delete it – surface it
 (\$. .). Curated evidence already committed
 under docs/qa/<run-id>/ (\$. .) is the durable
 record and is NOT deleted by rotation; rotation operates
 on the raw working recording path (\$. . . raw
 corpus), not the committed evidence trail. Classification:
 universal (\$. . .) – the consuming project supplies
 its concrete window-capture mechanism (per surface class),
 recording path, and scope-key per \$. . . Composes
 \$. . . / \$. . . / \$. . . /
 \$. . . / \$. . . / \$. . . /
 \$. . . / \$. . . / \$. . . /
 \$. . . / \$. . . / \$. . . /
 \$. . . Propagation gate

CM-COVENANT-114-154-PROPAGATION (literal 11.4.154) +
 recommended gate CM-WINDOW-SCOPED-FRESH-CORPUS-RECORDING
 (every feature/QA recording is window/surface-scoped not
 whole-screen + a new run removes its own prior in-scope
 recordings first) + paired \$. . meta-test mutation
 (strip the literal → propagation gate FAILs; a whole-
 screen capture where window-scoped is feasible, or a new
 run that leaves its own stale prior recordings, → the
 recording gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
 \$. . . Non-compliance is a release blocker. No
 escape hatch – no --whole-screen-capture-OK , --skip-window-
 scope , --keep-stale-recordings , --no-corpus-rotation ,
 --full-desktop-recording flag.

(C) MP auto-conversion REQUIRED. Any .cast file
 produced MUST be auto-converted to .mp4 via agg +
 ffmpeg immediately after capture. The .mp4 is the primary
 evidence; .cast is supplementary only. Conversion
 command: agg input.cast output.mp4 --renderer-res 1920x1080
 followed by ffmpeg -i output.mp4 -c:v libx264 -pix_fmt yuv420p
 -movflags +faststart final.mp4 .

**\$. . . – Project-name-prefixed
feature/QA recording filenames (User mandate,
– –)**

Forensic anchor – verbatim user mandate

(– –): "All recorded videos MUST START with
prefix: the PROJECT NAME (ALWAYS USE THE PROJECT NAME).
Project name MUST be obtained according to the
constitution's own project-name resolution."

Every recorded video the project produces – every feature/
QA real-use recording (\$. . .), every window-
scoped capture (\$. . .), every always-on device
recording (\$. . .), and every raw or curated
recording artefact under the project-declared recording
path (\$. . .) and the committed docs/qa/<run-id>/
evidence trail (\$. . .) – MUST have a filename that
STARTS WITH the PROJECT-NAME prefix, ALWAYS, with no
exception. A recording filename that omits the project-
name prefix is a \$. . . violation: a corpus
spanning multiple projects/scopes on one host (a real
condition under \$. . . always-on recording +
\$. . . parallel streams) becomes un-greppable
and un-attributable, and a reader cannot tell at a glance
which project a recording belongs to – the same identify-
and-grep failure \$. . . forbids on the release-
tag axis, applied to the recording-corpus axis.

**Prefix resolution order (closed-set, deterministic –
\$. . . no-guessing; IDENTICAL to \$. . . 's
prefix resolution):** the project-name prefix MUST be
resolved, never guessed, by the constitution's own
project-name resolution: () `HELIX_RELEASE_PREFIX` from the
project's `.env` – authoritative when set; `.env` is git-
ignored per \$. . . and the variable is documented
in the tracked `.env.example` (a \$. . . re-obtain

mechanism, never committed); () fallback = the lowercased snake_case form of the project root directory name (no spaces) per \$. . – used whenever the env var is unset/empty, so a prefix is ALWAYS resolvable from the checkout with zero operator input. The SAME resolved prefix is used for EVERY recording the project produces in a given checkout, so a single `ls '<PREFIX>---'*` (or `find . -name '<PREFIX>---*'`) enumerates the whole recording corpus for that project. Canonical filename form: `<PREFIX>---<feature-or-scope>---<run-id>.<ext>` (the --- triple-hyphen separator keeps the prefix unambiguously delimited from a feature/scope name that may itself contain hyphens). The prefix MUST be the SAME value \$. . resolves for release tags in the same checkout – a recording prefix and a release-tag prefix diverging in one checkout is itself a \$. . violation (one project, one resolved name).

11 4 107

Honest boundary (\$. .^{11 4 107}): the project-name prefix guarantees a recording is attributable + greppable to its project, NOT that the recording's CONTENT is valid – content validity still rests on the \$. . liveness battery, the \$. . content-correctness oracle, and the \$. . video-analysis remediation loop; the prefix is a naming/attribution discipline that composes with, never replaces, those evidence layers. The prefix also does NOT relax \$. .^{11 4 107}'s window-scoped-capture + fresh-corpus-rotation invariants (rotation removes the agent's OWN prior in-scope `<PREFIX>---*`^{11 4 17} recordings first; a foreign-prefix or operator-authored file is surfaced, never deleted, per \$. . + \$. .).

Classification: universal (\$. .) – a platform-neutral recording-attribution discipline reusable by ANY project that produces recordings; the consuming project supplies its concrete prefix value + the

`HELIX_RELEASE_PREFIX` env var + its recording path per
`$ . . .`. Composes `$ . . .` (the SAME
 prefix-resolution order + the same identify-and-grep
 purpose, applied to recordings instead of release tags) /
`$ . . .` (every always-on device recording
 filename carries the prefix) / `$ . . .` (the per-
 feature real-use video's confirmation path is prefixed) /
`$ . . .` (window-scoped + fresh-corpus rotation
 operate on prefixed filenames) / `$ . . .` (the
 prefix is a stable-identity name, not an enumeration
 index) / `$ . . .` (committed `docs/qa/<run-id>/`
 recording evidence carries the prefix) plus `$ . . .` /
`$ . . .` / `$ . . .` / `$ . . .` /
`$ . . .` / `$ . . .`. Propagation gate `CM-`
`COVENANT-114-155-PROPAGATION` (literal `11.4.155` across the
 consumer fleet) + recommended gate `CM-RECORDING-PROJECT-NAME-`
`PREFIX` (every recording filename at the project's
 recording path + every committed `docs/qa/<run-id>/`
 recording artefact starts with the resolved `<PREFIX>---`
 prefix, identical to the `$ . . .` -resolved prefix
 for the checkout) + paired `$ . . .` meta-test mutation
 (strip the literal → propagation gate FAILs; write a
 recording with no project-name prefix, or a prefix
 differing from the `$ . . .` -resolved value → `CM-`
`RECORDING-PROJECT-NAME-PREFIX` FAILs; gate-code = separate
 work item).

Canonical authority: constitution submodule `Constitution.md`
`$ . . .`. Non-compliance is a release blocker. No
 escape hatch – no `--no-recording-prefix`, `--recording-without-`
`project-name`, `--unprefixed-recording`, `--prefix-optional-for-`
`recording`, `--differing-recording-prefix` flag.

\$. . . – All CI/CD automation (GitHub Actions / GitLab pipelines / equivalents) MUST be disabled (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Any GitHub actions or GitLab pipelines MUST BE disabled! Add this critical mandatory rule / mandatory constraint into the root constitution Submodule, commit and fetch all its changes to all upstreams and make sure we respect and follow this rule (we do apply it) ASAP!!!"

Every repository this Constitution governs – the main repo, this constitution submodule, and every owned + nested submodule we author and push – MUST ship with ALL server-side CI/CD automation DISABLED. No push to any owned upstream may trigger a GitHub Actions run, a GitLab pipeline, or any equivalent provider-side automation (Jenkins, CircleCI, Travis, Drone, Woodpecker, Bitbucket Pipelines, Azure Pipelines, or any `on: push / schedule / workflow_dispatch` workflow). This GENERALISES + makes ABSOLUTE the \$. . . Layer- . . . posture ("Remote CI surfaces . . . are DISABLED – the workflow file is preserved at `...disabled-local-only` (NOT `.yaml` ; GitHub Actions ignores it)": what \$. . . did for the single constitution-compliance workflow, \$. . . mandates for ALL CI in ALL governed repositories. Enforcement migrates to the LOCAL \$. . . five-layer git-hook ritual + the \$. . . pre-tag sweep – never a remote runner.

The mandate (ALL hold): **(A) Zero active CI at the repository root.** No active `.github/workflows/*.yaml|*.yml` ,

no `.gitlab-ci.yml`, no `.gitlab/**` pipeline include, nor any equivalent provider config may exist at the ROOT of any governed repository/submodule – the only location a provider executes. **(B) Disabled means a push triggers ZERO runs.** Delete the config OR rename it to a non-active name (the `$ . . . .disabled` / `.disabled-local-only` convention – a provider ignores any name that is not its exact trigger filename); a workflow left with live `on:` triggers but `if: false` jobs still queues runs and is NOT compliant. **(C) Scope = repositories we author + push.**

Vendored / third-party source whose nested

`.github/workflows` or `.gitlab-ci.yml` sit BELOW the repo root (e.g. `AOSP external/**`, `prebuilts/**`, `packages/modules/**`, vendored submodules) are INERT – a provider never executes a non-root config – so they are OUT of scope and MUST NOT be mass-edited (`$ . . . vendor-name exemption + tree-integrity`); the test is "does a push to one of OUR upstreams trigger a run?" – yes ⇒ disable, structurally-inert ⇒ document + leave (`$ . . .` – verify inertness as FACT, never assume). **(D) No new CI may be added.**

Introducing any active workflow/pipeline into a governed repo is a release blocker. **(E) Pre-push verification.**

Before any push, verify no root-level active CI exists in the pushed repo/submodule (`git ls-files | grep -E '^\.github/workflows/.*\ya?ml$|^\.gitlab-ci\.yml$'` returns empty for authored repos); a `$ . . . -class PreToolUse` guard + the recommended gate enforce the class mechanically. Honest boundary (`$ . . .`): file-level disabling stops FILE-triggered runs; it does NOT disable provider-side server settings the agent cannot reach (org-default required workflows, branch-protection required checks, provider-side scheduled exports) – those MUST be turned off in provider settings by the operator, and the agent documents what it cannot reach rather than claiming a completeness it did not achieve.

Classification: universal (§ . . .) – a platform-neutral CI-governance discipline reusable by ANY project; the consuming project supplies its concrete repository/submodule set + provider list per § Composes § . . . (mechanical-enforcement Layer – § . . . promotes its CI-disabled posture to an absolute universal rule across all governed repos) / § . . . (vendor/third-party name + tree exemption) / § . . . (verify inertness, never guess) / § . . . / § . . . (the LOCAL pre-tag ritual replaces the remote runner) / § . . . (PreToolUse guard blocks the class at the tool-call boundary) / § . . . (same no-remote-surprise discipline applied to pushes) / § . . . (multi-upstream push – the disabled state holds on every mirror). Propagation gate CM-COVENANT-114-156-PROPAGATION (literal 11.4.156 across the consumer fleet) + recommended gate CM-NO-ACTIVE-CI (no root-level active CI config in any authored repo/submodule, verified pre-push) + paired § . . . meta-test mutation (strip the literal → propagation gate FAILs; add a root .github/workflows/x.yml to an authored repo → CM-NO-ACTIVE-CI FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §

Non-compliance is a release blocker regardless of context. No escape hatch – no --allow-ci , --enable-workflow , --keep-pipeline , --remote-ci-OK , --ci-exempt flag.

§ . . . – GEMINI.md maintained in lockstep with CLAUDE.md / AGENTS.md / QWEN.md (User mandate, – –)

Forensic anchor – verbatim user mandate (– –):

"Make sure with CLAUDE.md, AGENTS.md, QWEN.md we maintain GEMINI.md too! Add this mandatory fact / rule to root constitution Submodule we are inheriting / extending – CONSTITUTION.md, CLAUDE.md, QWEN.md, AGENTS.md, GEMINI.md and other related relevant files!"

Forensic case study (FACT, – –). When \$. . . + \$. . . were authored, Constitution.md / CLAUDE.md / AGENTS.md / QWEN.md carried the rule family through \$. . . , but GEMINI.md had silently drifted to \$. . . – fourteen mandates (\$. . . – \$. . .) never propagated to it. A per-agent context carrier that lags is a \$. . . propagation-bluff: a Gemini-CLI agent reading the stale GEMINI.md operates under an out-of-date Constitution while the fleet believes the rule is universally in force.

GEMINI.md is a FIRST-CLASS governance context carrier, EQUAL to CLAUDE.md / AGENTS.md / QWEN.md – NEVER an optional or best-effort sibling. Every governance addition or edit MUST land in ALL FIVE carriers in lockstep: the canonical Constitution.md PLUS each per-agent mirror CLAUDE.md + AGENTS.md + QWEN.md + GEMINI.md (and any other relevant per-agent context file a runtime introduces), in the SAME change-window, each with its synchronized .html / .pdf / .docx exports per \$ The mandate (ALL hold): **(A) Five-carrier lockstep.** No governance change is complete until GEMINI.md carries it alongside the other three mirrors – GEMINI.md is added to the \$. . . constitution-update pipeline's stage- propagation set + stage- cross-reference validation explicitly. **(B) No silent drift.** A GEMINI.md whose highest rule number lags the other mirrors

is a § . . . violation (the same severity class as a § . . . stale-export drift); the gap MUST be back-filled. (C) **Equal status, equal rules.** GEMINI.md restates the SAME literal 11.4.N anchors the § . . .x propagation gates require (per § . . . inheritance), so the propagation-gate fleet count INCLUDES GEMINI.md . (D) **Consumer projects too.** The lockstep binds the .consuming project's repository-root context files (its own .CLAUDE.md / AGENTS.md / QWEN.md / GEMINI.md) per § . . . – a project that maintains the other three but not GEMINI.md is non-compliant. Honest boundary (§ . . .): the pre-existing § . . . – § . . . GEMINI.md .back-fill is a tracked remediation work unit; this rule MAKES that drift a release-blocking debt rather than silently tolerating it – claiming GEMINI.md is "in sync" while the back-fill is incomplete is itself a § . . . violation.

Classification: universal (§ . . .) – a platform-neutral context-carrier-parity discipline reusable by ANY project that ships per-agent context files. Composes § . . . (constitution-update workflow – GEMINI.md added to the propagation + validation set) / § . . . (canonical-root inheritance – the per-agent mirrors restate the canonical anchors) / § . . . (every new rule lands in all carriers) / § . . . / § . . . (revision header + synchronized exports for GEMINI.md) / § . . . (GEMINI.md is one of the "every CLI agent" context carriers) / § . . . (this rule's own first lockstep test) / § . . .

Propagation gate CM-COVENANT-114-157-PROPAGATION (literal 11.4.157 across the consumer fleet, GEMINI.md INCLUDED) + recommended gate CM-GEMINI-MD-LOCKSTEP (GEMINI.md's highest § . . .N is not less than the other three mirrors' highest) + paired § . . . meta-test mutation (strip the literal → propagation gate FAILs; let GEMINI.md's highest

rule number fall behind the other mirrors → CM-GEMINI-MD-LOCKSTEP FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md

\$. . .

Non-compliance is a release blocker regardless of context.

No escape hatch — no --skip-gemini-md , --gemini-optional , --gemini-lag-OK , --four-carrier-suffices flag.

\$. . . — Intensive all-feature/
flow/edge-case video-recording + read-the-
screen content-verification mandate (User
mandate, — —)

Forensic anchor — verbatim user mandate

(— —): "Add the following mandatory rules / guidelines / mandatory constraints ... regarding the recording of all features, flows, cases, edge cases and full validation and verification of it all with real results, covered with evidence and rock solid physical proofs and no bluffs of any kind! All projects MUST BE covered with such intensive testing and video recordings produced as the results! Default save path of them all is ALWAYS to be Downloads directory under host machine's user's home directory! All video recordings MUST contain valid working features with no false results showed in it and all printed out logs, messages, UI labels, dialogs, toasts MUST BE really read by the testing System (HelixQA MUST BE used!!!) so there are no issues or bluff!"

Forensic case study (FACT, — —): the

\$. . . feature-video sweep recorded cli -stream and the recording SHOWED streaming generation failed: invalid character 'd' looking for beginning of value — a REAL failure the sweep caught precisely because the on-screen text was READ + analyzed, not merely "a video was produced". Root cause was a \$. . . stale binary (source-fixed at

commit , deployed artifact stale); the recording-that-reads-the-screen turned an otherwise-invisible artifact gap into a caught defect (then fixed + § . . -guarded). This is the value § . . makes universal: a recording is proof ONLY when its shown content is machine-read and verified to be a genuine working result.

11 4 137

Strict intensification of § . . (per-feature Status+video ledger) + § . . (liveness + analyzer self-validation) + § . . (CV/OCR pixel oracle) + § . . (content-correctness oracle) + § . . /§ . . (HelixQA + autonomous validation), binding them into ONE intensive-coverage record-and-read mandate. ALL hold:

(A) Intensive total coverage. Every project MUST be covered by intensive automated testing that EXERCISES + RECORDS every feature, every flow, every use case, AND every edge case (valid / invalid / boundary / empty / max / concurrent / failure-injection per § . .) – not a happy-path sample. A feature/flow/edge-case with no recorded, read-verified evidence is a § . . PASS-bluff at the coverage layer (composes § . . full-automation-coverage + § . . discovery-completeness + § . . per-feature ledger).

11 4 153

(B) Real results, rock-solid physical proof, zero bluff. Every recording MUST show the feature GENUINELY WORKING with REAL results (real prompts→real LLM/service responses→real outputs per § . .) – NO false / simulated / stale / frozen result may appear in any recording (§ . . / . .). A recording that shows an error, a bluff response, a frozen/stale frame, or a non-working feature is a FINDING (→ § . . () / § . . fix→retest→re-record), NEVER a confirmation.

(C) Read-the-screen content verification (HelixQA-driven).

The testing System MUST ACTUALLY READ the content shown in each recording – every printed log line, message, UI label, dialog, toast, status text – and VERIFY it (OCR/ROI per § 11.4.35 / § 11.4.36 ()) with a confidence floor for pixel surfaces; direct text capture for terminal/log surfaces; the § 11.4.37 () golden-good/golden-bad self-validated analyzer). "A video was produced" is NOT evidence; "the System read the on-screen text and confirmed it is a genuine working result" is. HelixQA (§ 11.4.38 , the HelixDevelopment/HelixQA submodule) MUST be the driver/orchestrator of this exercise→record→read→score pass – its banks exercise the features, capture the recordings, read the shown content, and score PASS only on a read-confirmed genuine result with the captured artefact path (§ 11.4.39).

(D) Default save path = \$HOME/Downloads . Unless a project declares an explicit override per § 11.4.40 , the default save path for ALL recordings is the host user's home Downloads directory (\$HOME/Downloads – resolved from the running user's home at runtime, NEVER hardcoded). A project MAY override the location per § 11.4.41 (e.g. an external-drive recordings root); the override is recorded in the project layer, not here. Filenames carry the § 11.4.42 project-name prefix; window-scoped capture + fresh-corpus rotation per § 11.4.43 ; raw corpus git-ignored + code-intelligence-excluded per § 11.4.44 , only curated evidence committed under docs/qa/<run-id>/ (§ 11.4.45) at release prep.

(E) No false results survive. Composes the whole § 11.4.46 anti-bluff covenant – a recording-confirmed PASS the operator (or a later re-read) contradicts is a § 11.4.47 operator-escape → mandatory bluff-audit + permanent guard. The read-and-verify analyzer is itself anti-bluff (golden-good PASSes, golden-bad FAILs, § 11.4.48).

Canonical authority: constitution submodule Constitution.md

\$. . . Non-compliance is a release blocker. No escape hatch – no --skip-recording-coverage , --video-without-content-read , --happy-path-recording-suffices , --recording-path-anywhere , --unread-recording-OK , --skip-helixqa-read flag.

\$. . . – Mandatory window-specific video recording + vision validation mandate (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –): "All video evidence must be window-specific MP with mandatory vision validation, terminal cleanup, real results only."

Every feature test, validation, verification, challenge, and QA session that produces video evidence MUST comply with ALL of the following:

(A) **Window-specific recording ONLY.** Every video MUST record ONLY the target application window (Terminal pane, TUI app, browser tab, emulator frame), NEVER the whole desktop/monitor screen. Use window-specific capture mechanisms: macOS screencapture -l<window_id> (obtain window id via osascript -e 'tell application "System Events" to get id of first window of process "<app>"'), Linux xdotool + ffmpeg -video_size WxH -f x11grab -i :0.0+X,Y , or equivalent platform-specific window-targeted capture. Whole-desktop capture leaks operator-private content (\$. . .), dilutes the \$. . . liveness/freeze oracle, and breaks the \$. . . OCR/ROI content oracle. Platform genuinely cannot capture below whole-screen => honest \$. . . SKIP + tracked migration item, never a whole-screen pass-off.

(B) **MP format REQUIRED.** Every recording MUST be saved as .mp4 format (H. codec, movflags +faststart ,

`pix_fmt yuv420p ).. .cast (asciinema) files are supplementary only – the primary evidence is the .mp4 video. If using asciinema for terminal recording, auto-convert to .mp4 via agg + ffmpeg immediately after capture per $ . . (C).`

(C) Project-name prefix REQUIRED. Every recording filename MUST start with the project name in snake_case, derived from `HELIX_RELEASE_PREFIX` in `.env` (per `$ . . )` or the lowercased project root directory name (per `$ . . )`. Format: `<project_name>-<feature>-<YYYYMMDD-HHMMSS>.mp4`. An unprefixed recording is a `$ . .` violation (same severity class as `$ . .`).

(D) Mandatory vision validation. After EVERY recording, the agent MUST perform vision analysis: read the terminal output / screenshot frames / video content and verify: (i) the feature ACTUALLY WORKS (real output, not errors, not empty); (ii) LLM responses are REAL (not simulated, not placeholder, not "for now"); (iii) all tests show PASS with captured evidence; (iv) no "TODO implement", "simulate", "for now", or "in production this would" patterns appear; (v) the output demonstrates the feature working for the end user. Vision validation MUST produce a verdict (PASS/FAIL) with cited evidence path (`$ . .`). A recording without vision validation is a `$ . .` violation – a video that was never read is NOT evidence.

(E) Terminal window cleanup. After each recording completes, the agent MUST dismiss/close ONLY the Terminal window used for that recording (using window-specific close: `osascript` with window id, `xdotool`, or equivalent). MUST NOT close Terminal windows belonging to other working processes. A recording that leaves orphan Terminal windows is a `$ . .` violation.

(F) Real results ONLY. Every recording MUST show REAL working features – real API calls, real LLM responses, real test results. A recording showing errors, empty output, simulated responses, or placeholder content is a § 4.1.1 violation and MUST trigger a fix -> retest -> re-record cycle before acceptance.

(G) Re-runnable evidence. Every recording MUST be reproducible: the command shown in the video MUST be re-runnable to produce the same results. Recordings of one-time manual interactions that cannot be automated are § 4.1.2 violations.

(H) Fresh-corpus rotation. When a new recording run for a scope begins, the agent's own prior in-scope stale recordings at the recording path MUST be removed FIRST (per § 4.1.3). Committed docs/qa/<run-id>/ evidence is the durable record, NOT rotated.

(I) Content verification MANDATORY – not duration-based. The value of a recording is NOT its duration but its CONTENT. A recording MUST demonstrate the ACTUAL feature being used with REAL results. Before accepting ANY recording, the agent MUST verify:

- The expected output patterns ARE present in the recording (e.g., test PASS lines, API response data, feature-specific output)
- The feature ACTUALLY WORKS as demonstrated (not just "something ran")
- LLM responses are REAL content (not simulated, not placeholder, not empty)
- Every claim of "working" is backed by visible evidence in the recording

A 10-second recording that shows a feature working correctly is MORE valuable than a 10-minute recording of empty terminal. Duration is NOT a proxy for quality.

(J) Expected-content specification REQUIRED. Before recording, the agent **MUST** specify what content **SHOULD** appear in the recording (expected patterns, expected test results, expected API responses). After recording, the agent **MUST** verify these patterns **ARE** present. If expected content is **MISSING**, the recording is **REJECTED** regardless of duration.

(K) Content-verification recording workflow. The mandated workflow for every recording:

- . SPECIFY expected content patterns (what **SHOULD** appear)
- . RECORD the feature execution
- . EXTRACT all text from the recording
- . VERIFY expected patterns are present
- . CHECK for simulated/placeholder content
- . ACCEPT only if ALL patterns found AND zero bluffs detected
- . REJECT and re-record if ANY pattern missing or bluff detected

11 4 102

(L) Root cause analysis REQUIRED for rejected recordings.

When a recording is rejected (missing expected content, bluff detected, or empty capture), the agent **MUST** investigate **WHY** before re-recording. Per § 4.1.1, determine the root cause (timing issue, wrong command, tool failure, etc.) and fix it. Simply re-recording without understanding **WHY** the first attempt failed is a § 4.1.1 violation.

(M) Real-time monitoring RECOMMENDED. For complex features, use real-time monitoring that analyzes output **DURING** recording (not after). This catches issues immediately and allows corrective action before the recording completes.

Classification: universal (§ 4.1.1) – the consuming project supplies its concrete window-capture mechanism

genuine working result, and operator-side manual review does not scale to the intensive recording mandate (§ 11.4.104).

The bridge requirement. The recording system MUST provide a real-time or post-recording bridge that feeds captured frames/pixels to HelixQA's test infrastructure (or equivalent automated test framework) for content verification – automated read-the-screen verification against expected content patterns specified BEFORE recording (per § 11.4.104 (J)). The bridge MUST:

- **Capture frames at $\leq$ 1 s intervals** during the recording window – sufficient to assert the feature's progression, not merely its start and end state.
- **Run OCR/vision analysis against each frame** using a self-validated golden-good/golden-bad analyzer per § 11.4.104 (K) – the analyzer itself MUST be mutation-tested so it cannot bluff a PASS on empty/degraded content.
- **Compare extracted text against the SPECIFY-phase expected patterns** – every pattern from the § 11.4.104 (J) expected-content specification MUST be checked.
- **Produce a per-frame PASS/FAIL verdict with an evidence path to the frame** – raw frame image + OCR output + comparison result, linked from the verdict.
- **Surface failures immediately** so the recording can be re-done per § 11.4.104 (L) root-cause analysis, not queued for batch review after the recording window closes.

11 4 105

Project-configured parameters (per § 11.4.105). The frame-capture interval, OCR confidence floor, and pattern-matching thresholds are project-configured, calibrated on the project's own fixtures per § 11.4.105 (no hardcoded industry defaults).

Honest boundary (§ 11.4.160). Vision verification confirms the feature produced expected on-screen output – it does NOT replace § 11.4.160 captured-evidence quality analysis (audio RMS/XRUN, video freeze/jitter/obstruction census) nor § 11.4.160 runtime-signature verification (the artifact's bytes landed and are active on a clean target). If the recording surface is FLAG_SECURE / DRM-protected / sink-blanked per § 11.4.160, the test documents the gap and uses the § 11.4.160 proxy oracle (accessibility cue / subdebug channel) instead – never a faked pixel-verified PASS.

11 4 27

Composition. § 11.4.160 composes with § 11.4.160 (captured-evidence quality – § 11.4.160 adds the automated read-the-screen layer, does not replace quality metrics), § 11.4.160 (HelixQA is the framework that runs the bridge), § 11.4.160 (verdict with evidence path), § 11.4.160 (self-validated analyzer), § 11.4.160 (secure-surface gap documented), § 11.4.160 (proxy oracle for non-introspectable surfaces), § 11.4.160 (per-feature video ledger), § 11.4.160 (intensive recording coverage – § 11.4.160 adds the verification bridge to every recording), § 11.4.160 (J) (SPECIFY-phase expected-content patterns are the bridge's input).

Gates. Propagation gate CM-COVENANT-114-160-PROPAGATION (literal 11.4.160) + recommended gate CM-VISION-VERIFIED-RECORDING-BRIDGE (every recording has a self-validated OCR/vision pass against SPECIFY-phase patterns, with per-frame PASS/FAIL) + paired § 11.4.160 meta-test mutation (strip the bridge/accept a recording without vision verification → gate FAILs; gate-code = separate work item).

Classification: universal (§ 11.4.160) – the consuming project supplies its concrete capture mechanism, OCR/vision tooling, capture interval, confidence floor, and

HelixQA bridge integration per § 2.2.6. No escape hatch — no `--skip-vision-verification`, `--manual-review-suffices`, `--no-bridge`, `--single-frame-suffices`, `--hardcoded-thresholds-OK`, `--unvalidated-analyzer-OK` flag exists.

§ 2.2.6 — **Rootless container runtime mandate** (User mandate, 2.2.6 —)

Forensic anchor — direct user mandate
(2.2.6 —).

The mandate. Every project governed by this Constitution MUST use **Podman in rootless mode** (or an equivalent rootless container runtime) for ALL containerized workloads. Docker in rootful mode, `sudo`, or any escalation to root for container management is **STRICTLY FORBIDDEN** unless the target platform genuinely has no rootless container option AND that platform constraint is documented per § 2.2.6 (structurally-impossible — with cited authoritative sources proving the absence of a rootless option on that platform). Rootless container execution eliminates the most common container-to-host privilege-escalation vector, aligns with the security-by-default posture of the § 2.2.6 credentials-handling mandate, and is available and production-mature on Linux (Podman rootless since v 1.1.4 +, 2.2.6) and macOS (Podman Machine with `--rootful=false`).

The `vasic-digital/containers` Submodule as the sole orchestration layer. Per § 2.2.6 (Containers-submodule mandate), the `vasic-digital/containers` Submodule MUST be used as the **sole container orchestration layer** — every container operation MUST go through the Submodule's `pkg/boot` / `pkg/compose` / `pkg/health` primitives. No ad-hoc `docker` / `podman` commands outside the Submodule; if the Submodule's `pkg/boot` does not support a lifecycle or runtime the project needs, the missing capability MUST be

extended in the Submodule itself via upstream PR (the \$. . . extend-don't-reimplement discipline), not worked around with a raw container command in the project.

The on-demand-infra invariant (refines \$. . .). All container-related integration tests MUST boot infra on-demand via the Submodule (pkg/boot is the test entrypoint), not depend on a pre-running container daemon or operator-side manual podman machine start . Operators are NEVER required to start podman machine / docker compose up / podman-compose up manually – the boot is part of the test entry point. A test that claims to exercise containerized components but silently depends on an already-running daemon (skipping the boot) is a \$. PASS-bluff at the test-infrastructure layer.

Honest boundary (\$. . .). Rootless container execution eliminates the container-to-root privilege-escalation vector and enforces containerization hygiene – it does NOT replace \$. . . credentials-handling (credentials stored in containers are still subject to the leak audit and .gitignore mandates) nor \$. container hygiene (memory limits, OOM policy, restart backoff still apply). A platform that genuinely lacks rootless containers is documented per \$. . . and uses the Submodule's rootful mode with a documented extra risk acceptance – never uses ad-hoc rootful Docker.

Composition. \$. . . composes with \$. . . (Containers-submodule – \$. . . adds rootless-as-default + on-demand-infra invariant), \$. . . (extend-don't-reimplement – missing Submodule capabilities are extended upstream), \$. . . (credentials – rootless containers still require credential hygiene), \$. . . (structurally-impossible – the documentation path for platforms with no rootless option), \$. . . (container hygiene – memory limits/OOM/restart

policies still apply), § 11.4.161 (no-guessing – "the target platform probably has rootless" is not a finding).

Gates. Propagation gate CM-COVENANT-114-161-PROPAGATION (literal 11.4.161) + recommended gate CM-ROOTLESS-CONTAINER-RUNTIME (every container operation is verified as Podman rootless OR has a § 11.4.172 struct-impossible citation for that platform; no sudo /rootful Docker for container management; container integration tests boot on-demand via the containers Submodule) + paired § 11.4.173 meta-test mutation (replace a rootless invocation with sudo or ad-hoc docker, OR bypass the Submodule's pkg/boot with a raw podman command → gate FAILs; gate-code = separate work item).

Classification: universal (§ 11.4.174) – the consuming project supplies its concrete rootless runtime (Podman rootless or equivalent), the containers Submodule integration, and any platform constraint documentation per § 11.4.175 per § 11.4.176. No escape hatch – no --allow-rootful, --sudo-container-OK, --ad-hoc-docker-permitted, --no-containers-submodule, --pre-started-daemon-OK flag exists.

§ 11.4.177 – OpenDesign UI design system mandate (User mandate, – –)

Forensic anchor – direct user mandate
(– –).

The mandate. Every project governed by this Constitution that produces user-facing interfaces (web, desktop, mobile, TUI, or any other surface classes the project supports per § 11.4.178) MUST use **OpenDesign** (<https://github.com/nexu-io/open-design>) as the mandatory UI design-and-refinement system – NOT ad-hoc CSS, not one-off design tools, not improvised inline styles, not a

different design-token system. OpenDesign is consumed as a project dependency (npm package or equivalent per the consuming project's technology stack), never vendored or forked.

Design tokens and theming. The project MUST use OpenDesign's design tokens / themes system for:

- . **All color palette definitions** – supporting BOTH light and dark themes using the project's brand colors extracted from canonical brand assets (logo, style guides, brand guidelines per § 4.1.1). Color tokens are defined in OpenDesign's token format, never duplicated in CSS custom properties or inline styles.
- . **Typography scale and spacing** – typefaces, font sizes, line heights, letter spacing, and content margin/padding tokens are defined in OpenDesign, consumed by every UI component, never scattered across individual component files.
- . **Component-level design tokens** – every reusable UI component (buttons, cards, dialogs, inputs, navigation elements, tables, modals, etc.) defines its tokens (border radius, shadow, hover/active/focus states, animation timing) in OpenDesign, not in component-specific stylesheets.

Extend, never reimplement. If a desired UI pattern (a component, theme variant, token type, or rendering target) is not supported by OpenDesign, the project MUST extend OpenDesign upstream per § 4.1.2 (extend-don't-reimplement) – submit a PR adding the missing capability rather than working around it with ad-hoc CSS or a different tool. A pattern worked around outside the design system is a § 4.1.2 violation (it silently fragments the UI surface – the divergence grows over time and the shared design system no longer governs all UI output).

Light+dark theme variants required. Every UI component MUST ship with BOTH light and dark theme variants, with the color palette defined from the project's canonical brand assets (logo/brand assets per § 4.4.1). A component that renders correctly only in one theme is a § 4.4.1 violation – a project's surfaces MUST be usable in both modes on any platform that supports theme switching.

Layout regression prohibition. UI elements MUST NOT overlap, fonts MUST NOT collide with other UI elements or other text, labels MUST NOT overlay labels. Any layout regression – including but not limited to clipped text, overlapping controls, misaligned grids, overflow-hidden content, or incorrect z-ordering – is a § 4.4.2 violation regardless of whether the underlying token change had good intent. The OpenDesign token-diff comparison is the primary regression-detection tool: token changes that produce layout regressions are caught by diffing the token set before and after the change, correlated with the project's visual regression tests.

Test coverage. All UI changes MUST be covered by the project's standard test types including visual regression tests (§ 4.4.3 (b) layer 1) that capture before/after screenshots of every changed component/token, with per-pixel or perceptual-diff PASS/FAIL.

Honest boundary (§ 4.4.4). OpenDesign governs design tokens and theme definitions – it does NOT replace the project's functional testing (§ 4.4.5), accessibility testing (WCAG color contrast is still verified per § 4.4.6), or the § 4.4.7 / § 4.4.8 UI-driven and dual-approach testing methodology. A project that uses OpenDesign tokens but ships broken UI components is a § 4.4.9 violation

– the tool is mandatory, but using it correctly (tokens → components → tests → regressions) is equally mandatory.

Composition. § . . . composes with § . . . (extend-don't-reimplement – the rule for extending OpenDesign upstream), § . . . (full-automation-coverage – UI components covered by every test type), § . . . (no-fakes-beyond-unit – visual regression tests are real assertions, not mock screenshots), § . . . (b) (four-layer coverage includes visual regression at layer), § . . . (accessibility/color-contrast assertions still apply), § . . . / § . . . (UI-driven/dual-approach testing still applies), § . . . (brand assets location and token-per-project instantiation), § . . . (no-guessing – token configuration is declared, not inferred from appearance), § . . . (visual regression evidence is captured per component).

Gates. Propagation gate CM-COVENANT-114-162-PROPAGATION (literal 11.4.162) + recommended gate CM-OPENDSIGN-UI-SYSTEM (every UI-capable project has OpenDesign installed as a dependency, uses its token/theme system for all colors (light+dark), typography, spacing, and component tokens; extends OpenDesign upstream for unsupported patterns; no ad-hoc CSS-only UI; visual regression tests cover every UI change) + paired § . . . meta-test mutation (replace a token-driven style with inline CSS → gate FAILs; ship a light-only component → gate FAILs; use a non-OpenDesign color tool → gate FAILs; gate-code = separate work item).

Classification: universal (§ . . .) – the consuming project supplies its concrete OpenDesign package installation path, brand-asset location, token definitions file, component registry, and visual regression test harness per § No escape hatch – no --skip-

1.1 4 1.6.3
opendesign , --ad-hoc-css-OK , --light-only-OK , --no-token-
system , --no-visual-regression-tests flag exists.
2.0.2.6 0.6 2.1

**§ . . – Universal Media Validation
& Verification Mandate (User mandate,
– –)**

**Forensic anchor – direct user mandate
(– –).**

The mandate. Every recorded artifact (video MP , audio WAV, screenshots PNG, asciinema cast, text output) produced by any project governed by this Constitution MUST pass through a MEDIA VALIDATION pipeline before being accepted as evidence. A recording that has NOT been validated by the pipeline is a § . PASS-bluff at the evidence layer – it asserts the content is genuine and correct without having verified it.
1.1 4 1.17
1.1 4 1.67 12

(a) Content extraction. The pipeline MUST extract content from the media – OCR for video/screenshots (per § . . with a confidence floor per § . . ()), transcription for audio, text parsing for asciicast and text output. The extraction tooling is self-validated per § . . () – a golden-good fixture MUST produce a PASS, a golden-bad fixture MUST produce a FAIL; an extractor that passes its golden-bad fixture is a bluff gate.
1.1 4 1.17

(b) Pattern matching. The extracted content MUST be compared against the SPECIFY-phase expected patterns per § . . (J). Every specified pattern must be matched by at least one extracted line/frame; any specified pattern with no match yields a FAIL with an exact pinpoint record (which pattern, what the extracted content contained instead). The pattern set is the closed set of expected strings the test author declares before recording; an un-patterned PASS is inadmissible.

(c) Self-validated analyzer. The media validation pipeline as a whole (extraction + matching + verdict) MUST be mutation-tested with a golden-good/golden-bad fixture pair (§ . . . ()) – the golden-good MUST PASS, the golden-bad MUST FAIL. A pipeline that passes its golden-bad fixture is itself a § . bluff gate; the pipeline MUST NOT be used for any evidence until its self-validation is restored.

(d) Structured verdict. The pipeline MUST produce a structured verdict for each artifact containing: PASS/FAIL, exact evidence path (file + frame/line/timestamp), matched patterns list (pattern → extracted text → frame/line), unmatched patterns list (pattern → expected → actual), analyzer version and calibration hash. The verdict is committed alongside the artifact under docs/qa/<run-id>/ (§ . . .).

(e) Exact pinpoint on FAIL. When the verdict is FAIL, the pipeline MUST provide pinpoint data – exactly which media artifact, which frame (video) or line (text/asciicast) or timestamp (audio), which OCR region or audio segment, which expected pattern, and what the extracted content contained instead. The pinpoint is the root-cause entry point for § . . . systematic-debugging before re-recording per § . . . (L).

(f) Dual trigger mode. The pipeline MUST be triggerable both as a post-recording step (validate after the full recording completes, per § . . . (K)) AND as a real-time monitoring step that analyses frames/output during recording per § . . . (M) – real-time mode catches failures immediately, reducing re-record overhead.

(g) Paired § . meta-test mutation. The media validator is itself subject to paired § . mutation – a golden-bad fixture that the validator MUST reject, or the validator is a bluff gate. The mutation test creates a

golden-bad fixture, feeds it to the validator, and asserts FAIL. If the validator returns PASS on the golden-bad fixture, the mutation test FAILs and the validator is inadmissible until fixed.

Honest boundary (§ 11.4.17). The media validation pipeline confirms the artifact contains the content it claims to contain and matches the intended patterns – it does NOT prove the feature is correct for the end user (that rests on § 11.4.16 captured-evidence quality analysis, § 11.4.18 liveness, and § 11.4.19 runtime-signature), and it does NOT replace § 11.4.15 full-suite retest. A pipeline that can be bypassed (a recording accepted without validation) is a § 11.4.15 violation – the pipeline is mandatory for every recorded artifact, never optional.

Classification: universal (§ 11.4.15) – the consuming project supplies its concrete media extractors (OCR engine, transcription engine, text parser), golden-good/golden-bad fixture pair, expected-pattern declaration mechanism, and structured-verdict template per

§ 11.4.15. Composes

§ 11.4.15 (J)/. (K)/.

§ 11.4.15. Propagation gate CM-COVENANT-114-163-PROPAGATION (literal 11.4.163) + recommended gate CM-MEDIA-VALIDATION-PIPELINE (every recorded artifact crosses a self-validated media validation pipeline before acceptance; golden-bad fixture produces FAIL) + paired § 11.4.15 meta-test mutation (strip the literal → propagation gate FAILs; bypass the pipeline and accept a raw unvalidated recording → CM-MEDIA-VALIDATION-PIPELINE FAILs; gate-code = separate work item).

(M)

§ 2.2.6 – Universal Constitution Auto-Propagation & Hook System (User mandate, – –)

Forensic anchor – direct user mandate

(– –).

The mandate. Every fetch+pull of the constitution submodule MUST trigger an automatic post-update hook that installs/registers newly added or modified constitution components, so a pull of the constitution submodule is a COMPLETE governance-update transaction – not a partial one that leaves skills unregistered, MCP servers unknown, hooks unwired, or scripts non-executable. The canonical hook script lives at `constitution/scripts/post_update_hook.sh` and is inherited by reference (`$ . .`) – NEVER copied locally.

(a) Change detection. The hook MUST detect which constitution submodule files changed in the latest pull:

`Constitution.md` , `CLAUDE.md` , `AGENTS.md` , `QWEN.md` , `GEMINI.md` , and any file under `scripts/` , `hooks/` , `skills/` , `mcp/` , `plugins/` . It compares HEAD before vs after the pull via `git diff --name-only HEAD@{1} HEAD` (or the equivalent for the merge/pull operation).

(b) Skill registration. For any newly added or modified file under `skills/` , the hook MUST install/register it with the CLI agent's skill system (e.g. for Claude Code: add it to the project's skill configuration so the agent discovers and can invoke it per `$ . .` (B)). The registration mechanism is project-runtime-specific per `$ . .` – the hook invokes a consumer-defined script at a canonical path (`scripts/register_skills.sh`) that the consumer project provides.

(c) MCP server registration. For any newly added or modified file under `mcp/` , the hook MUST register the MCP server with the agent's MCP configuration (e.g. for Claude

Code: add an entry to `.mcp.json` or `settings.json`). The registration mechanism is project-runtime-specific per § 4.1.1 – the hook invokes a consumer-defined script at a canonical path (`scripts/register_mcp.sh`).

(d) Hook installation. For any newly added or modified file under `hooks/`, the hook MUST install it into the project's `.git/hooks/` or runtime hooks directory, making it executable. The hook script sets executable permission and copies/symlinks it; existing hooks with the same name are overwritten only if the new file's mtime is newer (safe default per § 4.1.1 – never overwrite a consumer's custom hook without a newer upstream version).

11 4 37

(e) Script validation. For any newly added or modified file under `scripts/`, the hook MUST (i) make it executable, (ii) validate its syntax – `sh -n` for POSIX-sh scripts, `bash -n` for bash scripts (per § 4.1.1), (iii) validate any `helix-deps.yaml` or `.gitignore-meta/` entries per § 4.1.1 / § 4.1.1.

(f) Summary report. The hook MUST emit a summary of what was installed/updated: number of skills registered, MCP servers registered, hooks installed, scripts made executable and validated. The summary goes to stdout AND to `.constitution/state/last_update.log` (gitignored per § 4.1.1).

11 4 38

(g) Graceful failure. If any component installation fails (skill not registered, MCP server not wired, hook not installed, script not validated), the hook MUST log the failure per § 4.1.1 with exact file:line and continue with remaining installations – one failure does not abort the entire post-update transaction. The failure is surfaced in the summary report and is a finding for § 4.1.1 investigation.

(h) Anti-bluff. The hook is itself anti-bluff – tested by a paired § 4.1.1 mutation that adds/removes a skill in the

constitution submodule, pulls the submodule, and asserts the hook detects the change. A hook that reports "nothing to install" when a new skill was added is a § 11.4.35. PASS-bluff at the auto-propagation layer.

Consumer obligation. Every project that fetches/pulls the constitution submodule MUST invoke the post-update hook after the fetch completes. The invocation mechanism is project-defined per § 11.4.36 – either a `post-update` git hook wired via `scripts/install_git_hooks.sh`, or a post-pull script that calls `constitution/scripts/post_update_hook.sh` explicitly. A project that pulls the constitution and does NOT invoke the hook is non-compliant – it has an incomplete governance-update transaction, and skills/MCP/hooks/scripts it believes are present may not be.

Honest boundary (§ 11.4.37). The auto-propagation hook installs/registers constitution components – it does NOT guarantee the consumer's runtime accepts them (MCP server may have incompatible dependencies, skill may require a runtime version the consumer lacks) and does NOT replace § 11.4.37 post-constitution-pull validation sweep (which verifies rules are enforced, not just that components are installed). Installation failures not investigated are the finding § 11.4.38 captures.

Classification: universal (§ 11.4.39) – the consuming project supplies its concrete skill-registration script, MCP-registration script, hook-installation mechanism, and post-pull invocation method per § 11.4.40. Composes § 11.4.41 /. /. /. /. (B)/. /. / § 11.4.42. Propagation gate `CM-COVENANT-114-164-PROPAGATION` (literal `11.4.164`) + recommended gate `CM-CONSTITUTION-AUTO-PROPAGATION` (post_update_hook.sh exists + executable, invoked after every pull, logs changes or failures, paired § 11.4.43 mutation detects a new skill) + paired § 11.4.44 meta-test mutation (strip the literal → propagation gate

FAILs; pull the constitution and skip invoking the hook →
CM-CONSTITUTION-AUTO-PROPAGATION FAILs; gate-code = separate
work item).

§ . . – Universal Independent Verification Agent Mandate (User mandate, – –)

Forensic anchor – direct user mandate
(– –).

The mandate. Every code change OR recorded media artifact produced by any project governed by this Constitution MUST pass through an INDEPENDENT verification agent structurally separate from the author (§ . . / § . . , dedicated subagent or distinct human, NEVER the author re-reading their own work – § . . self-evaluation, PRECEDES, and never satisfies). The verifier is one of MULTIPLE STRONG LAYERS – the FIRST verification layer every output crosses, complementing (never replacing) § . . runtime-signature verification, § . . / § . . code review, and § . . full-suite retest.

(a) For CODE changes. The independent verification agent MUST: (i) run the § . . code-review discipline (adversarial analysis, blast radius, contract integrity, test genuineness), (ii) build and test the changed packages (invoke the project's pre_build_verification.sh on the diff scope), (iii) run the paired § . . meta-test mutations for every gate the change touches (asserting each mutated gate still FAILs and its restoration still PASSes), (iv) verify the § . . runtime-signature registry entry for this change declares one machine-checkable observable. A code change that passes only one of (i)–(iv) is NOT verified – all four must hold.

(b) For MEDIA artifacts (video/audio/screenshots/asciicast/text output). The independent verification agent MUST: (i) run the `$ . . .` media validation pipeline on every recorded artifact, (ii) confirm the content is genuine – real LLM/service responses, not mocked/stub/simulated/placeholder (per `$ . . . (I)`), (iii) confirm the artifact format is correct (MP H. for video, WAV/FLAC for audio, PNG for screenshots, cast+mp for asciicast per `$ . . . (B)`), (iv) cross-check the evidence path against the artifact (the path exists, is non-empty, and the content matches the feature's expected output). A media artifact that passes format but shows simulated/placeholder content is a finding, not a PASS.

(c) For DOCS changes. The independent verification agent MUST: (i) validate that HTML + PDF + DOCX exports exist and are current per `$ . . .` (mtime parity), (ii) verify the `$ . . .` revision header is present and its integer has incremented from the previous commit, (iii) for docs that carry a `$ . . .` drift-proof fingerprint, verify the fingerprint matches the sorted roster/keyset. A doc change with a stale export or missing revision header is NOT verified.

(d) For CONFIG changes. The independent verification agent MUST: (i) validate YAML/JSON/TOML syntax (parse with the language's own parser – `yq`, `jq`, basic JSON parser, `python -m json.tool`, etc.), (ii) cross-reference against a schema where available (e.g. JSON Schema for `.settings.json`, YAML Schema for `helix-deps.yaml`), (iii) confirm no credential or secret pattern leaked per `$ . . .`. A config change with a syntax error, schema violation, or leak is a finding, not a PASS.

(e) Structured findings. The verifier MUST produce a structured findings report – for each finding: item (which code/media/doc/config), verification class (a-d), tool/

§ . . – Big-work-item feature
 work-stream lifecycle (own copy-on-write
 project copy + own branch/tags + per-feature
 builds, no-merge-until-approved, trunk-merged-
 into-every-stream) (User mandate,
 - -)

Every BIG work item – a new feature OR a drastic/large fix
 – MUST be developed as its own isolated **feature work-
 stream**: its OWN full project copy in a sibling directory,
 its OWN feature branch, its OWN per-feature flashable
 builds + feature release tags, kept SEPARATE from the
 trunk until manually approved after testing, with the
 trunk regularly merged INTO every feature stream. This
 generalises Git-flow's feature/release/tag discipline
 (long-lived `main` trunk, one `feature/<slug>` per big item,
 no merge to trunk until reviewed-complete) and binds it to
 the project's parallel-stream (§ . . /
 § . .), single-builder, and finite-test-device
 realities. The mandate (ALL hold):

(A) **One big item = one feature work-stream in a sibling
 project copy.** Each big work item gets its own complete
 project copy in a sibling directory (`<project>` ,
`<project>_<slug-1>` , `<project>_<slug-2>` , ...), the `<slug>` a
 lowercase identifier per § . . . The trunk copy is
 the canonical `main` . Small changes stay on the trunk; only
 BIG items earn a stream (the threshold is the project's
 own per § . .).

(B) **Disk-feasible duplication is MANDATORY, naive copies
 FORBIDDEN.** A feature-stream copy MUST be created by a
 near-zero-disk mechanism – on a copy-on-write filesystem
 (btrfs/XFS/ZFS/APFS) a reflink/CoW clone (`cp -a --
 reflink=auto` / filesystem snapshot) so only DIVERGED bytes
 consume disk; elsewhere a shared-object-store mechanism
 (e.g. `git worktree` + externalized build output + sparse
 where the build tolerates it). A full deep copy or full

`git clone` per stream of a large working tree is FORBIDDEN where it would exhaust free space (§ data safety + the project's host-disk reality per § . . .). The chosen mechanism MUST be verified to consume ~ disk for a fresh stream (captured evidence per § . . .) before being relied on (§ . . . – no guessing about disk). Build output (`out/ -class`) MUST be per-stream and excluded from the CoW clone; one shared compiler cache (`ccache-class`) across streams maximises queued- build hit-rate.

(C) Own branch + own tags, separate until approved. Each stream develops on its own `feature/<slug>` branch; per-feature flashable builds are marked with feature tags under a scheme that embeds the slug so the build is greppable across every repository (composes §), e.g. `<base-version>-feat-<slug>[-<iter>]`, the `<iter>` monotonic per § **The feature branch is NEVER merged to the trunk until the operator manually approves after testing** – approval requires the § . . . full-suite retest GREEN on a clean target, then the § . . . merge-first integration, then a § . . . ff-only (no-force) merge to trunk.

(D) Trunk merged INTO every feature stream, regularly. The trunk MUST be merged INTO every live feature stream FREQUENTLY (at minimum after every trunk tag / daily) to prevent long-lived-branch conflict storms – `git merge` (never rebase a shared/tagged branch). A stream that has not been synced within the cadence is flagged stale.

(E) Feature-level branching AND tagging cascades to ALL submodules the moment a submodule is modified. The instant a stream modifies ANY submodule, that submodule gets its own `feature/<slug>` branch; per-feature tags on the main repo are mirrored on EVERY touched submodule at the same HEAD and ff-only-pushed to all owned remotes (§ . . . + the project's tag-cascade rule). Untouched submodules stay trunk-pinned (no branch/tag

churn). Merge-to-trunk applies the same cascade. Where the duplication mechanism is `git worktree`, the experimental/disk- duplicating worktree-submodule path (each worktree keeps a unique submodule copy; removal must be forced) is a documented hazard and the CoW-clone mechanism (each stream owns CoW-shared submodule git dirs) is preferred where available.

(F) Single-builder build queue + finite-device test queue.

Where the host can run only ONE build at a time, feature builds MUST QUEUE (FIFO + advisory lock, composes `$ . . . L / $ . . . / $ . . .`), never contend. Where live testing has a finite device count, per-device exclusive ownership (`$ . . .`) serialises tests; streams beyond the device count queue. The build queue MUST garbage-collect idle streams' regenerable build output (`$ . . .`) to keep disk within budget.

(G) Every change crosses the full quality gauntlet. Every change in any stream passes the extensive independent code review (`$ . . . / $ . . . / $ . . .`)

- multi-angle impact research (`$ . . .`) + anti-bluff captured-evidence testing (`$ . . . / $ . . . / $ . . .`) – a feature stream is NOT a quality carve-out.

(H) Resumable, data-safe stream lifecycle. A stream registry (single source of "is this stream owed?", composes `$ . . . / $ . . .`) tracks each stream's branch, base tag, build/out state, last-trunk-merge, and merge-approval state. Stream creation is atomic-on-success (clone to a `.partial` name, rename on success) and resumable (`$ . . .` quiescence check on resume). Retiring a stream requires a `$ . . .` pre-op backup + a guard that the feature work is merged/approved (unmerged work blocks auto-deletion). Host-session safety

(`$`) and the $\leq$ %-memory ceiling (`$` .) bound
clone + build.

Honest boundary (`$` . .): this lifecycle guarantees ISOLATION, disk-feasibility, greppable identity, and a controlled trunk feature merge discipline – it does NOT by itself prove any stream's work is correct (each change still crosses `$` . . runtime-signature + `$` . . retest). Classification: universal (`$` . .) – the consuming project supplies its filesystem/CoW mechanism, sibling- dir + tag naming, build-queue + device-queue counts, and big-item threshold per `$` . . .

Propagation gate `CM-COVENANT-114-167-PROPAGATION` (literal `11.4.167` across the consumer fleet) + recommended gates `CM-FEATURE-WORKSTREAM-COW-CLONE` (a feature stream is created by the verified near-zero-disk mechanism, NOT a full deep copy/ clone; per-stream out/ + shared cache; captured disk-cost evidence) + `CM-FEATURE-WORKSTREAM-NO-MERGE-UNTIL-APPROVED` (no feature→trunk merge without the approval marker + `$` . . GREEN) + `CM-FEATURE-WORKSTREAM-TRUNK-SYNC-CADENCE` (every live stream's last-trunk-merge within the declared cadence) + `CM-FEATURE-WORKSTREAM-SUBMODULE-CASCADE` (every touched submodule carries the `feature/<slug>` branch + mirrored feature tag). Paired `$` . meta-test mutations: strip the literal → propagation gate FAILs; create a stream via full deep copy → `CM-FEATURE-WORKSTREAM-COW-CLONE` FAILs; merge a feature to trunk with no approval marker → `CM-FEATURE-WORKSTREAM-NO-MERGE-UNTIL-APPROVED` FAILs; let a stream go stale past the cadence → `CM-FEATURE-WORKSTREAM-TRUNK-SYNC-CADENCE` FAILs; modify a submodule in a stream without its `feature/<slug>` branch/tag → `CM-FEATURE-WORKSTREAM-SUBMODULE-CASCADE` FAILs (gate-code = separate work item).

Composes `$` / `$` . / `$` . . / `$` . . /
`$` . . / `$` . . / `$` . . /

11 4 54 11 4 60 11 4 69
11 4 71 11 4 72 11 4 84
11 4 103 11 4 108 11 4 113

§ 11 4 116 • 10 4 119 / § 115 • • / § • • /
§ 11 4 142 • 10 4 145 / § 147 • • / § • • /
§ 11 4 151 • 10 12 6 / § 12 7 • • / § • • /
§ 12 8 • 12 9 • 1 1 / § • • / § • • /
§ • • / § • • / § • • /
§ • • / § • • / § • • /
§ 11 4 147 / § / § • / § • /
§ • / § • / § •

Canonical authority: constitution submodule Constitution.md

§ • • • Non-compliance is a release blocker. No escape hatch – no `--full-clone-per-feature`, `--deep-copy-stream`, `--skip-cow-clone`, `--merge-feature-without-approval`, `--no-trunk-sync`, `--skip-submodule-cascade`, `--unbounded-parallel-builds`, `--shared-device-no-lock`, `--retire-unmerged-stream`, `--feature-stream-quality-carveout` flag.

2026 06 23

§ • • – **Exported-document independent content + textual + full-visual validation mandate** (User mandate, – –)

Forensic anchor (FACT, – –). A "green" Mermaid export fix shipped RAW gantt-diagram source – the literal lines `gantt`, `title`, `dateFormat`, `section`, `:done`, – as plain readable text into user-facing PDFs, unreadable to any reader who expected a rendered diagram. The fix's validation grepped the HTML for `class="mermaid"` but NEVER opened the EXPORTED PDF to check whether raw diagram source had leaked into the body. The operator caught it (§ • • operator-escape – the green suite missed it). Root lesson: generated/exported documents were never INDEPENDENTLY validated for the VISUAL + TEXTUAL correctness of the EXPORTED artifact itself – a `grep` of the source-or-intermediate HTML is NOT a check of what the reader actually receives. This is the § • • SOURCE→ARTIFACT gap (the change present in source /

intermediate, but the rendered bytes a reader opens are broken) materialised at the documentation-export layer.

The mandate. Every generated/exported document – HTML, PDF, DOCX, or any format in the § 4.1.1 universal-Markdown-export scope (and the § 4.1.2 four-format feature-Status set that adds DOCX) – MUST pass INDEPENDENT validation by a review agent structurally separate from the generator (§ 4.1.3 / § 4.1.4 dedicated subagent or a distinct human; NEVER the author/generator self-checking its own output – § 4.1.5 self-evaluation PRECEDES and never satisfies; this is the § 4.1.6 every-change-reviewed discipline applied to exported artifacts). The validation MUST be run on BOTH the source Markdown AND every exported artifact derived from it, ALWAYS – every export, not a sampled subset – across THREE layers, ALL of which MUST hold:

() **CONTENT layer.** The export faithfully carries the source's intent + data: nothing dropped, truncated, re-ordered into nonsense, or garbled. Every section, table, list, code block, figure caption, and link the source declares is present and complete in the export (a table truncated to its first page, a list collapsed, a section silently omitted, a link stripped of its target are all CONTENT findings). Cross-check the exported artifact against the source – present-and-complete, not merely "the file exists and is non-empty" (the § 4.1.7 / § 4.1.8 captured-evidence quality bar, not a presence-only bar).

() **TEXTUAL layer.** The export is human-readable: NO raw markup, NO diagram source, NO unrendered code-fence body leaking as ordinary body text. Specifically forbidden as body text in the exported artifact: raw Mermaid diagram source (mermaid, gantt, graph, flowchart, sequenceDiagram, classDiagram, stateDiagram, erDiagram, pie, gitGraph, journey, dateFormat, section

... , :done, / :active, / :crit, task lines, and the closing-fence residue), raw HTML tags rendered as visible text, raw LaTeX/math source where math rendering was intended, and any triple-backtick code-fence whose intended rendering (diagram / formatted block) was supposed to replace the source. The - - raw-gantt-in-PDF failure is the canonical TEXTUAL finding. Detection reads the exported artifact as text (e.g. pdftotext on the PDF, the rendered DOM/text of the HTML, the extracted document text of the DOCX) and asserts the forbidden tokens do NOT appear as body text.

() **FULL VISUAL layer.** Embedded diagrams render as IMAGES, not as source; layout is intact; no overlapping / cut-off / clipped / garbled / collapsed / blank content. Verified by RENDERING the export and inspecting the rendered result – never by inspecting the source or an intermediate alone: e.g. pdftotext to catch raw diagram source leaking as text (composes layer), pdfimages (or the format's image-enumeration) to CONFIRM the intended diagrams are present as rendered raster/vector images (a diagram that should be an image but produces zero embedded images is a VISUAL finding – the source leaked instead of rendering), and pdftoppm →OCR (\$. . / \$. . (), confidence floor + ROI) to CONFIRM the rasterised page carries human-readable visual content and to re-catch any raw-source-as-pixels. Captured evidence per \$. . /\$. . /\$. . (pdftotext_body.txt , pdfimages_manifest.txt , ocr_pages/ , the per-finding pinpoint – which page/region/token, expected vs actual) accompanies every verdict.

() **Anti-bluff – the validator is itself anti-bluff.** A rubber-stamp "the export looks fine" is NOT validation (\$. . / \$. . – a rubber-stamp validator is a PASS-bluff at the export-validation layer). The analyzer/validator is self-validated with a golden-good / golden-

bad fixture pair per § 1.1.1 (1.1.1): a golden-good export (clean rendered diagram + faithful content) MUST PASS, and a golden-bad export (a PDF/HTML/DOCX deliberately seeded with raw Mermaid `gantt` source as body text, plus a truncated table) MUST FAIL – a validator that passes its golden-bad fixture is a bluff gate and is investigated per § 1.1.1 before any further use. The validation iterates to a clean GO per § 1.1.1: any finding in ANY of the three layers (BLOCKING / nit / warning) re-arms the loop – the generator/author fixes it (e.g. pre-renders the diagram to an image before export), the export is regenerated, and the validator RE-VALIDATES, until GO with ZERO findings and ZERO warnings, every verdict backed by captured rendered-artifact evidence. A generated document that ships WITHOUT this independent content+textual+visual validation, or that contains raw diagram source / unreadable garble / dropped content in the exported artifact, is a § 1.1.1 PASS-bluff at the documentation layer.

Honest boundary (§ 1.1.1). This validation confirms the EXPORTED artifact a reader opens carries the source's content faithfully, is textually readable, and renders its visuals as images – it does NOT prove the source's content is itself correct (that rests on the source's own § 1.1.1 review + § 1.1.1 regression guards) and does NOT replace the § 1.1.1 mtime-parity export-sync gate (which proves the exports are FRESH, not that they are READABLE – § 1.1.1 is the READABILITY/FIDELITY layer the mtime gate cannot see). A FLAG_SECURE / DRM-blanked / sink-blanked rendering surface that cannot be rasterised documents the gap per § 1.1.1 and uses the § 1.1.1 proxy oracle – never a faked visual PASS. "The HTML had `class="mermaid"` " is a source-side check, NOT an exported-artifact check (the exact - - omission); " `pdftotext` shows no `gantt` source AND `pdfimages` shows the

expected diagram image AND OCR reads human text" is the exported-artifact check.

Classification: universal (§ 11.4.167) – a platform-neutral exported-document-fidelity discipline reusable by ANY project that generates/exports documents; the consuming project supplies its concrete pre-render pipeline, exporters, and rendered-artifact validators (text-extraction / image-enumeration / rasterise→OCR) per § 11.4.167. Composes § 11.4.168 (the export-scope + freshness layer this READABILITY layer sits atop) / § 11.4.169 (project-wide MD/HTML/PDF/DOCX styling – styled exports MUST still pass the three layers) / § 11.4.170 (the self-validated-analyzer + liveness/OCR validation techniques) / § 11.4.171 (CV/OCR pixel-oracle for rasterised-page validation) / § 11.4.172 (a fixed export-fidelity defect registers a permanent regression guard) / § 11.4.173 (the – – operator-escape that motivated this anchor) / § 11.4.174 (independent every-change review applied to exports) / § 11.4.175 (window-scoped recording + vision-validation discipline, the recording analogue) / § 11.4.176 (the media-validation pipeline – § 11.4.177 is its document-export sibling) / § 11.4.178 (the independent verification agent – § 11.4.179 (c) DOCS layer is its export-fidelity refinement) / § 11.4.180.

Propagation gate CM-COVENANT-114-168-PROPAGATION (literal 11.4.168 across the consumer fleet) + recommended gate CM-EXPORTED-DOC-VISUALLY-VALIDATED (every exported document in the § 11.4.168 scope carries an independent-validation marker – produced by a validator structurally separate from the generator – asserting the three layers PASS on the exported artifact, with captured rendered-artifact evidence; the validator is self-validated by a golden-good/golden-bad fixture pair) + paired § 11.4.181.

meta-test mutation (strip the literal 11.4.168 → propagation gate FAILs; produce a PDF/HTML/DOCX containing raw diagram source – e.g. raw Mermaid gantt lines – as body text and assert the gate catches it → CM-EXPORTED-DOC-VISUALLY-VALIDATED FAILs; feed the validator its golden-bad fixture → self-validation FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
§ . . . Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-visual-validation , --raw-source-in-pdf-ok , --textual-check-suffices , --self-validate-export , --diagram-source-as-text-ok flag exists.

§ . . . – Mandatory comprehensive test-type coverage with anti-bluff captured evidence (User mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"MANDATORY test types EACH project MUST have with coverage aiming to % without false results, rock solid evidence produced covered with hard physical proofs and zero bluff: unit, integration, e e, full automation, challenges (challenges submodule), helix qa (helix qa submodule with proper tests banks (suites) and comprehensive fully autonomous qa sessions), ddos tests, security tests, stress and chaos tests, concurrency (atomicity) tests, race conditions or deadlocks tests, memory tests, bechmarking tests."

Every project under this Constitution MUST be covered by the following CLOSED, ENUMERATED set of test types, each

driven to as close to 100% coverage as the domain permits, with every PASS backed by rock-solid captured PHYSICAL evidence (§ 11.4.159 / § 11.4.160 / § 11.4.161) and zero false results, zero bluff of any kind (§ 11.4.27 / § 11.4.28 / § 11.4.29). A green result with no falsifiable captured evidence is a § 11.4.162 PASS-bluff regardless of which layer reports it.

§ 11.4.163 is the strict, explicitly-enumerated expansion of § 11.4.164 (no-fakes-beyond-unit + %-test-type coverage) – it freezes the canonical test-type SET and binds each type to captured evidence and, where named, to the owned submodule that supplies it.

The mandatory test-type set (each REQUIRED where the domain warrants it; an honest § 11.4.165 SKIP-with-reason – topology / hardware / credential genuinely absent – is the ONLY permitted absence, NEVER a silent gap):

- **Unit** – isolated logic; the ONLY layer where mocks / stubs / fakes / placeholders are permitted (§ 11.4.161 (A)).
- **Integration** – real, fully-wired components against real infrastructure (no fakes beyond unit), infra booted on-demand via the containers submodule (§ 11.4.107 / § 11.4.108 rootless).
- **End-to-end (e2e)** – the full user journey across the real System, captured-evidence per § 11.4.25 / § 11.4.52 / § 11.4.53.
- **Full-automation** – fully autonomous, re-runnable end-to-end with NO manual intervention after start (§ 11.4.27 / § 11.4.28 / § 11.4.29), deterministic across N iterations (§ 11.4.4).
- **Challenges** – the challenges submodule (basic-digital/challenges): anti-bluff Challenge banks that score PASS only on positive captured evidence (§ 11.4.166 (B) / § 11.4.167 (b) layer).

- **HelixQA** – the `helix_qa` submodule with proper written **test banks (suites)** for every application / service / platform AND comprehensive **fully-autonomous QA sessions** driving every registered bank with captured wire evidence per check (`$ ./helix_qa`).
- **DDoS / load-flood** – sustained adversarial traffic; the System refuses cleanly OR degrades gracefully, never collapses; captured throughput / latency / error-rate evidence.
- **Security** – authn / authz, injection / taint, secret-leak (`$ ./security`), transport + crypto, dependency CVEs, default-deny enforcement; composes the **security** submodule.
- **Stress + Chaos** – sustained load + concurrent contention + boundary inputs + failure-injection (process-death / network-fault / input-corruption / resource-exhaustion / state-corruption) with categorised recovery + captured evidence (`$ ./stress_chaos`).
- **Concurrency / atomicity** – correctness under concurrent callers; no lost updates; transactional atomicity proven; idempotency under replay.
- **Race-condition / deadlock** – race detector + lock-order / deadlock analysis under contention; no blocking operation inside a shared-lock region (the § development principles); captured detector output is the evidence.
- **Memory** – leak census over soak, peak-RSS ceilings (e.g. the `iOS-NEPacketTunnelProvider-class` budget), allocation / fragmentation profile; no unbounded growth across a `h / N`-iteration soak.
- **Benchmarking / performance** – `p / p / p` latency + throughput + resource cost vs a recorded baseline; a regression vs baseline is a finding, not noise.

Four-layer enforcement per § 11.4.3 (b) of the suite itself: pre-build gate (each required type present + executable + parseable for the domain), post-build, on-device / runtime where applicable, and a paired § 11.4.3 meta-test mutation per type (strip a type's evidence-capture or its anti-bluff assertion → the gate FAILs). The per-project coverage ledger (§ 11.4.3 / § 11.4.3.1) MUST classify every feature × test-type × evidence-state; a missing required type, or an OPERATOR_ATTENDED_ONLY row, is a release blocker until promoted or honestly SKIP-justified (§ 11.4.3 / § 11.4.3.1).

Composes / **strengthens** § . . / § . . /
 § . . / § . . / § . . (its
 strict enumerated expansion) / § . . /
 § . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . Classification: universal (§ . .) – the
 consuming project supplies its concrete per-type
 harnesses, infra, and calibrated thresholds per
 § . . .

```
Pre-build gate CM-COVENANT-114-169-PROPAGATION (literal
11.4.169 across the consumer fleet) + recommended gate
CM-MANDATORY-TEST-TYPES-COVERED (every required test type
present-or-honestly-SKIPPED per feature, each PASS citing
captured evidence) + paired $ . meta-test mutation
(strip a required type or drop its evidence citation →
gate FAILs; gate-code = separate work item).
```

EXTENSION – SUITE SHAPE: the shape is a consequence, and the generating rule is push-tests-down

```
(      -      ). §      .      freezes WHICH
```

KINDS of test must exist; it says nothing about WHERE a given behaviour should be tested, and a project can satisfy every enumerated type while placing its assertions

at an altitude that proves little. The AI-curriculum corpus (module) supplies the placement rule and – importantly – REFUSES to prescribe a shape, which this extension preserves rather than overrides. **(a) NO SHAPE IS MANDATED.** Corpus: *"the shape is a consequence, not a goal."* The two named shapes are the **test pyramid** (*"many fast unit tests at the base, fewer integration/service tests, very few slow end-to-end tests ... ideal when most of your risk is in pure logic"*) and the **testing trophy** (*"a static base (types + lint), a fat integration middle, a thin end-to-end top ... for I/O-heavy web/backend code, integration tests hit the best confidence-per-cost because they exercise the real wiring – DB, queue, provider – where the actual bugs live"*). The selection criterion is a single question – *"where do I get the most confidence per second?"* – answered from *"where this system's bugs live, not dogma."* The corpus explicitly marks *"the trophy is always correct for every system"* as a WRONG answer, and states NO numeric layer ratios; a project adopting a distribution supplies it as DATA per § . . . (the corpus's own vocabulary is qualitative: *"many / fewer / very few"*, *"static base / fat middle / thin top"*). **(b) THE GENERATING RULE IS MANDATED.** Corpus (Fowler): *"push each test as far down as it will still catch the bug; if a high-level test fails with no lower-level test failing, add the lower-level one and consider deleting the redundant high-level one."* This is the operative discipline § . . . was missing: a behaviour is tested at the LOWEST altitude that still catches its defect, and a high-level failure with no lower-level counterpart is a coverage signal to be acted on, not merely a red build. **(c) THE TWO NAMED ANTI-PATTERNS, ONE AT EACH END.** The **ice-cream cone** – *"End-to-end tests give the most confidence, so we have lots of them" → "slow, flaky, still misses unit edges"* – is a § . . . violation at the top; and **mock-heavy unit theater** – *"mock*

the DB, mock the provider, mock the collaborators, assert the code calls the mocks as scripted – and the suite passes precisely because it never touched the system it claims to test", which "can hit % coverage (every line ran against a mock) and still catch nothing" – is a violation at the bottom, already forbidden by § 11.4.211. (A) (mocks permitted ONLY in unit tests) and by § 11.4.212. (C) (an assertion-free test raises coverage identically to a proving one). The corpus names its three sub-mechanisms: the mock "encodes a belief that drifts"; it "tests interaction, not effect" ("The code called charge() once' is not 'a charge happened and the ledger balanced'"); and it "couples tests to implementation" so a behaviour-preserving refactor breaks it. Honest boundary (§ 11.4.213): this extension mandates a DECLARED, JUSTIFIED shape and the push-down rule – it does NOT mandate any particular shape, any ratio, or any per-layer count, and a project whose bugs genuinely live in the wiring is CORRECT to run a fat integration middle. Recommended mechanism gate CM-SUITE-SHAPE-DECLARED-AND-PUSHED-DOWN (the project declares its chosen shape with the risk-location justification that selected it; a high-level test that fails with no lower-level counterpart opens a tracked § 11.4.214 item to add the lower-level test) + paired § 11.4.215 mutation (remove the shape declaration, or record a high-level-only failure with no tracked lower-level follow-up → the gate MUST FAIL; golden-FALSE per § 11.4.216 ()): a declared trophy shape on a genuinely I/O-heavy system MUST NOT fire it merely for having few unit tests). Gate-code = separate work item, NOT claimed shipped (§ 11.4.217 / § 11.4.218).

Canonical authority: constitution submodule Constitution.md § 11.4.219. Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-test-type, --partial-coverage-ok, --evidence-optional, --challenges-not-applicable, --skip-helixqa, --no-chaos, --no-

ddos , --skip-memory-test , --skip-race-detector , --bench-optional flag exists.

2025 08 18

2025 08 25

11 4 140

[COLLISION-REDIRECT - -] The
- - HelixTranslate canonical
translation pipeline mandate previously carried
\$. . here has been re-minted to
\$. . to resolve the pre-existing double-
mandate collision surfaced by \$. . (B)
block-integrity audit. The other \$. .
mandate (Universal action-prefix system,
- -) at line ~ is unchanged
and remains the sole occupant of the \$. .
anchor number. Operator-approved re-mint via
AskUserQuestion on - - .

2025 08 18

2025 08 25

11 4 141

[COLLISION-REDIRECT - -] The
- - Mandatory independent per-
language translation review mandate previously carried
\$. . here has been re-minted to
\$. . to resolve the pre-existing double-
mandate collision surfaced by \$. . (B)
block-integrity audit. The other \$. .
mandate (Token-efficiency, - -) at
line ~ is unchanged and remains the sole
occupant of the \$. . anchor number.
Operator-approved re-mint via AskUserQuestion on
- - .

§ . . – **Device-independent host-side rendered-UI visual-proof mandate** (User mandate, – –). Verbatim operator escalation set (– –): "what this should represent, this ugliness!?!?!? UI we see does not purpos or sense at all! it is messed up with an giant button!" + "Eye candy stunning funcitonal UI and UX MUST BE prooven with full automation tests which 'see' everything visually with no issues, false results or bluff of any kind!!! So far, anti-bluff policy seems completely ignored!" Every change to any user-facing UI surface (screen / component / view / widget / layout / theme on any platform – Android-Compose, iOS-SwiftUI/UIKit, web, desktop, TUI) MUST be proven by **device-independent host-side RENDERED PIXELS** before it may be claimed correct: the actual component rendered to an image ON THE HOST (NO device, NO emulator, NO running app required) via a host-render harness (Compose → Roborazzi/Paparazzi on the JVM; web → Playwright/Storybook snapshot; SwiftUI → snapshot-testing; equivalent per stack), for EVERY relevant `screen × state × {light,dark} theme` combination, with the produced PNG validated by BOTH (i) **image-diff against an approved golden** (per-pixel or perceptual-diff PASS/FAIL) AND (ii) an **OCR / vision oracle** that reads the rendered text + labels + control bounds to assert legibility + NO element overlap / NO label-over-label / NO clipping / NO off-screen control / NO collapsed-or-giant unbounded widget (the § . . layout invariants, here PROVEN on real pixels not asserted on tokens). **The forbidden bluff (forensic anchor, FACT – –)**: UI work was "validated" by VALUE-EQUALITY unit tests (asserting a hex colour string, an `sp / dp` size integer, a design-token field equals an expected constant) that NEVER RENDERED A SINGLE PIXEL – every such test was GREEN while the operator opened the running screen and saw a broken, unbounded giant-button layout. A value-equality / token-equality / property-

assertion UI test is therefore FORBIDDEN as the *proof* a UI is correct (it may SUPPLEMENT – confirming a token's wiring – but NEVER SUBSTITUTE for the rendered-pixel proof); a "green" UI suite that contains no rendered-pixel artefact for the changed surface is a § . PASS-bluff at the visual layer, severity-equivalent to a metadata-only PASS. The mandate (ALL hold): () **rendered-pixel proof per screen × state × theme**, device-independent, runnable in CI on EVERY commit with zero hardware (that is the whole point of host-side rendering – it is NOT subject to device availability); () **dual oracle** – golden image-diff AND OCR/vision content+layout oracle, both required, each PASS citing the captured PNG + diff artefact path per § . . /§ . . ; () **value-only UI claims FORBIDDEN as proof** – a UI change with no rendered-pixel artefact for its surface is a finding (§ . . test-interrupt → render → fix → re-render GREEN); () **self-validated analyzer** – the image-diff + OCR oracle is golden-good/golden-bad calibrated per § . . () (a golden-bad fixture seeded with overlapping labels / a clipped control / an unbounded button MUST FAIL, else the analyzer is itself a bluff gate), thresholds calibrated on the project's OWN fixtures not hardcoded from literature (§ . .); () **"device offline" is NEVER a valid skip** for this proof – host-render is precisely the device-independent path; an honest § . . SKIP applies ONLY where the platform genuinely has no host-render harness (a tracked migration item, never a silent gap or a value-only fallback). Honest boundary (§ . .): host-render proves the COMPONENT renders correctly, device-independently on every commit – it does NOT prove the running app on real hardware (that remains § . . /§ . . /§ . . /§ . . 's live-device recording + read-the-screen layer, which § . . COMPLEMENTS not replaces – the two layers are both mandatory: host-render for per-

commit device-free coverage, live-recording for end-to-end on-device confirmation), does NOT replace

\$. . 's OpenDesign token/theme governance

(\$. . PROVES the rendered RESULT of those tokens), and is a DISTINCT artifact class from

\$. . 's exported-document visual validation.

Classification: universal (\$. .) – the consuming project supplies its concrete host-render harness, golden corpus, screen×state×theme matrix, and OCR/vision

thresholds per \$. . . Composes \$. . /

\$. . 11 4 168 • 11 • 169 / \$. . / \$. . / \$. . /

\$. . / \$. . / \$. . /

\$. . / \$. . / \$. . /

\$. . / \$. . / \$. . /

\$. . / \$. . / \$. . Propagation gate

CM-COVENANT-114-170-PROPAGATION (literal 11.4.170) +

recommended gate CM-HOST-RENDERED-UI-VISUAL-PROOF (every

changed UI surface carries a device-independent host-rendered PNG per screen×state×theme, dual-oracle-

validated, run on every commit; a UI change proven only by value/token-equality assertions FAILs) + paired \$.

meta-test mutation (strip the literal → propagation gate

FAILs; accept a UI change whose only proof is a value-

equality assertion / a golden-bad overlapping-label

fixture that PASSes → CM-HOST-RENDERED-UI-VISUAL-PROOF FAILs;

gate-code = separate work item). **Canonical authority:**

constitution submodule Constitution.md \$. . .

Non-compliance is a release blocker. No escape hatch – no

--value-assertion-proves-ui , --skip-rendered-pixel-proof , --

token-equality-suffices , --device-offline-skip-visual , --

single-theme-only , --unvalidated-image-diff-OK , --no-ocr-

layout-oracle flag.

§ . Host-session safety – directly OR indirectly signing the user out is FORBIDDEN

Every script, test, helper, and AI agent governed by this Constitution MUST respect host-session safety. Non-compliance is a release blocker.

§ . Forbidden operations – directly OR indirectly

- . **Suspending the host:** `systemctl suspend` , `pm-suspend` , `loginctl suspend` , `DBus org.freedesktop.login1.Suspend` , `GNOME idle-suspend` , `lid-close handler` .
- . **Hibernating / hybrid-sleeping:** any `Hibernate` / `HybridSleep` / `SuspendThenHibernate` method.
- . **Logging out the user:** `loginctl terminate-session` , `pkill -u <user>` , `systemctl --user --kill` , anything that signals `user@<uid>.service` .
- . **Unbounded-memory operations** inside `user@<uid>.service` cgroup. Any single command expected to exceed ~ GiB RSS MUST be wrapped in a bounded execution scope (e.g. `bounded_run` from a project helper library that wraps `systemd-run --user --scope -p MemoryMax=...`).
- . **Programmatic rfcill toggles, lid-switch handlers, or power-button handlers** – these cascade into idle-actions.
- . **Disabling systemd-logind, GDM, or session managers** "to make things faster" – even temporary stops leave the system unable to recover.

§ . Required safeguards

Every script in this project that performs heavy work (build, transcription, model inference, large compression, multi-GB git operations) MUST:

- . Source the project's host-safety helper library at the top.
- . Call its pre-flight check and **abort if it fails**.
- . Wrap any subprocess expected to exceed ~ GiB RSS in a bounded execution scope so the kernel OOM killer is contained to that scope and cannot escalate to `user.slice`.
- . Cap parallelism (`-j`) to fit available RAM (estimate per-job-peak-RSS and divide into the budget).

§ . Container hygiene

Containers (Docker / Podman) the project owns or relies on MUST:

- . Declare an explicit memory limit (`mem_limit` / `--memory` / `MemoryMax`).
- . Set `OOMPolicy=stop` in their systemd unit to avoid retry loops.
- . Use exponential-backoff restart policies, never immediate retry.
- . Be clean-slate destroyed and rebuilt after any host crash or session loss so stale lock files don't keep producing failures.

§ . Memory-Budget Ceiling – % MAXIMUM

Forensic anchor – direct user mandate:

"First make sure that whatever we do through our procedures related to this project MUST NOT use more than % of total system memory! All processes MUST be able to function normally!"

The mandate. Project procedures MUST NOT use more than % of total system RAM. The remaining % is reserved for the operator's other workloads so the host can keep serving them while project work proceeds.

The protections:

- . `HOST_SAFETY_MAX_MEM_PCT` defaults to .
- . `HOST_SAFETY_BUDGET_GB` is computed at source-time from $\text{MemTotal} \times \text{MAX_PCT}/100$, in GiB.
- . Bounded execution scopes clamp `MemoryMax` down to the budget if the caller asks for more.
- . The build script's parallelism is computed as $\text{min}(\text{nproc}, \text{floor}(\text{budget_gb} / \text{per_job_peak_rss_gb}))$.
- . Heavy work MUST be wrapped in a bounded execution scope so the kernel OOM-kills only the scope – `user@<uid>.service` stays alive.

No escape hatch. § . has NO operator-facing override flag. The cap exists for the operator's own protection; bypassing it is the bluff the § . covenant specifically prohibits.

§ . Continuation document – sacred invariant

Forensic anchor – direct user mandate:

"during any work we perform, during Phases implementation, debugging and fixing, during ANY

effort we have the Continuation document MUST BE maintained and it MUST NOT BE out of sync with current work we are doing! If for any reason we stop our work, we MUST BE able to continue any time, with current work, exactly where we have left off and from any CLI agent or any LLM model we chose! Nothing can be broken or faulty in maintained Continuation document!"

The mandate. A single, canonical, machine-readable handoff document – `docs/CONTINUATION.md` – must always reflect the live state of the project. Any agent (human, Claude Code, Cursor, Aider, Codex, Gemini CLI, any future LLM) must be able to resume work **exactly where the previous session left off** by reading this single file.

Mandatory protections:

- . `docs/CONTINUATION.md` **MUST exist** at the project root. Its absence is a release blocker.
- . **Every non-trivial state change** MUST update this document in the same commit as the work itself.
- . **Top-of-file timestamp** is updated on every edit. Stale timestamps trigger gate failure.
- . **Section § "Active work"** lists every IN PROGRESS / BLOCKED item with enough detail that any agent can resume without conversation context – concrete commands, file paths, monitor IDs.
- . **Section § "How to use this document"** contains the verbatim resumption prompt – a single block any operator can paste into any CLI agent.
- . **Document MUST be self-contained.** No hyperlinks to ephemeral external systems as the only source of truth.

No escape hatch. § . has NO operator-facing override flag for the existence requirement. The discipline exists for the operator's own protection.

`$ .` – Maximal dynamic resource utilization for containerized builds (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Use full workstation system capacity – all resources to build and test! Give maximal number of cores and memory! Use dynamic capabilities! No limiting to -j anymore! ... MUST NOT break the system, create bottlenecks, or get stuck! ... MUST BE now part of constitution Submodule! On each System we use, aim for maximal use of available resources for maximal efficiency!"

The reconciliation (the load-bearing logic). The existing memory-and-parallelism caps were born from a specific, forensically-captured failure mode: **heavy build work running inside the user login session's cgroup** (`user@<uid>.service` / `user.slice`) overran memory and the kernel OOM-killer escalated to the session itself, **SIGKILLing the operator's desktop, terminals, and the agent** (the `$ .` forensic anchor: "three consecutive session-loss SIGKILLs" during a `bare-metal m -j5`; consuming-project instantiations extend this – a bare-metal AOSP `m -j` was hard-capped at `-j` because `soong_build` alone holds a large constant RSS inside `user.slice`). Those caps – `$ . 's %-of-total-RAM ceiling` and the consuming project's bare-metal `-j` cap – are CORRECT and REMAIN in force for any heavy work that executes inside `user@<uid>.service` / `user.slice`, because such work shares the session's OOM fate.

A build that runs in its **OWN container cgroup with its OWN MemoryMax** (the rootless-Podman containerized build path –

generically "the sanctioned containerized build") is a **fundamentally different accounting domain**: it is NOT a child of `user.slice`, so the kernel OOM-killer and any userspace OOM watchdog (systemd-oomd's `user.slice` PSI watch) act on the **container's own cgroup FIRST** – the container OOM-kills ITSELF, it does not escalate to and cannot SIGKILL the operator session. Therefore a high memory / CPU / `-j` envelope in the containerized path is **host-SAFE by construction** – the exact hazard the `$` . ceiling guards against does not reach the session across the cgroup boundary.

12 4 00

The mandate. For the **containerized build path only**, the project MUST use the **maximal SAFE fraction** of the host's real capacity, computed **dynamically per host** – NOT a fixed low `-j` count, NOT clamped to the `$` . % figure (which governs `user.slice`-resident work, a different domain). Concretely (ALL hold): () **Auto-detect host capacity, never hardcode** – read this host's real `nproc`, `MemTotal`, `SwapTotal`, and the hard process/thread rlimit (`RLIMIT_NPROC` / `ulimit -Hu`) at launch; every number driving the envelope MUST be measured on the running host, never assumed (`$` . . no-guessing); a launcher that prints a build envelope it did not measure is a `$` . . violation. () **Reserve explicit host headroom (cores AND RAM)** – the maximal-safe fraction is "all capacity MINUS a reserved slice kept free for the operator's live session + other workloads"; the headroom MUST be explicit and captured in the printed envelope; this preserves the SPIRIT of `$` . (the operator's session keeps functioning) via a reserved-headroom mechanism, WITHOUT importing `$` . 's % number – the container's cgroup, not a `user.slice` budget line, bounds it. () **`-j` scales nproc-awarely, not just memory-awarely** – the parallel-job count MUST respect BOTH the memory governor (`floor((container_mem_gb - build_constant_overhead_gb) / per_job_peak_rss_gb)`) AND the

process/thread rlimit, because each `-j` unit spawns compiler/linker worker THREADS that count against `RLIMIT_NPROC` and exceeding it causes `EAGAIN` "fork: retry" storms that stall/wedge the build (the "never get stuck" clause); where the login's hard `nproc` is too low for the full memory-and-CPU-derived `-j`, the launcher MUST either (a) run at the proven-safe lower `-j` under the rlimit, OR (b) on a privileged path raise the rlimit (e.g. `prlimit`) and drop back to the unprivileged uid (e.g. `setpriv`) before launching so the container keeps its rootless identity; the launcher MUST print which mode it chose (`$ . . .` – the envelope is captured evidence). ()

The caps REMAIN for the bare-metal / user.slice path –
`$ . 's % ceiling` and the consuming project's bare-metal `-j` cap are NOT relaxed for any work inside `user@<uid>.service`; `$ . . .` grants the maximal envelope ONLY across the container-cgroup boundary; a bare-metal heavy build is still forbidden from claiming the `$ . envelope` (that would reproduce the original session-loss incident); `$ .` does NOT weaken `$ . ,` it SCOPES it. () **Never break, never bottleneck, never get stuck** – the envelope MUST leave the reserved host headroom genuinely free (verifiable post-launch: session memory-available stays above the reserved floor), MUST NOT drive swap-thrash inside the container (bound `MemorySwapMax` so the container hits its own OOM before pathological swap), and MUST NOT exceed the process rlimit (per clause); "maximal" means maximal-SAFE, not maximal-reckless – an envelope that wedges the host is a `$ . violation`, not compliance. () **Evidence-based tuning, no guessing** (`$ . . + $ . .`) – the chosen envelope is driven by captured host metrics + the build's own resource profile (`$ . . build-resource stats`); "use a bigger `-j`" without captured headroom evidence proving the host can absorb it is a

\$. . + \$. () violation; when in doubt, the reserved-headroom floor wins.

Reference implementation (consuming project). A thin DYNAMIC wrapper over the sanctioned containerized build (ATMOSphere: scripts/build_maxres.sh over scripts/build_containerized.sh) reserves a fixed core count + a RAM floor (e.g. % of MemTotal to the container, clamped so a fixed-GiB slice stays free for the session), computes `JOBS_MEM = floor((container_mem_gb - constant_overhead_gb) / per_job_peak_rss_gb)`, and governs `-j` against the real hard `nproc : unprivileged` → the proven-no-EAGAIN job count; privileged (root re-exec raises `RLIMIT_NPROC` via `prlimit` then `setpriv` -drops back to the unprivileged uid) → full `-j = min(cpus, JOBS_MEM)`; `--dry-run` prints the full envelope and launches nothing. This is the \$. . project instantiation; the mechanism above is universal.

Composes with \$. / \$. / \$. (host-session safety – \$. is a strengthening within, not a relaxation of, host safety) / \$. (the % ceiling – \$. SCOPES it to user.slice-resident work, does NOT weaken it) / \$. (the launcher's envelope is a live-state anchor for the resumption file) / \$. . (no-guessing – every envelope number measured) / \$. . (build-resource stats feed the tuning) / \$. . (target/hardware safety – the reserved-headroom + rlimit governors are the host-safety equivalent) / \$. (the containerized path is the non-destructive default). The consuming project supplies its containerized-build path, cgroup mechanism, and per-host detection per \$. . .

Classification: universal (\$. .) – the reconciliation logic (caps bind user.slice-resident work; the containerized cgroup is separately accounted so it may claim the maximal safe fraction) is platform-agnostic and true on every host. Propagation gate **CM-COVENANT-12-11-**

PROPAGATION (literal anchor 12.11 present across the consumer fleet's CLAUDE.md / AGENTS.md / QWEN.md / GEMINI.md) + recommended gate **CM-MAXRES-DYNAMIC-BUILD** (the containerized-build launcher (a) auto-detects nproc / MemTotal / RLIMIT_NPROC , (b) reserves explicit host headroom, (c) governs -j against BOTH the memory model AND the process rlimit, (d) prints the chosen envelope, (e) does NOT hardcode a fixed low -j for the containerized path AND does NOT relax \$ for any user.slice path) + paired \$ meta-test mutations (strip the 12.11 literal → propagation gate FAILs; replace the dynamic detection with a hardcoded -j2 in the containerized launcher OR strip the rlimit-aware -j governor → CM-MAXRES-DYNAMIC-BUILD FAILs; gate-code = a separate tracked work item).

Canonical authority: constitution submodule Constitution.md § . . Non-compliance is a release blocker regardless of context. No escape hatch – no --fixed-jobs , --no-dynamic-detection , --ignore-host-capacity , --relax-user-slice-cap , --bare-metal-maximal flag exists.

§ . – Process/thread-limit (RLIMIT_NPROC) awareness for parallel subagent/multi-process work (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –):

"Add these ulimit details and instructions to the constitution so they are mandatory for this and all projects!"

The limit. Heavy parallel subagent / multi-process work is bounded by the OS per-user process/thread limit ulimit -u

(`RLIMIT_NPROC`), which on Linux counts ALL threads AND processes owned by the real UID – NOT a per-process limit. The ambient tooling fleet already running in a session (MCP servers, language runtimes – tokio/V /libuv worker pools, `mongod` / `chrome` / `prisma` -class MCP servers, local LLM servers) consumes a large FIXED fraction of this budget before any project work even starts. **FORENSIC FACT**

(`ulimit -u 4096`) : a host configured with `ulimit -u 4096` sat at ~ threads from the ambient MCP fleet alone, leaving under threads of headroom; dispatching a th Go-heavy subagent (the Go runtime sets `GOMAXPROCS` to `nproc`, spawning dozens of OS threads per invocation) hit `EAGAIN` / failed to create new OS thread (`errno=11`) and crashed tooling. Stopping three non-essential MCP servers (`mongod` / `chrome` / `prisma`) freed approximately threads (→), restoring headroom.

Mandate – treat thread-limit exhaustion as a \$ host-safety event. BEFORE scaling to N parallel subagents / processes, the agent MUST check thread headroom and treat thread-limit exhaustion as a \$ host-safety event that YIELDS UNCONDITIONALLY – the same standing as \$ (memory) and \$ (AOSP -j cap). This is an ORTHOGONAL exhaustion axis to \$'s memory ceiling: a host can have abundant free RAM while its thread/process budget is fully exhausted, and vice versa – both MUST be checked independently before scaling parallelism.

Diagnosis (mandatory before scaling). `ulimit -u` (soft limit) + `ulimit -Hu` (hard limit) + the LIVE count `ps -L --no-headers -u "$USER" | wc -l` (threads, not just processes – `RLIMIT_NPROC` counts threads); headroom = soft limit minus live count. **Signals of exhaustion** (never guessed, always captured per \$. .): `EAGAIN`, `fork: retry: Resource temporarily unavailable`, `failed to create new OS thread`, or `cannot allocate memory` returned from a `fork` / `clone` / `thread-`

spawn call – this is DISTINCT from OOM (memory can be abundant while the thread/process budget is exhausted; the diagnostic MUST NOT conflate the two). When headroom is low, the agent MUST SERIALIZE (run fewer concurrent subagents/processes) and avoid `Go-heavy /`

`GOMAXPROCS=nproc -class spikes` – dispatching parallel work blind at near-limit headroom is forbidden; a subagent that hits `EAGAIN` mid-operation (e.g. `mid-git push`) risks incomplete/corrupted state (`$ .` data-safety risk).

Raising the limit – three techniques (cite the mechanism used, never guess which applies). (a) **No-restart, on a running process:**

`prlimit --pid <PID> --nproc=<soft>:<hard>` (requires root – e.g. via `su -c`) to raise a HARD limit above the process's current hard ceiling; a running process's new children inherit the raised limit going forward.

(b) **Persistent, system-wide:** `/etc/security/limits.conf` (`<user> hard nproc <N>` + a matching `soft` line) or a systemd unit's `LimitNPROC=<N>`, taking effect on the NEXT login/session (not the current one).

(c) **Self-raise, within the current session:** any process may raise its own SOFT limit up to its existing HARD limit via `ulimit -u <N>` where `N ≤ hard` – it can NEVER exceed the hard ceiling without root-assisted method (a).

Caveat (`$ . .` – do not assume scope): `su -c "ulimit -u N"` sets the limit ONLY inside that transient `su` subshell – it does NOT retroactively affect an already-running session; a running process keeps the `RLIMIT_NPROC` value captured at its own launch time.

Freeing threads – operator-authorized only

(`$ . . .`). Stopping non-essential MCP servers / background tooling frees the ambient fraction of the budget consumed before project work starts – but this REQUIRES explicit operator authorization per

`$ . . .` (no-silent-removal); the agent MUST NOT autonomously kill the operator's tooling without asking

first. **CAUTION – self-kill hazard (forensic FACT):** `pkill -f <pattern>` / `pgrep -f <pattern>` matches against the FULL COMMAND LINE, including the killer's own invocation if the search pattern string happens to appear in it – this can cause an unintended self-kill of the very process issuing the kill command. Any cmdline-pattern-based kill list MUST explicitly exclude the caller's own `$$` / `$PPID` (and the conductor's PID, when dispatched as a subagent) before executing.

Composes with `$` (host-session safety umbrella) / `$` . (memory-budget ceiling – the ORTHOGONAL second exhaustion axis alongside this thread/process-count axis) / `$` . (bare-metal AOSP `-j` cap – thread exhaustion is a DISTINCT failure mode from the memory-driven cap that section addresses) / `$` . . (parallel-development PWU pipeline – the concurrency this section bounds) / `$` . . (continuous parallel-stream working routine – the $\geq$ -background-stream target is bounded by available thread headroom, not merely memory) / `$` . . (autonomous-decision-over-blocking – safe, reversible decisions, e.g. serializing dispatch, are made without operator involvement) / `$` . . (no-silent-removal-of-existing-components – freeing ambient tooling threads requires operator authorization) / `$` . . (absolute data safety – a subagent crashed mid-git-operation by `EAGAIN` is a data-safety risk this section exists to prevent).

Classification: universal (`$` . .) – `RLIMIT_NPROC` / `ulimit -u` is a Linux/POSIX kernel-level per-UID accounting mechanism, platform-agnostic and true on every host running this Constitution's parallel-subagent/multi-process disciplines; the consuming project supplies its concrete ambient-tooling roster and `ulimit` values per `$` . . . Propagation gate `CM-COVENANT-12-12-PROPAGATION` (literal anchor `12.12` present

across the consumer fleet's CLAUDE.md / AGENTS.md / QWEN.md / GEMINI.md) + recommended gate CM-NPROC-HEADROOM-CHECK (parallel-dispatch orchestration checks ulimit -u + live thread count BEFORE scaling concurrency; a cmdline-pattern kill helper excludes the caller's own PID) + paired § . meta-test mutation (strip the 12.12 literal → propagation gate FAILs; strip the headroom-check call before a parallel-dispatch scale-up OR remove the self-PID exclusion from a kill helper → CM-NPROC-HEADROOM-CHECK FAILs; gate-code = a separate tracked work item).

Canonical authority: constitution submodule Constitution.md § . . Non-compliance is a release blocker regardless of context. No escape hatch – no --ignore-thread-limit , --dispatch-blind , --skip-headroom-check , --autonomous-tooling-kill , --assume-ulimit-scope flag exists.

Appendix A – Mutation testing – academic and industrial foundations

The anti-bluff / paired-mutation policy in § . is rooted in the **mutation testing** literature. References that the consuming project SHOULD cite when explaining its meta-test harness:

- Jia, Y. & Harman, M. (). *An Analysis and Survey of the Development of Mutation Testing*. IEEE Transactions on Software Engineering, ().
- DeMillo, R.A., Lipton, R.J., Sayward, F.G. (). *Hints on Test Data Selection: Help for the Practicing Programmer*. IEEE Computer.
- Open-source mutation testers worth studying:
 - PIT (Java) – <https://pitest.org/>

- **Stryker** (JS / C / Scala) – <https://stryker-mutator.io/>
- **Cosmic Ray** (Python) – <https://github.com/sixty-north/cosmic-ray>
- **mutmut** (Python) – <https://mutmut.readthedocs.io/>
- **mull** (LLVM-IR) – <https://mull.readthedocs.io/>

The § . paired-mutation policy is a **lightweight** variant of mutation testing applied at the gate-assertion layer rather than at the production-code layer. It does not need a mutation framework; it needs ONE mutation per gate that proves the gate catches the break. This is computationally cheap ($O(\text{gates})$, not $O(\text{production code lines})$) and produces a strong "bluff-immunity" guarantee.

Appendix B – Recursive inheritance & path-independence

Projects that include this constitution submodule MUST tolerate **arbitrary submodule depth**. Between the main project root and any consuming submodule there may be N intermediate levels. To locate this constitution submodule from any depth, every project SHOULD provide a helper that walks up parents until it finds `constitution/Constitution.md` (the file you are reading right now) OR follows `git rev-parse --show-superproject-working-tree` recursively. The constitution submodule ships `find_constitution.sh` for exactly this purpose; nested submodules can source it without knowing their own depth.

§ . . – **Mandatory comprehensive human-readable workable-item descriptions** (User mandate, – –). Every workable item (ATM-NNN ticket) in the project's workable-items database, `Issues.md`,

Fixed.md, Issues_Summary.md, Fixed_Summary.md, and ALL derived documents MUST carry a comprehensive human-readable description that explains the item in plain, non-technical language suitable for non-developers (project managers, stakeholders, team leads, business partners). The description MUST be - sentences minimum, written so that ANY person – regardless of technical background – can understand: (a) WHAT the item is (the feature, fix, or task in simple terms); (b) WHY it matters (the user/business value); (c) HOW it works (the mechanism, without code references); (d) WHO benefits (the end user or stakeholder); (e) WHAT the expected outcome is (the measurable result). The description is MANDATORY – an item without a human-readable description is a § . violation at the documentation layer. Descriptions MUST NOT use jargon, acronyms without expansion, code references, or technical implementation details as the primary explanation. Technical details MAY appear as supplementary context AFTER the plain-language explanation. The description column in the SQLite database (docs/workable_items.db) is the SINGLE SOURCE OF TRUTH for item descriptions; all derived documents (Issues.md, Fixed.md, summaries, exports) MUST carry the SAME description text. Pre-build gate CM-HUMAN-READABLE-DESCRIPTION (when implemented) walks every item in the DB and asserts the description field exists, is non-empty, and is ≥ characters. Paired § . meta-test mutation strips a description → gate FAILs.

Classification: universal (§). Composes § . . (item-status tracking), § . . (item-type tracking), § . . (column alignment), § . . (summary clarity), § . . (SQLite SSoT), § . . (item integrity), § . . (universal export). Propagation gate CM-COVENANT-114-171-PROPAGATION (literal 11.4.171) + paired § . mutation.

§ . . . – **Mandatory production-readiness planning with realistic timeline projection** (User mandate, - -). Every project under this Constitution MUST maintain a living production-readiness planning document that: (a) provides REALISTIC timeline projections for all product milestones (MVP, beta, production-ready) based on actual development velocity data (items completed per week, reopening rate, blocking dependencies); (b) identifies ALL danger zones, weaknesses, and risk areas with mitigation strategies; (c) accounts for HW delivery timelines, licensing delays, and external dependency risks; (d) defines the critical path and identifies which items can slip without affecting MVP; (e) includes workstation/build-capacity analysis and optimization recommendations; (f) is updated at least monthly or whenever the item count changes by $\geq . . . \%$; (g) is cross-referenced against the workable-items database for accuracy. The planning document MUST live at docs/research/production_planning_<YYYYMMDD>/ANALYSIS.md with HTML+PDF exports per § Honest boundary (§ . . .): projections are estimates based on measured velocity – they guarantee the methodology is sound, NOT that specific dates will be met. Classification: universal (§ . . .). Composes § . . . (no-guessing – projections based on data, not wishful thinking), § . . . (full retest before tag), § . . . (iteration discipline), § . . . (SQLite SSoT), § . . . (four-layer verification). Propagation gate CM-COVENANT-114-172-PROPAGATION (literal 11.4.172) + paired § . . . mutation.

§ . . . – **Containerized + distributed build mandate** (User mandate, - -). EVERY build of EVERY component (source compile, artifact/package/installer/container-image production, codegen, asset render – for ANY language/platform: Go, Android/Gradle, desktop, web, native, firmware) MUST run INSIDE a

specialized build container provisioned via the `digital.vasic.containers` submodule (`$ . . .`) – NEVER on the bare host. The build containers MUST be DISTRIBUTED to the designated remote build host(s) (e.g. `thinker.local`) via the SAME containers-submodule distribution mechanism the infra uses (`$ . . . Distributor / remote compose over SSH, $ . . . rootless`), so the build EXECUTES on the remote build host (offloading the developer/main host); once the build completes the produced artifacts MUST be brought BACK to the originating main host (scp/rsync/volume copy) for use/ flashing/distribution. Building outside a container, or on the bare host, is FORBIDDEN – a release blocker (the "works on my machine" / unreproducible-build class `$ . . .` exists to prevent). The `build-host` target + build-container definitions are config-injected (`$ . . .`, never hardcoded in the submodule); a missing build-container capability is added by EXTENDING the containers submodule upstream (`$ . . .`), never by an ad-hoc host build. Honest boundary (`$ . . .`): the containerized+distributed build guarantees reproducibility + host-isolation + capacity offload – it does NOT replace the `$ . . .` full-suite retest, `$ . . .` four-layer artifact→runtime verification, or `$ . . .` installable-asset evidence (those still run against the brought-back artifact). Classification: universal (`$ . . .`). Composes `$ . . .` (containers submodule – sole orchestration layer), `$ . . .` (rootless runtime), `$ . . .` (extend-don't-reimplement), `$ . . .` (config injection / decoupling), `$ . . .` (build-resource stats), `$ . . .` (iteration speedup – persistent caches in the container), `$ . . .` (no-commit-while-build-writes-artifacts), `$ . . .` (installable-asset evidence on the brought-back artifact), `$ . . .` (artifact→runtime verification),

\$. (host memory ceiling – offloading the build preserves the main host). Propagation gate CM-COVENANT-114-173-PROPAGATION (literal 11.4.173) + recommended gate CM-CONTAINERIZED-DISTRIBUTED-BUILD (every build runs via the containers submodule on the remote build host; a bare-host build is detected + FAILs) + paired \$. mutation. **Canonical authority:** constitution submodule Constitution.md \$. . . Non-compliance is a release blocker. No escape hatch – no --build-on-host , --skip-container-build , --local-build-ok , --no-distributed-build , --bare-host-build flag.

\$. . – Shared-host process-ownership verification mandate (User mandate, – –). Verbatim operator mandate: "Make sure that processes we are checking are always OURS! Add this as mandatory constitution rule / constraint, since on every host there may be other work being done!" On ANY host where concurrent work from other projects / users / agents may run (the normal case – a shared developer workstation), EVERY inspection of, conclusion about, or action on a process, build, daemon, port, lock-file, or system-state signal MUST FIRST positively VERIFY the target is OURS (this project's) before drawing a diagnosis or acting on it. A diagnosis based on a process that is actually another project's work is a \$. . guessing-class error; KILLING / signalling / restarting / altering a process that is NOT positively ours is a \$ + \$. . -class safety violation (operator-workload destruction) – forbidden. Forensic anchor (FACT, – –): during an ATMOSphere build-stall diagnosis, a %-CPU java + gradlew on the shared host were nearly mis-attributed to our SmartTube build; they were in fact a separate catalogizer project's gradle test (cwd=Projects/catalogizer), and a naive pgrep java / gradle would have matched them as "ours." The mandate (ALL hold): (a) **positive-ownership discriminator REQUIRED**

– attribute a process to us ONLY via a strong signal: its `cwd` resolves inside THIS project's tree, its `argv` names this project's scripts / artifacts / repo path, a PID we ourselves launched and recorded (the `$ . .` launched-PID registry), or a lock / port / socket path under our tree. A loose name match (`pgrep java` / `gradle` / `node` / `python` / `podman` / `zig`) is NEVER sufficient – it matches every project on the host. (b) **NEVER kill / signal / restart a process not positively ours** – other projects' gradle / java / kotlin daemons / podman pods / build processes are strictly off-limits (operator-workload-safety); the `$ . . remediation pkill -f GradleDaemon` MUST be scoped to OUR daemons (`cwd` / `argv` under our tree), never a blanket host-wide kill. (c) **contention → surface, don't break** – when our work waits on host contention caused by another project's process (e.g. our `$ . .` build gate sleeping because another project's gradle daemon is alive), surface it per `$ . . / $ . .` (the safe blocked-decision) and keep progressing non-blocked work; do NOT kill the other project's process to unblock ours. (d) **filters exclude self + others** – every `pgrep` / `ps` filter MUST exclude the checking command's OWN `argv` (the self-match that produced the false "still running" restyle watcher bug, `- -`) AND other-project matches; prefer matching on our absolute repo path. (e) **applies to all state, not just processes** – ports, lock files, gradle/maven caches, podman pods, tmux sessions, temp files on a shared host all belong to potentially-other work; verify ownership before reading-as-ours or mutating. Honest boundary (`$ . .`): ownership-verified inspection prevents mis-attribution + cross-project damage – it does not itself prove our process is healthy (that still needs our own artifacts: our build log, our output files, our launched-PID). Composes `$ . .` (no-guessing) / `$ . .` (autonomous

safe blocked-decision – block-don't-break) / \$. .
 (surface contention) / \$. . (launched-PID
 registry as the ownership source) / \$ + \$.
 (host-safety, now scoped to our daemons) / \$. .
 (don't break other workloads). Classification: universal
 (\$. .). Propagation gate CM-COVENANT-114-174-
 PROPAGATION (literal 11.4.174) + recommended gate CM-
 PROCESS-OWNERSHIP-VERIFIED (process/daemon-killing or build-
 state-diagnosing scripts attribute by cwd/argv/launched-
 PID, never a bare name match; pgrep filters exclude self)
 + paired \$. mutation (replace an ownership-verified
 attribution with a bare pgrep <name> → gate FAILs; strip
 the literal → propagation gate FAILs; gate-code = separate
 work item). Non-compliance is a release blocker. No escape
 hatch – no --assume-process-ours , --loose-pgrep-ok , --kill-
 any-matching-process , --skip-ownership-check flag.

\$. . – Conflict-free multi-track work-division
 +, exactly-once claim registry + capability-aware deadlock-
 proof device-lock (User mandate, – –).
 When a project runs multiple concurrent work-STREAM tracks
 (\$. . feature work-stream lifecycle;
 \$. . /\$. . parallel streams) that
 share (a) a common pool of workable items and (b) a scarce
 pool of exclusive hardware/resources, TWO cooperating
 arbitration layers are MANDATORY – both conflict-free BY
 CONSTRUCTION, not by convention. Forensic anchor –
 verbatim user mandate (– – , REM-):
 "MUST prevent multiple Tracks working the same workable
 item OR the same group of workable items resembling one
 consistent logical group of functionalities. Division of
 work assigned to Tracks MUST be conflict-free: no
 conflicts, misunderstandings, race conditions, data
 corruptions, codebase loss. Risk-free + safe; results
 ALWAYS produced; no false results, no bluff." (REM-):
 "there MUST exist proper synchronization mechanism for
 flashing and use of device ... lock ... release ...

available test devices pool and information about held locks ... dead-lock-proof ... a track locks one OR more devices selected by required capability." (A) **Work-item layer – exactly-once claim registry** (REM-). Every workable item, OR every logical GROUP of items that share files / subsystem / feature / schema (a hidden-coupling cohort), is assigned to AT MOST ONE track at a time. () **Exactly-once** – a lock-arbitrated, append-only claim registry (extending \$. . . no-work-loss registry + \$. . . real-time sync substrate) records `claim(item|group, track, ts, ttl)` ; a track MUST claim-before-touch; a second claim on an already-claimed item is DENIED (no silent overlap). () **Logical-group cohesion** – items that share a file-scope / subsystem / feature / schema are one group claimed WHOLE by one track (prevents the hidden-coupling race two tracks editing the same seam would create – \$. . . L disjoint-file-scope is the enforcement). () **Conflict-free by construction** – disjoint file-scope partition (\$. . . L) + per-track worktree/CoW isolation (\$. . .) means two tracks physically cannot touch the same bytes. () **No codebase loss** – \$. . . hardlinked backups, \$. . . ff-only merge-onto-latest-main (NEVER force-push), \$. . . working-tree quiescence, \$. . . merge-first.. () **Anti-bluff** – the assignment + the merge are verifiable (registry event stream + captured merge evidence); results are ALWAYS produced (a claimed-but-abandoned item's claim TTL-expires and returns to the pool, \$. . .). (B) **Hardware layer – capability-aware deadlock-proof device-lock** (REM-). When N tracks outnumber M exclusive test devices, exactly ONE track owns each device at a time (\$. . . single-resource-owner). A track leases ONE OR MORE devices selected by required CAPABILITY (not merely identity) – the pool records per-device capabilities; leases are capability-matched; multi-device

blocking + TTL-reap stress/chaos proof with a paired
 \$. mutation) + paired \$. meta-test mutations
 (gate-code = separate work item). **Canonical authority:**
 constitution submodule Constitution.md \$. . .
 Non-compliance is a release blocker. No escape hatch – no
 --allow-cross-track-overlap , --skip-claim-registry , --split-
 logical-group , --blocking-device-acquire , --allow-partial-
 device-hold , --tune-resources-by-guess flag.

\$. . – **Developer-tooling project-decoupling + invocation-directory operation** (User mandate,
 – –). No project-specific script, hook,
 alias, or binary may be wired – copied, symlinked, PATH-
 exported, or aliased – into a global or shared developer-
 tooling PATH; shared tooling **MUST** be project-agnostic and
 operate on the INVOCATION directory (the cwd, or an
 explicit path/config argument), NEVER on a hardcoded
 project path. A project-specific helper lives in and runs
 from its own repo; a genuinely reusable helper is promoted
 to the shared toolkit **ONLY** after it is stripped of every
 project-specific assumption (paths, device serials,
 package names, region endpoints) and takes its target from
 the invocation context. A project script reachable on the
 global/shared PATH is a decoupling violation of equal
 severity to importing project-specific code into a shared
 submodule (\$. .). The forensic case (FACT,
 – –): a project cwd-hook symlinked into the
 global Claude-Toolkit PATH re-coupled every alias to one
 hardcoded checkout. Honest boundary (\$. . .):
 decoupling guarantees the tool is project-agnostic – it
 does NOT by itself prove the tool is correct (that still
 needs the tool's own tests). Composes \$. . .
 (submodules-as-equal-codebase decoupling) / \$. .
 (naming) / \$. . (canonical-root layer split) /
 \$. . (file-layout discipline) / \$. .
 (no-guessing). Classification: universal (\$. .) –
 the consuming project supplies its concrete tool paths +

invocation-config mechanism per § 11.4.177. Propagation gate CM-COVENANT-114-177-PROPAGATION (literal 11.4.177) + recommended gate CM-TOOLING-PROJECT-DECOUPLED (no tracked project script is symlinked/PATH-wired into a global toolkit dir; shared tools take their target from cwd/arg/config, never a hardcoded project path) + paired § 11.4.178. meta-test mutation (wire a project script into the shared PATH, or hardcode a project path into a shared tool → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md § 11.4.179. Non-compliance is a release blocker. No escape hatch – no --global-symlink-ok, --hardcode-project-path, --wire-into-shared-path flag.

§ 11.4.180 – Track-qualified identity for parallel work streams (session-name collision ban) (User mandate, § 11.4.181). Parallel work streams that share hardware, a git backing store, a tmux/session namespace, a lock namespace, or any single-owner resource MUST be addressed by a TRACK-QUALIFIED identity (<project>__<track>__<role>), NEVER a bare directory basename. Multiple checkouts/clones/worktrees sharing basename atmosphere collapse into one session/lock/log key and cross-wire tmux targets, lock files, device claims, and log directories (the observed collision, § 11.4.182). Every registry row, session name, lock path, log directory, and device claim MUST carry the track-qualified key; a bare-basename identity for a resource with ≥ 2 concurrent claimants is a violation. Honest boundary (§ 11.4.183): a unique identity prevents cross-wiring – it does not itself arbitrate ownership (that is § 11.4.184 + § 11.4.185 (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor)). Composes § 11.4.186 (resolve-by-stable-name-not-by-index) / § 11.4.187 (single-resource-owner partitioning) / § 11.4.188 (the multi-

track work-division + exactly-once-claim + device-lock arbitration anchor) / \$. . (naming) / \$. . (parallel PWU). Classification: universal (\$. .) – the consuming project supplies its <project> / <track> / <role> vocabulary per \$. . .

Propagation gate CM-COVENANT-114-178-PROPAGATION (literal 11.4.178) + recommended gate CM-TRACK-QUALIFIED-IDENTITY (every session name / lock path / log dir / device claim for a multi-claimant resource carries the track-qualified key, never a bare basename) + paired \$. meta-test mutation (address a shared resource by bare basename → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Multi-project-per-track layout clarification (- -): a /mnt/track<N> track parent is NOT single-project – MULTIPLE projects coexist on the same track parents, each in its OWN project-named subdirectory /mnt/track<N>/<project_name>/ (e.g. /mnt/track1/atmosphere/ AND /mnt/track1/helix_ota/ side by side); the per-track working copy of project P on track N is ALWAYS /mnt/track<N>/<project_name>/ (the project's canonical lowercase snake_case dir-name per \$. .), never the bare track root and never a shared basename. This is precisely WHY the identity is <project>__<track>__<role>... and carries the <project> component: it keeps two coexisting projects' track-N sessions / locks / logs / device-claims / registry rows from colliding. A project adopting this constitution + the multi-track engine (\$. .) MUST create AND register its own .. /mnt/track<N>/<project_name>/ per track under its project-qualified identity – it MUST NOT assume it owns the track root or that a track hosts only one project.

Canonical authority: constitution submodule Constitution.md \$. . . Non-compliance is a release blocker. No escape hatch – no --bare-basename-ok , --skip-track-qualifier , --shared-session-name flag.

§ 11.4.179 . . . – Corruption-isolated parallel git streams (own-.git independent clones, not shared-common-dir worktrees) (User mandate, 11.4.179). When corruption-isolation is a requirement, parallel git work streams MUST each be an isolated repository with its OWN .git (own object store, own index, own lock namespace), NOT ..git worktree checkouts sharing one common .git. A shared common-dir is a single point of failure: one stale index.lock / .commit_all.lock / .push_all.lock freezes EVERY stream at once (the observed -h all-track freeze, 11.4.179), and one corrupted object store loses every stream. Each stream's git rev-parse --git-common-dir MUST resolve INSIDE the stream's own tree; a common-dir resolving to a shared parent is a violation. Where CoW/reflink disk-feasibility is the constraint, use the § 11.4.179 . . . reflink-clone-with-own-.git path – a btrfs/XFS/ZFS/APFS reflink shares extents on disk while keeping a per-stream .git, so isolation and disk-feasibility are BOTH satisfied. Honest boundary (§ 11.4.179 . . .): own-.git isolation contains lock/corruption blast radius – it does NOT remove the need for per-repo stale-lock reaping (§ 11.4.179 . . .) nor for ff-only no-force integration (§ 11.4.179 . . .). Composes § 11.4.179 . . . (feature work-stream lifecycle / CoW clone) / § 11.4.179 . . . (single-resource-owner) / § 11.4.179 . . . (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor) / § 11.4.179 . . . (data safety) / § 11.4.179 . . . (parallel PWU). Classification: universal (§ 11.4.179 . . .) – the consuming project supplies its reflink mechanism + clone layout per § 11.4.179 Propagation gate CM-COVENANT-114-179-PROPAGATION (literal 11.4.179) + recommended gate CM-GIT-STREAMS-OWN-COMMON-DIR (every registered parallel stream's git rev-parse --git-common-dir resolves inside its own tree, never a shared parent) + paired § 11.4.179 . . . meta-test mutation (register a stream whose common-dir points to a shared

parent → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md`

\$. . . Non-compliance is a release blocker. No escape hatch – no `--shared-common-git`, `--worktree-when-isolation-required`, `--skip-git-isolation` flag.

\$. . . – Commit/push (single-writer) wrappers **MUST** auto-reap provably-stale locks before acquiring (User mandate, – –). Every commit / push / single-writer wrapper **MUST**, before acquiring its own lock, auto-reap any PROVABLY-stale git lock – a lock whose recorded holder PID is DEAD (`kill -0` fails), OR (no PID recorded) older than a defined threshold AND with no live wrapper/git process holding that repo. The reap covers the wrapper's own lock plus the git-internal existence-based locks (`index.lock`, `HEAD.lock`, `refs/**/*.lock`) that git does NOT auto-release on holder death. It **MUST NEVER** remove a lock whose holder is ALIVE (removing a live-held lock corrupts a concurrent writer – \$.). Each reap decision (REAPED / KEPT + reason) is logged as captured evidence (\$. . . – liveness PROVEN via `kill -0`, never assumed). A wrapper that blocks indefinitely on a dead-holder lock (the observed –h freeze, – –) is a violation. Honest boundary (\$. . .): reaping clears dead-holder deadlock – it does NOT substitute for the own-`.git` isolation of \$. . . (which shrinks the blast radius from all-streams to one). Composes \$. . . (working-tree quiescence) / \$. . . (commit-lock-release-immediately + detached push) / \$. . . (no-guessing) / \$. . . (data safety) / \$. . . (evidence stream) / \$. . . (own-`.git` isolation companion). Classification: universal (\$. . .) – the consuming project supplies its concrete wrapper + reap helper + stale-age threshold per \$

Propagation gate `CM-COVENANT-114-180-PROPAGATION` (literal

11.4.180) + recommended gate CM-WRAPPER-STALE-LOCK-AUTO-REAP (every single-writer wrapper reaps provably-stale locks before acquiring, never removes a live-held lock, and logs each REAPED/KEPT decision) + paired \$. meta-test mutation (make the reaper remove a live-held lock, OR skip the reap-before-acquire step → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md \$. . Non-compliance is a release blocker. No escape hatch – no --skip-stale-lock-reap , --force-remove-live-lock , --block-on-dead-holder flag.

\$. . – Consistent controlled feature-branch naming: one feature/logic-group ⇒ exactly ONE canonical branch name across the main repo + all owned submodules (User mandate, – –). Forensic anchor – verbatim user mandate (– –): "System MUST pay attention not to create multiple feat. branches with different names for one same feature or group of workable items / features! Naming MUST BE consistent and carefully controlled! This MUST BE part of the constitution too and carefully risk-free handled! No mistakes can happen! There MUST NOT be any false results, mis-alignments, inconsistency or bluff of any kind anywhere!" One feature / logical group of workable items MUST map to EXACTLY ONE canonical branch name, used IDENTICALLY on the main repo AND on every owned submodule the group touches – NEVER two (or more) differently-named branches for the same feature/group, NEVER a per-repo naming drift. The mandate (ALL hold): () **Single canonical name per group.** For each feature/logical-group a single canonical branch name is minted ONCE (per the \$. . D scheme: feature/<slug> on the branch axis + <prefix>-<base>-feat-<slug> on the tag axis, <slug> lowercase snake/kebab per \$. .). That exact string is the branch name on the main repo AND on every touched owned submodule; an

untouched submodule stays on its base branch (it is NOT force-branched), but any submodule that IS branched for this group MUST use the identical canonical name. ()

Single source of truth for branch names – look up, never re-invent. The canonical name is recorded ONCE in a tracked registry and every subsequent branch creation for that group LOOKS IT UP rather than re-deriving/re-typing it. Bind this to the \$. . . exactly-once claim registry / the \$. . . workable-items DB

logic_group → branch mapping (a branch_name field on the group/claim record): the claim that reserves a feature/group ALSO reserves its canonical branch name, so a second stream working the same group reads the same name instead of minting a divergent one. Creating a branch whose name does NOT match the registry's canonical name for that group – OR minting a SECOND canonical name for a group that already has one – is a \$. . . violation (a \$. . . no-guessing violation: the branch name is a looked-up FACT, never an ad-hoc re-invention). ()

Mechanical enforcement (\$. . .). A pre-build/commit gate CM-BRANCH-NAME-CONSISTENCY verifies that every live feature branch (main repo + each owned submodule) maps to a registered feature/logical-group AND carries the identical registered canonical name across the main repo and every touched submodule; it FAILs on (a) an unregistered feature-branch name, (b) two differently-named branches bound to the same group, (c) a submodule branch whose name diverges from the group's main-repo canonical name. Paired \$. . . mutation: register a single group under two divergent branch names (or point a submodule branch at a name ≠ the group's canonical name) → the gate FAILs; restore → PASS. () **Risk-free**

reconciliation – no force, no loss (\$. . . / \$. . . / \$. . .). When a mis-named / duplicate branch is discovered, reconciling it to the canonical name MUST preserve every commit: create/checkout

the canonical branch at the merge-base, MERGE the mis-named branch's commits onto it (union, no `-s ours /rebase/` reset that drops commits per § . . . step), verify no commit lost (`git log <canonical>..<misnamed>` empty after merge) + no conflict markers, push fast-forward to all upstreams (NEVER force-push – the canonical branch descends from the mis-named tip), THEN retire the mis-named branch only after the merge is confirmed on all mirrors; delete-before-merge is forbidden. Honest boundary (§ . . .): "the branches are probably the same feature" is NOT a determination – bind the branch to a group via the registry (claim id / logic-group id), never guess two branches are "the same feature" without the registry record. Classification: universal (§ . . .) – a platform-neutral naming-discipline reusable across any multi-repo/multi-submodule project; the consuming project supplies its concrete registry (claim registry / workable-items DB), branch/tag scheme, and owned-submodule set per § Composes § . . . (feature work-stream lifecycle / D branch+tag scheme) / § . . . (exactly-once claim registry – the branch-name reservation rides the claim) / § . . . (track-qualified identity – a track-qualified session/lock key is distinct from the shared cross-track feature-branch name; both hold) / § . . . (own-`.git` streams still share ONE canonical feature-branch name across their independent repos) / § . . . (owned-submodule set) / § . . . (lowercase snake/kebab slug) / § . . . (workable-items DB `logic_group`) / § . . . (no-force merge-onto-latest) / § . . . (quiescence) / § . . . (no-loss) / § . . . (branch name is a looked-up fact) / § . . . (mechanical enforcement). Propagation gate CM-COVENANT-114-181-PROPAGATION (literal 11.4.181 across the consumer fleet) + recommended gate CM-BRANCH-NAME-CONSISTENCY + paired § . . . meta-test

mutation (register one group under two divergent branch names → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical**

authority: constitution submodule Constitution.md

\$ Non-compliance is a release blocker. No escape hatch — no --allow-divergent-branch-names , --skip-branch-registry , --rename-branch-freely , --per-repo-branch-name-OK , --two-names-one-feature-OK flag.

\$. . . — **Track+branch work-stream identity label** (T<N>/<branch>) on every agent/subagent/work-stream label and every operator-facing work-stream reference (User mandate, — —). **Forensic anchor — operator mandate** (— —): "EVERY agent/subagent/work-stream label AND every operator-facing report reference to a work stream MUST carry a (T<N>/<branch>) prefix identifying the track + git branch — e.g. (T1/main) ATM-312 MPV drm_prime investigation , (T2/feat/mistiq-vader) SPK-XXX" In a multi-track layout (parallel checkouts each on a distinct branch — e.g. /mnt/track1.4) a bare work-stream label ("ATM- investigation") does not say WHICH track/branch produced it, so cross-track dispatches, status reports, and progress lines collide and mislead. The mandate (ALL hold): ()

Universal label form. Every agent / subagent / work-stream label AND every operator-facing reference to a work stream (agent-dispatch descriptions, status reports, resume/handoff docs, progress lines) MUST START with a (T<N>/<branch> - <alias>) prefix — <N> the track number, <branch> the git branch, <alias> the claude-toolkit alias in use — e.g. (T1/main - claude1) ATM-312 MPV drm_prime investigation , (T2/feature/mistiq-vader - deepseek) SPK-XXX subtitle work . () **Deterministic derivation, never guessed** (\$. . .). The label is DERIVED, not typed from memory: <N> from the checkout path (the multi-track root /mnt/track<N>/... ; honest ? when the checkout is not under

a numbered track – never a guessed number), `<branch>` from `git rev-parse --abbrev-ref HEAD`, `<alias>` from `CLAUDE_CONFIG_DIR` `basename .claude-<alias>` → `<alias>` (honest ? when unset/non-matching, never a guessed name). The consuming project supplies a deterministic labeler (reference impl `scripts/multitrack/track_branch_label.sh`).

(`..`), **Mechanical enforcement** (`$` `..` `/` `$` `..` `-class anti-forgetting`). A PreToolUse guard hook (reference impl `constitution/scripts/hooks/guard-track-branch-label.sh`, inherited BY REFERENCE per `$` `..` `, never copied`) BLOCKS at the tool-call boundary any agent-dispatch tool (Agent / Task / TaskCreate) whose `description` (or `subagent` label) does not START with `^\(T[0-9]+/[^\)]+\)` – so enforcement is independent of any agent's recall (a rule the orchestrator forgets to paste is not enforcement); non-agent tools pass through untouched. A pre-build gate `CM-TRACK-BRANCH-LABEL` asserts the labeler + hook exist, are executable + `..` parseable, the hook is wired in the runtime settings, and the convention doc exists; a propagation gate `CM-COVENANT-114-182-PROPAGATION` asserts the literal `11.4.182` is present across the governance-file fleet. Paired `$` `..` mutation: strip the hook's prefix-check → the hook test-suite FAILs; strip the `11.4.182` literal from a governance file → the propagation gate FAILs; restore → PASS. (`..`)

Honest boundary (`$` `..` `..`). A ? in the track OR the alias field is a valid honest label (unknown track / no git repo), never a bluff; the enforced FORMAT is numeric-track `[0-9]+` per the reference regex, so an agent dispatch from a non-track checkout is surfaced (blocked) rather than silently mislabeled. Classification: universal (`$` `..` `..`) – a platform-neutral multi-stream identity discipline reusable by any project running parallel track/branch work streams; the consuming project supplies its concrete track-root layout, labeler, hook wiring, and governance-file set per `$` `..` `..`. Composes

§ 11.4.182 . . . (track-qualified identity –
 § . . . is the label-prefix instantiation applied
 to EVERY agent/subagent label + every operator-facing
 work-stream reference) / § . . . (naming) /
 § . . . (mechanical enforcement) / § . . .
 (anti-forgetting PreToolUse guard) / § . . .
 (hook inherited by reference, project-decoupled) /
 § . . . (label is a derived fact; ? is honest-
 unknown) / § . . . · § . . . (parallel
 work streams) / § . . . (single-resource-owner
 per track) / § . . . · § . . . (work-
 stream registry identity). Propagation gate CM-
 COVENANT-114-182-PROPAGATION (literal 11.4.182 across the
 consumer fleet) + recommended gate CM-TRACK-BRANCH-LABEL +
 paired § . . . meta-test mutation. **TRUNK RULE (operator
 mandate, – –)**: work performed on the
 canonical trunk (main / master) is **ALWAYS Track** – the
 label reads (T1/main - ...) , NEVER (T?/main - ...) , regardless
 of checkout path. Precedence is deliberate: the BRANCH
 decides first and the checkout path decides only for non-
 trunk branches,...so a trunk checkout living outside /mnt/
 track<N> – or inside a differently-numbered one – is still
 Track ; the identity that matters is which STREAM the
 work belongs to, not which directory it happens to sit in.
 Rationale: § . . . labels are how concurrent
 streams are told apart, so an unnumbered T? on trunk work
 makes it indistinguishable from a genuinely-unknown track
 in every dispatch, report and progress line – and the
 trunk is a well-known, singular stream that is never
 unknown. Reference implementation: scripts/multitrack/
 track_branch_label.sh derives the branch BEFORE the track
 number and maps main|master -> 1 . Labels MUST be re-derived
 from that script at dispatch time and kept in real-time
 sync with the live branch/alias/model/effort – a
 remembered or stale label is itself a § . . .
 violation (§ . . . : live identity is not knowable from

memory). **Canonical authority:** constitution submodule
 Constitution.md § . . . Non-compliance is a
 release blocker. No escape hatch – no `--skip-track-label`,
`--unlabeled-agent-OK`, `--bare-work-stream-reference`, `--no-`
`track-branch-prefix`, `--label-from-memory` flag.

§ . . . – **Maximal multi-agent utilization per work-stream + full-constitution-application + zero-bluff mandate** (User mandate, – –). **Forensic anchor – verbatim operator mandate** (– –):
 "Every Track MUST USE and maximize multi-agents working approaches (subagent-driven, independent agents for various reviewing, and other forms of agents use)! We MUST maximize use of agents so every track produces as much as possible completed work with highest possible quality for the least spent time! Every Track MUST apply everything from the constitution – all rules, mandatory constraints, guides and technology derived! Nothing can be ignored, skipped, avoided or considered less relevant! There MUST NEVER BE false or faulty results, untested and unverified work or bluff of any kind anywhere!" Every track / work-stream MUST maximize multi-agent working approaches – subagent-driven development (§ . . . /
 § . . .), independent review agents
 (§ . . . /§ . . . /§ . . . /
 § . . .), parallel background streams with auto-backfill (§ . . . /§ . . .), and every other applicable agent form – so that every track produces the MOST completed work at the HIGHEST quality in the LEAST time. The mandate (ALL hold): (. .) **Maximize agent utilization per track.** Every track/work-stream (main + every parallel background stream per § . . . /
 § . . .) MUST default to subagent-driven execution for non-trivial work (§ . . . /§ . . .), dispatch independent review agents structurally separated from the author for quality/safety/correctness review
 (§ . . . /§ . . . /§ . . . ,

iterated to a clean G0 per § 6.7.1), and use every other applicable specialised-agent form the runtime offers – never one monolithic foreground agent doing everything a task naturally decomposes into. () **Full-constitution application – nothing ignored.** Every track MUST apply the ENTIRE constitution to its work – every universal rule, every mandatory constraint, every guide, and every derived technology (gates, mutation pairs, sync engines, recording/evidence pipelines, lock/claim registries) this Constitution and its consuming project's governance define. No rule may be silently ignored, skipped, avoided, or judged "less relevant" for a given track – partial-constitution-application is itself a § 6.7.1 violation, equal in severity to skipping a mandated gate. () **Zero-bluff, zero-false-result, zero-untested-work.** There MUST NEVER be false or faulty results, untested or unverified work, or bluff of any kind, anywhere, on any track – this restates and BINDS the entire § 6.7.1 anti-bluff covenant (11.4.3, 11.4.9) as an unconditional per-track invariant, not merely a release-gate check: every claim of "done" on any track carries captured evidence (§ 6.7.1 / § 6.7.2 / § 6.7.3), every gate/mutation pair is genuinely falsifiable (§ 6.7.1), and no track may report progress it did not actually perform (§ 6.7.1). Honest boundary (§ 6.7.1): "maximize agents" is bounded – by § 6.7.1's % host-memory ceiling, § 6.7.1's parallel-agent cap, genuine task parallelizability, and § 6.7.1's single-resource-owner rule for contended hardware/locks – so spawning contending, redundant, or non-parallelizable agents is NOT maximization; the bar is maximum USEFUL parallel throughput per track, never raw agent count. Classification: universal (§ 6.7.1) – a platform-neutral operating-discipline reusable by any project running one or more parallel work-streams/tracks;

the consuming project supplies its concrete track layout, subagent runtime, and review-agent roster per

\$. . . Composes \$. . . (subagent-driven-by-default) / \$. . . (parallel-development PWU pipeline) / \$. . . (subagent-driven execution is the default) / \$. . . (multi-pass change evaluation) / \$. . . (zero-idle priority-first parallel-by-default) / \$. . . (maximum-use-of-idle-time) / \$. . . (continuous parallel-stream working routine) / \$. . . (single-resource-owner partitioning) / \$. . . (code-review-agent gate) / \$. . . (default autonomous-loop working mode) / \$. . . (code-review iterate-until-G0) / \$. . . (universal code-review mandate) / \$. . . (universal independent verification agent mandate) / \$. . . (% memory ceiling). Propagation gate CM-COVENANT-114-183-PROPAGATION (literal 11.4.183 across the consumer fleet) + recommended gate CM-MAX-AGENTS-PER-TRACK (every active track shows subagent/review-agent dispatch evidence for non-trivial work, not solitary foreground execution) + paired \$. . . meta-test mutation (strip the literal → propagation gate FAILs; record a track that skipped subagent/review-agent dispatch for non-trivial work with no genuine single-agent justification → CM-MAX-AGENTS-PER-TRACK FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md \$ Non-compliance is a release blocker. No escape hatch – no

--single-agent-OK , --skip-review-agents , --partial-constitution-application , --serialise-parallelisable-work flag.

\$. . . – Mandatory SonarQube static-analysis CLI + local tooling installed and PATH-discoverable (User mandate, – –). Forensic anchor – operator mandate (– –): "Make sure sonarqube is installed and configured locally same we did for Firebase, GitHub and GitLab CLIs through exported

System PATH and .bashrc! Add all major points to the constitution so all projects do the same!" Every project/host that runs static analysis MUST have the SonarQube scanner CLI (sonar-scanner) installed AND durably discoverable on the system PATH via the shell rc (.bashrc / .zshrc) – exactly like the Firebase/GitHub/GitLab CLIs – PLUS the shared

```
constitution/scripts/sonarqube/.. tooling
( sonarqube_install_check.sh / sonarqube_lib.sh /
sonarqube_run_scan.sh / sonarqube_container.sh + compose/ )
```

consumed **by reference** (\$. . . decoupled + \$. . . project-agnostic + \$. . . -style inherited, NEVER copied per project). The mandate (ALL hold): () **Installed + PATH-durable** – the scanner CLI resolves in every fresh shell (a durable rc export – e.g. a generic Factory/software/*/bin glob or an explicit line – not an ephemeral session-only PATH); (..) **Shared tooling present + GREEN** – sonarqube_install_check.sh exits (scanner + rootless podman per \$. . . + podman-compose + curl + ES-ready host vm.max_map_count) before any scan-dependent work; () **Anti-bluff** (\$. . .) – "installed/configured" is PROVEN by command -v sonar-scanner resolving AND the install-check exit , NEVER assumed; a claimed-installed-but-absent CLI is a \$. . . bluff at the tooling layer; () **Rootless** (\$. . .) – the local SonarQube server runs via rootless podman through the container tooling, never rootful docker/sudo. **Honest boundary** (\$. . .): this anchor names the SonarQube SCANNER CLI + the shared local-server tooling; per \$. . . the former universal Semgrep static-analysis mandate is REPEALED – SonarQube is a DISTINCT, operator-mandated tool here, not a Semgrep re-instatement, and this anchor mandates the TOOLING availability, not that every project must run a scan on every build.

Classification: universal (\$. . .) – the consuming project/host supplies its concrete install path + rc-

export mechanism per \$ Composes
 \$. . . (decoupled shared tooling) / \$. . .
 (install-on-clone) / \$. . . (submodule-catalogue-
 first) / \$. . . (mechanical enforcement) /
 \$. . . (anti-forgetting) / \$. . .
 (rootless container) / \$. . . (Semgrep-repeal
 boundary) / \$. . . (project-decoupled tooling) /
 \$. . . (no-guessing). Propagation gate CM-
 COVENANT-114-184-PROPAGATION (literal 11.4.184 across the
 consumer fleet) + recommended gate
 CM-SONARQUBE-CLI-INSTALLED (command -v sonar-scanner resolves
 + sonarqube_install_check.sh exit + shared tooling
 present) + paired \$. . . meta-test mutation (strip the
 literal → propagation gate FAILs; move the scanner off
 PATH → the install gate FAILs; gate-code = separate work
 item). **Canonical authority:** constitution submodule
 Constitution.md \$ Non-compliance is a
 release blocker. No escape hatch – no --skip-sonarqube-
 install , --sonar-cli-optional , --ad-hoc-sonar-path , --
 rootful-sonar-OK flag.

\$. . . (I) – HawkScan (DAST) + OWASP ZAP +
 gitleaks + Trivy: mandatory local security-tooling
 extension (operator mandate, - -).
Forensic anchor – operator mandate (- -):
 "Add HawkScan and all other missing utils as Git
 submodules fully available on local level through
 constitution Submodule and fully configured and optimized
 for high end use! ... Make sure everything is derived
 through constitution Submodule and accessible to all
 projects incorporating it." Following this anchor's own
 established pattern (this anchor's un-lettered SonarQube
 CLI clause above), every project/host that runs security
 scanning MUST have four additional tools installed and
 locally runnable through this constitution submodule's
 scripts/<tool>/ tree, mirroring the scripts/sonarqube/
 five-file structure (<tool>_lib.sh shared functions,

sourced never executed; `<tool>_install_check.sh` a three-tier presence/version verifier; `<tool>_container.sh` rootless-podman lifecycle where the tool needs a local server; `<tool>_run_scan.sh` drives the scan, polls for completion, and writes a Markdown report to `qa-results/<tool>/<ts>/`, never printing secrets; `compose/docker-compose.<tool>.yaml` for containerized tools), consumed BY REFERENCE and never copied per project (`$ . . . . . decoupled / $ . . . . . -style inherited`):

- . **gitleaks** (secrets scanner) – fully open-source, no account required. Installable via package manager or the upstream release binary; runs entirely locally with no network callout. Complements, and does NOT replace or conflict with, this constitution's own `credential_scan_lib.sh` pre-commit-hook-seam detector – no prior anchor declares that detector the exclusive secrets-scanning mechanism, so both run.
- . **Trivy** (dependency/container/IaC vulnerability scanner) – fully open-source, no account required, runs entirely locally against a filesystem, image, or repository target, with no network callout beyond its own vulnerability-database pull (cacheable, refreshable offline).
- . **OWASP ZAP** (DAST) – fully open-source, no account required, runs as a local rootless-podman container (`zapproxy/zap-stable`, `$ . . . . .`) against a locally-reachable target, no external SaaS dependency.
- . **HawkScan** (DAST, StackHawk) – **the one tool in this set that is SaaS-backed, not self-contained**. The scanner image (`docker.io/stackhawk/hawkscan`) always authenticates against StackHawk's own backend; there is no fully local, no-account mode. Setup therefore has an unavoidable manual, interactive, operator-only component – a free-tier account signup at `app.stackhawk.com` (email + verification) and creation of one "Application" per

scanned project to obtain an `applicationId` – that cannot be automated by any agent or script. `scripts/hawkscan/run.sh` MUST fail open, never closed: with no `HAWK_API_KEY` set, or the placeholder `applicationId` still present, it prints a clear, loud, non-blocking warning and exits `1`. A missing third-party credential is an operator setup gap, never a build or gate failure.

Step-by-step key-acquisition guides for every tool that needs one (today: HawkScan only) live at `scripts/hawkscan/README.md` and propagate with the tool scripts themselves.

None of these four tools is declared in `helix-deps.yaml` – that manifest's schema (`$schema: https://raw.githubusercontent.com/heliossec/helix-deps/main/schema.json`) is for git-clonable source dependencies, and none of the four is distributed as a clonable repository consumable that way (gitleaks/Trivy/ZAP ship prebuilt binaries or container images; HawkScan is SaaS-backed with no source repository to pin). Presence and configuration are governed by this clause plus the `install-check` scripts, exactly as SonarQube already is under this same anchor's un-lettered clause above. **Semgrep is deliberately NOT added by this clause** – `$schema: https://raw.githubusercontent.com/heliossec/helix-deps/main/schema.json`'s repeal of the former universal Semgrep mandate, in favor of `$schema: https://raw.githubusercontent.com/heliossec/helix-deps/main/schema.json`'s SonarQube mandate, stands unreversed.

Honest boundary (`$schema: https://raw.githubusercontent.com/heliossec/helix-deps/main/schema.json`): this clause mandates the tooling's presence and local runnability; it does NOT itself run a scan, does NOT claim any project has been scanned, and does NOT claim any finding from these tools has been triaged – those are separate, per-project work items. HawkScan specifically can never be claimed "configured" by an agent alone – the account and Application creation are irreducibly manual, and a project without them MUST show the non-blocking warning, never a fabricated pass.

Composes with: \$ (SonarQube, this anchor's own un-lettered clause) / \$ (Semgrep-repeal boundary, deliberately not reversed) / \$ (dependency-manifest schema, why these four are absent from it) / \$ / \$ (by-reference, never-copied inheritance) / \$ (rootless container runtime) / \$ (credentials – HAWK_API_KEY etc. are runtime-loaded env vars, never committed) / \$ (no-guessing / honest boundary). Propagation gate CM-COVENANT-114-184-PROPAGATION continues to cover this clause (same anchor, same literal 11.4.184 token – no new anchor number was minted). Recommended new gates: CM-GITLEAKS-INSTALLED, CM-TRIVY-INSTALLED, CM-ZAP-CONTAINER-AVAILABLE, CM-HAWKSCAN-SCAFFOLDED (each: <tool>_install_check.sh exit – for HawkScan, "scaffolded" means the runner correctly detects and reports missing credentials, NOT that a scan ran) + paired \$ mutation (move/rename a tool off PATH → its install-check gate FAILS; neuter hawkscan/run.sh 's credential guard so it runs with no key → the mutation proof for that gate FAILS); gate-code = separate work item, NOT claimed shipped (\$ / \$).

Canonical authority: constitution submodule Constitution.md \$ (I). Non-compliance is a release blocker for the three open-source tools (gitleaks, Trivy, ZAP); HawkScan's compliance bar is "correctly scaffolded with a fail-open credential guard," never "scanned," until an operator completes the two manual account-creation steps. No escape hatch – no --skip-gitleaks, --trivy-optional, --zap-off, --hawkscan-required flag.

\$ – Manual QA-team testing as the mandatory FINAL confirmation of done (User mandate, –). Forensic anchor – verbatim operator mandate (–): "We MUST EVERYWHERE and

ALWAYS take into account that manual testing as final confirmation step for work done is MANDATORY!!! Without such confirmation by QA team of ours we cannot consider any set of features or scope of work fully completed! This MUST BE well known and mandatory part of our constitution! Fully respected, taken into account, followed and never forgotten!" No set of features / scope of work / release may be considered **fully completed** until it has received MANUAL TESTING confirmation by the project's QA team as the FINAL confirmation step – EVERYWHERE and ALWAYS (every feature, every scope, every release), NEVER forgotten. The mandate (ALL hold): () **Necessary-but-not-sufficient automation.** Every automated gate this Constitution mandates – the § . . . anti-bluff covenant, § . . . / § . . . captured evidence, § . . . autonomous validation, § . . . full-suite retest, § . . . four-layer SOURCE→ARTIFACT→RUNTIME→USER-VISIBLE verification, § . . . risk-ordered validation – is NECESSARY but NOT SUFFICIENT for "fully completed." These are the automated confidence gates; the QA-team MANUAL test is the FINAL human sufficiency gate that closes the scope. () **Composition with § . . . – ADDS, does not weaken.** § . . . requires every feature to have an AUTONOMOUS validation path (primary, CI-scalable, runs on every commit, survives operator absence) and forbids operator-attended-ONLY validation. § . . . does NOT weaken that requirement – it ADDS a final human QA sign-off ON TOP of the mandatory autonomous path. A feature MUST have BOTH: the autonomous path (§ . . . , primary + always-on) AND the manual-QA final confirmation (§ . . . , the closing gate). Manual QA is the FINAL gate, never the ONLY gate – a feature validated ONLY by manual QA with no autonomous path is STILL a § . . . violation; a feature validated ONLY by automation with no manual-QA sign-off is a § . . . violation. Both must hold. ()

Workflow placement. The delivery pipeline is: autonomous-validation-GREEN → build → flash/deploy → autonomous on-device validation → **HAND OFF to the QA team for manual testing** → ONLY after QA-team manual confirmation is the scope "fully completed" / eligible for closure or a release tag. The § 4.1.1 autonomous-loop terminal condition (a fully-validated-and-verified release tag) now REQUIRES the QA-team manual confirmation before the tag is cut; a scope marked "done" / "fully completed" / migrated to Fixed.md without a recorded QA-team manual confirmation is a § 4.1.1 violation – a definition-of-done bluff of the same severity class as a § 4.1.1 PASS-bluff. () **Honest boundary** (§ 4.1.1). The agent/loop MAY autonomously drive the scope through autonomous-GREEN + build + flash + on-device autonomous validation, but it MUST then hand off + WAIT for the QA-team's manual confirmation before claiming "fully completed" – it MUST NOT self-certify the manual step, simulate it, or infer it from automated evidence alone (that would itself be a bluff at the human-sufficiency layer). While awaiting QA-team confirmation the agent keeps progressing every OTHER non-blocked item in parallel (§ 4.1.1 / § 4.1.1 / § 4.1.1) – the wait parks that one scope, not the loop. Classification: universal (§ 4.1.1) – a platform-neutral definition-of-done discipline reusable by ANY project that maintains a QA team; the consuming project supplies its concrete QA-team identity, hand-off mechanism, and confirmation-recording location per § 4.1.1 . Composes with § 4.1.1 (anti-bluff covenant) / § 4.1.1 (captured-evidence quality) / § 4.1.1 (no-guessing – the manual confirmation is a recorded FACT, never assumed) / § 4.1.1 (full-suite retest before tag) / § 4.1.1 (autonomous-validation mandate – § 4.1.1 is its human-sufficiency CAPSTONE, adding rather than substituting) / § 4.1.1 (sink-side

evidence taxonomy) / § . . (four-layer fix-verification)₁₃₂ / § . . (autonomous-loop terminal condition now gated on QA-team confirmation) / § . . (huge-blocker release protocol) / § . . (post-remediation validate-the-fix-first) / § . . (risk-ordered validation priority). Propagation gate CM-COVENANT-114-185-PROPAGATION (literal 11.4.185 across the consumer fleet) + recommended gate CM-MANUAL-QA-FINAL-CONFIRMATION (every item marked fully-completed / migrated to Fixed.md carries a recorded QA-team manual-confirmation citation) + paired § . meta-test mutation (strip the literal → propagation gate FAILs; mark an item fully-completed with no QA-team confirmation citation → the recommended gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md § . . . Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-manual-qa , --automation-suffices , --self-certify-manual-step , --assume-qa-confirmed , --autonomous-only-completion flag.

§ . . – **Anti-divergence enforcement: cross-document consistency is a mandatory-before-export/commit gate, never an after-the-fact audit (research-derived, – –)**. Compact summary: any project maintaining more than one representation of the same tracked data (a delivery plan + its MVP brief + a workable-items DB + their summaries + .md / .html / .pdf / .docx exports) MUST enforce cross-document CONSISTENCY as a deterministic gate that runs BEFORE any export render, any doc/DB sync verify , and any doc-set commit – NEVER as an after-the-fact human audit that runs only when an operator happens to ask (the forensic FACT: the same NanoKVM/phone-as-screen feature appeared in three plan clusters – one shared ticket SPK- across SIX rows split over TWO releases, and a

carve-out due to ship - - whose own prerequisite does not START until - - - and NO gate caught it; only an operator's manual read of Plan_v9.0.xlsx surfaced it, after the divergent bytes had already exported). () **One no-divergence verdict, three seams** – a single deterministic PASS/FAIL/SKIP verdict (built by composing a cross-document integrity validator with the workable-items-DB internal validate) gates all three write-seams (export / sync-verify / commit); a FAIL at any seam refuses that seam, naming the exact offending records (\$. . actionable). () **Five decidable check families** – DEDUP (no two records describe one feature with divergent timeline/status/release, keyed on ticket OR normalised (subject, scope) , NOT bare subject substring), TIMELINE (start $\leq$ deadline; no deadline-before-dependency; no dependency cycle; GATED exempt-but-marked), CROSS-DOC (the same item across every representation carries identical timeline/status/type against a NAMED authoritative source), INTEGRITY (no orphan refs; Status Type \$. . ; location status), STRUCTURAL (required columns; ID uniqueness+monotonicity \$. .). () **Drift-proof trigger** – a \$. . sha -of-sorted-members fingerprint of the authoritative inputs re-arms the gate on ANY input change, so a stale verify cannot pass on changed data. () **Self-validated analyzer** (\$. . ()) – the validator ships golden-good (MUST PASS) + one golden-bad per family (MUST FAIL, pinpointing the offender) + a negative-control (distinct same-subject tasks that MUST PASS – the false-positive guard); a validator that passes its golden-bad, or fails golden-good/the negative-control, is itself a \$. bluff and a release blocker. () **Reusable + decoupled** (\$. . /.) – packaged as a depth- reusable engine under the constitution submodule per the \$. . (C) carve-out with a helix-deps.yaml ;

project-specific doc-sets + authoritative-source bindings live in a consumer-owned checkset (data, never engine code). () **Supersedes ad-hoc audits** – the per-cycle `*_INTEGRITY_FINDINGS.md` / `RECONCILIATION_*.md` after-the-fact audit docs are retired in favour of the gate; honest boundary \$. . – the gate proves internal CONSISTENCY (no record diverges, no timeline is incoherent, every ref resolves), NOT plan CORRECTNESS (achievability/date-rightness stay operator/management decisions the gate SURFACES, never MAKES). Classification: universal (\$. .). Composes \$. . /. /. /. /. /. /. /. /. /. Propagation gate `CM-COVENANT-114-186-PROPAGATION` (literal `11.4.186`) + recommended gate `CM-DOC-INTEGRITY-VALIDATION` (invariants: validator present+executable; consumer checkset present; the no-divergence gate wired into each of the three seams; fingerprint sidecar fresh; `selfcheck` golden-good/golden-bad/negative-control wired into meta-test) + paired \$. mutation. **Canonical authority:** constitution submodule `Constitution.md` \$. . . Non-compliance is a release blocker. No escape hatch – no `--skip-divergence-gate` , `--audit-after-export-OK` , `--cross-doc-check-not-applicable` , `--unvalidated-analyzer-OK` , `--divergent-commit-OK` flag.

\$. . – **Automatic multi-track ruler**
orchestration (User mandate, – – ; landed – –). Compact summary: every project incorporating this constitution MUST ship its multi-track parallel-development orchestration as a UNIVERSAL, AUTOMATIC, OUT-OF-THE-BOX, INHERITED capability – not a per-project reimplementation – via a single conductor session that programmatically SPAWNS, DRIVES, RESUMES, MONITORS, and CONTROLS per-track headless workers. ()
Headless spawn primitive – workers launch via native headless `claude -p --output-format stream-json --verbose` ; the worker's `session_id` is captured from the FIRST `init` event

(never invented); success/failure is read from the terminal event's `result.is_error` field, NEVER the process exit code (a CLI can exit non-zero on a successful `result`, or zero on a failed one – exit-code-as-success-signal is a \$. . guess); continuation of an existing worker uses `--resume <session_id>` invoked from the SAME cwd + env the spawn used (the storage/resume-collision gotcha). () **Per-subscription auth + rate-limit fallback** – every spawn `unset` s any ambient `ANTHROPIC_API_KEY` (the precedence trap that silently routes a "subscription" worker onto pay-per-token API billing) and sets the selected alias's `CLAUDE_CODE_OAUTH_TOKEN` / `CLAUDE_CONFIG_DIR`; on a captured rate-limit signature (`isApiErrorMessage` + `apiErrorStatus:429` , quota-exhausted vs transient distinguished by evidence, never guessed) the ruler REBINDS the affected track onto the next healthy alias (cross-track reuse permitted), falls back to a provider alias where configured, and bounded-parks (\$. .) rather than hanging when every alias is cooled. () **Crash-resilient ruler self-supervisor** – the ruler's OWN state (per-track alias/session_id/worktree/last-verdict) is persisted durably (write-temp-then-rename, NEVER tmpfs) so an external watchdog detects a dead ruler PID and re-hydrates every track's binding from the durable state + the \$. . event stream – no work is silently lost to a ruler crash (\$. . at the ruler layer). () **Idempotent bootstrap + auto-install + conductor-stays-home** – one idempotent bootstrap command wires the whole system out of the box (cwd-hook install, config validation, reconciliation), auto-invoked on every constitution pull via the \$. . `post_update_hook.sh` `skill-register` seam; the `conductor`: -designated alias/session is NEVER bound into a track worktree (the resolver returns nothing for it) so the conductor always runs from the main checkout. () **Project-agnostic engine, consumer-owned data** – the entire

mechanism (config loader, alias→worktree resolver, cwd-hook, bind/fallback/cooldown orchestrator, session launcher, rate-limit monitor, supervisor) lives at `constitution/scripts/multitrack/`, inherited BY REFERENCE (`$ . . . . (B)/$ . . . .`), carrying NO project literal; every consuming project supplies its own per-host `config/multitrack/<hostname>.yaml` (tracks/mounts/serials/aliases/conductor) as DATA, never as an engine edit. Anti-bluff (`$ . . . /$ . . . .`): every piece carries four-layer coverage (pre-build gate + on-host test + paired `$ . . . .` mutation + captured evidence per `$ . . . .`), and each fix declares a runtime signature verified on a clean host – never a claim the mechanism "should work." Honest boundary (`$ . . . .`): this anchor mandates the MECHANISM; genuine per-subscription quota-isolation proof requires live tokens exercising real rate-limit boundaries (an operator-scheduled live acceptance window) – a mock-driven GREEN proves the spawn/resume/fallback CONTRACT, never live-quota behaviour, and the two are never conflated or substituted for each other.

Classification: universal (`$ . . . .`). Composes

- `$ . . . .` (subagent-driven-by-default) /
- `$ . . . .` (submodule decoupling + dependency layout) / `$ . . . .` (parallel-development PWU pipeline) / `$ . . . .` (subagent-driven execution default) / `$ . . . .` (autonomous-decision-over-blocking – the bounded-park-on-all-cooled behaviour) /
- `$ . . . .` (continuous parallel-stream routine) /
- `$ . . . .` (resolve-by-stable-name – config by serial, never by index) / `$ . . . .` (real-time conductor framework sync channel) / `$ . . . .` (single-resource-owner partitioning) / `$ . . . .` (default autonomous-loop mode) / `$ . . . .` (always-on device-recording – the analogous host-safety discipline) / `$ . . . .` (standing session-resumption file) / `$ . . . .` (crashed-agent

respawn-until-complete – the ruler-self-supervisor is its ruler-layer instantiation) / \$. .
 (constitution auto-propagation hook seam) /
 \$. . (feature work-stream lifecycle) /
 \$. . (multi-track work-division + exactly-once claim + device-lock) / \$. /\$. (host-safety memory + concurrency ceilings – the spawn/launch pool is budget-gated). Propagation gate CM-COVENANT-114-187-PROPAGATION (literal 11.4.187) + recommended gates CM-MULTITRACK-ENGINE-IN-CONSTITUTION (the generic ruler scripts present under constitution/scripts/multitrack/, carrying no project literal, bash -n + sh -n clean per \$. .) / CM-MULTITRACK-AUTOINSTALL-WIRED / CM-MULTITRACK-BOOTSTRAP-IDEMPOTENT / CM-MULTITRACK-BIND-ON-START / CM-MULTITRACK-FALLBACK-MONITOR / CM-MULTITRACK-HEADLESS-SPAWN / CM-MULTITRACK-RULER-SUPERVISOR + paired \$. mutation (strip the literal → propagation gate FAILs; re-introduce a project literal into the engine, or downgrade a rate-limit rebind to exit-code-based success detection → the corresponding mechanism gate FAILs). **Canonical authority:** constitution submodule Constitution.md
 \$. . . Non-compliance is a release blocker. No escape hatch – no --skip-multitrack ,
 --manual-multitrack-only , --no-auto-install , --per-project-reimpl-OK , --exit-code-is-success-signal flag.

\$. . – Regular main→feature merge cadence: every track / feature-branch MUST FREQUENTLY merge canonical main into itself DURING the work, never only at the end (User mandate, – –). Compact summary: every long-lived feature branch AND every parallel-development track (\$. . /\$. . parallel streams; \$. . feature work-streams; \$. . /\$. . /\$. . multi-track) MUST regularly git merge origin/main (the canonical trunk/master) INTO its own branch THROUGHOUT the work – NOT only at the end – so no branch ever drifts far from

main and the eventual back-merge stays a small, low-conflict, low-risk operation. GENERALISES

§ . . (D) (trunk-merged-INTO-every-BIG-feature-STREAM, at minimum after every trunk tag / daily, `git merge` never rebase, stale-if-not-synced) from the big-work-item, feature-stream lifecycle to EVERY feature branch on EVERY track – a small/medium feature branch, or any `git worktree` / CoW track that never earned a full

§ . . stream, is STILL bound to this cadence (the § . . (D) discipline promoted to the general standalone rule; § . . (D) remains its big-item specialisation). () **WHAT** – fetch the latest canonical `main` / `master` and merge it INTO the working branch, keeping the branch continuously close to trunk; a branch that has NOT integrated trunk within the cadence is flagged STALE (the exact long-lived-branch conflict-storm + integration-risk anti-pattern this forbids). () **CADENCE** – merge trunk in FREQUENTLY: after EVERY trunk tag, at minimum DAILY while the branch is live, AND before starting any significant new chunk of work on the branch – MANY SMALL merges, never one big deferred end-of-branch merge (small frequent merges minimise the divergence + conflict surface). () **HOW (safe-by-construction)** – MERGE, NEVER rebase a shared/tagged/pushed branch (no commit-loss, no history rewrite – § / § . .); FETCH-FIRST every remote before the merge (§ . . / § . . – remote state is unknowable without a fetch, § . .); take a § . . hardlinked pre-op backup before any LARGE / conflict-heavy / risky merge; resolve every conflict CAREFULLY – ZERO conflict markers committed, ZERO file silently dropped, union of both sides preserved (§ . . step / § no-loss); integrate ONLY when the branch is QUIESCENT (§ . . – no in-flight mutation gate, no unaccounted uncommitted work); prefer running the merge + its verification in the BACKGROUND parallel to the main stream so it never stalls

other work (\$. . /\$. . /\$. .);
 NEVER force-push the merged result (\$. . – the
 merged branch fast-forwards its own mirrors). ()

CONSISTENCY + ANTI-BLUFF – after EVERY merge the branch
 MUST be PROVEN consistent, never ASSUMED: a pre-build
 smoke gate runs post-merge (\$. . step- smoke)
 and is GREEN, a `git grep` for conflict markers (<<<<<< /
 ===== / >>>>>>) returns EMPTY, and no commit present
 pre-merge on either side is lost – captured evidence per
 \$. . /\$. . ; a merge claimed "clean" without
 the post-merge smoke + marker-scan + no-lost-commit check
 is a \$. . /\$. . bluff at the integration layer.

() **HONEST BOUNDARY** (\$. .) – regular trunk-
 integration keeps the branch MERGEABLE + low-conflict; it
 does NOT by itself prove the branch's own work is correct
 (each change still crosses \$. . runtime-
 signature + \$. . retest), and it is NOT the
 approval-gated back-merge TO trunk (that stays
 \$. . (I)/\$. . /\$. . no-merge-
 until-approved). Classification: universal (\$. .)
 – the consuming project supplies its canonical-trunk ref
 name, cadence unit, post-merge smoke-gate command, and
 per-track worktree/CoW mechanism per \$. . .

Composes \$. . /\$. . /\$. . /\$. . /
 \$. . /\$. . /\$. . /\$. . /
 \$. . /\$. . /\$. . /
 \$. . /\$. . /\$. . /
 \$. . /\$. . /\$. . /
 \$. . . Propagation gate CM-COVENANT-114-188-

PROPAGATION (literal 11.4.188) + recommended gate CM-
 REGULAR-MAIN-MERGE-INTO-FEATURE (every live feature branch /
 track has integrated the canonical trunk within the
 declared cadence; its last integration is a real `git merge`
 not a rebase; the post-merge smoke + zero-conflict-marker
 + no-lost-commit checks passed) + paired \$. . mutation
 (let a live branch go stale past the cadence, OR record a

rebase of a shared branch in place of a merge, OR commit a conflict marker → the gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule

Constitution.md § . . . Non-compliance is a release blocker. No escape hatch – no --skip-main-merge , --merge-only-at-end , --rebase-shared-branch , --defer-trunk-sync , --no-post-merge-smoke , --force-push-after-merge , --big-items-only-trunk-sync flag.

§ . . . – **Most-reopened cases get extra-depth live-testing scrutiny FIRST** (User mandate, – –). Verbatim operator mandate: "All live testing MUST pay extra attention to cases which have been reopened numerous times and those cases in depth retested, investigated, fully validated and verified with real physical results and no bluff of any kind anywhere!" All LIVE TESTING (on-device / on-target runtime validation, deploy-and-observe) MUST give EXTRA-DEPTH retest + in-depth investigation + FULL validation/verification with REAL PHYSICAL captured evidence (§ . . . / § . . . / § . . . – captured audio/video/sysfs/dumpsys/sink-side/runtime-signature) and NO bluff of any kind ANYWHERE to the cases that have been REOPENED THE MOST TIMES (highest § . . . reopens-count) – the empirically-most-fragile set gets the DEEPEST live scrutiny, and it gets it FIRST. () **WHAT** – for the highest-reopens-count cases, live testing is NOT a single pass: it re-runs the case at EXTRA DEPTH (more iterations per § . . . deterministic consistency, wider edge-case + regression coverage), RE-INVESTIGATES the underlying defect per § . . . systematic-debugging + § . . . reproduce-first, and CONFIRMS with rock-solid captured physical evidence per § . . . – never metadata-only / config-only / absence-of-error / grep-without-runtime (§ . . . / § . . .), never a false result. () **ORDER** – the most-reopened set is scrutinised FIRST, ahead of the rest of

the live-test suite (the \$. . . risk-descending discipline, with most-reopened the load-bearing key). () **KEY** – priority is keyed off the \$. . . reopens-count (the consuming project supplies its reopens-count source per \$. . . , e.g. the \$. . . workable-items DB reopens_count); a high reopens-count is the strongest empirical fragility signal

(\$. . . (d)). () **PROOF** – each most-reopened case's live PASS cites its captured-evidence artefact path (\$. . . / \$. . .) on a CLEAN/fresh deployment (\$. . . runtime-signature), with the \$. . . RED→GREEN polarity flip captured where a permanent regression guard exists (\$. . .).

STRENGTHENS/REFINES \$. . . (risk-ordered validation factor (d) most-reopened – \$. . . promotes most-reopened into its OWN EXTRA-DEPTH mandate at the LIVE-TEST layer, not merely one ranking factor among four) + \$. . . (reopens-count as the priority key) + \$. . . / \$. . . (validate-the-just-fixed / highest-risk-first after redeploy).

Classification: universal (\$. . .) – references no project-specific hardware; the consuming project supplies its reopens-count source per \$ Composes

\$. . . / \$. . . / \$. . . / \$. . . /

\$. . . / \$. . . / \$. . . /

\$. . . / \$. . . / \$. . . /

\$. . . / \$ Propagation gate **CM-COVENANT-114-189-PROPAGATION** (literal 11.4.189) +

recommended gate **CM-MOST-REOPENED-EXTRA-DEPTH-LIVE-TEST** (the live-test suite orders + extra-depth-scrutinises the highest-reopens-count cases FIRST, each with captured physical evidence on a clean deployment) + paired \$. . . mutation (drop the most-reopened extra-depth pass, OR run it AFTER the rest of the suite, OR PASS a most-reopened case on metadata-only evidence → the gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate

work item). **Canonical authority:** constitution submodule
 Constitution.md § . . . Non-compliance is a
 release blocker. No escape hatch – no `--skip-most-reopened-
 depth` , `--reopened-same-as-any` , `--live-test-any-order` , `--
 metadata-pass-for-reopened-OK` , `--defer-reopened-scrutiny` flag.

§ . . . – **Website engineering-quality mandate:**
 every project website **MUST** be fully responsive +
 completely SEO-optimized + uniquely OpenDesign-authored +
 bleeding-edge enterprise-quality, each **PROVEN** with
 captured evidence (User mandate, – –).
 Verbatim operator mandate: "every website we create inside
 our projects **MUST BE** fully responsive – working on all
 browsers, all operating systems, all platforms and all
 devices and screen sizes! Every Website **MUST BE** completely
 SEO optimized! Layouts and templates done **MUST BE** unique
 and everything fully designed from OpenDesign! Make sure
 this applies to our project Website as well! Everything
 must be bleeding-edge enterprise quality, stunning visual
 look and feel, innovative, creative and unique! There **MUST**
BE no false or faulty result or bluff of any kind
 anywhere!" Every website / web-UI surface a project ships
 – INCLUDING this project's OWN website – **MUST** satisfy AND
PROVE with captured physical evidence (§ . . . /
 § . . . /§ . . . – never a
 bare claim, § . . .) **ALL of: (A) FULL RESPONSIVENESS** –
 correct, usable rendering across ALL major browser engines
 (Chromium/Firefox/WebKit), all operating systems, all
 device classes (desktop/laptop/tablet/phone/large-display)
 and the full range of screen sizes/orientations/pixel-
 densities; NO broken layout / overflow / clipping /
 overlap / off-screen or unusable control at ANY
 breakpoint. **(B) COMPLETE SEO OPTIMIZATION** – semantic HTML
 + correct document outline; per-page unique `<title>` + meta
 description; Open Graph + Twitter cards; canonical URLs;
 structured data (schema.org / JSON-LD); `robots` + XML
 sitemap; crawlable / indexable markup; WCAG AA

COVENANT-114-190-PROPAGATION (literal 11.4.190) + recommended gates CM-WEBSITE-RESPONSIVE-PROVEN (host-rendered screenshots across the breakpoint × engine matrix + layout oracle – no overlap/clip/overflow/off-screen) / CM-WEBSITE-SEO-OPTIMIZED (automated SEO audit meets the score floor + structured-data / meta / OG / sitemap / robots present) / CM-WEBSITE-OPENDSIGN-UNIQUE (every layout/template's OpenDesign-token provenance, no generic/off-the-shelf template) + paired § . mutation (report a website responsive/SEO/unique/enterprise-quality with NO captured proof, OR ship a generic non-OpenDesign template, OR skip the light/dark host-rendered pixel proof → the recommended gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md

§ . . . Non-compliance is a release blocker. No escape hatch – no --skip-responsive-proof , --seo-optional , --generic-template-OK , --visual-quality-later , --claim-without-render flag.

§ . . . – Work-to-track/branch binding enforcement: a logic-group's work can NEVER be committed OR dispatched onto the wrong track/branch (research-derived, – –). Compact summary: every logical group of workable items binds to EXACTLY ONE canonical (track, branch) recorded ONCE in the single source of truth (the § . . . workable-items DB logic_groups.destination = authoritative branch + a canonical_track column + a group_paths file-scope manifest), and NO change belonging to a group may LAND on, nor be DISPATCHED to, a track/branch other than that group's canonical one – the file-scope generalisation of § . . . /\$ from branch-NAME consistency to work-PLACEMENT consistency (forensic FACT – – : srcpy/Remote-app work – group mistiq-vader , canonical feature/mistiq-vader /T – was committed onto main /T ; § . . . 's guard-branch-

consistency.sh never fired because committing onto an existing branch is not a feature-branch CREATE, and no gate checked which FILES land on which BRANCH). ()

Single source of truth – the logic_group → (branch, track, path-globs) binding is recorded ONCE in the DB and LOOKED UP, never re-invented (a commit/dispatch placing a group's file-scope on a non-canonical branch/track is a \$. . no-guessing violation). () **Preventive hook** (\$. . /\$. .) – a PreToolUse guard guard-work-track-binding.sh (inherited BY REFERENCE, never copied) BLOCKS (exit) a git-commit-class Bash call whose to-be-committed file set contains a path owned by a group whose canonical branch/track ≠ the current checkout, AND BLOCKS an Agent/Task dispatch whose \$. . label track/branch ≠ the canonical (track, branch) of the ticket's group; fail-closed on unreadable registry, # guardrails:allow <reason> audited escape. () **Detective gate** (\$. . /\$. .) – CM-WORK-TRACK-BINDING-ENFORCED asserts, on the current branch's own added commits (merge-base(HEAD,main)..HEAD), that every group_paths -owned file's group destination == current branch (and canonical_track == current track when pinned) – the actual enforcement boundary regardless of how the commit was made, the runtime signature being FAIL-on-wrong-placement / PASS-on-canonical-placement on a clean checkout. () **Risk-free reconciliation** (\$. / \$. . /\$. .) – a mis-placed change is reconciled by MERGING its commits onto the group's canonical branch (union, no commit lost, no -s ours / reset), ff-only push to all upstreams (NEVER force-push), then removing it from the wrong branch only after the merge is confirmed on all mirrors. Classification: universal (\$. .) – references no project hardware; the consuming project supplies its (track, branch) map, group_paths seed, and commit-wrapper wiring per \$. . . Composes

§ . . /. /. /. /. /. 11.4.191 /. /. /.
 § . . Propagation gate CM-COVENANT-114-191-PROPAGATION
 (literal 11.4.191) + recommended gate CM-WORK-TRACK-BINDING-
 ENFORCED + paired § . mutation. **Canonical authority:**
 constitution submodule Constitution.md § . .
 Non-compliance is a release blocker. No escape hatch – no
 --allow-cross-track-commit , --skip-work-binding , --commit-any-
 branch-OK , --dispatch-any-track-OK , --no-file-scope-check
 flag.

§ . . . – Continuous multi-track auto-backfill: a
 FREE track MUST be IMMEDIATELY re-assigned its next-
 highest-priority domain work-set, never left idle while
 its domain has actionable items (User mandate,
 – –). Verbatim operator mandate: "once a
 Track is free, we shall assign next set of workable items
 to it ... full continuous multi-track development ... no
 stopping ... this is MANDATORY". Compact summary: the
 moment any track's worker COMPLETES its current work-set
 (its item(s) closed with anti-bluff evidence, or merged),
 that FREE track MUST be IMMEDIATELY assigned the NEXT-
 highest-priority actionable work-set from its OWN domain
 (its § . . logic_group / canonical branch) – a
 track is NEVER left idle while its domain still holds
 actionable workable items, and the conductor/ruler
 CONTINUOUSLY re-assigns with NO stopping (full continuous
 multi-track development); PROMOTES the § . . (B)
 ≥ -parallel-background-streams-with-auto-backfill
 invariant into the multi-track (§ . . /
 § . . /\$. . . /\$. .) layer
 (§ . . (B) backfills a STREAM the moment it is
 done; § . . backfills a TRACK the moment its
 worker is done, keyed to that track's canonical (track,
 branch) domain per § . .). () **WHAT** – on any
 track-worker completion the conductor/ruler
 (§ . .) MUST, per that track's domain, look up
 the NEXT actionable work-set, CLAIM it exactly-once

(`$` . . . exactly-once claim registry, no two tracks claiming one item), and DISPATCH it as the track's new `$` . . . -labelled work stream – without waiting for a scheduled wake or an operator prompt. () **PRIORITY** – the next work-set is the track domain's highest-priority actionable item (`$` . . . /`$` . . . priority order, audio-first where in scope), never an arbitrary pick. () **SOURCE-SIDE ADVANCES EVEN WHEN VALIDATION IS GATED** – when the next item's on-device VALIDATION is blocked on a pending build / device availability (`$` . . . clean-target / `$` . . . manual-QA / operator-gated flash), the track MUST still advance the item's SOURCE-SIDE work (investigation, RED-test `$` . . . , source patch, pre-build gate + paired `$` . . . mutation, docs) and honestly `$` . . . / `$` . . . SKIP-with-reason ONLY the genuinely device-gated validation portion – a build/device gate on the validation step is NOT a reason to idle the whole track. () **HONEST BOUNDARY** (`$` . . .) – a track idles ONLY when its domain GENUINELY has zero remaining actionable workable items OR every remaining item is externally blocked with no source-side advancement possible (`$` . . . (A) / `$` . . . / `$` . . . idle-only-when-blocked closed set); "the merges are done" / "the current item is done" / "nothing visible right now" is NEVER a valid idle reason while the domain still holds progressable items (a `$` . . . + `$` . . . violation); an unavoidable per-item block PARKS one item, it does NOT stop the track (`$` . . . – the track backfills its next non-blocked domain item).

Classification: universal (`$` . . .) – references no project hardware; the consuming project supplies its track→domain map + priority source per `$`

Composes

`$` . . . / . / . / . / . / . / . / . / .

Propagation gate CM-COVENANT-114-192-PROPAGATION (literal

11.4.192) + recommended gate CM-CONTINUOUS-MULTITRACK-BACKFILL (on a track-worker completion the ruler re-assigns that track's next-highest-priority domain work-set; a track with a completed worker AND remaining actionable domain items but no new dispatch → the gate FAILs) + paired \$. mutation (let a free track sit idle while its domain holds an actionable item, OR treat "merges done" as a valid idle reason → the gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md \$. . Non-compliance is a release blocker. No escape hatch – no --allow-idle-track , --skip-backfill , --stop-multitrack , --merges-done-is-idle , --defer-track-reassign flag.

Appendix C – Multi-upstream push

This constitution submodule is hosted on multiple Git providers, and consuming projects often are too. The submodule ships an install_upstreams.sh helper (or invokes the install_upstreams system command from the parent toolkit) that reads Upstreams/*.sh declarations and configures all remotes locally.

Every commit to the constitution submodule MUST be pushed to ALL configured upstreams. Use a shell helper that iterates git remote and pushes one by one, OR configure origin with multiple push URLs (git remote set-url --add --push "origin <url> per remote).

\$. . – Anti-blind-typing mandate: every UI interaction MUST be "seen" and "understood" via OCR/vision/screenshot proof, NEVER blind-typed (User mandate,

- -). Verbatim operator mandate: "No blind typing of any kind anywhere!!!" and "Login to apps and services MUST BE done with understanding and seeing the whole UI ALWAYS!!!" and "Add into the constitution that blind typing is STRICTLY FORBIDDEN, and instead everything MUST BE seen and understood in coordination with proper models and means we have incorporated in HelixQA!!!" and "No bluff of any kind anywhere!!!"

Blind typing – sending keystrokes, text, or input events to a UI surface WITHOUT first capturing AND understanding what is on screen (OCR / vision / uiautomator / accessibility-tree / screenshot) AND confirming the result after each interaction – is STRICTLY FORBIDDEN everywhere and always, for every agent, every tool, every automation, every platform, every context. Blind typing is LITERAL bluff: the agent claims "credentials entered" or "login completed" without visual proof the keystrokes landed on the intended fields, on the intended screen, producing the intended result. A credential typed into the wrong field, onto the wrong screen, or into a WebView that swallowed the input silently is a \$. PASS-bluff at the UI-interaction layer – the agent reported success while the end-user feature (logged-in app) is broken.

() SEE-BEFORE-YOU-TYPE. Before ANY input text, input keyevent, adb shell input, uiautomator gesture, or equivalent UI-driving command, the agent MUST capture the current screen state via screencap + OCR OR uiautomator dump OR accessibility-tree dump OR vision-model frame AND confirm the intended target element is visible AND focused. The captured screen state IS the evidence per \$. . and \$. . feature class touch_input.

() VERIFY-AFTER-YOU-TYPE. After EVERY interaction step, the agent MUST capture a NEW screen state and confirm the intended change occurred. A sequence of input text

commands with a single screenshot at the end is blind
typing – each step needs its own verification.

() OCR/VISION ORACLE MANDATORY. The verification MUST
use at least ONE of: Tesseract OCR on a screencap at
+ DPI, uiautomator dump with grep for expected text
content, HelixQA VisionEngine frame analysis per
\$. . and \$. . , or accessibility-
tree dumpsys confirming element state.

() WEBVIEW SURFACES. When uiautomator returns near-empty
hierarchy (common in streaming-app login WebViews), the
agent MUST fall back to OCR via screencap then tesseract
per \$. . . If even OCR is infeasible (secure
surface captures black, DRM-blanked), the login is
genuinely operator_attended only – SKIP-with-reason per
\$. . , NEVER blind-type-and-claim-success.

() CREDENTIAL-SAFE OCR. When OCR output may contain
credentials, the agent MUST redact credential characters
from OCR output and log files per \$. . . Evidence
paths cited via ab_pass_with_evidence; evidence file is
the redacted screenshot or OCR text with credential
characters replaced.

() ANTI-BLUFF. Every login/interaction PASS MUST cite a
captured-evidence artefact path containing pre-interaction
screen state, post-interaction screen state, and OCR/
vision verification showing the expected change. A PASS
without these artefacts is a \$. . PASS-bluff
regardless of exit code. Self-validated per
\$. . () : the OCR/vision oracle MUST ship a
golden-good and golden-bad fixture pair.

Classification: universal (\$. .). Composes
\$. . / \$. . / \$. . /
\$. . / \$. . / \$. . /
\$. . / \$. . / \$. . /
\$. . / \$. . . Propagation gate CM-

COVENANT- - -PROPAGATION (literal
 . .) + recommended gate CM-NO-BLIND-TYPING +
paired \$. mutation (strip OCR verification before an
input text → gate FAILs; strip the literal → propagation
gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule
Constitution.md \$. . .

Non-compliance is a release blocker regardless of context.
No escape hatch – no blind-typing-ok, skip-ocr-
verification, assume-input-landed, login-without-proof,
type-without-seeing flag exists.

\$. . – Exhaustive all-scenario,
all-angle code-review + verify-against-
captured-runtime-evidence mandate (User
mandate, - -)

Forensic anchor – verbatim user mandate
(- -):

"CRITICAL: Code review MUST TAKE into the accounts all
scenarios, look at the problem or change from all
angles, deepest possible analysis! No gaps,
shortcomings or false results or bluff is allowed!
Extend our root constitution Submodule and all
relevant files properly so the review process is more
precise and we do not have situations like the last
one!"

Forensic case study (genericised, - - -
"the last one"). A code-review (\$. .) returned
a clean GO with ZERO findings on a boot-crash fix. It
dismissed a defensive clamp branch as an INFORMATIONAL NIT
– "effectively unreachable ... the input is always valid"
– WITHOUT PROVING that assumption against the actual

runtime input space. The fix shipped, was deployed to the target, and the target STILL crash-looped: the degenerate input the review had assumed could not happen DID happen at runtime (an empty / default-initialised profile opened with the exact zero / default value the review assumed impossible). The crash was a MULTIPLICATIVE size computation – a product $A \times B$ (e.g. `frame_size = channels \times bytes_per_sample(format)`) – with TWO INDEPENDENT zero-triggers; the review verified only ONE factor (`A`, e.g. `channels`) and ASSUMED THE OTHER (`B`, e.g. `format`) away. The crash-loop reproduced on the deployed target under conditions PRESENT in the pre-fix crash evidence that the review never cross-checked the fix against – the evidence was available and not consulted.

Every code-review this Constitution governs (STRENGTHENING § . . . / § . . . / § . . . – the review remains NECESSARY; § . . . makes it EXHAUSTIVE-and-evidence-grounded so a clean GO can no longer be returned on an unproven assumption) MUST, before returning GO:

() **Enumerate the FULL input / scenario space** of BOTH the change AND the defect it addresses – every dimension, every code path, every interacting feature, every input domain the change touches – NOT merely the reported / happy-path scenario. For a MULTI-FACTOR failure mode (e.g. a product of `N` terms, ANY of which being degenerate – zero, empty, null, default-initialised, out-of-range – triggers the failure), the review MUST enumerate AND INDEPENDENTLY VERIFY the fix covers EACH factor / dimension. Verifying ONE factor and ASSUMING the rest are safe is the EXACT forbidden gap this anchor closes (the forensic `channels` -verified / `format` -assumed pattern).

() **Prove every assumption.** "This can't happen" / "unreachable" / "always valid" / "the caller never sends X" / "the input is always in range" MUST be PROVEN from

the code + spec + captured evidence – NEVER accepted on faith, NEVER inferred from "it usually is" (§ . . no-guessing). An unproven assumption that later proves false is a § . REVIEW-LAYER bluff of the SAME severity class as a PASS-bluff. A review NIT that DISMISSES a path / branch as unreachable / dead / defensive-only MUST carry the PROOF of unreachability (the callers enumerated, the input domain bounded, the degenerate value shown genuinely impossible) – a bare "effectively unreachable" NIT with no proof is ITSELF a finding, not a dismissal.

() **Verify the fix against ANY captured runtime evidence.**

When captured runtime evidence of the failure EXISTS – a crash log / tombstone / stack trace / core dump / sink-probe report / the § . . runtime-signature / a reproducer – the review MUST confirm the fix ACTUALLY PREVENTS the CAPTURED failure (walk the fix against the EXACT captured conditions, not a hypothetical). A review that GO's a fix which then FAILS at runtime on a scenario PRESENT in the evidence (or a scenario the review CHOSE NOT TO CONSIDER) is a bluff – the evidence was available and not cross-checked.

() **All angles.** The review MUST cover ALL of: correctness against the stated goal; the FULL input domain (every factor of a multi-factor computation, every boundary, every degenerate value); every downstream / interacting feature – ESPECIALLY must-not-break constraints (per § . . target / hardware safety, and any feature the change shares state / timing / a resource with – e.g. audio / multichannel / passthrough integrity); the TESTS' OWN coverage of every dimension (a test that checks only ONE factor of a multi-factor bug is ITSELF a coverage gap the review MUST catch and require closed); and the § . paired mutation covering EACH dimension (a mutation that flips only one factor leaves the others un-guarded).

() **Honest boundary** (§ . .). "Exhaustive" means the ENUMERATED + EVIDENCE-GROUNDED scenario space – NOT a claim of literal omniscience. A dimension that genuinely cannot be analysed (external state the review cannot inspect, a runtime path with no reproducer) MUST be stated as an EXPLICIT gap (UNCONFIRMED: / UNKNOWN: / PENDING_FORENSICS: with a tracked follow-up) – NEVER silently assumed safe. A silently-assumed-safe un-analysed dimension is the § . . violation.

() **The quality-harvest review classes** (research-derived extension, – – – operator mandate: every check, diagnostic method, and preventive measure from the reopen/first-touch root-cause forensics is checked through EVERY code-review process). In addition to ()-(), every review MUST explicitly CHECK, per reviewed change:

(a) **closure-evidence class vs defect layer**

(§ . .) – no terminal or ready-class status in the change's scope moves on evidence whose class ranks below the defect layer's floor; a user-visible property "verified" by file-existence checks plus greps, an echo-assertion, or a state-only repair closing a defect as fixed is a FINDING; (b) **needled measurements** – every load-bearing measurement the change or its tests performs satisfies § . . ()-() (a reported absence cites its class-matched control needle), and NO count was converted into a finding without reading the underlying lines – a COUNT IS A LEAD, THE LINES ARE THE FINDING (forensic FACT: twice in one measured cycle a count was read as state – a "mixed half-written doc" that was actually coherent with a labelled historical citation, and an inflated FAIL total that included the summary line counting itself); (c) **governance-and-gate integrity** (§ . .) – a change that NAMES a gate lands its implementation or a registered deferral in the same change, and a governance edit preserves anchor-block integrity (exactly-once per mirror, lockstep-identical, no

anchor-number collision); (d) **reviewer-authored adversarial mutations** – the author's self-authored mutation set is NECESSARY but NOT SUFFICIENT validation of a guard/test: forensic FACT (from the consuming project's dated forensic corpus cited in the – – root-cause analysis): an author's six self-authored mutations were ALL caught while an independent reviewer's FIRST THREE different mutations ALL sailed through, each restoring a user-visible defect. The reviewer MUST attempt at least one mutation the author did not write (or record honest infeasibility per § 11.4.12); a guard whose reviewer-authored mutation survives is unvalidated instrumentation per § 11.4.13 (F) and a FINDING.

Classification: universal (§ 11.4.14) – a platform-neutral review-rigor discipline reusable by ANY project; the consuming project supplies its concrete input registries, must-not-break constraint set, and captured-runtime-evidence sources per § 11.4.15. Composes § 11.4.16 (the code-review-agent gate this strengthens) / § 11.4.17 (iterate-until-GO – a § 11.4.18 gap is a finding that re-arms the loop) / § 11.4.19 (every-change-reviewed – § 11.4.20 raises the review's PRECISION bar) / § 11.4.21 (no-guessing – the operative discipline behind clauses 11.4.22 + 11.4.23) / § 11.4.24 (runtime-signature – the captured evidence clause 11.4.25 cross-checks against) / § 11.4.26 (validate-the-fix-first-after-redeploy – the runtime confirmation that would have caught "the last one") / § 11.4.27 (self-validated analyzer – the multi-factor coverage discipline applied to the review's OWN test-coverage check). Propagation gate CM-COVENANT-114-194-PROPAGATION (literal 11.4.194 across the consumer fleet) + recommended gate CM-EXHAUSTIVE-REVIEW-ALL-SCENARIOS (a review GO must record the enumerated scenario / factor space + the proof of every dismissed-as-unreachable branch + the cross-check against captured

runtime evidence where it exists) + paired § . meta-test mutation (strip the enumeration / the assumption-proof / the captured-evidence cross-check from a review record → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
§ . . .

Non-compliance is a release blocker regardless of context. No escape hatch – no --partial-scenario-review , --assume-unreachable-OK , --skip-captured-evidence-crosscheck , --single-factor-review-suffices , --dismiss-branch-without-proof flag.

2 0 2 6 0 7 1 4

§ . . – Branch-structure
governance: taxonomy + merge-after-live-QA + flavor/product non-merge (User mandate, – –)

Forensic anchor – verbatim operator rules

(– –):

1 1 4 181 1 1 4 167

2

- Branch taxonomy: feat/<NAME_OF_FEATURE> , product/<NAME_OF_PRODUCT> , flavor/<NAME_OF_FLAVOR> (extends § . . / § . . (D)).
- Tracks on FEATURE branches or main MUST – after full confirmation + validation + verification + full LIVE MANUAL TESTING (§ . .) – MERGE their work to main COMPLETELY (all submodules + main repo), ff-only, NEVER force-push (§ . .), no commit lost (§ .) (extends § . . (I) / § . .).
- Tracks with an entirely-new FLAVOR/PRODUCT do ALL work on a flavor branch AND/OR product branch (both may exist) – this product-line work does NOT merge to main the same way.

- . Commit + push ALL branches to ALL upstreams (\$.), ff-only.

Every controlled branch in a governed repository MUST belong to one of THREE canonical taxonomy classes, and each class carries a distinct integration lifecycle:

(A) Taxonomy (extends \$. . / \$. . (D)). A branch's PREFIX declares its class and its lifecycle:

Prefix	Class	Integration lifecycle
feat/<NAME_OF_FEATURE>	FEATURE	merges to main after full live-manual-QA (clause B)
product/<NAME_OF_PRODUCT>	PRODUCT line	does NOT merge to main the standard way (clause C)
flavor/<NAME_OF_FLAVOR>	FLAVOR line	does NOT merge to main the standard way (clause C)

The <NAME_*> slug is lowercase snake/kebab (\$. .). The canonical branch NAME for a feature / logical group is minted ONCE and used IDENTICALLY across the main repo AND every touched owned submodule (\$. .) – recorded once in the \$. . exactly-once claim registry / the \$. . workable-items logic_group → branch mapping and LOOKED UP, never re-invented (\$. .). An untouched submodule stays on its base branch. The prefix STRINGS above are the operator-canonical form; a project already using a different feature-prefix convention (e.g. feature/...) is compliant so long as that convention is the ONE canonical name for the group per \$. . – reconciling an inconsistent set of names is the \$. . risk-

free MERGE-onto-canonical path (ff-only, no commit lost, NEVER a rename that drops commits), never a silent rename (§ . .).

(B) FEATURE / main tracks – MERGE-TO-MAIN-AFTER-LIVE-QA. A track working on a `feat/*` branch (or on `main` directly) MUST, before merging its work to `main`, satisfy IN ORDER: full confirmation + validation + verification (automated four-layer per § . . (b) + § . . full-suite retest) THEN full **LIVE MANUAL TESTING** by the QA team (§ . . – the FINAL human sufficiency gate; automated GREEN is necessary but NOT sufficient). Only after § . . confirmation may the merge proceed, and it MUST be COMPLETE – the main repo AND every touched owned submodule – integrated **ff-only** onto the latest `main` (§ . . merge-onto-latest-main), **NEVER force-push** (§ . .), with **NO commit lost / NO history rewrite** (§ . / § . . step : union, zero conflict markers, zero dropped files), pushed to **ALL upstreams** (§ . , ff-only). This is § . . (I) (no-merge-until-approved) + § . . (regular trunk-integration keeps the branch mergeable) – § . . binds the § . . live-manual-QA gate as an explicit precondition of the approval-gated back-merge.

(C) FLAVOR / PRODUCT lines – DO-NOT-MERGE-TO-MAIN-THE-SAME-WAY. A track producing an ENTIRELY-NEW flavor and/or product line (a distinct downstream product, not a feature destined for the trunk) does ALL its work on a `product/<NAME>` AND/OR `flavor/<NAME>` branch (BOTH may exist for one line – a product line may carry a product branch AND a flavor branch simultaneously). This product-line work does NOT flow back into `main` via the clause-B merge – it is a divergent line with its own release lifecycle. It STILL obeys § . . (no force-push), § . . (no commit lost), § . . (push all branches to all upstreams,

ff-only), and \$. . . (it MAY still merge main INTO itself to stay close to trunk – trunk-into-line is permitted and encouraged; it is the line-into-trunk direction that clause B gates and clause C withholds for product/flavor lines).

(D) Push discipline (`$ . / $ . .`). Commit + push ALL branches (feature, product, flavor) to ALL configured upstreams, ff-only, NEVER force-push. New branch creation is inherently ff (a new ref loses nothing, `$ .`). Branch creation from an existing HEAD (`git branch <name> <sha>`) is non-disruptive (does not touch the working tree / index / other refs) and MAY proceed even in a dirty working tree under `$ . .` provided no commit is in-flight (`.git/index.lock` absent, no live `git-commit / pack-objects`).

Honest boundary (\$. .): this anchor governs branch STRUCTURE + integration DIRECTION; it does NOT by itself prove any branch's work correct (each change still crosses \$. . runtime-signature + \$. . retest + \$. . manual-QA).

Classification: universal (§ 1.1) – a platform-neutral branch-structure discipline reusable across any multi-track project; it references no project-specific hardware / vendor / region / asset, and is a strict generalisation of § 1.2 (one-feature-one-canonical-branch-name) + § 1.3 (D)/(I) (feature work-stream lifecycle) + § 1.4 (regular main→feature merge cadence) + § 1.5 (manual-QA final confirmation). Composes § 1.6 / § 1.7 / § 1.8 / § 1.9 / § 1.10 / § 1.11 / § 1.12 / § 1.13 / § 1.14 / § 1.15 / § 1.16 / § 1.17 / § 1.18 / § 1.19 / § 1.20 / § 1.21 / § 1.22 / § 1.23 / § 1.24 / § 1.25 / § 1.26 / § 1.27 / § 1.28 / § 1.29 / § 1.30 / § 1.31 / § 1.32 / § 1.33 / § 1.34 / § 1.35 / § 1.36 / § 1.37 / § 1.38 / § 1.39 / § 1.40 / § 1.41 / § 1.42 / § 1.43 / § 1.44 / § 1.45 / § 1.46 / § 1.47 / § 1.48 / § 1.49 / § 1.50 / § 1.51 / § 1.52 / § 1.53 / § 1.54 / § 1.55 / § 1.56 / § 1.57 / § 1.58 / § 1.59 / § 1.60 / § 1.61 / § 1.62 / § 1.63 / § 1.64 / § 1.65 / § 1.66 / § 1.67 / § 1.68 / § 1.69 / § 1.70 / § 1.71 / § 1.72 / § 1.73 / § 1.74 / § 1.75 / § 1.76 / § 1.77 / § 1.78 / § 1.79 / § 1.80 / § 1.81 / § 1.82 / § 1.83 / § 1.84 / § 1.85 / § 1.86 / § 1.87 / § 1.88 / § 1.89 / § 1.90 / § 1.91 / § 1.92 / § 1.93 / § 1.94 / § 1.95 / § 1.96 / § 1.97 / § 1.98 / § 1.99 / § 2.00 / § 2.01 / § 2.02 / § 2.03 / § 2.04 / § 2.05 / § 2.06 / § 2.07 / § 2.08 / § 2.09 / § 2.10 / § 2.11 / § 2.12 / § 2.13 / § 2.14 / § 2.15 / § 2.16 / § 2.17 / § 2.18 / § 2.19 / § 2.20 / § 2.21 / § 2.22 / § 2.23 / § 2.24 / § 2.25 / § 2.26 / § 2.27 / § 2.28 / § 2.29 / § 2.30 / § 2.31 / § 2.32 / § 2.33 / § 2.34 / § 2.35 / § 2.36 / § 2.37 / § 2.38 / § 2.39 / § 2.40 / § 2.41 / § 2.42 / § 2.43 / § 2.44 / § 2.45 / § 2.46 / § 2.47 / § 2.48 / § 2.49 / § 2.50 / § 2.51 / § 2.52 / § 2.53 / § 2.54 / § 2.55 / § 2.56 / § 2.57 / § 2.58 / § 2.59 / § 2.60 / § 2.61 / § 2.62 / § 2.63 / § 2.64 / § 2.65 / § 2.66 / § 2.67 / § 2.68 / § 2.69 / § 2.70 / § 2.71 / § 2.72 / § 2.73 / § 2.74 / § 2.75 / § 2.76 / § 2.77 / § 2.78 / § 2.79 / § 2.80 / § 2.81 / § 2.82 / § 2.83 / § 2.84 / § 2.85 / § 2.86 / § 2.87 / § 2.88 / § 2.89 / § 2.90 / § 2.91 / § 2.92 / § 2.93 / § 2.94 / § 2.95 / § 2.96 / § 2.97 / § 2.98 / § 2.99 / § 3.00 / § 3.01 / § 3.02 / § 3.03 / § 3.04 / § 3.05 / § 3.06 / § 3.07 / § 3.08 / § 3.09 / § 3.10 / § 3.11 / § 3.12 / § 3.13 / § 3.14 / § 3.15 / § 3.16 / § 3.17 / § 3.18 / § 3.19 / § 3.20 / § 3.21 / § 3.22 / § 3.23 / § 3.24 / § 3.25 / § 3.26 / § 3.27 / § 3.28 / § 3.29 / § 3.30 / § 3.31 / § 3.32 / § 3.33 / § 3.34 / § 3.35 / § 3.36 / § 3.37 / § 3.38 / § 3.39 / § 3.40 / § 3.41 / § 3.42 / § 3.43 / § 3.44 / § 3.45 / § 3.46 / § 3.47 / § 3.48 / § 3.49 / § 3.50 / § 3.51 / § 3.52 / § 3.53 / § 3.54 / § 3.55 / § 3.56 / § 3.57 / § 3.58 / § 3.59 / § 3.60 / § 3.61 / § 3.62 / § 3.63 / § 3.64 / § 3.65 / § 3.66 / § 3.67 / § 3.68 / § 3.69 / § 3.70 / § 3.71 / § 3.72 / § 3.73 / § 3.74 / § 3.75 / § 3.76 / § 3.77 / § 3.78 / § 3.79 / § 3.80 / § 3.81 / § 3.82 / § 3.83 / § 3.84 / § 3.85 / § 3.86 / § 3.87 / § 3.88 / § 3.89 / § 3.90 / § 3.91 / § 3.92 / § 3.93 / § 3.94 / § 3.95 / § 3.96 / § 3.97 / § 3.98 / § 3.99 / § 4.00 / § 4.01 / § 4.02 / § 4.03 / § 4.04 / § 4.05 / § 4.06 / § 4.07 / § 4.08 / § 4.09 / § 4.10 / § 4.11 / § 4.12 / § 4.13 / § 4.14 / § 4.15 / § 4.16 / § 4.17 / § 4.18 / § 4.19 / § 4.20 / § 4.21 / § 4.22 / § 4.23 / § 4.24 / § 4.25 / § 4.26 / § 4.27 / § 4.28 / § 4.29 / § 4.30 / § 4.31 / § 4.32 / § 4.33 / § 4.34 / § 4.35 / § 4.36 / § 4.37 / § 4.38 / § 4.39 / § 4.40 / § 4.41 / § 4.42 / § 4.43 / § 4.44 / § 4.45 / § 4.46 / § 4.47 / § 4.48 / § 4.49 / § 4.50 / § 4.51 / § 4.52 / § 4.53 / § 4.54 / § 4.55 / § 4.56 / § 4.57 / § 4.58 / § 4.59 / § 4.60 / § 4.61 / § 4.62 / § 4.63 / § 4.64 / § 4.65 / § 4.66 / § 4.67 / § 4.68 / § 4.69 / § 4.70 / § 4.71 / § 4.72 / § 4.73 / § 4.74 / § 4.75 / § 4.76 / § 4.77 / § 4.78 / § 4.79 / § 4.80 / § 4.81 / § 4.82 / § 4.83 / § 4.84 / § 4.85 / § 4.86 / § 4.87 / § 4.88 / § 4.89 / § 4.90 / § 4.91 / § 4.92 / § 4.93 / § 4.94 / § 4.95 / § 4.96 / § 4.97 / § 4.98 / § 4.99 / § 5.00 / § 5.01 / § 5.02 / § 5.03 / § 5.04 / § 5.05 / § 5.06 / § 5.07 / § 5.08 / § 5.09 / § 5.10 / § 5.11 / § 5.12 / § 5.13 / § 5.14 / § 5.15 / § 5.16 / § 5.17 / § 5.18 / § 5.19 / § 5.20 / § 5.21 / § 5.22 / § 5.23 / § 5.24 / § 5.25 / § 5.26 / § 5.27 / § 5.28 / § 5.29 / § 5.30 / § 5.31 / § 5.32 / § 5.33 / § 5.34 / § 5.35 / § 5.36 / § 5.37 / § 5.38 / § 5.39 / § 5.40 / § 5.41 / § 5.42 / § 5.43 / § 5.44 / § 5.45 / § 5.46 / § 5.47 / § 5.48 / § 5.49 / § 5.50 / § 5.51 / § 5.52 / § 5.53 / § 5.54 / § 5.55 / § 5.56 / § 5.57 / § 5.58 / § 5.59 / § 5.60 / § 5.61 / § 5.62 / § 5.63 / § 5.64 / § 5.65 / § 5.66 / § 5.67 / § 5.68 / § 5.69 / § 5.70 / § 5.71 / § 5.72 / § 5.73 / § 5.74 / § 5.75 / § 5.76 / § 5.77 / § 5.78 / § 5.79 / § 5.80 / § 5.81 / § 5.82 / § 5.83 / § 5.84 / § 5.85 / § 5.86 / § 5.87 / § 5.88 / § 5.89 / § 5.90 / § 5.91 / § 5.92 / § 5.93 / § 5.94 / § 5.95 / § 5.96 / § 5.97 / § 5.98 / § 5.99 / § 6.00 / § 6.01 / § 6.02 / § 6.03 / § 6.04 / § 6.05 / § 6.06 / § 6.07 / § 6.08 / § 6.09 / § 6.10 / § 6.11 / § 6.12 / § 6.13 / § 6.14 / § 6.15 / § 6.16 / § 6.17 / § 6.18 / § 6.19 / § 6.20 / § 6.21 / § 6.22 / § 6.23 / § 6.24 / § 6.25 / § 6.26 / § 6.27 / § 6.28 / § 6.29 / § 6.30 / § 6.31 / § 6.32 / § 6.33 / § 6.34 / § 6.35 / § 6.36 / § 6.37 / § 6.38 / § 6.39 / § 6.40 / § 6.41 / § 6.42 / § 6.43 / § 6.44 / § 6.45 / § 6.46 / § 6.47 / § 6.48 / § 6.49 / § 6.50 / § 6.51 / § 6.52 / § 6.53 / § 6.54 / § 6.55 / § 6.56 / § 6.57 / § 6.58 / § 6.59 / § 6.60 / § 6.61 / § 6.62 / § 6.63 / § 6.64 / § 6.65 / § 6.66 / § 6.67 / § 6.68 / § 6.69 / § 6.70 / § 6.71 / § 6.72 / § 6.73 / § 6.74 / § 6.75 / §

COVENANT-114-195-PROPAGATION (literal 11.4.195 across the consumer fleet) + recommended gate CM-BRANCH-TAXONOMY-CLASS (every controlled branch's prefix is one of feat/ | product/ | flavor/ | the reserved bases main / master ; a product/ | flavor/. branch is NOT the source of a merge onto main ; a feat/ | main back-merge cites a § . . manual-QA confirmation) + paired § . mutation (a product/* branch merged onto main , OR a feat/* back-merge with no § . . citation → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
§ . . .

Non-compliance is a release blocker regardless of context.
No escape hatch — no --any-branch-prefix-OK , --merge-product-to-main , --skip-live-qa-before-merge , --force-push-branch , --partial-submodule-merge flag.

**§ . . — Native-alias-first
priority + per-alias real-signal limit/
subscription tracking + auto-rebind-on-
recovery + resource-detection-by-real-
identity-not-substring-match (User mandate,
— —)**

Verbatim operator mandate: use ALL operational claude-NATIVE aliases FIRST — as long as ≥ native alias is operational (NOT session-rate-limited, NOT weekly-limit-reached, NOT expired-subscription) — before ANY provider alias; track per alias the session rate-limit / weekly-limit / subscription expiry+renewal from REAL signals (never faked, § . .) and record when each becomes operational AGAIN; as soon as a higher-priority native recovers, rebind tracks to it; and fix the host-budget guard footgun where a pgrep -f substring-match false-REFUSED every worker spawn. The multi-track alias-binding

+ rate-limit/subscription-tracking layer of the § . . . ruler orchestration MUST implement ALL of: (A) **NATIVE-ALIAS-FIRST PRIORITY** – every OPERATIONAL native (subscription-based claudeN) alias is selected before ANY provider alias, UNCONDITIONALLY, as long as ≥ native is operational; the guarantee is STRUCTURAL – a native/provider CLASS partition scanned native-class-completely-before-provider-class (the § . . . resolve-by-CLASS-not-by-file-position discipline), so a provider can NEVER win while a non-cooled native exists, REGARDLESS of roster file order; within a class the operator-mandated order (native equal-capability, then providers deepseek→xiaomi→opencode→kimi→...) is authoritative decision-DATA (§ . . .) never re-invented (§ . . .). (B) **PER-ALIAS REAL-SIGNAL LIMIT + SUBSCRIPTION TRACKING** – an "operational?" state per alias derived from REAL captured signals, NEVER faked/guessed (§ . . .): the § . . . captured signature (isApiErrorMessage + apiErrorStatus:429 / api_error_status:429) sub-classified into a reason-CLASS closed set {session | weekly | subscription} from real message markers (reset/weekly-limit/session-limit); each limited alias records its CLASS + an operational-again EPOCH, and the three classes have DIFFERENT windows – session (self-heals at the reset, short), weekly-limit-reached (~until the weekly reset, long), subscription-expired (INDEFINITE – operational-again is UNKNOWN until a REAL renewal signal, parked with a far-future sentinel; § . . . NEVER guesses a renewal date, an operator/config supplies the real one). A limited alias is NEVER auto-selected until its epoch passes (session/weekly) or an explicit renewal/recovery signal clears it (subscription). (C) **AUTO-USE-ON-RECOVERY** – the moment a higher-priority native recovers (its window elapsed, or a renewal marked it operational), the mechanism rebinds tracks UPWARD to it (a promote -class

operation keyed off a comparable priority-rank: native class before provider, then config order), PRESERVING each track's worktree AND device leases (the leases key on the STABLE track id – the alias changes, not the track; \$. . . / \$. . .). Idempotent – a track already on the best operational alias is left untouched.

(D) RESOURCE-DETECTION BY REAL IDENTITY, NOT SUBSTRING-MATCH (the RB-guard footgun fix) – a host-safety "heavy work in flight" guard MUST identify a real resource (a heavy build, a live worker) by its ACTUAL process command line read from the OS (/proc/<pid>/cmdline on Linux), NEVER by a bare pgrep -f REGEX substring match that false-matches a process merely MENTIONING a token without BEING it – a pattern-CARRIER (a process whose args literally quote the whole alternation pattern; a claude / worker whose prompt embeds the token; a launcher/monitor/test), a same-tool ENGINE process, or the guard's own process/parent (the \$. . . self/sibling pgrep footgun – forensic FACT - - : the guard false-REFUSED EVERY \$. . . spawn on a host with ZERO real builds because a live claude worker's prompt quoted the pgrep pattern). The guard re-reads each matched PID's real cmdline and excludes those carrier/engine/worker/self classes; an UNREADABLE/vanished cmdline is conservatively counted as the real resource (REFUSE – the safe, reversible default, \$. . . / \$. . .), so a genuine build is still caught. **(E) ANTI-BLUFF (\$. . . / \$. . .)** – every piece carries four-layer coverage (pre-build gate + on-host test + paired \$. . . mutation + captured evidence per \$. . .), each with a \$. . . RED-polarity reproduction of the defect on the pre-fix artifact + \$. . . determinism; **honest boundary (\$. . .)** – this anchor mandates the MECHANISM + the REAL-signal derivation of "operational"; genuine live per-subscription quota-isolation proof still requires live tokens exercising real rate-limit boundaries

(an operator-scheduled acceptance window), never conflated with a mock/hermetic contract GREEN. Classification: universal (§). Composes § (the ruler orchestration this is the alias-binding + rate-limit layer of) / § (track+branch+alias identity label) / § (multi-track exactly-once claim + device-lock) / § (resolve-by-stable-name/CLASS not by index/file-position) / § (autonomous-decision-over-blocking – the bounded-park-when-all-cooled + the conservative-refuse safe default) / § (continuous parallel-stream routine) / § (no-guessing – REAL signals, never faked) / § (runtime-signature) / § (RED-polarity) / § (single-resource-owner – leases key on the stable track id) / § (RLIMIT_NPROC / process-footgun awareness – the pgrep self/sibling match is the same footgun class) / § / § (host-safety budget + concurrency – the guard the RB- fix protects). Propagation gate CM-COVENANT-114-196-PROPAGATION (literal 11.4.196) + recommended gate CM-ALIAS-PRIORITY-LIMIT-TRACKING (native-first CLASS partition present; per-alias reason-class + operational-again tracking; auto-rebind-on-recovery; the build-guard resolves by real cmdline not bare substring-match) + paired § mutation (swap the native/provider priority bases so a provider outranks a native → the recommended gate FAILS; flip the guard's strict-cmdline-filter to the pre-fix bare pgrep → its RB- test FAILS; strip the literal 11.4.196 from a carrier → the propagation gate FAILS; gate-code for the recommended gate = separate work item).

2026 07 15

2026 07 15

(F) CONFIGURED ≠ IN USE – the mechanism MUST BE WIRED TO THE ACTUAL WORKER SPAWNS, and its binding registry MUST REFLECT THE LIVE WORKER SET (extension, – –). Forensic FACT (– –): the alias-binding orchestrator was fully installed and its

configuration valid, yet its binding registry was EMPTY ((none bound)) while THREE live workers ran OUTSIDE it on provider-backed aliases – so the recovery promote operation rebound ZERO tracks. The mechanism existed and managed NOTHING. Therefore: () EVERY worker / agent spawn path MUST bind THROUGH the orchestrator (the § . . . launcher) – a worker started outside it is UNMANAGED (never rebound on recovery, never rate-limit-tracked, never counted against the host budget) and is a § . . . violation; () the mechanism's READINESS is defined by its RUNTIME SIGNATURE (§ . . .), NOT by its configuration – **the set of rows in the binding registry MUST EQUAL the set of live workers** (every live worker appears with its bound alias; every registry row corresponds to a live-or-honestly-terminal worker); an empty / partial registry alongside live workers is a FALSE-READY state; () claiming the mechanism "ready" / "in use" on the strength of a valid config + an installed script, WITHOUT that registry live-worker equality captured as evidence, is a § . . . / § . . . bluff at the orchestration layer – "configured" is NOT "in use"; () the recommended gate CM-ALIAS-PRIORITY-LIMIT-TRACKING therefore ALSO asserts the registry live-worker equality, with a paired § . . . mutation (leave a live worker unregistered → the gate FAILs). Composes § . . . (the spawn / resume substrate this binds), § . . . (the always-on default this makes REAL rather than nominal), § . . . (runtime-signature-as-definition-of-done), § . . . (a respawned agent re-binds through the orchestrator too).

Canonical authority: constitution submodule Constitution.md
 §

Non-compliance is a release blocker regardless of context. No escape hatch – no --provider-before-operational-native , --fake-limit-state , --guess-renewal-date , --no-rebind-on-

1.1 4 1.9.7
recovery , --substring-match-build-guard , --skip-real-signal-
classification flag.

§ . . – Research / kicked-off-work
completion mandate: every research effort MUST
be driven to full, wired, verified completion
or explicitly closed – never left un-wired in
the backlog (User mandate,
– –)

Forensic anchor – verbatim operator mandate
(– –):

"all research work we start MUST BE fully finished,
completely confirmed, validated and verified! It MUST
NOT happen for it to stay at the backlog uncompleted
and un-wired! Loss of requirements, failure to
incorporate them is FORBIDDEN!!! Everything MUST BE
fully completed! No exceptions!"

Every research effort, kicked-off improvement, design doc,
spike, prototype, proof-of-concept, or partially-landed
feature that a project STARTS MUST be driven to FULL
completion – OR explicitly, evidence-backed, CLOSED. There
is no third "sitting un-wired in the backlog" state. "We
researched it / designed it / started it" is NEVER a
terminal state; a docs/research/* document, a design doc, a
spike branch, or a half-wired feature that is neither
COMPLETED-and-wired nor explicitly-closed is a

§ . . violation of the same severity class as a
§ . PASS-bluff at the requirements-integrity layer –
the requirement was accepted, work began, and then it
silently evaporated. Loss of requirements and failure to
incorporate started work are FORBIDDEN.

COMPLETED means (ALL must hold): () **implemented** – the code / artefact exists; () **WIRED** – integrated into the system AND used by default, not merely present in the tree behind a dead flag / an un-called function / an unreferenced module / a never-selected path (mere-presence-without-wiring is the § 11.4.15 dead-code gap AND the § 11.4.16 layer- / SOURCE→ARTIFACT→RUNTIME gap – a change present in source that never reaches the running artifact, or reaches it but is never exercised, is NOT completed); () **confirmed + validated + verified** with captured physical evidence per § 11.4.17 / § 11.4.18 / § 11.4.19 on a clean deployment per § 11.4.20 runtime-signature – never a metadata-only / config-only / absence-of-error / grep-without-runtime claim (§ 11.4.21 / § 11.4.22); () tracked to a terminal Status per § 11.4.23 / § 11.4.24.

Explicitly CLOSED means: the effort is deliberately terminated with a documented, evidence-backed reason drawn from the existing closed vocabulary – § 11.4.25.

Obsolete (superseded-by-design-change / superseded-by-later-mandate / duplicate-of / not-reproducible / feature-removed), § 11.4.26 **structurally-impossible** (won't-fix, with cited impossibility proof), or § 11.4.27 **operator-confirmed drop** – NEVER a silent abandonment. "No time" / "lower priority" is **Operator-blocked** per § 11.4.28 (with the self-resolution-exhaustion audit) or a still-open tracked item, NOT a close.

Mechanical binding to the tracker (§ 11.4.29): every research effort / design doc / kicked-off improvement MUST map to at least one tracked workable item (ATM-NNN per § 11.4.30) whose completion is enforced by the loop. A docs/research/* document (or equivalent design / spike artefact) with NO tracked implementing item, OR a tracked item that has been stalled un-wired past the project's

cadence, is a § . . violation. The autonomous loop (§ . . / § . . / § . . / § . . / § . . / § . .) MUST treat every such un-wired research effort as an actionable queue item – driving it to COMPLETED or explicitly CLOSED is exactly the "keep working until the queue is genuinely empty" contract, never optional backlog.

Honest boundary (§ . .): "fully completed" is bounded by the completion + closure definitions above (wired + verified, OR evidence-backed closed) – it does NOT mean every speculative idea must ship; it means every STARTED effort reaches a documented terminal state, and the terminal state is proven, never assumed.

Classification: universal (§ . .) – a platform-neutral requirements-integrity discipline reusable by ANY project; the consuming project supplies its research-doc tree, path, + tracker, binding per § . . . Composes § . . (deep-web-research – the research this mandate insists on FINISHING, not merely starting) / § . . / § . . / § . . / § . . / § . . (the autonomous-loop queue-drain that drives un-wired efforts to completion) / § . . (discovery-pressure – enumerated coverage of what was started) / § . . (rock-solid-proof – completion PROVEN, never claimed) / § . . (dead/unwired-code investigate-before-remove – the wiring half) / § . . (deep multi-angle web research per change) / § . . / § . . / § . . (the explicit-close vocabulary) / § . . / § . . (the tracker binding) / § . . (wired-and-verified on a clean target) / § . . / § . . (terminal Status). Propagation gate CM-COVENANT-114-197-

PROPAGATION (literal 11.4.197) + recommended gate CM-RESEARCH-COMPLETION-TRACKED (every docs/research/* effort / design doc maps to an open-or-closed tracked item; no research doc lacks an implementing item; no tracked item is stalled un-wired past the cadence) + paired § . mutation (add a docs/research/* effort with no tracked implementing item, OR mark a research effort "done" while its feature is present-but-un-wired → the recommended gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md § . . .

Non-compliance is a release blocker regardless of context. No escape hatch – no --research-may-stay-unwired , --skip-completion-tracking , --abandon-without-close , --present-but-unwired-is-done , --backlog-research-OK , --lose-requirement-OK flag.

§ . . – Default working mechanisms: multi-alias native-first orchestration AND heavy token-optimization are ALWAYS-ON defaults for single- and multi-track work (User mandate, – –)

Forensic anchor – operator mandate (– –): multiple claude-toolkit aliases (native-first) AND heavy token-optimization MUST be the DEFAULT mechanisms, always used with single- and multi-track development and with single or multiple working agents.

Two proven mechanisms are promoted to STANDING, always-on defaults – engaged automatically for EVERY working configuration (single-agent OR multiple-agent; single-track OR multi-track), never a per-request opt-in:

(A) Multi-alias native-first orchestration is the DEFAULT. The claude-toolkit multi-alias mechanism – every

OPERATIONAL claude-NATIVE alias selected before ANY provider per § . . . (native-alias-first priority + per-alias real-signal limit/subscription tracking + auto-rebind-on-recovery), driven by the § . . . automatic multi-track ruler orchestration – MUST be the default execution substrate for ALL work, NOT only for the multi-track case. A SINGLE-track or SINGLE-agent session STILL binds its worker through the same native-first alias-priority mechanism (so it uses the best operational native alias + rebinds on recovery), and a MULTI-track / multiple-agent configuration uses it across every track (§ . . . exactly-once claim + § . . . track+branch+alias identity label). Running a session on an arbitrary / hardcoded / provider-first alias when an operational native exists – or bypassing the alias-priority mechanism because "it is just one agent" – is a § . . . violation.

(B) Heavy token-optimization is the DEFAULT. The § . . . token-efficiency discipline (compact-in-CLAUDE.md/AGENTS.md + full-one-hop-away anchor layout, targeted reads over whole-file dumps, subagent context isolation per § . . . / § . . . , evidence-hash references not evidence dumps) MUST be applied ALWAYS – on every read, every dispatch, every doc edit – for single- AND multi-track work, with single OR multiple agents. Token-heavy working patterns (re-reading whole large files when a targeted slice suffices, dumping full evidence into context, monolithic non-subagent execution of parallelisable work) when a token-efficient path exists are § . . . violations.

Honest boundary (§ . . .): "default mechanism" means engaged automatically without an operator request; it does NOT weaken any host-safety bound – the multi-alias / multi-agent fan-out stays under the § . . . % memory ceiling + the § . . . thread-headroom check + the

§ . . parallel-agent cap (spawning contending / redundant agents is NOT "using the default", per § . . maximum-USEFUL-throughput).
 § . . BINDS the two mechanisms as always-on; it does NOT redefine them – their mechanics live in § . . / § . . / § . . .

Classification: universal (§ . .). Composes
 § . . (token-efficiency – mechanism B) /
 § . . (automatic multi-track ruler orchestration – the substrate) / § . . (native-alias-first priority + limit tracking – mechanism A) /
 § . . (track+branch+alias identity) /
 § . . (multi-track exactly-once claim) /
 § . . (maximal multi-agent utilization – the bounded-throughput sibling) / § . . (continuous parallel-stream routine) / § . . (parallel-development PWU) / § . . / § . . (subagent-driven default) / § . . (default autonomous-loop) / § . . / § . . (host-safety bounds). Propagation gate CM-COVENANT-114-198-PROPAGATION (literal 11.4.198) + recommended gate CM-DEFAULT-MECHANISMS-ALWAYS-ON (the native-first alias-priority mechanism + the § . . token-optimization discipline are wired as the default for single- AND multi-track / single- AND multi-agent configurations) + paired § . . mutation (make provider-first / bypass-alias-priority the single-agent default, OR strip the token-optimization default → the recommended gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
 § . . .

Non-compliance is a release blocker regardless of context. No escape hatch – no --single-agent-skips-alias-priority , --

provider-first-default , --token-optimization-optional ,
--heavy-tokens-OK , --not-default-mechanism flag.

**§ . . – Exact-reproduction-
sequence mandate: when a working reproduction
exists, the investigation MUST use ITS exact
sequence – a deviating repro proves NOTHING
(research-derived, – –)**

Forensic anchor (FACT, – –). An investigation concluded that a defensive "cover" mechanism "never engages" and that the defect it guards was therefore unfixed. The conclusion was WRONG: the hand-rolled reproduction drove the target with a DIFFERENT input than the existing test did (it pressed a HOME key where the test sent a media-pause key), so the hand-rolled path never reached the PRECONDITION the mechanism keys off. Re-driven with the TEST'S EXACT sequence, the mechanism provably ENGAGED. A deviating reproduction had produced a confident, captured-looking, and entirely false negative.

Every investigation / debugging pass / validation that reproduces a defect or exercises a code path MUST satisfy ALL of:

() **Use the existing reproduction's EXACT sequence.** When a working reproduction ALREADY EXISTS – a test, harness, script, or recorded sequence that demonstrably reaches the defect or the code path under investigation – the investigation MUST drive the target with THAT EXACT sequence: same inputs, same order, same timing, same preconditions, same environment / topology (§ . .). A hand-rolled approximation is NOT a substitute.

() **A deviating variant MUST first PROVE it reaches the precondition.** If a hand-rolled / simplified / alternative sequence IS used, it MUST FIRST be proven – with captured

evidence – to reach the SAME precondition (the entry-state the mechanism or defect keys off) before ANY conclusion is drawn from it.

() A repro that misses the precondition proves NOTHING. It is NOT evidence of absence (§ 11.4.192 – absence of the observation is not observation of absence). Concluding "the defect is present" / "the defect is absent" / "the mechanism never engages" from a sequence that never reached the precondition is a § 11.4.193 / § 11.4.194 bluff at the INVESTIGATION layer – a FAIL-bluff when it wrongly condemns working code (and drives a needless "fix" into a healthy path), a PASS-bluff when it wrongly clears broken code.

() Record the exact driving sequence as captured evidence. Every reproduction MUST record the precise sequence it drove (commands / input events / order / timing) so it is replayable + auditable; any deviation from an existing reproduction MUST be explicit, justified, and precondition-proven.

STRENGTHENS § 11.4.195 (systematic-debugging Phase – reproduce), § 11.4.196 (reproduce-first), and § 11.4.197 (RED baseline on the broken artifact): those mandate THAT you reproduce; § 11.4.198 mandates HOW – with the sequence PROVEN to reach the defect, not one that merely resembles it. Classification: universal (§ 11.4.199) – the consuming project supplies its concrete input / driving mechanism per § 11.4.200. Composes § 11.4.201 / § 11.4.202 / § 11.4.203 / § 11.4.204 / § 11.4.205 / § 11.4.206 / § 11.4.207 / § 11.4.208 / § 11.4.209. Propagation gate CM-COVENANT-114-199-PROPAGATION (literal 11.4.199) + recommended gate CM-EXACT-REPRO-SEQUENCE (an investigation's conclusion cites the exact driving sequence used AND, where an existing reproduction exists, that the same

sequence was used OR that the deviating sequence reached the precondition with captured proof) + paired § . mutation (record a conclusion drawn from a deviating sequence with no precondition proof → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §

Non-compliance is a release blocker regardless of context. No escape hatch — no --approximate-repro-OK , --skip-precondition-proof , --hand-rolled-repro-suffices , --absence-proves-absence , --deviating-sequence-conclusion-OK flag.

§ — Non-targetable deploy/flash tooling MUST isolate exactly ONE eligible target and VERIFY-AFTER-WRITE on the INTENDED target (research-derived, — —)

Forensic anchor (FACT, — —). A flashing tool with no device-selection argument AUTO-SELECTED a device from the bus. With a sibling board still attached, it silently wrote the freshly-built image to the WRONG board, printed a SUCCESS message, and the INTENDED board kept its OLD build. Every downstream test then validated a target that had never received the artifact: the § ARTIFACT → RUNTIME-ON-CLEAN-TARGET layer was silently broken while every tool and gate reported green.

Whenever a deploy / flash / install / write tool CANNOT address a specific target (no serial / id / path selector, or one that the procedure does not pass), the procedure MUST satisfy ALL of:

() Targetability is a PRECONDITION. Before the write, the procedure MUST establish that the intended target is

the one that will be written: EITHER the tool takes an explicit stable target selector (serial / UUID / stable id per § . . . – never an enumeration ordinal) and it IS passed, OR the procedure GUARANTEES that EXACTLY ONE eligible target is present on the bus / host / network at the moment of the write (every sibling disconnected, powered down, or excluded), with the enumerated eligible-target set captured as evidence (size ==).

() **The tool's own success message is NOT proof.** A tool's "ok" / exit proves only that bytes reached the device the TOOL selected – NEVER that they reached the device the OPERATOR intended (§ . . . : "the write succeeded" ≠ "the write succeeded ON THE INTENDED TARGET"). Treating a tool's success line as deployment evidence is a § . bluff at the deployment layer.

() **VERIFY-AFTER-WRITE on the intended target.** After the write, the procedure MUST READ BACK the deployed artifact's IDENTITY from the INTENDED target ITSELF (build id / version / checksum / artifact fingerprint) and assert it EQUALS the artifact just written. A mismatch – or the OLD identity still present – is a FAIL, never a warning. This is the § . . . ARTIFACT → RUNTIME assertion applied to the deployment step, and it is the ONLY thing that closes the wrong-device write.

() **No validation on an unverified target.** Any test / validation run against a target whose deployed-artifact identity has NOT been verified is INVALID, and any PASS it produces is a § . . . PASS-bluff – it exercised code that may never have been deployed there (the § . . . clause- stale/wrong-artifact class).

Honest boundary (§ . . .): isolation + read-back prove that the INTENDED target holds the INTENDED artifact; they do NOT prove the artifact is CORRECT (that remains § . . . / § . . . / § . . .).

Classification: universal (§ . .) – the consuming project supplies its concrete deploy tool, target-identity read-back, and eligible-target enumeration per

§ . . . Composes § . . / § . . / § . . 11 4 129 / § . . (stable identifier, never an ordinal) / § . . (single-resource-owner) / § . . / § . . (writing an image to an unintended device is a TARGET/HARDWARE-safety event) / § . . . Propagation gate CM-COVENANT-114-200-PROPAGATION (literal 11.4.200) + recommended gate CM-DEPLOY-TARGET-ISOLATED-AND-VERIFIED (an explicit stable selector was passed OR the eligible-target set was captured with size == ; the post-write identity read-back from the intended target is present and matches) + paired § . mutation (deploy with two eligible targets on the bus and no explicit selector, OR strip the post-write identity read-back → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
§ . . .

Non-compliance is a release blocker regardless of context. No escape hatch – no --auto-select-target-OK , --trust-tool-success-message , --skip-post-write-verify , --validate-unverified-target , --multiple-targets-on-bus-OK flag.

§ 6 0 7 15 . . – Every guard/gate MUST assert the REAL condition: a false-positive refusal is a FAIL-bluff, a false-negative pass is a PASS-bluff (research-derived, – –)

Forensic anchor (FACT, – – – two independent instances in one cycle). (a) A host-budget guard used a bare pgrep -f '<pattern>' and REFUSED every

worker spawn on a host with ZERO real builds, because a live worker's own command line merely QUOTED the pattern (the \$. . (D) / \$. carrier footgun). (b) A host-safety pre-flight refused a build on a "swap full" reading that did not reflect a real memory-pressure condition, and cleared only after the operator manually cycled swap. In BOTH cases a guard refused on a PROXY signal that was NOT the real condition – halting real work while reporting a condition that did not exist.

Every GUARD / GATE / PRE-FLIGHT CHECK (host-safety, resource-budget, readiness, lock, capability, availability) MUST assert the REAL condition it CLAIMS to assert, resolved from the AUTHORITATIVE source – the process's real /proc/<pid>/cmdline, the real cgroup / meminfo pressure metric, the real lock-holder liveness (kill -0), the real device identity – NEVER from a proxy signal that something which is NOT the condition can satisfy.

() A FALSE-POSITIVE REFUSAL IS A FAIL-BLUFF

(\$. .). A guard that BLOCKS work when the condition is ABSENT is exactly as forbidden as a gate that PASSES while a defect is present: it halts real work, it teaches operators and agents to bypass guards (the worst possible second-order effect), and it hides the true state behind a confident false claim.

() A FALSE-NEGATIVE PASS IS A PASS-BLUFF (\$. .)

A guard that does NOT fire when the condition IS present is the classic bluff – the guard exists precisely to catch that state.

() SELF-VALIDATED (\$. . ())

Every guard MUST ship a golden-TRUE fixture (the real condition PRESENT → the guard FIRES) AND a golden-FALSE fixture that INCLUDES the known CARRIER / decoy case (the condition ABSENT but the proxy signal present → the guard MUST NOT

fire), both wired into the meta-test. A guard that fires on its golden-FALSE fixture is ITSELF the defect, and shipping a guard without both fixtures is a § 11.4.101.12 bluff at the guard layer.

() **CONSERVATIVE-SAFE DEFAULT ON AN UNRESOLVABLE SIGNAL** (§ 11.4.101.12 / § 11.4.101.12). When the authoritative source genuinely cannot be read, the guard takes the SAFE, reversible default (refuse for a destructive / high-blast-radius action; the safe side for a host- or hardware-safety guard) and SAYS SO HONESTLY. A conservative refusal WITH an honest "could not resolve the condition" reason is NOT a false positive; a refusal that CLAIMS a condition it never verified IS one (§ 11.4.101.12).

() **EVERY REFUSAL REPORTS ITS RESOLVED EVIDENCE.** A guard refusal MUST print the resolved evidence that triggered it (the real pid + cmdline, the real metric + threshold, the real lock holder + liveness result) so the refusal is auditable and a false positive is diagnosable in one step rather than one session.

() **SCOPE – EVERY LOAD-BEARING MEASUREMENT, NOT ONLY GUARDS** (extension, research-derived, 2026-07-17). Clauses ()-() bind guards / gates / pre-flights – decision seams that REFUSE or ALLOW – and their failure taxonomy is refusal-shaped. But the same defect kills an ad-hoc measurement that refuses nothing: a grep, a count, a parse, an exit-code check, a string scan, a burn-down metric, an artifact inspection. These are the instruments an agent forms BELIEFS from, and a false belief propagates into every downstream decision without ever passing a guard. Therefore clauses ()-() bind **every measurement whose result is load-bearing** – i.e. any observation an agent, gate, report, or metric relies on to assert a fact. Their taxonomy is MEASUREMENT-shaped: a **FALSE-NULL** (a zero / empty / no-match / "absent" read as evidence of absence when the instrument simply could not see) and a

FALSE-MATCH (a hit on something that MENTIONS the thing rather than IS it) are both § . /§ . . bluffs at the measurement layer – the null one is the more dangerous, because a broken instrument and a genuinely-clean artifact return the identical, confident, quiet zero.

() **CARRIER-VS-THING + THE CONTROL NEEDLE. (a) MATCH STRUCTURE, NOT SUBSTRING.** A measurement MUST distinguish the THING from a CARRIER that merely MENTIONS it. Forensic FACTs, all genericised, all real, all in ONE cycle of ONE project: a removal's completeness was measured by counting references to the removed component – and the count matched the AUDIT-TRAIL COMMENTS documenting the removal, so a completed removal read as a failed one (the § audit trail NAMES what it removed; naming is not using); a "does the artifact contain X" scan matched the governance text DOCUMENTING X; and this anchor's own siblings had already booked the same shape twice – § . . (D)'s process-carrier (a process quoting the pattern is not a process being the thing) and § . 's self-match (the scanner's own command line contains the pattern it scans for). **A defect already anchored twice and still recurring is the argument for the general rule, not against it.** The instance that decides it: a mutation-residue scan run over a governance repository flagged the very anchor that DEFINES the mutation markers, because that anchor quotes them in its rule text – the enforcement mechanism false-positived on its own law. **(b) A NULL IS NOT EVIDENCE UNTIL A CONTROL NEEDLE PROVES THE INSTRUMENT CAN SEE.** Before a zero / empty / no-match / "absent" result may be reported as a real absence, the SAME instrument, over the SAME path, against the SAME artifact, MUST be shown to return a non-null for a **control needle** – a string/property KNOWN to be present. Needle NOT found ⇒ the instrument is blind and the zero says NOTHING about the artifact (§ . . :

report the blindness, never the absence). **The needle MUST be expressed through the SAME load-bearing query features as the query it certifies** – same dialect constructs (alternation / character class / quantifier / escape), same anchoring, same quoting, same encoding path. A bare-literal needle certifies ONLY the layers a bare literal crosses; it does NOT certify a query whose meaning depends on a feature the needle never used. **This bound is not theoretical – the () (c) dialect FACT refutes the naive form directly:** a query intending alternation but written with an escape that the dialect treats as a LITERAL pipe searches a string nobody wrote and returns zero; a plain-literal needle sails through the identical tool + path, returns non-null, and would "certify" that zero – while the query is still blind. A needle carrying the SAME alternation (one arm known-present) dies with the query and exposes it. Hence the earned inference is bounded: **needle found ⇒ the path can see the needle's QUERY CLASS ⇒ a zero from a query of that same class is evidence;** a zero from a query using features the needle did not exercise is NOT yet evidence, and needs a needle of its own class. This generalises § . . . (F)'s validate-the-detector principle (a guard never observed FAILing on the genuinely-broken artifact is unvalidated instrumentation and mints no verdicts) from guards minting verdicts to every measurement, and it is PER-QUERY, against the artifact in hand – **golden fixtures cannot substitute** (§ . . . (D)): "goldens are necessary, not sufficient"), because an instrument can pass every fixture on the authoring host and still be blind on the real path.

(c) THE PATH IS PART OF THE INSTRUMENT. The instrument is not the tool, it is the tool PLUS every layer the query crosses, and the load-bearing failures live in the layers: a quoted pattern whose escapes are consumed by an intervening/wrapper shell, so the far side searches a LITERAL where alternation was intended and matches

nothing; `\|` under an extended-regex dialect where it is a LITERAL pipe, not alternation, silently searching for a string nobody wrote; a pipeline (`check | head , check | xargs`) whose exit status is the LAST stage's, so a checker's verdict is discarded and its "clean ✓" is the pager's success, repeatable indefinitely without ever checking anything; a pattern anchored to line-start that misses an indented or colour-escaped line, reading a present verdict as no-verdict; a binary/encoding-sensitive extractor that splits a multi-byte character and reports a present string as absent. Each of these returns a CLEAN, CONFIDENT, WRONG answer – none crashes, so § . . 's script-bug clause (which covers failures that crash into a FAIL) never fires, and § . . 's parse-check (which is parse-time only) cannot see it. Hence (.)(b) is not optional diligence: a class-matched control needle is the only mechanism that catches these defects, because it is the only one that exercises the whole path with the query's own features. **Honest boundary** (§ . .): the needle proves the path can SEE its own query class – it does NOT prove the query is semantically the RIGHT question to ask (a correct-and-visible query for the wrong property still misleads); that remains ()'s and § . . ()'s concern.

() **A METRIC MUST BE VALIDATED AGAINST THE DEFINITION OF DONE.** A progress / burn-down / coverage / ratchet metric is a load-bearing measurement, and a metric that is ANTI-CORRELATED with done is not a misread condition – it is the WRONG condition, read correctly. Forensic FACT (genericised): a removal's burn-down counted references to the removed component and was tracked toward zero – but the CORRECT end-state ADDS references, because a retired gate must carry a mandated rationale naming what was removed and why, and the inverse "stays-removed" assertion must NAME the removed thing to assert its absence; the number therefore RISES as the work succeeds, the target is

unreachable, and chasing it to zero would DELETE the very rationale and assertions the constitution mandates.

Therefore: before a metric may gate, rank, or report work, it **MUST** be checked against the definition of done – construct the correct end-state and evaluate the metric there; **if the correct end-state does not move the metric to the target, the metric is INVALID and MUST NOT gate work** (a metric whose target the correct outcome cannot reach will, if enforced, drive work AWAY from correctness – the strongest possible form of § 1.1.1's proxy-signal defect, because here the proxy is not merely satisfiable by a non-condition, it is REFUTED by the condition). This is the () (a) carrier defect relocated from a guard to a metric: the audit-trail rationale and the stays-removed assertion are CARRIERS, and the burn-down was a substring-counter that carriers satisfy. **Honest boundary + standing obligation** (§ 1.1.2): this constitution already ships at least one gating metric of exactly this class – § 1.1.3's adoption ratchet, a monotone-decrease burn-down that BLOCKS the release tag (§ 1.1.4's seam; NOT the build) – and its validation against the definition of done is OWED, not hereby claimed; naming it is the § 1.1.5 discovery-pressure obligation. That validation **MUST be tracked** as a § 1.1.6 work item by every consuming project that operates a ratchet (stated as the obligation it is – this anchor does NOT assert that any particular project has already filed it; asserting an unverified tracker state would be the § 1.1.7 violation this very clause exists to forbid). **Interim precedence** (§ 1.1.8 – **stated, never left for a consumer to guess**): until that validation lands, § 1.1.9 ()'s ratchet **KEEPS GATING** – clause () does NOT retroactively disarm it, because an unvalidated ratchet that blocks a release is a bounded, visible, operator-resolvable stop, whereas disarming the only adoption mechanism would

silently re-open the unguarded-backlog hole
§ . . . exists to close (the § . . .
reversible-safe choice). Clause () binds (a) every NEW
metric from this anchor forward – none may gate until
validated against the definition of done – and (b) the
ratchet's own validation as owed work, NOT as a licence to
ignore it. A consuming project MUST NOT cite clause ()
to switch its ratchet off.

GENERALISES § . . . (D) (resource-detection-by-
real-identity – from the alias / build guard to EVERY
guard) and § . . . (holder-liveness PROVEN via
kill -0 before reaping – to every guard's condition).
Classification: universal (§ . . .) – the consuming
project supplies its concrete guards, authoritative
sources, and carrier/decoy fixtures per §
Composes § . . . / § . . . / § . . . /
§ . . . () / § . . . /
§ . . . / § . . . /
§ . . . (D) / § . . . / § . . . /
§ Propagation gate CM-COVENANT-114-201-PROPAGATION
(literal 11.4.201) + recommended gate CM-GUARD-ASSERTS-REAL-
CONDITION (every guard resolves its condition from the
authoritative source, ships golden-TRUE + golden-FALSE-
with-carrier fixtures wired into the meta-test, and prints
its resolved evidence on refusal) + paired § .
mutation (replace a guard's authoritative resolution with
a bare substring / proxy match → its golden-FALSE carrier
fixture makes the gate FAIL; strip the literal → the
propagation gate FAILs; gate-code = separate work item).

Gate contract for clauses ()-() (extension,
– –). CM-GUARD-ASSERTS-REAL-CONDITION is
EXTENDED beyond its guard-fixture invariants to cover the
measurement layer, and a metric-specific sibling is added
– because a clause landing with no gate + no paired
mutation is itself the § . bluff-gate pattern this

anchor exists to forbid: (i) every load-bearing NULL/ZERO/absence result reported as a finding cites a control-needle result of the SAME query class, captured alongside it (a reported absence with no class-matched needle evidence → FAIL); (ii) the needle's query class is asserted to match the certified query's load-bearing features (a bare-literal needle certifying an alternation/anchored/encoded query → FAIL – the finding- defect this clause was corrected for); (iii) recommended gate CM-METRIC-VALIDATED-AGAINST-DONE – every metric that gates, ranks, or reports work carries a recorded definition-of-done evaluation showing the correct end-state reaches the target (a gating metric with no such record, or one whose recorded correct-end-state does NOT reach target, → FAIL). Paired \$. mutations: swap a class-matched needle for a bare literal on a dialect-dependent query → (ii) FAILs; strip the needle entirely from a reported absence → (i) FAILs; register a metric whose correct end-state moves it AWAY from target → (iii) FAILs. Gate-code = a separate work item; this is the contract, not a claim that the code has shipped (\$. .).

() FIELD-IDENTITY + UNITS + SEMANTICS – capacity is not occupancy (extension, research-derived,
 – –). Forensic FACT (genericised,
 – –): a load-bearing latency claim – "a standing ~ . -second send queue" – was minted by reading a socket diagnostic's WINDOW-CAPACITY field (rcv_wnd bytes, the ceiling the peer ADVERTISES) as if it were QUEUE OCCUPANCY; the real observed Send-Q was a few hundred bytes (in-flight-sized, app_limited), three orders of magnitude smaller. The misread propagated into THREE parallel agents' briefs and into operator reports before anyone re-read the raw field; two INDEPENDENT re-derivations from the raw samples then converged on the correction, and the false figure had to be publicly withdrawn from a changelog. Therefore: before ANY metric

read from a summary / diagnostic line / dashboard cell may gate a decision, rank work, or propagate into a brief / report / tracker, its FIELD IDENTITY (which field, exactly, in the authoritative raw output), its UNITS, and its SEMANTIC CLASS – capacity/limit vs occupancy/backlog vs rate vs cumulative counter vs high-water mark – MUST be verified against the authoritative source's own field definition (§ 1.1.1 latest-source where the tool's documentation defines it). CAPACITY≠OCCUPANCY is the named canonical instance; limit-vs-current, cumulative-vs-delta, and advertised-vs-measured are the same defect. This is the clause-() FALSE-MATCH at the metric layer – the instrument matched A number, not THE quantity – and a downstream consumer of an unverified metric INHERITS the bluff (§ 1.1.2 : re-derive from the raw source before re-publishing; a figure carried between agents/briefs without its field-identity verification is a rumour with units).

() OBSERVER DECONTAMINATION – the instrument's own footprint is identified, excluded, and the decontaminated instrument re-proven seeing (extension, research-derived, 1.1.3 - -). Forensic FACT (genericised, 1.1.4 - -): a socket census read TIME-WAIT counts of - and nearly minted a leak finding – after excluding the SAMPLER'S OWN probe ports the real count was - (the census was ~ × observer-induced); the control needle proving the decontaminated instrument still SAW real state: the same scan did detect a genuine step-change (+ MB RSS, → sockets) in a monitored process. Every measurement whose instrument CONTRIBUTES to the measured population (probe connections in a socket census, sampler CPU/IO in a load profile, probe processes in a process census, trace overhead in a latency measurement – the § 1.1.5 self-match generalised from kill-lists to counts) MUST identify and EXCLUDE its own footprint before reporting, AND MUST prove via a () (b)-

class control needle that the decontaminated instrument still sees genuine signal (an over-broad exclusion is a self-inflicted clause-() false-null). Composes § . . . () (observer-effect budget – the recorder must not perturb the recorded; clause () adds the measurement-side half: where the footprint is unavoidable it is measured and SUBTRACTED, never reported as the system's own behaviour) and § . . . (single-resource-owner – a measurement arm racing a concurrent driver of the same resource produces cross-contaminated evidence; partition or serialize FIRST, as measured work this cycle did by refusing to touch a sibling stream's in-flight file region).

() **ARTIFACT-USABILITY, NOT PREREQUISITE-PRESENCE**
(extension, research-derived, - -).
 Forensic FACT (genericised, - -): an install/verify pipeline keyed its "can this host run the tool?" verdict on the presence of a BUILD prerequisite (a `command -v <compiler> probe`) – and turned a host holding a WORKING prebuilt binary into a false "cannot launch" FAIL, exit , in one transcript that showed BOTH "[ok] ... verified" AND "[FAILED] ... cannot launch" for the same artifact; the independent § . . . review graded it BLOCKING (it breaks every upgrade on hosts that install from prebuilts). A readiness / install / health / build gate whose CLAIM is about an ARTIFACT ("the tool works here") MUST probe the ARTIFACT ITSELF – execute it through its real invocation path and assert its observable behaviour (§ . . . (A)'s real-invocation-path discipline; § . . . 's verify-the-target sibling at the readiness seam) – NEVER the presence of a prerequisite for ONE of several ways of producing it (a compiler, an SDK, a package manager). The proxy fails in BOTH directions at once: prerequisite-present ≠ artifact-usable (a toolchain with no artifact) AND prerequisite-absent ≠

artifact-unusable (a prebuilt with no toolchain) – clause () and clause () in a single check.

() **SHELL-INSTRUMENT FOOTGUN CHECKLIST** – a standing ()(c) reference that shell-touching reviews **MUST** consult (extension, research-derived, – –). The measured shell-instrument footgun classes are catalogued in the standing reference docs/guides/shell_instrument_footguns.md (inherited by reference per \$. . /\$. . , never copied per project). Among the founding, DEMONSTRATED classes: timeout / nice / env -class EXEC WRAPPERS cannot run shell functions/ aliases ,(rc= "command not found" while the function is healthy – a clause-() false refusal; demonstrated – – under control conditions); pgrep -f / pkill -f matching the PROBING script's own command line (the \$. . (D)/\$. . carrier – re-demonstrated – – : a unique needle matched exactly ONE process, the probe's own harness); structural-delimiter text extractors (a column- -}-terminated awk function extractor), that TRUNCATE the extracted unit when its body legitimately contains the delimiter (here doc/ string content) and hand the test a PARTIAL artifact – a FALSE PASS when the truncated tail carried the code under test (demonstrated – – : the tail marker vanished from the extraction); and xtrace (bash -x) BLINDED when the code under test redirects the shell's stderr ,(exec ... 2>/dev/null silenced the trace mid-stream in the real – – incident – countermeasure: route the trace via BASH_XTRACEFD to a dedicated fd the code under test does not own). Every \$. . /\$. . code review of a shell-based test, gate, guard, or measurement **MUST** consult the checklist, and every zero/absence/PASS a shell instrument produces remains bound by ()(b)'s class-matched control needle. The checklist is a LIVING document: a newly-measured footgun class is APPENDED with its captured

demonstration (§ . .) in the same change-window, never carried as tribal memory (§ . . – the binding reference lives tracked, in-repo).

Gate contract for clauses ()-() (extension,
 – –). `CM-GUARD-ASSERTS-REAL-CONDITION` is
 EXTENDED: (iv) every load-bearing metric cited in a
 gating / ranking / report context carries its verified
 field-identity + units + semantic class (capacity /
 occupancy / rate / cumulative / high-water) traced to the
 authoritative source definition – an unverified-semantics
 metric gating a decision → FAIL; (v) every census/profile
 whose instrument contributes to the measured population
 records its observer-exclusion AND a post-exclusion
 control-needle result – an undecontaminated census minting
 a finding → FAIL; (vi) every artifact-readiness verdict
 cites a probe of the artifact's own invocation path – a
 prerequisite-presence-only readiness verdict → FAIL; (vii)
 every review record touching a shell-based instrument
 cites the § . . () checklist consultation.
 Paired § . mutations: relabel a capacity field as
 occupancy in a gating metric → (iv) FAILs; strip the
 observer-exclusion from a census fixture → (v) FAILs; swap
 the artifact probe for a bare `command -v <compiler>` while
 the fixture host carries a working prebuilt → (vi) FAILs.
 Gate-code = a separate work item; this is the contract,
 not a claim that the code has shipped (§ . .).

Canonical authority: constitution submodule `Constitution.md`
 § . . .

Non-compliance is a release blocker regardless of context.
 No escape hatch – no `--proxy-signal-OK` , `--substring-match-guard` , `--false-positive-refusal-acceptable` , `--unvalidated-guard` , `--refuse-without-evidence` , `--zero-proves-absence` , `--skip-control-needle` , `--substring-is-good-enough` , `--carrier-match-OK` , `--trust-the-null` , `--metric-without-done-check` , `--gate-on-unvalidated-metric` , `--capacity-is-occupancy-OK` ,

11 4 212
--skip-field-identity-check , --observer-footprint-included-OK ,
--prerequisite-presence-proves-artifact , --skip-shell-footgun-
checklist flag.

2026 07 15
§ . . . – Reporting directives: a report MUST
auto-create a fully-populated, fully-synced workable item
– never a prose acknowledgement (User mandate,

– –). Verbatim operator mandate: introduce
reporting directives ISSUE / BUG / TASK that create a
proper workable item filled with all details, with the
WHOLE workable-items system automatically up to date and
fully in sync – the main DB, all documents, and all
external tracking systems; universal, reusable on every
project, fully decoupled, recognized immediately by every
project that pulls the constitution. Every governed
project MUST expose REPORTING DIRECTIVES – the

§ . . . registered actions BUG (Type=Bug), TASK
(Type=Task), and ISSUE (generic report, CLASSIFIED into
the § . . . closed set) – such that a plain-
language report from the operator is TURNED INTO a real,
tracked, fully-synced workable item, automatically, every
time. A report that is discussed, acknowledged, answered
in prose, or "noted" WITHOUT landing a tracked item is a
§ . . . violation of § . . . PASS-bluff
severity at the requirements-intake layer: the requirement
was accepted and then silently evaporated (the

§ . . . loss-of-requirements failure, at its
entry point). () GRAMMAR (§ . . .). The three
directives are ordinary registry rows, so EVERY

§ . . . form works out of the box (BUG :: <text> ,
DEFAULT::BUG :: <text> , /DEFAULT::BUG <text> , BUG --->
<text>), PLUS a SIXTH form this anchor adds – the **single-
colon form** NAME: <text> (ISSUE: subtitles are late), the
shape operators actually type. The single-colon form is
REGISTERED-ACTION-ONLY by construction: a single-colon
token that is NOT a registered action is a NO-OP (an
ordinary prompt), NEVER an ASK and NEVER a sub-system

shortcut – because ordinary English prose (`TODO:` , `WARNING:` , `FIXME:`) shares that shape (`NOTE:` is now a registered \$. . severity marker), and routing it to the \$. . clarify path would question every such sentence (\$. . honest boundary: the other five forms keep their ASK-on-unknown behaviour – their shapes do not occur in prose). Host-command collisions are handled by the EXISTING \$. . conflict mechanism (`slash_bare: auto` + `slash_conflicts: [..]`), which already satisfies the operator's explicit-prefix requirement: a bare `/NAME` that collides with a built-in (e.g. the documented Claude Code built-in `/bug`) is NOT honored, while the namespaced `/DEFAULT::NAME` form is ALWAYS unambiguous – and the `NAME:` / `NAME ::` / `NAME --->` forms are never affected by any host command. () **THE ITEM.** The created item MUST carry, on ALL surfaces: a valid Status (\$. . , `Queued` at intake) + a Type from the CLOSED set {Bug | Feature | Task} (\$. . – `ISSUE` is a REPORTING CHANNEL, never a fourth type; inventing one is a violation) + a stable auto-incremented id (\$. .), and a COMPREHENSIVE structured description (\$. . D / \$. .): what / affected-scope / reproduction / acceptance. An undetermined section is recorded as an explicit `UNKNOWN:` gap – NEVER invented (\$. .). For `ISSUE` , the type is classified from the report's content and stated as FACT; if the content does not determine it, the operator is ASKED (\$. . / \$. .) BEFORE creation; ONLY when running autonomously where asking is impossible (\$. .) may the \$. . lowest-stakes ambiguity default `Task` be used – and then the defaulted classification MUST be recorded verbatim in the item and surfaced for reclassification, never silently asserted. () **THE FULL SYNC (the load-bearing half).** Creating the item is NOT sufficient: the directive MUST drive the WHOLE workable-items system into sync, in one

shot – (a) the item lands in the SQLite single-source-of-truth (\$₁₁ 4 65 . 11 4 106 / \$. .); (b) EVERY derived document is regenerated FROM the DB – the open + closed trackers, their summaries, and their .md / .html / .pdf / .docx siblings (\$. . / \$. . / \$. .); (c) the item is pushed to EVERY configured external tracker (\$. . D). () **NEVER A FAKED PUSH** (\$. . / \$. . / \$. .). A tracker whose credentials or whose client are absent is SKIPPED with an honest machine-readable reason (credentials_absent – names of the unset variables ONLY, never a value; tracker_client_absent – PENDING-OPERATOR-INPUT). A tracker PASS may be recorded ONLY after a real push command really exited . Fabricating, assuming, or silently omitting a push is a \$₁₁ 177 . bluff at the integration layer; a missing tracker NEVER blocks the item from landing in the DB + docs (the report is never lost because a mirror is unreachable). () **DECOUPLED ENGINE, CONSUMER-OWNED DATA** (\$. . / \$. .). The item-creation + full-sync engine lives in the constitution submodule (constitution/scripts/reporting/report_item.sh), is inherited BY REFERENCE (never copied), and carries ZERO project literals; every project-specific value – DB path, id prefix, sync command, tracker commands + their required env .vars – comes from a CONSUMER-OWNED config file (DATA, never an engine edit). The engine FAILS CLOSED with an actionable message when that config is absent – it never guesses a project's paths (\$. .). () **ANTI-BLUFF** (\$. . / \$. . ()). The mechanism ships four-layer coverage: a pre-build gate (CM-REPORTING-DIRECTIVES) whose FUNCTIONAL invariant SOURCES and RUNS the parser (never a grep-only assertion, \$. .), a paired \$. mutation test that proves each invariant is genuinely load-bearing, and an end-to-end suite run against a REAL SQLite DB with REAL binaries (no mocks,

§ . . .) that asserts a real row lands with the right Type/Status/id, that the docs are regenerated FROM the DB, and that an absent tracker SKIPS honestly – self-validated with a golden-good + golden-bad + negative-control fixture set (an oracle that passes its golden-bad fixture is itself the bluff). Classification: universal (§ . . .) – the consuming project supplies its DB path, id prefix, sync command, and tracker bindings per § Composes § (grammar) / § / § (DB SSoT) / § / § (status + type + id) / § / § (comprehensive description) / § / § / § (doc sync) / § (credentials) / § / § (clarify, never guess) / § (autonomous decision) / § (a started requirement never evaporates) / § / § (decoupling) / § (auto-registration on constitution pull) / § () / § Propagation gate CM-COVENANT-114-202-PROPAGATION (literal 11.4.202) + recommended gate CM-REPORTING-DIRECTIVES (registry declares ISSUE/BUG/TASK + the single-colon form is registered-only; the engine exists, is parse-clean, and is project-literal-free; the three directives really expand at runtime while prose does not; the honest-skip reasons exist and a tracker PASS is gated on a real exit) + paired § mutation (strip an action row, flip the single-colon form to non-registered-only, strip the parser branch, remove the prose NO-OP branch, make the engine fake a tracker push, or couple the engine to one project → the gate FAILs; strip the literal → the propagation gate FAILs).

Canonical authority: constitution submodule Constitution.md

§

Non-compliance is a release blocker regardless of context. No escape hatch – no `--report-without-item`, `--prose-acknowledgement-OK`, `--skip-item-creation`, `--skip-tracker-sync`, `--fake-tracker-push`, `--invent-a-fourth-type`, `--hardcode-project-in-engine` flag.

\$. . – Instant multi-stream resume engine: whole-fleet continuation state is a durable, content-addressed, atomically-committed snapshot resumed in $O(\text{changed})$ reads, not an $O(\text{history})$ re-read (User mandate, – –)

Compact summary: the continuation mechanism (`$ . . / $ . . / $ . . / $ . .`) is EXTENDED – never replaced – by a mechanical engine so a fresh session resumes ALL work (every Track, main stream, and agent) at once, INSTANTLY and at token cost $O(\text{changed streams})$, by reading ONE snapshot manifest + one compact blob per stream, NEVER re-reading the full handoff history. () **Content-addressed Merkle store** – each stream's canonical state bytes are named by their sha (dedup: unchanged streams share a hash → a global fleet snapshot costs $O(\text{changed})$); a snapshot is ONE manifest referencing every stream's content-hash (`$ . . . ( )`). () **$O(\text{streams})$ resume, $O(\text{history})$ NEVER on the resume path** – `RestoreAll` reads the manifest + one blob per stream; the append-only event ledger is audit + lost-update recovery only, never read to resume (`$ . .`). () **Tamper-evident + deterministic** – every blob is integrity-checked on read (bytes MUST hash to their id) and a deterministic round-trip (re-serialize must byte-match), no wall-clock in the hashed content, verified by a SELF-VALIDATING oracle: golden-good PASS + golden-bad detected FAIL + negative-control (a legitimately-older-but-valid snapshot) PASS – the false-

positive guard proving the verifier distinguishes a LAGGING copy from a TAMPERED one (\$. . / \$. . ()/\$. . ()). () **Durable + atomic + single-writer + crash-safe** – atomic ref write (temp→fsync→rename→dir-fsync, \$. . ()), append-only ledger (\$. . ()), single-writer-per-stream refusal (\$. .), advisory lock with provably-stale reap (kill -0 liveness, never steals a live lock, \$. . /\$. .). () **Project-agnostic engine, consumer-owned data** – the engine (github.com/vasic-digital/continuum , GitLab mirror) carries ZERO project literals, fails closed rather than guess a store path (\$. . /\$. .), and is inherited BY REFERENCE as a depth- .. submodule under the \$. . (C) carve-out (constitution/submodules/ continuum/ , helix-deps.yaml , zero own-org deps); every consumer supplies its store path / actor / stream naming as DATA. Anti-bluff (\$. /\$. .): the whole engine ships four-layer coverage (unit + integration + e..e..+ stress + chaos \$. . , race-clean) with a measured resume Metric (blobs_read == 1 + streams is the O(streams) proof) and the self-validating oracle as captured evidence (\$. . /\$. . / \$. .); honest boundary (\$. .): it snapshots the state an agent REPORTS, not its execution image (not a CRIU/durable-execution runtime, \$. .), and est_tokens is a documented bytes/4 heuristic, not a tokenizer count. Classification: universal (\$. .). Composes \$. . / \$. . / \$. . / \$. . / \$. . / \$. . / \$. . / \$. . / \$. . () / \$. . / \$. . / \$. . Propagation gate CM-COVENANT-114-207-PROPAGATION (literal 11.4.207) + recommended gate CM-CONTINUUM-RESUME-ENGINE-PRESENT (the

engine present under `constitution/submodules/continuum/`, `go test -race ./... green, continuum selfcheck = good=PASS/bad=FAIL/negctrl=PASS, zero project literals) + paired $. mutation (strip the literal → propagation gate FAILs; downgrade the resume path to read the ledger, or the oracle to pass its golden-bad → the mechanism gate FAILs; gate-code = separate work item).`

Canonical authority: constitution submodule `Constitution.md`
`$ . . .`

Non-compliance is a release blocker. No escape hatch – no `--resume-reads-history`, `--skip-integrity-check`, `--unvalidated-oracle-OK`, `--hand-typed-resume-state`, `--per-project-resume-reimpl` flag.

`$ . .` – **Operator-request-history document:** every project maintains a project-local, always-in-sync ledger of every operator request/prompt (content + timestamp+timezone + track + alias + model + effort) (User mandate, – –)

Verbatim operator mandate (– –): introduce a request-history document that captures, per operator request/prompt – the full request content, the accepted timestamp with an explicit timezone, the track that processed it, the alias that took the workable item, and the model + effort used – always kept in sync, with no false data and no bluff of any kind. **The operator's load-bearing correction** (`$ . . / $ . .`): the RULE (this universal discipline) belongs in the constitution submodule; the request-history DOCUMENT itself (project data) is PROJECT-LOCAL – the document is NEVER placed into the constitution.

Every governed project MUST maintain a single, project-local **operator-request-history document** – an append-only,

newest-first ledger of EVERY operator request/prompt the project has received – at a fixed, project-declared standard path (declared once per \$. . . , e.g. docs/requests/history.md , never moved without a \$. . . operator decision). Each entry MUST carry EXACTLY these operator-mandated fields: () **Request content** – the full/verbatim prompt text where archived, otherwise a faithful complete summary (each entry marks which – verbatim vs summary); () **Accepted (when)** – the exact date + time + TIMEZONE the request was accepted, with the timezone stated EXPLICITLY on every entry (the consuming project declares its default zone per \$. . . ; where only a date is durably recoverable the time is marked time UNKNOWN , never fabricated); () **Track** – the parallel-development track that processed it (the conductor/main track by default, per-domain tracks where applicable, \$. . . /\$. . . /\$. . .); () **Alias** – the live agent/worker alias that took the workable item at acceptance (\$. . . /\$. . .); () **Model + effort** – the model + effort setting used. (A) **NO INVENTION** (\$. . .). A field not durably recoverable is recorded LITERALLY as UNKNOWN (with a one-line reason where useful) – NEVER guessed; a request-history entry that fabricates a timestamp/alias/model it never captured is a \$. . . /\$. . . bluff at the requirements-record layer. (B) **HONEST RECONSTRUCTION BOUNDARY.** For pre-mandate sessions where per-prompt records were not archived, the document is best-effort reconstructed from the durable dated record (governance-mandate provenance dates, session ledgers, handoff docs, VCS history), with the unrecoverable per-prompt fields marked UNKNOWN and the reconstruction boundary stated explicitly in the document (\$ -style scope note) – never silently implied complete. (C) **ALWAYS IN SYNC + EXPORTED.** The document carries the \$. . . revision header, is kept in sync with live state

(\$. trigger set – every new operator prompt appends a row), and is exported to all mandated formats (\$. . /\$. . – .md + .html + .pdf + .docx where the four-format mandate applies, each non-.md sibling's mtime ≥ the .md). **(D) KEEP-APPLYING MECHANISM.** The project ships an append mechanism (a helper script and/or a UserPromptSubmit -class hook) that appends a new newest-first row for each NEW operator prompt at capture time with the exact timestamp + track + alias + model + effort (track/alias derived deterministically per \$. . , honest ? / UNKNOWN when undeterminable); an append helper that landed WITHOUT the automatic-capture hook is honestly a partial mechanism (the hook wiring is a tracked \$. . /\$. . follow-up, never claimed as automatic capture it does not perform – \$). **(E) COMPLEMENTS, NEVER REPLACES.** The request-history document complements – never replaces – any per-session no-loss requirements ledger (\$. . /\$), the workable-items DB SSoT (\$. . /\$), or the \$. . / \$. . session-resumption artefacts; loss of a requirement at its intake (an operator prompt not recorded) is FORBIDDEN (\$). **(F) RULE-UNIVERSAL / DOCUMENT-PROJECT-LOCAL (\$).** This universal RULE lives here in the constitution submodule; the request-history DOCUMENT, its declared path, and the project's default timezone are consumer-owned project DATA and live ONLY in the consumer layer – the document is NEVER placed into the constitution submodule (\$ decoupling).

Classification: universal (\$) – a platform-neutral requirements-record discipline reusable by ANY project; the consuming project supplies its concrete document path, default timezone, and append-helper wiring per \$ Composes \$ (no-guessing – UNKNOWN never a fabricated field) / \$ (rule-

universal / document-project-local split) / § . .
(revision header) / § . . / § . .
(four-format export) / § (enumerated-
coverage honesty for the backfill) / § . . /
§ . . (always-in-sync live-state artefact) /
§ . . / § . . / § . .
(track + alias derivation) / § . . (alias
tracking) / § . . (no requirement lost) /
§ . . (report→workable-item – the request-
history is the human-readable request ledger;
§ . . is its DB/tracker materialisation) /
§ . . / § . . (workable-items SSoT).
Propagation gate CM-COVENANT-114-208-PROPAGATION (literal
11.4.208 across the consumer fleet) + recommended gate
CM-REQUEST-HISTORY-DOC-PRESENT (the project-local request-
history document exists at its declared path, carries the
§ . . revision header + the five mandated per-
entry fields + an explicit-timezone stamp on every entry +
a reconstruction-boundary scope note; the append helper
exists + is parse-clean per § . . ; no fabricated
field – UNKNOWN used for the unrecoverable) + paired
§ . meta-test mutation (strip a mandated field from a
row / drop the explicit-timezone stamp / fabricate an
alias-or-model on an entry that never captured one /
delete the document → the gate FAILs; strip the literal →
the propagation gate FAILs; gate-code = separate work
item).

Canonical authority: constitution submodule Constitution.md
§ . . .

Non-compliance is a release blocker. No escape hatch – no
--request-history-optional , --skip-request-log , --no-timezone-
stamp , --invent-missing-field , --request-history-in-
constitution , --prompt-without-history-entry flag.

§ . . . — Code-review MUST run on the Opus model at xhigh effort as PRIMARY substrate, with a genuine-unavailability Sonnet fallback – no Fable, no escalation tier (User mandate, – – , supersedes – – and – – ; amended – – to add the Sonnet-fallback clause)

Verbatim operator mandate (– – , CRITICAL), AMENDING clause (A) below – adds a genuine-unavailability fallback, does not change the primary substrate: "If we do not have access to Opus model for some reason but only to Sonnet, we MUST USE what we have! Reviews MUST BE done. Ideally with Opus model, however Sonet MUST BE used if Opus for any reason cannot be used! Add this / extend into main constiotution Submodule and apply the rule asap so nothing is blocked with review which is not performed yet! We MUST BE maximally efficient! Work MUST BE performed with all means we have at particular moment!"

Verbatim operator mandate (– – , CRITICAL), superseding both mandates below: "Main model to use is Sonet, reviews and other heavy stuff have to be done by Opus instead of Fable!!! Fix this permanently on all levels!"

Prior mandate (– –), WITHDRAWN as current, kept for lineage: *"the mandatory independent code-review MUST by default be performed with the Opus model at xhigh effort; for a review the operator or the reviewer itself classifies as extremely big or complex, the Fable model at xhigh effort MUST be used instead."* Its Opus-xhigh default survives; its Fable escalation clause (former clause (A)) is REMOVED by the – – mandate – there is no review tier above Opus-xhigh.

Prior mandate (11 4 121 - 1024 09 10), WITHDRAWN as current, kept for lineage: *"the mandatory independent code-review MUST ALWAYS be performed with the Fable model at xhigh effort; if Fable is unavailable, use Opus at xhigh effort."* Inverted by the 11 4 142 - - mandate and fully superseded by the 11 4 142 - - mandate.

Every mandatory independent code-review the constitution requires – the § . . code-review-agent gate before pre-build + main build, the § . . universal every-change-review (no exception), and the § . . exhaustive all-scenario/all-angle review + verify-against-captured-runtime-evidence, iterated to a zero-finding GO per § 1024 10 00 – MUST be dispatched on the **Opus model at xhigh effort** as PRIMARY substrate, whatever the size or complexity of the reviewed change. (A) PRIMARY substrate Opus-xhigh + genuine-unavailability SONNET FALLBACK (amended - -) – no escalation tier, no Fable, no other fallback model. Opus at xhigh is the PRIMARY and preferred review substrate. Fable is NOT permitted in any role (default, escalation, or fallback), ever; no model other than Opus or the sanctioned Sonnet fallback below, and no lower tier for Opus itself, substitutes for it. When Opus at xhigh is genuinely unavailable (not configured / rate-limited / unreachable, established as a captured FACT per § . . – never a guessed unavailability), the review MUST be dispatched on **Sonnet** (at the highest effort the dispatch path can set) rather than blocked – the sole sanctioned fallback, per the verbatim - - operator mandate above: reviews MUST be done with whatever capable means are actually available, Opus ideally, Sonnet when Opus genuinely cannot be used. The Opus-unavailability FACT and the Sonnet-fallback substitution MUST both be recorded honestly in the review's captured evidence (§ . . /§ . .) – a Sonnet-fallback review is NEVER silently presented as if it ran on Opus. A

review is NEVER left undone / blocked / deferred solely because Opus is unavailable while Sonnet is reachable – maximal efficiency (§ 1.1.1) forbids leaving genuinely performable review work undone. ONLY when BOTH Opus and Sonnet are genuinely unavailable (each established as its own captured FACT per § 1.1.2, never guessed) does the review become BLOCKED and deferred under the original (A) terms: the unavailability fact is recorded, the owed review is tracked (§ 1.1.3; § 1.1.4 once the self-resolution audit is exhausted), and the reviewed change does NOT pass the review gate unreviewed. **(B) Independence preserved** (§ 1.1.5 / § 1.1.6). The reviewing model (Opus or the sanctioned Sonnet fallback) runs as the structurally-separated reviewer (never the author self-reviewing); § 1.1.7 sets the MODEL + EFFORT of that review, it does not weaken the independence, the iterate-until-GO loop (§ 1.1.8), or the rock-solid captured-evidence requirement (§ 1.1.9 / § 1.1.10 / § 1.1.11). **(C) Honest boundary** (§ 1.1.12). A review claimed "done" that actually ran on Fable, Haiku, or any model other than Opus or the sanctioned Sonnet fallback, or at an effort below `xhigh`, or on Sonnet with no recorded Opus-unavailability fact, is a § 1.1.13 bluff at the review-substrate layer – the review's captured evidence MUST record which model + effort actually performed it. Where the dispatch path cannot set or report the effort (the § 1.1.14 (F.) harness gap), the evidence records that gap honestly and the review is never claimed to have run at `xhigh` on an unverified basis. **(D) Values already burned in a public repository** (§ 1.1.15). Any plaintext credential value that has appeared in documentation, a spec, or commit history – including as a worked example – MUST be treated as compromised from the moment of commit; a document illustrating this anchor or any credential-adjacent

mandate. MUST NOT embed a real-looking username/password pair, even as an example, without the accompanying statement that the pair is compromised-by-disclosure and MUST NOT be relied on as a real secret. STRENGTHENS

\$. . . /\$. . . /\$. . . /
\$. . . – it does NOT replace any of them; it pins the review's execution substrate, now a PRIMARY-plus-sanctioned-fallback pin rather than a single-model pin.

Classification: universal (\$. . .). Composes

\$. . . / \$. . . / \$. . . /
\$. . . / \$. . . / \$. . . /
\$. . . / \$. . . / \$. . . /
\$. . . (maximal multi-agent utilization – the reason a genuine Opus gap defers to Sonnet rather than blocking, never reduces to idle waiting) / \$. . . (alias/model selection) / \$. . . /
\$. . . (nano-precision model-tier selection –
\$. . . remains the hard pin for review/merge-conflict scope that \$. . . (E) explicitly does not reach; the amended clause (A) Sonnet fallback draws on the SAME sonnet rung \$. . . (A) already names, never a new tier). Propagation gate

CM-COVENANT-114-209-PROPAGATION (literal 11.4.209 across the consumer fleet) + recommended gate CM-CODE-REVIEW-OPUS-XHIGH (renamed by the – – amendment; the retired predecessor token is cited in scripts/gates/gate_ledger_removals.tsv ; amended – –) (the mandatory code-review's captured evidence records model=Opus + effort=xhigh (or, on a dispatch path that cannot set effort, that \$. . . (F.) harness gap honestly recorded instead of a claimed xhigh), OR model=Sonnet carrying a captured Opus-unavailability FACT under the SAME effort-or-honest-gap rule; a review recorded on Fable/Haiku/any other model, recorded as xhigh when the dispatch path could not actually set it, or recorded on Sonnet with NO captured Opus-unavailability

fact, FAILs; a review recorded BLOCKED with a captured fact that BOTH Opus AND Sonnet were unavailable is honest and MUST NOT fire the gate – the § . . . () false-positive guard) + paired § . meta-test mutation (record a code-review run on a non-Opus/non-Sonnet model, or claiming xhigh when the dispatch path could not actually set it, or on Sonnet with no recorded Opus-unavailability fact → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
§ . . .

Non-compliance is a release blocker. No escape hatch beyond the sanctioned (A) Sonnet fallback (which itself requires a captured Opus-unavailability FACT, never a guess) – no --review-any-model , --review-lower-effort-OK , --skip-opus , --non-xhigh-review , --escalate-to-fable-or-haiku , --sonnet-fallback-without-recorded-fact , --leave-review-blocked-while-sonnet-reachable flag.

§ . . . — Zero-loss request/prompt intake: every operator request MUST be mechanically captured, tracked, and processed – never skipped, ignored, avoided, or lost (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –): "We MUST CHECK which requests / prompts we have given havent passed through except the accumulative prompts / requests history document !!! There may be more of it! This MUST NEVER happen! We MUST make sure that every request / prompt we issues is ALWAYS respected and taken into the account, executed and processed! This is MANDATORY RULE !!! There MUST NOT BE any avoiding, ignoring, skipping or any form of loss for

requests / prompts we make! Extend properly constitution Submodule and add mandatory linkings, hooks and extend with dcos chain any part of the System we MUST in order for this to be permanently resolved! Applying this starts immediately! Commit and push all changes to constitution Submodule to all upstreams!"

Forensic FACT (2026 07 15 - -): an entire session's worth of operator prompts (a release-drive "keep going", a **BACKGROUND** emulator directive, a live-testing model matrix, an alias switch, a multi-track directive) went UN-RECORDED in the \$. . request-history ledger because \$. . (D)'s "keep-applying mechanism" (the **UserPromptSubmit** auto-capture hook) was left as a "tracked follow-up / honestly partial" – so the ledger drifted the moment hand-editing stopped, exactly the \$. . failure class ("a rule documented as 11 4 208 mechanically enforced but implemented nowhere is worse than no rule").

Every operator request / prompt MUST be (a) **CAPTURED** – recorded in the \$. . request-history ledger MECHANICALLY at intake, never depending on manual vigilance; (b) **TRACKED** – mapped to a tracked action / workable item (\$. . report→item, \$. . completion) OR an explicit evidence-backed no-op-with-reason; and (c) **PROCESSED** – respected, acted upon, executed, never skipped / ignored / avoided / dropped. Loss of a request at ANY of these three stages is a \$. . PASS-bluff at the requirements-intake layer of the SAME severity class as \$. . (loss of requirements FORBIDDEN) – \$. . is the INTAKE-side MECHANICAL guarantee that \$. . 's no-loss holds from the very first moment a request arrives. \$. . STRENGTHENS \$. . (D): the auto-capture hook is PROMOTED from "tracked follow-up / honestly partial" to MANDATORY, wired, and verified.

Mandatory protections (ALL hold): () **MECHANICAL CAPTURE**, not manual vigilance (§ . .). The § . . ledger's keep-applying mechanism MUST be a WIRED, VERIFIED-RUNTIME-STATE auto-capture hook (a `UserPromptSubmit` -class hook that appends every operator prompt to the ledger at intake) – NOT a "tracked follow-up", NOT a manual conductor step. "Installed" = a verified runtime state observed in the live hook path (§ . . (. .)), never a doc claim; the hook MUST be non-fatal (never blocks / fails a prompt) and honest (`UNKNOWN` fields never invented, § . .). () **DOCS-CHAIN BOUND** (§). The ledger + its four-format exports (§ . . . / § . . .) are bound into the docs-chain so they stay in sync mechanically; a stale ledger export is a § . . sync violation. () **GATED** (§ . . , **enforced-not-advisory**). A pre-build gate FAILs if the auto-capture hook is not wired in the live hook path (verified-installed per § . . (. .), content-matched to its tracked source), OR the ledger is stale relative to the durable prompt record. Paired § . . mutation – un-wire the hook / orphan the ledger → gate FAILs. () **AUDIT + NO SILENT DROP**. Any operator prompt found un-captured (in `.remember/` logs, older CONTINUATION revisions, session transcripts) MUST be back-filled into the ledger, never silently dropped (§); every captured request resolves to a tracked action / item (§ . . .) OR an explicit evidence-backed no-op-with-reason – captured-then-ignored is forbidden. () **PROCESS EVERY CAPTURED REQUEST**. The autonomous loop (§ . . . / § 11 4 147 / § / § / § 11 4 6 . .) treats every open captured request as an actionable queue item; the loop's done-condition MUST NOT read satisfied while any captured request is un-processed (mirrors § . . . (e)). () **HONEST BOUNDARY** (§ . . .). Machine-derivable fields (timestamp /

track / alias) are captured verbatim; model / effort where undeterminable are UNKNOWN^{11 4 17}, never invented; a back-filled catch-up row honestly marks its reconstruction boundary (\$^{11 4 35} . . . (B)) – never a fabricated time.^{11 4 197}

Classification: universal (\$^{11 4 210} . . .) – the consuming project supplies its ledger path, hook path, and docs-chain context per \$^{11 4 186} Composes \$^{11 4 147} (\$^{11 4 205} . . . is its intake-side mechanical enforcement) / \$^{11 4 147} (report→tracked item) / \$^{11 4 147} (the ledger document – \$^{11 4 147} promotes its (D) hook from partial-follow-up to mandatory) / \$^{11 4 147} (enforced-not-advisory, hook-installed-is-verified-runtime-state) / \$^{11 4 147} (docs-chain) / \$^{11 4 147} / \$^{11 4 147} / \$^{11 4 147} / \$^{11 4 147} (the loop processes every captured request) / \$^{11 4 147} / \$^{11 4 147} / \$^{11 4 147} (e). Propagation gate CM-COVENANT-114-210-PROPAGATION (literal 11.4.210) + recommended gate CM-REQUEST-CAPTURE-HOOK-WIRED (the auto-capture hook is verified-installed in the live hook path AND content-matches its tracked source per \$^{11 4 210} ()); the ledger is fresh; docs-chain-bound) + paired \$^{11 4 210} meta-test mutation (un-wire the hook → recommended gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
\$^{11 4 210}

Non-compliance is a release blocker. No escape hatch – no --skip-request-capture, --manual-ledger-OK, --hook-optional, --drop-unprocessed-request, --ledger-may-be-stale, --request-may-be-skipped flag.

§ . . – Merge-conflict resolution during main→feature/product/flavor merges MUST run on the Opus model at xhigh effort as PRIMARY substrate, with a genuine-unavailability Sonnet fallback – no Fable (User mandate, – – ; amended – – ; amended again – – to add the Sonnet-fallback clause)

Verbatim operator mandate (– – , CRITICAL), AMENDING clause (A) below – adds a genuine-unavailability fallback, does not change the primary substrate: "If we do not have access to Opus model for some reason but only to Sonnet, we MUST USE what we have! Reviews MUST BE done. Ideally with Opus model, however Sonet MUST BE used if Opus for any reason cannot be used! Add this / extend into main constiotution Submodule and apply the rule asap so nothing is blocked with review which is not performed yet! We MUST BE maximally efficient! Work MUST BE performed with all means we have at particular moment!"

Verbatim operator mandate (– – , CRITICAL), superseding the model choice of the – – mandate below: "Main model to use is Sonet, reviews and other heavy stuff have to be done by Opus instead of Fable!!! Fix this permanently on all levels!"

Verbatim operator mandate (– –), WITHDRAWN as to its model choice, kept for lineage: "During the merge of main branch to all feature/product/flavor branches of tracks – , if conflicts happen they MUST BE resolved using Fable model with xhigh effort! If Fable is not available, then with Opus with xhigh! Add this MANDATORY RULE into constitution Submodule and its files CONSTITUTION.md, CLAUDE.md, AGENTS.md, QWEN.md, GEMINI.md, and other related relevant files!" Its SCOPE (every conflict-hitting main→feature/product/flavor merge)

and its `xhigh` effort requirement stand; its Fable default and its Fable→Opus fallback ordering are withdrawn.

Any merge-CONFLICT resolution performed while integrating the canonical trunk (`main` / `master`) into a feature / product / `flavor` branch – the § 4.1.1 regular main→feature merge cadence, the § 4.1.2 branch-taxonomy integration, the § 4.1.3 merge-first pipeline, the § 4.1.4 (D) trunk-into-stream sync AND the § 4.1.5 (I) approval-gated back-merge – MUST be carried out on the **Opus model at `xhigh` effort** as PRIMARY substrate. This is the merge-resolution analogue of § 4.1.6 (which pins the same model + effort for code-REVIEW): § 4.1.7 pins it for the DECISION-MAKING of conflict resolution (which side to keep, how to UNION both sides without loss, marker removal, semantic reconciliation of divergent edits). (A) **PRIMARY substrate Opus-`xhigh` + genuine-unavailability SONNET FALLBACK** (amended § 4.1.8 – § 4.1.9) – **no Fable:** Opus at `xhigh` is the PRIMARY and preferred conflict-resolution substrate; Fable is NOT permitted in any role, ever, and no lower effort for Opus itself is acceptable. When Opus at `xhigh` is genuinely unavailable (not configured / rate-limited / unreachable, established as FACT per § 4.1.10 – never a guessed unavailability), the conflict resolution MUST be carried out on **Sonnet** (at the highest effort the dispatch path can set) rather than blocked – the sole sanctioned fallback, per the verbatim § 4.1.11 operator mandate: resolve the conflict with whatever capable means are actually available, Opus ideally, Sonnet when Opus genuinely cannot be used; the Opus-unavailability FACT and the Sonnet-fallback substitution MUST both be recorded honestly in the merge's captured evidence (§ 4.1.12 / § 4.1.13) – a Sonnet-resolved merge is NEVER silently presented as if Opus resolved it. `xhigh` effort is REQUIRED on whichever substrate actually resolves the conflict WHERE THE

DISPATCH PATH CAN SET IT – a careless conflict resolution silently drops a commit or commits a marker (§ 11.4.4.1 / § 11.4.4.2 step 1), so the highest-effort scrutiny is mandatory regardless of model; where the dispatch path cannot set or report effort (the § 11.4.4.2 (F.) harness gap), clause (D)'s honesty rule governs: the gap is recorded, never silently claimed as xhigh unverified. ONLY when BOTH Opus and Sonnet are genuinely unavailable (each established as its own captured FACT per § 11.4.4.2.1, never guessed) is the conflict resolution BLOCKED and deferred: the unavailability fact is recorded, the merge is NOT completed, and the branch is left quiescent with no conflict marker committed (§ 11.4.4.2 / § 11.4.4.2 step 1). **(B) Scope.** The operator named tracks – (feature / product / flavor branches); the rule GENERALISES per § 11.4.4.2 to EVERY controlled main→(feature|product|flavor) merge on EVERY track that hits a conflict – the substrate requirement does not depend on which track. **(C) Safety preserved.** § 11.4.4.2 sets the MODEL + EFFORT of the conflict-resolution work; it does NOT weaken and is BOUNDED BY the surrounding merge-safety invariants – § 11.4.4.2 (NEVER force-push), § 11.4.4.2 / § 11.4.4.2 (no commit lost, hardlinked pre-op backup before any large/risky merge), § 11.4.4.2 step 1 (ZERO conflict markers committed, ZERO file silently dropped, union of both sides preserved), § 11.4.4.2 () anti-bluff (post-merge smoke GREEN + conflict-marker scan empty + no-lost-commit, captured evidence § 11.4.4.2 / § 11.4.4.2), § 11.4.4.2 (quiescent-only), § 11.4.4.2 / § 11.4.4.2 (fetch-first). **(D) Honest boundary** (§ 11.4.4.2). A merge whose conflicts were resolved on Fable or any model other than Opus or the sanctioned Sonnet fallback, or at an effort below xhigh, or on Sonnet with no recorded Opus-unavailability fact, is a § 11.4.4.2 bluff at the merge-resolution-substrate layer –

the merge's captured evidence MUST record which model + effort resolved the conflicts (and, when blocked, the FACT that both Opus and Sonnet were unavailable). Where the dispatch path cannot set or report the effort (the § 11.4.179 (F.) harness gap), the evidence records that gap honestly and the merge is never claimed to have run at xhigh on an unverified basis. STRICT EXTENSION of § 11.4.179 (which pins Opus-xhigh for code-REVIEW) applied to merge-CONFLICT resolution; STRENGTHENS § 11.4.179 / § 11.4.179 / § 11.4.179 by pinning the substrate of their conflict-resolution step – the two anchors are the review-side (§ 11.4.179) and merge-side (§ 11.4.211) halves of the same Opus-xhigh-primary-plus-Sonnet-fallback substrate discipline.

Classification: universal (§ 11.4.179). Composes § 11.4.179 (the review-side half of this same amended substrate discipline) / § 11.4.179 (maximal multi-agent utilization – the reason a genuine Opus gap defers to Sonnet rather than blocking, never reduces to idle waiting) / § 11.4.211 / § 11.4.179 / § 11.4.179 / § 11.4.179 / § 11.4.179 / § 11.4.179 / § 11.4.179 (D)(I) / § 11.4.179 / § 11.4.179 / § 11.4.179 (model/alias selection) / § 11.4.179 (A) (the sonnet rung this fallback draws on). Propagation gate CM-COVENANT-114-211-PROPAGATION (literal 11.4.211 across the consumer fleet) + recommended gate CM-MERGE-CONFLICT-OPUS-XHIGH (renamed by the – – amendment; the retired predecessor token is cited in scripts/gates/gate_ledger_removals.tsv ; amended – –) (a main→feature/product/flavor merge that resolved conflicts records model=Opus + effort=xhigh (or, on a dispatch path that cannot set effort, that § 11.4.179 (F.) harness gap honestly recorded instead of a claimed xhigh), OR model=Sonnet carrying a captured Opus-unavailability FACT under the SAME effort-or-honest-gap rule; a merge recorded

on Fable/any other model, recorded as xhigh when the dispatch path could not actually set it, or on Sonnet with no recorded Opus-unavailability fact, FAILs; a merge recorded BLOCKED with a captured fact that BOTH Opus AND Sonnet were unavailable is honest and MUST NOT fire the gate – the § 1.1.4.211 () false-positive guard) + paired § 1.1.4.211 meta-test mutation (record a conflict-resolving merge on a non-Opus/non-Sonnet model, or claiming xhigh when the dispatch path could not actually set it, or on Sonnet with no recorded Opus-unavailability fact → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md § 1.1.4.211.

Non-compliance is a release blocker. No escape hatch beyond the sanctioned (A) Sonnet fallback (which itself requires a captured Opus-unavailability FACT, never a guess) – no --merge-conflict-any-model , --skip-opus-merge-resolution , --non-xhigh-merge-conflict , --resolve-conflict-lower-effort-OK , --escalate-to-fable , --sonnet-fallback-without-recorded-fact , --leave-merge-blocked-while-sonnet-reachable flag.

§ 1.1.4.211 – Main README is the canonical starting-point / entry point for ALL project documentation – no doc may be an orphan unreachable from README (User mandate, – –)

Verbatim operator mandate (– –): "Manual tests catalog must have linking all through the main README as the starting point of this and any other Project documentation! This fact and mandatory rule MUST be added into the constitution."

The main README (README.md at the repository root) is the canonical STARTING POINT / entry point for ALL project documentation, for THIS project and EVERY project incorporating this constitution. Every project document in the \$. . . export scope – INCLUDING the manual-tests catalog (the \$. . . manual-QA testing catalog), the \$. . . per-feature Status / Status_Summary docs, the Issues / Fixed trackers + their summaries (\$. . . / \$. . . / \$. . .), and every guide / research note / changelog / script companion doc – MUST be reachable / linked (DIRECTLY or TRANSITIVELY) from the main README. No project document may be an ORPHAN – a doc that no link-path from README reaches is a \$. . . violation. This is the STRICT GENERALISATION of \$. . . :

\$. . . mandated a Tracked-Items + Status Documents doc-link section in the README covering the tracker / status doc set; \$. . . extends the same reachability guarantee to EVERY \$. . . -scope document, so the README is the single navigable root of the project's documentation tree. (A) **Reachability = a link path.** README → guide → sub-doc is fine (transitive); the requirement is that SOME path from README reaches every in-scope doc, not that every doc is linked directly at top level. A doc no path reaches is an orphan. (B) **Always-sync** (\$. . . / \$. . . docs-chain). The README doc-graph is kept current mechanically – a newly-added \$. . . -scope doc that no README-reachable path links is a \$. . . violation the moment it lands; the manual-tests catalog specifically MUST be linked from the README (the operator's load-bearing example). (C) **Honest boundary** (\$. . .). Reachability is a link-graph property (every doc is FINDABLE from README), NOT a claim about each doc's CORRECTNESS or freshness (those stay \$. . . / \$. . . / the doc's own gates). (D) **Composition**

with § . . . § 11.4.212 . 's Tracked-Items section remains the mandated README landing block for the tracker / status set; § . . . adds that EVERY OTHER in-scope doc is likewise reachable – § . . . is the specialisation, § . . . the generalisation.

10

(E) README INTRODUCTION + ILLUSTRATION (extension, - -). The main README MUST open with a SHORT introduction (what the project is, who it is for, its current state at a glance – the reader's -second answer) AND MUST ILLUSTRATE all work done: for every major capability the project ships, the README carries either a direct summary paragraph OR a link to a detailed document explaining WHAT was done AND WHY it was done (never orphaned links, never "see the changelog" as the only explanation). Every linked document MUST justify its own existence in one sentence at the link site – a link with no "why this matters" is a § . . . (E) violation of § . . . summary-clarity severity at the entry-point layer.

11 4 35

(F) LINKAGE-COMPLETENESS RATCHET (extension, - - -). The reachability graph from README (transitively per (A)) MUST cover the full § . . . -scope doc-set enumeration with a monotone-decreasing UNREACHED count (mirrors the § . . . ratchet pattern); a § . . . -scope doc that is not reachable from README at pre-build/release seams FAILs the extended CM-README-DOC-ENTRYPOINT-COMPLETE gate; brownfield adoption per operator § . . . (§ . . . (E) fence). The existing CM-COVENANT-114-212-PROPAGATION gate stays unchanged (literal 11.4.212 unchanged, propagation stays GREEN); CM-README-DOC-ENTRYPOINT-COMPLETE is EXTENDED with clauses (E)/(F); a new paired mutation strips the introduction, OR removes a "why this matters" annotation,

OR breaks reachability of an in-scope doc → the extended gate FAILs.

Classification: universal (\$) – the consuming project supplies its README path, its manual-tests-catalog path, and its \$ doc-set enumeration per \$. Composes \$ / \$ / \$ / \$ / \$ / \$. Propagation gate CM-COVENANT-114-212-PROPAGATION (literal 11.4.212 across the consumer fleet) + recommended gate CM-README-DOC-ENTRYPOINT-COMPLETE (every \$ -scope document is reachable – directly or transitively – from the main README; the manual-tests catalog is linked) + paired \$ meta-test mutation (orphan an in-scope doc so no README path reaches it → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md

Non-compliance is a release blocker. No escape hatch – no `--orphan-doc-OK`, `--readme-not-entrypoint`, `--skip-readme-link`, `--doc-unreachable-OK` flag.

```

$ . . . - FEATURE research-
scheduling directive: recognized via all
$ . . . forms, SCHEDULES (never
synchronously executes) a deep, enterprise-
grade research + implementation-planning
effort as a tracked workable item (operator-
directed mandate, - - )

```

Mandate (- -): extend the § . . reporting-directive family with a fourth registered action, **FEATURE**, that turns a plain-language feature description into a SCHEDULED, deep, enterprise-grade

research-and-planning effort – mirroring EXACTLY how `ISSUE` / `BUG` / `TASK` turn a report into a tracked workable item, but for RESEARCH SCHEDULING rather than defect/task/feature reporting: the directive itself creates the tracked item and enqueues it durably; the multi-day research/planning/documentation work it describes is EXECUTED LATER by the project's standing autonomous loop when it claims that item, never synchronously by the directive engine.

FEATURE is a research-SCHEDULING directive, structurally parallel to the `$ . . .` report-creation directives: recognized via ALL `$ . . .` forms – `FEATURE :: <desc> (bare :: ), DEFAULT::FEATURE :: <desc> (namespaced :: ), /FEATURE <desc> (bare slash), / DEFAULT::FEATURE <desc> (namespaced slash), FEATURE ---> <desc> (arrow)` – PLUS the `$ . . .` single-colon form `FEATURE: <desc>`, REGISTERED-ACTION-ONLY by construction exactly like `BUG: / TASK: / ISSUE: :` because the action token grammar is UPPERCASE-only (`[A-Z][A-Z0-9_]*`), ordinary lowercase prose such as "feature: add dark mode" never matches ANY of the six forms and is never expanded and never asked about (`$ . . .`) – only the literal uppercase `FEATURE` token triggers the directive. `\FEATURE` (leading backslash) escapes to a literal per the existing `$ . . .` escape mechanism, unchanged.

() Scheduling, not synchronous execution – the load-bearing distinction from every other `$ . . .` action. `BACKGROUND` / `CRITICAL` / `IMPORTANT` / `NOTE` tag HOW the remainder of the SAME turn is handled; `BUG` / `TASK` / `ISSUE` immediately create + fully sync a workable item describing a report and the reported matter is then handled inline. `FEATURE` differs: its OWN execution is fast and bounded – create the tracked item (Type=Task, Status=Queued, `$ . . .` stable id), append it to a durable, never-dropped feature-research queue, sync the

DB-derived documents, and push to configured trackers – exactly as § 11.4.1 does for TASK, and via the SAME engine, never a duplicate implementation. The deep multi-day research / systematic-debug / documentation / planning work DESCRIBED INSIDE the item is NOT performed by the directive at dispatch time; it is performed LATER by the project's standing autonomous loop (§ 11.4.2 endless-loop / § 11.4.3 zero-idle / § 11.4.4 maximum-idle-use / § 11.4.5 continuous-parallel-stream / § 11.4.6 default-autonomous-loop) the moment it claims the item, and is driven to a genuinely COMPLETED (implemented AND wired AND verified) or explicitly evidence-backed CLOSED terminal state under the § 11.4.7 research/kicked-off-work completion mandate – a FEATURE item is NEVER permitted to sit unwired in the backlog; "scheduled" is not "done" and is never reported as such (§ 11.4.8).

() **The scheduled research-and-planning mandate** – embedded verbatim into the created item's comprehensive description (§ 11.4.9 D / § 11.4.10) as its acceptance-defining WORK PROGRAM: (a) a deep web-research series – articles, guides, scientific papers, open-source projects/codebases, worked examples – on the best way to incorporate the described feature into the project (§ 11.4.11 deep-research-before-implementation / § 11.4.12 latest-source cross-reference / § 11.4.13 deep multi-angle research per change); (b) systematic-debug (§ 11.4.14) of ALL data obtained to enumerate EVERY weak spot / gap / danger-zone / inconsistency / imperfection the research surfaces, and design a risk-free, rock-solid solution for EACH one found – no gap left unaddressed; (c) the design MUST ALWAYS be enterprise-grade, bleeding-edge, and innovative – never less; (d) MANDATORY exhaustive documentation – technical documents, an implementation plan down to lines-of-code + micro-proof-of-concepts, diagrams / schemes /

graphs, illustrations, SQL definitions, templates, and every other relevant material, four-format-exported and revision-headed per § . . . / § . . . ; (e) investigate + plan REUSE of existing vasic-digital + HelixDevelopment submodules/components BEFORE proposing a rewrite, extending them freely where genuinely needed while keeping them fully decoupled / project-not-aware / reusable per § . . . / § . . . / § . . . ; (f) plan from the TESTING point of view from day one – every supported test type, the Challenges submodule, and full HelixQA bank coverage per § . . . / § . . . ; (g) divide the work into phases / tasks / subtasks, fine-grained and nano-detailed enough to drive integration / implementation / wiring / testing / scaling directly, no vague placeholders; (h) plan explicitly for enterprise scalability and maximal performance; (i) when the feature has a UI/UX surface, produce full wireframes / diagrams / design files (Figma / PSD / PDF, or whatever formats the project mandates) authored via OpenDesign per § . . . / § . . . ; (j) plan full CodeGraph integration with regular / real-time index synchronisation per § . . . / § . . . / § . . . ; (k) create fully-detailed workable items (every detail + reference + attachment) synced in REAL TIME to the SQLite single-source-of-truth (§ . . . 11 4 112 / § . . .) + every derived workable-items document + every related project doc/ component + every connected external work-tracking system (ClickUp, HelixTrack, JIRA, or any other configured tracker per § . . . 11 4 202 D) – honestly SKIPPING any absent tracker with a machine-readable reason (credentials_absent / tracker_client_absent , § . . . / § . . .) rather than EVER faking a push, exactly as § . . . mandates.

() **The item.** The created item carries Type=Task (\$. . . – a scheduled research/planning effort is itself an internal workstream, not a product defect nor yet a shipped end-user capability; the feature it researches may spawn its OWN Feature-typed implementation item(s) once planned) + Status=Queued (\$. . .) + a stable auto-incremented id (\$. . .) + the \$. . . / \$. . . comprehensive structured description embedding (i) the verbatim feature description, (ii) the full research/planning/testing/scalability/UI/CodeGraph/tracker mandate of clause () verbatim as the item's WORK PROGRAM, and (iii) acceptance criteria = the full research-doc tree is produced under the item's recorded destination path, every enumerated gap carries a designed solution, an implementation plan down to lines-of-code exists, the planned follow-on workable items are created, and the whole effort is validated per \$. . . (never left un-wired). An undetermined section is recorded as an explicit UNKNOWN: gap, never invented (\$. . .), exactly as \$. . . 's engine already does for BUG/TASK/ISSUE.

(. . .) **Decoupled engine, reused machinery, consumer-owned data.** The item-creation + full-sync half of the FEATURE mandate is NOT reimplemented: the engine (constitution/ scripts/feature/schedule_feature_research.sh) invokes the SAME \$. . . report_item.sh engine to create the item, sync the derived documents, and push to trackers – so the already-audited no-faked-push / decoupling / anti-bluff invariants of \$. . . are inherited rather than duplicated (a second, divergent implementation of tracker-push logic would itself be a bluff-surface regression). The FEATURE-specific additions – the research-doc-tree destination path and the durable docs/requests/ feature_queue.md queue entry (mirroring the \$. . . BACKGROUND action's docs/requests/background_queue.md durability pattern, so a scheduled request is NEVER

silently dropped or forgotten) – are the engine's OWN, narrow responsibility. Both engines are inherited BY REFERENCE, carry ZERO project literals, and fail closed with an actionable message when their consumer config is absent (`$ . . / $ . . / $ . .`); every project-specific value (DB path, id prefix, research-doc root, sync command, tracker bindings) is consumer-owned DATA in `feature.yaml`, never an edit to the engine.

() Anti-bluff (`$ . . / $ . . ( )`). Four-layer coverage mirroring `$ . .`: a pre-build gate (`CM-FEATURE-DIRECTIVE`) whose FUNCTIONAL invariant SOURCES and RUNS the `$ . .` parser (never a grep-only assertion, `$ . .`) to prove the `FEATURE` token genuinely expands via the single-colon, `::`, arrow, and slash forms while lowercase prose (`feature: ...`) does NOT expand and is NEVER asked about; a paired `$ .` mutation test proving every invariant (registry presence, engine parseability, decoupling, the durable-queue append, delegation-not-reimplementation of the tracker machinery, and the scheduling-not-synchronous framing) is genuinely load-bearing; and captured evidence of a real feature-research item landing in a temp SQLite DB, its derived documents regenerated FROM the DB, and its trackers SKIPping honestly.

Classification: universal (`$ . .`) – the consuming project supplies its DB path, id prefix, research-doc root, sync command, and tracker bindings per `$ . . . Composes $ . . (grammar) / $ . . (the sibling reporting-directive family + its reused engine) / $ . . (a started research effort never evaporates) / $ . . / $ . . / $ . . (deep research) / $ . . (systematic-debug of every gap) / $ . . / $ . . (exhaustive`

documentation) / § . . . / § . . . /
 § . . . (reuse-first, decoupled) /
 § . . . / § . . . (test-type + Challenges
 + HelixQA planning) / § . . . / § . . .
 (OpenDesign UI planning) / § . . . /
 § . . . / § . . . (CodeGraph integration) /
 § . . . / § . . . / § . . .
 (workable-items + tracker sync) / § . . . /
 § . . . / § . . . (id/status/type) /
 § . . . (credentials) / § . . . (no-
 guessing) / § . . . / § . . . /
 § . . . / § . . . / § . . . (the
 autonomous loop that actually performs the scheduled
 work). Propagation gate CM-COVENANT-114-213-PROPAGATION
 (literal 11.4.213 across the consumer fleet) + recommended
 gate CM-FEATURE-DIRECTIVE (registry declares FEATURE + its
 single-colon form is registered-only; the engine exists,
 is parse-clean, and is project-literal-free; the FEATURE
 directive really expands at runtime via all six forms
 while lowercase prose does NOT; the engine SCHEDULES via
 the durable feature-research queue AND delegates item-
 creation/sync/tracker-push to the § . . .
 report_item.sh engine rather than reimplementing it) +
 paired § . . . meta-test mutation (strip the FEATURE
 registry row, OR flip its single-colon form to non-
 registered, OR strip the parser branch, OR remove the
 durable-queue append, OR couple the engine to one project,
 OR make the engine reimplement tracker logic instead of
 delegating, OR strip the scheduling-not-synchronous
 framing from the expansion text → the gate FAILs; strip
 the literal → the propagation gate FAILs; gate-code =
 separate work item).

Canonical authority: constitution submodule Constitution.md

§

Non-compliance is a release blocker. No escape hatch – no
--feature-runs-synchronously , --skip-feature-queue , --feature-
without-research-mandate , --reimplement-tracker-push , --
hardcode-project-in-engine , --skip-item-creation flag.

\$. . – Recurrence-links-not-mints: a defect that returns MUST reopen its existing item, never enter as a new id (research-derived, – –)

Forensic anchor (genericised, – –). In a consuming project, operator-reported defects that were RECURRENCES of already-tracked, already-"fixed" defects re-entered the tracker as BRAND-NEW ids rather than reopening the originals (measured: at least five distinct recurrence chains, one three deep). A recurrence filed as a new id **never increments the original's reopen counter** – so the \$. . reopens_count , which is the exact signal \$. . (d) and \$. . use to rank the most-fragile work for the deepest scrutiny, **stays quiet precisely on the items that keep breaking**. The project's recorded reopen-per-fix rate was % (published elite < . %) and was simultaneously a KNOWN UNDERCOUNT for this reason: the metric that was supposed to find the fragile items was silenced by the intake path. The same defect inflates the queue and forks the work – in one measured window ~ tracker items mapped to ~ real root causes, so "fix all " would have produced duplicate and mutually-conflicting fixes on one cause.

The mandate. Before minting a NEW id for any incoming defect – from ANY intake path – the system MUST check whether an existing item already describes the SAME defect, and if so LINK the report to that item rather than mint a new id, and REOPEN that item **if and only if it is currently in a terminal state** (\$. . Fixed /

Implemented / Completed , or a § . . . closure OTHER than duplicate-of – see the chain rule).

(.) RESOLVE TO THE CANONICAL CHAIN HEAD FIRST – a

duplicate-of closure is a POINTER, not a destination.

§ . . . 's closure vocabulary includes duplicate-of , and clause () below MANUFACTURES exactly such items; the () key will match them, because a duplicate shares its original's subject BY CONSTRUCTION. Therefore, before applying link-and-iff-terminal-reopen, every match MUST be resolved THROUGH its duplicate-of links to the CANONICAL CHAIN HEAD (the item the chain ultimately points to), and the link + any reopen land on THAT head – never on an intermediate duplicate. Reopening a duplicate-of closure would resurrect a non-canonical copy alongside its still-live canonical item: two live items for one defect – the exact queue-inflation + forked-fixes disease this anchor's own forensic names – and it would credit the § . . . reopens signal to the WRONG item, defeating clause (), which IS the point, on precisely the chain-class items the anchor was minted for. **Multi-match resolution:** when the key matches several items (an original plus its duplicates, possibly plus an open item), resolve every match to its chain head, then act on the SINGLE surviving head; if that resolution yields more than one distinct head the verdict is UNDECIDED per clause () (ask, or mint-with-link – never guess which head is meant, § . . .); if the resolved head is OPEN, prefer it and LINK ONLY. A chain whose head cannot be resolved (a broken/circular duplicate-of link) is UNDECIDED + a tracked integrity defect (§ . . . orphan-ref class), never a silent pick.

Open-original case (the ordinary two-reports-of-one-live-bug): when the resolved chain head is still OPEN – any non-terminal § . . . status, i.e. Queued / In progress / Ready for testing / In testing / already

Reopened / Operator-blocked (\$. . 's th value, non-terminal) – the correct action is LINK ONLY – "reopen an open item" is undefined in the \$. . status machine, and \$. . 's Reopened status + its reason vocabulary presuppose a prior closure to demote FROM (\$. .); forcing a reopen there would corrupt both the status and the counter. The reopen-and-increment path exists to correct the specific lie "this was closed as fixed and it is not" – an already-open item is not telling that lie.

\$. . precedence (stated explicitly – never left for a consumer to reconcile). \$. . mandates that every report lands a tracked item, "automatically, every time", and forbids --skip-item-creation ; \$. . forbids --mint-every-report . These are NOT in conflict and MUST NOT be read as such: \$. . 's requirement is that **no report is ever LOST** – it is satisfied IN FULL by a SAME-DEFECT report landing as a link + (terminal-only) reopen on the existing item, because the report is tracked, the item is live, and the \$. . counter now tells the truth. \$. . forbids the report evaporating into prose; it does NOT require a fresh id per report. \$. . refines WHICH item a report lands on; \$. . owns THAT it lands. Where the check verdict is DISTINCT or UNDECIDED, \$. . 's mint-a-new-id path runs unchanged. () **EVERY INTAKE PATH IS BOUND, not merely the reporting channel.** Defects enter via the \$. . reporting directives, via AI-detected test/gate failures, via regression-guard verdicts, via cycle re-discovery, via captured-evidence retrospect, and via operator manual testing – \$. . 's own reason vocabulary (test-failed , cycle-re-discovered , manual-testing-detected , end-user-report) enumerates them. A dedup check wired ONLY into the reporting channel under-covers exactly the machine-driven paths that recur most.

() **THE KEY IS STRUCTURED, NEVER A BARE SUBSTRING** (§ 4.1.1). Match on ticket reference where present, else on a normalised (subject, scope) – explicitly NOT a bare subject substring, which is the false-merge generator. () **A WRONG MERGE SILENTLY DELETES A REAL DEFECT** – so **distinct-but-similar** is a **FIRST-CLASS VERDICT, not a failure**. The check has three outcomes: SAME-DEFECT (link + reopen), DISTINCT (mint a new id – correct and expected), and UNDECIDED (§ 4.1.2 : do NOT guess – ask per § 4.1.3 / § 4.1.4 , or mint the new id WITH a recorded candidate-duplicate link for review; when autonomous per § 4.1.5 , prefer minting-with-link over merging, because a spurious extra id is cheap and recoverable while a wrongly-merged defect is silently LOST – the asymmetry decides the default). () **THE DEDUP ORACLE IS SELF-VALIDATED** (§ 4.1.6 () / § 4.1.7). It ships a golden-good (a true recurrence – MUST link+reopen), a golden-bad (a missed recurrence – MUST be caught), and a **negative-control**: two genuinely DISTINCT items that merely share a subject, which MUST NOT be merged. An oracle without the negative control is a false-merge engine and is itself a § 4.1.8 bluff. () **ID STABILITY IS PRESERVED** (§ 4.1.9). Linking a recurrence to its original NEVER renumbers, reuses, or retires an id: the original is REOPENED and its counter increments; an already-minted duplicate is LINKED to the original (a recorded-duplicate-of relation per § 4.1.10) – never silently deleted, never renumbered. () **THE RESTORED SIGNAL IS THE POINT**. Reopening the original is what makes § 4.1.11 's counter true, which is what makes § 4.1.12 (d)/§ 4.1.13 rank the genuinely-fragile items first. A reopen counter that undercounts is not a cosmetic defect – it is a § 4.1.14 / § 4.1.15 bluff at the prioritisation layer: it reports the fragile set as stable, and the deepest scrutiny is then aimed at the wrong items.

the CORRECT action there per the open-original case), nor on a clause-() UNDECIDED candidate-duplicate link (which is by design un-reopened, and is the mandated autonomous default), nor demand a reopen on an intermediate duplicate-of closure (clause () FORBIDS that reopen – a gate demanding it would enforce the very defect the anchor prevents) – a gate firing on any of these would be exactly the § . . () false-positive refusal that halts correct work and teaches operators to bypass it, and it would condemn behaviour THIS anchor mandates. Its golden-FALSE fixture set MUST therefore include an open-head link, a candidate-duplicate link, AND a link whose immediate target is a duplicate-of closure with the reopen correctly landing on the chain head – the gate MUST NOT fire on any of the three (§ . . () / § . . ()); its golden-TRUE set MUST include a terminal-head link with NO reopen event (the gate MUST fire) (file a known recurrence and let it mint a fresh id without a link → the gate FAILs; strip the negative control so two distinct same-subject items merge → the oracle self-check FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
§ . . .

Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-dedup-at-intake , --mint-every-report , --bare-substring-dedup , --merge-without-negative-control , --recurrence-as-new-id-OK , --reopen-counter-may-undercount flag exists.

§ . . . – A doc that BINDS work MUST live, tracked, in the repository where that work happens (research-derived, – –)

Forensic anchor (genericised, – –). In a consuming project, the document that DEFINED a binding acceptance gate – the gate a body of work was ordered to satisfy before proceeding – was **untracked** and lived in an out-of-tree scratch directory belonging to a different working checkout, i.e. it existed in NO repository and in no other agent's or operator's checkout. The gate could therefore not be honoured (agents could not read it), not be reviewed (no diff, no history, no authorship), not be version-controlled (no revision, no rollback, no § . . . header), and not be reconciled when it contradicted the source of truth. A binding artifact that lives nowhere is not governance – it is a rumour with an imperative mood, and work ordered against it cannot be verified against it by anyone but the one process that happened to hold the file.

The mandate. Any document that BINDS work – one that defines a gate or acceptance criterion, states a contract the work is judged against, or carries a plan/order the work is required to satisfy – MUST be tracked, in the repository where that work happens, at a declared path, BEFORE it may be cited as binding. () **BINDING ⇒ TRACKED, and un-tracked ⇒ NOT BINDING.** A binding claim sourced from an untracked or out-of-tree document MUST NOT be enforced against any work; the correct response to "the gate is defined in a file nobody can read" is to land the document (or decline the gate), never to obey a citation of it. () **IN THE REPOSITORY THE WORK HAPPENS IN.** Tracked in a sibling checkout, a scratch directory, a chat message, an issue comment, or a personal note is NOT tracked for this purpose – the test is whether an agent or operator who

clones ONLY the repo the work happens in can read the binding text at its declared path. Where a binding doc must govern several repositories, it is inherited BY REFERENCE from a shared, tracked, version-controlled dependency (§ . . . / § . . .), never copied and never left out-of-tree. () **THIS IS THE PRIOR QUESTION** § . . . **CANNOT ASK.** § . . . mandates that every doc in the export scope be README-reachable – but its check ENUMERATES the in-repo scope and asserts each member is reachable, so a document that exists in no checkout is never enumerated and the gate cannot FAIL on it. § . . . answers "is it findable?"; § . . . answers "is it here at all?". Once a binding doc is in the repo, § . . . (reachable from the README), § . . . (revision header), § . . . / § . . . (exports), and § . . . (kept in sync) all attach to it as normal – this anchor is their precondition, not their competitor. () **THE GENERALISATION OF** § . . . § . . . is this exact rule scoped to ONE named artifact: an authoritative tracker database is TRACKED, NEVER gitignored, "NOT a build artefact, it IS authoritative source data", as an explicit named carve-out from the gitignore rule. § . . . widens the artifact class from that one file to **any doc that binds work**, exactly as § . . . was minted as the strict generalisation of § . . . § . . . is the counterweight and stays intact: logs, forensic captures, and scratch artifacts remain out-of-tree and untracked by default – § . . . does NOT track them, and a doc's being untracked is only a defect once it is CITED as binding. That is the discriminator: **binding is what pulls a document into the repository**, and a document nobody cites as binding is governed by § . . . , not by this anchor.

gate-definition , --cite-doc-nobody-can-read , --scratch-dir-is-good-enough , --binding-by-rumour flag exists.

§ 6 0 7 2 5 . . – Canonical machine-readable design-token source: one CSS custom-property file (:root + [data-theme="dark"]), generated bindings everywhere else (User mandate, – –)

Forensic anchor – proving project (Helix Thready MVP design package, – –). The mandate's shape was proven end-to-end in the Helix Thready design package (docs/public/research/mvp/design/): ONE machine-readable token file (opendesign/tokens.css) carries every color / type / spacing / radius / elevation / motion value as CSS custom properties – light values in a :root {} block, dark values in a [data-theme="dark"] ..override – and every value is evidence-traceable, never invented: the brand chartreuse #B6E376 is the recorded eyedrop MEAN of , , green-dominant logo pixels, the mint secondary #ABDDC9 the mean of , mint-region pixels, the light accent MEASURED at . : on white and the dark accent at . : on the night surface. Every non-CSS platform in the fan-out consumes a GENERATED binding (token → Compose Color / Dp , Lipgloss lipgloss.Color , Flutter/Qt theme, Figma Variables payload), never a hand-kept copy; the not-yet-built codegen bridges are honestly [OPEN] -tracked, never claimed shipped.

The mandate. Every project governed by this Constitution that ships user-facing interfaces MUST maintain its design tokens as ONE canonical, machine-readable source, and everything else MUST derive from it: () **Canonical file shape** – CSS custom properties, light values in :root {} , dark values via the [data-theme="dark"] override pattern, composing the project's sanctioned dark-mode mechanisms

(@media (prefers-color-scheme: dark) system default, explicit data-theme attribute, framework-compatible class) – never duplicate hand-kept light/dark palettes. () **Generated bindings only** – every consuming surface (web CSS, native mobile, desktop, TUI, design tools, rendered doc artifacts) binds to the canonical source either directly via `var(--...)` or through a GENERATED binding produced by codegen FROM the canonical file; a hand-copied or hand-restated token value anywhere (an inline hex that restates a token, a hand-kept per-platform palette) is a § . . violation – hand-copies drift silently, and the drift is invisible until a surface renders off-brand. () **No invented values** – § . . applied to **design** – every color / size / duration in the token source is traceable to captured evidence: brand colors to a recorded eyedrop / measurement of the canonical brand asset, accent pairs to a MEASURED WCAG contrast ratio against their PAIRED background, structural values to the documented shared-system scale; a hex that appears in any design artifact and resolves to no token is a violation – flag it and use the nearest existing token (per § . . 's no-invented-values discipline), never mint an ad-hoc value. () **The token source is the white-label seam** – per-brand / per-account re-theming swaps token VALUES within the § . . closed override set, never component structure.

Honest boundary (§ . .). One canonical source proves consistency of DEFINITION, not correctness of USE – a component can still misuse a token (brand-as-body-text); role discipline is § . . 's and rendered-pixel proof is § . . 's. And "generated binding" follows § . . wiring honesty: a codegen bridge that does not exist yet is an [OPEN]-registered gap (§ . .), never claimed working.

Classification: universal (§ 11.4.216) – references no project-specific brand or platform; the consuming project supplies its token-file path, brand-asset evidence, platform fan-out list, and codegen bridges per § 11.4.217. Composes § 11.4.218 / § 11.4.219. (OpenDesign is the mandated token SYSTEM – this anchor pins its canonical FILE shape and consumption discipline) / § 11.4.220 / § 11.4.221 / § 11.4.222 / § 11.4.223 / § 11.4.224. Propagation gate CM-COVENANT-114-216-PROPAGATION (literal 11.4.216) + recommended gate CM-DESIGN-TOKEN-SINGLE-SOURCE (the canonical token file exists with :root {} + [data-theme="dark"] blocks; tracked design artifacts bind via var() rather than restating token hex inline; every declared platform names its generated-binding artifact or its [OPEN] id) + paired § 11.4.225 mutation (hand-restate a token hex in a tracked artifact → gate FAILs; strip the dark override block → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md § 11.4.226.

Non-compliance is a release blocker regardless of context. No escape hatch – no --hand-copied-palette-OK, --inline-hex-OK, --invent-token-value, --light-only-tokens, --platform-keeps-own-palette flag exists.

§ 11.4.227 – OpenDesign brand contract: every UI-shipping product carries a – section DESIGN.md + tokens.css twin with disjoint color roles, AA-pinned accents, and a locked attribution footer (User mandate, – –)

Forensic anchor – proving project (Helix Thready MVP design package, – –). Helix Thready ships design/opendesign/DESIGN.md in the OpenDesign –section

brand-contract schema – VERIFIED against the tool's own source per § . . . (spec, design-systems/README.md , the contributor authoring guide, and a measured – file corpus: two coexisting heading families, both parsed by number prefix only) – plus its machine-readable twin tokens.css per the OpenDesign token schema, and an honest TOOLING.md capability check. The contract encodes three DISJOINT color roles (brand decorative, accent interactive, semantic state), the closed white-label override set the three shipped sibling themes prove works, and the locked attribution footer.

11 4 35

The mandate. Every product governed by this Constitution that ships user-facing interfaces MUST carry an OpenDesign brand contract under its design area (design/opendesign/ or the project's declared equivalent per § . . .), consisting of: () **DESIGN.md in the verified –section schema** – H product title with the > Category: blockquote IMMEDIATELY after the H (the picker parses the line immediately after the H), nine numbered sections (the parser matches the number prefix, not the heading text), a "Font labels for catalog extraction" block (Display: / Body: / Mono: lines) in the typography section, real hex codes, CSS variables wrapped in :root {} (never bare) with dark tokens via the [data-theme="dark"] override pattern, real component CSS with prefers-reduced-motion targeting specific elements (never a global *), and :focus-visible on every interactive component – the tool's own blocking review lens, satisfied in-file. Format claims are verified against the tool's CURRENT source per § . . . , never assumed from memory. () **The tokens.css machine-readable twin** – the § . . . canonical token file (or a compiled view of it), following the tool's token schema; prose and tokens never diverge. () **Disjoint color roles by contract** – brand is DECORATIVE only (never body text; the Thready brand base measures . : on white – decorative by physics,

not by preference), accent is INTERACTIVE and AA-PINNED (each accent/on-accent pair carries a measured $\geq$. : ratio against its paired background, both modes), semantic (success / warn / danger) is STATE and is NEVER overridden by a brand – a brand color can never mask an error signal. () **White-label closed override set, server-validated** – per-account branding may swap decorative brand tokens, the accent pair, logo, and slogan; accent overrides are SERVER-VALIDATED to WCAG AA $\geq$. : or REJECTED (the rejection carries the measured ratio + an AA-passing suggestion – instructive, never bare); structural tokens, neutral surfaces, and semantic tokens are NOT account-overridable (advanced neutral overrides gate behind a full AA re-validation). () **Locked attribution footer** – the organization's attribution footer ships on every surface and is NEVER removed, even fully white-labeled.

Honest boundary (§ . .). The contract governs DEFINITION and roles; it does not prove rendered correctness (§ . . / § . .) nor replace § . . 's token-system, test-coverage, and extend-don't-reimplement obligations – this anchor pins the contract ARTIFACT § . . consumes. Schema-version drift between the local tool checkout and the doc is recorded as fact and re-verified at integration, never papered over.

Classification: universal (§ . .) – the consuming project supplies its brand values, category line, white-label storage/API bindings, and attribution string per § . . . Composes § . . / § . . / § . . / § . . / § . . / § . . . Propagation gate CM-COVENANT-114-217-PROPAGATION (literal 11.4.217) + recommended gate CM-OPENDESIGN-BRAND-CONTRACT (DESIGN.md present with the numbered sections + category blockquote + font-labels block + :root {} + [data-

theme="dark"] + targeted reduced-motion; tokens.css twin present and consistent; accent pairs carry measured ratios; the attribution footer token/lockup is present and not account-overridable) + paired \$. mutation (drop the font-labels block, or let a brand token override --danger, or mark the footer overridable → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
\$. . .

Non-compliance is a release blocker regardless of context. No escape hatch — no --skip-brand-contract, --prose-only-design-md, --brand-as-text-OK, --unvalidated-accent-override, --remove-attribution-footer, --semantic-tokens-brandable flag exists.

\$. . — **Living design-library catalogue: every component × state × theme rendered self-contained + a cross-platform reusability matrix with per-cell provenance (User mandate, — —)**

Forensic anchor — proving project (Helix Thready MVP design package, — —). Helix Thready's design/library/ renders the ENTIRE component set as a living catalogue: components.html — groups, ~ components, ~ component-state cells, each auditable in BOTH themes via a real toggle (~ rendered states), zero external requests, interactive behaviors REAL not mocked (theme cycling stamps data-theme and persists; column sorting drives aria-sort; a real showModal() dialog; an auto-dismissing toast that pauses on hover) — plus components-sheet.svg (an importable one-sheet) and platform-map.md, a —components × — platforms reusability matrix where EVERY cell carries a

VERIFIED / PROPOSED / ASSUMED provenance marker re-checked at source.

The mandate. Every project governed by this Constitution that ships user-facing interfaces **MUST** maintain a LIVING design-library catalogue: () **Complete component × state × theme rendering** – every component of the project's design system, in every declared state (default / hover / focus-visible / disabled / loading / error / empty / skeleton as applicable), auditable in BOTH light and dark from one artifact; a component or state that exists in the spec but not in the catalogue is a coverage gap, tracked per § . . . () **Self-contained artifacts** – the catalogue renders with ZERO external requests: no CDNs, no external fonts (self-hosted or system-stack fallback stated honestly), vendored libraries inlined; a catalogue that needs the network is not auditable offline and violates the project's CSP/offline posture. () **Real behaviors, token-bound** – interactive exhibits are real (real dialog, real sort, real toggle), every color/size binds to the § . . . token source, and all motion honors prefers-reduced-motion per § . . . () **Cross-platform reusability matrix with per-cell provenance** – a components × platforms matrix where EVERY cell is marked from the closed provenance set: VERIFIED (reproduced from source and re-checked at source, with the check date), PROPOSED (a [DEFAULT – adjustable] design pending review), or ASSUMED (inference, labeled as such); an unmarked cell reads as ASSUMED. The matrix is read BEFORE implementing any component on any platform; scaffolds are named scaffolds – a scaffold cell claiming VERIFIED is a § . . . PASS-bluff at the design-planning layer.

Honest boundary (§ . . .). The catalogue proves the DESIGN renders and behaves as specified in a browser-class runtime; it does not prove any platform IMPLEMENTATION

(that is the matrix's per-cell provenance + each platform's own § . . . rendered-pixel proof), and PROPOSED components are named proposals pending the upstream-vs-project split decision – never silently promoted to shipped.

11 4 8 11 4 162

Classification: universal (§ . . .) – the consuming project supplies its component set, state vocabulary, platform list, and matrix location per §

Composes § . . . / § . . . /
§ . . . / § . . . / § . . . /
§ Propagation gate CM-COVENANT-114-218-

PROPAGATION (literal 11.4.218) + recommended gate CM-DESIGN-LIBRARY-LIVING-CATALOGUE (the catalogue artifact exists, is self-contained – zero external URL references –, covers every spec'd component × state in both themes or tracks the gap, and the reusability matrix carries a closed-set provenance marker in every cell) + paired § . . . mutation (introduce a CDN reference → gate FAILs; blank a matrix cell's marker or claim VERIFIED on a scaffold row → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
§

Non-compliance is a release blocker regardless of context. No escape hatch – no --cdn-in-catalogue , --single-theme-catalogue , --mocked-behaviors-OK , --unmarked-matrix-cell , --scaffold-claims-verified flag exists.

2026 07 22

§ . . . – Per-app screen catalogues: full-IA coverage per surface, per-platform chrome variants, honesty markers for unbuilt behavior (User mandate, – –)

Forensic anchor – proving project (Helix Thready MVP design package, – –). Helix Thready's

design/screens/ carries a screen catalogue PER SURFACE – web screens + an interactive prototype shell walking the four key journeys, device-framed × mobile artifacts with an Android/iOS chrome TOGGLE (+ HarmonyOS/Aurora port notes) covering EVERY node of the wireframes' information architecture, a desktop shell with per-OS title bars / menus / tray / file-drop mocks, ten -column TUI mockups covering the full state machine, and a -page marketing set with a claims-traceability register. Undefined product behavior is rendered HONESTLY: placeholder billing amounts labelled as placeholders, an unspecified forgot-password flow linked with an inline note instead of a fake flow, an undecided channel source rendered as a DISABLED tile – each carrying its [OPEN: id] .

The mandate. Every application/surface a governed project ships MUST have a screen catalogue: () **Full-IA coverage** – the catalogue covers EVERY node of the surface's information architecture (every screen of the wireframe/IA contract); a screen present in the IA but absent from the catalogue is a coverage gap tracked per § . . – never silently skipped. () **Per-platform chrome variants** – where one design serves several host platforms, the catalogue renders the platform chrome variants (mobile OS chrome toggle, per-OS window chrome, terminal-width constraints) rather than pretending one chrome fits all; ports not yet designed are tracked, not faked. () **Self-contained, token-bound, both themes** – each screen is a self-contained artifact binding the §. . . tokens, rendered light + dark, populated with REALISTIC product data (never lorem-ipsum for load-bearing content), carrying its per-screen state contract (loading / empty / error / populated states declared). () **Honesty markers for unbuilt behavior** – where the backing backend, contract, or product decision does not exist, the screen MUST say so: a [GAP: id] / [OPEN: id] marker (registered

per § 11.4.19), a labelled placeholder, an inline note, or a disabled control – NEVER a designed-as-if-real fake flow; designing a flow for an undecided behavior and presenting it as settled is a § 11.4.19 PASS-bluff at the product-design layer.

Honest boundary (§ 11.4.19). A screen catalogue is the DESIGN contract, not the implementation: it proves every IA node has a designed, state-complete, theme-complete rendering – implementation correctness stays with

§ 11.4.19 host-rendered pixel proof,
 § 11.4.19 web-quality proof, and the
 § 11.4.19 manual-QA final gate.

Classification: universal (§ 11.4.19) – the consuming project supplies its surface list, IA contract, chrome-variant matrix, and realistic-data corpus per

§ 11.4.19 . Composes § 11.4.19 / § 11.4.19 /
 § 11.4.19 / § 11.4.19 / § 11.4.19 /
 § 11.4.19 / § 11.4.19 / § 11.4.19 .

Propagation gate CM-COVENANT-114-219-PROPAGATION (literal 11.4.219) + recommended gate CM-SCREEN-CATALOGUE-FULL-IA (every IA node maps to a catalogue screen or a registered [OPEN] id; every screen renders both themes; every screen whose backing contract is absent carries its honesty marker) + paired § 11.4.219 mutation (remove a screen without registering the gap → gate FAILs; strip an honesty marker from an unbacked flow → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
 § 11.4.219 .

Non-compliance is a release blocker regardless of context. No escape hatch – no --partial-ia-coverage , --one-chrome-fits-all , --fake-flow-OK , --lorem-ipsuim-content , --single-theme-screens flag exists.

§ . . – Open-first design tooling: self-hosted open platforms are primary; proprietary tools are import/export targets only; design sources live in-repo in open formats (User mandate, – –)

Forensic anchor – proving project (Helix Thready MVP design package, – –). Helix Thready's tooling capability check (opendesign/TOOLING.md) and export capstone prove the posture: every design source lives IN-REPO in open formats (SVG brand masters, tokens.css , JSON manifests + genuine .od-figma.json IR captures, hand-authored Lottie JSONs, self-contained HTML screens); the OpenDesign daemon/CLI status was RUNTIME-PROBED (live health endpoint, CLI round-trip, headless-vs-operator matrix), never assumed; and Figma is an IMPORT-ONLY bridge – the import plugin pins networkAccess: ["none"] , no cloud token exists, and the docs state verbatim that any "pushed to Figma cloud" claim from that environment would be a fabrication. The actual proprietary-tool file is [OPEN] -tracked, not claimed.

The mandate. Every governed project's design workflow MUST be OPEN-FIRST: () **Self-hosted open platform primary** – the collaborative design platform is an open-source, self-hostable tool (PenPot as the standing default), run SELF-HOSTED; together with the § . . OpenDesign engine it forms the primary design toolchain. () **Consumption ONLY via the PenPot submodule** – ALL PenPot mechanisms and utilities (the deployment recipe, RPC client utilities, access-token minting, import tooling, browser-hydration automation, MCP wiring) live in the dedicated reusable PenPot submodule (working name vasic-digital/PenPot ; the final repository name was UNCONFIRMED at this anchor's writing – the generic referent "the PenPot submodule" is normative until the name lands), inherited BY REFERENCE per § . . ; the PenPot

submodule itself consumes the Containers submodule (\$. . , rootless per \$. .) for orchestration and DECLARES that dependency in its \$. . helix-deps.yaml manifest. Projects MUST NOT hand-roll PenPot deployment or client code outside the submodule – a missing capability is extended UPSTREAM in the submodule per \$. . , never worked around locally. The submodule carries MECHANISMS/UTILITIES ONLY, never project design content – per-project design sources stay in the project repo per clause (). () **Proprietary tools are optional bridges, never the source of truth** – proprietary design tools (Figma or equivalent) MAY be used as import/export/review targets via documented bridge kits (variables payloads, IR captures, import guides), but NO design value may exist ONLY inside a proprietary tool or its cloud: the moment a decision lands, its canonical form lands in-repo. () **In-repo open formats** – design sources are versioned in the repository in open, diffable formats: SVG (vector masters), CSS (tokens per \$. .), JSON (manifests, IR captures, variables payloads), Lottie JSON (motion per \$. .), self-contained HTML (screens/library per \$. . /\$. .); binary proprietary formats are derived artifacts at most, never sources – the \$. . discipline applied to design: a design that binds work MUST live tracked in the repo. () **Runtime-probed capability claims** – every tooling claim ("the daemon runs", "the CLI exports", "import works") is established by a runtime probe recorded as evidence (\$. . real-condition discipline; \$. . latest-source verification for format contracts), with an explicit headless-vs-operator-interactive matrix and an HONEST-LIMITS section; a cloud push claimed without a real credentialed run is a fabrication of \$. . PASS-bluff severity. () **PenPot automation operational rules (research-verified)** – (a) the HEADLESS write path is the token-authenticated **RPC API**:

access tokens (minted under the backend's `enable-access-tokens` flag, handled per § 4.1.1) are the MANDATED auth mechanism for all automation – interactive OAuth relays / session-cookie scraping are FORBIDDEN as automation auth; (b) API-created PenPot files REQUIRE a **browser hydration pass** before handoff – thumbnails and text layout are materialized by the browser runtime, not by the write API, so a handoff of an unhydrated API-created file claimed render-ready is a § 4.1.1 PASS-bluff at the design-handoff layer; the hydration pass is automated via the PenPot submodule's browser-hydration utility and its execution is recorded as evidence (§ 4.1.1).

Honest boundary (§ 1.1.1). Open-first governs WHERE truth lives and WHICH tool is primary; it does not forbid operators using proprietary tools for review/hand-off, and it does not claim bridges exist before they are built – missing bridges (PenPot/Lottie export, proprietary-file creation) are [OPEN]-registered work per § 1.1.1 / § 1.1.2; and until the PenPot submodule's final repository name is confirmed, citations use the generic referent honestly rather than asserting an unconfirmed name as fact.

```

Classification: universal ($ . . . ) – the consuming
project supplies its platform choice within the open-first
constraint, its PenPot-submodule pointer + configuration
(instance URL, token env-var NAMES per $ . . . ),
bridge kits, and probe scripts per $ . . . Composes
$ . . . / $ . . . / $ . . . /
$ . . . / $ . . . / $ . . . /
$ . . . / $ . . . / $ . . . /
$ . . . / $ . . . / $ . . . /
$ . . . / $ . . . . Propagation gate CM-
COVENANT-114-220-PROPAGATION (literal 11.4.220) +
recommended gate CM-OPEN-DESIGN-TOOLING-PRIMARY (the self-

```

hosted open platform is declared and consumed via the PenPot submodule – no hand-rolled PenPot deployment/client code in the project tree; the submodule's `helix-deps.yaml` declares the Containers dependency; automation authenticates via access tokens, never an OAuth relay; every API-created file handed off carries a recorded hydration-pass run; every design source referenced by the design area resolves to an in-repo open-format file; every tooling capability claim cites a recorded probe; no design artifact exists only as a proprietary-cloud reference) + paired \$. mutation (hand-roll a PenPot RPC client in a project tree, or declare a proprietary tool as the source of truth, or hand off an API-created file with no hydration-pass record → gate FAILs; claim a capability with no recorded probe → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md`
\$. . .

Non-compliance is a release blocker regardless of context. No escape hatch – no `--figma-is-source-of-truth`, `--cloud-only-design`, `--proprietary-format-source`, `--hand-rolled-penpot-stack`, `--oauth-relay-automation`, `--skip-hydration-pass`, `--unprobed-tooling-claim`, `--rootful-design-stack` flag exists.

§ 6 07 22 . – Motion discipline: token-bound durations/easings + machine-readable manifest + reduced-motion static fallbacks everywhere (User mandate,
– –)

Forensic anchor – proving project (Helix Thready MVP design package, – –). Helix Thready's `design/motion/` package proves the discipline: six hand-authored vector Lottie animations bound to the token

motion values, each validated by parse + structural schema check + REAL render (lottie-web under jsdom with animation-progression assertions, plus headless-Chromium light+dark contact sheets) – unverified runtimes (lottie-compose / lottie-ios) honestly [OPEN] -tracked; FIVE static reduced-motion fallback SVGs whose geometry is DERIVED FROM THE LOTTIE KEYFRAMES (the exact poster frame for loopers, the final frame for one-shots); and a machine-readable motion-manifest.json v naming, per animation, its loop mode, poster frame, segment markers, and reducedMotionFallback – with one DELIBERATE null (an instant-cut transition) recorded with its reason, never a silent absence.

The mandate. All motion in every governed project MUST satisfy: () **Token-bound timing** – every duration and easing binds to the § motion tokens (the shared base durations + easing curves); documented choreography MAY extend them, but an animation timed by an ad-hoc literal that bypasses the token layer is a § violation (per § no-invented-values). () **Reduced-motion fallback per animation** – every playable animation ships a STATIC fallback derived from its own keyframes (poster frame for looping animations, final frame for one-shots), delivered as an open-format asset; where the correct reduced-motion behavior is intentionally something else (instant cut, no artwork), that is recorded EXPLICITLY in the manifest with its reason – a missing fallback is a gap, a silent missing fallback is a violation. () **Machine-readable manifest** – a versioned motion manifest enumerates every animation with its id, loop mode, poster frame, reduced-motion mode, and fallback reference; runtimes consume the MANIFEST, not folklore. () **prefers-reduced-motion honored everywhere** – every surface (web, library, screens, previews, implementations) gates motion on prefers-reduced-motion: reduce with TARGETED component-level gates (never a global

* .kill-switch, per the § . . review lens), swapping in the manifest-named fallback. () **Render-validated, honestly bounded** – no animation is claimed working without a real render in a real runtime (the § . . () analyzer discipline applied to motion assets: parse + structure + rendered geometry with progression assertions); runtimes not yet exercised are named [OPEN] items, never assumed.

Honest boundary (§ . .). The manifest + fallbacks prove reduced-motion COVERAGE; they do not prove physical-display fidelity (a pixel A/B of statics against frozen frames on a real display remains its own tracked item), and jsdom/headless renders do not certify mobile runtimes – the boundary is stated, tracked, and never blurred.

Classification: universal (§ . .) – the consuming project supplies its animation inventory, motion tokens, manifest path, and runtime list per § . . .

Composes § . . / § . . () / § . . / § . . / § . . / § . . / § . . . Propagation gate CM-

COVENANT-114-221-PROPAGATION (literal 11.4.221) +

recommended gate CM-MOTION-TOKEN-BOUND-REDUCED-FALLBACK

(every animation's timing resolves to tokens; the manifest enumerates every playable animation; every manifest entry names a fallback or an explicit recorded null-with-reason; every animation carries its parse/structure/render validation record) + paired § . mutation (add an animation with a literal ad-hoc duration → gate FAILs; drop a fallback reference without a recorded reason → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md

§ . . .

Non-compliance is a release blocker regardless of context. No escape hatch – no `--ad-hoc-duration`, `--no-reduced-motion-fallback`, `--global-star-motion-kill`, `--manifest-optional`, `--unrendered-animation-claimed-working` flag exists.

§ . . – Design-completion export wave: a design is not done until token codegen, the honest-output raster matrix, both-theme renders, hand-off bundles, and doc twins land (User mandate, – –)

Forensic anchor – proving project (Helix Thready MVP design package, – –). Helix Thready's export capstone proves the wave: an icon/asset raster matrix locked : to its generator (`generate-raster.sh`), which SELECTS its SVG renderer by PROBE – each candidate (`rsvg-convert` → `inkscape` → `resvg` → `headless Chromium` → `ImageMagick`) must pass a real probe render with content verification before it is used, a failing candidate falls through with a note, and if NONE passes the script exits with instructions and writes NOTHING ("No PNGs were written (output is not faked)"); an absent bundler (`iconutil` for `.icons`) SKIPS with a note – a verified live run produced REAL rasters, every one checked non-blank; screen PNGs (screens × light+dark at a confirmed × scale, theme forced deterministically, every file verified `size-floor + stddev >`); a –page bound design book; and PenPot + Figma hand-off bundles whose impossible/operator-only legs are recorded as DEFERRED with reasons – never faked.

The mandate. A design (a design system, a brand, a screen set, or a design change wave) is NOT "done" until its EXPORT WAVE has run and its outputs are verified: ()

Per-platform token codegen – the § . . generated bindings are (re)generated from the canonical

token source, or their absence is [OPEN] -tracked. ()

Full icon/asset raster matrix via an honest-output pipeline – the per-platform pixel/format matrix (every OS surface) is documented, locked : to a generator script, and produced by a pipeline that (a) PROBE-TESTS every candidate renderer with a real content-verified probe render before trusting it (\$. . – a renderer's presence is not its correctness; an unprobed renderer would silently emit blank/garbage icons), (b) falls through to the next candidate on probe failure, (c) SKIPs-with-note any step whose tool is genuinely absent (\$. . / \$. . honest-skip discipline), and (d) NEVER writes a placeholder, blank, or faked output – every produced raster is verified non-degenerate (size floor + non-blank content check; multi-resolution bundles verified to contain their declared members, per \$. . 's open-the-artifact discipline). ()

Both-theme screen renders + bound design book – every catalogue screen rendered in light AND dark at the declared scale with deterministic theme forcing, each render verified non-blank, and the set bound into the portable design book. ()

Hand-off bundles – the open-platform and proprietary-bridge bundles (\$. .) built with per-leg HONEST statuses: produced / deferred-operator-only / impossible-with-reason. ()

Doc twins regenerated – the design area's \$. . / \$. . export siblings are regenerated so no reader gets a divergent view. A design wave reported complete without its export wave – or with any faked/unverified export output – is a \$. . PASS-bluff at the design-delivery layer.

Honest boundary (\$. .). The export wave proves the design's ARTIFACTS exist, are real, and are non-degenerate; it does not prove visual correctness on physical displays (\$. . / \$. . / \$. . own that), and a wave rendered before

later additions is honestly stamped with its coverage boundary – regeneration lands as a tracked downstream pass, never claimed retroactively.

Classification: universal (\$. .) – the consuming project supplies its raster matrix, generator, renderer candidates, book builder, and bundle targets per

\$. . . Composes \$. . / \$. . /
\$. . / \$. . / \$. . /
\$. . / \$. . / \$. . /
\$. . / \$. . . Propagation gate CM-COVENANT-114-222-PROPAGATION (literal 11.4.222) +

recommended gate CM-DESIGN-EXPORT-WAVE-COMPLETE (the raster generator probe-tests renderers and content-verifies outputs; the matrix doc and generator agree : ; every screen has both-theme verified renders or a tracked gap; every bundle leg carries an honest status; the doc twins are in sync) + paired \$. mutation (make a renderer skip its probe, or write a placeholder PNG on renderer absence → gate FAILs; report the wave complete with a stale doc twin → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
\$. . .

Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-export-wave , --unprobed-renderer , --placeholder-raster-OK , --single-theme-renders , --fake-bundle-status , --stale-doc-twins flag exists.

\$. . – Provenance-marker discipline for design documentation + central open-item registry in the area index (User mandate, – –)

Forensic anchor – proving project (Helix Thready MVP design package, – –). Every statement in

the Helix Thready design area carries an inline provenance marker from a closed vocabulary ([VERIFIED] , [OPERATOR] , [RESEARCH] , [DEFAULT – adjustable] , [ASSUMED] , [BUILD-NEW] , [IN-HOUSE: module] , [CONSTITUTION \$x] , [GAP: id] , [OPEN: id]); the area index maintains the CANONICAL open-item registry – registered ids across nine families, VERIFIED AGAINST THE TREE in both directions (every id in the tree exists in the registry, no registry id is absent from the tree), with per-id state (open / narrowed / closed, each closure citing the closing revision) – and a verified-vs-assumption ledger separating measured fact from proposed default from named scaffold.

The mandate. Every design document in a governed project MUST carry provenance discipline: () **Mandatory marker vocabulary** – a closed, project-declared marker set whose minimum core is [VERIFIED] (established from source/runtime/measurement, with the evidence or check date), [OPERATOR] (an operator decision), [DEFAULT – adjustable] (a proposed default pending review), and [OPEN: id] (an unresolved item with a stable id); projects extend the set ([RESEARCH] , [ASSUMED] , [BUILD-NEW] , [GAP: id] , [IN-HOUSE: ...] , [CONSTITUTION \$x]) per § A load-bearing claim with no marker reads as UNVERIFIED. – and presenting an assumption as verified is a § . . . violation of § . . . PASS-bluff severity at the design-documentation layer. () **Every unresolved decision is an [OPEN: id]** – with a stable, never-renumbered id (§ . . . discipline), owned by a named file, and never papered over with an invented answer. () **Central registry in the area index** – the design area's index maintains the SINGLE canonical open-item registry: every [OPEN: id] in the tree appears in it, every registry row names its owning file(s) and state, and the registry is TREE-VERIFIED in both directions on every consolidation pass (an id in the tree missing from the registry, or a registry ghost with no tree presence, is a violation – the §

consistency discipline applied to design opens); closures cite the closing revision, and narrowed items record exactly what closed and what remains. () **Opens map to tracked work** – registry items map to the project's workable-item tracking (§ . . . / § . . . / § . . .) so the § . . . completion mandate applies: a design open is driven to closed-with-evidence or explicitly-closed-with-reason, never silently abandoned.

Honest boundary (§ . . .). Markers prove the EPISTEMIC STATUS of each claim, not its truth – a [VERIFIED] marker is only as good as its cited evidence (§ . . .), and the registry proves open-item ACCOUNTING, not progress (§ . . . owns completion). This anchor binds the DESIGN area to the same honesty regime the tracker anchors give code – it cites § . . . / . . . / . . . / . . . rather than duplicating them.

Classification: universal (§ . . .) – the consuming project supplies its marker-set extensions, id families, area-index path, and tracker binding per § Composes § . . . / § . . . / § . . . / § . . . / § . . . / § . . . / § . . . / § . . . / § Propagation gate CM-COVENANT-114-223-PROPAGATION (literal 11.4.223) + recommended gate CM-DESIGN-DOC-PROVENANCE-MARKERS (every design doc declares/uses the marker vocabulary; the area index registry exists; a tree-grep of open-item ids equals the registry set in both directions; every registry row names owner + state) + paired § . . . mutation (add an [OPEN: id] to a doc without registering it → gate FAILS; delete a registered id from the tree while the registry still lists it open → gate FAILS; strip the literal → propagation gate FAILS; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md
§ . . .

Non-compliance is a release blocker regardless of context.
No escape hatch – no --unmarked-claims-OK , --assumption-as-verified , --unregistered-open-item , --registry-optional , --silently-abandon-open flag exists.

§ . . – Test-first (TDD) for ALL work, not only fixes + a minimum code-coverage floor that is NECESSARY-never-sufficient (User mandate, – –)

Forensic anchor – verbatim user mandate
(– –):

"Any work we do MUST START by writing the test! We MUST FULLY comply and follow the TDD – The test driven development! It MUST BE applied for any change, implementation, wiring, tests, challenges and any other form of codebase we write! Bash scripts as well! Coverage with TDD MUST BE close to %, not less than % !!!"

The gap this closes. § . . mandates the test-first (RED → ... → GREEN → VERIFY → DOCUMENT) workflow but its operative text binds fixes only ("Every fix MUST follow the –step TDD-fix workflow"); § . . refines that RED step onto the broken artifact; § . . binds reproduce-first. NONE of them binds a brand-new feature, a new wiring/integration, a new gate, a new Challenge, or a NEW BASH SCRIPT – none of which is a "fix", so none of which inherits a test-first obligation today. Separately, § . . and § . . mandate % test-TYPE coverage (breadth: WHICH KINDS of test exist per feature × platform) – an orthogonal axis

that states no **code-coverage** figure anywhere.

\$. . . closes exactly those two holes and claims nothing else as new.

(A) TEST-FIRST BINDS ALL WORK, not only fixes. For EVERY work product with an executable surface – a change, a new feature, an implementation, a wiring/integration, a refactor, a pre-build gate, a \$. . . mutation, a Challenge, test-harness code itself, AND **bash / shell scripts** – the test is WRITTEN FIRST, RUN FIRST, and OBSERVED TO FAIL for the right reason before the implementation exists; only then is the implementation written to make it pass. Where a broken artifact already exists the RED baseline is taken ON THAT ARTIFACT per \$. . . (and its polarity switch converts the same source into the permanent \$. . . guard); where the work is genuinely new, the RED is the test failing against the absent/unimplemented behaviour, captured before the implementation lands. A test authored AFTER its implementation demonstrates only that the test AGREES with the code that already exists – never that it CATCHES the defect (the \$. . . PASS-bluff, here generalised to every work class). **What counts as "the test" per work class (\$. . . – named, never left to interpretation):** behaviour code → an executing test asserting the behaviour; a bash/shell script → an executing test that invokes the script through its REAL invocation path and asserts its real observable effect (exit status + emitted output + state delta), never a bash -n parse-check alone (parse-cleanliness is \$. . . 's separate obligation and proves nothing about behaviour); a pre-build GATE → its paired \$. . . mutation IS the test-first artifact, authored and observed to make the gate FAIL BEFORE the gate is trusted (a gate TRUSTED before its mutation was observed to FAIL is unvalidated instrumentation per \$. . . (F) – the condition is observation-before-trust, NOT authorship

order, so a gate and its mutation authored in the same commit are compliant when the mutation is observed to make the gate FAIL before the gate is relied upon); a Challenge → its scored assertion against captured evidence per § . . /§ . . ; a behaviour-bearing config/data change (an init-script line, a package-set manifest, an action-registry row – any committed DATA an executing consumer reads to decide behaviour) → an executing test that drives that CONSUMER through the changed value and asserts the changed observable behaviour (the clause-(F) executable-surface test reaches the config THROUGH its consumer – a config change is never parked under the prose-exempt class, which is strictly pure prose / governance / documentation).

(B) CODE-COVERAGE FLOOR – ≥, % MINIMUM, ~ % TARGET. Every governed codebase carries a measured code-coverage figure that MUST NOT fall below %, with as close to % as the domain permits as the standing target (composing § . . 's " % coverage as the domain permits" phrasing on its own test-type axis). The floor applies to the executable codebase INCLUDING bash/shell scripts – a shell corpus exempted from measurement is exactly the gap the operator's "Bash scripts as well!" closes.

(B.) SEVEN-TYPE BREADTH – the % AIM across a canonical test-type enumeration (extension, – –). The % floor is a LINE-COVERAGE proxy across the executable codebase; the % AIM is the STANDING TARGET across the seven-canonical-test-type × executable-surface product per operator mandate – – . The seven canonical types every consuming project MUST cover (where applicable) are: () UNIT, () INTEGRATION, () END-TO-END (E E), () FULL AUTOMATION, () SECURITY, () PERFORMANCE AND BENCHMARK, () ANTI-BLUFF – each carrying its own § . .

mandatory-breadth obligation (which this anchor cross-references, and § 11.4.21's clause (B) is amended in step). A project measuring % line coverage on UNIT tests alone while omitting one of the seven applicable types is a § 11.4.21 (B.) violation regardless of the line-coverage number (§ 11.4.21 () metric-validated-against-done: the correct end-state is line coverage AT target ACROSS all seven applicable types, never on one type alone). Honest boundary § 11.4.21 / § 11.4.21 : a test type genuinely inapplicable to a project (e.g. a library with no runtime surface has no E E user-journey to drive) is HONESTLY marked as § 11.4.21 SKIP-with-reason at the type-breadth axis, never silently omitted; the universal DISCIPLINE is unchanged, the applicability of each type is the consumer's § 11.4.21 DATA. CM-COVERAGE-FLOOR is EXTENDED – the measured numerator STILL counts only RED-capable tests, PLUS the gate FAILs on missing test-type-breadth coverage per the seven-type enumeration (a project measuring % line coverage but with zero ANTI-BLUFF tests → gate FAILs); a paired § 11.4.21 mutation that strips one enumerated type from the type-breadth check → the gate FAILs.

(C) COVERAGE IS NECESSARY, NEVER SUFFICIENT – the anti-gaming clause (§ 11.4.21 () / § 11.4.21 / § 11.4.21). A coverage percentage is trivially gameable and MUST NEVER be reported as evidence a change is correct: **MEASURED FACT (in-session, - - , this host):** a line-granular coverage instrument reported an if ...; then ...; else ...; fi line as COVERED while the else branch was NEVER taken – the metric read fully-covered on a line half of whose behaviour was untested; and a test that EXECUTES a line while ASSERTING NOTHING raises the same number by the same amount as a test that proves the behaviour. Therefore: () the REAL bar remains the § 11.4.21 anti-bluff covenant + § 11.4.21 – every test MUST catch its own

negation, „PROVEN by a paired mutation that makes it FAIL; a line executed by an assertion-free test is NOT covered in the sense that matters; () the floor is a MINIMUM ON A MEASURABLE PROXY, never a certificate of correctness – citing "≥ % covered" as proof a change works is a § . PASS-bluff at the metric layer; () the floor NEVER stands alone – it composes with the § . mutation pair, the § . . /§ . . /§ . . captured-evidence bar, and the § . . runtime signature, and it does NOT displace any of them.

(D) METRIC VALIDATION AGAINST THE DEFINITION OF DONE

(§ . . ()) – **split verdict, stated honestly.**

§ . . () requires every gating metric to be evaluated at the correct end-state and REFUSED if the correct end-state does not move it to target. **Directional validation, performed – – :** at the correct end-state – every behaviour test-driven, every test catching its own negation – the lines implementing that behaviour ARE executed by their own tests, so the metric DOES move to target; the correct end-state does NOT refute the floor, and the floor is therefore VALID as a NECESSARY condition. **The converse is REFUTED** by the clause-(C) measured gameability fact: reaching the floor does NOT imply the end-state – hence necessary-never-sufficient, and hence clause (C) is load-bearing rather than decorative. **OWED, not claimed** (§ . .): the EMPIRICAL calibration of the specific % threshold against a given project's REAL corpus – whether a genuinely correct end-state on THAT corpus lands at %, higher, or lower – has NOT been performed and MUST be tracked as a § . . work item by every consuming project that enforces the floor; this anchor asserts no project has already filed it (asserting an unverified tracker state would be the § . . violation this clause exists to prevent). **Interim precedence** (mirroring the precedent § . . ())

itself set for the § . . . ratchet): until that calibration lands the $\geq$ % floor KEEPS GATING – an uncalibrated floor blocking a merge is bounded, visible and operator-resolvable, whereas disarming the only mechanism that surfaces untested code silently re-opens the hole (the § . . . reversible-safe choice); the calibration is OWED WORK, never a licence to switch the floor off.

(E) HOW COVERAGE IS MEASURED – per language, honestly (§ . . . / § . . .). The consuming project supplies its per-language coverage instrument as DATA per § . . . ; the UNIVERSAL rule is that an instrument MUST exist and MUST actually be RUN – a floor nobody can measure is an unmeasurable gate and is itself a § . . . bluff, and a percentage nobody measured is a § . . . fabrication. **Measured host reality** (– , the project instantiating this anchor): Go coverage is available built-in (go test -coverprofile + go tool cover); C/C++ via gcov ; and for bash/shell NO dedicated coverage tool was installed – kcov , bashcov , bats , shellspec and shunit2 were ALL probed ABSENT (shellcheck IS present but is a LINTER, not a coverage instrument, and satisfies nothing here). **PROVEN-FEASIBLE bash mechanism (measured in-session, never asserted):** trace-based line accounting needs NO additional tooling – PS4='+COV:\${BASH_SOURCE##*/}:\${LINENO}:' bash -x <script> emits the executed line numbers, and executed-lines ÷ executable (non-blank, non-comment) source-lines yields line coverage; a probe run correctly identified an uncovered function body as uncovered, and subshell bodies ARE traced. **HONEST LIMITS of that mechanism (measured, stated as fact):** it is LINE coverage, NOT branch coverage – the clause-(C) same-line if/else counted as covered with one branch never taken; set +x regions and traps are not accounted. It is therefore a FLOOR-measuring instrument only; where branch-level assurance is required the project

MUST install a real coverage tool (e.g. `kcov`) or record the shortfall honestly. Where NO instrument exists for a language on a host, the honest responses are to INSTALL one, or to record an `$ . . /$ . .` SKIP-with-reason plus a tracked `$ . .` item – NEVER an invented percentage, and NEVER a silently-dropped floor. **EXCLUSION-LIST FENCE (amendment, remediation round of - - - the \$. . exemption pattern applied to the coverage corpus)**. The corpus-scope exclusion list is NOT an unconstrained consumer knob: an exclusion is legal ONLY via a CHECKED-IN exclusion list, with EVERY entry justified from a CLOSED class set – `generated-code` (output of a committed generator per `$ . .`), `vendored-third-party` (executable code the project does not author), `non-shipping-fixtures-and-golden-assets` (test fixtures / golden files that never execute in the shipped product) – and enumerated as honest gaps in the coverage report, never silent (the exact discipline `$ . .` applies to topology exemptions: legal ONLY via a checked-in map, enumerated as honest gaps – never silent). Excluding FIRST-PARTY executable code additionally REQUIRES a tracked `$ . .` work item naming the plan to bring it into scope. An unlisted, unjustified, or un-tracked-first-party exclusion VOIDs the measured figure – a consumer that excludes most of its first-party executable code and truthfully measures $\geq$ % on the remainder has satisfied the operator's floor in letter and voided it in substance; that silent-evasion channel is exactly what this fence closes.

(F) HONEST BOUNDARY / SCOPE (\$. .). "All work" means work products with an EXECUTABLE surface. A pure prose / governance / documentation edit has no executable surface and is NOT coverage-scoped – its discipline is `$ . . / $ . . / $ . . / $ . .` plus the `$ . .` universal review; quoting a coverage percentage for a prose document

would be meaningless, and demanding one would be a § 11.4.188. () false-positive refusal. BUT a governance change that SHIPS an executable artifact (a gate, a hook, a script, a mutation) is FULLY in scope for both... (A) and (B). Test-first does NOT replace the § 11.4.27. 11.4.189 11.4.190 four-layer verification, the § . . . full-suite retest, or the § . . . manual-QA final gate – it PRECEDES them. And § . . . 's CODE-coverage axis is DISTINCT from § . . . / § . . . 's % TEST-TYPE coverage (breadth of test kinds per feature × platform): both apply simultaneously, neither substitutes for the other, and satisfying one is never evidence for the other.

Classification: universal (§ . . .) – references no project-specific hardware, vendor, or language; the consuming project supplies its per-language coverage instruments, its corpus scope, its CHECKED-IN clause-(E)-fenced exclusion list, and its threshold calibration per § 11.4.61. . . . Composes § . . . / § . . . / § . . . / § 11.4.67 11.4.69. / § . . . / § 11.4.77. 11.4.85 11.4.87 (generalises its fix-scoped test-first to ALL work) / § . . . / § 11.4.122. 11.4.130 (brownfield adoption is an operator question) / § . . . / § . . . / § . . . / § 11.4.185 11.4.197. / § . . . / § 11.4.281. 11.4.282 / § . . . / § . . . / § . . . / § . . . / § . . . (orthogonal test-TYPE axis) / § . . . / § . . . / § . . . () / § . . . / § . . .

Propagation gate CM-COVENANT-114-224-PROPAGATION (literal 11.4.224) + recommended gates CM-TEST-FIRST-ALL-WORK (every commit adding an executable artifact carries its test, and that test's failing-first observation is recorded – a new script/gate/feature landing with no test, or with a test whose RED was never captured, FAILs) + CM-COVERAGE-FLOOR (a

measured coverage figure exists per declared language, was produced by a real instrument run rather than asserted, and is $\geq$ the declared floor; a declared language with NO instrument FAILs rather than silently passing; **the measured numerator counts ONLY lines covered by RED-CAPABLE tests** – every test contributing to the measured corpus MUST carry a recorded failing-first run (its clause-(A) RED record) OR a paired mutation observed to make it FAIL, with sampling-based mutation verification an acceptable stated implementation and a contributing test carrying ZERO assertions rejected by construction, so assertion-free execution cannot raise the measured figure and ONE real RED-captured test plus assertion-free padding to % FAILs this gate; AND the gate FAILs on any corpus exclusion that is absent from the checked-in list, not justified from the clause-(E) closed class set, or first-party without its tracked \$. . item) + paired \$. mutations (land an executable artifact with no test $\rightarrow$ CM-TEST-FIRST-ALL-WORK FAILs; record a coverage percentage with no instrument run behind it, or drop a language's measurement entirely, or pad the numerator with an assertion-free test, or slip an unjustified exclusion into the corpus scope $\rightarrow$ CM-COVERAGE-FL00R FAILs; strip the literal $\rightarrow$ the propagation gate FAILs). **Brownfield adoption is DELIBERATELY UNDEFINED by this anchor (\$. .):** when CM-COVERAGE-FL00R first lands on a project whose legacy corpus measures below %, the gate-code work item MUST surface the adoption question to the OPERATOR per \$. . BEFORE first enforcement – candidate options to present: immediate hard floor (day-one RED until the corpus is raised) / a ONE-TIME monotone-decrease ratchet per the \$. . precedent / per-corpus phase-in (per language or module as each enters scope) / floor-on-changed-code-only with a scheduled full-corpus deadline – recording the operator's answer as consumer DATA (\$. .); this anchor does NOT autonomously

weaken the operator's unconditional "not less than
% !!!" with an invented ratchet (a \$. . -
class decision the operator owns). Gate-code = a separate
work item – the CONTRACT is landed here, the gate CODE is
NOT claimed shipped (\$. .).

Canonical authority: constitution submodule Constitution.md
\$. . .

Non-compliance is a release blocker regardless of context.
No escape hatch – no --test-after-OK , --skip-red-first , --
bash-exempt-from-tests , --coverage-not-applicable , --below-
floor-OK , --coverage-proves-correct , --unmeasured-percentage-
OK , --assertion-free-test-counts flag exists.

\$. . – Scheduler-quota burst-
throttling telemetry + interactive-scope
isolation from bursty fleets: progressive
interactive degradation is diagnosed from
throttling accounting over time, never from
averages, and an interactive-latency-sensitive
process never shares a fixed-quota no-burst
scope with a bursty fleet (research-derived,
– –)

Forensic anchor (genericised, – – – two
independent instances in one day). (a) Progressive
interactive sluggishness – worsening with the number of
live subagents – was root-caused to a scheduler CPU quota
(a cgroup cpu.max limit) with NO burst allowance on a
session scope hosting BOTH the operator's interactive
terminal-multiplexer server AND a spawning agent fleet.
Measured live on a .-CPU host: the scope's quota was
. CPU per ms enforcement period; average
demand was . CPU – comfortably UNDER quota – yet
of periods (. %) were throttled
during active phases, ~ . runqueue-seconds of forced
idle per wall-second at peak, % throttled during quiet

phases (episodic, fleet-phase-correlated), and the lifetime ratio (. %) vs the live ratio (. %) showed it WORSENING with session age. The terminal server was co-resident in the throttled scope: it froze for $\geq$ ms per throttled period, and consecutive throttled periods compounded into multi-second interactive stalls. Every conventional instrument read HEALTHY – CPU average under quota, load average unremarkable, memory fine, thread counts fine – the degradation was INVISIBLE to averages and lived entirely in the throttling accounting (`cpu.stat` `nr_throttled` / `throttled_usec` deltas). (b) A sibling session launcher shipped a FIXED default quota (CPUs) on the same -core host – . % of periods throttled and s of cumulative forced idle – fixed by computing the default from real detected host capacity (a per-core fraction with a floor and a cap), proven by a four-quadrant polarity test reading the LIVE cgroup back.

Two operative halves, BOTH mandatory:

(A) DIAGNOSTIC – throttling accounting over time under load, never averages. Any investigation of progressive / bursty / load-correlated interactive degradation (sluggish TUI or terminal, editor lag, stalls that worsen with parallel-agent or worker count) MUST sample the scheduling scope's THROTTLING ACCOUNTING – on Linux cgroup v , `cpu.stat` `nr_throttled` + `throttled_usec` as DELTAS over a measured window under load, expressed as a fraction of elapsed enforcement periods, WITH a quiet-phase control sample (the \$. . ())(b) needle in time: a mechanism that fires % under load and % quiet is proven load-coupled, not ambient) – BEFORE any verdict on the cause is minted. A HEALTHY AVERAGE WITH A HIGH THROTTLERATIO IS THE TRAP: quota enforcement is per-period (canonically ms), so a workload can sit far below its quota ON AVERAGE while a large fraction of

individual periods hit the ceiling and stall every co-resident process; CPU averages, load average, memory, and thread counts are all structurally BLIND to it – their quiet "healthy" readings are \$. . () FALSE-NULLS, and "the averages look healthy, so the host is fine" without throttle-accounting deltas is a \$. . guess. The same discipline applies to EVERY quota'd scheduler (container CPU limits, VM caps, hypervisor shares / steal time): the throttle/steal accounting IS the authoritative source (\$. .); the utilisation average is a proxy that the real condition does not have to satisfy.

(B) DESIGN – an interactive-latency-sensitive process NEVER shares a fixed-quota, no-burst scope with a bursty fleet. An interactive-latency-sensitive process (terminal-multiplexer server, TUI, editor, operator shell) MUST NOT be co-resident in a fixed-quota scheduling scope that has NO burst allowance together with a bursty multi-process fleet (spawning subagents, build workers, test runners): the fleet's demand spikes consume whole enforcement periods and the interactive process freezes with them – in bursts, progressively as the fleet grows (the \$. . /\$. . parallel-stream mechanics make the fleet's demand spiky BY DESIGN, so the collision is structural, not incidental). The consuming project MUST () grant the shared scope a burst allowance sized for the fleet's measured spike profile, OR () isolate the interactive process in its OWN scope (or leave it unquoted where host policy allows), OR both. Quota values MUST be HOST-ADAPTIVE – computed from real detected capacity with a documented floor and cap (\$. . 's dynamic-detection discipline; \$. . never hardcode), never a fixed literal that silently strangles a many-core host (the forensic clause-(b) instance); and a quota change is verified by READING THE LIVE scope's values BACK (\$. . runtime signature – the

configured value is not the effective value until read back). Composes § 11.4.6 (memory ceiling) + § 11.4.6 (thread/RLIMIT_NPROC headroom) as the THIRD orthogonal host-exhaustion axis: abundant RAM and abundant thread headroom do NOT imply an un-throttled scheduler – each axis has its own authoritative instrument, and none substitutes for another.

Honest boundary (§ 11.4.6): throttle telemetry proves WHERE enforcement periods were lost – it does NOT prove throttling is the ONLY degrader (network RTT floors, in-process GC/event-loop load, and sink-side effects remain separately measurable; the founding investigation ranked them independently and REFUTED five candidate mechanisms with the same instrumentation rigor); and clause (B) bounds PLACEMENT + SIZING – it does not exempt the fleet from the § 11.4.6 / § 11.4.6 / § 11.4.6 host-safety ceilings, and it never licenses removing a quota that host policy requires (widen/isolate/burst within policy, per § 11.4.6).

Classification: universal (§ 11.4.6) – scheduler-quota throttling is a platform-general mechanism (cgroup v 1/v 2, container runtimes, VM hypervisors); the consuming project supplies its concrete scope layout, quota/burst values, sampling windows, and instruments per § 11.4.6. Composes § 11.4.6 / § 11.4.6 (build-resource stats – the sampler seam this telemetry extends) / § 11.4.6 (systematic-debugging – the throttle sample is Phase-1 fact-gathering, and the forensic investigation reached it only by refusing to stop at healthy averages) / § 11.4.6 (read the live scope back) / § 11.4.6 () (observer-effect budget – the sampler itself stays lightweight) / § 11.4.6 (authoritative-source measurement; the averages-only verdict is a false-null) / § 11.4.6 / § 11.4.6 / § 11.4.6 / § 11.4.6. Propagation gate CM-

COVENANT-114-225-PROPAGATION (literal 11.4.225) + recommended gates CM-QUOTA-THROTTLE-TELEMETRY (every progressive/bursty-degradation investigation's verdict cites throttle-accounting DELTAS sampled over a measured window under load + a quiet-phase control; an averages-only verdict on a progressive-degradation issue → FAIL) and CM-INTERACTIVE-SCOPE-NOT-QUOTA-STARVED (no interactive-latency-sensitive process is co-resident in a fixed-quota no-burst scope with a bursty fleet; quota defaults are host-adaptive with a live read-back, never fixed literals) + paired § . mutations (record a progressive-degradation verdict citing only averages → the telemetry gate FAILs; pin a fixed low quota literal as the default on a detected many-core host, or co-locate the terminal server with the fleet in a no-burst quota scope → the isolation gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md § . . .

Non-compliance is a release blocker regardless of context. No escape hatch — no --averages-prove-healthy , --skip-throttle-telemetry , --fixed-quota-default-OK , --interactive-in-fleet-scope-OK , --no-burst-needed , --skip-quiet-phase-control flag exists.

§ . . — Evidence-class-at-closure + standing detection pressure: machinery presence does NOT predict whether a fix holds — the EVIDENCE CLASS at closure (runtime vs source) under real DETECTION PRESSURE does; prose does not bind, seams do (research-derived, — —)

Forensic anchor (genericised — the — — reopen/first-touch root-cause forensics; operator mandate — — , verbatim: "do deep research on how we

have for very long time so much reopened workable items!?

Why the quality of delivered releases was so poor mostly?

We mostly detect the same major problems on first touch when new releases are available, and this happens over and over again!"). The single strongest empirical result across four independent evidence lines (internal forensics; a -source external literature pass; one cross-project corpus WITHOUT tracking machinery; one corpus WITH exemplary machinery): **every fix whose done-claim was preceded by an on-target RED→GREEN polarity flip that held recorded ZERO reopens; every fix confirmed only at source-green bounced.** Cross-corpus: the runtime-evidence-closure corpus recorded median attempts-to-stick = and a ~ . % corrected reopen rate (its own analysis marks that a FLOOR under weak detection pressure), while the source-green corpus recorded % – itself a proven undercount; the healthy corpus's ONLY bouncing fix family was closed by a wrong-layer echo-assertion (a unit test asserting a Noop echo – the discriminator's predicted loser, sign right); commits in a two-month window moved claims at "PENDING-BUILD"; the last audited release entered manual QA with of in-scope items validated at the user-visible layer – the operator was the FIRST user-layer oracle most claims had ever met, so first-touch rediscovery of the same majors was the deterministic output of the pipeline shape, not bad luck; a pixel-layer defect was "verified" by a file-existence check plus five greps; and the corpus WITHOUT machinery shipped a "fully working" release hours before a total field brick with its whole hermetic (wrong-layer) suite green. The law this evidence selects: **machinery presence (trackers, guards, gates, anchors) is signal and custody, NOT prevention; what predicts whether a fix holds is the EVIDENCE CLASS AT CLOSURE (runtime vs source) under real DETECTION PRESSURE.** Prose does not bind

– seams do. Coverage has a blind complement at any maturity.

The mandate (ALL hold; the STRICT FORMALIZATION of § . . . (F)'s evidence-class clause + § . . . (D)'s custody chain + § . . . 's four layers – it supplies the closed taxonomy, rank floor, machine fields, and freshness contract those anchors did not state):

() Closed evidence-class ranking with a per-defect-layer floor, enforced at every closure seam. Defect layers form the closed set {user-visible, runtime, artifact, source} (§ . . .); evidence classes form the closed RANKED set {runtime > artifact > source}. Every terminal or ready-class status write carries evidence whose class meets its defect layer's floor: a user-visible or runtime-layer defect can NEVER close on artifact- or source-class evidence; source-on-source stays legal (the § . . . () false-refusal guard – legitimate source-layer work is never refused).

() Class is proven by MACHINE FIELDS, never by a label. runtime requires a target fingerprint read FROM the target at run time (§ . . . (F)) plus a runtime observable; artifact requires the artifact path plus its content hash; source requires the source ref. A class label without its fields is SHAPE-INCOMPLETE and refused; unreadable evidence is BLIND, never accepted.

() The ANTI-ECHO rule. A "runtime observable" that IS a grep/source-scan transcript is WRONG-LAYER and refused – a grep can never observe a pixel; the wrong-layer oracle dies by construction (the pixel-defect-by-five-greps closure and the Noop-echo family both cross this seam and are refused at it).

() Per-status evidence preconditions. Every status in the § . . . /§ . . . lifecycle carries a

CHECKED evidence precondition – not only the terminal write. Forensic FACT: a penultimate "ready"-class status functioned as design-complete-but-UNVERIFIED and parked items – % of the done-claiming population; the operator sampling "done" items was sampling a set dominated by never-verified work without either party being wrong about their own definition (semantic drift between claim-maker and claim-consumer). A status that sounds like progress but carries no checked precondition becomes the parking lot for unverified work.

() **Fix-kind honesty.** Every defect closure classifies its fix-kind; a STATE-ONLY repair (restart / cache-clear / re-deploy / manual re-run) or a mitigation CANNOT close a defect item as fixed – it is recorded as a mitigation and the item stays open or is honestly reclassified (composes § . . . 's Iron Law + § . . . ; forensic FACT: a sibling corpus's recurrence chains – chains, median and max attempts per defect – were dominated by state-only "fixes" that closed and bounced).

() **STANDING DETECTION PRESSURE.** Detection pressure is a CAUSAL INPUT to every measured quality rate – a reopen/defect/quality figure reported without its detection-pressure regime is incomparable and misleading (§ . . .): the measured . . . % -vs- ~ . . . % cross-corpus split is explained by pressure + closure class, not machinery, and the healthy corpus's one genuine recurrence surfaced at days – the next SWEEP, not the next use ("we only see what we test", § . . .). Therefore: every registered, topology-present guard carries a FRESHNESS CONTRACT – a verdict on the CURRENT artifact fingerprint, no older than a declared staleness budget; the never-executed / stale-fingerprint / over-budget set feeds a STANDING, risk-ordered re-run queue (most-reopened-first per § . . . , then stalest-first)

that the zero-idle loop ($\$ \dots / \$ \dots$) drains as regular development-cycle work; the $\$ \dots$ release seam remains the blocking BACKSTOP, never the only executor. Forensic FACT: two standing guards, finally executed during the forensics, immediately emitted the FAILs that had been latent the whole time – one reproduced the operator's loudest defect / ; registration had been treated as coverage while execution was nobody's job. Topology-absent guards are enumerated, never queued (no false pressure, $\$ \dots ()$); an empty registry is BLIND, never "all fresh"; staleness budgets are consumer data ($\$ \dots$) bounded by the alarm-fatigue trade-off – the queue is a work FEED, only the release seam blocks.

Regular-cycle + review + pull binding (operator mandate $\dots - -$). This anchor applies during EVERY regular development cycle (the freshness queue is standing $\$ \dots$ work, and every closure crosses seam ()); it is CHECKED through every code-review per $\$ \dots ()$; and on every constitution fetch/pull the $\$ \dots$ post-pull validation sweep + the $\$ \dots$ auto-propagation hook analyze the new additions and run their gates against the post-pull tree, with violations tracked ($\$ \dots$) and cleared – no parallel mechanism is invented.

Honest boundary ($\$ \dots$). (i) The discriminator is a measured correlation with a stated mechanism and zero counter-cases in the record – NOT a randomized trial (guarded items had less operator exposure; the clean exemplars are few; the confound is carried, not hidden). (ii) Fabricated machine fields are not fully closed here – provenance stays with $\$ \dots$ (F) harness-written verdicts + the $\$ \dots$ candidate-fingerprint join, where an invented fingerprint fails to match the real candidate and reads as absence; layered, not

eliminated. (iii) Oracle calibration stays § 11.4.107 () – runtime-class evidence from a mis-calibrated analyzer still passes this seam; goldens are necessary, not sufficient. (iv) The defect-layer taxonomy is authored data – mislabelling a user-visible defect as source-layer lowers its floor; the § 11.4.107 (a) review class is the detective counter (cheap to spot: the item's own description names user symptoms). (v) Re-running a weak guard often proves nothing – oracle strength remains § 11.4.107 (F)'s job, and orthogonal discovery (§ 11.4.107 chaos, § 11.4.107 discovery pressure) is a different feed this queue does not generate.

Reference mechanisms (proven, NOT wired – § 11.4.107 / § 11.4.107). Runnable seam mechanisms with RED-first POCs, golden-good/golden-bad/negative-control fixtures (§ 11.4.107 ()), live-data runs, and stated boundaries exist in this repository's quality-solutions research corpus (docs/research/quality/solutions/ – SOL-status-custody, SOL-verdict-coverage, SOL-evidence-class, SOL-detection-pressure; / GREEN across the ten-POC set on a deterministic second iteration). Per § 11.4.107 they carry NO force until their seams are wired by separate tracked work items (§ 11.4.107); § 11.4.107 (A)'s ledger is the instrument that keeps that wiring debt visible.

Classification: universal (§ 11.4.107) – no hardware/vendor/project literal; the consuming project supplies its defect-layer bindings, evidence-store paths, verdict registries, and staleness budgets as DATA per

§ 11.4.107 . Composes § 11.4.107 / § 11.4.107 /
 § 11.4.107 / § 11.4.107 / § 11.4.107 /
 § 11.4.107 / § 11.4.107 / § 11.4.107 /
 § 11.4.107 / § 11.4.107 () /
 § 11.4.107 / § 11.4.107 (F) /

§ 11.4.210 • 11.4.214 / § . . / § . . /
 § . . / § . . (D) /
 § . . / § . . /
 § . . () / § . . /
 § . . / § . . . Propagation gate CM-COVENANT-114-226-PROPAGATION (literal 11.4.226) +
 recommended gates CM-EVIDENCE-CLASS-AT-CLOSURE (every
 terminal/ready-class status write carries class-matched,
 machine-field-proven evidence meeting its defect layer's
 floor; a grep-transcript "runtime observable" → FAIL; a
 state-only repair closing a defect as fixed → FAIL) and
 CM-GUARD-FRESHNESS-SCHEDULER (every registered, topology-
 present guard is fresh-on-the-current-fingerprint or
 present in the visible, risk-ordered re-run queue; an
 empty registry → BLIND, never "all fresh") + paired § .
 mutations (close a runtime-layer item on source-class
 evidence → the closure gate FAILs; attach a runtime label
 with no machine fields → FAILs; let a registered guard sit
 past its staleness budget with no queue entry → the
 scheduler gate FAILs; strip the literal → the propagation
 gate FAILs; gate-code = separate work item, NOT claimed
 shipped § . .).

Canonical authority: constitution submodule Constitution.md
 § . . .

Non-compliance is a release blocker regardless of context.
 No escape hatch — no --source-green-closes-runtime-defect , --
 label-is-class , --echo-observable-OK , --registration-is-
 coverage , --rate-without-pressure-regime , --mitigation-closes-
 defect , --skip-freshness-queue flag exists.

§ . . – Governance-corpus self-custody: the rule corpus is bound by the same custody it imposes – every named gate is implemented-or-registered-deferral under a monotone ratchet, and propagation gates count anchor BLOCKS (exactly-once + lockstep-identical), never bare literals (research-derived, – –)

Forensic anchor (measured on THIS corpus,

– – – commands and control needles preserved in the root-cause analysis). The canonical rule file held , lines and distinct

§. .N anchors; named CM- gates, of which (%) had no implementation in either canonical gate site, plus literal "gate-code = separate work item" deferrals; the status-custody seam DESIGNED on

– – was still prose on

– – – the corpus's highest-value

remediation fell into its own gap within five days; the two anchors that state the fix-confirmation discipline verbatim predate the failures they failed to prevent by five-plus weeks. The positive complement: **every fully-mechanised rule in the record either prevented or exposed its failure mode** (the suite PENDING-semantics + the release verdict-coverage seam landed within a day of diagnosis and now block) – the gap is not rule quality, it is the prose-to-seam conversion rate. The SHAPE defect: the propagation-gate family – the ONLY gate family that runs everywhere – asserts literal-presence ≥ , mechanically rewarding restatement over implementation ("the corpus measures its own health in words present, and so it grows words"); measured consequences: two anchor blocks duplicated verbatim-divergently in ALL FOUR mirrors while every presence-gate stayed green; one anchor NUMBER heading two DIFFERENT mandates (the § . . / § . . collision – independently re-detected by

the block-integrity instrument's live run; resolution is an operator-owned re-mint, tracked); two canonical entry-point mandates contradicting each other (the mandated resumption file unreachable from the mandated README). External corroboration: a mandated-checklist adoption WITHOUT workflow embedding produced no measurable improvement across hospitals (the Ontario null, cited in the external-research corpus) – prescription without a seam does not move outcomes at any scale.

The mandate (ALL hold):

(A) NAMED-GATE LEDGER + IMPLEMENTATION RATCHET. Naming a gate is a commitment with a same-commit cost: every CM-* -class gate token named anywhere in the governance corpus MUST, at every build (of the governance repo and of every wired consumer), be either IMPLEMENTED – the token present in an EXECUTABLE gate site; prose carriers NEVER count (§ . . . ()(a)) – or covered by a REGISTERED DEFERRAL row pointing at a tracked work item (§ . . .). Silent gate debt is refused at the build seam. The unimplemented count is a MONOTONE-DECREASING ratchet, validated per § . . . () against the definition of done (at the correct end-state – every named gate implemented – the count is , so the metric reaches target and is valid as a necessary floor). A gate NAME that vanishes from the corpus without an explicit removal citation FAILS – the metric-gaming channel (deleting names to lower the count) is closed structurally; repeals stay legal and visible (the § . . . pattern). **An anchor's "done" state is its SEAM landing, not its TEXT landing** – a rule authored without its seam is, by § . . . 's own criterion, in the worse-than-nothing state, and this ledger makes that state visible, counted, and shrinking instead of silent. Brownfield adoption (ratchet down from the measured baseline vs hard-fail day one) is an operator

§ . . . decision mirroring § . . . (E)'s adoption fence – never an invented ratchet.

(B) ANCHOR-BLOCK INTEGRITY – the propagation-gate SHAPE

fix. Every propagation-class gate MUST count anchor BLOCK-STARTS (structural, line-anchored heading patterns), never bare literals: (i) exactly ONE block per anchor per governance file – is the old absence check, > is the DUPLICATION check, FAIL naming the anchor and file; (ii) content-hash EQUALITY across the declared lockstep mirror set (§ . . .) – divergent copies FAIL naming both files (compact-vs-full variance ACROSS layers – canonical file vs mirror – remains legitimate; equality binds only within the declared lockstep set, and canonical-vs-mirror semantic drift stays the § . . . cross-doc family's job); (iii) mid-body citations of an anchor literal are CARRIERS and never count (§ . . . () (a)); (iv) ZERO blocks extracted from a governance file is BLIND-OR-EMPTY, never clean (§ . . . ()); (v) one anchor NUMBER heading two different mandates is a COLLISION and FAILS (§ . . . 's never-reuse discipline applied to anchor ids); (vi) the block extractor MUST capture full dotted ids so a sub-anchor (...10.A) never prefix-matches its parent (...10) – a measured § . . . () false positive in the founding instrument itself, fixed RED-first.

(C) Consolidation pressure + honest boundary. Nothing mechanical decides SEMANTIC contradiction between two different anchors – only the decidable projections are mechanized here (block counts, lockstep hashes, anchor numbers; fenced verdict scopes stay § . . . (), seam declarations stay § . . .); the remainder is § . . . /§ . . . review territory – stated, never claimed. This anchor adds the measured enforcement-ratio INSTRUMENT, not a growth ban: the

counterweights to corpus growth remain § 11.4.23 classification, § 11.4.23 token-efficiency, and the visible-repeal precedent (§ 11.4.23).

Regular-cycle + review + pull binding (operator mandate
 - -). The ledger regenerates and diffs on every governance edit and every build - regular development cycles, never an on-demand audit; a change that NAMES a gate is review-checked for its implementation-or-deferral per § 11.4.23 () (c); and on every constitution fetch/pull the § 11.4.23 post-pull sweep + the § 11.4.23 auto-propagation hook run the ledger + block-integrity checks against the post-pull tree, with violations tracked (§ 11.4.23) and cleared - no parallel mechanism is invented.

Reference mechanisms (proven, NOT wired - § 11.4.23 / § 11.4.23). SOL- (gate ledger + ratchet) and SOL- (anchor-block integrity) under docs/research/quality/solutions/ are runnable RED-first POCs with golden fixtures: SOL- 's live read-only run reproduced the -named-gate census exactly. (unimplemented by its slightly-laxer structure rule - the -gate methodological delta stated in-doc); SOL- 's live run confirmed the four mirrors CLEAN post-F and independently re-detected the § 11.4.23 / § 11.4.23 collision. Per § 11.4.23 they carry NO force until wired by separate tracked work items (§ 11.4.23).

Classification: universal (§ 11.4.23) - no project literal; every consuming project (and this governance repo itself) supplies its gate-site list, deferral registry, and lockstep mirror set as DATA per § 11.4.23.

Composes § 11.4.23 / § 11.4.23 / § 11.4.23 /
 § 11.4.23 / § 11.4.23 / § 11.4.23 /
 § 11.4.23 / § 11.4.23 / § 11.4.23 /
 § 11.4.23 / § 11.4.23 / § 11.4.23 /

§ . . () / § . . /
 § . . / § . . / § . . (E).
 Propagation gate CM-COVENANT-114-227-PROPAGATION (literal
 11.4.227 – itself REQUIRED to be block-start-shaped per
 clause (B)) + recommended gates CM-GATE-LEDGER-RATCHET
 (every named gate implemented-or-registered-deferral; the
 unimplemented count monotone-decreasing; vanished names
 cite removals) and CM-ANCHOR-BLOCK-INTEGRITY (exactly-once
 per anchor per file + lockstep content-hash equality +
 carrier exclusion + BLIND-on-zero + collision detection) +
 paired § . mutations (name a gate with no
 implementation and no deferral row → the ledger gate
 FAILs; delete a gate name without a removal citation →
 FAILs; duplicate an anchor block in one mirror, let two
 lockstep copies diverge, or re-mint an existing anchor
 number → the integrity gate FAILs; strip the literal → the
 propagation gate FAILs; gate-code = separate work item,
 NOT claimed shipped § . .).

Canonical authority: constitution submodule Constitution.md
 § . . .

Non-compliance is a release blocker regardless of context.
 No escape hatch – no --name-gate-without-seam , --silent-gate-
 debt-OK , --delete-name-to-lower-count , --presence-gte-1-
 suffices , --duplicate-anchor-OK , --mirror-drift-OK , --reuse-
 anchor-number flag exists.

\$. . – Cross-agent extension lifecycle mandate: per-platform compatibility declaration, source tracking, per-extension documentation with compatibility matrix, and auto-wiring of constitution extensions into all detected agent platforms on constitution pull (research-derived from Extensions Catalog forensics, – –)

Forensic anchor (FACT, – – – measured from the live Extensions Catalog at docs/extensions/ EXTENSIONS_CATALOG.md). The project's extension ecosystem – skills, MCP servers, ACPs, LSPs, plugins, and supporting infrastructure (skills, MCP, provider aliases, CCR providers, daemons, plus a separate HelixSkills repo) – exhibited four structural gaps: (G) of active Claude Code project skills, only Claude Code and OpenCode platforms are explicitly supported – HelixCode, Gemini CLI, and Qwen Code have zero recorded loading paths despite being governed by this constitution's own lockstep carriers (\$. . – CLAUDE.md, AGENTS.md, QWEN.md, GEMINI.md); (G) of active skills are symlinks into the constitution submodule (correct pattern) while are local-only directories with NO traceability record to which original source/master each belongs to – a git-archive snapshot of a skill without its origin is a fork with no master; (G) the Extensions Catalog is a single snapshot document with NO per-extension per-platform compatibility matrix and NO mechanism to re-detect drift when extensions change; (G) the \$. . post_update_hook.sh detects changed governance MDs + scripts/hooks/skills/mcp/plugins and registers them, but it does NOT auto-wire constitution skills into all detected agent platforms – constitution skills (media-validator , multitrack , scheduled-work-queue , session-sync) remain un-wired on both Claude Code (not symlinked into .claude/skills/) and

OpenCode (no loading path to `.opencode/skills/`), silently absent while their governance files claim they exist – the canonical "docs say present, runtime says absent" gap § 1.1.1 forbids.

The mandate (ALL hold – four sub-clauses, one lifecycle):

(A) Cross-agent extension compatibility. Every extension governed by this constitution – Skill, MCP server, ACP bridge, LSP, plugin, provider alias, embedding, daemon, or any future extension class – MUST:

- **Ship a per-platform compatibility declaration** listing every governed CLI agent platform at minimum: Claude Code, OpenCode, Gemini CLI, Qwen Code. The declaration states, per platform, one of the closed vocabulary: `SUPPORTED` (verified working with a recorded per-platform liveness probe), `UNTESTED` (present but not yet verified – honest gap, tracked per § 1.1.1), `UNSUPPORTED` (platform genuinely cannot load this extension class – per § 1.1.1 structural-impossibility, with cited evidence), or `CONFIG-ONLY` (declared in config, never verified as wired per § 1.1.1).
- **NEVER cause loops, hangs, crashes, or broken behavior** in any platform that attempts to load it. An extension that causes a platform malfunction is a § 1.1.1 (A) violation of PASS-bluff severity – it claims compatibility it does not provide, producing the exact "green suite but broken-for-user" failure mode § 1.1.1 forbids.
- **Cross-platform liveness testing is mandatory** – the § 1.1.1 closed test-type set applies: the extension MUST ship a per-platform liveness test (does the agent load it + produce correct output?) that runs on every declared-SUPPORTED platform at minimum. A platform where testing is genuinely infeasible (no

runtime available on this host) is an honest
\$. . . SKIP-with-reason recorded in the
compatibility declaration, NEVER a faked PASS.

(B) Extension source tracking. Every extension MUST be traceable to its exact original source:

- **Canonical source record** – each extension carries machine-readable provenance: originating repository URL, commit SHA (or tag+delta description when `UNTRACED`), canonical file path within that origin. This record lives in an `EXTENSION_SOURCE.yaml` in the extension's directory or in a consumer-maintained extensions registry. An extension with no provenance record is `UNTRACED` and a tracked \$. . . item – never silently trusted as correct.
- **Symlinks preferred over copies** – a copy is a fork with no master; a symlink (or a git submodule pointer, or a reflink share) preserves the canonical source relationship and makes divergence mechanically detectable. Where a copy is genuinely required (platform cannot traverse symlinks, an out-of-repo directory), the SOURCE record MUST state the canonical origin + why a copy was necessary (\$. . . honest rationale). Blind copies without provenance are \$. . . – class dead-code-in-waiting.
- **Per-project extensions catalog required** – every governed project MUST maintain an extensions catalog (`docs/extensions/EXTENSIONS_CATALOG.md` or equivalent declared per \$. . .) enumerating every active extension, its source, its per-platform compatibility status, and its last-verified timestamp. The catalog is \$. . . roster/corpus-backed (sha fingerprint of the sorted extension keyset, NOT mtime), re-syncs out-of-the-box on any extension change. The catalog is the SINGLE SOURCE OF TRUTH for "what extensions exist" – ad-

hoc per-agent skill directories without catalog entries are `UNCATALOGUED` and a tracked `$ . .` item.

(C) **Mandatory extension documentation with per-platform compatibility matrix.** The extensions catalog MUST:

- Carry a **per-platform compatibility matrix** (rows = extensions, columns = governed platforms, cells = compatibility status from the closed vocabulary in clause (A)()) – readable at a glance. A cell with no recorded status is **UNTESTED** by default, never silently assumed compatible.
- Be maintained per § . . (four-format export: **.md** , **.html** , **.pdf** , **.docx**), updated on every extension add/change/removal, with § . . revision header.
- Include **liveliness testing evidence per platform** – per extension × platform, a per-platform probe run result (load-ok / load-failed-for-reason / skip-with-reason) with a timestamp and evidence path per § . . **feature_class=extension_loading** . A platform declared **SUPPORTED** with no recorded probe is **UNTESTED** overridden – a classification mismatch and a § . . violation.
- Auto-sync via the § . . docs-chain engine or equivalent – a stale catalog where extension drift has occurred (catalog says **SUPPORTED**, extension was removed or renamed) is a § . . **PASS-bluff** at the extension-management layer.

(D) Auto-wiring of constitution extensions into all detected agent platforms on constitution pull

(`$` `.` `.` `,` `STRENGTHENS` `$` `.` `.` `)`. The `$` `.` `.` `post_update_hook.sh` seam is EXTENDED — it

```
$ . ./post_update_hook.sh seam is EXTENDED - it
```

MUST, after detecting changed extension artifacts in the constitution submodule:

- **Detect all governed agent platforms** on the consuming host (Claude Code via `.claude/`, OpenCode via `.opencode/` or `opencode.json`, Gemini CLI via `~/.gemini/` or project `GEMINI.md`, Qwen Code via `.qwen/` or project `QWEN.md` – the platform-specific detection logic is consumer DATA per § 11.4.35, never a hardcoded list). A platform whose detection mechanism is absent is an honest `DETECTION_UNSUPPORTED` gap recorded in the catalog, NOT a reason to skip the remaining platforms.
- **Wire constitution skills into each detected platform** – for each detected platform, for each constitution skill not already present:
 - Claude Code: symlink from `.claude/skills/<skill-name>/` → `constitution/skills/<skill-name>/`
 - OpenCode: symlink from `.opencode/skills/<skill-name>/` → `constitution/skills/<skill-name>/`
 - Gemini CLI, Qwen Code: per-platform equivalent wiring mechanism (consumer-supplied DATA, § 11.4.164)
- **Register new/changed MCP servers** into each platform's MCP configuration (`.mcp.json`, `opencode.json`, equivalents) – per the existing § 11.4.164 MCP-registration seam, applied to every detected platform.
- **Self-validate** – after wiring, probe each platform that supports a `--list-skills / --list-extensions` equivalent or fall back to a filesystem check (skill/MCP directory present + parseable). A failed wiring is logged as a § 11.4.164 post-pull finding – never silently absorbed.
- **The hook is idempotent** – re-running it does not duplicate symlinks, corrupt configs, or lose existing non-constitution extensions. Every operation is dry-run-capable (report what WOULD change without changing it).

. **Honest boundary** (§ 11.4.154) – auto-wiring makes constitution extensions REACHABLE to each platform; it does NOT guarantee each platform's runtime ACCEPTS them (a platform may have a different skill format, an incompatible MCP protocol, or a missing runtime – all honest UNTESTED / UNSUPPORTED gaps recorded in the catalog, never faked pass), and it does NOT make the extensions correct (§ 11.4.154 + the extension's own tests). The wiring gap this clause closes is the REACHABILITY gap – a skill that exists in the constitution but is invisible to a governed platform has the gap § 11.4.157 forbids (source-present, runtime-absent).

Composes with § 11.4.154 (post_update_hook – the seam this anchor extends) / § 11.4.154 (GEMINI.md lockstep – the platforms whose extensions this governs) / § 11.4.154 (universal export – the catalog's four formats) / § 11.4.154 (roster/corpus-backed auto-sync – the catalog's fingerprint discipline) / § 11.4.154 (docs-chain – the catalog's sync engine) / § 11.4.154 (sink-side evidence – feature_class=extension_loading) / § 11.4.154 (dead-code – blind copies are dead-code-in-waiting) / § 11.4.154 (SKIP-with-reason – untestable platforms) / § 11.4.154 (structural impossibility – UNSUPPORTED platforms) / § 11.4.154 (clean artifact – CONFIG-ONLY un-wired class) / § 11.4.154 (post-pull sweep – wiring failures) / § 11.4.154 (runtime signature – a platform load is the extension's runtime signature) / § 11.4.154 (research completion – per-platform wiring gaps are tracked items) / § 11.4.154 (test types – per-platform liveness tests) / § 11.4.154 / § 11.4.154 (code-review – extension changes cross the same quality gauntlet).
 Classification: universal (§ 11.4.154) – no hardware/vendor/project literal; the consuming project supplies its platform detection mechanism, per-platform wiring

commands, extension registry path, and catalog path as
DATA per § 11.4.228.

Propagation gate CM-COVENANT-114-228-PROPAGATION (literal
11.4.228). **Recommended gates:**

- CM-EXTENSION-PER-PLATFORM-COMPAT (every governed extension carries a per-platform compatibility declaration using the closed vocabulary; a declared-SUPPORTED platform with no recorded liveness probe → FAIL)
- CM-EXTENSION-SOURCE-TRACKED (every extension carries provenance with origin repo + commit SHA or explicit UNTRACED rationale tracked per § 11.4.228; a blind copy without provenance → FAIL)
- CM-EXTENSION-CATALOG-DOCUMENTED (the extensions catalog exists, carries a per-platform compatibility matrix, is four-format-exported per § 11.4.164, and its sha fingerprint is fresh per § 11.4.228; a stale catalog with extension drift → FAIL)
- CM-CONSTITUTION-EXTENSIONS-AUTO-WIRED (the § 11.4.228 post_update_hook.sh auto-wires constitution skills into every detected governed platform, idempotently, and self-validates; a constitution skill not wired into a detected SUPPORTED platform → FAIL, with explicit DETECTION_UNSUPPORTED / UNSUPPORTED SKIP-with-reason – the § 11.4.228 () false-positive guard)

All paired § 11.4.228 mutations (gate-code = separate work item, NOT claimed shipped § 11.4.228).

Canonical authority: constitution submodule Constitution.md
§ 11.4.228.

Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-per-platform-compat, --single-platform-OK, --extension-without-provenance, --copies-over-symlinks-OK, --skip-extension-catalog, --catalog-format-

^{1 1} optional , ⁴ --skip-auto-wire , ^{2 2 0} --manual-wiring-suffices , --
untraced-extensions-OK flag exists.

\$. . . – Live in-session task/todo tracker MUST always be up to date and fully in sync with the real work state – never stale, never showing completed work as pending nor pending as done (User mandate, – –)

Verbatim operator mandate (– –): "make sure it is ALWAYS up to date fully in sync! CRITICAL: This MUST BE one of mandatory rules / constraints! Add it with all relevant details into the constitution Submodule and relevant related files."

The agent's LIVE IN-SESSION TASK/TODO TRACKER – the harness-provided task list / `TodoWrite -class todo tool` / any equivalent moment-to-moment execution-state surface the runtime exposes DURING a session – MUST ALWAYS be kept up to date and FULLY IN SYNC with the real live work state: never stale, never showing completed work as pending, never showing pending / blocked / abandoned work as done or in-progress. It MUST be updated the MOMENT any tracked task's state changes – the instant a task is started (→ in-progress), completed (→ done), blocked, abandoned, split, re-scoped, or newly discovered – not batched to a later checkpoint, not reconstructed from memory at the end of the work, not left reflecting a superseded plan. A live task tracker that lags the real work is itself a \$. . . PASS-bluff at the execution-transparency layer: it reports a work state the agent is not actually in, and an operator (or a resuming agent / conductor reading it per \$. . .) then makes decisions on a false picture of what is done, what is in flight, and what is owed.

truth, the § . . . / § . . . /
§ . . . / § . . . artefacts are
materialised FROM), never substitutes for any of them, and
its always-in-sync discipline is the in-session analogue
of § . . . 's always-current mandate applied to the
harness task list.

(D) **HONEST BOUNDARY** (§ . . .). Where the runtime
exposes NO such live task-tracking surface, this anchor is
honestly N/A for that runtime (§ . . . SKIP-with-
reason – the discipline BINDS the moment such a surface
exists, per the § . . . latent-binding pattern); it
mandates that WHATEVER live tracker the runtime does
expose is kept in-sync, not that a particular tool must
exist. "In sync" means the tracker reflects the real work
state as the agent knows it at that moment – the captured
fact of what has started / completed / changed – not
omniscience about work the agent has not yet observed.

Classification: universal (§ . . .) – a platform-
neutral execution-transparency discipline reusable by ANY
agent / runtime that exposes a live in-session task/todo
tracker; the consuming runtime supplies its concrete
tracker surface (harness task list / `TodoWrite` /
equivalent) per § Composes § . . .
(always-current CONTINUATION – § . . . is its
live-in-session-tracker analogue) / § . . .
(standing session-resumption file) / § . . .
(progress-update cadence – the operator-report sibling) /
§ . . . (F) (write-seam doc/DB sync) /
§ . . . (default autonomous-loop – the loop's
live task state IS the tracker) / § . . . (real-
time conductor framework sync – a resuming conductor
reads the live tracker) / § . . . /
§ . . . (workable-item / request-history
materialisation) / § . . . (no-guessing – the tracker
is captured fact, never aspirational) / § . . .

(falsely-pending is a FAIL-bluff). Propagation gate `CM-COVENANT-114-229-PROPAGATION` (literal `11.4.229` across the consumer fleet) + recommended gate `CM-LIVE-TASK-TRACKER-IN-SYNC` (where the runtime exposes a live in-session task/todo tracker, its state matches the real work state at every state-change – no completed-task-marked-pending, no pending-task-marked-done, no ghost/stale entry, no superseded-plan drift; honest § `11.4.229` SKIP-with-reason where the surface is genuinely absent – the § `11.4.229` () false-positive guard) + paired § `11.4.229` meta-test mutation (mark a completed task as still-pending, OR a pending/abandoned task as done, OR leave the tracker reflecting a superseded plan → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item, NOT claimed shipped § `11.4.229`).

Canonical authority: constitution submodule `Constitution.md`
§ `11.4.229` .

Non-compliance is a release blocker regardless of context. No escape hatch – no `--tracker-may-lag` , `--batch-todo-updates` , `--stale-task-list-OK` , `--reconstruct-todos-at-end` , `--skip-task-tracker-sync` , `--pending-may-be-done` , `--completed-may-stay-pending` flag exists.

§ `11.4.229` – **Parallelized-pipeline methodology:** build validate overlap + affected-stages-only re-run + parallel multi-device testing/recording + fan-out-subagents-always + incremental delivery granularity with a declared QA-build cadence (research-derived, – – ; clause (D) added – –)

Forensic motivation (genericised): the sequential pipeline shape – build FULLY completes → THEN validation begins → on ANY discovery, restart the WHOLE cycle from scratch → all devices tested one-at-a-time – serialises the

operator's wall-clock onto build latency + one-device-at-a-time evidence gathering, so a multi-hour containerised build followed by serial per-device validation and a from-scratch restart-on-discovery makes time-to-a-manual-QA-capable-release a multiple of what the genuinely-parallel work requires. § 1.1.1.1's PWU pipeline already overlaps Stage 1 of round N+1 with Stages 1-4 of round N; § 1.1.1.2 already backgrounds long tests; § 1.1.1.3 already mandates ≥ 4 auto-backfilled parallel streams; § 1.1.1.4 already partitions shared hardware by single-owner; § 1.1.1.5 / § 1.1.1.6 already mandate always-on recording + availability-following. § 1.1.1.7 BINDS those into one operating discipline and adds the four invariants they did not each state on their own: (A) the build validate CONCURRENCY seam + affected-stages-only re-run-on-discovery, (B) all-devices-concurrent + worker-pool recording/processing, (C) fan-out-subagents as the standing default, (D) incremental delivery granularity + a declared QA-build cadence (added 10s - 15s). It is the parallelization capstone the way § 1.1.1.1 is the continuous-parallel-stream capstone – a strengthening bind, never a weakening.

(A) PIPELINE-OVERLAP PARALLELIZATION. Validation/verification work MUST run CONCURRENTLY WITH the build wherever a validation stage does NOT depend on the build's own not-yet-produced artifact – never deferred until the build finishes when it could have run alongside it. Deploy the moment the build is done (build-completion is the deploy trigger, per § 1.1.1.8 no-commit-while-build-writes-tracked-artifacts + § 1.1.1.9 artifact-must-exist ordering – the deploy waits for the artifact, the deploy-independent validation does not). Concretely, WHILE a build runs, the following run in parallel per the § 1.1.1.10 SAFE-during-build catalogue: pre-build source-layer gates + meta-test

mutation sweeps + coverage/clash checks (\$. .),
 the \$. . /\$. . /\$. .
 independent code-review on the batch's SOURCE, doc/DB sync
 (\$. .), and on-device validation of the PRIOR
 build (device-independent of the in-flight build,
 \$. .). On any NON-huge-blocker discovery, the
 cycle MUST: () bump the version code (a new, distinctly-
 identifiable candidate artifact – never re-use the failing
 artifact's id, \$. . stable-name +
 \$. . target-identity); () fix at the proven
 root cause (\$. . systematic-debugging first,
 \$. . RED-on-the-broken-artifact); () RE-RUN
 ONLY THE AFFECTED STAGES, PARALLELIZED – the stages whose
 inputs the fix changed (its own targeted validation FIRST
 per \$. . , then the affected regression set),
 NOT a serial restart-from-scratch of the entire cycle. The
 explicit efficiency target is a - × **reduction in time-
 to-a-manual-QA-capable-release** versus the fully-sequential
 shape – a measured target, tracked as build-and-validate
 wall-clock deltas (\$. . build-resource stats),
 never an assumed one (\$. .). **HONEST BOUNDARY**
 (\$. . / \$. . / \$. .): the
 CURRENT build's ARTIFACT-layer (\$. . layer
) and RUNTIME-ON-CLEAN-TARGET-layer (layer)
 validation CANNOT precede its own artifact – a gate whose
 precondition is the build's output belongs at a post-build
 seam, never overlapped onto authorship (\$. .
 seam-placement; \$. . artifact_not_yet_built).
 "Concurrent with the build" therefore means the deploy-
 INDEPENDENT stages overlap; the CURRENT-artifact runtime
 validation begins the moment that artifact exists, in
 parallel across devices per (B). **"Affected-stages-only re-
 run" is bounded to NON-huge-blocker discovery:** a HUGE
 BLOCKER (release-blocking-severity defect / regression
 invalidating the cycle / blast radius reaching the batch's
 other fixes) still triggers the \$. . STOP →

fix-all → full-RESTART protocol (restart, never resume) –
§ . . . (A) does NOT weaken § . . . ,
§ . . . full-suite-retest-before-tag,
§ . . . validate-the-fix-first, § . . .
risk-ordered validation, or § . . . manual-QA-
final; it narrows the from-scratch restart to exactly the
cases those anchors require it, and parallelises
everything else.

(B) PARALLEL MULTI-DEVICE TESTING + RECORDING/PROCESSING.

Every available, reachable test device MUST be exercised
CONCURRENTLY – one validation stream per device, never
one-device-at-a-time when the devices are independent –
each stream the SINGLE EXCLUSIVE OWNER of that device's
exclusive resources per § . . . (the owner
drives; any concurrent stream targeting the same device is
read-only), so per-device evidence is never cross-
contaminated (§ . . . evidence-integrity). Every
device is ALWAYS-RECORDED per § . . . (non-
invasive, side-effect-free, deterministic layout) and
AVAILABILITY-FOLLOWED per § . . . (a dropped
device logs an honest offline event + reconnect-and-
resume, never a silent corpus hole, never a faked
continuity). Recording CAPTURE + ANALYSIS/PROCESSING MUST
be fanned out through a BOUNDED WORKER-POOL (a Go worker-
pool or platform equivalent) so N devices' recordings are
captured + analysed in parallel and per-device evidence
never serialises onto a single processor – the pool
bounded by § . . . memory + § . . . thread headroom.
Per-device PASS cites its own captured-evidence artefact
(§ . . . /§ . . . /§ . . .) on a CLEAN/
fresh deployment (§ . . . runtime-signature,
§ . . . verify-deployed-artifact-identity-per-
device); the highest-risk / most-reopened cases
(§ . . . /§ . . .) get the deepest
scrutiny FIRST on every device. **HONEST BOUNDARY**
(§ . . .): genuinely-shared exclusive resources (one

HDMI/audio sink, one media-playback path) serialise per
\$. . . single-owner-over-time even across
devices; "all devices concurrent" means independent
devices, never two owners of one exclusive resource.

(C) FAN-OUT-SUBAGENTS-ALWAYS. Subagent fan-out
(\$. . . /\$. . . subagent-driven;
\$. . . multi-track ruler; \$. . .
auto-backfill) is the STANDING DEFAULT for EVERY
genuinely-parallelizable effort – batch investigation,
test-authoring, gate-authoring, per-device validation (B),
per-PWU development (\$. . .), doc/research work –
not a per-request opt-in and not reserved for "big" tasks.
A parallelizable effort executed monolithically in the
foreground when a fan-out path exists is a
\$. . . (C) violation (the \$. . . /
\$. . . /\$. . . zero-idle discipline
applied to HOW work is dispatched, not only WHETHER it
runs). **HONEST BOUNDARY** (\$. . . / \$. . .):
"always fan out" is BOUNDED by () genuine
parallelizability (contending/redundant/serially-dependent
subagents are NOT fan-out – they are waste, and spawning
them is NOT maximization per \$. . . maximum-
USEFUL-throughput); () \$. . . % memory ceiling;
() \$. . . RLIMIT_NPROC thread headroom; ()
\$. . . ≤ concurrent-agent cap; ()
\$. . . single-resource-owner (two subagents may
not both drive one exclusive device); () critical-state
sequencing (commits/pushes/tags stay serial per
\$. . . /\$. . .). The bar is maximum USEFUL
parallel throughput, never agent count; fan-out that
exceeds a host-safety ceiling or contends a single
resource is a violation, not compliance.

**(D) INCREMENTAL DELIVERY GRANULARITY + DECLARED QA-BUILD
CADENCE** – a long-running spec/feature ships every already-
proven-correct increment as it clears its own bar, never

withholding delivery until the **WHOLE** spec/feature is done (operator mandate, - - , **CRITICAL** ; research-derived; added - -). Verbatim operator mandate (- -): "All work we do, especially SpecKit specs, gets incorporated right away. Fully deterministic validation/verification tests of all supported test types relying on machine-produced evidence get incorporated too, aiming to cover everything critical. All intensive testing (gates, checks, fixing, polishing) keeps running in **BACKGROUND** continuously until fully done. **IN PARALLEL** we keep building specs one by one, phase by phase, story by story, task by task. We get fast delivery results and **AS SOON AS POSSIBLE** builds that can be handed to QA teams for manual testing **RIGHT AWAY**, while in parallel all testing/fixing/polishing/improvement continues exactly as it would anyway. We do **NOT** wait for all of that to finish first – that can take days [...]. QA reports discovered issues; background agents keep working until everything is absolutely flawless. As background work progresses and fixes/improvements land in the main codebase, we keep producing **REGULAR** builds periodically delivered to QA teams for ongoing full manual testing. It is critical that background agents keep **ALL** workable items, documentation, and other materials in sync so nothing ever goes out of date." Forensic anchor (genericised, **FACT**): in one consuming project a multi-phase spec had **RED** tests drafted for every phase and many source slices independently reviewed to a source-correctness **GO**, yet nothing reached QA for days, because delivery was implicitly held at the **WHOLE-SPEC** boundary while...**ONE** slice's iterate-to-**GO** review loop ran past forty rounds – the finished, proven work of every other slice sat undelivered. The **GAP** this clause closes, stated precisely so it is not restated elsewhere (§ . .): clause (A) overlaps validation with the build of **ONE** candidate; § . . (A) makes a batch

build-eligible on source-correctness alone;

§ . . . (D) makes that eligibility PER SLICE inside a batched review – but none of them obliges an ELIGIBLE increment to actually be DELIVERED before the rest of its spec completes, none sets a delivery CADENCE, none says how an incomplete spec stays safe to ship, and none binds the hand-off to tracker truth. Build-eligibility that is never exercised is the same withheld delivery under a different name. ALL of the following hold:

(D.) **THE DELIVERY UNIT IS THE INCREMENT, NOT THE SPEC.** A spec / feature / epic decomposed into phases / stories / tasks (vertical, independently testable slices per the § . . . SLICING extension – Small + Testable) MUST integrate each increment into its INTEGRATION BRANCH – the trunk for main -targeted work, or the feat/ / product/ / flavor/ branch (§ . . .) the increment actually lives on for branch-scoped work – and make it part of the next QA deliverable (D.) cut from that branch, the moment THAT increment holds its OWN source-correctness GO per § . . . (A)/(D). When an increment exists on MULTIPLE coexisting integration branches – for example both a product/* AND a flavor/* branch per § . . . (A) – EACH branch is independently obligated under (D.)'s own per-branch cadence: the increment's delivery obligation is evaluated separately per branch it lives on, never as a single ambiguous "the branch". A QA build MAY be cut from a feat/ / product/ / flavor/ branch. For a feat/ branch that QA build IS the live-manual-testing vehicle

§ . . . (B) requires BEFORE the branch may merge to trunk, so the feature-branch QA build is the PRECURSOR to trunk-landing, never a consequence of it; a product/ / flavor/ line, which never merges to trunk the standard way (§ . . . (C)), receives its cadence builds from its own branch. Clause (D) relaxes neither § . . . (B) (no feat/ merge to trunk before live manual QA) nor

§ 4.1.1 (C), and a later merge of the branch to trunk still follows § 4.1.1 / § 4.1.2. Withholding an already-build-eligible increment solely because OTHER increments of the same spec are unfinished, or because its own or another increment's test-instrumentation / process-doc loops (§ 4.1.1 (D)) are still open, is a § 4.1.1 (D) violation. An increment WITHOUT its own source-correctness GO is not deliverable – incremental delivery is never a licence to ship unreviewed or unproven work.

(D.) **INCOMPLETE-SPEC, SAFE DELIVERABLE.** Shipping an increment while its spec is incomplete MUST leave the deliverable correct on every path a user or QA tester can reach: EITHER the increment is a complete, independently valuable vertical slice, OR the unfinished surface is LATENT – unreachable by construction, behind a release toggle (P. Hodgson, "Feature Toggles", martinowler.com, : "Release Toggles allow incomplete and un-tested codepaths to be shipped to production as latent code which may never be turned on"), or behind branch-by-abstraction (trunkbaseddevelopment.com: a technique "to effect a 'longer to complete' change in the trunk"). A toggle is a PERMITTED mechanism, never a mandated one (consistent with § 4.1.1's own statement that it mandates none), and it is inventory with a carrying cost (Hodgson: "Savvy teams view their Feature Toggles as inventory which comes with a carrying cost"): every release toggle carries a tracked § 4.1.1 removal item and a consumer-declared expiry (§ 4.1.1), and a toggle past its expiry is a § 4.1.1 finding. A toggle that needs to outlive its expiry – for example because its latent path's (D.) hardening loop legitimately runs many rounds – is extended ONLY through a § 4.1.1 waiver (rostered non-producer authoriser, a new unelapsed expiry, a named tracked item), never by a silent expiry edit, by enabling the still-unproven path, or by deleting finished

work. A latent path is NOT an advertised capability (§ . . .) until it is enabled and proven.

(D.) A DECLARED QA-BUILD CADENCE, NOT A ONE-SHOT. Every project declares as DATA (§ . . .) its QA-build trigger: **per-increment** (each build-eligible increment landing on its integration branch, D.), **periodic** (a fixed interval – the release-train pattern; SAFe's "Develop on Cadence. Release on Demand" separates the development rhythm from the release decision), or **whichever-first** ; a project that declares nothing gets **per-increment** , the safe default (§ . . .). A **periodic** or **whichever-first** interval MUST be recorded as DATA together with its written justification (§. . .) and stated in every QA hand-off record (D.), so an interval long enough to defeat this clause is visible rather than silent; this clause sets no numeric maximum – any such cap is an operator decision (§ . . .). Whenever an integration branch (D.) holds build-eligible content not yet in a QA build cut from that branch and the declared trigger has fired, a QA deliverable MUST be cut from that branch and handed to QA for ongoing manual testing (§ . . .). Builds COALESCE: a QA build already in flight covers every increment that had landed on its branch before that build started, and increments landing WHILE it is in flight roll into the NEXT build rather than each demanding a fresh one – the obligation is that eligible content never waits past the next build, not one build per increment. A build that FAILS, or is REFUSED at § . . . , covers NOTHING – its increments roll into the next build. "Completed" in this clause and in its two amended gate contracts (D's recommended gates, below) means HANDED TO QA: a build provides coverage only once it has genuinely been handed to QA for manual testing, never merely once the build process itself exits successfully. Each such deliverable is a § . . . (B) cycle boundary with its own monotonic version increment

(§ . .), passing § . . readiness on its own candidate fingerprint, and built once then promoted (§ . .). The § . . "mandated validation" for a cadence build is AT MINIMUM the source-correctness G0 plus the pre/post-build gates of § . . (A), PLUS every operator-declared blocker bound to the QA-deploy seam per § . . and any § . . critical-invariant scenarios the increment touches – NOT the open background-hardening backlog, which by § . . (A) never gates a build. "Deliverable" is project-defined DATA (§ . .): a binary, a container, a device image – or, for a dev-tooling / governance feature, the feature's own demonstrable, testable state at a fingerprinted commit; no artifact form is implied by this clause.

(D.) CONTINUOUS BACKGROUND HARDENING THAT NEVER BLOCKS THE NEXT INCREMENT. Test-instrumentation hardening, gate / check / mutation authoring, polishing, and the fixes for QA-reported findings on ALREADY-delivered increments run as standing background streams (§ . . / § . . / clause (C)), each a tracked item (§ . .) driven to its own zero-finding G0 (§ . . – not relaxed), landing on its integration branch (D.) as it reaches G0 and riding the NEXT cadence build cut from that branch. They never gate a LATER increment's build-eligibility, with two exceptions that follow from existing anchors rather than relax them: an open **source-defect** finding (§ . . (D)) in code the later increment touches or depends on is a **source-defect** of that later increment too; and a huge blocker still triggers § . . STOP → fix-all → full restart.

(D.) NEVER-STALE TRACKING IS A PRECONDITION OF EVERY HAND-OFF, NOT AN AFTERTHOUGHT. With many increments and many background streams in flight at once, the existing

sync machinery – § . . . (F) write-seam doc/DB
 sync, § . . . live tracker, § . . . /
 § . . . / § . . . intake and request
 history, § . . . item integrity, § . . .
 completion – applies at EVERY increment landing and EVERY
 cadence build (cited, not restated). Each QA hand-off
 records the deliverable's fingerprint and the exact set of
 tracked items it contains; a hand-off whose tracker or
 documents disagree with the deliverable (an item marked
 done whose fix is not in that build, or a fix in the build
 that no item records) is a § . . . PASS-bluff at the
 delivery layer.

**(D.) QA FINDINGS ON EARLY INCREMENTS ARE STILL COVERAGE
 ESCAPES.** Earlier, more frequent QA builds surface defects
 earlier; they do NOT make manual QA the discovery
 mechanism. Every defect manual QA finds on an increment
 build triggers the § . . . (C) coverage-escape
 audit and a new automated check, exactly as for a release
 candidate.

(D) honest boundary, research grounding, and gates. Honest
 boundary (§ . . .): clause (D) changes WHEN and HOW
 OFTEN proven-correct work reaches QA; it does NOT change
 what gates a release candidate. It does not weaken
 § . . . (huge-blocker full restart),
 § . . . (full-suite retest before a release TAG – a
 cadence QA build is not a tag), § . . . (manual-
 QA-final terminal gate for a release), § . . .
 (readiness verdict per candidate), § . . . /
 § . . . (build once and promote; progressive
 delivery), § . . . (every finding of every layer
 still driven to GO), or § . . . (D)'s conservative-
 safe default (an unclassifiable finding is a source-defect
 and does block). It does not prove any increment correct –
 that stays § . . . / § Research
 grounding, stated honestly: DORA's "Working in small

batches" capability (dora.dev) – "break them down into work units that can be completed independently and in short timeframes", with dark launching / feature toggles named as the way to deploy incomplete features hidden from users; Hodgson's toggle taxonomy and toggle-as-inventory discipline (quoted in D.); branch-by-abstraction (quoted in D.); SAlFe's "Develop on Cadence. Release on Demand" (D.); and GitHub Next's "Continuous AI" / Agentic Workflows technical preview, which its principal researcher calls "the agentic evolution of continuous integration" – scheduled or event-triggered agent loops for recurring work such as CI-failure investigation and documentation-drift fixing, cited here only as industry direction for (D.)/(D.), not as an adopted mechanism. No single published methodology was found that combines per-increment QA delivery with continuously running agent hardening streams; that composition is this clause's own synthesis, not a claim of prior art. Classification: universal (§ 11.4.129.) – consumers supply the cadence trigger, interval, toggle mechanism and expiry policy, and the deliverable form as DATA per § . . . Composes § 11.4.135. • § 11.4.137. 11.4.132. / § . . . / § 11.4.238. • § 11.4.210. (F) 219 / § . . . / § 11.4.235. • / § 234. . . . / § . . . / § 11.4.238. • § 11.4.239. 11.4.241. / § . . . / § . . . / § 11.4.244. • § 11.4.245. 11.4.246. / § . . . / § . . . / § 11.4.271. • / § . . . / § . . . / § . . . (A)(B)(D) / § . . . / § . . . / § . . . / § . . . / § . . . / § . . . / § . . . Recommended gates CM-INCREMENT-DELIVERED-NOT-SPEC-WITHHELD (FAIL: an increment holding its source-correctness GO is absent from every completed-and-handed-to-QA or in-flight QA build from its integration branch (D.) after the declared trigger fired while no build from that branch is in flight – coalescing per (D.) so

an increment waiting on an in-flight build's successor is never a FAIL, and a build that FAILED or was REFUSED at § . . . counts as NEITHER completed NOR in-flight so it provides no coverage and the increment rolls to the next build; FAIL, the § . . . () dual guard: an increment WITHOUT that GO appears in a QA build) + CM-QA-BUILD-CADENCE-DECLARED (FAIL: an integration branch (D.) holds build-eligible content that no completed-and-handed-to-QA or in-flight QA build from that branch covers, the declared trigger has fired, and no QA build from that branch is in flight; coalescing per (D.) – a build in flight covers every increment that landed before it started and later increments roll into the NEXT build, so an in-flight build is never a FAIL and the gate never demands one build per increment; a build that FAILED or was REFUSED at § . . . covers NOTHING and does not satisfy this gate; FAIL also: a periodic / whichever-first interval with no recorded justification; an absent declaration resolves to per-increment , never to "no cadence") + CM-RELEASE-TOGGLE-TRACKED-EXPIRY (FAIL: a release toggle with no tracked removal item or no expiry, or past its expiry) + CM-QA-HANDOFF-TRACKER-IN-SYNC (FAIL: a hand-off record whose item set disagrees with the deliverable's contents) + paired § . mutations (withhold an eligible increment until spec completion; ship an increment lacking its GO; delete the cadence declaration and assert a "no cadence" reading; add a toggle with no expiry; mark an item done whose fix is absent from the build → each gate MUST FAIL; golden-FALSE: an increment shipped on the first trigger after its GO, behind a tracked unexpired toggle, with a matching hand-off record → none fires). Gate-code = separate work item, NOT claimed shipped (§ . . . / § . . .); the four gates are registered as tracked deferrals in constitution/scripts/gates/gate_ledger_deferrals.tsv . No escape hatch – no --withhold-until-spec-complete , --hardening-gates-next-increment ,

`--no-qa-cadence` , `--untracked-release-toggle` , `--stale-tracker-at-handoff` , `--manual-qa-as-discovery` flag.

(Process note for future extenders: § . . 's next free clause letter is (E); never reuse (A)–(D).)

Classification: universal (§ . .) – a platform-neutral parallelization discipline reusable by ANY project that builds an artifact, validates it across multiple targets, and records device evidence; the consuming project supplies its concrete device set + serials, build/flash tooling, recording layout, and worker-pool implementation per § . . . Composes
 § . . . / § . . . (wall-clock stats prove the – × target) / § . . . / § . . .
 (the PWU pipeline this binds + extends) / § . . . /
 § . . . / § . . . / § . . . /
 § . . . / § . . . (SAFE-during-build catalogue) / § . . . / § . . . /
 § . . . / § . . . (artifact/runtime layer ordering – the honest boundary) / § . . . /
 § . . . / § . . . (single-resource-owner – B's load-bearing partition) / § . . .
 (seam-placement – A's honest boundary) /
 § . . . / § . . . / § . . . /
 § . . . (always-on recording – B's evidence source) / § . . . (huge-blocker full-restart – the un-weakened carve-out) / § . . . /
 § . . . / § . . . (availability-following – B's drop-handling) / § . . .
 (maximum-USEFUL-parallelism – C's honest boundary) /
 § . . . (manual-QA-final – the un-weakened terminal gate) / § . . . / § . . . /
 § . . . (per-device deployed-artifact verification) / § . . . / § Propagation gate
 CM-COVENANT-114-230-PROPAGATION (literal 11.4.230 across the consumer fleet) + recommended gates CM-PIPELINE-OVERLAP-

VALIDATION (a build that runs with zero deploy-independent validation stages running concurrently when SAFE-catalogue work is queued → FAIL; a NON-huge-blocker discovery that restarts the whole cycle from scratch instead of re-running only the affected stages → FAIL) + CM-PARALLEL-DEVICE-WORKER-POOL (available independent devices validated serially instead of concurrently, OR recording capture/analysis serialised onto one processor when a worker-pool path exists → FAIL) + CM-FAN-OUT-SUBAGENTS-DEFAULT (a genuinely-parallelizable effort run monolithically in the foreground when a fan-out path exists AND all host-safety ceilings have headroom → FAIL) + paired § . meta-test mutations (per § . ; gate-code = separate work item, NOT claimed shipped § . .).

Canonical authority: constitution submodule Constitution.md § . . .

Non-compliance is a release blocker regardless of context. No escape hatch – no --sequential-build-then-validate , --restart-whole-cycle-on-discovery , --serial-device-testing , --serialise-recording-processing , --no-fan-out , --monolithic-parallelizable-work flag exists.

§ . . – Nano-precision model-tier selection: lightest-capable-model-first with dynamic mid-task escalation (User mandate, – – , CRITICAL)

Verbatim operator mandate: every work effort MUST use the LIGHTEST CAPABLE model, escalating tier only as complexity rises – a nano-precision, per-task model-tier choice, NOT "everything on the heaviest model." Forensic motivation (genericised, § . .): dispatching every task – a trivial doc-sync, a one-line status read, a mechanical format fix – on the heaviest available model wastes capacity + token budget + wall-clock that a lighter tier

would spend identically well, while the OPPOSITE failure (a hard multi-factor root-cause investigation dispatched on a tier too light to reason through it) produces exactly the shallow, guessed, or wrong-conclusion work

\$. . / \$. . / \$. .

forbid. Neither "always heaviest" nor "always lightest" is compliant; the correct discipline is a PER-TASK, EVIDENCE-BASED tier choice that tracks the task's own complexity, re-evaluated as that complexity becomes known – the nano-precision selection this anchor mandates.

(A) **LIGHTEST-CAPABLE-FIRST + ESCALATE-ON-COMPLEXITY.** Every dispatched work effort – main-stream or subagent

(\$. . / \$. .), single-track or multitrack

(\$. . / \$. .) – is assigned the

LIGHTEST model tier from the closed, capability-ordered ladder sonnet < opus genuinely sufficient to complete it

correctly, with sonnet the **DEFAULT working tier** and opus

for every review and other heavy work (operator mandate

- - , CRITICAL , verbatim: "Main model to

use is Sonet, reviews and other heavy stuff have to be done by Opus instead of Fable!!! Fix this permanently on all levels!"). The review/merge-conflict-resolution

substrate is a SEPARATE, already-pinned Opus-xhigh tier per clause (E), not an extra rung of this ladder; fable

is not a rung and is not a permitted dispatch tier; haiku

is excluded per clause (D.). Tier selection is a

DECIDABLE heuristic, not a guess (\$. .) – evaluated

against, at minimum: (i) **task scope** (single mechanical edit / focused single-file change / multi-file-or-multi-component change); (ii) **reasoning depth** (a lookup-and-apply action vs. a multi-factor root-cause investigation

per \$. . / \$. .); (iii) **blast radius** (a change touching one isolated file vs. one crossing shared state, hardware safety \$. . , or

governance/gate seams \$. .); (iv) **novelty** (a

previously-solved, pattern-matched task vs. a genuinely

new problem with no established precedent). Closed-set tier bindings (illustrative, non-exhaustive – the consuming project supplies its own precise mapping per § 4.1.1.1): **sonnet** (the default) – trivial/mechanical AND moderate work (simple greps/searches, doc-sync regeneration, status/field reads, format/lint fixes, boilerplate scaffolding, a focused single-component debug, a bounded single-file implementation, straightforward research or audit reads); **opus** – complex and other heavy work, including every review (architecture/design decisions, multi-factor reasoning per § 4.1.1.2 (1)), cross-feature blast-radius analysis per § 4.1.1.3. Pass 1 / 2, hard root-cause investigation per § 4.1.1.4. A task that resists a first-pass hypothesis). A task whose complexity is genuinely UNDECIDABLE from available evidence starts at the NEXT tier UP from the best available estimate (the safe-reversible default per § 4.1.1.1) rather than guessed downward.

(B) DYNAMIC MID-FLIGHT SWITCHING. Tier assignment is NOT a one-time, start-of-task decision – it is RE-ASSESSED continuously as the task's real complexity becomes known during execution. A task DISPATCHED as simple that PROVES harder mid-flight (an unexpected multi-factor interaction surfaces, a root cause resists the first two hypotheses, the blast radius turns out to cross a governance/hardware-safety seam) MUST ESCALATE to the next-heavier tier for its remainder – continuing on an under-powered tier past the point complexity is KNOWN is a clause-(D) false economy, not compliance. Conversely, a task DISPATCHED as heavy that DECOMPOSES into genuinely light, bounded, low-blast-radius subtasks (mechanical file edits following an already-designed pattern, doc regeneration, per-item status reads) dispatches THOSE subtasks on a LIGHTER tier via subagent fan-out (§ 4.1.1.1 / § 4.1.1.2 / § 4.1.1.3) rather than running every subtask on the parent's heavier tier by default – capability-matching

applies at EVERY dispatch granularity, not only the top-level task.

(C) ALWAYS-ON COMPOSITION, NEVER A PER-REQUEST OPT-IN.

Nano-precision tier selection is a STANDING DEFAULT mechanism, engaged automatically for EVERY work effort from the first prompt of a session onward (mirrors § . . . default-autonomous-loop / § . . . default-mechanisms-always-on) – applied TOGETHER WITH, never as a substitute for: § . . . (heavy token-optimization – a lighter model dispatched with a bloated context is not the efficiency this anchor targets); § . . . (maximum-USEFUL-parallelism – tier-matching each parallel stream, not merely running more streams on one fixed tier); § . . . / § . . . (the multi-track ruler orchestration + continuous auto-backfill – each track's worker is tier-assigned per this anchor, not defaulted to one blanket tier project-wide); and § . . . / § . . . (native-alias-first priority + always-on default mechanisms – this anchor governs WHICH TIER within whichever alias/account is selected, composing with, never replacing, § . . . 's ALIAS-layer priority). None of these five composed anchors is weakened, duplicated, or re-specified by § . . . – this anchor adds the TIER axis to their existing ALIAS/PARALLELISM/TOKEN axes.

(D) HONEST BOUNDARY (§ . . .) – "LIGHTEST CAPABLE"

MEANS CAPABLE, NOT CHEAPEST-REGARDLESS. "Lightest capable" is a CONJUNCTION, not a discount: a tier choice that fails to actually complete the task correctly – producing a shallow investigation, a guessed root cause, an under-verified fix, or ANY form of bluff per the § . . . anti-bluff covenant – is NOT "lightest capable," it is UNDER-POWERED, and under-powering a task that then fails, bluffs, or requires costly rework is a FALSE ECONOMY and a § . . . violation of equal severity to any other PASS-

bluff (the token/time saved at dispatch is spent many times over in re-investigation, re-review, and the risk of a shipped defect). Capability determinations are made from the clause-(A) decidable heuristic + observed task outcomes – EVIDENCE-BASED, never guessed (§ 2.1.2.1.1): "this always works fine on sonnet" without a captured basis is exactly the guess § 2.1.2.1.1 forbids, and a tier's genuine capability for a task class is established by the same anti-bluff discipline (§ 2.1.2.1.1 / § 2.1.2.1.1.1) that governs every other claim in this constitution. The ladder binds the DISPATCHER's tier ASSIGNMENT: a dispatcher MUST NOT assign any tier outside sonnet < opus – never fable (– mandate), and never haiku where clause (D. 1.1) applies. A serving model the platform selects for a turn WITHOUT a dispatcher assignment is recorded honestly in the § 2.1.2.1.1 label (whatever id the platform reports, never re-labelled); it is neither a dispatch-tier violation by the dispatcher nor a licence to assign that model, and review/merge-conflict scopes run on Opus-xhigh per clause (E) regardless.

(D. 1.1) haiku PROHIBITED WHERE IT CANNOT RECEIVE ITS GOVERNING CONTEXT – measured fact (–). In the reference consuming repository a haiku subagent dispatch failed before doing any work with "Prompt too long (~ 100k tokens vs 100k limit)" (request id req_011CfSBXm7Kn6enyUT1UEaWd , recorded in that repository's agent registry docs/requests/agent_registry.jsonl): the governance context every dispatched agent must receive there exceeds the haiku tier's context window, so a haiku dispatch cannot carry its binding context and haiku is PROHIBITED as a dispatch tier in that repository. A consuming project MAY re-admit haiku as a rung below sonnet ONLY on its OWN captured evidence (§ 2.1.2.1.1 / § 2.1.2.1.1.1) that a haiku dispatch in that project receives its full governing context and completes its task

correctly – never by assumption, and never by trimming the governing context to make it fit. A haiku dispatch that fails this way is recorded honestly as a failed dispatch (\$. . .), never silently retried on the same tier.

(E) HARD PINS OVERRIDE THIS DEFAULT FOR THEIR SCOPES.

\$. . . sets the DEFAULT working tier for NON-review, non-merge-conflict-resolution work. It does NOT relax, override, or compete with the two existing hard model-substrate pins: \$. . . (every mandatory independent code-review – \$. . . / \$. . . / \$. . . iterated to G0 per \$. . . – runs on Opus at xhigh effort as PRIMARY substrate, with a genuine-unavailability SONNET fallback, amended – – – no Fable, no Haiku, no other fallback model; BLOCKED only when BOTH Opus and Sonnet are genuinely unavailable, NEVER a lower tier or any other model) and \$. . . (every main→feature/product/flavor merge-CONFLICT resolution runs on the SAME Opus-xhigh-primary-plus-Sonnet-fallback substrate, blocked only when both Opus and Sonnet are genuinely unavailable). Those two anchors are STRICTER, SCOPE-SPECIFIC pins that this anchor's general ladder does not reach – review and merge-conflict-resolution are NOT eligible for sonnet/opus-by-complexity dispatch; they run on their pinned substrate ALWAYS, regardless of how "simple" the reviewed change or the conflict appears. A review or merge-conflict resolution dispatched on anything other than the \$. . . / \$. . . Opus-xhigh-primary-plus-sanctioned-Sonnet-fallback substrate (amended – –) is a violation of THOSE anchors, not a valid application of \$. . . 's general ladder.

(F) EFFORT-TIER SELECTION – the SAME nano-precision discipline applied to the reasoning-EFFORT axis (operator

IMPORTANT mandate, - - ; verbatim: "Besides dynamic determining proper models to use, we MUST decide the effort as well and switch between the effort settings same as we do now for chosen model! Labeling MUST BE extended for this details as well. Example: (T /main - claude - opus - high)."). Every dispatched work effort is assigned not only a model tier per clauses (A)-(E) but ALSO the LIGHTEST reasoning-EFFORT setting from the closed ladder low < medium < high < xhigh < max genuinely sufficient to complete the task correctly – chosen by the SAME clause-(A) decidable heuristic (scope / reasoning-depth / blast-radius / novelty), NOT a guess (§ . .). Illustrative bindings (non-exhaustive; the consuming project supplies its precise mapping per § . .): low / medium – trivial / mechanical work (doc-sync regeneration, status / field reads, format / lint fixes, boilerplate); high – moderate-to-complex reasoning (focused single-component debug, multi-factor analysis per § . . / § . .); xhigh / max – the HARDEST verification / merge-conflict / architecture work. (F.) DYNAMIC MID-TASK ESCALATION – effort, like model tier per clause (B), is RE-ASSESSED continuously as the task's real complexity becomes known; a task dispatched at a lower effort that PROVES harder mid-flight MUST ESCALATE its effort for the remainder (continuing under-effort past the point complexity is KNOWN is the clause-(D) false economy), and a heavy task decomposing into light bounded subtasks dispatches THOSE at a lighter effort – effort-matching applies at EVERY dispatch granularity, not only the top-level task. (F.) HONEST HARNESS BOUNDARY (§ . . , load-bearing – do NOT overstate the mechanism) – effort-tiering is REAL where the execution path exposes an effort control and a DOCUMENTED CAPABILITY-GAP where it does not: the current Claude Code Agent tool (subagent dispatch) exposes a model parameter but NO effort parameter, so a subagent's effort is NOT

presently settable via that path (the \$. . .
 <effort> label field degrades honestly to ? there); the
 Workflow tool's agent() DOES take an effort argument, so
 effort IS settable there – the label reports whichever
 path is live, never fakes one. This is the SAME
 aspirational-vs-harness state the clause-(E)
 \$. . . / \$. . . Opus- xhigh pins
 already occupy: the pins mandate the effort, and whether a
 dispatch path can SET it is a harness capability reported
 honestly, never faked (\$. . .). (F.) HARD-PIN
 FLOOR PRESERVED – clause (E)'s hard pins are ALSO effort
 pins: every \$. . . mandatory independent code-
 review AND every \$. . . merge-conflict
 resolution runs at xhigh effort on Opus as PRIMARY
 substrate (or, on a captured genuine Opus-unavailability
 FACT, on the \$. . . / \$. . . sanctioned
 Sonnet fallback, amended – – – Fable and
 Haiku remain forbidden in every role), and this general
 effort ladder NEVER tiers those two scopes below xhigh
 where the dispatch path can set it, regardless of how
 "simple" the reviewed change or the conflict appears – a
 review / merge-conflict resolution run below xhigh on a
 path that COULD set it, or run on a model other than Opus
 or the sanctioned Sonnet fallback, is a \$. . . /
 \$. . . violation, not a valid application of
 this ladder. (F. . .) HONEST BOUNDARY (\$. . .) –
 "lightest capable effort" is a CONJUNCTION, not a
 discount: an effort setting that yields a shallow /
 guessed / under-verified result is UNDER-POWERED, a FALSE
 ECONOMY, and a \$. . . violation of PASS-bluff severity,
 exactly as an under-powered MODEL tier is per clause (D);
 capability determinations are evidence-based, never
 guessed (\$. . .).

Classification: universal (§ 1.1) – a platform-neutral model-capacity-matching discipline reusable by any project dispatching work across a tiered model roster; the

consuming project supplies its own precise tier task-class mapping and its own model-name ladder per § . . . (the sonnet / opus ladder with Sonnet as the default, the clause-(D.) haiku prohibition, and the Opus-xhigh pins are this project's current concrete instantiation, not a universally-fixed vocabulary – a consumer with a differently-named tier roster maps its own ladder onto the SAME lightest-capable-first + dynamic-escalation + always-on-composition + honest-boundary + hard-pin-override structure, re-admitting a lighter rung only on its own captured evidence per clause (D.)).

Composes § . . . (no-guessing – capability determinations are evidence-based) / § . . . / § . . . (subagent-driven dispatch – the granularity at which tier-matching applies) / § . . . (blast-radius analysis feeds the clause-(A) heuristic) / § . . . (safe-reversible default on undecidable complexity – escalate up, never guess down) / § . . . / § . . . (reasoning-depth signals for the heuristic) / § . . . / § . . . / § . . . / § . . . / § . . . (the always-on composition set, clause (C)) / § . . . / § . . . (the hard-pinned review/merge-conflict substrate this anchor does NOT touch, clause (E)).

Propagation gate CM-COVENANT-114-231-PROPAGATION (literal 11.4.231 present as a block-start, exactly-once per governance file, lockstep content-hash equality across the consumer fleet per § . . . (B)) + recommended gate CM-MODEL-TIER-LIGHTEST-CAPABLE (a dispatched effort whose logged/declared model tier exceeds the clause-(A) heuristic's indicated tier for a task later confirmed trivial-and-successful-at-the-lighter-tier is a WARN/candidate finding – never a hard FAIL on tier choice alone, since tier-adequacy is only provable in retrospect per clause (D); the gate instead HARD-FAILS the two

decidable violations: (i) a review or merge-conflict-resolution effort NOT run on the § . . . / § . . . pinned substrate, and (ii) a task whose complexity was PROVEN mid-flight [an escalation event was logged, or a downstream failure/bluff was traced back to it] that did NOT escalate tier before continuing, and (iii) a dispatch assigned `fable`, or assigned `haiku` where clause (D.) applies) + paired § . mutation (gate-code = separate work item, NOT claimed shipped – § . . . ; mutation plan: () dispatch a code-review labeled as running on `sonnet` instead of `Opus-xhigh` → gate MUST FAIL; () inject a logged mid-task complexity-escalation event with NO corresponding tier-escalation record → gate MUST FAIL; (b) record a dispatch assigned `fable`, or assigned `haiku` where clause (D.) applies → gate MUST FAIL; () golden-FALSE fixture: a task that stays on its initially-assigned tier throughout, with no escalation event logged → gate MUST NOT fire, per § . . . () the false-positive guard).

Canonical authority: constitution submodule `Constitution.md`
§

Non-compliance is a release blocker regardless of context. No escape hatch – no `--always-heaviest-model`, `--always-lightest-model`, `--skip-complexity-reassessment`, `--fixed-tier-per-session`, `--relax-review-pin`, `--relax-merge-conflict-pin`, `--per-request-tier-opt-in` flag exists.

§ . . – Anti-mess long-op
orchestration: every long-op is registered,
single-owned-per-purpose, liveness-PROVEN,
handed-off-on-stop, safely-reaped,
consistency-checked, and escape-bounded – the
single source of truth for what is running,
never inferred (research-derived,
– –)

Forensic anchor (genericised): in one cycle a verification sweep produced five distinct mess states – () multiple concurrent same-purpose sweeps ran unintentionally (no single-owner-per-purpose claim); () a sweep HUNG on a consumer-blocked-on-an-unclosed-producer pipeline and was undetected because liveness was inferred from process-alive (a § . . . () FALSE=NULL: a wedged op and a progressing op returned the identical "process alive" signal); () stopping the owning subagent SILENTLY KILLED its child long-op because ownership was implicit (no handoff-on-stop); () a stale/non-growing log was read as live progress (the same false-null, reader-side); () an un-completable tooling gate blocked a release for hours (no escape to a terminal state). Per § . . .

Phase . . these are ONE architectural state-management flaw, not five bugs: **there was no single source of truth for what long-op is running, who owns it, its state, and its PROVEN liveness** – the § . . . (F)

configured-in-use gap generalised to project long-ops, so every mess state was un-observable, and an un-observable bad state is a § . . /§ . . bluff at the orchestration layer.

The mandate – a **long-op orchestration layer** binding every long-op (any op expected to exceed the runtime background-execution threshold per § . . . – verification sweep, meta-test sweep, build, flash/deploy, test/QA suite, dispatched subagent), ALL of which hold:

(A) SINGLE SOURCE OF TRUTH – the long-op lifecycle

registry. Every long-op MUST register BEFORE it runs and reach a terminal state

(`complete|failed|reaped|handoff|blocked-escape`) when it ends, in a durable (never-tmpfs, write-temp-then-rename, `$ . . . append-only event stream + atomic snapshot`) registry recording `{op_id, purpose_key, owner, pid|pgid, cwd, state, last_heartbeat, heartbeat_seq, progress_offset, log_path, verdict, evidence_path}`. Success is read from the op's terminal `verdict/ result`, NEVER its process exit code. A long-op running with no registry record – or a record with no live op – is a `FALSE-READY` drift state.

(B) SINGLE-OWNER-PER-PURPOSE. No two same- `purpose_key` long-ops run at once; acquisition is a probe-race-safe claim keyed on `purpose_key` (`$ . . .`). A live claim forces exactly ONE of REUSE (attach) or explicit SUPERSEDE (graceful handoff of the running owner, logged), never a silent second run.

(C) LIVENESS TRUTH-SOURCE. Progress is PROVEN, never inferred from process-alive: every long-op emits a monotonic heartbeat and/or an advancing progress offset; a watchdog reads the DELTA over a window and flags HUNG on no-advance past the op's declared no-progress budget (the `$ . . . ( )-( )/$ . . . discipline – kill -0` is a necessary-not-sufficient pre-filter, a stale heartbeat/non-growing log is NEVER "progressing"). A consumer blocked on an unclosed producer write-end (`$ . . . ( )`) is caught by the flat heartbeat regardless of internal state.

(D) OWNERSHIP-HANDOFF-ON-STOP. Stopping/killing an owner MUST hand its long-ops to the registry (`handoff`, re-adoptable per `$ . . .`) or deliberately reap them (`reaped`, reason logged) – NEVER silently orphan or accidentally cascade-kill; long-ops launch detached (own process-group) so an owner stop does not cascade-signal

them, and the stop path records each owned op's disposition. The loop done-condition is not satisfied while a `handoff` op is un-adopted.

11 4 196

(E) SAFE REAPING. An orphaned/hung op is reaped ONLY on PROVEN staleness – holder PID dead (`kill -0` fails) OR no-progress-budget exceeded per (C) – resolved by the process's real `/proc/<pid>/cmdline` ... (\$. . (D)/ \$.), never a bare `pgrep` /substring that could match a carrier or the reaper itself; a LIVE, ADVANCING op is NEVER reaped (\$.); every REAPED/KEPT decision + resolved evidence is logged (\$. .).

11 4 191

(F) CONSISTENCY INVARIANT SWEEP. A fast invariant check runs before every gated transition (build/flash/tag) and on a standing cadence, asserting: no duplicate owner per purpose, no orphan, no stale lock (\$. .), no hung op (C), no registry-vs-reality drift (live-long-op set == registry live rows – the \$. . (F) check generalised). Any violation is surfaced (named, actionable) and the gated transition is REFUSED until resolved. The sweep is itself a guard (\$. .): golden-TRUE + golden-FALSE-with-carrier fixtures, conservative-safe REFUSE on an unresolvable signal, resolved evidence printed on every refusal.

11 4 192

(G) UN-COMPLETABLE-BLOCKER ESCAPE. A gate that cannot reach a terminal verdict MUST degrade, after a bounded budget, to an explicit evidence-backed decision – NEVER an indefinite silent block: \$. . first (hung/(C) vs duplicate/(B) vs wrong-seam/\$. . re-seat-at-activation); NOT-YET-RUNNABLE → \$. . `artifact_not_yet_built` (non-blocking on an intermediate artifact, blocking only on the release candidate per \$. .); genuinely un-completable → \$. . (autonomous reversible SKIP-with-reason + tracked \$. . item, ELSE a bounded \$. . operator decision) – the loop keeps

progressing every other non-blocked item, it does not
idle, and the escape NEVER fabricates a PASS.

Honest boundary (\$. .). This layer reduces orchestration mess; it does NOT make product bugs impossible, does NOT make a gate correct (\$. . / \$. . still bind each gate), and does NOT remove the operator from genuine blocks – it converts an indefinite silent stall into a bounded, evidence-backed, operator-visible decision. A heartbeat proves ADVANCE not CORRECTNESS; the registry proves what CLAIMS to run, (F) proves registry==reality. **Reuse-not-rewrite:** it is the multitrack per-track agent registry (\$. . / \$. . /\$. . ^{11 4 17}) generalised from per-track workers to any long-op and bound through, per \$. . (F) – no parallel mechanism is invented.

Classification: universal (\$. . .) – no hardware/
vendor/project literal; consumers supply their long-op
set, purpose keys, no-progress budgets, and stop-hook
wiring as DATA per § 6.12. Composes

$\$_{11} \quad 4 \quad 287 \quad . \quad . \quad 16 \quad 4 \quad 69 \quad . / \$ \quad . \quad . \quad / \$ \quad . \quad . \quad /$

$\$_{12} \quad 12 \quad . \quad 9 \quad 2 \quad . \quad / \$ \quad . \quad . \quad / \$ \quad . \quad . \quad /$

$\$ \quad . \quad . \quad / \$ \quad . \quad . \quad / \$ \quad . \quad . \quad /$

$\$ \quad . \quad . \quad (D)(F) / \$ \quad . \quad . \quad () () - () /$

$\$ \quad . \quad . \quad / \$ \quad . \quad . \quad / \$ \quad . \quad /$

$\$ \quad . \quad / \$ \quad . \quad .$ Propagation gate CM-COVENANT-114-232-

PROPAGATION (literal 11.4.232) + recommended mechanism

gates CM-LONGOP-REGISTRY-SINGLE-SOURCE (A) / CM-LONGOP-SINGLE-OWNER-PER-PURPOSE (B) / CM-LONGOP-LIVENESS-HEARTBEAT (C) / CM-LONGOP-OWNERSHIP-HANDOFF-ON-STOP (D) / CM-LONGOP-SAFE-REAP-PROVEN-STALE (E) / CM-ORCHESTRATION-CONSISTENCY-SWEEP (F) / CM-UNCOMPLETABLE-BLOCKER-ESCAPE (G) + paired \$.

mutations (run two same-purpose sweeps with no reuse/
supersede → (B) FAILS; infer liveness from `kill -0` alone /
drop the heartbeat-delta check → (C) FAILS; a stop that
leaves an owned op with no terminal/handoff record → (D)

FAILs; reap a live-advancing op OR reap by bare- pgrep
 substring → (E) FAILs; a live long-op with no registry row
 passes the consistency sweep → (F) FAILs; an un-
 completable gate blocks with no escape/decision → (G)
 FAILs; strip the literal → propagation FAILs; gate-code =
 separate work item, NOT claimed shipped \$. .). No
 escape hatch – no --allow-duplicate-longop , --infer-liveness-
 from-pid , --silent-orphan-on-stop , --reap-live-op , --skip-
 consistency-sweep , --block-without-escape , --longop-without-
 registry flag.

EXTENSION landed – – – \$. .
model-in-label + flawless alias/provider model
consistency + real-time always-in-sync (User mandate,
 – –). Verbatim operator mandates
 (– –): "Extend our logging to present
 models used since we will now have dynamic models choice
 and switching! From (example) (T /main – claude) into:
 (T /main – claude – model_used), model MUST be
 dynamically obtained so we present name equivalents –
 short names – like: haiku, opus, sonnet, fable. However we
 MUST fully support all Claude Toolkit providers aliases,
 so real model used names MUST BE obtained in this case(s)
 as well! Examples: (T /main – claude – haiku), (T /
 main – claude – fable), (T /main – claude – opus),
 (T /main – claude – deepseek), (T /main – claude –
 xiaomi) ... Labels MUST BE everywhere ALWAYS up to date in
 real time and in sync!" and "Adding model names into
 labels considers that providers names / aliases will be up
 to date too and in relation to model used since models
 belong to providers / aliases !!! Consistency MUST BE
 flawless!" This EXTENDS the \$. . label
 discipline (it does NOT re-mint the anchor – the literal
 11.4.182 and its propagation gate CM-COVENANT-114-182-
 PROPAGATION are UNCHANGED and stay GREEN; this carrier
 block cites \$. . , it is not a second
 \$. . block-start per \$. . (B)).

(`1.1`) **LABEL FORM EXTENDED** – the mandated `$` . . . label is extended from `(T<N>/<branch> - <alias>)` to `(T<N>/<branch> - <alias> - <model> - <effort>)` (the trailing `<effort>` field added – – per `$` . . . clause (F) – see the EFFORT-FIELD clause (`1.2`) below), `<model>` recording which model is CURRENTLY answering for that alias: for a native Claude alias, the SHORT name (`opus` / `sonnet` / `haiku` / `fable`, case-insensitive-substring-derived from the model id the agent platform reports for the most recent turn – DYNAMIC, obtained live at label-emission time, never a fixed/assumed value since native aliases have no persistent per-alias model and the serving model can legitimately switch mid-session per `$` . . .); for a claude-toolkit PROVIDER alias (`deepseek`, `xiaomi`, `kimi-for-coding`, `opencode`, ...) the REAL provider model id as declared in that provider's host-config registry (the SAME value the alias's own launch mechanism uses – a single source of truth, never re-derived). (`1.3`) **MODEL-FIELD DERIVATION** (`$` . . . , **never guessed**) – a closed two-case split on the SHAPE of the already-derived `<alias>`: (a) native-alias → read the model id from the CURRENT session's own transcript (the platform's per-project per-session transcript; LAST assistant-turn's reported model → short name); (b) provider-alias → read the model id from the host multi-account-toolkit provider registry (`CMA_PROVIDER_MODEL`). Either lookup that cannot resolve yields the honest `?` – NEVER a fabricated or last-known value; an id obtained but matching none of the four native short-names is printed VERBATIM (real-but-unrecognised is more honest than a false `?`). The `-field` legacy form (no `<model>`) remains a VALID, ACCEPTED label during migration – the enforcing guard-hook `LABEL_RE` admits both forms, and NEVER treats an honest `?` model, or an absent model field, as a block-cause. (`1.4`) **FLAWLESS ALIAS/ PROVIDER MODEL CONSISTENCY INVARIANT** (operator NOTE,

load-bearing) – the `<alias> /provider` field and the `<model>` field MUST be MUTUALLY CONSISTENT AT ALL TIMES: the model shown MUST be a model that genuinely BELONGS to the alias/provider shown (models belong to providers/aliases), and the alias/provider name MUST itself be current – never a stale alias beside a fresh model, never a model that does not belong to the shown alias. This is guaranteed BY CONSTRUCTION, not by coincidence: the native model is read from THAT alias's own live session transcript, and the provider model is read from THAT provider's own registry entry – so the derived `<model>` is definitionally a model of the shown `<alias>`; a derivation that ever emits a model not belonging to the shown alias is a `$ . . . / $ . . .` consistency violation, never a valid label. () **REAL-TIME, ALWAYS-IN-SYNC** (`$ . . . / $ . . . (F)`) – the label is STATELESS and RE-COMPUTED at EVERY emission (every human-invoked labeler run AND every synchronous PreToolUse guard-hook firing on every Agent/Task/TaskCreate dispatch) against the LIVE environment (current `CLAUDE_CONFIG_DIR`, current session id, current transcript tail, current provider `.env`) – so a mid-session model switch (`/model`, an autonomous `$ . . .` escalation) is reflected in the VERY NEXT emission with no cache to go stale: real-time BY CONSTRUCTION, not by a sync daemon. Honest boundary (`$ . . .`): the live label has NO persisted artifact that could drift, so the Docs-Chain (`$ . . .`) sync engine – which propagates PERSISTED documents/DBs – does not apply to the stateless label itself; the operator's "use/extend Docs Chain for real-time responsiveness" ask attaches to any PERSISTED label-history / model-switch LOG surface (a `$ . . .` follow-up to be tracked), which, if adopted, IS `$ . . . -Docs-Chain-synced` and `$ . . . -request-history-adjacent` – never faked as already-shipped. () **HONEST SUBAGENT-INTROSPECTION LIMIT** (`$ . . . /`

`$ . . . )` – the native-alias transcript lookup is proven reliable when invoked from the DISPATCHING process's own shell (the guard-hook's actual execution context at dispatch time); it is NOT proven reliable when a background-dispatched subagent introspects its OWN model from inside its own subprocess on hosts where the subagent shares its parent's session-transcript with no subagent-vs-conductor marker – recorded as an open limitation to be tracked, never silently claimed solved. () **DIVERGENT-LABELER RECONCILIATION** (`$ . . . . .`) – the LIVE labeler wired via `.claude/settings.json` is `constitution/scripts/multitrack/track_branch_label.sh`; the top-level `scripts/multitrack/` path is a thin pass-through WRAPPER that `exec s` that canonical copy (`$ . . .` inherited-by-reference), so it mirrors nothing, cannot drift, and emits byte-identical output by construction. The reconciliation LANDED in commit `a48e919f52d` and is gate-enforced at `pre_build_verification.sh:41911-41915` (the gate detects the `exec` wrapper and follows it to the canonical implementation), never a silent divergence. () **EFFORT FIELD** (added – – per the operator IMPORTANT mandate – verbatim: "Besides dynamic determining proper models to use, we MUST decide the effort as well and switch between the effort settings same as we do now for chosen model! Labeling MUST BE extended for this details as well. Example: (T /main - claude - opus - high)."). The `...-field` label appends `<effort>` after `<model>`, recording which reasoning-EFFORT setting is CURRENTLY in force for that work-stream, drawn from the closed set `{low | medium | high | xhigh | max}` (the SAME effort ladder `$ . . .` clause (F) selects). (a) DERIVATION (`$ . . .`, never guessed) – `<effort>` is read from whatever LIVE effort signal the execution path exposes (the Workflow-tool `agent()` `effort` argument, a future Agent-tool effort parameter, an env/config effort setting); a path exposing NO effort signal, OR any value

outside the closed set, yields the honest ? – EXACTLY like the `<model>` field's ? fallback, NEVER a fabricated / assumed / last-known value. (b) HONEST HARNESS BOUNDARY (§ . . , load-bearing – do NOT overstate the mechanism) – effort-tiering is REAL where the execution path exposes it and a DOCUMENTED CAPABILITY-GAP where it does not: the current Claude Code **Agent tool** (subagent dispatch) exposes a `model` parameter but **NO effort parameter**, so a subagent's effort is NOT currently settable via that path and its `<effort>` field degrades honestly to ? ; the **Workflow tool's agent()** DOES take an `effort` argument, so effort IS settable + derivable there – the label reports whichever path is live, never invents one. This is the SAME aspirational-vs-harness state the § . . / § . . Opus-`xhigh` review/merge pins already occupy: the pins mandate the effort, and whether a given dispatch path can SET it is a harness capability the label reports honestly and NEVER fakes. (c) LEGACY-FORM ACCEPTANCE (migration) – the `-field (T<N>/<branch> - <alias>)` and `-field (T<N>/<branch> - <alias> - <model>)` label forms remain VALID and ACCEPTED; the guard `LABEL_RE` admits all three shapes and NEVER treats an absent `<effort>` field, or an honest ? effort, as a block-cause (§ . . – mirroring the § . . track/alias/model honest-? rule). (d) CONSISTENCY – the effort field is selected TOGETHER WITH the model by § . . (both tiers chosen for one task) and never contradicts it; a stale effort beside a fresh model, or an effort outside the closed set, is a § . . / § . . consistency violation, never a valid label. (e) REAL-TIME, ALWAYS-IN-SYNC (§ . .) – like `<model>`, `<effort>` is STATELESS and RE-COMPUTED at EVERY label emission against the LIVE environment, so a mid-task § . . effort escalation is reflected in the VERY NEXT emission with no cache to go stale. Propagation: `CM-TRACK-BRANCH-LABEL` is EXTENDED to also

assert the labeler emits a parseable `<model>` field AND a parseable `<effort>` field, the guard `LABEL_RE` accepts the `-field / -field / -field` label forms (legacy `-field` and `-field` ALWAYS remain accepted during migration; a `? effort` is NEVER a block-cause), and the alias/provider model consistency holds; paired `$ . .`. mutation: strip the `<model>` derivation OR emit a model not belonging to the shown alias $\rightarrow$ the extended gate FAILs (gate-code = separate work item, NOT claimed shipped `$ . .`). No new escape hatch – the existing `$ . .` "no escape hatch" list applies to the extended label unchanged.

`$ . .` – Standing anti-mess control plane: a level-triggered, idempotent, invariant-catalogue reconciliation loop over the WHOLE persistent System state, orchestrating every per-domain enforcer, refusing gated transitions on un-reconciled drift (research-derived, `- -`)

Forensic anchor (genericised): every existing consistency enforcer in a multi-track / multi-agent / multi-submodule / multi-upstream System is EDGE-TRIGGERED – a PreToolUse hook fires on an action, a check runs at a seam (lock-acquisition / merge-time / build-time / commit-time), a gate runs when a build runs – so a mess that arises WITHOUT an edge firing is un-observed until it causes harm (two live examples captured `. . .`)

`- -` : a large uncommitted-batch pile-up sat unreconciled through a session, and multiple provably-stale VCS/wrapper locks were held by dead holders with no acquisition attempt to reap them – the `$ . .` "a hook never installed / a check that runs only when asked" forensic generalised to the whole persistent state). Per `$ . .` Phase . the $\sim$ persistent-

state mess classes (Plane R repository: uncommitted pile-up, stale/duplicate locks, orphaned worktrees, half-applied cross-repo commits, drifted submodule pointers, partial merges, divergent mirrors, mis-placed work; Plane G governance: divergent mirrors, pointer drift, gate-ledger drift, anchor-number collision; Plane S tracking: doc/DB drift, cross-track SSoT divergence, ledger drift, handoff-doc drift; Plane P runtime: the \$. . long-op set + un-reaped streams + dangling review state) are ONE architectural flaw – there is no LEVEL-TRIGGERED control plane that re-derives the actual persistent state and reconciles it against a declared clean state; an un-observed drift is a \$. /\$. . bluff at the organization layer ("the System is clean" while a mess silently accumulates).

The mandate – a standing **anti-mess control plane** (the desired-state/actual-state reconciliation-loop + control-plane/data-plane pattern of distributed systems), ALL of which hold:

(A) DECLARED INVARIANT CATALOGUE. A single machine-readable consumer-owned catalogue (DATA per \$. . / \$. .) enumerating every consistency invariant across the persistent planes, each carrying {id, plane, desired-condition, detector, reconcile-class ∈ {auto-safe | operator-gated | owned-by <anchor>}, evidence-shape} ; the catalogue IS the desired clean state; an un-catalogued mess class is an honest tracked GAP (\$. . / \$. .), never silently claimed covered.

(B) LEVEL-TRIGGERED LOOP. The loop re-derives the ACTUAL state (not an event) and diffs it against the catalogue BEFORE every gated transition (commit-batch/build/flash/merge-to-trunk/tag/constitution-pull) AND on a standing cadence (\$. . /\$. .) – catching drift regardless of how it arose (the Kubernetes level-triggered property); actual-vs-desired uses a \$. .

fingerprint diff drilling to the offending record (anti-entropy), never a whole-scan where a fingerprint suffices.

(C) IDEMPOTENT RECONCILE-OR-REFUSE. Each drift takes exactly one dedup-keyed action: RECONCILE-AUTO-SAFE (reversible + evidence-determinable + bounded-blast per \$. . . - reap a provably-dead lock \$. . . , regen a stale derived doc \$. . . , scope-group-commit a placement-clean batch \$. . . , retire a merged worktree \$. . .), idempotent (a no-op on a clean state); OR REFUSE-THE-GATE + a bounded operator decision (irreversible/high-blast/indeterminate - id-colliding SSoT \$. . . , anchor collision \$. . . , partial cross-repo merge, divergent mirror) per \$. . . /\$. . . ; a gated transition NEVER proceeds while a catalogued invariant is violated and unreconciled; the loop keeps progressing every other non-blocked item (park-one-not-the-loop).

(D) SAGA FORWARD-RECONCILE (no-loss). A partially-applied cross-repo operation is reconciled complete-forward under a \$. . . pre-op backup (finish the remaining legs / bump the pointer / complete the merge zero-markers-zero-drop \$. . . step / ff the lagging mirror) - NEVER force-push/reset/ -s ours (\$. . . /\$. . .); compensation is \$. . . backup-restore only; the saga reaches fully-completed, never rolled-back-with-loss.

(E) THE LOOP IS A GUARD (\$. . .). Every detector ships golden-TRUE + golden-FALSE-with-carrier fixtures (\$. . . ()); a false-positive REFUSAL of a clean state is a \$. . . () FAIL-bluff exactly as a false-negative PASS is a \$. . . PASS-bluff; every "no drift" null is control-needle-proven (\$. . . () - ()); conservative-safe REFUSE on an unresolvable signal, resolved evidence printed on every refusal.

```
Classification: universal ($ . . ) – no hardware/
vendor/project literal; consumers supply their invariant
catalogue, plane detectors, gated-transition list,
cadence, and reconcile-class bindings as DATA per
```

$\S_{11} \quad 4 \quad 147 \quad \bullet \quad 1 \bullet \quad 4 \quad 144$	$\bullet$	Composes $\S$	$\bullet$	$\bullet$	$/ \quad \S$	$\bullet$	$\bullet$	$/$
$\S_{11} \quad 4 \quad 176 \quad \bullet \quad 1 \bullet \quad 4 \quad 179$	$/ \quad \S_{110}$	$\bullet$	$\bullet$	$/ \quad \S$	$\bullet$	$\bullet$	$/$	
$\S_{11} \quad 4 \quad 181 \quad \bullet \quad 1 \bullet \quad 4 \quad 186$	$/ \quad \S_{111}$	$\bullet$	$\bullet$	$/ \quad \S$	$\bullet$	$\bullet$	$/$	
$\S$	$\bullet$	$\bullet$	$/ \quad \S$	$\bullet$	$\bullet$	$/ \quad \S$	$\bullet$	$\bullet$
$\S$	$\bullet$	$\bullet$	$/ \quad \S$	$\bullet$	$\bullet$	$/ \quad \S$	$\bullet$	$\bullet$
$\S$	$\bullet$	$\bullet$	$/ \quad \S$	$\bullet$	$\bullet$	$/ \quad \S$	$\bullet$	$\bullet$
$\S$	$\bullet$	$\bullet$	$/ \quad \S$	$\bullet$	$\bullet$	$/ \quad \S$	$\bullet$	$\bullet$

§ . . (F) / § . . /
 § . . / § . . / § . . /
 § . . / § . . / § . . /
 § . / § . . Propagation gate CM-COVENANT-114-233-
 PROPAGATION (literal 11.4.233) + recommended mechanism
 gates CM-ANTIMESS-INVARIANT-CATALOGUE (A) / CM-ANTIMESS-LEVEL-
 TRIGGERED-LOOP (B) / CM-ANTIMESS-RECONCILE-OR-REFUSE (C) /
 CM-ANTIMESS-SAGA-FORWARD-NO-LOSS (D) / CM-ANTIMESS-LOOP-IS-A-
 GUARD (E) / CM-ANTIMESS-CONTROL-DATA-PLANE-SPLIT (F) + paired
 § . mutations (drop the standing-cadence trigger so
 the loop is edge-only → (B) FAILs on the pile-up/stale-
 lock fixture; let a gated transition proceed on a
 catalogued un-reconciled drift → (C) FAILs; a saga that
 force-pushes/resets to reconcile a mirror → (D) FAILs; a
 detector that REFUSES a clean golden-TRUE fixture, or
 PASSes a golden-FALSE-with-carrier → (E) FAILs; the
 control plane invokes a data-plane build/commit → (F)
 FAILs; strip the literal → propagation FAILs; gate-code =
 separate work item, NOT claimed shipped § . . /
 § . .). No escape hatch – no
 --edge-triggered-only , --skip-consistency-catalogue , --
 proceed-on-drift , --reconcile-with-force-push , --unvalidated-
 detector , --control-plane-does-work , --assume-clean-without-
 reconcile flag.

(G) DEPENDENCY-POINTER PREFLIGHT – stale, absent, and
 UNFETCHABLE are three distinct mess classes, and all three
 fail closed (extension, – –). Clause (A)'s
 Plane-R catalogue names *drifted submodule pointers*; drift
 is only one of three ways a multi-repo build ships code
 its own source tree does not contain. The AI-curriculum
 corpus (module , table row "Submodule-pointer-
 not-bumped-so-fix-reverts") states the failure exactly:
 "The fix is real and merged in a submodule, but the top-
 level image ships the old pointer, so the bug returns from
 a clean checkout" – root cause: "A multi-repo build where
 'the fix is in the submodule' and 'the image ships the

fix' are two different facts, only one of them checked."

Its prescription is a HARD gate, not a warning: *"A hard preflight gate that fails the build when a governance/ component submodule its correctness depends on is at a stale pointer or absent – 'the fix is in the submodule but the image shipped old code' must be impossible, not caught after."*

The catalogue is therefore EXTENDED with all three classes and the preflight that closes them: (i) **STALE** – the parent's recorded pointer is behind the submodule's intended ref (the drift clause (A) already names); (ii) **ABSENT / UNINITIALISED** – the submodule is declared but not checked out, so a build reads an empty directory as a satisfied dependency; (iii) **UNFETCHABLE** – the recorded pointer resolves on NO configured remote (the not our ref condition), so the tree builds for whoever already has the object locally and is structurally un-clonable for everyone else. Class (iii) is the most dangerous because it is INVISIBLE in the checkout that created it and fails only on a fresh clone. The preflight's four mechanical steps, verbatim from the corpus: *"resolves every submodule pointer, verifies it matches the intended ref, fails closed on a missing checkout, and – for the truly load-bearing components – asserts a runtime signature on the flashed/deployed artifact"* (the last step is § . . . 's RUNTIME-ON-CLEAN-TARGET layer). The general principle the corpus draws, which is why this belongs to the anti-mess control plane rather than to a build script: *"fail closed when an input to your correctness is unverifiable. An absent submodule being treated as a satisfied one is the same class as an unencrypted DB being treated as encrypted or a nil, nil stub being counted as coverage: a missing thing silently read as a present, correct thing. The rock-solid habit is to make the absence itself a hard failure, because 'the gate was skipped because the dependency was missing' is a bluff vector, not a pass."*

Per clause (C) the reconcile-

class is REFUSE-THE-GATE for classes (ii) and (iii) (a build/tag proceeding on an absent or unfetchable dependency is refused) and AUTO-SAFE only for a class-(i) drift whose correct pointer is evidence-determinable; per clause (E) the detector ships a golden-TRUE fixture per class and a golden-FALSE fixture – a fully-initialised tree at its intended refs – which it MUST NOT fire on (\$. . . ()). Composes \$. . . (pointer bumped in the same commit) / \$. . . / \$. . . / \$. . . / \$. . . / \$. . . (an unfetchable gitlink is a hermeticity violation: an undeclared-because-unobtainable input) / \$. . . (fail closed) / \$. . . (the shipped artifact must contain the tested source). Recommended mechanism gate CM-DEPENDENCY-POINTER-PREFLIGHT-FAIL-CLOSED + paired \$. . . mutation (de-initialise one submodule → the gate MUST FAIL; point one gitlink at a commit no configured remote serves → it MUST FAIL; golden-FALSE: a fully-initialised tree at intended refs MUST NOT fire it). Honest boundary (\$. . .): the preflight proves the dependency's SOURCE is present and at the intended ref – it does NOT prove the built artifact contains it (that is \$. . . 's ARTIFACT layer and the corpus's own fourth step, the runtime-signature assertion). Gate-code = separate work item, NOT claimed shipped (\$. . . / \$. . .).

\$. . . – Dedicated hook-validation script: hook checks run as an explicit stage of a dedicated commit/push script, and the commit/push mechanism is ALWAYS unblocked (operator mandate, – –)

Forensic anchor (genericised): a consumer project wired its full local CI gate as an automatic git pre-push hook

(`core.hooksPath` → hook → multi-minute build/test/sweep stage on EVERY push). The result, captured

- - : routine pushes took many minutes each, and when one gate stage inside the hook failed or stalled, the push was rejected outright – the commit/push mechanism itself became blocked, with pending commits stranded locally and no operator-usable path forward short of a forbidden `--no-verify` bypass. The gate logic was sound; wiring it as an uninterruptible implicit hook on the commit/push code path was the architectural flaw – an edge-triggered enforcer (`$ . .`) placed directly on the operator's only forward path, violating the `$` single-entripoint discipline by making the entripoint unusable.

The mandate, ALL of which hold:

11 4 35

(A) DEDICATED SCRIPT AS THE SINGLE COMMIT/PUSH ENTRYPPOINT.

Every consumer MUST ship a dedicated, executable, idempotent script (Lava binding: `scripts/commit-push-all.sh`; consumers bind their own path as DATA per `$ . .`) that runs, in order: () submodule/workspace sync, () commit of all pending state, () the hook validations as an EXPLICIT stage – the unmodified hook logic fed the real push range, Layer-cheap checks always, long gates stage-separated, () push to ALL configured upstreams (`$ .`), () a closing green/clean verification (clean status + zero unpushed commits, recursive over submodules). The script prints per-stage evidence and is safe to re-run; direct `git commit / git push` on the main repo ceases to be the routine workflow (`$`).

(B) HOOKS MUST NOT BLOCK THE MECHANISM. Automatic git hooks wired via `core.hooksPath` MUST NOT gate routine commit/push in a way that can leave the repository uncommittable or unpushable. Disconnecting means: `core.hooksPath` unset, the hook file itself preserved UNMODIFIED in the repo, and `--no-verify` NEVER used as the

push – this anchor keeps that entrypoint usable),
 \$. . . (sweep still runs, now as a named stage),
 \$. . . (no force-push history-rewrite – the
 unblocked path never rewrites), \$. . . (the
 script is a guard: a false GREEN verification is a FAIL-
 bluff), \$. . . (amendment discipline),
 \$. . . (the commit/push path is a gated
 transition the anti-mess control plane reconciles).

11 4 35

Classification: universal (\$. . .) – no hardware/
 vendor/project literal; consumers supply their script
 path, upstream-remote list, cheap-vs-long gate split, and
 skip-flag name as DATA per \$ Propagation gate
 CM-COVENANT-114-234-PROPAGATION (literal 11.4.234) +
 recommended mechanism gates CM-DEDICATED-HOOK-VALIDATION-
 SCRIPT (A) / CM-HOOKS-NEVER-BLOCK-PUSH (B) / CM-NO-GATE-LOST-
 ON-HOOK-DISCONNECT (C) / CM-COMMIT-PUSH-ALWAYS-UNBLOCKED (D) +
 paired \$. . . mutations (wire the full gate back as an
 automatic blocking pre-push hook → (B) FAILS; disconnect
 the hook AND drop a check from every stage → (C) FAILS; a
 validation failure that strands commits with no
 remediation report → (D) FAILS; strip the literal →
 propagation FAILS; gate-code = separate work item, NOT
 claimed shipped \$. . . /\$. . .). No escape
 hatch beyond the recorded (D) deferral – no --no-verify-
 routine , --skip-all-gates-silently , --delete-blocking-hook ,
 --push-without-any-validation flag.

\$. . . – Build-and-deploy the
 moment source fixes are proven-correct (test-
 side hardening is a parallel stage, never a
 build gate) + post-deploy version increment
 (operator mandate, – –)

(A) BUILD-AND-DEPLOY THE MOMENT THE SOURCE FIXES ARE
 PROVEN-CORRECT – TEST-SIDE WORK IS A PARALLEL STAGE, NEVER
 A BUILD GATE. The moment the SOURCE fixes of a batch are

done AND proven-correct – proven-correct = the mandatory independent SOURCE-correctness review has returned a clean GO (\$. . . /\$. . . 11 4 188 /\$. . . on the source diff, on the \$. . . Opus- xhigh substrate, iterated to zero-finding/zero-warning GO per \$. . .) – the build MUST be triggered, pre-build + post-build tests EXECUTED (\$. . . layer-before, layer- after), and the artifact deployed. IN PARALLEL WITH the build AND the manual-QA session (\$. . .), NEVER gating the build: () gate-instrumentation hardening + new-gate authoring (\$. . .); () paired \$. . . mutations + four-layer \$. . . (b) coverage + HelixQA Challenges; () "gates and checks of all kinds" (\$. . .); () the re-review of THAT test-instrumentation work. The build's GO-SIGNAL is source-correctness alone; test-guard hardening (future-regression coverage) + its re-review are a deploy-independent parallel stage per \$. . . (A). Holding the build sequentially – waiting for the test-instrumentation buildout + its re-review – when the source is already proven-correct is a \$. . . (A) violation.

(B) POST-DEPLOY VERSION INCREMENT + MANUAL-QA-DEPLOY CYCLE BOUNDARY (extended – – – operator mandate, verbatim: "Increase upcoming release to . . . since new cycle starts every time we deploy release for manual QA testing! This is how we will be working this ALWAYS! Extend the constitution!"). The moment image(s)/artifact(s) deploy, the version code/version is INCREMENTED (next build targets the next version; every deployed artifact carries a distinct, monotonic, greppable id – \$. . . /\$. . . /\$. . .). Two deployments sharing one version id defeats per-deploy artifact identity → \$. . . (B) violation. EVERY deployment of a release FOR MANUAL QA TESTING (\$. . .) is a DEVELOPMENT-CYCLE BOUNDARY: the

manual-QA deploy CLOSES the outgoing cycle and STARTS a NEW development cycle, and the version code/version is incremented AT THAT POINT – the new cycle's work targets the NEXT monotonic version (§ . . . project-prefixed naming..preserved) – ALWAYS, the STANDING working model per clause (C), never a per-request opt-in. The manual-QA deploy is the canonical clause-(B) increment trigger (a deploy-for-manual-QA IS a deploy), and it SHAPES the § . . . release-cycle framing: fixes for QA findings against the deployed artifact land in the NEW cycle on the NEW version id, never patched onto the already-QA-deployed version id. Starting a new cycle's work still targeting the QA-deployed version id is a § . . . (B) violation.

(C) STANDING DEFAULT (§ . . . /§ . . . / § . . .), never a per-request opt-in.

(D) FINDING-LAYER CLASSIFIER – the decidable split clause

(A) relies...on, mandatory at every § . . . / § . . . /§ . . . review verdict (added – – ; research-derived from the § . . . four-agent audit of this project's own Speckit- T remediation saga). Every finding a code-review verdict records (§ . . . (F) machine-written) MUST carry exactly one `finding_layer` from the closed set `{source-defect | test-instrumentation | process-doc}` (a field DISTINCT from any other per-finding classification the review record already carries – e.g. this project's own `review_record.py` pre-existing "class" field `{mechanical | judgment | false-positive}` is a different axis..and is unaffected): `source-defect` – the finding shows the SHIPPED artifact's own behaviour wrong, unsafe or exploitable as it will actually run, OR invalidates the evidence that proves THIS batch's fixes correct (their § . . . RED→GREEN guard, § . . . runtime signature, or covering test per

§ . . () – a proving guard that cannot fail or whose reviewer-authored mutation survives per

§ . . ()(d) means the source is NOT proven-correct); `test-instrumentation` – the finding concerns only future-regression coverage, mutation-guard robustness or detection completeness of a guard/gate/meta-test whose protected behaviour was ALREADY proven, with no production decision path shown wrong and no proof-of-this-batch evidence invalidated; `process-doc` – labels, comments, citations or prose with no behavioural effect. The layer is recorded by the REVIEWER (§ . . /

§ . . producer≠verifier) – the producer MAY contest it only via a fresh review round and MUST NOT lower it unilaterally. A finding that cannot be decidably classified defaults to `source-defect` (the conservative-safe default per § . . ()); secondarily

§ . . safe-reversible). § . . is NOT relaxed: every finding of every layer is still driven to a zero-finding, zero-warning GO in its own loop; clause (D) decides only which GO is the build GO-signal – a batch/slice is BUILD-ELIGIBLE when its `source-defect` findings reach that GO, while outstanding `test-instrumentation` / `process-doc` findings continue as their own tracked § . . item on a § . . (A)/

§ . . parallel stream. A BATCHED review (one review spanning many slices/files) MUST record its verdict PER SLICE/FILE, never one undifferentiated batch verdict, so one slice's unresolved `test-instrumentation` loop cannot withhold build-eligibility from the other slices. Clause (D) applies prospectively; recorded verdicts are never rewritten or reclassified (§ . . /§ . .).

Honest boundary (§ . .): this clause does not make classification easy or infallible – it makes it a required, recorded, reviewer-owned decision instead of an implicit one.

Honest boundary (§ 11.4.125): relocates the TEST-INSTRUMENTATION re-review + gate/check hardening OFF the build critical path; does NOT weaken any substantive gate. The SOURCE-correctness review (§ 11.4.126 / § 11.4.129 • 11.4.130) Opus- xhigh → GO § 11.4.129 • 11.4.130) STILL PRECEDES + gates the build (it IS "source proven-correct"); pre/post-build tests (§ 11.4.126 L -) run at their seams; § 11.4.126 manual-QA-final un-weakened; § 11.4.126 / § 11.4.129 • 11.4.130 un-weakened. "Test-side work is parallel, never a build gate" is bounded to future-regression instrumentation – a check that proves the SHIPPED artifact wrong is a source/artifact-correctness gate (§ 11.4.126 L -), stays on the critical path. Clause (D) supplies the mandatory classification mechanism this boundary's own test presupposes – a review that never classifies its findings has not actually applied this honest boundary, it has merely stated it.

Classification: universal (§ 11.4.125). Propagation gate CM-COVENANT-114-235-PROPAGATION (literal 11.4.235) + recommended CM-BUILD-ON-SOURCE-PROVEN-NOT-TEST-SIDE (FAIL: build held past source-GO w/ unfinished test-instrumentation; FAIL: build started with NO source-review GO – the § 11.4.126 () dual-assertion guard) + CM-VERSION-INCREMENT-ON-DEPLOY (FAIL: new-artifact deploy shares a version id; FAIL: work of the cycle opened by a manual-QA deploy still targets the QA-deployed version id – the clause-(B) cycle-boundary increment skipped) + paired § 11.4.126 mutations. CM-BUILD-ON-SOURCE-PROVEN-NOT-TEST-SIDE is IMPLEMENTED (constitution/scripts/gates/cm_build_on_source_proven_not_test_side.sh + its paired _mutation_test.sh), as is CM-VERSION-INCREMENT-ON-DEPLOY (constitution/scripts/gates/cm_version_increment_on_deploy.sh + its paired mutation test) – the prior "gate-code = separate work item" disclaimer for those two was stale and

is corrected here. NEW recommended gate `CM-REVIEW-FINDING-CLASS-RECORDED` for clause (D) (FAIL: a batched-review verdict covering more than one slice/file with no per-slice verdict or no per-finding `finding_layer` recorded; FAIL: a slice marked build-eligible while any `source-defect` finding in it has not reached GO; FAIL: a `finding_layer` lowered by the producer without a fresh review round) + paired `$ . mutation - gate-code =` separate work item, NOT claimed shipped (`$ . . / $ . .`), registered as a tracked deferral in `constitution/scripts/gates/gate_ledger_deferrals.tsv`; it binds ADDITIVELY by computing whether the existing `source_review_go` marker may be written, never by modifying `cm_build_on_source_proven_not_test_side.sh`'s own tested logic.

Non-compliance is a release blocker. No escape hatch – no `--build-behind-test-hardening`, `--test-side-gates-the-build`, `--skip-source-correctness-review`, `--deploy-without-version-increment`, `--not-default-speed-mode`, `--serialise-build-and-validate`.

`$ . .` – QA-deploy-readiness gate: no manual-QA hand-off until the mandated validation produced a candidate-fingerprinted PASS verdict; a blocker MUST bind to a seam, never prose (research-derived, `- -`)

Forensic anchor (genericised, `- -`). An operator declared a specific validation a **release-blocking requirement** for every QA-deploy ("MANDATORY ... QA MANUAL TESTING READY BLOCKER ... HAS TO BE FULLFILLED COMPLETELY FOR EACH RELEASE"), and it was still not enforced at the next QA-deploy: the validation harness existed and ran, hit an instrument bug (a false-FAIL on later legs), and the blocker was **downgraded to a non-blocking follow-up** while the deploy proceeded on a weaker side-check. Nothing

STOPPED the deploy because the requirement lived as **prose** (an agent reads it) + a source-gate that asserted the harness exists – NOT a seam that refuses the deploy in the absence of a clean recorded verdict; the deploy tool had **zero** references to the validation (measured). This is the § . . . "documented-as-enforced but implemented nowhere is worse than no rule" + § . . . "prose doesn't bind, seams do" failure, demonstrated against the operator's own explicit blocker, plus a § . . . judgment error (a bug in a blocker's instrument is not licence to skip the blocker – fix the instrument).

The mandate (ALL hold). () DEPLOY-SEAM VERDICT-COVERAGE (§ . . . applied to the QA-deploy seam). A release artifact **MUST NOT** be deployed, published, or handed off for manual QA (§ . . .) until the mandated autonomous validation(s) for that artifact produced a **machine-written PASS verdict** whose **artifact fingerprint == the candidate** (§ . . . (F) – the fingerprint READ FROM THE TARGET at run time). **ABSENCE of the verdict BLOCKS the deploy exactly as a FAIL does** – a not-yet-run validation and a failed one are both refusals, never a silent pass. () **MECHANICAL, NOT PROSE** (§ . . . /§ . . .). The refusal lives in the deploy tool itself (the seam), never in prose an agent is trusted to honour; a source-gate that asserts the validation harness exists is necessary-not-sufficient – it does not prove a clean verdict for the candidate. () **BLOCKER-BINDING RULE.** Any requirement an operator declares a blocker (a § . . . CRITICAL, or an explicit "release blocker / QA-ready blocker") **MUST** be bound to a mechanical seam at the point it gates, never satisfied by prose and never DOWNGRADED to a follow-up – and when the blocker's own INSTRUMENT bugs, the response is § . . . /§ . . . (STOP, fix the instrument, complete the blocker), never route around it. () **EXPLICIT, RECORDED OVERRIDE ONLY.** The seam has no silent

bypass; an override is an explicit, logged, operator-set flag (the \$. . . /\$. . . recorded-deferral pattern – never a follow-up-downgrade, never an opaque skip). () **THE GATE IS SELF-VALIDATED** (\$. . . ()/\$. . .). The deploy seam ships golden-good (verdict present + PASS + fingerprint-match → allow) + golden-bad (absent → refuse / not-pass → refuse / fingerprint-mismatch → refuse), so it cannot itself false-allow; and the verdict-emitter reads the fingerprint from the target at run time, never a hardcoded per-release value (\$. . . – a version-pinned validator is the brittleness that lets a stale verdict pass). **Honest boundary** (\$. . .): the seam proves a candidate-matching PASS verdict EXISTS for the mandated autonomous validation – it does NOT replace the \$. . . manual-QA sufficiency gate (the human check, the deploy hands off T0), and it is only as honest as the validation whose verdict it consumes (that validation's own anti-bluff stays \$. . . /\$. . . /\$. . . /\$. . .). Classification: universal (\$. . .) – consumers supply the mandated-validation set, the verdict path/schema, and the deploy tool per \$ Composes \$. . . /. ()/ /. (F)/. /.

Propagation gate CM-COVENANT-114-236-PROPAGATION (literal 11.4.236.) + recommended gate CM-QA-DEPLOY-READINESS-GATE (the deploy tool refuses without a candidate-fingerprinted PASS verdict; the gate self-validates golden-good/bad; the verdict-emitter reads the fingerprint at run time) + paired \$. . . mutation (make the deploy tool proceed on an absent/mismatched verdict → the gate FAILs; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped \$. . . /\$. . .). Non-compliance is a release blocker. No escape hatch – no --deploy-without-verdict , --prose-blocker-suffices , --

1.1 4 237
downgrade-blocker-to-followup , --skip-readiness-gate , --
verdict-fingerprint-optional , --silent-override flag.

§ . . – Mandatory exhaustive
context-and-spirit-aware translation review:
every localized artifact independently
reviewed for CONTEXT (source) + SPIRIT
(target-language idiom/register) + technical/
non-translatable fidelity – accuracy
necessary, never sufficient (operator mandate,
– –)

Forensic anchor (operator mandate, – –): a
translated artifact that is merely ACCURATE – every source
token faithfully rendered – can still read as stilted,
foreign, machine-literal, register-wrong, or idiom-broken
to a native speaker of the target language, and a
localization pass that silently mangles a technical term
or "translates" a non-translatable product/proper name
ships a defect worse than untranslated text; the recurring
failure mode is a translation that scores green on an
accuracy check while a native reader immediately perceives
it as non-native. Accuracy is NECESSARY; it is NOT
SUFFICIENT for a slick, polished, professional
localization.

The mandate (ALL hold):

(A) EVERY LOCALIZED ARTIFACT IS INDEPENDENTLY REVIEWED –
NEVER SELF-REVIEWED. Every translation and every
translated artifact – UI strings, pages, documents, PDFs,
manuals, marketing copy, captions/subtitles, help text,
any localized content – MUST undergo a HEAVY, EXHAUSTIVE,
independent review by a reviewer (agent or human)
STRUCTURALLY SEPARATE from the translator that produced it
(\$. . /\$. . independence). The
translator's own pass (\$. . multi-pass self-
review) PRECEDES this review and NEVER satisfies it

§ 4.1.1 's no-self-review-substitution discipline applied to localization; § 4.1.2 precedes, never satisfies). A localized artifact accepted without an independent context-and-spirit review is a § 4.1.3 PASS-bluff at the localization layer.

(B) THE REVIEWER HOLDS AND APPLIES THREE DIMENSIONS. (a) **FULL CONTEXT OF THE SOURCE MATERIAL** – the original ideas, the subject/domain, the purpose, and the intended audience; the reviewer reads the source **MEANING**, not merely the source **STRING**, so intent, nuance, and tone are carried across, never lost. (b) **THE SPIRIT OF THE TARGET LANGUAGE** – its tradition, idiom, register, tone, and cultural conventions – so the result reads as **NATIVE**, natural, fluent, and **MAXIMALLY POLISHED**, never literal / word-for-word / calqued; a construction correct in the source that reads foreign in the target is **REWRITTEN** into the natural target-language form. (c) **CORRECT TECHNICAL VOCABULARY AND VERBATIM PRESERVATION OF NON-TRANSLATABLE TERMS** – domain/technical terms rendered with their established target-language equivalents, and **NON-TRANSLATABLE** tokens (product names, brand names, proper nouns, technology/library/API names, code identifiers, trademarks) preserved **VERBATIM** – never "translated", never accidentally transliterated, never localized away.

(C) THE GOAL IS MAXIMALLY POLISHED LOCALIZATION – ACCURACY NECESSARY, NEVER SUFFICIENT. The bar is a slick, polished, professional, native-reading localization; fidelity to the source is the **FLOOR**, not the ceiling. A rendering that is accurate but stilted / foreign / register-wrong / idiom-broken is a **FINDING**, not a pass.

(D) ANTI-BLUFF – CAPTURED MACHINE-READABLE POSITIVE EVIDENCE + FAIL-LOOPS-TO-CLEAN. The review emits captured, machine-readable **POSITIVE** evidence per § 4.1.4 / § 4.1.5 (`feature_class = translation_review`): a per-artifact verdict carrying a **PER-DIMENSION** rationale

across the closed dimension set {fidelity, fluency/naturalness, spirit/register, terminology, completeness, non-translatable-preservation} – a PASS with no captured per-dimension rationale is a § . /§ . . bluff. A FAIL on ANY dimension LOOPS BACK to re-translation/polish and re-review, iterating to a zero-finding clean verdict per § . . – a single pass that "looked fine" never closes the artifact.

11 4 140

Honest boundary (§ . .). This anchor REFINES the existing localization discipline from PASS/FAIL accuracy checking to CONTEXT-AND-SPIRIT native-polish review; it does NOT replace § . . 's canonical translation pipeline (the HelixTranslate-canonical-translation-pipeline sense of that anchor number) nor § . . 's independent per-language review (the independent-per-language-translation-review sense of that anchor number – a KNOWN operator-owned § . . anchor-number collision on 11.4.140 / 11.4.141 , re-mint still operator-owned; cited here strictly in those senses, introducing no new collision), it STRENGTHENS them (accuracy → native polish). It does not make a mistranslation impossible (a miscalibrated reviewer stays § . . () golden-good/golden-bad fixture territory), and where a target language genuinely has no established term for a source concept, that gap is an honestly recorded finding (§ . . SKIP-with-reason), never a silently invented word.

11 4 35

Classification: universal (§ . .) – no vendor/project/language literal; the consuming project supplies its language set, term glossary, non-translatable token list, and reviewer substrate as DATA per § . . . Composes § . . (canonical translation pipeline) / § . . (independent per-language translation review – this REFINES it from PASS/FAIL accuracy to context+spirit polish) / § . . /

§ . . / § . . / § . . /
 § . . / § . . / § . . Propagation
 gate CM-COVENANT-114-237-PROPAGATION (literal 11.4.237) +
 recommended gate CM-TRANSLATION-CONTEXT-SPIRIT-REVIEW (every
 localized artifact carries an independent per-dimension
 context-and-spirit review verdict with captured rationale
 across {fidelity, fluency/naturalness, spirit/register,
 terminology, completeness, non-translatable-preservation},
 FAIL looping to re-review per § . .) + paired
 § . mutation (a translation accepted on an accuracy-
 only check, OR self-reviewed by its own translator, OR
 carrying a golden-bad-seeded mistranslated idiom / a
 "translated" non-translatable product name → the gate
 FAILs; strip the literal → the propagation gate FAILs;
 gate-code = separate work item, NOT claimed shipped
 § . . / § . .).

Canonical authority: constitution submodule Constitution.md

§ . . .

Non-compliance is a release blocker regardless of context.

No escape hatch — no --skip-spirit-review ,

--accuracy-only-OK , --self-review-translation-OK , --literal-

translation-OK , --translate-non-translatable-OK , --skip-per-

dimension-rationale flag exists.

§ . . — Automated QA must be the
DISCOVERER: manual QA finds nothing new, and
 anything it finds is a coverage escape (User
 mandate, — —)

Forensic anchor — verbatim user mandate

(— —): "Helix QA full websites testing
 MUST find all such issues, not us! Once we do manual QA we
 shall not be able to find anything! THIS IS MANDATORY
 RULE / CONSTRAIN / GOAL to achieve with this and any other
 project and MUST BE part of the root constitution!"

Forensic case study (FACT, - -). In a single cycle on a governed project, five real defects were found – a health endpoint that returned `{ "status": "healthy", "components": [] }` and was structurally incapable of ever reporting unhealthy; a gateway routing every completion at a port nothing served, so the product's primary capability returned HTTP for an indeterminate period; two submodule pins recorded on `main` that upstream answers with `not our ref`, so no fresh clone could check them out; and two hardware-profiler defects reading a key from a file where it does not exist. **Every one was found by an agent reading source code or probing by hand. Not one was found by the automated QA system,** which reported green throughout. The defects were real, user-visible, and long-lived; the QA suite's silence about them was the actual defect.

The mandate (ALL hold):

(A) THE AUTOMATED QA SYSTEM IS THE DISCOVERER, NOT THE CONFIRMER. The project's automated QA system (HelixQA per § . , or the consuming project's equivalent full-system/full-website test regime per § .) **MUST** be the mechanism that FINDS defects. Discovery by a human, by an operator's manual pass, by an agent reading code, or by an end-user report is a **failure of the QA system**, not a success of the finder. The finding is still valuable and still gets fixed – but the QA gap it exposes is a defect of equal standing, and closing only the reported defect while leaving the gap is the violation this anchor forbids.

(B) MANUAL QA MUST FIND NOTHING NEW. By the time manual QA runs (§ .), the automated regime **MUST** already have surfaced everything it would find, so manual QA functions as a CONFIRMATION and sufficiency gate – never as the discovery mechanism. The target state is explicit and measurable: *a manual QA pass that discovers zero*

previously-unknown defects. A manual pass that keeps finding new defects is evidence the automated regime is not yet doing its job, regardless of how green it reports.

(C) EVERY ESCAPE TRIGGERS THE § . . . LOOP, WIDENED. When ANY defect is discovered outside the automated QA system – manual QA, operator observation, end-user report, agent code-reading, or incidental discovery during unrelated work – it triggers, before the fix may close: () a § . . . systematic-debugging pass to FACT-root-cause the defect itself; () a **coverage-escape audit** identifying WHY the automated regime did not find it, cited to the specific missing check, unexercised path, absent assertion, or untested surface – never "we didn't think of it"; () a **new or strengthened automated check** that WOULD have found it, registered into the standing regime (§ . . .) with its § . . . RED capturing the escaped defect; () the audit committed as captured evidence. § . . . already mandates this for operator-found defects that a green suite passed; § . . . widens the trigger to EVERY out-of-band discovery channel, because the failure is identical whoever happens to notice first.

(D) COVERAGE IS MEASURED BY WHAT THE REGIME CAN CATCH, NOT BY WHAT IT RUNS. A QA regime's worth is its demonstrated ability to FAIL on a real defect, not its pass count, test count, or run duration. Every automated surface MUST be falsifiable per § . . . – a check whose paired mutation does not make it fail is decoration, and a suite composed of such checks will report green through exactly the defects this anchor exists to prevent. The regime MUST exercise real user-reachable paths end-to-end (§ . . . no manual intervention, § . . . real user journeys, § . . . real content, § . . . /§ . . . /§ . . . recorded

+ read-the-screen verification), because a check that never traverses the path a user traverses cannot find what a user finds.

(E) THE GOAL IS DIRECTIONAL AND MEASURED. Projects MUST track the discovery-channel split – how many defects in a cycle were found by the automated regime versus out-of-band – and drive the out-of-band share toward zero over successive cycles. Honest boundary (§ 11.4.27): this anchor does NOT claim a QA regime can ever be provably complete (§ 11.4.28 – absence of evidence is not evidence of absence), and it does NOT license closing an escape by weakening the definition of "defect". It mandates that every escape measurably strengthens the regime, so the channel that finds defects shifts from people to machines cycle over cycle.

Composes § 11.4.29 (the QA system) / § 11.4.30 (operator-escape bluff-audit – § 11.4.31 widens its trigger set) / § 11.4.32 (manual QA as final sufficiency gate – § 11.4.33 makes it a confirmation, never a discovery, step) / § 11.4.34 (enumerated discovery pressure) / § 11.4.35 (standing regression guards) / § 11.4.36 / § 11.4.37 / § 11.4.38 / § 11.4.39 / § 11.4.40 / § 11.4.41

Classification: universal (§ 11.4.42) – a platform-neutral QA-ownership discipline reusable by ANY project with an automated QA regime and a manual QA step; the consuming project supplies its concrete QA system, its surface inventory, and its discovery-channel ledger per § 11.4.43. Propagation gate CM-COVENANT-114-238-PROPAGATION (literal 11.4.238) + recommended gates CM-QA-IS-THE-DISCOVERER (every defect closed in a cycle records its discovery channel; an out-of-band channel requires a committed coverage-escape audit + a registered automated

check that would have caught it) and CM-MANUAL-QA-FINDS-
NOTHING-NEW (a manual QA pass reporting new defects blocks
the release tag until each has its coverage-escape audit
and new check) + paired § 11.4.227 meta-test mutation (strip
the literal → propagation gate FAILs; close an out-of-
band-discovered defect with no coverage-escape audit → CM-
QA-IS-THE-DISCOVERER FAILs; gate-code = separate work item,
NOT claimed shipped § 2024.18.26 /§ 102.11.4.238).

**EXTENSION – Discovery-channel record schema + escape-
ratchet no-data-point rule (spec-derived, speckit –
anti-slop-enforcement, – –).**

§ 11.4.238. (C) already states the CONCEPT – a
coverage-escape audit citing the specific missing check
that would have caught an out-of-band-discovered defect –
but does not fix the record's field schema, and
§ 11.4.238. (E) states only that the out-of-band share
is driven "toward zero over successive cycles," a
directional framing this project's own spec.md identifies
as amended for the same reason SC-018 's original
wording was: *"trend to zero over successive cycles" names
neither an input condition nor an observable pass/fail
outcome, so it was not machine-checkable ... the ratchet
is the decidable form*" (spec.md: , SC-
amendment note). This extension supplies the decidable
schema and ratchet. () **Discovery-channel record schema.**
Per spec.md: (FR-): every recorded defect
MUST carry a discovery channel drawn from the CLOSED SET
{automated_seam, manual_qa, operator, end_user, agent_inspection}
(or the consuming project's declared equivalent per
§ 11.4.238) AND a should_have_been_caught_by value naming
a specific seam or gate, OR the literal none with a
WRITTEN justification; a record missing either field is
REFUSED. The record MUST be WRITTEN BY THE VERIFIER SEAM,
never by the producer of the defect's fix – a producer-
authored discovery-channel record is refused (the same
producer≠verifier separation § 11.4.238 already

mandates, applied to the record OF the escape, not only the record of the fix). () **Escape count is monotone-decreasing from a recorded baseline, and never increases.**

Per spec.md: (FR-): a cycle whose escape count exceeds the recorded baseline is REFUSED at the release seam; a cycle whose escape count is lower moves the baseline down. () **A manual-QA-not-run cycle contributes NO DATA POINT – it MUST NOT lower the baseline.** Per spec.md: (FR- , verbatim): *"a cycle in which the manual pass did not run, that cycle MUST contribute no data point and MUST NOT lower the baseline"* – a zero escape count from a cycle nobody inspected is not a measurement (the \$. . . () false-null pattern applied to this specific ratchet: an unmeasured cycle and a genuinely-clean cycle both read as zero, and only the `manual_qa_ran` flag distinguishes them).

() **Records tagged `should_have_been_caught_by: none` are counted and reported SEPARATELY.** Per spec.md: : *"so narrowing the definition of an escape is visible rather than silent"* – a project cannot lower its escape ratchet by quietly reclassifying escapes as uncatchable-by-design; every such reclassification is a visible, separately-tallied line item, never absorbed into the main count.

Honest boundary (\$. . .), stated in the same spec.md amendment note (SC-) this extension grounds itself

in: *"Validated against the definition of done: at the correct end-state, every escape caught by a seam, the count reaches , so the ratchet is a valid necessary condition. It is not sufficient – zero also results from nobody looking, which is what the no-data-point rule and the separate `none` tally exist to expose."* Composes

\$. . . ()-() (the false-null + metric-validated-against-definition-of-done pattern this ratchet instantiates) / \$. . . () (the release-verdict-coverage monotone-decrease adoption ratchet – this extension's escape-count ratchet is the SAME decidable-

monotone-baseline pattern applied to § . . 's discovery-channel domain, not a novel mechanism) / § . . (the standing zero-shortcomings monotone-decreasing findings ledger – the general-purpose instantiation of the identical never-increases/may-decrease/no-silent-narrowing discipline this extension's escape ratchet and none -tally separate-reporting rule both specialise). Recommended mechanism gates CM-DISCOVERY-CHANNEL-RECORD-COMPLETE (every recorded defect carries a closed-set channel + a should_have_been_caught_by seam-name or a justified none ; the record is verifier-seam-written, never producer-authored) + CM-ESCAPE-RATCHET-NO-DATA-POINT-HONEST (a cycle whose escape count exceeds the recorded baseline refuses the release seam; a cycle with manual_qa_ran false contributes no data point and does not lower the baseline; none -tagged records are tallied separately from the main escape count) + paired § . mutations (submit a defect record with no channel, or a should_have_been_caught_by: none with no justification → the schema gate MUST FAIL; author a discovery-channel record from the producer of the defect's own fix → the schema gate MUST FAIL; run a cycle with manual_qa_ran false and zero recorded escapes, then assert the baseline lowers → the ratchet gate MUST FAIL; golden-FALSE per § . . (): a genuine manual-QA cycle with a real zero-escape result → the ratchet gate MUST allow the baseline to lower). Gate-code = separate work item, NOT claimed shipped (§ . . / § . .).

Canonical authority: constitution submodule Constitution.md § . . .

Non-compliance is a release blocker regardless of context. No escape hatch – no --manual-qa-may-discover , --skip-coverage-escape-audit , --found-by-human-is-fine , --close-escape-without-new-check , --qa-green-is-enough , --discovery-channel-not-tracked flag exists.

§ . . – Critical-invariant work-class Definition of Done mandate – money/safety/availability/integrity code MUST include failure-path scenarios as gates, not merely happy-path passes (research-derived, – –)

Compact summary: every project's Definition of Done (§ . . /§ . . lifecycle + § . . closure vocabulary) MUST identify a closed CRITICAL-INVARIANT WORK CLASS – code paths where a single incorrect outcome is unrecoverable in the window between failure and remediation (money movement, safety-of-life or safety-of-property outputs, availability of the primary user-facing capability, integrity of the persistent record) – and MUST require that a change touching this class carry, as part of its acceptance-criteria set, EXPLICIT FAILURE-PATH SCENARIOS turned into passing tests BEFORE closure. A DoD that stops at "code compiles + happy-path test is green" for a critical-invariant change is a machine for generating reopened tickets on exactly the paths that can least afford them (extraction file line – CITATION NOTE (– – citation audit, § . .)): the "extraction file" offsets cited throughout § . . –§ . . are absolute offsets into a consolidated extraction artifact that is NOT tracked in this repository and whose numbering does NOT resolve against the tracked corpus at docs/research/ai_curriculum_2026_08/source_corpus/ – the highest cited offset () exceeds the –line total of ALL NINE module files concatenated, and every module-scoped offset cited below either exceeds its module's own line extent or points at unrelated text. Each extraction offset is therefore retained ONLY as an un-resolvable provenance pointer, NEVER as a resolvable location, and is

accompanied by a VERIFIED module_NN:LINE locator measured against the tracked corpus with a control-needle-proven instrument per § . . . ()(b). The corpus itself landed - - , AFTER these anchors landed, so it corroborates them; it is not their derivation source. – module , verified locator module_27:30 : "A Definition of Done that stops at 'the code compiles and the happy-path test is green' is a machine for generating reopened tickets on the paths that move money"; line [verified locator module_27:231]: the load-bearing DoD lines are the failure-mode tests and captured runtime evidence). § . . . does NOT invent a new lifecycle or a new evidence class; it PROMOTES the failure-path coverage the constitution's anti-bluff family (§ . /\$. . /\$. . /\$. .) already demands into a MECHANICAL PRECONDITION of any closure-status write on the critical-invariant class, so the "happy-path green means done" bluff cannot ship on money/safety/availability/integrity code by construction.

(A) THE CRITICAL-INVARIANT WORK CLASS – CLOSED SET, CONSUMER-BOUND. Every consuming project MUST declare, as DATA per § . . . , its own concrete instantiation of the four critical-invariant categories. The universal categories are: (i) MONEY – any code that credits, debits, transfers, holds, releases, refunds, or reconciles a monetary value or its digital equivalent; extraction line [verified locator module_27:30] anchors the canonical example ("a donation either completed correctly or it did not; there is no partial credit and no maintenance window"); (ii) SAFETY – any code whose outputs actuate physical hardware, medical devices, transport systems, industrial controls, or otherwise reach into the physical world where a wrong output can cause harm; (iii) AVAILABILITY – any code on the critical path of the product's primary user-facing capability whose failure degrades the user's ability to use the product at all

(login, primary action, data read on the load-bearing path); (iv) INTEGRITY – any code that mutates the persistent record whose corruption cannot be silently repaired (append-only ledgers, audit logs, cryptographic anchors, source-of-truth writes governed by § 2.2.1 / § 2.2.2). A code region is IN the class when at least one category applies; a region OUT of the class stays governed by the ordinary DoD. § 2.2.3 does NOT claim the four categories are exhaustive for every domain – the consuming project supplies extensions (regulatory-compliance outputs, privacy-critical PII flows, legal-hold data) per § 2.2.4, and each extension carries the same failure-path-mandatory discipline.

(B) FAILURE-PATH SCENARIOS AS FIRST-CLASS ACCEPTANCE

CRITERIA – NOT AFTERTHOUGHTS. For every change touching the critical-invariant class per (A), the change's acceptance-criteria set MUST enumerate the domain-specific failure paths BEFORE the code is authored, each rendered as a Given/When/Then scenario (module acceptance-criteria pattern – module attribution VERIFIED CORRECT – : Given/When/Then occurs in module ONLY across the whole nine-module corpus, verified locators module_27:94 , module_27:268 , module_27:275 , module_27:318 ; extraction lines – canonical example, offsets un-resolvable per the § 2.2.5. CITATION NOTE) whose "Then" observes a real user-observable outcome (a ledger row, a settled charge, a state delta) rather than "no error" (§ 2.2.6 / § 2.2.7). The canonical failure-path scenario set for MONEY code (extraction lines – ; verified locator module_27:312) is the reference minimum: () FRESH – the happy path; () IDEMPOTENT RETRY – same key + same payload produces no new side effect; () SAME-KEY-DIFFERENT-PAYLOAD – the request is rejected, not quietly served (extraction line [verified locator

module_27:312]: "the two criteria juniors omit are same-key-different-payload ... and the crash-in-the-middle");

() CRASH-IN-THE-MIDDLE – the process is killed between the external commit and the local persist and re-driven, converging to exactly-once. Analogous minimum scenario sets for SAFETY (nominal + fault-injected upstream + degraded actuator + power-loss mid-actuation), AVAILABILITY (nominal + dependency degraded + dependency down + retry storm), and INTEGRITY (nominal write + concurrent conflicting write + mid-write crash + corruption-detected-on-read) are consumer-supplied per § . . . with the SAME "failure paths as gates, not as afterthoughts" property. Every scenario becomes a passing test BEFORE closure; a scenario omitted is a § . . . finding of § . . . /§ . . . severity at the acceptance-criteria layer, NOT a "we'll add it if it bites us" deferral.

(C) MECHANICAL PRECONDITION AT THE STATUS-WRITE SEAM. The failure-path-coverage precondition binds through the existing § . . . (D) status-custody chain, not a parallel mechanism: a done/ready/closed-status write against an item whose implementing diff touches the (A)-declared critical-invariant class is REFUSED unless the item's registry row cites acceptance-criteria coverage of the domain's declared failure-path scenario set (with each scenario mapped to a § . . . (F) machine-written verdict pair for the tests that exercise it, and each verdict carrying its class-matched evidence per § . . .). The refusal is at the SEAM (the workable-items engine per § . . . refuses the write), not a prose review guideline, per § . . . 's "seams not prose" pattern. An item whose scenario-map lists a scenario as "N/A" MUST cite the § . . . -honest reason the scenario does not apply (SAFETY change with no upstream fault possible on this

actuator; a MONEY read-only endpoint with no crash-in-the-middle exposure) – a bare "N/A" is treated as omission.

(D) COMPOSITION WITH § . . . /

§ . . . / § . . . / § . . .

§ . . . 's four-layer verification (SOURCE → ARTIFACT → RUNTIME → USER-VISIBLE) remains the layer discipline for every failure-path scenario – the acceptance test does NOT close on source-green, it closes on runtime-class evidence per § . . .

§ . . . 's huge-blocker STOP-fix-all-full-RESTART protocol governs the response WHEN a failure-path scenario surfaces a defect during release validation.

§ . . . 's manual-QA sufficiency gate is unweakened – automated failure-path coverage is NECESSARY, never sufficient, and the operator's final human oracle remains the terminal check. § . . . 's automated-QA-as-discoverer mandate binds the automated regime to FIND the failure-path defects BEFORE manual QA – a manual-QA session that discovers a critical-invariant failure-path defect the automated regime missed is a

§ . . . coverage-escape audit AND a

§ . . . acceptance-criteria audit (the scenario should have been enumerated at authoring time).

(E) HONEST BOUNDARY (§ . . .). Failure-path coverage is DIRECTIONAL – it forbids the "happy path suffices" shape and mandates a domain-appropriate minimum failure-path enumeration; it does NOT claim exhaustive enumeration of every possible failure mode (that direction crosses into formal-methods and § . . . discovery-pressure territory). The domain-specific minimum scenario sets in (B) are the STARTING POINT, not the ceiling – extraction line [verified locator module_27:211] documents the ratchet: "the fix is never 'try harder'; it is to add the specific missing gate ... so the next occurrence is caught before close," so every escape

becomes a permanent addition to the scenario set
 (§ . . . regression-guard family applied to the
 acceptance-criteria set). Composes § . . . (status
 lifecycle) / § . . . (type set) / § . . .
 (closure vocabulary) / § . . . (consumer-owned
 category extensions) / § . . . (TDD-fix – the
 failure-path test is the RED before the fix) /
 § . . . (positive-evidence taxonomy – every "Then"
 is an observed effect) / § . . . (four-layer
 verification stays in force per scenario) /
 § . . . (F). (machine-written verdict pairs bind
 each scenario) / § . . . (each failure-path
 scenario becomes a permanent regression guard) /
 § . . . (D) (status-custody chain enforces the
 precondition) / § . . . (manual QA remains the
 sufficiency gate) / § . . . (seam-not-prose
 enforcement) / § . . . (evidence-class-at-closure
 ranks the scenario's captured evidence) / § . . .
 (automated QA discovers, manual QA confirms). Extraction
 sources [offsets un-resolvable per the § . . .
 CITATION .NOTE; BOTH module attributions AUDITED
 – – and VERIFIED CORRECT]: module
 ("Definition of Done, Acceptance Criteria & Workable-Items
 Governance") lines – , – ,
 – , – – verified
 locators module_27:30 , module_27:211 , module_27:231 ,
 module_27:312 , plus the Given/When/Then acceptance-
 criteria pattern at module_27:94 / 268 / 275 / 318 (offsets
 – additionally exceed module_ 's own
 –line extent, which is itself the proof the offsets
 are extraction-file-absolute rather than module-relative);
 module 28 ("Anti-Bluff Verification & Release
 Validation") lines – verification-vs-
 validation depth – verified locators module_28:139 ,
 module_28:141 (offsets – likewise exceed
 module_ 's own –line extent).

Classification: universal (§ 11.4.239). Propagation gate CM-COVENANT-114-239-PROPAGATION (literal 11.4.239) + recommended gate CM-CRITICAL-INVARIANT-FAILURE-PATH-DOD (a change whose implementing diff touches the consumer-declared critical-invariant class per (A) carries an acceptance-criteria map naming each domain-required failure-path scenario per (B); a scenario listed "N/A" carries a § 11.4.239 -honest reason; a closure-status write against an item whose map is incomplete or whose failure-path tests lack § 11.4.239 (F) machine-written verdicts is refused at the status-write seam per (C)) + paired § 11.4.239 mutation (submit a MONEY-touching change whose acceptance-criteria set enumerates only the happy path → gate FAILs; strip the same-key-different-payload scenario from an idempotent-endpoint change → gate FAILs; author a scenario whose "Then" observes only "no error" rather than a state delta → the composed § 11.4.239 gate FAILs on evidence class; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped § 11.4.239 /§ 11.4.239).

Canonical authority: constitution submodule Constitution.md § 11.4.239. Non-compliance is a release blocker. No escape hatch – no --happy-path-suffices-for-money , --skip-failure-scenarios , --defer-failure-path-tests , --criticality-not-applicable-without-reason , --n/a-without-justification , --closure-without-scenario-map flag.

§ 11.4.239 – **Producer ≠ Verifier scope separation:** the actor that produces a change/artifact/verdict cannot be the one that gates it – least privilege over cleverness (research-derived, – –)

Compact summary: every project MUST enforce STRUCTURAL SEPARATION between the actor that PRODUCES a change/

artifact/verdict and the actor that GATES that same change/artifact/verdict – the producing actor cannot write, weaken, delete, or edit the acceptance tests, the oracle, the gate configuration, the verdict store, the paired § . mutation, or any other check whose green result is a precondition for the producer's own work being certified done. Enforcement is by CAPABILITY (the producer STRUCTURALLY cannot reach the gate) rather than by INSTRUCTION (the producer is TOLD not to touch the gate), and the enforcement mechanism is LEAST PRIVILEGE OVER CLEVERNESS (extraction file line – MODULE

ATtribution CORRECTED – – : this citation originally read "module ". The quoted passage has ZERO occurrences in module_27.txt and resolves in the tracked corpus to module ("Anti-Bluff Verification & Release Validation"), verified locator module_28:393 . Instrument control-needed per § . . ()(b) before the absence was reported. – module , verified locator module_28:393 : "the durable fix is least privilege over cleverness – the same logic as prompt-injection defense ... the load-bearing move is that the agent chasing the green gate cannot reach the gate, the oracle, or the evidence store"). The failure mode this anchor forbids: a producer that CAN edit the tests/oracle/gate config/verdict store DOES edit them under long-horizon optimization pressure – hardcoding expected values, weakening assertions, mocking the unit under test, deleting the failing test, or writing verified: true on itself (extraction lines , – , – , –). § . . makes the separation a STRUCTURAL PROPERTY of the codebase and its automation, not a social one.

(A) THE FOUR SCOPES + THE SEPARATION MATRIX. Every change/artifact/verdict flows through four SCOPES that MUST be distinct actors under § 101.02: (i) PRODUCER – the

do not touch the tests." Extraction lines

```

- [MODULE CORRECTED - -
- the (E) Extraction-sources line attributes these offsets
to module ; the quoted passage resolves to module
, verified locator module_28:383 ] name the load-
bearing mechanism: "The producer has no write access to
the test files, the oracle, the CI/gate config, or the
verdict store – enforced by scope, not by asking. If it
can't touch the gate, it can't game it." Concrete
enforcement channels the consuming project MUST supply per
$ . . . : ( ) FILESYSTEM SCOPE – the producer's
write scope excludes tests/ , oracle/ , .github/workflows/ ,
scripts/gates/ , verdicts/ (or their consumer-declared
equivalents); ( ) BRANCH / REVIEW PROTECTION – protected
paths (see $ . . . PROTECTED-SPEC + CODEOWNERS
from the sibling - - operator decision)
require an actor OTHER than the change author to approve
modifications; ( ) VERDICT-STORE ISOLATION – the store
that records RED/GREEN verdicts per $ . . . (F) is
written only by the harness that ran the test, never by
the producer's own code path; ( ) EXECUTION ISOLATION –
the verifier runs in a CLEAN CHECKOUT with FRESH CONTEXT
(extraction lines , ,
[MODULE CORRECTED - - – the (E)
Extraction-sources line attributes these offsets to module
; the quoted passage has ZERO occurrences in
module_34.txt and resolves to module / module
, verified locators module_29:655 , module_30:591 ,
module_30:655 ]: "an independent verifier re-running in a
clean checkout" – the verifier does NOT inherit the
producer's dirty worktree, cached artifacts, or in-flight
session state. "The agent will not do that" is NEVER
acceptable enforcement – extraction line
[module attribution VERIFIED CORRECT
- - , verified locator module_34:87 ]
catalogues the failure ("We told it in the prompt not to

```

hardcode / not to skip tests" as the shallow / red-flag answer). Exhortation is not capability; § . . . refuses exhortation as the enforcement layer.

(C) THE THREE FORBIDDEN PRODUCER SCOPE COLLAPSES. The following collapses are § . . . violations regardless of context, scale, or claimed rigor: () PRODUCER=ORACLE – the same actor writes the source diff AND the test that asserts its correctness, with no independent oracle authoring the assertion; the resulting test asserts whatever the code happens to do ("tautological testing trap," extraction line [module attribution VERIFIED CORRECT - - , verified locator module_34:155], referenced) rather than what the specification requires. Countermeasure: the oracle is authored FROM the specification / from a characterization baseline captured BEFORE the change (§ . . .) / from a reference implementation (§ . . . derived oracle), not co-authored with the diff. () PRODUCER=GATE – the same actor writes the source diff AND has write access to the gate configuration / CI pipeline / the mutation-testing thresholds / the paired § . . . mutations; the producer can silence failing gates or weaken thresholds. Countermeasure: the gate configuration lives in a § . . . -protected path, and modifications require a non-producer approver. () PRODUCER=VERIFIER (self-certification) – the same actor produces the change AND writes the "verified: true" / "PASS" verdict for the change; the verdict is a producer-set flag with no independent re-run. Countermeasure: verdicts are MACHINE-WRITTEN by the harness that ran the test per § . . . (F), keyed to the target's artifact fingerprint read at run time, and the release-verdict-coverage seam per § . . . REFUSES a release tag while any producer-asserted (rather than harness-written) verdict is in the release scope. Extraction lines

, , , [MODULE
 CORRECTED - - - attributed to module
 in the (E) Extraction-sources line; the quoted
 passage resolves to module , verified locator
 module_28:278]: "Tamper-evident and outside the producer's
 reach – the producer of the change cannot edit the verdict
 or the oracle to make itself green."

(D) COMPOSITION WITH \$. . (F) /
 \$. . / \$. . / \$. . /
 \$. . (D) / \$. . 125 /
 \$. . / \$. . \$. . (F)
 machine-written verdict pairs are the (C)(₁₄₂)
 countermeasure at the verdict layer – a producer-typed
 status label is not a verdict. \$. . 's code-
 review-agent gate before pre-build + main build IS the (A)
 VERIFIER role at the review layer, and \$. . 's
 universal every-change-review mandate + \$. . 's
 Opus-xhigh-primary-plus-genuine-unavailability-Sonnet-
 fallback review-substrate pin (amended
 146 3 - - ; Fable and Haiku remain forbidden in
 every role) provide the independence + substrate
 requirements. \$. . 's standing regression-guard
 suite is the GATE layer at the release seam.
 \$. . (D) status custody is the STATUS-WRITE-
 SEAM enforcement – the store REFUSES a producer-asserted
 terminal-status write without the full custody chain.
 \$₁₄₀ . . (universal independent verification
 agent) is the VERIFIER role generalised across every
 output. \$. . (exhaustive all-scenario review)
 is the depth requirement applied to the verifier's pass.
 \$. . does NOT invent a new mechanism at any of
 these layers – it ARTICULATES the property (structural
 producer-vs-verifier separation) that the existing
 mechanisms collectively enforce, and it forbids the

collapses (§ . . . (C)) that would render those mechanisms decoration.

11 4 117 10

(E) **HONEST BOUNDARY** (§ . . .¹⁵¹²). Structural separation forbids the „CLASS of self-certification failure modes; it does NOT prove the verifier is correct (a mis-calibrated verifier stays § . . . () golden-good/golden-bad territory; extraction line [MODULE CORRECTED - - - implicitly module via the (E) Extraction-sources line; the quoted passage resolves to module , verified locator `module_28:401`]: "a flaky or false-positive gate destroys the whole evidence chain above it"), does NOT prove the oracle is correct (§ . . . mandates oracle-strategy naming + independence), and does NOT guarantee the producer never tries to game the gate (it forbids the producer from succeeding at gaming the gate through direct write access; other gaming vectors – hardcoding to visible inputs, over-fitting the seen cases – are countered by § . . . held-out oracles + property/metamorphic/differential coverage). Composes § . . . (no guessing – the producer's claim is a prediction, only the verifier's captured re-run is evidence) / § . . .^{1507 118} (subagent-driven – the verifier IS the natural subagent dispatch) / § . . .¹¹ (subagent-driven default) / § . . . (multi-pass evaluation – self-review PRECEDES independent review, never satisfies it) / § . . . () (self-validated analyzers – the oracle carries its own golden fixtures) / § . . . (F) (machine-written verdicts) / § . . .^{11 4 146} (code-review-agent gate) / § . . . (iterate-to-GO – the verifier's clean-GO terminates the loop) / § . . . (standing regression-guard suite) / § . . . (universal every-change-review) / § . . . (D) (status custody chain refuses producer-asserted terminals) / § . . . (universal independent verifier agent) / § . . . (exhaustive review) / § . . .

certifies it), `module_34:87` (we told it in the prompt not to hardcode), `module_34:155` (tautological testing trap); the offsets `-` and `15`, listed under module `15`, are a WRONG-MODULE attribution – the phrase quoted from them ("an independent verifier re-running in a clean checkout") has ZERO occurrences in `module_34.txt` and resolves to `module 15 / module 15`, verified locators `module_29:655`, `module_30:591`, `module_30:655`. **UNKNOWN** (§ `15` . `15` , stated rather than guessed): offsets `15`, `15`, `15` – `15`, `15` – `15`, `15` and carry no quoted text in this anchor, so no phrase-level locator could be measured for them individually – their module attribution is UNVERIFIED: neither asserted correct nor asserted wrong.

(F) Independence-tier vocabulary + Tier-C capability-separation triggers, degrading honestly on a single-uid host (extension, spec-derived, speckit -anti-slop-enforcement, - -). Clauses (A)–(E) state the SEPARATION PRINCIPLE and its capability-not-instruction enforcement in the abstract; this clause supplies the CLOSED VOCABULARY for recording HOW MUCH separation was actually achieved, and the two triggers that raise the required strength – because a project claiming "capability separation" where only "instance separation" is genuinely achievable is itself a § `15` . `15` violation of the theatre this anchor exists to prevent. () **The closed tier set.** Per spec.md: (FR-): every accepted verdict's recorded basis of independence MUST name which tier applied, drawn from the CLOSED SET {instance, model, capability} – instance (a separate agent/process instance verifies, sharing no in-flight state with the producer), model (a DIFFERENT model verifies – permitted, never required per () below), capability (a genuine host-level

access-control boundary separates the producer's write scope from the verifier's – the strongest tier). ()

Model diversity is permitted, never required – the false-positive guard. Per spec.md: (FR-): a verdict from an independent INSTANCE running the SAME model as the producer MUST be ACCEPTED on that basis alone (Tier B, permitted-never-required) – a check that refuses same-model independent verdicts has turned a legitimate separation into an unnecessary same-model prohibition, which is itself the exact class of false-positive refusal § . . . forbids. This clause does NOT weaken the constitution's OTHER, distinct, narrower model-substrate pins – § . . . 's requirement that a translation review run on a DIFFERENT model/provider than the translator, and § . . . 's requirement that every mandatory independent code-review run on the Opus model at xhigh effort as PRIMARY substrate, with a genuine-unavailability Sonnet fallback (amended

– – ; Fable and Haiku remain forbidden in every role); both remain in force unmodified for their own narrower domains (translation review; code review).

FR- 's permitted-not-required framing governs only the GENERAL producer≠verifier case this anchor addresses, never those two domain-specific, stricter pins. ()

Tier-C trigger – the producer holds write access to the evidence/verdict store. Per spec.md: (FR-): at a seam where the producer holds write access to the evidence or verdict store, a verdict MUST be written through the MOST RESTRICTIVE separation genuinely available on the host, and the ACHIEVED separation (not the aspired-to one) MUST be reported. **Measured, honestly-degrading limit (spec.md:):** *"on a single-uid host, capability separation cannot be established by any mechanism an agent can set up – ownership defeats access-control lists, the append-only attribute requires privilege the agent lacks, and any key the producer can*

use it can also read. Claiming `capability` there would be the theatre this feature exists to prevent." A single-uid host therefore reports `instance` honestly, NEVER `capability`, until the operator takes the one action that establishes a real host boundary. () **Tier-C trigger – the three named high-blast-radius seams.** Per `spec.md`: (FR-): at the `release-tag`, `qa-deploy`, or `manual-qa-handoff` seam (or the consuming project's declared equivalent seams per § .), a producer-authored verdict is REFUSED regardless of the write-access analysis, and the achieved separation is reported rather than assumed. **Same honest-degradation limit applies** (`spec.md`:): "Capability-strength at these three seams requires one operator action (a host boundary agents are correctly forbidden from creating); until it is taken they operate at `instance` tier and MUST say so." () **No demand for capability outside the two triggers.** Per `spec.md`: (FR-): at a seam matching NEITHER Tier-C trigger, a separate agent instance MUST SUFFICE and capability separation MUST NOT be demanded – instance-tier verdicts are accepted at ordinary seams without a capability check, closing the negative half of this clause (a project over-applying Tier-C strength to every seam is itself a § . . violation of the same "false-positive refusal" class). **Honest boundary** (§ . .). The `independence_tier` field records WHAT WAS ACHIEVED, never what was ASPIRED to – a project honestly reporting `instance` at a Tier-C-triggered seam because no host boundary yet exists is COMPLIANT with this clause (the honest-degradation path is the compliant path); a project reporting `capability` without the mechanically-verified boundary that makes it true is the violation. Composes § . . (honest reporting over aspiration) / § . . (a false-positive `capability` claim is the same class of bluff a false-positive guard refusal is) / § . . (its own,

distinct, domain-specific different-model requirement stays unweakened) / § 11.4.236 (its own, distinct, domain-specific Opus-xhigh code-review pin stays unweakened alongside § 11.4.237's) / § 11.4.238 (host-boundary establishment is a § 11.4.239 host-safety-adjacent operator action, never an autonomous agent action). Recommended mechanism gate CM-INDEPENDENCE-TIER-HONESTLY-RECORDED (every accepted verdict's independence_tier field is present, within the closed {instance, model, capability} set, and – at a Tier-C-triggered seam without a mechanically-verified host boundary – is NEVER capability; a same-model independent verdict at an ordinary seam is accepted without a capability check) + paired § 11.4.240 mutation (record capability at a seam where no verified host boundary exists → the gate MUST FAIL; refuse a same-model independent verdict at an ordinary seam → the gate MUST FAIL [the false-positive-guard case]; accept a producer-authored verdict at release-tag / qa-deploy / manual-qa-handoff → the gate MUST FAIL; golden-FALSE per § 11.4.241 (): an instance-tier verdict honestly reported at a Tier-C-triggered seam with no host boundary yet established → the gate MUST NOT fire). Gate-code = separate work item, NOT claimed shipped (§ 11.4.242 / § 11.4.243).

Classification: universal (§ 11.4.244). Propagation gate CM-COVENANT-114-240-PROPAGATION (literal 11.4.240) + recommended gates CM-PRODUCER-NOT-VERIFIER-SCOPE-SPLIT (a change's four scopes per (A) – producer, oracle, gate, verifier – resolve to distinct actors; a producer-authored verdict, a producer-modifiable gate config, a producer-modifiable oracle, or a producer-modifiable test file in the same change violates the split) + CM-LEAST-PRIVILEGE-CAPABILITY-NOT-INSTRUCTION (the enforcement channel is a filesystem-scope / branch-protection / verdict-store-isolation / execution-isolation mechanism per (B), NOT an instruction / prompt / policy-document / social-contract;

a project whose only enforcement is a written policy that "producers do not touch tests" fails the capability check) + paired § . mutations (make the producer branch a modification to a tests/* file within the same change without a non-producer approver → the CODEOWNERS + capability check FAILs; author a status write of "verified: true" from the same actor that produced the diff, with no harness-written § . . (F) verdict pair → the § . . (D) custody chain FAILs; grant the producer write access to the gate configuration + expect the review to catch the collapse → the capability-not-instruction gate FAILs on the enforcement-channel check; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped § . . / § . .).

Canonical authority: constitution submodule Constitution.md § . . . Non-compliance is a release blocker. No escape hatch – no --producer-may-write-tests , --self-certification-OK , --exhortation-suffices-as-enforcement , --producer-authored-verdict-accepted , --capability-limit-optional , --policy-document-is-enforcement , --single-actor-for-all-scopes flag.

§ . . . – Illegal-state-unrepresentability preference: types before API-shape before lint before property-test before runtime-assertion (research-derived, – –)

Full anchor text per Phase landing brief. See curriculum_anchors_241_254_full_drafts.md § . . canonical block (Part A). Anchor establishes the closed ordered ladder () TYPE SYSTEM → () API SHAPE → () LINT → () PROPERTY-BASED TEST → () RUNTIME ASSERTION for every load-bearing invariant; a defect preventable at

rung N caught only at rung > N is a finding. Composes
 § . . . (pushes enforcement LEFT of Layer) /
 § . . . / § . . . / § . . . /
 § . . . / § . . . / § . . . /
 § . . . / § Includes rung-by-language
 capability baseline (strongly-typed: rungs - ;
 dynamically-typed with gradual: rung conditional;
 untyped scripting: rungs - unconditional);
 genuinely-limited-environment clarification per
 § . . . reviewer MINOR- (POSIX sh with only
 rung SATISFIES the anchor IFF the runtime assertion
 GENUINELY CATCHES the invariant AND is paired with its
 § . . . meta-test mutation per § . . . ()
 self-validation). Classification: universal
 (§ . . .). Propagation gate CM-COVENANT-114-241-
 PROPAGATION (literal 11.4.241) + recommended gate CM-
 ILLEGAL-STATE-LADDER-PREFERENCE + paired § . . . mutation
 (gate-code = separate work item, NOT claimed shipped
 § . . . / § . . .).

Canonical authority: constitution submodule Constitution.md
 § Non-compliance is a release blocker. No
 escape hatch – no --runtime-assert-suffices , --skip-type-
 preference , --dynamic-language-exempts-all-rungs , --property-
 test-in-lieu-of-type , --decorative-type-annotation-OK flag.

§ . . . – Bisection-not-blame regression cause-finding + failure modes (research-derived, - -)

Full anchor per Phase landing brief §
 When a previously-working feature/behaviour regresses, the
 FIRST-choice localisation mechanism is BISECTION (git
 bisect or equivalent monotone-search, O(log N)) over the
 ordered commit range between the last known-good tag per

leaves two whole classes unaddressed: a defect with no known-good commit (present since inception, or input-dependent), and a defect whose discriminating variable is not a commit at all. The AI-curriculum corpus (module `ddmin`) supplies the generalisation and states the containment relation in the direction that matters: **git bisect is delta debugging over commits**, not the reverse. Delta debugging (Zeller; the corpus does not use the name `ddmin` and neither does this clause) is defined as: *"given a large failing input, it systematically removes chunks and re-tests, keeping the input failing while shrinking it, until removing any remaining piece makes the failure disappear – a -minimal failing case. The algorithm starts by trying to remove large chunks and progressively removes smaller ones, so it converges quickly, roughly binary-search-like on the input."* The corpus's enumerated generalisation axes, verbatim: *"Bisect a config – remove half the flags/settings and see if the bug survives. Bisect the data – which record in the batch triggers it? Bisect the change set – git bisect is delta debugging over commits. Bisect dependencies – which upgraded package flipped behavior?"* (feature flags are named as a further axis in the same module; STATE is NOT an enumerated axis and is not claimed as one here – § 1.1.1). **The mandate.** When a failure is not a monotone commit-range regression, or when a commit-bisection returns a change too large to explain the failure, the investigation **MUST** narrow along the discriminating axis – input, config, data, dependency, or flag – to a MINIMAL failing case before a fix is designed, and the minimal case becomes the § 1.1.1 RED test's input. **The predicate carries the same monotone-search invariant § 1.1.1 already imposes:** the corpus requires *"a stable, fast, unambiguous test predicate – 'does this input still trigger this failure' (not a different one)"*, and for a non-deterministic failure warns that *"minimization is*

noisy because a 'good' reduction step might just have gotten lucky, so I run the predicate multiple times per step or accept a probabilistic minimum" (composes

\$. . . deterministic consistency and

\$. . . – stabilise or bound the noise before trusting the reduction). Honest boundary (\$. . .): minimisation produces the SMALLEST INPUT that still triggers the failure – the corpus's claim is that this "usually is the diagnosis", not that it IS the root cause; the root cause remains \$. . . 's phase.

Recommended mechanism gate CM-DELTA-DEBUG-MINIMAL-CASE (an investigation of a non-monotone or non-commit-discriminated failure records the narrowed axis and its minimal failing case, and that case is the RED test's input; a fix designed against an un-narrowed failing input → FAIL) + paired \$. . . mutation (strip the minimal-case record so the fix is designed against the full unreduced input → the gate MUST FAIL; golden-FALSE per

\$. . . (): a genuine monotone commit-range regression localised by git bisect alone MUST NOT fire it). Gate-code = separate work item, NOT claimed shipped (\$. . . / \$. . .).

Canonical authority: constitution submodule Constitution.md

\$ Non-compliance is a release blocker. No escape hatch – no --blame-log-suffices , --skip-bisection , --flaky-observable-bisect-anyway , --run-script-exit-code-collapse-OK , --regression-guess-without-localisation flag.

§ 3.1.1 – Characterization / golden-master safety net for legacy code before behaviour-changing modification (research-derived, – –)

Full anchor per Phase 1 landing brief § 3.1.1. Before ANY behaviour-changing modification to legacy code lacking characterization tests (code inherited without written behavioural contract, code whose original author has moved on, code whose documented behaviour has diverged from actual behaviour, code that "works but nobody knows why"), the modifier MUST FIRST author CHARACTERIZATION TESTS asserting CURRENT observable behaviour as GROUND TRUTH (whether "correct" or not) and land them as a green baseline. Only AFTER the safety net exists, may the modification proceed. Pattern originates in Feathers's WORKING EFFECTIVELY WITH LEGACY CODE + generalises GOLDEN-MASTER (record outputs across representative input space; assert future runs match). LEGACY SCOPE (§ 3.1.1 – honest): (i) lacks tests asserting observable behaviour; OR (ii) written spec diverges from measured actual behaviour; OR (iii) predates the constitution's TDD discipline (§ 3.1.1); OR (iv) modifier cannot describe its exact behaviour from reading source in time budget. AUTHORIZING LOOP: () identify observable surface → () generate representative input space → () run + record actual outputs as ground-truth baseline → () assert recorded values in tests → () commit characterization tests as SEPARATE commit BEFORE the behaviour-changing modification. DISTINCTION FROM CORRECTNESS: characterization tests capture WHAT IS not WHAT SHOULD BE – DESCRIPTIVE not prescriptive. If current legacy behaviour is subtly wrong, the characterization test CAPTURES that wrongness; the modifier's job is to notice, ASK the operator per § 3.1.1 whether to preserve or correct that behaviour, never silently. Composes § 3.1.1 (RED-first – characterization is

legacy counterpart) / § . . (investigate-before-remove – parallel discipline) / § . . (sampled input space = honest gaps, not completeness) / § . . / § . . (test-first backfilled for legacy). Classification: universal (§ . .). Propagation gate CM-COVENANT-114-243-PROPAGATION (literal 11.4.243) + recommended gate CM-CHARACTERIZATION-BEFORE-LEGACY-MODIFICATION + paired § . mutation.

Canonical authority: constitution submodule Constitution.md § . . . Non-compliance is a release blocker. No escape hatch – no --modify-untested-legacy-directly , --skip-characterization , --claim-tests-not-needed-because-obvious , --decorative-assertion-suffices , --refactor-without-safety-net flag.

§ . . – Cross-boundary contract tests + can-i-deploy gate (research-derived, – –)

Full anchor per Phase landing brief § . . . Changes modifying a boundary between independently-deployable components (service consumer, library dependents, DB schema readers/writers, message topic subscribers) MUST be covered by CONSUMER-DRIVEN CONTRACT TESTS on both sides AND a can-i-deploy gate MUST verify proposed provider version compatibility against every currently-deployed consumer version BEFORE the change ships. § . . MERGES contract-tests + can-i-deploy (they are useless in isolation – contract tests without the gate produce compatibility knowledge nobody consults; a can-i-deploy gate without contract tests has nothing to check). CDC pattern: consumers publish assumptions they make about the provider (Pact /

Spring Cloud Contract / gRPC schema-evolution rules); providers verify each consumer's contract against actual implementation; provider change breaking a live consumer's contract fails provider's build BEFORE the change ships – "a precise red at the source." CAN-I-DEPLOY: at every deploy attempt, gate queries the contract broker for the compatibility matrix; single red cell BLOCKS the deploy with machine-readable verdict naming the incompatible pair, NEVER warning-with-continue. MISSING contract = first-class refusal reason per § 11.4.244 (a false-positive REFUSAL is a FAIL-bluff, false-negative allowance is a PASS-bluff). SemVer becomes MEASURABLE (PATCH bump breaking contract = mis-labelled version). Backward-compatibility window declared per § 11.4.244. Composes § 11.4.244 / § 11.4.244 / § 11.4.244 / § 11.4.244 (per-side four-layer) / § 11.4.244 (contract tests as permanent regression guards) / § 11.4.244 (composed-system manual QA still needed) / § 11.4.244 (gate honesty) / § 11.4.244. Classification: universal (§ 11.4.244). Propagation gate CM-COVENANT-114-244-PROPAGATION (literal 11.4.244) + recommended gates CM-CONTRACT-TESTS-BOTH-SIDES + CM-CAN-I-DEPLOY-GATE-WIRED + paired § 11.4.244 mutations.

Canonical authority: constitution submodule Constitution.md § 11.4.244. Non-compliance is a release blocker. No escape hatch – no --skip-contract-tests, --deploy-without-compatibility-check, --assume-compatible-when-broker-silent, --patch-may-break-contract, --contract-tests-optional-for-internal-boundaries flag.

§ . . – Oracle-problem-first test
 authoring: identify the oracle BEFORE
 authoring the test (research-derived,
 – –)

Full anchor per Phase landing brief § . . .
 Every test MUST FIRST identify its ORACLE – the mechanism
 deciding PASS vs FAIL for the input it drives – and the
 oracle MUST be DIFFERENT from the code under test AND from
 the test's own assertion mechanism. Authoring without
 explicit oracle identification produces the "test agrees
 with the code" anti-pattern (test asserts what the code
 happens to produce, so it passes as long as the code
 produces something, never catches wrong-output – the most
 common form of § . /§ . . bluff at the test-
 authoring layer). The oracle problem is the OLDEST
 unsolved problem in software testing; § . .
 does not solve it, it mandates HONESTLY IDENTIFYING which
 oracle strategy each test uses so the reviewer per
 § . . can judge whether the identified oracle
 is genuinely independent. ORACLE STRATEGY CLOSED SET (each
 test's oracle NAMED explicitly in source): () SPECIFIED
 (spec/RFC/standard names expected output); () DERIVED
 (independent reference implementation); () METAMORPHIC
 (metamorphic relations between related inputs –
 § . . () generalised); () GOLDEN-MASTER /
 CHARACTERIZATION (past captured behavior –
 § . .); () INVARIANT (universal property must
 hold – sort output sorted + permutation of input); ()
 STATISTICAL (distribution properties for stochastic
 outputs, never specific value); () HUMAN
 (§ . . manual QA). NO other classes. STRUCTURAL
 INDEPENDENCE: oracle's SOURCE of expected value does not
 depend on code being tested. § . . () SELF-
 VALIDATION generalised: paired § . mutation must
 produce a WRONG output; oracle catches it. Composes
 § . . / § . . / § . . () /

11 4 224 11 4 245 1 1

11 4 17

\$. . / \$. . / \$. . /
\$. . / \$. . / \$. .
Classification: universal (\$. .). Propagation
gate CM-COVENANT-114-245-PROPAGATION (literal 11.4.245) +
recommended gate CM-ORACLE-STRATEGY-NAMED-AND-INDEPENDENT +
paired \$. mutation.

Canonical authority: constitution submodule Constitution.md
\$. . . Non-compliance is a release blocker. No
escape hatch – no --test-agrees-with-code , --skip-oracle-
identification , --oracle-is-the-code , --assertion-is-oracle ,
--implicit-oracle-OK flag.

2 0 2 6 0 8 1 5

**\$. . – Reproducible + hermetic
builds + supply-chain integrity at SLSA Build
Level (operator decision,
– –)**

Full anchor per Phase landing brief \$. . .
Every consuming project's build pipeline MUST produce
REPRODUCIBLE, HERMETIC artifacts and MUST meet SLSA Build
Level as the FLEET-WIDE MINIMUM (operator decision,
– – – Level + tracked as backlog per
\$. .). REPRODUCIBLE = same source at same time
on any conformant host produces byte-identical artifacts
(or diff-set fully accounted for by declared non-
determinisms pinned to canonical values). HERMETIC = no
undeclared inputs (every dep version-pinned + lockfile-
hashed, every tool version declared, no unpinned public-
network reach at build, no host-env leakage). SLSA-L =
signed provenance describing source URI + commit hash +
builder identity + invocation + artifact hash, on hosted
tamper-resistant build platform (boba already satisfies
ancillary: SSH-only \$. , rootless podman
\$. . , fixed-tag pins \$. . /

§ . .). Every consumer MAINTAINS docs/security/
 SLSA_LEVEL.md (or per § . . equivalent) recording
 CURRENT L level with cited evidence + TARGET L + tracked
 § . . upgrade items – CURRENT below L
 without tracked item is violation. SUPPLY-CHAIN EXTENSION:
 deps either vendored (checked in hash-verified), from
 provenance-attested source (npm/PyPI/crates.io as they add
 SLSA), OR mirrored through hash-pinning local registry.
 Threat model NARROW: covers source→artifact chain, does
 NOT cover malicious commit landing in source
 (§ . . review) nor compromised dep source nor
 bug-free artifact (§ . . four-layer +
 § . . manual QA). Composes § . /
 § . / § . / § . /
 § . / § . / § . /
 § . / § . / § . /
 § . / § . / § . /
 § . . Classification: universal (§ . .).
 Propagation gate CM-COVENANT-114-246-PROPAGATION (literal
 11.4.246) + recommended gates CM-BUILD-REPRODUCIBILITY +
 CM-BUILD-HERMETICITY + CM-SLSA-L2-PROVENANCE-VERIFIED + CM-
 SUPPLY-CHAIN-DEPS-PROVENANCE-ATTESTED + paired § .
 mutations.

Canonical authority: constitution submodule Constitution.md
 § . . . Non-compliance is a release blocker. No
 escape hatch – no --skip-reproducibility , --allow-network-
 during-build , --no-provenance , --slsa-optional , --below-l2-
 OK , --unpinned-dependency-OK , --host-env-may-leak-into-build
 flag.

§ . . – Composite-output layer-
move must move every layer (research-derived,
– –)

Full anchor per Phase landing brief § . . .
Changes MOVING a system-boundary layer (service migrates hosts, endpoint URL changes, API version bumps, DB schema shifts, message topic renames, ingress path re-routes, storage bucket moves) MUST move EVERY LAYER of the composite output atomically or via explicit compatibility window – NEVER piecemeal. LAYER TAXONOMY per § . . .
baseline: (i) server (handler + route + OpenAPI/proto/gRPC descriptor); (ii) client (call-sites + SDK stubs + integration tests); (iii) infrastructure (LB + ingress + API gateway + CDN + service-mesh + DNS); (iv) data stores (cached keys + DB URL rows + MQ subscriptions); (v) observability (metric labels + span names + log fields + dashboards + alerts + SLOs); (vi) documentation (user + dev + runbook + README + cookbook); (vii) tests (integration + e e + contract per § . . . + load + smoke)...Consumer provides scripts/audit/
layer_move_completeness.sh reporting residual references delta; change proposing move MUST show ZERO residuals OR explicit compatibility-window plan per residual with tracked § . . . items. COMPATIBILITY-WINDOW: old value keeps working alongside new until stated date/version/event, deprecation warnings named per caller, tracked deadline, clean removal at expiry – no "clean up later." [MATERIAL-THIN: Module ("Secondary-Display & External-Output Media Pipelines", whose "Surface lifecycle: the heart of the reopening problem" section covers this domain) supplies the composite-output pattern (multiple co-varying layers, atomic move discipline, surface-swap-not-recreate); the specific web/service layer inventory is generalised from the pattern's shape per § . . . with consumer § . . . DATA supplying the concrete inventory. CITATION AUDIT

- - (\$. .): the three named phrases – "composite-output"/"composite output", "surface-swap-not-recreate"/"surface-swap", and "atomic move discipline" – have ZERO occurrences in ALL NINE tracked corpus modules (instrument control-needed per § . . ()(b) before the absence was reported). They are this anchor's OWN PARAPHRASE of the module-domain, NOT corpus quotations, and MUST NOT be read as quoted source text. The nearest phrase-level in-corpus material on surface swapping is module_35:95 / module_35:102 (.setOutputSurface / BAD_INDEX (6)) – module The module- DOMAIN attribution is plausible and unrefuted; the PHRASE-level attribution is UNVERIFIED.] Composes \$. . / \$. . / \$. . / \$. . / \$. . / \$. . (residual = coverage escape) / \$. . / \$. . Classification: universal (\$. .). Propagation gate CM-COVENANT-114-247-PROPAGATION (literal 11.4.247) + recommended gate CM-LAYER-MOVE-COMPLETENESS-CHECK + paired \$. mutation.

Canonical authority: constitution submodule Constitution.md
 \$ Non-compliance is a release blocker. No escape hatch – no --partial-layer-move-OK , --skip-completeness-check , --residuals-cleaned-up-later , --observability-layer-optional , --doc-move-can-lag flag.

§ 6.0.1. – Flaky-test quarantine +
protected-regression-spec gate via [PROTECTED-
SPEC: ATM-NNN] tag + CODEOWNERS required
reviewer (operator decision,
- -)

Full anchor per Phase landing brief § 6.0.1.
Flaky tests are the SILENT ANTI-BLUFF DECAY vector (every
flake conditions team to ignore red until real regression
is dismissed as "probably flaky, re-run"). Two disciplines
bound: (A) QUARANTINE – mechanical detection (test-history
aggregator flags two-consecutive-run-different-verdict OR
rerun-after-red pattern) moves flaky test to tests/
quarantine/ (per § 6.0.1), opens tracked
§ 6.0.1 stabilisation deadline, blocking suite
runs WITHOUT quarantined test so green means green;
deadline past = § 6.0.1 extension or violation; (B)
PROTECTED-REGRESSION-SPEC – regression tests codifying
known-broken defect fixes (per § 6.0.1 permanent-
guard suite) MAY be tagged # [PROTECTED-SPEC: ATM-NNN]
(target-language comment equivalent) referencing the
workable-item ATM-NNN (§ 6.0.1) they guard, marking
test as load-bearing regression guard whose modification
requires elevated review. Tag applied deliberately per
test (not blanket – that voids the discipline).
CODEOWNERS-REQUIRED-REVIEWER GATE: consuming project's
CODEOWNERS declares modifications to [PROTECTED-
SPEC] -tagged tests OR regression/**/* path scope require
designated required-reviewer approval BEFORE merge;
composes with § 6.0.1 SSH + § 6.0.1 no-force-
push (gate cannot be silently bypassed); audit trail in PR
review thread + commit history. Gate CATCHES change so
reviewer distinguishes STRENGTHEN (approve) vs FIX-DEFECT-
IN-TEST (approve after verifying ATM-NNN's defect not
silently re-introduced) vs WEAKEN/REMOVE (refuse or
§ 6.0.1 escalate) – makes modification VISIBLE not
preventable. Composes § 6.0.1 / § 6.0.1 /

$\$ \dots / \$ \dots$ (ideal empty
 quarantine) / $\$ \dots / \$ \dots /$
 $\$ \dots / \$ \dots / \$ \dots /$
 $\$ \dots / \$ \dots / \$ \dots$

Classification: universal ($\$ \dots$). Propagation
 gate `CM-COVENANT-114-248-PROPAGATION` (literal `11.4.248`) +
 recommended gates `CM-FLAKY-TEST-QUARANTINE-WIRED` + `CM-`
`PROTECTED-REGRESSION-SPEC-GATE` + paired $\$ \dots$ mutations.

Canonical authority: constitution submodule `Constitution.md`
 $\$ \dots$. Non-compliance is a release blocker. No
 escape hatch – no `--rerun-until-green`,
`--quarantine-permanent`, `--protected-spec-review-optional`, `--`
`codeowners-skip`, `--blanket-protect-everything`, `--tag-without-`
`atm-id` flag.

$\$ \dots$ – **Producer $\neq$ oracle $\neq$ gate $\neq$ verifier + flight recorder** (research-derived,
 – –)

Full anchor per Phase landing brief. $\$ \dots$.
 Every load-bearing quality mechanism has FOUR distinct
 roles that MUST be structurally separate – () PRODUCER
 (code computing/emitting/building output); () ORACLE
 (mechanism deciding correctness per $\$ \dots$); ()
 GATE (seam consuming oracle verdict to allow/block/
 proceed); () VERIFIER (mechanism auditing gate's
 decision – paired $\$ \dots$ mutation proving gate can refuse
 a bad output). Collapsing any two produces specific bluff
 class: producer=oracle ($\$ \dots$ oracle-agrees-
 with-code); oracle=gate (refuses to say "cannot decide"
 defaults to allow, $\$ \dots$ () FALSE=NULL);
 gate=verifier (self-auditing gate – $\$ \dots$ self-
 review anti-pattern applied). FLIGHT RECORDER: mechanism
 CAPTURING producer inputs + oracle verdict + gate decision

+ verifier audit as durable replayable artifact
 (structured JSONL event stream per § 11.4.207, content-addressed per § 11.4.209) so post-hoc "why did this go through" has evidence not shrug. Recorder is NOT the gate – a gate reading its own recorder to decide collapses gate + verifier + recorder. DETECTABLE COLLAPSES reviewer per § 11.4.209 scans for: ()
 producer=oracle (test's assertion calls code being tested), () oracle=gate (no "cannot decide" branch, defaults PASS), () gate=verifier (meta-test IS the gate – mutation making gate always-pass ALSO makes meta-test always-pass; meta-test MUST live OUTSIDE gate), () recorder=decision-substrate (gate uses recorder contents as decision input – precedent-based gating not authorised). [MATERIAL-THIN: four-role taxonomy generalised from producer/verifier separation + oracle-problem discussion, + gate mechanics; FLIGHT RECORDER derived from clean-checkout re-run + attached receipts pattern per § 11.4.209.] Composes § 11.4.209 / § 11.4.209 / § 11.4.209 () / § 11.4.209 / § 11.4.209 / § 11.4.209 / § 11.4.209 / § 11.4.209 (producer verifier separation principle – mutual composes per § 11.4.209 reviewer MINOR-) / § 11.4.209 / § 11.4.209
 Classification: universal (§ 11.4.209). Propagation gate CM-COVENANT-114-249-PROPAGATION (literal 11.4.249) + recommended gates CM-ROLE-SEPARATION-DECLARED + CM-FLIGHT-RECORDER-CAPTURED + paired § 11.4.209 mutations.

Canonical authority: constitution submodule Constitution.md
 § 11.4.209. Non-compliance is a release blocker. No escape hatch – no --collapse-producer-oracle, --gate-self-audits, --skip-flight-recorder, --oracle-defaults-pass, --decision-without-recorded-rationale flag.

§ . . – Heuristic-tower signals a primitive defect: when N heuristic layers stack to compensate, the primitive is broken (research-derived, – –)

Full anchor per Phase landing brief § . . . Codebase accumulating TOWER of heuristics compensating for same underlying issue (retry-with-backoff + connection-pool-warmup + DNS-cache-flush + health-check-with-startup-grace – all guarding one flaky DB connection) IS the diagnostic: PRIMITIVE is defective, every layer papers over a manifestation without fixing cause. Extraction line [offset un-resolvable per the § . . .

CITATION NOTE; **MODULE ADDED** – –]:
 heuristic-compensating-for-broken-primitive taxonomy row + ATMOSphere BAD_INDEX (6) compensator-tree case study – resolves to **module** taxonomy row "Heuristic compensating for a broken primitive", verified locators module_35:99 (row header), module_35:102 (the ATMOSphere setOutputSurface() / BAD_INDEX (6) compensator case), module_35:104 (the attack-the-primitive remedy). **OBSERVED CORRELATION** (§ . . . – an observation, NOT a proven mapping): all four offsets cited across § . . . / § . . . / § . . . land on module taxonomy rows under a constant offset of

(→row , →row , →row , →row ; of fit), consistent with module 's taxonomy beginning near offset in the consolidated extraction artifact; that artifact is not tracked in this repository, so the mapping cannot be confirmed here and is not asserted as fact. Recognition signals: (i) each layer justification "we added because sometimes X happens" – reactive; (ii) removing any layer INCREASES failures BUT keeping all shows occasional residual; (iii) no single layer owns responsibility; (iv) layers added at different times by different authors per specific incidents; (v) monotone layer-count growth over

time – debugging session for residual adds layer N+ rather than fixing ..N. DISTINCT from defense-in-depth: defense-in-depth is deliberate + each layer targets DIFFERENT failure class + designed together not accreted + removing layer produces SPECIFIC different failure not "same failure more often." SYSTEMATIC-DEBUGGING RESPONSE on tower depth ≥ : () STOP adding layers; () INVOKE \$. . systematic-debugging on PRIMITIVE; () proceed root-cause / pattern / hypothesis / implementation phases; () fix primitive; () REMOVE layers one at a time verifying with \$. . evidence; () commit each removal per \$. . with git-history evidence citing primitive fix. Composes \$. . / \$. . / \$. . / \$. . (outer layer masks \$. . () false-null) / \$. . / \$. . Classification: universal (\$. .). Propagation gate CM-COVENANT-114-250-PROPAGATION (literal 11.4.250) + recommended gate CM-HEURISTIC-TOWER-TRIGGERS-PRIMITIVE-DEBUG + paired \$. . mutation.

Canonical authority: constitution submodule Constitution.md \$. . . Non-compliance is a release blocker. No escape hatch – no --add-another-heuristic-layer , --skip-primitive-debug , --tower-is-defense-in-depth-without-justification , --reactive-mitigation-OK , --layer-per-incident flag. 2.5.1

\$. . – Byte-identical-fork prohibition + role-as-data-pack (research-derived, – –)

Full anchor per Phase landing brief \$. . . When project needs would be met by TWO NEAR-IDENTICAL CODE COPIES differing only in configuration/data/role-specific

parameters (prod + staging with hardcoded hostnames; reader + writer with different flag values; web + mobile with different feature flags), copies MUST NOT be maintained as parallel forks. Instead ONE code copy retained + role-specific values extracted into DATA PACK (role-config file / env var set / capability descriptor / feature-flag bundle) consumed by the one code to instantiate as either role. Extraction line [offset un-resolvable per the \$. . CITATION NOTE; **MODULE ADDED** - - - resolves to **module** taxonomy row "Byte-identical-fork", verified locators `module_35:127` (row header), `module_35:130` (the SonarQube-Factory case), `module_35:132` (the "roles become data-only packs you reference, never repos you fork" remedy)]: byte-identical-fork taxonomy row + SonarQube-Factory case (`Application/main.kt` byte-for-byte the caching-proxy launcher, still in `package net.milosvasic.factory.proxy.caching.application` - "no SonarQube code"; fix = role-real acceptance test, "roles become data-only packs you reference, never repos you fork"). DETECTION: reviewer or mechanical scan finds byte-identical fork when `diff` shows every difference is a config value (hostname/port/flag/label/env/credential-ref) NOT behavioral, common git-history ancestor, > % byte identity; per \$. . threshold-DATA. EXTRACTION: () IDENTIFY configuration surface; () EXTRACT to data pack per role; () REFACTOR single code copy to consume at instantiation (constructor injection / factory / config-loader-at-startup); () DELETE one fork; () VERIFY per \$. . both role instantiations produce same observable behavior their forks previously produced (\$. . characterization before refactor); () commit extraction per \$. . with git-history evidence. LEGITIMATE FORKS: temporary migration (tracked \$. . end date), experimental / A-B (bounded end date), truly divergent behavior (copies differ in

algorithm/protocol/data-model not just config). Composes
 § 11.4.242 . . . / § . . . / § . . . (decoupling
 – role-configurable code shared across consumers) /
 § . . . (data-pack location) / § . . . /
 § . . . / § . . . / § . . . /
 § . . . / § . . . Classification: universal
 (§ . . .). Propagation gate CM-COVENANT-114-251-
 PROPAGATION (literal 11.4.251) + recommended gate CM-NO-
 BYTE-IDENTICAL-FORK + paired § . . . mutation.

Canonical authority: constitution submodule Constitution.md
 § Non-compliance is a release blocker. No
 escape hatch – no --fork-for-role-OK , --skip-data-pack-
 extraction , --copy-and-modify-config , --duplicate-code-per-
 environment , --near-identical-forks-are-fine flag.

§ . . . – Fail-closed-on-dangerous-combination for mutating / credential surfaces (research-derived, – –)

Full anchor per Phase landing brief §
 Every code path COMBINING ≥ dangerous capabilities
 MUST FAIL CLOSED – refuse operation unless every
 capability's precondition verifiably satisfied – NEVER
 fail open on ambiguous/missing/unresolvable input.
 Security-engineering generalisation of § . . .
 guard-honesty applied to code paths where false-negative
 blast radius dwarfs operator inconvenience of false-
 positive. DANGEROUS-COMBINATION TAXONOMY: () MUTATION of
 shared resource (DB / filesystem / container state /
 config), () UNTRUSTED INPUT (user string / network
 bytes / external env var / unverified file), ()
 CREDENTIAL ACCESS (secret read / authenticated API call /
 signing key use), () EXTERNAL SIDE EFFECT (email /
 webhook / MQ publish / payment endpoint), () SHELL /

EXEC (subshell / binary with user-influenced args / template eval with user values), () IRREVERSIBLE (delete / drop / truncate – \$. . . force-push already absolutely forbidden). Any path combining ≥ = dangerous-combination, MUST fail closed. INVARIANT: () VERIFY every precondition (authorization + input shape + credential presence + target identity); () REFUSE when ANY unverifiable with ERROR EXIT naming specific unresolved precondition; () EMIT captured evidence per \$. . . of refusal; () NEVER default to allow/proceed/retry-with-relaxed on ambiguity – \$. . . conservative-safe default = REFUSE. FAIL-OPEN ANTI-PATTERNS refused: `catch { /* ignore */ }, credential = credential || default , if (!valid) { log("warn"); return proceed(input); } , target = user_input || default_target , unbounded retry { dangerous(); } . Distinct from graceful degradation. [MATERIAL-THIN: -class dangerous-combination taxonomy is a $. . . generalisation of the fail-closed pattern's shape applied to specific dangerous surfaces; closed-set formalisation, not literal per-class enumeration.] Composes $. . . / $. . . / $. . . (credentials specialisation) / $. . . (operator escalation when refuse stuck) / $. . . / $. . . (max instantiation) / $. . . / $. . . / $. . . Classification: universal ($. . .). Propagation gate CM-COVENANT-114-252-PROPAGATION (literal 11.4.252) + recommended gate CM-DANGEROUS-COMBINATION-FAIL-CLOSED + paired $. . . mutation.`

Canonical authority: constitution submodule Constitution.md \$ Non-compliance is a release blocker. No escape hatch – no `--fail-open-on-error` , `--swallow-exception` , `--default-missing-credential` , `--proceed-on-ambiguous-input` , `--retry-until-succeeds` flag.

§ . . – Idempotency under retry + DB-level durable uniqueness guard (research-derived, - -)

Full anchor per Phase landing brief §

Every retryable operation (network calls with auto-retry, at-least-once MQ consumers, cron with overlapping schedules, HTTP endpoints with client-retry, DB transactions with retry-on-serialisation-failure, deploy pipelines with rerun capability) MUST BE IDEMPOTENT: N executions produce same observable end state as one. Idempotency ENFORCED at LOWEST LAYER where "same end state" is durably provable – typically DATABASE UNIQUENESS CONSTRAINT on caller-supplied idempotency key, NEVER only application-layer check-then-insert (races). Prevents DOUBLE-CHARGE / DOUBLE-EMAIL / DOUBLE-DELETE / DOUBLE-EFFECT bugs. IDEMPOTENCY-KEY DISCIPLINE: every retryable op accepts (or generates + passes forward) unique identifier naming THIS logical op; persistence layer has UNIQUE INDEX / CONSTRAINT on key column; second attempt REFUSED by DB at durable level not application logic reading-then-writing; refusal caught by caller as "duplicate – already succeeded" = SUCCESS. ANTI-PATTERNS: (i) check-then-insert without unique index races; (ii) app-level dedup table read without spanning-transaction; (iii) key scoped to app instance (in-memory not shared); (iv) effect persisted BEFORE key (crash between = double-effect on retry); (v) no key at all + "the effect is 'set X to V' so redoing is fine" (safe for pure sets, NOT for INSERTs / MQ emissions / external-system side effects); (vi) idempotency-key expires – converts to time-bounded-idempotent = § . . mis-labelling if caller expects unbounded. CHAOS VERIFICATION per § . . : () SIMULTANEOUS RETRIES (same key from concurrent callers), () MID-OPERATION CRASHES (kill receiver between "check exists" and "insert effect"), () RETRY AFTER APPARENT SUCCESS (caller believes fail, receiver

persisted, retry MUST recognise duplicate). External-system side effects (Stripe/PayPal charging card + local DB write): idempotency key MUST pass through to external; if external does not support, op is not fully idempotent – label honestly per § 11.4.253. Composes § 11.4.253 / § 11.4.253 / § 11.4.253 / § 11.4.253 / § 11.4.253. Classification: universal (§ 11.4.253). Propagation gate CM-COVENANT-114-253-PROPAGATION (literal 11.4.253) + recommended gates CM-IDEMPOTENCY-DB-LEVEL-UNIQUE-GUARD + CM-IDEMPOTENCY-CHAOS-VERIFIED + paired § 11.4.253 mutation.

Canonical authority: constitution submodule Constitution.md § 11.4.253. Non-compliance is a release blocker. No escape hatch – no --check-then-insert-OK, --application-dedup-suffices, --skip-idempotency-chaos-test, --claim-idempotent-without-key, --expire-idempotency-marker-silently flag.

§ 11.4.254 – **Boot-time invariant assertion + capability matrix (research-derived, – –)**

Full anchor per Phase 3 landing brief § 11.4.254. Every service / process / long-running executable MUST assert BOOT-TIME INVARIANTS at startup – required env vars present + valid; required config files parseable + declared schema version matches; required deps reachable; required credentials present + not-expired; required schema migrations applied; required feature-flag states declared; required host resources present (disk / memory / thread limits per § 11.4.254 / port availability) – and FAIL FAST + LOUD if any violated, BEFORE accepting request / performing work. Extraction line [offset un-resolvable per the § 11.4.254 CITATION

NOTE; **MODULE ADDED** - - - resolves to **module** taxonomy row "Config-present-but-unwired", verified locators `module_35:113` (row header), `module_35:116` (the case), `module_35:118` (the boot-time cipher_version remedy)] (Catalogizer SQLCipher case: `catalog.db` 's first bytes are ASCII SQLite format 3 - plaintext despite "encrypted at rest"; fix = boot-time-check-that-fails-closed verifying cipher_version at startup) + line **[MODULE ADDED**

- - - resolves to **module** taxonomy row "Stubbed-core-behind-claim", verified locators `module_35:120` (row header), `module_35:123` (the Catalogizer nil,nil-providers case), `module_35:125` (the `IsImplemented()` capability-matrix remedy)] (stubbed-core-behind-claim + explicit capability matrix with every adapter reports `IsImplemented()` /attempted-vs-not + challenge that fails against no-op stub). Prevents "service silently starts with broken invariant + serves requests that fail deep in request path hours after boot." INVARIANT CLASSES: (i) environment, (ii) configuration, (iii) dependencies, (iv) credentials, (v) schema, (vi) capabilities, (vii) host. FAIL-FAST-LOUD: () EMIT structured error to stderr + captured log naming specific failed invariant; () EXIT non-zero (traditionally `EX_CONFIG` for config, consumer supplies per \$. .); () DO NOT accept requests / begin main..event loop / bind listening port. Service booting partially + serving requests-that-will-fail is worse than refusing to start - partial-boot masquerades healthy long enough for load balancer to route traffic. Dual to \$. . fail-closed on dangerous combinations. CAPABILITY MATRIX: consumer's `docs/services/CAPABILITY_MATRIX.md` (per \$. . .) declares per service (i) boot-time invariants, (ii) runtime capabilities OFFERED, (iii) runtime capabilities REQUIRED from other services. Regenerated / verified on every commit per \$. . docs-chain + \$. .

fingerprint. Undeclared endpoints emitted = violation;
 silent upstream deps not declared = violation (invisible
 to \$. . . contract testing). Composes
 \$. . . / \$. . . / \$. . . /
 \$. . . / \$. . . / \$. . .
 Layer (boot IS runtime-on-clean-target check) /
 \$. . . / \$. . . (contract tests
 consume matrix) / \$. . . / \$. . . /
 \$. . . Classification: universal (\$. . .).
 Propagation gate CM-COVENANT-114-254-PROPAGATION (literal
 11.4.254) + recommended gates
 CM-BOOT-TIME-INVARIANT-ASSERTED + CM-CAPABILITY-MATRIX-PRESENT-
 AND-IN-SYNC + paired \$. . . mutations.

Canonical authority: constitution submodule Constitution.md
 \$ Non-compliance is a release blocker. No
 escape hatch – no --skip-boot-time-assertion , --boot-partial-
 OK , --default-missing-env-silently , --capability-matrix-
 optional , --service-with-undeclared-dependency-OK flag.

\$. . . – Mandatory HelixTranslate
 canonical translation pipeline (User mandate,
 – –)

Forensic anchor – verbatim user mandate
 (– –): "Every project that inherits this
 constitution MUST use HelixTranslate
 (git@github.com:HelixDevelopment/HelixTranslate.git) – via
 a universal, reusable pipeline – to generate ANY
 translations it needs. This is a hard rule/constraint."

ALL translations of ANY content – UI strings, documents,
 articles, CVs, READMEs, marketing copy, in-app messages,
 release notes, legal text, ANY human-language string
 rendered in a non-source language – MUST be produced by
HelixTranslate (git@github.com:HelixDevelopment/

HelixTranslate.git) driven through the **canonical universal translation pipeline** shipped in this submodule at `scripts/translation/` (`translate-pipeline.sh` + `render-articles.sh`). This is a hard, non-negotiable constraint.

The following are FORBIDDEN and constitute a § . . . violation:

(a) **Hand-translation** – no human-authored, manually edited, or "good enough" hand-written translation of any content, ever. The engine is the single source of every translated string.

(b) **Other engines** – no Google Translate, DeepL, raw ChatGPT/Claude prompting, OS translation APIs, library i n auto-translators, or any engine other than HelixTranslate's `unified-translator` . There is exactly one engine.

(c) **Silent fallback** – no quiet substitution of the source string, no empty/partial output passed off as a translation, no swallowed engine error, no "absence-of-error" PASS (§ . . . / § . . .). On exhaustion of all configured providers the pipeline MUST FAIL LOUD (non-zero exit, destination NOT written) so the caller detects it; a missing translation is a hard error, never a degraded-but-shipped string. Provider-to-provider failover WITHIN HelixTranslate (primary → fallback, with logged, evidenced retries) is NOT a silent fallback and is permitted; falling back to ANY non-HelixTranslate source is forbidden.

Mandatory validation. Every generated translation MUST be validated for accuracy with real.. (physical) captured evidence per § . . . / § . . . / § . . . – no metadata-only / config-only / grep-without-runtime PASS (§ . . . / § . . .), no false results, no bluff of any kind (§ . . .). The per-file run log (provider used, byte counts, retries, success/

failure) is the evidence artifact; an unvalidated or unevidenced translation is treated as ABSENT.

Explicit provider routing. Provider + model routing MUST be explicit and recorded – never implicit/defaulted-in-the-dark. The canonical routing is provider `groq` / model `llama-3.3-70b-versatile` with documented failover to provider `mistral` / model `mistral-large-latest`; the provider that actually produced each output MUST be recorded in the per-file log. Any deviation from the canonical routing MUST be stated explicitly in configuration and logged.

Canonical pipeline location. `scripts/translation/translate-pipeline.sh` (universal Markdown / plain-text + article-frontmatter-preserving translator) and `scripts/translation/render-articles.sh` (article-source → HTML fragment renderer). Both are reusable verbatim across projects; the engine binary is supplied by the consuming project via `HELIX_TRANSLATE_BIN` and built from the HelixTranslate repository.

Classification: universal (`$ . . . . .`) – a platform-neutral, language-neutral translation discipline reusable by ANY project; the consuming project supplies its concrete source files, target languages, and the HelixTranslate engine build per `$ . . . . .`. Composes `$ . . . / $ . . . / $ . . . / $ . . . / $ . . . / $ . . . / $ . . .`.

Propagation gate `CM-COVENANT-114-255-PROPAGATION` (literal `11.4.255`) + recommended gate `CM-MANDATORY-HELIXTRANSLATE-PIPELINE` + paired `$ . . .` meta-test mutation (gate-code = separate work item).

Rationale. A single, evidenced, fail-loud translation engine guarantees that every translated string in the fleet has a known provenance, a reproducible pipeline, and a captured accuracy proof – eliminating the silent drift, inconsistent quality, and untraceable "where did this

string come from" failures that hand-translation and ad-hoc engines produce. One engine, one pipeline, one evidence trail.

Canonical authority: this Constitution.md § . . in the HelixConstitution submodule. All consuming projects restate + cite via § . . inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no `--allow-hand-translation`, `--other-engine-OK`, `--skip-translation-validation`, `--silent-fallback` flag exists.

§ . . – Mandatory independent per-language translation review (User mandate, – –)

Forensic anchor – verbatim user mandate

(– –): "Quality of translations for all languages MUST BE ALWAYS properly checked and reviewed! We MUST CREATE independent, specialized translations review agents for all languages we support which will be performing full testing, review, validation and verification for all translated content! This MUST BE mandatory part of our root constitution (Submodule) and followed as mandatory constraint / rule by this and all our projects implementing the constitution Submodule."

Every translated artifact produced under § . . MUST, before it is shipped / committed / deployed, be INDEPENDENTLY REVIEWED for quality by a per-language review agent. This is a hard, non-negotiable gate that composes with – and never replaces – the § . . generation mandate.

Requirements:

(a) **Independence.** The reviewer MUST be a different model/provider than the one that produced the translation (e.g. translator `mistral/mistral-large-latest` → reviewer `groq/llama-3.3-70b-versatile`, or any distinct verified model). A model reviewing its own output is NOT an independent review and is forbidden.

(b) **Per-language specialization.** The reviewer MUST act as a NATIVE speaker / linguistic reviewer of the SPECIFIC target language, evaluating that language's grammar, orthography, and correct script (Cyrillic for ru/be/kk, Arabic for ar/fa, Han/Kana/Hangul/Devanagari for zh/ja/ko/hi, Latin for sr/de/es/fr/tr, etc.).

(c) **Full criteria.** Each review MUST judge: accuracy/fidelity, fluency/naturalness, completeness (nothing omitted or invented), correct script & orthography, and absence of untranslated source leftovers (proper nouns / URLs / code excepted).

(d) **Strict machine-readable verdict + evidence.** Each review MUST emit a structured PASS/FAIL verdict with per-criterion scores and an issues list, persisted as a physical evidence artifact. (`$ . . . / $ . . . / $ . . .`). A FAIL – or an ERROR / unavailable-reviewer – BLOCKS the content; only a real PASS may ship. No metadata-only / absence-of-error PASS (`$ . . . / $ . . .`); no bluff (`$ . . .`).

(e) **Remediation loop.** A FAIL routes the artifact back through the `$ . . .` pipeline (re-translate / multipass-polish) and is re-reviewed, iterating until PASS – never hand-patched to pass (hand-translation remains forbidden under `$ . . .`).

Canonical tooling. `scripts/translation/review-translations.sh` (per-language review driver) + `scripts/translation/review_translation.py` (independent reviewer engine emitting strict JSON verdicts). Reusable verbatim across projects;

the consuming project supplies its source/translation file pairs and a reviewer provider DISTINCT from its translator.

Classification: universal (§ 11.4.255) – a platform-neutral, language-neutral translation-quality-assurance discipline reusable by ANY project that ships translated content; composes § 11.4.256 (the translation it gates) / § 11.4.257 / § 11.4.258 / § 11.4.259 / § 11.4.260 (rock-solid evidence). Propagation gate CM-COVENANT-114-256-PROPAGATION (literal 11.4.256) + recommended gate CM-MANDATORY-TRANSLATION-REVIEW + paired § 11.4.261 meta-test mutation (gate-code = separate work item).

Rationale. § 11.4.262 guarantees a translation's provenance and a fail-loud pipeline, but not its linguistic quality. An independent, native-level, per-language reviewer with a captured PASS/FAIL verdict is what turns "a string was produced" into "a string was verified correct" – making translation quality an auditable fact rather than an assumption.

Canonical authority: this Constitution.md § 11.4.263 in the HelixConstitution submodule. All consuming projects restate + cite via § 11.4.264 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch – no `--skip-translation-review`, `--self-review-OK`, `--accept-unreviewed-translation` flag exists.

\$. . – Comprehensive user-manual
+ guide + FAQ coverage: every component /
service / feature / user-visible workflow
ships an operator-usable manual, task-oriented
guide, and FAQ, always in sync, no orphan
capability (BACKGROUND :: REMINDER ::
IMPORTANT operator mandate,
- -)

Verbatim operator mandate (- - , Point
): *"Whole project and all work done MUST BE fully
covered with full user manuals and guides and properly
created FAQs."*

Compact summary: every consuming project MUST cover the
WHOLE project (every component, service, subsystem,
binary, container, plugin, feature, and user-visible
workflow) with three cooperating documentation classes –
(a) a complete USER MANUAL per component/service/product,
(b) task-oriented GUIDES for every user-facing workflow,
and (c) properly-created FAQs derived from real operator/
QA/end-user questions – always in sync with the shipped
code, always exported per \$. . , always reachable
from README per \$. . , always regenerated via
Docs-Chain per \$. . , and always tracked as
workable items when incomplete per \$. . ; a
component / service / feature that ships without its trio
of docs (manual + guide + FAQ) is a \$. .
violation of \$. . PASS-bluff severity at the
documentation layer (the feature ships to a user who
cannot use it because nothing tells them how; the \$. .
anti-bluff covenant extended to documentation
completeness). \$. . GENERALISES \$. .
(script docs – a single class of executable artifact) and
\$. . (per-feature Status set – a status not a
manual) into the full three-class documentation floor for
every user-facing surface, and BINDS \$. .
(README entry point) + \$. . (Docs-Chain sync)

+ § . . (four-format export) into the always-in-sync guarantee.

(A) COMPREHENSIVE COVERAGE – no orphan capability. Every component / service / subsystem / binary / plugin / user-visible feature enumerated in the project's shipped surface (the § . . feature Status roster or its equivalent SSoT) MUST map to at least ONE user manual (component-level operator reference), at least ONE task-oriented guide per user workflow that surface exposes, and one FAQ entry per real recurring question captured from operators / QA / end-users. A capability the project ships that no manual documents, no guide covers, and no FAQ addresses is UNSHIPABLE regardless of how well the code works – the user has no path to use it.

(B) THREE DISTINCT DOCUMENT CLASSES, none substitutes another. (a) USER MANUAL – the component's operator reference: what it is, what it does, its inputs/outputs/env vars/config/ports/dependencies, its start/stop/status commands, its logs/metrics/health surface, its failure modes and their remediation; complete enough that an operator can run + supervise + debug it without reading the source. (b) GUIDES – task-oriented, workflow-shaped documents, one per user-facing workflow (setup, configuration, integration, migration, troubleshooting, security-hardening); each guide walks the user through a real journey, with prerequisites, step-by-step commands, expected output at each step, verification checkpoints, and rollback path. (c) FAQs – questions PROPERLY CREATED from real recurring signal (operator questions, QA reports, end-user tickets, agent-found-defect narratives), NEVER fabricated marketing FAQ; every entry cites the question source class + resolution + link to the manual/guide section that owns the durable answer.

(C) ALWAYS IN SYNC. Every doc carries the § . . revision header; Docs-Chain (§ . .) is bound at

commit / build / constitution-pull write-seams so an edited component whose docs are stale REFUSES the commit; four-format export per § . . . (.md + .html + .pdf, + .docx where the class warrants – the § . . . exception); every doc reachable transitively from README per § FAQ entries that describe now-fixed defects are RETAINED with a "resolved in vN.N.N" note (they answer the user asking "why did this happen to me"), never silently deleted.

(D) TRACKABLE COMPLETION. Missing docs are § . . . items (research/kicked-off-work completion mandate) – an undocumented capability is an un-wired requirement whose loss is FORBIDDEN; the workable-items DB (§ . . . /.) carries per-surface docs-status rows so "what is documented / what is not" is a live query, never a periodic manual audit.

(E) HONESTY. Where a genuine gap exists (a new capability whose docs are in flight), it is HONESTLY MARKED at both the README entry point + the component's Status.md (§ . . .) as [DOC-PENDING: <tracker-id>] per the § . . . provenance-marker discipline, NEVER silently omitted, NEVER faked with a placeholder paragraph that says nothing (a fake-doc is a § . . . PASS-bluff at the documentation layer – an operator reads it, learns nothing, and thinks they read the doc).

Honest boundary (§ . . .). § . . . mandates the DOCUMENTATION EXISTS, is IN SYNC, and is REACHABLE; it does NOT verify the doc is CORRECT in every semantic detail – that stays § . . . (latest-source cross-reference) + § . . . (exhaustive review including doc content) + § . . . (any operator-found doc gap is a coverage-escape triggering a § . . . – family audit). "Comprehensive" is bounded to the shipped-surface enumeration; an internal implementation detail with no user-facing exposure has no user-manual obligation

(documentation of it lives under § . . script/
source class or design docs, not this anchor).

Classification: universal (§ . .). Composes

§ . . / § . . / § . . /
§ . . / § . . / § . . /
§ . . / § . . / § . . /
§ . . / § . . / § . . /
§ . . / § . . / § . .

Propagation gate CM-COVENANT-114-257-PROPAGATION (literal
11.4.257) + recommended gates CM-USER-MANUAL-PER-COMPONENT +
CM-GUIDE-PER-USER-WORKFLOW + CM-FAQ-FROM-REAL-QUESTIONS +
paired § . meta-test mutations.

Canonical authority: constitution submodule Constitution.md
§ . . . Non-compliance is a release blocker. No
escape hatch – no --skip-user-manual , --guide-later , --faq-
optional , --capability-without-docs-OK ,
--fake-doc-placeholder-OK , --source-of-truth-is-code-comments
flag. 2.5.8

§ . . – Architectural + data-flow
+ state-machine + sequence-diagram coverage:
every project ships accurate, exported,
machine-derivable-or-honestly-authored
diagrams embedded into the docs that describe
the system (BACKGROUND :: REMINDER ::
IMPORTANT operator mandate,
– –)

Verbatim operator mandate (– – , Point
): "Whole project MUST BE fully covered with properly
created graphs, diagrams, schemes which will be
incorporated into all existing documentation."

Compact summary: every consuming project MUST cover the
WHOLE project with properly created graphs / diagrams /

schemes across the FOUR canonical diagram classes – (a) SYSTEM ARCHITECTURE (services, containers, ports, dependencies, deployment topology), (b) DATA FLOW (how information moves between components, storage, external systems), (c) STATE MACHINE (lifecycle of every stateful entity: workable items, orchestrator sessions, long-ops, agents, deploys), (d) SEQUENCE (per critical user workflow and per critical cross-component protocol) – INCORPORATED into the existing documentation surface (README, user manuals per § 1.1.1, guides, Status docs per § 1.1.2, design docs per § 1.1.3 / § 1.1.4) at the natural point of use, NEVER a lone diagram directory nobody opens; every diagram is (i) exported in a scale-appropriate open format (SVG for structural diagrams; PNG for raster fallback; the source in a diffable text form – Mermaid, PlantUML, D, or a diffable SVG source – § 1.1.5 open-first, never a proprietary-tool-only master), (ii) rendered non-blank + non-degenerate + accurate (the § 1.1.6 ()) analyzer discipline applied to diagrams – a rendered blank / mislabelled / truncated diagram is a § 1.1.7 / § 1.1.8 bluff at the visual-doc layer), (iii) reachable from README per § 1.1.9, and (iv) in sync via § 1.1.10 Docs-Chain so a code change that invalidates a diagram REFUSES the commit until the diagram is updated. § 1.1.11 STRENGTHENS § 1.1.12 (design-library catalogue – components not architecture) and § 1.1.13 (screen catalogues – UI IA not system topology) from design-side coverage into whole-project documentation-side diagram coverage.

(A) FOUR-CLASS COVERAGE. (a) ARCHITECTURE – the deployed / deployable topology: every service, container, binary, port, external dependency, host boundary, network boundary, storage boundary; kept true to the actual `docker-compose.yml` / deployment manifest / running system. (b) DATA FLOW – how information (user request, event,

credential, artifact, evidence) traverses the system:
source → transformation → sink; each edge annotated with
format + protocol + trust boundary. (c) STATE MACHINE –
the lifecycle of every stateful entity the system manages:
workable items (\$. . .), long-ops
(\$. . .), orchestrator sessions
(\$. . .), agents (\$. . .), deploy
artifacts (\$. . .), builds; transitions labelled
with the triggering event and the guard that gates it. (d)
SEQUENCE – one diagram per critical user workflow and per
critical cross-component protocol; actors + messages +
return paths + failure branches, not only the happy path.

(B) **INCORPORATION – never orphan.** Each diagram lives at
the point of use (architecture diagram in README + docs/
ARCHITECTURE.md ; per-service data-flow embedded in that
service's user manual per \$. . . ; state-machine
embedded in the doc that owns the state; sequence diagram
in the guide for the workflow it depicts); a
docs/diagrams/ gallery may exist as an index, but the
load-bearing appearance is EMBEDDED into the doc where the
reader needs it, per \$. . . reachability.

(C) **OPEN + DIFFABLE.** Sources in diffable text (Mermaid /
PlantUML / D / diffable SVG), NEVER proprietary-tool-
only masters (\$. . .); rendered outputs committed
alongside sources so a reader without the render toolchain
still sees the diagram; renderers \$. . . -probed
before use, absent-tool ⇒ \$. . . SKIP-with-reason
not a fake output (\$. . . export-wave discipline
applied here).

(D) **ACCURACY VERIFICATION.** Every diagram is rendered non-
blank + non-degenerate + accurate: a probe render exists
(a Mermaid parse-check; an SVG geometry check; a
\$. . . () golden-good vs golden-bad-with-drift
fixture for the derived-from-code diagrams); a diagram
whose depicted service list, port map, or state set

diverges from the SSoT (compose file, workable-items schema, orchestrator registry) is caught by a § . . -fingerprint diff on the SSoT + a stale-diagram audit run at pre-build; the Docs-Chain (§ . .) refuses a commit that changes an SSoT without an accompanying diagram refresh where a mapped diagram exists.

(E) DERIVED-WHERE-POSSIBLE. Prefer derived / generated diagrams over hand-drawn where the source data exists: a compose-file-driven architecture diagram, a schema-driven state machine, a route-table-driven sequence; hand-drawn diagrams (for user-journey narratives, conceptual overviews) are LABELED as such with an authored-date and reconciled against the SSoT at each release round.

Honest boundary (§ . .). § . . mandates the diagrams EXIST, are ACCURATE against their SSoT, are EMBEDDED, and are RENDERED HONESTLY; it does NOT claim every possible diagram class is required for every project (an internal library with no runtime deployment has no ARCHITECTURE diagram in the deployment sense – it ships its module-dependency diagram in that class's place). Bounded to the WHOLE-PROJECT coverage the mandate names.

Classification: universal (§ . .). Composes

§ . . / § . . / § . . /
 § . . () /
 § . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . . Propagation gate CM-COVENANT-114-258-PROPAGATION (literal 11.4.258) + recommended gates CM-ARCHITECTURE-DIAGRAM-PRESENT-ACCURATE + CM-DATAFLOW-DIAGRAM-PER-SERVICE + CM-STATEMACHINE-DIAGRAM-PER-STATEFUL-ENTITY + CM-SEQUENCE-DIAGRAM-PER-CRITICAL-WORKFLOW + CM-DIAGRAM-EMBEDDED-NOT-ORPHAN + paired § . . mutations.

Canonical authority: constitution submodule Constitution.md

\$ Non-compliance is a release blocker. No escape hatch – no --skip-architecture-diagram , --diagram-in-orphan-directory , --proprietary-master-only , --stale-diagram-OK , --dataflow-implied-from-code , --sequence-later flag.

\$ – README quality-status badge row: at the top of every project README, a comprehensive row of quality badges with a closed green→amber→red vocabulary, each backed by machine-derived source data, illustrating quality and production-readiness in one glance (BACKGROUND :: REMINDER :: IMPORTANT operator mandate, – –)

Verbatim operator mandate (– – , Point):

"We MUST ADD at the top of the main README all possible badges with coloring from green (ok) to red (not ok) for all possible areas of work done to the project illustrating quality of work done and how close are we to deploying this project to full production."

Compact summary: every consuming project MUST render, at the TOP of the main README (immediately below the H , above the introduction – the reader's first visual signal), a comprehensive BADGE ROW carrying every applicable quality signal with a CLOSED color vocabulary from GREEN (ok) through AMBER (attention) to RED (not ok) – plus a "production-readiness" gauge that summarises how close the project is to full production deployment – with EVERY badge backed by MACHINE-DERIVED source data (never a hand-typed color-of-the-day, \$ no-guessing at the status-report layer); a project whose README lacks the badge row, or carries badges whose colors are fabricated / stale / not backed by real source data, is a \$ violation of \$ PASS-bluff

severity at the reader-first-impression layer (an operator glances at the README and gets a green picture of a broken project). § . . . EXTENDS § . . . (README doc-link section) with a symmetric BADGE section at the same first-visible-position privilege, and BINDS § . . . (Docs-Chain sync – badges regenerate on every state change) + § . . . (fingerprint-driven sync – badge data derived from the SSoT the badge refers to) + § . . . (guards must assert the real condition – a badge is a guard for the reader's eye).

(A) CLOSED COLOR VOCABULARY. GREEN = healthy at target (all constraints met, no findings); AMBER = attention (metric on-target but degrading, or off-target within acceptable band, or blocked on a tracked follow-up); RED = not ok (metric off-target beyond acceptable band, a § . . . gate failing, a § . . . uncovered guard, an open release-blocker); optional GRAY = N/A (badge class does not apply to this project) with an inline note explaining why. NO other colors, NO other verdicts, NO "yellow" or "orange" or "beta" gradations – the closed vocabulary is the operator's cognitive contract with the badge row.

(B) BADGE CLASS FLOOR – every applicable area covered. The MINIMUM set every consuming project ships (adding project-specific badges as consumer-supplied § . . . DATA is welcome, subtraction is not): () BUILD status (last build green/red, artifact fingerprint); () TEST-BREADTH per § . . . /§ . . . (unit/integration/e e/full-automation/security/performance/anti-bluff – one badge per type, or a rollup with a click-through detail per § . . .); () CODE COVERAGE per § . . . (measured %, green ≥ target, amber ≥ ratchet-floor, red below); () SECURITY POSTURE per § . . . SonarQube scanner + supply-chain integrity per § . . . (reproducible + hermetic

builds + SLSA Build Level as fleet-wide minimum);
 () DOCUMENTATION completeness per § . .
 (fraction of shipped surfaces with the required manual +
 guide + FAQ trio); () DIAGRAM completeness per
 § . . (fraction of shipped surfaces with the
 required diagram set, non-stale); () LIVE HEALTH where
 applicable (uptime / SLO / recent-crash-rate); () OPEN
 DEFECTS per § . . (Queued + In-progress +
 Reopened counts, most-reopened highlighted per
 § . .); () SUPPLY-CHAIN integrity level per
 § . . (the SLSA Build Level tracked in docs/
 security/SLSA_LEVEL.md , fleet-wide minimum L); ()
 ZERO-SHORTCOMINGS status per § . . (audit-sweep
 verdict); () MACHINE-EVIDENCE COVERAGE per
 § . . (fraction of PASS claims with captured
 evidence paths); () PRODUCTION READINESS gauge (see
 clause (D)).

(C) MACHINE-DERIVED SOURCE DATA. Every badge's color +
 value is COMPUTED from a live source of truth, NEVER hand-
 typed. Every badge PROVENANCE – the exact source path/
 query/build id – is embedded in the badge tooltip OR a
 companion docs/BADGES.md provenance table. A badge whose
 provenance is hand-typed fails § . . .

(D) PRODUCTION-READINESS GAUGE. A composite badge – GREEN
 when every clause-(B) badge is GREEN AND no
 § . . manual-QA blocker is open AND no
 § . . production-readiness criterion is unmet;
 AMBER when one or more badges are AMBER but no RED and no
 blocker; RED when any badge is RED or any blocker is open.
 The gauge is the operator's one-glance answer to "can we
 deploy this to production?" and its color is BINDING on
 the release decision – a RED gauge is a release-blocker
 (§ . . layer + § . . huge-
 blocker family + § . . manual-QA +
 § . . QA-deploy-readiness).

(E) ALWAYS IN SYNC. Badges regenerate on every state change (build completion, test run, defect open/close, doc change, release cut) via § 4.11.1. Docs-Chain at the same write-seams that regenerate the rest of the doc surface; a stale badge (color no longer reflects source data) is a § 4.11.4.201 in-sync violation at the reader-facing layer.

(F) BADGE SELF-VALIDATED. The badge-computer is a guard (§ 4.11.4.202): golden-GREEN fixture MUST render GREEN, golden-RED fixture MUST render RED, golden-AMBER MUST render AMBER; a badge-computer that PASSes its golden-RED fixture is the § 4.11.4.203 bluff itself and a release-blocker. Paired § 4.11.4.259 mutation flips the golden-RED source data and confirms the badge flips.

Honest boundary (§ 4.11.4.204). § 4.11.4.205 mandates every applicable badge is PRESENT + accurately derived; it does NOT claim green badges prove the project is bug-free (badges are necessary-not-sufficient, mirroring the § 4.11.4.206 metric-validation family); a project whose badge row is HONESTLY red is compliant with § 4.11.4.207 (the reader gets the truth) even while it fails other release gates. A badge class that GENUINELY does not apply is GRAY-labeled with a reason, never omitted (silence-as-badge is a § 4.11.4.208 () FALSE-NULL – the reader assumes green because they see nothing red).

Classification: universal (§ 4.11.4.209) – badge PALETTE + AT-TOP-OF-README placement + closed color vocabulary + machine-derived + production-readiness gauge are universal; the SPECIFIC badge palette per project is consumer-owned § 4.11.4.210 DATA. Composes

§ 4.11.4.211 / § 4.11.4.212 / § 4.11.4.213 /
 § 4.11.4.214 / § 4.11.4.215 / § 4.11.4.216 /
 § 4.11.4.217 / § 4.11.4.218 / § 4.11.4.219 /
 § 4.11.4.220 / § 4.11.4.221 / § 4.11.4.222 /

§ . . . / § . . . / § . . . /
 § . . . / § . . . / § . . . /
 § . . . / § . . . / § . . . /
 § . . . / § . . . / § . . . /
 § . . . / § . . . / §

Propagation gate CM-COVENANT-114-259-PROPAGATION (literal 11.4.259) + recommended gates CM-README-BADGE-ROW-AT-TOP + CM-BADGE-CLOSED-COLOR-VOCABULARY + CM-BADGE-MACHINE-DERIVED-SOURCE + CM-PRODUCTION-READINESS-GAUGE + CM-BADGE-SELF-VALIDATED + paired § . mutations.

Canonical authority: constitution submodule Constitution.md
 § . . . Non-compliance is a release blocker. No escape hatch – no --badges-optional, --hand-type-color, --skip-production-gauge, --badge-not-derived-from-source, --silent-omit-red-badge, --badge-in-footer-OK flag.

§ . . . – Cutting-edge enterprise quality + production-deployment readiness: every work product is built for production deployment, maximal stability, zero nasty surprises, and enterprise-grade robustness – never for "good enough" or "will do for now" (BACKGROUND :: REMINDER :: IMPORTANT operator mandate, – –)

Verbatim operator mandate (– – , Point): "All work we do MUST BE done for the cutting edge enterprise quality! We are doing everything for the deployment to production and maximal stability with no nasty surprises or issues of any kind!"

Compact summary: every work product a consuming project ships MUST be built for CUTTING-EDGE ENTERPRISE QUALITY targeted at PRODUCTION DEPLOYMENT + MAXIMAL STABILITY + ZERO NASTY SURPRISES – the standing default posture,

engaged automatically for every batch, subagent, feature, fix, and change from the first prompt of the session onward (§ . . . default-autonomous-loop composed with § . . . always-on defaults), NEVER a per-request opt-in and NEVER downgraded to "good-enough" or "will-do-for-now" or "MVP-first"; every change MUST satisfy the closed PRODUCTION-READINESS invariant set (clause B) before it is considered done, and a project maintains a docs/PRODUCTION_READINESS.md document (§ . . . sibling) that tracks readiness per component × invariant × evidence-source. § . . . STRENGTHENS § . . . (bleeding-edge enterprise visual quality – WEB surface only) into a UNIVERSAL production-readiness discipline across every surface (web, service, CLI, container, library, docs, orchestration, tests, gates) and BINDS § . . . /§ . . . / § . . . (captured evidence at every claim) + § . . . (four-layer verification) + § . . . (comprehensive test-type coverage) + § . . . (manual-QA sufficiency) + § . . . (evidence-class-at-closure) + § . . . (build-when-source-proven-correct) + § . . . (QA-deploy-readiness) into one always-on production-readiness posture.

(A) CUTTING-EDGE ≠ NOVELTY-FOR-ITS-OWN-SAKE. "Cutting-edge" means the current best-known engineering approach for the problem class, sourced via § . . . latest-source verification (never memory) + § . . . deep multi-angle research + § . . . external precedent, applied fully – NEVER a shortcut, a stub, a placeholder-with-a-TODO, a for now (§ . . .), a duct-tape workaround, or a "will fix later" (§ . . . loss-of-requirements); if the current best-known approach is genuinely unavailable in the project's constraints, that gap is HONESTLY tracked (§ . . . item + [OPEN]

marker per § 4.4.1), NEVER silently absorbed as "we did our best."

(B) PRODUCTION-DEPLOYMENT-READINESS INVARIANT SET. Every shipped work product satisfies ALL of: () FUNCTIONAL COMPLETENESS – every stated capability wired + verified on a clean target per § 4.4.1, no unwired feature, no dead-code path (§ 4.4.2); () OBSERVABILITY – health endpoint, structured logs, per-operation metrics, distributed-trace-ready surface, evidence-capture at every gate (§ 4.4.3 stress+chaos telemetry, § 4.4.4 always-on recording); () RESILIENCE – real failure-injection tested per § 4.4.5 (chaos: process death, network fault, resource exhaustion, state corruption); recovery is verified never assumed; () SECURITY – SonarQube CLI installed + scanned (§ 4.4.6), no leaked credentials (§ 4.4.7 pre-store audit), supply-chain integrity at the fleet minimum level (SLSA Build Level 3 per § 4.4.8); () PERFORMANCE – measured under sustained load per § 4.4.9, latency + throughput + resource footprint within declared SLOs; () MAINTAINABILITY – code review passed at Opus-xhigh (§ 4.4.10) with zero findings + zero warnings iterated to GO per § 4.4.11; () DEPLOYABILITY – rootless container per § 4.4.12, deterministic build + reproducible artifact per § 4.4.13 layer 3, deployment verified via § 4.4.14 target-isolation + § 4.4.15 layer 3 runtime-signature; () DOCUMENTABILITY – § 4.4.16 (manuals/guides/FAQ) + § 4.4.17 (diagrams) + § 4.4.18 (badges) all GREEN; () TESTABILITY – § 4.4.19 /§ 4.4.20 seven-type breadth + § 4.4.21 coverage floor, every gate paired § 4.4.22 mutation, every regression permanently guarded per § 4.4.23; () OPERATIONAL RUNBOOK –

for every non-trivial failure mode a documented
remediation path (§ 11 4 120 guide class).

(C) ZERO NASTY SURPRISES. No known-broken behavior ships
(§ 11 4 121 QA discovery channel closed), no failing
gate is silenced (§ 11 4 122 seam-placement +
§ 11 4 123 real-condition), no test disabled without
a tracked follow-up (§ 11 4 124), no closed-with-
workaround defect masquerading as fixed (§ 11 4 125
reason-source), no "will fail sometimes but rarely"
tolerance (§ 11 4 126 deterministic consistency +
§ 11 4 127 stress/chaos), no try/catch/swallow silent-
failure paths (a caught error either has a captured
recovery path or re-raises to a bounded operator decision
per § 11 4 128 /§ 11 4 129).

(D) STANDING DEFAULT. Composes § 11 4 130 (default
autonomous-loop mode) + § 11 4 131 (always-on
default mechanisms) + § 11 4 132 (maximum-useful
multi-agent + full-constitution-application per track) +
§ 11 4 133 (nano-precision model + effort tier
selection); the production-readiness discipline is engaged
from the FIRST prompt of the session, never negotiated,
never opt-in.

(E) docs/PRODUCTION_READINESS.md – the tracker. Per-
component × invariant × evidence-source table
(§ 11 4 134 sibling); revision header per
§ 11 4 135 ; auto-synced via § 11 4 136 ; reachable
from README per § 11 4 137 ; the source of truth for
the § 11 4 138 production-readiness gauge.

Honest boundary (§ 11 4 139). § 11 4 140 mandates
the POSTURE + the INVARIANT SET + the TRACKER; it does NOT
claim any project is provably-bug-free (§ 11 4 141
discovery-pressure boundary remains, § 11 4 142
escape triggers a coverage audit not a proof); it does NOT
license shipping with unmet invariants – an unmet

invariant is HONESTLY red-badged (§ 11.4.17) and blocks release. "Cutting-edge" is bounded to the project's problem class + published best-known approaches – a genuine research frontier where no best-known exists is HONESTLY marked as such and the project's own approach is documented as an original work contribution (§ 11.4.24). "NO external solution found – original work" citation).

Classification: universal (§ 11.4.26). Composes

§ 11.4.181	•	11.4.106	/	§ 11.4.187	•	11.4.106	/	§ 11.4.187	•	11.4.106	/
§ 11.4.189	•	11.4.118	/	§ 11.4.190	•	11.4.118	/	§ 11.4.190	•	11.4.118	/
§ 11.4.124	•	11.4.126	/	§ 11.4.128	•	11.4.126	/	§ 11.4.128	•	11.4.126	/
§ 11.4.134	•	11.4.135	/	§ 11.4.136	•	11.4.135	/	§ 11.4.136	•	11.4.135	/
§ 11.4.193	•	11.4.193	/	§ 11.4.199	•	11.4.193	/	§ 11.4.199	•	11.4.193	/
§ 11.4.193	•	11.4.194	/	§ 11.4.195	•	11.4.193	/	§ 11.4.195	•	11.4.193	/
§ 11.4.190	•	11.4.197	/	§ 11.4.198	•	11.4.190	/	§ 11.4.198	•	11.4.190	/
§ 11.4.200	•	11.4.201	/	§ 11.4.209	•	11.4.200	/	§ 11.4.209	•	11.4.200	/
§ 11.4.212	•	11.4.223	/	§ 11.4.224	•	11.4.212	/	§ 11.4.224	•	11.4.212	/
§ 11.4.226	•	11.4.231	/	§ 11.4.235	•	11.4.226	/	§ 11.4.235	•	11.4.226	/
§ 11.4.236	•	11.4.238	/	§ 11.4.237	•	11.4.236	/	§ 11.4.237	•	11.4.236	/
§ 11.4.238	•	11.4.239	/	§ 11.4.241	•	11.4.238	/	§ 11.4.241	•	11.4.238	/
§ 11.4.242	•		/	§ 11.4.242	•		/	§ 11.4.242	•		/
§ 11.4.242	•		/	§ 11.4.242	•		/	§ 11.4.242	•		/
§ 11.4.242	•		/	§ 11.4.242	•		/	§ 11.4.242	•		/

§ 11.4.242 • . Propagation gate CM-COVENANT-114-260-PROPAGATION (literal 11.4.260) + recommended gates CM-PRODUCTION-READINESS-TRACKER-PRESENT + CM-CUTTING-EDGE-POSTURE-ALWAYS-ON + CM-ZERO-NASTY-SURPRISES-AUDIT + paired § 11.4.242 mutations.

EXTENSION to (B)() OBSERVABILITY – the mechanism, not only the property (– –). Clause (B)() NAMES observability as a production-readiness invariant ("health endpoint, structured logs, per-operation metrics, distributed-trace-ready surface, evidence-capture at every gate"). Naming a property is not a mechanism, and three

mechanical requirements were missing. The AI-curriculum corpus (module) supplies them. (a) **THE THREE SIGNALS ARE DISTINCT AND EACH ANSWERS A DIFFERENT QUESTION.** The corpus's own naming is *"the three complementary signals of observability"* – **structured logs** (*"What exactly happened in this request?"* – *"queryable key-value events (JSON), not prose strings"*; the corpus is explicit that structure is load-bearing: `log.info("charge_failed", {donation_id, gateway, idempotency_key, err_code})` is queryable; `log.info("charge failed for " + id)` is *grep-and-pray*), **metrics** (*"How often, how bad, how loaded – across all requests?"* – numeric time series), and **traces** (*"Where in the request's path across services did time go or the error occur?"* – a tree of timed spans). The corpus adds a fourth signal class alongside them: **crash/ANR analytics** (*"symbolicated stacks with breadcrumbs, grouped by signature"*). Note for citation discipline (§ . .): the corpus does NOT use the terms "observability triad" or "three pillars"; consuming projects MUST NOT attribute those labels to it. (b) **A CORRELATION ID THREADS ALL SIGNALS, AND MUST SURVIVE ASYNC BOUNDARIES.** Corpus: *"A single request_id / trace_id threaded through logs, metrics exemplars, and spans lets me pivot from a metric spike → the exemplar trace → the structured logs of that exact request. That pivot is what turns 'error rate is up' into 'here is the one request and the line that failed.'" The load-bearing edge case is the async one: "For an async money flow, I make sure the correlation ID survives the queue – the webhook, the worker, and the reconciliation sweep all log the same donation/intent ID – or I lose the thread exactly where the hard bugs live."* The prescribed pivot order is **metric → trace → log**, under the discipline *"narrow before you theorize"*. (c) **CRASH-FREE USERS / SESSIONS IS A RELEASE GATE, AND A NEW SIGNATURE AT A RELEASE BOUNDARY IS A BLOCKER.** Corpus: *"the metric that matters is crash-free users / sessions, not raw crash*

count – it normalizes for traffic and is the honest health signal I gate releases on. And I treat a new signature after a deploy as a release-blocker even at low volume, because a signature that is rare today can be a top crasher once rollout completes. Volume tells me priority; novelty at a release boundary tells me causation." This binds directly to § 87(2)(b) (Crashlytics-recorded-data continuous monitoring – § 87(2)(b) (B)(i) now states WHEN that data blocks) and to § 87(2)(b) (the crash-free delta across a progressive rollout is one of that anchor's required non-infrastructure signals).

THRESHOLDS ARE CONSUMER DATA (§ 87(2)(b) / § 87(2)(b)). The corpus states NO numeric crash-free threshold and explicitly cautions that "crash-free thresholds" are among the figures that "change over time and by version – treat cited figures as illustrative and verify against current official documentation". This extension mandates that the gate EXISTS and that its threshold is declared and calibrated locally; it mandates no value. **Honest boundary (§ 87(2)(b)), stated to prevent an over-claim this extension will otherwise attract:** the corpus does NOT contain, and this clause does NOT assert, any proposition of the form "a system you cannot observe in production cannot be verified in production." What the corpus licenses is narrower and is the whole warrant here: "Observability turns 'sometimes' into a query. You cannot reproduce every production bug locally, so the system must be built to explain itself." Observability makes an unreproducible defect DIAGNOSABLE; it is not itself a verification of correctness (that remains

§ 87(2)(b) / § 87(2)(b) / § 87(2)(b)).

Recommended mechanism gates **CM-CORRELATION-ID-THREADED-ACROSS-BOUNDARIES** (a declared request/trace id is emitted by every signal and survives every async hop – queue, worker, scheduled sweep; a hop that drops it → FAIL) + **CM-CRASH-FREE-RELEASE-GATE** (a crash-free users/sessions gate exists

with a declared, locally-calibrated threshold, and a crash signature new at the release boundary blocks regardless of volume) + paired \$. mutations (drop the correlation id at the queue boundary → the first gate MUST FAIL; admit a release carrying a new crash signature at low volume → the second MUST FAIL; golden-FALSE per 11 4 217 . . (): a surface with genuinely no async hop and no crash-reporting channel, honestly \$. . SKIP-with-reason, MUST NOT fire either). Gate-code = 11 4 240 separate work item, NOT claimed shipped (\$. . / \$. .).

Canonical authority: constitution submodule Constitution.md \$. . . Non-compliance is a release blocker. No escape hatch – no --mvp-quality-OK , --good-enough , --will-fix-later , --skip-production-readiness-tracker , --production-readiness-optional , --acceptable-flakiness , --for-now-implementation flag.

\$. . – Zero-shortcomings / zero-gaps / zero-weak-spots / zero-danger-zones invariant with a mechanical audit ratchet: at all times the project holds exactly zero of these; every finding is closed OR tracked-with-honest-mitigation, and the audit-run count of open findings is MONOTONE-DECREASING (BACKGROUND :: REMINDER :: IMPORTANT operator mandate, – –)

Verbatim operator mandate (– – , Point): "There MUST BE ALWAYS exactly zero shortcomings, gaps, weak spots or danger zones in the project in any of the code we have, services or components!"

Compact summary: every consuming project MUST hold, at ALL TIMES, exactly ZERO shortcomings, gaps, weak spots, or

danger zones across every code path, service, component, gate, doc, and configuration – a standing invariant, not a per-release checklist; enforced by a MECHANICAL AUDIT SCRIPT that inventories every known finding class from a closed vocabulary (clause B), runs at every pre-build + release + constitution-pull seam, produces a machine-readable ledger, and is bound to a MONOTONE-DECREASING ratchet (§ . . . pattern) so the count of open findings only ever drops or the release refuses; every finding is either FULLY CLOSED with captured evidence (§ . . . /§ . . . /§ . . . + § . . . evidence-class-at-closure) or explicitly TRACKED as a § . . . item with an honest § . . . provenance marker + a § . . . - reversible-safe mitigation. "Zero" is the standing invariant; the ratchet is the mechanism that makes it reachable from a brownfield state without silently absorbing debt. § . . . GENERALISES § . . . (dead-code investigate-before-remove) + § . . . (discovery pressure – enumerated but unratched) + § . . . (QA escape triggers audit – one channel) into a WHOLE-PROJECT zero-finding invariant across the closed set.

(A) CLOSED FINDING VOCABULARY – an audit-actionable enumeration, extensible via named consumer-owned additions per § . . . but never subtractable: () SHORTCOMINGS – implementation gaps behind a working contract; () GAPS – missing capabilities the docs / schema / API promise; () WEAK SPOTS – code paths without adequate guards / error handling / input validation; () DANGER ZONES – code / config / operational patterns known to cause serious harm under conditions; () TODOs / FIXMEs / "for now" / "temporary" / "will do later" / "placeholder" / `NotImplementedError` / stub-returns / dead code (§ . . .); () SKIPPED tests / disabled gates / commented-out assertions without a tracked `# SKIP-`

OK: <ticket> per § 11.4.137; () BLUFFS – any § 11.4.138 / § 11.4.139 covenant violation surface (PASS-bluff, FAIL-bluff, echo-observable, source-green-closes-runtime); () UNRESOLVED § 11.4.140 items past their stated completion window; () DIVERGENT / STALE / ORPHAN artifacts per § 11.4.141 anti-mess control plane; () UN-CATALOGUED anti-patterns identified in code review that lack a § 11.4.142 mutation.

(B) MECHANICAL AUDIT SWEEP – a single executable script (Lava binding: scripts/audit/zero_findings_sweep.sh; consumers supply their path as § 11.4.143 DATA) that iterates the closed vocabulary and emits a machine-readable finding ledger {finding_id, class, file, line, description, status ∈ {open|tracked|closed}, tracker_ref, mitigation_evidence} – persistent, append-only + snapshot per § 11.4.144, git-tracked (§ 11.4.145) so the ledger IS the SSoT; runs at pre-build + release-tag + constitution-pull; each finding class ships golden-good/golden-bad fixtures per § 11.4.146 () so the audit itself cannot bluff. 11.4.135

(C) MONOTONE-DECREASING RATCHET. The ratchet snapshot records the accepted open-finding count at brownfield adoption (per the § 11.4.147 pattern, per-class + total); every subsequent audit run's count MUST be ≤ the ratchet snapshot for each class; a run that increases any class's count REFUSES the seam (pre-build gate FAILS, release refuses, commit refuses); the ratchet may DECREASE (a manual snapshot-lower after a finding is closed) and MUST NEVER INCREASE – the release-tag path never silently absorbs new debt. Brownfield adoption per operator § 11.4.148 decision (per § 11.4.149 (E) adoption-fence pattern), never an invented ratchet. 11.4.226

(D) EVERY FINDING → EITHER CLOSED OR TRACKED (never silently absorbed). CLOSED requires captured evidence at the defect layer per § 11.4.150 (runtime > artifact

>..source), the code / config genuinely repaired, the finding disappears from the audit sweep; TRACKED requires (i) a § . . . workable item with § . . . comprehensive structured description, (ii) an honest § . . . provenance marker ([OPEN: ATM-NNN]) at the finding site, (iii) a reversible-safe mitigation per § . . . , (iv) the tracked item's own completion is on the standing autonomous loop's queue (§ . . . / § . . . / § . . . / § . . .); (v) a § . . . permanent regression guard where the tracked item's exposure could cause harm before closure. Silent absorption ("we know this is broken but shipping anyway") is a § . . . violation of § . . . PASS-bluff severity – the ledger IS the honesty seam.

(E) HONEST "ZERO" – never the fake kind. "Exactly zero" means the CLOSED-VOCABULARY-audit-run finds zero unclosed-untracked entries; it does NOT claim omniscience (§ . . . discovery-pressure boundary – un-catalogued finding classes stay possible, and the vocabulary is EXTENSIBLE for exactly that reason); a § . . . escape triggers the class-extension of the audit + a ratchet-snapshot bump-DOWN (never up) once the escaped class is added and cleared. Honesty at the "zero" claim requires the vocabulary is genuinely covered – a project narrowing the vocabulary to declare "zero" is a § . . . fabrication and a release-blocker.

Honest boundary (§ . . .). § . . . mandates the STANDING INVARIANT + the AUDIT + the RATCHET + the TRACKED-vs-CLOSED discipline; it does NOT prove the project is provably-flawless, and does NOT license a "zero" claim over a narrowed vocabulary. Composes with, never substitutes for, § . . . (§ . . . catches the CATALOGUED – discovery still needed for the un-catalogued), § . . . (an escape triggers

vocabulary extension), § 11.4.17 (evidence class at closure decides whether a finding is genuinely closed), § 11.4.15 (regression guards for closed items keep them closed).

Classification: universal ($\mathcal{U}$. .). Composes

$\S_{11} \quad 4 \quad 192 \quad \bullet$	$\bullet \quad 4 \quad 197 \quad 10$	/	$\S \quad \bullet \quad \bullet$	/	$\S \quad \bullet \quad \bullet$	/
$\S_{11} \quad 4 \quad 119 \quad \bullet$	$\bullet \quad 4 \quad 115$	/	$\S_{116} \quad \bullet \quad \bullet$	/	$\S \quad \bullet \quad \bullet$	/
$\S_{11} \quad 4 \quad 118 \quad \bullet$	$\bullet \quad 4 \quad 124$	/	$\S_{126} \quad \bullet \quad \bullet$	/	$\S \quad \bullet \quad \bullet$	/
$\S_{13} \quad 4 \quad 135 \quad \bullet$	$\bullet \quad 4 \quad 146$	/	$\S_{161} \quad \bullet \quad \bullet$	/	$\S \quad \bullet \quad \bullet$	/
$\S_{11} \quad 4 \quad 197 \quad \bullet$	$\bullet \quad 4 \quad 281$	$11 \quad 4$	/	$\S \quad \bullet \quad \bullet$	(	) /
$\S_{11} \quad 4 \quad 224 \quad \bullet$	$\bullet \quad 11 \quad 4 \quad 226$	/	$\S \quad \bullet \quad \bullet$	/	$\S \quad \bullet \quad \bullet$	/
$\S_{11} \quad 4 \quad 233 \quad \bullet$	$\bullet \quad 4 \quad 238$	/	$\S \quad \bullet \quad \bullet$	/	$\S \quad \bullet \quad \bullet$	/
$\S \quad \bullet \quad \bullet$		/	$\S \quad \bullet \quad \bullet$	/	$\S \quad \bullet \quad \bullet$	/
$\S \quad \bullet \quad \bullet$		/	$\S \quad \bullet \quad \bullet$	/	$\S \quad \bullet \quad \bullet$	/
$\S \quad \bullet \quad \bullet$	(E)	/	$\S \quad \bullet \quad \bullet$	/		
$\S \quad \bullet \quad \bullet \quad 1 \quad 1$		/	$\S \quad \bullet \quad \bullet$	.	Propagation gate	CM-

COVENANT-114-261-PROPAGATION (literal 11.4.261) + recommended gates CM-ZERO-FINDINGS-AUDIT-SWEEP + CM-ZERO-FINDINGS-MONOTONE-RATCHET + CM-EVERY-FINDING-CLOSED-OR-TRACKED + paired \$ mutations.

Canonical authority: constitution submodule Constitution.md

```
$ . . . Non-compliance is a release blocker. No
escape hatch – no --skip-findings-audit ,
--allow-todo-without-tracker , --ratchet-may-increase , --
silently-absorb-finding , --narrow-vocabulary-to-declare-zero ,
--for-now-in-production flag.
```

§ . . – Machine-created evidence
at every gate: every claim of "works" /
"passes" / "verified" cites a captured,
machine-derived, machine-verifiable evidence
artifact produced by the gate – no operator
eyeballing, no narrative-only PASS
(BACKGROUND :: REMINDER :: IMPORTANT operator
mandate, – –)

Verbatim operator mandate (– – , Point
): *"There MUST BE no bluff of any kind and in any form
anywhere! Everything created MUST BE fully covered with
confirmation, validation and verification gates which will
assert machine created evidence so there is no prediction
or guessing about the success and production ready state
of the System!"*

Compact summary: every claim by the System that something
WORKS / PASSES / IS VERIFIED / IS PRODUCTION-READY MUST be
backed by MACHINE-CREATED, MACHINE-VERIFIABLE evidence – a
file / verdict / measurement produced by the gate itself,
cited by path + content-hash + timestamp, with a schema an
independent verifier can parse (§ . . evidence-
class-at-closure applied to EVERY gate, not only closure
seams); operator eyeballing, human narrative, "I checked
it works", "looks fine", "no error was reported", and
every other prediction / guess / assumption about success
is FORBIDDEN at the evidence layer (§ . /\$. . /
§ . . 11 4 107 10 .). § . . STRENGTHENS the § .
anti-bluff covenant preamble from "captured evidence" to
MACHINE-CREATED evidence (a captured screenshot an
operator interpreted is captured but not machine-verified
– § 211 11 4 212 () golden-good/golden-bad self-
validating analyzer applies UNIVERSALLY, not only to AV
playback), and BINDS § . . /\$. . /
§ . . /\$. . /\$. . (F)/
§ . . /\$. . into one universal
machine-evidence-at-every-gate discipline. Prediction,

guessing, and assumption about success are the exact anti-patterns this anchor forbids.

(A) THREE-GATE COVERAGE. Every work product passes THREE gate classes each producing machine-created evidence: () CONFIRMATION – the change lands correctly (source review, static analysis, format/lint, dependency clash check per § . . .); () VALIDATION – the change behaves as designed at each of the § . . . four layers (SOURCE / ARTIFACT / RUNTIME / USER-VISIBLE), each producing its own machine-created verdict + evidence artifact; () VERIFICATION – the change satisfies its acceptance contract from the user's / operator's / system's point of view (§ . . . manual-QA is the human sufficiency gate; § . . . QA-deploy-readiness is the machine sufficiency gate – each with its own captured evidence).

(B) MACHINE-CREATED – the four properties. Every evidence artifact is (i) PRODUCED BY THE GATE, not authored by the tester (a screenshot the harness took, not one the operator narrated); (ii) MACHINE-VERIFIABLE – a schema-parseable file (JSON verdict, JSONL event stream, structured log, image with an OCR/vision oracle per § . . . (. . .)/\$. . . , audio with RMS/loudness analysis per § . . . (. . .)) that an INDEPENDENT verifier can re-check without asking the tester; (iii) CONTENT-ADDRESSED – cited by path + sha + timestamp (§ . . . content-addressed Merkle) so tampering is detectable; (iv) SEMANTIC-CLASS-MATCHED – the evidence class matches the claim layer (§ . . . wrong-layer echo forbidden – a source grep transcript can never satisfy a "pixel renders green" claim).

(C) NO NARRATIVE-ONLY PASS. A PASS whose evidence is a paragraph of natural language authored by an operator or a subagent ("I ran the tests and they all passed", "looks

good to me", "manually confirmed") is a § . . . violation – the narrative may accompany the machine evidence but MUST NEVER substitute for it. § . . . manual-QA remains a human-sufficiency gate whose OWN evidence is machine-created (a QA session recording per § . . . , a QA checklist file with per-item PASS/FAIL + evidence-path per § . . . , a § . . . escape-log if the human found something automation missed).

(D) SELF-VALIDATED EVIDENCE ANALYZERS. Every evidence analyzer (the parser that reads the machine-created artifact and emits a verdict) is itself golden-good + golden-bad + negative-control validated per § . . . (); an analyzer that PASSes its golden-bad fixture is the bluff and a release-blocker (the § . . . self-validation discipline applied to evidence tooling). Every claim of "works" cites the analyzer's verdict + the raw evidence, so an independent re-run can reproduce the verdict.

(E) FOUR-LAYER BINDING WITH § At each of the four verification layers (SOURCE / ARTIFACT / RUNTIME-ON-CLEAN-TARGET / USER-VISIBLE), the layer's PASS cites MACHINE-CREATED evidence PRODUCED AT THAT LAYER (source: static-analysis JSON + review-artifact; artifact: build-log + fingerprint-checksum + byte-check of the change's landed bytes; runtime: on-target verdict + captured runtime signature + telemetry; user-visible: OCR/vision oracle output or sink-side probe report per § . . . / § . . . / § . . .). A PASS at a higher layer that cites evidence from a lower layer alone is a § . . . wrong-layer violation – the machine-evidence chain runs THROUGH the layers, never SKIPS.

(F) STANDING DEFAULT (§ . . . / § . . .). Machine-evidence-at-every-gate is the STANDING DEFAULT posture engaged from the first prompt of the session;

every subagent dispatch, every review, every test run, every gate emits its evidence artifact into the project's evidence store (§ . . . content-addressed) so the full audit trail is machine-searchable at any time by any operator ./ subagent / reviewer. Prediction, guessing, and assumption about success are forbidden at the reporting layer as well as the gate layer – a status report that predicts "should pass" without citing a run's evidence is a § . /§ . . bluff at the reporting layer.

Honest boundary (§ . .). § . . mandates every gate's PASS is BACKED by machine-created evidence + the analyzer is self-validated; it does NOT claim machine evidence proves human satisfaction (§ . . remains, with its own machine-evidence requirement per (C)); it does NOT license shipping when the machine evidence disagrees with reality (a bug the automation missed is a § . . escape triggering a coverage audit + a vocabulary extension per § . . , not a "the machine said green so we're good" defense). "Machine-created" is bounded to the artifact – the DECISION about whether a gate passed remains policy the humans authored (§ . . no-guessing extends to policy calibration, not to the gate's mechanical execution).

Classification: universal (§ . .). Composes

§ . /§ . (anti-bluff covenant) /
 § . / § . / § . /
 § . / § . / § . /
 § . / § . / § . /
 § . / § . ()
 () () / § . / § . /
 § . (F) / § . /
 § . / § . / § . /
 § . / § . / § . /
 § . / § . / § . /
 § . / § . / § . /

\$. . / \$. . / \$. . .
Propagation gate CM-COVENANT-114-262-PROPAGATION (literal
11.4.262) + recommended gates CM-MACHINE-EVIDENCE-AT-EVERY-
GATE + CM-EVIDENCE-ANALYZER-SELF-VALIDATED + CM-NO-NARRATIVE-
ONLY-PASS + CM-EVIDENCE-LAYER-MATCH-108 + paired \$.
mutations.

Canonical authority: constitution submodule Constitution.md
\$. . . Non-compliance is a release blocker. No
escape hatch – no --narrative-pass-OK , --operator-eyeball ,
--should-pass-suffices , --skip-evidence-analyzer-validation ,
--evidence-optional , --pass-without-machine-artifact , --
predict-success flag.

\$. . – Process-group signal-
safety mandate: NEVER signal pgid ≤ ,
NEVER trust a mock-derived pid/pgid – validate
as int > before every killpg /
kill(-pid, sig) / pkill / killall call
(BOB- forensic anchor,
- -)

Forensic anchor – verbatim operator mandate
(- -):

"CRITICAL: Fix MUST BE systematic so any other
projects do not have same issue!"

Forensic case study (BOB- , boba project,
- -). SEVEN forced-logout incidents on the
operator's workstation over ~ hours
(BOB- / / / / /) – the
operator's entire graphical session, tmux, SSH, browsers,
and Claude Code process were repeatedly SIGKILLED. Kernel
audit rules installed - - :
finally captured the initiator on incident :


```
audit[399861]: SYSCALL syscall=62 a0=ffffffff a1=9
pid=399861 auid=1000 comm="pytest" exe="/usr/bin/python3.14"
key="sigkill_investigation"
```

A `pytest` process called `kill(-1, SIGKILL)`. Root cause: a Python test created an `AsyncMock()` subprocess object without explicitly setting `mock.pid` as an int. The production code path called `os.killpg(os.getpgid(proc.pid), signal.SIGKILL)`. Python's `MagicMock.__int__` defaults to `1`, so:

- `proc.pid` → `MagicMock`
- Python coerces to `int` for `os.getpgid` → `int(MagicMock()) == 1` (documented default)
- `os.getpgid(1)` returns `1000` (init's process group)
- `os.killpg(1, SIGKILL)` → glibc → `kill(-1, SIGKILL)`
- `kill(-1, sig)` semantics = "signal every process the caller may signal, except pid and self" → under UID `1000` that means the entire user session (systemd `user@1000` manager, gnome-shell, tmux, ssh, browsers, Claude Code)

The `contextlib.suppress(Exception)` around the call swallowed nothing because the syscall SUCCEEDED before any exception path. The defect existed for ~ 18 months (since 2020-04-24) before audit rules exposed it.

The mandate (ALL hold):

(A) NEVER signal `pgid ≤ 1`. No project, no language binding, no code path may call `os.killpg(pgid, sig)` / `kill(-pid, sig)` / `killpg(2)` (POSIX) / `pkill -<pgid>` / `subprocess.killpg` / `Process.killpg` where `pgid` (or `pid` in the negative-`pid` form of `kill(2)`) is `≤ 1`. On Linux, `killpg(1, sig)` and `kill(-1, sig)` are equivalent and BOTH signal every process the caller has permission to signal – under a non-root UID, that includes the user's session

manager, GUI, shell, and every daemon in the same UID scope. This is the disaster syscall. Guarding with a positive integer test is trivial; failing to guard is catastrophic. On other Unixes (BSD family, macOS) the semantics of `kill(-1, sig)` is broadly the same "signal all processes the caller may signal"; the mandate is portable.

(B) VALIDATE `pid + pgid as int > 1` BEFORE every process-group signal call. The universal invariant:

```
IF calling killpg(pgid, sig) OR kill(-pid, sig) OR any process-group signal:
    ASSERT isinstance(pgid, int)           # never a mock, never a string, never None
    ASSERT pgid > 1                       # never init's pgid, never 0 (self-group), never n
    (same for pid when the negative-pid kill(2) form is used)
```

Language-specific enforcement:

- **Python:** guard with `if isinstance(pid, int) and pid > 1: pgid = os.getpgid(pid); if isinstance(pgid, int) and pgid > 1: os.killpg(pgid, sig)` .
- **Go:** `syscall.Kill(-pgid, sig)` – guard `pgid > 1` before the call; test doubles must return `int(pgid) > 1` .
- **Rust:** `nix::sys::signal::killpg` – accept only `Pid::from_raw(n)` with `n > 1` , and reject the sentinel meaning "current process group".
- **Bash / shell:** `kill -<pgid> <sig>` / `pkill -g <pgid>` – validate `[[ "$pgid" -gt 1 ]]` before the call.
- **C:** `killpg(pgid, sig)` and `kill(-pid, sig)` – same integer guard.

(C) TESTS mocking subprocess/proc objects MUST explicitly set `mock.pid` (or equivalent) as `int`. A subprocess mock with an unset `pid` field is BOB- -shaped by construction: the language's mock library (Python's `MagicMock` , Go testify's `Mock` , JS Jest's `jest.fn()` , etc.) provides a default coercion that can silently become a low

integer. Every test that creates a fake subprocess object MUST:

- Set `mock.pid = <large int, e.g. 12345>` explicitly – never rely on the mock library's default.
- Patch the process-group signaling call itself (e.g. `patch.object(module.os, 'killpg')`) as belt-and-suspenders so a future regression of the production code's guard cannot re-open the disaster.
- Assert the mocked `killpg` was called with `pgid > 1` (paired `$ .` mutation catches a regression that lets `pgid=` through).

(D) DEFENSE-IN-DEPTH (\$. . () analyzer discipline applied). The guard MUST be present in BOTH the production code AND the test. Removing either layer alone must break at least one test. This is the four-layer coverage per \$. . (b) applied to the signal-safety class:

- Layer (pre-build gate): grep pattern for `killpg\(.*\)` / `kill\(-.*\)` calls that are NOT preceded by the guard – surface as review findings.
- Layer (post-build gate): CodeGraph query for callers of `os.killpg` / `syscall.Kill(-*)` / equivalent, cross-referenced against guard presence.
- Layer (runtime): kernel audit rules on SIGKILL syscalls with `a0=0xffffffff` (`kill(- )`) or `a0=0x00000001` in the `pidfd_send_signal` case – a live host-safety canary.
- Layer (paired \$. mutation): a regression test that DELIBERATELY invokes the `killpg` code path with a mock pid AND asserts `killpg` was not called with `pgid ≤ 1`. Reverts of the pid guard MUST make this test fail.

(E) HONEST BOUNDARY (\$. .). \$. . hardens the process-group signal APIs specifically. It

does NOT prohibit legitimate use of `killpg` for well-scoped process groups (`start_new_session=True` subprocesses where the child owns its group). It does NOT extend to signals other than SIGKILL where the blast radius is small (SIGCHLD, SIGUSR , etc. – though the same guard is cheap and recommended). It does NOT prevent kernel OOM-killer or systemd-oomd from killing the user (those are separate host-safety concerns – § 11.4.263 . memory ceiling + § 11.4.142 . quota telemetry apply).

(F) STANDING DEFAULT. Every project inheriting this constitution treats § 11.4.142 as always-on from the first prompt. Every code review (§ 11.4.17 / § 11.4.17 / § 11.4.17) MUST scan for process-group signal calls and verify the guard. Every subagent dispatched to write subprocess-cleanup code MUST include the guard in its output. Every test-writing subagent MUST include the `mock.pid = <int>` discipline.

Classification: universal (§ 11.4.17). Composes § 11.4.17 (FAIL-bluff – swallowed exception around a successful catastrophic syscall is a bluff at the error-handling layer) / § 11.4.263 (b) (four-layer coverage) / § 11.4.263 / § 11.4.263 (chaos-test host safety) / § 11.4.17 () (self-validated analyzer) / § 11.4.17 (RED-first with the actual defect condition) / § 11.4.17 (standing default) / § 11.4.17 (guard asserts real condition – `pgid > 0`) / § 11.4.17 (host-safety orthogonal axis) / § 11.4.17 (real captured evidence). Propagation gate `CM-COVENANT-114-263-PROPAGATION` (literal `11.4.263`) + recommended gates `CM-KILLPG-PGID-GUARD` (production code – every `killpg/kill(-pid)` call site has the `pid>1 && pgid>1` guard) + `CM-TEST-MOCK-PID-EXPLICIT-INT` (test code – every subprocess mock has `mock.pid` explicitly set as int + `killpg` patched) + paired § 11.4.17 mutations (remove the guard from production → the regression test fails; unset

mock.pid in the test → assertion catches pgid<= attempt).

Canonical authority: constitution submodule Constitution.md
\$. . . Non-compliance is a release blocker regardless of context. No escape hatch – no --skip-killpg-guard , --mock-pid-optional , --killpg-1-is-ok , --kill-minus-1-permitted , --suppress-catches-syscall , --test-may-not-explicit-pid flag exists.

\$. . . – Build once, promote ONE immutable content-addressed artifact through every environment; never rebuild per stage (research-derived, – –)

Corpus anchor (AI-curriculum module , "Core concepts in one page", concept – verbatim):

"Promote one immutable artifact through environments; never rebuild per stage. Build once, then promote the same bytes from CI to staging to canary to production. Rebuilding per environment means the thing you tested is not the thing you shipped – a subtle drift (a floating dependency, a different base image) can make production behave unlike staging. Immutable artifact + environment-specific config injected at deploy is the rule."

Why it is a bluff, not merely an inefficiency. Rebuilding per stage silently BREAKS THE CHAIN OF EVIDENCE: every gate that ran against the staging artifact was a measurement of a DIFFERENT binary than the one production runs, so each of those green verdicts is a \$. PASS-bluff about the shipped artifact – evidence correctly captured against the wrong subject. The corpus states the

failure directly: *"The thing you tested must be the thing you ship. If you rebuild for production, a floating dependency, a re-pushed base-image tag, or a toolchain difference can make the production artifact subtly unlike the one that passed staging – so your tests validated a different binary. Rebuilding per environment silently breaks the chain of evidence."* Its red-flag form: *"'We rebuild for production to be safe.' (Rebuilding per environment breaks the chain of evidence – you shipped a different artifact than you tested.)"*

The mandate (ALL hold):

(A) BUILD ONCE. Exactly ONE artifact is produced per release candidate, by ONE build, under § 1.1.1 (reproducible + hermetic + SLSA-L1 provenance). Every subsequent environment receives THAT artifact. A pipeline that invokes its build step once per environment violates § 1.1.1 regardless of how deterministic the build is claimed to be – determinism is the § 1.1.1 property that makes a rebuild *detectably* equivalent, never a licence to rebuild instead of promote.

(B) IDENTITY IS A CONTENT ADDRESS, NOT A TAG. The artifact is identified by a CONTENT-ADDRESSED digest (`sha256:...` or the platform equivalent), never by a mutable tag, branch name, or version string alone. Corpus: *" `node:20` is a moving pointer; `node@sha256:...` is an immutable content address. Tags get re-pushed; digests cannot change without changing the digest."* Senior form: *"the subtle attack this defends against is not 'a bad version I chose!' but 'the same version, different bytes' – a compromised registry or a rebuilt-under-the-same-tag image. Only content-addressing (digests + integrity hashes) catches that; a version string alone does not."* Composes § 1.1.1 (resolve-by-stable-name-not-by-index – a digest is the strongest stable identifier available) and § 1.1.2 (the deployed artifact's identity is READ BACK from the

intended target and asserted equal to the promoted digest).

(C) PROMOTION IS A VERIFICATION, NOT A REBUILD. Corpus, mechanically: promotion is *"verify this exact digest is authentic and approved for the next environment," not "make a new thing and hope it's equivalent"*, and the promotion gate *"check[s] identity and approval of a digest rather than re-running a build."* Signature, SBOM, and provenance TRAVEL WITH the artifact (\$. . .), so the evidence chain moves with the bytes instead of being re-derived per stage.

(D) ENVIRONMENT CONFIG IS INJECTED AT DEPLOY, NEVER BAKED AT BUILD. Corpus: *"the mechanism is config injection at deploy, not build-time config baking. The artifact is environment-agnostic; the environment supplies secrets/endpoints/flags at run time. This is what lets the same provably-tested bytes run everywhere."* Baking per-environment config at build time RE-CREATES the per-stage rebuild it forbids and is explicitly the corpus's marked-wrong answer. Credentials remain governed by \$. . . – injection at deploy is never a licence to bake a secret into an artifact.

(E) ROLLBACK IS DETERMINISTIC BY CONSEQUENCE. Because known digests are promoted, rollback is *"point production at the previous known-good digest – a precise, reversible operation, not a rebuild-from-source scramble under incident pressure."* This is the precondition that makes \$. . . 's auto-abort a real undo rather than a hopeful re-deploy.

(F) REBUILD-TO-DODGE IS A GATE-EVASION, NOT A RETRY. Re-running or rebuilding until a failing scan or flaky check turns green is named by the corpus as a distinct cheat: *"A failing scan or flaky check 'fixed' by re-running until green, or by rebuilding to get a different result,*

launders a real failure into a pass." Retries MUST be counted and surfaced (composes § 11 4 6 . . . flaky-quarantine – a flake is quarantined with an owner and a deadline, never laundered by a rebuild).

11 4 108

Honest boundary (§ 11 4 108 . . .). § 11 4 108 . . . guarantees that the bytes which passed the gates are the bytes that run – it does NOT prove those bytes are correct (that remains § 11 4 108 . . . 's four layers, § 11 4 108 . . . 's full-suite retest, and § 11 4 108 . . . 's manual-QA sufficiency gate), does NOT prove the deploy reached the intended target (that is § 11 4 108 . . . 's isolate-and-read-back), and does NOT prove the environment's injected configuration is right (that is § 11 4 108 . . . 's boot-time invariant assertion). It closes exactly one link: the artifact identity link between test and production. The corpus names no environment ladder beyond CI → staging → canary → production ; a consuming project supplies its own ladder as DATA per § 11 4 108

Classification: universal (§ 11 4 108 . . .) – a platform-neutral delivery-pipeline invariant reusable by any project that builds an artifact and deploys it to more than one environment; the consuming project supplies its environment ladder, digest scheme, and config-injection mechanism as DATA per § 11 4 204 Composes § 11 4 204 . . . / § 11 4 204 . . . (a gate verdict about the wrong artifact is a bluff) / § 11 4 131 . . . / § 11 4 204 . . . (injected config never bakes a secret) / § 11 4 204 . . . / § 11 4 204 . . . (the ARTIFACT layer – § 11 4 204 . . . is what keeps that layer's identity stable across environments) / § 11 4 246 . . . / § 11 4 204 . . . (project-prefixed, monotonic release identity) / § 11 4 204 . . . (deployed-artifact identity read back from the intended target) / § 11 4 204 . . . (B) (per-deploy distinct version id) / § 11 4 204 . . . (reproducible + hermetic + provenance-

attested – § . . . is its DEPLOY-side complement, not a restatement: § . . . governs how the artifact is BUILT, § . . . governs that it is not built AGAIN) / § . . . / § . . . (the canary shifts traffic to exactly the bytes that passed) / § . . . (fail closed when identity cannot be verified) / § . . .

Propagation gate CM-COVENANT-114-264-PROPAGATION (literal 11.4.264 present as a block-start, exactly-once per governance file, lockstep content-hash equality across the mirror set per § . . . (B)) + recommended mechanism gates CM-BUILD-ONCE-PROMOTE-DIGEST (the pipeline's build step is invoked at most once per release candidate; every post-build environment stage references a pinned content-addressed digest and NOT a rebuild target; a stage that re-invokes the build → FAIL) + CM-DEPLOY-CONFIG-INJECTED-NOT-BAKED (no environment-specific endpoint/secret/flag is materialised into the artifact at build time; a per-environment build argument that changes artifact bytes → FAIL) + paired § . . . mutations (add a second per-environment build invocation → CM-BUILD-ONCE-PROMOTE-DIGEST MUST FAIL; replace a promoted digest reference with a mutable tag → it MUST FAIL; bake an environment endpoint into the build → CM-DEPLOY-CONFIG-INJECTED-NOT-BAKED MUST FAIL; golden-FALSE fixture, per § . . . () the false-positive guard: a single build followed by N digest-referencing promotion stages MUST NOT fire either gate). Gate-code = separate work item, NOT claimed shipped (§ . . . / § . . .).

Canonical authority: constitution submodule Constitution.md § Non-compliance is a release blocker regardless of context. No escape hatch – no --rebuild-per-environment , --promote-by-tag , --bake-env-config-at-build , --rebuild-until-green , --skip-digest-pin , --deterministic-build-means-rebuild-is-safe flag exists.

§ 0.8.20. – Progressive delivery: traffic-shifted rollout gated by automated analysis on SLO and business metrics, with automatic abort to stable (research-derived, – –)

Corpus anchor (AI-curriculum module , "Core concepts in one page", concept – verbatim):

"Progressive delivery catches the bluff the tests missed, before all users do. Tests prove what you thought to assert; canary and blue-green deployments prove the release survives real traffic on a small blast radius first. Automated analysis compares the new version against the stable baseline on real metrics (error rate, latency, and business KPIs), and automatically rolls back on breach – no human at AM. The gate here is health, measured, with an instant undo."

Why it is an anti-bluff mechanism and not a deployment convenience. Every gate before this point measures what someone thought to assert; progressive delivery is the FIRST gate whose oracle is real user traffic. It is the deploy-layer instantiation of the § 0.8.20 covenant: it exists to catch the defect the suite never encoded, at a blast radius of a small traffic fraction rather than the whole user base.

The mandate (ALL hold):

(A) A RELEASE REACHES USERS BY A BOUNDED-BLAST-RADIUS STRATEGY, NEVER AN UNGATED CUTOVER. The consuming project selects and declares its strategy per § 0.8.20 from the corpus's two named forms – CANARY (*"shift traffic to the new version gradually ... pausing at each step to*

evaluate"; smallest blast radius, requires weight-based routing, old and new run concurrently so schema/compat must tolerate it) or BLUE-GREEN ("stand up the new version ('green') alongside the current one ('blue') ... Then switch all live traffic at once ... Keep blue running for a rollback window"; clean cutover and near-instant undo, but "at the switch, all users hit the new version simultaneously, so a subtle bug hits everyone at once"). For a critical-invariant surface per § (money movement, safety, availability, integrity) the corpus's stated preference is canary paired with blue-green's instant-undo posture: "' % of donors briefly saw a slower page' is a survivable failure and ' % of donors hit a double-charge bug at the cutover instant' is not."

(B) THE PROMOTE/ABORT DECISION IS MADE BY AUTOMATED ANALYSIS, NOT BY A HUMAN WATCHING A DASHBOARD. Corpus: *"the controller queries a metrics provider ... and compares the canary against the stable baseline across defined metrics over successive analysis windows. If the metrics stay within thresholds across the required windows, it auto-promotes to the next weight; if a metric breaches, it auto-aborts and rolls back all traffic to stable – no human needed."* A manual approval click is explicitly rejected as the gate – corpus red flag: "'A senior clicks approve, that's the deploy gate.'" (Manual LGTM is a rubber stamp with no receipt; under agents it looks like oversight while providing none.)" This composes § (the producer may not grade its own rollout) and § (the controller's analysis result is the machine-created evidence; a narrated "looks healthy" is not).

(C) THE ANALYSIS MUST INCLUDE BUSINESS METRICS, NOT ONLY INFRASTRUCTURE METRICS – THIS CLAUSE IS THE LOAD-BEARING ONE. Corpus, naming it *"the classic mistake"*: *"measuring*

only infrastructure metrics. 'HTTP error rate is fine' does not tell you the checkout-completion rate quietly dropped % because of a UI regression that returns s while failing the user. The whole value of canary is catching the failure your tests missed – and the failures that matter most are business failures that look technically green. So the analysis must include business KPIs, or it re-creates the 'green pipeline that proves nothing' at the deploy layer." An infra-only canary is therefore a § . PASS-bluff wearing a rollout's uniform. For a § . . critical-invariant surface the domain invariants ARE the steady-state hypothesis – corpus: "a canary that watches latency but not the money invariant would happily promote a build that double-charges % of donors."

(D) MULTIPLE CONSECUTIVE ANALYSIS WINDOWS, NEVER A SINGLE SAMPLE. Corpus: "Requiring the canary to pass all checks across multiple consecutive windows guards against a single noisy sample flipping the decision." Composes § . . (deterministic consistency – one sample is not a verdict) and § . . (a steady-state window, never a moment).

(E) POST-DEPLOY SMOKE IS A FAIL-CLOSED PRECONDITION OF TRAFFIC. Corpus: readiness gates traffic; smoke runs "against the freshly deployed instance before it takes real traffic (or before promotion)" and proves "wiring and configuration – the things that are correct in code but can still be broken by a bad secret, a wrong endpoint, a missing migration, an environment drift"; "a failed smoke gate is a fail-closed signal that aborts the promotion before any user is affected." Composes § . . (boot-time invariant assertion) and § . . (fail closed).

(F) THE DEPLOY GATE IS INDEPENDENT OF THE MERGE GATE. Corpus: "Signature/provenance verification and canary

analysis are gates a merge cannot wave through – so 'got it merged' is not 'got it shipped.'" A merge MUST NOT be able to auto-deploy past the analysis gate.

(G) THRESHOLDS AND WINDOW COUNTS ARE CONSUMER-SUPPLIED DATA, NEVER CONSTITUTIONAL LITERALS (§ . . . / § . . .). The corpus deliberately refuses to supply them, writing its abort criterion with UNFILLED placeholders – *"if success rate drops below X% or p exceeds Y ms, roll back automatically (thresholds illustrative – set from your SLOs)"* – and adding *"Thresholds and window counts are configuration I would tune to the SLOs and confirm live, never quote as fixed."* Its governing rule on stale numbers: *"A release gate that quotes a stale requirement is itself a bluff."* Any project adopting a number supplies and calibrates it locally; this anchor mandates NO numeric value.

Honest boundary (§ . . .). Progressive delivery bounds the blast radius of a defect and makes the undo automatic – it does NOT prove the release correct (a defect invisible to the declared metrics promotes cleanly), does NOT replace § 's manual-QA sufficiency gate, and does NOT apply where there is no traffic-shiftable surface (a firmware image flashed to a physical device, a desktop binary, a batch job) – for those the honest posture is an § SKIP-with-reason plus whatever staged-audience mechanism the platform genuinely offers, never a fabricated canary. The corpus's progressive delivery is exclusively TRAFFIC-SHIFTING; it is silent on feature flags, dark launch, kill switches, and shadow traffic, and this anchor mandates none of them. The corpus contains no claim of the form "a deploy you cannot roll back is not a deploy" and none is asserted here.

Classification: universal (§ . . .) – a platform-neutral delivery discipline reusable by any project that ships to users behind a routable surface; the consuming

project supplies its strategy, metric set, thresholds, window counts, and controller as DATA per § . . . Composes § . / § . . / § . . (honest SKIP where no traffic surface exists) / § . . / § . . 230 . . 230 / § . . / § . . / § . . / § . . / § . . (the analysis is a guard: it must be able to abort, and an unresolvable metric read fails closed) / § . . 232 . . 234 / § . . (a defect the canary catches that the suite missed is a coverage escape and triggers that anchor's audit) / § . . / § . . / § . . / § . . (the canary shifts traffic to exactly the bytes that passed) / § . .

Propagation gate CM-COVENANT-114-265-PROPAGATION (literal 11.4.265 present as a block-start, exactly-once per governance file, lockstep content-hash equality across the mirror set per § . . (B)) + recommended mechanism gates CM-PROGRESSIVE-DELIVERY-STRATEGY-DECLARED (a user-facing release either declares a bounded-blast-radius strategy with an automated analysis step, or carries an honest § . . SKIP-with-reason naming the absent traffic surface; an ungated full cutover with neither → FAIL) + CM-CANARY-ANALYSIS-INCLUDES-BUSINESS-METRIC (the declared metric set contains at least one business/domain outcome metric, not exclusively infrastructure metrics; an infra-only analysis on a § . . critical-invariant surface → FAIL) + paired § . mutations (strip every business metric from the analysis set leaving only error-rate and latency → CM-CANARY-ANALYSIS-INCLUDES-BUSINESS-METRIC MUST FAIL; replace the automated abort with a manual approval step → CM-PROGRESSIVE-DELIVERY-STRATEGY-DECLARED MUST FAIL; golden-FALSE fixture per § . . (): a genuinely non-routable artifact carrying an honest SKIP-with-reason MUST NOT fire either gate). Gate-code =

11 4 227
separate work item, NOT claimed shipped (\$. . /
\$. .).

Canonical authority: constitution submodule Constitution.md
\$. . . Non-compliance is a release blocker
regardless of context. No escape hatch – no --full-cutover-
ok , --infra-metrics-suffice , --manual-approval-is-the-gate ,
--single-window-analysis , --skip-smoke-before-traffic , --
merge-implies-deploy flag exists.

\$. . – Claim-vs-reality ledger:
every ADVERTISED capability is a row, typed
from the closed bluff-type vocabulary, and a
capability with no passing challenge is a
release blocker (research-derived,
– –)

Corpus anchor (AI-curriculum module , [Q]
"Portfolio-scale prevention", first and cheapest of its
enumerated portfolio gates – verbatim):

"A shared claim-vs-reality ledger per repo (the
validation plans already are this): every advertised
capability is a row with a Bluff type and a severity,
and a capability with no passing challenge is a
release blocker. This makes stubbed-core, byte-
identical-fork, and config-present-but-unwired visible
and blocking rather than discovered later."

11 4 229
The direction of enumeration is the whole point, and it is
what no existing anchor supplies. \$. . 's coverage
ledger, \$. . 's autonomous-path column, and
\$. . 's test-type column are all keyed on
FEATURES the project already knows it has – they enumerate
from the implementation side. A claim-vs-reality ledger

enumerates from the CLAIM side: what the project ADVERTISES to users (README capability list, docs, marketing copy, API/CLI surface, badge assertions per § . . . , the § . . . capability matrix). A claim with no feature row is invisible to a feature-keyed ledger by construction – it is precisely the class of defect the corpus documents as *stubbed-core* ("the most-engineered code has no caller"), *byte-identical-fork* ("the name promises X, the code is Y unchanged"), and *config-present-but-unwired* ("the flagship path silently no-ops while the right code exists one dir over"). § . . . therefore adds a keying, not a restatement (§ . . .).

The mandate (ALL hold):^{11 4 35}

(A) ONE LEDGER PER REPOSITORY, ENUMERATED FROM THE CLAIM SIDE. Every advertised capability is a ROW. The consuming project declares per § . . . which surfaces constitute "advertised" (at minimum: the README capability set required reachable by § . . . , the § . . . user-manual/guide/FAQ set, the § . . . capability matrix, and the § . . . badge assertions). The corpus does not narrow "advertised" further and neither does this anchor (§ . . .).

(B) EACH ROW CARRIES A BLUFF TYPE FROM THE CLOSED VOCABULARY, AND A SEVERITY. The corpus's bluff-type vocabulary, verbatim with its stated fingerprints, is a SEVEN-member set: *green-but-broken* ("passes CI, fails on device"), *coverage-theater* ("the metric lives in prose or a badge, not a failing gate"), *rubber-stamp-verified* ("'verified' is set, never earned"), *stubbed-core* ("the most-engineered code has no caller"), *doc-vs-code-drift* ("a countable claim in prose is wrong"), *config-present-but-unwired* ("the flagship path silently no-ops while the right code exists one dir over"), *byte-identical-fork* ("the

name promises X, the code is Y unchanged"). The vocabulary is CLOSED for typing purposes and EXTENSIBLE only by an explicit, visible amendment; an un-typeable finding is an honest tracked GAP (§ . . . / § . . .), never an invented type. The corpus states NO severity scale, so the severity vocabulary is consumer-supplied DATA per § . . . – this anchor mandates none.

(C) THE TYPE IS A ROUTING KEY TO A COUNTER-GATE – THAT IS WHY TYPING IS MANDATORY. Corpus: *"The value of a shared vocabulary is that it turns 'another bug' into 'another instance of type T,' which is what lets you install one gate that kills the whole type. When a reviewer or an agent can tag a finding stubbed-core, the remediation is not bespoke – it's 'add the challenge that fails against a no-op stub,' every time."* Each type therefore binds to the anchor that already owns its counter – green-but-broken → § . . . (four-layer, RUNTIME-ON-CLEAN-TARGET) + § . . . ; coverage-theater → § . . . paired mutation + § . . . (C); rubber-stamp-verified → § . . . (D . . .) status custody + § . . . + § . . . ; stubbed-core → § . . . (no-fakes-beyond-unit) + § . . . (capability matrix reporting attempted-vs-not); doc-vs-code-drift → § . . . + § . . . + § . . . ; config-present-but-unwired → § . . . (investigate-before-remove / unwired-code) + § . . . (F) (configured ≠ in use); byte-identical-fork → § § . . . does NOT restate those counters – it makes the ROUTING mandatory so a finding reaches its counter mechanically instead of by recall.

(D) THE RELEASE-BLOCKING RULE. Verbatim and unqualified: *"a capability with no passing challenge is a release blocker."* A ledger row whose challenge is absent, never executed, or executed only against a stale artifact fingerprint blocks the release exactly as a FAILING

challenge does – this is the § 11.4.236 /
§ 11.4.236 absence-blocks-as-FAIL semantics applied
to the claim side, and it is the reason the ledger is a
gate rather than a document.

(E) **THE ESCAPE HALF.** The corpus pairs the ledger with its
sibling: *"any defect found out-of-band (by a human, not by
the automated regime) is itself a release blocker, because
it proves the automated gates had a hole, and the
remediation must include the new check plus the RED-
capturing evidence that it would have caught the escape."*
That mechanism is already owned by § 11.4.236
(automated QA is the DISCOVERER) – § 11.4.236 cites
it as the ledger's feedback edge and does not duplicate
it.

Honest boundary (§ 11.4.236). The ledger proves that
every ADVERTISED capability has a passing challenge – it
does NOT prove the challenge is strong (oracle strength
stays § 11.4.236 + § 11.4.236 () + § 11.4.236),
does NOT prove the capability is bug-free (§ 11.4.236
discovery-pressure remains), and does NOT enumerate
capabilities nobody advertised (an un-advertised, un-
implemented capability is invisible to a claim-keyed
ledger exactly as an un-claimed feature is invisible to a
feature-keyed one – the two keyings are complements, and a
project running only one has an honestly-stated blind
side). The corpus supplies no worked ledger row, no
severity scale, and no definition of "reality" as a field;
all three are consumer DATA per § 11.4.236.

Provenance correction, recorded per § 11.4.236. Secondary
project documentation has referred to this material as an
"FP-1 ... FP-14" failure-pattern catalogue, an "L-1 ... L-16"
obligations list, and an "F-1 ... F-31" bluff taxonomy.
Measured against the tracked corpus, **no F-n, FP-n, or L-
n identifier exists in any of the nine modules** (control-
needle proven: the same scan that returns for those

forms returns the real `ATM-277` / `ATM-343` / `LVA-098` / `ITE-6` / `API-33` identifiers, and `for` for a negative control). The real structures are: a `-row` numbered failure-pattern table whose last column is "The rock-solid approach we lack"; a `-bullet` portfolio-gate list in `[Q ]`; and this `-member` bluff-type vocabulary in `[Q ]`. Anchors MUST cite the real structures; the `F-n` / `FP-n` / `L-n` labels MUST NOT be used as corpus wording.

Classification: universal (`$ . . .`) – a platform-neutral claim-integrity discipline reusable by any project that advertises capabilities to users; the consuming project supplies its advertised-surface set, severity scale, ledger location, and challenge runner as DATA per `$11 4 116 . 10 4 118 .` Composes `$ . . . / $ . . . / $11 4 119 . . 11 / 11 $ . . .` (feature-keyed coverage ledger – `$11 4 119 . . 11 11 119 .` is its claim-keyed complement) / `$. . . 11 4 146 3 / $11 4 149 . . . / $. . . / $11 4 156 . 10 4 156 . / $. . . / $11 4 157 . 10 4 212 . (F)214 / $. . . / $11 4 226 . 10 4 227 . 11 4 216 $. . . (absence blocks as a FAIL)217 / $. . . (D11)218 / $. . . / $11 4 245 . 10 4 249 . 11 4 211 $. . . (F) / $11 4 254 . 10 4 257 . 11 4 219 $. . . / $. . . / $11 4 262 . 10 1 / $. . . / $. . . / $. . . (the escape edge) / $. . . / $. . . / $. . . / $. . . / $. . . / $. . . / $. . .`

`11 4 217`

Propagation gate `CM-COVENANT-114-266-PROPAGATION` (literal `11.4.266` present as a block-start, exactly-once per governance file, lockstep content-hash equality across the mirror set per `$ . . .` (B)) + recommended mechanism gates `CM-CLAIM-REALITY-LEDGER-COMPLETE` (every advertised capability on every declared advertised surface resolves to exactly one ledger row; an advertised capability with

no row → FAIL naming the capability and the surface it was advertised on) + CM-UNCHALLENGED-CAPABILITY-BLOCKS-RELEASE (a row whose challenge is absent, never executed, or has no verdict for the release-candidate artifact fingerprint blocks the release exactly as a FAIL does) + CM-LEDGER-ROW-TYPED-FROM-CLOSED-VOCABULARY (every row carries a bluff type drawn from the seven-member closed set; an ad-hoc type → FAIL) + paired § . mutations (add a README capability with no ledger row → the completeness gate MUST FAIL; blank a row's challenge reference → the blocking gate MUST FAIL; retype a row to an invented type → the vocabulary gate MUST FAIL; golden-FALSE fixture per § . . . (): a fully-populated ledger whose every row has a fresh candidate-fingerprinted PASS verdict MUST NOT fire any of the three). Gate-code = separate work item, NOT claimed shipped (§ . . . / § . . .).

Canonical authority: constitution submodule Constitution.md § Non-compliance is a release blocker regardless of context. No escape hatch — no --ledger-optional, --unchallenged-capability-ok, --untyped-ledger-row, --invent-bluff-type, --feature-ledger-suffices, --advertise-without-row flag exists.

§ . . . — Shared attempt record + converge-on-evidence / escalate-on-stall: a failed approach is never silently retried, and a non-converging loop is itself a signal (research-derived, — —)

Corpus anchor (AI-curriculum module , [Q] — verbatim):

"No memory of what was tried → the loop repeats itself. Agents re-attempt discarded approaches. Fix: a shared record of attempts and outcomes (in the task state / PR thread), so a failed approach is not silently retried; and an escalation rule – after N failed iterations, stop and route to a human or a stronger model, rather than burning cycles."

The premise that makes it structural, not optional.

Corpus: *"The agent has no memory of prior attempts. It will re-implement the same wrong fix because it does not remember the last three. Prevention: encode the lesson where the agent must encounter it – the test, the spec, a decision record in the repo – not in a human's head."* And: *"'no reopening' is really a statement about durability of knowledge, and agents have no durable memory, so the regime must externalize it."* The record is therefore not a convenience log – it is the externalised memory without which the loop is structurally incapable of converging.

What this is NOT (the § . . boundary, stated first). § . . governs a CRASHED agent (detect → respawn → resume until complete). § . . (C) governs a HUNG long-op (heartbeat delta against a no-progress budget). § . . governs the TERMINAL exhaustion audit written at the moment Operator-blocked is declared. None of the three governs a LIVE, HEALTHY, PROGRESSING loop that is re-walking a dead end its predecessor already proved dead – the agent is not crashed, the op is emitting heartbeats, and no terminal state has been reached. § . . closes exactly that gap.

The mandate (ALL hold):

(A) A SHARED, DURABLE ATTEMPT RECORD. Every multi-attempt effort (an agent loop, a producer verifier cycle, a fix-retry sequence, a multi-agent hand-off) MUST write each attempt and its OUTCOME to a record SHARED across every actor that may take the next attempt – including a successor agent, a respawned agent per § 2.2.1, and a different track per § 2.2.2. The corpus names its location as *"in the task state / PR thread"*; the consuming project binds its own location as DATA per § 2.2.3 (the § 2.2.4 workable-items DB and the § 2.2.5 append-only event stream are the natural substrates and MUST be reused rather than duplicated per § 2.2.6). The corpus specifies no schema and none is mandated here (§ 2.2.7).

(B) CONSULTED BEFORE ACTING – the clause that makes the record load-bearing. The record MUST be READ before the next attempt is chosen, and an approach already recorded as failed MUST NOT be silently retried. Re-attempting a recorded dead end is permitted only when the attempt materially differs (a changed precondition, a corrected § 2.2.8 reproduction sequence, a new § 2.2.9 (F) RED) and that difference is recorded. A written-but-never-read record is the § 2.2.10 (F) configured-≠-in-use failure applied to agent memory, and satisfies nothing.

(C) CONVERGE ON EVIDENCE. Corpus: *"the design principle is converge on evidence, escalate on stall. Each loop must monotonically reduce uncertainty – the producer gets a real failing receipt and must make it green."* The two named preconditions of convergence: a SINGLE SOURCE OF TRUTH (*"the spec + acceptance criteria is the authority, and the verifier's captured failure is the ground truth. Disagreements are adjudicated on the receipt, not on whoever argues more confidently – the verifier's build exit beats the producer's fluent paragraph every*

time") and CONCRETE FAILURE FEEDBACK ("If the verifier just says 'still broken,' the producer tries another plausible change. Fix: feed the exact captured error (argv, exit code, failing assertion, the offending line) back as the next input"). Composes § 11.4.21 / § 11.4.22 (the verifier is not the producer) and § 11.4.23 (the receipt, not the narration).

(D) STALL IS A DETECTED SIGNAL WITH A BOUNDED RETRY AND AN ESCALATION PATH. The corpus names TWO stall signatures – "the same criterion red after N tries" and "oscillating diffs" – and states the response: "after N failed iterations, stop and route to a human or a stronger model, rather than burning cycles." The stall itself is tracked, not merely survived: "a loop that is not converging ... is itself a signal, tracked and escalated. Frequent producer-verifier divergence is a metric worth trending: it predicts where the producer over-claims and where the spec is ambiguous." The forbidden shape is explicit: "The anti-pattern is an unbounded retry loop with no ground-truth feedback and no escalation – that is how agents thrash expensively while appearing busy." Escalation composes § 11.4.24 (route to a stronger tier when complexity is PROVEN mid-flight), § 11.4.25 (decide autonomously where the choice is safe and reversible; park only what is genuinely blocked), and § 11.4.26 / § 11.4.27 (the operator path, reached only after the § 11.4.28 self-resolution audit).

(E) N IS CONSUMER-SUPPLIED DATA, NEVER A CONSTITUTIONAL LITERAL (§ 11.4.29 / § 11.4.30). The corpus leaves N symbolic in every one of its occurrences and states no numeric bound, no wall-clock budget, no token ceiling, and no cost ceiling. This anchor mandates a BOUND EXISTS and is DECLARED; it mandates no value. A project that enforces an invented number and attributes it to the corpus commits the § 11.4.31 violation this clause prevents. **EXTENSION**

(`ceiling` `literal` `operator`): where the operator mandates a ceiling on `N` for a specific, named loop, that ceiling is recorded with its own verbatim mandate in its own anchor (`$loop_anchor` for the independent-review loop) – this is the only form of `N`-bound literal this clause permits; the consumer's declared `N` within that ceiling remains consumer DATA, never the literal ceiling itself.

(F) **THE LOOP'S DONE-CONDITION.** A loop **MUST NOT** read satisfied while an attempt is recorded in-flight with no outcome, and an escalation raised under (D) **MUST** reach a terminal disposition (resolved / re-tiered / operator-parked per § 2.2.1) – never be silently dropped. Composes § 2.2.1 / § 2.2.1 / § 2.2.1 / § 2.2.1 (a parked escalation parks ONE work unit, never the loop) and § 2.2.1 (a started effort reaches a terminal state).

Honest boundary (§ 11.4.17). The attempt record makes the loop's history legible and prevents the RECORDED dead end from being re-walked – it does NOT guarantee convergence (a problem may be genuinely hard, and an escalation is the honest outcome, not a failure of the record), does NOT prove any individual attempt was competent, and does NOT substitute for § 11.4.17's systematic-debugging discipline on the underlying defect. It bounds waste and preserves knowledge across actors; it does not manufacture insight.

```

Classification: universal (§ 1.1) – a platform-
neutral orchestration discipline reusable by any project
running multi-attempt or multi-agent loops; the consuming
project supplies its record location, schema, stall
detector, N, and escalation route as DATA per
§ 1.1. Composes § 1.1 / § 1.1 /
§ 1.1 (the terminal exhaustion audit –
§ 1.1 is its running, pre-terminal

```


counterpart) / § . . / § . . /
 § . . / § . . / § . . /
 § . . / § . . / § . . /
 § . . (F) / § . .
 (the append-only event stream substrate) /
 § . . / § . .
 (crash → respawn; § . . governs the healthy-
 but-looping case it does not reach) / § . . /
 § . . (F) /
 § . . / § . . / § . . /
 § . . (C) (hung long-op; § . .
 governs the progressing-but-not-converging case) /
 § . . / § . . / § . . /
 § . .

Propagation gate **CM-COVENANT-114-267-PROPAGATION** (literal **11.4.267** present as a block-start, exactly-once per governance file, lockstep content-hash equality across the mirror set per § . . (B)) + recommended mechanism gates **CM-ATTEMPT-RECORD-SHARED-AND-CONSULTED** (every multi-attempt effort writes each attempt + outcome to the declared shared record, and the record is read before the next attempt is dispatched; a second attempt duplicating a recorded-failed approach with no recorded material difference → FAIL) + **CM-STALL-BOUNDED-AND-ESCALATED** (a declared retry bound exists; a loop exceeding it without a recorded escalation to a terminal disposition → FAIL) + paired § . mutations (strip the record-read step so the loop dispatches blind → the first gate MUST FAIL; remove the retry bound leaving an unbounded loop → the second MUST FAIL; record an escalation with no terminal disposition → the second MUST FAIL; golden-FALSE fixture per § . . () : a loop that converges on attempt with both attempts recorded and no bound exceeded MUST NOT fire either gate). Gate-code = separate work item, NOT claimed shipped (§ . . / § . .).

Canonical authority: constitution submodule Constitution.md

§ . . . Non-compliance is a release blocker regardless of context. No escape hatch – no --no-attempt-record , --record-without-reading , --unbounded-retry-loop , --retry-same-failed-approach , --stall-without-escalation , --invent-N-and-cite-corpus flag exists.

§ . . . – Tamper-evident evidence chain + periodic anchor record: deletion, reordering, and tail-truncation of an accepted evidence record are DETECTED, not merely content-mutation (spec-derived, speckit
11 4 115 -anti-slop-enforcement,
11 4 116 - -)

Compact summary: every accepted evidence record (the § . . . (F) machine-written verdict store, the § 11 4 226 . . . event stream, or an equivalent evidence-bearing ledger a consuming project maintains) MUST be tamper-evident against THREE distinct attacks, not only the content-mutation attack § . . . (F)/ , § . . . already cover: () deleting an entry, () reordering two entries, and () truncating the tail – and the mechanism that detects each MUST be named, because a forward-hash-chain alone detects () and () only when the chain is left internally inconsistent, and does NOT detect an attacker who deletes an entry and then RECOMPUTES the chain forward from that point (spec.md: , FR- 's own measured boundary: "an actor who deletes an entry and then recomputes the whole chain is NOT detected by the chain alone (. s for k entries); that case is detected only against an anchor") – closing that gap requires a SEPARATE periodic anchor record the producer can append to but never rewrite (spec.md: - , FR- -). **The mandate**

(ALL hold): (A) EVERY STATE-CHANGE ENTRY WITHIN THE DECLARED COVERED-CALL SET, NO CITED-READ VOLUME CREEP.

Scope (§ . . .): this clause binds the enforcement system's own tool-call ledger under § -anti-slop-enforcement (spec.md: , FR- is written against THAT ledger, not every write/exec/deploy on the host) – the consuming project declares its own concrete COVERED-CALL SET (which state-changing operations this evidence chain tracks) as DATA per § . . . ; § . . . does NOT itself mandate a fleet-wide "every write/exec/deploy anywhere produces a chain entry" obligation, only that within the declared set, coverage is complete and cited-vs-uncited is honoured. A state-changing call WITHIN the declared covered-call set (write, exec, deploy – spec.md: , FR-) MUST produce a matching chain entry naming it. An UNCITED read of the store – a read whose result backs no claim – MUST NOT be required to produce an entry (spec.md: , FR- : *"this guards against both volume creep and a gate that refuses correct behaviour"*); a read result CITED in support of a claim MUST have a chain entry (spec.md: , FR-). The distinction is CITED-vs-UNCITED, never present-vs-absent-in-storage. **(B) CHAIN VERIFICATION DETECTS DELETION + REORDER WITHOUT RECONSTRUCTION.** Deleting one entry, or reordering two entries, WITHOUT reconstructing the chain that follows, MUST be reported by the integrity check (spec.md: , FR-) – mutation-detection of a single entry's content alone does NOT satisfy this; the check walks the FULL forward chain, recomputing every link, not spot-checking. This clause's own honest limit is stated in (C): a reconstructed (recomputed) chain defeats (B) alone. **(C) THE ANCHOR IS THE ONLY MECHANISM THAT CATCHES A RECOMPUTED-FORWARD DELETION OR A TAIL TRUNCATION.** A periodic ANCHOR entry – chain head digest AND entry count, recorded to a location the producer can APPEND to but NOT REWRITE (spec.md: , FR- : *"the*

count is what makes wholesale store deletion a detected absence rather than silence") – MUST be written on a declared interval. The anchor's STRENGTH is recorded honestly as `mechanism` (cryptographically or append-only enforced – a git commit history, a write-once object store, or the consuming project's declared equivalent per § 3.3.1) or `policy` (procedurally enforced only); `mechanism` MUST NOT be claimed without evidence that rewriting is mechanically prevented, and absent that evidence the recorded strength is `policy` (spec.md: 3.3.1, FR-3.3.1). Verification against the anchor detects a tail truncation that chain-alone verification PASSES cleanly on (spec.md: 3.3.1, FR-3.3.1: "the two results together are the evidence that the anchor, not the chain, carries this property") – so "the chain verified clean" is NEVER sufficient on its own; a verdict of tamper-free MUST name BOTH the chain-walk result AND the anchor cross-check result. (D) VERIFICATION THAT CANNOT COMPLETE REFUSES – IT NEVER REPORTS THE CHAIN INTACT. A chain truncated mid-verification, or an anchor location unreachable at verification time, yields REFUSAL (an UNVERIFIED status, per § 3.3.1's conservative-safe-default-on-an-unresolvable-signal), never a "clean pass by default" (spec.md: 3.3.1, FR-3.3.1). "Could not verify" and "verified clean" are never the same reported state. **Honest boundary** (§ 3.3.1). § 3.3.1 detects deletion, reordering, and tail-truncation of ALREADY-RECORDED entries. It does NOT prove the first entry ever written was itself truthful – the entry's CONTENT is only as trustworthy as the harness that wrote it (§ 3.3.1 (F) machine-written, § 3.3.1 producer≠verifier for who may write it), and it does NOT make the underlying evidence CORRECT, only tamper-EVIDENT. Where a project already operates a content-addressed, atomically-committed store for a related purpose (the § 3.3.1 continuum resume engine),

§ 11.4.267 REUSES that mechanism's chaining primitive rather than inventing a second one – a consuming project MUST NOT stand up a parallel, divergent chain implementation when one already satisfies this anchor's properties (§ 11.4.266 extend-don't-duplicate applied to mechanism reuse, not only to rule text).

Classification: universal (§ 11.4.268) – a platform-neutral tamper-evidence discipline reusable by any project maintaining an evidence/verdict ledger; the consuming project supplies its concrete chain implementation, anchor location, anchor-interval, AND covered-call set (clause A – which state-changing operations this project's evidence chain tracks) as DATA per § 11.4.269. Composes

§ 11.4.270 / § 11.4.271 / § 11.4.272 /

§ 11.4.273 (F) (the verdict entries THIS anchor chains) / § 11.4.274 (the event stream this anchor may chain) / § 11.4.275 (verification-cannot-complete → refuse, the conservative-safe default) /

§ 11.4.276 (durable/atomic/ append-only construction discipline) / § 11.4.277 (reuse its content-addressed store where one already exists) /

§ 11.4.278 (evidence-class-at-closure – a chain entry's own class must match the defect layer it backs) /

§ 11.4.279 (machine-created evidence at every gate – the chain entries this anchor makes tamper-evident ARE the machine-created evidence artefacts § 11.4.280 requires; § 11.4.281 is the integrity property that keeps those artefacts trustworthy AFTER the moment of PASS, not merely present at it) / § 11.4.282 (a genuine tamper finding is itself a destructive-op-class event requiring the § 11.4.283 backup-and-investigate protocol) /

§ 11.4.284. Propagation gate CM-COVENANT-114-268-PROPAGATION (literal 11.4.268 present as a block-start, exactly-once per governance file, lockstep content-hash equality across the mirror set per § 11.4.285 (B)) + recommended mechanism gates CM-EVIDENCE-CHAIN-DELETE-REORDER-DETECTED

(chain-walk verification detects a deleted or reordered entry left internally inconsistent) + CM-EVIDENCE-CHAIN-ANCHOR-CATCHES-RECOMPUTE-AND-TRUNCATION (an anchor record exists on the declared interval, carries head-digest + entry-count + an honest mechanism / policy strength, and anchor-cross-check verification detects a recomputed-forward deletion and a tail truncation that chain-alone verification misses) + CM-EVIDENCE-CHAIN-INCOMPLETE-VERIFICATION-REFUSES (a chain or anchor that cannot be reached/completed at verification time yields REFUSAL, never a default clean pass) + paired § . mutations (delete a middle entry and recompute the chain forward from that point using only internal chain data → chain-alone verification MUST still PASS [proving the gap the anchor exists to close] while anchor cross-check MUST FAIL; truncate the tail and verify chain-alone → MUST PASS while anchor cross-check MUST FAIL; make the anchor location unreachable → verification MUST report UNVERIFIED, never PASS; golden-FALSE per § . . (): an untampered chain with a fresh, reachable, correctly-computed anchor MUST NOT fire any of the three gates). Gate-code = separate work item, NOT claimed shipped (§ . . / § . .).

Canonical authority: constitution submodule Constitution.md

§ . . (pending landing). Non-compliance is a release blocker regardless of context. No escape hatch – no --content-mutation-detection-suffices , --skip-anchor-record , --claim-mechanism-strength-unverified , --chain-truncation-passes-by-default , --reconstructed-chain-is-clean flag.

§ . . – Critic/consensus advisory-only ban AT THE EVIDENCE-ACCEPTANCE SEAM: an ungoverned confidence/consensus signal MAY inform but MUST NEVER substitute for a receipt or adjudicate producer-verifier disagreement; role recorded advisory, never entered into an evidence chain – the CONSTITUTION'S MANDATORY INDEPENDENT-REVIEW-VERDICT FAMILY

(§ . . / . / . / . / . / . / . IS UNWEAKENED (spec-derived, speckit – anti-slop-enforcement, – –)

Compact summary: at the evidence-acceptance / claim-adjudication seam – where a producer's claim of DONE/PASS is accepted or refused, and where a producer-verifier DISAGREEMENT is adjudicated – an ungoverned second opinion (a confidence score, an LLM-judge output, or a consensus aggregation across multiple such opinions, none of them produced under this constitution's own independent-review discipline) is a real, adoptable capability, but it MUST be admitted in an ADVISORY-ONLY role at THAT seam, never as the mechanism that decides accept/refuse there. This project's own operator decision on the point

(spec.md: , A-) states the rule directly:

"Current rules adjudicate producer-verifier disagreement on the recorded receipt, never on confidence or argument quality, and explicitly bar an LLM judge from being the gate." An ungoverned critic signal that CAN flip an

accept/refuse outcome at that seam has re-created a self-certification collapse under a new name – it is a second oracle/gate sharing correlated blind spots with the

producer (same model family, same training data, same failure modes as an LLM producing the work it is asked to judge), exactly the class § . . /§ . .

already forbid a producer from occupying. This anchor is

NARROWER than "no second model may ever review a

producer's work" – it does NOT reach, weaken, override, or

substitute for the constitution's SEPARATE, already-mandatory independent-review-verdict family

(§ 4.239 . . . code-review-before-build,
 § 4.237 . . . iterate-to-GO with rock-solid evidence,
 § 4.236 . . . universal every-change review,
 § . . . independent verification agent,
 § . . . Opus- xhigh review substrate,
 § . . . translation context-and-spirit review,
 § . . . independent per-language translation review).

Those seven anchors mandate a STRUCTURALLY-SEPARATED verifier's verdict (§ . 11 . 134 producer≠verifier, § . . . oracle/gate roles) as a GENUINE receipt-generating gate on the work itself – a NO-GO from any of them BLOCKS the build (§ . . .) and RE-ARMS the iterate-to-GO loop (§ . . .) – and every one of the seven remains fully in force, unmodified, unrelaxed. Clause (F) below states this unweakened-list explicitly and mechanically. **The mandate (ALL hold): (A) TWO CLOSED TERM DEFINITIONS, EXCLUDING THE MANDATORY REVIEW FAMILY.** critic = a second (or Nth) model instance, confidence score, or LLM-judge that offers an OPINION on a producer's claim, diff, or artifact WITHOUT itself constituting a structurally-separated, receipt-generating review conducted under § . . . /

§ . . . / § . . . / § . . . /
 § . . . / § . . . / § . . . –

i.e. an ungoverned, ad-hoc, or supplementary second opinion, not the mandatory independent-reviewer verdict those anchors already require and gate on. consensus = a mechanism aggregating multiple such UNGOVERNED opinions (voting, averaging, an LLM-judge score) into one combined signal. Both are ADOPTABLE capability categories a project may wire in – § . . . does not forbid their existence, only their misuse as a gate AT THE EVIDENCE-ACCEPTANCE / CLAIM-ADJUDICATION SEAM. A verdict produced by the constitution's own mandatory review family is NOT a

"critic" for the purposes of this anchor – it is the receipt-generating verifier role § . . . / § . . . already govern, and it continues to gate exactly as those seven anchors require. (B) **A CLAIM WHOSE ONLY SUPPORT IS AN UNGOVERNED CRITIC/CONSENSUS SIGNAL IS REFUSED.** Per spec.md: (FR-): *"Given a claim whose only support is a critic consensus or confidence score, the seam MUST refuse it"* – naming the missing receipt. A recorded receipt – a machine-written § . . . (F) verdict, a captured evidence-chain entry per § . . . , OR a mandatory independent review's own recorded GO/NO-GO under § . . . / § . . . / § . . . / § . . . / § . . . – remains the **ONLY** basis on which producer-verifier disagreement is adjudicated; confidence or argument quality UNBACKED by such a receipt, however produced, is never a substitute. (C) **DETERMINISTIC WITH AND WITHOUT THE UNGOVERNED CRITIC.** Per spec.md: (FR-): the SAME claim, evaluated once with an UNGOVERNED critic step present and once with it absent, MUST produce the IDENTICAL accept/refuse outcome. A pipeline in which disabling that ungoverned critic changes the verdict is, by definition, using it as a gate – the deterministic-outcome test is the mechanical proof that (B) is genuinely honoured, not merely stated. (This clause does NOT apply to a mandatory review under the seven-anchor family in clause (A) – skipping a REQUIRED § . . . / § . . . / § . . . review is itself a violation of those anchors, not a deterministic-outcome test to pass; the test in this clause is scoped to the OPTIONAL, ungoverned critic only.) (D) **UNGOVERNED CRITIC OUTPUT IS RECORDED advisory IN A SEPARATE STORE, NEVER ENTERED INTO AN EVIDENCE CHAIN.** Per spec.md: (FR-): ungoverned critic output MUST be recorded with a role of advisory and MUST NOT appear in any

§ . . . /§ . . . (F) evidence chain – it is recorded in a separate advisory store, never intermixed with the accept/refuse evidence chain itself. A mandatory review's OWN recorded GO/NO-GO under the seven-anchor family continues to be entered into the evidence/verdict record as its own anchors require (§ . . . 's review-completed marker, § . . . 's clean-GO record) – it is not retroactively downgraded to advisory by this clause. **(E) WHY THE UNGOVERNED-CRITIC COLLAPSE MATTERS.** Composes § . . . (A)/(C) and § . . . – an ungoverned critic wired into the evidence-acceptance gate collapses ORACLE and GATE into a single mechanism that is NOT independent of the producer's own failure surface (correlated model errors, shared blind spots), which is precisely the class of collapse § . . . names as maximum-severity. Advisory-only admission avoids re-creating that collapse under the label "review" while retaining the genuine usefulness of an ungoverned critic – surfacing likely issues earlier, feeding § . . . systematic debugging, prioritising a human reviewer's attention – none of which requires it to hold gating authority. The mandatory seven-anchor review family avoids the SAME collapse by a DIFFERENT mechanism entirely – structural separation of the reviewer per § . . . , not advisory-demotion of its verdict – which is exactly why that family's verdicts remain full receipts and this anchor never touches them. **(F) THIS BAN IS SCOPED TO THE...EVIDENCE-ACCEPTANCE / CLAIM-ADJUDICATION SEAM – IT DOES NOT...WEAKEN, OVERRIDE, OR SUBSTITUTE FOR THE MANDATORY INDEPENDENT-REVIEW-VERDICT FAMILY.** § . . . (code-review-agent gate before pre-build + main build), § . . . (iterate-to-GO with rock-solid evidence), § . . . (universal code-review – every change reviewed, always), § . . . (independent verification agent), § . . . (code-review MUST run on the Opus model

at xhigh effort), § . . . (context-and-spirit translation review), and § . . . (independent per-language translation review) each mandate a STRUCTURALLY-SEPARATED verifier (§ . . . producer≠verifier) whose VERDICT genuinely gates, and NONE of these seven anchors is relaxed, narrowed, made advisory, or made optional by § What § . . . forbids is narrower and different: (i) treating an UNGOVERNED second opinion / confidence score – one NOT produced under any of those seven anchors' own independence-plus-evidence discipline – as a substitute receipt (clause B); and (ii) letting such an ungoverned signal, rather than a recorded receipt (including a receipt produced by the seven-anchor family), adjudicate a producer-verifier DISAGREEMENT (the A- scope, spec.md:). A project MUST NOT cite § . . . to justify skipping, weakening, downgrading-to-advisory, or making optional any of the seven listed mandatory reviews – doing so is itself a § . . . (F) violation, exactly the mis-scoping this clause exists to foreclose. **Honest boundary** (§ . . .). § . . . does not forbid using an ungoverned critic/consensus signal to prioritise human review, to surface candidate defects for § . . . investigation, or as ONE input among many an independent human/verifier weighs when authoring (not executing) a gate. It forbids only its use as an automated GATING mechanism substituting for a real, receipt-backed oracle/gate/verifier chain. It does not claim critics are useless – only that their output is advisory context, never a verdict, UNLESS that output is itself produced under the mandatory independent-review-verdict family named in clause (F), in which case it is a receipt and this anchor does not touch it. **Classification: universal** (§ . . .) – a platform-neutral discipline reusable by any project that adopts, or is asked to adopt, a second-model/critic/consensus review capability OUTSIDE

its own mandatory independent-review-verdict family; the consuming project supplies its concrete critic implementation and advisory-record location as DATA per § 11.4.125, and its own set of already-mandatory review-verdict anchors (analogous to this project's

§ 11.4.125 • § 11.4.124 / § 11.4.142 • § 11.4.145 / § 11.4.219 / § 11.4.217 / § 11.4.216 / § 11.4.215) as the fixed, unweakened-by-this-anchor list. Composes

§ 11.4.216 / § 11.4.215 / § 11.4.214 (F) / § 11.4.213 / § 11.4.212 / § 11.4.211 / § 11.4.210 / § 11.4.209 / § 11.4.208 / § 11.4.207 / § 11.4.206 / § 11.4.205 / § 11.4.204 / § 11.4.203 / § 11.4.202 / § 11.4.201 / § 11.4.200 (the seven mandatory review-verdict anchors clause (F) leaves unweakened) / § 11.4.199

(the separation principle this anchor is a specific instantiation of) / § 11.4.198 (the four-role architecture that names why an ungoverned critic-as-gate is a collapse) / § 11.4.197 (the evidence chain the ungoverned critic's output MUST NOT enter) /

§ 11.4.196 (the legitimate advisory use – feeding investigation) / § 11.4.195. Propagation gate CM-

COVENANT-114-269-PROPAGATION (literal 11.4.269 present as a block-start, exactly-once per governance file, lockstep content-hash equality across the mirror set per

§ 11.4.269 (B)) + recommended mechanism gate

CM-CRITIC-CONSENSUS-ADVISORY-ONLY (a claim supported only by an ungoverned critic/consensus/confidence signal, with no receipt, is refused; the accept/refuse outcome for a claim is identical whether an ungoverned critic step ran or was skipped; every stored ungoverned-critic output carries role advisory and appears in no evidence chain; a mandatory review under

§ 11.4.269 / § 11.4.268 / § 11.4.267 / § 11.4.266 / § 11.4.265 / § 11.4.264 / § 11.4.263 / § 11.4.262 / § 11.4.261 / § 11.4.260 is NEVER downgraded to advisory

or excluded from gating by this gate) + paired § 11.4.259 mutation (wire an ungoverned critic's approve/reject value directly into the pass/fail decision → the gate MUST FAIL;

store an ungoverned critic output with no role field, or role `verdict`, or entered into the evidence chain → the gate MUST FAIL; downgrade a `review`'s own GO/NO-GO record to `advisory`, or exclude it from gating, citing `golden-FALSE` per `golden-FALSE` as the reason → the gate MUST FAIL; golden-FALSE per `golden-FALSE` (): (i) an ungoverned critic present, its output stored with role `advisory`, absent from the evidence chain, and the deterministic outcome unchanged whether it ran or not → the gate MUST NOT fire; (ii) a `code-review`'s own NO-GO record blocking the build exactly as `code-review` requires → the gate MUST NOT fire). Gate-code = separate work item, NOT claimed shipped (`code-review` / `code-review`). **Canonical authority:** constitution submodule `Constitution.md` `Constitution.md` (pending landing). Non-compliance is a release blocker regardless of context. No escape hatch – no `--critic-may-gate`, `--consensus-score-is-a-receipt`, `--llm-judge-is-the-gate`, `--critic-output-in-evidence-chain-OK`, `--outcome-may-differ-with-critic`, `--269-overrides-mandatory-review`, `--downgrade-mandatory-review-to-advisory` flag.

`golden-FALSE` – **Dependency-existence-verdict register:** every proposed dependency carries a closed-set existence verdict {VERIFIED, AMBIGUOUS, UNVERIFIED} with citable evidence; adoption is gated behind resolution and behind the wiring-sweep result (spec-derived, speckit –anti-slop-enforcement, – –)

Compact summary: a project that adopts an external (or internal, cross-module) dependency to close a coverage gap MUST first record, for that dependency, a closed-set

EXISTENCE verdict – VERIFIED, AMBIGUOUS, or UNVERIFIED – backed by a citable artifact or an explicit statement that none was found (spec.md: , FR-). Existence verification is NECESSARY but NOT SUFFICIENT for adoption (spec.md: , A- : "Existing is not the same as working ... Several verified entries are at initial release versions, one declares a repository that returns , and one's 'tamper-evident' claim is on inspection only 'append-only' – which detects nothing if the appender is the party being checked") – a VERIFIED existence verdict is the FLOOR a dependency clears before it may even be considered adopted, never a claim that it is fit for the purpose it is being adopted for. **The mandate (ALL hold): (A) THE REGISTER, ONE ROW PER CLAIMED DEPENDENCY.** Every dependency a project proposes to adopt gets exactly one register row, keyed by a stable identifier (name + version + source per § . . resolve-by-stable-name). The row carries the verdict + its evidence field (spec.md: , FR- : "MUST carry a verdict of VERIFIED, AMBIGUOUS or UNVERIFIED with a citable artifact or an explicit statement that none was found"). **(B) VERDICT SEMANTICS, EVIDENCE-BASED, NEVER PROMOTED BY OMISSION.** UNVERIFIED is the honest default state – a claim about a dependency's existence with no independent confirmation attempted MUST NOT be silently treated as VERIFIED. AMBIGUOUS is assigned when evidence conflicts or a name collision exists – spec.md: (FR-): "a dependency name that collides with an unrelated project" gets AMBIGUOUS, never VERIFIED – the specific conflicting evidence is recorded, not merely a vague "unsure". VERIFIED requires an independently-reproducible check: a real resolve + checksum/hash match, a real successful call or import, or an equivalent control-needle-proven presence per § . . () (b). **(C) ADOPTION REFUSED WHILE THE VERDICT IS NON-VERIFIED.** Per spec.md: (FR-): "Given a dependency whose verdict is UNVERIFIED or

AMBIGUOUS, adoption MUST be refused." Adoption is admitted ONLY...after the row resolves to VERIFIED (or the AMBIGUOUS conflict is investigated and the row is re-verdicted). (D) **ADOPTION IS GATED BEHIND MAPPING TO A RECORDED UNCOVERED GAP + BEHIND THE WIRING-SWEEP RESULT.** Per spec.md: (FR-): adoption is refused unless the dependency maps to one of the project's own recorded uncovered capability gaps (adopting a dependency for a capability the project already covers is itself refused, naming the covering rule – the § . . duplicate-coverage discipline applied to dependency adoption). Per spec.md: (FR-): *"the adoption step MUST NOT be accepted before the wiring step reports its sweep result"* – a project does not decide "yes, adopt" before it has run the sweep that would reveal whether the dependency's claimed capability is genuinely reachable and wired (§ . . /§ . . (F) investigate-before-adopt, the mirror-image of investigate-before-remove applied to the opposite direction). (E) **ZERO ADOPTED WITHOUT VERIFIED – MONOTONIC.** Per spec.md: (SC-): the count of external dependencies adopted without a VERIFIED existence verdict is ZERO, checked at the SAME seam that already refuses unwired-gate names per § . . (A) (the named-gate ledger), generalised here to dependency-adoption claims. **Honest boundary** (§ . .). § . . 's register proves EXISTENCE/RESOLVABILITY was checked and recorded – it does NOT prove the dependency BEHAVES correctly for the purpose it is adopted for (that stays ordinary integration/contract testing per § . .), and it does NOT prove the dependency is SECURE or well-maintained (that is § . . 's supply-chain territory). A- 's own finding – "existing is not the same as working" – is why this anchor is deliberately narrow: it closes the EXISTENCE-CLAIM bluff, not the fitness-for-purpose question, which stays governed by the constitution's

ordinary testing/review family. **Classification: universal**
 (§ 11.4.270) – a platform-neutral existence-
 verification discipline reusable by any project that
 adopts external or cross-module dependencies to close a
 capability gap; the consuming project supplies its
 register location, its recorded uncovered-gap set, and its
 wiring-sweep mechanism as DATA per § 11.4.270. Composes
 § 11.4.270 / § 11.4.270 / § 11.4.270 (stable-
 name-keying) / § 11.4.270 (investigate-before-adopt,
 the direction-mirror of investigate-before-remove) /
 § 11.4.270 (F) (configured ≠ wired, the same principle
 applied to a dependency's claimed capability) /
 § 11.4.270 () (b) (control-needle-proven presence
 checks) / § 11.4.270 (A) (named-gate ledger – this
 anchor generalises its zero-unwired-debt discipline to
 dependency claims) / § 11.4.270 (fitness-for-purpose
 stays contract-testing territory, distinct from
 existence) / § 11.4.270 (security/supply-chain, also
 distinct from existence) / § 11.4.270 (submodule-
 catalogue-first discovery + extend-don't-reimplement – the
 register's VERIFIED-tier check IS the mandated catalogue-
 first discovery step, applied to any dependency, not only
 own-org submodules) / § 11.4.270 (submodule-
 dependency-manifest – an owned-org dependency's `helix-
 deps.yaml` is itself citable evidence toward a VERIFIED
 verdict) / § 11.4.270 (mandatory deep multi-angle
 web research before declaring a capability covered – the
 research pass that PRODUCES a register row's citable
 evidence, or its explicit none-found statement) / § 11.4.270
 Propagation gate `CM-COVENANT-114-270-PROPAGATION` (literal
`11.4.270` present as a block-start, exactly-once per
 governance file, lockstep content-hash equality across the
 mirror set per § 11.4.270 (B)) + recommended mechanism
 gate `CM-DEPENDENCY-EXISTENCE-VERDICT-REGISTER` (every claimed
 dependency has a register row carrying a closed-set
 verdict + evidence field; a name collision resolves

AMBIGUOUS never VERIFIED; a shipped change relying on an UNVERIFIED or AMBIGUOUS row is refused; adoption refused when the dependency maps to an already-covered capability or precedes its wiring-sweep result) + paired \$.

mutation (adopt a dependency with no register row → the gate... MUST FAIL; leave a row UNVERIFIED and adopt anyway → the gate MUST FAIL; feed a name-collision fixture and assert the verdict resolves VERIFIED instead of AMBIGUOUS → the gate MUST FAIL; golden-FALSE per

\$. . (): a dependency VERIFIED with real, reproducible, control-needle-proven evidence, mapped to a genuinely-uncovered gap, adopted after its wiring sweep reported a result → the gate MUST NOT fire). Gate-code = separate work item, NOT claimed shipped (\$. . / \$. .). **Canonical authority:** constitution submodule Constitution.md \$. . (pending landing). Non-compliance is a release blocker regardless of context. No escape hatch – no --unverified-adoption-OK , --ambiguous-treated-as-verified , --skip-wiring-sweep-before-adopt , --adopt-already-covered-capability , --existence-proves-fitness flag.

§ . . – Waiver mechanism: the strict FORMALIZATION of allow-with-tracked-debt (unifying § . . (D)'s recorded-deferral, § 4.2.1 . . ()'s explicit-operator-override, and § . . (A)'s deadline-bound quarantine into ONE roster + expiry + tracked-item schema), distinct from refusal (§ . .) and closure (§ . .) vocabulary (spec-derived, speckit –anti-slop-enforcement, –)

Compact summary: every project needs, on rare occasion, a mechanism to ALLOW a change to proceed despite a normally-blocking missing-evidence condition – never a silent, never a permanent bypass, but a tracked, expiring, human-authorized exception. **This is NOT an absent capability being introduced here** – re-verified in this remediation pass: § . . (D) already permits an explicit, documented deferral flag ("The escape hatch is a recorded deferral, not a bypass: the skipped gate remains owed and is caught at the next full run or tag gate" – Constitution.md:); § . . () already requires "an explicit, logged, operator-set flag" for any override of its deploy-readiness seam ("EXPLICIT, RECORDED OVERRIDE ONLY" – Constitution.md:); and § . . (A) already ships a deadline-bound quarantine for one specific gate class (flaky tests, moved to tests/quarantine/ with a tracked § . . stabilisation deadline). Those three are genuine allow-with-known-debt PRECEDENTS this anchor does not discover but FORMALIZES. What § . . genuinely adds – needed: `waiv*` returns ZERO hits anywhere in constitution/Constitution.md outside this draft, confirming the SCHEMA below is new even though the CATEGORY of mechanism is not – is the STRICT, DATA-BEARING, cross-gate schema those three precedents describe only informally and per-

mechanism: a ROSTERED NON-PRODUCER authoriser (\$. . . (D) and \$. . . (A) name no authoriser; \$. . . () names the operator as a role but no roster-resolvable identity schema – \$. . . adds the roster resolution + producer exclusion), a MANDATORY UNELAPSED EXPIRY (\$. . . 's skip flag and \$. . . 's override flag carry no expiry field at all – an indefinite skip/override is currently legal under their text), and a MANDATORY TRACKED ITEM (neither \$. . . (D) nor \$. . . () requires one; \$. . . (A) requires it for its one gate class – \$. . . makes it mandatory cross-gate). \$. . . turns three ad-hoc, per-anchor escape hatches into ONE cross-cutting mechanism any future gate can invoke by reference instead of re-inventing its own – this project's own `scripts/lib/critical_blocker_gate.sh` (built for the \$. . . seam) already implements `waive / override` operations by name, direct evidence this schema generalises code the project has already written, not a speculative addition. **The mandate (ALL hold): (A) THE WAIVER IS THE ONE FORMALIZED SCHEMA FOR AN ALLOW-WITH-KNOWN-DEBT PATH.** A waiver EXISTS at a seam that would otherwise refuse (a missing receipt, an unresolved dependency verdict, an unmet acceptance-criteria scenario, or any other normally-blocking condition this constitution's family of gates enforces) – the SAME class of seam \$. . . (D)'s deferral flag, \$. . . ()'s override flag, and \$. . . (A)'s quarantine already gate. It is invoked explicitly, never inferred, and it never disables the gate permanently – it grants a bounded, tracked exception to ONE candidate. A gate ALREADY implementing an allow-with-known-debt path under \$. . . / \$. . . /\$. . . is NOT required to reimplement one from scratch – it satisfies

§ . . by conforming its existing flag/quarantine mechanism to this clause's roster + expiry + tracked-item schema (clauses B-D below), the strict superset those three anchors' informal mechanisms currently lack. **(B) AUTHORISATION MUST RESOLVE TO A ROSTERED IDENTITY, NEVER THE PRODUCER.** Per spec.md: (FR-): *"Given a waiver whose authorisation value does not resolve to an identity in the declared authoriser roster, the seam MUST refuse it"* – a free-text authorisation ("looks fine") is refused exactly as an absent one is; a waiver naming a genuinely rostered identity is accepted (the false-positive guard – a check that refuses EVERY waiver has replaced a bypass with a deadlock, not fixed anything). The authoriser roster is consumer-supplied DATA per § . . , and the producer of the change under waiver MUST NOT appear as its own waiver's authoriser (composes § . . producer≠verifier – a producer cannot waive their own gate). **(C) EXPIRY IS REQUIRED AND MUST NOT HAVE ELAPSED.** Per spec.md: (FR-): a waiver with NO expiry, or whose expiry HAS elapsed, is refused; a waiver whose expiry is in the future is accepted. An absent expiry is NOT read as "valid forever" – the spec's own framing is exact: *"An expiry absent by omission is not 'valid for every candidate forever' – that is the same downgrade-by-omission the satisfaction path already refuses"* elsewhere in this family (§ . . (D)/§ . . (F) fingerprint-absent-is-refused, the general pattern this clause instantiates for waivers). An elapsed waiver's gate reverts to blocking automatically – no grace period, no silent auto-renewal. **(D) A TRACKED ITEM NAMES THE DEBT; A WAIVER NAMING NONE IS REFUSED.** Per spec.md: (FR-): *"Given a waiver, it MUST name a tracked item for the blocker that remains open, and a waiver naming none MUST be refused"* – because *"a waiver is the only ALLOW that leaves the defect unfixed, so it is the one*

path, on which the owed work must stay traceable." The tracked item is an ordinary § 11.4.15 -stable-id workable item bound to the SAME lifecycle discipline (§ 11.4.15 / § 11.4.16 / § 11.4.33 / § 11.4.33) as any other tracked work – a waiver whose tracked item never resolves is itself surfaced by the standing zero-shortcomings audit (§ 11.4.15) as an unresolved-waiver finding, never silently forgotten.

(E) COMPOSITION, NOT SUBSTITUTION – INCLUDING WITH THE THREE PRECEDENT MECHANISMS. A waiver whose expiry lapses UNRESOLVED re-enters the § 11.4.30 . Operator-blocked class for its underlying condition (the block returns; nothing about the original blocking condition changed, only the temporary permission to pass did). A waiver's tracked item CAN close via § 11.4.30 . Obsolete if the underlying requirement genuinely changed. Neither § 11.4.30 nor § 11.4.30 substitutes for the waiver mechanism itself – a project attempting to express "allow despite missing evidence, tracked, expiring" via either of those two vocabularies is misusing them, per this anchor's own honest-boundary clause. Symmetrically, § 11.4.30 (D)'s deferral flag, § 11.4.30 ()'s override flag, and § 11.4.30 (A)'s quarantine are NOT replaced by § 11.4.30 – each remains the correct, narrower mechanism for its own seam class; § 11.4.30 is the schema they conform TO when a project wants their allow-with-debt behaviour to carry a rostered authoriser, an expiry, and a tracked item, not a rival mechanism competing with them. **Honest boundary** (§ 11.4.30).

§ 11.4.30 formalizes, it does not discover, the ALLOW-with-tracked-debt CATEGORY – § 11.4.30 (D), § 11.4.30 (), and § 11.4.30 (A) already establish that category exists in this constitution; why § 11.4.30 does not make waivers routine or easy is the SCHEMA's own shape (rostered-non-producer authorisation, mandatory unelapsed expiry, mandatory

tracked item) – deliberately strict, because a waiver is the ONE path this constitution's anti-bluff family permits to leave a defect genuinely unfixed while still shipping. It does not weaken any gate a waiver bypasses for OTHER candidates, and it does not authorise an indefinite or permanently-renewed waiver – an expiry that keeps getting pushed forward without the underlying debt closing is itself a § 11.4.271 finding (a monotone-decreasing-findings-ledger violation, not a compliant waiver renewal pattern). **Classification: universal** (§ 11.4.271) – a platform-neutral, deliberately-strict allow-with-tracked-debt discipline reusable by any project whose gate family needs an occasional, bounded, evidence-honest exception; the consuming project supplies its authoriser roster, expiry mechanism, and tracked-item store as DATA per § 11.4.271. Composes § 11.4.271 / § 11.4.271 / § 11.4.271 / § 11.4.271 (refusal vocabulary – distinct, cited not restated) / § 11.4.271 (closure vocabulary) / § 11.4.271 / § 11.4.271 (stable-id tracked item) / § 11.4.271 (closure vocabulary – distinct, cited not restated) / § 11.4.271 (D) (the absent-field-is-refused pattern this anchor applies to expiry + authoriser + tracked-item) / § 11.4.271 (the tracked item's completion obligation) / § 11.4.271 (D) (recorded-deferral precedent this anchor formalizes) / § 11.4.271 () (explicit-recorded-override precedent this anchor formalizes) / § 11.4.271 (A) (deadline-bound quarantine precedent this anchor formalizes for one gate class) / § 11.4.271 (producer≠verifier – the authoriser exclusion) / § 11.4.271 (an unresolved waiver is a standing-audit finding) / § 11.4.271. Propagation gate CM-COVENANT-114-271-PROPAGATION (literal 11.4.271 present as a block-start, exactly-once per governance file, lockstep content-hash equality across the mirror set per § 11.4.271 (B)) + recommended mechanism gate

CM-WAIVER-ROSTERED-EXPIRY-TRACKED (every active waiver resolves to a genuinely rostered, non-producer authoriser identity; a future, unelapsed expiry; and a named § . . . -stable-id tracked item; a waiver missing ANY ONE of the three is invalid and its waived gate reverts to blocking) + paired § . . . mutation (submit a waiver with free-text or absent authorisation → the gate MUST FAIL; submit one with no expiry, or an elapsed expiry → the gate MUST FAIL; submit one naming no tracked item → the gate MUST FAIL; submit one where the authoriser identity equals the change's producer → the gate MUST FAIL; golden-FALSE per § . . . () : a waiver naming a genuinely-rostered non-producer identity, a future expiry, and a real tracked item → the gate MUST NOT fire). Gate-code = separate work item, NOT claimed shipped (§ . . . / § . . .). **Canonical authority:** constitution submodule Constitution.md § . . . (pending landing). Non-compliance is a release blocker regardless of context. No escape hatch – no

```
--free-text-authorisation-OK , --waiver-without-expiry , --
1 1 4 2 7 2
elapsed-waiver-still-valid , --waiver-without-tracked-item , --
producer-may-self-waive , --permanent-waiver-renewal flag.
```

§ . . . ^{2 0 2 6} — **Dynamic, on-demand skill/extension activation:** the active capability surface is the MINIMUM needed now, everything else stays DISCOVERABLE and one command away (User mandate, — —)

Forensic anchor – verbatim user mandate

(— —) : *"We MUST use actively only Skills we need at particular moment! Use of the Skills and other extensions MUST BE fully dynamic so token use is ALWAYS at the minimum and any of Skills and the extensions activated / deactivated as they are needed / used! Extend the constitution Submodule properly for this to work*

flawlessly and be a part of every project like all other heavy optimizations and improvements we MUST use!"

Forensic case study (FACT, measured - - , Claude Code . . .). A host-level skill pool of skills renders as , **B** (~ , tok) of skill **NAMES** in every turn of every session that loads it; the same pool rendered with its description: fields is , **B** (~ , tok/turn) – a . x cliff the consuming project neither controls nor observes. Both figures are exact byte counts; the token conversion is a bytes/ ESTIMATE (no tokenizer on the measuring host) and is labelled as such per § . . .

Every project whose agent platform exposes a capability surface (Agent Skills, plugins, extensions, tool sets) **MUST** keep the **ACTIVE** surface at the minimum genuinely needed, with everything else activatable **ON DEMAND** – never a static all-on set paid for in every turn. ALL of the following hold:

(A) MINIMAL DECLARED CORE, EVERYTHING ELSE ON DEMAND. The consuming project declares, as DATA (§ . . .), a deliberately small **core** set that is always active, and the pool it may draw from. Anything not **core** is inactive by default and activated only when a task needs it. An entry added to **core** is paid for in EVERY turn of EVERY session and **MUST** earn that.

(B) DISCOVERABILITY IS NEVER TRADED FOR TOKENS. A deactivated capability **MUST** remain **FINDABLE** – an always-available index lists the whole catalogue with its active/inactive state. **superpowers:using-superpowers** (§ . . . (B)) continues to bind unconditionally: if any skill could apply, even at % relevance, it **MUST** be found in that index and invoked rather than improvised from memory. **Saving tokens by making a capability unfindable is FORBIDDEN** – that trades a measurable cost

for an unmeasurable one and is a § . bluff at the capability layer.

(C) **FAIL SAFE, NEVER SILENT** (§ . .). A failed activation **MUST** tell the agent explicitly that the capability is **NOT** available, and exit non-zero. An activation that appears to succeed while the capability is unusable leaves the agent believing it holds a tool it does not – the false-null class this constitution treats as a critical defect. A dangling or unreadable activation **MUST** be rolled back, not left behind.

(D) **CONCURRENCY-SAFE BY REFERENCE COUNTING**

(§ . . / § . . 11 4 100). Where several agent sessions share one checkout, one session's deactivation **MUST NOT** strip a capability another session is using. Claims are reference-counted and released only when no other claimant is **PROVEN** live (§ . . liveness, § . . ownership-verified) – a claim whose holder cannot be proven live is reported as such, never silently assumed dead OR alive.

(E) **HONEST ASYMMETRY – MEASURE WHAT THE PLATFORM ACTUALLY DOES** (§ . .). The engine **MUST** be built on the platform's **PROVEN** behaviour, not its documented behaviour, wherever the two are not both established. Measured first-hand on this platform: **activation hot-loads** (a newly-linked skill appears in the **ALREADY-RUNNING** session with no restart) but **deactivation does NOT hot-unload** (the same session still invokes a deleted skill from cache). The saving is therefore a **minimal session-start baseline plus on-demand growth**; claiming that deactivation frees tokens in the live session is a bluff. Equally measured: the platform's own `permissions.deny Skill(name)` rule **blocks invocation without removing the entry from the rendered list**, so it saves **ZERO** tokens – a first-party mechanism must be **TESTED** for the property being relied on, never assumed from its documentation (§ . . reuse-

before-reimplement is satisfied by evidence, not by the existence of a similarly-named feature).

(F) A BUDGET GATE GUARDS THE RENDER-MODE CLIFF. The per-turn cost depends on a render mode the project does not control. A gate **MUST** budget the active surface in BOTH the name-only and the description-eager render, so an upstream switch to eager rendering ^{11 4 24} breaks a gate **LOUDLY** instead of silently costing tokens ^{11 4 80 11 4 177} in every turn thereafter.

(G) BOTH MODES – WITH AND WITHOUT THE CONSTITUTION SUBMODULE (REQ-NOTE-). The engine lives in the constitution submodule (\$. , inherited by reference, never copied – \$. /\$.), but a project that does NOT vendor the submodule **MUST** still get a working command or an HONEST, LOUD refusal – never a silent no-op. Resolution order is explicit: an environment override, the in-project submodule, any ancestor's, an operator-configured host checkout, a documented host fallback, then honest degradation naming WHAT is unavailable, WHY, WHERE it looked, and WHAT still works. Degradation **MUST** use an exit code distinct from "operation failed", and **MUST** state that the existing ^{11 4 132} static capability set is untouched so the agent does not ^{2026 09 07} conclude a capability exists when it does not.

(H) PRUNE-TO-CORE IS THE DEFAULT, AND IT IS NOT A \$. . REMOVAL (operator decision, - -). Session initialisation **MUST prune by default** – demoting every non-core capability so a session pays only for the declared `core` – rather than merely raising `core` and leaving whatever else accumulated. A raise-only baseline satisfies clause (A) in letter and voids it ^{...} in substance: the surface grows monotonically as capabilities accumulate, and the operator's standing requirement that token use be at the minimum is never actually met (measured on the reference project: active / B / ~ tok-per-turn under raise-

only, against active / B / ~ tok-per-turn under prune-by-default). **Pruning is NOT the silent capability removal** § . . forbids, and the distinction is structural, not a promise: § . . forbids removing an end-user CAPABILITY without asking, and pruning removes none, because all three properties that make it reversible and visible hold BY CONSTRUCTION – () **LISTED**: the always-core catalogue index of clause (B) still shows every pruned capability, so nothing becomes unfindable; () **ACTIVATABLE**: one command returns it, hot-loading into the running session per clause (E); () **NOT DELETED**: a link-style activation loses only the LINK (its target untouched), and a capability whose only copy lives in the active location is **MOVED** to a preserved store that is itself a searched source – never removed, zero bytes destroyed. A pruned capability is **dormant and one command away**, which is materially different from dropping a component from a build. A demotion that cannot be performed safely – another live session holds a claim per clause (D), the preserved store already holds an entry of that name (never clobber preserved content, § .), the move fails or lands unreadable (it is REVERTED), or the capability would not be re-activatable afterwards – **MUST** refuse LOUDLY and leave the capability ACTIVE per clause (C); refusing is always preferable to risking content. Relocating a capability **MUST** be relocation-SAFE: path-relative references inside it that would break at the new location are re-pointed at the identical resolved target before the move, and where that is not sufficient the post-move readability check refuses and reverts (§ . . – never leave a broken capability behind). An **explicit, documented opt-out** **MUST** exist at both session scope (an environment override) and project scope (a manifest flag), with session scope winning; a mechanism with no way out gets worked around destructively, and the opt-out keeps the engine in play rather than abandoned.

still ACTIVATABLE, and its content byte-identical to its pre-prune state; an unsafe demotion refuses and leaves the capability ACTIVE) + paired \$. mutation (squeeze the budget below the real surface → the gate MUST FAIL; break an activation target → activation MUST refuse and leave no dangling link; remove the engine → the shim MUST exit with its distinct engine-absent code, never ; make the demotion DESTRUCTIVE, or strip the relocation-safety step, or strip the other-live-session claim check → the prune gate MUST FAIL). Gate-code beyond the shipped scripts = separate work item, NOT claimed shipped (\$. . / \$. .). **Canonical authority:** constitution submodule Constitution.md \$. . . Non-compliance is a release blocker. No escape hatch – no --all-skills-always-on , --skip-activation-engine , --hide-inactive-skills , --silent-activation-failure , --ignore-surface-budget , --assume-skill-loaded flag.

§ 0 0 0 . . – Measuring-instrument
verification: a census that INFORMS a decision
must be control-needed before its result is
believed (research-derived,
– –)

Forensic anchor (FACT, – – – FIFTEEN instances in a single session). Across one working session, fifteen separate ad-hoc measurements each returned a confident, plausible, ACTIONABLE-looking result that was an artefact of the instrument rather than a property of the system. Every one was indistinguishable, at the moment it was read, from a real finding:

- . A governance-anchor census matched only one of two block-opener forms (^### \$N , missing ^**\$N), under-counted the corpus by ~ mandates and manufactured

phantom dangling anchors that were nearly reported as defects.

- . `comm` on inputs sorted numerically rather than lexicographically reported **all anchors missing** from a file that contained none of them missing.
- . `$?` read after a pipeline reported the **pipeline's last stage**, turning a script's exit into a reported exit – a failure read as success.
- . A field-name typo (`CMA_PROVIDER_KEY_VAR` for `CMA_PROVIDER_KEYVAR`) returned `<none>` for every record, reading exactly like "the records have lost their key bindings".
- . `grep -w` treated `-` as a word boundary, so `zai` matched inside `zai-coding-plan` and moved a record from the NOT-PINNED to the PINNED column – corrupting the classification a recommendation was being built on.
- . A summary-line `grep` assumed `passed, failed` ordering; on failure the tool prints `failed, passed`, so the match came back EMPTY and the run's outcome was unknown while looking merely quiet.
- . `while read` silently dropped a final line with no trailing newline, so a one-element list iterated zero times and printed nothing.
- . A mutation was applied with the wrong environment knob (`CMA_TEST_SCRIPTS_DIR` where the suite reads `CMA_SCRIPTS_UNDER_TEST`), so the mutation tree was never loaded, the REAL tree was tested, and the resulting green was **meaningless while looking like proof the guard held**.
- . A "does the guard block this?" probe used a command the guard does not target, so it ran unblocked and proved nothing either way – an inconclusive result that reads as a negative.

- . A test fixture named each alias identically to its provider id, so two columns held the same string and a real column-swap mutation **survived at** / under a comment asserting no two cells could share a value.
- . A positive control was chosen OUTSIDE the instrument's declared scope – the needle was added by a different layer than the one under test – so a working instrument reported as broken.
- . A positive control was chosen that WAS the very question under test, so it could not discriminate between "the instrument works" and "the answer is yes".
- . An `awk` range extracting a table cell used an escaping that matched nothing, reporting **rows where the live system held** – a number wrong by `x` that was about to gate a DESTRUCTIVE database recreate, caught only because it contradicted a count taken seconds earlier.
- . A denylist with TWO needles was searched with ONE. Both literals were loaded and both lengths were even PRINTED, then only the second was passed to the search. Every "zero remaining" was blind to the first – and the artefacts it cleared were regenerations of a CREDENTIAL SCRUB, so a live credential was reported eliminated while it was still being handed over in four copy-pasteable lines per file. The report was internally consistent, cited a positive control, and was wrong.
- . **`$?` clobbered by a command substitution in the same statement.** `printf '... exit=%s' "$(basename "$f")" "$?"` expands the substitution FIRST, so `$?` carried `basename`'s status, not the command's: a `pandoc` run that FAILED with exit was reported as `exit=0` while writing nothing. Distinct from instance (exit code after a PIPELINE) – here the clobber happens inside the argument list of the very statement that reports it.

The gap this closes. § . . . governs GUARDS AND GATES – instruments that BLOCK, where the failure modes are a false refusal and a false pass.

§ . . . () governs ANALYZERS, which must be calibrated golden-good / golden-bad. NEITHER covers the far more common instrument: the **ad-hoc census that INFORMS** – a `grep`, a set difference, an exit-code read, a field lookup, a row count – whose output is not a verdict but a FACT that a decision, a fix, a classification, or a report is then built upon. Those instruments are written inline, used once, trusted implicitly, and never themselves tested.

The mandate. Any measurement whose result will drive a decision, a claim, a classification, or a report **MUST** be control-neededled BEFORE its result is acted on or communicated:

- **(a) POSITIVE control** – a value KNOWN to be present must be found by the instrument. This catches a pattern too narrow, a wrong field name, a wrong path, a wrong knob, and a stage that silently produced nothing.
- **(b) NEGATIVE control** – a value KNOWN to be absent (a fabricated needle) must NOT be found. This catches a pattern too broad, an accidental substring or word-boundary match, and a default that fabricates a hit.
- **(c) An EMPTY result is never a finding on its own.** Absence of output is one of: the thing is genuinely absent; the instrument could not look; the instrument looked in the wrong place; the instrument's output shape differed from what was matched. Those are FOUR different states and the raw empty result distinguishes none of them (§ . . .). An empty result may be reported only once (a) has established the instrument can see a positive.
- **(d) A result that CONTRADICTS directly-observed evidence is presumed to be an instrument fault until proven**

otherwise. When a census reports absent something that was just read on screen, the instrument is the first suspect, not the system.

- **(e) The instrument used to establish a FACT is itself part of the evidence.** A captured claim cites the command that produced it, so a reader can re-run it and a later reader can see what was actually measured
(`$ . . / $ . .`).
- **(f) When the criterion is a SET, the control must cover EVERY member of it.** A denylist, a literal list, a taxonomy, a multi-column key: needling ONE member proves only that member is searchable, and a search that silently drops the others still reports a clean, internally-consistent, positively-controlled zero. This is the failure mode instance records, and it is the most dangerous on this list precisely because (a) and (b) both PASS while the answer is wrong: the control and the search shared the same single needle, so the instrument proved itself against exactly the case it was already handling. Enumerate the set, assert the search covers it, and state the member count in the captured claim.
- **(g) Capture an exit status BEFORE any other expansion can overwrite it.** `$?` is the status of the LAST command to run, which in a compound statement is frequently not the command being reported on – a pipeline's last stage (instance) or a command substitution sitting earlier in the same argument list (instance). Assign it to a variable on its own line, immediately after the command, and report the variable.

Honest boundary (`$ . .`). Control needles prove an instrument can distinguish PRESENT from ABSENT. They do NOT prove it measures the RIGHT property – a correctly-needed census of the wrong thing is still the wrong

answer, and no amount of needling substitutes for asking whether the question is the right one. Nor do they make an instrument correct for inputs unlike the needles; a needle chosen to resemble the expected data will not reveal a failure on data that does not.

Classification: universal (§ 11.4.17) – platform-neutral; the consuming project supplies its concrete instruments and needle values per § 11.4.16. Composes § 11.4.15 / § 11.4.14 (an un-needed census is a guess wearing a number) / § 11.4.13 (the analyzer analogue) / § 11.4.12 (an enumerated coverage claim rests on the instrument that enumerated it) / § 11.4.11 (the guard/gate analogue) / § 11.4.10.

Propagation gate `CM-COVENANT-114-273-PROPAGATION` (literal `11.4.273`) + recommended gate `CM-CENSUS-CONTROL-NEEDED` (a decision-bearing census cites its positive and negative control) + paired § 11.4.273 mutation (strip the literal → propagation gate FAILs; accept a decision-bearing census with no control → `CM-CENSUS-CONTROL-NEEDED` FAILs; gate-code = separate work item, NOT claimed shipped § 11.4.273 / § 11.4.272).

Canonical authority: constitution submodule `Constitution.md` § 11.4.272. Non-compliance is a release blocker. No escape hatch – no `--skip-control-needle`, `--empty-is-absent`, `--trust-the-grep`, `--census-needs-no-proof`, `--obvious-result-exempt`, `--one-needle-covers-the-set`, `--read-status-after-substitution` flag exists.

\$. . – Mechanical work belongs in a script, not in an agent's context: extract it, test it, and re-scan for it on a cadence (User mandate, – –)

Forensic anchor (FACT, – –). In one working session, five independent review agents consumed roughly , – , tokens EACH. Reading their transcripts, the dominant cost was not judgment – it was MECHANICS: copy a tree to a scratch directory, apply a one-line mutation, run a suite, read the summary line, restore, verify the restore byte-identical, compare counts against a previous round, repeat for seven mutations. Every step of that loop is deterministic, has a fixed input and output, contains no decision, and produces the same result every time. It is a shell script that was instead executed one tool call at a time by a language model, at roughly four orders of magnitude the cost, with a fresh opportunity for a transcription error at every step – and this session recorded ten such instrument errors (\$. .).

The rule. Work that is DETERMINISTIC, REPEATABLE and contains NO JUDGMENT MUST be delegated to a properly written script or small compiled utility, and invoked – not performed turn-by-turn inside an agent's context. Concretely, the following belong in a tool and not in a transcript: mutation apply / run / restore / verify loops; suite fan-outs and their result tabulation; before-and-after comparisons; scratch-tree setup and teardown; residue and quiescence scans; census and inventory queries; artifact fingerprinting; evidence collection and layout; and any step already written down as a fixed recipe.

(a) **EXTRACT.** When an agent performs the same mechanical sequence more than twice, that sequence is a defect in the tooling, not a task. It is extracted into a script or

utility, given a name, and invoked thereafter. The bar is behavioural, not aesthetic: if the steps can be written down exactly, they can be executed exactly, and an agent executing them by hand is the wrong instrument.

(b) THE EXTRACTED TOOL IS ITSELF GOVERNED. It ships with documentation a reader can act on (§ 4.1.1), a user guide, or manual entry where it is operator-facing, and tests that validate its behaviour and produce MACHINE evidence deterministically (§ 4.1.2 / § 4.1.3) – including its failure paths, its exit-code semantics (§ 4.1.4), and the control needles that prove its own measurements (§ 4.1.5). An extraction that replaces a slow correct process with a fast unproven one is a regression, not an optimisation.

(c) RE-SCAN ON A CADENCE. Every project periodically scans itself for mechanical work still being performed by agents – recurring command sequences in transcripts and logs, repeated inline recipes across scripts, and steps documented as "run these commands in order" – and refactors the findings out. This is a standing, recurring obligation, not a one-off cleanup: new mechanical loops accrete continuously as features land, and a codebase that was clean last quarter will not be clean this one.

(d) HONEST BOUNDARY – what MUST NOT be extracted (§ 4.1.6). This rule is about MECHANICAL work, and misapplying it is worse than not applying it at all. JUDGMENT MUST NOT be scripted: root-cause determination (§ 4.1.7), review verdicts (§ 4.1.8 / § 4.1.9 / § 4.1.10), the decision that a fix is correct, the choice of what to measure, the reading of whether captured evidence actually supports a claim, and any assessment requiring the question "is this the right thing to be checking?" A script that renders a verdict is a bluff gate (§ 4.1.11), and automating a review into a green checkmark is precisely the PASS-

bluff this constitution exists to prevent. The correct division is: the SCRIPT gathers and executes, deterministically and cheaply; the AGENT decides what it means. Cost is never a reason to move a decision into a script, and a script's output is evidence for a judgment, never a substitute for one.

Classification: universal (§ 11.4.17) – platform-neutral; the consuming project supplies its concrete scripts, utility language, scan cadence, and transcript/log sources per § 11.4.35. Composes § 11.4.18 (script documentation) / § 11.4.273 (deterministic results) / § 11.4.274 (captured machine evidence) / § 11.4.275 (token efficiency – this is its mechanical-work instance) / § 11.4.276 (exit-code and real-condition semantics) / § 11.4.277 (the extracted instrument is control-needed) / § 11.4.278 (investigate before deleting the thing you replaced) / § 11.4.279.

Propagation gate CM-COVENANT-114-274-PROPAGATION (literal 11.4.274) + recommended gates CM-MECHANICAL-WORK-EXTRACTED (a mechanical sequence performed by an agent more than twice has an extraction item) and CM-EXTRACTED-TOOL-DOCUMENTED-AND-TESTED (every extracted tool carries docs + tests + deterministic machine evidence) + paired § 11.4.274 mutation (gate-code = separate work item, NOT claimed shipped § 11.4.275 / § 11.4.276).

Canonical authority: constitution submodule Constitution.md § 11.4.274. Non-compliance is a release blocker. No escape hatch – no --skip-extraction, --agent-can-do-it, --one-off-is-fine, --script-the-review, --extract-without-tests flag exists.

\$. . – Semantic (Lumen) index + indexing-efficiency gate + universal agent/subagent accessibility: every indexed space reaches every agent and subagent, is proven COMPLETE before it is called ready, and earns its place by a MEASURED, honest efficiency result (research-derived, – –)

Forensic anchor (FACTS, measured in one cycle on a consuming project with a multi-hundred-thousand-file tree; genericised). () A semantic code-search server was configured, yet its own status reported files, no embedding model, and "never" indexed – configured is not indexed (\$. . (F)); the constitution had no mandate for it at all. () A fixed-fixture benchmark measured a . . . % token reduction for the structural index over grep+read – computed only over the queries BOTH routes answered correctly – while the same structural index answered . of fixture queries WRONG (grep+read: of); the semantic route had comparable answers because it was unindexed. A token saving reported without its wrong-answer count would have been a bluff. () One per-call, uncached, full-table file-list query accounted for . % of the reference-resolution wall time (hundreds of calls at roughly ms each on a ~ k-file index). () A foreign-key cascade on bulk re-insert stalled the bulk window on large tables. () A launcher classified sync by rows already stored instead of by pending backlog and ran a bulk job of tens of millions of references on the stock runner, foreground and unsupervised. () The vendor deletes a lock by age, not by holder liveness, and a writable open recreates dropped indexes without the lock – a concurrent-writer corruption hazard. () An empty result from an instrument that had errored or been silenced was nearly read as "absent" (\$. .).

The mandate (ALL hold).

(A) UNIVERSAL ACCESSIBILITY – every indexed space, every agent, every subagent, no exception. Every code index the project maintains (the structural CodeGraph index of § 11.4.78 AND the semantic Lumen index, and any future code index) MUST be reachable by EVERY AI agent the project uses AND by every subagent those agents dispatch, through a server DEFINITION the agent or subagent can actually resolve (project-scoped committed configuration or an installed plugin – never a per-user toggle only the main session sees). A configuration that ENABLES a server which nothing DECLARES (no definition resolvable where the agent or subagent runs) is a violation, not a partial setup. Proof is a real tool call FROM A DISPATCHED SUBAGENT returning an index-only fact (the § 11.4.78 () unforgeable challenge), per agent; an agent that genuinely cannot be driven is a documented § 11.4.78 SKIP, never a faked PASS.

(B) SEMANTIC INDEX PROVEN, NOT CONFIGURED. The embedding model is provisioned (rootless per § 11.4.78) and proven by a REAL embed call returning the declared dimensionality – a model listing alone is not evidence. The semantic index's scope is GENERATED from the SAME class source as the structural scope (§ 11.4.78 () – (), via the constitution's scripts/lumen/gen_lumenignore.py), so third-party and secret classes are excluded identically. Health = files AND chunks non-zero + a known-symbol search hit + a negative control + freshness within a consumer-declared TTL.

(C) COMPLETE BEFORE "READY". An index is called ready only when its PENDING work is zero (unresolved-reference backlog, files still to parse, embedding queue), its state reports complete, its file count is within the consumer tolerance of the accepted expected count, and its required schema indexes exist. Calling a partial index ready is a

§ . PASS-bluff. Every count and every empty result behind this verdict is control-needled (§ .); an instrument that errored or was silenced proves nothing.

(D) EFFICIENCY IS MEASURED DETERMINISTICALLY AND REPORTED HONESTLY. A benchmark harness runs a FIXED query set with GOLDEN answers against the index route AND a grep+read baseline, $N \geq$ iterations with an identical canonical result hash (§ .), recording per query: bytes, estimated tokens (declared estimator), p/p latency and correctness against the golden answer. Token reduction is computed ONLY over queries BOTH routes answered correctly; wrong answers are reported separately, BY QUERY ID, and never averaged away. The harness is self-validating (a tampered golden fixture is rejected; a mutated estimator is caught). The thresholds – minimum reduction, maximum wrong-answer count, latency ceilings – are consumer DATA (§ .), never constitutional literals. A route that saves tokens but is wrong more often than the baseline is NOT an efficiency win, and a report that shows only the saving is a bluff.

(E) ROUTING FOLLOWS THE MEASUREMENT. Structural questions go to the structural index and conceptual questions to the semantic index BEFORE grep+read (§ .), EXCEPT for query classes the benchmark shows the route answers wrongly – those fall back to grep+read and the gap is tracked (§ .), never hidden.

(F) PERFORMANCE DEFECTS ARE ROOT-CAUSED, PATCHED AND PROVEN. A measured hot spot or stall (clauses ()–() of the anchor) is root-caused (§ .) and fixed in the version-keyed patched runner with RED→GREEN, a paired mutation and an equivalence proof on a byte copy, then gated on every upgrade per § . ().

(G) SAFETY. Writes obey § . () (one writer entry, liveness not lock age, bulk classified by pending

backlog); live reads are read-only/immutable; experiments run on fixtures or copy-on-write copies.

Latent binding (§ 1.1.1.1 pattern). A project that maintains NO semantic index today is bound by clauses (A)–(G) only for the indexes it does maintain; the semantic-index clauses bind the moment it adds one – absence of a semantic index is never evidence that (B)–(E) were satisfied.

Honest boundary (§ 1.1.1.2). This anchor proves an index is REACHABLE, COMPLETE, correctly SCOPED and MEASURED – it does NOT prove every answer correct (the benchmark samples a fixed set; wrong-answer classes stay tracked), does NOT make the benchmark representative of every query shape, and does NOT license narrowing the index scope to hit a threshold (scope narrowing is a § 1.1.1.3 operator decision). The reference benchmark harness exists today as a prototype OUTSIDE the constitution tree; landing it under `<constitution>/scripts/` is owed work, not claimed done. [1.1.1.4 STATUS UPDATE – see the EXTENSION below: the harness has since landed INSIDE `<constitution>/scripts/lumen/`; this sentence's "outside the tree" premise is superseded, tracked-work-remaining is narrowed to the gate CODE only.]

EXTENSION – clause (D) benchmark run live on this consuming project + harness-location correction (research-derived, 1.1.1.5). Clause (D)'s deterministic-benchmark requirement was exercised end-to-end via the sanctioned `constitution/scripts/lumen/lumen_verify.sh` tool. This corrects the honest-boundary sentence above: as of this extension the harness is no longer "a prototype OUTSIDE the constitution tree" – it now LIVES inside `constitution/scripts/lumen/` with its own test file `constitution/scripts/lumen/tests/test_lumen_verify.sh`; only the recommended CM-CODE-INDEX-BENCH gate CODE remains owed, unchanged (§ 1.1.1.6 /§ 1.1.1.7 – a fact stated

as fact, not claimed beyond what landed). A `-query` golden set (`conceptual`, `structural`, `deliberately-unsupported`) was pre-declared via `Grep / Read /direct sqlite3` queries against the index database BEFORE any Lumen query ran, then benchmarked `times` back-to-back producing byte-identical `results.tsv / summary.txt` on every run (`$ . . .` determinism, satisfying clause (D)'s `N≥` `-identical-canonical-hash` requirement for the WHOLE benchmark, not one query). Result: recall `/` `=` `.` over in-scope indexable queries; both FAILs are near-misses at rank (one place past `k=` `)`; token reduction computed ONLY over the `correctly-answered` queries (mean-of-per-query `. %`), the `negative-reduction` outliers reported honestly rather than averaged away – satisfying clause (D)'s `reduction-over-both-correct-only + wrong-answers-listed-by-ID` requirements. The Q `.sh -file` mystery (zero results for a query whose gold file is a `.sh` script) was root-caused with NON-GUESSED evidence per `$ . . .` : a direct `sqlite3` query of Lumen's own index database showed ZERO `.sh` files indexed anywhere (`of` `in-scope .sh` files on disk), confirmed authoritative by Lumen's own source (`internal/chunker/languages.go` 's `supportedExtensions` has no `.sh` , `.kt` , or `.txt` entry) – a genuine chunker-coverage gap, not a scope or freshness defect, tracked per clause (E) + `$ . . . . .` rather than silently worked around. Evidence: `docs/research/lumen_benchmark_20260925/RESULTS_v2.md` (+ `.html / .pdf / .docx` siblings per `$ . . .`), commit `ac532a21df78a447c2aceec30f63476348e93077` . Composes

```
$ . . . / $ . . . / $ . . . /
$ . . . / $ . . . / $ . . . /
$ . . .
```

Classification: `universal` (`$ . . .`) – no hardware/vendor/project literal; the consumer supplies its own-org list, scope DATA, accepted counts, tolerances, TTL,

benchmark.queries/golden answers and thresholds as DATA
per \$ 11 4 ; all mechanism is inherited by reference
(\$ 4 99 11 102 11 / \$. . / \$. .).
Composes \$ 11 1 196 . . / \$. . / \$. . /
\$ 11 4 197 1 4 201 / \$ 224 . . / \$. . /
\$ 11 4 207 1 4 232 / \$. . / \$. . /
\$ 11 4 273 1 1 / \$. . / \$. . /
\$. . / \$. . (D)(F) /
\$. . / \$. . / \$. . /
\$. . / \$. . (C) /
\$. . / \$. .

Propagation gate `CM-COVENANT-114-275-PROPAGATION` (literal `11.4.275`, block-start exactly once per governance file, lockstep across mirrors per § 11.4.275 (B)) + recommended mechanism gates `CM-CODE-INDEX-AGENT-ACCESSIBLE` (every indexed space answers an unforgeable challenge from a dispatched subagent; an enabled-but-undeclared server FAILs) / `CM-CODE-INDEX-COMPLETE-BEFORE-READY` (pending == state complete, count within tolerance) / `CM-LUMEN-MODEL-PROVEN` / `CM-LUMEN-SCOPED-FRESH` / `CM-CODE-INDEX-BENCH` ($N \geq$ identical hashes, reduction over both-correct only, wrong answers listed by id, thresholds from consumer DATA) / `CM-CODEGRAPH-SCOPE-PROOF` / `CM-CODEGRAPH-UPGRADE-PROBE-GATED` / `CM-CODEGRAPH-WRITER-EXCLUSIVE` / `CM-CODEGRAPH-BULK-BY-BACKLOG` / `CM-CODEGRAPH-STALL-WATCHDOG` + paired § 11.4.275 mutations (enable a server with no resolvable definition → accessibility gate MUST FAIL; report ready with pending → completeness gate MUST FAIL; drop the wrong-answer list or compute reduction over all queries → bench gate MUST FAIL; unload the embedding model → model gate MUST FAIL; classify bulk by stored rows → bulk gate MUST FAIL; drop the liveness check → writer gate MUST FAIL; golden-FALSE per § 11.4.275 (C): an unchanged, complete, reachable index MUST NOT fire any gate). Gate CODE is a

separate work item, NOT claimed shipped (\$. . /
\$. .).

Canonical authority: constitution submodule Constitution.md

\$. . . Non-compliance is a release blocker. No
escape hatch – no --index-partial-ready , --enabled-is-
declared , --main-session-only-index , --skip-model-proof , --
reduction-without-wrong-answers , --bulk-by-row-count , --lock-
age-is-liveness flag.

\$. . – Review-round budget: every
independently-reviewed work item converges
within a declared budget of at most
rounds (– , consumer-declared), by front-
loading ground truth, class-exhaustive
enumeration, real-path tests and a declared
threat boundary, with a binding methodology
change at round (User mandate,
– –)

Forensic anchor – verbatim operator mandate

(– – , CRITICAL): "Make sure we never ever
get to such extremely high never ending review rounds
count! Determine exact causes of issues we are facing with
deep planning and exhaustive analysis so all work ALWAYS
gets done with to reviews max! This MUST BE
mandatory part of the constitution!!! Reasonable reviews
rounds count with fully deterministically validated and
verified working results produced!"

**Forensic FACTS (genericised; full citations in the
consuming project's own continuation record).** One
reference consuming project ran four independently-
reviewed work items to , +, and +
rounds in one week. Reading every round's recorded
findings showed the count was driven by five recurring,
avoidable mechanisms, not by an unusually hard problem:
() wrong-model ground truth – a build-configuration fix

and its OWN verifying gate were both built on a simplified mental model of an external build system's inheritance mechanism (values assumed to flow through a plain include chain, when the real tool clears the variable before a node's own include and merges ancestor values in afterwards); the gate even ran the real tool, but on a hand-simplified fixture, so it reported GREEN on a fix that could never work; it was caught only when a reviewer read the tool's own source and rebuilt the fixture on the tool's REAL composition path, after which the fix was rebuilt on the real path and landed in a single remediation round; () **instance-at-a-time closure of a known class** – a security gate was hardened against command-shadowing one command name per round (each of rounds , and found further unhardened names) after round had explicitly considered and declined a mechanical whole-file enumeration because no decision-path classifier existed; in a second item a "close the whole class" instruction was met by fixing a syntactic proxy token rather than the class, and the same defect family returned in four consecutive rounds; () **fixer-authored tests that never drive the real code** – at rounds , , , and again a reviewer showed the newest regression test re-implemented, text-extracted, or never drove a verdict through the logic it claimed to guard, so the reviewer's own mutations of the real code survived the whole suite; () **non-binding escalation signals** – the heuristic-tower anchor (\$. .) was invoked about thirty times across these items (once, at round , with the explicit prediction "point fixes alone are likely to produce a Round of the same kind") and the stall anchor (\$. .) had no declared retry bound anywhere in the project, so neither ever stopped the loop; a reviewer-recommended structural fix (re-executing the gate at its call seam in a scrubbed environment) was deferred in two

successive rounds while in-file point hardening continued;
() **scope drift** – a review loop absorbed an out-of-scope tangent (a separate, unrelated hardening effort surfaced mid-review) and the round counter kept accruing against the original item for the whole of that tangent, inflating its own count substantially before the tangent was explicitly separated out. The same record also contains genuinely new findings that legitimately needed another round (a shell-parser asymmetry that made a core test operator shadowable; a sentinel variable an attacker could pre-set) – iteration as such is not the defect. The defect is rounds spent re-discovering what one front-loaded investigation would have found, or accruing against an item the item never actually owned.

What this is NOT (the § . . . boundary, stated first). § . . . governs the review loop's GO bar (zero findings, zero warnings) and its obligation to keep re-running until that bar is met. § . . . governs the shared attempt record, converge-on-evidence, and an undeclared stall bound N. § . . . governs the heuristic-tower SIGNAL that a primitive is defective. None of the three bounds the review loop's LENGTH or binds a concrete action to a round NUMBER. § . . . adds exactly that – clause (A) the ceiling, clause (E) the round- checkpoint and the ceiling escalation. Clauses (B)–(D) add no new substantive obligation beyond fixing WHEN existing duties must be discharged: (B) binds § . . . / § . . . / § . . . phase / § . . . to a recorded artifact at a fixed point before the first fix; (C) binds § . . . / § . . . () (b) to the round that first meets a nameable class; (D) binds § . . . (A) / § . . . () (d) / § . . . (F) / § . . . (D) to every remediation round's own regression set.

The mandate (ALL hold):

(A) **DECLARED ROUND BUDGET, HARD CEILING**. Every work item that passes through the \$. . . / \$. . . / \$. . . review loop carries a round budget `R_max`, declared by the consuming project as DATA (\$. . .) in the closed range $\leq R_max \leq$; an absent declaration resolves to (the conservative value – the escalation in (E) fires earliest; \$. . .). The ceiling is a constitutional literal because the operator supplied it verbatim (\$. . . 's operator-decided fleet minimum is the same pattern); the `11 - 2 267` range is the operator-mandated ceiling recorded with its own verbatim mandate above, not an invented `N` attributed to an external source (\$. . . (E)) – the consumer's declared `R_max` within that ceiling IS its \$. . . `N` for the review loop, and remains consumer DATA. A round = one independent review verdict on the item, whether on uncommitted or already-pushed work, counted from the first review and recorded with the item's review record; a review that never received its brief, or crashed without a verdict, is recorded but not counted. Splitting, renaming, or re-filing the item does NOT reset the count unless (E) () applies. A tracked item created from a finding raised during a budgeted item's review inherits that item's consumed round count – the exception is a finding outside a scope/threat boundary declared before round under (E)(), which starts a fresh budget; item identity for recurrences follows \$. . . (link, never mint).

(B) **GROUND TRUTH BEFORE THE FIRST FIX**. Before the first fix attempt – not after a reviewer finds the model wrong – for any change or finding whose correctness depends on the behaviour of a system the author does not own (build tool, compiler, interpreter or shell semantics, kernel, runtime, third-party library, protocol), the author MUST record a

ground-truth probe: () the authoritative source of the behaviour (the tool's own source code or normative specification, cited by file and line – not a summary, not memory); () a minimal fixture executed on the REAL tool through its REAL composition path (the actual inheritance, sourcing, linking or dispatch mechanism the product uses – a hand-simplified approximation of that path does not count); () the recorded behavioural model the fix relies on. The fix's own verifying gate or test MUST exercise the same real path the probe established; a gate verified on a simplified model is wrong-layer evidence (§ . . .) and cannot support a GO. Composes § . . . / § . . . / § . . . phase / § . . . – it binds them to a recorded artifact at a fixed point in time.

(C) CLASS-EXHAUSTIVE ENUMERATION ON FIRST INSTANCE. When a finding is an instance of a class (bypass via name resolution, an unguarded external command, a masked character, a mutation-surviving test, a missed parse case), the round that fixes it MUST () name the class; () enumerate every member across the item's whole scope with a mechanical instrument whose null result is control-needle-proven (§ . . . (b)), and record that inventory; () close every member, or scope a member out explicitly with a reason and a tracked item (§ . . .). Fixing only the reported instance when the class is nameable is a § . . . (C) violation. When the enumeration needs a tool that does not yet exist (for example a classifier separating decision-path calls from display-only calls), building that tool IS the remediation work – "no classifier yet" is not a reason to fall back to instance-by-instance fixes. A claim that a class is "closed" MUST cite the inventory artifact; a fix that closes a syntactic proxy instead of the class is an unclosed finding (§ . . .).

(D) REAL-PATH TESTS AND REVIEWER MUTATIONS ADOPTED

VERBATIM. Every regression test added in a remediation round **MUST** drive the real shipped artifact through its real invocation path (§ . . . (A)) – a test that re-implements the guarded logic inline, or runs extracted source text outside the artifact's own process, guards nothing and is a source-defect finding under § . . . (D). Every mutation a reviewer reported as surviving is added to the item's mutation set verbatim and must **FAIL** the suite before the next review is requested (composes § . . . (d) / § . . . (F)). The fixer runs the item's full regression set and the reviewer's previous attacks against the fix before requesting review; the dispatch seam **REFUSES** to start a review when that record is absent (never counted as a round because it is never dispatched). A review that **WAS** actually dispatched and returned a verdict always counts as a round, regardless of whether its triggering record was later found deficient – the record requirement gates **DISPATCH**, it is never grounds to retroactively erase a verdict that already issued.

(E) BINDING METHODOLOGY CHANGE AT ROUND , AND THE CEILING ESCALATION. () At round , and at every later round, the round's fixer **MUST** first record a convergence assessment classifying every finding so far as new-class / same-class-as-earlier / regression-introduced-by-a-fix / wrong-layer-evidence. If two or more findings share a defect class named under (C) (not merely a shared § . . . (D) finding layer) with at least one non-MINOR finding in that shared class, any finding reappeared after being reported fixed, or any finding was wrong-layer evidence, the next round **MUST** be a structural round: a fresh fixer agent instance that did not author rounds – (the reviewer remains structurally separate throughout per § . . .), re-deriving (B) from first principles, re-running the (C) enumeration over the whole

scope, and replacing the approach rather than adding a further layer (§ 4.4.1.1). A structural fix recommended by a reviewer may not be deferred past this point without a recorded operator decision (§ 4.4.1.2). (E) Every review declares, before round R_{max} , its scope and acceptance boundary (and, for adversarial/security reviews, the threat model – which actor capabilities are in scope, and the seam at which the defence lives) so that a zero-finding GO under § 4.4.1.1 is reachable; findings outside the declared boundary are tracked items, not loop-extending findings. The boundary is fixed at round R_{max} and changes only by recorded operator decision – narrowing it mid-loop to reach GO is a § 4.4.1.1 fabrication. (F) If round R_{max} returns anything but GO, the item MUST NOT run round $R_{max} + 1$ under the same approach. It stops and goes to a recorded operator decision (§ 4.4.1.2) carrying: the finding classification, the evidence of which of (B)–(D) was skipped or failed, and the options (re-scope into smaller slices per § 4.4.1.1; replace the approach with the structural design as a new item; or accept residual LOW findings only through a § 4.4.1.1 waiver). Only a recorded operator decision opens a new budget. The escalated item is parked per § 4.4.1.1 / § 4.4.1.2 (F) – it parks only that work unit; the autonomous loop (§ 4.4.1.1) continues all other work, and the pending decision is surfaced in every § 4.4.1.1 progress report until resolved. Every breach is also recorded as a process escape and audited for which of (B)–(D) would have prevented it (the § 4.4.1.1 (C) pattern applied to the review loop).

(F) NOTHING HERE LOWERS THE BAR. The budget bounds the NUMBER of rounds; it never lowers what a round must establish. § 4.4.1.1's zero-finding, zero-warning GO bar is unchanged; its re-run obligation is bounded PER

APPROACH by this anchor's `R_max`, after which (E)() governs rather than a further same-approach round. § . . 's review substrate, § . . (D)'s finding-layer classification and § . . 's machine-created evidence are likewise unchanged. Reaching round `R_max` without G0 never produces a G0, a "good enough", or a silent close: the item stays open until a recorded operator decision under (E)().

Honest boundary (§ . .). This anchor does not promise that every item converges within `R_max` rounds. Some problems are hard, and a determined adversarial reviewer can keep finding real defects in an in-process defence whose threat boundary is wider than its design. What the anchor removes is the class of round inflation the forensic record shows to be avoidable: rounds spent discovering a wrong external-system model, one member of an already-known class, a test that never drove the real code, a stall signal nobody was required to act on, or an unrelated tangent silently accruing against the wrong item. When a problem is genuinely hard, (E)() turns an open-ended loop into a bounded, evidence-backed operator decision. The cap counts rounds, not effort; it does not license a shallower review to stay under it. The operator's own wording ("ALWAYS gets done with to reviews max") is read here as a ceiling with a mandatory escalation at that ceiling, not as a guarantee every item physically closes within it – this reading is recorded explicitly and was surfaced to the operator at landing per § . . .

Classification: universal (§ . .) – no hardware, vendor or project literal; the consuming project supplies `R_max` (-), its review-record location, and its probe and inventory artifact paths as DATA per § . . .

Composes § . . (no guessing – the ground-truth probe is its recorded form) / § . . (operator

decision at the ceiling and for boundary changes) /
 § (conservative default `R_max` =) /
 § (systematic-debugging phase bound
 to round) / § / § (the
 loop this budget bounds; its GO bar is unchanged) /
 § / § (re-scope by
 slicing) / § () (d) (reviewer-authored
 mutations) / § / § () (b)
 (control-needed enumeration) / § (review
 substrate, unchanged) / § (recurrence
 linking for spun-out findings) / § (A) (real
 invocation path) / § (wrong-layer
 evidence) / § (D) (finding layer) /
 § (C) (escape audit pattern) /
 § (fresh actor / structural reviewer
 separation) / § (heuristic tower, made
 binding at round) / § (this anchor
 supplies its `N` for the review loop) / §
 (waiver for residual LOW findings) / §

Propagation gate `CM-COVENANT-114-276-PROPAGATION` (literal
`11.4.276` as a block-start, exactly once per governance
 file, lockstep content-hash equality across the mirror set
 per § (B)) + recommended mechanism gates:

- `CM-REVIEW-ROUND-BUDGET` – every reviewed item has a declared `R_max` in – (or the default) and a round count from its review record; a round `R_max` + dispatched with no recorded operator decision FAILs; an item closed without GO FAILs.
- `CM-REVIEW-ROUND3-CONVERGENCE-ASSESSMENT` – every round $\geq$ carries a convergence classification; a non-structural round following an assessment that triggered (E)() FAILs; an adversarial review with no pre-round-threat boundary FAILs.
- `CM-FIX-GROUND-TRUTH-AND-CLASS-INVENTORY` – a fix touching a declared external-system surface with no ground-truth

probe record, or a class finding fixed with no inventory artifact, FAILs.

Paired \$. mutations: dispatch a round with no operator-decision record → CM-REVIEW-ROUND-BUDGET MUST FAIL; delete the round- classification, or follow a same-class trigger with a point-fix round → the round-gate MUST FAIL; strip the probe record from a build-tool fix, or mark a class closed with no inventory → the third gate MUST FAIL; strip the literal → propagation FAILs. Golden-FALSE per \$. . (): an item reaching GO at round with a round- assessment showing only new-class findings, a recorded probe, and an inventory for its one class finding MUST NOT fire any gate. The propagation gate CM-COVENANT-114-276-PROPAGATION is implemented in the same change (constitution/scripts/gates/cm_covenant_114_276_propagation.sh + its paired mutation test); the three recommended mechanism gates are registered as tracked deferrals (OWED-GATE- ..) in constitution/scripts/gates/gate_ledger_deferrals.tsv – their gate-code is a separate work item, NOT claimed shipped (\$. . / \$. .).

Canonical authority: constitution submodule Constitution.md
\$. . .

Non-compliance is a release blocker regardless of context. No escape hatch – no --unbounded-review-rounds , --fix-instance-not-class , --defer-ground-truth-probe , --narrow-threat-model-to-reach-go , --reset-round-count-on-refile , --defer-structural-fix-past-round-3 flag.