

SpecKit-^{0 0 4} + Superpowers Development Cycle Analysis

SpecKit- + Superpowers Development Cycle: An Exhaustive Analysis

Revision: 3

Created: 2026-09-27T17:33:54Z

Last modified: 2026-09-28T16:22:39Z

Status: active

Scope: this document analyses the SpecKit “004-fast-dev-cycles” feature cycle — its methodology, its real-time execution, its verified measurements, and the governance anchors it exercises — as observed and independently re-verified against the live project state at the moment of writing. Every specific figure below was checked against a real, current source (a tracked file, the agent registry, the operator request-history ledger, or a live command this document’s author ran itself) at authoring time; anywhere a figure could **not** be independently confirmed, it is marked `UNVERIFIED` rather than presented as settled fact, per this project’s own §11.4.6 no-guessing mandate — which this document is itself bound by.

Revision-2 note (2026-09-28): Revision 1 (committed 2026-09-27T22:39Z, constitution commit `c380a44`) covered the Foundational batch (`T001` – `T014`) through T014’s round-10 review, honestly reporting the batch as **not yet GO** — two of three round-10 fixes still unconfirmed at that time. This revision extends the document forward through a full additional day of real, committed development: T014 in fact reached a clean GO (after two **further** review rounds this revision’s own predecessor did not yet know about), Phase 3 (US1)’s RED tests and several implementation tasks landed, a genuine concurrency-race incident was root-caused

and forward-fixed, a substantial body of unrelated already-complete work was discovered sitting staged-and-uncommitted and was recovered, and four [SERIAL] wiring tasks are mid-flight as parallel drafts at the moment of this writing. Sections 1–9 below are **preserved verbatim from Revision 1** (they remain an accurate historical record of the Foundational-batch cycle as it stood at that time) except where a Revision-2 correction note is inserted inline; the new material — timeline, process lessons, measured figures, diagrams, and the current honest status — is in [Section 10](#), added at the end of this document precisely so a reader can see what was already here versus what this revision adds, per this project’s own §11.4.226 evidence-class discipline (never silently rewrite a prior claim; supersede it visibly, with its own new evidence).

Revision-3 note (2026-09-28, later the same day): this revision was requested by the same operator directive Revision 2’s own §11.4.140 quote already cites (docs/requests/history.md entries R-2026-09-27-202244 and R-2026-09-28-112940 , identical text, dispatched twice — the second time via the background queue as BG-20260928-2117), continuing the same “keep this document current” mandate rather than a new, distinct request. It extends the document forward through roughly seven more hours of real, committed work in the **parent** ATMOSphere repository (this constitution submodule itself received no further fastcycle-tool commits in that window beyond what Revision 2 already covered) — the Foundational batch’s own tasks.md checkboxes for T007 – T014 were discovered still unchecked despite Section 10.3’s own proof that the batch reached GO and was committed hours earlier (a real, independently-confirmed checkbox/reality drift, distinct from — but the same class of defect as — the source-present/runtime-absent bluff Section 10 already documents for T020); Phase 3 (US1)’s reconciliation-by-verification pass converted several structurally-broken tests (permanent-fail assertions with no flippable done-signal, and one test that “bluffed green” on an artifact-layer fact standing in for an unproven runtime-layer claim) into real, control-needle-proven assertions; T029 / T030 / T032 (three of the four [SERIAL] wiring tasks Revision 2 left as draft-only) were reviewed, applied, tested, and committed by the conductor; and a large, independently-researched QA/PM findings-triage effort (BG-20260928-1956 , 47 findings, 4 parallel subagents + one serialized single-writer DB pass) landed as 37 files and +38,850/-7,830 lines in the parent repository, exercising the exact fan-out-research + single-writer-serialization pattern this document’s own Section 3 already names as this cycle’s live instance of superpowers:subagent-driven-development . As

with Revision 2, sections 1–10 are **preserved verbatim** except for small inline correction notes; the new material is in [Section 11](#), appended at the end for the same §11.4.226 reason Revision 2’s own note states. One item is reported here as genuinely still **PENDING** at authoring time, per this document’s own no-guessing discipline: two read-only reconciliation subagents (dispatched 2026-09-28T16:16Z and 16:19Z, per the agent registry, keys `f23ce7750cea7245` and `1f72b97fcc1ced00`) were independently investigating the Phase 1 and Phase 2 checkbox states *at the same moment* this section was authored, and neither had produced a real `complete` event in the registry by the time this section’s own investigation concluded — this document does not, and cannot honestly, report their outcome.

Table of contents

1. [Executive summary](#)
2. [SpecKit workflow methodology](#)
3. [Superpowers methodology as applied here](#)
4. [Detailed narrative of the Foundational-batch \(T001–T014\) development + review cycle](#)
5. [Measurements, performance and metrics analysis](#)
6. [Diagrams and schemes](#)
7. [Constitution anchors applied](#)
8. [Lessons learned and recommendations](#)
9. [Verification methodology and honest gaps](#)
10. [Update — 2026-09-28 \(Revision 2\)](#)
 - [10.1 Executive summary of the delta](#)
 - [10.2 Timeline of everything landed since Revision 1](#)
 - [10.3 T014 reaches GO: rounds 11 and 12, closing Revision 1’s open gap](#)
 - [10.4 Process lessons learned today](#)
 - [10.5 Measured figures \(Revision 2\)](#)
 - [10.6 Diagrams \(Revision 2\)](#)

- 10.7 Current honest status and what remains
- 10.8 Verification methodology for this revision

11. Update — 2026-09-28, later (Revision 3)

- 11.1 Executive summary of the delta
- 11.2 Timeline of everything landed since Revision 2
- 11.3 The Phase 1 / Phase 2 checkbox-reality gap — an in-flight, honestly-unresolved finding
- 11.4 ATM-1097: a pandoc-specific YAML-misdetection bug, isolated to one code path
- 11.5 BG-20260928-1956: fan-out research + single-writer serialization, worked example
- 11.6 Reconciliation-by-verification: three more worked examples from Phase 3 (US1)
- 11.7 Measured figures (Revision 3)
- 11.8 Diagrams (Revision 3)
- 11.9 Constitution anchors newly exercised
- 11.10 Current honest status and what remains
- 11.11 Verification methodology for this revision

1

. Executive summary

SpecKit-004 (“fast-dev-cycles”) is a feature specification, implementation plan and task breakdown living at `specs/004-fast-dev-cycles/` in the parent ATMOSphere repository. It implements workable item **ATM-1041** (Type: Task), whose operator brief (quoted verbatim in `spec.md`) is that development iterations on this project have become “extremely slow”, that the reopen rate on closed items is “still extremely high”, and that the project must root-cause both problems by measurement — never by assertion — and then re-engineer the gate/iteration loop to be at least an order of magnitude faster **without losing any of the anti-bluff quality guarantees** this project’s constitution already mandates.

The spec resolves into **8 user stories**, **25 functional requirements** (FR-001 .. FR-025), and **10 measurable success criteria** (SC-001 .. SC-010). The implementation plan (plan.md , Revision 2) decomposes this into **68 plan tasks** (T-A01 .. T-H08 , verified: T-A01 through T-H08 is the full span present in plan.md , confirmed by an exact count of 68 distinct task ids) spanning Phases A–H, itself informed by **46 root-cause candidates** (RC-01 .. RC-46) and **39 recorded decisions** (DEC-01 .. DEC-39) in research.md . The task breakdown (tasks.md , Revision 2) turns this plan into **195 executable, test-first tasks** (T001 .. T195) across **11 phases** — Setup, Foundational, then one phase per user story, then a final Polish/Verification phase — each phase gated by an explicit, evidence-typed checkpoint that pauses for operator approval before the next phase may start.

As of the moment this document was authored, the project has completed:

- **Phase 1 (Setup, T001–T006) and Phase 2 (Foundational, T007–T014)** — implemented against the real tracked tree, currently uncommitted (git status shows the entire constitution/scripts/fastcycle/ directory as untracked, ?? scripts/fastcycle/), pending the T014 independent-review gate reaching a clean GO. This is the deepest, most rigorously reviewed part of the cycle so far, and is the subject of Section 4 below.
- **RED-test drafting** for parts of several later user-story phases (US2/Phase 4 and US4/Phase 6 and US6/Phase 8, per the agent registry’s own task-id citations) performed in **isolated scratch copies**, explicitly **not yet applied to the real tracked tree** — a distinction this document is careful to preserve throughout, because conflating “drafted in scratch” with “implemented on the real tree” would itself be the exact kind of bluff this project’s constitution forbids.

The Foundational batch has now been through **two full rounds of mandatory independent adversarial review** (T014, Opus at xhigh effort, per §11.4.209) — round 9 returned **NO-GO** with 3 BLOCKING + 1 IMPORTANT + 4 MINOR findings; the fixes for those findings then **triggered a real-execution discovery** that no review round had itself caught: a live, host-wide SIGTERM / SIGKILL hazard in one of the batch’s own test drivers, root-caused and fixed with a one-line awk guard. Round 10 then returned a **second NO-GO**, confirming every round-9 fix held and finding **9 new issues** (1 BLOCKING + 3 IMPORTANT + 5 MINOR) that round 9 could not have found because they did not yet exist in the pre-round-9 code. As of the moment of writing, **one** of the three parallel fix agents dispatched against round

10's findings has genuinely completed and been independently re-verified by the conductor (the `host_guard.sh` fix: its test suite grew from 125/125 to 131/131 passing assertions); **two** remain in flight with no confirmed-real completion event yet recorded. This document reports that state honestly rather than assuming completion.

2

. Speckit workflow methodology

This project drives non-trivial development work through the Speckit slash-command pipeline (`/speckit.specify` → `/speckit.clarify` → `/speckit.plan` → `/speckit.tasks` → `/speckit.superspec.execute` , with `/speckit.analyze` and `/speckit.checklist` available at any point for cross-artefact consistency checks). Each stage produces a durable, version-controlled artefact under `specs/<NNN>-<slug>/` , and every later stage is required to trace back to the earlier ones rather than re-deriving scope independently.

. The artefact chain, as actually produced
for -fast-dev-cycles

Artefact	Lines	Purpose
spec.md	269	User scenarios, functional requirements (FR-001 .. FR-025), success criteria (SC-001 .. SC-010), out-of-scope declarations, assumptions. Carries a Clarifications section recording 4 operator answers to scope-narrowing questions asked during / speckit.clarify (review-step depth, consumer-project reach, baseline stratification, and how binding the “10x” target is).
plan.md	1,850	The Phased Implementation Plan (68 tasks T-A01 .. T-H08 across Phases A–H) plus a Phase 2 hand-off section. This is where root causes are turned into concrete engineering tasks.

Artefact	Lines	Purpose
research.md	verified: contains root-cause candidates through RC-46 and recorded decisions through DEC-39 (both counts independently re-derived from the file's own text, not merely cited from another document)	The multi-pass deep-web-research log FR-003 mandates, plus the root-cause classification (CONFIRMED / REFUTED / UNDETERMINED) FR-002 mandates.
data-model.md	contains 25 ### -level headings (entity groups plus their sub-sections; tasks.md itself cites "14 entity groups" as the entity-group count proper)	The Key Entities from spec.md (Cycle Record, Root Cause, Gate, Affected-Set, Escape Mechanism, Handoff Record, Verification Report) expanded into concrete schemas.
contracts/	13 distinct contracts (verified: 52 files in the directory ÷ 4 export formats each = 13; e.g. catch-set-comparison-harness , affected-set-and-verdict-cache , agent-registry-and-handoff , common-conventions — the last defining the shared C-001 .. C-007 conventions every tool in the batch must follow)	Interface contracts each tool in the plan must satisfy — inputs, outputs, exit-code conventions, self-validation requirements — authored before the tool exists.

Artefact	Lines	Purpose
<code>quickstart.md</code>	—	10 independent-test scenarios, one per user story, phrased as “can be fully tested by ...” acceptance procedures.
<code>checklists/tasks-coverage.md</code>	—	The mechanical join proving every <code>FR-*</code> / <code>SC-*</code> maps to a user story and every plan task (<code>T-A01 ..</code>) maps to at least one executable task (<code>T001 ..</code>), so no requirement or plan task is silently dropped when the plan is decomposed into tasks.
<code>tasks.md</code>	476 (frontmatter through the final checkpoint)	195 executable tasks (<code>T001 .. T195</code>) across 11 phases, each task line carrying inline <code>[P]</code> / <code>[TDD]</code> / <code>[REVIEW]</code> / <code>[SUBAGENT]</code> / <code>[SERIAL]</code> markers and its owning plan-task id(s) and <code>FR</code> / <code>SC</code> id(s), so the task ↔ requirement join is re-derivable from the file alone without consulting the checklist.

. The phase / dependency structure

`tasks.md` organises its 195 tasks into 11 phases with hard task-id boundaries (independently re-derived from the file's own `## Phase` headings and task-line ids, not merely copied from a summary line):

Phase	Task range	Story	Goal
1	T001–T006	Setup	Skeleton, consumer-owned data, tool preflight, child workable items
2	T007–T014	Foundational (blocking)	Shared library every contract relies on: C-001 exit codes, C-002 canonical JSON, C-003 determinism, C-004 control needles, C-005 self-validation triple, C-007 host safety
3	T015–T049	US1 MVP	Know exactly where cycle time goes — instrumentation + frozen baseline
4	T050–T085	US2	Gates that are far faster and lose nothing (catch-set comparison, verdict cache, affected-set selection, sharding, batch bisection, gate audit, flake ledger, backstop)

Phase	Task range	Story	Goal
5	T086–T103	US3	Stop reopening — escape-mechanism classification, evidence-class-gated closure refusal, sibling-instance search, recurrence linking
6	T104–T121	US4	Minimal tokens per completed item — governance-subset selection, content-addressed evidence reuse
7	T122–T140	US5	Reliable orchestration — truthful agent registry, crash/quota handoffs, custody sweep
8	T141–T153	US6	A host that does not throttle the work — measured resource attribution, backup-first clean-up, unweakened host-safety limits

Phase	Task range	Story	Goal
9	T154–T165	US7	Everything lands in the constitution, clean and pushed — recursive clean/push verification, propagation/lockstep checks
10	T166–T177	US8	Every consuming project gets the speed-up — enumerate, audit, migrate or honestly NOT-MIGRATED
11	T178–T195	Polish/Verification	Final gate: every SC-001 .. SC-010 gets an evaluated verdict with evidence

Every phase ends with a **Checkpoint** line stating its own pass condition and its own evidence class (`artifact` / `runtime` / `source`), and every checkpoint in `tasks.md` ends with the literal instruction “**Pause for operator approval**” — the plan does not treat phase completion as self-authorizing; an operator gate sits between every phase and the next.

The dependency shape is not a flat sequence: Phase 2 (Foundational) blocks every user-story phase (nothing in Phases 3–10 may start before its checkpoint passes), Phase 3 (US1, the measurement MVP) gates every later phase that claims a saving (its own task `T049` is an explicit **re-planning checkpoint** that recomputes every root-cause’s measured share from the newly frozen baseline before any saving-claiming task may proceed), and Phases 4–8 (US2, US3, US4, US5, US6) are designed to run as **parallel streams** once Phase 3’s baseline exists. Phase 9 (US7 — constitution codification + recursive clean/push verification) depends on the *merged* mechanisms from all of those streams; Phase 10 (US8 — consumer migration) depends on Phase 9’s codification.

. The test-first / subagent-driven / review-gated discipline encoded in every task line

Every task line in `tasks.md` follows the literal format `- [ ] [ID] [P?] [Story] [markers] Description (plan T-Xnn; FR/SC ids)`, and the markers are not decorative — each is a load-bearing constraint the task's own executor must obey:

- **[TDD]** — the failing test is written and *observed* failing for the right reason (its output stored under `qa-results/fastcycle/<story>/red/` with the candidate fingerprint) **before** the implementation exists; a test first observed GREEN is explicitly refused as a guard. Every such test additionally ships its paired §1.1 mutation and, for any null/absent/clean/hit result, a class-matched control needle per §11.4.273 (Section 7 below quotes this anchor's operative text).
- **[REVIEW]** — the task's output is accepted only through the independent review gate on the pinned tier: **Opus at xhigh effort, never Fable, never a lower tier or effort**, in a fresh context on a clean checkout, framed to disprove completion, iterated to a zero-finding zero-warning GO. Related changes are batched into one review request of ordered $\leq \sim 400$ -line slices with a machine pre-check pack attached, so speed comes from *fewer, better-prepared* reviews — never from a thinner or skipped review.
- **[SUBAGENT]** — delegable to a parallel subagent whose scope is *exactly* the files named on its own task line (plus its own fixtures directory); it never edits a contention path, never commits or pushes, and its completion claim is treated as a **prediction** until the conductor re-runs its tests from a clean checkout — the producer is never its own verifier. A **[TDD]** subagent that writes the RED test is never the same subagent that implements the fix that turns it GREEN.
- **[SERIAL]** — touches a single-owner contention path (the workable-items database, the agent-registry writer, the pre-build verification script, the meta-test mutation registry, the regression-guard registry, the commit wrapper, the doc-export sync script, the flash script, `.claude/settings.json`, every constitution anchor and its lockstep mirrors, and every commit/push) — executed by the conductor only, one at a time, never concurrently with another **[SERIAL]** task on the same path.

This four-marker discipline is *how* the spec's own `FR-023` clarification ("keep the review requirement unchanged ... speed comes from batching related changes per review and from machine pre-checks so reviews pass first time") becomes an executable constraint rather than a paragraph of prose — a distinction this project's own governance repeatedly insists on (see §11.4.205 / §11.4.227 in Section 7: a rule authored without its enforcement seam is, per this project's own doctrine, *worse* than no rule, because it stops the scrutiny that would otherwise have caught its absence).

. Superpowers methodology as applied here

`superpowers:systematic-debugging` and `superpowers:using-superpowers` are not invoked as abstract discipline in this cycle — they are visible, concretely, in the shape of the real work performed, and this section shows that work rather than describing the skills in the abstract.

. The Iron Law in practice: the `descendants()` investigation

The clearest, most consequential instance this cycle is the investigation of what the T014 round-8 review had flagged only as “suspected but unconfirmed host disruption” in `test_run_all.sh`'s `R8b-D6` sub-test (the test asserting the mutation driver's own cleanup-signal-handling test fixture, `drvwrap.sh`, dies within a fixed wall-clock bound after receiving `TERM`). Rather than treating the symptom — “the test sometimes takes longer than its bound” — as something to patch by loosening the bound or retrying, the investigating agent applied the systematic-debugging arc in full:

1. **Root-cause phase — read the real failure, reproduce it.** The agent traced the exact mechanism: `descendants()`, the transitive-closure helper the mutation driver's cleanup path uses to find every process it must signal, builds its parent map from `ps -eo pid=,ppid=` output with **no de-duplication** — a later line for the same pid key silently overwrites an earlier one. The `R8b-D6` fixture's fake `ps` shim deliberately appends synthetic lines *after* the real

output, one of which claims PID 1 is a child of the fixture's own fake root. Because PID 1 is the universal ancestor of (almost) every process on a Linux host, once that fake line overwrote `descendants()`'s internal mapping for PID 1, the closure computation transitively pulled in **905 of 907** real process IDs on the host — essentially the whole process table.

2. **Pattern phase — classify, don't guess.** The agent explicitly refused to classify the symptom as `flaky` (which this project's own governance treats as a forbidden vocabulary word absent captured evidence, §11.4.7): “the defect is deterministic and always present; only its timing outcome varies with host state” — host load (measured at the time: load average ~31–36 on a 64-core host, 53 processes directly reparented to PID 1) is a real, *amplifying* factor on how long the resulting collateral signal storm takes to settle, but it is not the root cause, and the report is explicit that this conclusion is a measured fact, not an assumption.
3. **Hypothesis phase — falsifiable, and proven both ways.** The stated hypothesis (“a later duplicate `ps` line for pid 1 poisons the closure, causing real host-wide signal traffic that scales with how many processes are reparented to PID 1”) was proven by a **side-effect-free** measurement run twice (905 pids on the unpatched closure, 1 pid — just the synthetic root — on the patched closure, over the identical input stream), and independently corroborated by a **live** demonstration: a separately process-grouped, unrelated harness process, run alongside the unpatched fixture purely as an observer, received roughly 100 real `SIGTERM` deliveries over about 10 seconds and was killed outright once by the collateral blast — direct, non-speculative proof the hazard is real and currently live in the tracked test file, not theoretical.
4. **Implementation phase — the minimal fix at the proven cause, never the symptom.** The fix is a single `awk` boolean added to the existing closure guard: a candidate is admitted to the transitive-closure “visited” set only if it is a bare non-negative integer string strictly greater than 1 — mirroring the guard `kill_tree()` already applied at *signal time*, moved *into the closure computation itself* so a poisoned parent-map entry can no longer be exploited as an expansion vector. This was verified RED (905 pids) → GREEN (1 pid) in an isolated scratch copy, then 3 clean full-fixture runs against the patched copy, then confirmed to introduce **no regression** against the sibling unit test that independently exercises the same pair of functions through a different fixture. Only after all of that was the fix even proposed for application to the real tracked

file — and it was applied by the conductor, never autonomously by the investigating subagent, because it touched an already-committed, load-bearing test file gating the whole feature's promotion.

This is precisely the “Iron Law: NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST” this project's constitution states for

`superpowers:systematic-debugging` (§11.4.102) — and it is also a textbook instance of this project's §11.4.250 anchor (“heuristic-tower signals a primitive defect”), inverted usefully: rather than *adding* a compensating heuristic on top of a symptom, the investigation found and removed the actual defective primitive.

. **using-superpowers in practice: skill-discovery before improvisation**

The same investigating agent, and the review agents around it, consistently reached for the project's own existing tooling (`lib/host_guard.sh` 's live resource probes, the `triple_harness.sh` self-validation-triple machinery, the already-registered paired mutations in `scripts/testing/meta_test_false_positive_proof.sh`) rather than improvising ad-hoc equivalents — the observable signature of `superpowers:using-superpowers` 's “survey available skills before acting” discipline applied to in-repo tooling as much as to Claude Code's own skill marketplace.

. **A second, smaller but equally instructive instance: `r7f2`**

Separately from the `descendants()` investigation, a different intermittent signal — internally tracked as `r7f2` , a `pid` -related `HUP` signal delivery issue — was investigated across **11 reproduction attempts, including deliberately induced heavy host load**, and was closed as a genuine non-reproducible flake rather than either (a) silently ignored or (b) mis-filed as a fixed defect on the strength of a single non-reproduction. The round-9 independent-review report itself notes plainly, “ `r7f2 pid-HUP` passed once. I did not chase that flake further ” — an honest, evidence-scoped statement rather than a claim of resolution, exactly the discipline §11.4.6 (no-guessing) and §11.4.7 (demotion requires same-conditions evidence) require.

. Detailed narrative of the Foundational-batch (T — T) development + review cycle

. What the Foundational batch is

Phase 2 of `tasks.md` (tasks `T007` — `T014`) builds the shared library every later contract in the plan depends on: `C-001` (exit-code conventions), `C-002` (canonical JSON with a `body_hash` excluding volatile fields), `C-003` (determinism — a `--determinism-check` flag that runs a tool twice and asserts identical output), `C-004` (control needles — a class-matched, known-present probe that must prove an instrument can see before its silence is trusted, directly implementing this project's own §11.4.273 anchor), `C-005` (the self-validation triple: golden-good must PASS, golden-bad must FAIL and pinpoint, negative-control must PASS), and `C-007` (host safety — live `nproc / ulimit -u / ps -L thread-count/ MemAvailable` probes that clamp requested parallelism and refuse to ever *raise* a limit).

Independently verified file inventory (re-derived at authoring time by listing `constitution/scripts/fastcycle/` and excluding test fixtures and the Python `__pycache__` directory, neither of which is part of the batch's reviewed source surface):

File	Lines
lib/fc_common.py	390
lib/fc_common.sh	9
lib/host_guard.sh	220
tests/check_deps.sh	89
tests/run_all.sh	93
tests/test_check_deps.sh	327
tests/test_fc_common_red.sh	752
tests/test_foundational_mutations_descendants_filter.sh	62
tests/test_foundational_mutations_kill_tree.sh	79
tests/test_foundational_mutations.sh	111
tests/test_host_guard_red.sh	507
tests/test_run_all.sh	409
tests/test_triple_harness_red.sh	431
tests/lib/triple_harness.sh	177
Total (14 files)	3,656

(One of the round-10 review dispatch notes cited “14 files, ~3594 lines” — this document’s own independently re-run `wc -l` at authoring time returns 3,656; the ~62-line difference is consistent with the intervening `test_foundational_mutations.sh` split and the B1/B2/B3/I1 fixes landing between that citation and this document’s own count, and is reported here as the honestly re-measured current figure rather than silently repeating the earlier estimate.)

Beyond these 14 reviewed source/test files, the batch also carries **69** `triple_harness` self-validation fixture directories (each holding a `golden-good` / `golden-bad` / `negative-control` triple, per `C-005`) and **10** `host_guard`

fixture shims (fake `ps / id / nproc / awk` binaries used to drive `host_guard.sh` through resource-exhaustion scenarios it cannot safely be driven through on the real host) — both counts independently re-derived by directory listing at authoring time.

. The T review-gate mandate

Task `T014` is the batch's mandatory independent review: "batching one request, slices ≤ 400 lines ... by a separate instance on the pinned tier — Opus `xhigh` only, never Fable, never a lower tier or effort — fresh context, clean checkout, framed to disprove completion; iterate to a zero-finding, zero-warning GO" (`tasks.md` line 86, quoted verbatim). This directly instantiates §11.4.209 (Section 7 quotes its operative text in full) and §11.4.134's iterate-until-GO discipline — a review that stops at the *first* clean pass is explicitly not sufficient; the batch must survive a *fresh*, adversarially-framed re-review after every remediation round.

. Round : NO-GO, and the fixes applied

Round 9's independent Opus review (dispatched fresh, clean checkout, framed to disprove completion) returned **NO-GO** with 3 BLOCKING, 1 IMPORTANT, and 4 MINOR findings, quoted here verbatim from the operator request-history ledger's own persisted subagent hand-back report (`docs/requests/history.md`, entry `R-2026-09-27-173243`):

- **B1** — "The `descendants()` fix broke the M7 mutation, so the T013 driver fails." A prior fix (adding the closure-membership filter described in Section 3.1 above, at the point where round 9 first inspected it) had removed pid 1 from ever entering the visited set — but the existing kill-tree unit test *relied* on pid 1 reaching `kill_tree()`'s own `[ "$p" -gt 1 ]` guard to exercise that guard at all. With the closure filter now excluding pid 1 upstream, the guard's own TERM/KILL checks were never exercised, and a deliberately weakened guard (`M7`, the paired mutation meant to catch exactly this regression) sailed through uncaught — measured directly: "with M7 applied ... the `kill_tree` test still passes (`rc=0`), so the mutation is not caught." A bluff gate, caught before it shipped.

- **B2** — the *new* closure-membership check itself had **zero** paired mutation coverage: removing it entirely still left every test passing, meaning a future regression reintroducing the host-wide SIGTERM hazard (Section 3.1) would have nothing to catch it.
- **B3** — a separately-discovered, pre-existing defect: `lib/host_guard.sh` could **never** refuse work, even when a resource limit was already exhausted — floored to `N=1` unconditionally (measured: `FC_GUARD_ACTIVE_AGENTS=6` with 40 agents already active still returned `N=1 REASON=none`), directly violating this batch's own `C-007` contract and the spec's own edge-case requirement that work degrade “with a named reason, never ... weakening a limit.”
- **I1** — the 60% memory ceiling was only checked when the caller passed `--per-job-mem-kb` , and the batch's own caller (`test_foundational_mutations.sh`) did not pass it — measured: a fixture at 99% memory already used still returned `N=16 REASON=none` , silently skipping the ceiling entirely.
- **M1–M4** — missing `.gitkeep` markers on the Phase-1 skeleton directories (a fresh clone would not reproduce the T002 layout); stale wording in the meta-test banner and a contract comment; `triple_harness.sh` output not following the `C-002` canonical-JSON shape; and warning-level hardening items (a bash job-control stderr leak, a theoretical `pgid -reuse` race window in `fc_common.py` 's process-group reaping).

Every finding above was fixed with its own paired-mutation regression coverage before the batch was resubmitted:

- **B1 + B2** were fixed by *splitting* the combined kill-tree test into two independent tests, each stubbing out the *other* function — `test_foundational_mutations_kill_tree.sh` now stubs `descendants()` entirely and proves `kill_tree()` 's own `-gt 1` guard in isolation; the new sibling `test_foundational_mutations_descendants_filter.sh` stubs `ps` and calls `descendants()` directly, proving the closure filter in isolation with a real fake root as its positive control. This is exactly the fix described in Section 3.1 — and it is also where the *real-execution* discovery of the live host-wide signalling hazard happened: applying this split and then actually *running* the full suite for the first time (rather than trusting the split's own logic) is what surfaced the `descendants()` de-duplication bug, which no review round, static read, or `bash -n` check had caught.

- **B3 + I1** were fixed by adding an `exhaust()` construct to `host_guard.sh` that returns `N=0` with a named `<limit>_exhausted` reason whenever headroom, the agent cap, or the memory budget is at or below zero — reserving the “floor to 1” behaviour only for a genuine positive budget — and by making the memory-ceiling read unconditional rather than gated behind an optional flag.
- **M3 + M4** were fixed by documenting `triple_harness.sh`’s `C-002` exemption (it is inherited shared machinery per the contract’s own inheritance-scoping text, not a violation), adding `disown "$!"` after the offending backgrounded job in `test_check_deps.sh` (root-caused to bash’s own job-notification mechanism reporting an un-disowned `setsid` job as “Killed” on the parent’s stderr), and documenting the `fc_common.py` `pgid -reuse` window via docstring (judged genuinely negligible: a sub-instruction kernel-reap race with no constructible test and no available `pidfd -group-signal` primitive on this host).

. Round : a second NO-GO, and the current state of remediation

A fresh Opus `xhigh` review — deliberately re-framed to specifically hunt for any further instance of the “function lost across `exec`” bug class the `descendants()` fix had exposed (Section 8.1 below names this pattern explicitly) — first **confirmed** every round-9 fix held: the full `run_all.sh` suite ran clean (`run_all: 8 run, 0 failed` , 9m38s wall-clock), all 8 foundational mutations were caught, `host_guard` passed 125/125 assertions, `fc_common` passed 229/229 assertions, and an explicit, deliberate search for any *other* instance of the export-across- `exec` bug class found none. It then returned a **second NO-GO** with genuinely new findings — new because they did not exist, or were not yet independently confirmed, before round 9’s fixes landed:

- **BLOCKING** — `host_guard.sh`’s new `exhaust()` logic had a precedence bug: a confirmed `*_exhausted` reason on one limit could be silently **overridden** by an `unreadable_*` reason on a *different* limit, reverting the result to the wrong `N=1` instead of the correct `N=0` — reproduced three separate ways against the real file.
- **IMPORTANT x3** — that exhausted-vs-unreadable precedence had zero paired mutation coverage (deleting it left the test suite still green); several `host_guard` tests had, as a side effect of the I1 fix making the memory check

unconditional, come to depend on the *live* host's actual memory headroom staying under 60% used — which would genuinely fail during, for example, this project's own documented 42 GB containerised AOSP build; and the R8b-D6 “the driver signalled the real tree” assertion in `test_run_all.sh` was found to pass even when the fixture's `kill()` wrapper refuses every signal, because its log-write happens *before* the gating check — meaning only the 25-second wall-clock bound, not the intended signal-delivery assertion, had actually been catching the regression all along.

- **MINOR x5** — three of the B3/I1/kill-tree fixes independently confirmed catchable by hand still had no *registered* paired mutation; `host_guard.sh` 's own header text contradicted its actual floor-to-1 code (and the test that locked in that contradiction); the `fc_common.py` `pgid -reuse` docstring's reasoning was incomplete (a real, portable fix exists via `os.waitid(..., os.WNOWAIT)` , not requiring the whole approach to change); three of `test_fc_common_red.sh` 's orphan-detection checks had no positive control per this project's own §11.4.273; and the pre-check pack still carried 160 `shellcheck` info/style items across 8 files, of which exactly **2** were genuine findings (`SC2086` at `test_check_deps.sh:32` , plus two instances of `SC2143`) among a majority of false-positive `SC2329` s.

Three parallel fix agents were then dispatched against these findings, scoped to disjoint file sets so none could contend with another:

Fix agent	Scope	Status at authoring time
Fix host_guard exhausted-vs-unreadable bug	lib/host_guard.sh + tests/test_host_guard_red.sh	Confirmed complete, independently re-verified by the conductor (not merely trusting the agent's own report): all three original repro commands now correctly return <code>N=0</code> with the right exhausted reason preserved; the full local suite grew from 125/125 to 131/131 passing (6 new paired tests added, mutation-verified against the pre-fix file: exactly the 5 fix-discriminating tests fail on the unpatched file and 1 insurance test still passes on it); the 3 live-host-memory-dependent test blocks were pinned against a <code>\$MEMOK</code> override rather than left dependent on the actual live host's memory state.

Fix agent	Scope	Status at authoring time
Fix fc_common pgid-reuse + orphan-check control	lib/fc_common.py + tests/test_fc_common_red.sh	In flight — dispatched; no independently-confirmed real-completion event recorded in the agent registry as of the moment this document was authored.
Fix R8b-D6 kill-log gate + shellcheck cleanup	tests/test_run_all.sh + shellcheck items across check_deps.sh / triple_harness.sh / test_check_deps.sh / test_triple_harness_red.sh	In flight — dispatched; no independently-confirmed real-completion event recorded as of the moment this document was authored.

This document deliberately does **not** claim the Foundational batch has reached a clean round-11 GO, because that has not happened yet at the time of writing — two of the three round-10 fixes remain unconfirmed, and no fresh round-11 independent review has been dispatched. Reporting otherwise would be exactly the kind of “reports green while the work is not actually done” bluff this project’s entire governance apparatus exists to make impossible.

Revision-2 correction note (2026-09-28): the batch DID go on to reach a clean GO — but not at round 11. A round 11 (dispatched after this section was originally written) found a **new BLOCKING defect** the two rounds above never touched (a silent-fabricated-verdict hazard under inherited `SIGCHLD=SIG_IGN` , unrelated to the `descendants()` /kill-tree story this section tells), and a subsequent round 12 found one further IMPORTANT finding (a positive-control gap in round 11's own new test) before the batch finally reached GO and was committed. See [Section 10.3](#) for the full, independently-verified account — reported there rather than edited in here, so this section remains an accurate record of what was truly known at the moment it was written.

. Measurements, performance and metrics analysis

Every figure in this section was checked against a real, current artefact at authoring time (a tracked source file this document's author read directly, the agent registry, or the operator request-history ledger's own persisted verbatim subagent reports) — never estimated. Where the task brief that requested this document offered a figure this document's author could **not** independently confirm, that is stated explicitly rather than silently repeated.

. Test-suite pass counts (verified via the agent registry's own captured runs, not re-run by this document's author to avoid resource contention with the live fix agents – see Section)

Suite	Count at round-10 review	Count after the host_guard round-10 fix	Independently re-verified?
test_host_guard_red.sh	125/125	131/131 (+6 new paired tests)	Yes — by the conductor, re-running the 3 original repro commands directly and the full suite, per the registry's own complete event for key 23b531299eda6d38
test_fc_common_red.sh	229/229	unchanged as of writing (its own round-10 fix is still in flight)	Yes, at the 229/229 figure — cited identically in both the round-10 review's own report and its persisted verbatim hand-back in docs/requests/history.md
Full run_all.sh sweep (all test_*.sh files)	run_all: 8 run, 0 failed, 9m38s wall-clock	not yet re-run since the host_guard fix, as of writing	Yes — cited in the round-10 review's registry completion note and again verbatim in its persisted docs/requests/history.md report

Suite	Count at round-10 review	Count after the host_guard round-10 fix	Independently re- verified?
<code>test_foundational_mutations.sh</code> (paired mutations) 5 2	8 caught / 8 (<code>M1</code> – <code>M8</code>)	unchanged	Yes — the file itself (read directly by this document's author) currently registers exactly 8 <code>mut</code> calls and 5 <code>ctl</code> calls

. **Mutation-coverage growth:**
`test_foundational_mutations.sh`

Directly verified by reading the tracked file: it currently registers **5 controls** (`fc` ,
`th` , `hg` , `km` , `df`) and **8 paired mutations** (`M1` – `M8`), for **13 total registered checks**. The file's own in-source comment states the pre-split shape had a single combined kill-tree/descendants test (one control, `M7` alone covering both functions together) — consistent with the B1/B2 fix (Section 4.3) having grown the registration from 4 controls + 7 mutations (11 total) to the current 5 controls + 8 mutations (13 total) by splitting one previously-combined check into two independently-mutation-covered ones. This is a directly observable, concrete instance of §11.4.201's control- needle discipline being applied to the project's *own* meta-test suite, not merely to product code.

. **Parallel vs. serial mutation-suite wall-clock** (directly quoted from the tracked file's own header comment, `test_foundational_mutations.sh` lines –)

“measured 2026-09-27 on the 64-CPU dev host at load average ~40: the previous serial form took **425 s** (12 full test runs) ... the parallel form measured **68 s** at load average ~38”

This is a real, ~6.25x wall-clock reduction for this one suite, achieved by running the 5 controls and 8 mutation jobs concurrently on separate `mktemp` tree copies, bounded by the same `lib/host_guard.sh` parallelism clamp the batch itself builds and tests — i.e. the Foundational batch's own tooling was used to speed up the Foundational batch's own test suite, which is a direct, self-referential proof that the tooling works as intended.

. The `descendants()` regression's own measured numbers

- Unpatched transitive closure over the `R8b-D6` fixture's poisoned input: **905 pids out of 907** input lines admitted to the visited set.
- Patched closure over the identical input: **1 pid** (only the synthetic fixture root).
- The fixed wall-clock bound the regression could exceed: **25 seconds** (250 polling iterations at a 0.1-second step).
- Measured post-fix driver death time: **~0.2 seconds** (`i=2` on the same 0.1-second-step polling loop) — i.e. the fix did not merely bring the driver back *under* its 25-second bound, it collapsed the driver's cleanup time to roughly 1% of that bound.
- **A figure this document's author could NOT independently verify:** the task brief that requested this document cited an approximate “~125–126 seconds elapsed per run before the fix, ~99–100 seconds after” for `test_run_all.sh`'s own overall wall-clock, framed as roughly a 20% reduction. No tracked artefact, agent-registry entry, or persisted request-history report available to this document's author at authoring time contains that specific pair of figures — the `qa-results/fastcycle/` evidence directory the plan designates for exactly this kind of captured timing data does not yet exist on disk (confirmed: `find` returns nothing under it). Rather than repeat an unverified number as settled fact, this document reports the two numbers it *could* verify (905 → 1 pids; 25 s cap vs. ~0.2 s actual) as the real, directly measured performance evidence for this fix, and marks the 125–126 s / 99–100 s pair `UNVERIFIED` per §11.4.6.

. Review-round effectiveness

Round	Verdict	Findings	Notable
9	NO-GO	3 BLOCKING + 1 IMPORTANT + 4 MINOR (8 total)	Found on the batch's state at that time; two of the three BLOCKING findings (B1, B2) concerned a fix that had <i>itself</i> just landed, and the third (B3) was a wholly separate pre-existing defect the review discovered independently
10	NO-GO	1 BLOCKING + 3 IMPORTANT + 5 MINOR (9 total)	Found on a smaller diff than round 9 reviewed (only the round-9 remediation, not the whole batch afresh) — every one of these 9 findings is a <i>different</i> defect class than any of round 9's 8, confirming the fixes did not merely relocate the same bug

Two independent, adversarially-reframed review rounds over the same 14-file, ~3,656-line batch found **17 distinct findings** total, with **zero overlap** in defect identity between the two rounds. This is the concrete, measured argument for why this project's §11.4.134 mandatory-iterate-to-GO discipline is load-bearing rather than a formality: a single review pass — even a rigorous one — would, on this evidence, have missed at minimum the entire round-10 finding set, including a ⁵ ⁶ BLOCKING defect (the exhausted-vs-unreadable precedence bug) capable of silently letting a 7th agent start work under a cap of 6.

. **shellcheck triage: a worked example against mass-suppression**

The pre-check pack attached to round 10's review carried **160** `shellcheck` `info/` style items across 8 files. Of these, exactly **2** were confirmed genuine findings — `SC2086` (unquoted variable expansion) at `test_check_deps.sh:32` and two instances of `SC2143` (a use- `grep -q` -instead-of-piping-to- `wc -l` style note) — with the remainder being false-positive `SC2329` (“function appears unused”) notices, a known class of false positive for functions invoked only indirectly (via `trap`, sourcing, or dynamic dispatch) that `shellcheck`'s static analysis cannot see. This is a direct, worked demonstration of this project's own doctrine against mass-suppressing linter warnings (this project's §11.4 family forbids blanket-disabling a check class rather than triaging each finding on its merits) — 158 of 160 ⁵ ⁷ items were correctly identified as noise *by individually checking each*, not by disabling the rule class that produced them.

. **Fixture-corpus scale (a proxy for how exhaustively `c-005`'s self-validation triple has been exercised in this batch alone)**

- **69** independent `triple_harness` fixture directories, each carrying a `golden-good` / `golden-bad` / `negative-control` triple — directly counted by directory listing at authoring time.
- **10** `host_guard` fixture shims (fake `ps` / `id` / `nproc` / `awk` binaries) driving resource-exhaustion scenarios that cannot safely be constructed against the real host.

. Diagrams and schemes

. Speckit- phase / dependency graph

flowchart TD

```
P1["Phase 1: Setup<br/>T001-T006"] --> P2["Phase 2: Foundational<br/>T007-T014<br/>(blocks every story phase)"]
```

```
P2 --> P3["Phase 3: US1 - Know where<br/>time goes (MVP)<br/>T015-T049"]
```

```
P3 --> RP["T049 re-planning checkpoint<br/>(recompute RC-01..RC-46 shares<br/>from the frozen baseline)"]
```

```
RP --> P4["Phase 4: US2 - Fast,<br/>lossless gates<br/>T050-T085"]
```

```
RP --> P5["Phase 5: US3 - Stop<br/>reopening<br/>T086-T103"]
```

```
RP --> P6["Phase 6: US4 - Minimal<br/>tokens per item<br/>T104-T121"]
```

```
RP --> P7["Phase 7: US5 - Reliable<br/>orchestration<br/>T122-T140"]
```

```
RP --> P8["Phase 8: US6 - Host that<br/>does not throttle<br/>T141-T153"]
```

```
P4 --> P9["Phase 9: US7 - Everything<br/>lands in constitution,<br/>clean and pushed<br/>T154-T165"]
```

```
P5 --> P9
```

```
P6 --> P9
```

```
P7 --> P9
```

```
P8 --> P9
```

```
P9 --> P10["Phase 10: US8 - Every<br/>consuming project migrated<br/>T166-T177"]
```

```
P10 --> P11["Phase 11: Polish and<br/>Verification<br/>T178-T195<br/>(every SC-001..SC-010 verdict)"]
```

```
style P2 fill:#f9d,stroke:#333
```

```
style RP fill:#ffd,stroke:#333
```

```
style P11 fill:#dfd,stroke:#333
```

. Sequence diagram: the T review cycle as it actually ran

```
sequenceDiagram
    participant C as Conductor
    participant R9 as Reviewer (round 9, Opus xhigh)
    participant FA as Fix agents (B1/B2, B3/I1, M3/M4)
    participant EX as Real execution (run_all.sh)
    participant R10 as Reviewer (round 10, Opus xhigh)
    participant FB as Fix agents (round 10, x3 parallel)

    C->>R9: Dispatch fresh, clean checkout,<br/>framed to disprove completion
    R9-->>C: NO-GO: 3 BLOCKING + 1 IMPORTANT + 4 MINOR
    C->>FA: Dispatch parallel fixes on disjoint files
    FA-->>C: B1+B2 fixed by splitting kill_tree/descendants tests
    FA-->>C: B3+I1 fixed (host_guard exhaust() + unconditional mem check)
    FA-->>C: M3+M4 fixed (C-002 exemption doc, disown, docstring)
    C->>EX: Run full suite for the first time post-split
    EX-->>C: R8b-D6 fails: driver still alive after 25s
    Note over C,EX: Real execution finds what no review round<br/>had yet caught
    C->>C: Systematic-debugging root-cause investigation<br/>(isolated scratch, zero real-tree edits)
    Note over C: descendants() de-dup bug found:<br/>905/907 pids poisoned, RED to GREEN in scratch
    C->>C: Apply the proven one-line awk fix<br/>to the real tracked file
    C->>R10: Dispatch fresh, re-framed to hunt for<br/>the same bug class elsewhere
    R10-->>C: Confirms round-9 fixes hold and the new fix holds,<br/>finds NO further instance of the bug class
    R10-->>C: NO-GO (round 10): 1 BLOCKING + 3 IMPORTANT + 5 MINOR<br/>(9 findings, none overlapping round 9's 8)
    C->>FB: Dispatch 3 parallel fixes on disjoint files
    FB-->>C: host_guard fix: CONFIRMED complete, 125/125 -> 131/131
    Note over FB,C: fc_common fix: in flight, no confirmed completion yet
    Note over FB,C: R8b-D6 kill-log + shellcheck fix: in flight, no confirmed completion yet
```

. State diagram: the descendants() regression's lifecycle

stateDiagram-v2

[*] --> Correct: Original test_foundational_mutations.sh
(no closure-membership filter,
single combined kill_tree/descendants test)

Correct --> FilterAdded: Round-9-triggered fix adds
the pid-greater-than-1 closure filter

FilterAdded --> MaskedRegression: Filter starves kill_tree's OWN
-gt-1 guard of pid 1
(M7 mutation now uncaught - a bluff gate)

MaskedRegression --> DiscoveredByReview: Round 9 review measures it directly:
M7-weakened mutation still passes rc=0

DiscoveredByReview --> SplitTests: Fix: split into two independent tests,
each stubbing out the OTHER function

SplitTests --> RealExecutionDiscovery: First real full-suite run after the split
surfaces R8b-D6: driver still alive after 25s

RealExecutionDiscovery --> RootCaused: Systematic-debugging investigation:
de-dup bug in descendants() awk closure,
905/907 pids poisoned (measured, isolated scratch)

RootCaused --> Fixed: One-line awk guard applied to the
real tracked file (905 -> 1 pids, RED to GREEN)

Fixed --> IndependentlyReVerified: Round-10 review deliberately re-hunts
for the same bug class: finds none

IndependentlyReVerified --> [*]: Confirmed stable (round 10 GO on this specific defect,
batch overall still NO-GO on 9 unrelated findings)

. File-scope partitioning for the round-parallel fan-out

```

flowchart LR
    subgraph Conductor
        R["Round-10 review verdict:<br/>1 BLOCKING + 3 IMPORTANT + 5 MINOR"]
    end
    R --> A["Fix agent A<br/>Scope: lib/host_guard.sh<br/>+ tests/test_host_guard_red.sh<br/>Findings: BLOCKING-1, IMPORTANT-2,<br/>+IMPORTANT-3, MINOR-6"]
    R --> B["Fix agent B<br/>Scope: lib/fc_common.py<br/>+ tests/test_fc_common_red.sh<br/>Findings: MINOR-7, MINOR-8"]
    R --> D["Fix agent C<br/>Scope: tests/test_run_all.sh<br/>+ shellcheck items across<br/>4 other files<br/>Findings: IMPORTANT-4, MINOR-9"]
    A -.no shared files.- B
    B -.no shared files.- D
    A -.no shared files.- D
    A --> CA["Conductor independently<br/>re-verifies: 131/131,<br/>+CONFIRMED complete"]
    B --> CB["Status at authoring time:<br/>in flight, unconfirmed"]
    D --> CD["Status at authoring time:<br/>in flight, unconfirmed"]

```

. Constitution anchors applied

Every anchor below is quoted or precisely summarised from `Constitution.md` directly^{11 4 6} (line numbers cited are as found in this document's own real-time read of the file at authoring time), not paraphrased from memory.

§ . . – No-guessing mandate (`Constitution.md:632`)

Forbids likely / probably / maybe / might / possibly / seems / appears to and synonyms when describing causes; either prove a cause with captured forensic evidence and state it as fact, or explicitly mark `UNCONFIRMED:` / `UNKNOWN:` / `PENDING_FORENSICS:` with a tracked follow-up. **Applied throughout this cycle:** the `descendants()` investigation states its host-load conclusion as measured fact

with the measurement cited, not as a guess; this document itself marks the one figure (Section 5.4) it could not verify as UNVERIFIED rather than repeating it as settled.

§ . . – Deterministic Consistency Mandate (Constitution.md:4319)

“Not a single feature, System or application flow ... MUST NOT partially work, or sometimes work and sometimes not ... results MUST BE consistent and successful without exception.” Operationalised as N-iteration deterministic outcome (a test PASSing once is not sufficient — the same test must PASS identically across repeated runs against the same input). **Applied directly:** C-003 's -- determinism-check (run twice, assert identical body_hash) is exactly this anchor turned into a mechanical contract every new tool in the batch must satisfy, and the descendants() fix's own verification protocol explicitly ran 3 clean repeats before being trusted.

§ . . / § . . / § . . (C) – Parallel subagent fan-out on disjoint file scopes

§11.4.58 (parallel-development methodology, Constitution.md:5331) establishes the Parallel Work Unit pipeline and its 4-layer lock hierarchy so disjoint-scope work runs fully concurrently; §11.4.94 (Constitution.md:7892) makes zero-idle, priority-first, parallel-by-default the standing operating mode; §11.4.230(C) (Constitution.md:10969) makes subagent fan-out the standing default for any genuinely parallelizable effort, bounded by host-safety ceilings and single-resource-owner discipline. **Applied directly:** every [SUBAGENT] / [P] -marked task in tasks.md , and concretely the 3-way round-10 fix fan-out (Section 4.4 / Figure 6.4), each agent scoped to disjoint files with no contention, exactly matching this triad's structure.

§ . . / § . . (F) – Systematic-debugging + machine-written RED→GREEN verdicts

§11.4.102 (Constitution.md:8175) mandates automatic activation of `superpowers:systematic-debugging` on any spotted issue, with the Iron Law “NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST.” §11.4.115(F) (extension within Constitution.md:8678) requires RED and GREEN verdicts to be machine-written, never prose, carrying the item id, guard identity, polarity, exit code, the target’s artifact fingerprint read from the target at run time, iteration count, and evidence whose class matches the defect’s layer. **Applied directly:** the `descendants()` investigation is Section 3.1’s worked example of the Iron Law in full; every [TDD] task’s RED evidence is stored under `qa-results/fastcycle/<story>/red/` with the candidate fingerprint, per `tasks.md`’s own Test Discipline section.

§ . . – Code-review iterate-until-GO + rock-solid-evidence mandate (Constitution.md:9102)

A review returning ANY finding — BLOCKING, IMPORTANT, or MINOR — must be re-run after remediation, and re-run again after that, until it returns a clean GO with ZERO new findings and ZERO warnings; a single pass that “addressed the findings” is explicitly insufficient, because the fixes themselves can introduce new findings. **Applied directly:** this is precisely what happened across rounds 9 and 10 (Section 5.5) — round 10 found 9 entirely new findings on a *smaller* diff than round 9 reviewed, which is the anchor’s own predicted failure mode (“a fix-A-creates-B failure mode”) observed happening in real time on this batch.

§ . . / ATM– D – Crashed/ async-agent registry hygiene (Constitution.md:9407)

Every asynchronously launched agent must be tracked through its full lifecycle so a crash never loses or silently mis-records its work; a crash ≠ done. **Applied directly, and its own known limitation honestly surfaced by this project’s own tooling:** the agent registry (`docs/requests/agent_registry.jsonl`) currently has a documented defect, internally tracked as ATM-858 D1 , where an asynchronously

dispatched agent's `complete` event fires at *launch* time rather than at real completion — every dispatch entry cited in Section 4.4 explicitly annotates “`'complete'` above fired at `async launch`, ATM-858 D1 ” and relies on a *separate*, later, genuinely-real completion event (or the honest absence of one, as for the two still-in-flight round-10 fix agents) to establish ground truth. This document's own Section 4.4 table was built by applying exactly that distinction.

§ 4.4.1 — Exact-reproduction-sequence mandate

When a working reproduction exists, the investigation must use its EXACT sequence — a deviating repro proves nothing. **Applied directly:** the `descendants()` investigation deliberately reproduced the R8b-D6 fixture's *exact* fake-`ps` shim sequence (not an approximation of it) before drawing any conclusion, and explicitly declined to re-trigger a third real host-wide signal storm purely for redundant confirmation once the mechanism was already pinned by two side-effect-free measurements plus one live signal-capture — an application of the anchor's spirit (use the real sequence; do not manufacture unnecessary risk once sufficient rigor is already achieved).

§ 4.4.2 — Every guard/gate MUST assert the REAL condition

A false-positive refusal is a FAIL-bluff exactly as a false-negative pass is a PASS-bluff; every guard ships golden-TRUE + golden-FALSE-with-carrier fixtures.

Applied directly: C-005's self-validation triple (golden-good/golden-bad/negative-control) *is* this anchor's mandated fixture shape, applied to every new tool in the batch; the round-9 B1 finding (“mutation NOT caught — a bluff gate”) is a direct instance of this anchor's core failure mode being caught and fixed within this very cycle.

§ . . / § . . —
Opus- xhigh -pinned review/merge-conflict
substrate (Constitution.md:10574 /
Constitution.md:10606)

Every mandatory independent code-review, and every main → feature merge-conflict resolution, must run on the Opus model at xhigh effort, ALWAYS — no Fable, no lower tier, no fallback model; genuine unavailability BLOCKS the review rather than substituting a lower tier. **Applied directly:** both T014 review rounds in this cycle were dispatched explicitly labelled “Opus^{11 4 263} xhigh^{11 4 273}”, per the agent-registry entries this document read directly (Section 4).

§ . . / § . . — Process-
group signal-safety + control-needle-proven
measurements (Constitution.md:11518 /
Constitution.md:11808)

§11.4.263 forbids ever signalling a pgid / pid ≤ 1 and requires validating a value as a real integer > 1 before every killpg / kill(-pid, sig) call. §11.4.273 requires any census/measurement whose result informs a decision to be control-needed — a positive control proving the instrument can see, and a negative control proving it does not over-match — before its result is trusted. **Applied directly, and this is the exact fault class this whole story is about:** the descendants() bug was a §11.4.263-class hazard (unrelated, unrelated real processes were signalled with kill -s TERM / KILL because a non-integer/≤1 guard existed at signal time but not at the earlier closure- membership stage); the fix restates the same pid > 1 guard at the point of closure construction. And the mutation-testing methodology in test_foundational_mutations.sh itself explicitly implements §11.4.273’s “a sed that matched nothing cannot read as caught” control (mut_job() ’s cmp -s check before trusting any mutation’s caught/uncaught verdict) — the project applying its own anti-bluff-measurement doctrine to its own measurement tooling.

. Lessons learned and recommendations

. The generic shell lesson: a helper function can survive `exec` differently than the caller expects

The `descendants()` regression is not, in its most general form, a bug about process closures specifically — it is an instance of a broader, genuinely useful lesson for anyone authoring bash test fixtures that use `exec`: functions and environment variables `export -f`’d before an `exec` **do** survive it (because `exec` replaces the process image but the shell’s own exported state carries forward), but a plain, non-exported function or variable defined only in the calling shell **does not** — because `exec` replaces that shell’s entire process image, and nothing of the old, un-exported state exists afterward for the new image to inherit. A test fixture that defines a helper function, then does `exec bash "$1"` as its last line, will silently lose that helper for every subsequent invocation inside the new process image unless the helper was `export -f`’d first — and the failure mode is not a crash or an obvious error; it is “the helper is not found, the function call that depends on it silently no-ops or falls through to a different code path, and the symptom that eventually surfaces (if any) can be arbitrarily far removed from the actual cause in both code location and wall- clock time.” This is exactly the shape of pitfall this project’s own §11.4.201(12) “shell-instrument footgun checklist” (a standing reference this cycle’s own governance already maintains) exists to catalogue — and this cycle’s own concrete instance (the `descendants()` / `kill()` -wrapper/ `drvwrap.sh` interaction) is a strong candidate to be added to that catalogue as a fresh, cycle-native worked example, generalising past the specific process-closure framing to the underlying “exported vs. non-exported state across `exec`” mechanism.

. Real execution finds what review, by itself, cannot

The single most consequential fact this cycle demonstrates is that **the `descendants()` de-duplication defect was invisible to two rounds of rigorous, adversarially-framed static/logical review** — round 9’s own review, which *introduced* the fix that (unintentionally) exposed the underlying defect, still did

not itself catch the defect; it took an actual, real execution of the full test suite, on real process-table data, to surface it as a concrete “still alive after 25 seconds” failure. This is the empirical argument for why this project’s constitution insists on *both* mandatory independent review *and* real, runtime, on-target execution as co-equal, non-substitutable layers (§11.4.108’s four-layer fix-verification discipline, cited by name in `spec.md`’s own Workable-Item Linkage section) — a code review, however rigorous, reasons about what the code *should* do; only real execution proves what it *actually* does under real conditions.

. **Adversarial reframing between review rounds is load-bearing, not ceremonial**

Section 5.5’s finding that round 10 discovered 9 entirely new, non-overlapping findings on a *smaller* diff than round 9 reviewed is a direct, measured argument against treating “the previous round’s findings are fixed” as equivalent to “the batch is done.” Each review round in this cycle was deliberately re-framed — round 10 was explicitly briefed to hunt for a specific bug class the previous round’s fix had exposed — rather than merely re-checking the same finding list. This project’s recommendation, borne out by this cycle’s own evidence, is that every re-review dispatch should continue to carry an explicit “and specifically hunt for X” framing derived from what the *previous* round’s fixes changed, not merely a generic “please re-check” instruction.

. **Honest partial-completion reporting is itself a discipline worth preserving**

At the moment of authoring, one of three round-10 fixes is confirmed complete and two are not. This document’s own Section 4.4 table, and its Section 5.1 table, were both constructed to reflect that asymmetry precisely — never rounding “2 of 3 still in flight” up to “the batch is basically done.” This is recommended as a standing practice for any status report produced mid-cycle in this project: report the confirmed state, name the unconfirmed state explicitly, and never let a document’s tone imply more certainty than its underlying evidence supports.

. Verification methodology and honest gaps of this document itself

This document was produced by directly reading the following live artefacts at authoring time, rather than from memory or from the task brief that requested it (which is itself explicitly framed, and treated here, as a starting dossier requiring independent confirmation, not a source of truth):

- `specs/004-fast-dev-cycles/spec.md` (full text, 269 lines), `plan.md` (task-id and phase-boundary extraction via targeted `grep`, given its 1,850-line size), `research.md` and `tasks.md` (full text plus targeted structural greps for phase/task-id boundaries).
- Direct `wc -l` and directory-listing counts of every file named in Section 4.1 and Section 5.7, run against the real, current tracked tree.
- `docs/requests/agent_registry.jsonl` (2,921 lines at the time of reading), filtered for every entry mentioning this cycle's work, cross-checked against its own documented `ATM-858 D1 async-completion-timing` defect so that a launch-time `complete` event was never mistaken for a real one.
- `docs/requests/history.md` (the operator request-history ledger), whose persisted verbatim subagent hand-back reports were the primary source for Section 3.1's and Section 4's most detailed factual claims, quoted directly rather than paraphrased wherever the exact wording carries load-bearing detail.
- `Constitution.md` (11,900 lines), read at the exact line numbers cited in Section 7 for each anchor, rather than reproduced from a compact summary.

One deliberate methodological choice: this document's author did **not** independently re-run the fastcycle test suite (`run_all.sh` or the individual `test_*.sh` files) against the live tree, even though doing so is read-only and would not modify any source file. This is because, at authoring time, three fix agents were actively working inside `constitution/scripts/fastcycle/` (the exact directory this dispatch was told to treat as fully disjoint from this task's own scope), and `lib/host_guard.sh` 's own live resource probes (the very thing under test) read the *shared* host's current process count, thread headroom, and memory availability — launching an additional heavy test run concurrently would have measurably perturbed the host state those agents' own verification runs depend on, in a project whose own §12.6/§12.8/§12.12 host-safety discipline exists specifically to prevent exactly that class of self-inflicted interference. The pass/fail figures

reported in Section 5 are therefore cited from the fix and review agents' own captured, registry- and ledger-recorded runtime evidence — which is itself the project's own mandated RED → GREEN and `run_all.sh /mutation-suite` invocation discipline — rather than from a fresh run this document's author performed itself.

Known gap, honestly recorded rather than silently worked around: per this project's own §11.4.212 anchor ("main README is the canonical starting-point ... no doc may be an orphan unreachable from README"), this document should ultimately be linked, directly or transitively, from `constitution/README.md`. This dispatch's own scope was explicitly restricted to new files under `docs/optimization/` only, and `README.md` was, at the moment of writing, a file this dispatch was told other in-flight work might also be touching (it appeared among the pre-existing uncommitted changes present in the working tree before this dispatch began) — so no edit to `README.md` was made here. Wiring this document into the README's own doc-link section is recorded here as an explicitly owed, tracked follow-up (per this project's own §11.4.197 research/kicked-off-work completion mandate — a started requirement is never silently left un-wired), to be done as its own small, disjoint, reviewed change once the README's current uncommitted state is known to be settled.

1 0 2 0 2 6 0 9 2 8

2

. Update — — (Revision)

This section was requested directly by the operator via the project's own §11.4.140 action-prefix mandate — the originating request, quoted verbatim from the operator request-history ledger (`docs/requests/history.md` , entry `R-2026-09-28-112940` , 2026-09-28 11:29:08 MSK), reads: *"REMINDER —> Has this been done, committed and pushed? —> BACKGROUND :: IMPORTANT: Do exhaustive analisys of all work we are doing now via SpecKit and Superpowers with all details, proper summary, fully structured content divided into sections, additional explanations, measurements, analisys of performance and metrics changes, graphs, diagarams and schemes and put it all into constitution Submodule under docs/optimization/ dir."* This section is that update, produced by directly reading the live git history of both the parent repository and the constitution submodule, the live

specs/004-fast-dev-cycles/tasks.md task list, the agent registry, the operator request-history ledger, and the real on-disk evidence artefacts every implementation task below cites — exactly the same verification discipline Sections 1–9 above already establish for Revision 1.

. Executive summary of the delta

Since Revision 1 was committed (constitution c380a44 , 2026-09-27T22:39Z), this cycle has landed **31 real commits** across the two repositories this feature spans — **12** in the constitution submodule, **19** in the parent ATMOSphere repository — none of them reverted, none of them force-pushed (§11.4.113 forbids that absolutely), all independently re-verified before landing per this project's own §11.4.240 producer≠verifier discipline. The headline changes:

1. **T014 (the Foundational batch's mandatory review gate) finally reached a clean GO.** Revision 1 left this genuinely open — two of three round-10 fixes were unconfirmed at that document's authoring time. This revision can report, with full evidence, that the batch went through **two further review rounds** (11 and 12) neither of which Revision 1 could have known about, the second of which found a genuinely new class of defect (a silent-fabricated-test-verdict hazard under an inherited SIGCHLD=SIG_IGN signal disposition — a defect class distinct from anything Revision 1 described), and the batch was committed to the real tracked tree (b72361b) with a final state of run_all: 9 run, 0 failed , **15/15** mutations caught, test_fc_common_red.sh **233/233**, test_host_guard_red.sh **131/131**. See [Section 10.3](#).
2. **Phase 3 (US1, the measurement MVP) has substantially progressed.** All 12 of US1's RED tests (T015 – T026) are landed and independently verified. Seven implementation tasks have landed with real, machine- written GO verdicts: T028 (fc_timer.sh), T034 (review_record.py), T036 (dispatch_stamp.sh), T038 (transcript_ingest.py), T039 (anchor_citations.py), T041 (cycle_report.py), T046 (plan_struct_check.py causes) — several of these required two independent review rounds with real, fixed findings before reaching GO, continuing exactly the pattern Section 5.5 documented for T014. T043 (select_sample.py + baseline_replay.sh) is implemented and survived a genuine concurrency-race incident (Section 10.4(a)); its own independent review is **in flight** as of this writing.

3. **A genuine concurrency-race incident happened, was root-caused, and was forward-fixed** — running `scripts/commit_all.sh` (a broad working-tree stage) while a subagent was still actively writing a file swept in that file's mid-write state into an unrelated commit. This is reported in full, including the fix and the operational lesson recorded from it, in [Section 10.4\(a\)](#) — it is **not** glossed over, because this project's own §11.4.6 mandate forbids treating a genuine incident as anything other than what it was.
4. **A substantial, unrelated body of already-complete work (ATM-1014, 42 real defect fixes) was discovered sitting fully staged but never committed**, and was recovered rather than left idle — real mtime evidence puts it staged for **at least ~42 hours** before this session found and committed it. See [Section 10.4\(c\)](#).
5. **Four [SERIAL] wiring tasks (T029 – T032) are mid-flight as parallel draft subagents at the moment of this writing**. Per each task's own [SERIAL] marker (Section 2.3 above), only the conductor may apply, test, and commit these — the four subagents were dispatched to produce *proposed patches only*, written to a scratch evidence directory, never touching the real contention-path files. All four have produced substantial draft output (patches + notes, real file sizes documented in Section 10.5), but **none has yet been reviewed, applied, tested, or committed by the conductor** — this document reports that state exactly as it is, not as “done.”

Honest overall status, stated plainly: SpecKit-004 is **not complete**. Phase 3 (US1) alone still has T029 – T033 , T035 , T037 , T040 , T042 , and T044 – T049 remaining (draft-only or not-yet-started), and Phases 4 through 11 — US2 through US8 plus the final Polish/Verification phase, the overwhelming majority of the 195-task plan — have not been started at all. See [Section 10.7](#) for the precise breakdown.

. Timeline of everything landed since Revision

Every commit below is cited by its real, independently-verified hash. Times are as recorded by `git show -s --format=%cI` at authoring time (the committer's local clock, +05:00). “Repo” distinguishes the constitution submodule from the parent

ATMOSphere repository, since this feature's implementation deliberately spans both (the tool source lives in the submodule per this project's own §11.4.28 decoupling discipline; the plan, tasks, and device-tree wiring live in the parent).

#	Time	Repo	Commit	What landed
1	01:26	constitution	b72361b	Foundational batch (T007 – T014) landed to the real tracked tree — T014 GO after rounds 9–12 (see 10.3)
2	—	parent	ada22d27782	Bump submodule pointer for the Foundational batch
3	—	parent	4f8c0fddee6	Phase 3 US1 RED tests T016 – T022 , T024 , T026 (T015 / T023 / T025 still in flight at that point)
4	—	constitution	7ce30dc	Bump submodule pointer, mirroring #3
5	—	parent	a2d9f7106e1	Phase 3 US1 — all 12 RED tests (T015 – T026) confirmed DONE and independently verified
6	—	parent	f2a290af13e	Correction: a prior claim that constitution 7871f2c was pushed was found false — a real commit_all.sh cascade-push gap, documented rather than silently re-asserted
7	—	constitution	005c8a1	Phase 3 US1 RED tests T016 – T022 , T024 , T026 (constitution-side landing matching #3)

#	Time	Repo	Commit	What landed
8	—	constitution	7871f2c	T023 RED test (cycle_report.py); missing_record / empty_window fixtures for cycle_report landed incidentally in the prior commit due to concurrent subagent writes — noted honestly as harmless orphans until this commit referenced them
9	—	parent	bc2339a8895	T027 DONE — all 10 US1 mutation-observations complete (7 verified-flip + 3 honest-gap)
10	—	parent	19bc3fe657e	SpecKit-004 Implementation subsection begun — T028 / T034 / T039 / T041 dispatched
11	—	parent	c5c4825c1c3	T028 (fc_timer.sh) + T041 (cycle_report.py) — both independently re-verified GO
12	—	constitution	7c2e1d5	Constitution-side landing matching #11
13	—	parent	d62abd751a1	T039 / T034 fix passes + T046 dispatched; T036 's missing-RED- test gap documented
14	—	parent	3ec1f6724d5	T046 landed + independently re- verified, its review dispatched; T039 / T034 fix passes still in flight
15	—	parent	17c4e44d634	T039 / T034 fix passes independently re-verified by the conductor, a fresh batched review dispatched

#	Time	Repo	Commit	What landed
16	—	parent	5686e31a15b	T046 review returns NO-GO (1 BLOCKING — a malformed-row silent drop), reproduced + a fix pass dispatched
17	—	parent	502f477e9c4	T039 (anchor_citations.py) + T034 (review_record.py) — both GO after a second Opus- xhigh review round
18	—	constitution	0c4fbb7	Constitution-side landing matching #17
19	—	parent	03add525954	T046 GO on its round-2 review — T039 + T034 + T046 all fully accepted; T036 implementation in flight
20	—	parent	23c67ef724c	T036 's genuinely-missing RED test authored (a real spec gap, closed) + T036 's implementation dispatched; T046 's fix pass independently re-verified, its round-2 review dispatched
21	—	constitution	50f9ba8	Constitution-side landing matching #20
22	—	parent	0913903567d	T038 (credential-sensitive) implementation dispatched with a thorough baseline investigation; T036 's review still in flight
23	—	parent	847f8691428	T036 landed + independently re-verified, its review dispatched

#	Time	Repo	Commit	What landed
24	13:02	constitution	ba51e21	T036 (dispatch_stamp.sh) landed and reached GO after 2 review rounds (3 BLOCKING findings in round 1, all fixed + independently re-verified)
25	13:03	constitution	2ed3489	T038 (transcript_ingest.py , 552+ lines) + T020 's RED test — GO after round 1 (credential-safety independently re-verified by conductor AND reviewer, both from scratch)
26	—	constitution	8a42520	§11.4.65/§11.4.73 four-format doc-export sync (stale companion docs regenerated project-wide)
27	—	constitution	ae9184a	Bump submodule pointer for #24–#26
28	13:18	constitution	c33c1d4	T043 's baseline_replay.sh race-recovery fix — see Section 10.4(a)
29	13:27	parent	3e1ea0fa484	ATM-1014 recovery — an unrelated, already-complete 42-defect-fix body of work found staged-uncommitted, recovered — see Section 10.4(c)
30	13:29	parent	72d918172e0	T003 (config/fastcycle/ Setup files) + a real §11.4.65 INCLUDED-scope gap fix in sync_all_markdown_exports.sh — see Section 10.4(b)

#	Time	Repo	Commit	What landed
31	—	parent	2ea6897a147	Bump submodule pointer to c33c1d4 (a bump this revision's author independently confirmed had been missed by the prior commit)

(Timestamps are shown only where independently re-confirmed against `git show -s --format=%cI` at authoring time; several early-session commits landed in rapid succession without this document's author re-querying each one's exact clock time individually — the **ordering** above is the real, verified `git log` ordering in each repository, which is the load-bearing fact for the narrative.)

As of this writing, four further items are dispatched and in flight (none show a real `complete` event in the agent registry as of authoring time — see [Section 10.7](#) for why a registry `complete` event is not itself trusted at face value, per this project's own documented `ATM-858 D1` defect):

- T029 draft — wiring `fc_timer.sh` into `pre_build_verification.sh`
- T030 draft — wiring stage timers into `commit_all.sh` + per-document rows into `sync_all_markdown_exports.sh`
- T031 draft — adding `requested_at` / `applied_at` timing to the tracker writer's `db.go`
- T032 draft — per-mutation start/end/verdict rows in `meta_test_false_positive_proof.sh`
- T043¹² 's own independent review¹ (Opus `xhigh`)

. T reaches G0: rounds and
, closing Revision 's open gap

Revision 1 (Section 4.4) reported T014 at round 10's NO-GO, with two of three round-10 fix agents still unconfirmed. This revision reports what actually happened next, read directly from the landing commit's own message (`b72361b` , quoted and summarised here, not paraphrased from memory):

Round 11 (dispatched after the two remaining round-10 fixes were confirmed complete) returned **1 BLOCKING + 2 IMPORTANT + 3 MINOR** — a genuinely new set of findings, in a defect class Revision 1's own narrative never touched:

- **BLOCKING-1:** `fc_common.py`'s `_reap_group / _leader_exited_unreaped` logic **silently fabricated a fake** `returncode=0` **verdict for a genuinely FAILING command**, whenever the calling process had inherited `SIGCHLD=SIG_IGN` from its own parent — under that disposition, the kernel auto-reaps child processes, so `Popen.wait() / .poll()` raise `ECHILD`, and the pre-fix code's exception handling silently returned `0` instead of surfacing the real, unknown exit state. This is precisely the class of silent-false-PASS defect this project's entire §11.4 anti-bluff covenant exists to make structurally impossible — found and fixed on the batch's *own* test-verdict-reporting machinery, not on product code, which is arguably the highest-leverage place such a defect could hide. Fixed by resetting `SIGCHLD` to `SIG_DFL` before every spawn in `_run_bounded`, plus a belt-and-suspenders `_leader_reap_honest()` that fails closed (a `NO_HONEST_VERDICT` sentinel, never a fabricated `0`) if the reset is somehow bypassed.
- **IMPORTANT-2:** round 10's own correct `exhaust() / unreadable()` fix (Revision 1, Section 4.4's B3/I1 fixes) had, as a side effect, left ~19 `test_host_guard_red.sh` checks silently dependent on the real host's live memory/thread state — exactly the class of hazard Revision 1's own round-10 entry flagged as a concern (Section 4.4's own IMPORTANT ×3 bullet). Confirmed and fixed by pinning `$MEMOK / $AMPLE` on every affected line.
- **IMPORTANT-3:** `kill_tree()`'s root guard `( [ "$r" -gt 1 ] || return 0` — the *only* thing preventing a bad root argument from `TERM / KILL`-ing the entire host process tree, directly the §11.4.263 hazard this whole story is about) had **no direct test of its own** — a mutation of it survived every existing check. Fixed with a new `test_foundational_mutations_killtree_root_guard.sh` + a newly-registered mutation, `M12`.
- **MINOR ×3:** a `shellcheck` -count-accuracy nit, a needle-coverage gap in `test_fc_common_red.sh`'s leak-check (which, while being fixed, itself surfaced a **real bug in the needle**: bash's own single-statement `-c` exec-optimization silently discarded the needle's marker from the running process's `argv`), and round-10's own changes plus the BLOCKING-1 fix having landed

with no registered paired mutations of their own (fixed: M13 / M14 / M15 + a tightened marker-based “caught” verification, backfilled onto all 15 registered mutations).

Round 12 (GO, but not zero-finding — one IMPORTANT, two MINOR, fixed before commit):

- **IMPORTANT:** the new `killtree_root_guard.sh` test round 11 had just added carried **no positive control** — an entirely broken `kill_tree()` extraction (verified by deliberately injecting a syntax error) still reported “ok” on all three pathological-root checks and exited `0`, a complete pass-bluff on the *newest* test in the suite. Fixed by adding a genuinely-valid-root positive-control case, run *before* the absence checks, proving the extraction+shim+invocation pipeline itself actually works.
- **MINOR** (fixed): `test_fc_common_red.sh` leaked two temp directories per invocation via a bare `tempfile.mkdtemp()` with no cleanup — fixed with `shutil.rmtree(..., ignore_errors=True)` at both call sites.
- **MINOR** (tracked, deliberately not fixed this round): `fc_common.py` registers no `SIGTERM` / `SIGINT` handler for standalone invocation, so an external signal delivered to a standalone run could skip its `finally` cleanup block. The commit message records this as “largely mitigated in this codebase’s actual usage” (the outer `kill_tree` signals the whole process-group snapshot independent of `fc_common`’s own cleanup) but “a real gap for standalone use,” explicitly flagged for a dedicated follow-up rather than silently dropped, per this project’s own §11.4.197 mandate.

Final landed state (from the commit’s own summary line, independently consistent with the per-file counts this document’s Revision 1 already established): `run_all:` 9 run, 0 failed . **15/15** registered mutations caught. `test_fc_common_red.sh` **233/233**. `test_host_guard_red.sh` **131/131**.

Two things about this sequence are worth stating explicitly, because they are the concrete, measured continuation of Revision 1’s own Section 5.5 argument (“adversarial reframing between review rounds is load-bearing, not ceremonial”): first, **round 11’s own BLOCKING finding is a defect class none of rounds 9 or 10 found**, on a file (`fc_common.py`) that both of those earlier rounds had already reviewed — confirming, on a *third* independent data point beyond the two Revision 1 already reported, that each fresh review round genuinely finds defects the prior

rounds structurally could not. Second, round 12's own IMPORTANT finding was a **missing positive control on round 11's own brand-new test** — meaning the review discipline here is reviewing not only the product code but its *own* prior remediation work with the same rigor, exactly the “the fixes themselves can^{10 4} introduce new findings” failure mode Revision 1's Section 7 (§11.4.134) already predicted, now observed a second time on a different file.

. Process lessons learned today

(a) `commit_all.sh` and a live subagent write are a genuine race — and it happened. While this session's own parent-pointer bump ran `scripts/commit_all.sh` detached (per §11.4.88's own default-async discipline), T043's implementer subagent was **still actively writing** `scripts/fastcycle/cycle/baseline_replay.sh` to disk. `commit_all.sh` 's own normal `git add -A` -style staging — which is *by design*, not a defect in the tool, since that is how it is meant to capture a full working tree — swept in that file's **current, mid-write state** at that exact moment and committed it (`ae9184a`). `select_sample.py` , already in its final state at the same moment, was unaffected (independently confirmed byte-identical before/after). The swept-in draft delegated `replay --determinism-check` to `fc_common.py` 's generic byte-identical- `body_hash` mechanism — but `replay` 's own report embeds real wall-clock timing fields (`start_ns` / `end_ns` / `duration_ms` / `median_ms`), which are inherently non-deterministic between two runs, so the generic mechanism reported “nondeterministic” on a perfectly healthy commit's very first real test (caught by the implementer's own testing, not assumed correct). The commit that lands the actual fix (`c33c1d4`) makes `--determinism-check` compare **only** the two runs' `verdict_set` fields as sets, never the timing fields — matching T024 's own RED-test contract wording exactly. **The operational lesson, recorded verbatim in that commit's own message and repeated here because it is genuinely worth carrying forward:** `scripts/commit_all.sh` (or any command doing broad working-tree staging) must **never** run while a subagent is still actively writing files anywhere in the same tree — wait for the subagent's own completion report first, or a race like this one recurs. This is not a defect in `commit_all.sh` ; it is a genuine gap in this session's own operating discipline, forward-fixed per §11.4.113 (force-

push is absolutely forbidden — the only correct response to a bad commit already landed is a forward fix, never a history rewrite), and recorded honestly here rather than minimised.

(b) A pre-existing anti-bluff mechanism caught a real, previously-invisible scope gap — exactly as designed.

`config/` was never added to `sync_all_markdown_exports.sh` 's `build_candidate_list()` directory walk, so `config/fastcycle/REGENERATION.md` (T004 's §11.4.77 regeneration declaration — a real Markdown document meant to be read, unlike its `.tsv` / `.yaml` siblings the exporter never globs) had no HTML/PDF/DOCX siblings. This project's own §11.4.74 Layer 2 sibling-parity check **correctly refused** a `commit_all.sh` commit over the file for exactly that reason — while `--paths` independently refused to render it too (“outside the §11.4.65 INCLUDED scope”), the same two-gate deadlock the exporter's own scope-addition comment already documents and resolves the same way (§11.4.120): widen the scope. The fix (landed in 72d918172e0) adds `config/**/*.md` to both the documented INCLUDED-scope comment and the real directory walk, matching the pre-existing `specs/**/*.md` precedent exactly. Verified before commit: a targeted `--paths` run against the specific file succeeded (real, non-degenerate siblings created); a full `--check-only` sweep confirmed **all 2,723 tracked Markdown files project-wide** were then in sync (0 out of sync); a second, pre-existing file newly brought into the widened scope (`config/multitrack/ROADMAP_1.2.0-dev.md`) was already in sync independently — no regression, no duplicate work. This is a direct, concrete instance of this project's own anti-bluff tooling doing exactly what it exists to do: refuse to silently proceed past a real, previously- unnoticed structural gap.

(c) A substantial body of already-complete, verified work sat staged and uncommitted for a measured, real duration.

While investigating an unrelated blocker, this session found `check_terminal_skip_exit0.sh` 's enumerator had a real defect — its detector enumerated wired tests via a frozen two-runner literal, so a third real runner shape (`run_device_test_nohup`) matched nothing, meaning 5 wired scripts were never scanned at all while the tool printed “clean (344 wired scripts scanned).” Once fixed, the widened enumeration found **42 real terminal SKIP-then-exit-0 defects** across `device/rockchip/rk3588/tests/*.sh` (a terminal `SKIP` that exits `0` is tallied `PASS` at the release gate — exactly the class of PASS-bluff this project's anti-bluff discipline exists to prevent). This entire body of work — the detector fix, a new standing black-box guard (`test_terminal_skip_exit0_enumeration_red.sh` , 103/103 legs independently

re-run before commit), a real end-to-end paired-mutation proof in `meta_test_false_positive_proof.sh` , and all 42 individual defect fixes — was **already fully written and complete**, but had never been committed. **Directly measured, real evidence of the gap**: the new standing guard's on-disk mtime is `2026-09-26 19:24:18 +0500` ; it was committed (`3e1ea0fa484`) at `2026-09-28T13:27:22+05:00` — a real, verified gap of **1 day 18 hours 3 minutes (≈42.05 hours)**. A companion artefact directory found alongside it (`qa-results/agent_custody_20260926/partial_after_quota/`) is consistent with this work having been produced during a separate 2026-09-26 agent-custody investigation (documented in this project's own memory index as the “respawn-all-dead-agents audit”) and left staged when that investigation's own agent was interrupted — this attribution is recorded here as a **reasonable, evidence-consistent inference**, not a confirmed fact, per this document's own §11.4.6 discipline. Once found, the work was independently re-verified (a sample fix, `test_idle_dim_follows_settings.sh` , spot-checked to confirm its diff was *exactly* the documented exit-0-to-exit-2 conversion and nothing else) and committed intact rather than left idle. This is directly relevant to SpecKit-004's own subject matter — a real, measured instance of the “reopen rate is high, iterations are slow” pattern the whole feature exists to fix, here manifesting as complete work sitting invisible and undelivered for nearly two days, discovered only by chance during unrelated investigation rather than by any systematic sweep.

(d) Two minor stray artefacts were found and cleaned up during investigation — a file literally named `$FC_OUT` at the repository root (a shell-quoting mistake from earlier `plan_struct_check.py` testing) and a 0-byte `constitution/docs/workable_items.db` (a `cwd` /path-resolution mistake). Both were deleted; neither had ever been committed. Minor in isolation, but recorded here as exactly the kind of drift careful, evidence-based investigation catches² that a less careful pass would have missed or, worse, accidentally committed.

. Measured figures (Revision)

All figures below were read directly from a real, current on-disk artefact at authoring time — never estimated, per this document's own Section 9 discipline continued into this revision.

T032's own real measurement of `meta_test_false_positive_proof.sh` 's structure (from the T032 draft agent's own `NOTES.md` , itself produced by direct reading and `grep` , not assumption — a genuinely more precise figure than Revision 1's own approximate citation elsewhere in this document):

Metric	Value	How measured
Total file length	28,813 lines	Direct line count, read at authoring time
Distinct mutation-testing idioms in the file	3	Structural read: 4 shared driver functions, one fully-inlined bespoke banner-delimited pattern (the dominant one), and one case explicitly noted dead-code-with-zero-call-sites
Call sites of <code>mutate_gate_selftest_or_red()</code>	65	<code>grep -c '^mutate_gate_selftest_or_red\\\$'</code>
Call sites of <code>mutate_anchor_gate()</code>	5	Same method
Call sites of <code>mutate_gate_via_fixture()</code>	0 (confirmed dead code, not something this investigation invented)	Same method, cross-checked against a broader unanchored grep to rule out a too-strict anchor missing real call sites
Total mutation labels flowing through the 4 shared helpers	70	65 + 5 (the 6th, <code>mutate_gate_direct</code> 's one call, is an internal fallback <i>inside</i> <code>mutate_gate_selftest_or_red</code> , not an independent top-level call, so instrumenting the function itself already covers it)

Metric	Value	How measured
Raw count of the file's other, bespoke <code>^echo; echo "- banner</code> lines	1,053	Direct grep, with 3 independently spot-checked representative blocks confirming the pattern holds (2 of the 1,053 are non-mutation banners — a “Baseline check” banner and one other — leaving 1,051 genuine per-mutation banner-delimited blocks)

T043's own real measurement of the live 90-day baseline stratification (from `qa-results/fastcycle/us1/green/T043/select_90d_2026-08-23.json` , read directly by this document's author):

Type	Available (90-day window, 2026-05-25 .. 2026-08-23)	Selected	Below required minimum?
Bug	81	5	No
Task	34	5	No
Feature	2	2	Yes (below <code>min_per_type=5</code>)

Of the 30 items in the full 90-day sample, **11** (37%) are marked `excluded_from_duration` — 1 for a `retroactive-registration` data-quality issue (`ATM-742` , an `Opened→closure` `db-write` gap under 60 seconds) and 10 for a `bulk-import-cluster` on `2026-08-15` (`qa-results/spk_import_20260815` , threshold 10) — confirming `T044` 's own planned tracker-data-quality repair pass (Section 2.2, Phase 3 task list) still has real, concrete, measured work ahead of it rather than being a speculative placeholder.

Note on a figure this document's author could NOT independently confirm: the task brief that requested this update also referenced disk-safety findings (a sibling worktree measuring ~49 GB, `/tmp` `tmpfs` at over 80% used) as if they were part of `T043` 's own evidence. On direct investigation, this document's author traced those specific figures to a **different, earlier artefact** — an `R6` host-resource reconnaissance report (`specs/004-fast-dev-cycles/research/`

R6_consumers_and_host.md , persisted in the operator request-history ledger as entry R-2026-09-26-181854) dated **2026-09-26**, two days before this session's own work and unrelated to tasks.md 's T043 (select_sample.py + baseline_replay.sh). Per this document's own §11.4.6 no-guessing discipline, those figures are **not** repeated here as part of today's T043 measurements — attributing a stale, differently-sourced figure to today's work would itself be exactly the kind of misattribution this project's anti-bluff covenant forbids. The real host-headroom figure this session's own conductor *did* check before dispatching today's 4th parallel agent (per §12.12, quoted from docs/CONTINUATION.md 's own live narrative) was **1,008 of 262,144 threads live, 147 GB of 251 GB memory free** — ample headroom, and a genuinely different, current measurement from the stale R6 figures the brief referenced.

T029's own real measurement of pre_build_verification.sh 's section-banner structure: the T029 draft agent's own evidence directory (qa-results/fastcycle/us1/green/T029_draft/extension_phase/) contains candidates_banner_sites_147.tsv — a directly-enumerated list of **147** candidate wrap sites (banner + gate-level sites combined) the draft's own extension-phase tooling (fc_wrap_scan.py) identified in the live file, independently of tasks.md 's own line-text estimate of “94 section-banner wrappers” (that figure describes only the top-level log_section() banners; the draft's own count additionally covers individual gate-level timer sites inside those sections, which is why the two numbers are not the same measurement and are not in conflict).

. Diagrams (Revision)

10.6.1 — Full timeline of today's 31 commits, grouped by task

```

flowchart TD
    subgraph Foundational["T014 Foundational batch"]
        F1["b72361b<br/>T007-T014 landed<br/>T014 GO after rounds 9-12"]
    end
    subgraph US1RED["Phase 3 (US1) RED tests"]
        R1["T015-T026<br/>all 12 RED tests landed<br/>+ independently verified"]
        R2["T027<br/>10/10 mutation-observations"]
    end
    subgraph US1Impl["Phase 3 (US1) implementation"]
        I1["T028 fc_timer.sh<br/>+ T041 cycle_report.py<br/>GO round 1"]
        I2["T039 anchor_citations.py<br/>+ T034 review_record.py<br/>NO-GO round 1 -> fix -> GO round 2"]
        I3["T046 plan_struct_check.py causes<br/>NO-GO round 1 -> fix -> GO round 2"]
        I4["T036 dispatch_stamp.sh<br/>NO-GO round 1 (3 BLOCKING) -> fix -> GO round 2"]
        I5["T038 transcript_ingest.py<br/>+ T020 RED test<br/>GO round 1 (credential-safe)"]
        I6["T043 select_sample.py<br/>+ baseline_replay.sh<br/>race-recovery fix landed<br/>review in flight"]
    end
    subgraph Unrelated["Discovered + recovered (unrelated to Speckit-004)"]
        U1["ATM-1014<br/>42 real defect fixes<br/>staged ~42h, recovered"]
    end
    subgraph ScopeFix["Tooling scope-gap fix"]
        S1["T003 config/fastcycle/<br/>+ sync_all_markdown_exports.sh<br/>>INCLUDED-scope fix"]
    end
    subgraph Draft["In flight now (draft-only, SERIAL)"]
        D1["T029 draft: pre_build_verification.sh"]
        D2["T030 draft: commit_all.sh + exporter"]
        D3["T031 draft: db.go timing"]
        D4["T032 draft: meta_test per-mutation rows"]
    end

    F1 --> R1 --> R2 --> I1 --> I2 --> I3 --> I4 --> I5 --> I6
    I6 --> U1 --> S1
    S1 --> D1
    S1 --> D2
    S1 --> D3
    S1 --> D4

    style F1 fill:#dfd,stroke:#333

```

```

style U1 fill:#ffd,stroke:#333
style I6 fill:#ffd,stroke:#333
style D1 fill:#fdd,stroke:#333
style D2 fill:#fdd,stroke:#333
style D3 fill:#fdd,stroke:#333
style D4 fill:#fdd,stroke:#333

```

10.6.2 — Sequence diagram: the `commit_all.sh` / T043 concurrency-race incident (Section 10.4(a))

```

sequenceDiagram
    participant C as Conductor
    participant CA as commit_all.sh (detached, parent-pointer bump)
    participant T43 as T043 implementer subagent
    participant Fix as Race-recovery fix (c33c1d4)

    C->>T43: Dispatch T043 implementation<br/>(select_sample.py +
baseline_replay.sh)
    par T043 still writing
        T43->>T43: Actively writing baseline_replay.sh to disk
        and commit_all.sh running detached
            C->>CA: Launch parent-pointer-bump commit, detached
            CA->>CA: git add -A style staging<br/>(by design, captures full
tree)
        end
        Note over CA,T43: RACE: commit_all.sh sweeps in<br/>
>baseline_replay.sh's CURRENT mid-write state
        CA-->>C: Commits ae9184a (unintentionally includes<br/>the mid-write
draft)
        T43-->>C: Hands back completed implementation,<br/>own testing catches
determinism-check bug<br/>on the swept-in draft
        C->>C: Root-cause: generic body_hash mechanism<br/>compares wall-clock
timing fields<br/>(inherently non-deterministic)
        C->>Fix: Land the real fix:<br/>compare ONLY verdict_set fields, never
timing
        Fix-->>C: bash -n + shellcheck clean,<br/>py_compile clean, live
selfcheck confirms
        Note over C: Operational lesson recorded:<br/>never run commit_all.sh
while a subagent<br/>is actively writing in the same tree
        C->>C: T043 (both files) still awaits<br/>its own independent Opus-
xhigh review

```

10.6.3 — File-scope dependency graph: why T029 – T032 must run [SERIAL] and cannot be parallelized onto the real tree

```
flowchart LR
    subgraph SharedContentionPaths["Single-owner contention paths (§11.4.58 L3)"]
        PB["pre_build_verification.sh<br/>(the release gate itself)"]
        CM["commit_all.sh<br/>(the single commit/push entrypoint)"]
        DB["workable-items/db.go<br/>(single tracker writer)"]
        MT["meta_test_false_positive_proof.sh<br/>(28,813 lines, the mutation registry)"]
    end
    T29["T029 draft<br/>wire fc_timer.sh into PB"] --> PB
    T30["T030 draft<br/>wire timers into commit_all.sh<br/>+ exporter rows"] --> CM
    T31["T031 draft<br/>add requested_at/applied_at<br/>to db.go"] --> DB
    T32["T032 draft<br/>per-mutation start/end/verdict rows"] --> MT

    T29 -. disjoint from .- T30
    T30 -. disjoint from .- T31
    T31 -. disjoint from .- T32
    T29 -. disjoint from .- T32

    Note1["Each task's SCOPE is disjoint from the other 3<br/>(no two touch the same file) -- so all 4 CAN<br/>draft in parallel as subagents (and did)."]
    Note2["But EACH task's TARGET is a single-owner<br/>contention path per §11.4.58 L3 -- so the<br/>APPLY+TEST+COMMIT step of each MUST run<br/>>serially, conductor-only, never delegated,<br/>never concurrent with another SERIAL task<br/>even on a different file."]

    style PB fill:#fdd,stroke:#333
    style CM fill:#fdd,stroke:#333
    style DB fill:#fdd,stroke:#333
    style MT fill:#fdd,stroke:#333
```

. Current honest status and what remains

Honestly complete and committed (as of the moment of this writing): Phase 1 (Setup, partially — T003 landed, T001 / T002 / T004 – T006 status not re-audited in this revision); Phase 2 (Foundational, T007 – T014) — fully complete,

T014 GO, committed; Phase 3 (US1) RED tests T015 – T026 plus the T027 mutation-observation pass — fully complete; seven Phase 3 implementation tasks (T028 , T034 , T036 , T038 , T039 , T041 , T046) — fully complete with real GO verdicts.

In flight, not yet landed — reported honestly as such, not as done: T043 (`select_sample.py` + `baseline_replay.sh`) — implemented, survived the race-recovery fix, its own independent review dispatched and running as of authoring time, no verdict yet. T029 – T032 (the four [SERIAL] wiring tasks) — draft patches produced by parallel subagents, sitting in scratch evidence directories, **not yet reviewed, applied, tested, or committed by the conductor** (per each task’s own [SERIAL] marker, only the conductor may perform that step).

Not yet started, within Phase 3 (US1) alone: T033 (the full clean- checkout meta-test run, background, alone on the host per §12.8 — the “runtime evidence; long pole” the task’s own text names it), T035 (the `occurred_at` schema migration on `item_history`), T037 (wiring `dispatch_stamp.sh` into the registry writer + `.claude/settings.json`), T040 (build/deploy/QA event emitters), T042 (hand-verifying ≥ 3 items’ stage figures against `cycle_report.py` ’s own output), T044 (the tracker data-quality repair the Section 10.5 stratification data above shows is genuinely needed — 3 status desyncs, 6 exact-duplicate history rows, 3 Reopened-without-prior-closure annotations), T045 (collecting and freezing the stratified baseline itself — the actual deliverable US1 exists to produce), T047 (registering the US1 paired mutations + the tracker status-consistency permanent guard), T048 (the batched independent review of *all* US1 changes together), and T049 (the re-planning checkpoint that gates every later phase claiming a saving).

Not started at all: Phases 4 through 11 — US2 (fast, lossless gates), US3 (stop reopening), US4 (minimal tokens per item), US5 (reliable orchestration), US6 (a host that does not throttle the work), US7 (everything lands in the constitution, clean and pushed), US8 (every consuming project migrated), and the final Polish/ Verification phase. That is **68 of the plan’s 195 tasks accounted for in this revision (35%)**, with **127 tasks (65%) not yet touched** — a figure computed directly from the task ranges named in `tasks.md` (Sections 2.1–2.2 of this document) and the landed-or-in-flight set enumerated above, not estimated.

This is stated plainly, and repeated from Section 10.1, because this project's own governance (§11.4.126, quoted in this document's Section 7) explicitly forbids letting a status report's tone imply more completion than its underlying evidence supports — and a reader arriving at this section directly deserves the same honest floor the rest of this document holds itself to.

. Verification methodology for this revision

This revision was produced by the same discipline Section 9 already establishes for Revision 1, applied to the delta: direct `git log / git show` reads of both repositories (never a summary trusted at face value), direct reads of every cited evidence artefact under `qa-results/fastcycle/us1/green/` (including computing the real byte counts, line counts, and stratification figures in Section 10.5 rather than repeating a prior agent's own citation of them), a direct read of `docs/CONTINUATION.md` 's live top-of-file status block (which is itself the authoritative in-session source for what is genuinely still in flight versus landed), and a direct grep of `docs/requests/history.md` for the originating operator request and the commits' own persisted subagent hand-back reports. This document's own author did **not** independently re-run the fastcycle test suite against the live tree for this revision either, for the same reason Section 9 gives for Revision 1: doing so while `T029 – T032` 's draft agents and `T043` 's review were genuinely still in flight in the same host would have been a real, avoidable perturbation of the shared host state those agents' own verification depends on, and this project's own §12.6/§12.8/§12.12 host-safety discipline exists specifically to prevent exactly that.

One figure this revision could not independently confirm and does not repeat as fact: the task brief that requested this update characterized certain disk-safety findings (a ~49 GB sibling worktree, `/tmp` tmpfs over 80% used) as belonging to today's `T043` work. As documented in [Section 10.5](#), this document's author traced those figures to a different, earlier (2026-09-26) research artefact and does not present them as part of today's delta — per §11.4.6, an unconfirmed or misattributed figure is marked as such rather than silently carried forward.

One correction applied to this document's own prior revision: Section 4.4's original closing paragraph (“this document deliberately does not claim the Foundational batch has reached a clean round-11 GO”) is preserved verbatim, with a Revision-2 correction note inserted immediately after it pointing here — per this

project's own §11.4.226 evidence-class discipline, a prior honest “not yet known” statement is never silently edited into a false “was always known”; it is superseded visibly, with the new evidence that resolves it.

1 1

2 0 2 6 0 9 2 8

3

. Update — — , later (Revision)

This section was requested by the same operator directive Revision 2's own §11.4.140 quote already cites — the request was issued **twice**, verbatim, in the operator request-history ledger (docs/requests/history.md , entries R-2026-09-27-202244 , 2026-09-27 20:22:44 MSK, and R-2026-09-28-112940 , 2026-09-28 11:29:40 MSK): *“Do exhaustive analisys of all work we are doing now via SpecKit and Superpowers with all details, proper summary, fully structured content divided into sections, additional explanations, measurements, analisys of performance and metrics changes, graphs, diagarams and schemes and put it all into constitution Submodule under docs/optimization/ dir. Create full proper Markdown in depth documentation and exports into all mandatory formats! Make sure as soon as it is done everything is fully committed and pushed to all upstreams and let us know when we can review it fully on GitHub! This is extreely urgent and important!”* — dispatched a third time this session via the durable background queue as BG-20260928-2117 (docs/requests/background_queue.md:515), whose entry this revision's own author verified directly. This section is that continuation, produced by directly reading the live git history of both repositories, the live specs/004-fast-dev-cycles/ tasks.md task list, the live docs/workable_items.db (read-only sqlite3 queries), the agent registry, and the operator request-history ledger — the exact verification discipline Sections 9 and 10.8 already establish, applied to today's further delta.

. Executive summary of the delta

Since Revision 2 was committed (constitution 391c4ee , 2026-09-28T13:44:52+05:00, the commit Revision 2's own header cites), roughly seven further hours of real, committed work landed in the **parent** ATMOSphere repository — this constitution submodule received no further fastcycle-tool source

commits in that specific window (its own `HEAD` at authoring time, `acdcad4a9f9f1a9c2e0b9e4c07a999b1abaeba2f`, carries only the Foundational-batch and Phase-3 fastcycle work Sections 4 and 10 already document, plus this revision's own doc-export commit once it lands). The headline changes, each independently re-verified below rather than repeated from a prior agent's own claim:

1. **All four of Revision 2's [SERIAL] draft-only wiring tasks — `T029`, `T030`, `T031`, `T032` — were reviewed, applied, tested, and committed by the conductor.** `T031` (`db.go` `requested_at` / `applied_at` tracker-writer timing instrumentation) landed first, independently, in the constitution submodule (`5cada1d`, `2026-09-28T14:15:01+05:00`, **197 insertions** — this revision's own author independently confirmed both that this commit is a real ancestor of the submodule's current `HEAD` and that `db.go` genuinely contains the claimed `requested_at` / `applied_at` fields, via a direct `grep`, rather than trusting the commit message alone). `T030` + `T032` then landed together in one parent-repo commit (`88171be130c`, 11 files, **+3,044/-220** lines — independently re-confirmed via `git show --stat`, not repeated from the commit message's own prose), wiring `fc_timer.sh` stage timers into `scripts/commit_all.sh` (+406 lines) and per-mutation verdict rows into `scripts/testing/meta_test_false_positive_proof.sh` (+2,193 lines).
2. **A real, currently-observable checkbox/reality gap was found — and remains open at authoring time.** `tasks.md`'s own checkboxes for the Foundational batch (`T007` — `T014`) still read `[ ]` (unchecked) at this revision's authoring moment, despite Section 10.3 of this very document already proving — and this revision independently re-confirming via `git merge-base --is-ancestor` — that the batch reached a clean GO and was committed to the real tracked tree (`b72361b`) hours earlier. Two dedicated, read-only reconciliation subagents were dispatched to investigate the Phase 1 and Phase 2 checkbox state concurrently with this section's own authoring; neither had produced a `complete` registry event by the time this investigation concluded, so this document reports the gap as **found** and the reconciliation as **PENDING**, never as resolved. See [Section 11.3](#).
3. **A pandoc-specific export bug (`ATM-1097`) was found, root-cause-investigated, and worked around — but not yet fixed at its true source.** `workable-items` `export`'s internal pandoc invocation fails with exit 64 on `docs/Issues.md` ("YAML parse exception ... while scanning an alias",

triggered near line 5952 of 1,375+ lines) because an *earlier*, still-unidentified bare `---` line elsewhere in the file puts pandoc's `yaml_metadata_block` auto-detection heuristic into a stuck parsing state. A working, verified, but not-yet-wired workaround (`pandoc --from=markdown-yaml_metadata_block`) restored fresh HTML/PDF/DOCX siblings; the investigation's own interesting, genuinely useful finding — that the *identical* underlying `docs_chain` engine's own `md-to-html / html-to-pdf` transform chain regenerates the same file cleanly with no error — proves the bug is specific to one code path (`workable-items export` 's own internal pandoc call), not a universal pandoc/Markdown defect. See [Section 11.4](#).

4. **A large, independently-researched QA/PM findings-triage effort landed as a strong worked example of the fan-out-research + single-writer-serialization pattern this document's own Section 3 already names.**

`BG-20260928-1956` (47 findings, 4 parallel read-only research subagents split by finding prefix, one later serialized single-writer DB-write pass per §11.4.206) landed as **37 files, +38,850/-7,830 lines** (independently re-confirmed via `git show --stat` , commit `3dc7be47172` — the current `HEAD` of the parent repository at authoring time), creating 41 new workable items, augmenting **5 named** existing items despite the summary's own header claiming “6 existing items augmented” (a real, unresolved off-by-one discrepancy in the source document itself, reported honestly rather than silently corrected), and cross-referencing one multi-item finding to 5 already-open items rather than minting a duplicate. See [Section 11.5](#).

5. **Three further Phase 3 (US1) tests were converted from structurally broken to genuinely test-first via a “reconciliation-by-verification” methodology**

— confirming real on-disk/runtime behaviour against each task's stated acceptance property before marking it done, rather than trusting either a checkbox's absence or its presence. `T020` 's own test was found to have *bluffed green* — its final assertion verdicted a property `HOLDS` purely because a file existed on disk (an artifact-layer fact) while the same test's own earlier block had already empirically proven, by real invocation, that the runtime-layer property it claimed did **not** hold — precisely the §11.4.108/§11.4.226 source-present-runtime-absent bluff class this document's own Section 7 already names. See [Section 11.6](#).

Current real task-count snapshot (measured directly, `grep -c` against the live `specs/004-fast-dev-cycles/tasks.md` , at 2026-09-28T16:20Z): **30 of 195 tasks checked** `[x]` (15.4%), **165 unchecked** `[ ]` (84.6%). This is stated here as an honest, time-stamped snapshot only — Section 11.3 immediately below explains, with direct evidence, why this raw figure is known to **understate** real completion (the Foundational batch's 8 tasks are demonstrably done but unchecked), and this document does not attempt to project what the figure will read by the time a reader sees it, since the reconciliation this section documents was still running at authoring time.

. Timeline of everything landed since Revision

Every commit below is independently re-confirmed by this revision's own author via direct `git show` / `git log` , not repeated from a prior commit message alone.

#	Time (+05:00)	Repo	Commit	What landed
1	13:18	constitution	c33c1d4	T043 's race-recovery fix (already covered by Revision 2 §10.4(a); carried here as the timeline's own starting point)
2	13:27	parent	3e1ea0fa484	ATM-1014 recovery (already covered by Revision 2 §10.4(c))
3	13:29	parent	72d918172e0	T003 Setup files + §11.4.65 scope-gap fix (already covered by Revision 2 §10.4(b)); this revision independently re-confirms this is the earliest commit touching specs/004-fast-dev-cycles/tasks.md in the parent repository's own history (<code>git log --reverse</code>), consistent with tasks.md having only just entered this repository's own tracked history at that point
4	14:15	constitution	5cada1d	T031 — db.go requested_at / applied_at tracker-writer timing instrumentation, 197 insertions — independently confirmed landed and present in the live file

#	Time (+05:00)	Repo	Commit	What landed
5	17:43	parent	46910f486f0	feat(specs-004): add tasks.md with T029/T031/T034/T036/T038/T039/T041/T043/T046 complete — the file's own first tracked commit already carried T001 – T006 (Phase 1) checked [x] ; T007 – T014 (Phase 2) already [] at this point
6	18:56	parent	bddfa7d1f91	chore(specs-004): mark T015/T019/T021/T022/T024/T025/T026 complete — T015 / T024 / T025 converted from non-functional prose-stub/duplicate-logic tests to real subprocess-invocation assertions, self-validated via deliberate mutation; T019 / T021 / T022 / T026 confirmed already fully implemented by earlier parallel work with zero code changes needed

#	Time (+05:00)	Repo	Commit	What landed
7	19:05	parent	572dd6c1c34	Files ATM-1055 per §11.4.197: the T042 hand-verification subagent found a genuine <code>cycle_report.py</code> bug (<code>flag_status_desync()</code> 's bare string <code>!=</code> comparison between the short <code>item_history.event_type</code> form and the §11.4.33 long closure-vocabulary form of <code>items.status</code> , which are <i>by design</i> never byte-identical on a correctly-closed item — firing a false <code>STATUS_DESYNC</code> flag on essentially every genuinely-closed item) while independently hand-verifying 3 real items' timing figures and finding zero divergence on the timing itself
8	19:07	parent	a576d2fbca9	T016 / T017 reconciliation — same T015 -class structural defect (prose-only echo-block acceptance criteria with no polarity switch) found and fixed with real control-needle-proven assertions, both correctly <code>FAIL</code> pending T030
9	19:11	parent	e73c1238145	T016 / T017 marked complete after conductor reconciliation, independently re-verified: both exit 1 (correctly RED)

#	Time (+05:00)	Repo	Commit	What landed
10	19:20	parent	b2e02df55c7	T020 's own test found to bluff green on the T036 / T037 wiring gap — see Section 11.6
11	19:30	parent	8383d37cbbc	T023 (test_cycle_report_red.sh) — 4 “NOT YET IMPLEMENTED” prose stubs converted to real gating assertions with a real mutation-kill proof — see Section 11.6
12	20:17	parent	88171be130c	T030 + T032 landed (+3,044/-220, 11 files); T015 / T029 independently re-confirmed fully GREEN via two independent methods
13	20:50	parent	3dc7be47172	BG-20260928-1956 — 47 QA/PM findings triaged into 41 new + 5-named/6-claimed augmented workable items (Section 11.5's own honest discrepancy note), 37 files, +38,850/-7,830 lines (current parent HEAD at authoring time)

As of this writing, two further items are dispatched and genuinely in flight (neither shows a real complete event in the agent registry, per the agent registry's own tail read at 2026-09-28T16:19:59Z):

- Reconciliation of SpecKit-004 **Phase 1** (T001 – T006) — dispatched 16:16:14Z , in-flight 16:16:16Z , registry key f23ce7750cea7245
- Reconciliation of SpecKit-004 **Phase 2** (T007 – T014) — dispatched 16:19:10Z , in-flight 16:19:59Z , registry key 1f72b97fcc1ced00

. The Phase / Phase checkbox-reality gap – an in-flight, honestly-unresolved finding

This subsection exists specifically because this project's own §11.4.6 no-guessing mandate and §11.4.226 evidence-class discipline forbid reporting a finding's outcome before that outcome genuinely exists — and, at the moment of authoring, it does not yet exist. What follows is exactly what this revision's own author could and did independently verify, and nothing more.

The finding, independently confirmed.

`specs/004-fast-dev-cycles/tasks.md` 's own checkbox state for Phase 2 (the Foundational batch, `T007 – T014`) currently reads **8 of 8 unchecked** (`[ ]`), verified by a direct `grep -n` against the tracked file at authoring time. This is a real, currently-observable checkbox/reality gap: this document's own Section 10.3 (Revision 2, landed hours earlier) already proves, with full evidence — round-11 and round-12 review findings, the landing commit's own final summary line (`run_all: 9 run, 0 failed , 15/15 mutations, test_fc_common_red.sh 233/233, test_host_guard_red.sh 131/131`) — that the Foundational batch reached a clean GO and was committed to the real tracked tree. This revision independently re-confirmed the *committed* half of that claim itself, from scratch, rather than trusting Section 10.3's own prior citation: `git merge-base --is-ancestor b72361b HEAD` , run inside the constitution submodule, returns true (`b72361b` **is** an ancestor of the submodule's current `HEAD`), and the parent repository's own gitlink for `constitution` at its current `HEAD` resolves to `acdcad4a9f9f1a9c2e0b9e4c07a999b1abaeba2f` — a strict descendant of `b72361b` . The Foundational batch is genuinely, verifiably landed on the real tree; its own task file's checkboxes simply do not say so.

Phase 1's own state, reported honestly rather than assumed. `T001 – T006` currently read **6 of 6 checked** (`[x]`) in the live file. This revision's author traced this specific file's own tracked history in the parent repository back to its **very first** commit (`46910f486f0` , `2026-09-28T17:43:17+05:00` — independently confirmed via `git log --reverse --format="%h %cI %s"` against the file, which returns exactly this commit as the earliest) and found `T001 – T006` already checked `[x]` at that first commit, before any of the reconciliation work this section documents began. This revision's author explicitly does **not** know, and does not claim to know, whether this reflects the specific Phase-1 reconciliation subagent

(registry key `f23ce7750cea7245` , dispatched `16:16:14Z` , described in this section's own timeline) having already independently confirmed the same thing and found nothing to correct, or whether Phase 1's checkboxes were simply already accurate from the moment the file first entered this repository's tracked history and that subagent's own investigation is still in progress toward the same conclusion. Either is consistent with the evidence available to this document's author; neither is asserted as fact.

Why this matters, and what is genuinely still open. The Phase 2 gap this subsection documents is not a cosmetic bookkeeping issue — it is the exact class of drift this project's own SpecKit-004 feature (Section 1's executive summary) exists to fix: a status artefact silently diverging from real completion state, in **either** direction (a checkbox can lag real completion, as demonstrated here for Phase 2). A directly-observed sibling instance of the *opposite* direction — this document's own author found it while reading `docs/CONTINUATION.md` for this section's own verification (Section 11.11) — is `docs/CONTINUATION.md`'s own top-of-file "RESUME-HERE ADDENDUM 4" block, timestamped `20:44Z` , which states "Phase D commit is the very next action"; the commit that block describes (`BG-20260928-1956`) is, at this revision's own authoring moment, already the parent repository's current `HEAD` (`3dc7be47172` , Section 11.2's own timeline row 13) — meaning that specific status block's own "next action" text is itself now stale relative to the real committed state, an instance of the identical status-vs-reality drift this subsection is about, observed from the opposite direction. This is noted here as a directly-observed fact, not investigated further as part of this revision's own scope — fixing it belongs to whichever agent next regenerates `docs/CONTINUATION.md`'s own top block. The operative discipline this session's own work demonstrates, and which this document recommends as a standing practice (extending Section 8's own lessons-learned register): **periodic reconciliation-by- verification, never reconciliation-by-trust** — a checkbox is marked `[x]` only after its task's own stated acceptance property is independently re-confirmed against real, current, on-disk or runtime evidence, and a checkbox left `[ ]` past its task's real completion is treated with the same seriousness as a checkbox wrongly marked `[x]` too early, because both are the identical failure mode (status divergent from reality) observed from opposite directions. This document reports the Phase 2 gap as **FOUND** and the two dispatched reconciliations as **PENDING** — not fixed, not confirmed, not assumed — as of `2026-09-28T16:20Z` .

. ATM- : a pandoc-specific YAML-mis-detection bug, isolated to one code path

While regenerating `docs/Issues.md` 's HTML/PDF/DOCX siblings as part of the BG-20260928-1956 closing work (Section 11.5), `workable-items export --db docs/workable_items.db --out-dir docs` 's own internal pandoc invocation failed with:

```
Error parsing YAML metadata at .../docs/Issues.md (line 5952,
column 1): YAML parse exception at line 1, column 1, while
scanning an alias: did not find expected alphabetic or numeric
character (exit 64)
```

This revision's author read `ATM-1097` 's own tracked description directly (`sqlite3 docs/workable_items.db "SELECT description FROM items WHERE atm_id='ATM-1097';"` , read-only) rather than paraphrasing a prior agent's summary of it, and confirms the investigation's own real findings, exactly as recorded there:

- Line 5952 is a bare `---` line — an ordinary Markdown horizontal rule, one of 116 such lines in the file. Isolating that specific line plus its surrounding ~18 lines into a standalone test file and running pandoc directly against it **succeeds** with no error — proving this specific `---` is not, in isolation, the trigger.
- This means pandoc's own `yaml_metadata_block` reader extension is being put into an open/unclosed YAML-parsing state by some *earlier* bare `---` line elsewhere in the file that pandoc's own heuristic misidentifies as the *start* of a YAML front-matter block — and the failure only surfaces once the (still-open) parse reaches content, near line 5952, that cannot validly continue as YAML (the “while scanning an alias” wording is consistent with pandoc's YAML parser hitting a bare or misplaced `*` character it attempted to interpret as a YAML alias reference).
- **A genuinely useful, independently interesting finding, worth stating explicitly:** the *same* underlying `docs_chain` engine's own `md-to-html / html-to-pdf` transform chain (§11.4.106, the mechanism this project's own broader export/sync tooling is built on) regenerates the identical `docs/Issues.md` file **cleanly, with no error**, when invoked through its own pipeline

rather than through `workable-items export` 's own internal pandoc call. This is direct, positive evidence that the defect is specific to *one particular code path*'s pandoc invocation (a flag or reader-mode difference between the two callers), not a defect in pandoc's YAML-detection heuristic in general, nor in `docs/Issues.md` 's own Markdown content, which the item's own description independently confirms is "correct and complete — the failure is pandoc-side only."

The working workaround applied, verified, and honestly scoped as not-yet-

wired: `pandoc --from=markdown-yaml_metadata_block ...` (the extension explicitly disabled) for the HTML and DOCX legs, followed by `weasyprint --stylesheet constitution/styles/default-pdf.css docs/Issues.html docs/Issues.pdf` for the PDF leg. Verified per `ATM-1097` 's own recorded evidence: exit 0 on all three, no YAML-parse error, non-degenerate output (`html=2.45MB` , `pdf=2.97MB` , `docx=828KB` , versus the prior stale copies), §11.4.73-styled. The item's own description is explicit that this workaround is **not yet wired** into `scripts/testing/sync_all_markdown_exports.sh` or `workable-items export` itself — both still use the extension-enabled pandoc default and will regress to the exit-64 failure on their very next invocation against `docs/Issues.md` — and that the true root-cause bisection (identifying the specific *earlier* --- line responsible) remains open, tracked, and explicitly deferred pending evidence that the flag change is safe across the rest of the `docs/*.md` corpus (not yet checked for any file other than `Issues.md`). `ATM-1097` remains

11	5	20260928	1956
----	---	----------	------

 Status: Queued , Type: Bug , exactly as its own tracker row shows — this document does not, and the item itself does not, claim it closed.

. BG- - : fan-out research + single-writer serialization, worked example

`BG-20260928-1956` is a strong, concrete continuation of this document's own Section 3 ("Superpowers methodology as applied here") — a real instance of `superpowers:subagent-driven-development` 's fan-out-then-serialize shape, worth documenting as its own case study rather than only as a line in Section 11.2's timeline table, because the shape of the work is itself the methodological point.

The task. The project's Project-Management and QA teams delivered two source documents under `docs/planning/2026.09/issues_01/` : a 47-row Russian-language findings registry (`Comments_list.xlsx` , three sheets) and a 321-paragraph narrative summary with a proposed 6-step work order (`Current_Comments_to_codebase.docx`). Every one of the 47 findings required independent research (deep git history, direct current- `HEAD` source verification, cross-reference against the live `docs/workable_items.db` , product-coverage determination, effort estimation) before it could be either filed as a new workable item, matched to an existing one, or cross-referenced without duplication.

The fan-out. Four parallel subagents, split by the findings registry's own real `No` -column prefix (never an arbitrary even split — the partition follows the data's own natural grouping), each independently researched its own disjoint slice:

Group	Scope	Count	Evidence file
A	Audio (AU- *)	15	qa-results/qa_pm_20260928/ group_a_audio.json
B	Security + Authorization (SEC- * / AUTH- *)	14	qa-results/qa_pm_20260928/ group_b_security_auth.json
C	Build/Platform/ Bootloader/ Kernel/ DeviceTree/ BoardBinding/ OTA/Repo (BLD- * / OTA- * / REP- *)	11	qa-results/qa_pm_20260928/ group_c_build_platform.json
D	Process/ Licenses/ AppSigning/ SpeakerConfig/ HDMI/ CameraMic/ Compatibility (LIC- * / PRD- *)	7	qa-results/qa_pm_20260928/ group_d_process_license_product.json

15 + 14 + 11 + 7 = 47 — every one of the 47 findings accounted for, none silently dropped, independently re-confirmed by this revision’s own author by summing the group’s own stated counts directly from BG-20260928-1956_SUMMARY.md rather than trusting the “47/47” headline alone.

The serialization. The four parallel research streams’ own outputs were entirely read-only against the shared tracker — no research subagent itself wrote to docs/workable_items.db . The actual database writes (41 new items minted, 6 existing items augmented via a read-then-concatenate update that never discarded prior content, 1 finding cross-referenced to 5 already-open items rather than duplicated) were performed by a **single, serialized writer pass**, following this project’s own

§11.4.206 “exactly one writer per shared single-source-of-truth entity” discipline — the same discipline that governs every other write to this tracker throughout the whole Speckit-004 cycle this document documents (Section 2.3’s [SERIAL] marker convention, applied here to a QA-intake workflow rather than a fastcycle-tool task).

The outcome, independently re-confirmed — with one honest discrepancy surfaced, not smoothed over. This revision’s author independently ran `sqlite3 docs/workable_items.db "SELECT COUNT(*) FROM items;"` (read-only) at authoring time and it returns **921** — not the **920**

BG-20260928-1956_SUMMARY.md ’s own “Outcome” section states as the post-write total (itself claimed there as a delta from a pre-write 879, matching “41 new items” exactly: $879 + 41 = 920$). The real, current, independently-measured count is one item higher than the summary’s own claimed figure. This revision’s author does **not** know, and does not guess, the specific cause of that +1 — a reasonable, evidence-consistent possibility is that at least one further item (this revision’s own Section 11.4 names ATM-1097 , and Section 11.2 names ATM-1055 , both filed in connection with this same work) was minted after the summary document’s own count was taken and before this revision’s own query — but this is recorded explicitly as an inference, never as a confirmed fact, per §11.4.6. The summary’s own further claim of a successful `workable-items validate` pass with “all invariants satisfied” is recorded here as the summary’s own claim, not independently re-run by this revision’s author (re-running a full validate pass was judged an unnecessary, non-disjoint perturbation of the shared tracker state while this document’s own investigation and the two Phase 1/Phase 2 reconciliation subagents were concurrently active against overlapping artefacts — the same §12.6/§12.8/§12.12 host-safety reasoning Section 9 and Section 10.8 already give for not independently re-running the live fastcycle test suite).¹

. Reconciliation-by-verification: three more worked examples from Phase (US)

Three commits today extend this document’s own running theme (Section 8.2: “real execution finds what review, by itself, cannot”) with a distinct but related lesson — **a test that has never genuinely been run to completion, or that verdicts on the**

wrong evidence layer, is itself a defect indistinguishable from the product defect it was meant to catch, until someone actually reconciles the test's own claimed behaviour against real, current evidence.

T020 — a test that bluffed green on an artifact-layer fact.

`test_token_attribution_red.sh` 's own final assertion for its "PART A" property verdicted `HOLDS` purely because `dispatch_stamp.sh` **existed as a file on disk** (`T036` had landed) — an artifact-layer fact (§11.4.108 layer 2) standing in for the runtime-layer invariant (layer 3) the property actually claims: that a dispatch lacking an `item=` field is correctly *rejected*, not silently accepted. The same test's own earlier "A-real" block had already, by a real invocation of the real, unmodified `agent_registry_writer.sh`, empirically proven the **opposite** — `item=` -less dispatches are still silently accepted today with zero item attribution (`FOUND=0`) — directly contradicting its own later `HOLDS` verdict two lines below. Confirmed by a direct `grep : zero references to dispatch_stamp anywhere in agent_registry_writer.sh or .claude/settings.json` — `T037` (the `[SERIAL]` task that would actually wire `dispatch_stamp.sh` in) is genuinely unwired, matching `dispatch_stamp.sh` 's own in-source header comment ("NOT YET WIRED (deliberately, by design — `T037`'s job)"). The fix does not merely patch the one broken assertion: it requires the property's verdict to satisfy **both** a structural wiring reference (`grep for dispatch_stamp in the writer + settings`) **and** cross-validation against the real `A-real` result already captured, with a distinct third branch for "wiring reference present but not taking effect" (a source-only reference with no observed behavioural change) — self-validated via a temporary, backed-up, restored one-line no-op comment injected into `agent_registry_writer.sh`, confirming that new branch correctly fires before the fix was trusted.

T023 — four prose stubs converted to real gating assertions, with a real mutation-kill proof. `test_cycle_report_red.sh` carried 4 blocks reading "NOT YET IMPLEMENTED" where a real assertion was meant to sit. All 4 were replaced with real gating assertions invoking `cycle_report.py` against all 4 of its own fixtures (`missing_record`, `empty_window`, `golden_atm953` — run against the **live** database via `--item ATM-953`, and `negative_control_all_present`), each exact-match-diffed against the fixture's own `expected*` field. The fix's own self-validation went further than merely making the assertions pass: it ran **2 control needles** (an empty-window discrimination check against a real 132-record window;

a comparator discrimination check between two fixtures' outputs) **plus a real mutation-kill proof** — deliberately mutating the tool's own stage-selection logic to compare the wrong stage, confirming exactly one finding fired as a result, then cleaning the mutation up immediately — proving the new assertions genuinely catch the defect class they are meant to gate, rather than merely agreeing with whatever the tool currently happens to output (the exact “test agrees with code” anti-pattern this project's own §11.4.245 anchor names). Independently re-verified before commit: 15/15 pass, exit 0, `bash -n clean`.

T016 / T017 — the same T015 -class structural defect, found and fixed by the same pattern proven three times earlier the same session. Both tests

asserted a permanent, un-flippable fact with no polarity switch: T016

(`test_fc_commit_stage_timer_red.sh`) asserted `fc_timer.sh` was **absent**, a claim T028 had already made permanently false weeks/hours earlier; its two named acceptance criteria (one TSV row per executed dry-run-commit stage; one push row per remote with a read-back tip) were prose-only echo blocks, never real assertions. The fix adds a `RED_MODE` polarity switch

(`FC_TIMER_COMMIT_RED_MODE` , default `0` since T028 's landing is now permanent, `=1` retained as an audit-only escape hatch) and replaces the two echo blocks with real, control-needle-proven assertions (a `fc_timer_start` / `fc_timer_end` call-site check plus a TSV-row count; an `ls-remote` read-back plus a push-row count) — both of which, at the moment they were written, correctly `FAIL` , because T030 (the task that would actually wire the timers in) had not yet landed, and both are explicitly designed to flip `PASS` the instant T030 does land (which it since has, per Section 11.2's own timeline — this revision did not independently re-confirm whether T016 / T017 have themselves since been re-run post- T030 and flipped; that re-confirmation is recorded here as a genuinely owed, tracked follow-up rather than assumed). This is the same reconciliation pattern this session's own commit history records having already been “proven 3x earlier” against T015 / T024 / T025 (Section 11.2's own timeline, commit `bddfa7d1f91` , 18:56) — a fourth and fifth independent instance of the identical defect class (a permanently-failing or permanently-passing assertion with no way to ever observe the property it claims), found by the same reconciliation-by-verification discipline each time.

. Measured figures (Revision)

Every figure below was read directly from a real, current artefact at authoring time — never estimated — per this document's own Section 9 discipline, continued through Revision 2 (Section 10.8) and into this revision.

Metric	Value	How measured
tasks.md checkboxes checked ([x])	30 / 195 (15.4%)	grep -c '^s*- \[x\]' specs/ 004-fast-dev-cycles/ tasks.md , at 2026-09-28T16:20Z
tasks.md checkboxes unchecked ([])	165 / 195 (84.6%)	grep -c '^s*- \[\]' specs/ 004-fast-dev-cycles/ tasks.md , same moment
T031 diffstat (db.go timing instrumentation)	1 file, +197/-1	git show --stat 5cada1d , constitution submodule
T030 + T032 combined diffstat	11 files, +3,044/-220	git show --stat 88171be130c
BG-20260928-1956 diffstat	37 files, +38,850/-7,830	git show --stat 3dc7be47172
BG-20260928-1956 findings triaged	47 / 47 (15 + 14 + 11 + 7)	Direct group-count sum, BG-20260928-1956_SUMMARY.md
BG-20260928-1956 new workable items minted	41 (ATM-1056 – ATM-1096)	BG-20260928-1956_SUMMARY.md ; the summary's own claimed 879 → 920 delta live-rechecked — see the honest +1 discrepancy in Section 11.5
docs/ workable_items.db live item count, this revision's authoring moment	921	sqlite3 docs/ workable_items.db "SELECT COUNT(*) FROM items;" , run directly — 1 higher than the summary's own claimed 920, unexplained-but-plausible per Section 11.5

Metric	Value	How measured
BG-20260928-1956 existing items augmented	Summary header states 6 ; only 5 named in the summary's own text	BG-20260928-1956_SUMMARY.md — a real, unresolved discrepancy in the source document, reported honestly rather than silently corrected
ATM-1097 reproduction line	line 5952 of docs/ Issues.md	ATM-1097 's own tracked description, sqlite3 read
ATM-1097 bare --- line count in docs/Issues.md	116	ATM-1097 's own tracked description
Foundational batch (T007 — T014) landed-but- unchecked	8 / 8 tasks checkbox- [] despite a confirmed committed GO	Direct grep of tasks.md + git merge-base --is-ancestor b72361b HEAD (true)
Phase 1 (T001 — T006) checkbox state	6 / 6 checked [x] at this file's own first tracked commit	git log --reverse on tasks.md , earliest commit 46910f486f0
In-flight reconciliation agents, this section's own subject	2 , neither complete	Agent registry tail, keys f23ce7750cea7245 (Phase 1) and 1f72b97fcc1ced00 (Phase 2)

A figure this document's author could NOT independently confirm: whether T016 / T017 's two assertions (Section 11.6) have themselves been re-run since T030 landed and confirmed to flip PASS as designed. No tracked artefact or agent registry entry available to this revision's author at authoring time records that specific re-run event — marked UNVERIFIED / an owed follow-up rather than assumed, per §11.4.6.

. Diagrams (Revision)

11.8.1 — Sequence diagram: the BG-20260928-1956 fan-out-research + single-writer-serialization pattern

```
sequenceDiagram
    participant C as Conductor
    participant A as Research subagent A (AU-*, 15)
    participant B as Research subagent B (SEC-*/AUTH-*, 14)
    participant D as Research subagent C (BLD-*/OTA-*/REP-*, 11)
    participant E as Research subagent D (LIC-*/PRD-*, 7)
    participant W as Single-writer DB pass (§11.4.206)
    participant DB as workable_items.db

    C->>A: Dispatch, read-only research scope: group A
    C->>B: Dispatch, read-only research scope: group B
    C->>D: Dispatch, read-only research scope: group C
    C->>E: Dispatch, read-only research scope: group D
    par All four run fully concurrently, no shared write path
        A-->>C: group_a_audio.json (15 findings researched)
    and
        B-->>C: group_b_security_auth.json (14 findings researched)
    and
        D-->>C: group_c_build_platform.json (11 findings researched)
    and
        E-->>C: group_d_process_license_product.json (7 findings researched)
    end
    Note over C: 15+14+11+7 = 47/47 accounted for, none dropped
    C->>W: Dispatch ONE serialized writer pass<br/>over all 4 research outputs
    W->>DB: 41 new items minted
    W->>DB: 5 named existing items augmented<br/>(summary claims 6 -- honest discrepancy, 11.5)<br/>(read-then-concatenate, never discard)
    W->>DB: 1 finding cross-referenced to<br/>5 already-open items (no duplicate)
    DB-->>W: summary claims 920 (was 879);<br/>live re-query measures 921 (11.5)
    W-->>C: Regenerate docs/Issues.md + docs/Fixed.md from the DB
    Note over C,DB: Exactly ONE writer touched the DB at any moment --<br/>>the four research streams never contended with each other<br/>or with the write path
```

11.8.2 — Phase/task-completion state overview (Revision 3, honest snapshot)

flowchart TD

```
P1["Phase 1: Setup<br/>T001-T006<br/>6/6 checked in tasks.md<br/>(Phase-1 reconciliation IN FLIGHT<br/>to confirm this reading is real)"]
P2["Phase 2: Foundational<br/>T007-T014<br/>0/8 checked in tasks.md<br/>YET independently confirmed<br/>COMMITTED + G0 (b72361b)<br/>-- a real checkbox/reality gap<br/>(Phase-2 reconciliation IN FLIGHT)"]
P3RED["Phase 3 (US1) RED tests<br/>T015-T026: all 12 landed<br/>independently verified"]
P3IMPL["Phase 3 (US1) implementation<br/>T028/T034/T036/T038/T039/T041/T046: G0<br/>T029/T030/T031/T032: applied + committed today<br/>T043: implemented, review in flight"]
P3REM["Phase 3 (US1) remaining<br/>T033/T035/T037/T040/T042/T044/T045/<br/>T047/T048/T049: not started"]
P4T011["Phases 4-11 (US2..US8 + Polish)<br/>NOT STARTED<br/>127 of 195 total plan tasks"]
```

P1 -.honest uncertainty.-> P2

P2 --> P3RED --> P3IMPL --> P3REM --> P4T011

style P1 fill:#ffd,stroke:#333

style P2 fill:#fdd,stroke:#333

style P3RED fill:#dfd,stroke:#333

style P3IMPL fill:#ffd,stroke:#333

11 9 style P3REM fill:#eee,stroke:#333

style P4T011 fill:#eee,stroke:#333

. Constitution anchors newly exercised

Every anchor already quoted in Section 7 continued to apply throughout today's own work (§11.4.6, §11.4.50, §11.4.58/§11.4.94/§11.4.230(C), §11.4.102/§11.4.115(F), §11.4.134, §11.4.147/ATM-858 D1, §11.4.201, §11.4.209/§11.4.211) — this subsection records only the anchors this revision's own new material exercises concretely for the first time in this document.

§11.4.197 — Research / kicked-off-work completion mandate. ATM-1055

(Section 11.2's own timeline, commit 572dd6c1c34 , 19:05) is a direct instance: a genuine bug discovered as a *side effect* of unrelated hand-verification work was filed as its own tracked item rather than fixed inline and forgotten, or silently left unrecorded — "loss of requirements, failure to incorporate them is FORBIDDEN," per the anchor's own operative text.

§11.4.206 — Shared single-source-of-truth: exactly ONE writer.

BG-20260928-1956 's own fan-out-then-serialize shape (Section 11.5, Figure 11.8.1) is a direct, textbook application: four independent research streams produced data concurrently with zero write contention, and exactly one serialized pass performed every write to the shared `workable_items.db` — never a second concurrent writer.

§11.4.226 — Evidence-class-at-closure. T020 's own bug (Section 11.6) is a precise, concrete instance of this anchor's core failure mode: a runtime-layer claim ("an item-less dispatch is rejected") verdicted `HOLDS` on artifact-layer evidence alone (a file's mere existence on disk) — exactly the class of defect this anchor exists to make structurally visible rather than silently trusted.

§11.4.245 — Oracle-problem-first test authoring. T023 's own mutation-kill proof (Section 11.6) directly demonstrates this anchor's own operative distinction: an assertion is only a genuine test if it can be shown to disagree with a deliberately broken implementation, not merely agree with whatever the current implementation outputs — proven here by literally breaking the tool on purpose, confirming exactly one finding fired, and restoring it.

. Current honest status and what remains

Stated plainly, continuing exactly the discipline Section 10.7 already establishes:

SpecKit-004 is still not complete. Relative to Section 10.7's own accounting:

- T029 , T030 , T031 , T032 — all four of Revision 2's `[SERIAL]` draft-only wiring tasks — have moved to **landed and committed**.
- T043 's own independent review remains **in flight**, with no confirmed verdict recorded as of this writing (unchanged from Section 10.7's own report).
- T033 , T035 , T037 , T040 , T042 , T044 , T045 , T047 , T048 , T049 — the remainder of Phase 3 (US1) — remain **not started**, unchanged from Section 10.7.
- Phases 4 through 11 (US2 through US8 plus the final Polish/Verification phase) remain **not started at all**, unchanged from Section 10.7 — still **127 of the plan's 195 tasks** untouched.

- The Phase 1 / Phase 2 checkbox-reconciliation finding this revision's own Section 11.3 documents is, as of this writing, **genuinely open** — found, not resolved.

This is repeated here, in this revision's own closing section, for the identical reason Section 10.7 gives for repeating it there: this project's own §11.4.126 governance explicitly forbids letting a status report's tone imply more completion than its underlying evidence supports, and a reader arriving at this section directly deserves the same honest floor the rest of this document holds itself to.

. Verification methodology for this revision

This revision was produced by the same discipline Sections 9 and 10.8 already establish, applied to today's own delta: direct `git log / git show --stat` reads of both repositories (diffstat figures independently re-run, never taken from a commit message's own prose alone); a direct, read-only `sqlite3` query of the live `docs/workable_items.db` for ATM-1097 's own full tracked description and the DB's own item count; a direct `grep -c` of the live `specs/004-fast-dev-cycles/tasks.md` for the real, current checkbox counts; a direct `git merge-base --is-ancestor` check (run fresh by this revision's own author, not trusted from Section 10.3's own prior citation) to independently re-confirm the Foundational batch's own committed state; and a direct read of the agent registry's own tail to establish, honestly, that the two Phase 1/Phase 2 reconciliation subagents this section documents had **not** produced a `complete` event as of the moment this investigation concluded. This revision's own author did **not** independently re-run the fastcycle test suite, the full `workable-items validate` pass, nor any other write-adjacent operation against the live, shared tree — for the identical §12.6/§12.8/§12.12 host-safety reason Sections 9 and 10.8 already give: two read-only reconciliation subagents and this document's own investigation were genuinely concurrent against overlapping shared state (the tracker database, the task file, the constitution submodule) at authoring time, and an additional heavy or write-adjacent operation launched from this revision would have been an avoidable perturbation of exactly the shared evidence those other, independently-dispatched agents' own verification depends on.

One figure this revision could not independently confirm and does not repeat as settled fact: whether T016 / T017 's two control-needle-proven assertions have themselves been re-run and observed to flip PASS following T030 's landing (Section 11.6's own closing note; Section 11.7's own "figure this document's author could NOT independently confirm" entry). Per §11.4.6, this is recorded as an owed, tracked follow-up rather than assumed complete merely because its precondition (T030 landing) has itself since occurred.

The reconciliation this section's own subject describes was, by design, still unresolved at the moment this section was completed. This document does not, and per its own governing discipline cannot, state whether the Phase 1 and Phase 2 checkbox reconciliations concluded that the checkboxes were already correct, required correction, or surfaced some further finding neither this revision's own author nor its dispatching conductor anticipated. A future revision, once those two subagents' own real, independently-verified completion events exist, is the correct and only honest place to report that outcome.