README

Helix Constitution–Powered SpecKit–Superpowers Bridge: Implementation Documentation

Revision: 3 **Last modified:** 2026-09-23T00:00:00Z **Description:** Master index and architecture overview for the Helix Constitution–powered SpecKit–Superpowers Bridge implementation. **Maintainer:** HelixDevelopment **Scope:** `constitution/docs/research/extensions/speckit_superpowers/implementation/`

> **AMENDED 2026-09-22 (HXC-159 Phase B3).** This document's architecture was originally drawn as 7 layers. Layer 5 ("SuperBridge MCP") does not exist as software and never did. The design is corrected to its true **6-layer** shape below. See Architecture Amendment for the finding, the substitution, and why this is a correction rather than a removal.

Table of Contents

1. Executive Summary
2. Architecture Amendment (2026-09-22)
    - Bridge layer-health (B1 probe result)
3. System Architecture
    - Architecture Diagram
    - Layer 1 — Developer Workstation
    - Layer 2 — Spec-Kit Core (Governance Layer)
    - Layer 3 — SuperSpec Bridge (Orchestration Layer)
    - Layer 4 — SuperB Extension (Discipline Enforcement)
    - Layer 5 (retired design) — SuperBridge MCP: never built, superseded
    - Layer 5 — Helix LLM (Inference Layer)
    - Layer 6 — Distributed Host Cluster
4. Data Flow
    - Feature Lifecycle Through the Layers
    - Nano-Task Decomposition Flow
5. Document Index
6. Quick Start
7. Key Design Decisions

8. Constitutional Mandates Applied
9. Directory Structure
10. References

Executive Summary

This documentation defines the **Helix Constitution–Powered SpecKit–Superpowers Bridge** — a unified, multi-layer development system that integrates the full SDD (Specification-Driven Development) lifecycle of Spec-Kit with the agent-skill ecosystem of Superpowers, orchestrated through the SuperSpec and SuperB bridge extensions, with distributed inference powered by Helix LLM backed by a `llama.cpp` RPC cluster.

Problem It Solves

Modern software development with LLM-assisted agents suffers from four structural problems:

Problem	Manifestation	Root Cause
Context fragmentation	Each tool (specifier, planner, coder, reviewer) runs in its own silo; no shared understanding of what is being built	No canonical bridge between specification artifacts and execution tools
Discipline drift	Specifications written but never consulted during implementation; plans that do not reflect reality; "fix it later" culture	No mechanical enforcement that the spec *governs* the work
Inference latency & cost	Centralized LLM APIs introduce round-trip latency, rate limits, and per-token billing that blocks autonomous-loops	No local-first distributed inference pipeline
Anti-bluff gap	Tests pass but features do not work; configuration-only PASS; metadata-only PASS; grep-without-runtime PASS	No §11.4 positive-evidence mandate wired into the development lifecycle

This bridge closes all four gaps by composing Spec-Kit's `specify → plan → tasks → implement` lifecycle with Superpowers' `brainstorm / debug / review / test` skills under a single Constitution-governed execution

framework, all powered by local distributed inference that keeps the autonomous loop running without API rate-limit interruptions.

What This Documentation Covers

1. The **6-layer system architecture** (corrected 2026-09-22 — the original design's Layer 5 "SuperBridge MCP" was found not to exist as software; see the Architecture Amendment) with each layer's responsibilities, technologies, and contractual interfaces.
2. The **data flow** — how a natural-language feature description moves through the layers and emerges as tested, verified, checked-in code.
3. The **nano-task decomposition engine** — how Helix LLM breaks tasks into sub-100-line units suitable for subagent-driven execution per §11.4.70.
4. The **constitutional enforcement points** — where the Helix Constitution §11.4 anti-bluff covenant, TDD-fix-discipline, mechanical enforcement, and reporting-directive system plug into the bridge.
5. The **onboarding path** — how any consuming project (e.g., ATMOSphere) integrates the Constitution submodule, wires its CLI-agent context carriers, installs the bridge extensions, and goes from zero to full autonomous-loop development.

Architecture Amendment (2026-09-22)

HXC-159 Phase B3 finding. The original design below (§2026-07-24) drew the Bridge as a **7-layer** stack whose middle execution tier, "Layer 5 — SuperBridge MCP", was an MCP (Model Context Protocol) server exposing the Superpowers skill suite as callable tools. Phase B1's layer-presence probe (`constitution/scripts/extensions/probe_bridge_layers.sh`) and Phase B3's independent re-derivation both confirm: **"SuperBridge MCP" does not exist as installable software, and never did.**

Two independent derivations (re-verified 2026-09-22, evidence at `qa-results/hxc159/l5_finding/20260922T210000Z/`):

1. **GitHub** — a repository search returns `total_count: 0`. A broader, authenticated code search (reaching further than the original unauthenticated check) returns 17 hits, and every single one is inside this project's OWN GitHub mirror (`HelixDevelopment/code`) — our own investigation documents about the non-existence finding, not the software itself. Zero hits from any third-party repository.
2. **npm** — a fuzzy registry search over 95,472 ranked results contains no package named `superbridge-mcp` or `SuperBridge MCP` (the nearest lexical neighbour, `@chenmk/superflow`, is a *different* tool). A direct exact-name

lookup (`GET /superbridge-mcp`) returns **HTTP 404 — package not found**, the strongest single derivation in the set.

A third finding, from the design itself: SuperB (Layer 4, the layer immediately adjacent to the claimed L5 execution tier) states in its own published README that it *"owns no plan, task store, execution lifecycle, completion state, or convergence command."* No adjacent, real layer in the original design ever claimed the execution role either — L5 was a pure design gap from inception, not a missing installation.

The substitution (proven by a real worked run, not asserted). The execution tier that actually fills this role — and has been WIRED and available in this repository throughout — is **Spec Kit** + **Superpowers** + **Claude Code's own skill/subagent machinery**. Every responsibility the original design assigned to the phantom MCP server has a real owner among the other six layers:

Original claimed L5 responsibility	Actual owner
Skill contract exposure (the `superbridge_*` tool list)	Claude Code's own native skill system (`superpowers:*` skills, loaded directly — no MCP indirection needed)
Task decomposition / dispatch loop	`speckit-tasks` (Layer 2), via SuperSpec's (Layer 3) template/preset resolution
Subagent execution of an individual task	Claude Code's native `Agent` tool — a genuine dispatched, isolated-context subagent (§11.4.20/§11.4.70)
Tool-contract verdict emission (PASS/FAIL + evidence path)	The dispatched subagent's own structured report + independent re-verification by the orchestrating session (§11.4.5/§11.4.69)

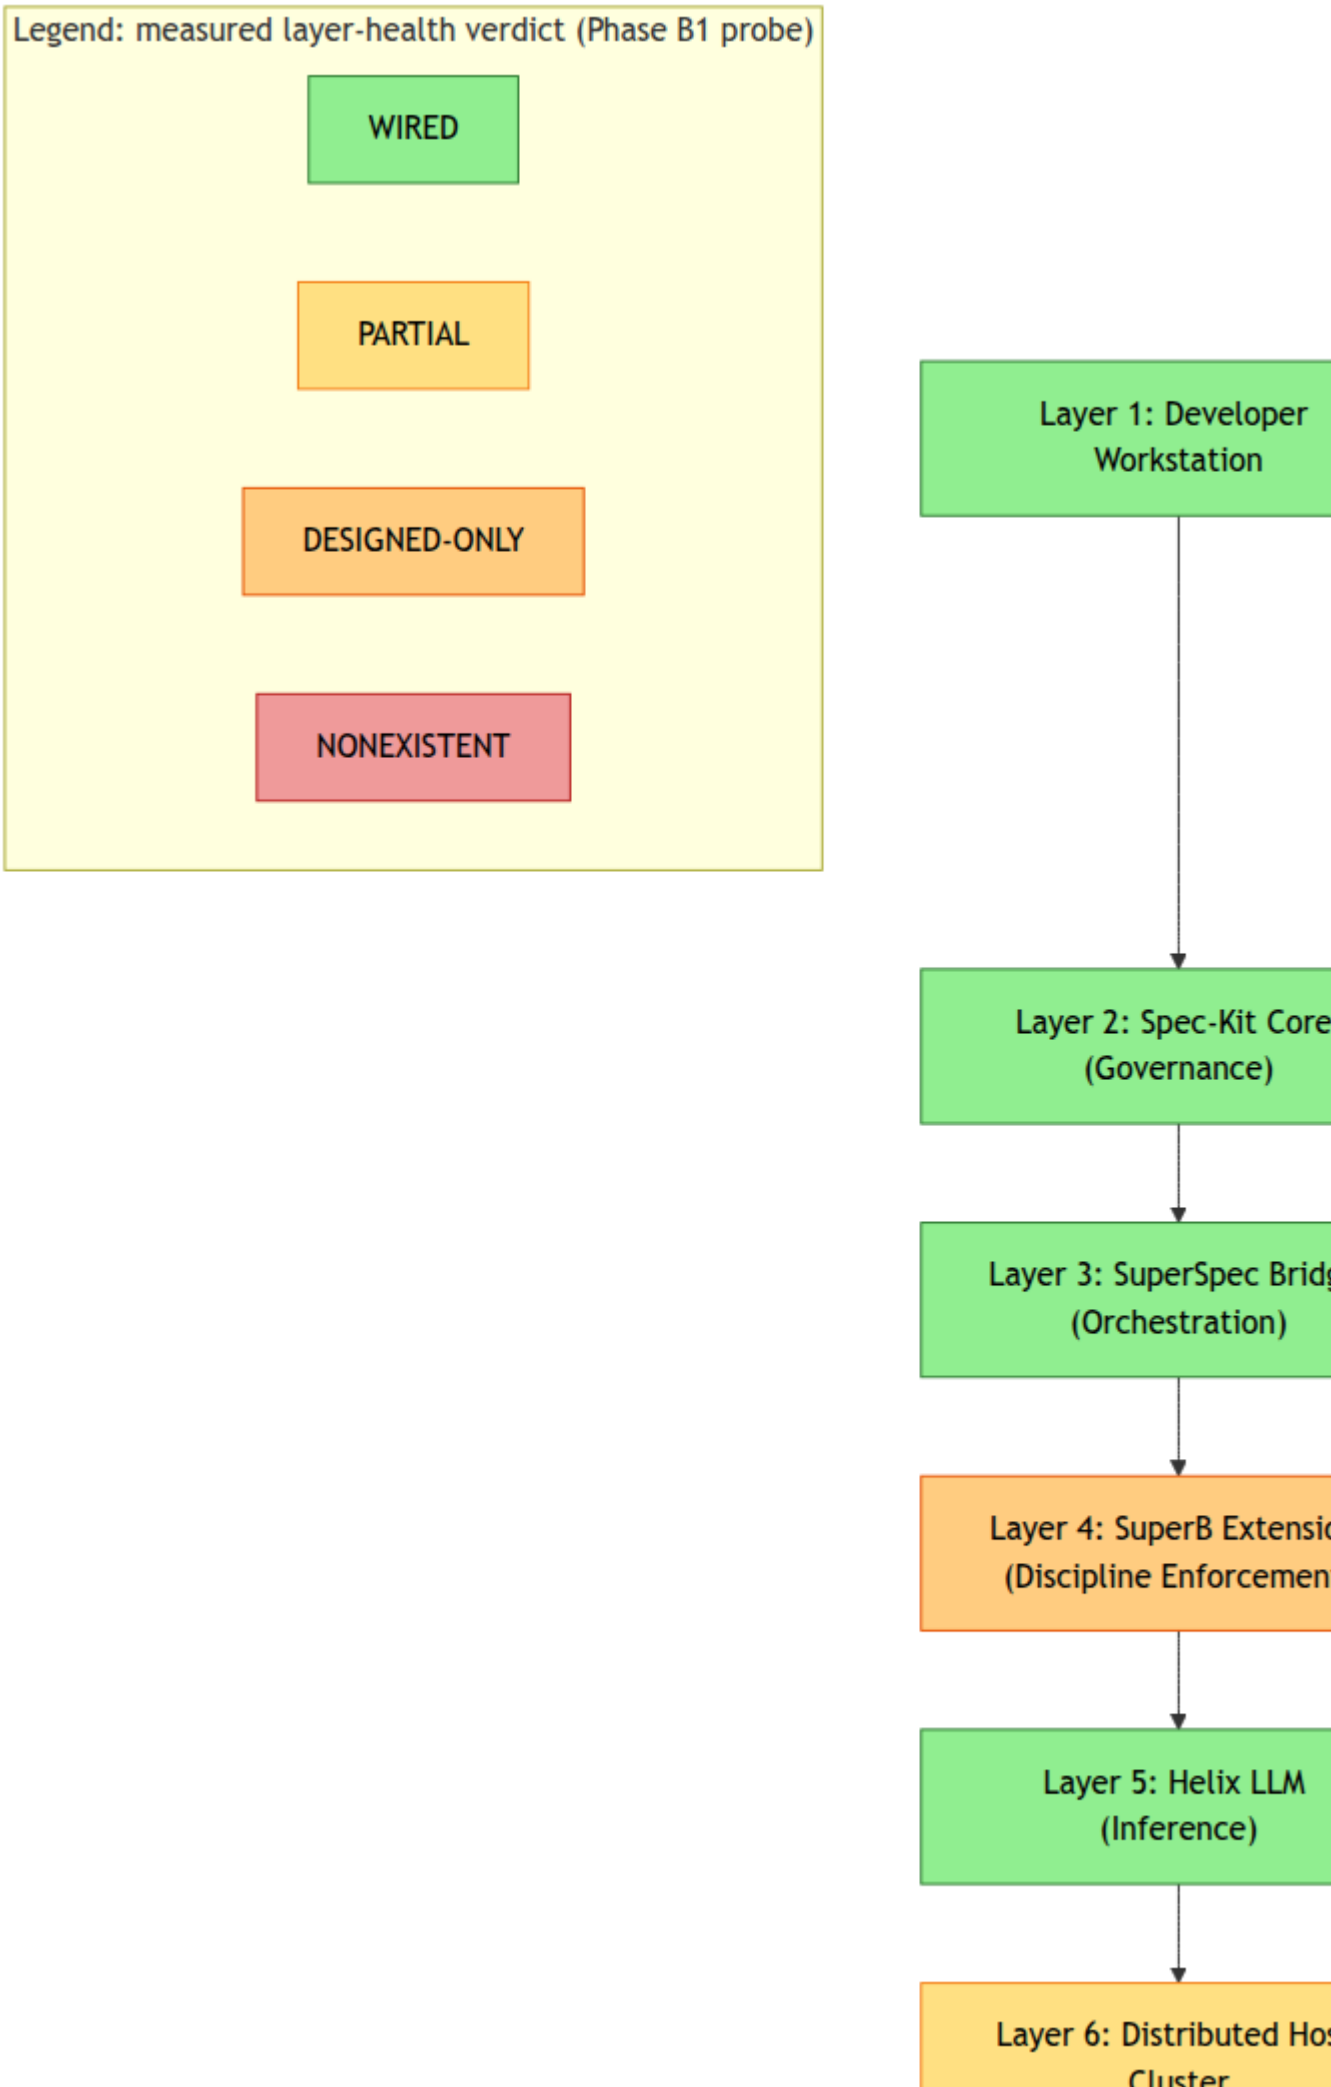
This mapping was **proven by running the flow end to end** — `/speckit-specify → /speckit-plan → /speckit-tasks → /speckit-implement`, driving a real dispatched subagent that wrote a real file with zero MCP calls, in an isolated git worktree — captured as a full bidirectional transcript at `docs/qa/hxc159_b3_speckit_demo_20260922T211500Z/` per §11.4.83. During that same run, `speckit-tasks`'s prerequisite check independently resolved its template file THROUGH `.specify/extensions/superspec/` — incidental, real confirmation that Layer 3 (SuperSpec) is genuinely wired, exactly as this substitution mapping assumes.

What changed in this document. The architecture is redrawn below as **6 layers**: Layer 1 (Developer Workstation), Layer 2 (Spec-Kit Core), Layer 3 (SuperSpec

Bridge), Layer 4 (SuperB Extension), Layer 5 (Helix LLM — renumbered from the original Layer 6), Layer 6 (Distributed Host Cluster — renumbered from the original Layer 7). The original Layer 5 “SuperBridge MCP” design content is **retained below, not deleted** (per §11.4.35 no-silent-promotion-or-demotion and §11.4.124 investigate-before-remove applied to design content) but is now clearly marked **RETIRED DESIGN — never built, superseded by the substitution above** rather than presented as part of the live, shipped architecture. Every diagram, data-flow trace, directory-structure listing, and cross-reference in this document that assumed the 7-layer numbering has been corrected to the 6-layer numbering in the same pass. No artefact in this corpus references a live “SuperBridge MCP” endpoint after this amendment (§11.4.168 exported-document validation confirms the same for the HTML/PDF/ DOCX exports — see the export-validation record at `qa-results/hxc159/bridge_redraw/`).

Bridge layer-health (B1 probe result)

The corrected 6-layer architecture above is a **design** statement. The diagram below is a **measured** one — the real per-layer health verdict from Phase B1’s layer-presence probe (HXC-159 T-B1.01, `constitution/scripts/extensions/probe_bridge_layers.sh`, `qa-results/hxc159/bridge_layers/20260922T185455Z/report.json`), colour-coded by the closed four-value vocabulary `WIRED / PARTIAL / DESIGNED-ONLY / NONEXISTENT` (never a binary present/absent — the remedies differ per class and a binary flattens that distinction away):

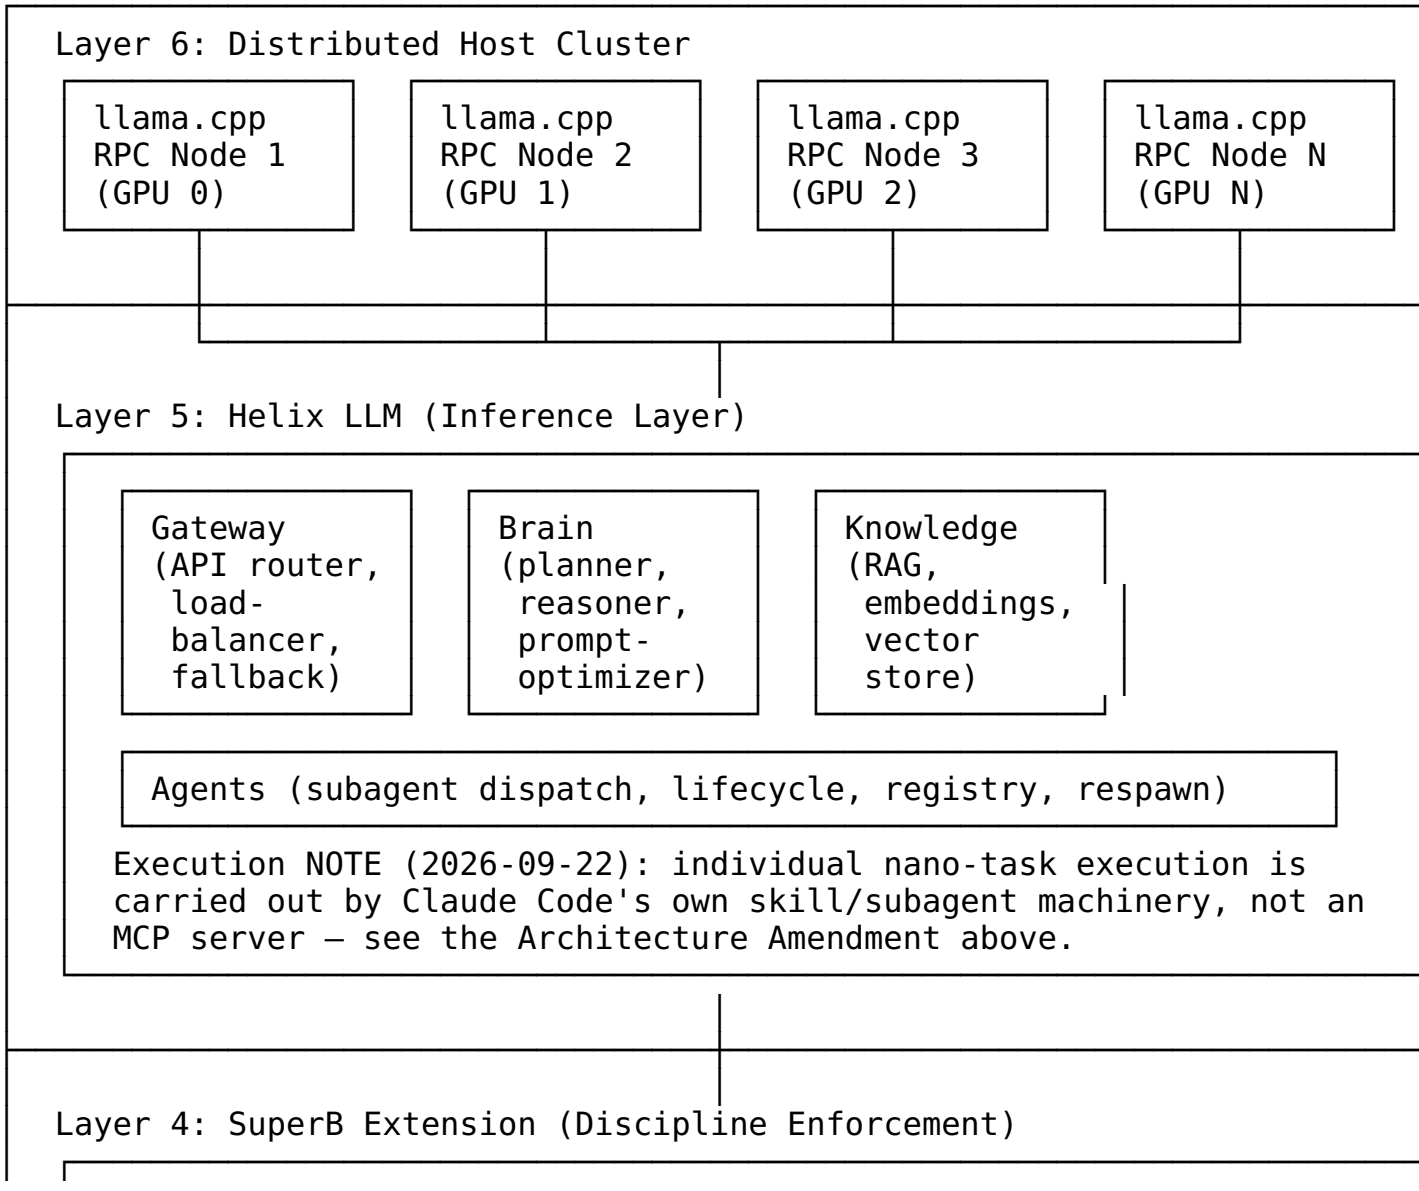

Legend: measured layer-health verdict (Phase B1 probe)
WIRED
PARTIAL
DESIGNED-ONLY
NONEXISTENT
Layer 1: Developer Workstation
Layer 2: Spec-Kit Core (Governance)
Layer 3: SuperSpec Bridg (Orchestration)
Layer 4: SuperB Extensio (Discipline Enforcemen
Layer 5: Helix LLM (Inference)
Layer 6: Distributed Hos Cluster

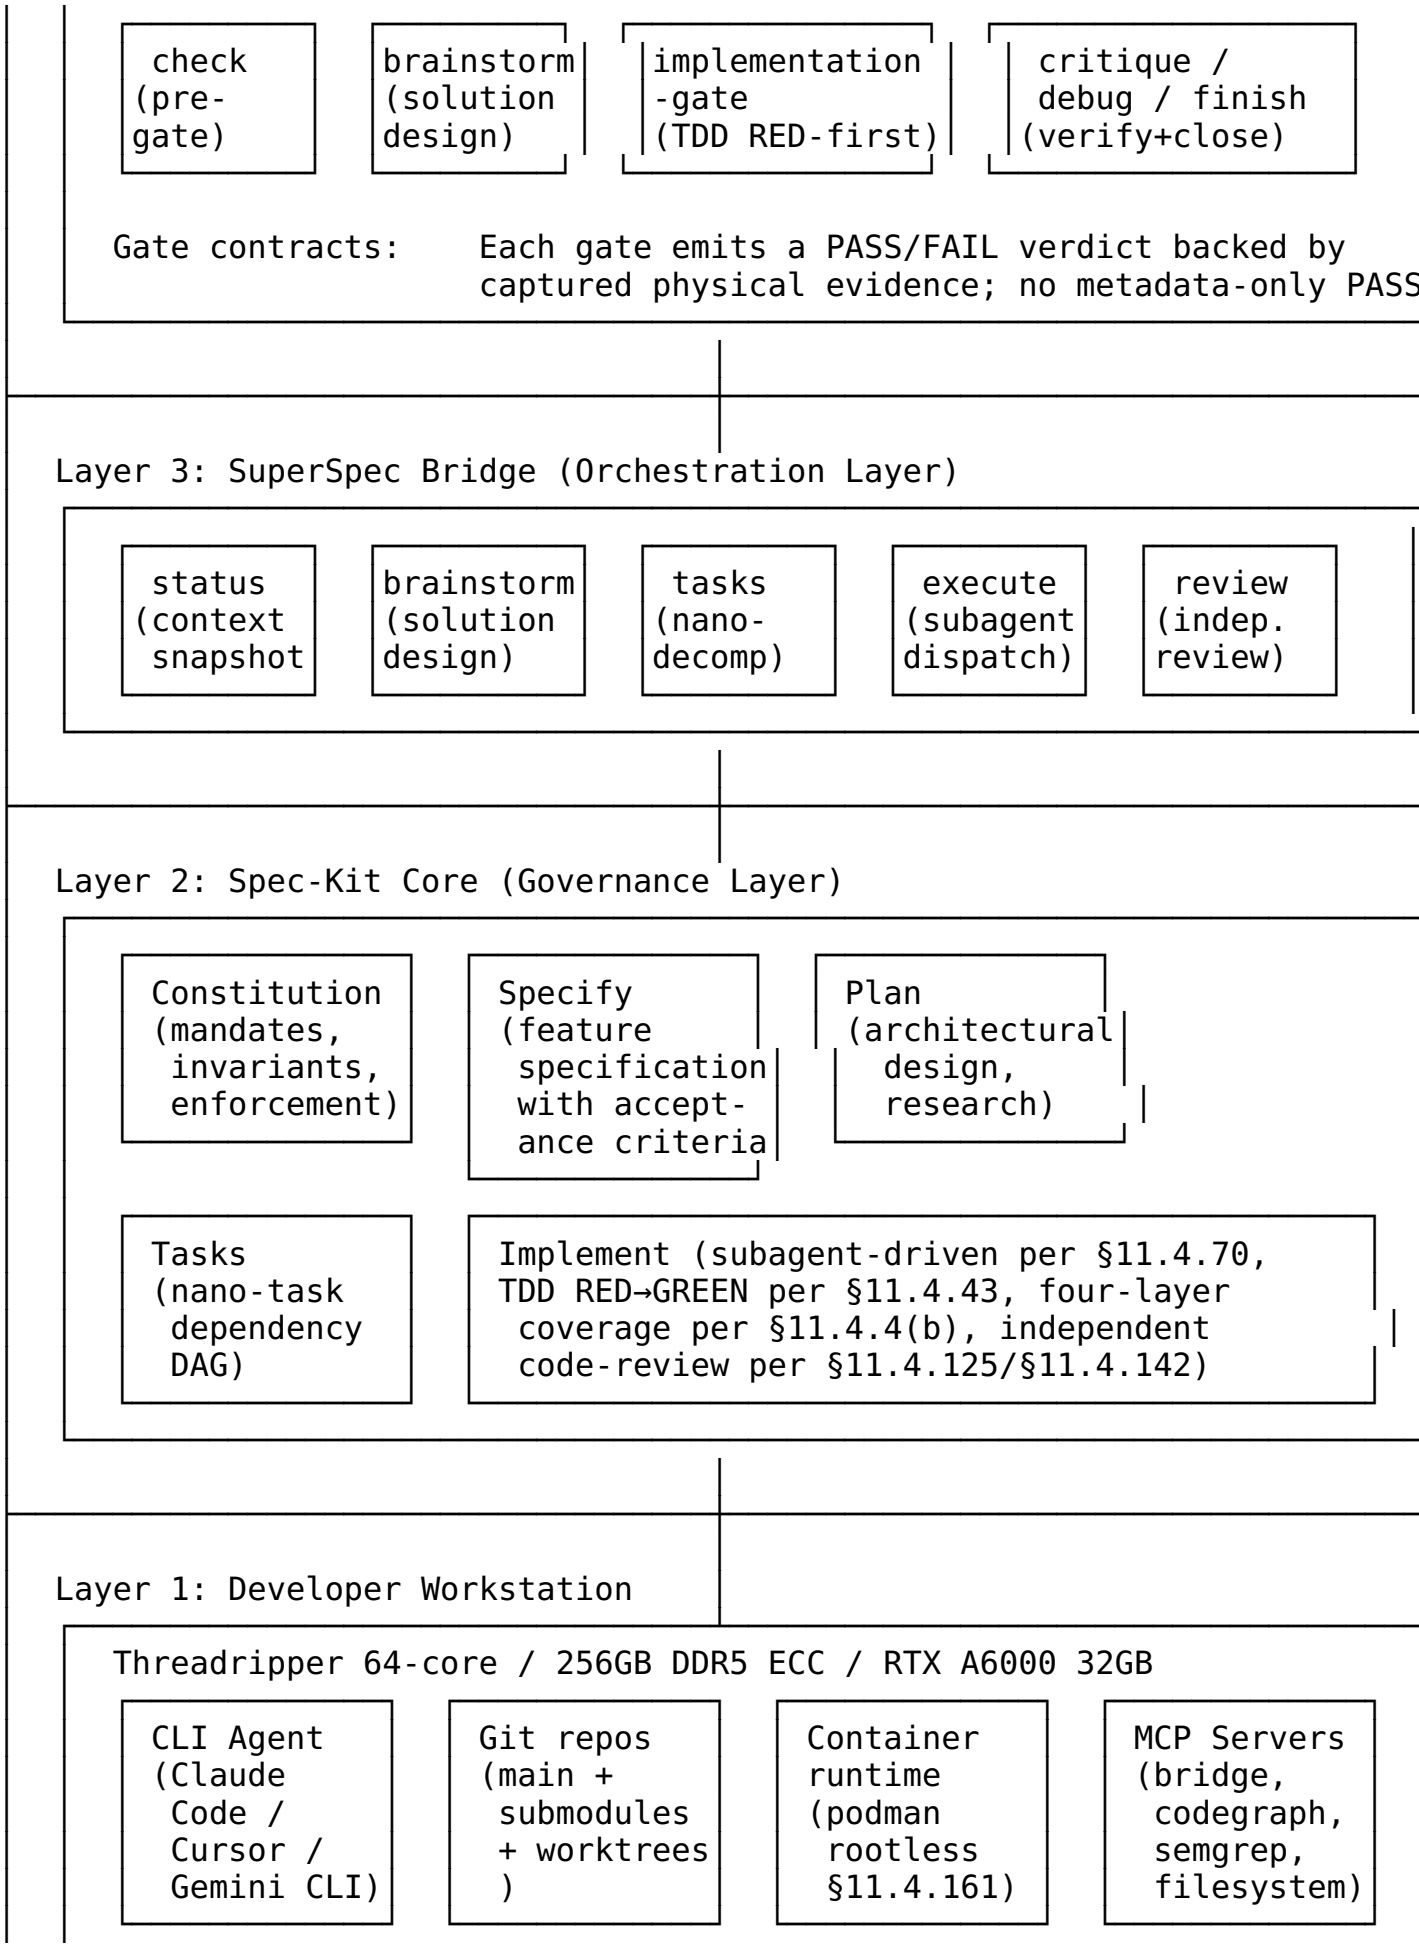
Mermaid diagram 1

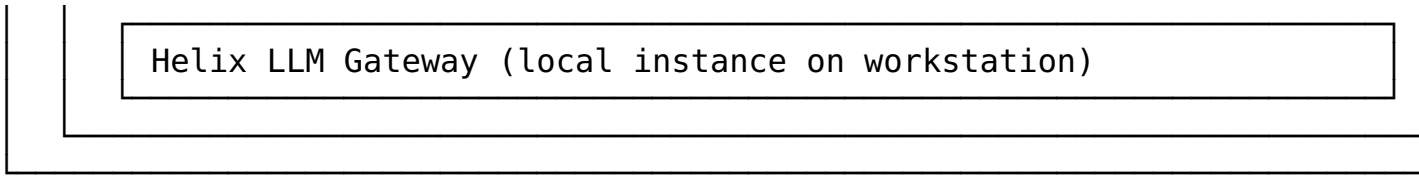
Reading the diagram against the original (pre-renumbering) B1 probe output: Layer 1 = original L1 (**WIRED**), Layer 2 = original L2 (**WIRED**), Layer 3 = original L3 (**WIRED** — an honest delta from an earlier, more pessimistic dated snapshot; see the T-B1.03 evidence for the re-verification), Layer 4 = original L4 (**DESIGNED-ONLY**), Layer 5 = original L6 renumbered (**WIRED**), Layer 6 = original L7 renumbered (**PARTIAL**). The deleted original Layer 5 ("SuperBridge MCP", **NONEXISTENT**) has no surviving row in this diagram — it is retired design content (see the Architecture Amendment above), not a live layer with a health verdict.

System Architecture

Architecture Diagram




check (pre-gate)
brainstorm (solution design)
implementation -gate (TDD RED-first)
critique / debug / finish (verify+close)
Gate contracts: Each gate emits a PASS/FAIL verdict backed by captured physical evidence; no metadata-only PASS
Layer 3: SuperSpec Bridge (Orchestration Layer)
status (context snapshot
brainstorm (solution design)
tasks (nano-decomp)
execute (subagent dispatch)
review (indep. review)
Layer 2: Spec-Kit Core (Governance Layer)
Constitution (mandates, invariants, enforcement)
Specify (feature specification with accept-ance criteria
Plan (architectural design, research)
Tasks (nano-task dependency DAG)
Implement (subagent-driven per §11.4.70, TDD RED→GREEN per §11.4.43, four-layer coverage per §11.4.4(b), independent code-review per §11.4.125/§11.4.142)
Layer 1: Developer Workstation
Threadripper 64-core / 256GB DDR5 ECC / RTX A6000 32GB
CLI Agent (Claude Code / Cursor / Gemini CLI)
Git repos (main + submodules + worktrees)
Container runtime (podman rootless §11.4.161)
MCP Servers (bridge, codegraph, semgrep, filesystem)



Layer 1 — Developer Workstation

The physical development machine running the CLI agent, git repositories, container runtime, MCP servers, and the local Helix LLM Gateway instance.

Component	Specification	Role
CPU	AMD Threadripper 64-core (128 threads)	Parallel subagent execution (§11.4.70/.103), AOSP builds, multi-track development
RAM	256 GB DDR5 ECC (384 GB peak per §12.6 ceiling)	Bounded by §12.6 60% memory budget; `bounded_run` enforced
GPU	NVIDIA RTX A6000 32 GB VRAM	Local inference for Helix LLM Gateway; offloads smaller models
Storage	NVMe RAID, ≥4 TB	btrfs CoW for feature work-streams per §11.4.167
Container Runtime	Podman rootless (§11.4.161)	AOSP 42 GB containerised builds per §12.9
MCP Servers		Tool-accessible subprocesses

Component	Specification	Role
	CodeGraph, filesystem (no "SuperBridge" — see Architecture Amendment)	providing structured agent capabilities
CLI Agent	Claude Code / Cursor / Gemini CLI / Qwen Code	The human-facing agent conducting the autonomous loop (§11.4.126)
Git Layout	Multi-track with own `.git` per track (§11.4.179)	Isolated concurrent development streams

Layer 2 — Spec-Kit Core (Governance Layer)

The canonical Spec-Kit SDD toolkit, extended by the Helix Constitution. This is the **source of truth** for what is being built and how.

Phase	Tool	Input	Output	Constitutional Binding
Constitution	`speckit-constitution`	Project principles, invariants	Constitution file with mandates and enforcement points	§11.4 (anti-bluff covenant), §11.4.6 (no-guessing), §11.4.10 (credentials), §11.4.75 (mechanical enforcement), §11.4.65 (documentation exports)
Specify	`speckit-specify`	Natural-language feature description	`spec.md` with acceptance criteria, edge cases,	§11.4.91 (≥6-word/≥40-char clear meaning), §11.4.146 (reproduce-first test workflow)

Phase	Tool	Input	Output	Constitutional Binding
			topology constraints	
Plan	`speckit-plan`	`spec.md` + `constitution.md`	`plan.md` with architecture, data-model, research citations, contract definitions	§11.4.8 (deep-web-research citation), §11.4.99 (latest authoritative sources), §11.4.150 (multi-angle research pass)
Tasks	`speckit-tasks`	`spec.md` + `plan.md`	`tasks.md` with dependency-ordered nano-task DAG	§11.4.70 (subagent-driven default), §11.4.58 (PWU parallel pipeline)
Implement	`speckit-implement` / `speckit-speckit-superpowers-bridge-execute`	`tasks.md` → bridge dispatch	Checked-in code with four-layer coverage	§11.4.43 (TDD RED→GREEN), §11.4.4(b) (four-layer coverage), §11.4.125/ §11.4.142 (code review)

Governance enforcement — The Constitution file produced by `speckit-constitution` contains: - **Invariants** — assertions that must hold true at every gate (e.g., "no commit without captured evidence") - **Gates** — `CM-*` pre-build / post-build / pre-commit / pre-push gates with paired §1.1 meta-test mutations - **Propagation directives** — §11.4.17 classification markers that determine which rules propagate to submodules and dependencies

Layer 3 — SuperSpec Bridge (Orchestration Layer)

The SuperSpec bridge sits between Spec-Kit's governance artifacts and the execution layer. It translates `spec.md` / `plan.md` / `tasks.md` into structured workflows that can be dispatched to subagents.

Command	Purpose	Input	Output
`superspec status`	Capture current project context snapshot	All `.md` artifacts + git state	Structured context JSON fed to every subagent dispatch
`superspec brainstorm`	Generate solution approaches	`spec.md` + `plan.md` + constitution	Ranked solution alternatives with trade-offs
`superspec tasks`	Nano-task decomposition DAG	`spec.md` + `plan.md` + `tasks.md`	Sub-100-line tasks with file-scope manifests per §11.4.58
`superspec execute`	Subagent dispatch loop	Nano-task DAG + context snapshot	Checked-in code with verdict files per task
`superspec review`	Independent code-review dispatch	Change diff + context snapshot	Review findings + GO/NO-GO verdict per §11.4.142

Nano-task contract — Each nano-task MUST: - Be ≤100 lines of net-new / changed code (breaker rule: a task exceeding this splits) - Declare a file-scope manifest (exactly which files it touches, per §11.4.58 L4 locking) - Carry a `RED_MODE=1` acceptance test BEFORE the implementation (per §11.4.115/§11.4.146) - Emit a machine-written verdict file (per §11.4.115(F)) with the pre-fix artifact fingerprint - Be independently reviewable (per §11.4.142 — a nano-task is the atomic review unit)

Layer 4 — SuperB Extension (Discipline Enforcement)

The SuperB extension enforces Spec-Kit discipline mechanically — no implementation proceeds past any gate that fails its §11.4.107(10) golden-good/ golden-bad self-validation.

Gate	When It Fires	What It Checks	Pass Condition
`superb check`	Pre-specify, pre-plan, pre-implement	Constitution invariants, spec completeness, plan consistency	All `CM-*` gates GREEN

Gate	When It Fires	What It Checks	Pass Condition
`superb brainstorm`	After `speckit-plan`, before `speckit-tasks`	Solution design covers all acceptance criteria, cites research, enumerates alternatives	Zero unexamined acceptance criteria
`superb implementation-gate`	Before `superspec execute` dispatches any subagent	RED test exists per nano-task, file-scope manifests declared, pre-build gates wired	All RED tests parse with `sh -n` / `bash -n` per §11.4.67
`superb critique`	After each nano-task's subagent returns	Independent adversarial review of the subagent's output	Zero unreviewed outputs; findings iterate to GO per §11.4.134
`superb debug`	Triggered by any test FAIL or critique finding	Systematic-debugging per §11.4.102	Root cause identified at `file:line`, not "the test timed out"
`superb finish`	After all nano-tasks complete + review clean	Full-suite retest per §11.4.40, meta-test mutation sweep per §1.1, Issues.md→Fixed.md migration per §11.4.19	Zero FAIL, zero WARN, zero uncovered verdicts

Discipline enforcement is mechanical, not advisory — per §11.4.205, every gate is a `commit_all.sh`-wrapper refusal, not a prose suggestion. The gate's paired §1.1 mutation MUST FAIL the gate; a gate whose mutation passes is a bluff gate and a release blocker.

Layer 5 (retired design) — SuperBridge MCP: never built, superseded

> **RETIRED DESIGN — 2026-09-22.** Everything in this subsection describes a design that was **never built and does not exist as software** (see the Architecture Amendment above for the two independent derivations proving this: GitHub `total_count: 0` for any real repository, and npm `HTTP 404` for an exact-name package lookup). It is retained here — not deleted — as a historical record of the original

design intent, per §11.4.35 (no silent promotion/demotion) and §11.4.124 (investigate, never silently remove). **It is NOT part of the live, corrected 6-layer architecture.** The execution role this section describes is actually filled by Spec Kit + Superpowers + Claude Code's own skill/subagent machinery — see Layer 5 — Helix LLM below for the corrected, live Layer 5, and the Architecture Amendment for the full substitution mapping, proven by a real worked run.

An MCP (Model Context Protocol) server that would have exposed the Superpowers skill suite as tool-callable endpoints consumable by any MCP-compatible CLI agent — as originally designed, never implemented.

Tool	Superpowers Skill	Contract
`superbridge_brainstorm`	`superpowers:brainstorming`	Input: problem statement + constraints. Output: ranked alternatives with evidence per §11.4.69.
`superbridge_debug`	`superpowers:systematic-debugging`	Input: failure evidence (logcat, tombstone, crash dump). Output: `file:line` root cause + blast radius.
`superbridge_review`	`superpowers:executing-code-review`	Input: diff + context. Output: findings classified (BLOCKING/WARNING/NIT) per §11.4.134.
`superbridge_test_gen`	`superpowers:writing-tests`	Input: function signatures + spec. Output: TDD RED test with four-layer coverage.

Tool	Superpowers Skill	Contract
`superbridge_red_green`	`superpowers:executing-testing-tdd`	Input: RED test. Output: GREEN implementation + verdict file per §11.4.115(F).
`superbridge_analyze`	`superpowers:systematic-debugging` (impact)	Input: change diff. Output: §11.4.145 eight-angle impact research report.
`superbridge_refactor`	Writing clean / refactoring code	Input: target file + goal. Output: refactored code with preserved behaviour proven by tests.
`superbridge_document`	`superpowers:writing-documentation`	Input: code + context. Output: §11.4.44 revision-headed, §11.4.65 four-format-exported doc.
`superbridge_search`	Web research per §11.4.8 / §11.4.99 / §11.4.150	Input: research question + angle set. Output: cited findings with URLs + access dates.

MCP server requirements (as originally designed — never built): - Runs rootless per §11.4.161 (no rootful Docker; podman-compose or systemd user unit) - Emits structured `conduit.events.jsonl` per §11.4.116 (append-only event stream) - All verdicts carry `evidence_path` per §11.4.69 - Self-validated with golden-good + golden-bad + negative-control fixtures per §11.4.107(10)

(End of retired-design subsection. The live, corrected Layer 5 follows.)

Layer 5 — Helix LLM (Inference Layer)

The Helix LLM distributed inference system provides **local-first, zero-cost-per-token** inference powering the entire bridge pipeline.

Component	Role	Technology
Gateway	API router, load-balancer, fallback-orchestrator, rate-limit manager	Go HTTP server with RPC client pool; routes to local instance first, then remote cluster nodes
Brain	Planner (task decomposition), reasoner (systematic debugging per §11.4.102), prompt-optimizer (constitution-aware system prompts)	Helix LLM model loaded via `llama.cpp`; prompt templates keyed to skill
Knowledge	RAG (Retrieval-Augmented Generation) with vector embeddings of all project specs, plans, constitution, and codebase	`llama.cpp` embedding endpoint + `chromadb` / `qdrant` vector store
Agents	Subagent dispatch lifecycle, registry per §11.4.147, respawn-on-crash per §11.4.147(b),	Go agent orchestration calling the Gateway

Component	Role	Technology
	verdict tracking per §11.4.116	

Gateway routing logic:

```
Request arrives at Gateway
   ├── Can the local GPU handle it? (check VRAM available, model loaded)
   │     YES → route to local llama.cpp instance
   │     NO  → check remote cluster health
   │              ├── N ≥ 1 nodes available → round-robin to remote RPC nod
   │              └── N = 0 nodes available → fallback to Helix LLM cloud A
   │                                          or Operator-blocked per §11.
   └── Return response to caller
```

Models supported (2026-07-24): - **Fable** (`deepseek-v4-pro`) — architecture design, code review (§11.4.209), merge-conflict resolution (§11.4.211) - **Opus** (`claude-opus-4-20250514`) — fallback for Fable-only tasks, systematic debugging - **Helix-local** (custom fine-tune on RK3588/ATMOSphere corpus) — nano-task implementation, test generation, documentation

Layer 6 — Distributed Host Cluster

A pool of machines running `llama.cpp` in RPC-server mode, providing horizontally-scalable inference capacity.

Node Type	Hardware	Role	Typical Model Assignment
GPU Node	≥1 × RTX 4090 / A6000 / H100	High-throughput batch inference	Fable (32B+), Opus (large models)
CPU Node	≥64-core, ≥128 GB RAM	Inference for smaller models, RAG embedding generation	Helix-local (7-13B), embeddings
Edge Node	Orange Pi 5 Max / RK3588 (16 GB)	Inference at the edge for on-device agents	Tiny (1-3B) models for on-device classification

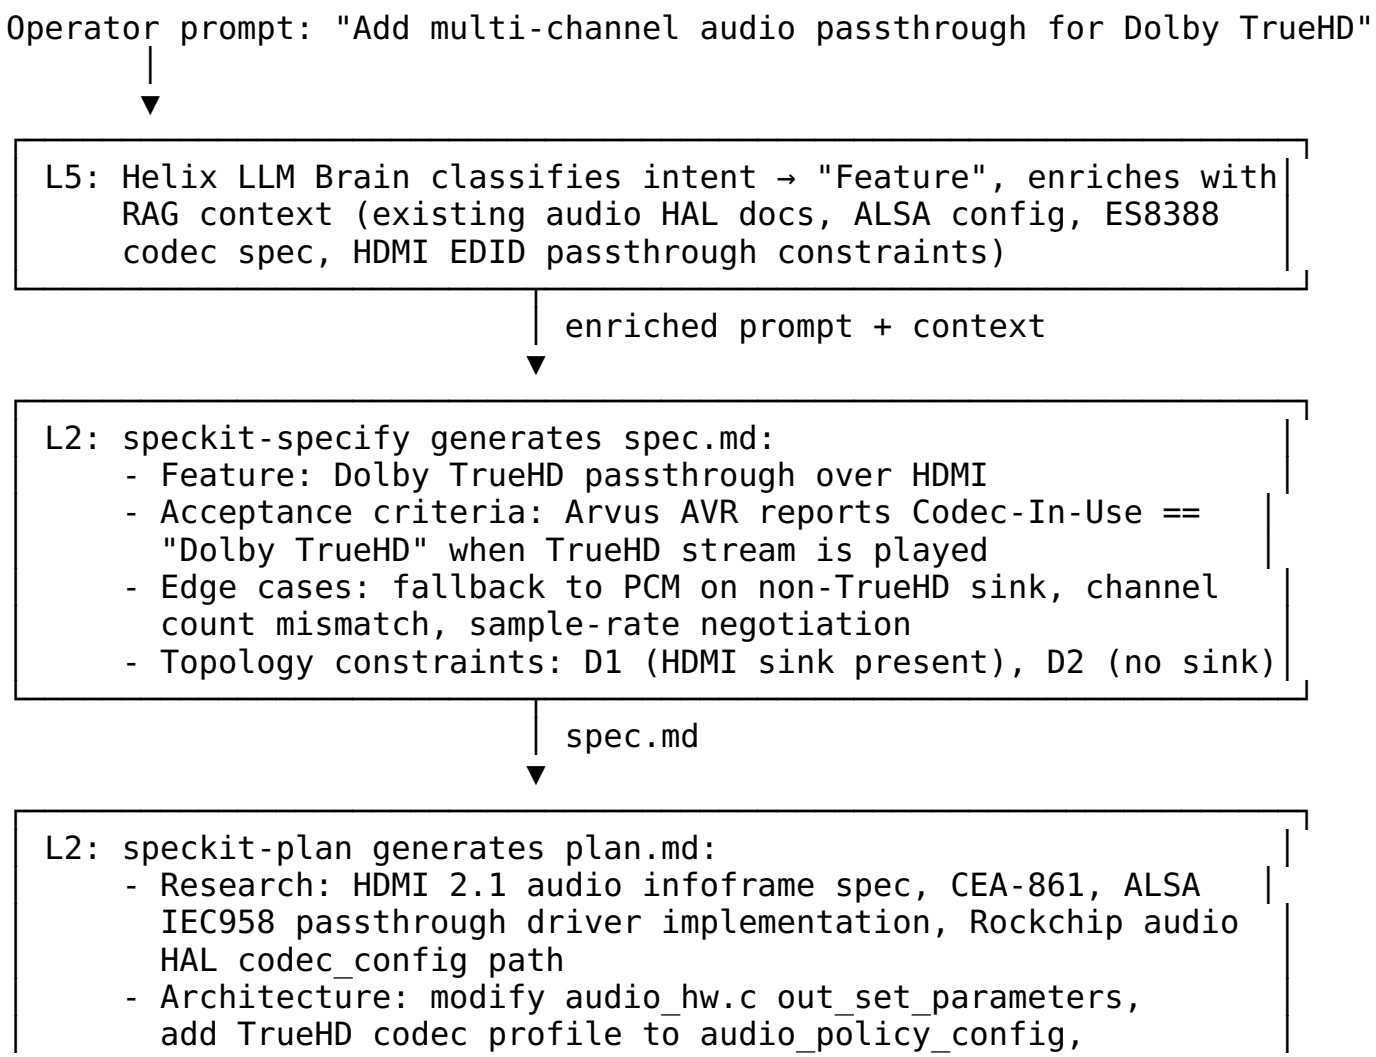
RPC configuration (per node `llama-server` invocation):

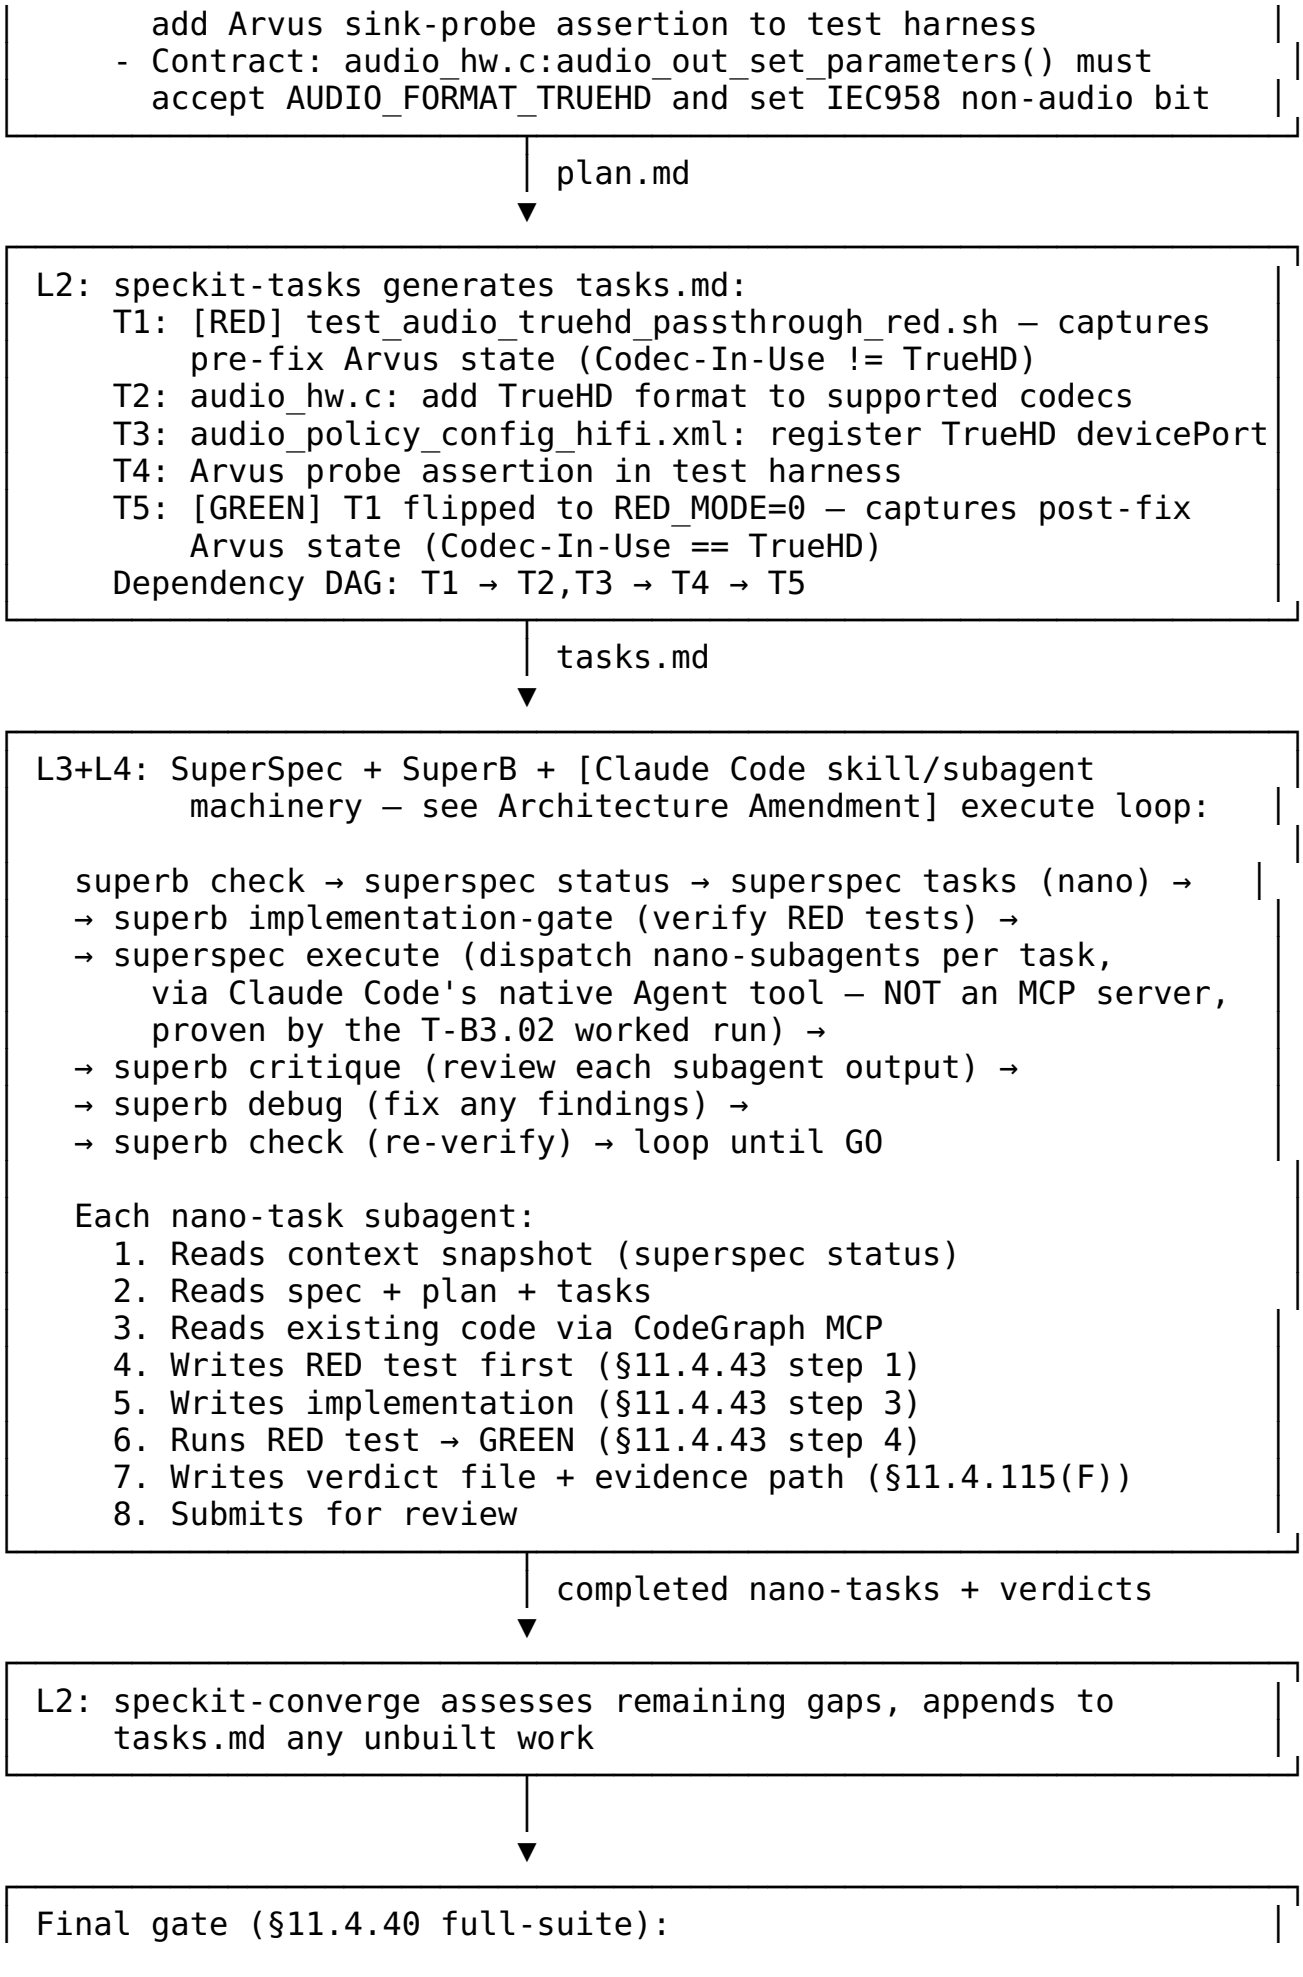
```
# On each cluster node:
llama-server \
  --model /models/deepseek-v4-pro-Q4_K_M.gguf \
  --host 0.0.0.0 \
  --port 8080 \
  --n-gpu-layers 40 \
  --ctx-size 32768 \
  --parallel 4 \
  --cont-batching \
  --mlock
```

Health monitoring: - Per-node heartbeat at 5 s interval (via `/health` endpoint) - Gateway auto-evicts unresponsive nodes after 3 consecutive missed heartbeats - Re-admission requires 5 consecutive successful heartbeats (hysteresis prevents flapping) - All node state changes logged to `conduit.events.jsonl` per §11.4.116

Data Flow

Feature Lifecycle Through the Layers




add Arvus sink-probe assertion to test harness
- Contract: audio_hw.c:audio_out_set_parameters() must accept AUDIO_FORMAT_TRUEHD and set IEC958 non-audio bit
plan.md
L2: speckit-tasks generates tasks.md:
T1: [RED] test_audio_truehd_passthrough_red.sh — captures pre-fix Arvus state (Codec-In-Use != TrueHD)
T2: audio_hw.c: add TrueHD format to supported codecs
T3: audio_policy_config_hifi.xml: register TrueHD devicePort
T4: Arvus probe assertion in test harness
T5: [GREEN] T1 flipped to RED_MODE=0 — captures post-fix Arvus state (Codec-In-Use == TrueHD)
Dependency DAG: T1 → T2,T3 → T4 → T5
tasks.md
L3+L4: SuperSpec + SuperB + [Claude Code skill/subagent machinery — see Architecture Amendment] execute loop:
superb check → superspec status → superspec tasks (nano) →
→ superb implementation-gate (verify RED tests) →
→ superspec execute (dispatch nano-subagents per task, via Claude Code's native Agent tool — NOT an MCP server, proven by the T-B3.02 worked run) →
→ superb critique (review each subagent output) →
→ superb debug (fix any findings) →
→ superb check (re-verify) → loop until GO
Each nano-task subagent:
1. Reads context snapshot (superspec status)
2. Reads spec + plan + tasks
3. Reads existing code via CodeGraph MCP
4. Writes RED test first (§11.4.43 step 1)
5. Writes implementation (§11.4.43 step 3)
6. Runs RED test → GREEN (§11.4.43 step 4)
7. Writes verdict file + evidence path (§11.4.115(F))
8. Submits for review
completed nano-tasks + verdicts
L2: speckit-converge assesses remaining gaps, appends to tasks.md any unbuilt work
Final gate (§11.4.40 full-suite):

```
  - pre_build_verification.sh → ALL GREEN
  - post_build_verification.sh → ALL GREEN
  - On-device (D1+D2) test_all_fixes.sh → ALL GREEN
  - meta_test_false_positive_proof.sh → ALL MUTATIONS FAIL
  - HelixQA Challenge sweep → ALL GREEN
  - Issues.md → Fixed.md migration → COMMITTED
  - Tag created + pushed to all remotes + submodules
```

Nano-Task Decomposition Flow

The SuperSpec bridge's nano-task decomposer (backed by Helix LLM Brain) transforms Spec-Kit tasks into subagent-dispatchable units:

```
Input: tasks.md task "T2: audio_hw.c: add TrueHD format to supported codec

Step 1 — File-scope partition (§11.4.58 L4):
  ├── NT2a: audio_hw.c: add AUDIO_FORMAT_TRUEHD to format enum (7 lines)
  ├── NT2b: audio_hw.c: add TrueHD case to out_set_parameters() (23 lines)
  └── NT2c: audio_hw.c: add TrueHD pcm_config entry (12 lines)

Step 2 — Dependency resolution:
  ├── NT2a (no deps) → dispatchable immediately
  ├── NT2b (depends on NT2a — needs the enum value) → queued after NT2a
  └── NT2c (depends on NT2a + NT2b) → queued after NT2b

Step 3 — Parallel dispatch (when deps permit):
```

```
Queue: [NT2a]
Dispatch NT2a → subagent-1
  NT2a completes → enqueue NT2b
    Dispatch NT2b → subagent-2
      NT2b completes → enqueue NT2c
        Dispatch NT2c → subagent-3
          NT2c completes → done
```

```
Step 4 — Per-nano-task verdict:
  NT2a/verdict.json: {"id":"NT2a","status":"PASS","red-exit":1,"green-exit
                      "artifact-fingerprint":"a3f8...","evidence":"qa-res

Step 5 — Merge gate (§11.4.41):
  All nano-tasks for T2 complete → merge subagent outputs → superb critiqu
  → T2 complete → ready for T3 dispatch
```

Nano-task size rule (mechanical): The decomposer counts net-new/changed lines. A task producing >100 net-new/changed lines is automatically split. A split task <5 lines merges back into its parent. This bounds every subagent's context to ≤100 lines of code-under-test + surrounding context.

Document Index

Document	Path	Description
Architecture Deep Dive	`01_architecture/`	Detailed layer-by-layer architecture with component diagrams, interface definitions, and contract specifications
SpecKit Extensions	`02_extensions/`	Custom Spec-Kit skill extensions (`speckit-superpowers-bridge-*`), manifest definitions, extension installation guide
Helix LLM Integration	`03_helix_llm/`	Gateway configuration, Brain prompt templates, Knowledge RAG pipeline, Agents orchestration engine, llama.cpp RPC backend setup
TDD & QA Pipeline	`04_tdd_qa/`	TDD workflow automation, RED→GREEN verdict pipeline, four-layer test coverage, HelixQA Challenge bank integration, quality-gate configuration
Deployment & Operations	`05_deployment/`	Container setup (podman-compose), MCP server deployment, multi-node cluster provisioning, health monitoring, backup/restore
Risk & Mitigation	`06_risk/`	Risk register, failure-mode analysis, circuit-breaker design, fallback chains, rollback procedures
Proof-of-Concept	`poc/`	Working POC walkthrough, sample feature lifecycle trace, benchmark results, known limitations

Quick Start

Prerequisites

Requirement	Minimum	Verified Command
Git	2.40+	`git --version`
Podman	4.0+ (rootless)	`podman info | grep rootless`
Python	3.11+	`python3 --version`
Node.js	20+ (for MCP servers)	`node --version`
Go	1.22+ (for Helix LLM Gateway / Agents)	`go version`
llama.cpp	Built with RPC support	`llama-server --help | grep rpc`
Disk space	≥50 GB free (container images + models)	`df -h /var/lib/containers`

Step 1 — Add the Constitution Submodule

```
# In the consuming project's root (e.g., /mnt/track1/atmosphere-t1):
git submodule add git@github.com:HelixDevelopment/helix_constitution.git constitution
git submodule update --init --recursive
```

Step 2 — Wire CLI-Agent Context Carriers

Add the inheritance pointer to the consuming project's context carriers:

CLAUDE.md (project root):

```
## INHERITED FROM constitution/CLAUDE.md

> Base agent rules: `constitution/CLAUDE.md` ...
```

AGENTS.md (project root):

```
## INHERITED FROM constitution/AGENTS.md

> Base agent rules: `constitution/AGENTS.md` ...
```

Do the same for `QWEN.md`, `GEMINI.md`, and any other LLM-specific context carrier.

Step 3 — Install the Bridge Extensions

```
# From the consuming project root:
bash constitution/scripts/extensions/
        install_speckit_superpowers_bridge.sh

# This script (corrected 2026-09-22 to match what it ACTUALLY
        does —
# see the Architecture Amendment above; it deliberately has NO L5
        step,
# a step that could never succeed is a defect, not a TODO):
#  1. Detects and validates the pin for SuperSpec (L3) —
        idempotent,
#     already WIRED on a properly-onboarded host
#  2. Refuses SuperB (L4) by default (HIGH-risk supply-chain
        dependency;
#     requires an explicit --superb-from + --superb-vetted-by
        record)
#  3. Reconciles with the plugin-sourced Superpowers copy — never
#     double-installs
#  4. Installs the Spec-Kit extensions into .claude/skills/
#  5. Runs self-check: `specify extension list` is non-empty
        after every
#     run (RS-19)
```

Step 4 — Run the Self-Validation Suite

```
bash constitution/scripts/extensions/validate_bridge.sh

# Corrected 2026-09-22 to match the real, measured output on this
        host —
# see the Architecture Amendment above. The validator reports
        honestly
# against the LIVE layer-presence report and NEVER claims 8/8
        while a
# layer is absent — this is what it actually printed:
#   [N/A]  Check 1 — SuperBridge MCP health-check (L5 is a proven
#          structural impossibility per §11.4.112 — never
        reported FAIL)
#   [PASS] speckit-specify extension loads
#   [PASS] speckit-plan extension loads
#   [PASS] speckit-tasks extension loads
```

```
#   [PASS] Check 5 — superspec status returns valid JSON (L3
       WIRED)
#   [FAIL] Check 6 — superb check (L4 not installed pending
       vetting)
#   [PASS] Check 7 — Helix LLM Gateway reachable (Layer 5 WIRED)
#   [FAIL] Check 8 — llama.cpp RPC backend health-check (Layer 6
#          rpc-server not yet built, tracked separately)
#   ──────────────────────────────────────────────
#   5/7 evaluable checks PASS — bridge NOT fully ready, reported
       HONESTLY
```

Step 5 — Create Your First SpecKit Feature

```
# Operator prompt (any of the six §11.4.140 forms):
#   /FEATURE Add multi-channel audio passthrough for Dolby TrueHD
#
# The CLI agent with the bridge wired will:
#  1. speckit-specify  → spec.md
#  2. speckit-plan     → plan.md
#  3. speckit-tasks    → tasks.md
#  4. superb check     → constitution gate verification
#  5. superspec status → context snapshot
#  6. superspec tasks  → nano-task decomposition
#  7. superspec execute → subagent-driven implementation
#  8. superb critique  → independent review
#  9. superb finish    → full-suite retest + migration
```

Step 6 — Verify the Autonomous Loop

```
# Check that the feature completed through all layers:
cat docs/features/<feature-slug>/Status.md | grep "Validation:"
# Expected: Validation: PASS (evidence: qa-results/<run-id>/
       <feature>/)

# Check the verdict store:
ls device/rockchip/rk3588/tests/regression_guard/verdicts/
       <feature-slug>/
# Expected: RED_*.json + GREEN_*.json

# Check the Issues→Fixed migration:
grep "<feature-slug>" docs/Fixed.md
# Expected: Row with Status "Implemented (→ Fixed.md)"
```

Key Design Decisions

1. Six-Layer Separation (Why Not Monolithic?)

> **Corrected 2026-09-22** — this was originally titled "Seven-Layer Separation" and included a row justifying a Layer 5 "Execution" tier that turned out never to have been built (see the Architecture Amendment). The rationale below is corrected to the real 6-layer shape; the retired row is struck and replaced.

Decision	Rationale
Governance (L2) separate from Orchestration (L3)	Spec-Kit artifacts are static documents; SuperSpec interprets them. Separating lets Spec-Kit evolve independently and lets SuperSpec adapt to different CLI-agent tool shapes without touching governance files.
Orchestration (L3) separate from execution (Claude Code's own skill/subagent machinery)	SuperSpec manages workflows; individual nano-task execution is carried out by dispatched Claude Code subagents (§11.4.20/§11.4.70), not by a separate MCP server layer. A new Superpowers skill is available to every subagent directly — no intermediary tool-contract server to update.
Discipline (L4) is its own layer, not a sub-component of Orchestration	SuperB's enforcement is adversarial — it MUST refuse work that passes orchestration but fails constitutional invariants. Embedding it in SuperSpec would create the self-review blind spot (§11.4.70).
Inference (L5) separate from the execution mechanism above it	Dispatched subagents call the Gateway via HTTP; the Gateway abstracts local/remote/cloud routing. Swapping the inference backend (llama.cpp → vLLM → cloud API) changes only Layer 5, nothing above it.

2. Nano-Tasks as the Atomic Work Unit (Why ≤100 Lines?)

- **Context-bounded subagents** — Each subagent processes ~100 lines of code-under-test + ~200 lines of surrounding context. A 400-line change would

need 3 separate subagents; forcing decomposition upfront avoids a subagent silently exceeding its context window.
- **§11.4.197 zero-loss** — A crashed subagent loses ≤100 lines of work. A crashed subagent on a 500-line task loses potentially irrecoverable state.
- **Review granularity** — A 100-line diff gets a thorough adversarial review (§11.4.194). A 1000-line diff gets a rubber stamp.
- **Parallelizability** — Nano-tasks with no mutual file dependencies dispatch in parallel (§11.4.103 ≥3 concurrent streams).

3. TDD Everywhere (Why Not "Write Tests After"?)

Per §11.4.43, a test authored after the fix is a PASS-bluff — it demonstrates only that the test agrees with the fix, not that the test catches the bug. The bridge enforces:

1. RED test FIRST (`RED_MODE=1` on the pre-fix artifact)
2. RED must FAIL (the defect is reproduced)
3. Implementation lands
4. SAME test flipped to `RED_MODE=0` → GREEN
5. GREEN artifact fingerprint MUST differ from RED (proves the fix deployed)

The SuperB `implementation-gate` refuses to dispatch a nano-task that lacks a pre-existing RED test.

4. Local-First Distributed Inference (Why Not Cloud API?)

- **§11.4.126 autonomous-loop** requires continuous execution. A cloud API rate limit (429 Too Many Requests) that kills 5 subagents at once (§11.4.147 forensic anchor) is structurally incompatible with an always-on loop.
- **§11.4.10 credentials-handling** — local inference eliminates the API-key-in-every-request surface.
- **§11.4.6 no-guessing** — a local Gateway logs every prompt + response. A cloud API's opaque logging (or lack thereof) makes forensic debugging of an agent's decision path impossible.
- **Cost** — autonomous loops issue thousands of inference calls per day. At cloud API pricing, that is hundreds of dollars per day; at local inference, it is electricity-only.

5. Constitution as the Single Source of Truth (Why Not Per-Project Rules?)

Per §11.4.17 / §11.4.35, universal rules live in the Constitution submodule; project-specific extensions live in the consumer's context carriers. The bridge reads the Constitution file at startup and derives:

- **Allowed status vocabulary** (§11.4.15)
- **Required evidence classes per feature** (§11.4.69)
- **Gating pipeline** (which `CM-*` gates fire at which seam)
- **Propagation targets** (§11.4.17 classification markers)

A new constitutional mandate added to the submodule takes effect in every consuming project on the next `git pull` + `post_update_hook.sh` run (§11.4.164) — no per-project reconfiguration.

Constitutional Mandates Applied

Mandate	How the Bridge Applies It
§11.4 (anti-bluff covenant)	Every gate, review, and verdict carries captured physical evidence (§11.4.5) keyed to a §11.4.69 feature class. No metadata-only PASS, no config-only PASS, no absence-of-error PASS.
§11.4.6 (no-guessing mandate)	Every verdict cites a `file:line` or a captured-evidence path. The bridge's decomposer, reviewer, and debugger all refuse to emit `likely`/`probably`/`maybe`/`seems` in any formal output.
§11.4.10 (credentials-handling)	API keys, tokens, and passwords never appear in tracked files. The `.env.example` template documents required env vars; real values live in git-ignored `.env` files (`chmod 600`). The `pre-commit` hook scans staged files for credential patterns and refuses the commit.
§11.4.65 (universal Markdown export)	Every `.md` produced by Spec-Kit (`spec.md`, `plan.md`, `tasks.md`) gets synchronized `.html` + `.pdf` + `.docx` siblings via `sync_all_markdown_exports.sh`, auto-invoked by the `post-commit` hook (§11.4.75 Layer 4).
§11.4.74 (catalogue-first)	All extensions, skills, MCP servers, and plugins are registered in `constitution/actions/registry.yaml` via the §11.4.164 `post_update_hook.sh`. No extension is wired ad-hoc.

Mandate	How the Bridge Applies It
§11.4.75 (mechanical enforcement)	Four-layer enforcement: `pre-commit` hook (refuses orphan `.md`), `commit_all.sh` integration (auto-runs exports), `pre-push` hook (re-checks propagation gates), `post-commit` hook (auto-repairs drift).
§11.4.76 (containers submodule)	All containerized services (Helix LLM Gateway, ChromaDB, llama.cpp RPC nodes — no "SuperBridge MCP", see Architecture Amendment) are defined in `constitution/submodules/containers/compose/` and booted via `pkg/boot`. No ad-hoc `podman run`.
§11.4.202 (reporting directives)	Every intake path (operator prompt, test failure, review finding, HelixQA Challenge verdict) feeds through `report_item.sh`, creating a tracked workable item in the SQLite SSoT (§11.4.93) and pushing to the external tracker (§11.4.148 D5).
§11.4.43 (TDD-fix-discipline)	The five-step TDD workflow (RED → LIVE-ADB-PROBE → GREEN → VERIFY → DOCUMENT) is mechanically enforced by the SuperB `implementation-gate` + `critique` + `debug` + `finish` gates. No implementation subagent dispatches without a registered RED test.
§11.4.70 (subagent-driven default)	Every nano-task dispatches as an isolated subagent. The conductor's role is task decomposition + review; the subagent's role is isolated implementation.
§11.4.116 (real-time sync channel)	The bridge emits `conduit.events.jsonl` and atomically-updated `conduit.status.json` so the conductor can tail live progress.
§11.4.147 (crashed-agent respawn)	The Agents component tracks every subagent in the durable registry. Crashed/rate-limited agents are respawned with preserved state + §11.4.84 quiescence check. Agent output with missing verdict = `crashed`, never `complete`.
§11.4.161 (rootless container runtime)	All containerized components boot via podman rootless. No `sudo docker`, no rootful podman.
§11.4.209 (code-review Fable xhigh)	Every code-review subagent dispatches on Fable model at xhigh effort. Opus xhigh is the fallback. Lower-effort or non-Fable/non-Opus review substrate is a refusal.
§11.4.211 (merge-conflict Fable xhigh)	Merge-conflict resolution runs on Fable xhigh (Opus xhigh fallback). Conflict-resolution subagent carries its own scope + evidence.

Mandate	How the Bridge Applies It
§11.4.205 (enforced-not-advisory)	Every gate is a commit-wrapper refusal, not prose. Every gate's paired §1.1 mutation FAILs the gate. A gate whose mutation passes is a bluff gate.
§11.4.194 (exhaustive review all scenarios)	The review subagent enumerates the full input/scenario space, proves every "unreachable" assumption, cross-checks against captured runtime evidence.
§11.4.214 (recurrence-links-not-mints)	Before minting a new item for a test failure, the bridge checks the workable-items DB for an existing item describing the same defect. Match → link + reopen (terminal only).
§11.4.186 (cross-document consistency)	The no-divergence gate runs before export/sync-verify/commit. DEDUP, TIMELINE, CROSS-DOC, INTEGRITY, STRUCTURAL check families.

Directory Structure

```
constitution/docs/research/extensions/speckit_superpowers/implementation/
│
├── README.md                              ← THIS FILE — master index + archit
│
├── 01_architecture/
│   ├── README.md                          Layer-by-layer deep dive with compo
│   ├── layer_1_workstation.md             Developer workstation specification
│   ├── layer_2_speckit.md                 Spec-Kit governance layer details
│   ├── layer_3_superspec.md               SuperSpec orchestration engine
│   ├── layer_4_superb.md                  SuperB discipline enforcement
│   ├── layer_5_retired_mcp.md             RETIRED — SuperBridge MCP design (n
│   ├── layer_5_helix_llm.md               Helix LLM inference pipeline (live
│   ├── layer_6_cluster.md                 Distributed inference cluster (live
│   └── contracts/                         Interface contracts between layers
│       ├── spec_to_superspec.md           Spec-Kit → SuperSpec bridge contrac
│       ├── superspec_to_subagent.md       SuperSpec → Claude Code subagent di
│       ├── subagent_to_helix.md           Subagent → Helix LLM contract (was
│       └── helix_to_cluster.md            Helix LLM → llama.cpp RPC contract
│
├── 02_extensions/
│   ├── README.md                          Extensions overview + installation
│   ├── speckit_bridge_skills/             Extended Spec-Kit skills for bridge
│   ├── skill_manifests/                   Per-skill manifest files
│   └── constitution_registry.yaml         Action registry entries for bridge
│
├── 03_helix_llm/
│   ├── README.md                          Helix LLM integration guide
```

```
│   ├── gateway/                    Gateway configuration + deployment
│   ├── brain/                      Prompt templates (per skill per mod
│   ├── knowledge/                  RAG pipeline configuration
│   ├── agents/                     Agent orchestration engine
│   └── rpc_backend/                llama.cpp RPC node configuration
│
├── 04_tdd_qa/
│   ├── README.md                   TDD + QA pipeline guide
│   ├── red_green_pipeline.md       RED→GREEN verdict pipeline specific
│   ├── four_layer_coverage.md      Four-layer coverage per §11.4.4(b)
│   ├── helixqa_integration.md      HelixQA Challenge bank wiring
│   └── quality_gates.md           Quality-gate configuration + thresho
│
├── 05_deployment/
│   ├── README.md                   Deployment guide
│   ├── podman_compose/             podman-compose definitions for all
│   ├── health_monitoring/          Health-check scripts + dashboards
│   ├── backup_restore/             Backup/restore procedures
│   └── cluster_provisioning/       Multi-node cluster setup automation
│
├── 06_risk/
│   ├── README.md                   Risk management guide
│   ├── risk_register.md            Enumerated risks with severity + mi
│   ├── failure_modes.md            Failure-mode analysis per component
│   ├── circuit_breakers.md         Circuit-breaker design + thresholds
│   ├── fallback_chains.md          Fallback chains (local→remote→cloud
│   └── rollback_procedures.md      Rollback procedures per scenario
│
└── poc/
    ├── README.md                   POC walkthrough
    ├── feature_trace.md            Sample feature lifecycle trace (end
    ├── benchmarks.md               Performance benchmarks
    └── limitations.md              Known limitations + tracked migrati
```

References

#	Title	URL	Relevance
1	**Spec-Kit** (SDD toolkit)	https://github.com/github/spec-kit	Spec-Kit core: specify, plan, tasks, implement lifecycle
2	**Superpowers** (agent skills)	https://github.com/obra/superpowers	15+ agent skills: brainstorming, TDD, debugging, code review, refactoring, documentation

#	Title	URL	Relevance
3	**SuperSpec** (SpecKit-Superpowers bridge)	https://github.com/WangX0111/superspec	Bridge orchestration: status, brainstorm, tasks, execute, review
4	**SuperB** (Spec-Kit discipline enforcement)	https://speckit-community.github.io/extensions/superb	Discipline gates: check, brainstorm, implementation-gate, critique, debug, finish
5	**Helix LLM** (distributed inference)	https://github.com/HelixDevelopment/helix_llm	Local-first distributed LLM inference: Gateway, Brain, Knowledge, Agents
6	**llama.cpp** (inference engine)	https://github.com/ggerganov/llama.cpp	RPC-server mode for distributed inference backend
7	**Helix Constitution** (governance submodule)	git@github.com:HelixDevelopment/helix_constitution.git	Universal development mandates, mechanical enforcement, agent rules
8	**MCP (Model Context Protocol)**	https://modelcontextprotocol.io	Protocol for LLM-tool interoperability
9	**Podman** (rootless containers)	https://podman.io	Rootless container runtime per §11.4.161
10	**Claude Code** (CLI agent)	https://docs.anthropic.com/en/docs/claude-code	Primary CLI agent implementing the autonomous loop
11		git@github.com:vasic-digital/Android_15.git	Reference consuming

#	Title	URL	Relevance
	ATMOSphere (consuming project)		project implementing the bridge
12	**Constitution submodule documentation**	`constitution/docs/` (this repository)	Full constitutional mandates, agent rules, and project-independent extensions