

11 4 103 Multi Track And Parallelism

2026 05 29

§ . . – Continuous parallel-stream
working routine (User mandate,
– –)

Short tag: `continuous-parallel-stream-routine` .

Forensic anchor — verbatim user mandate (2026-05-29):

"Do this working approach continuously and make it part of regular working routine and add it to the Constitution Submodule documented fully."

This anchor promotes the proven multi-stream operating pattern — repeatedly demonstrated effective across prior autonomous rounds — from an ad-hoc practice into the project's **standing default working routine**. §11.4.87 told the agent to keep working in an endless loop; §11.4.94 told it never to sit idle when something is progressable; §11.4.88 / §11.4.89 detached push and test execution; §11.4.96 catalogued what is safe to do during a long build. §11.4.103 binds those into a single always-on routine and adds the load-bearing invariant the prior anchors only implied: **the main work stream MUST always stay FREE**. It is the conductor's standing operating mode, not a per-request opt-in.

Strengthened (User mandate, 2026-05-29): the parallel-stream floor is raised from ≥ 2 to ≥ 3 and an **auto-backfill rule** is added — the moment any one stream is FULLY done, a new stream MUST immediately start and take its place, so the count NEVER drops below 3 while actionable workable items remain. Verbatim mandate: *"make sure there are always at all times at least 3 parallel working streams parallel with main work stream as long there are workable items we can do! as soon as one working stream is fully done, another new MUST start work and take its place."*

(A) Main stream stays FREE — ALL commit AND push run detached, never blocking. Every commit operation AND every push operation MUST run in the background (detached via `nohup ... > <log> 2>&1 & + disown`, per §11.4.88's commit-lock-release-immediately + detached-push seam). The conductor's main work stream MUST NOT block on a commit landing or a push completing — the moment the local commit is durable, the main stream returns to the priority queue while the push fans out to every upstream in the background. Blocking the main stream on a push (or a slow mirror) is a §11.4.103 violation (the exact 2026-05-26 §11.4.88 forensic harm — a 5-hour flock hold — generalised to the routine layer).

(B) ≥3 parallel background streams at all times whenever actionable work exists + auto-backfill. Whenever the queue holds three or more non-contending actionable items, the conductor MUST run **at least three** parallel background streams alongside the main stream (subagent-driven per §11.4.70 / §11.4.20, isolated per the §11.4.58 PWU worktree model + §11.4.84 working-tree quiescence). **Auto-backfill rule (User mandate, 2026-05-29):** the moment any one stream is FULLY done (its workable item closed with anti-bluff evidence per sub-rule (E), or merged), a new stream MUST IMMEDIATELY start and take its place by claiming the next-highest-priority non-contending item from the queue — the active-stream count NEVER drops below 3 while actionable workable items remain. Dropping below 3 active streams while three-plus disjoint actionable items wait — or serialising work that could run in parallel — is a §11.4.103 violation. The count may fall below 3 ONLY when there genuinely are no remaining non-contending actionable workable items OR every remaining item is externally blocked (per §11.4.94 / §11.4.97 / §11.4.101 idle-only-when-blocked). The §12.6 60% memory ceiling and the §11.4.58 6-concurrent-agent cap bound the stream count from above (so the standing band is 3–6 background streams).

(C) Most-critical + most-visible first; audio always top. Streams are filled in strict priority order — most-critical and most-user-visible items first — and **audio fixes are always the apex priority** per §11.4.72. When the conductor faces a choice of which item to dispatch onto a free stream, the highest-severity / most-user-visible / audio item wins; research, refactors, and low-stakes tasks fill remaining capacity only after the critical set is dispatched (per §11.4.42 priority-ordered iteration).

(D) Safe-during-build scope only while a heavy build runs. While the 42 GB containerised AOSP rebuild (§12.9) — or any heavy `gradle` / `m -j` build — is in flight, the parallel background streams **MUST** restrict themselves to the §11.4.96 SAFE catalogue: investigation / forensic analysis, documentation, test-authoring, gate + paired-mutation authoring, read-only on-device probes, submodule edits, web research, workable-items DB ops, and backgrounded pre-build / meta-test runs. The streams **MUST NOT** launch a second `gradle` / `m -j` / containerised build concurrent with the running one (§12.8 concurrency-hardening — a second heavy JVM/Soong build storms the host and risks the §12.7 OOM cascade). UNSAFE-during-build operations (§11.4.96 (α)-(θ)) are deferred to between-build windows.

(E) Heavy anti-bluff on every closure. Every item closed by a parallel stream **MUST** carry real captured evidence and survive the anti-bluff covenant family — root cause proven or explicitly marked `UNCONFIRMED:` / `UNKNOWN:` / `PENDING_FORENSICS:` per §11.4.6, captured-evidence per §11.4.5 + §11.4.69, deterministic consistency per §11.4.50, paired §1.1 meta-test mutations, and the structured root-cause investigation of §11.4.102 the moment any problem surfaces. A stream that closes an item on metadata-only / configuration-only / absence-of-error / grep-without-runtime evidence is a §11.4 PASS-bluff regardless of how parallel the routine looked.

(F) Idle ONLY when genuinely externally blocked. The conductor idles **ONLY** when every queued item is genuinely blocked on an external result it cannot accelerate (hardware access, network upstream, an in-flight build / test / push completion) **OR** the operator issues STOP **OR** §12 host-session-safety demands it (per §11.4.94(A) + §11.4.97). "Nothing visible to do" with progressable non-contending items in the queue is **NEVER** a valid idle reason. Posing a question and going idle while non-blocked items remain is a §11.4.94 + §11.4.97 + §11.4.103 violation; an unavoidable block parks one work unit, not the whole loop (§11.4.101 maximize-progress-while-blocked).

Composes with §11.4.58 (parallel-development PWU pipeline — §11.4.103 is its standing-routine framing), §11.4.70 / §11.4.20 (subagent-driven default — the parallel streams **ARE** subagents), §11.4.72 (audio top-priority — sub-rule (C)), §11.4.87 (endless-loop autonomous work — §11.4.103 is the always-on operating mode that anchor governs), §11.4.88 (background-push — sub-rule (A) is its routine-layer generalisation), §11.4.89 (background test execution — backgrounded

tests are one stream class), §11.4.94 (zero-idle priority-first parallel-by-default — sub-rules (B)+(C)+(F) bind it), §11.4.96 (safe-parallel-work-with-long-build catalogue — sub-rule (D) cites it), §11.4.97 (maximum-use-of-idle-time — sub-rule (F)), §11.4.101 (autonomous-decision-over-blocking — maximize-progress-while-blocked), §11.4.102 (systematic-debugging on any spotted problem inside any stream), §11.4.42 (priority-ordered iteration), §11.4.84 (working-tree quiescence for concurrent subagent commits), §12.6 / §12.7 / §12.8 / §12.9 (host-session safety bounds stream count + forbids a second heavy build), §9.2 (data-safety for any destructive op a stream would attempt).

Classification: universal (§11.4.17) — continuous parallel-stream working with a free main stream, priority-first dispatch, and anti-bluff closures is a reusable operating discipline for ANY project on ANY runtime that supports background execution + subagent delegation; nothing in the routine is hardware-, vendor-, or region-specific.

4-layer coverage per §11.4.4(b). Propagation gate `CM-COVENANT-114-103-PROPAGATION` enforces the literal anchor `11.4.103` across the canonical consumer fleet (parent + owned-submodule `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Paired §1.1 meta-test mutation strips the `11.4.103` literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.103 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--block-main-stream`, `--synchronous-commit`, `--synchronous-push`, `--single-stream-only`, `--skip-parallel-streams`, `--serialise-actionable-work`, `--idle-without-queue-survey` flag exists. The 2026-05-29 user mandate is unambiguous: continuous parallel-stream working with a free main stream is the regular working routine, not a per-request opt-in.

§ . . — **Big-work-item feature work-stream lifecycle** (own copy-on-write project copy + own branch/tags + per-feature builds, no-merge-until-approved, trunk-merged-into-every-stream) (User mandate, — —)

Every BIG work item — a new feature OR a drastic/large fix — MUST be developed as its own isolated **feature work-stream**: its OWN full project copy in a sibling directory, its OWN feature branch, its OWN per-feature flashable builds + feature release tags, kept SEPARATE from the trunk until manually approved after testing, with the trunk regularly merged INTO every feature stream. This generalises Git-flow's feature/release/tag discipline (long-lived `main` trunk, one `feature/<slug>` per big item, no merge to trunk until reviewed-complete) and binds it to the project's parallel-stream (§11.4.58/§11.4.103), single-builder, and finite-test-device realities. The mandate (ALL hold):

(A) One big item = one feature work-stream in a sibling project copy. Each big work item gets its own complete project copy in a sibling directory (`<project>` , `<project>_<slug-1>` , `<project>_<slug-2>` , ...), the `<slug>` a lowercase identifier per §11.4.29. The trunk copy is the canonical `main` . Small changes stay on the trunk; only BIG items earn a stream (the threshold is the project's own per §11.4.35).

(B) Disk-feasible duplication is MANDATORY, naive copies FORBIDDEN. A feature-stream copy MUST be created by a near-zero-disk mechanism — on a copy-on-write filesystem (btrfs/XFS/ZFS/APFS) a reflink/CoW clone (`cp -a --reflink=auto` / filesystem snapshot) so only DIVERGED bytes consume disk; elsewhere a shared-object-store mechanism (e.g. `git worktree` + externalized build output + sparse where the build tolerates it). A full deep copy or full `git clone` per stream of a large working tree is FORBIDDEN where it would exhaust free space (§9 data safety + the project's host-disk reality per §11.4.35). The chosen mechanism MUST be verified to consume ~0 disk for a fresh stream (captured evidence per §11.4.69) before being relied on (§11.4.6 — no guessing about disk). Build output (`out/` -class) MUST be per-stream and excluded from the CoW clone; one shared compiler cache (ccache-class) across streams maximises queued- build hit-rate.

(C) Own branch + own tags, separate until approved. Each stream develops on its own `feature/<slug>` branch; per-feature flashable builds are marked with feature tags under a scheme that embeds the slug so the build is greppable across every repository (composes §11.4.151), e.g. `<base-version>-feat-<slug>[ -<iter>]`, the `<iter>` monotonic per §11.4.54. **The feature branch is NEVER merged to the trunk until the operator manually approves after testing** — approval requires the §11.4.40 full-suite retest GREEN on a clean target, then the §11.4.41 merge-first integration, then a §11.4.113 ff-only (no-force) merge to trunk.

(D) Trunk merged INTO every feature stream, regularly. The trunk MUST be merged INTO every live feature stream FREQUENTLY (at minimum after every trunk tag / daily) to prevent long-lived-branch conflict storms — `git merge` (never rebase a shared/tagged branch). A stream that has not been synced within the cadence is flagged stale.

(E) Feature-level branching AND tagging cascades to ALL submodules the moment a submodule is modified. The instant a stream modifies ANY submodule, that submodule gets its own `feature/<slug>` branch; per-feature tags on the main repo are mirrored on EVERY touched submodule at the same HEAD and ff-only-pushed to all owned remotes (§11.4.113 + the project's tag-cascade rule). Untouched submodules stay trunk-pinned (no branch/tag churn). Merge-to-trunk applies the same cascade. Where the duplication mechanism is `git worktree`, the experimental/disk- duplicating worktree-submodule path (each worktree keeps a unique submodule copy; removal must be forced) is a documented hazard and the CoW-clone mechanism (each stream owns CoW-shared submodule git dirs) is preferred where available.

(F) Single-builder build queue + finite-device test queue. Where the host can run only ONE build at a time, feature builds MUST QUEUE (FIFO + advisory lock, composes §11.4.58 L1 / §12.7 / §12.8), never contend. Where live testing has a finite device count, per-device exclusive ownership (§11.4.119) serialises tests; streams beyond the device count queue. The build queue MUST garbage-collect idle streams' regenerable build output (§11.4.77) to keep disk within budget.

(G) Every change crosses the full quality gauntlet. Every change in any stream passes the extensive independent code review (§11.4.142/§11.4.125/§11.4.134)

- multi-angle impact research (§11.4.145) + anti-bluff captured-evidence testing (§11.4 / §11.4.5 / §11.4.69 / §11.4.107) — a feature stream is NOT a quality carve-out.

(H) Resumable, data-safe stream lifecycle. A stream registry (single source of "is this stream owed?", composes §11.4.116 / §11.4.147) tracks each stream's branch, base tag, build/out state, last-trunk-merge, and merge-approval state. Stream creation is atomic-on-success (clone to a `.partial` name, rename on success) and resumable (§11.4.84 quiescence check on resume). Retiring a stream requires a §9.2 pre-op backup + a guard that the feature work is merged/approved (unmerged work blocks auto-deletion). Host-session safety (§12) and the ≤60%-memory ceiling (§12.6) bound clone + build.

Honest boundary (§11.4.6): this lifecycle guarantees ISOLATION, disk-feasibility, greppable identity, and a controlled trunk ↔ feature merge discipline — it does NOT by itself prove any stream's work is correct (each change still crosses §11.4.108 runtime-signature + §11.4.40 retest). Classification: universal (§11.4.17) — the consuming project supplies its filesystem/CoW mechanism, sibling- dir + tag naming, build-queue + device-queue counts, and big-item threshold per §11.4.35.

Propagation gate `CM-COVENANT-114-167-PROPAGATION` (literal `11.4.167` across the consumer fleet) + recommended gates `CM-FEATURE-WORKSTREAM-COW-CLONE` (a feature stream is created by the verified near-zero-disk mechanism, NOT a full deep copy/ clone; per-stream out/ + shared cache; captured disk-cost evidence) + `CM-FEATURE-WORKSTREAM-NO-MERGE-UNTIL-APPROVED` (no feature → trunk merge without the approval marker + §11.4.40 GREEN) + `CM-FEATURE-WORKSTREAM-TRUNK-SYNC-CADENCE` (every live stream's last-trunk-merge within the declared cadence) + `CM-FEATURE-WORKSTREAM-SUBMODULE-CASCADE` (every touched submodule carries the `feature/<slug>` branch + mirrored feature tag). Paired §1.1 meta-test mutations: strip the literal → propagation gate FAILs; create a stream via full deep copy → `CM-FEATURE-WORKSTREAM-COW-CLONE` FAILs; merge a feature to trunk with no approval marker → `CM-FEATURE-WORKSTREAM-NO-MERGE-UNTIL-APPROVED` FAILs; let a stream go stale past the cadence → `CM-FEATURE-`

WORKSTREAM-TRUNK-SYNC-CADENCE FAILs; modify a submodule in a stream without its `feature/<slug> branch/tag` → `CM-FEATURE-WORKSTREAM-SUBMODULE-CASCADE` FAILs (gate-code = separate work item).

Composes §9 / §9.2 / §11.4.6 / §11.4.17 / §11.4.28 / §11.4.29 / §11.4.35 / §11.4.40 / §11.4.41 / §11.4.54 / §11.4.58 / §11.4.66 / §11.4.69 / §11.4.71 / §11.4.77 / §11.4.84 / §11.4.103 / §11.4.108 / §11.4.113 / §11.4.116 / §11.4.119 / §11.4.125 / §11.4.142 / §11.4.145 / §11.4.147 / §11.4.151 / §12 / §12.6 / §12.7 / §12.8 / §12.9 / §1.1.

Canonical authority: constitution submodule `Constitution.md` §11.4.167. Non-compliance is a release blocker. No escape hatch — no `--full-clone-per-feature`, `--deep-copy-stream`, `--skip-cow-clone`, `--merge-feature-without-approval`, `--no-trunk-sync`, `--skip-submodule-cascade`, `--unbounded-parallel-builds`, `--shared-device-no-lock`, `--retire-unmerged-stream`, `--feature-stream-quality-carveout` flag.

§11.4.176 — Conflict-free multi-track work-division + exactly-once claim registry + capability-aware deadlock-proof device-lock (User mandate, 2026-07-02). When a project runs multiple concurrent work-STREAM tracks (§11.4.167 feature work-stream lifecycle; §11.4.58/§11.4.103 parallel streams) that share (a) a common pool of workable items and (b) a scarce pool of exclusive hardware/resources, TWO cooperating arbitration layers are MANDATORY — both conflict-free BY CONSTRUCTION, not by convention. Forensic anchor — verbatim user mandate (2026-07-02, REM-17): "MUST prevent multiple Tracks working the same workable item OR the same group of workable items resembling one consistent logical group of functionalities. Division of work assigned to Tracks MUST be conflict-free: no conflicts, misunderstandings, race conditions, data corruptions, codebase loss. Risk-free + safe; results ALWAYS produced; no false results, no bluff." (REM-02): "there MUST exist proper synchronization mechanism for flashing and use of device ... lock ... release ... available test devices pool and information about held locks ... dead-lock-proof ... a track locks one OR more devices selected by required capability." **(A) Work-item layer — exactly-once claim registry (REM-17).** Every workable item, OR every logical GROUP of items that share files / subsystem / feature / schema (a hidden-coupling cohort), is assigned to AT MOST ONE track at a time. (1) **Exactly-once** — a lock-arbitrated, append-only claim registry (extending §11.4.147 no-work-loss registry + §11.4.116 real-time sync substrate) records `claim(item|group, track, ts, ttl)`; a track MUST claim-before-touch; a second claim on an already-claimed item is DENIED

(no silent overlap). (2) **Logical-group cohesion** — items that share a file-scope / subsystem / feature / schema are one group claimed WHOLE by one track (prevents the hidden-coupling race two tracks editing the same seam would create — §11.4.58 L3 disjoint-file-scope is the enforcement). (3) **Conflict-free by construction** — disjoint file-scope partition (§11.4.58 L3) + per-track worktree/CoW isolation (§11.4.167) means two tracks physically cannot touch the same bytes. (4) **No codebase loss** — §9.2 hardlinked backups, §11.4.113 ff-only merge-onto-latest-main (NEVER force-push), §11.4.84 working-tree quiescence, §11.4.41 merge-first. (5) **Anti-bluff** — the assignment + the merge are verifiable (registry event stream + captured merge evidence); results are ALWAYS produced (a claimed-but-abandoned item's claim TTL-expires and returns to the pool, §11.4.147). (B) **Hardware layer — capability-aware deadlock-proof device-lock (REM-02)**. When N tracks outnumber M exclusive test devices, exactly ONE track owns each device at a time (§11.4.119 single-resource-owner). A track leases ONE OR MORE devices selected by required CAPABILITY (not merely identity) — the pool records per-device capabilities; leases are capability-matched; multi-device leases are supported. The arbiter is DEADLOCK-FREE by construction — all four Coffman conditions broken: hold-and-wait ELIMINATED (multi-device acquire is ALL-OR-NOTHING inside one critical section — never a partial hold while waiting); circular-wait ELIMINATED (acquire is NON-BLOCKING — wins the whole set at once or returns EBUSY immediately, never waits on another track's device); no-preemption RELAXED (lease TTL + heartbeat + reap — a crashed holder's lease auto-expires, §11.4.147, so a dead track never holds the pool hostage); the only lock ever held is a short bounded flock over the atomic check-and-claim critical section, released before any device work. Registry = append-only JSONL event stream + atomically-rewritten snapshot (§11.4.116). (C) **Universal decoupling + evidence-based resource-tuning (REM-03)**. Both layers are universal, decoupled, project-not-aware, reusable (§11.4.17/§11.4.28/§11.4.35) — the consuming project injects its claim-store, device pool, and metric sources at runtime; nothing project-specific lands in the canonical root. Resource-tuning decisions (parallel-stream count, per-stream memory budget §12.6, device-lease breadth, build-queue depth §11.4.167) MUST be EVIDENCE-BASED — driven by captured metrics (§11.4.24 / §11.4.5), never guessed (§11.4.6). Composes §11.4.6 / §11.4.17 / §11.4.24 / §11.4.28 / §11.4.35 / §11.4.41 / §11.4.58 / §11.4.84 / §11.4.85 / §11.4.103 / §11.4.113 / §11.4.116 / §11.4.119 / §11.4.147 / §11.4.167 / §12.6 / §9.2. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-176-PROPAGATION (literal 11.4.176 across the consumer fleet) + recommended gates

CM-WORK-DIVISION-EXCLUSIVE-CLAIM + CM-NO-CROSS-TRACK-SCOPE-OVERLAP + CM-LOGICAL-GROUP-COHESION + CM-DEVICE-LOCK-DEADLOCK-FREE (device-lock arbiter ships an all-or-nothing + non-blocking + TTL-reap stress/chaos proof with a paired §1.1 mutation) + paired §1.1 meta-test mutations (gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.176. Non-compliance is a release blocker. No escape hatch — no --allow-cross-track-overlap , --skip-claim-registry , --split-logical-group , --blocking-device-acquire , --allow-partial-device-hold , --tune-resources-by-guess flag.

§11.4.178 — Track-qualified identity for parallel work streams (session-name collision ban) (User mandate, 2026-07-04). Parallel work streams that share hardware, a git backing store, a tmux/session namespace, a lock namespace, or any single-owner resource MUST be addressed by a TRACK-QUALIFIED identity (<project>__<track>__<role>), NEVER a bare directory basename. Multiple checkouts/clones/worktrees sharing basename atmosphere collapse into one session/lock/log key and cross-wire tmux targets, lock files, device claims, and log directories (the observed collision, 2026-07-04). Every registry row, session name, lock path, log directory, and device claim MUST carry the track-qualified key; a bare-basename identity for a resource with ≥ 2 concurrent claimants is a violation. Honest boundary (§11.4.6): a unique identity prevents cross-wiring — it does not itself arbitrate ownership (that is §11.4.119 + **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor)). Composes §11.4.111 (resolve-by-stable-name-not-by-index) / §11.4.119 (single-resource-owner partitioning) / **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor) / §11.4.29 (naming) / §11.4.58 (parallel PWU). Classification: universal (§11.4.17) — the consuming project supplies its <project> / <track> / <role> vocabulary per §11.4.35. Propagation gate CM-COVENANT-114-178-PROPAGATION (literal 11.4.178) + recommended gate CM-TRACK-QUALIFIED-IDENTITY (every session name / lock path / log dir / device claim for a multi-claimant resource carries the track-qualified key, never a bare basename) + paired §1.1 meta-test mutation (address a shared resource by bare basename → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Multi-project-per-track layout clarification (2026-07-10):** a /mnt/track<N> track parent is NOT single-project — MULTIPLE projects coexist on the same track parents, each in its OWN project-named subdirectory /mnt/track<N>/<project_name>/ (e.g. /mnt/track1/atmosphere/ AND /mnt/

track1/helix_ota/ side by side); the per-track working copy of project P on track N is ALWAYS /mnt/track<N>/<project_name>/ (the project's canonical lowercase snake_case dir-name per §11.4.29), never the bare track root and never a shared basename. This is precisely WHY the identity is

<project>__<track>__<role> and carries the <project> component: it keeps two coexisting projects' track-N sessions / locks / logs / device-claims / registry rows from colliding. A project adopting this constitution + the multi-track engine (§11.4.187) MUST create AND register its own /mnt/track<N>/<project_name>/ per track under its project-qualified identity — it MUST NOT assume it owns the track root or that a track hosts only one project. **Canonical authority:** constitution submodule Constitution.md §11.4.178. Non-compliance is a release blocker. No escape hatch — no --bare-basename-ok , --skip-track-qualifier , --shared-session-name flag.

§11.4.182 — Track+branch work-stream identity label (T<N>/<branch>) on every agent/subagent/work-stream label and every operator-facing work-stream reference (User mandate, 2026-07-05). Forensic anchor — operator mandate (2026-07-05): "EVERY agent/subagent/work-stream label AND every operator-facing report reference to a work stream MUST carry a

(T<N>/<branch>) prefix identifying the track + git branch — e.g. (T1/main) ATM-312 MPV drm_prime investigation , (T2/feat/mistiq-vader) SPK-XXX" In a multi-track layout (parallel checkouts each on a distinct branch — e.g. /mnt/track1..4) a bare work-stream label ("ATM-312 investigation") does not say WHICH track/branch produced it, so cross-track dispatches, status reports, and progress lines collide and mislead. The mandate (ALL hold): (1) **Universal label form.** Every agent / subagent / work-stream label AND every operator-facing reference to a work stream (agent-dispatch descriptions, status reports, resume/handoff docs, progress lines) MUST START with a (T<N>/<branch> - <alias>) prefix — <N> the track number, <branch> the git branch, <alias> the claude-toolkit alias in use — e.g. (T1/main - claude1) ATM-312 MPV drm_prime investigation , (T2/feature/mistiq-vader - deepseek) SPK-XXX subtitle work . (2) **Deterministic derivation, never guessed (§11.4.6).** The label is DERIVED, not typed from memory: <N> from the checkout path (the multi-track root /mnt/track<N>/... ; honest ? when the checkout is not under a numbered track — never a guessed number), <branch> from git rev-parse --abbrev-ref HEAD , <alias> from CLAUDE_CONFIG_DIR basename .claude-<alias> → <alias> (honest ? when unset/non-matching, never a guessed name). The

consuming project supplies a deterministic labeler (reference impl `scripts/multitrack/track_branch_label.sh`). (3) **Mechanical enforcement (§11.4.75 / §11.4.109-class anti-forgetting)**. A PreToolUse guard hook (reference impl `constitution/scripts/hooks/guard-track-branch-label.sh` , inherited BY REFERENCE per §11.4.177, never copied) BLOCKS at the tool-call boundary any agent-dispatch tool (Agent / Task / TaskCreate) whose `description` (or `subagent label`) does not START with `^(T[0-9]+/[^)]+\\)` — so enforcement is independent of any agent's recall (a rule the orchestrator forgets to paste is not enforcement); non-agent tools pass through untouched. A pre-build gate `CM-TRACK-BRANCH-LABEL` asserts the labeler + hook exist, are executable + parseable, the hook is wired in the runtime settings, and the convention doc exists; a propagation gate `CM-COVENANT-114-182-PROPAGATION` asserts the literal `11.4.182` is present across the governance-file fleet. Paired §1.1 mutation: strip the hook's prefix-check → the hook test-suite FAILs; strip the `11.4.182` literal from a governance file → the propagation gate FAILs; restore → PASS. (4) **Honest boundary (§11.4.6)**. A `?` in the track OR the alias field is a valid honest label (unknown track / no git repo), never a bluff; the enforced FORMAT is numeric-track `[0-9]+` per the reference regex, so an agent dispatch from a non-track checkout is surfaced (blocked) rather than silently mislabeled. Classification: universal (§11.4.17) — a platform-neutral multi-stream identity discipline reusable by any project running parallel track/branch work streams; the consuming project supplies its concrete track-root layout, labeler, hook wiring, and governance-file set per §11.4.35. Composes §11.4.178 (track-qualified identity — §11.4.182 is the label-prefix instantiation applied to EVERY agent/subagent label + every operator-facing work-stream reference) / §11.4.29 (naming) / §11.4.75 (mechanical enforcement) / §11.4.109 (anti-forgetting PreToolUse guard) / §11.4.177 (hook inherited by reference, project-decoupled) / §11.4.6 (label is a derived fact; `?` is honest-unknown) / §11.4.58 · §11.4.103 (parallel work streams) / §11.4.119 (single-resource-owner per track) / §11.4.116 · §11.4.147 (work-stream registry identity). Propagation gate `CM-COVENANT-114-182-PROPAGATION` (literal `11.4.182` across the consumer fleet) + recommended gate `CM-TRACK-BRANCH-LABEL` + paired §1.1 meta-test mutation. **TRUNK RULE (operator mandate, 2026-07-28)**: work performed on the canonical trunk (`main / master`) is **ALWAYS Track 1** — the label reads (`T1/main - ...`) , NEVER (`T?/main - ...`) , regardless of checkout path. Precedence is deliberate: the BRANCH decides first and the checkout path decides only for non-trunk branches, so a trunk checkout living outside `/mnt/`

track<N> — or inside a differently-numbered one — is still Track 1; the identity that matters is which STREAM the work belongs to, not which directory it happens to sit in. Rationale: §11.4.182 labels are how concurrent streams are told apart, so an unnumbered T? on trunk work makes it indistinguishable from a genuinely-unknown track in every dispatch, report and progress line — and the trunk is a well-known, singular stream that is never unknown. Reference implementation:

scripts/multitrack/track_branch_label.sh derives the branch BEFORE the track number and maps main|master -> 1 . Labels MUST be re-derived from that script at dispatch time and kept in real-time sync with the live branch/alias/model/effort — a remembered or stale label is itself a §11.4.182 violation (§11.4.6: live identity is not knowable from memory). **Canonical authority:** constitution submodule Constitution.md §11.4.182. Non-compliance is a release blocker. No escape hatch — no --skip-track-label , --unlabeled-agent-OK , --bare-work-stream-reference , --no-track-branch-prefix , --label-from-memory flag.

§11.4.183 — Maximal multi-agent utilization per work-stream + full-constitution-application + zero-bluff mandate (User mandate, 2026-07-07).

Forensic anchor — verbatim operator mandate (2026-07-07): "Every Track MUST USE and maximize multi-agents working approaches (subagent-driven, independent agents for various reviewing, and other forms of agents use)! We MUST maximize use of agents so every track produces as much as possible completed work with highest possible quality for the least spent time! Every Track MUST apply everything from the constitution — all rules, mandatory constraints, guides and technology derived! Nothing can be ignored, skipped, avoided or considered less relevant! There MUST NEVER BE false or faulty results, untested and unverified work or bluff of any kind anywhere!" Every track / work-stream MUST maximize multi-agent working approaches — subagent-driven development (§11.4.20/§11.4.70), independent review agents (§11.4.125/§11.4.142/§11.4.165/§11.4.134), parallel background streams with auto-backfill (§11.4.58/§11.4.103), and every other applicable agent form — so that every track produces the MOST completed work at the HIGHEST quality in the LEAST time. The mandate (ALL hold): (1) **Maximize agent utilization per track.** Every track/work-stream (main + every parallel background stream per §11.4.58/§11.4.103) MUST default to subagent-driven execution for non-trivial work (§11.4.20/§11.4.70), dispatch independent review agents structurally separated from the author for quality/safety/correctness review (§11.4.125/§11.4.142/§11.4.165, iterated to a clean GO per

§11.4.134), and use every other applicable specialised-agent form the runtime offers — never one monolithic foreground agent doing everything a task naturally decomposes into. (2) **Full-constitution application — nothing ignored.** Every track MUST apply the ENTIRE constitution to its work — every universal rule, every mandatory constraint, every guide, and every derived technology (gates, mutation pairs, sync engines, recording/evidence pipelines, lock/claim registries) this Constitution and its consuming project's governance define. No rule may be silently ignored, skipped, avoided, or judged "less relevant" for a given track — partial-constitution-application is itself a §11.4.183 violation, equal in severity to skipping a mandated gate. (3) **Zero-bluff, zero-false-result, zero-untested-work.** There MUST NEVER be false or faulty results, untested or unverified work, or bluff of any kind, anywhere, on any track — this restates and BINDS the entire §11.4 anti-bluff covenant (§11.4/.1/.2/.5/.6/.7/.50/.69/.107) as an unconditional per-track invariant, not merely a release-gate check: every claim of "done" on any track carries captured evidence (§11.4.5/§11.4.69/§11.4.107), every gate/mutation pair is genuinely falsifiable (§1.1), and no track may report progress it did not actually perform (§11.4.126). Honest boundary (§11.4.6): "maximize agents" is bounded — by §12.6's 60% host-memory ceiling, §11.4.58's parallel-agent cap, genuine task parallelizability, and §11.4.119's single-resource-owner rule for contended hardware/locks — so spawning contending, redundant, or non-parallelizable agents is NOT maximization; the bar is maximum USEFUL parallel throughput per track, never raw agent count. Classification: universal (§11.4.17) — a platform-neutral operating-discipline reusable by any project running one or more parallel work-streams/tracks; the consuming project supplies its concrete track layout, subagent runtime, and review-agent roster per §11.4.35. Composes §11.4.20 (subagent-driven-by-default) / §11.4.58 (parallel-development PWU pipeline) / §11.4.70 (subagent-driven execution is the default) / §11.4.92 (multi-pass change evaluation) / §11.4.94 (zero-idle priority-first parallel-by-default) / §11.4.97 (maximum-use-of-idle-time) / §11.4.103 (continuous parallel-stream working routine) / §11.4.119 (single-resource-owner partitioning) / §11.4.125 (code-review-agent gate) / §11.4.126 (default autonomous-loop working mode) / §11.4.134 (code-review iterate-until-GO) / §11.4.142 (universal code-review mandate) / §11.4.165 (universal independent verification agent mandate) / §12.6 (60% memory ceiling). Propagation gate `CM-COVENANT-114-183-PROPAGATION` (literal `11.4.183` across the consumer fleet) + recommended gate `CM-MAX-AGENTS-PER-TRACK` (every active track shows subagent/review-agent dispatch evidence for non-trivial work, not solitary foreground execution) + paired §1.1 meta-test mutation (strip the literal

→ propagation gate FAILs; record a track that skipped subagent/review-agent dispatch for non-trivial work with no genuine single-agent justification → CM-MAX-AGENTS-PER-TRACK FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.183. Non-compliance is a release blocker. No escape hatch — no `--single-agent-OK` , `--skip-review-agents` , `--partial-constitution-application` , `--serialise-parallelisable-work` flag.

§11.4.187 — Automatic multi-track ruler orchestration (User mandate, 2026-07-04; landed 2026-07-09). Compact summary: every project incorporating this constitution MUST ship its multi-track parallel-development orchestration as a UNIVERSAL, AUTOMATIC, OUT-OF-THE-BOX, INHERITED capability — not a per-project reimplementation — via a single conductor session that programmatically SPAWNS, DRIVES, RESUMES, MONITORS, and CONTROLS per-track headless workers. (1) **Headless spawn primitive** — workers launch via native headless `claude -p --output-format stream-json --verbose` ; the worker's `session_id` is captured from the FIRST `init` event (never invented); success/failure is read from the terminal event's `result.is_error` field, NEVER the process exit code (a CLI can exit non-zero on a successful `result` , or zero on a failed one — `exit-code-as-success-signal` is a §11.4.6 guess); continuation of an existing worker uses `--resume <session_id>` invoked from the SAME cwd + env the spawn used (the storage/resume-collision gotcha). (2) **Per-subscription auth + rate-limit fallback** — every spawn `unset` s any ambient `ANTHROPIC_API_KEY` (the precedence trap that silently routes a "subscription" worker onto pay-per-token API billing) and sets the selected alias's `CLAUDE_CODE_OAUTH_TOKEN` / `CLAUDE_CONFIG_DIR` ; on a captured rate-limit signature (`isApiErrorMessage + apiErrorStatus:429` , quota-exhausted vs transient distinguished by evidence, never guessed) the ruler REBINDS the affected track onto the next healthy alias (cross-track reuse permitted), falls back to a provider alias where configured, and bounded-parks (§11.4.101) rather than hanging when every alias is cooled. (3) **Crash-resilient ruler self-supervisor** — the ruler's OWN state (per-track alias/session_id/worktree/last-verdict) is persisted durably (write-temp-then-rename, NEVER tmpfs) so an external watchdog detects a dead ruler PID and re-hydrates every track's binding from the durable state + the §11.4.116 event stream — no work is silently lost to a ruler crash (§11.4.147 at the ruler layer). (4) **Idempotent bootstrap + auto-install + conductor-stays-home** — one idempotent bootstrap command wires the whole system out of the box (cwd-hook

install, config validation, reconciliation), auto-invoked on every constitution pull via the §11.4.164 `post_update_hook.sh` skill-register seam; the

`conductor`: -designated alias/session is NEVER bound into a track worktree (the resolver returns nothing for it) so the conductor always runs from the main checkout. (5) **Project-agnostic engine, consumer-owned data** — the entire mechanism (config loader, alias → worktree resolver, cwd-hook, bind/fallback/cooldown orchestrator, session launcher, rate-limit monitor, supervisor) lives at `constitution/scripts/multitrack/`, inherited BY REFERENCE (§11.4.28(B)/§11.4.177), carrying NO project literal; every consuming project supplies its own per-host `config/multitrack/<hostname>.yaml` (tracks/mounts/serials/aliases/conductor) as DATA, never as an engine edit. Anti-bluff (§11.4/§11.4.108): every piece carries four-layer coverage (pre-build gate + on-host test + paired §1.1 mutation + captured evidence per §11.4.69), and each fix declares a runtime signature verified on a clean host — never a claim the mechanism "should work." Honest boundary (§11.4.6): this anchor mandates the MECHANISM; genuine per-subscription quota-isolation proof requires live tokens exercising real rate-limit boundaries (an operator-scheduled live acceptance window) — a mock-driven GREEN proves the spawn/resume/fallback CONTRACT, never live-quota behaviour, and the two are never conflated or substituted for each other. Classification: universal (§11.4.17). Composes §11.4.20 (subagent-driven-by-default) / §11.4.28 (submodule decoupling + dependency layout) / §11.4.58 (parallel-development PWU pipeline) / §11.4.70 (subagent-driven execution default) / §11.4.101 (autonomous-decision-over-blocking — the bounded-park-on-all-cooled behaviour) / §11.4.103 (continuous parallel-stream routine) / §11.4.111 (resolve-by-stable-name — config by serial, never by index) / §11.4.116 (real-time conductor ↔ framework sync channel) / §11.4.119 (single-resource-owner partitioning) / §11.4.126 (default autonomous-loop mode) / §11.4.128 (always-on device-recording — the analogous host-safety discipline) / §11.4.131 (standing session-resumption file) / §11.4.147 (crashed-agent respawn-until-complete — the ruler-self-supervisor is its ruler-layer instantiation) / §11.4.164 (constitution auto-propagation hook seam) / §11.4.167 (feature work-stream lifecycle) / §11.4.176 (multi-track work-division + exactly-once claim + device-lock) / §12.6/§12.8 (host-safety memory + concurrency ceilings — the spawn/launch pool is budget-gated). Propagation gate `CM-COVENANT-114-187-PROPAGATION` (literal `11.4.187`) + recommended gates `CM-MULTITRACK-ENGINE-IN-CONSTITUTION` (the generic ruler scripts present under `constitution/scripts/multitrack/`, carrying no project literal, `bash -n + sh -n` clean per §11.4.67) / `CM-MULTITRACK-`

AUTOINSTALL-WIRED / CM-MULTITRACK-BOOTSTRAP-IDEMPOTENT / CM-MULTITRACK-BIND-ON-START / CM-MULTITRACK-FALLBACK-MONITOR / CM-MULTITRACK-HEADLESS-SPAWN / CM-MULTITRACK-RULER-SUPERVISOR + paired §1.1 mutation (strip the literal → propagation gate FAILs; re-introduce a project literal into the engine, or downgrade a rate-limit rebind to exit-code-based success detection → the corresponding mechanism gate FAILs). **Canonical authority:** constitution submodule Constitution.md §11.4.187. Non-compliance is a release blocker. No escape hatch — no `--skip-multitrack`, `--manual-multitrack-only`, `--no-auto-install`, `--per-project-reimpl-OK`, `--exit-code-is-success-signal` flag.

§11.4.191 — Work-to-track/branch binding enforcement: a logic-group's work can NEVER be committed OR dispatched onto the wrong track/branch (research-derived, 2026-07-10). Compact summary: every logical group of workable items binds to EXACTLY ONE canonical (track, branch) recorded ONCE in the single source of truth (the §11.4.93 workable-items DB `logic_groups.destination` = authoritative branch + a `canonical_track` column + a `group_paths` file-scope manifest), and NO change belonging to a group may LAND on, nor be DISPATCHED to, a track/branch other than that group's canonical one — the file-scope generalisation of §11.4.181/§11.4.182 from branch-NAME consistency to work-PLACEMENT consistency (forensic FACT 2026-07-10: `srccpy/Remote-app` work — group `mistiq-vader`, canonical `feature/mistiq-vader` /T2 — was committed onto `main` /T1; §11.4.181's `guard-branch-consistency.sh` never fired because committing onto an existing branch is not a feature-branch CREATE, and no gate checked which FILES land on which BRANCH). (1) **Single source of truth** — the `logic_group` → (branch, track, path-globs) binding is recorded ONCE in the DB and LOOKED UP, never re-invented (a commit/dispatch placing a group's file-scope on a non-canonical branch/track is a §11.4.6 no-guessing violation). (2) **Preventive hook (§11.4.109/§11.4.177)** — a PreToolUse guard `guard-work-track-binding.sh` (inherited BY REFERENCE, never copied) BLOCKS (exit 2) a git-commit-class Bash call whose to-be-committed file set contains a path owned by a group whose canonical branch/track ≠ the current checkout, AND BLOCKS an Agent/Task dispatch whose §11.4.182 label track/branch ≠ the canonical (track, branch) of the ticket's group; fail-closed on unreadable registry, `# guardrails:allow <reason>` audited escape. (3) **Detective gate (§11.4.110/§11.4.108)** — `CM-WORK-TRACK-BINDING-ENFORCED` asserts, on the current branch's own added commits (`merge-`

`base(HEAD,main)..HEAD`), that every `group_paths` -owned file's group destination == current branch (and `canonical_track` == current track when pinned) — the actual enforcement boundary regardless of how the commit was made, the runtime signature being FAIL-on-wrong-placement / PASS-on-canonical-placement on a clean checkout. (4) **Risk-free reconciliation (§9.2/§11.4.113/§11.4.84)** — a mis-placed change is reconciled by MERGING its commits onto the group's canonical branch (union, no commit lost, no `-s ours /reset`), ff-only push to all upstreams (NEVER force-push), then removing it from the wrong branch only after the merge is confirmed on all mirrors. Classification: universal (§11.4.17) — references no project hardware; the consuming project supplies its `(track, branch)` map, `group_paths` seed, and commit-wrapper wiring per §11.4.35.

Composes

§11.4.6/.28/.29/.35/.58/.75/.84/.93/.95/.109/.110/.113/.119/.167/.176/.177/.178/.181/.182/

§9.2. Propagation gate `CM-COVENANT-114-191-PROPAGATION` (literal `11.4.191`) + recommended gate `CM-WORK-TRACK-BINDING-ENFORCED` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.191. Non-compliance is a release blocker. No escape hatch — no `--allow-cross-track-commit`, `--skip-work-binding`, `--commit-any-branch-OK`, `--dispatch-any-track-OK`, `--no-file-scope-check` flag.

§11.4.192 — Continuous multi-track auto-backfill: a FREE track MUST be IMMEDIATELY re-assigned its next-highest-priority domain work-set, never left idle while its domain has actionable items (User mandate, 2026-07-11).

Verbatim operator mandate: "once a Track is free, we shall assign next set of workable items to it ... full continuous multi-track development ... no stopping ... this is MANDATORY". Compact summary: the moment any track's worker COMPLETES its current work-set (its item(s) closed with anti-bluff evidence, or merged), that FREE track MUST be IMMEDIATELY assigned the NEXT-highest-priority actionable work-set from its OWN domain (its §11.4.191 `logic_group` / canonical branch) — a track is NEVER left idle while its domain still holds actionable workable items, and the conductor/ruler CONTINUOUSLY re-assigns with NO stopping (full continuous multi-track development); PROMOTES the §11.4.103(B) ≥ 3 -parallel-background-streams-with-auto-backfill invariant into the multi-track (§11.4.176/§11.4.178/§11.4.179/§11.4.187) layer (§11.4.103(B) backfills a STREAM the moment it is done; §11.4.192 backfills a TRACK the moment its worker is done, keyed to that track's canonical `(track, branch)` domain per §11.4.191). (1) **WHAT** — on any track-worker completion the conductor/ruler

(§11.4.187) MUST, per that track's domain, look up the NEXT actionable work-set, CLAIM it exactly-once (§11.4.176 exactly-once claim registry, no two tracks claiming one item), and DISPATCH it as the track's new §11.4.182-labelled work stream — without waiting for a scheduled wake or an operator prompt. (2)

PRIORITY — the next work-set is the track domain's highest-priority actionable item (§11.4.42/§11.4.72 priority order, audio-first where in scope), never an arbitrary pick. (3) **SOURCE-SIDE ADVANCES EVEN WHEN VALIDATION IS GATED** —

when the next item's on-device VALIDATION is blocked on a pending build / device availability (§11.4.108 clean-target / §11.4.185 manual-QA / operator-gated flash), the track MUST still advance the item's SOURCE-SIDE work (investigation, RED-test §11.4.115, source patch, pre-build gate + paired §1.1 mutation, docs) and honestly §11.4.3/§11.4.52 SKIP-with-reason ONLY the genuinely device-gated validation portion — a build/device gate on the validation step is NOT a reason to idle the whole track. (4) **HONEST BOUNDARY (§11.4.6)** — a track idles ONLY

when its domain GENUINELY has zero remaining actionable workable items OR every remaining item is externally blocked with no source-side advancement possible (§11.4.94(A)/§11.4.97/§11.4.101 idle-only-when-blocked closed set); "the merges are done" / "the current item is done" / "nothing visible right now" is NEVER a valid idle reason while the domain still holds progressable items (a

§11.4.94+§11.4.97 violation); an unavoidable per-item block PARKS one item, it does NOT stop the track (§11.4.101 — the track backfills its next non-blocked domain item). Classification: universal (§11.4.17) — references no project

hardware; the consuming project supplies its track → domain map + priority source per §11.4.35. Composes

§11.4.42/.58/.72/.94/.97/.101/.103/.167/.176/.178/.179/.182/.187. Propagation gate

CM-COVENANT-114-192-PROPAGATION (literal 11.4.192) + recommended gate

CM-CONTINUOUS-MULTITRACK-BACKFILL (on a track-worker completion the ruler re-assigns that track's next-highest-priority domain work-set; a track with a completed worker AND remaining actionable domain items but no new dispatch → the gate FAILs) + paired §1.1 mutation (let a free track sit idle while its domain holds an actionable item, OR treat "merges done" as a valid idle reason → the gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.192. Non-compliance is a release blocker. No escape hatch — no `--allow-idle-track`, `--skip-backfill`, `--stop-multitrack`, `--merges-done-is-idle`, `--defer-track-reassign` flag.

Appendix C – Multi-upstream push

This constitution submodule is hosted on multiple Git providers, and consuming projects often are too. The submodule ships an `install_upstreams.sh` helper (or invokes the `install_upstreams` system command from the parent toolkit) that reads `Upstreams/*.sh` declarations and configures all remotes locally.

Every commit to the constitution submodule MUST be pushed to ALL configured upstreams. Use a shell helper that iterates `git remote` and pushes one by one, OR configure `origin` with multiple push URLs (`git remote set-url --add --push origin <url>` per remote).

\$. . . — Parallelized-pipeline methodology: build validate overlap + affected-stages-only re-run + parallel multi-device testing/recording + fan-out-subagents-always (research-derived, — —)

Forensic motivation (genericised): the sequential pipeline shape — build FULLY completes → THEN validation begins → on ANY discovery, restart the WHOLE cycle from scratch → all devices tested one-at-a-time — serialises the operator's wall-clock onto build latency + one-device-at-a-time evidence gathering, so a multi-hour containerised build followed by serial per-device validation and a from-scratch restart-on-discovery makes time-to-a-manual-QA-capable-release a multiple of what the genuinely-parallel work requires. §11.4.58's PWU pipeline already overlaps Stage 1 of round N+1 with Stages 4–5 of round N; §11.4.89 already backgrounds long tests; §11.4.103 already mandates ≥3 auto-backfilled parallel streams; §11.4.119 already partitions shared hardware by single-owner; §11.4.128/§11.4.144 already mandate always-on recording + availability-following. §11.4.230 BINDS those into one operating discipline and adds the three invariants they did not each state on their own: (A) the build ↔ validate CONCURRENCY seam + affected-stages-only re-run-on-discovery, (B) all-devices-concurrent + worker-pool

recording/processing, (C) fan-out-subagents as the standing default. It is the parallelization capstone the way §11.4.103 is the continuous-parallel-stream capstone — a strengthening bind, never a weakening.

(A) PIPELINE-OVERLAP PARALLELIZATION. Validation/verification work **MUST** run **CONCURRENTLY WITH** the build wherever a validation stage does **NOT** depend on the build's own not-yet-produced artifact — never deferred until the build finishes when it could have run alongside it. Deploy the moment the build is done (build-completion is the deploy trigger, per §11.4.121 no-commit-while-build-writes-tracked-artifacts + §11.4.108 artifact-must-exist ordering — the deploy waits for the artifact, the deploy-independent validation does not). Concretely, **WHILE** a build runs, the following run in parallel per the §11.4.96 **SAFE-during-build** catalogue: pre-build source-layer gates + meta-test mutation sweeps + coverage/clash checks (§11.4.110), the §11.4.125/§11.4.142/§11.4.194 independent code-review on the batch's **SOURCE**, doc/DB sync (§11.4.106), and on-device validation of the **PRIOR** build (device-independent of the in-flight build, §11.4.119). On any **NON-huge-blocker** discovery, the cycle **MUST**: (1) bump the version code (a new, distinctly-identifiable candidate artifact — never re-use the failing artifact's id, §11.4.111 stable-name + §11.4.200 target-identity); (2) fix at the proven root cause (§11.4.102 systematic-debugging first, §11.4.115 **RED-on-the-broken-artifact**); (3) **RE-RUN ONLY THE AFFECTED STAGES, PARALLELIZED** — the stages whose inputs the fix changed (its own targeted validation **FIRST** per §11.4.130, then the affected regression set), **NOT** a serial restart-from-scratch of the entire cycle. The explicit efficiency target is a **4–5× reduction in time-to-a-manual-QA-capable-release** versus the fully-sequential shape — a measured target, tracked as build-and-validate wall-clock deltas (§11.4.24 build-resource stats), never an assumed one (§11.4.6). **HONEST BOUNDARY (§11.4.6 / §11.4.108 / §11.4.120):** the **CURRENT** build's **ARTIFACT-layer** (§11.4.108 layer 2) and **RUNTIME-ON-CLEAN-TARGET-layer** (layer 3) validation **CANNOT** precede its own artifact — a gate whose precondition is the build's output belongs at a post-build seam, never overlapped onto authorship (§11.4.120 seam-placement; §11.4.69

`artifact_not_yet_built`). "Concurrent with the build" therefore means the deploy-**INDEPENDENT** stages overlap; the **CURRENT-artifact** runtime validation begins the moment that artifact exists, in parallel across devices per (B). **"Affected-stages-only re-run" is bounded to NON-huge-blocker discovery:** a **HUGE BLOCKER** (release-blocking-severity defect / regression invalidating the cycle / blast radius reaching the batch's other fixes) still triggers the §11.4.129 **STOP** → fix-

all → full-RESTART protocol (restart, never resume) — §11.4.230(A) does NOT weaken §11.4.129, §11.4.40 full-suite-retest-before-tag, §11.4.130 validate-the-fix-first, §11.4.132 risk-ordered validation, or §11.4.185 manual-QA-final; it narrows the from-scratch restart to exactly the cases those anchors require it, and parallelises everything else.

(B) PARALLEL MULTI-DEVICE TESTING + RECORDING/PROCESSING. Every available, reachable test device MUST be exercised CONCURRENTLY — one validation stream per device, never one-device-at-a-time when the devices are independent — each stream the SINGLE EXCLUSIVE OWNER of that device's exclusive resources per §11.4.119 (the owner drives; any concurrent stream targeting the same device is read-only), so per-device evidence is never cross-contaminated (§11.4.119 evidence-integrity). Every device is ALWAYS-RECORDED per §11.4.128 (non-invasive, side-effect-free, deterministic layout) and AVAILABILITY-FOLLOWED per §11.4.144 (a dropped device logs an honest offline event + reconnect-and-resume, never a silent corpus hole, never a faked continuity). Recording CAPTURE + ANALYSIS/PROCESSING MUST be fanned out through a BOUNDED WORKER-POOL (a Go worker-pool or platform equivalent) so N devices' recordings are captured + analysed in parallel and per-device evidence never serialises onto a single processor — the pool bounded by §12.6 memory + §12.12 thread headroom. Per-device PASS cites its own captured-evidence artefact (§11.4.5/§11.4.69/§11.4.107) on a CLEAN/fresh deployment (§11.4.108 runtime-signature, §11.4.200 verify-deployed-artifact-identity-per-device); the highest-risk / most-reopened cases (§11.4.132/§11.4.189) get the deepest scrutiny FIRST on every device. **HONEST BOUNDARY (§11.4.6):** genuinely-shared exclusive resources (one HDMI/audio sink, one media-playback path) serialise per §11.4.119 single-owner-over-time even across devices; "all devices concurrent" means independent devices, never two owners of one exclusive resource.

(C) FAN-OUT-SUBAGENTS-ALWAYS. Subagent fan-out (§11.4.20/§11.4.70 subagent-driven; §11.4.187 multi-track ruler; §11.4.192 auto-backfill) is the STANDING DEFAULT for EVERY genuinely-parallelizable effort — batch investigation, test-authoring, gate-authoring, per-device validation (B), per-PWU development (§11.4.58), doc/research work — not a per-request opt-in and not reserved for "big" tasks. A parallelizable effort executed monolithically in the foreground when a fan-out path exists is a §11.4.230(C) violation (the §11.4.94/§11.4.97/§11.4.103 zero-idle discipline applied to HOW work is dispatched, not only

WHETHER it runs). **HONEST BOUNDARY (§11.4.6 / §11.4.183)**: "always fan out" is BOUNDED by (1) genuine parallelizability (contending/redundant/serially-dependent subagents are NOT fan-out — they are waste, and spawning them is NOT maximization per §11.4.183 maximum-USEFUL-throughput); (2) §12.6 60% memory ceiling; (3) §12.12 `RLIMIT_NPROC` thread headroom; (4) §11.4.58 ≤ 6 concurrent-agent cap; (5) §11.4.119 single-resource-owner (two subagents may not both drive one exclusive device); (6) critical-state sequencing (commits/pushes/tags stay serial per §11.4.88/§11.4.113). The bar is maximum USEFUL parallel throughput, never agent count; fan-out that exceeds a host-safety ceiling or contends a single resource is a violation, not compliance.

Classification: universal (§11.4.17) — a platform-neutral parallelization discipline reusable by ANY project that builds an artifact, validates it across multiple targets, and records device evidence; the consuming project supplies its concrete device set + serials, build/flash tooling, recording layout, and worker-pool implementation per §11.4.35. Composes §11.4.20 / §11.4.24 (wall-clock stats prove the 4–5× target) / §11.4.40 / §11.4.58 (the PWU pipeline this binds + extends) / §11.4.69 / §11.4.70 / §11.4.88 / §11.4.89 / §11.4.94 / §11.4.96 (SAFE-during-build catalogue) / §11.4.97 / §11.4.102 / §11.4.103 / §11.4.108 (artifact/runtime layer ordering — the honest boundary) / §11.4.110 / §11.4.115 / §11.4.119 (single-resource-owner — B's load-bearing partition) / §11.4.120 (seam-placement — A's honest boundary) / §11.4.121 / §11.4.125 / §11.4.126 / §11.4.128 (always-on recording — B's evidence source) / §11.4.129 (huge-blocker full-restart — the un-weakened carve-out) / §11.4.130 / §11.4.132 / §11.4.144 (availability-following — B's drop-handling) / §11.4.183 (maximum-USEFUL-parallelism — C's honest boundary) / §11.4.185 (manual-QA-final — the un-weakened terminal gate) / §11.4.187 / §11.4.192 / §11.4.200 (per-device deployed-artifact verification) / §12.6 / §12.12. Propagation gate `CM-COVENANT-114-230-PROPAGATION` (literal `11.4.230` across the consumer fleet) + recommended gates `CM-PIPELINE-OVERLAP-VALIDATION` (a build that runs with zero deploy-independent validation stages running concurrently when SAFE-catalogue work is queued → FAIL; a NON-huge-blocker discovery that restarts the whole cycle from scratch instead of re-running only the affected stages → FAIL) + `CM-PARALLEL-DEVICE-WORKER-POOL` (available independent devices validated serially instead of concurrently, OR recording capture/analysis serialised onto one processor when a worker-pool path exists → FAIL) + `CM-FAN-OUT-SUBAGENTS-DEFAULT` (a genuinely-parallelizable effort run monolithically in the

foreground when a fan-out path exists AND all host-safety ceilings have headroom → FAIL) + paired §1.1 meta-test mutations (per §1.1; gate-code = separate work item, NOT claimed shipped §11.4.6).

Canonical authority: constitution submodule `Constitution.md` §11.4.230.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--sequential-build-then-validate ,  
--restart-whole-cycle-on-discovery , --serial-device-testing , --  
1 1 4 2 3 1  
serialise-recording-processing , --no-fan-out , --monolithic-  
parallelizable-work flag exists.
```

2 0 2 6 0 7 2 6

§ . . — Nano-precision model-tier
selection: lightest-capable-model-first with
dynamic mid-task escalation (User mandate,
– – , CRITICAL)

Verbatim operator mandate: every work effort MUST use the LIGHTEST CAPABLE model, escalating tier only as complexity rises — a nano-precision, per-task model-tier choice, NOT "everything on the heaviest model." Forensic motivation (genericised, §11.4.6): dispatching every task — a trivial doc-sync, a one-line status read, a mechanical format fix — on the heaviest available model wastes capacity + token budget + wall-clock that a lighter tier would spend identically well, while the OPPOSITE failure (a hard multi-factor root-cause investigation dispatched on a tier too light to reason through it) produces exactly the shallow, guessed, or wrong-conclusion work §11.4.6 / §11.4.102 / §11.4.194 forbid. Neither "always heaviest" nor "always lightest" is compliant; the correct discipline is a PER-TASK, EVIDENCE-BASED tier choice that tracks the task's own complexity, re-evaluated as that complexity becomes known — the nano-precision selection this anchor mandates.

(A) LIGHTEST-CAPABLE-FIRST + ESCALATE-ON-COMPLEXITY. Every dispatched work effort — main-stream or subagent (§11.4.20/§11.4.70), single-track or multitrack (§11.4.187/§11.4.192) — is assigned the LIGHTEST model tier from the closed, capability-ordered ladder `sonnet < opus` genuinely sufficient to complete it correctly, with `sonnet` the **DEFAULT working tier** and `opus` for every review and other heavy work (operator mandate 2026-09-26, CRITICAL , verbatim: "Main model to use is Sonet, reviews and other heavy stuff have to be

done by Opus instead of Fable!!! Fix this permanently on all levels!"). The review/merge-conflict-resolution substrate is a SEPARATE, already-pinned Opus-xhigh tier per clause (E), not an extra rung of this ladder; `fable` is not a rung and is not a permitted dispatch tier; `haiku` is excluded per clause (D.1). Tier selection is a DECIDABLE heuristic, not a guess (§11.4.6) — evaluated against, at minimum: (i) **task scope** (single mechanical edit / focused single-file change / multi-file-or-multi-component change); (ii) **reasoning depth** (a lookup-and-apply action vs. a multi-factor root-cause investigation per §11.4.102/§11.4.194); (iii) **blast radius** (a change touching one isolated file vs. one crossing shared state, hardware safety §11.4.133, or governance/gate seams §11.4.227); (iv) **novelty** (a previously-solved, pattern-matched task vs. a genuinely new problem with no established precedent). Closed-set tier bindings (illustrative, non-exhaustive — the consuming project supplies its own precise mapping per §11.4.35): `sonnet` (the default) — trivial/mechanical AND moderate work (simple greps/searches, doc-sync regeneration, status/field reads, format/lint fixes, boilerplate scaffolding, a focused single-component debug, a bounded single-file implementation, straightforward research or audit reads); `opus` — complex and other heavy work, including every review (architecture/design decisions, multi-factor reasoning per §11.4.194(1), cross-feature blast-radius analysis per §11.4.92 Pass 2/3, hard root-cause investigation per §11.4.102 that resists a first-pass hypothesis). A task whose complexity is genuinely UNDECIDABLE from available evidence starts at the NEXT tier UP from the best available estimate (the safe-reversible default per §11.4.101) rather than guessed downward.

(B) DYNAMIC MID-FLIGHT SWITCHING. Tier assignment is NOT a one-time, start-of-task decision — it is RE-ASSESSED continuously as the task's real complexity becomes known during execution. A task DISPATCHED as simple that PROVES harder mid-flight (an unexpected multi-factor interaction surfaces, a root cause resists the first two hypotheses, the blast radius turns out to cross a governance/hardware-safety seam) MUST ESCALATE to the next-heavier tier for its remainder — continuing on an under-powered tier past the point complexity is KNOWN is a clause-(D) false economy, not compliance. Conversely, a task DISPATCHED as heavy that DECOMPOSES into genuinely light, bounded, low-blast-radius subtasks (mechanical file edits following an already-designed pattern, doc regeneration, per-item status reads) dispatches THOSE subtasks on a

LIGHTER tier via subagent fan-out (§11.4.20/§11.4.70/§11.4.183) rather than running every subtask on the parent's heavier tier by default — capability-matching applies at EVERY dispatch granularity, not only the top-level task.

(C) ALWAYS-ON COMPOSITION, NEVER A PER-REQUEST OPT-IN. Nano-precision tier selection is a STANDING DEFAULT mechanism, engaged automatically for EVERY work effort from the first prompt of a session onward (mirrors §11.4.126 default-autonomous-loop / §11.4.198 default-mechanisms-always-on) — applied TOGETHER WITH, never as a substitute for: §11.4.141 (heavy token-optimization — a lighter model dispatched with a bloated context is not the efficiency this anchor targets); §11.4.183 (maximum-USEFUL-parallelism — tier-matching each parallel stream, not merely running more streams on one fixed tier); §11.4.187/§11.4.192 (the multi-track ruler orchestration + continuous auto-backfill — each track's worker is tier-assigned per this anchor, not defaulted to one blanket tier project-wide); and §11.4.196/§11.4.198 (native-alias-first priority + always-on default mechanisms — this anchor governs WHICH TIER within whichever alias/account is selected, composing with, never replacing, §11.4.196's ALIAS-layer priority). None of these five composed anchors is weakened, duplicated, or re-specified by §11.4.231 — this anchor adds the TIER axis to their existing ALIAS/PARALLELISM/TOKEN axes.

(D) HONEST BOUNDARY (§11.4.6) — "LIGHTEST CAPABLE" MEANS CAPABLE, NOT CHEAPEST-REGARDLESS. "Lightest capable" is a CONJUNCTION, not a discount: a tier choice that fails to actually complete the task correctly — producing a shallow investigation, a guessed root cause, an under-verified fix, or ANY form of bluff per the §11.4 anti-bluff covenant — is NOT "lightest capable," it is UNDER-POWERED, and under-powering a task that then fails, bluffs, or requires costly rework is a FALSE ECONOMY and a §11.4 violation of equal severity to any other PASS-bluff (the token/time saved at dispatch is spent many times over in re-investigation, re-review, and the risk of a shipped defect). Capability determinations are made from the clause-(A) decidable heuristic + observed task outcomes — EVIDENCE-BASED, never guessed (§11.4.6): "this always works fine on `sonnet` " without a captured basis is exactly the guess §11.4.6 forbids, and a tier's genuine capability for a task class is established by the same anti-bluff discipline (§11.4.5/§11.4.69/§11.4.107) that governs every other claim in this constitution. The ladder binds the DISPATCHER's tier ASSIGNMENT: a dispatcher MUST NOT assign any tier outside `sonnet < opus` — never `fable` (2026-09-26 mandate), and never `haiku` where clause (D.1) applies. A serving

model the platform selects for a turn WITHOUT a dispatcher assignment is recorded honestly in the §11.4.182 label (whatever id the platform reports, never re-labelled); it is neither a dispatch-tier violation by the dispatcher nor a licence to assign that model, and review/merge-conflict scopes run on Opus-xhigh per clause (E) regardless.

(D.1) haiku PROHIBITED WHERE IT CANNOT RECEIVE ITS GOVERNING CONTEXT — measured fact (2026-09-26). In the reference consuming repository a haiku subagent dispatch failed before doing any work with "Prompt too long (~421k tokens vs 200k limit)" (request id req_011CfSBXm7Kn6enyUT1UEaWd , recorded in that repository's agent registry docs/requests/agent_registry.jsonl): the governance context every dispatched agent must receive there exceeds the haiku tier's context window, so a haiku dispatch cannot carry its binding context and haiku is PROHIBITED as a dispatch tier in that repository. A consuming project MAY re-admit haiku as a rung below sonnet ONLY on its OWN captured evidence (§11.4.6/§11.4.35) that a haiku dispatch in that project receives its full governing context and completes its task correctly — never by assumption, and never by trimming the governing context to make it fit. A haiku dispatch that fails this way is recorded honestly as a failed dispatch (§11.4.147), never silently retried on the same tier.

(E) HARD PINS OVERRIDE THIS DEFAULT FOR THEIR SCOPES. §11.4.231 sets the DEFAULT working tier for NON-review, non-merge-conflict-resolution work. It does NOT relax, override, or compete with the two existing hard model-substrate pins: §11.4.209 (every mandatory independent code-review — §11.4.125/§11.4.142/§11.4.194 iterated to GO per §11.4.134 — runs on **Opus at xhigh effort**, ALWAYS — no Fable, no fallback model; genuine Opus-unavailability BLOCKS the review, NEVER a lower tier, another model, or a lower effort) and §11.4.211 (every main → feature/product/flavor merge-CONFLICT resolution runs on the SAME Opus-xhigh substrate, blocked — never substituted — on unavailability). Those two anchors are STRICTER, SCOPE-SPECIFIC pins that this anchor's general ladder does not reach — review and merge-conflict-resolution are NOT eligible for sonnet/opus-by-complexity dispatch; they run on their pinned substrate ALWAYS, regardless of how "simple" the reviewed change or the conflict appears. A review or merge-conflict resolution dispatched on anything other than the §11.4.209/§11.4.211 Opus-xhigh pin is a violation of THOSE anchors, not a valid application of §11.4.231's general ladder.

(F) EFFORT-TIER SELECTION — the SAME nano-precision discipline applied to the reasoning-EFFORT axis (operator IMPORTANT mandate, 2026-07-26; verbatim: "Besides dynamic determining proper models to use, we **MUST** decide the effort as well and switch between the effort settings same as we do now for chosen model! Labeling **MUST BE** extended for this details as well. Example: (T1/main - claude4 - opus - high)."). Every dispatched work effort is assigned not only a model tier per clauses (A)-(E) but ALSO the LIGHTEST reasoning-EFFORT setting from the closed ladder

low < medium < high < xhigh < max genuinely sufficient to complete the task correctly — chosen by the SAME clause-(A) decidable heuristic (scope / reasoning-depth / blast-radius / novelty), NOT a guess (§11.4.6). Illustrative bindings (non-exhaustive; the consuming project supplies its precise mapping per §11.4.35):

low / medium — trivial / mechanical work (doc-sync regeneration, status / field reads, format / lint fixes, boilerplate); high — moderate-to-complex reasoning (focused single-component debug, multi-factor analysis per §11.4.102 / §11.4.194); xhigh / max — the HARDEST verification / merge-conflict / architecture work.

(F.1) DYNAMIC MID-TASK ESCALATION — effort, like model tier per clause (B), is RE-ASSESSED continuously as the task's real complexity becomes known; a task dispatched at a lower effort that PROVES harder mid-flight **MUST ESCALATE** its effort for the remainder (continuing under-effort past the point complexity is KNOWN is the clause-(D) false economy), and a heavy task decomposing into light bounded subtasks dispatches THOSE at a lighter effort — effort-matching applies at EVERY dispatch granularity, not only the top-level task. (F.2) HONEST HARNESS BOUNDARY (§11.4.6, load-bearing — do NOT overstate the mechanism) — effort-tiering is REAL where the execution path exposes an effort control and a DOCUMENTED CAPABILITY-GAP where it does not: the current Claude Code **Agent tool** (subagent dispatch) exposes a `model` parameter but **NO effort parameter**, so a subagent's effort is NOT presently settable via that path (the §11.4.182 `<effort>` label field degrades honestly to `?` there); the

Workflow tool's `agent()` DOES take an `effort` argument, so effort IS settable there — the label reports whichever path is live, never fakes one. This is the SAME aspirational-vs-harness state the clause-(E) §11.4.209 / §11.4.211 Opus- xhigh pins already occupy: the pins mandate the effort, and whether a dispatch path can SET it is a harness capability reported honestly, never faked (§11.4.182). (F.3) HARD-PIN FLOOR PRESERVED — clause (E)'s hard pins are ALSO effort pins: every §11.4.209 mandatory independent code-review AND every §11.4.211 merge-

conflict resolution runs at `xhigh` effort on Opus (no fallback model), and this general effort ladder NEVER tiers those two scopes below `xhigh` regardless of how "simple" the reviewed change or the conflict appears — a review / merge-conflict resolution run below `xhigh` is a §11.4.209 / §11.4.211 violation, not a valid application of this ladder. (F.4) HONEST BOUNDARY (§11.4.6) — "lightest capable effort" is a CONJUNCTION, not a discount: an effort setting that yields a shallow / guessed / under-verified result is UNDER-POWERED, a FALSE ECONOMY, and a §11.4 violation of PASS-bluff severity, exactly as an under-powered MODEL tier is per clause (D); capability determinations are evidence-based, never guessed (§11.4.6).

Classification: universal (§11.4.17) — a platform-neutral model-capacity-matching discipline reusable by any project dispatching work across a tiered model roster; the consuming project supplies its own precise tier ↔ task-class mapping and its own model-name ladder per §11.4.35 (the `sonnet` / `opus` ladder with Sonnet as the default, the clause-(D.1) `haiku` prohibition, and the Opus-`xhigh` pins are this project's current concrete instantiation, not a universally-fixed vocabulary — a consumer with a differently-named tier roster maps its own ladder onto the SAME lightest-capable-first + dynamic-escalation + always-on-composition + honest-boundary + hard-pin-override structure, re-admitting a lighter rung only on its own captured evidence per clause (D.1)). Composes §11.4.6 (no-guessing — capability determinations are evidence-based) / §11.4.20 / §11.4.70 (subagent-driven dispatch — the granularity at which tier-matching applies) / §11.4.92 (blast-radius analysis feeds the clause-(A) heuristic) / §11.4.101 (safe-reversible default on undecidable complexity — escalate up, never guess down) / §11.4.102 / §11.4.194 (reasoning-depth signals for the heuristic) / §11.4.126 / §11.4.141 / §11.4.183 / §11.4.187 / §11.4.192 / §11.4.196 / §11.4.198 (the always-on composition set, clause (C)) / §11.4.209 / §11.4.211 (the hard-pinned review/merge-conflict substrate this anchor does NOT touch, clause (E)). Propagation gate `CM-COVENANT-114-231-PROPAGATION` (literal `11.4.231` present as a block-start, exactly-once per governance file, lockstep content-hash equality across the consumer fleet per §11.4.227(B)) + recommended gate `CM-MODEL-TIER-LIGHTEST-CAPABLE` (a dispatched effort whose logged/declared model tier exceeds the clause-(A) heuristic's indicated tier for a task later confirmed trivial-and-successful-at-the-lighter-tier is a WARN/candidate finding — never a hard FAIL on tier choice alone, since tier-adequacy is only provable in retrospect per clause (D); the gate instead HARD-FAILS the two decidable violations: (i) a review or merge-conflict-resolution

effort NOT run on the §11.4.209/§11.4.211 pinned substrate, and (ii) a task whose complexity was PROVEN mid-flight [an escalation event was logged, or a downstream failure/bluff was traced back to it] that did NOT escalate tier before continuing, and (iii) a dispatch assigned `fable`, or assigned `haiku` where clause (D.1) applies) + paired §1.1 mutation (gate-code = separate work item, NOT claimed shipped — §11.4.6; mutation plan: (1) dispatch a code-review labeled as running on `sonnet` instead of Opus-xhigh → gate MUST FAIL; (2) inject a logged mid-task complexity-escalation event with NO corresponding tier-escalation record → gate MUST FAIL; (2b) record a dispatch assigned `fable`, or assigned `haiku` where clause (D.1) applies → gate MUST FAIL; (3) golden-FALSE fixture: a task that stays on its initially-assigned tier throughout, with no escalation event logged → gate MUST NOT fire, per §11.4.201(1) the false-positive guard).

Canonical authority: constitution submodule `Constitution.md` §11.4.231.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--always-heaviest-model , --always-lightest-model ,
--skip-complexity-reassessment , --fixed-tier-per-session , --relax-
1 1 4 2 3 2
review-pin , --relax-merge-conflict-pin , --per-request-tier-opt-in
flag exists.
```

§ . . — Anti-mess long-op orchestration: every long-op is registered, single-owned-per-purpose, liveness-PROVEN, handed-off-on-stop, safely-reaped, consistency-checked, and escape-bounded — the single source of truth for what is running, never inferred (research-derived, — —)

Forensic anchor (genericised): in one cycle a verification sweep produced five distinct mess states — (1) multiple concurrent same-purpose sweeps ran unintentionally (no single-owner-per-purpose claim); (2) a sweep HUNG on a consumer-blocked-on-an-unclosed-producer pipeline and was undetected because liveness was inferred from process-alive (a §11.4.201(6) FALSE-NULL: a wedged op and a progressing op returned the identical "process alive" signal); (3) stopping the owning subagent SILENTLY KILLED its child long-op because ownership was implicit (no handoff-on-stop); (4) a stale/non-growing log was read as live progress (the same false-null, reader-side); (5) an un-completable tooling gate blocked a

release for hours (no escape to a terminal state). Per §11.4.102 Phase 4.5 these are ONE architectural state-management flaw, not five bugs: **there was no single source of truth for what long-op is running, who owns it, its state, and its PROVEN liveness** — the §11.4.196(F) configured-in-use gap generalised to project long-ops, so every mess state was un-observable, and an un-observable bad state is a §11.4/§11.4.1 bluff at the orchestration layer.

The mandate — a **long-op orchestration layer** binding every long-op (any op expected to exceed the runtime background-execution threshold per §11.4.89 — verification sweep, meta-test sweep, build, flash/deploy, test/QA suite, dispatched subagent), ALL of which hold:

(A) SINGLE SOURCE OF TRUTH — the long-op lifecycle registry. Every long-op MUST register BEFORE it runs and reach a terminal state (`complete|failed|reaped|handoff|blocked-escape`) when it ends, in a durable (never-tmpfs, write-temp-then-rename, §11.4.116 append-only event stream + atomic snapshot) registry recording `{op_id, purpose_key, owner, pid|pgid, cwd, state, last_heartbeat, heartbeat_seq, progress_offset, log_path, verdict, evidence_path}` . Success is read from the op's terminal verdict/ `result` , NEVER its process exit code. A long-op running with no registry record — or a record with no live op — is a FALSE-READY drift state.

(B) SINGLE-OWNER-PER-PURPOSE. No two same- `purpose_key` long-ops run at once; acquisition is a probe-race-safe claim keyed on `purpose_key` (§11.4.119). A live claim forces exactly ONE of REUSE (attach) or explicit SUPERSEDE (graceful handoff of the running owner, logged), never a silent second run.

(C) LIVENESS TRUTH-SOURCE. Progress is PROVEN, never inferred from process-alive: every long-op emits a monotonic heartbeat and/or an advancing progress offset; a watchdog reads the DELTA over a window and flags HUNG on no-advance past the op's declared no-progress budget (the §11.4.201(6)-(7)/ §11.4.89 discipline — `kill -0` is a necessary-not-sufficient pre-filter, a stale heartbeat/non-growing log is NEVER "progressing"). A consumer blocked on an unclosed producer write-end (§11.4.201(12)) is caught by the flat heartbeat regardless of internal state.

(D) OWNERSHIP-HANDOFF-ON-STOP. Stopping/killing an owner MUST hand its long-ops to the registry (`handoff` , re-adoptable per §11.4.147) or deliberately reap them (`reaped` , reason logged) — NEVER silently orphan or accidentally cascade-kill; long-ops launch detached (own process-group) so an owner stop does not cascade-signal them, and the stop path records each owned op's disposition. The loop done-condition is not satisfied while a `handoff` op is un-adopted.

(E) SAFE REAPING. An orphaned/hung op is reaped ONLY on PROVEN staleness — holder PID dead (`kill -0` fails) OR no-progress-budget exceeded per (C) — resolved by the process's real `/proc/<pid>/cmdline` (§11.4.196(D)/§12.12), never a bare `pgrep` /substring that could match a carrier or the reaper itself; a LIVE, ADVANCING op is NEVER reaped (§9.2); every REAPED/KEPT decision + resolved evidence is logged (§11.4.180).

(F) CONSISTENCY INVARIANT SWEEP. A fast invariant check runs before every gated transition (build/flash/tag) and on a standing cadence, asserting: no duplicate owner per purpose, no orphan, no stale lock (§11.4.180), no hung op (C), no registry-vs-reality drift (live-long-op set == registry live rows — the §11.4.196(F) check generalised). Any violation is surfaced (named, actionable) and the gated transition is REFUSED until resolved. The sweep is itself a guard (§11.4.201): golden-TRUE + golden-FALSE-with-carrier fixtures, conservative-safe REFUSE on an unresolvable signal, resolved evidence printed on every refusal.

(G) UN-COMPLETABLE-BLOCKER ESCAPE. A gate that cannot reach a terminal verdict MUST degrade, after a bounded budget, to an explicit evidence-backed decision — NEVER an indefinite silent block: §11.4.102 first (hung/(C) vs duplicate/(B) vs wrong-seam/§11.4.120 re-seat-at-activation); NOT-YET-RUNNABLE → §11.4.69 `artifact_not_yet_built` (non-blocking on an intermediate artifact, blocking only on the release candidate per §11.4.135); genuinely un-completable → §11.4.101 (autonomous reversible SKIP-with-reason + tracked §11.4.197 item, ELSE a bounded §11.4.66 operator decision) — the loop keeps progressing every other non-blocked item, it does not idle, and the escape NEVER fabricates a PASS.

Honest boundary (§11.4.6). This layer reduces orchestration mess; it does NOT make product bugs impossible, does NOT make a gate correct (§11.4.108/ §11.4.201 still bind each gate), and does NOT remove the operator from genuine blocks — it converts an indefinite silent stall into a bounded, evidence-backed, operator-visible decision. A heartbeat proves ADVANCE not CORRECTNESS; the registry proves what CLAIMS to run, (F) proves registry==reality. **Reuse-not-**

rewrite: it is the multitrack per-track agent registry (§11.4.147/§11.4.116/§11.4.187) generalised from per-track workers to any long-op and bound through, per §11.4.196(F) — no parallel mechanism is invented.

Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their long-op set, purpose keys, no-progress budgets, and stop-hook wiring as DATA per §11.4.35. Composes §11.4.89 / §11.4.101 / §11.4.116 / §11.4.119 / §11.4.120 / §11.4.135 / §11.4.147 / §11.4.180 / §11.4.187 / §11.4.196(D) (F) / §11.4.201(1)(6)-(12) / §11.4.207 / §11.4.69 / §12.6 / §12.12 / §9.2. Propagation gate `CM-COVENANT-114-232-PROPAGATION` (literal `11.4.232`) + recommended mechanism gates `CM-LONGOP-REGISTRY-SINGLE-SOURCE` (A) / `CM-LONGOP-SINGLE-OWNER-PER-PURPOSE` (B) / `CM-LONGOP-LIVENESS-HEARTBEAT` (C) / `CM-LONGOP-OWNERSHIP-HANDOFF-ON-STOP` (D) / `CM-LONGOP-SAFE-REAP-PROVEN-STALE` (E) / `CM-ORCHESTRATION-CONSISTENCY-SWEEP` (F) / `CM-UNCOMPLETABLE-BLOCKER-ESCAPE` (G) + paired §1.1 mutations (run two same-purpose sweeps with no reuse/supersede → (B) FAILs; infer liveness from `kill -0` alone / drop the heartbeat-delta check → (C) FAILs; a stop that leaves an owned op with no terminal/handoff record → (D) FAILs; reap a live-advancing op OR reap by bare- `pgrep` substring → (E) FAILs; a live long-op with no registry row passes the consistency sweep → (F) FAILs; an un-completable gate blocks with no escape/decision → (G) FAILs; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6). No escape hatch — no `--allow-duplicate-longop` , `--infer-liveness-from-pid` , `--silent-orphan-on-stop` , `--reap-live-op` , `--skip-consistency-sweep` , `--block-without-escape` , `--longop-without-registry` flag.

EXTENSION landed 2026-07-26 — §11.4.182 model-in-label + flawless alias/provider ↔ model consistency + real-time always-in-sync (User mandate, 2026-07-26). Verbatim operator mandates (2026-07-26): "Extend our logging to present models used since we will now have dynamic models choice and switching! From (example) (T1/main - claude4) into: (T1/main - claude4 - model_used), model MUST be dynamically obtained so we present name equivalents - short names - like: haiku, opus, sonnet, fable. However we MUST fully support all Claude Toolkit providers aliases, so real model used names MUST BE obtained in this case(s) as well! Examples: (T1/main - claude4 - haiku), (T1/main - claude4 - fable), (T1/main - claude4 - opus), (T1/main - claude4 - deepseek), (T1/main - claude4 - xiaomi) ... Labels MUST BE everywhere ALWAYS up to date in real time and in sync!" and "Adding model names into labels considers that providers names / aliases will be

up to date too and in relation to model used since models belong to providers / aliases !!! Consistency MUST BE flawless!" This EXTENDS the §11.4.182 label discipline (it does NOT re-mint the anchor — the literal 11.4.182 and its propagation gate CM-COVENANT-114-182-PROPAGATION are UNCHANGED and stay GREEN; this carrier block cites §11.4.182, it is not a second §11.4.182 block-start per §11.4.227(B)). (1) **LABEL FORM EXTENDED** — the mandated §11.4.182 label is extended from (T<N>/<branch> - <alias>) to (T<N>/<branch> - <alias> - <model> - <effort>) (the trailing <effort> field added 2026-07-26 per §11.4.231 clause (F) — see the EFFORT-FIELD clause (7) below), <model> recording which model is CURRENTLY answering for that alias: for a native Claude alias, the SHORT name (opus / sonnet / haiku / fable, case-insensitive-substring-derived from the model id the agent platform reports for the most recent turn — DYNAMIC, obtained live at label-emission time, never a fixed/assumed value since native aliases have no persistent per-alias model and the serving model can legitimately switch mid-session per §11.4.231); for a claude-toolkit PROVIDER alias (deepseek, xiaomi, kimi-for-coding, opencode, ...) the REAL provider model id as declared in that provider's host-config registry (the SAME value the alias's own launch mechanism uses — a single source of truth, never re-derived). (2) **MODEL-FIELD DERIVATION (§11.4.6, never guessed)** — a closed two-case split on the SHAPE of the already-derived <alias> : (a) native-alias → read the model id from the CURRENT session's own transcript (the platform's per-project per-session transcript; LAST assistant-turn's reported model → short name); (b) provider-alias → read the model id from the host multi-account-toolkit provider registry (CMA_PROVIDER_MODEL). Either lookup that cannot resolve yields the honest ? — NEVER a fabricated or last-known value; an id obtained but matching none of the four native short-names is printed VERBATIM (real-but-unrecognised is more honest than a false ?). The 3-field legacy form (no <model>) remains a VALID, ACCEPTED label during migration — the enforcing guard-hook LABEL_RE admits both forms, and NEVER treats an honest ? model, or an absent model field, as a block-cause. (3) **FLAWLESS ALIAS/PROVIDER ↔ MODEL CONSISTENCY INVARIANT (operator NOTE, load-bearing)** — the <alias> / provider field and the <model> field MUST be MUTUALLY CONSISTENT AT ALL TIMES: the model shown MUST be a model that genuinely BELONGS to the alias/provider shown (models belong to providers/aliases), and the alias/provider name MUST itself be current — never a stale alias beside a fresh model, never a model that does not belong to the shown alias. This is guaranteed BY CONSTRUCTION,

not by coincidence: the native model is read from THAT alias's own live session transcript, and the provider model is read from THAT provider's own registry entry — so the derived `<model>` is definitionally a model of the shown `<alias>` ; a derivation that ever emits a model not belonging to the shown alias is a §11.4.6 / §11.4.182 consistency violation, never a valid label. (4) **REAL-TIME, ALWAYS-IN-SYNC (§11.4.229/§11.4.106(F))** — the label is STATELESS and RE-COMPUTED at EVERY emission (every human-invoked labeler run AND every synchronous PreToolUse guard-hook firing on every Agent/Task/TaskCreate dispatch) against the LIVE environment (current `CLAUDE_CONFIG_DIR` , current session id, current transcript tail, current provider `.env`) — so a mid-session model switch (`/model` , an autonomous §11.4.231 escalation) is reflected in the VERY NEXT emission with no cache to go stale: real-time BY CONSTRUCTION, not by a sync daemon. Honest boundary (§11.4.6): the live label has NO persisted artifact that could drift, so the Docs-Chain (§11.4.106) sync engine — which propagates PERSISTED documents/DBs — does not apply to the stateless label itself; the operator's "use/extend Docs Chain for real-time responsiveness" ask attaches to any PERSISTED label-history / model-switch LOG surface (a §11.4.197 follow-up to be tracked), which, if adopted, IS §11.4.106-Docs-Chain-synced and §11.4.208-request-history-adjacent — never faked as already-shipped. (5) **HONEST SUBAGENT-INTROSPECTION LIMIT (§11.4.6/§11.4.197)** — the native-alias transcript lookup is proven reliable when invoked from the DISPATCHING process's own shell (the guard-hook's actual execution context at dispatch time); it is NOT proven reliable when a background-dispatched subagent introspects its OWN model from inside its own subprocess on hosts where the subagent shares its parent's session-transcript with no subagent-vs-conductor marker — recorded as an open limitation to be tracked, never silently claimed solved. (6) **DIVERGENT-LABELER RECONCILIATION (§11.4.227)** — the LIVE labeler wired via `.claude/settings.json` is `constitution/scripts/multitrack/track_branch_label.sh` ; the top-level `scripts/multitrack/` path is a thin pass-through WRAPPER that `exec` s that canonical copy (§11.4.177 inherited-by-reference), so it mirrors nothing, cannot drift, and emits byte-identical output by construction. The reconciliation LANDED in commit `a48e919f52d` and is gate-enforced at `pre_build_verification.sh:41911-41915` (the gate detects the `exec` wrapper and follows it to the canonical implementation), never a silent divergence. (7) **EFFORT FIELD (added 2026-07-26 per the operator IMPORTANT mandate — verbatim: "Besides dynamic determining proper models to use, we MUST**

decide the effort as well and switch between the effort settings same as we do now for chosen model! Labeling **MUST BE** extended for this details as well. Example: (T1/main - claude4 - opus - high)."). The 5-field label appends `<effort>` after `<model>`, recording which reasoning-EFFORT setting is CURRENTLY in force for that work-stream, drawn from the closed set `{low | medium | high | xhigh | max}` (the SAME effort ladder §11.4.231 clause (F) selects). (a) DERIVATION (§11.4.6, never guessed) — `<effort>` is read from whatever LIVE effort signal the execution path exposes (the Workflow-tool `agent()` `effort` argument, a future Agent-tool effort parameter, an env/config effort setting); a path exposing NO effort signal, OR any value outside the closed set, yields the honest `?` — EXACTLY like the `<model>` field's `?` fallback, NEVER a fabricated / assumed / last-known value. (b) HONEST HARNESS BOUNDARY (§11.4.6, load-bearing — do NOT overstate the mechanism) — effort-tiering is REAL where the execution path exposes it and a DOCUMENTED CAPABILITY-GAP where it does not: the current Claude Code **Agent tool** (subagent dispatch) exposes a `model` parameter but **NO effort parameter**, so a subagent's effort is NOT currently settable via that path and its `<effort>` field degrades honestly to `?`; the **Workflow tool's** `agent()` DOES take an `effort` argument, so effort IS settable + derivable there — the label reports whichever path is live, never invents one. This is the SAME aspirational-vs-harness state the §11.4.209 / §11.4.211 Opus- `xhigh` review/merge pins already occupy: the pins mandate the effort, and whether a given dispatch path can SET it is a harness capability the label reports honestly and NEVER fakes. (c) LEGACY-FORM ACCEPTANCE (migration) — the 3-field `(T<N>/<branch> - <alias>)` and 4-field `(T<N>/<branch> - <alias> - <model>)` label forms remain VALID and ACCEPTED; the guard `LABEL_RE` admits all three shapes and NEVER treats an absent `<effort>` field, or an honest `?` effort, as a block-cause (§11.4.6 — mirroring the §11.4.182 track/alias/model honest-`?` rule). (d) CONSISTENCY — the effort field is selected TOGETHER WITH the model by §11.4.231 (both tiers chosen for one task) and never contradicts it; a stale effort beside a fresh model, or an effort outside the closed set, is a §11.4.6 / §11.4.182 consistency violation, never a valid label. (e) REAL-TIME, ALWAYS-IN-SYNC (§11.4.229) — like `<model>`, `<effort>` is STATELESS and RE-COMPUTED at EVERY label emission against the LIVE environment, so a mid-task §11.4.231 effort escalation is reflected in the VERY NEXT emission with no cache to go stale. Propagation: `CM-TRACK-BRANCH-LABEL` is EXTENDED to also assert the labeler emits a parseable `<model>` field

AND a parseable `<effort>` field, the guard `LABEL_RE` accepts the 3-field / 4-field / 5-field label forms (legacy 3-field and 4-field ALWAYS remain accepted during migration; a `?` effort is NEVER a block-cause), and the alias/provider ↔ model consistency holds; paired §1.1 mutation: strip the `<model>` derivation OR emit a model not belonging to the shown alias → the extended gate FAILs (gate-code = separate work item, NOT claimed shipped §11.4.6). No new escape hatch — the existing §11.4.182 "no escape hatch" list applies to the extended label unchanged.

§ . . — Standing anti-mess control plane: a level-triggered, idempotent, invariant-catalogue reconciliation loop over the WHOLE persistent System state, orchestrating every per-domain enforcer, refusing gated transitions on un-reconciled drift (research-derived, — —)

Forensic anchor (genericised): every existing consistency enforcer in a multi-track / multi-agent / multi-submodule / multi-upstream System is EDGE-TRIGGERED — a PreToolUse hook fires on an action, a check runs at a seam (lock-acquisition / merge-time / build-time / commit-time), a gate runs when a build runs — so a mess that arises WITHOUT an edge firing is un-observed until it causes harm (two live examples captured 2026-07-26: a large uncommitted-batch pile-up sat unreconciled through a session, and multiple provably-stale VCS/wrapper locks were held by dead holders with no acquisition attempt to reap them — the §11.4.205 "a hook never installed / a check that runs only when asked" forensic generalised to the whole persistent state). Per §11.4.102 Phase 4.5 the ~19 persistent-state mess classes (Plane R repository: uncommitted pile-up, stale/duplicate locks, orphaned worktrees, half-applied cross-repo commits, drifted submodule pointers, partial merges, divergent mirrors, mis-placed work; Plane G governance: divergent mirrors, pointer drift, gate-ledger drift, anchor-number collision; Plane S tracking: doc/DB drift, cross-track SSoT divergence, ledger drift, handoff-doc drift; Plane P runtime: the §11.4.232 long-op set + un-reaped streams + dangling review state) are ONE architectural flaw — there is no LEVEL-TRIGGERED control plane that re-derives the actual persistent state and reconciles it against a declared clean state; an un-observed drift is a §11.4/§11.4.1 bluff at the organization layer ("the System is clean" while a mess silently accumulates).

The mandate — a standing **anti-mess control plane** (the desired-state/actual-state reconciliation-loop + control-plane/data-plane pattern of distributed systems), ALL of which hold:

(A) DECLARED INVARIANT CATALOGUE. A single machine-readable consumer-owned catalogue (DATA per §11.4.28/§11.4.35) enumerating every consistency invariant across the persistent planes, each carrying `{id, plane, desired-condition, detector, reconcile-class ∈ {auto-safe | operator-gated | owned-by <anchor>}, evidence-shape}` ; the catalogue IS the desired clean state; an un-catalogued mess class is an honest tracked GAP (§11.4.6/§11.4.197), never silently claimed covered.

(B) LEVEL-TRIGGERED LOOP. The loop re-derives the ACTUAL state (not an event) and diffs it against the catalogue BEFORE every gated transition (commit-batch/build/flash/merge-to-trunk/tag/constitution-pull) AND on a standing cadence (§11.4.94/§11.4.103) — catching drift regardless of how it arose (the Kubernetes level-triggered property); actual-vs-desired uses a §11.4.86 fingerprint diff drilling to the offending record (anti-entropy), never a whole-scan where a fingerprint suffices.

(C) IDEMPOTENT RECONCILE-OR-REFUSE. Each drift takes exactly one dedup-keyed action: RECONCILE-AUTO-SAFE (reversible + evidence-determinable + bounded-blast per §11.4.101 — reap a provably-dead lock §11.4.180, regen a stale derived doc §11.4.106, scope-group-commit a placement-clean batch §11.4.191, retire a merged worktree §11.4.167), idempotent (a no-op on a clean state); OR REFUSE-THE-GATE + a bounded operator decision (irreversible/high-blast/indeterminate — id-colliding SSoT §11.4.206, anchor collision §11.4.227, partial cross-repo merge, divergent mirror) per §11.4.101/§11.4.66; a gated transition NEVER proceeds while a catalogued invariant is violated and un-reconciled; the loop keeps progressing every other non-blocked item (park-one-not-the-loop).

(D) SAGA FORWARD-RECONCILE (no-loss). A partially-applied cross-repo operation is reconciled complete-forward under a §9.2 pre-op backup (finish the remaining legs / bump the pointer / complete the merge zero-markers-zero-drop §11.4.41 step 3 / ff the lagging mirror) — NEVER force-push/reset/ -s ours (§11.4.113/§9.2); compensation is §9.2 backup-restore only; the saga reaches fully-completed, never rolled-back-with-loss.

(E) THE LOOP IS A GUARD (§11.4.201). Every detector ships golden-TRUE + golden-FALSE-with-carrier fixtures (§11.4.107(10)); a false-positive REFUSAL of a clean state is a §11.4.201(1) FAIL-bluff exactly as a false-negative PASS is a §11.4 PASS-bluff; every "no drift" null is control-needle-proven (§11.4.201(6)-(12)); conservative-safe REFUSE on an unresolvable signal, resolved evidence printed on every refusal.

(F) CONTROL-PLANE / DATA-PLANE SEPARATION. The control plane DETECTS + RECONCILES state and NEVER performs the data-plane work (the actual build/commit/merge/long-op — §11.4.232 + the build/flash/commit wrappers); the two never merge; the control plane's own state is durable (§11.4.205(6) temp → fsync → rename, never tmpfs, may reuse §11.4.207).

Honest boundary (§11.4.6). Converges the CATALOGUED invariants; does NOT make an un-catalogued mess class impossible (extensible catalogue, never "provably mess-free"), does NOT make a gate correct (§11.4.108/§11.4.201), does NOT auto-fix irreversible drift (operator-gated §11.4.101/§11.4.66), and is eventual-consistency-at-every-gate-and-cadence NOT a real-time daemon (§11.4.205(7)).

Reuse-not-rewrite: it is the level-triggered COMPOSER over the existing per-domain enforcers (§11.4.180/§11.4.186/§11.4.106/§11.4.191/§11.4.206/§11.4.227/§11.4.232/§11.4.147/§11.4.167), inheriting each's detector — no parallel enforcer is invented (§11.4.227). It composes §11.4.232 (the RUNTIME data-plane long-op orchestrator — §11.4.232 clause (F) is Plane P of this catalogue) as the plane it does not touch.

Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their invariant catalogue, plane detectors, gated-transition list, cadence, and reconcile-class bindings as DATA per §11.4.35. Composes §11.4.6 / §11.4.26 / §11.4.28 / §11.4.32 / §11.4.84 / §11.4.86 / §11.4.101 / §11.4.106 / §11.4.108 / §11.4.116 / §11.4.121 / §11.4.147 / §11.4.164 / §11.4.167 / §11.4.176 / §11.4.179 / §11.4.180 / §11.4.181 / §11.4.186 / §11.4.191 / §11.4.196(F) / §11.4.201 / §11.4.205 / §11.4.206 / §11.4.207 / §11.4.227 / §11.4.232 / §12.6 / §9.2 / §2.1. Propagation gate CM-COVENANT-114-233-PROPAGATION (literal 11.4.233) + recommended mechanism gates CM-ANTIMESS-INVARIANT-CATALOGUE (A) / CM-ANTIMESS-LEVEL-TRIGGERED-LOOP (B) / CM-ANTIMESS-RECONCILE-OR-REFUSE (C) / CM-ANTIMESS-SAGA-FORWARD-NO-LOSS (D) / CM-ANTIMESS-LOOP-IS-A-GUARD (E) / CM-ANTIMESS-CONTROL-DATA-PLANE-SPLIT (F) + paired §1.1 mutations (drop the standing-cadence trigger so the loop is edge-only

→ (B) FAILs on the pile-up/stale-lock fixture; let a gated transition proceed on a catalogued un-reconciled drift → (C) FAILs; a saga that force-pushes/resets to reconcile a mirror → (D) FAILs; a detector that REFUSES a clean golden-TRUE fixture, or PASSes a golden-FALSE-with-carrier → (E) FAILs; the control plane invokes a data-plane build/commit → (F) FAILs; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227). No escape hatch — no `--edge-triggered-only` , `--skip-consistency-catalogue` , `--proceed-on-drift` , `--reconcile-with-force-push` , `--unvalidated-detector` , `--control-plane-does-work` , `--assume-clean-without-reconcile` flag.

(G) DEPENDENCY-POINTER PREFLIGHT — stale, absent, and

UNFETCHABLE are three distinct mess classes, and all three fail closed

(extension, 2026-08-20). Clause (A)'s Plane-R catalogue names *drifted submodule pointers*; drift is only one of three ways a multi-repo build ships code its own source tree does not contain. The AI-curriculum corpus (module 35, table row 5

"Submodule-pointer-not-bumped-so-fix-reverts") states the failure exactly: *"The fix is real and merged in a submodule, but the top-level image ships the old pointer, so the bug returns from a clean checkout"* — root cause: *"A multi-repo build where 'the fix is in the submodule' and 'the image ships the fix' are two different facts, only one of them checked."* Its prescription is a HARD gate, not a warning: *"A hard preflight gate that fails the build when a governance/component submodule its correctness depends on is at a stale pointer or absent — 'the fix is in the submodule but the image shipped old code' must be impossible, not caught after."* The catalogue is therefore EXTENDED with all three classes and the preflight that closes them: **(i)**

STALE — the parent's recorded pointer is behind the submodule's intended ref (the drift clause (A) already names); **(ii) ABSENT / UNINITIALISED** — the submodule is declared but not checked out, so a build reads an empty directory as a satisfied dependency; **(iii) UNFETCHABLE** — the recorded pointer resolves on NO configured remote (the `not our ref` condition), so the tree builds for whoever already has the object locally and is structurally un-clonable for everyone else.

Class (iii) is the most dangerous because it is INVISIBLE in the checkout that created it and fails only on a fresh clone. The preflight's four mechanical steps, verbatim from the corpus: *"resolves every submodule pointer, verifies it matches the intended ref, fails closed on a missing checkout, and — for the truly load-bearing components — asserts a runtime signature on the flashed/deployed artifact"* (the last step is §11.4.108's RUNTIME-ON-CLEAN-TARGET layer). The

general principle the corpus draws, which is why this belongs to the anti-mess control plane rather than to a build script: *"fail closed when an input to your correctness is unverifiable. An absent submodule being treated as a satisfied one is the same class as an unencrypted DB being treated as encrypted or a `nil, nil` stub being counted as coverage: a missing thing silently read as a present, correct thing. The rock-solid habit is to make the absence itself a hard failure, because 'the gate was skipped because the dependency was missing' is a bluff vector, not a pass."* Per clause (C) the reconcile-class is REFUSE-THE-GATE for classes (ii) and (iii) (a build/tag proceeding on an absent or unfetchable dependency is refused) and AUTO-SAFE only for a class-(i) drift whose correct pointer is evidence-determinable; per clause (E) the detector ships a golden-TRUE fixture per class and a golden-FALSE fixture — a fully-initialised tree at its intended refs — which it MUST NOT fire on (§11.4.201(1)). Composes §11.4.26 (pointer bumped in the same commit) / §11.4.28 / §11.4.31 / §11.4.32 / §11.4.36 / §11.4.108 / §11.4.167 / §11.4.246 (an unfetchable gitlink is a hermeticity violation: an undeclared-because-unobtainable input) / §11.4.252 (fail closed) / §11.4.264 (the shipped artifact must contain the tested source). Recommended mechanism gate `CM-DEPENDENCY-POINTER-PREFLIGHT-FAIL-CLOSED` + paired §1.1 mutation (de-initialise one submodule → the gate MUST FAIL; point one gitlink at a commit no configured remote serves → it MUST FAIL; golden-FALSE: a fully-initialised tree at intended refs MUST NOT fire it). Honest boundary (§11.4.6): the preflight proves the dependency's SOURCE is present and at the intended ref — it does NOT prove the built artifact contains it (that is §11.4.108's ARTIFACT layer and the corpus's own fourth step, the runtime-signature assertion). Gate-code = separate work item, NOT claimed shipped (§11.4.6 / §11.4.227).