

Project Lifecycle And Release

§ 11.4.8 – Deep-web-research-before-implementation

Before designing a non-trivial fix, before implementing a new feature, before declaring an architectural choice — perform deep web research to verify the chosen approach is informed by current state-of-the-art. The research surface includes:

1. **Official documentation** of the platform, framework, language, or service.
2. **Vendor technical guides** for the hardware / cloud / library.
3. **Open-source codebases** that already solved analogous problems.
4. **Coding tutorials + technical articles**.
5. **Issue trackers** for the relevant projects.

Operative rule. A fix that re-invents a wheel (or reproduces a known-broken pattern) when the open-source community has already solved the problem is a §11.4 violation by omission.

Documentation requirement. Every non-trivial fix's commit message (or accompanying entry in the project's Issues / Fixed file) MUST cite the research sources that informed it: at least one external link OR the literal "NO external solution found — original work".

§ 11.4.9 – Batch-source-fixes-before-rebuild

When working through a multi-defect closure cycle, all source-side fixes that DO NOT require runtime validation to design MUST be landed BEFORE the next artifact rebuild. The anti-pattern this mandate eliminates is `Fix A → rebuild → flash → cycle → fix B → rebuild → ...` which serializes operator time onto rebuild latency.

Exceptions where a fix DOES require interim rebuild MUST be documented in the fix's commit message as `REQUIRES_REBUILD: <reason>`. The default is the batch path; `rebuild-now` is the operator-authorized exception.

§ . . — Subagent-driven-by-default mandate (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

"Make sure we ALWAYS WORK EVERYTHING subagents-driven if it is possible or applicable!"

When operating in a multi-agent runtime (Claude Code with subagents, Cursor with task-runners, Aider with sub-sessions, or any CLI tool that supports delegated execution), the primary agent MUST default to subagent delegation for any task that satisfies AT LEAST ONE of:

1. **Multi-step scope** — three or more discrete editing / file-creation / verification phases. Hand-off boundaries are where stalls happen; pre-planned subagent decomposition avoids them.
2. **Parallelisable work** — two or more independent tasks with no shared mutable state. Parallel subagents finish wall-clock-faster than sequential foreground work AND insulate the primary agent's context window from the volume of file reads.
3. **Long-running diagnostic loops** — repeated probe / re-test cycles, build-flash-cycle loops, soak tests. Subagents survive context compression boundaries that would otherwise truncate progress.
4. **Domain-specific or specialised workflows** — code review, security audit, infrastructure change review, performance triage, documentation propagation across N files. Specialised subagents (`code-reviewer` , `iac-reviewer` , `general-purpose`) bring pre-curated tool sets and disciplines the primary agent would re-derive from scratch.

The primary agent SHOULD only do foreground work when AT LEAST ONE of:

- The task is a single edit or single file read with no follow-up.

- The task requires conversational clarification with the operator mid-execution (subagents cannot ask the operator questions).
- The task is gating critical state (a commit, a push, a tag cascade) that must be sequenced before the next operator action.
- The task is so quick (under ~30 seconds of execution) that subagent spin-up overhead exceeds the work itself.

Anti-stall discipline. Subagents have watchdog timeouts in most runtimes (Claude Code: ~600 s of no-progress). When delegating, the parent agent **MUST**: (a) scope tightly (no "complete everything" prompts — break into 4-6 tightly-scoped tasks); (b) require **checkpoint commits** after each major task so partial progress is preserved even on stall; (c) explicitly prohibit destructive operations the subagent might attempt to "be helpful" (`git reset --hard` , `git push --force` , `--no-verify`); (d) provide explicit discipline pointers (§11.4.6 no-guessing, §11.4.10 credentials, §11.4.17 classification).

Anti-bluff applied. A subagent that returns "completed" without captured evidence in commits / outputs / artifacts is a §11.4 PASS-bluff at the multi-agent layer. The parent agent **MUST** verify subagent claims against repo state (`git log` , `git status` , post-completion gate runs) before treating the work as landed.

Coordination. When parallel subagents touch overlapping files, the parent agent **MUST** partition the work so subagents work on **non-overlapping files**. The parent's `commit_all.sh --auto-cascade` naturally bundles both subagents' dirty files into atomic commits via `git add -A` — exploit this for "forward-reference + provider" patterns where one subagent creates files another subagent references.

No escape hatch. Operating exclusively in the foreground when subagent delegation is feasible burns operator wall-clock time, risks context-window overflow on large tasks, and forfeits the parallelism that multi-agent runtimes were designed to provide. Pre-build gate `CM-SUBAGENT-DELEGATION-AUDIT` (when implemented per consuming project) scans recent session transcripts for multi-step foreground work that should have been delegated; paired mutation enforces the gate is not a bluff.

§ . . — Full-suite retest before
release tag mandate (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-05-17):

"Please, do complete retests with all existing tests (multi-hours efforts) when all new workable items are done, fixed, polished and verified. Time is essential! We should already have this in our (root) Constitution, CLAUDE.MD and AGENTS.MD."

Operative rule. A release tag (any `vX.Y.Z` / `X.Y.Z-suffix` identifier MUST NOT be created until a **COMPLETE retest with ALL existing tests** has been executed on a clean baseline AFTER every workable item in the batch is done, fixed, polished, and individually verified. A "spot-check" retest that runs only the tests directly touched by the batch is FORBIDDEN — it misses interaction defects between the batch's fixes and previously- stable code paths.

The complete retest comprises:

1. **Pre-build verification full sweep** — every gate in the project's pre-build script runs; 0 FAIL required.
2. **Post-build verification full sweep** — every gate in the project's post-build script runs against the freshly-assembled image; 0 FAIL required.
3. **On-device 4-phase cycle** (e.g. `test_all_fixes.sh` or project equivalent) on **EVERY owned device** (the full topology at that point in time). Each phase (Immediate Fresh Flash / After Reboot / After Factory Reset / After Final Reboot) MUST complete; no phase may be skipped. Phase summaries captured.
4. **Meta-test full mutation sweep** — every paired mutation in `meta_test_false_positive_proof.sh` (or project equivalent) executed; every gate proven non-bluff.
5. **Test bank full sweep** — if a Challenge-driven test bank exists, every Challenge in the bank covered by a full QA session (not a sub-set).

6. **Issues.md / Fixed.md state audit** — no item silently demoted; every Reopened item has §11.4.34 Reopened-Details ; every closure has captured-evidence per §11.4.5; no Status: value outside the §11.4.15 closed-set.
7. **CONTINUATION.md sync check** — per §12.10, document reflects current state at the moment of tagging.

Time is essential. A complete retest is typically a **12–48 hour elapsed effort** depending on parallelism, device count, and meta-test mutation count. This is NOT optional and NOT abbreviated. Operators should plan release cadence with this duration in mind, NOT skip the retest to ship faster. Skipping the retest is the exact "tests passed but feature broken" failure mode §11.4 specifically prohibits.

Composition with §11.4.4. Per-fix retest (`test_all_fixes.sh` run after each individual fix lands) is STILL mandatory per §11.4.4. §11.4.40 is the **additional final integrity check** that runs after the BATCH is complete — catching interaction defects that individual per-fix retests cannot see in isolation.

Composition with §11.4.7. The complete retest is the authoritative captured-evidence baseline for any closure of items present in the batch. If a §11.4.7-promoted item PASSED its per-fix retest but FAILs the full-suite retest, the closure MUST be reverted and the item moved back to `In progress` — captured-evidence-contradicts under same-conditions per §11.4.7.

Composition with §11.4.39. Per-feature on-device end-user validation runs as part of step 3 (on-device cycle) — every feature's wrapper-test fires during the full-suite retest. The two mandates are complementary: §11.4.39 covers individual feature validation breadth; §11.4.40 covers release-time integration depth.

Gate CM-FULL-SUITE-RETEST-MANDATE . Pre-tag gate inspects the release-candidate state and verifies that within the last 72 hours of the candidate commit, evidence exists of (a) full pre-build + post-build run, (b) on-device 4-phase cycle on every owned device, (c) meta-test full sweep, (d) Issues.md/Fixed.md audit pass. Evidence captured in `docs/changelogs/<tag>.md` per §11.4.4(c) requirements. Paired mutation (§1.1): strip `Full-suite retest evidence` block from a changelog → gate FAILs.

Classification: universal (per §11.4.17) — every consuming project's release procedure. No escape hatch — there is no `--skip-full-retest` flag, no `--quick-release` mode. Operators who feel time-pressured to skip the full retest should instead delay the release until the retest completes; shipping unverified code is worse than delayed shipping.

§ 11.4.4 — Iteration-discipline mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

"Are we clear about working iterations and sorting priorities for testing and fixing? We first fix top and middle critical priorities → We flash and test all of these → If nothing is broken and no new issues are reported by users or found we run full system testing (many hours for everything) → if anything is still broken or new issues are reported we get back to fixing → Cycle repeats. Make sure this is absolutely clear and mandatory foundation of our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD! We MUST WORK in this manner until otherwise has been told."

Operative rule. Project work proceeds in priority-ordered iteration cycles. Each cycle is a closed loop of five steps that MUST run in order; advancing past any step before its acceptance condition is met is a §11.4 bluff equivalent to skipping the step entirely.

The five-step cycle:

1. **Priority-ordered fix selection.** Open the project's issue tracker, filter to items whose `**Status:**` is one of `Queued` | `In progress` | `Reopened` | `operator-blocked` (per §11.4.15 closed-set), sort by severity DESC then intra-group criticality DESC (per §11.4.12 sort order), select ONLY items classified `Top critical` or `Middle critical` for the current batch. `Low` items are explicitly DEFERRED until the critical batch is closed.
2. **Batch implementation.** Land source-side fixes for the selected batch per §11.4.9 (batch-source-fixes-before-rebuild) — every non-runtime-dependent fix lands BEFORE the next rebuild, every fix gets the §11.4.4 four-layer coverage (pre-build gate + post-build gate + on-device test + paired mutation).

3. **Smoke-test gate.** After flash, run the SMOKE bundle: (a) anti-bluff baseline (`meta_test_false_positive_proof.sh` for gates touched by this batch), (b) every `test_*.sh` whose name matches a fix in this batch, (c) the critical-path regression probe (boot + launcher + WiFi + audio + secondary-display routing core), (d) the §11.4.39 per-feature on-device validation for any feature touched. Estimated runtime: <30 minutes on the parallel device fleet. Smoke MUST emit zero FAIL and zero unclassified WARN.
4. **Full-system-test gate.** ONLY if (a) smoke is GREEN AND (b) no new operator/user report has surfaced during the batch's validation window, run the §11.4.40 complete retest (all 7 steps including 12–48 h on-device 4-phase cycle on every owned device). If smoke FAILED OR a new report arrived, the full-system test is FORBIDDEN; loop back to step 1 with the new evidence added to the priority queue.
5. **Release-readiness or loop.** If full-system test is GREEN AND no new report arrived during it, the batch is release-ready (the §11.4.40 tagging procedure may proceed pending operator authorization). Otherwise loop back to step 1.

Priority taxonomy (binds existing severity convention). `Top critical` = severity `C` (Critical per Issues.md convention) AND intra-group criticality `5`. `Middle critical` = severity `C` with intra-group `1 – 4` OR severity `M` (Major) with any intra-group score. `Low` = severity `L` per the existing `[C/M/L]` taxonomy documented in §11.4.12. Severity `WARN` items are treated as `Low` for iteration-discipline purposes.

Cycle repeats. The five-step loop continues, fed by the priority queue, until the operator explicitly authorizes a release. There is no upper bound on iteration count; "we've cycled enough" is not a Constitution-recognised stopping condition. Time is essential per §11.4.40 BUT the cycle MUST NOT be truncated to ship faster — that is exactly the bluff §11.4 forbids.

Composition. §11.4.4 (per-fix retest inside step 2) + §11.4.7 (closures require same-conditions positive evidence; step 4 is authoritative baseline) + §11.4.9 (source-side batching inside step 2) + §11.4.34 (Reopened items attribute By: AI / User) + §11.4.40 (the multi-hour full-suite retest IS step 4 — §11.4.42 is the meta-loop conductor binding the instruments).

Anti-bluff coupling. Per §11.4.2 + §11.4.5, every smoke-test and full-system-test PASS MUST carry positive captured evidence of user-visible behaviour. Tests AND HelixQA Challenges bound equally — a Challenge that scores PASS without applicable analysis is a §11.4 PASS-bluff.

No escape hatch. No `--skip-priority-batch` , `--skip-smoke` , `--full-suite-only` , or `--release-without-loop` flag exists. Subagents performing autonomous work default to the §11.4.42 path. Operators who feel time-pressured to short-circuit the loop should add capacity (more devices in parallel, more agent slots) rather than skip steps. Shipping under-validated code is worse than delayed shipping.

Gate `CM-COVENANT-114-42-PROPAGATION` mirrors `CM-COVENANT-114-40-PROPAGATION` 1:1 — every CLAUDE.md / AGENTS.md in the covenant file set carries the §11.4.42 anchor. Paired mutation strips the anchor literal from one consumer file → gate FAILs.

Classification: ^{11 4 46}universal (per §11.4.17) — every consuming project's iteration workflow.
_{2026 05 18}

§ . . — `Validate-recent-work-before-post-flash-tests` mandate (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-05-18):

"We see that on newly flashed devices we execute all post flash tests. Make sure we do execute them all after we validate and verify all recent work from the Issues and Continuation docs first! Once all new / latest work is fully validated and verified, no new issues found, nothing is discovered broken or faulty, then we run all post-flash tests! If validation and verification of recent work fails for any reason, and we must go back to fix anything that popped up we do not have to run post-flash tests! We MUST save time! All post-flash tests can be executed only after recent work on features and issues is fully validated and confirmed with complete anti-bluff policy enforced! Add this so it is absolutely clear into our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD! Start following all these rules IMMEDIATELY! Iff required abort all post-flash testing if it is in progress!"

Operative rule. After every device flash, the orchestrator MUST first run a recent-work validation pass before any full post-flash suite (`test_all_fixes.sh` or project equivalent). The validation pass enumerates recent work from three docs (Issues.md / Fixed.md / CONTINUATION.md §3), maps each recent item to its on-device test binding, runs ONLY those tests, classifies each result per §11.4.6, and demands 100% green before the full suite is permitted.

Definition of "recent work" (closed-set):

- docs/Issues.md actionable headings with `**Status:**` $\in$ {In progress, Ready for testing, Reopened} (per §11.4.15 closed-set).
- docs/Fixed.md headings closed within `--max-age-days N` (default 7).
- docs/CONTINUATION.md §3 "Active work" sub-headings listed as IN PROGRESS or BLOCKED.

Recent-work test binding. Each recent item's on-device test is discovered via §-letter or Fix # match against the project's `test_*.sh` filenames OR an explicit `Test:` line in the item heading. Items without an on-device test are §11.4.4 four-layer-coverage violations (a fix without an on-device test is forbidden) — flagged as VALIDATION-PASS-BLOCKED until the test lands.

Authorization sequence (machine-enforced):

1. `scripts/testing/recent_work_validate.sh --device <serial>` runs; writes `/data/local/tmp/.recent_work_validated` on the device IFF all targeted tests return PASS with captured evidence.
2. The full suite (`test_all_fixes.sh`) checks for the marker file at entry; absent or stale (older than the last flash boot epoch) → refuses to proceed with exit code 11.
3. Marker is invalidated automatically on reboot — the orchestrator stores the device boot epoch in the marker and re-validates.

Anti-bluff during validation. Each recent-work item that landed a fix in this batch MUST have a paired §11.4.43 RED test (captured before the fix) AND the same test now GREEN (captured after the fix). A GREEN with no prior RED is itself the bluff §11.4 forbids — the gate FAILs.

Honest classification of validation FAILs. Per §11.4.6: PRODUCT defect (real regression — block full-suite, surface to operator), TEST-INFRA defect (test-script bug — fix per §11.4.1 at source layer, re-run), BLUFFY-THRESHOLD (e.g. SLO floor wrong — re-baseline), or UNCONFIRMED (not reproducible in isolation — tag per §11.4.7 PENDING_CYCLE_RETEST, NEVER silently demote).

Composition. §11.4.4 (STOP-on-discovery) + §11.4.6 (no- guessing) + §11.4.7 (demotion-evidence) + §11.4.40 (full-suite gate) + §11.4.42 (iteration discipline — smoke-test ≡ recent-work- validation) + §11.4.43 (RED test for each recent-work item IS the validation test) + §11.4.44 (revision header determines recency) + §12.10 (CONTINUATION.md source-of-truth for §3 Active work).

Gates:

- `CM-COVENANT-114-46-PROPAGATION` — anchor literal present in every CLAUDE.md / AGENTS.md across the covenant file set.
- `CM-AF-RECENT-WORK-VALIDATION-GATE` — asserts `scripts/testing/recent_work_validate.sh` exists + executable + sources the anti-bluff library + the full-suite orchestrator contains the entry-gate check.
- `CM-AF-VALIDATION-ARTIFACT-FILE` — asserts the marker file path `/data/local/tmp/.recent_work_validated` is literal-identical across helper + orchestrator + propagation-block + this Constitution section.

Paired mutations strip the anchor / remove the entry-gate / flip the marker path → respective gate FAILs.

No escape hatch. No `--skip-validation`, `--full-suite-always`, `--ignore-recent-work` flag is permitted. The full-suite-without-validation pattern is the precise "tests passed but feature broken" failure mode §11.4 specifically prohibits. The 60–90 min validation pass is ALWAYS faster than 4–6 h of full-suite chasing a defect.

Classification: universal²⁰²⁶⁻⁰⁵⁻¹⁸ (per §11.4.17).

§ 11.4.17 – Firebase Data Review Mandate (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-18):

"We MUST regularly before every bigger working round check Firebase information gathered so far: Crashlytics (all fatals, non-fatals and ANRs), Analytics data and Performance data. For anything problematic depending on severity proper workable items (issues) MUST be created (Issues, Issues_Summary docs) with full references (links) to original data on Firebase. We MUST make sure we do not create for same issues multiple entries, so we MUST distinguish between same issue manifested in different forms (stacktraces)! Everything MUST BE comprehensive and in-depth level of details so any changes we do (or fixes) do solve the original root problems!"

The mandate. Before every "bigger working round" (pre-build, pre-flash, pre-tag, daily, post-deployment burn-in) the operator/ loop MUST execute the Firebase review pass via `scripts/firebase/review_round.sh`. The pass queries Crashlytics (fatals + non-fatals + ANRs) + Analytics + Performance, classifies each finding by the §11.4.47 severity table, dedup-maps to existing Issues.md entries via the three-tier algorithm, and drafts new Issue entries for unrecognised findings. Skipping the pass is a §11.4 PASS-bluff: Firebase IS the captured evidence from real end-user devices; ignoring it is the precise "tests pass, feature broken in the wild" failure mode §11.4 prohibits.

Five mandatory elements (ALL must hold):

1. **Trigger cadence (5 trigger types).** Pre-build (blocking, <24 h freshness), pre-flash (blocking, <24 h), pre-tag (blocking, <6 h), daily 09:00 UTC default (non-blocking sweep), post-deployment burn-in T+24 h (non-blocking).
2. **Three-source query.** Crashlytics + Analytics + Performance — ALL three. Skipping one source is a §11.4 PASS-bluff (Firebase reports a regression on the skipped axis and we never see it).
3. **Issues.md output.** Every "problematic" finding MUST map to either (a) a new Issues.md entry, OR (b) a recognised existing entry via the dedup algorithm. Both paths produce a full Firebase Console URL in the entry's metadata so any future agent can re-fetch the raw data without rebuilding context.
4. **Three-tier deduplication.** Tier 1 (exact Firebase Issue-ID match) → Tier 2 (stacktrace-similarity cluster hash: top-3 non-generic frames normalised + crash class → SHA-256 first 8 hex chars) → Tier 3 (monthly operator merge review). Multiple Issues.md entries pointing at the same underlying root cause is a §11.4 violation (operator chases ghosts; root cause stays unfixed).
5. **Comprehensive root-cause analysis.** Every Firebase-sourced Issue carries the §11.4.4(a) systematic-debugging output — Phase 1 evidence, Phase 2 pattern, Phase 3 root-cause hypothesis (UNCONFIRMED until §11.4.43 RED test reproduces), Phase 4 fix direction. Comment-only entries ("crash in App X") are PASS-bluff stubs and FAIL the Issue-xref gate.

Severity classification: Crashlytics FATAL impactedUsers > 10 last 7d → Critical; 1–10 → Major; 0 but historic → Low. NON-FATAL

100/day → Major; 10–100 → Low; < 10 → Informational. ANR 1% sessions → Major; 0.1–1% → Low. Performance KPI regression 20% → Major; 10–20% → Low. Analytics funnel-drop > 50% → Major; 25–50% → Low; feature_used regression > 75% → Major. Composes with §11.4.16 Type assignment.

Gates:

- CM-COVENANT-114-47-PROPAGATION — anchor present across every CLAUDE.md / AGENTS.md in the covenant file set.

- CM-AF-FIREBASE-REVIEW-CADENCE —
scripts/firebase/review_round.sh exists, executable, contains the 7-stage pipeline literals + dedup algorithm + severity table.
- CM-AF-FIREBASE-ISSUE-XREF — every Issues.md entry whose
Source: is Firebase Crashlytics / Firebase Analytics /
Firebase Performance carries **Firebase Issue IDs:** + **Firebase
URL:** + (**Stacktrace Cluster Hash:** for Crashlytics OR **KPI:**
for Performance OR **Funnel:** for Analytics).

Paired mutations strip the anchor / rename the review_round function / remove the Firebase-URL field from the template → respective gate FAILs.

Composition. §11.4.4 (STOP-on-discovery), §11.4.4(a) (systematic-debugging), §11.4.6 (no-guessing — Firebase data IS captured evidence), §11.4.7 (demotion-evidence — a previously- Fixed Issue resurfacing in Firebase IS positive captured evidence the fix didn't hold; reopen attribution = "AI: captured-evidence-contradicts" per §11.4.34), §11.4.10 (credentials — never log the bearer token), §11.4.12 (Issues_Summary sync), §11.4.14 (cleanup of /tmp/firebase_review_* work-dirs), §11.4.15 (Status: Queued), §11.4.16 (Type per severity), §11.4.34 (Reopened-Details on Firebase-resurface), §11.4.42 (implicit step 2 inclusion), §11.4.43 (RED test per stacktrace before fix lands), §11.4.44 (revision header on every Firebase-sourced Issue), §11.4.45 (Firebase Status.md at docs/firebase/status/Status.md), §11.4.46 (validation pass consults latest Firebase delta).

No escape hatch. No --skip-firebase-review ,
--no-issue-from- firebase , --firebase-review-not-applicable flag is permitted. The operator MAY filter by minimum severity (--severity-min major) to reduce noise but the pass itself MUST execute.

Classification: universal (per §11.4.17) — every consuming project that ships to real end-user devices benefits from the same review cadence and the same dedup algorithm.

§ . . — Autonomous-Validation Mandate
(User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-18):

"Make sure we have full automation tests which will do all this work in full automation! IMPORTANT: Make sure that all existing tests and Challenges do work in anti-bluff manner — they MUST confirm that all tested codebase really works as expected! We had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can't be used! This MUST NOT be the case and execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users of the product! This MUST BE part of Constitution of our project, its CLAUDE.MD and AGENTS.MD if it is not there already, and to be applied to all Submodules's Constitution, CLAUDE.MD and AGENTS.MD as well."

Why this anchor exists. §11.4.25 (full-automation coverage) and §11.4.27 (no-fakes-beyond-unit-tests + 100%-test-type-coverage) and §11.4.39 (per-feature on-device end-user validation) collectively mandated that every user-facing feature has automation tests with captured runtime evidence. They did NOT explicitly forbid operator-attended-only validation paths — i.e., a feature whose ONLY proof of working state requires a human to physically present at the device, drive UI manually, and observe outcomes. Phase 39.§CN exposed the gap: §CN's Kinopoisk 5.1 EAC3 closure was structurally fixed (libavutil bundled in APEX, six decoders register) yet end-user-validation degraded to "operator drives remote control while test polls" because Kinopoisk's TV-Compose UI is not accessibility-instrumented and uiautomator returns near-empty hierarchy. That validation path is operator-attended, non-reproducible in CI, and cannot prove deterministic consistency (§11.4.50) across N iterations. The User mandate elevates "full automation" from "preferred" to "mandatory": every user-facing PASS MUST have at least one captured-evidence path that does NOT require operator presence.

Operative rule (4 mandatory elements):

1. **Every user-facing feature MUST have at least one autonomous validation path.** "Autonomous" means: the validation runs end-to-end via `adb shell` + scripted automation, produces captured runtime evidence per §11.4.5, and reaches a PASS / FAIL verdict WITHOUT a human present to drive UI, observe screen, or make decisions. Operator-attended tests are SUPPLEMENTARY, never PRIMARY. A feature whose ONLY validation path is operator-attended is

a §11.4.52 violation regardless of how carefully the operator's captures match §11.4.5 evidence requirements — the path does not scale to CI, does not run on every commit, does not survive operator unavailability, and produces the exact "tests pass but feature doesn't work for users" failure mode §11.4 specifically forbids.

2. **Acceptable autonomous-validation paths** (any one or more — non-exhaustive):

- **Programmatic instrumentation APK** — a small platform-signed app that exercises the under-test surface directly via SDK APIs (e.g., `MediaCodec.createDecoderByName`, `MediaCodec.configure`, `AudioTrack` writes) and writes a structured JSON result file to a discoverable path. Closes UI-instrumentation gaps where the target app is not accessibility-instrumented.
- **Headless intent dispatch + state poll** — `am start --es /`
`am broadcast --es` to dispatch the under-test code path, then poll `dumpsys`, `/proc/<pid>/maps`, `media.metrics`, `/sys/...`, kernel sysfs, or the integration's network probe for the expected state transition.
- **ADB-driven uiautomator** — applicable ONLY if the target app exposes accessibility nodes the dump can resolve. MUST verify via `uiautomator dump | grep -c clickable=true ≥ 1` before claiming uiautomator coverage is real. A near-empty hierarchy is itself the captured-evidence proving UI-driven automation is INFEASIBLE for that target and the test MUST fall back to instrumentation-APK or headless-intent paths.
- **Network-side sink probe** — for features that surface on a network-accessible peripheral (e.g., Arvus HDMI dashboard at `http://<sink>/`, Sonos REST API, AirPlay receiver), the sink's own report counts as autonomous out-of-band evidence per §11.4.13.
- **HelixQA autonomous QA session** — for projects that have HelixQA fully incorporated per §11.4.27, the orchestrator's end-to-end session output counts as autonomous evidence.

3. **Per-feature classification + tracking.** Every entry in the project's coverage ledger (§11.4.25) MUST classify each feature's autonomous-validation path as one of: `AUTONOMOUS_VERIFIED` (path exists + last-run-green per §11.4.46), `AUTONOMOUS_DESIGNED` (path exists but RED per §11.4.43),

`OPERATOR_ATTENDED_ONLY` (path requires human presence — release blocker until promoted to one of the other states), or `NOT_APPLICABLE` (e.g., a CLI tool that has no UI surface and the unit + integration tests already constitute autonomous proof). `OPERATOR_ATTENDED_ONLY` rows MUST cite a tracked work item per §11.4.15 + §11.4.16 that designs the autonomous-path migration. Closing the work item requires the migration to land, not just the operator capture.

4. **Anti-bluff for autonomous paths themselves.** An autonomous path that returns PASS without exercising the user-visible behaviour is a §11.4.52 violation just as severe as the original §11.4 PASS-bluff pattern. The autonomous path MUST produce positive captured evidence per §11.4.5 (audio: channel count + sample rate + glitch census; video: frame count + routing target
- frame health + obstruction census). A grep on a metadata field that doesn't prove behaviour is not an autonomous validation — it is a metadata bluff dressed up as automation. Paired meta-test mutations (§1.1) MUST catch the autonomous path's degradation to metadata-only or grep-only.

Composition.

- §11.4.25 — full-automation-coverage mandate. §11.4.52 is the strict expansion of invariant 1 + 2: "anti-bluff posture" and "proof of working capability" cannot be carried by operator-attended-only evidence — they MUST be carried by at least one autonomous path.
- §11.4.27 — no-fakes-beyond-unit-tests + 100%-test-type-coverage. §11.4.52 is the operational layer that closes the gap between "the project has full-automation tests" and "every PASS has an autonomous evidence path".
- §11.4.39 — per-feature on-device end-user validation. §11.4.52 refines: the on-device scenario MUST itself be autonomously drivable.
- §11.4.43 — TDD-fix-discipline. The autonomous path MUST exist in the RED state BEFORE the fix lands, then go GREEN after. An autonomous path authored after the fix is a §11.4 PASS-bluff.
- §11.4.48 — UI-driven video testing. §11.4.52 establishes that when uiautomator returns a near-empty hierarchy on the target app, UI-driven IS infeasible and the test MUST fall back to instrumentation-APK or headless-intent. The fallback is REQUIRED — not optional.

- §11.4.49 — dual-approach testing. The Intent variant in a dual-approach pair often IS the autonomous path when the UI variant requires operator assistance.
- §11.4.50 — deterministic consistency. Autonomous paths run N iterations naturally; operator-attended paths cannot satisfy the N-iteration discipline at scale.
- §11.4.51 — live-ADB-first maximization. The instrumentation-APK path is itself `LIVE_ADB_TESTABLE` (`gradle build + adb install -r`
 - `run + read JSON result`), so live-probe IS the design loop for the autonomous path.

Pre-build gates:

- `CM-COVENANT-114-52-PROPAGATION` — anchor literal §11.4.52 present in all 5 canonical files (`constitution/Constitution.md` , `constitution/CLAUDE.md` , `constitution/AGENTS.md` , plus the parent consumer's `CLAUDE.md + AGENTS.md`) PLUS the full per-consumer propagation (parent project + every owned submodule
 - nested submodules + HelixQA dependencies — same scan pattern as `CM-COVENANT-114-51-PROPAGATION`).
- `CM-AF-AUTONOMOUS-PATH-PER-FEATURE` — every entry in the consuming project's coverage ledger (§11.4.25) has a non-empty autonomous-classification column with value in `{AUTONOMOUS_VERIFIED, AUTONOMOUS_DESIGNED, OPERATOR_ATTENDED_ONLY, NOT_APPLICABLE}` . `OPERATOR_ATTENDED_ONLY` rows MUST cite a tracked migration work item. (Project-side gate — consuming projects implement the ledger surface in their own pre-build test suite.)

Paired mutations (per §1.1):

- Strip §11.4.52 literal from `constitution/Constitution.md` → `CM-COVENANT-114-52-PROPAGATION` FAILs.
- Inject a feature row with `OPERATOR_ATTENDED_ONLY` and NO tracked migration work item → `CM-AF-AUTONOMOUS-PATH-PER-FEATURE` FAILs.

No escape hatch. No `--allow-operator-attended-only` , `--skip-autonomous-path` , `--manual-validation-suffices` flag exists. The discipline exists because the user mandate is unambiguous: "execution of tests and

Challenges MUST guarantee the quality, the completion and full usability by end users". Operator-attended-only validation cannot guarantee that property at the release-gate layer — only autonomous paths can.

Classification: universal (per §11.4.17). Applies to every project consuming the constitution submodule. Each consumer adapts the autonomous-validation path catalogue to its own technology surface (Android: instrumentation APK + uiautomator + headless intent; Web: Playwright headless + API probe; CLI: subprocess

- stdout/stderr assertion; mobile native: device farm headless runner) but the mandate to provide at least one autonomous path per user-facing feature is universal.

2026-05-19

§ . . – Blocker-resolution interactive-clarification mandate (User mandate, – –)

Forensic anchor — direct user mandate (verbatim, 2026-05-19):

"If any blockers which can be resolved with interactive response ever happen again, perform in depth research on options doable by your side and how much inputs from us you really need, then create options and present them to us. After we answer, preferably you will be unblocked and be able to continue work on blocked items. Let us make this main approach when such situations (blockers) do happen!"

When any task is blocked (waiting on operator decision, hardware access, external-service authorization, ambiguous scope, or any input the agent cannot mechanically derive), the agent **MUST** follow this five-step discipline before idling or asking free-form questions:

1. **In-depth research on agent-side options.** Identify the maximum scope that can land unilaterally without operator input — research the codebase, upstream documentation, existing tooling, the current state, what files/components are involved, what tests/gates already exist. Surface every agent-actionable path even if it requires a small operator confirmation.
2. **Calculate minimum-viable operator input.** Reduce the question set to the smallest closed set of operator-only decisions (preference, authorization, scope-bounding, hardware availability, schedule). Anything that can be researched **MUST NOT** be asked.
3. **Construct 2–4 mutually-exclusive options.** Each option carries (a) a short label, (b) a one-line trade-off description, (c) an explicit statement of what the agent will do *after that answer*. One option marked "Recommended" with rationale. Options **MUST** genuinely differ — restating one option as three close variants is itself a §11.4.66 violation.
4. **Present via the platform's interactive question mechanism.** On Claude Code that is `AskUserQuestion` (max 4 questions, each with 2–4 options; supports multi-select for non-mutually-exclusive sets). Other platforms (Copilot CLI, Codex, Gemini CLI) use their equivalents. **NEVER** inline free-text "what would you like?" when interactive options would do — that wastes operator attention.
5. **After the operator answers, resume work without additional round-trips.** Every option's promised action **MUST** be sufficient to unblock — if a follow-up clarifying question is needed, the prior options were insufficiently researched and that's itself a §11.4.66 violation. The contract is: ask once, unblock, continue.

Composes with:

- §11.4.6 (no-guessing — interactive options replace agent guessing about operator preference)

- §11.4.7 (demotion-evidence — operator preference is captured- evidence for direction changes)
- §11.4.40 (full-suite retest — uninterrupted work post-answer means the test cycle resumes cleanly without context drift)
- §11.4.41 (merge-first — when operators make conflicting choices across sessions, the merge-first discipline preserves both)
- §11.4.42 (iteration-discipline — interactive-question loop is built into the per-cycle methodology)
- §11.4.52 (autonomous-validation — most blockers can be researched to autonomous-path level, leaving only true operator preference to interactive-ask)

Mandatory protections:

- **No silent waiting.** If blocked, ask. If genuinely no operator input is needed, do not ask — proceed.
- **No bulk-text questions when interactive options would do.** Free-text prompts are reserved for truly open-ended ideation, never for closed-set decisions.
- **Each interactive question minimizes operator cost** — the operator should be able to answer all questions in under 30 seconds total.
- **Recommended option is highlighted explicitly** — the operator should know what the agent would default to if no answer comes.
- **Out-of-band channels (slack, voice, email) NEVER substitute** — the interactive mechanism creates an audit trail that free-text channels do not.

Pre-build gate `CM-COVENANT-114-66-PROPAGATION` enforces the anchor literal is present across the 42-file consumer fleet (parent CLAUDE.md / AGENTS.md + 10 owned-submodule pairs + HelixQA-owned pairs). Paired meta-test mutation strips the literal → gate FAILs. No escape hatch — no `--skip-ask` , `--silent-wait` , `--free-form-only` flag is permitted.

Canonical authority: constitution submodule `Constitution.md` §11.4.66.

Non-compliance is a release blocker regardless of context.

§ . . — Subagent-Driven Execution Is The Default (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"Always do if possible Subagent-driven! Add this into our root (constitution Submodule) Constitution.md, CLAUDE.md and AGENTS.md. This should be the default choice ALWAYS!"

When executing implementation plans authored via `superpowers:writing-plans` (or any equivalent task-decomposed execution flow), the **default execution model is subagent-driven** per `superpowers:subagent-driven-development`. Inline execution via `superpowers:executing-plans` is permitted ONLY when (a) the task is trivial AND fits in a single sub-300-line edit, OR (b) the operator explicitly requests inline execution at brainstorm-handoff time per the `superpowers:writing-plans` skill's explicit-execution-mode prompt.

Why subagent-driven is the default:

1. **Isolated context per task** — each subagent gets a fresh context window scoped to its task, avoiding the conductor's context bloat that degrades reasoning quality after ~100 k tokens.
2. **Two-stage review naturally enforced** — fresh subagent does the work, conductor reviews, optional second subagent verifies. This composes with §11.4.4 four-layer coverage + §11.4.43 TDD-fix- discipline.
3. **Parallel-PWU compatible** (§11.4.58) — subagents are the unit of parallelism in the PWU pipeline; making them the default aligns the writing-plans handoff with the parallel-development methodology.
4. **Anti-bluff seam** (§11.4) — the conductor's review of a subagent's output is structurally separated from the work itself, eliminating self-review blind spots.
5. **Survives operator absence** — subagents resume from their on-disk plan + spec inputs; the conductor session can be lost without losing in-flight task state.

Composition with §11.4.4 (four-layer coverage), §11.4.6 (no- guessing — subagent's captured output IS the evidence), §11.4.42 (iteration discipline), §11.4.43 (TDD-fix), §11.4.50 (deterministic consistency — the subagent's

deterministic exit code), §11.4.51 (LIVE_ADB_FIRST — subagents classify rebuild-requirement before commit), §11.4.58 (parallel-development PWU — subagents ARE the parallel work units).

No escape hatch. `--inline-execution-required`, `--no-subagents`, `--monolithic-execution` are NOT permitted flags. Operator may request inline at brainstorm-handoff time per `superpowers:writing-plans` skill; absent that, subagent-driven is default. Skipping subagent-driven for non-trivial work without recorded operator authorisation is itself a §11.4 PASS-bluff (the conductor self-reviews, the anti-bluff seam disappears, and the forensic class of incident this Constitution catalogues becomes mechanically possible again).

Captured-evidence enforcement. Pre-build gate `CM-COVENANT-114-70-SUBAGENT-DEFAULT-PROPAGATION` enforces this anchor literal in every CLAUDE.md / AGENTS.md across parent + 10 owned submodules + nested submodules + HelixQA dependencies (the ~44-file consumer fleet). Paired §1.1 meta-test mutation strips the anchor literal → gate FAILs.

Canonical authority: constitution submodule `Constitution.md` §11.4.70. Default applies to every `superpowers:writing-plans` skill terminal handoff and every multi-task execution flow with independent subtasks.

Non-compliance is a release blocker regardless of context.

§ . . — Audio Top-Priority Mandate (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"Make sure all fixes for audio are always top priority in main working stream!"

The conductor (main working stream — whether a Claude Code session, an AI agent, or a human operator) MUST treat audio fixes as the highest-priority class on the serial dispatch queue. Audio fixes include but are not limited to §EU HDMI rejection, §EM multichannel HDMI, §ET Arvus integration, §EV/§EW/§EX D3 audio

defects, HiFi (Fix #74 pcm_config_deep, #76 HDMI/SPDIF resilience), AC3 (Fix #106 FFmpeg Codec2), ES8388 (Fix #103/§104/#105 mutual-output + DAPM), multichannel LPCM (Fix #112), and every future audio-stack improvement.

Composition with §11.4.42 (iteration-discipline mandate) + §11.4.58 (parallel-development methodology). Per §11.4.42, the conductor selects work items in priority order; this anchor pins audio at the apex of that order. Per §11.4.58 PWU pipeline, parallel BACKGROUND subagents (research, refactors, infrastructure documentation) MAY run concurrently with audio work — they do NOT preempt audio on the main-stream serial dispatch queue.

Operative rule: any time the conductor faces a choice between dispatching an audio task vs a non-audio task on the SAME serial resource (most commonly the pre_build_verification.sh + meta_test_false_positive_proof.sh edit-and-commit critical path), the audio task wins.

No escape hatch — there is no "but this non-audio task is faster" or "but this research is more interesting" override. The discipline exists because audio-stack regressions are user-perceptible and high-impact (D3 silent post-flash is a release blocker), while research and refactors can wait.

Canonical authority: constitution submodule `Constitution.md` §11.4.72.

Non-compliance is a process violation regardless of context.

§ . . — Iteration-speedup discipline mandate (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-22):

"How can we speed-up this whole development and fixing process? ... Do not forget to all speed optimizations critical rules and mandatory constraints MUST BE all added into our root (constitution Submodule) Constitution.md, CLAUDE.md, AGENTS.md and QWEN.md and all other relevant constitution Submodules files!"

Iteration cycle time is a first-order quality enabler. A 40-min rebuild cycle that catches one defect per cycle bounds the project's defect-discovery rate at 1 defect per 40 min; a 15-min cycle bounds it at $\sim 3\times$ that rate. Slow cycles are a §11.4 PASS-bluff at the velocity layer — the project ships fewer validated features per unit of operator time than it should.

The 2026-05-22 forensic session that produced this anchor witnessed (a) 2 subagent watchdog stalls each ~ 10 min wasted, (b) 1 stalled `commit_all.sh` ~ 10 min lost on submodule cascade prompt, (c) 1 wrong-call-site speculative kernel patch (§GQ.2) burning a full 45 min rebuild cycle before forensic evidence pointed to the correct site (§GT Fix B-1), (d) 1 stale `git index.lock` from a killed process blocking the next commit, (e) 1 auto- `git gc` maintenance colliding with urgent commit. Each line item cost between 5 and 45 minutes of wall-clock that should have been productive iteration time. The aggregate impact for that single morning was ~ 3 hours of operator wait dilated onto a job whose intrinsic compute was ~ 90 minutes.

The mandate. Every consuming project's build / test / commit / debug pipeline MUST adopt the following speedup disciplines AS MANDATORY. Each is independently enforceable:

(A) Phase forensic before any speculative source patch

Before applying ANY non-trivial source patch (kernel, framework, native code, HAL, ASoC), the agent MUST first complete Phase 1 of `superpowers:systematic-debugging` — read the error/log/call-chain end-to-end and identify a FACT-grade root cause. Speculative patches without Phase 1 evidence are §11.4.6 (no-guessing) violations AND §11.4.82 violations. A wrong-call-site rebuild cycle wastes 30-90 min of compute; the alternative (15-30 min of focused source reading) is a net positive every time. Phase 1 is mandatory regardless of operator time pressure — "we don't have time for forensics, just rebuild" is the §11.4.82 anti-pattern.

(B) Live-ADB-First (or live-equivalent) before any rebuild

Per §11.4.51 — strengthened by §11.4.82 to a release-blocker mandate. $\sim 70\%$ of audio / UI / boot-script / config changes are LIVE_ADB_TESTABLE (push + `setprop` + `pm install -r` + 30-second validation cycle). Only kernel /

framework Java/AIDL / native C++-in-APEX / sepolicy / init.rc / ro.* / Android.bp / XML-resource-overlays / codec-XML-in-APEX changes are genuinely REQUIRES_REBUILD. Skipping the live-probe step for a LIVE_ADB_TESTABLE change is a §11.4.82 violation — costs operator 45 min of rebuild for what could have been 30 seconds of `adb push` .

(C) Pre-flight before launching rebuild orchestrators

Before any `nohup bash scripts/<rebuild>.sh &` invocation, the agent MUST run a ~30-second pre-flight that verifies (i) target device(s) reachable via ADB, (ii) downstream sinks (Arvus AVR, HDMI sink, Bluetooth peers) reachable as the test expects, (iii) host has sufficient memory and disk per §12.6 budget, (iv) no stale lock files in `.git/` or build directories, (v) no orphan git/build processes from prior killed iterations. A rebuild orchestrator launched against a broken precondition wastes 45 min discovering at the test stage what 30 seconds of pre-flight would have caught.

(D) Persistent build caches outside containers

Containerized AOSP / kernel builds MUST persist `ccache` AND any equivalent build-system intermediate cache (Soong cache, sccache, Gradle daemon state) on the host via bind-mount, NOT inside the ephemeral container layer. Without this, every container restart drops the cache and forces a cold rebuild (~40 min for AOSP+kernel on a fast workstation). With it, incremental rebuilds where only one driver / one app changed drop to 5-15 min (ccache hit rates 80-95%). This is a one-time setup of ~15 min for ~25-32 min savings per subsequent cycle.

(E) Module-only rebuild for loadable-module-only changes

When a patch touches ONLY one or more files inside `kernel/.../drivers/*/foo.c` AND the affected drivers are built as loadable modules (`CONFIG_*=m` in the active defconfig), the rebuild MUST be `make -C kernel M=<driver-path> modules` (~2-5 min) NOT a full kernel rebuild (~25 min). A `flash_kernel_only.sh` or equivalent helper writes the rebuilt module(s) to the running device without re-flashing system.img. This saves 15-20 min per cycle for the ~10-30% of kernel patches that qualify. Built-in drivers (`CONFIG_*=y`) do not qualify — full rebuild required for those.

(F) Parallel multi-device testing

When the project owns more than one validation device (D3, D4, ... in topology terms), the test orchestrator **MUST** run autonomous validation cycles on **EVERY** device in parallel via background processes, with separate `qa-results/<TS>/<device-tag>/` output directories per device. This catches per-device topology defects (per §11.4.3) one cycle earlier instead of waiting for a device-N-focused re-cycle. Direct wall-clock savings = 0 (parallel); amortized savings = 5-15 min per cycle when a device-specific defect surfaces.

(G) Subagent scope discipline + worktree isolation

Subagents dispatched per `superpowers:subagent-driven-development` **MUST** be:

- **Scope-bounded** to ≤ 30 min of focused work with intermediate output emitted as the subagent progresses (the canonical watchdog stalls at 600s of no-output; if a subagent writes its findings file incrementally it survives even when blocked on an external dependency).
- **Worktree-isolated by default** via the Agent tool's `isolation: "worktree"` parameter when supported, so parallel subagents do not contend on the same working tree and so a subagent crash never corrupts the conductor's state.
- **Single-responsibility** — one investigation OR one fix-implementation OR one doc draft, not multiple. The empirical pattern from the §GT investigation is that 8-task subagents stall while 2-3-task subagents complete reliably.

(H) Lock-file + stale-process hygiene

Killed or crashed agents leave behind `.git/index.lock`, `.git/.commit_all.lock`, `.lock` files in build directories, and orphan child processes (`git add -A`, `git pack-objects`, `git gc`). The next agent **MUST** detect these on session start (or after operator-visible crash) and clean them — never wait silently for a never-arriving release. Auto `git-gc` **MUST** be disabled in repos with concurrent multi-agent work (`git config gc.auto 0`) — its background invocation is a guaranteed lock-contention source.

(I) Cycle telemetry per §11.4.24 build-resource stats

Every iteration cycle MUST record (i) commit hash of the source under test, (ii) wall-clock of build + flash + test phases separately, (iii) the speedup-discipline flag set (which of A-H applied), (iv) outcome (PASS / FAIL / partial / stall). Aggregated weekly, this exposes which speedup gives the biggest empirical return on the specific project, and which is being neglected. The §11.4.24 build-resource stats tracker is the canonical telemetry pipeline; §11.4.82 extends its row schema with iteration-discipline columns.

Composes with. §11.4.4 (test-interrupt-on-discovery + retest), §11.4.6 (no-guessing — speculative patches forbidden), §11.4.9 (batch-source-fixes-before-rebuild), §11.4.20 / §11.4.70 (subagent-driven), §11.4.24 (build-resource stats), §11.4.42 (iteration-discipline conductor loop), §11.4.43 (TDD-fix discipline — RED-test enables faster cycle reads), §11.4.50 (deterministic consistency — caches don't break determinism if ccache config disables time-sensitive headers), §11.4.51 (LIVE-ADB-first maximisation — strengthened to mandate here), §11.4.52 (autonomous-validation — every cycle's verdict is automation-produced), §11.4.58 (parallel-development PWU — parallel devices and parallel subagents are §11.4.58 instances), §12.7 (-j2 host-safety cap remains — speedups MUST stay within memory budget), §107 (end-user usability — fast iteration = fewer cycles → fewer defects shipped).

Pre-build gate. CM-ITERATION-SPEEDUP-DISCIPLINE (when implemented in consuming project): audits the most recent N iteration cycles for telemetry entries citing which of (A)-(I) applied; flags any cycle where none of (A)-(I) was active as a candidate for review. Paired §1.1 meta-test mutation: strip the speedup-flag column from a synthetic telemetry row → gate FAILs.

Operational discipline. When a consuming project ships an iteration speedup (e.g. lands a persistent-cache bind-mount, a parallel-device orchestrator, a module-only rebuild helper), the change MUST be classified as universal (per §11.4.17 — most are) OR project-specific (rare — e.g. a script that hardcodes D3 / D4 / Arvus IP 192.168.4.185). Universal speedups belong in this constitution submodule's helper inventory (or in a separate Containers submodule per §11.4.76); project-specific wrappers stay in the consuming repo.

Classification: universal (per §11.4.17). Applies to every project under this Constitution. The disciplines are language-/platform-/build-system-agnostic.

Canonical authority: constitution submodule `Constitution.md` §11.4.82.

Non-compliance is a release blocker. No escape hatch — no `--skip-phase1-forensic`, `--no-pre-flight`, `--rebuild-everything-always`, `--unlimited-subagent-scope`, `--ignore-locks`, `--no-telemetry` flag exists. The disciplines exist because their absence has been forensically demonstrated to cost the operator hours per session.

2 0 2 6 0 5 2 6

§ . . — Endless-loop autonomous work + zero-idle agent dispatch + anti-bluff testing mandate (User mandate, - -)

Forensic anchor — verbatim User mandate (2026-05-26):

"when we say to you that all work MUST BE continued in endless loop until there is no any open items, no unfinished workable items from our Issues docs, or from Continuation document or any unfinished work by agents at the moment it means that you will ALWAYS run until there is nothing workable elft to be done in endless loop, fullyautonomously. You will spawn agents or agents-driven work whenever that is possible or required! Not a single agent or main work stream will sit idle except if it waits for the results of something - some execution and similar! All work MUST BE always covered with comprehensive tests - all supported tests for the project which produce real proofs for the contexts and functionalities they are testing! Bluff of any kind is not allowed and all work and tests MUST WORK completely as proofs (phisical) driven and in complete anti-bluff manner!"

When an operator instructs an AI agent to "continue in endless loop fully autonomously" — or any semantically-equivalent phrasing — the agent MUST interpret this as a HARD-CONTRACT execution covenant covering five non-negotiable obligations:

(A) Endless-loop continuation. The agent MUST continue working in the autonomous-loop framework until ALL of the following are simultaneously TRUE:

- docs/Issues.md Status-column has zero In progress / Ready for testing / In testing / Reopened entries (§11.4.15 closed-set values).
- docs/CONTINUATION.md §3 "Active work" section is empty.
- No background subagent is mid-execution (TaskList reports all tasks completed or pending with pending items NOT being prerequisite-blockers for anything actionable).
- No external dependency is in-flight (build, flash, push, sync).

If any of those conditions is FALSE, the agent MUST continue working — by claiming the next item from the priority queue, by dispatching a subagent for parallelisable work, or by polling a background task that will be the wake signal. An agent that schedules a wake-up and ends the turn while Issues.md still has actionable items, OR while CONTINUATION.md §3 still has active work, OR while a subagent is still working productively, is in violation.

(B) Zero-idle agent dispatch. When the agent identifies parallelisable work that does not contend on the same file-scope as the main work stream, the agent MUST dispatch a background subagent rather than serialising the work. The main work stream + every background subagent operate concurrently. "Wait for the results of something" is the ONLY acceptable reason for an agent (main or subagent) to be idle. The §11.4.20 + §11.4.58 + §11.4.70 subagent-driven covenants compose with this anchor.

(C) Comprehensive test coverage with real (physical) proofs. Every workable item closed in the loop MUST land four-layer test coverage per §11.4.4(b) — pre-build gate + on-device test + paired §1.1 meta-test mutation + HelixQA Challenge bank entry — and every test PASS MUST cite a captured-evidence artefact per §11.4.5 + §11.4.69. The phrase "physical proofs" is interpreted as: captured audio (`tinycap` WAV with RMS + `ffprobe` channels), captured video (screen recording + `ffprobe` frames + recording-analyzer event match), captured network state (`dumpsys` + sink-side probe per §11.4.13), captured UI state (`uiautomator` dump per §11.4.48), captured sysfs state (`/sys/...` / `/proc/...` snapshots). Metadata-only PASS, configuration-only PASS, absence-of-error PASS, grep-without-runtime PASS are all critical defects regardless of the green summary line.

(D) Anti-bluff guarantee end-to-end. The §11.4 covenant family (§11.4.1 FAIL-bluffs / §11.4.2 recorded-evidence / §11.4.6 no-guessing / §11.4.7 demotion-evidence / §11.4.27 no-fakes-beyond-unit / §11.4.50 deterministic-consistency / §11.4.52 autonomous-validation / §11.4.68 audio sink-side / §11.4.69 universal sink-side / §11.4.83 docs/qa/ end-user evidence) is the operative truth-discipline of the endless-loop covenant. Every closure narrative **MUST** cite captured-evidence; every demotion **MUST** cite same-conditions retest; every status transition **MUST** cite the deterministic-consistency baseline. **Tests AND Challenges (HelixQA) are bound equally** — a Challenge that scores PASS on a non-functional feature is the same class of defect as a unit test that does. The User mandate explicitly invokes the historical forensic anchor: "we had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can't be used! This **MUST NOT** be the case and execution of tests and Challenges **MUST** guarantee the quality, the completion and full usability by end users of the product!"

(E) Termination criteria. The endless-loop terminates **ONLY** when:

1. ALL conditions in (A) are simultaneously TRUE (zero open items + zero active CONTINUATION work + zero subagents in flight + zero external dependencies in-flight); OR
2. The operator explicitly issues a `STOP / END LOOP / pause autonomous work` instruction; OR
3. The host-session-safety covenant (§12 + §12.6 + §12.7 + §12.8 + §12.9 + §12.10) demands suspension to protect the operator's session; OR
4. The agent has scheduled a wake-up to resume against a known-future-actionable signal (e.g. CI completion, build completion, hardware-attended phase).

No escape hatch. No `--idle-OK` , `--skip-endless-loop` , `--bluff-permitted-for-this-task` , `--metadata-only-test-suffices` , `--no-physical-proof-required` flag exists. The covenant is uniform across every iteration.

Composes with §11.4 (anti-bluff covenant — historical forensic anchor of the project), §11.4.1 (FAIL-bluffs), §11.4.2 (recorded-evidence), §11.4.4 (four-layer coverage), §11.4.5 (audio + video 5-layer quality), §11.4.6 (no-guessing), §11.4.7 (demotion-evidence), §11.4.20 + §11.4.58 + §11.4.70 (subagent-driven defaults),

§11.4.27 (no-fakes-beyond-unit + 100% test-type coverage), §11.4.42 (iteration-discipline), §11.4.43 (TDD-fix), §11.4.50 (deterministic-consistency), §11.4.52 (autonomous-validation), §11.4.68 (audio sink-side), §11.4.69 (universal sink-side), §11.4.83 (docs/qa/ end-user evidence), §11.4.85 (stress + chaos), §11.4.86 (roster/corpus auto-sync), §12 family (host-session safety boundaries), §12.10 (CONTINUATION-document maintenance — the source-of-truth state the loop checks against).

Pre-build gate `CM-COVENANT-114-87-PROPAGATION` enforces this anchor literal in every `CLAUDE.md` / `AGENTS.md` / `QWEN.md` across the canonical fleet (parent + consumer submodules). Paired meta-test mutation strips the literal `11.4.87` → gate FAILs.

Canonical authority: this Constitution.md §11.4.87 in the HelixConstitution submodule (`git@github.com:HelixDevelopment/HelixConstitution.git`) — inherited per §11.4.35 by every consuming project's repo-root `CLAUDE.md` / `AGENTS.md` / `QWEN.md`.

Non-compliance is a release blocker regardless of context. Skipping the loop, idling without a waited-on signal, or accepting bluff-evidence PASS is severity-equivalent to a §11.4 PASS-bluff at the project-execution layer.

§ . . — Background test execution mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-27):

"Any tests we are executing, especially long test cycles, MUST BE performed in background in parallel with main work stream! This MUST NOT block our capabilities to work on queued workable items (and tackle main Issues document(s)). Make sure we strictly follow this starting now! Main work stream can be blocked or sit iddle only if absolutely needed and if it depends hard on results of some background execution."

Forensic incident. 2026-05-27 the conductor invoked `bash device/rockchip/rk3588/tests/pre_build_verification.sh` synchronously with a `timeout 360` foreground wrapper — even though the operator had previously mandated §11.4.87

endless-loop zero-idle. The 6-7 minute test run blocked the entire main work stream from progressing on §JV/§JW/§JX/§JY scaffolding + further Issues.md drains. This is the exact anti-pattern §11.4.87(B) prohibits at the test-execution layer (analogous to §11.4.88 at the push-execution layer).

The mandate.

(A) Long-running tests run detached. Any test cycle expected to exceed ~30 seconds — `pre_build_verification.sh` , `meta_test_false_positive_proof.sh` , `test_all_fixes.sh` , `recent_work_validate.sh` , HelixQA Challenge banks, on-device 4-phase cycles, full-suite retests per §11.4.40, `audio_loop_supervisor.sh`, `dual_display_record.sh` harnesses, `verify-all-constitution-rules.sh` — MUST be spawned via `nohup ... > <log> 2>&1 & + disown` and `<log>` placed under a known directory (e.g. `qa-results/<test_id>_<timestamp>.log`) for later inspection.

(B) Main stream proceeds on queued items. The conductor MUST immediately return to the §11.4.42 priority queue (open Issues.md items, queued PWUs, doc-sync work, audit work, anything not depending hard on the in-flight test's exit code). "Wait for results" is the ONLY acceptable idle reason per §11.4.87(B) — and even then, the conductor checks back on the test result via short polling (`tail` / `mtime` / `exit-status` file) rather than blocking on `wait` .

(C) Hard-dependency gating. A subsequent step that genuinely requires the test's exit code (e.g. `commit_all.sh` refuses if meta-test is in flight per §11.4.84; release-tag gate per §11.4.40 step 4) MUST express that dependency explicitly — poll the `exit-status` file or check `pgrep -af <test-name>` before proceeding. If the test is still running and the next step needs its result, surface that to the operator per §11.4.66 with interactive options.

(D) Failure surface + cleanup. Backgrounded test failures land in their `<log>` files. The next autonomous-loop tick MUST check that directory (analogue of §11.4.88(D)). Test-script-level cleanup (per §11.4.14 + §11.4.84 + this anchor's restore discipline) is non-negotiable — `trap '<cleanup>' EXIT INT TERM` in every long-running test that mutates global state (mid-mutation backup files, sysfs nodes, ALSA states, etc.).

(E) Pre-empt the blocking pattern. The conductor MUST NOT invoke `bash <long-test>.sh` synchronously when a backgrounded variant suffices. Foreground execution is permitted ONLY when: (a) the test completes in <30 s, OR

(b) the operator explicitly requests foreground execution for a specific step. Synchronous execution of a multi-minute test absent operator authorisation is a §11.4.89 violation regardless of how much "context-cache benefit" the foreground path appears to offer.

(F) Concurrency control. Multiple backgrounded tests of the SAME script MUST serialise via per-script flock (e.g.

`.git/.pre_build.lock` , `.git/.meta_test.lock`). Concurrent invocations of DIFFERENT tests run in parallel.

Composes with §11.4.42 (iteration discipline — backgrounding tests recovers the operator's wait-time per §11.4.82(F) parallel multi-device testing), §11.4.66 (interactive clarification — surface hard-dependencies as options instead of blocking), §11.4.82 (iteration-speedup discipline — §11.4.89 is the dominant blocker class at the test layer, analogous to §11.4.88 at the push layer), §11.4.84 (working-tree quiescence — backgrounded mutation tests still need quiescence checks), §11.4.85 (stress + chaos — long stress runs are quintessential §11.4.89 candidates), §11.4.87 (endless-loop zero-idle — §11.4.89 is the implementation seam at the test-execution layer that makes §11.4.87(B) genuinely zero-idle for the test blocker class), §11.4.88 (background-push — §11.4.89 is the symmetric anchor for tests).

Pre-build gate `CM-COVENANT-114-89-PROPAGATION` enforces this anchor literal across the canonical fleet. Pre-build gate

`CM-BACKGROUND-TEST-EXECUTION-WIRED` verifies that long-running test invocations in helper scripts use `nohup ... & + disown` pattern (or equivalent). Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

(G) Subagent long-op discipline + conductor-owns-inherently-long-ops (extension, 2026-07-15). Forensic FACT (2026-07-15): background subagents were repeatedly KILLED by the runtime's no-progress stream watchdog (~600 s with no tool-call output) because each had run ONE synchronous long operation inside a single tool call (a full pre-build sweep / a long test suite / a device flash / a hung network call) — the host and the network were ruled out as causes; the watchdog fired on the SILENCE, not on any failure. Therefore: (1) a SUBAGENT MUST NOT run any single synchronous operation expected to exceed the runtime's no-progress watchdog budget (conservatively ~5 min) — it MUST background the operation per clause (A) (`nohup ... & + disown`) and POLL with short, frequent tool calls (each poll is the progress the watchdog observes); (2) inherently-long

operations that a subagent CANNOT usefully background-and-poll (full artifact builds, full test suites, device flashes / deploys) are the CONDUCTOR's to run — the conductor (main session) carries no such watchdog and MUST own them, dispatching subagents only for the parallelisable slices around them; (3) a subagent killed by the stall watchdog is a CRASH per §11.4.147 — its work is NOT lost, NOT done, NOT forgotten: the registry entry flips to `crashed` and the unit is `RESPAWNED` until `complete` ; a stall-kill silently absorbed as a completion is a §11.4.147 + §11.4 bluff. Honest boundary (§11.4.6): the watchdog budget is a `RUNTIME` property — read it from the runtime, never guess it; when it is unknown, treat any operation > ~5 min as long.

Canonical authority: this Constitution.md §11.4.89 in the HelixConstitution submodule (`git@github.com:HelixDevelopment/HelixConstitution.git`).

Non-compliance is a release blocker. Synchronous long-test execution (without operator authorisation) is severity-equivalent to a §11.4 PASS-bluff at the project-execution layer. No escape hatch beyond explicit per-invocation operator authorisation.

2026 05 27

§ . . – Zero-idle priority-first
parallel-by-default operating mode (User
mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-27):

"We MUST NEVER sit iddle / wait or sleep if there is possibility for us to work on something like in this case! Continue all work further in parallel if possible fully autonomously! We MUST not waste time since work scope we have is huge! ... Always check if there is a possibility to work on something while we are not working actively on something! Pick always by priority - most critical workable items and other tasks MUST BE done first! ... Everything we can should be done fully in background, should not affect stability of the System or create problems, should be done using subagents-driven approach! ... Stay still / iddle if nothing is left to be done at all or waiting for something that is blocking us / you!!!"

§11.4.94 is the **operating-mode reaffirmation** that binds §11.4.87 (endless-loop), §11.4.20 / §11.4.70 (subagent-driven), §11.4.42 (iteration-priority), §11.4.58 (parallel PWU), §11.4.72 (audio top-priority), §11.4.82 (iteration-speedup), §11.4.88 (background-push), §11.4.89 (background-test) into a SINGLE always-on enforcement contract:

The mandate (binding on the conductor + every subagent + every helper script):

(A) Idle-only-when-genuinely-blocked. The conductor MAY sleep / wait / be idle ONLY when (i) every priority-queued workable item is genuinely blocked on external dependency (operator hardware action, network upstream, hardware-attended forensic, build/test completion that the conductor cannot accelerate), OR (ii) the operator has issued an explicit STOP, OR (iii) host-session-safety demands it per §12. "I don't see what to do next" is NEVER a valid idle reason — the conductor MUST recheck the priority queue per §11.4.42 step 1 and §11.4.87(A) termination criteria.

(B) Always check parallel-work feasibility at every pause point. Before any wake-up schedule / before any "waiting for X" / before any sleep call, the conductor MUST: (1) survey the priority queue per §11.4.42 SCOPE LOCK + §11.4.72 audio-first + §11.4.87(B) zero-idle, (2) identify ALL items that can progress without contending with the in-flight blocker, (3) dispatch them in parallel per §11.4.20 / §11.4.70 (subagent-driven if non-trivial) + §11.4.58 (PWU disjoint scope) + §11.4.89 (background long tests). Only AFTER step (3) returns "no parallel work available" may the conductor schedule a wake/sleep.

(C) Priority order MANDATORY at every pick. Per §11.4.42 priority axis + §11.4.72 audio-top-priority axis. Pick the highest-Severity / highest-priority item the conductor can autonomously progress. Lower-priority items wait until higher-priority queues drain or block on external dependency.

(D) Subagent-driven by default for non-trivial scope. Per §11.4.20 + §11.4.70. The conductor remains the integration + commit + push seam (per §11.4.58 + §11.4.84 quiescence + §11.4.88); subagents execute the analyses + scaffolds + tests + migrations in parallel. Inline execution permitted only when the task is trivial (single-file edit, sub-300-line diff) OR the operator explicitly requested it.

(E) Background by default for long-running work. Per §11.4.85 stress + §11.4.89 background test + §11.4.88 background push. Any operation expected to exceed ~30 s wall-clock MUST run detached (nohup + disown) with log under `qa-results/<op-id>_<ts>.log`. The conductor returns to the priority queue immediately and polls back when needed.

(F) Stability-preserving — no compromise. Parallel work MUST NOT compromise system stability or create new problems. Composes with §11.4.92 multi-pass change-evaluation — every parallel branch's output passes 5-pass before integration. Composes with §11.4.84 working-tree quiescence — concurrent subagents in same checkout coordinate via lockfile or git worktree. Composes with §12.6 / §12.7 / §12.8 / §12.9 host-session safety — parallel work respects 60% memory ceiling and -j2 AOSP cap.

(G) Status updates as work progresses. Per operator mandate "Let us know how goes this mandatory approach". The conductor surfaces parallel-work catalogue + in-flight branches + recent closures in compact summary form whenever the operator asks "how are we progressing" OR at natural milestone boundaries (after each constitution-anchor landing, after each commit batch, after each subagent return).

Anti-pattern explicitly forbidden:

- Scheduling a wake-up without first surveying the parallel-work queue per (B).
- Returning "nothing to do — sleeping" when non-trivial priority items remain that the conductor could autonomously progress.
- Serialising work that could run in parallel without justification.
- Picking lower-priority items while higher-priority ones remain progressable.
- Doing work foreground that should run via subagent or background.
- Letting parallel work introduce instability or regression.

Pre-build gate `CM-COVENANT-114-94-PROPAGATION` enforces this anchor literal across the canonical fleet. Pre-build gate `CM-PARALLEL-WORK-AUDIT` (when implemented) audits the last N wake-up schedule decisions for surfaced parallel-work-queue evidence — wake-ups without the survey are §11.4.94 violations. Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Composes with §11.4.20 (subagent-driven) / §11.4.42 (iteration discipline) / §11.4.58 (parallel-development PWU) / §11.4.70 (subagent-driven default) / §11.4.72 (audio top-priority) / §11.4.82 (iteration-speedup) / §11.4.84 (working-tree quiescence) / §11.4.85 (stress + chaos) / §11.4.87 (endless-loop zero-idle) / §11.4.88 (background-push) / §11.4.89 (background-test) / §11.4.92 (multi-pass evaluation) / §12.6 + §12.7 + §12.8 + §12.9 (host-session safety).

Canonical authority: this Constitution.md §11.4.94 in the HelixConstitution submodule.

Non-compliance is a release blocker. Conductor idle without exhausting the parallel-work queue is severity-equivalent to a §11.4 PASS-bluff at the project-execution layer — it implies the conductor has progressed the project when it has not. No escape hatch beyond the (A) genuinely-blocked / explicit-STOP / host-safety triad.

2026 05 27

§ . . — **Maximum-use-of-idle-time**
mandate + progress-update cadence (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-05-27):

"keep it working, we should do as much as possible, if not it all but as much as we can as long as there is iddle time! it MUST be used! make sure that it is clear in constitution! keep us updated about all progress and all phisycal proofs and gathered data as you progress through all open workable items!"

§11.4.97 strengthens §11.4.87 (endless-loop) + §11.4.94 (zero-idle priority-first) + §11.4.96 (safe-parallel-with-build catalogue) with an **explicit maximum-use-of-idle-time mandate** AND a **progress-update cadence mandate**:

(A) Maximum-use idle-time mandate. Every minute of conductor idle time during which (i) the conductor has work it could autonomously progress AND (ii) is NOT genuinely blocked on external dependency (operator hardware action, network upstream, host-session safety, build/test completion) is a §11.4.97 violation. "As

much as possible, if not it all but as much as we can" is the operative phrasing — the conductor MUST dispatch work continuously through the entire idle window, not just sample-check at scheduled wakes.

(B) Progress-update cadence mandate. The conductor MUST emit operator-facing progress updates at the following natural milestone boundaries (no operator prompt required):

- Every commit that lands (HEAD advance) — 1-line "what just landed".
- Every subagent return — 1-line "subagent X returned, integrated".
- Every constitutional anchor landed — 1-line "§X.Y.Z propagated to 6 remotes".
- Every gathered physical proof — 1-line "captured evidence at qa-results/".
- Every milestone closure (Issues → Fixed migration, Obsolete classification, etc.) — 1-line.

Progress updates are CONCISE (1-3 lines each); the operator's mental model stays current without conductor overhead.

(C) Continuous physical-proof gathering mandate. Per §11.4.5 + §11.4.6 + §11.4.69 — every autonomous-fixable item the conductor closes MUST gather positive captured-evidence: ADB probes, dumphsys captures, screencaps, ALSA hw_params, sysfs reads, file-content hashes, runtime metric snapshots. The evidence is committed alongside the closure narrative. Per §11.4.93 SQLite-SSoT, the evidence path goes into the `item_history.evidence_path` column when DB integration lands.

(D) Per-anchor composition. §11.4.97 binds together: §11.4.5 (captured-evidence quality), §11.4.6 (no-guessing — every progress claim cites evidence), §11.4.13 (sink-side evidence when applicable), §11.4.20 (subagent-driven default for parallel scope), §11.4.27 (no-fakes-beyond-unit — closures cited with real captured evidence), §11.4.42 (priority queue), §11.4.50 (deterministic consistency — N=3 where applicable), §11.4.52 (autonomous validation), §11.4.69 (sink-side taxonomy), §11.4.70 (subagent-driven), §11.4.72 (audio top-priority), §11.4.83 (docs/qa transcript), §11.4.85 (stress + chaos), §11.4.87 (endless-loop), §11.4.88 (background-push), §11.4.89 (background-test), §11.4.94 (zero-idle), §11.4.96 (safe-parallel catalogue).

(E) Idle-only-when-genuinely-blocked. Same closed-set as §11.4.94(A): operator STOP, external dependency the conductor cannot accelerate, host-session-safety demand. Per §11.4.96 the safe-during-build catalogue (A-K) covers what's progressable; the conductor MUST exhaust the catalogue before scheduling sleep.

Operator-facing visibility. §11.4.97's progress-update cadence is the contract by which the operator stays informed without prompting. The conductor does NOT wait for operator status checks — milestone updates emit autonomously.

Pre-build gate `CM-COVENANT-114-97-PROPAGATION` enforces this anchor literal across the canonical fleet. Pre-build gate `CM-IDLE-TIME-AUDIT` (when implemented) audits the conductor's wake/sleep schedule across recent sessions for unjustified idle periods. Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Composes with every anchor in §11.4 (this is the operating-mode capstone):
§11.4.5 + §11.4.6 + §11.4.13 + §11.4.20 + §11.4.27 + §11.4.42 + §11.4.50 +
§11.4.52 + §11.4.69 + §11.4.70 + §11.4.72 + §11.4.83 + §11.4.85 + §11.4.87 +
§11.4.88 + §11.4.89 + §11.4.94 + §11.4.96.

Canonical authority: this Constitution.md §11.4.97 in the HelixConstitution submodule.

Non-compliance is a release blocker. Conductor idle when progressable work remains, OR conductor failing to emit milestone progress updates, is severity-equivalent to a §11.4 PASS-bluff at the operating-mode layer.

2026 05 28

§ . . – Autonomous-decision-over-blocking mandate (User mandate, – –)

Short tag: `autonomous-decision-over-blocking` .

Forensic anchor — verbatim user mandate (2026-05-28):

"when working in endless working loop fully autonomously try to decide most properly about points which would block execution and wait for us. If we haven't answered now work would be blocked whole night! If possible and if that will not cause any issues make proper and most reliable and safe decision so we achieve maximal efficiency and work gets fully done!"

The mandate. When operating in an autonomous / endless-loop mode (per §11.4.87), the agent **MUST** minimize operator-blocking and instead make the safe, reliable, reversible decision itself — so work is NOT stalled (e.g. overnight) waiting for input. The §11.4.87 endless-loop covenant told the agent to keep working; §11.4.101 tells it **HOW** to handle the decision points that would otherwise force it to stop and wait. A wrong bias toward blocking-when-it-was-safe-to-decide wastes operator time (the forensic harm: "blocked whole night"); a wrong bias toward deciding-when-it-should-have-blocked risks irreversible harm. The closed-set decision rule below is the precise boundary between the two.

Decision rule (closed-set — the agent MAY proceed autonomously when ALL hold):

(a) the action is **reversible** OR has a captured pre-op backup per §9.2 (hardlinked `.git` mirror, recorded refs/tags/submodule pointers); (b) the agent can **determine the safe choice from captured evidence** per §11.4.6 (no guessing — `LIKELY` / `probably` / `seems` is NOT a safe-choice determination; the evidence must state the correct choice as fact, or the condition is not met); (c) a wrong choice's **blast radius is bounded AND recoverable** — the damage is local, undoable, and does not cascade beyond the work unit; (d) it **composes with the existing covenants** — anti-bluff §11.4 (no PASS-bluff to skip the decision), host-session-safety §12 (no decision that suspends / logs-out / overruns the §12.6 memory budget), data-safety §9 (no destructive op without the §9 protocol).

Block-only-when rule (the agent MUST BLOCK — ask via the §11.4.66 interactive mechanism — ONLY when ALL of the following hold): the action is **irreversible AND high-blast-radius AND the safe choice cannot be determined from evidence**. Canonical block-required examples: external-account state the agent cannot inspect (registration, certification, billing portals), hardware the agent cannot physically access, destructive repository ops without a backup, force-push (which independently requires §9.2 + §11.4.41 authorisation), spending money or

sending messages / data to third parties. `operator-blocked` per §11.4.21 remains the last-resort status and is reached only after this rule fires AND the §11.4.21 self-resolution-exhaustion audit is complete.

Maximize-progress-while-blocked rule. When a block IS unavoidable, the agent MUST still maximize progress on every NON-blocked item in parallel per §11.4.87 + §11.4.94 (zero-idle priority-first) rather than idling — the blocked decision parks one work unit, it does not pause the loop. Posing the §11.4.66 question and continuing other work is the correct behaviour; posing the question and going idle is a §11.4.94 + §11.4.97 violation.

Composes with §11.4.6 (no-guessing — sub-rule (b) IS the no-guessing test applied to the decision itself), §11.4.21 (operator-blocked is the last-resort status the block-rule routes into, after self-resolution exhaustion), §11.4.40 (a release-tag decision is irreversible-and-high-blast-radius → always blocks per the block-rule), §11.4.41 (force-push is an explicit block-required example), §11.4.66 (the interactive-clarification mechanism the block-rule uses), §11.4.87 (the endless-loop covenant §11.4.101 refines — §11.4.87 says keep working, §11.4.101 says how to clear the decision points), §11.4.94 (zero-idle — maximize-progress-while-blocked), §9.2 (backup is the reversibility substitute in sub-rule (a)), §12 (host-safety bounds every autonomous decision).

Classification: universal (§11.4.17) — autonomous-decision-vs-block boundary is a reusable discipline for ANY project running an agent in an endless / unattended loop, hardware- and domain-agnostic.

4-layer coverage per §11.4.4(b). Propagation gate `CM-COVENANT-114-101-PROPAGATION` enforces the literal anchor `11.4.101` across the canonical consumer fleet (parent + owned-submodule `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Paired §1.1 meta-test mutation strips the `11.4.101` literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this `Constitution.md` §11.4.101 in the `HelixConstitution` submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--always-block-on-decision` , `--never-decide-autonomously` , `--skip-decision-rule` , `--block-without-self-resolution` flag exists. The

2026-05-28 user mandate is unambiguous: in autonomous mode the agent makes the safe reversible decision and keeps working — it does NOT block the operator overnight on a choice it had the evidence to make safely.

2026 05 29

§ . . — Mandatory systematic-debugging activation + always-loaded skill-discovery + plugin-dependency availability (User mandate, - -)

Short tag: `systematic-debugging-always-on` .

Forensic anchor — verbatim user mandate (2026-05-29):

"Make sure that we ALWAYS trigger / start the `/superpowers:systematic-debugging` skills when any issues happen! If this is possible to activate and use in this situations out of the box when we spot problems / issues / bugs / misalignments / inconsistencies we MUST activate the skill(s) and make strongest efforts in full in depth analysis / debugging and determine root causes of all problem or obtain relevant data and information we need! ... we MUST make sure that `/using-superpowers` skill is ALWAYS loaded, applied and used! All dependencies (plugins) that Claude Code or other market places are offering MUST BE installed if these are not already available for loading and use!"

This anchor binds three cooperating invariants. Together they make the difference between an agent that **guesses-and-retries** (symptom-patching, "transient", "flaky", "probably a race") and an agent that **investigates-to-root-cause-first** mechanical, every single time a problem surfaces. The §11.4.6 no-guessing mandate said *do not state an unproven cause*; §11.4.102 says *the moment a problem appears, the very first action is to launch the structured root-cause investigation skill — not to propose a fix*.

(A) Mandatory systematic-debugging activation on ANY spotted problem. On ANY spotted issue, bug, test failure, gate failure, regression, misalignment, inconsistency, unexpected behaviour, or contradiction-between-evidence-and-expectation, the agent MUST activate the structured root-cause-investigation skill

— `superpowers:systematic-debugging` on Claude Code, or the platform-equivalent structured-debugging discipline on any other runtime — **BEFORE proposing, writing, or applying any fix**. This is the **Iron Law: NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST**. The investigation **MUST** run its full four-phase arc: (1) **root-cause phase** — read the actual error / failing assertion / divergent evidence, reproduce deterministically, gather the facts the failure exposes; (2) **pattern phase** — identify what class of defect this is, search for the same pattern elsewhere in the codebase; (3) **hypothesis phase** — form a falsifiable hypothesis about the cause, then prove or disprove it with captured evidence (never `LIKELY` / `probably` / `seems` per §11.4.6); (4) **implementation phase** — only after the cause is a **FACT**, design the fix against the proven root cause. Guess-and-retry, symptom-patching, "let me just try changing this and see", and re-running a failed test hoping it passes ("it's probably transient / flaky") **WITHOUT** a completed root-cause investigation are themselves §11.4.102 violations. Real-world signal this anchor closes: declaring a failure `transient` / `flaky` / `intermittent` / `probably a timing issue` **without captured forensic evidence proving that classification** is simultaneously a §11.4.6 (no-guessing) violation AND a §11.4.7 (demotion-evidence) violation — §11.4.102 makes the corrective response mechanical: the classification is forbidden until the systematic-debugging arc has captured the evidence that proves it.

(B) Mandatory always-loaded skill-discovery discipline (`using-superpowers`). The skill-discovery meta-skill — `superpowers:using-superpowers` on Claude Code, or the platform-equivalent skill-index / capability-discovery discipline on any other runtime — **MUST** be loaded and applied at session start and consulted before any task is begun. Its operative rule: **before acting on ANY request, survey the available skills/capabilities; if ANY skill could apply to the task at hand — even at a 1% chance of relevance — that skill MUST be invoked rather than the agent improvising from memory**. This guarantees the agent reaches for the purpose-built discipline (systematic-debugging for a bug, the brainstorming skill before creative work, the test-driven-development skill before writing a fix, etc.) instead of ad-hoc behaviour. Skipping skill-discovery at session start, or improvising a task that a loaded skill was built to handle, is a §11.4.102 violation.

(C) Mandatory plugin / dependency availability. Every skill plugin, marketplace package, or capability dependency that the project relies on — offered by Claude Code, its skill marketplaces, or the equivalent plugin/extension ecosystem of any other runtime the project uses — **MUST** be installed and loadable on the host

BEFORE the dependent work proceeds. A missing plugin/dependency that blocks a mandated skill (e.g. the `superpowers` plugin absent so `systematic-debugging` cannot launch) is a **release-blocker** until it is installed and confirmed loadable. Install mechanism (generic; runtime supplies the concrete command): the runtime's own plugin/marketplace install path — for Claude Code, the in-session `/plugin` marketplace flow (add the marketplace, install the plugin, confirm the skill appears in the available-skills list); for other runtimes, the documented package/extension installer (npm/pip/cargo/the editor's extension manager). The availability check is itself anti-bluff: a plugin claimed "installed" without the dependent skill actually appearing in the loaded-skills/capabilities list is a §11.4.6 (no-guessing) violation — confirm by observing the skill in the live capability list, not by assuming the install succeeded (install-command exit 0 ≠ skill loadable, mirroring the §11.4.80 `npm exit 0 ≠ working binary` lesson).

(D) AUTOMATIC activation — no operator prompt required (User mandate, 2026-07-13). Verbatim operator mandate (2026-07-13): "Make sure we automatically turn on the systematic-debugging plugin and skills every time we work on any issue! Add this as MANDATORY RULE to our root constitution Submodule and all of its relevant files!" The clause-(A) systematic-debugging activation is AUTOMATIC and the STANDING DEFAULT — the agent MUST auto-activate `superpowers:systematic-debugging` (or the platform-equivalent structured-debugging discipline) on ANY spotted issue / bug / test failure / gate failure / regression / misalignment / inconsistency / unexpected behaviour WITHOUT waiting for an operator prompt, request, or per-issue confirmation. It is engaged the MOMENT an issue is detected — the operator MUST NOT have to ask for it, request it, or re-enable it per issue or per session (composes §11.4.126 default-autonomous-loop, which promotes the standing-default working mode from the first prompt of a session; §11.4.102(D) applies that same always-on-by-default discipline to systematic-debugging activation). Before any such investigation proceeds, the `superpowers` plugin + the `superpowers:systematic-debugging` AND `superpowers:using-superpowers` skills MUST be confirmed installed + loadable per clause (C) — observed in the live capability list, never assumed (§11.4.6 anti-bluff). "I did not activate systematic-debugging because the operator did not ask for it" is itself a §11.4.102(D) violation: clause (A) is NOT a prompt-gated option, it is a standing mechanical default engaged automatically the instant a problem surfaces. §11.4.102(D) STRENGTHENS clause (A) — it makes the

activation trigger prompt-independent and standing-default; it does not weaken it: every clause-(A) four-phase-arc and Iron-Law ("NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST") requirement continues to apply in full.

Composes with §11.4.4 (test-interrupt-on-discovery — §11.4.4 says STOP the cycle the moment a defect is discovered; §11.4.102 says the FIRST action after stopping is to launch systematic-debugging), §11.4.6 (no-guessing — §11.4.102(A) is the procedural enforcement that produces the facts §11.4.6 demands; the forbidden `LIKELY` / `transient` / `flaky` vocabulary is exactly what guess-and-retry emits), §11.4.7 (demotion-evidence — a FAIL → lower-severity demotion requires same-conditions evidence, which is precisely what the systematic-debugging arc captures), §11.4.8 (deep-web-research — runs inside the hypothesis phase), §11.4.43 (TDD-fix — the RED test is written against the root cause the investigation proved), §11.4.70 (subagent-driven — a systematic-debugging investigation is a natural subagent dispatch), §11.4.82(A) (iteration-speedup — already mandates Phase-1 forensic before any speculative patch; §11.4.102 generalises that from "source patch" to ANY fix and adds the always-loaded skill-discovery + plugin-availability invariants), §11.4.92 (multi-pass change-evaluation — the investigation feeds Pass 1 + Pass 4), §11.4.126 (default-autonomous-loop — §11.4.102(D)'s always-on-by-default systematic-debugging activation mirrors §11.4.126's from-first-prompt standing-default working mode).

Classification: universal (§11.4.17) — root-cause-before-fix, skill-discovery-before-task, and dependency-availability are reusable disciplines for ANY project on ANY runtime that exposes a skill/plugin ecosystem; the concrete skill names (`superpowers:systematic-debugging` , `superpowers:using-superpowers`) are the Claude Code instantiation, with the generic "or platform-equivalent" wording carrying the rule to other runtimes per §11.4.81 cross-platform-parity reasoning.

4-layer coverage per §11.4.4(b). Propagation gate `CM-COVENANT-114-102-PROPAGATION` enforces the literal anchor `11.4.102` across the canonical consumer fleet (parent + owned-submodule `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Paired §1.1 meta-test mutation strips the `11.4.102` literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.102 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-systematic-debugging`, `--guess-and-retry-OK`, `--symptom-patch-permitted`, `--skip-skill-discovery`, `--plugin-optional`, `--missing-plugin-is-warning` flag exists. The 2026-05-29 user mandate is unambiguous: the moment a problem is spotted the structured root-cause investigation skill activates, `using-superpowers` is always loaded, and every mandated plugin/dependency is installed and loadable — there is no path by which a fix is proposed before the cause is proven.

2026 06 03

§ . . — **No-silent-removal-of-existing-components-without-operator-confirmation** mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-03):

"Never ever remove any application, system component or service from already existing codebase / System without interactively asked question to us! THIS IS MANDATORY RULE / CONSTRAINT!"

Forensic case study (this project, FACT — captured). During the 1.1.8-dev burn-down, two shipped capabilities — F2 (an Apple-TV-class application) and F4 (a Huawei HMS / Mobile-Services component) — were removed from the existing System WITHOUT first asking the operator. The operator reversed both removals. A removal that the operator has to discover and reverse after the fact is a defect of the same severity class as a §11.4 PASS-bluff: the System silently lost a user-facing capability the operator never agreed to drop.

The mandate. No application, system component, service, package, feature, driver, module, library, prebuilt asset, or any other already-existing capability of the existing codebase / shipped System may be removed (deleted, dropped from the package set, disabled-into-non-shipping, un-bundled, de-listed, or otherwise made unavailable to the end user) WITHOUT FIRST interactively asking the operator and receiving an EXPLICIT keep-or-remove decision. The question MUST be posed through the platform's interactive clarification mechanism per §11.4.66

(`AskUserQuestion` on Claude Code; the documented interactive prompt on other runtimes) — NEVER a free-text "should I remove X?" buried in narrative, NEVER a silent removal justified post-hoc, NEVER an autonomous removal decision. A silent removal is a **release blocker** regardless of how well-intentioned the rationale (deduplication, "it was broken anyway", "geo-restricted", "incompatible", "superseded") — the operator decides, the agent asks.

What counts as a removal (non-exhaustive). Deleting an app/APK/binary from the build's package set (e.g. a project's `PRODUCT_PACKAGES` / `device.mk` / equivalent manifest), removing a service from the init/boot/service-registry set, dropping a kernel module / driver / config from the shipping configuration, unbundling a prebuilt asset, deleting a submodule or a submodule's shipped output, removing a feature flag that gated a live capability, or any edit whose NET EFFECT is "an end-user capability that shipped before no longer ships." Adding, replacing-with-equivalent-or-better (where the operator has approved the replacement), or fixing a capability is NOT a removal. When uncertain whether an edit constitutes a removal, treat it AS a removal and ask (per §11.4.6 no-guessing + §11.4.101 block-only-when-truly-needed — removal of an existing user-facing capability is high-blast-radius and MUST be operator-confirmed, never autonomously decided).

Composition. §11.4.122 composes with §11.4.66 (interactive-clarification — the removal question is asked via the platform interactive mechanism, not free-text), §11.4.101 (autonomous-decision-over-blocking — a removal of an existing user-facing capability is exactly the irreversible-ish, high-blast-radius class §11.4.101 RESERVES for an operator question rather than an autonomous decision; §11.4.122 makes the "ask the operator" outcome mandatory for this class, NOT optional), §11.4.90 (Obsolete classification — the tracked path to actually DROP a capability is: ask the operator → operator approves → mark the item `obsolete` (→ `Fixed.md`) with `Obsolete-Details` reason `feature-removed` AND an operator-approval citation → then remove; the removal never precedes the operator's yes), §11.4.112 (structural-impossibility won't-fix — even a capability proven structurally impossible to keep is closed via the §11.4.90 path with operator awareness, not silently deleted), §11.4.6 (no-guessing — "the operator probably wants this gone" is the exact guess forbidden), §11.4.40 / §11.4.42 (full-suite-retest / iteration-discipline — a removal landed without the operator's yes fails the release gate).

Enforcement. Propagation gate `CM-COVENANT-114-122-PROPAGATION` enforces the literal anchor `11.4.122` across the consumer fleet (every `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Recommended gate `CM-NO-SILENT-COMPONENT-REMOVAL` (a diff-driven detector that flags a removal of a shipped capability — a deletion from the project's package/service/module manifest set — without a matching operator-approval citation in the commit / tracker entry, per the §11.4.110 change-impact-clash pattern; a removal whose tracker item is `Obsolete` (→ `Fixed.md`) with reason `feature-removed` + operator-approval citation passes). Paired §1.1 meta-test mutations: strip the `11.4.122` literal → propagation gate FAILs; introduce a manifest removal with no operator-approval citation → `CM-NO-SILENT-COMPONENT-REMOVAL` FAILs. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§11.4.17) — "never silently remove an existing end-user capability without operator confirmation" is a platform-neutral discipline reusable by ANY project that ships a set of user-facing capabilities; the consuming project supplies its concrete capability-manifest paths (package set, service registry, module config) per §11.4.35.

Canonical authority: this `Constitution.md` §11.4.122 in the `HelixConstitution` submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
11 4 126
--remove-without-asking , --silent-removal , --autonomous-removal-OK ,
--dedup-removal-exempt , --it-was-broken-anyway flag exists.
2026 06 04
```

§ . . — Default autonomous-loop
working mode from first prompt (User mandate,
- -)

Forensic anchor — verbatim user mandate (2026-06-04):

"Make sure that you continue work in endless fully autonomous loop, do not stop until new fully validated and verified version (tag) is created and published (all submodules and main repo) or IN A CASE OF some other main stream work until it is fully completed with all side work streams and nothing else is left in our working queue! THIS MUST BE ALWAYS the default working mode without us asking you! We tend to achieve ABSOLUTE EFFICIENCY, with this and all other projects which will incorporate this MANDATORY RULE / CONSTRAINT!!! This way of (your) working will be ALWAYS applied / followed / executed / fully respected, as soon as we assign / send first request (prompt) in the session! This stops only if we explicitly say so or nothing is left to be done in current working scope (release that will come / upcoming version)!!! Any mimicking (imitation) of this behavior / rules / mandatory constraints, false results or any kind of bluff(s) is ABSOLUTELY FORBIDDEN!!!!"

The endless fully-autonomous loop is the **DEFAULT working mode**, engaged automatically the moment the operator sends the FIRST request / prompt of a session — the operator MUST NOT have to ask for it, request it, restate it, or re-enable it per session. §11.4.87 framed the endless-loop covenant as an explicit-instruction opt-in ("continue in endless loop fully autonomously" or a semantically-equivalent phrasing); §11.4.126 is the **capstone** that promotes the same covenant to always-on: from the first prompt onward, every agent operates in the §11.4.87 loop discipline as the standing default, with §11.4.94 zero-idle, §11.4.97 maximum-idle-use, §11.4.101 autonomous-decision-over-blocking, and §11.4.103 continuous-parallel-stream all engaged by default — no per-session activation handshake.

The continuation contract (the loop continues until ONE of the two terminal conditions holds): (A) **Release scope** — a new, fully-validated-and-verified version (tag) is created AND published across all owned submodules AND the main repo to all configured remotes (per §2.1 multi-upstream push + §11.4.40 full-suite-retest-before-tag + §11.4.113 absolute-no-force-push merge-onto-latest-main). (B) **Non-release main-stream scope** — the main-stream work goal is fully completed AND every side work stream is done AND the working queue holds nothing left to

do for the current scope. Until ONE of (A) or (B) holds, the agent MUST keep working — it claims the next priority item, dispatches the next parallel stream, or progresses every non-blocked item, per §11.4.42 / §11.4.72 / §11.4.94 / §11.4.103.

The loop STOPS ONLY on: (1) the operator explicitly saying so (STOP / pause / end); (2) nothing left to do in the current working scope — the upcoming release / the current main-stream goal — with the queue genuinely empty per the (A)/(B) terminal conditions; (3) a §12 host-session-safety demand (the loop yields to host safety unconditionally). Idle-while-blocked parks one work unit, it does not stop the loop — the agent keeps progressing every non-blocked item in parallel per §11.4.101 + §11.4.94 + §11.4.97.

Goal — ABSOLUTE EFFICIENCY. The discipline exists to drive the project to completion with no operator-side restart overhead, no idle gaps, no stop-and-wait round-trips. It applies to this project AND every project that incorporates this Constitution (universal per §11.4.17).

Anti-bluff — actually do the work, never imitate it. Mimicking / imitating this loop behaviour, narrating continuation without performing it, fabricating progress, or emitting false / bluff results of ANY kind is ABSOLUTELY FORBIDDEN — this composes the entire §11.4 anti-bluff covenant family (§11.4 / §11.4.1 / §11.4.2 / §11.4.5 / §11.4.6 / §11.4.50 / §11.4.69 / §11.4.107). The agent MUST genuinely perform the continuous work and capture positive evidence for every closure; a report that claims the loop ran while no real work / no captured evidence was produced is a §11.4 PASS-bluff at the operating-mode layer, the severity class this mandate specifically prohibits.

Classification: universal (§11.4.17) — a platform-neutral operating-mode discipline reusable by ANY project; the consuming project supplies its concrete release-tag / publish mechanism + working-queue source per §11.4.35. Composes with §11.4.87 (the endless-loop covenant — §11.4.126 promotes it from opt-in to always-on default) / §11.4.94 (zero-idle priority-first parallel-by-default) / §11.4.97 (maximum-idle-use + progress-update cadence) / §11.4.101 (autonomous-decision-over-blocking) / §11.4.103 (continuous parallel-stream routine) / §11.4.66 (interactive clarification ONLY when a decision is genuinely operator-blocked per §11.4.101) / §11.4.6 (no-guessing — terminal conditions determined from captured queue/tracker state, never assumed) / §11.4.40 (release-tag full-suite retest is the (A) terminal gate) / §11.4.42 (iteration discipline is the loop body) / §11.4.72 (audio-first priority within the loop) / §11.4.113 (publish via merge-onto-latest-main, no force-

push) / §2.1 (multi-upstream publish) / §12 (host-session safety overrides the loop). Propagation gate CM-COVENANT-114-126-PROPAGATION (literal 11.4.126 across the consumer fleet) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.126 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

11 4 127
--ask-before-continuing , --single-turn-only , --not-default-loop , --mimic-OK flag exists.
2026 06 06

§ . . — Session-handoff resumption-prompt mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-06):

"make sure that in situations like this now when new session is needed you ALWAYS prepera such sentence - which will be valid for particular moment and the phase of the project and enough for work to continue."

When the agent determines a fresh session is needed (context-window limits, performance degradation) OR the operator asks whether a new session is needed / requests a handoff, the agent MUST ALWAYS prepare and provide a ready-to-paste **resumption prompt valid for that EXACT moment and project phase** — self-contained enough that pasting it into a fresh session resumes work with ZERO loss. The agent MUST proactively offer the resumption prompt the moment it recognises the need (not only when asked). Two variants on demand: a SHORT first-sentence variant ("Read <handoff docs> , then continue <terminal goal> ...") AND a FULL detailed block. The prompt MUST: (1) point to the live handoff doc(s) — .remember/remember.md if present + docs/CONTINUATION.md per §12.10 — and instruct reading them FIRST + git fetch --all ; (2) state the current PHASE + the immediate NEXT action + the terminal goal; (3) embed the exact current-state anchors (build IDs / artifact MD5, device/target serials, commit HEAD, in-flight PIDs + log paths, captured-evidence paths); (4) restate the binding constraints (anti-bluff §11.4 covenant, no-force-push §11.4.113, exact version/

naming, hardware/target gotchas); (5) be MOMENT-VALID — reflect the actual live state, NEVER a generic template. The handoff doc(s) MUST be brought current BEFORE the prompt is given (§12.10). A handoff that omits the resumption prompt, or gives a stale / generic one, is a §11.4.127 violation.

Classification: universal (§11.4.17) — a platform-neutral session-continuity discipline reusable by ANY project worked by context-bounded agents; the consuming project supplies its concrete handoff-doc paths + state anchors per §11.4.35. Composes with §12.10 (CONTINUATION doc — the resumption prompt is its operator-facing entry point) / §11.4.6 (no-guessing — the embedded state is captured fact, not assumed) / §11.4.66 (interactive clarification — "do you need a new session?" is itself a §11.4.66 decision point) / §11.4.87 / §11.4.103 / §11.4.126 (the fresh session CONTINUES the always-on autonomous loop — the resumption prompt is HOW the loop survives a context reset). Propagation gate CM-COVENANT-114-127-PROPAGATION (literal 11.4.127 across the consumer fleet) + recommended gate CM-HANDOFF-RESUMPTION-PROMPT-PRESENT + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.127 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

`--skip-handoff-prompt` , `--generic-prompt-OK` , `--no-resumption-sentence` , `--handoff-without-state` ^{1 1 4 1 2 9} ^{2 0 2 6 0 6 0 6} flag exists.

§ . . — Huge-blocker release protocol (User mandate, — —)

Forensic anchor — direct user mandate (2026-06-06): when a huge blocker is discovered during release validation we MUST stop all testing, return to fixing ALL discovered issues, process all recorded data from the last session, land rock-solid fixes, author NEW validation+verification tests of ALL supported test types, rebuild, reflash, and RESTART the full validation+verification of every fix/change from the last release tag to now — on both devices in parallel, recorded, with real physical captured proofs and no bluff.

When a HUGE BLOCKER (a release-blocking-severity defect — a core user-facing capability broken, a regression that invalidates the in-flight validation cycle, or a defect whose blast radius reaches the batch's other fixes) is discovered during release validation, the agent **MUST** execute the following protocol in order, with NO partial / spot-check shortcut: (1) **STOP all testing** immediately on every device (the §11.4.4 test-interrupt-on-discovery STOP, escalated to release granularity — a validation cycle that continues past a huge blocker produces "all green" summaries on a known-broken build, the exact §11.4 PASS-bluff). (2) **Return to fixing ALL discovered issues** — not only the huge blocker but every defect surfaced this cycle; investigate each to root cause per §11.4.102 systematic-debugging, isolate regressions against the last known-good tag per §11.4.114. (3) **Process all recorded data from the last session** — analyse the §11.4.128 raw-corpus slice for the cycle (this is the §11.4.128(3) release-prep analyse-trigger) so no defect the recordings already captured is missed. (4) **Land rock-solid fixes** — per §11.4.123 (rock-solid proof or deep research), §11.4.43/§11.4.115 (RED-on-broken-artifact then GREEN), §11.4.9 batch-source-fixes-before-rebuild. (5) **Author NEW validation+verification tests of ALL supported test types** — per §11.4.27 (unit / integration / e2e / full-automation / security / stress / chaos / performance / UI / UX / Challenges / HelixQA) + §11.4.85 stress+chaos, each anti-bluff with captured evidence + paired §1.1 mutation. (6) **Rebuild + reflash** — full rebuild (NOT module-only) for retest-baseline integrity, deploy to a CLEAN target per §11.4.108 (no stale-overlay shadow). (7) **RESTART the full validation+verification from the last release tag to now** — re-validate EVERY fix/change landed since the last tag (per §11.4.40 full-suite-retest-before-tag — the huge-blocker discovery invalidated the in-flight cycle, so the cycle restarts from the beginning, NOT resumes), running on both/all owned devices IN PARALLEL per §11.4.103/§11.4.119 (single-resource-owner partitioning), every device's run RECORDED per §11.4.128, producing REAL physical captured proofs per §11.4.5/§11.4.69/§11.4.107 with no bluff of any kind. This anchor BINDS the existing release-cycle anchors for the huge-blocker case — it adds the STOP → fix-all → process-recordings → new-tests-all-types → rebuild → reflash → full-restart sequencing and the "restart, never resume" rule, citing §11.4.4 + §11.4.40 + §11.4.108 + §11.4.114/§11.4.115 + §11.4.123 + §11.4.128 rather than duplicating their text.

Classification: universal (§11.4.17) — a platform-neutral release-discipline reusable by ANY project that runs a release-validation cycle; the consuming project supplies its concrete device set, build/flash mechanism, and last-known-good tag per

§11.4.35. Composes with §11.4.4 (test-interrupt-on-discovery — §11.4.129 is its release-granularity escalation) / §11.4.40 (full-suite retest before tag — §11.4.129 mandates the RESTART, not resume) / §11.4.42 (iteration discipline) / §11.4.9 (batch-source-fixes-before-rebuild) / §11.4.27 (all-test-types coverage) / §11.4.85 (stress+chaos) / §11.4.102 (systematic-debugging) / §11.4.108 (clean-target rebuild+reflush+runtime-signature) / §11.4.114 + §11.4.115 (regression isolation + RED-on-broken-artifact) / §11.4.123 (rock-solid proof) / §11.4.128 (process the recorded corpus) / §11.4.103 + §11.4.119 (parallel recorded multi-device).

Propagation gate `CM-COVENANT-114-129-PROPAGATION` (literal `11.4.129` across the consumer fleet) + recommended gate `CM-HUGE-BLOCKER-FULL-RESTART` (a release-blocker discovery during validation triggers a from-the-last-tag full restart, not a resume) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.129 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--resume-after-blocker , --spot-validate-after-fix ,  
1 1 4 1 3 0  
--skip-recording-analysis , --skip-new-tests , --module-only-after-  
blocker , --single-device-restart flag exists.  
2 0 2 6 0 6 0 6
```

§ . . — Post-remediation validate-the-fix-FIRST-after-redeploy (User mandate, — —)

Forensic anchor — direct user mandate (2026-06-06): when a blocker discovered during release validation is fixed and a new artifact is produced + redeployed, we MUST first re-test the SPECIFIC last-failing features and validate+verify the just-incorporated fixes BEFORE proceeding to the broader / full validation — a first fix attempt may not work, may be incomplete, or may regress again, so confirming the targeted fix on the new artifact first prevents wasting a full multi-hour cycle on a still-broken fix.

When a blocker / critical failure is discovered during release validation and then FIXED, and a new artifact is produced (rebuild / new flashing image / redeploy) and the target is reflashed / redistributed / updated, the agent MUST execute the following ordering with NO shortcut: (1) **Re-test the SPECIFIC last-failing**

features FIRST — run the targeted guard tests for exactly the defects this fix addressed, BEFORE any broader / full-suite validation. (2) **Validate + verify the just-incorporated fixes with real captured evidence** — the RED test that reproduced the defect on the pre-fix artifact (per §11.4.115) MUST now flip GREEN at `RED_MODE=0` on the new artifact, AND the fix's declared runtime signature (per §11.4.108) MUST verify on the CLEAN target the redeploy produced. Metadata-only / config-only / absence-of-error / grep-without-runtime PASS are forbidden (§11.4 / §11.4.1) — the proof is captured per §11.4.5 / §11.4.69 / §11.4.107 / §11.4.123. (3) **Only after the targeted fix is CONFIRMED working on the new artifact** does the agent proceed to the §11.4.40 full retest of every fix/change from the last release tag to now.

Rationale. A first fix attempt may not work, may be incomplete, or may regress again under the new artifact. Confirming the targeted fix FIRST catches a fix-did-not-take case immediately — instead of discovering it hours later at the END of a full multi-hour validation cycle (and then having to restart the whole cycle per §11.4.129). The targeted re-test is cheap; the full cycle is expensive; ordering cheap-confirmation-first is the §11.4.82 iteration-speedup discipline applied to the post-blocker reflash. This is the §11.4.46 recent-work-validation gate specialised for the post-blocker-reflash case: the §11.4.129 huge-blocker protocol's step (7) full-restart MUST be preceded by this targeted-fix-first confirmation phase.

Honest boundary (§11.4.6): "the fix probably took" is not "the fix took" — the targeted RED → GREEN flip + runtime-signature verification on the new artifact is the captured proof; proceeding to the full cycle without it is a §11.4.130 violation. A targeted re-test that still FAILs on the new artifact is a finding — return to the §11.4.114/§11.4.115 isolate → RED → fix loop and produce another artifact; do NOT proceed to the full cycle on a still-broken fix.

Classification: universal (§11.4.17) — a platform-neutral release-discipline reusable by ANY project that fixes a blocker, redeploys, and re-validates; the consuming project supplies its concrete targeted-guard-test set, build/flash mechanism, and clean-deployment / runtime-signature observables per §11.4.35. Composes with §11.4.4 (test-interrupt-on-discovery — the blocker that triggered the fix) / §11.4.40 (full-suite retest before tag — §11.4.130 is the targeted-first phase that GATES the full retest) / §11.4.46 (validate-recent-work-before-post-flash-tests — §11.4.130 is its post-blocker-reflash specialisation) / §11.4.108 (runtime-signature-as-definition-of-done on the clean target the redeploy produced) / §11.4.114 (regression isolation

— a still-failing targeted re-test re-enters the isolation loop) / §11.4.115 (RED-baseline + polarity-switch — the RED test flips GREEN on the new artifact) / §11.4.123 (rock-solid proof — captured evidence, never "probably took") / §11.4.129 (huge-blocker protocol — §11.4.130 is its post-reflash targeted-confirmation phase that precedes the step-7 full-restart) / §11.4.82 (iteration-speedup — cheap-confirmation-first). Propagation gate `CM-COVENANT-114-130-PROPAGATION` (literal `11.4.130` across the consumer fleet) + recommended gate `CM-FIX-FIRST-AFTER-REDEPLOY` (after a post-blocker redeploy, the targeted last-failing-feature guard tests run + pass on the new artifact before the full retest starts) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.130 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-targeted-retest`, `--full-cycle-first`, `--assume-fix-took`, `--validate-fix-at-end`, `--skip-red-green-flip-on-new-artifact` flag exists.

§ . . — Standing session-resumption
file mandate (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-06-07): "Make this markdown a standard file which will be written EVERY TIME when we need fresh session out of the box! It MUST BE always up to date and in sync so whenever new session is created all we have to do is just point to it!"

Every project MUST maintain a SINGLE canonical, always-current **session-resumption file** at a fixed, project-declared standard path (the project declares the exact path once per §11.4.35 and never moves it without a §11.4.66 operator decision). This file is the OUT-OF-THE-BOX entry point for any fresh session: creating a new session requires ONLY pointing the new agent at this one file. §11.4.131 promotes §11.4.127 (which mandated PREPARING a resumption prompt on demand) into a STANDING, version-controlled ARTIFACT — the resumption prompt is no longer an ephemeral in-message reply but a persistent file that is ALWAYS present and ALWAYS in sync.

(A) **Existence + fixed path.** The file MUST exist at the declared canonical path at all times; its location is encoded as a literal path string in the project-layer instantiation (§11.4.35) and never silently moved.

(B) **Always written + always synced.** The file MUST be (re)written whenever a fresh session is needed OR whenever the live project state materially changes (new HEAD, build/artifact id, phase transition, device/target state, in-flight job, blocking decision) — the §12.10 CONTINUATION.md trigger set applied to the resumption file. A stale resumption file (describing a prior HEAD / phase / artifact) is a §11.4.131 violation of the same severity class as a §12.10 stale-CONTINUATION violation: it silently breaks the zero-loss-resumption guarantee.

(C) **Content (composes §11.4.127).** The file MUST carry both the §11.4.127 SHORT (one-paste first-sentence) AND FULL (paste-ready block) variants; point to the live handoff docs read FIRST (`.remember/remember.md` if present + `docs/CONTINUATION.md` per §12.10) + `git fetch` ; embed the exact live-state anchors (HEAD commit, build/artifact ids + checksums, device/target serials, in-flight PIDs + log paths, captured-evidence paths); state the current PHASE + immediate NEXT action + terminal goal; and restate the binding constraints (anti-bluff §11.4, no-force-push §11.4.113, exact version/naming, hardware/target gotchas). It MUST be MOMENT-VALID, never a generic template (§11.4.6).

(D) **Export + freshness.** The file is in §11.4.65 scope → synchronized `.html` / `.pdf` siblings refreshed on every update; a §11.4.44 revision header tracks its monotonic revision + last-modified timestamp.

(E) **Out-of-the-box resumption.** A fresh session, given ONLY this file's path, MUST be able to fully resume with zero additional context. The file IS the session bootstrap.

Composes §12.10 (CONTINUATION.md live-state — §11.4.131 is the dedicated paste-ready-prompt sibling of it) / §11.4.127 (resumption-prompt discipline promoted to a standing file) / §11.4.65 (universal export) / §11.4.44 (revision header) / §11.4.6 (moment-valid, never generic) / §11.4.66 (path change is an operator decision) / §11.4.126 (default-loop continuity across sessions).

Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-131-PROPAGATION` (literal `11.4.131` across the consumer fleet) + recommended gate `CM-SESSION-RESUMPTION-FILE-PRESENT` (the declared file exists + carries

SHORT+FULL + revision header + is fresh vs HEAD) + paired §1.1 meta-test mutation (strip the file / make it stale → gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.131 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
11 4 132
--skip-resumption-file , --ephemeral-prompt-only ,
--stale-resumption-OK , --generic-template-OK flag exists.
2026 06 07
```

§ . . — Risk-ordered validation
priority mandate (User mandate,
- -)

Forensic anchor — verbatim user mandate (2026-06-07): "We MUST ALWAYS first test and validate features, functionalities and fixes/changes that have been worked most recently, the ones which were most problematic, which have the most chance to crash or break again, the ones which have been re-opened the most times! Then, after we validate and verify all this with real (physical) proofs and hard evidence, with no false results and bluffs of any kind, we continue with all other existing tests in the test suites! This IS MANDATORY."

Tests, validations, and verifications MUST run in **RISK-DESCENDING order** — the highest-risk set FIRST, and only AFTER that set is fully GREEN with real (physical) captured evidence does the remainder of the suite run. The risk ranking is computed from a closed set of factors, highest-risk first:

(a) **Most-recently-worked** — features / functionalities / fixes / changes touched most recently (the freshest code is the least-exercised in the wild and the most likely to carry an undiscovered defect).

(b) **Historically most-problematic** — the items with the longest defect history, the most prior fixes, the most prior failures.

(c) **Highest crash/break/regress likelihood** — items whose blast radius, complexity, or dependency surface gives them the greatest chance to crash or break again.

(d) **Most-reopened** — items with the highest reopen count per §11.4.55 (a high reopens-count is the strongest empirical signal of fragility — a defect that keeps coming back).

The highest-risk set, ordered by (a)–(d), is validated FIRST. Each item in that set MUST pass with real (physical) captured evidence per §11.4.5 / §11.4.69 / §11.4.107 — no metadata-only / config-only / absence-of-error / grep-without-runtime PASS (§11.4 / §11.4.1), no false results, no bluff of any kind (§11.4.6). ONLY AFTER the entire highest-risk set is GREEN with captured proof does the agent proceed to all other existing tests in the test suites. Running the full suite in arbitrary order, or running lower-risk tests before the highest-risk set is GREEN, is a §11.4.132 violation: it spends the operator's scarce validation window on the items least likely to fail while deferring the items most likely to fail.

§11.4.132 REFINES / STRENGTHENS §11.4.130 (post-remediation validate-the-fix-FIRST-after-redeploy — §11.4.130 says re-test the just-fixed items first; §11.4.132 generalises "first" to the full risk-ordered set) + §11.4.46 (validate-recent-work-before-post-flash-tests — §11.4.46 gates the full suite on recent-work GREEN; §11.4.132 adds the explicit risk-ordering WITHIN the recent / high-risk set) + §11.4.42 (iteration-discipline priority order — §11.4.132 applies the same priority discipline to VALIDATION ordering, not just implementation ordering). The risk ranking is the validation-layer analogue of the §11.4.42 / §11.4.72 implementation-layer priority queue.

Classification: universal (§11.4.17) — a platform-neutral validation-ordering discipline reusable by ANY project; the consuming project supplies its concrete recency / problematic-history / reopen-count sources (e.g. the §11.4.93 workable-items DB `reopens_count` + `last_modified` columns) per §11.4.35. Composes §11.4.4 / §11.4.5 / §11.4.6 / §11.4.7 / §11.4.40 / §11.4.42 / §11.4.46 / §11.4.50 / §11.4.55 / §11.4.69 / §11.4.107 / §11.4.130. Propagation gate `CM-COVENANT-114-132-PROPAGATION` (literal `11.4.132`) + recommended gate `CM-RISK-ORDERED-VALIDATION-PRIORITY` + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.132 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-risk-ordering`, `--any-order-OK`, `--suite-order-fixed` flag exists.

§ . . – Target-System + hardware safety mandate (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-06-08): "Make sure that all changes we do to the System are ALWAYS safe for the System itself and for the hardware the system runs on! This is MANDATORY."

Every change to the TARGET system — firmware, kernel, init / boot scripts, drivers, sysfs / devfreq / voltage / clock / thermal / regulator register writes, partition layout / bootloader / U-Boot, HAL, framework, prebuilts, device configuration, or any equivalent on a non-Android target — MUST ALWAYS be safe for BOTH: (a) **the target System itself** — the change MUST NOT brick the device, induce a boot-loop, corrupt user / system data, or render the device unrecoverable; AND (b) **the hardware the System runs on** — the change MUST NOT drive any hardware-control surface beyond its safe electrical / thermal / voltage / clock envelope, and MUST NOT damage panels, storage, radios, regulators, or any other physical component.

Concrete obligations (ALL hold):

1. **Reversible-first.** Prefer reversible changes; when a change is irreversible or high-blast-radius (bootloader / U-Boot / partition layout / one-way fuse / OTP), verify it against a known-good value (MD5 / checksum / golden manifest) BEFORE applying, and capture a pre-op backup of any state the change overwrites (per §9.2). A wrong irreversible write that bricks the device is the worst-case the mandate exists to prevent.
2. **No unverified hardware-control writes.** NEVER write an unverified value to a hardware-control sysfs node / register / firmware field that governs voltage, clock frequency, regulator output, thermal-throttle threshold, current limit, or any other physical-safety parameter. Every such value MUST be within the component's datasheet / vendor-documented safe range; the safe range MUST be established as FACT (cited datasheet / vendor reference per §11.4.8 + §11.4.99), never guessed (§11.4.6). Pinning a max OPP / disabling a throttle / raising a regulator output without proving the cooling design and electrical envelope tolerate it is forbidden.

3. **Thermal / performance changes respect the cooling design.** Any change that raises sustained power, frequency, or voltage (e.g. forcing a performance governor, pinning the top OPP, disabling idle / thermal management) **MUST** be validated against the device's actual cooling capacity — captured thermal evidence (sustained-load temperature $\leq$ the documented limit) is the proof; "it boots" is not proof it is thermally safe.
4. **Flashing uses the sanctioned tool + verified image.** Deploying firmware **MUST** use the project's sanctioned flashing tool against a freshly-built, integrity-verified image — never an ad-hoc partition write, never a stale / unverified artifact, never a tool / image whose provenance is unestablished.
5. **Unprovable-safety $\Rightarrow$ blocked.** A change whose safety for the target System and its hardware cannot be established from captured evidence **MUST** be treated as UNSAFE and blocked — proceed only after the safe outcome is determined as FACT (§11.4.6), the change is reversible OR backed up (§11.4.101 reversible-first + §9.2), and (when the safe choice still cannot be determined) the operator is asked via §11.4.66 / §11.4.101's block-only-when rule. "Probably safe" is the exact guess this mandate forbids.

Distinct from §12 host-session safety. §12 (and §12.6 / §12.7 / §12.8 / §12.9 / §12.10) protects the DEVELOPER's HOST machine and session (no host suspend / logout / OOM cascade / unbounded-memory build). §11.4.133 protects the TARGET device the System is built FOR and runs ON. Both apply; neither weakens the other. A build step can be host-safe (bounded memory, no host suspend) yet still ship a target-unsafe change (an over-voltage regulator write) — §11.4.133 is the gate for the latter.

Classification: universal (§11.4.17) — a platform-neutral target-safety discipline reusable by ANY project that produces firmware / kernel / driver / device-config changes deployed onto target hardware; the consuming project supplies its concrete hardware-control surfaces, datasheet-safe ranges, known-good bootloader / image hashes, and sanctioned flashing tool per §11.4.35. Composes §12 (host-session safety — sibling-not-overlapping protection scope) / §11.4.6 (no-guessing — the safe range / known-good value is FACT, never assumed) / §11.4.101 (reversible-first + autonomous-decision-over-blocking — irreversible high-blast-radius hardware writes are the class that escalates to an operator question) / §11.4.108 (the change's runtime signature on a clean target proves it landed AND is safe) / §11.4.123 (rock-solid proof — safety claims need captured

evidence, never a metadata-only assertion). Propagation gate `CM-COVENANT-114-133-PROPAGATION` (literal `11.4.133` across the consumer fleet) + recommended gate `CM-TARGET-HARDWARE-SAFETY` + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.133 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

`--unsafe-hardware-write`, `--skip-system-safety`, `--brick-risk-accepted` flag exists.

2026 06 11

§ . . — Mandatory deep multi-angle web research per change/issue, before declaring fixed or structural (User mandate, - -)

Forensic anchor — verbatim user mandate (2026-06-11):

"For every single issue we fix or improvement we make — besides tight systematic-debugging, fixing, code review by independent agents, comprehensive tests — we MUST ALWAYS do deep web research from various angles! No matter how big/small/simple/complex, dig deep the internet for articles, technical documentation, APIs, open-source code — EVERYTHING that can help make the best possible solution OR confirm we don't have a more serious problem we're unaware of! We MUST do everything possible to STOP the constant issue-reopening and finally start closing items as fixed+working! Do this ALWAYS in parallel with the main work stream! Add the mandatory rules into the constitution Submodule CONSTITUTION.md, CLAUDE.md, AGENTS.md, QWEN.md and other relevant files. Commit + push to all upstreams. Apply ASAP to every workable item."

Forensic case study (FACT, 2026-06-11). A subtitle-on-secondary-display goal had been verdicted `won't-fix: structurally-impossible` (the §11.4.112 secure-surface pixel-blanking reasoning). A subsequent DEEP multi-angle web

research pass OVERTURNED that verdict — it surfaced the real, workable OCR-of-the-PRIMARY-display path (read the caption on the introspectable primary surface rather than the FLAG_SECURE secondary), proving the goal was NOT structurally impossible, only mis-classified for want of research. A single research pass from a different angle converted a "we can't" into a shipped capability. This is the exact reopen-cycle the mandate exists to break: items churn between Fixed ↔ Reopened, or sit mis-labelled `structural / won't-fix`, because the solution-space (or the confirmation that the closure is genuinely safe) was never deep-researched from enough angles.

For EVERY fix / improvement / change / closure — no matter how big, small, simple, or complex — the agent MUST, IN ADDITION to tight systematic-debugging (§11.4.102), the fix itself, independent-agent code review (§11.4.125 / §11.4.142), and comprehensive multi-layer tests (§11.4.4(b) / §11.4.40), perform a DOCUMENTED **deep multi-angle web research pass** that digs the internet from VARIOUS angles — articles, official + vendor technical documentation, API references, standards, issue trackers, maintainer guidance, reusable open-source code — to BOTH (i) discover the best possible solution AND (ii) confirm there is NOT a more serious underlying problem the team is unaware of. This BINDS + STRENGTHENS the existing research/proof anchors into one mandatory, reopen-breaking, run-in-parallel discipline; it does not re-author them:

- **(A) No closure-as-fixed/structural WITHOUT a documented deep-research pass.** No issue may be marked `Fixed / Implemented / Completed` (§11.4.33) OR classified `structurally-impossible won't-fix` (§11.4.112) UNTIL a deep multi-angle research pass is documented for it. §11.4.8 already requires deep research before designing a non-trivial fix and §11.4.123 already requires it when validation is unclear; §11.4.150 makes the pass UNCONDITIONAL — required for EVERY workable item, however trivial, at the closure AND structural-verdict gates — precisely because the forensic case proved a "trivial-looking" structural verdict was wrong.
- **(B) Multiple angles, not a single lookup.** The pass MUST interrogate the problem from several distinct angles (e.g. official-docs / standards / known-bug-trackers / alternative-approach / failure-mode / security / performance / platform-constraint / open-source-precedent — the consuming project picks the

applicable subset, ≥ 2 genuinely-distinct angles), so a single-source confirmation-bias miss (the §11.4.145 LLM-confirmation-bias failure mode) cannot pass as research. A one-link drive-by is NOT a deep multi-angle pass.

- **(C) Confirm-no-bigger-problem, not just find-a-fix.** The pass MUST explicitly seek evidence that the chosen fix does NOT mask a deeper defect AND that the closure (or the `structural` verdict) is genuinely safe — the "do we actually have a more serious problem we're unaware of?" check. Absence of a found bigger problem is recorded as a positive finding only when the search that would have surfaced it is enumerated (composes §11.4.118 discovery-completeness — "we found nothing worse" requires "here is what we searched").
- **(D) Latest-source + cited.** Sources MUST be the LATEST authoritative versions per §11.4.99 (never training-data / memory), each cited by URL + access date in the item's research artefact AND the closure commit footer (`Deep-research <date>: <urls>` OR the literal `NO external solution found – original work` per §11.4.8 when the search genuinely yields nothing). Stale or memory-sourced "research" is a §11.4.99 + §11.4.150 violation.
- **(E) Reopen-breaking is the PURPOSE.** The explicit goal is to STOP the constant Fixed $\leftrightarrow$ Reopened churn and finally close items as fixed-AND-working: a reopen (§11.4.34 / §11.4.55) whose root cause a deep multi-angle pass would have surfaced is a §11.4.150 miss; and a `structural` / `won't-fix` verdict reached without the pass is forbidden (the case study's overturned verdict is the canonical anti-pattern). Composes §11.4.7 (demotion-evidence — a reopen IS a demotion, and the research is part of the demotion-prevention) + §11.4.112 (structural-impossibility now REQUIRES the deep-research pass + its cited authorities before the verdict is earned).
- **(F) ALWAYS in parallel with the main stream.** The research pass MUST run as a background, subagent-driven work stream (per §11.4.70 / §11.4.20 / §11.4.103 continuous-parallel-stream + §11.4.89 background execution) concurrent with the main fix/build/test work — it NEVER serialises the main stream or stalls the loop (§11.4.94 / §11.4.97 / §11.4.126). A blocked research angle parks that angle, it does not pause the closure pipeline (§11.4.101).

- **(G) Apply ASAP to every workable item.** The discipline applies retroactively + going forward to EVERY workable item (§11.4.93 / §11.4.95 SSoT); items closed without a documented pass are re-audited as part of the §11.4.40 / §11.4.42 release-gate sweep.

Honest boundary (§11.4.6): a deep-research pass reduces the unknown-unknown surface and breaks the most common reopen causes — it does NOT prove zero remaining defects (§11.4.118) and does NOT replace §11.4.108 four-layer runtime-signature verification, §11.4.125 / §11.4.142 independent review, or §11.4.40 full-suite retest. It is one of MULTIPLE STRONG LAYERS — the research layer that every fix/closure/structural-verdict additionally crosses. "We probably don't have a bigger problem" without the enumerated multi-angle search is a guess (§11.4.6), never a finding.

Classification: universal (§11.4.17) — a platform-neutral research-and-closure discipline reusable by ANY project; the consuming project supplies its concrete research corpora, angle set, item tracker, and closure-commit-footer convention per §11.4.35. Composes §11.4.8 (deep-web-research-before-implementation — §11.4.150 makes it unconditional + multi-angle + closure-gating) / §11.4.99 (latest-source verification of the cited authorities) / §11.4.123 (rock-solid-proof-or-deep-research — §11.4.150 raises "when unsure" to "always") / §11.4.118 (discovery-completeness — the confirm-no-bigger-problem search is enumerated) / §11.4.145 (independent multi-angle impact-research — §11.4.150 is its external-research sibling, adversarial against confirmation bias) / §11.4.125 / §11.4.142 (independent review — research precedes, never substitutes) / §11.4.7 (demotion/reopen evidence) / §11.4.112 (structural-impossibility now requires the cited-authorities pass) / §11.4.34 / §11.4.55 (reopen attribution — the churn this breaks) / §11.4.70 / §11.4.20 / §11.4.89 / §11.4.103 (parallel background subagent stream) / §11.4.40 / §11.4.42 (release-gate re-audit) / §11.4.93 / §11.4.95 (per-item SSoT). Propagation gate `CM-COVENANT-114-150-PROPAGATION` (literal `11.4.150`) + recommended gate `CM-DEEP-RESEARCH-PER-ISSUE` (every closed/structural-verdicted item carries a documented multi-angle deep-research artefact + cited-source closure footer, run in parallel, before the closure/structural verdict is accepted) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; close an item or reach a `structural` verdict with no documented multi-angle research artefact / cited footer → `CM-DEEP-RESEARCH-PER-ISSUE` FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.150. Non-compliance is a release blocker. No escape hatch — no `--skip-deep-research`, `--single-source-suffices`, `--trivial-change-no-research`, `--close-without-research`, `--structural-without-research`, `--serialise-research`, `--research-from-memory-OK` flag.

§ 11.4.151 — Project-prefixed release-tag/version-naming mandate (User mandate, 2026-06-12)

Forensic anchor — verbatim user mandate (2026-06-12):

"Every release tag and version name we create — on the main repository and on every Submodule we own — MUST be prefixed with the project's release prefix, e.g. `myproject-1.0.0-dev-0.0.1`. The prefix MUST come from `HELIX_RELEASE_PREFIX` in our `.env` if it is set, otherwise from the lowercased project root directory name. The SAME prefix MUST be used across the main repo and all owned Submodules in one release so a release is greppable across every repository."

Every release tag AND every version name created on the main repository AND on every owned-by-us submodule (per §11.4.28) MUST be prefixed with the project's release prefix, in the form `<PREFIX>-<version>` (canonical example: `myproject-1.0.0-dev-0.0.1`). The prefix makes a release identifiable + greppable across every repository the release spans, so a single `git tag -l '<PREFIX>-*'` (or equivalent) enumerates the whole release surface in one query.

Prefix resolution order (closed-set, deterministic — §11.4.6 no-guessing):

1. **`HELIX_RELEASE_PREFIX` from the project's `.env`** — authoritative when set. `.env` is git-ignored per §11.4.30; the variable MUST be documented in the tracked `.env.example` placeholder so a fresh clone knows the contract (the regeneration-mechanism discipline of §11.4.77 applies — `.env` is re-obtained, never committed).

2. **Fallback = the lowercased snake_case form of the project root directory name** (no spaces), per §11.4.29 — used whenever `HELIX_RELEASE_PREFIX` is unset/empty. The fallback is deterministic from the checkout, so a release prefix is always resolvable without operator input.

The prefix MUST be IDENTICAL across the main repo and all owned

submodules within a single release — a release that tags the main repo

`<PREFIX>-1.2.0` while tagging an owned submodule `<OTHER>-1.2.0` (or unprefixed) is a §11.4.151 violation: the cross-repo grep no longer enumerates the release. Version codes increment monotonically within the prefix

(`<PREFIX>-...-0.0.1` → `<PREFIX>-...-0.0.2` → ...), never reset, never skipped — same append-only discipline as §11.4.54's ticket ids applied to the version axis.

Honest boundary (§11.4.6): the prefix guarantees a release is identifiable + uniform across every repository, NOT that the release contents are correct — the tag is still only created after the §11.4.40 full-suite retest GREEN and reaches every upstream via the §11.4.113 merge-onto-latest-main path (NEVER a force-push), fanned out per §2.1.

Classification: universal (§11.4.17) — a platform-neutral release-naming discipline reusable by ANY project; the consuming project supplies its concrete prefix value + the `HELIX_RELEASE_PREFIX` env var per §11.4.35. Composes §2.1 (multi-upstream push — the prefixed tag fans out to every remote) / §11.4.29 (lowercase snake_case — the fallback prefix follows it) / §11.4.30 (`.env` git-ignored — the authoritative prefix source) / §11.4.40 (full-suite retest before the tag is created) / §11.4.113 (absolute no-force-push — the prefixed tag reaches upstreams via merge-onto-latest-main) / §11.4.126 (the release-scope terminal condition is a published, prefixed tag). Propagation gate `CM-COVENANT-114-151-PROPAGATION` (literal `11.4.151` across the consumer fleet) + recommended gate `CM-RELEASE-PREFIX-NAMING` (every release tag/version on the main repo + every owned submodule carries the resolved `<PREFIX>-` prefix, identical across the release) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; create an unprefixed or differing-prefix release tag → `CM-RELEASE-PREFIX-NAMING` FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.151. Non-compliance is a release blocker. No escape hatch — no `--no-release-prefix`, `--unprefixed-tag`, `--prefix-optional`, `--differing-submodule-prefix` flag.

2026 06 13

§ . . — Crashlytics-recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-13):

"For every project that has Firebase Crashlytics enabled / wired, we MUST continuously monitor ALL of the Crashlytics-recorded data — crashes, ANRs, performance traces, and non-fatals — systematically debug each, fix and improve, and cover everything with validation and verification tests. This MUST be checked regularly, with no false results and no bluff of any kind!"

Every project that has Firebase Crashlytics enabled / wired (the SDK linked into a shipping artifact, crash + non-fatal + ANR reporting active) MUST treat the Crashlytics console as a first-class captured-evidence channel from real end-user devices and continuously process every datum it records. Crashlytics IS the user-visible-behaviour signal §11.4 exists to protect: a recorded crash / ANR / performance regression / non-fatal is a real end user hitting a broken feature, and a recorded issue left unmonitored, undebugged, unfixed, or uncovered-by-test is the precise "tests pass while the feature is broken in the wild" failure mode the §11.4 covenant forbids — an open Crashlytics issue with a green test suite is a §11.4 PASS-bluff at the field-telemetry layer.

§11.4.152 STRENGTHENS + COMPLETES §11.4.47 (Firebase data review mandate): §11.4.47 owns the periodic REVIEW + dedup + Issue-creation pass that surfaces Crashlytics/Analytics/Performance findings and maps them to tracker entries; §11.4.152 owns what happens to each surfaced item AFTER it is recorded — the systematic-debug → fix/improve → regression-test-coverage lifecycle that drives it to a proven, regression-immune closure. §11.4.47 finds it; §11.4.152 fixes it and proves it stays fixed. Both are mandatory; neither substitutes for the other.

The mandate (ALL must hold):

1. **Continuous monitoring of ALL four Crashlytics surfaces.** The monitoring pass MUST cover fatal crashes, ANRs, performance traces / regressions, AND non-fatals — all four. Skipping any one surface is a §11.4 PASS-bluff (a regression surfaces on the skipped axis and the team never sees it — the §11.4.47 three-source lesson extended to the non-fatal + ANR + performance axes). Non-fatals in particular are the silent class: every catch / fallback / recovered-error path that records a non-fatal is a real degraded user experience and MUST be triaged, not ignored because the app did not crash.
2. **Systematic-debugging of each recorded issue (reproduce-before-fix).** Every Crashlytics-recorded issue selected for fix MUST go through the §11.4.102 systematic-debugging arc — Iron Law: NO FIX WITHOUT ROOT-CAUSE INVESTIGATION FIRST — root-cause → pattern → falsifiable hypothesis proven with captured evidence → fix designed against the proven cause. The §11.4.115 RED-baseline-on-the-broken-artifact reproduces the recorded defect (the recorded stacktrace / ANR trace / non-fatal context IS the §11.4.5/§11.4.69 captured evidence of defect-present) BEFORE the fix; a fix without a prior reproducing test is a §11.4.43/§11.4.123 violation (no falsifiability evidence the fix addresses the root cause vs masks the symptom). Where the recorded issue cannot be reproduced, the deep-research-before-declaring-untestable path (§11.4.123 / §11.4.150) is mandatory before any "cannot reproduce" classification.
3. **A fix / improvement for every confirmed issue.** Every confirmed Crashlytics issue (severity per the §11.4.47 classification table) MUST receive a real fix or improvement landed against its proven root cause (§11.4.9 / §11.4.43 / §11.4.108 four-layer verification on a clean artifact). "Acknowledged in the console" is not a fix; muting / ignoring a recurring issue without a tracked rationale is a §11.4.6 / §11.4.90 violation.
4. **Validation + verification regression-test coverage per closed issue.** Every Crashlytics issue closed by a fix MUST, in the SAME commit as the fix (§11.4.43 DOCUMENT step), register a permanent §11.4.135 regression guard into the standing suite — a §11.4.115 polarity test whose `RED_MODE=1` captures the recorded defect on the pre-fix artifact and whose `RED_MODE=0` is the standing GREEN guard asserting the defect is ABSENT. The guard MUST exercise the same user-reachable code path the Crashlytics stacktrace / trace identifies, with

rock-solid captured evidence per §11.4.5 / §11.4.69 / §11.4.107 / §11.4.123. A Crashlytics issue marked resolved WITHOUT a falsifiable regression test is FORBIDDEN — it is the canonical recurrence vector (the issue silently returns months later, invisible to the green suite), the §11.4.138 operator-escape class applied to field telemetry. Each closure carries the §11.4.34 demotion-evidence + a closure-log entry recording the Crashlytics issue id / console URL, the root-cause analysis, the fix commit, and the validation + verification test paths.

5. **Regular cadence — checked repeatedly, never once.** The monitoring + debug + fix + cover cycle MUST fire on a regular cadence, reusing the §11.4.47 five-trigger set (pre-build / pre-flash / pre-distribute / pre-tag blocking; daily + post-deployment-burn-in non-blocking) so newly-recorded issues are caught while they are cheap. A one-time sweep that is never repeated is a §11.4.152 violation; field telemetry accrues continuously and so must its processing.
6. **No false results, no bluff.** Every verdict in the cycle is captured evidence per §11.4 / §11.4.1 / §11.4.6 — a "no new Crashlytics issues" claim requires the enumerated monitored-surfaces + queried-window as proof (§11.4.118 absence-of-evidence is not evidence-of-absence); a "fixed" claim requires the RED → GREEN flip on the artifact (§11.4.115 / §11.4.130); a regression guard whose paired §1.1 mutation does not make it FAIL is itself a bluff gate.

Honest boundary (§11.4.6): processing every recorded issue reduces the known-field-defect surface and breaks the silent-recurrence vector — it does NOT prove zero remaining field defects (a defect no user has hit yet is not yet recorded), and it does NOT replace §11.4.108 runtime-signature verification or §11.4.40 full-suite retest. It is the field-telemetry-driven layer that complements them; an issue is "closed" only when its regression guard is GREEN on a clean artifact, never when the console mark is flipped.

Classification: universal (§11.4.17) — a platform-neutral field-telemetry-processing discipline reusable by ANY project that wires Firebase Crashlytics (or, by the same shape, any equivalent crash/ANR/perf/non-fatal field-reporting backend); the consuming project supplies its concrete Crashlytics project handle, console-access credential (§11.4.10, never logged), severity table, and per-issue closure-log path per §11.4.35. The reference project-level instantiation is a consuming project's §6.O (Crashlytics-Resolved Issue Coverage Mandate — per-issue validation test + Challenge test +

`.lava-ci-evidence/crashlytics-resolved/<date>-<slug>.md` closure log) +

§6.AC (Comprehensive Non-Fatal Telemetry Mandate — every catch / fallback / recovered-error path records a non-fatal with triage context), which together are the project-specific embodiment of this universal clause. Composes §11.4 / §11.4.1 / §11.4.5 / §11.4.6 / §11.4.9 / §11.4.34 / §11.4.40 / §11.4.43 / §11.4.47 / §11.4.69 / §11.4.90 / §11.4.102 / §11.4.107 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.130 / §11.4.135 / §11.4.138 / §11.4.150 / §1.1. Propagation gate `CM-COVENANT-114-152-PROPAGATION` (literal `11.4.152` across the consumer fleet) + recommended gate `CM-CRASHLYTICS-ISSUE-FULLY-COVERED` (every closed Crashlytics issue carries a systematic-debug root-cause record + a registered §11.4.135 regression guard + a closure-log entry citing the console issue id/URL + the validation/verification test paths; a console-resolved issue with no falsifiable regression guard FAILs) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a Crashlytics issue resolved with no registered regression guard / no closure log → `CM-CRASHLYTICS-ISSUE-FULLY-COVERED` FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.152. Non-compliance is a release blocker. No escape hatch — no `--skip-crashlytics-monitoring`, `--monitor-crashes-only`, `--skip-non-fatals`, `--resolve-without-regression-test`, `--console-mark-is-fixed`, `--monitor-once`, `--mute-without-rationale` flag.

§11.4.172 — Mandatory production-readiness planning with realistic timeline projection (User mandate, 2026-06-29). Every project under this Constitution MUST maintain a living production-readiness planning document that: (a) provides REALISTIC timeline projections for all product milestones (MVP, beta, production-ready) based on actual development velocity data (items completed per week, reopening rate, blocking dependencies); (b) identifies ALL danger zones, weaknesses, and risk areas with mitigation strategies; (c) accounts for HW delivery timelines, licensing delays, and external dependency risks; (d) defines the critical path and identifies which items can slip without affecting MVP; (e) includes workstation/build-capacity analysis and optimization recommendations; (f) is updated at least monthly or whenever the item count changes by $\geq 10\%$; (g) is cross-referenced against the workable-items database for accuracy. The planning document MUST live at `docs/research/production_planning_<YYYYMMDD>/ANALYSIS.md` with HTML+PDF exports per §11.4.65. Honest boundary (§11.4.6): projections are estimates based on measured velocity — they guarantee the methodology is sound, NOT that specific dates will be met. Classification: universal

(§11.4.17). Composes §11.4.6 (no-guessing — projections based on data, not wishful thinking), §11.4.40 (full retest before tag), §11.4.42 (iteration discipline), §11.4.93 (SQLite SSoT), §11.4.108 (four-layer verification). Propagation gate CM-COVENANT-114-172-PROPAGATION (literal 11.4.172) + paired §1.1 mutation.

§11.4.185 — Manual QA-team testing as the mandatory FINAL confirmation of done (User mandate, 2026-07-07). Forensic anchor — verbatim operator mandate (2026-07-07): "We MUST EVERYWHERE and ALWAYS take into account that manual testing as final confirmation step for work done is MANDATORY!!! Without such confirmation by QA team of ours we cannot consider any set of features or scope of work fully completed! This MUST BE well known and mandatory part of our constitution! Fully respected, taken into account, followed and never forgotten!" No set of features / scope of work / release may be considered **fully completed** until it has received MANUAL TESTING confirmation by the project's QA team as the FINAL confirmation step — EVERYWHERE and ALWAYS (every feature, every scope, every release), NEVER forgotten. The mandate (ALL hold): (1) **Necessary-but-not-sufficient automation.** Every automated gate this Constitution mandates — the §11.4 anti-bluff covenant, §11.4.5/§11.4.69 captured evidence, §11.4.52 autonomous validation, §11.4.40 full-suite retest, §11.4.108 four-layer SOURCE → ARTIFACT → RUNTIME → USER-VISIBLE verification, §11.4.132 risk-ordered validation — is NECESSARY but NOT SUFFICIENT for "fully completed." These are the automated confidence gates; the QA-team MANUAL test is the FINAL human sufficiency gate that closes the scope. (2) **Composition with §11.4.52 — ADDS, does not weaken.** §11.4.52 requires every feature to have an AUTONOMOUS validation path (primary, CI-scalable, runs on every commit, survives operator absence) and forbids operator-attended-ONLY validation. §11.4.185 does NOT weaken that requirement — it ADDS a final human QA sign-off ON TOP of the mandatory autonomous path. A feature MUST have BOTH: the autonomous path (§11.4.52, primary + always-on) AND the manual-QA final confirmation (§11.4.185, the closing gate). Manual QA is the FINAL gate, never the ONLY gate — a feature validated ONLY by manual QA with no autonomous path is STILL a §11.4.52 violation; a feature validated ONLY by automation with no manual-QA sign-off is a §11.4.185 violation. Both must hold. (3) **Workflow placement.** The delivery pipeline is: autonomous-validation-GREEN → build → flash/deploy → autonomous on-device validation → **HAND OFF to the QA team for manual testing** → ONLY after QA-team manual confirmation is the scope "fully completed" / eligible for closure or a release tag. The §11.4.126 autonomous-loop

terminal condition (a fully-validated-and-verified release tag) now **REQUIRES** the QA-team manual confirmation before the tag is cut; a scope marked "done" / "fully completed" / migrated to Fixed.md without a recorded QA-team manual confirmation is a §11.4.185 violation — a definition-of-done bluff of the same severity class as a §11.4 PASS-bluff. (4) **Honest boundary (§11.4.6)**. The agent/loop MAY autonomously drive the scope through autonomous-GREEN + build + flash + on-device autonomous validation, but it **MUST** then hand off + **WAIT** for the QA-team's manual confirmation before claiming "fully completed" — it **MUST NOT** self-certify the manual step, simulate it, or infer it from automated evidence alone (that would itself be a bluff at the human-sufficiency layer). While awaiting QA-team confirmation the agent keeps progressing every **OTHER** non-blocked item in parallel (§11.4.87/§11.4.94/§11.4.97/§11.4.101) — the wait parks that one scope, not the loop. Classification: universal (§11.4.17) — a platform-neutral definition-of-done discipline reusable by ANY project that maintains a QA team; the consuming project supplies its concrete QA-team identity, hand-off mechanism, and confirmation-recording location per §11.4.35. Composes with §11.4 (anti-bluff covenant) / §11.4.5 (captured-evidence quality) / §11.4.6 (no-guessing — the manual confirmation is a recorded **FACT**, never assumed) / §11.4.40 (full-suite retest before tag) / §11.4.52 (autonomous-validation mandate — §11.4.185 is its human-sufficiency **CAPSTONE**, adding rather than substituting) / §11.4.69 (sink-side evidence taxonomy) / §11.4.108 (four-layer fix-verification) / §11.4.126 (autonomous-loop terminal condition now gated on QA-team confirmation) / §11.4.129 (huge-blocker release protocol) / §11.4.130 (post-remediation validate-the-fix-first) / §11.4.132 (risk-ordered validation priority). Propagation gate `CM-COVENANT-114-185-PROPAGATION` (literal `11.4.185` across the consumer fleet) + recommended gate `CM-MANUAL-QA-FINAL-CONFIRMATION` (every item marked fully-completed / migrated to Fixed.md carries a recorded QA-team manual confirmation citation) + paired §1.1 meta-test mutation (strip the literal → propagation gate **FAILs**; mark an item fully-completed with no QA-team confirmation citation → the recommended gate **FAILs**; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.185. Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-manual-qa`, `--automation-suffices`, `--self-certify-manual-step`, `--assume-qa-confirmed`, `--autonomous-only-completion` flag.

§ . . – Non-targetable deploy/flash tooling **MUST isolate exactly ONE eligible target and VERIFY-AFTER-WRITE on the INTENDED target (research-derived, – –)**

Forensic anchor (FACT, 2026-07-15). A flashing tool with no device-selection argument AUTO-SELECTED a device from the bus. With a sibling board still attached, it silently wrote the freshly-built image to the WRONG board, printed a SUCCESS message, and the INTENDED board kept its OLD build. Every downstream test then validated a target that had never received the artifact: the §11.4.108 ARTIFACT → RUNTIME-ON-CLEAN-TARGET layer was silently broken while every tool and gate reported green.

Whenever a deploy / flash / install / write tool CANNOT address a specific target (no serial / id / path selector, or one that the procedure does not pass), the procedure MUST satisfy ALL of:

(1) Targetability is a PRECONDITION. Before the write, the procedure MUST establish that the intended target is the one that will be written: EITHER the tool takes an explicit stable target selector (serial / UUID / stable id per §11.4.111 — never an enumeration ordinal) and it IS passed, OR the procedure GUARANTEES that EXACTLY ONE eligible target is present on the bus / host / network at the moment of the write (every sibling disconnected, powered down, or excluded), with the enumerated eligible-target set captured as evidence (size == 1).

(2) The tool's own success message is NOT proof. A tool's "ok" / exit 0 proves only that bytes reached the device the TOOL selected — NEVER that they reached the device the OPERATOR intended (§11.4.6: "the write succeeded" ≠ "the write succeeded ON THE INTENDED TARGET"). Treating a tool's success line as deployment evidence is a §11.4 bluff at the deployment layer.

(3) VERIFY-AFTER-WRITE on the intended target. After the write, the procedure MUST READ BACK the deployed artifact's IDENTITY from the INTENDED target ITSELF (build id / version / checksum / artifact fingerprint) and assert it EQUALS the artifact just written. A mismatch — or the OLD identity still present — is a FAIL, never a warning. This is the §11.4.108 ARTIFACT → RUNTIME assertion applied to the deployment step, and it is the ONLY thing that closes the wrong-device write.

(4) No validation on an unverified target. Any test / validation run against a target whose deployed-artifact identity has NOT been verified is INVALID, and any PASS it produces is a §11.4 PASS-bluff — it exercised code that may never have been deployed there (the §11.4.108 clause-4 stale/wrong-artifact class).

Honest boundary (§11.4.6): isolation + read-back prove that the INTENDED target holds the INTENDED artifact; they do NOT prove the artifact is CORRECT (that remains §11.4.108 / §11.4.40 / §11.4.130). Classification: universal (§11.4.17) — the consuming project supplies its concrete deploy tool, target-identity read-back, and eligible-target enumeration per §11.4.35. Composes §11.4.6 / §11.4.46 / §11.4.108 / §11.4.111 (stable identifier, never an ordinal) / §11.4.119 (single-resource-owner) / §11.4.130 / §11.4.133 (writing an image to an unintended device is a TARGET/HARDWARE-safety event) / §11.4.139. Propagation gate `CM-COVENANT-114-200-PROPAGATION` (literal `11.4.200`) + recommended gate `CM-DEPLOY-TARGET-ISOLATED-AND-VERIFIED` (an explicit stable selector was passed OR the eligible-target set was captured with size == 1; the post-write identity read-back from the intended target is present and matches) + paired §1.1 mutation (deploy with two eligible targets on the bus and no explicit selector, OR strip the post-write identity read-back → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.200.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--auto-select-target-OK`, `--trust-tool-success-message`, `--skip-post-write-verify`, `--validate-unverified-target`, `--multiple-targets-on-bus-OK` flag.

§ . . — Instant multi-stream resume engine: whole-fleet continuation state is a durable, content-addressed, atomically-committed snapshot resumed in O(changed) reads, not an O(history) re-read (User mandate, — —)

Compact summary: the continuation mechanism (§12.10 / §11.4.127 / §11.4.131 / §11.4.205) is EXTENDED — never replaced — by a mechanical engine so a fresh session resumes ALL work (every Track, main stream, and agent) at once, INSTANTLY and at token cost O(changed streams), by reading ONE snapshot

manifest + one compact blob per stream, NEVER re-reading the full handoff history.

(1) **Content-addressed Merkle store** — each stream's canonical state bytes are named by their sha256 (dedup: unchanged streams share a hash → a global fleet snapshot costs $O(\text{changed})$); a snapshot is ONE manifest referencing every stream's content-hash (§11.4.205(4)). (2) **$O(\text{streams})$ resume, $O(\text{history})$ NEVER on the resume path** — `RestoreAll` reads the manifest + one blob per stream; the append-only event ledger is audit + lost-update recovery only, never read to resume (§11.4.127). (3) **Tamper-evident + deterministic** — every blob is integrity-checked on read (bytes MUST hash to their id) and a deterministic round-trip (re-serialize must byte-match), no wall-clock in the hashed content, verified by a SELF-VALIDATING oracle: golden-good PASS + golden-bad detected FAIL + negative-control (a legitimately-older-but-valid snapshot) PASS — the false-positive guard proving the verifier distinguishes a LAGGING copy from a TAMPERED one (§11.4.201/§11.4.107(10)/§11.4.206(3)). (4) **Durable + atomic + single-writer + crash-safe** — atomic ref write (temp → fsync → rename → dir-fsync, §11.4.205(6)), append-only ledger (§11.4.205(6)), single-writer-per-stream refusal (§11.4.206), advisory lock with provably-stale reap (`kill -0 liveness`, never steals a live lock, §11.4.180/§9.2). (5) **Project-agnostic engine, consumer-owned data** — the engine (`github.com/vasic-digital/continuum` , GitLab mirror) carries ZERO project literals, fails closed rather than guess a store path (§11.4.6/§11.4.177), and is inherited BY REFERENCE as a depth-1 submodule under the §11.4.28(C) carve-out (`constitution/submodules/continuum/` , `helix-deps.yaml` , zero own-org deps); every consumer supplies its store path / actor / stream naming as DATA. Anti-bluff (§11.4/§11.4.108): the whole engine ships four-layer coverage (unit + integration + e2e + stress + chaos §11.4.85, race-clean) with a measured resume Metric (`blobs_read == 1 + streams` is the $O(\text{streams})$ proof) and the self-validating oracle as captured evidence (§11.4.5/§11.4.69/§11.4.107); honest boundary (§11.4.6): it snapshots the state an agent REPORTS, not its execution image (not a CRIU/durable-execution runtime, §11.4.112), and `est_tokens` is a documented `bytes/4` heuristic, not a tokenizer count. Classification: universal (§11.4.17). Composes §12.10 / §11.4.127 / §11.4.131 / §11.4.205 / §11.4.116 / §11.4.147 / §11.4.180 / §11.4.201 / §11.4.206 / §11.4.107(10) / §11.4.28 / §11.4.177 / §11.4.85 / §9.2. Propagation gate `CM-COVENANT-114-207-PROPAGATION` (literal `11.4.207`) + recommended gate `CM-CONTINUUM-RESUME-ENGINE-PRESENT` (the engine present under `constitution/submodules/continuum/` , `go test -race ./... green, continuum selfcheck =`

good=PASS/bad=FAIL/negctrl=PASS, zero project literals) + paired §1.1 mutation (strip the literal → propagation gate FAILs; downgrade the resume path to read the ledger, or the oracle to pass its golden-bad → the mechanism gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.207.

Non-compliance is a release blocker. No escape hatch — no `--resume-reads-history`, `--skip-integrity-check`, `--unvalidated-oracle-OK`, `--hand-typed-resume-state`, `--per-project-resume-reimpl` flag.

§ . . — Operator-request-history document: every project maintains a project-local, always-in-sync ledger of every operator request/prompt (content + timestamp+timezone + track + alias + model + effort) (User mandate, — —)

Verbatim operator mandate (2026-07-15): introduce a request-history document that captures, per operator request/prompt — the full request content, the accepted timestamp with an explicit timezone, the track that processed it, the alias that took the workable item, and the model + effort used — always kept in sync, with no false data and no bluff of any kind. **The operator's load-bearing correction (§11.4.35 / §11.4.17):** the RULE (this universal discipline) belongs in the constitution submodule; the request-history DOCUMENT itself (project data) is PROJECT-LOCAL — the document is NEVER placed into the constitution.

Every governed project MUST maintain a single, project-local **operator-request-history document** — an append-only, newest-first ledger of EVERY operator request/prompt the project has received — at a fixed, project-declared standard path (declared once per §11.4.35, e.g. `docs/requests/history.md`, never moved without a §11.4.66 operator decision). Each entry MUST carry EXACTLY these operator-mandated fields: **(1) Request content** — the full/verbatim prompt text where archived, otherwise a faithful complete summary (each entry marks which — verbatim vs summary); **(2) Accepted (when)** — the exact date + time + TIMEZONE the request was accepted, with the timezone stated EXPLICITLY on every entry (the consuming project declares its default zone per §11.4.35; where only a date is durably recoverable the time is marked `time UNKNOWN`, never fabricated); **(3) Track** — the parallel-development track that processed it (the conductor/main track

by default, per-domain tracks where applicable, §11.4.176/§11.4.178/§11.4.182); **(4) Alias** — the live agent/worker alias that took the workable item at acceptance (§11.4.182/§11.4.196); **(5) Model + effort** — the model + effort setting used. **(A) NO INVENTION (§11.4.6)**. A field not durably recoverable is recorded LITERALLY as `UNKNOWN` (with a one-line reason where useful) — NEVER guessed; a request-history entry that fabricates a timestamp/alias/model it never captured is a §11.4/§11.4.1 bluff at the requirements-record layer. **(B) HONEST RECONSTRUCTION BOUNDARY**. For pre-mandate sessions where per-prompt records were not archived, the document is best-effort reconstructed from the durable dated record (governance-mandate provenance dates, session ledgers, handoff docs, VCS history), with the unrecoverable per-prompt fields marked `UNKNOWN` and the reconstruction boundary stated explicitly in the document (§0-style scope note) — never silently implied complete. **(C) ALWAYS IN SYNC + EXPORTED**. The document carries the §11.4.44 revision header, is kept in sync with live state (§12.10 trigger set — every new operator prompt appends a row), and is exported to all mandated formats (§11.4.65/§11.4.73 — `.md` + `.html` + `.pdf` + `.docx` where the four-format mandate applies, each non-`.md` sibling's mtime ≥ the `.md`). **(D) KEEP-APPLYING MECHANISM**. The project ships an append mechanism (a helper script and/or a `UserPromptSubmit` -class hook) that appends a new newest-first row for each NEW operator prompt at capture time with the exact timestamp + track + alias + model + effort (track/alias derived deterministically per §11.4.182, honest ? / `UNKNOWN` when undeterminable); an append helper that landed WITHOUT the automatic-capture hook is honestly a partial mechanism (the hook wiring is a tracked §11.4.197/§12.10 follow-up, never claimed as automatic capture it does not perform — §11.4.6). **(E) COMPLEMENTS, NEVER REPLACES**. The request-history document complements — never replaces — any per-session no-loss requirements ledger (§11.4.197/§11.4.202), the workable-items DB SSoT (§11.4.93/§11.4.95), or the §12.10/§11.4.131 session-resumption artefacts; loss of a requirement at its intake (an operator prompt not recorded) is FORBIDDEN (§11.4.197). **(F) RULE-UNIVERSAL / DOCUMENT-PROJECT-LOCAL (§11.4.35)**. This universal RULE lives here in the constitution submodule; the request-history DOCUMENT, its declared path, and the project's default timezone are consumer-owned project DATA and live ONLY in the consumer layer — the document is NEVER placed into the constitution submodule (§11.4.28 decoupling).

Classification: universal (§11.4.17) — a platform-neutral requirements-record discipline reusable by ANY project; the consuming project supplies its concrete document path, default timezone, and append-helper wiring per §11.4.35.

Composes §11.4.6 (no-guessing — UNKNOWN never a fabricated field) / §11.4.35 (rule-universal / document-project-local split) / §11.4.44 (revision header) / §11.4.65 / §11.4.73 (four-format export) / §11.4.118 (enumerated-coverage honesty for the backfill) / §11.4.131 / §12.10 (always-in-sync live-state artefact) / §11.4.176 / §11.4.178 / §11.4.182 (track + alias derivation) / §11.4.196 (alias tracking) / §11.4.197 (no requirement lost) / §11.4.202 (report → workable-item — the request-history is the human-readable request ledger; §11.4.202 is its DB/tracker materialisation) / §11.4.93 / §11.4.95 (workable-items SSoT). Propagation gate CM-COVENANT-114-208-PROPAGATION (literal 11.4.208 across the consumer fleet) + recommended gate CM-REQUEST-HISTORY-DOC-PRESENT (the project-local request-history document exists at its declared path, carries the §11.4.44 revision header + the five mandated per-entry fields + an explicit-timezone stamp on every entry + a reconstruction-boundary scope note; the append helper exists + is parse-clean per §11.4.67; no fabricated field — UNKNOWN used for the unrecoverable) + paired §1.1 meta-test mutation (strip a mandated field from a row / drop the explicit-timezone stamp / fabricate an alias-or-model on an entry that never captured one / delete the document → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.208.

Non-compliance is a release blocker. No escape hatch — no --request-history-optional, --skip-request-log, --no-timezone-stamp, --invent-missing-field, --request-history-in-constitution, --prompt-without-history-entry flag.

§ 2.0 2.6 0.7 1.5 — Zero-loss request/prompt intake: every operator request MUST be mechanically captured, tracked, and processed — never skipped, ignored, avoided, or lost (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-07-15): "We MUST CHECK which requests / prompts we have given havent passed through except the accumulative prompts / requests history document !!! There may be more of it! This

MUST NEVER happen! We MUST make sure that every request / prompt we issues is ALWAYS respected and taken into the account, executed and processed! This is MANDATORY RULE !!! There MUST NOT BE any avoiding, ignoring, skipping or any form of loss for requests / prompts we make! Extend properly constitution Submodule and add mandatory linkings, hooks and extend with dcos chain any part of the System we MUST in order for this to be permanently resolved! Applying this starts immediately! Commit and push all changes to constitution Submodule to all upstreams!"

Forensic FACT (2026-07-15): an entire session's worth of operator prompts (a release-drive "keep going", a BACKGROUND emulator directive, a live-testing model matrix, an alias switch, a multi-track directive) went UN-RECORDED in the §11.4.208 request-history ledger because §11.4.208(D)'s "keep-applying mechanism" (the UserPromptSubmit auto-capture hook) was left as a "tracked follow-up / honestly partial" — so the ledger drifted the moment hand-editing stopped, exactly the §11.4.205 failure class ("a rule documented as mechanically enforced but implemented nowhere is worse than no rule").

Every operator request / prompt MUST be (a) **CAPTURED** — recorded in the §11.4.208 request-history ledger MECHANICALLY at intake, never depending on manual vigilance; (b) **TRACKED** — mapped to a tracked action / workable item (§11.4.202 report → item, §11.4.197 completion) OR an explicit evidence-backed no-op-with-reason; and (c) **PROCESSED** — respected, acted upon, executed, never skipped / ignored / avoided / dropped. Loss of a request at ANY of these three stages is a §11.4 PASS-bluff at the requirements-intake layer of the SAME severity class as §11.4.197 (loss of requirements FORBIDDEN) — §11.4.210 is the INTAKE-side MECHANICAL guarantee that §11.4.197's no-loss holds from the very first moment a request arrives. §11.4.210 STRENGTHENS §11.4.208(D): the auto-capture hook is PROMOTED from "tracked follow-up / honestly partial" to MANDATORY, wired, and verified.

Mandatory protections (ALL hold): (1) **MECHANICAL CAPTURE, not manual vigilance (§11.4.205).** The §11.4.208 ledger's keep-applying mechanism MUST be a WIRED, VERIFIED-RUNTIME-STATE auto-capture hook (a UserPromptSubmit -class hook that appends every operator prompt to the ledger at intake) — NOT a "tracked follow-up", NOT a manual conductor step. "Installed" = a verified runtime state observed in the live hook path (§11.4.205(3)), never a doc claim; the hook MUST be non-fatal (never blocks / fails a prompt) and honest

(UNKNOWN fields never invented, §11.4.6). (2) **DOCS-CHAIN BOUND (§11.4.106)**. The ledger + its four-format exports (§11.4.65/§11.4.73) are bound into the docs-chain so they stay in sync mechanically; a stale ledger export is a §11.4.106 sync violation. (3) **GATED (§11.4.75, enforced-not-advisory)**. A pre-build gate FAILs if the auto-capture hook is not wired in the live hook path (verified-installed per §11.4.205(3), content-matched to its tracked source), OR the ledger is stale relative to the durable prompt record. Paired §1.1 mutation — un-wire the hook / orphan the ledger → gate FAILs. (4) **AUDIT + NO SILENT DROP**. Any operator prompt found un-captured (in .remember/ logs, older CONTINUATION revisions, session transcripts) MUST be back-filled into the ledger, never silently dropped (§11.4.197); every captured request resolves to a tracked action / item (§11.4.202) OR an explicit evidence-backed no-op-with-reason — captured-then-ignored is forbidden. (5) **PROCESS EVERY CAPTURED REQUEST**. The autonomous loop (§11.4.126 / §11.4.87 / §11.4.94 / §11.4.97 / §11.4.103) treats every open captured request as an actionable queue item; the loop's done-condition MUST NOT read satisfied while any captured request is un-processed (mirrors §11.4.147(e)). (6) **HONEST BOUNDARY (§11.4.6)**. Machine-derivable fields (timestamp / track / alias) are captured verbatim; model / effort where undeterminable are UNKNOWN , never invented; a back-filled catch-up row honestly marks its reconstruction boundary (§11.4.208(B)) — never a fabricated time.

Classification: universal (§11.4.17) — the consuming project supplies its ledger path, hook path, and docs-chain context per §11.4.35. Composes §11.4.197 (§11.4.210 is its intake-side mechanical enforcement) / §11.4.202 (report → tracked item) / §11.4.208 (the ledger document — §11.4.210 promotes its (D) hook from partial-follow-up to mandatory) / §11.4.205 (enforced-not-advisory, hook-installed-is-verified-runtime-state) / §11.4.106 (docs-chain) / §11.4.126 / §11.4.87 / §11.4.94 / §11.4.97 / §11.4.103 (the loop processes every captured request) / §11.4.6 / §11.4.75 / §11.4.147(e). Propagation gate CM-COVENANT-114-210-PROPAGATION (literal 11.4.210) + recommended gate CM-REQUEST-CAPTURE-HOOK-WIRED (the auto-capture hook is verified-installed in the live hook path AND content-matches its tracked source per §11.4.205(3); the ledger is fresh; docs-chain-bound) + paired §1.1 meta-test mutation (un-wire the hook → recommended gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.210.

Non-compliance is a release blocker. No escape hatch — no `--skip-request-capture`, `--manual-ledger-OK`, `--hook-optional`, `--drop-unprocessed-request`, `--ledger-may-be-stale`, `--request-may-be-skipped` flag.

§ 11.4.211 — Merge-conflict resolution during main→feature/product/flavor merges MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, — —)

Verbatim operator mandate (2026-07-16): "During the merge of main branch to all feature/product/flavor branches of tracks 2 - 4, if conflicts happen they MUST BE resolved using Fable model with xhigh effort! If Fable is not available, then with Opus with xhigh! Add this MANDATORY RULE into constitution Submodule and its files CONSTITUTION.md, CLAUDE.md, AGENTS.md, QWEN.md, GEMINI.md, and other related relevant files!"

Any merge-CONFLICT resolution performed while integrating the canonical trunk (`main / master`) into a feature / product / flavor branch — the §11.4.188 regular main → feature merge cadence, the §11.4.195 branch-taxonomy integration, the §11.4.41 merge-first pipeline, the §11.4.167(D) trunk-into-stream sync AND the §11.4.167(I) approval-gated back-merge — MUST be carried out on the **Fable model at xhigh effort**. This is the merge-resolution analogue of §11.4.209 (which pins the same model + effort for code-REVIEW): §11.4.211 pins it for the DECISION-MAKING of conflict resolution (which side to keep, how to UNION both sides without loss, marker removal, semantic reconciliation of divergent edits). **(A) Fallback (closed, ordered):** ONLY when Fable is genuinely unavailable (not configured / rate-limited / unreachable, established as FACT per §11.4.6 — never a guessed unavailability) the conflict resolution MUST run on **Opus at xhigh effort**; a lower-effort or a non-Fable/non-Opus model is NOT an acceptable substrate for resolving merge conflicts. `xhigh` effort is REQUIRED in both cases — a careless conflict resolution silently drops a commit or commits a marker (§9 / §11.4.41 step 3), so the highest-effort scrutiny is mandatory. **(B) Scope.** The operator named tracks 2–4 (`feature / product / flavor` branches); the rule GENERALISES per §11.4.17 to EVERY controlled main → (feature|product|flavor) merge on EVERY track that hits a conflict — the substrate requirement does not depend on which track. **(C) Safety preserved.** §11.4.211 sets the MODEL + EFFORT of the conflict-

resolution work; it does NOT weaken and is BOUNDED BY the surrounding merge-safety invariants — §11.4.113 (NEVER force-push), §9 / §9.2 (no commit lost, hardlinked pre-op backup before any large/risky merge), §11.4.41 step 3 (ZERO conflict markers committed, ZERO file silently dropped, union of both sides preserved), §11.4.188(4) anti-bluff (post-merge smoke GREEN + conflict-marker scan empty + no-lost-commit, captured evidence §11.4.5/§11.4.69), §11.4.84 (quiescent-only), §11.4.37/§11.4.71 (fetch-first). **(D) Honest boundary (§11.4.6).** A merge whose conflicts were resolved on a weaker model / lower effort than Fable-xhigh (or Opus-xhigh on genuine Fable-unavailability) is a §11.4 bluff at the merge-resolution-substrate layer — the merge's captured evidence MUST record which model + effort resolved the conflicts (and, on fallback, the FACT of Fable's unavailability). STRICT EXTENSION of §11.4.209 (which pins Fable-xhigh / Opus-xhigh for code-REVIEW) applied to merge-CONFLICT resolution; STRENGTHENS §11.4.188 / §11.4.195 / §11.4.41 by pinning the substrate of their conflict-resolution step — the two anchors are the review-side (§11.4.209) and merge-side (§11.4.211) halves of the same Fable-xhigh substrate discipline.

Classification: universal (§11.4.17). Composes §11.4.209 / §11.4.188 / §11.4.195 / §11.4.41 / §11.4.113 / §9 / §9.2 / §11.4.6 / §11.4.167(D)(I) / §11.4.84 / §11.4.37 / §11.4.71 / §11.4.196 (model/alias selection). Propagation gate `CM-COVENANT-114-211-PROPAGATION` (literal `11.4.211` across the consumer fleet) + recommended gate `CM-MERGE-CONFLICT-FABLE-XHIGH` (a main → feature/product/flavor merge that resolved conflicts records model=Fable + effort=xhigh, OR model=Opus + effort=xhigh with a recorded Fable-unavailability fact) + paired §1.1 meta-test mutation (record a conflict-resolving merge on a lower-effort or non-Fable/non-Opus model without a Fable-unavailability fact → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.211.

Non-compliance is a release blocker. No escape hatch — no `--merge-conflict-any-model`, `--skip-fable-merge-resolution`, `--non-xhigh-merge-conflict`, `--resolve-conflict-lower-effort-OK`, `--fallback-without-fable-unavailability-proof` flag.

§ . . — **FEATURE** research-scheduling directive: recognized via all § . . forms, **SCHEDULES** (never synchronously executes) a deep, enterprise-grade research + implementation-planning effort as a tracked workable item (operator-directed mandate, — —)

Mandate (2026-07-16): extend the §11.4.202 reporting-directive family with a fourth registered action, **FEATURE**, that turns a plain-language feature description into a **SCHEDULED**, deep, enterprise-grade research-and-planning effort — mirroring EXACTLY how **ISSUE** / **BUG** / **TASK** turn a report into a tracked workable item, but for **RESEARCH SCHEDULING** rather than defect/task/feature reporting: the directive itself creates the tracked item and enqueues it durably; the multi-day research/planning/documentation work it describes is **EXECUTED LATER** by the project's standing autonomous loop when it claims that item, never synchronously by the directive engine.

FEATURE is a research-**SCHEDULING** directive, structurally parallel to the §11.4.202 report-creation directives: recognized via ALL §11.4.140 forms —

FEATURE :: <desc> (bare ::), **DEFAULT::FEATURE** :: <desc> (namespaced ::), /**FEATURE** <desc> (bare slash), /**DEFAULT::FEATURE** <desc> (namespaced slash), **FEATURE** ---> <desc> (arrow) — PLUS the §11.4.202 single-colon form **FEATURE**: <desc>, **REGISTERED-ACTION-ONLY** by construction exactly like **BUG**: / **TASK**: / **ISSUE**: : because the action token grammar is UPPERCASE-only ([A-Z][A-Z0-9_]*), ordinary lowercase prose such as "feature: add dark mode" never matches ANY of the six forms and is never expanded and never asked about (§11.4.6) — only the literal uppercase **FEATURE** token triggers the directive. **\FEATURE** (leading backslash) escapes to a literal per the existing §11.4.140 escape mechanism, unchanged.

(1) Scheduling, not synchronous execution — the load-bearing distinction from every other §11.4.140 action. **BACKGROUND** / **CRITICAL** / **IMPORTANT** / **NOTE** tag HOW the remainder of the SAME turn is handled; **BUG** / **TASK** / **ISSUE** immediately create + fully sync a workable item describing a report and the reported matter is then handled inline. **FEATURE** differs: its OWN execution is fast and bounded — create the tracked item (Type=Task, Status=Queued, §11.4.54 stable id), append it to a durable, never-dropped feature-research queue, sync the

DB-derived documents, and push to configured trackers — exactly as §11.4.202 does for TASK , and via the SAME engine, never a duplicate implementation. The deep multi-day research / systematic-debug / documentation / planning work DESCRIBED INSIDE the item is NOT performed by the directive at dispatch time; it is performed LATER by the project's standing autonomous loop (§11.4.87 endless-loop / §11.4.94 zero-idle / §11.4.97 maximum-idle-use / §11.4.103 continuous-parallel-stream / §11.4.126 default-autonomous-loop) the moment it claims the item, and is driven to a genuinely COMPLETED (implemented AND wired AND verified) or explicitly evidence-backed CLOSED terminal state under the §11.4.197 research/kicked-off-work completion mandate — a FEATURE item is NEVER permitted to sit un-wired in the backlog; "scheduled" is not "done" and is never reported as such (§11.4.6).

(2) The scheduled research-and-planning mandate — embedded verbatim into the created item's comprehensive description (§11.4.148 D2 / §11.4.171) as its acceptance-defining WORK PROGRAM: (a) a deep web-research series — articles, guides, scientific papers, open-source projects/codebases, worked examples — on the best way to incorporate the described feature into the project (§11.4.8 deep-research-before-implementation / §11.4.99 latest-source cross-reference / §11.4.150 deep multi-angle research per change); (b) systematic-debug (§11.4.102) of ALL data obtained to enumerate EVERY weak spot / gap / danger-zone / inconsistency / imperfection the research surfaces, and design a risk-free, rock-solid solution for EACH one found — no gap left unaddressed; (c) the design MUST ALWAYS be enterprise-grade, bleeding-edge, and innovative — never less; (d) MANDATORY exhaustive documentation — technical documents, an implementation plan down to lines-of-code + micro-proof-of-concepts, diagrams / schemes / graphs, illustrations, SQL definitions, templates, and every other relevant material, four-format-exported and revision-headed per §11.4.65 / §11.4.73; (e) investigate + plan REUSE of existing vasic-digital + HelixDevelopment submodules/components BEFORE proposing a rewrite, extending them freely where genuinely needed while keeping them fully decoupled / project-not-aware / reusable per §11.4.28 / §11.4.74 / §11.4.177; (f) plan from the TESTING point of view from day one — every supported test type, the Challenges submodule, and full HelixQA bank coverage per §11.4.27 / §11.4.169; (g) divide the work into phases / tasks / subtasks, fine-grained and nano-detailed enough to drive integration / implementation / wiring / testing / scaling directly, no vague placeholders; (h) plan explicitly for enterprise scalability and maximal performance;

(i) when the feature has a UI/UX surface, produce full wireframes / diagrams / design files (Figma / PSD / PDF, or whatever formats the project mandates) authored via OpenDesign per §11.4.162 / §11.4.190; (j) plan full CodeGraph integration with regular / real-time index synchronisation per §11.4.78 / §11.4.79 / §11.4.80; (k) create fully-detailed workable items (every detail + reference + attachment) synced in REAL TIME to the SQLite single-source-of-truth (§11.4.93 / §11.4.95) + every derived workable-items document + every related project doc/ component + every connected external work-tracking system (ClickUp, HelixTrack, JIRA, or any other configured tracker per §11.4.148 D5) — honestly SKIPPING any absent tracker with a machine-readable reason (`credentials_absent` / `tracker_client_absent` , §11.4.10 / §11.4.6) rather than EVER faking a push, exactly as §11.4.202 mandates.

(3) The item. The created item carries Type=Task (§11.4.16 — a scheduled research/planning effort is itself an internal workstream, not a product defect nor yet a shipped end-user capability; the feature it researches may spawn its OWN Feature-typed implementation item(s) once planned) + Status=Queued (§11.4.15) + a stable auto-incremented id (§11.4.54) + the §11.4.148 / §11.4.171 comprehensive structured description embedding (i) the verbatim feature description, (ii) the full research/planning/testing/scalability/UI/CodeGraph/tracker mandate of clause (2) verbatim as the item's WORK PROGRAM, and (iii) acceptance criteria = the full research-doc tree is produced under the item's recorded destination path, every enumerated gap carries a designed solution, an implementation plan down to lines-of-code exists, the planned follow-on workable items are created, and the whole effort is validated per §11.4.197 (never left un-wired). An undetermined section is recorded as an explicit `UNKNOWN: gap, never invented` (§11.4.6), exactly as §11.4.202's engine already does for BUG/TASK/ISSUE.

(4) Decoupled engine, reused machinery, consumer-owned data. The item-creation + full-sync half of the FEATURE mandate is NOT reimplemented: the engine (`constitution/scripts/feature/schedule_feature_research.sh`) invokes the SAME §11.4.202 `report_item.sh` engine to create the item, sync the derived documents, and push to trackers — so the already-audited no-faked-push / decoupling / anti-bluff invariants of §11.4.202 are inherited rather than duplicated (a second, divergent implementation of tracker-push logic would itself be a bluff-surface regression). The FEATURE-specific additions — the research-doc-tree destination path and the durable `docs/requests/feature_queue.md` queue entry (mirroring the §11.4.140 `BACKGROUND` action's `docs/requests/`

`background_queue.md` durability pattern, so a scheduled request is NEVER silently dropped or forgotten) — are the engine's OWN, narrow responsibility. Both engines are inherited BY REFERENCE, carry ZERO project literals, and fail closed with an actionable message when their consumer config is absent (§11.4.6 / §11.4.28 / §11.4.177); every project-specific value (DB path, id prefix, research-doc root, sync command, tracker bindings) is consumer-owned DATA in `feature.yaml`, never an edit to the engine.

(5) Anti-bluff (§11.4 / §11.4.107(10)). Four-layer coverage mirroring §11.4.202: a pre-build gate (`CM-FEATURE-DIRECTIVE`) whose FUNCTIONAL invariant SOURCES and RUNS the §11.4.140 parser (never a grep-only assertion, §11.4.108) to prove the `FEATURE` token genuinely expands via the single-colon, `::`, arrow, and slash forms while lowercase prose (`feature: ...`) does NOT expand and is NEVER asked about; a paired §1.1 mutation test proving every invariant (registry presence, engine parseability, decoupling, the durable-queue append, delegation-not-reimplementation of the tracker machinery, and the scheduling-not-synchronous framing) is genuinely load-bearing; and captured evidence of a real feature-research item landing in a temp SQLite DB, its derived documents regenerated FROM the DB, and its trackers SKIPping honestly.

Classification: universal (§11.4.17) — the consuming project supplies its DB path, id prefix, research-doc root, sync command, and tracker bindings per §11.4.35.

Composes §11.4.140 (grammar) / §11.4.202 (the sibling reporting-directive family + its reused engine) / §11.4.197 (a started research effort never evaporates) / §11.4.8 / §11.4.99 / §11.4.150 (deep research) / §11.4.102 (systematic-debug of every gap) / §11.4.65 / §11.4.73 (exhaustive documentation) / §11.4.28 / §11.4.74 / §11.4.177 (reuse-first, decoupled) / §11.4.27 / §11.4.169 (test-type + Challenges + HelixQA planning) / §11.4.162 / §11.4.190 (OpenDesign UI planning) / §11.4.78 / §11.4.79 / §11.4.80 (CodeGraph integration) / §11.4.93 / §11.4.95 / §11.4.148 (workable-items + tracker sync) / §11.4.54 / §11.4.15 / §11.4.16 (id/status/type) / §11.4.10 (credentials) / §11.4.6 (no-guessing) / §11.4.87 / §11.4.94 / §11.4.97 / §11.4.103 / §11.4.126 (the autonomous loop that actually performs the scheduled work). Propagation gate `CM-COVENANT-114-213-PROPAGATION` (literal `11.4.213` across the consumer fleet) + recommended gate `CM-FEATURE-DIRECTIVE` (registry declares `FEATURE` + its single-colon form is registered-only; the engine exists, is parse-clean, and is project-literal-free; the `FEATURE` directive really expands at runtime via all six forms while lowercase prose does NOT; the engine SCHEDULES via the durable feature-research queue AND delegates item-creation/

sync/tracker-push to the §11.4.202 `report_item.sh` engine rather than reimplementing it) + paired §1.1 meta-test mutation (strip the `FEATURE` registry row, OR flip its single-colon form to non-registered, OR strip the parser branch, OR remove the durable-queue append, OR couple the engine to one project, OR make the engine reimplement tracker logic instead of delegating, OR strip the scheduling-not-synchronous framing from the expansion text → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.213.

Non-compliance is a release blocker. No escape hatch — no `--feature-runs-synchronously`, `--skip-feature-queue`, `--feature-without-research-mandate`, `--reimplement-tracker-push`, `--hardcode-project-in-engine`, `--skip-item-creation` flag.

§ . . — Live in-session task/todo tracker MUST always be up to date and fully in sync with the real work state — never stale, never showing completed work as pending nor pending as done (User mandate, — —)

Verbatim operator mandate (2026-07-25): "make sure it is ALWAYS up to date fully in sync! CRITICAL: This MUST BE one of mandatory rules / constraints! Add it with all relevant details into the constitution Submodule and relevant related files."

The agent's LIVE IN-SESSION TASK/TODO TRACKER — the harness-provided task list / `TodoWrite` -class todo tool / any equivalent moment-to-moment execution-state surface the runtime exposes DURING a session — MUST ALWAYS be kept up to date and FULLY IN SYNC with the real live work state: never stale, never showing completed work as pending, never showing pending / blocked / abandoned work as done or in-progress. It MUST be updated the MOMENT any tracked task's state changes — the instant a task is started (→ in-progress), completed (→ done), blocked, abandoned, split, re-scoped, or newly discovered — not batched to a later checkpoint, not reconstructed from memory at the end of the work, not left reflecting a superseded plan. A live task tracker that lags the real work is itself a §11.4 PASS-bluff at the execution-transparency layer: it reports a work

state the agent is not actually in, and an operator (or a resuming agent / conductor reading it per §11.4.116) then makes decisions on a false picture of what is done, what is in flight, and what is owed.

(A) EVERY STATE-CHANGE UPDATES IT, IMMEDIATELY. The moment a task starts, completes, changes status, or is discovered, the tracker is updated to match — exactly one task marked in-progress at a time where the surface models that single-active-task convention, no completed task left marked pending, no pending task left marked done, no ghost task lingering for work that was dropped (dropped work is closed with a reason, never silently left "in-progress"). Deferring the update ("I'll mark it all done at the end") is precisely the staleness this anchor forbids.

(B) NEVER STALE, NEVER INVERTED (§11.4.6). The tracker's content is CAPTURED FACT about the real work state — never a guess, never aspirational, never a plan the actual work has since diverged from. Showing completed work as still-pending, or showing pending / blocked / abandoned work as done or in-progress, is a §11.4 / §11.4.1 bluff of PASS-bluff severity (falsely-done) or FAIL-bluff severity (falsely-pending); a tracker that has drifted from reality is corrected the MOMENT the drift is observed, never carried forward.

(C) DISTINCT ARTIFACT — COMPLEMENTS, NEVER REPLACES. The live in-session task/todo tracker is the EPHEMERAL, moment-to-moment execution-state surface WITHIN a single session — a DIFFERENT artifact from each of: §12.10's version-controlled CONTINUATION document, §11.4.131's standing committed session-resumption FILE, §11.4.97's operator-facing milestone progress-update cadence, §11.4.106(F)'s doc/DB sync enforced at the commit / build / constitution-pull WRITE-SEAMS, and §11.4.202 / §11.4.208's workable-item DB and operator-request-history ledgers. Those anchors bind COMMITTED / CROSS-SESSION / OPERATOR-REPORT / TRACKED-WORK artefacts; NONE of them binds the live in-session task/todo tracker itself — §11.4.229 closes exactly that gap. It COMPLEMENTS them (a correct live tracker is the truth the §12.10 / §11.4.131 / §11.4.202 / §11.4.208 artefacts are materialised FROM), never substitutes for any of them, and its always-in-sync discipline is the in-session analogue of §12.10's always-current mandate applied to the harness task list.

(D) HONEST BOUNDARY (§11.4.6). Where the runtime exposes NO such live task-tracking surface, this anchor is honestly N/A for that runtime (§11.4.3 SKIP-with-reason — the discipline BINDS the moment such a surface exists, per the §11.4.96 latent-binding pattern); it mandates that WHATEVER live tracker the

runtime does expose is kept in-sync, not that a particular tool must exist. "In sync" means the tracker reflects the real work state as the agent knows it at that moment — the captured fact of what has started / completed / changed — not omniscience about work the agent has not yet observed.

Classification: universal (§11.4.17) — a platform-neutral execution-transparency discipline reusable by ANY agent / runtime that exposes a live in-session task/todo tracker; the consuming runtime supplies its concrete tracker surface (harness task list / `TodoWrite` / equivalent) per §11.4.35. Composes §12.10 (always-current CONTINUATION — §11.4.229 is its live-in-session-tracker analogue) / §11.4.131 (standing session-resumption file) / §11.4.97 (progress-update cadence — the operator-report sibling) / §11.4.106(F) (write-seam doc/DB sync) / §11.4.126 (default autonomous-loop — the loop's live task state IS the tracker) / §11.4.116 (real-time conductor ↔ framework sync — a resuming conductor reads the live tracker) / §11.4.202 / §11.4.208 (workable-item / request-history materialisation) / §11.4.6 (no-guessing — the tracker is captured fact, never aspirational) / §11.4.1 (falsely-pending is a FAIL-bluff). Propagation gate `CM-COVENANT-114-229-PROPAGATION` (literal `11.4.229` across the consumer fleet) + recommended gate `CM-LIVE-TASK-TRACKER-IN-SYNC` (where the runtime exposes a live in-session task/todo tracker, its state matches the real work state at every state-change — no completed-task-marked-pending, no pending-task-marked-done, no ghost/stale entry, no superseded-plan drift; honest §11.4.3 SKIP-with-reason where the surface is genuinely absent — the §11.4.201(1) false-positive guard) + paired §1.1 meta-test mutation (mark a completed task as still-pending, OR a pending/abandoned task as done, OR leave the tracker reflecting a superseded plan → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6).

Canonical authority: constitution submodule `Constitution.md` §11.4.229.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--tracker-may-lag` , `--batch-todo-updates` , `--stale-task-list-OK` , `--reconstruct-todos-at-end` , `--skip-task-tracker-sync` , `--pending-may-be-done` , `--completed-may-stay-pending` flag exists.

§ . 2026 07 27 — Build-and-deploy the moment source fixes are proven-correct (test-side hardening is a parallel stage, never a build gate) + post-deploy version increment (operator mandate, — —)

(A) BUILD-AND-DEPLOY THE MOMENT THE SOURCE FIXES ARE PROVEN-CORRECT — TEST-SIDE WORK IS A PARALLEL STAGE, NEVER A BUILD GATE.

The moment the SOURCE fixes of a batch are done AND proven-correct — proven-correct = the mandatory independent SOURCE-correctness review has returned a clean GO (§11.4.125/§11.4.142/§11.4.194 on the source diff, on the §11.4.209 Fable- xhigh substrate, iterated to zero-finding/zero-warning GO per §11.4.134) — the build MUST be triggered, pre-build + post-build tests EXECUTED (§11.4.108 layer-1 before, layer-2 after), and the artifact deployed. IN PARALLEL WITH the build AND the manual-QA session (§11.4.185), NEVER gating the build: (1) gate-instrumentation hardening + new-gate authoring (§11.4.110); (2) paired §1.1 mutations + four-layer §11.4.4(b) coverage + HelixQA Challenges; (3) "gates and checks of all kinds" (§11.4.96); (4) the re-review of THAT test-instrumentation work. The build's GO-SIGNAL is source-correctness alone; test-guard hardening (future-regression coverage) + its re-review are a deploy-independent parallel stage per §11.4.230(A). Holding the build sequentially — waiting for the test-instrumentation buildout + its re-review — when the source is already proven-correct is a §11.4.235(A) violation.

(B) POST-DEPLOY VERSION INCREMENT + MANUAL-QA-DEPLOY CYCLE BOUNDARY (extended 2026-07-28 — operator mandate, verbatim: "Increase upcoming release to 0.1.5 since new cycle starts every time we deploy release for manual QA testing! This is how we will be working this ALWAYS! Extend the constitution!").

The moment image(s)/artifact(s) deploy, the version code/version is INCREMENTED (next build targets the next version; every deployed artifact carries a distinct, monotonic, greppable id — §11.4.151/§11.4.108/§11.4.200). Two deployments sharing one version id defeats per-deploy artifact identity → §11.4.235(B) violation. EVERY deployment of a release FOR MANUAL QA TESTING (§11.4.185) is a DEVELOPMENT-CYCLE BOUNDARY: the manual-QA deploy CLOSES the outgoing cycle and STARTS a NEW development cycle, and the version code/version is incremented AT THAT POINT — the new cycle's work targets the NEXT monotonic version (§11.4.151 project-prefixed naming preserved) — ALWAYS, the STANDING working model per clause (C), never a per-

request opt-in. The manual-QA deploy is the canonical clause-(B) increment trigger (a deploy-for-manual-QA IS a deploy), and it SHAPES the §11.4.126 release-cycle framing: fixes for QA findings against the deployed artifact land in the NEW cycle on the NEW version id, never patched onto the already-QA-deployed version id. Starting a new cycle's work still targeting the QA-deployed version id is a §11.4.235(B) violation.

(C) STANDING DEFAULT (§11.4.126/§11.4.230/§11.4.103), never a per-request opt-in.

Honest boundary (§11.4.6): relocates the TEST-INSTRUMENTATION re-review + gate/check hardening OFF the build critical path; does NOT weaken any substantive gate. The SOURCE-correctness review (§11.4.125/§11.4.142/ §11.4.194 Fable- xhigh §11.4.209 → GO §11.4.134) STILL PRECEDES + gates the build (it IS "source proven-correct"); pre/post-build tests (§11.4.108 L1-2) run at their seams; §11.4.185 manual-QA-final un-weakened; §11.4.129/§11.4.40 un-weakened. "Test-side work is parallel, never a build gate" is bounded to future-regression instrumentation — a check that proves the SHIPPED artifact wrong is a source/artifact-correctness gate (§11.4.108 L1-2), stays on the critical path.

Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-235-PROPAGATION (literal 11.4.235) + recommended CM-BUILD-ON-SOURCE-PROVEN-NOT-TEST-SIDE (FAIL: build held past source-GO w/ unfinished test-instrumentation; FAIL: build started with NO source-review GO — the §11.4.201(1) dual-assertion guard) + CM-VERSION-INCREMENT-ON-DEPLOY (FAIL: new-artifact deploy shares a version id; FAIL: work of the cycle opened by a manual-QA deploy still targets the QA-deployed version id — the clause-(B) cycle-boundary increment skipped) + paired §1.1 mutations; gate-code = separate work item (§11.4.6/ §11.4.227).

Non-compliance is a release blocker. No escape hatch — no --build-behind-test-hardening , --test-side-gates-the-build , --skip-source-correctness-review , --deploy-without-version-increment , --not-default-speed-mode , --serialise-build-and-validate .

§ . . — QA-deploy-readiness gate: no manual-QA hand-off until the mandated validation produced a candidate-fingerprinted PASS verdict; a blocker MUST bind to a seam, never prose (research-derived, — —)

Forensic anchor (genericised, 2026-07-31). An operator declared a specific validation a **release-blocking requirement** for every QA-deploy ("MANDATORY ... QA MANUAL TESTING READY BLOCKER ... HAS TO BE FULLFILLED COMPLETELY FOR EACH RELEASE"), and it was still not enforced at the next QA-deploy: the validation harness existed and ran, hit an instrument bug (a false-FAIL on later legs), and the blocker was **downgraded to a non-blocking follow-up** while the deploy proceeded on a weaker side-check. Nothing STOPPED the deploy because the requirement lived as **prose** (an agent reads it) + a source-gate that asserted the harness *exists* — NOT a seam that refuses the deploy in the absence of a clean recorded verdict; the deploy tool had **zero** references to the validation (measured). This is the §11.4.205 "documented-as-enforced but implemented nowhere is worse than no rule" + §11.4.226 "prose doesn't bind, seams do" failure, demonstrated against the operator's own explicit blocker, plus a §11.4.120 judgment error (a bug in a blocker's instrument is not licence to skip the blocker — fix the instrument).

The mandate (ALL hold). (1) **DEPLOY-SEAM VERDICT-COVERAGE (§11.4.135 applied to the QA-deploy seam).** A release artifact MUST NOT be deployed, published, or handed off for manual QA (§11.4.185) until the mandated autonomous validation(s) for that artifact produced a **machine-written PASS verdict** whose **artifact fingerprint == the candidate** (§11.4.115(F) — the fingerprint READ FROM THE TARGET at run time). **ABSENCE of the verdict BLOCKS the deploy exactly as a FAIL does** — a not-yet-run validation and a failed one are both refusals, never a silent pass. (2) **MECHANICAL, NOT PROSE (§11.4.205/§11.4.226).** The refusal lives in the deploy tool itself (the seam), never in prose an agent is trusted to honour; a source-gate that asserts the validation harness *exists* is necessary-not-sufficient — it does not prove a clean verdict for the candidate. (3) **BLOCKER-BINDING RULE.** Any requirement an operator declares a blocker (a §11.4.140 CRITICAL, or an explicit "release blocker / QA-ready blocker") MUST be bound to a mechanical seam at the point it gates, never satisfied by prose and never DOWNGRADED to a follow-up — and when the blocker's own INSTRUMENT bugs, the response is §11.4.4/§11.4.120 (STOP, fix the instrument, complete the

blocker), never route around it. (4) **EXPLICIT, RECORDED OVERRIDE ONLY.** The seam has no silent bypass; an override is an explicit, logged, operator-set flag (the §11.4.234/§11.4.135 recorded-deferral pattern — never a follow-up-downgrade, never an opaque skip). (5) **THE GATE IS SELF-VALIDATED (§11.4.107(10)/§11.4.201).** The deploy seam ships golden-good (verdict present + PASS + fingerprint-match → allow) + golden-bad (absent → refuse / not-pass → refuse / fingerprint-mismatch → refuse), so it cannot itself false-allow; and the verdict-emitter reads the fingerprint from the target at run time, never a hardcoded per-release value (§11.4.111 — a version-pinned validator is the brittleness that lets a stale verdict pass). **Honest boundary (§11.4.6):** the seam proves a candidate-matching PASS verdict EXISTS for the mandated autonomous validation — it does NOT replace the §11.4.185 manual-QA sufficiency gate (the human check the deploy hands off TO), and it is only as honest as the validation whose verdict it consumes (that validation's own anti-bluff stays §11.4.5/§11.4.69/§11.4.107). Classification: universal (§11.4.17) — consumers supply the mandated-validation set, the verdict path/schema, and the deploy tool per §11.4.35. Composes §11.4.4/.107(10)/.108/.111/.115(F)/.120/.126/.135/.146(D3)/.185/.201/.205/.226/.234. Propagation gate `CM-COVENANT-114-236-PROPAGATION` (literal `11.4.236`) + recommended gate `CM-QA-DEPLOY-READINESS-GATE` (the deploy tool refuses without a candidate-fingerprinted PASS verdict; the gate self-validates golden-good/bad; the verdict-emitter reads the fingerprint at run time) + paired §1.1 mutation (make the deploy tool proceed on an absent/mismatched verdict → the gate FAILs; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227). Non-compliance is a release blocker. No escape hatch — no `--deploy-without-verdict`, `--prose-blocker-suffices`, `--downgrade-blocker-to-followup`, `--skip-readiness-gate`, `--verdict-fingerprint-optional`, `--silent-override` flag.

§ . . – Cutting-edge enterprise quality + production-deployment readiness: every work product is built for production deployment, maximal stability, zero nasty surprises, and enterprise-grade robustness – never for "good enough" or "will do for now" (BACKGROUND :: REMINDER :: IMPORTANT operator mandate, – –)

Verbatim operator mandate (2026-08-15, Point 6): *"All work we do MUST BE done for the cutting edge enterprise quality! We are doing everything for the deployment to production and maximal stability with no nasty surprises or issues of any kind!"*

Compact summary: every work product a consuming project ships MUST be built for CUTTING-EDGE ENTERPRISE QUALITY targeted at PRODUCTION DEPLOYMENT + MAXIMAL STABILITY + ZERO NASTY SURPRISES — the standing default posture, engaged automatically for every batch, subagent, feature, fix, and change from the first prompt of the session onward (§11.4.126 default-autonomous-loop composed with §11.4.198 always-on defaults), NEVER a per-request opt-in and NEVER downgraded to "good-enough" or "will-do-for-now" or "MVP-first"; every change MUST satisfy the closed PRODUCTION-READINESS invariant set (clause B) before it is considered done, and a project maintains a docs/PRODUCTION_READINESS.md document (§11.4.153 sibling) that tracks readiness per component × invariant × evidence-source. §11.4.260 STRENGTHENS §11.4.190 (bleeding-edge enterprise visual quality — WEB surface only) into a UNIVERSAL production-readiness discipline across every surface (web, service, CLI, container, library, docs, orchestration, tests, gates) and BINDS §11.4.5/§11.4.69/§11.4.107 (captured evidence at every claim) + §11.4.108 (four-layer verification) + §11.4.169 (comprehensive test-type coverage) + §11.4.185 (manual-QA sufficiency) + §11.4.226 (evidence-class-at-closure) + §11.4.235 (build-when-source-proven-correct) + §11.4.236 (QA-deploy-readiness) into one always-on production-readiness posture.

(A) CUTTING-EDGE ≠ NOVELTY-FOR-ITS-OWN-SAKE. "Cutting-edge" means the current best-known engineering approach for the problem class, sourced via §11.4.99 latest-source verification (never memory) + §11.4.150 deep multi-angle research + §11.4.8 external precedent, applied fully — NEVER a shortcut, a stub, a placeholder-with-a-TODO, a for now (§11.4.124), a duct-tape workaround, or a

"will fix later" (§11.4.197 loss-of-requirements); if the current best-known approach is genuinely unavailable in the project's constraints, that gap is HONESTLY tracked (§11.4.197 item + [OPEN] marker per §11.4.223), NEVER silently absorbed as "we did our best."

(B) PRODUCTION-DEPLOYMENT-READINESS INVARIANT SET. Every shipped work product satisfies ALL of: (1) FUNCTIONAL COMPLETENESS — every stated capability wired + verified on a clean target per §11.4.108, no unwired feature, no dead-code path (§11.4.124); (2) OBSERVABILITY — health endpoint, structured logs, per-operation metrics, distributed-trace-ready surface, evidence-capture at every gate (§11.4.85 stress+chaos telemetry, §11.4.128 always-on recording); (3) RESILIENCE — real failure-injection tested per §11.4.85 (chaos: process death, network fault, resource exhaustion, state corruption); recovery is verified never assumed; (4) SECURITY — SonarQube CLI installed + scanned (§11.4.184), no leaked credentials (§11.4.10 pre-store audit), supply-chain integrity at the fleet minimum level (SLSA Build Level 2 per §11.4.246); (5) PERFORMANCE — measured under sustained load per §11.4.85, latency + throughput + resource footprint within declared SLOs; (6) MAINTAINABILITY — code review passed at Fable-xhigh (§11.4.209) with zero findings + zero warnings iterated to GO per §11.4.134; (7) DEPLOYABILITY — rootless container per §11.4.161, deterministic build + reproducible artifact per §11.4.108 layer 2, deployment verified via §11.4.200 target-isolation + §11.4.108 layer 3 runtime-signature; (8) DOCUMENTABILITY — §11.4.257 (manuals/guides/FAQ) + §11.4.258 (diagrams) + §11.4.259 (badges) all GREEN; (9) TESTABILITY — §11.4.27/§11.4.169 seven-type breadth + §11.4.224 coverage floor, every gate paired §1.1 mutation, every regression permanently guarded per §11.4.135; (10) OPERATIONAL RUNBOOK — for every non-trivial failure mode a documented remediation path (§11.4.257 guide class).

(C) ZERO NASTY SURPRISES. No known-broken behavior ships (§11.4.238 QA discovery channel closed), no failing gate is silenced (§11.4.120 seam-placement + §11.4.201 real-condition), no test disabled without a tracked follow-up (§11.4.197), no closed-with-workaround defect masquerading as fixed (§11.4.34 reason-source), no "will fail sometimes but rarely" tolerance (§11.4.50 deterministic consistency + §11.4.85 stress/chaos), no try/catch/swallow silent-failure paths (a caught error either has a captured recovery path or re-raises to a bounded operator decision per §11.4.101/§11.4.66).

(D) STANDING DEFAULT. Composes §11.4.126 (default autonomous-loop mode) + §11.4.198 (always-on default mechanisms) + §11.4.183 (maximum-useful multi-agent + full-constitution-application per track) + §11.4.231 (nano-precision model + effort tier selection); the production-readiness discipline is engaged from the FIRST prompt of the session, never negotiated, never opt-in.

(E) docs/PRODUCTION_READINESS.md — the tracker. Per-component × invariant × evidence-source table (§11.4.153 sibling); revision header per §11.4.44; auto-synced via §11.4.106; reachable from README per §11.4.212; the source of truth for the §11.4.259 production-readiness gauge.

Honest boundary (§11.4.6). §11.4.260 mandates the POSTURE + the INVARIANT SET + the TRACKER; it does NOT claim any project is provably-bug-free (§11.4.118 discovery-pressure boundary remains, §11.4.238 escape triggers a coverage audit not a proof); it does NOT license shipping with unmet invariants — an unmet invariant is HONESTLY red-badged (§11.4.259) and blocks release. "Cutting-edge" is bounded to the project's problem class + published best-known approaches — a genuine research frontier where no best-known exists is HONESTLY marked as such and the project's own approach is documented as an original work contribution (§11.4.8 "NO external solution found — original work" citation).

Classification: universal (§11.4.17). Composes §11.4.5 / §11.4.6 / §11.4.8 / §11.4.10 / §11.4.27 / §11.4.34 / §11.4.44 / §11.4.50 / §11.4.66 / §11.4.69 / §11.4.85 / §11.4.99 / §11.4.101 / §11.4.106 / §11.4.107 / §11.4.108 / §11.4.118 / §11.4.120 / §11.4.124 / §11.4.126 / §11.4.128 / §11.4.134 / §11.4.135 / §11.4.150 / §11.4.153 / §11.4.161 / §11.4.169 / §11.4.183 / §11.4.184 / §11.4.185 / §11.4.190 / §11.4.197 / §11.4.198 / §11.4.200 / §11.4.201 / §11.4.209 / §11.4.212 / §11.4.223 / §11.4.224 / §11.4.226 / §11.4.231 / §11.4.235 / §11.4.236 / §11.4.238 / §11.4.257 / §11.4.258 / §11.4.259 / §11.4.261 / §11.4.262. Propagation gate CM-COVENANT-114-260-PROPAGATION (literal 11.4.260) + recommended gates CM-PRODUCTION-READINESS-TRACKER-PRESENT + CM-CUTTING-EDGE-POSTURE-ALWAYS-ON + CM-ZERO-NASTY-SURPRISES-AUDIT + paired §1.1 mutations.

EXTENSION to (B)(2) OBSERVABILITY — the mechanism, not only the property (2026-08-20). Clause (B)(2) NAMES observability as a production-readiness invariant (*"health endpoint, structured logs, per-operation metrics, distributed-trace-ready surface, evidence-capture at every gate"*). Naming a property is not a mechanism, and three mechanical requirements were missing.

The AI-curriculum corpus (module 33) supplies them. **(a) THE THREE SIGNALS ARE DISTINCT AND EACH ANSWERS A DIFFERENT QUESTION.** The corpus's own naming is *"the three complementary signals of observability"* — **structured logs** (*"What exactly happened in this request?"* — *"queryable key–value events (JSON), not prose strings"*; the corpus is explicit that structure is load-bearing: *" log.info("charge_failed", {donation_id, gateway, idempotency_key, err_code}) is queryable; log.info("charge failed for " + id) is grep-and-pray"*), **metrics** (*"How often, how bad, how loaded — across all requests?"* — numeric time series), and **traces** (*"Where in the request's path across services did time go or the error occur?"* — a tree of timed spans). The corpus adds a fourth signal class alongside them: **crash/ANR analytics** (*"symbolicated stacks with breadcrumbs, grouped by signature"*). Note for citation discipline (§11.4.6): the corpus does NOT use the terms "observability triad" or "three pillars"; consuming projects MUST NOT attribute those labels to it. **(b) A CORRELATION ID THREADS ALL SIGNALS, AND MUST SURVIVE ASYNC BOUNDARIES.** Corpus: *"A single request_id / trace_id threaded through logs, metrics exemplars, and spans lets me pivot from a metric spike → the exemplar trace → the structured logs of that exact request. That pivot is what turns 'error rate is up' into 'here is the one request and the line that failed.'"* The load-bearing edge case is the async one: *"For an async money flow, I make sure the correlation ID survives the queue — the webhook, the worker, and the reconciliation sweep all log the same donation/intent ID — or I lose the thread exactly where the hard bugs live."* The prescribed pivot order is **metric** → **trace** → **log**, under the discipline *"narrow before you theorize"*. **(c) CRASH-FREE USERS / SESSIONS IS A RELEASE GATE, AND A NEW SIGNATURE AT A RELEASE BOUNDARY IS A BLOCKER.** Corpus: *"the metric that matters is crash-free users / sessions, not raw crash count — it normalizes for traffic and is the honest health signal I gate releases on. And I treat a new signature after a deploy as a release-blocker even at low volume, because a signature that is rare today can be a top crasher once rollout completes. Volume tells me priority; novelty at a release boundary tells me causation."* This binds directly to §11.4.152 (Crashlytics-recorded-data continuous monitoring — §11.4.260(B)(2) now states WHEN that data blocks) and to §11.4.265 (the crash-free delta across a progressive rollout is one of that anchor's required non-infrastructure signals). **THRESHOLDS ARE CONSUMER DATA (§11.4.6 / §11.4.35).** The corpus states NO numeric crash-free threshold and explicitly cautions that *"crash-free thresholds"* are among the figures that *"change over time and by version — treat cited figures as illustrative and verify against current official documentation"*. This extension

mandates that the gate EXISTS and that its threshold is declared and calibrated locally; it mandates no value. **Honest boundary (§11.4.6), stated to prevent an over-claim this extension will otherwise attract:** the corpus does NOT contain, and this clause does NOT assert, any proposition of the form "a system you cannot observe in production cannot be verified in production." What the corpus licenses is narrower and is the whole warrant here: *"Observability turns 'sometimes' into a query. You cannot reproduce every production bug locally, so the system must be built to explain itself."* Observability makes an unreproducible defect DIAGNOSABLE; it is not itself a verification of correctness (that remains §11.4.108 / §11.4.185 / §11.4.262). Recommended mechanism gates `CM-CORRELATION-ID-THREADED-ACROSS-BOUNDARIES` (a declared request/trace id is emitted by every signal and survives every async hop — queue, worker, scheduled sweep; a hop that drops it → FAIL) + `CM-CRASH-FREE-RELEASE-GATE` (a crash-free users/sessions gate exists with a declared, locally-calibrated threshold, and a crash signature new at the release boundary blocks regardless of volume) + paired §1.1 mutations (drop the correlation id at the queue boundary → the first gate MUST FAIL; admit a release carrying a new crash signature at low volume → the second MUST FAIL; golden-FALSE per §11.4.201(1): a surface with genuinely no async hop and no crash-reporting channel, honestly §11.4.3 SKIP-with-reason, MUST NOT fire either). Gate-code = separate work item, NOT claimed shipped (§11.4.6 / §11.4.227).

Canonical authority: constitution submodule `Constitution.md` §11.4.260. Non-compliance is a release blocker. No escape hatch — no `--mvp-quality-OK`, `--good-enough`, `--will-fix-later`, `--skip-production-readiness-tracker`, `--production-readiness-optional`, `--acceptable-flakiness`, `--for-now-implementation` flag.

§ . . – Zero-shortcomings / zero-gaps / zero-weak-spots / zero-danger-zones invariant with a mechanical audit ratchet: at all times the project holds exactly zero of these; every finding is closed OR tracked-with-honest-mitigation, and the audit-run count of open findings is MONOTONE-DECREASING (BACKGROUND :: REMINDER :: IMPORTANT operator mandate, – –)

Verbatim operator mandate (2026-08-15, Point 7): *"There MUST BE ALWAYS exactly zero shortcomings, gaps, weak spots or danger zones in the project in any of the code we have, services or components!"*

Compact summary: every consuming project MUST hold, at ALL TIMES, exactly ZERO shortcomings, gaps, weak spots, or danger zones across every code path, service, component, gate, doc, and configuration — a standing invariant, not a per-release checklist; enforced by a MECHANICAL AUDIT SCRIPT that inventories every known finding class from a closed vocabulary (clause B), runs at every pre-build + release + constitution-pull seam, produces a machine-readable ledger, and is bound to a MONOTONE-DECREASING ratchet (§11.4.135 pattern) so the count of open findings only ever drops or the release refuses; every finding is either FULLY CLOSED with captured evidence (§11.4.5/§11.4.69/§11.4.107 + §11.4.226 evidence-class-at-closure) or explicitly TRACKED as a §11.4.197 item with an honest §11.4.223 provenance marker + a §11.4.101-reversible-safe mitigation. "Zero" is the standing invariant; the ratchet is the mechanism that makes it reachable from a brownfield state without silently absorbing debt. §11.4.261 GENERALISES §11.4.124 (dead-code investigate-before-remove) + §11.4.118 (discovery pressure — enumerated but unratched) + §11.4.238 (QA escape triggers audit — one channel) into a WHOLE-PROJECT zero-finding invariant across the closed set.

(A) CLOSED FINDING VOCABULARY — an audit-actionable enumeration, extensible via named consumer-owned additions per §11.4.35 but never subtractable: (1) SHORTCOMINGS — implementation gaps behind a working contract; (2) GAPS — missing capabilities the docs / schema / API promise; (3) WEAK SPOTS — code paths without adequate guards / error handling / input validation; (4) DANGER ZONES — code / config / operational patterns known to cause serious harm under conditions; (5) TODOs / FIXMEs / "for now" /

"temporary" / "will do later" / "placeholder" / `NotImplementedError` / stub-returns / dead code (§11.4.124); (6) SKIPPED tests / disabled gates / commented-out assertions without a tracked `# SKIP-OK: <ticket>` per §11.4.3; (7) BLUFFS — any §11.4/§11.4.1 covenant violation surface (PASS-bluff, FAIL-bluff, echo-observable, source-green-closes-runtime); (8) UNRESOLVED §11.4.197 items past their stated completion window; (9) DIVERGENT / STALE / ORPHAN artifacts per §11.4.233 anti-mess control plane; (10) UN-CATALOGUED anti-patterns identified in code review that lack a §1.1 mutation.

(B) MECHANICAL AUDIT SWEEP — a single executable script (Lava binding: `scripts/audit/zero_findings_sweep.sh` ; consumers supply their path as §11.4.35 DATA) that iterates the closed vocabulary and emits a machine-readable finding ledger `{finding_id, class, file, line, description, status ∈ {open|tracked|closed}, tracker_ref, mitigation_evidence}` — persistent, append-only + snapshot per §11.4.116, git-tracked (§11.4.95) so the ledger IS the SSoT; runs at pre-build + release-tag + constitution-pull; each finding class ships golden-good/golden-bad fixtures per §11.4.107(10) so the audit itself cannot bluff.

(C) MONOTONE-DECREASING RATCHET. The ratchet snapshot records the accepted open-finding count at brownfield adoption (per the §11.4.135 pattern, per-class + total); every subsequent audit run's count MUST be $\leq$ the ratchet snapshot for each class; a run that increases any class's count REFUSES the seam (pre-build gate FAILs, release refuses, commit refuses); the ratchet may DECREASE (a manual snapshot-lower after a finding is closed) and MUST NEVER INCREASE — the release-tag path never silently absorbs new debt. Brownfield adoption per operator §11.4.66 decision (per §11.4.224(E) adoption-fence pattern), never an invented ratchet.

(D) EVERY FINDING → EITHER CLOSED OR TRACKED (never silently absorbed). CLOSED requires captured evidence at the defect layer per §11.4.226 (runtime > artifact > source), the code / config genuinely repaired, the finding disappears from the audit sweep; TRACKED requires (i) a §11.4.197 workable item with §11.4.148 comprehensive structured description, (ii) an honest §11.4.223 provenance marker (`[OPEN: ATM-NNN]`) at the finding site, (iii) a reversible-safe mitigation per §11.4.101, (iv) the tracked item's own completion is on the standing autonomous loop's queue (§11.4.87/§11.4.94/§11.4.97/§11.4.103/§11.4.126); (v) a §11.4.135 permanent regression guard where the tracked item's exposure could

cause harm before closure. Silent absorption ("we know this is broken but shipping anyway") is a §11.4.261 violation of §11.4 PASS-bluff severity — the ledger IS the honesty seam.

(E) HONEST "ZERO" — never the fake kind. "Exactly zero" means the CLOSED-VOCABULARY-audit-run finds zero unclosed-untracked entries; it does NOT claim omniscience (§11.4.118 discovery-pressure boundary — un-catalogued finding classes stay possible, and the vocabulary is EXTENSIBLE for exactly that reason); a §11.4.238 escape triggers the class-extension of the audit + a ratchet-snapshot bump-DOWN (never up) once the escaped class is added and cleared. Honesty at the "zero" claim requires the vocabulary is genuinely covered — a project narrowing the vocabulary to declare "zero" is a §11.4.6 fabrication and a release-blocker.

Honest boundary (§11.4.6). §11.4.261 mandates the STANDING INVARIANT + the AUDIT + the RATCHET + the TRACKED-vs-CLOSED discipline; it does NOT prove the project is provably-flawless, and does NOT license a "zero" claim over a narrowed vocabulary. Composes with, never substitutes for, §11.4.118 (§11.4.261 catches the CATALOGUED — discovery still needed for the un-catalogued), §11.4.238 (an escape triggers vocabulary extension), §11.4.226 (evidence class at closure decides whether a finding is genuinely closed), §11.4.135 (regression guards for closed items keep them closed).

Classification: universal (§11.4.17). Composes §11.4.3 / §11.4.6 / §11.4.15 / §11.4.16 / §11.4.34 / §11.4.44 / §11.4.66 / §11.4.87 / §11.4.94 / §11.4.95 / §11.4.97 / §11.4.101 / §11.4.103 / §11.4.107(10) / §11.4.113 / §11.4.115 / §11.4.116 / §11.4.118 / §11.4.124 / §11.4.126 / §11.4.135 / §11.4.148 / §11.4.161 / §11.4.197 / §11.4.201 / §11.4.223 / §11.4.224(E) / §11.4.226 / §11.4.233 / §11.4.238.

Propagation gate `CM-COVENANT-114-261-PROPAGATION` (literal `11.4.261`) + recommended gates `CM-ZERO-FINDINGS-AUDIT-SWEEP` + `CM-ZERO-FINDINGS-MONOTONE-RATCHET` + `CM-EVERY-FINDING-CLOSED-OR-TRACKED` + paired §1.1 mutations.

Canonical authority: constitution submodule `Constitution.md` §11.4.261. Non-compliance is a release blocker. No escape hatch — no `--skip-findings-audit`, `--allow-todo-without-tracker`, `--ratchet-may-increase`, `--silently-absorb-finding`, `--narrow-vocabulary-to-declare-zero`, `--for-now-in-production` flag.

2026 08 20

§ . . — Build once, promote ONE immutable content-addressed artifact through every environment; never rebuild per stage (research-derived, — —)

Corpus anchor (AI-curriculum module 31, "Core concepts in one page", concept 4 — verbatim):

"Promote one immutable artifact through environments; never rebuild per stage. Build once, then promote the same bytes from CI to staging to canary to production. Rebuilding per environment means the thing you tested is not the thing you shipped — a subtle drift (a floating dependency, a different base image) can make production behave unlike staging. Immutable artifact + environment-specific config injected at deploy is the rule."

Why it is a bluff, not merely an inefficiency. Rebuilding per stage silently BREAKS THE CHAIN OF EVIDENCE: every gate that ran against the staging artifact was a measurement of a DIFFERENT binary than the one production runs, so each of those green verdicts is a §11.4 PASS-bluff about the shipped artifact — evidence correctly captured against the wrong subject. The corpus states the failure directly: *"The thing you tested must be the thing you ship. If you rebuild for production, a floating dependency, a re-pushed base-image tag, or a toolchain difference can make the production artifact subtly unlike the one that passed staging — so your tests validated a different binary. Rebuilding per environment silently breaks the chain of evidence."* Its red-flag form: *"'We rebuild for production to be safe.' (Rebuilding per environment breaks the chain of evidence — you shipped a different artifact than you tested.)"*

The mandate (ALL hold):

(A) BUILD ONCE. Exactly ONE artifact is produced per release candidate, by ONE build, under §11.4.246 (reproducible + hermetic + SLSA-L2 provenance). Every subsequent environment receives THAT artifact. A pipeline that invokes its build

step once per environment violates §11.4.264 regardless of how deterministic the build is claimed to be — determinism is the §11.4.246 property that makes a rebuild *detectably* equivalent, never a licence to rebuild instead of promote.

(B) IDENTITY IS A CONTENT ADDRESS, NOT A TAG. The artifact is identified by a CONTENT-ADDRESSED digest (`sha256:...` or the platform equivalent), never by a mutable tag, branch name, or version string alone. Corpus: "*node:20 is a moving pointer; node@sha256:... is an immutable content address. Tags get re-pushed; digests cannot change without changing the digest.*" Senior form: "*the subtle attack this defends against is not 'a bad version I chose' but 'the same version, different bytes' — a compromised registry or a rebuilt-under-the-same-tag image. Only content-addressing (digests + integrity hashes) catches that; a version string alone does not.*" Composes §11.4.111 (resolve-by-stable-name-not-by-index — a digest is the strongest stable identifier available) and §11.4.200 (the deployed artifact's identity is READ BACK from the intended target and asserted equal to the promoted digest).

(C) PROMOTION IS A VERIFICATION, NOT A REBUILD. Corpus, mechanically: promotion is "*verify this exact digest is authentic and approved for the next environment,*" not "*make a new thing and hope it's equivalent*", and the promotion gate "*check[s] identity and approval of a digest rather than re-running a build.*" Signature, SBOM, and provenance TRAVEL WITH the artifact (§11.4.246), so the evidence chain moves with the bytes instead of being re-derived per stage.

(D) ENVIRONMENT CONFIG IS INJECTED AT DEPLOY, NEVER BAKED AT BUILD. Corpus: "*the mechanism is config injection at deploy, not build-time config baking. The artifact is environment-agnostic; the environment supplies secrets/endpoints/flags at run time. This is what lets the same provably-tested bytes run everywhere.*" Baking per-environment config at build time RE-CREATES the per-stage rebuild it forbids and is explicitly the corpus's marked-wrong answer. Credentials remain governed by §11.4.10 — injection at deploy is never a licence to bake a secret into an artifact.

(E) ROLLBACK IS DETERMINISTIC BY CONSEQUENCE. Because known digests are promoted, rollback is "*point production at the previous known-good digest — a precise, reversible operation, not a rebuild-from-source scramble under incident pressure.*" This is the precondition that makes §11.4.265's auto-abort a real undo rather than a hopeful re-deploy.

(F) REBUILD-TO-DODGE IS A GATE-EVASION, NOT A RETRY. Re-running or rebuilding until a failing scan or flaky check turns green is named by the corpus as a distinct cheat: *"A failing scan or flaky check 'fixed' by re-running until green, or by rebuilding to get a different result, launders a real failure into a pass."* Retries MUST be counted and surfaced (composes §11.4.248 flaky-quarantine — a flake is quarantined with an owner and a deadline, never laundered by a rebuild).

Honest boundary (§11.4.6). §11.4.264 guarantees that the bytes which passed the gates are the bytes that run — it does NOT prove those bytes are correct (that remains §11.4.108's four layers, §11.4.40's full-suite retest, and §11.4.185's manual-QA sufficiency gate), does NOT prove the deploy reached the intended target (that is §11.4.200's isolate-and-read-back), and does NOT prove the environment's injected configuration is right (that is §11.4.254's boot-time invariant assertion). It closes exactly one link: the artifact identity link between test and production. The corpus names no environment ladder beyond `CI → staging → canary → production` ; a consuming project supplies its own ladder as DATA per §11.4.35.

Classification: universal (§11.4.17) — a platform-neutral delivery-pipeline invariant reusable by any project that builds an artifact and deploys it to more than one environment; the consuming project supplies its environment ladder, digest scheme, and config-injection mechanism as DATA per §11.4.35. Composes §11.4 / §11.4.1 (a gate verdict about the wrong artifact is a bluff) / §11.4.6 / §11.4.10 (injected config never bakes a secret) / §11.4.35 / §11.4.108 (the ARTIFACT layer — §11.4.264 is what keeps that layer's identity stable across environments) / §11.4.111 / §11.4.151 (project-prefixed, monotonic release identity) / §11.4.200 (deployed-artifact identity read back from the intended target) / §11.4.235(B) (per-deploy distinct version id) / §11.4.246 (reproducible + hermetic + provenance-attested — §11.4.264 is its DEPLOY-side complement, not a restatement: §11.4.246 governs how the artifact is BUILT, §11.4.264 governs that it is not built AGAIN) / §11.4.248 / §11.4.265 (the canary shifts traffic to exactly the bytes that passed) / §11.4.252 (fail closed when identity cannot be verified) / §1.1.

Propagation gate `CM-COVENANT-114-264-PROPAGATION` (literal `11.4.264` present as a block-start, exactly-once per governance file, lockstep content-hash equality across the mirror set per §11.4.227(B)) + recommended mechanism gates `CM-BUILD-ONCE-PROMOTE-DIGEST` (the pipeline's build step is invoked at most once per release candidate; every post-build environment stage references a

pinned content-addressed digest and NOT a rebuild target; a stage that re-invokes the build → FAIL) + CM-DEPLOY-CONFIG-INJECTED-NOT-BAKED (no environment-specific endpoint/secret/flag is materialised into the artifact at build time; a per-environment build argument that changes artifact bytes → FAIL) + paired §1.1 mutations (add a second per-environment build invocation → CM-BUILD-ONCE-PROMOTE-DIGEST MUST FAIL; replace a promoted digest reference with a mutable tag → it MUST FAIL; bake an environment endpoint into the build → CM-DEPLOY-CONFIG-INJECTED-NOT-BAKED MUST FAIL; golden-FALSE fixture, per §11.4.201(1) the false-positive guard: a single build followed by N digest-referencing promotion stages MUST NOT fire either gate). Gate-code = separate work item, NOT claimed shipped (§11.4.6 / §11.4.227).

Canonical authority: constitution submodule Constitution.md §11.4.264. Non-compliance is a release blocker regardless of context. No escape hatch — no --rebuild-per-environment, --promote-by-tag, --bake-env-config-at-build, --rebuild-until-green, --skip-digest-pin, --deterministic-build-means-rebuild-is-safe flag exists.

§_{26 08 20}. — Progressive delivery:
traffic-shifted rollout gated by automated
analysis on SLO and business metrics, with
automatic abort to stable (research-derived,
— —)

Corpus anchor (AI-curriculum module 31, "Core concepts in one page", concept 5 — verbatim):

"Progressive delivery catches the bluff the tests missed, before all users do. Tests prove what you thought to assert; canary and blue-green deployments prove the release survives real traffic on a small blast radius first. Automated analysis compares the new version against the stable baseline on real metrics (error rate, latency, and business KPIs), and automatically rolls back on breach — no human at 3 AM. The gate here is health, measured, with an instant undo."

Why it is an anti-bluff mechanism and not a deployment convenience. Every gate before this point measures what someone thought to assert; progressive delivery is the FIRST gate whose oracle is real user traffic. It is the deploy-layer instantiation of the §11.4 covenant: it exists to catch the defect the suite never encoded, at a blast radius of a small traffic fraction rather than the whole user base.

The mandate (ALL hold):

(A) A RELEASE REACHES USERS BY A BOUNDED-BLAST-RADIUS

STRATEGY, NEVER AN UNGATED CUTOVER. The consuming project selects and declares its strategy per §11.4.35 from the corpus's two named forms — CANARY (*"shift traffic to the new version gradually ... pausing at each step to evaluate"*; smallest blast radius, requires weight-based routing, old and new run concurrently so schema/compat must tolerate it) or BLUE-GREEN (*"stand up the new version ('green') alongside the current one ('blue') ... Then switch all live traffic at once ... Keep blue running for a rollback window"*; clean cutover and near-instant undo, but *"at the switch, all users hit the new version simultaneously, so a subtle bug hits everyone at once"*). For a critical-invariant surface per §11.4.239 (money movement, safety, availability, integrity) the corpus's stated preference is canary paired with blue-green's instant-undo posture: *"'1% of donors briefly saw a slower page' is a survivable failure and '100% of donors hit a double-charge bug at the cutover instant' is not."*

(B) THE PROMOTE/ABORT DECISION IS MADE BY AUTOMATED ANALYSIS,

NOT BY A HUMAN WATCHING A DASHBOARD. Corpus: *"the controller queries a metrics provider ... and compares the canary against the stable baseline across defined metrics over successive analysis windows. If the metrics stay within thresholds across the required windows, it auto-promotes to the next weight; if a metric breaches, it auto-aborts and rolls back all traffic to stable — no human needed."* A manual approval click is explicitly rejected as the gate — corpus red flag: *"A senior clicks approve, that's the deploy gate." (Manual LGTM is a rubber stamp with no receipt; under agents it looks like oversight while providing none.)"*

This composes §11.4.240 (the producer may not grade its own rollout) and §11.4.262 (the controller's analysis result is the machine-created evidence; a narrated "looks healthy" is not).

(C) THE ANALYSIS MUST INCLUDE BUSINESS METRICS, NOT ONLY INFRASTRUCTURE METRICS — THIS CLAUSE IS THE LOAD-BEARING ONE.

Corpus, naming it *"the classic mistake"*: *"measuring only infrastructure metrics."*

'HTTP error rate is fine' does not tell you the checkout-completion rate quietly dropped 2% because of a UI regression that returns 200s while failing the user. The whole value of canary is catching the failure your tests missed — and the failures that matter most are business failures that look technically green. So the analysis must include business KPIs, or it re-creates the 'green pipeline that proves nothing' at the deploy layer." An infra-only canary is therefore a §11.4 PASS-bluff wearing a rollout's uniform. For a §11.4.239 critical-invariant surface the domain invariants ARE the steady-state hypothesis — corpus: *"a canary that watches latency but not the money invariant would happily promote a build that double-charges 1% of donors."*

(D) MULTIPLE CONSECUTIVE ANALYSIS WINDOWS, NEVER A SINGLE

SAMPLE. Corpus: *"Requiring the canary to pass all checks across multiple consecutive windows guards against a single noisy sample flipping the decision."* Composes §11.4.50 (deterministic consistency — one sample is not a verdict) and §11.4.107 (a steady-state window, never a moment).

(E) POST-DEPLOY SMOKE IS A FAIL-CLOSED PRECONDITION OF TRAFFIC.

Corpus: readiness gates traffic; smoke runs *"against the freshly deployed instance before it takes real traffic (or before promotion)"* and proves *"wiring and configuration — the things that are correct in code but can still be broken by a bad secret, a wrong endpoint, a missing migration, an environment drift"; "a failed smoke gate is a fail-closed signal that aborts the promotion before any user is affected."* Composes §11.4.254 (boot-time invariant assertion) and §11.4.252 (fail closed).

(F) THE DEPLOY GATE IS INDEPENDENT OF THE MERGE GATE. Corpus:

"Signature/provenance verification and canary analysis are gates a merge cannot wave through — so 'got it merged' is not 'got it shipped.'" A merge MUST NOT be able to auto-deploy past the analysis gate.

(G) THRESHOLDS AND WINDOW COUNTS ARE CONSUMER-SUPPLIED

DATA, NEVER CONSTITUTIONAL LITERALS (§11.4.6 / §11.4.35). The corpus deliberately refuses to supply them, writing its abort criterion with UNFILLED placeholders — *"if success rate drops below X% or p99 exceeds Y ms, roll back automatically (thresholds illustrative — set from your SLOs)"* — and adding *"Thresholds and window counts are configuration I would tune to the SLOs and*

confirm live, never quote as fixed." Its governing rule on stale numbers: "A release gate that quotes a stale requirement is itself a bluff." Any project adopting a number supplies and calibrates it locally; this anchor mandates NO numeric value.

Honest boundary (§11.4.6). Progressive delivery bounds the blast radius of a defect and makes the undo automatic — it does NOT prove the release correct (a defect invisible to the declared metrics promotes cleanly), does NOT replace §11.4.185's manual-QA sufficiency gate, and does NOT apply where there is no traffic-shiftable surface (a firmware image flashed to a physical device, a desktop binary, a batch job) — for those the honest posture is an §11.4.3 SKIP-with-reason plus whatever staged-audience mechanism the platform genuinely offers, never a fabricated canary. The corpus's progressive delivery is exclusively TRAFFIC-SHIFTING; it is silent on feature flags, dark launch, kill switches, and shadow traffic, and this anchor mandates none of them. The corpus contains no claim of the form "a deploy you cannot roll back is not a deploy" and none is asserted here.

Classification: universal (§11.4.17) — a platform-neutral delivery discipline reusable by any project that ships to users behind a routable surface; the consuming project supplies its strategy, metric set, thresholds, window counts, and controller as DATA per §11.4.35. Composes §11.4 / §11.4.1 / §11.4.3 (honest SKIP where no traffic surface exists) / §11.4.5 / §11.4.6 / §11.4.35 / §11.4.50 / §11.4.69 / §11.4.107 / §11.4.185 / §11.4.201 (the analysis is a guard: it must be able to abort, and an unresolvable metric read fails closed) / §11.4.235 / §11.4.238 (a defect the canary catches that the suite missed is a coverage escape and triggers that anchor's audit) / §11.4.239 / §11.4.240 / §11.4.252 / §11.4.254 / §11.4.262 / §11.4.264 (the canary shifts traffic to exactly the bytes that passed) / §1.1.

Propagation gate `CM-COVENANT-114-265-PROPAGATION` (literal `11.4.265` present as a block-start, exactly-once per governance file, lockstep content-hash equality across the mirror set per §11.4.227(B)) + recommended mechanism gates `CM-PROGRESSIVE-DELIVERY-STRATEGY-DECLARED` (a user-facing release either declares a bounded-blast-radius strategy with an automated analysis step, or carries an honest §11.4.3 SKIP-with-reason naming the absent traffic surface; an ungated full cutover with neither → FAIL) + `CM-CANARY-ANALYSIS-INCLUDES-BUSINESS-METRIC` (the declared metric set contains at least one business/domain outcome metric, not exclusively infrastructure metrics; an infra-only analysis on a §11.4.239 critical-invariant surface → FAIL) + paired §1.1 mutations (strip every business metric from the analysis set leaving only error-rate and latency → `CM-`

CANARY-ANALYSIS-INCLUDES-BUSINESS-METRIC MUST FAIL; replace the automated abort with a manual approval step → CM-PROGRESSIVE-DELIVERY-STRATEGY-DECLARED MUST FAIL; golden-FALSE fixture per §11.4.201(1): a genuinely non-routable artifact carrying an honest SKIP-with-reason MUST NOT fire either gate). Gate-code = separate work item, NOT claimed shipped (§11.4.6 / §11.4.227).

Canonical authority: constitution submodule Constitution.md §11.4.265. Non-compliance is a release blocker regardless of context. No escape hatch — no --full-cutover-ok, --infra-metrics-suffice, --manual-approval-is-the-gate, --single-window-analysis, --skip-smoke-before-traffic, --merge-implies-deploy flag exists.

§ . . — Claim-vs-reality ledger:
every ADVERTISED capability is a row, typed from
the closed bluff-type vocabulary, and a
capability with no passing challenge is a
release blocker (research-derived,
— —)

Corpus anchor (AI-curriculum module 35, [Q11] "Portfolio-scale prevention", first and cheapest of its enumerated portfolio gates — verbatim):

"A shared claim-vs-reality ledger per repo (the validation plans already are this): every advertised capability is a row with a Bluff type and a severity, and a capability with no passing challenge is a release blocker. This makes stubbed-core, byte-identical-fork, and config-present-but-unwired visible and blocking rather than discovered later."

The direction of enumeration is the whole point, and it is what no existing anchor supplies. §11.4.25's coverage ledger, §11.4.52's autonomous-path column, and §11.4.169's test-type column are all keyed on FEATURES the project already knows it has — they enumerate from the implementation side. A claim-vs-reality ledger enumerates from the CLAIM side: what the project ADVERTISES to users (README capability list, docs, marketing copy, API/CLI surface, badge assertions

per §11.4.259, the §11.4.254 capability matrix). A claim with no feature row is invisible to a feature-keyed ledger by construction — it is precisely the class of defect the corpus documents as *stubbed-core* ("the most-engineered code has no caller"), *byte-identical-fork* ("the name promises X, the code is Y unchanged"), and *config-present-but-unwired* ("the flagship path silently no-ops while the right code exists one dir over"). §11.4.266 therefore adds a keying, not a restatement (§11.4.227).

The mandate (ALL hold):

(A) ONE LEDGER PER REPOSITORY, ENUMERATED FROM THE CLAIM SIDE.

Every advertised capability is a ROW. The consuming project declares per §11.4.35 which surfaces constitute "advertised" (at minimum: the README capability set required reachable by §11.4.212, the §11.4.257 user-manual/guide/FAQ set, the §11.4.254 capability matrix, and the §11.4.259 badge assertions). The corpus does not narrow "advertised" further and neither does this anchor (§11.4.6).

(B) EACH ROW CARRIES A BLUFF TYPE FROM THE CLOSED VOCABULARY, AND A SEVERITY.

The corpus's bluff-type vocabulary, verbatim with its stated fingerprints, is a SEVEN-member set: *green-but-broken* ("passes CI, fails on device"), *coverage-theater* ("the metric lives in prose or a badge, not a failing gate"), *rubber-stamp-verified* ("verified" is set, never earned"), *stubbed-core* ("the most-engineered code has no caller"), *doc-vs-code-drift* ("a countable claim in prose is wrong"), *config-present-but-unwired* ("the flagship path silently no-ops while the right code exists one dir over"), *byte-identical-fork* ("the name promises X, the code is Y unchanged"). The vocabulary is CLOSED for typing purposes and EXTENSIBLE only by an explicit, visible amendment; an un-typeable finding is an honest tracked GAP (§11.4.6 / §11.4.197), never an invented type. The corpus states NO severity scale, so the severity vocabulary is consumer-supplied DATA per §11.4.35 — this anchor mandates none.

(C) THE TYPE IS A ROUTING KEY TO A COUNTER-GATE — THAT IS WHY TYPING IS MANDATORY.

Corpus: "The value of a shared vocabulary is that it turns 'another bug' into 'another instance of type T,' which is what lets you install one gate that kills the whole type. When a reviewer or an agent can tag a finding *stubbed-core*, the remediation is not bespoke — it's 'add the challenge that fails against a no-op stub,' every time." Each type therefore binds to the anchor that already owns its counter — *green-but-broken* → §11.4.108 (four-layer,

RUNTIME-ON-CLEAN-TARGET) + §11.4.262; coverage-theater → §1.1 paired mutation + §11.4.224(C); rubber-stamp-verified → §11.4.146(D3) status custody + §11.4.240 + §11.4.249; stubbed-core → §11.4.27 (no-fakes-beyond-unit) + §11.4.254 (capability matrix reporting attempted-vs-not); doc-vs-code-drift → §11.4.106 + §11.4.186 + §11.4.227; config-present-but-unwired → §11.4.124 (investigate-before-remove / unwired-code) + §11.4.196(F) (configured ≠ in use); byte-identical-fork → §11.4.251. §11.4.266 does NOT restate those counters — it makes the ROUTING mandatory so a finding reaches its counter mechanically instead of by recall.

(D) THE RELEASE-BLOCKING RULE. Verbatim and unqualified: *"a capability with no passing challenge is a release blocker."* A ledger row whose challenge is absent, never executed, or executed only against a stale artifact fingerprint blocks the release exactly as a FAILING challenge does — this is the §11.4.135 / §11.4.236 absence-blocks-as-FAIL semantics applied to the claim side, and it is the reason the ledger is a gate rather than a document.

(E) THE ESCAPE HALF. The corpus pairs the ledger with its sibling: *"any defect found out-of-band (by a human, not by the automated regime) is itself a release blocker, because it proves the automated gates had a hole, and the remediation must include the new check plus the RED-capturing evidence that it would have caught the escape."* That mechanism is already owned by §11.4.238 (automated QA is the DISCOVERER) — §11.4.266 cites it as the ledger's feedback edge and does not duplicate it.

Honest boundary (§11.4.6). The ledger proves that every ADVERTISED capability has a passing challenge — it does NOT prove the challenge is strong (oracle strength stays §11.4.245 + §11.4.107(10) + §1.1), does NOT prove the capability is bug-free (§11.4.118 discovery-pressure remains), and does NOT enumerate capabilities nobody advertised (an un-advertised, un-implemented capability is invisible to a claim-keyed ledger exactly as an un-claimed feature is invisible to a feature-keyed one — the two keyings are complements, and a project running only one has an honestly-stated blind side). The corpus supplies no worked ledger row, no severity scale, and no definition of "reality" as a field; all three are consumer DATA per §11.4.35.

Provenance correction, recorded per §11.4.6. Secondary project documentation has referred to this material as an " FP-1 ... FP-14 " failure-pattern catalogue, an " L-1 ... L-16 " obligations list, and an " F-1 ... F-31 " bluff taxonomy. Measured

against the tracked corpus, **no F-n , FP-n , or L-n identifier exists in any of the nine modules** (control-needle proven: the same scan that returns 0 for those forms returns the real `ATM-277 / ATM-343 / LVA-098 / ITE-6 / API-33` identifiers, and 0 for a negative control). The real structures are: a 14-row numbered failure-pattern table whose last column is "The rock-solid approach we lack"; a 6-bullet portfolio-gate list in [Q11]; and this 7-member bluff-type vocabulary in [Q2]. Anchors **MUST** cite the real structures; the `F-n / FP-n / L-n` labels **MUST NOT** be used as corpus wording.

Classification: universal (§11.4.17) — a platform-neutral claim-integrity discipline reusable by any project that advertises capabilities to users; the consuming project supplies its advertised-surface set, severity scale, ledger location, and challenge runner as DATA per §11.4.35. Composes §11.4 / §11.4.1 / §11.4.6 / §11.4.25 (feature-keyed coverage ledger — §11.4.266 is its claim-keyed complement) / §11.4.27 / §11.4.35 / §11.4.52 / §11.4.106 / §11.4.108 / §11.4.115(F) / §11.4.118 / §11.4.124 / §11.4.135 (absence blocks as a FAIL) / §11.4.146(D3) / §11.4.169 / §11.4.186 / §11.4.196(F) / §11.4.197 / §11.4.212 / §11.4.224 / §11.4.226 / §11.4.227 / §11.4.236 / §11.4.238 (the escape edge) / §11.4.240 / §11.4.245 / §11.4.249 / §11.4.251 / §11.4.254 / §11.4.257 / §11.4.259 / §11.4.262 / §1.1.

Propagation gate `CM-COVENANT-114-266-PROPAGATION` (literal `11.4.266` present as a block-start, exactly-once per governance file, lockstep content-hash equality across the mirror set per §11.4.227(B)) + recommended mechanism gates `CM-CLAIM-REALITY-LEDGER-COMPLETE` (every advertised capability on every declared advertised surface resolves to exactly one ledger row; an advertised capability with no row → FAIL naming the capability and the surface it was advertised on) + `CM-UNCHALLENGED-CAPABILITY-BLOCKS-RELEASE` (a row whose challenge is absent, never executed, or has no verdict for the release-candidate artifact fingerprint blocks the release exactly as a FAIL does) + `CM-LEDGER-ROW-TYPED-FROM-CLOSED-VOCABULARY` (every row carries a bluff type drawn from the seven-member closed set; an ad-hoc type → FAIL) + paired §1.1 mutations (add a README capability with no ledger row → the completeness gate **MUST** FAIL; blank a row's challenge reference → the blocking gate **MUST** FAIL; retype a row to an invented type → the vocabulary gate **MUST** FAIL; golden-FALSE fixture per §11.4.201(1): a fully-populated ledger whose every row has a fresh candidate-fingerprinted PASS verdict **MUST NOT** fire any of the three). Gate-code = separate work item, **NOT** claimed shipped (§11.4.6 / §11.4.227).

Canonical authority: constitution submodule `Constitution.md` §11.4.266. Non-compliance is a release blocker regardless of context. No escape hatch — no `--ledger-optional`, `--unchallenged-capability-ok`, `--untyped-ledger-row`, `--invent-bluff-type`, `--feature-ledger-suffices`, `--advertise-without-row` flag exists.

§ . . — Shared attempt record + converge-on-evidence / escalate-on-stall: a failed approach is never silently retried, and a non-converging loop is itself a signal (research-derived, — —)

Corpus anchor (AI-curriculum module 34, [Q15] — verbatim):

"No memory of what was tried → the loop repeats itself. Agents re-attempt discarded approaches. Fix: a shared record of attempts and outcomes (in the task state / PR thread), so a failed approach is not silently retried; and an escalation rule — after N failed iterations, stop and route to a human or a stronger model, rather than burning cycles."

The premise that makes it structural, not optional. Corpus: *"The agent has no memory of prior attempts. It will re-implement the same wrong fix because it does not remember the last three. Prevention: encode the lesson where the agent must encounter it — the test, the spec, a decision record in the repo — not in a human's head."* And: *"no reopening' is really a statement about durability of knowledge, and agents have no durable memory, so the regime must externalize it."* The record is therefore not a convenience log — it is the externalised memory without which the loop is structurally incapable of converging.

What this is NOT (the §11.4.227 boundary, stated first). §11.4.147 governs a CRASHED agent (detect → respawn → resume until `complete`). §11.4.232(C) governs a HUNG long-op (heartbeat delta against a no-progress budget). §11.4.21 governs the TERMINAL exhaustion audit written at the moment `operator-blocked` is declared. None of the three governs a LIVE, HEALTHY,

PROGRESSING loop that is re-walking a dead end its predecessor already proved dead — the agent is not crashed, the op is emitting heartbeats, and no terminal state has been reached. §11.4.267 closes exactly that gap.

The mandate (ALL hold):

(A) A SHARED, DURABLE ATTEMPT RECORD. Every multi-attempt effort (an agent loop, a producer ↔ verifier cycle, a fix-retry sequence, a multi-agent hand-off) MUST write each attempt and its OUTCOME to a record SHARED across every actor that may take the next attempt — including a successor agent, a respawned agent per §11.4.147, and a different track per §11.4.176. The corpus names its location as *"in the task state / PR thread"*; the consuming project binds its own location as DATA per §11.4.35 (the §11.4.93 workable-items DB and the §11.4.116 append-only event stream are the natural substrates and MUST be reused rather than duplicated per §11.4.227). The corpus specifies no schema and none is mandated here (§11.4.6).

(B) CONSULTED BEFORE ACTING — the clause that makes the record load-bearing. The record MUST be READ before the next attempt is chosen, and an approach already recorded as failed MUST NOT be silently retried. Re-attempting a recorded dead end is permitted only when the attempt materially differs (a changed precondition, a corrected §11.4.199 reproduction sequence, a new §11.4.115(F) RED) and that difference is recorded. A written-but-never-read record is the §11.4.196(F) configured-≠-in-use failure applied to agent memory, and satisfies nothing.

(C) CONVERGE ON EVIDENCE. Corpus: *"the design principle is converge on evidence, escalate on stall. Each loop must monotonically reduce uncertainty — the producer gets a real failing receipt and must make it green."* The two named preconditions of convergence: a SINGLE SOURCE OF TRUTH (*"the spec + acceptance criteria is the authority, and the verifier's captured failure is the ground truth. Disagreements are adjudicated on the receipt, not on whoever argues more confidently — the verifier's build exit 1 beats the producer's fluent paragraph every time"*) and CONCRETE FAILURE FEEDBACK (*"If the verifier just says 'still broken,' the producer tries another plausible change. Fix: feed the exact captured error (argv, exit code, failing assertion, the offending line) back as the next input"*). Composes §11.4.240 / §11.4.249 (the verifier is not the producer) and §11.4.262 (the receipt, not the narration).

(D) STALL IS A DETECTED SIGNAL WITH A BOUNDED RETRY AND AN ESCALATION PATH.

The corpus names TWO stall signatures — *"the same criterion red after N tries"* and *"oscillating diffs"* — and states the response: *"after N failed iterations, stop and route to a human or a stronger model, rather than burning cycles."* The stall itself is tracked, not merely survived: *"a loop that is not converging ... is itself a signal, tracked and escalated. Frequent producer–verifier divergence is a metric worth trending: it predicts where the producer over-claims and where the spec is ambiguous."* The forbidden shape is explicit: *"The anti-pattern is an unbounded retry loop with no ground-truth feedback and no escalation — that is how agents thrash expensively while appearing busy."* Escalation composes §11.4.231 (route to a stronger tier when complexity is PROVEN mid-flight), §11.4.101 (decide autonomously where the choice is safe and reversible; park only what is genuinely blocked), and §11.4.66 / §11.4.21 (the operator path, reached only after the §11.4.21 self-resolution audit).

(E) N IS CONSUMER-SUPPLIED DATA, NEVER A CONSTITUTIONAL LITERAL (§11.4.6 / §11.4.35).

The corpus leaves N symbolic in every one of its occurrences and states no numeric bound, no wall-clock budget, no token ceiling, and no cost ceiling. This anchor mandates a BOUND EXISTS and is DECLARED; it mandates no value. A project that enforces an invented number and attributes it to the corpus commits the §11.4.6 violation this clause prevents.

(F) THE LOOP'S DONE-CONDITION. A loop MUST NOT read satisfied while an attempt is recorded in-flight with no outcome, and an escalation raised under (D) MUST reach a terminal disposition (resolved / re-tiered / operator-parked per §11.4.101) — never be silently dropped. Composes §11.4.87 / §11.4.94 / §11.4.97 / §11.4.126 / §11.4.147 (a parked escalation parks ONE work unit, never the loop) and §11.4.197 (a started effort reaches a terminal state).

Honest boundary (§11.4.6). The attempt record makes the loop's history legible and prevents the RECORDED dead end from being re-walked — it does NOT guarantee convergence (a problem may be genuinely hard, and an escalation is the honest outcome, not a failure of the record), does NOT prove any individual attempt was competent, and does NOT substitute for §11.4.102's systematic-debugging discipline on the underlying defect. It bounds waste and preserves knowledge across actors; it does not manufacture insight.

Classification: universal (§11.4.17) — a platform-neutral orchestration discipline reusable by any project running multi-attempt or multi-agent loops; the consuming project supplies its record location, schema, stall detector, `N`, and escalation route as DATA per §11.4.35. Composes §11.4.6 / §11.4.20 / §11.4.21 (the terminal exhaustion audit — §11.4.267 is its running, pre-terminal counterpart) / §11.4.35 / §11.4.66 / §11.4.70 / §11.4.87 / §11.4.93 / §11.4.94 / §11.4.97 / §11.4.101 / §11.4.102 / §11.4.115(F) / §11.4.116 (the append-only event stream substrate) / §11.4.126 / §11.4.146 / §11.4.147 (crash → respawn; §11.4.267 governs the healthy-but-looping case it does not reach) / §11.4.176 / §11.4.187 / §11.4.196(F) / §11.4.197 / §11.4.199 / §11.4.231 / §11.4.232(C) (hung long-op; §11.4.267 governs the progressing-but-not-converging case) / §11.4.240 / §11.4.249 / §11.4.262 / §1.1.

Propagation gate `CM-COVENANT-114-267-PROPAGATION` (literal `11.4.267` present as a block-start, exactly-once per governance file, lockstep content-hash equality across the mirror set per §11.4.227(B)) + recommended mechanism gates `CM-ATTEMPT-RECORD-SHARED-AND-CONSULTED` (every multi-attempt effort writes each attempt + outcome to the declared shared record, and the record is read before the next attempt is dispatched; a second attempt duplicating a recorded-failed approach with no recorded material difference → FAIL) + `CM-STALL-BOUNDED-AND-ESCALATED` (a declared retry bound exists; a loop exceeding it without a recorded escalation to a terminal disposition → FAIL) + paired §1.1 mutations (strip the record-read step so the loop dispatches blind → the first gate MUST FAIL; remove the retry bound leaving an unbounded loop → the second MUST FAIL; record an escalation with no terminal disposition → the second MUST FAIL; golden-FALSE fixture per §11.4.201(1): a loop that converges on attempt 2 with both attempts recorded and no bound exceeded MUST NOT fire either gate). Gate-code = separate work item, NOT claimed shipped (§11.4.6 / §11.4.227).

Canonical authority: constitution submodule `Constitution.md` §11.4.267. Non-compliance is a release blocker regardless of context. No escape hatch — no `--no-attempt-record`, `--record-without-reading`, `--unbounded-retry-loop`, `--retry-same-failed-approach`, `--stall-without-escalation`, `--invent-N-and-cite-corpus` flag exists.