

Fixtures: transcript_ingest (T , plan T-A)

RED-set fixtures for `constitution/scripts/fastcycle/tests/test_token_attribution_red.sh` (T020). Guards T036 (`$FC/tokens/dispatch_stamp.sh`) and T038 (`$FC/tokens/transcript_ingest.py`) — neither exists yet; this test and its fixtures are written and observed to fail against that absence (§11.4.224, common-conventions.md C-005). Producer≠Verifier (§11.4.240): this directory and the test script are the RED artefact only; `dispatch_stamp.sh` and `transcript_ingest.py` are separate, later, implementation tasks owned by other agents.

Real-schema provenance (\$. . — nothing here is invented)

Every field shape below was read from this project's OWN live Claude Code transcripts on 2026-09-28 (session `8824e088-62a6-4d4f-9f6c-2320055ad033` and sibling sessions under `~/.claude-claude5/projects/-mnt-track1-atmosphere-t1/`), never assumed:

- Top-level (non-sidechain) transcript records: `parentUuid` , `isSidechain` , `type` (`user` | `assistant` |...), `uuid` , `timestamp` , `sessionId` , `message` . An assistant record's `message` carries `model` , `id` , `type` , `role` , `stop_reason` , `content` , and (when the API call actually returned) a `usage` object with `input_tokens` , `output_tokens` , `cache_creation_input_tokens` , `cache_read_input_tokens` (plus provider-internal fields such as `output_tokens_details` , `server_tool_use` , `cache_creation` , `service_tier` , `iterations` , `speed` — irrelevant to token attribution and omitted from these minimal fixtures).
- A rate-limited/rejected turn's `message.usage` is present with EVERY count explicitly `0` (captured verbatim: session `0099c079-62e3-4846-a2a6-eacb7d42a85c` , line 49, `error: "rate_limit"` , `isApiErrorMessage:`

true , `apiErrorStatus: 429` , all four usage counts `0`) — this is the REAL shape `missing_usage_block/usage_present_zero.jsonl` is modelled on: a genuinely-zero, fully-instrumented turn, never to be confused with a turn that carries no `usage` key at all.

- Async agent dispatch: the PARENT transcript's dispatching `assistant` turn contains a `tool_use` content block (`type: "tool_use"` , `name: "Agent"` , `input: {description, prompt}`); the immediately-following `user` turn (the tool result) carries `toolUseResult: {isAsync: true, status: "async_launched", agentId, description, resolvedModel, outputFile, ...}` (captured verbatim from a genuine live async dispatch, `agentId a3769936d2ace3124` , session `8824e088-62a6-4d4f-9f6c-2320055ad033` , 2026-09-28 — the SAME real payload shape `scripts/hooks/test_agent_registry_dispatch_writer_red.sh` 's own A8 case is built from, cited there as `agentId ac39c9d3e8933edd3` , 2026-09-27).
- **Real captured dispatch description today carries NO `item=<ATM-nnnn>` token anywhere** — verbatim captured example (session `8824e088-62a6-4d4f-9f6c-2320055ad033` , `toolUseResult.description`): `"(T1/main - claude5 - sonnet - high) T018 RED test review_record"` . This is the direct evidence behind property (a) in the test.
- Subagent transcript on-disk convention: `<CLAUDE_CONFIG_DIR>/projects/<project-slug>/<parent_session_id>.jsonl` (the parent/CLI transcript) alongside `<parent_session_id>/subagents/agent-<agentId>.jsonl` (the subagent's OWN transcript) — directly observed via `<repo>/.../tasks/<agentId>.output` symlinks resolving to exactly that path. A subagent transcript record carries `agentId` directly (not only derivable from the file name) plus `isSidechain: true` and (for its first `user` record) `promptId` . This is the real join key `transcript_ingest.py` (T038)'s "session → agent → item" keying (plan.md T-A06) needs: `sessionId` (session) → `agentId` (agent, matching both the subagent record's own field AND the parent's `toolUseResult.agentId` AND the `agent-<agentId>.jsonl` filename) → the `item=<ATM-nnnn>` token embedded in the parent's stamped `description` (item, once T036/T037 land).

Fixture map

Path	Property under test	Shape
<code>cache_hits/session.jsonl</code>	cache-hit tokens are counted, not dropped	one <code>user</code> turn + two <code>assistant</code> turns: turn 1 is cache-cold (<code>cache_creation_input_tokens=1200</code> , <code>cache_read_input_tokens=0</code>), turn 2 is a genuine cache HIT reading the same 1200 tokens back (<code>cache_read_input_tokens=1200</code>)
<code>miss-ing_usage_block/us-age_absent.jsonl</code>	a turn with NO <code>usage</code> key MUST be reported missing/ UN- MEASURED (common- conven- tions.md term), never coerced to 0	<code>assistant</code> <code>message</code> object has no "usage" key at all
<code>miss-ing_usage_block/us-age_present_zero.jsonl</code>	negative control (§11.4.201(6) false-null guard, C-004 control needle) — a genuinely- zero, fully-in- strumented turn MUST NOT be reported missing	<code>assistant</code> <code>message.usage</code> present, all four counts explicitly 0 (modelled on the real rate-limited-turn shape above)

Path	Property under test	Shape
sub-agent_attribution/parent_session.jsonl + sub-agent_attribution/parent_session/subagents/agent-fixturet020attr01.jsonl	a subagent's token usage attributes to its parent dispatch's item	parent dispatches with description containing item=ATM-9999 (the future T036-stamped tag) and toolUseResult.agentId = "fixturet020attr01" ; the subagent transcript (real on-disk path convention) carries the SAME agentId and its own real usage (cache_read_input_tokens=500 , cache_creation_input_tokens=8000)
credential_leak/session.jsonl	a planted secret-shaped string in message content never reaches the telemetry DB (§11.4.10)	user + assistant turns whose text embeds the literal marker FASTCYCLE-T020-PLANTED-MARK-ER-7f3a9c2e1b8d4f6091ab34cd , clearly labelled synthetic; a genuine usage block accompanies the assistant turn so the fixture also exercises normal ingest

Credential-scanner safety (verified, not assumed)

Every fixture file above was scanned with this repository's OWN

helix_cred_scan_file (constitution/scripts/hooks/credential_scan_lib.sh) before being checked in, with a positive control run first (a synthetic AKIA -shaped AWS key, confirmed CAUGHT) proving the scanner instrument was alive for that run — all six fixture files scanned clean. The planted marker in credential_leak/session.jsonl deliberately avoids every keyword the scanner's HELIX_CRED_VALUE_PATTERN matches (password|passwd|secret|api[_-]?key|access[_-]?token|auth[_-]?token|client[_-]?secret adjacent to : / =) and every literal prefix it matches (AKIA , ghp_ , gho_ , github_pat_ , xox[baprs]- , sk- , AIza , -----BEGIN ... PRIVATE KEY-----) — it is a synthetic marker STRING representing "sensitive message-

body content that must never reach a metrics-only DB", not a credential-SHAPED assignment the scanner is designed to catch (that is a different, already-covered feature).